

O'REILLY®

4th Edition

Includes Azure
Service Fabric

Programming WCF Services

DESIGN AND BUILD MAINTAINABLE
SERVICE-ORIENTED SYSTEMS



Juval Löwy &
Michael Montgomery

Programming WCF Services

Programming WCF Services is the authoritative, bestselling guide to Microsoft's unified platform for developing modern, service-oriented applications on Windows. Hailed as the definitive treatment of WCF, this guide provides unique insight, rather than documentation, to help you learn the topics and skills you need for building maintainable, extensible, and reusable WCF-based applications.

Authors Juval Löwy—one of the world's top .NET experts—and Michael Montgomery have revised this edition to include the productivity-enhancing features of .NET Framework 4.6, along with the latest WCF ideas and techniques. By teaching you the why and the how of WCF programming, this book will help you master WCF and make you a better software engineer.

- Learn WCF's architecture and essential building blocks, including key concepts such as reliability and transport sessions
- Use built-in features such as service contracts, instance and concurrency management, transactions, queued services, and security
- Increase the quality of your WCF services by using design options, tips, and best practices in Löwy's ServiceModelEx framework
- Understand the rationale behind particular design decisions, and rarely understood aspects of WCF development
- Learn why Azure Service Fabric is *the* killer app for modern DevOps

Juval Löwy is the founder of IDesign, specializing in system and project design. He has mentored many architects across the world in his Master Classes. Juval participated in the Microsoft strategic design reviews for C# and WCF. Microsoft recognized him as a Software Legend, one of the world's top experts and industry leaders.

Michael Montgomery is a master architect with IDesign, specializing in distributed system design and microservices. Monty has published many articles on architecture, always sharing his unique "From the Field" perspective. He speaks regularly at .NET and industry-related events.

MICROSOFT

US \$59.99

CAN \$68.99

ISBN: 978-1-491-94483-7



5 5 9 9 9
9 781491 944837



Twitter: @oreillymedia
facebook.com/oreilly

4TH EDITION

Programming WCF Services

Juval Lowy & Michael Montgomery

Beijing • Boston • Farnham • Sebastopol • Tokyo

O'REILLY®

Programming WCF Services

by Juval Lowy and Michael Montgomery

Copyright © 2016 Juval Lowy and Michael Montgomery. All rights reserved.

Printed in the United States of America.

Published by O'Reilly Media, Inc., 1005 Gravenstein Highway North, Sebastopol, CA 95472.

O'Reilly books may be purchased for educational, business, or sales promotional use. Online editions are also available for most titles (<http://safaribooksonline.com>). For more information, contact our corporate/institutional sales department: 800-998-9938 or corporate@oreilly.com.

Editors: Brian MacDonald and Rachel Roumeliotis

Interior Designer: David Futato

Production Editor: Melanie Yarbrough

Cover Designer: Randy Comer

Proofreader: Jasmine Kwityn

Illustrator: Rebecca Demarest

Indexer: WordCo Indexing Services

December 2015: Fourth Edition

Revision History for the Fourth Edition

2015-11-04: First Release

See <http://oreilly.com/catalog/errata.csp?isbn=9781491944837> for release details.

The O'Reilly logo is a registered trademark of O'Reilly Media, Inc. *Programming WCF Services*, the cover image, and related trade dress are trademarks of O'Reilly Media, Inc.

While the publisher and the author have used good faith efforts to ensure that the information and instructions contained in this work are accurate, the publisher and the author disclaim all responsibility for errors or omissions, including without limitation responsibility for damages resulting from the use of or reliance on this work. Use of the information and instructions contained in this work is at your own risk. If any code samples or other technology this work contains or describes is subject to open source licenses or the intellectual property rights of others, it is your responsibility to ensure that your use thereof complies with such licenses and/or rights.

978-1-491-94483-7

[LSI]

Table of Contents

Preface.....	xiii
1. WCF Essentials.....	1
What Is WCF?	1
Services	3
Service Execution Boundaries	4
WCF and Location Transparency	5
Addresses	5
TCP Addresses	6
HTTP Addresses	7
IPC Addresses	7
MSMQ Addresses	8
WebSocket Addresses	8
UDP Addresses	8
Contracts	8
The Service Contract	9
Hosting	13
IIS Hosting	14
Self-Hosting	15
WAS Hosting	21
Custom Hosting in IIS/WAS	21
Choosing a Host	22
Bindings	23
The Common Bindings	24
Choosing a Binding	26
Additional Bindings	27
Using a Binding	30
Endpoints	30

Administrative Endpoint Configuration	31
Programmatic Endpoint Configuration	35
Default Endpoints	37
The Configure() Method	40
Metadata Exchange	43
Metadata over HTTP-GET	44
The Metadata Exchange Endpoint	47
The Metadata Explorer	54
More on Behavior Configuration	56
Client-Side Programming	58
Generating the Proxy	59
Administrative Client Configuration	64
Programmatic Client Configuration	71
The WCF-Provided Test Client	72
Programmatic Versus Administrative Configuration	75
Configuration Policy	75
WCF Architecture	76
Host Architecture	77
Working with Channels	78
The InProcFactory Class	80
Transport-Level Sessions	84
Transport Session and Binding	84
Transport Session Termination	85
Reliability	86
Bindings, Reliability, and Ordered Messages	87
Configuring Reliability	88
Requiring Ordered Delivery	89
2. Service Contracts.....	93
Operation Overloading	93
Contract Inheritance	96
Client-Side Contract Hierarchy	97
Service Contract Factoring and Design	101
Contract Factoring	101
Factoring Metrics	103
Contract Queries	105
Programmatic Metadata Processing	106
The MetadataHelper Class	109
3. Data Contracts.....	115
Serialization	115
.NET Serialization	117

The WCF Formatters	119
Data Contract via Serialization	122
Data Contract Attributes	123
Importing a Data Contract	125
Data Contracts and the Serializable Attribute	129
Inferred Data Contracts	130
Composite Data Contracts	131
Data Contract Events	133
Shared Data Contracts	136
Data Contract Hierarchy	137
Known Types	138
Service Known Types	140
Multiple Known Types	141
Configuring Known Types	142
Data Contract Resolvers	143
Objects and Interfaces	155
Data Contract Equivalence	157
Serialization Order	158
Versioning	160
New Members	161
Missing Members	162
Versioning Round-Trip	166
Data Contract Factoring and Design	169
Data Contract Design	169
Data Contract Factoring	169
Operation Factoring	181
The Data Transfer Object Pattern	183
Enumerations	184
Delegates and Data Contracts	186
Generics	187
Collections	191
Concrete Collections	192
Custom Collections	194
The CollectionDataContract Attribute	194
Referencing a Collection	196
Dictionaries	196
4. Instance Management.....	201
Behaviors	201
Per-Call Services	203
Benefits of Per-Call Services	204
Configuring Per-Call Services	204

Per-Call Services and Transport Sessions	206
Designing Per-Call Services	206
Choosing Per-Call Services	209
Per-Session Services	209
Configuring Private Sessions	210
Sessions and Reliability	215
The Session ID	216
Session Termination	218
Singleton Service	218
Initializing a Singleton	220
Choosing a Singleton	223
Demarcating Operations	223
Instance Deactivation	226
Configuring with ReleaseInstanceMode.None	227
Configuring with ReleaseInstanceMode.BeforeCall	228
Configuring with ReleaseInstanceMode.AfterCall	229
Configuring with ReleaseInstanceMode.BeforeAndAfterCall	229
Explicit Deactivation	230
Using Instance Deactivation	231
Durable Services	231
Durable Services and Instance Management Modes	232
Instance IDs and Durable Storage	233
Explicit Instance IDs	234
Instance IDs in Headers	236
Context Bindings for Instance IDs	239
Automatic Durable Behavior	245
Throttling	252
Configuring Throttling	255
5. Operations.....	261
Request-Reply Operations	261
One-Way Operations	262
Configuring One-Way Operations	262
One-Way Operations and Reliability	263
One-Way Operations and Sessionful Services	263
One-Way Operations and Exceptions	264
Callback Operations	266
The Callback Contract	267
Client Callback Setup	268
Service-Side Callback Invocation	272
Callback Connection Management	275
The Duplex Proxy and Type Safety	278

The Duplex Factory	281
Callback Contract Hierarchy	283
Events	285
Streaming	289
I/O Streams	289
Streaming and Binding	290
Streaming and Transport	291
6. Faults.....	293
Error Isolation and Decoupling	293
Error Masking	294
Channel Faulting	295
Fault Propagation	299
Fault Contracts	301
Fault Debugging	305
Faults and Callbacks	312
Error-Handling Extensions	315
Providing a Fault	316
Handling a Fault	319
Installing Error-Handling Extensions	322
The Host and Error Extensions	325
Callbacks and Error Extensions	329
7. Transactions.....	333
The Recovery Challenge	333
Transactions	334
Transactional Resources	335
Transaction Properties	335
Transaction Management	337
Resource Managers	341
Transaction Propagation	341
Transaction Flow and Bindings	341
Transaction Flow and the Operation Contract	343
One-Way Calls	345
Transaction Protocols and Managers	345
Protocols and Bindings	346
Transaction Managers	348
Transaction Manager Promotion	350
The Transaction Class	352
The Ambient Transaction	352
Local Versus Distributed Transactions	353
Transactional Service Programming	354

Setting the Ambient Transaction	354
Transaction Propagation Modes	357
Voting and Completion	365
Transaction Isolation	369
Transaction Timeout	371
Explicit Transaction Programming	373
The TransactionScope Class	373
Transaction Flow Management	375
Non-Service Clients	383
Service State Management	384
The Transaction Boundary	385
Instance Management and Transactions	386
Per-Call Transactional Services	387
Per-Session Transactional Services	391
Transactional Durable Services	406
Transactional Behavior	408
Transactional Singleton Service	415
Instancing Modes and Transactions	418
Callbacks	419
Callback Transaction Modes	420
Callback Voting	422
Using Transactional Callbacks	422
8. Concurrency Management.	427
Instance Management and Concurrency	428
Service Concurrency Modes	428
ConcurrencyMode.Single	429
ConcurrencyMode.Multiple	429
ConcurrencyMode.Reentrant	433
Instances and Concurrent Access	436
Per-Call Services	436
Sessionful and Singleton Services	437
Resources and Services	438
Deadlocked Access	438
Deadlock Avoidance	439
Resource Synchronization Context	441
.NET Synchronization Contexts	441
The UI Synchronization Context	445
Service Synchronization Context	451
Hosting on the UI Thread	452
A Form as a Service	458
The UI Thread and Concurrency Management	461

Custom Service Synchronization Contexts	464
The Thread Pool Synchronizer	464
Replacing the I/O Completion Thread Pool	469
Thread Affinity	471
Priority Processing	473
Callbacks and Client Safety	477
Callbacks with ConcurrencyMode.Single	478
Callbacks with ConcurrencyMode.Multiple	479
Callbacks with ConcurrencyMode.Reentrant	479
Callbacks and Synchronization Contexts	480
Callbacks and the UI Synchronization Context	480
Callback Custom Synchronization Contexts	484
Asynchronous Calls	488
Requirements for an Asynchronous Mechanism	488
Proxy-Based Asynchronous Calls	489
Asynchronous Invocation	492
Polling or Waiting for Completion	494
Continuations	497
One-Way Asynchronous Operations	508
Asynchronous Error Handling	512
Asynchronous Calls and Transactions	515
Synchronous Versus Asynchronous Calls	516
Asynchronous Service Operations	517
9. Queued Services	543
Disconnected Services and Clients	543
Queued Calls	544
Queued Calls Architecture	545
Queued Contracts	545
Configuration and Setup	546
Transactions	554
Delivery and Playback	554
Service Transaction Configuration	556
Nontransactional Queues	559
Instance Management	560
Per-Call Queued Services	561
Sessionful Queued Services	563
Singleton Service	566
Concurrency Management	567
Throttling	568
Delivery Failures	569
The Dead-Letter Queue	570

Time to Live	571
Configuring the Dead-Letter Queue	571
Processing the Dead-Letter Queue	573
Playback Failures	578
Poison Messages	579
Poison Message Handling	579
Queued Versus Connected Calls	584
Requiring Queuing	586
The Response Service	587
Designing a Response Service Contract	588
Client-Side Programming	592
Queued Service-Side Programming	595
Response Service-Side Programming	597
Transactions	598
The HTTP Bridge	602
Designing the Bridge	602
Transaction Configuration	603
Service-Side Configuration	604
Client-Side Configuration	606
10. Security.....	609
Authentication	609
Authorization	610
Transfer Security	611
Transfer Security Modes	612
Transfer Security Mode Configuration	614
Transport Security and Credentials	617
Message Security and Credentials	617
Identity Management	618
Overall Policy	618
Scenario-Driven Approach	619
Intranet Application Scenario	620
Securing the Intranet Bindings	620
Constraining Message Protection	628
Authentication	629
Identities	632
The Security Call Context	633
Impersonation	635
Authorization	643
Identity Management	650
Callbacks	651
Internet Application Scenario	652

Securing the Internet Bindings	652
Message Protection	655
Authentication	660
Using Windows Credentials	662
Using the ASP.NET Providers	663
Identity Management	673
Business-to-Business Application Scenario	673
Securing the Business-to-Business Bindings	674
Authentication	674
Authorization	677
Identity Management	679
Host Security Configuration	679
Anonymous Application Scenario	680
Securing the Anonymous Bindings	680
Authentication	681
Authorization	681
Identity Management	681
Callbacks	682
No Security Scenario	682
Unsecuring the Bindings	682
Authentication	683
Authorization	683
Identity Management	683
Callbacks	683
Scenarios Summary	684
Declarative Security Framework	684
The SecurityBehaviorAttribute	685
Host-Side Declarative Security	693
Client-Side Declarative Security	695
Security Auditing	702
Configuring Security Audits	703
Declarative Security Auditing	706
11. The Azure Service Fabric.....	709
Why Azure Service Fabric	709
An Integrated Platform	710
Design Consistency	710
Programming Model Consistency	711
Future Compatible	711
Preparing for the Service Fabric	713
ServiceModelEx.ServiceFabric	713
Configuring the Service Fabric	714

IT Operations	714
Development	716
Both	718
Programming the Service Fabric	723
The Service Programming Model	724
Reliable Services	724
Reliable Actors	734
State Management	737
Hosting	745
The Client Programming Model	746
Modeling the Service Fabric	756
Modeling Approach	757
Modeling Reliable Services	758
Modeling Reliable Actors	761
Modeling Hosting	770
Testing	771
A. Introduction to Service Orientation	779
B. Headers and Contexts	795
C. Discovery	819
D. Publish-Subscribe Service	883
E. Generic Interceptor	911
F. WCF Coding Standard	929
G. ServiceModelEx Catalog	943
Index	969

Preface

On the software time scale, WCF is ancient. In August 2001, I first learned the details of a Microsoft effort to rewrite COM+ using managed code. Nothing much happened after that. Then, during a C# 2.0 Strategic Design Review in July 2002, the remoting program manager outlined in broad strokes plans to rework remoting into something that developers should actually use. At the same time, Microsoft was also working on incorporating the new security specs for web services into the ASMX stack and actively working with others on drafting a score of additional web services specs.

In July 2003, I was given access to a new transactional infrastructure that improved on the deficiencies in transactional .NET programming. At the time, there was no cohesive programming model that unified these distinct technologies. Toward the end of 2003, I was privileged to be invited to join a small team of outside industry experts and to participate in the strategic design review of a new development platform codenamed *Indigo*. Some of the smartest and nicest people I ever worked with were part of that team. Over the next two to three years, Indigo went through some three generations of programming models. The final declarative, endpoint-driven object model debuted in early 2005, was stabilized by August of that year, and was named the Windows Communication Foundation (WCF). WCF was released in November 2006 as part of .NET 3.0.

On the software relevancy scale, WCF is a titan. It has all the correct elements required to build a modern software system including the foundational pattern of all modern software; an extensible, interception-based pipeline. It is the basis for the upcoming Azure Service Fabric and it integrates with the Service Fabric as-is. In fact, the previous three editions of this book showed how to implement all the elements of the Service Fabric already. As such, WCF is as paramount as ever.

WCF has several facets, and is the only technology today that offers interoperability, productivity and extensibility. It is the ultimate interoperability solution, an implementation of a long list of industry standards. To the distributed application developer, it is the easiest way of making remote calls and even queued calls. To the system

developer, it is the next generation of productivity-oriented features, such as transactions and hosting, that provide off-the-shelf plumbing for applications. To the application developer, it is a declarative programming model for structuring applications. And to the architect, it is a tool for building service-oriented applications. WCF is, in actuality, all of those, simply because it was designed that way—to be the unified offering of Microsoft’s disparate technologies.

To me, WCF is a superior development platform, which to a large extent subsumed raw .NET programming. All .NET developers should use WCF, regardless of their application types, sizes, or industry domains. WCF is a fundamental technology that provides an easy and clean way to generate services and applications in compliance with what I regard as sound design principles. WCF was engineered from the ground up to simplify application development and deployment and to lower the overall cost of ownership. WCF services allow you to build service-oriented applications, from standalone desktop applications to web-based applications and services to high-end Enterprise applications.

How This Book Is Organized

This book covers the topics and skills you need to design and develop service-oriented WCF-based applications, illustrating how to take advantage of built-in features such as service hosting, instance management, concurrency management, transactions, disconnected queued calls, security, and the new Windows Azure Service Fabric. But in fact, this book is not about WCF. It is about modern software engineering, using WCF as a medium to express these design ideas. The various concepts discussed in this book transcend technologies. So while the book shows you how to use these features in WCF, it focuses on the “why” and on the rationale behind particular design decisions. You’ll learn about not only WCF programming and the related system issues, but also relevant design options, tips, best practices, and pitfalls. I approach almost every topic and aspect from a software engineering standpoint, because my objective is to make you not just a WCF expert, but also a better software engineer. Armed with the insights this text provides, you can engineer your applications for maintainability, extensibility, reusability, and productivity.

This fourth edition has provided me with several opportunities: first, to catch up with WCF in .NET 4.6 with its new features of parallel processing and configuration. Second, I wanted to present the upcoming Azure Service Fabric, which is a fundamentally disruptive technology because of the sort of applications it allows developers to build and how it unifies development and operations. Third, I have had a few more years’ worth of WCF techniques, ideas, and helper classes, as well as improvement of the ideas I had in the previous editions. I believe this new material will make this edition valuable even to readers of the third edition.

This book avoids many implementation details of WCF and largely confines its coverage to the possibilities and practical aspects of using WCF: how to apply the technology and how to choose among the available design and programming models. It makes the most of what .NET 4.6 and the Service Fabric has to offer, and in some respects is an advanced C# book as well.

In addition, the book contains many useful utilities, tools, and helper classes I have written, collectively known as *ServiceModelEx*. My tools, helper classes, and attributes aim at increasing your productivity and the quality of your WCF services. *ServiceModelEx* is literally a small framework that sits on top of WCF and compensates for some oversights in its design. *ServiceModelEx* also simplifies and automates certain tasks. This book is as much about my tools, ideas, and techniques as it is about native WCF, and my framework also demonstrates how you can extend WCF. Many readers have told me that aside from the explanations in this book, *ServiceModelEx* is the most valuable asset the book offers. I have also kept to my guideline that, in principle, readers should not have to use all (or any part) of *ServiceModelEx*. In practice, *ServiceModelEx* is your WCF power tools collection. You can also use each helper class, utility, or framework individually, as there are few, if any, interdependencies.

Each chapter addresses a single topic and discusses it in depth. However, the chapters often rely on those that precede them, so you should read the book in order.

Here is a brief summary of the chapters and appendixes in this book:

Chapter 1, WCF Essentials

This first chapter starts by explaining what WCF is, then describes essential WCF concepts and building blocks (such as addresses, contracts, bindings, endpoints, hosting, and clients) and key concepts such as reliability and transport sessions. The chapter includes a discussion of the WCF architecture, which is really the linchpin of all that follows in the subsequent chapters. This chapter assumes that you understand the basic motivation and benefit of service orientation. If that is not the case, you should first read [Appendix A](#). Even if you are already familiar with the basic concepts of WCF, I recommend that you give this chapter at least a cursory reading, not only to ensure that you have a solid foundation, but also because some of the helper classes and terms introduced here will be used and extended throughout the book.

Chapter 2, Service Contracts

[Chapter 2](#) is dedicated to the topic of designing and working with service contracts. First, it covers some useful techniques for service contract overloading and inheritance, as well as some advanced techniques. The chapter also discusses how to design and factor contracts that cater to reuse, maintainability, and extensibility. It ends by showing you how to interact programmatically with the metadata of the exposed contracts at runtime.

Chapter 3, Data Contracts

Chapter 3 deals with how the client and the service can exchange data without ever actually sharing the data type itself or using the same development technology. In this chapter, you will see how to deal with some interesting real-life issues, such as data versioning, and how to pass collections of items. This chapter also extends the contract design and factoring discussion of **Chapter 2** to data contracts.

Chapter 4, Instance Management

This chapter answers the question of which service instance handles which client's request. WCF supports several service instance management, activation, and lifetime management techniques, and your choices will have drastic implications for scalability, performance, the programming model, and the business workflow. This chapter presents the rationale behind each of the instance management modes, offers guidelines on when and how to best use them, and also addresses some related topics, such as durability and throttling.

Chapter 5, Operations

Chapter 5 deals with the types of operations clients can invoke on a service and related design guidelines, such as how to improve on and extend the basic WCF offering to support callback setup and teardown, manage callback ports and channels, and provide for type-safe duplex proxies.

Chapter 6, Faults

This chapter discusses the best practices of error handling, enabling you to decouple the client's error handling from the service's. When required, the chapter shows how services can report errors and exceptions back to their clients, since constructs such as exceptions and exception handling are technology-specific and should not transcend the service boundary. This chapter also demonstrates how you can extend and improve on WCF's basic error-handling mechanism.

Chapter 7, Transactions

This chapter begins by explaining the motivation for transactions in general, then discusses the many aspects of transactional services: the transaction management architecture, transaction propagation configuration, the declarative transaction support offered by WCF, and how clients can create transactions. The chapter ends by discussing relevant design guidelines such as transactional service state management and instancing modes.

Chapter 8, Concurrency Management

Chapter 8 first describes the powerful yet simple declarative way WCF offers for managing concurrency and synchronization, both for the client and the service. The chapter then presents more advanced aspects of concurrency management,

such as callbacks, reentrancy, thread affinity, and synchronization context, best practices and guidelines for avoiding deadlocks, and asynchronous call management. The chapter also explores the Task-based Asynchronous Pattern (TAP) and how this style of asynchrony relates to WCF, highlighting best practice techniques, as well as, discussing potential pitfalls.

Chapter 9, Queued Services

Chapter 9 shows how clients can queue up calls to services, thus enabling asynchronous, disconnected work. The chapter starts by showing how to set up and configure queued services, then focuses on aspects such as transactions, instance management, and failures and their impact on both the business model of the service and its implementation. It also presents techniques for streamlining queues, call management, and several original design ideas (such as a queued response service).

Chapter 10, Security

This chapter demystifies service-oriented security by breaking down this multi-faceted task into its basic elements, such as message transfer, authentication, and authorization. It also demonstrates how to provide security for key scenarios such as intranet and Internet applications. Finally, it presents my framework for declarative WCF security, designed to automate security setup and to considerably simplify managing security.

Chapter 11, The Azure Service Fabric

Chapter 11 introduces you to the Azure Service Fabric. This chapter starts by briefly explaining the value and rationale behind the new platform and continues by presenting the Service Fabric's many capabilities. The chapter ends by presenting `ServiceModelEx.ServiceFabric`, our operational Service Fabric that simplifies how you prepare for the Service Fabric and allows you to deploy Service Fabric code to any Windows environment that can host WCF services.

Appendix A, Introduction to Service Orientation

This appendix is designed for readers who want to understand what service orientation is all about: it presents my take on service orientation and puts it in a concrete context. The appendix defines service-oriented applications (as opposed to mere architecture) and the services themselves and examines the benefits of the methodology. It then presents the principles of service orientation and augments the abstract tenets with a few more practical points required by most applications. In this appendix, I also share my perspective on where SOA and WCF are heading.

Appendix B, Headers and Contexts

This appendix introduces two distinct techniques for enabling the client to pass out-of-band parameters to the service, resulting in a custom logical context: you

will see how to use either the message headers or the context binding to achieve this goal. This appendix also presents my helper classes, which greatly simplify and encapsulate the required programming. These helper classes and custom contexts are used in several places in the book.

Appendix C, Discovery

This appendix starts by describing the basic offering of service discovery and availability announcements introduced in .NET 4.0. Discovery simplifies service and client deployment and management, and allows for great volatility across time and deployment sites. The appendix then shows some simple techniques and helper classes you can use to streamline the programming model. The appendix ends with our technique for TCP-based discovery that mimics regular WCF discovery, benefiting from discovery even where UDP is disallowed.

Appendix D, Publish-Subscribe Service

Appendix D presents several techniques for implementing a publish-subscribe event management solution. It starts with my framework, which lets you develop a publishing and a subscription service in, at most, one line of code. The appendix ends with a discovery-enabled publish-subscribe solution that requires no explicit subscriptions steps.

Appendix E, Generic Interceptor

This appendix presents a general-purpose extensible framework for intercepting calls to your WCF services. It walks through the technique and thought process behind such an extension and shows two examples of how to utilize this simple yet powerful and useful technique.

Appendix F, WCF Coding Standard

Appendix F is basically a consolidated list of all the best practices and dos and don'ts mentioned throughout this book. The standard is all about the “how” and the “what,” not the “why.” The rationale behind it is found in the rest of the book. The standard also uses the terms and helper classes discussed in this book.

Appendix G, ServiceModelEx Catalog

The final appendix presents a catalog of the 100 or so public helper types (excluding internal helper types) of ServiceModelEx mentioned in the book, arranged by categories and techniques, with a short description of each.

Introducing Michael “Monty” Montgomery

In 2009 at my annual Architect’s Master Class I met Monty, who at the time was the chief architect of a large enterprise. Over the next few years we kept in touch since Monty was one of the best architects I have ever met. Monty has a natural knack for service-oriented design and the related system issues. He also proved to be a WCF expert. In 2012 he finally made the jump and joined IDesign. The reason I wanted

him had almost nothing to do with his technical expertise. Monty has a fundamental appreciation for what it takes to make a developer team of a broad spectrum of acumen productive. Over the years Monty has mentored many architects on how to achieve just that objective, so he was a natural fit to help with this fourth edition of the WCF ‘Bible’. Together we bounced new and old ideas off each other, culminating in the new techniques you read in this book.

With the addition of a second author, I wanted to keep the original author’s voice from the previous three editions when referring to personal opinions or contributions. As with the previous editions, when the text mentions “I” or “my,” it refers to Juval. When it says “we” or “our,” it refers to both Juval and Monty.

Some Assumptions About the Reader

I assume that you, the reader, are an experienced developer and that you are comfortable with object-oriented concepts such as encapsulation and inheritance. I will take advantage of your existing understanding of object and component technology and terminology, and port that knowledge to WCF. You should ideally have a fair understanding of .NET and know C# 4.5. Although the book uses C# for the most part, it is just as pertinent to Visual Basic developers.

What You Need to Use This Book

To use this book, you will need .NET 4.5 or .NET 4.6 and Visual Studio 2015 or 2016. For the Service Fabric appendix, you will need the Windows Azure Service Fabric SDK. Unless I explicitly mention otherwise, the contents apply to any platform that can run WCF. You may also install additional Windows components, such as MSMQ and IIS.

Conventions Used in This Book

The following typographical conventions are used in this book:

Italic

Used for technical terms, URLs, email addresses, filenames, and file extensions.

Constant width

Used for code samples, statements, namespaces, classes, assemblies, interface directives, operators, attributes, and reserved words.

Constant width bold

Used for emphasis in code samples.

Constant width italic

Shows text that should be replaced with user-supplied values or by values determined by context.



This element signifies a tip or suggestion.



This element signifies a general note.



This element indicates a warning or caution.

Whenever I wish to make a point in a code sample, I do so with the static `Assert` method of the `Debug` class:

```
int number = 1+2;
Debug.Assert(number == 3);
```

The `Assert` method accepts a Boolean value and throws an exception when it is false.

The book follows the recommended naming guidelines and coding style available at [iDesign](#). Whenever it deviates from that standard, it is likely the result of space or line-length constraints. As for naming conventions, I use “Pascal casing” for public member methods and properties; this means the first letter of each word in the name is capitalized. For local variables and method parameters I use “camel casing,” in which the first letter of each word in the name is capitalized, with the exception of the first word. The names of private member variables are prefixed with `m_`:

```
class SomeClass
{
    int m_Number;

    public int Number
        {get;set;};
}
```

I use ellipses between curly braces to indicate the presence of code that is necessary but unspecified:

```
class SomeClass  
{...}
```

In the interest of clarity and space, code examples often do not contain all the `using` statements needed to specify all the namespaces the examples require; instead, such examples include only the new namespaces introduced in the preceding text.

Safari® Books Online



Safari Books Online is an on-demand digital library that delivers expert **content** in both book and video form from the world's leading authors in technology and business.

Technology professionals, software developers, web designers, and business and creative professionals use Safari Books Online as their primary resource for research, problem solving, learning, and certification training.

Safari Books Online offers a range of **plans and pricing** for **enterprise, government, education**, and individuals.

Members have access to thousands of books, training videos, and prepublication manuscripts in one fully searchable database from publishers like O'Reilly Media, Prentice Hall Professional, Addison-Wesley Professional, Microsoft Press, Sams, Que, Peachpit Press, Focal Press, Cisco Press, John Wiley & Sons, Syngress, Morgan Kaufmann, IBM Redbooks, Packt, Adobe Press, FT Press, Apress, Manning, New Riders, McGraw-Hill, Jones & Bartlett, Course Technology, and hundreds **more**. For more information about Safari Books Online, please visit us [online](#).

How to Contact Us

Please address comments and questions concerning this book to the publisher:

O'Reilly Media, Inc.
1005 Gravenstein Highway North
Sebastopol, CA 95472
800-998-9938 (in the United States or Canada)
707-829-0515 (international or local)
707-829-0104 (fax)

We have a web page for this book, where we list errata, examples, and any additional information. You can access this page at <http://bit.ly/prog-wcf-services>.

You can also contact the author at: <http://www.idesign.net>.

The author has posted a comprehensive code library on the IDesign website with more than 150 downloads related to WCF essentials, contract design, instance man-

agement, operations and calls, faults, transactions, concurrency, queuing, security, and the service bus. The downloads articulate many of the code snippets in this book in a working fashion. However, because of the large number of the downloads, the maintenance involved, and the fact that these secondary, accompanying examples do not themselves appear in the book, they are provided separately from the official sources.

To comment or ask technical questions about this book, send email to bookquestions@oreilly.com.

For more information about our books, courses, conferences, and news, see our website at <http://www.oreilly.com>.

Find us on Facebook: <http://facebook.com/oreilly>

Follow us on Twitter: <http://twitter.com/oreillymedia>

Watch us on YouTube: <http://www.youtube.com/oreillymedia>

Acknowledgments

I would not have been able to come to terms with WCF in its early days without the constant support of and interaction with the WCF (then Indigo) program managers. I am especially grateful to Steve Swartz, one of the WCF architects, not just for his knowledge and insight, but also for his patience with me and those long IM sessions. Thanks go to Yasser Shohoud, Doug Purdy, and Shy Cohen for the fascinating strategic design reviews and to Krish Srinivasan for his almost philosophical approach to engineering. Working with you guys has been the best part of learning WCF and a privilege in its own right. The following WCF program managers also shared their time and helped clarify WCF: Andy Milligan, Brian McNamara, Eugene Osovetsky, Kenny Wolf, Kirill Gavrylyuk, Max Feingold, Michael Maruchek, Mike Vernal, and Steve Millet. Thanks also to the group manager, Angela Mills.

Outside Microsoft, I am grateful to Nicholas Paldino for his help. This is the sixth book that Nick has reviewed for me, and by now I am hooked. I simply cannot imagine going to print without Nick's insight, meticulous scrutiny, and diligent help, to say nothing about Nick's technical competence, which is nothing short of awe-striking and humbling. The books are always a superior product as a result of his editing and feedback.

I dedicate the book to the IDesign Alumni. A unique post-class benefit of attending one of the IDesign Master Classes is joining this private forum where you can discuss practicing our techniques, share observation and insights, and ask pertinent questions. The members of the forum are all world-class developers, architects, and project leads, and the discussion amongst this group is profound, insightful, and of the highest quality. Many of the ideas and techniques in this book were first tried on

and trialed by the fire of the Alumni. The best emails I get each day are the discussion of the Alumni forum, and the level of mutual involvement, empathy, excellence, professional integrity and responsiveness is second to none. It is truly inspiring that a community of mostly virtual peers can care about their trade and each other that much. I thank you for every time you made us stop and think about the best way of answering a question, a smarter way of solving an Alumni's design problem, or the need for an improved technique. It is people like you that we write books for.

WCF Essentials

This chapter describes the essential concepts and building blocks of Windows Communication Foundation (WCF) and its architecture enabling you to build simple services. You will learn the basic terms regarding addresses, bindings, contracts, and endpoints; see how to host a service; learn how to write a client; understand related topics, such as in-process (in-proc) hosting reliability, and transport sessions; and see how to utilize WCF in Visual Studio. Even if you are already familiar with the basic concepts of WCF, I recommend that you give this chapter at least a cursory reading, not only to ensure that you have a solid foundation, but also because some of the helper classes and terms introduced here will be used and extended throughout the book.

What Is WCF?

WCF is a software development kit for developing and deploying services on Windows (I will describe what a service is in the next section). But WCF is much more—it is literally a better .NET. WCF provides a runtime environment for your services, enabling you to expose Common Language Runtime (CLR) types as services and to consume other services as CLR types. Although in theory you could build services without WCF, in practice, building services is significantly easier with WCF. WCF is Microsoft's implementation of a set of industry standards defining service interactions, type conversions, marshaling, and the management of various protocols. Consequently, WCF provides interoperability between services.

WCF provides developers with the essential off-the-shelf plumbing required by almost all applications, and as such, it greatly increases productivity. The first release of WCF (as part of .NET 3.0) provided many useful facilities for developing services, such as hosting, service instance management, asynchronous calls, reliability, transaction management, disconnected queued calls, and security. The second release of

WCF (as part of .NET 3.5) provided additional tools and extended the original offering with additional communication options. The third release (as part of .NET 4.0) included configuration changes, a few extensions, and the new features of discovery (discussed in [Appendix C](#)) and routers (not discussed in this book). With WCF's fourth release (as part of .NET 4.5), WCF has several new simplification features and additional bindings, including UDP and WebSocket bindings. WCF has an elegant extensibility model you can use to enrich the basic offering. In fact, WCF itself is written using this extensibility model. This book is dedicated to exploring these aspects and features.

WCF is part of .NET 4.5, so it can run only on operating systems that support it. Presently, this list consists of Windows XP and later and Windows Server 2003 and later.

Most of the WCF functionality is included in a single assembly called *System.ServiceModel.dll*, located in the `System.ServiceModel` namespace.

Azure Service Bus

Microsoft originally designed the Azure Service Bus to solve the broad and challenging connectivity issues of consuming web services over the Internet.

The Azure Service Bus has since established itself as an indispensable tool for the modern developer who is creating connected systems on any platform. The Azure Service Bus has grown into a fully fledged connectivity platform with a diverse set of features.

Recognizing that the challenges of providing service connectivity are not limited to the Internet, Microsoft has released an Intranet version of the Azure Service Bus named *Service Bus for Windows Server*. This version of the Service Bus targets the on-premises versions of Windows Server and SQL Server.

As a connectivity platform, the Service Bus offers a wide array of options for distributed computing solutions. These include service relay features, event publishing and subscription, persistent queues, and topics (queued events). The Service Bus also supports a modern, federated security model through claim-based access control. The Service Bus now supports a variety of programming models including classic WCF, REST-based and even Advanced Message Queuing Protocol (AMQP). From this set of building blocks, developers can create even more sophisticated patterns that enhance the scalability, throughput, discoverability, availability, and security of their solutions. The previous edition of this book devoted an in-depth chapter and a number of appendix sections to the Service Bus, along with several extensions and original enhancements.

We consider the Service Bus nothing short of a fundamentally disruptive, essential platform for modern distributed computing. There is no question the Service Bus

now merits a book in its own right to fully explore its many features and patterns in depth. We have therefore removed the chapters on the Service Bus from this edition. The material from those chapters will appear in greater and up-to-date detail in our forthcoming book dedicated to the Service Bus platform.

Services

A *service* is a unit of functionality exposed to the world. In that respect, it is the next evolutionary step in the long journey from functions to objects to components to services. *Service orientation* (SO) is an abstract set of principles and best practices for building service-oriented applications. [Appendix A](#) provides a concise overview and outlines the motivation for using this methodology. The rest of this book assumes you are familiar with these principles. A *service-oriented application* aggregates services into a single logical application, similar to the way a component-oriented application aggregates components and an object-oriented application aggregates objects, as shown in [Figure 1-1](#).

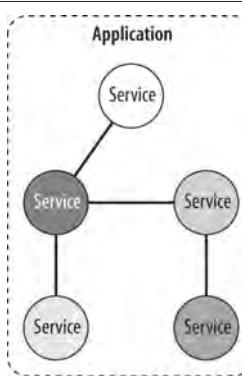


Figure 1-1. A service-oriented application

The services can be local or remote, can be developed by multiple parties using any technology, can be versioned independently, and can even execute on different timelines. Inside a service, you will find concepts such as languages, technologies, platforms, versions, and frameworks, yet between services, only prescribed communication patterns are allowed.

The *client* of a service is merely the party consuming its functionality. The client can be literally anything—for instance, a Windows Forms, WPF, Silverlight, or Windows Store App class, an ASP.NET page, or another service.

Clients and services interact by sending and receiving messages. Messages may be transferred directly from the client to the service or be sent via an intermediary such

as the Azure Service Bus. With WCF, messages are SOAP messages. These messages are independent of transport protocols—unlike web services, WCF services may communicate over a variety of transports (not just HTTP). WCF clients may interoperate with non-WCF services, and WCF services can interact with non-WCF clients. That said, if you develop both the client and the service, you can typically construct the application so that both ends require WCF in order to utilize WCF-specific advantages.

Because the making of the service is opaque from the outside, a WCF service typically exposes *metadata* describing the available functionality and possible ways of communicating with the service. The metadata is published in a predefined, technology-neutral way, such as using Web Services Description Language (WSDL) over HTTP-GET or an industry standard for metadata exchange over any protocol. A non-WCF client can import the metadata to its native environment as native types. Similarly, a WCF client can import the metadata of a non-WCF service and consume it as native CLR classes and interfaces.

Service Execution Boundaries

With WCF, the client never interacts with a service directly, even when dealing with a local, in-memory service. Instead, the client always uses a proxy to forward calls to the service. The proxy exposes the same operations as the service, plus some proxy-management methods.

WCF allows the client to communicate with a service across all execution boundaries. On the same machine, the client can consume services in the same app domain, across app domains in the same process, or across processes (see [Figure 1-2](#)).

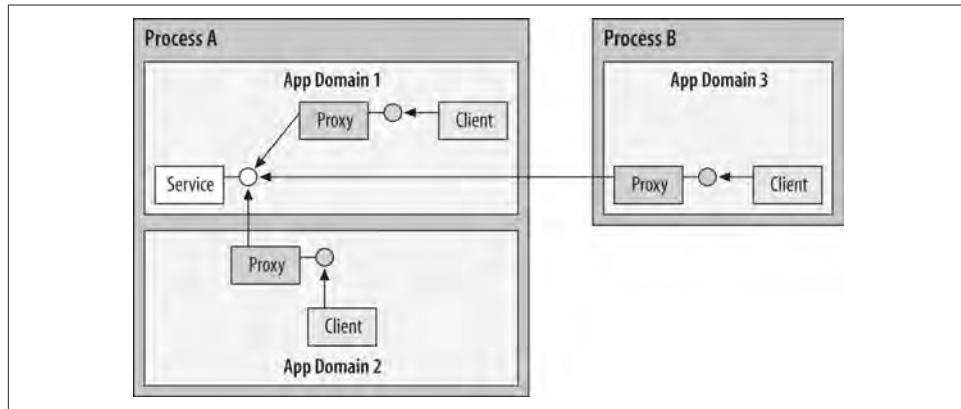


Figure 1-2. Same-machine communication using WCF

Across machine boundaries ([Figure 1-3](#)), the client can interact with services in its intranet or across the Internet.

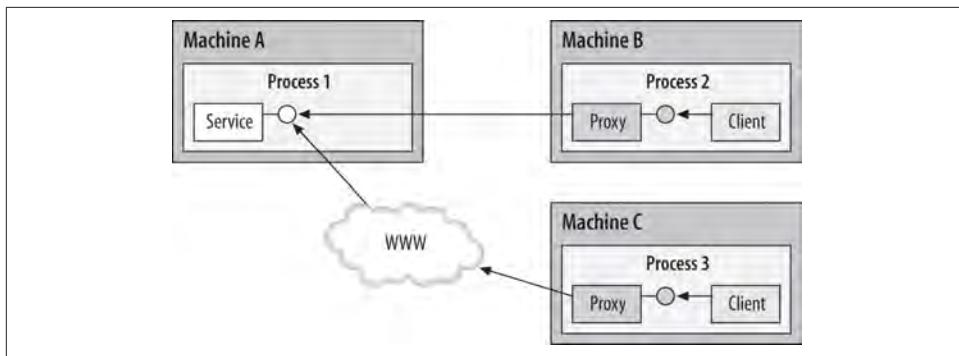


Figure 1-3. Cross-machine communication using WCF

WCF and Location Transparency

In the past, distributed computing technologies such as DCOM and .NET remoting aspired to provide the same programming model to the client regardless of whether the object was local or remote. In the case of a local call, the client used a direct reference, and when dealing with a remote object, the client used a proxy. The problem with trying to use the local programming model as the remote programming model was that there is much more to a remote call than an object with a wire. Complex issues such as lifecycle management, reliability, state management, and security reared their heads, making the remote programming model significantly more complex. Numerous problems arose, all because the remote object was trying to be what it is not—a local object.

WCF also strives to provide the client with the same programming model regardless of the location of the service. However, the WCF approach is the exact opposite: it takes the remote programming model of instantiating and using a proxy and uses it even in the most local case. Because all interactions are done via a proxy, requiring the same configuration and hosting, WCF maintains the same programming model for the local and remote cases; thus, it not only enables you to switch locations without affecting the client, but also significantly simplifies the application programming model. Another important benefit of always using a proxy is that it enables WCF to intercept the calls and add its value, as you will see later on.

Addresses

In WCF, every service is associated with a unique address. The address provides two important elements: the location of the service and the transport protocol, or *transport scheme*, used to communicate with the service. The location portion of the address indicates the name of the target machine, site, or network; a communication port, pipe, or queue; and an optional specific path, or *Uniform Resource Identifier*

(URI). A URI can be any unique string, such as the service name or a globally unique identifier (GUID).

Out of the box, WCF supports the following transport schemes:

- HTTP/HTTPS
- TCP
- IPC
- MSMQ
- Service Bus
- WebSocket
- UDP

Addresses always have the following format:

[base address]/[optional URI]

The base address is always in this format:

[transport]://[machine or domain][:optional port]

Here are a few sample addresses:

```
http://localhost:8001  
http://localhost:8001/MyService  
net.tcp://localhost:8002/MyService  
net.pipe://localhost/MyPipe  
net.msmq://localhost/private/MyQueue  
net.msmq://localhost/MyQueue  
ws://localhost/MyService  
soap.udp://localhost:8081/MyService
```

The way to read an address such as:

`http://localhost:8001`

is like this: “Using HTTP, go to the machine called `localhost`, where on port 8001 someone is waiting for my calls.”

If there is also a URI, as in:

`http://localhost:8001/MyService`

the address will read as follows: “Using HTTP, go to the machine called `localhost`, where on port 8001 someone called `MyService` is waiting for my calls.”

TCP Addresses

TCP addresses use `net.tcp` for transport and typically include a port number, as in:

```
net.tcp://localhost:8002/MyService
```

When a port number is not specified, the TCP address defaults to port 808:

```
net.tcp://localhost/MyService
```

It is possible for two TCP addresses (from the same host, as discussed later in this chapter) to share a port:

```
net.tcp://localhost:8002/MyService  
net.tcp://localhost:8002/MyOtherService
```

TCP-based addresses are used throughout this book.



You can configure TCP-based addresses from different service hosts to share a port.

HTTP Addresses

HTTP addresses use `http` for transport and can also use `https` for secure transport. You typically use HTTP addresses with outward-facing Internet-based services, and you can specify a port as shown here:

```
http://localhost:8001
```

If you do not specify the port number, it defaults to 80 (and port 443 for HTTPS). As with TCP addresses, two HTTP addresses from the same host can share a port, even on the same machine.

HTTP-based addresses are also used throughout this book.

IPC Addresses

Inter-process communication (IPC) addresses use `net.pipe` for transport, to indicate the use of the Windows named pipe mechanism. In WCF, services that use IPC can only accept calls from the same machine. Consequently, you must specify either the explicit local machine name or `localhost` for the machine name, followed by a unique string for the pipe name:

```
net.pipe://localhost/MyPipe
```

You can open a named pipe only once per machine, so it is not possible for two named pipe addresses to share a pipe name on the same machine.

IPC-based addresses are used throughout this book.



The IPC address format as provided by Microsoft is incorrect, indicating the mechanism instead of the protocol. The correct scheme format should have been `net.ipc` instead of `net.pipe`, much like the TCP address uses `net.tcp` rather than `net.socket`.

MSMQ Addresses

MSMQ addresses use `net.msmq` for transport, to indicate the use of the Microsoft Message Queue (MSMQ). You must specify the queue name. When you're dealing with private queues, you must also specify the queue type, but you can omit that for public queues:

```
net.msmq://localhost/private/MyService  
net.msmq://localhost/MyService
```

[Chapter 9](#) is dedicated to making queued calls.

WebSocket Addresses

WebSocket addresses are unique in that they are asymmetrical between client and service. The client uses `ws` for transport and `wss` for secure transport, while the service always uses `http` or `https` respectively. WebSocket addresses are required when you need callbacks over the Internet, and you can specify a port as shown here:

```
ws://localhost:8080
```

If you do not specify the port number, a WebSocket defaults to the standard HTTP port 80 (and port 443 for wss or HTTPS). As with TCP addresses, two WebSocket addresses on the same machine from the same host can share a port.

UDP Addresses

UDP addresses use `soap.udp` for transport, to indicate SOAP over UDP is being used. You can also specify a port as shown here:

```
soap.udp://localhost:8081
```

Contracts

In WCF, all services expose *contracts*. The contract is a platform-neutral and standard way of describing what the service does. WCF defines four types of contracts:

Service contracts

Describe which operations the client can perform on the service. Service contracts are the subject of [Chapter 2](#), but they are used extensively in every chapter in this book.

Data contracts

Define which data types are passed to and from the service. WCF defines implicit contracts for built-in types such as `int` and `string`, but you can easily define explicit opt-in data contracts for custom types. [Chapter 3](#) is dedicated to defining and using data contracts, and subsequent chapters make use of data contracts as required.

Fault contracts

Define which errors are raised by the service and how the service handles and propagates errors to its clients. [Chapter 6](#) is dedicated to defining and using fault contracts.

Message contracts

Allow the service to interact directly with messages. Message contracts can be typed or untyped and are useful in interoperability cases when another party has already dictated some explicit (typically proprietary) message format. This, however, is by no means the usual case for common WCF applications, so this book makes no use of message contracts. Unless you are required to leverage the flexibility, power, and extensibility of message contracts, you should avoid them, as they add no value, but do add complexity. In many cases, the desire to use message contracts indicates a need for a custom application context, which you can address using custom headers (a useful alternative technique used throughout this book). For more on message headers, see [Appendix B](#).

The Service Contract

The `ServiceContractAttribute` is defined as:

```
[AttributeUsage(AttributeTargets.Interface|AttributeTargets.Class,
    Inherited = false)]
public sealed class ServiceContractAttribute : Attribute
{
    public string Name
    {get;set;}
    public string Namespace
    {get;set;}
    //More members
}
```

This attribute allows you to define a service contract. You can apply the attribute on an interface or a class, as shown in [Example 1-1](#).

Example 1-1. Defining and implementing a service contract

```
[ServiceContract]
interface IMyContract
{
```

```

[OperationContract]
string MyMethod(string text);

//Will not be part of the contract
string MyOtherMethod(string text);
}
class MyService : IMyContract
{
    public string MyMethod(string text)
    {
        return "Hello " + text;
    }
    public string MyOtherMethod(string text)
    {
        return "Cannot call this method over WCF";
    }
}

```

The `ServiceContract` attribute maps a CLR interface (or inferred interface, as you will see later) to a technology-neutral service contract. The `ServiceContract` attribute exposes a CLR interface (or a class) as a WCF contract independently of that type's visibility. The type visibility has no bearing on WCF, because visibility is a CLR concept. Applying the `ServiceContract` attribute on an internal interface exposes that interface as a public service contract, ready to be consumed across the assembly boundary. Without the `ServiceContract` attribute, the interface is not visible to WCF clients, in line with the service-oriented tenet that service boundaries should be explicit. To enforce that tenet, all contracts must explicitly opt in: only interfaces (or classes) decorated with the `ServiceContract` attribute will be considered WCF contracts; other types will not.

In addition, none of the members of the type will ever be part of the contract when using the `ServiceContract` attribute. You must explicitly indicate to WCF which methods to expose as part of the WCF contract using the `OperationContractAttribute`, defined as:

```

[AttributeUsage(AttributeTargets.Method)]
public sealed class OperationContractAttribute : Attribute
{
    public string Name
    {get;set;}
    //More members
}

```

You can apply the `OperationContract` attribute only on methods, not on properties, indexers, or events, which are CLR concepts. WCF only understands *operations*—logical functions—and the `OperationContract` attribute exposes a contract method as a logical operation to perform as part of the service contract. Other methods on the interface (or class) that do not have the `OperationContract` attribute will not be part

of the contract. This enforces explicit service boundaries and maintains an explicit opt-in model for the operations themselves. In addition, a contract operation cannot use object references as parameters: only primitive types or data contracts are allowed.

Applying the `ServiceContract` attribute

WCF lets you apply the `ServiceContract` attribute on an interface or on a class. When you apply it on an interface, some class needs to implement that interface. In general, you use plain C# or VB to implement the interface, and nothing in the service class code pertains to it being a WCF service:

```
[ServiceContract]
interface IMyContract
{
    [OperationContract]
    string MyMethod();
}
class MyService : IMyContract
{
    public string MyMethod()
    {
        return "Hello WCF";
    }
}
```

You can use implicit or explicit interface implementation:

```
class MyService : IMyContract
{
    string IMyContract.MyMethod()
    {
        return "Hello WCF";
    }
}
```



Because the client can never use the service class directly and must always go through a proxy, using explicit interface implementation is less important in WCF than it is in regular .NET programming.

A single class can support multiple contracts by deriving and implementing multiple interfaces decorated with the `ServiceContract` attribute:

```
[ServiceContract]
interface IMyContract
{
    [OperationContract]
    string MyMethod();
```

```

}
[ServiceContract]
interface IMyOtherContract
{
    [OperationContract]
    void MyOtherMethod();
}
class MyService : IMyContract,IMyOtherContract
{
    public string MyMethod()
    {...}
    public void MyOtherMethod()
    {...}
}

```

There are, however, a few implementation constraints on the service implementation class. You should avoid parameterized constructors, because WCF will only use the default constructor. Also, although the class can use internal properties, indexers, and static members, no WCF client will ever be able to access them.

WCF also lets you apply the `ServiceContract` attribute directly on the service class, without defining a separate contract first:

```

//Avoid
[ServiceContract]
class MyService
{
    [OperationContract]
    string MyMethod()
    {
        return "Hello WCF";
    }
}

```

Under the covers, WCF will infer the contract definition. You can apply the `OperationContract` attribute on any method of the class, be it private or public.



Avoid using the `ServiceContract` attribute directly on the service class. Always define a separate contract so that you can both consume it independently of the class and have other classes implement it.

Names and namespaces

You can and should define a namespace for your contract. The contract namespace serves the same purpose in WCF as it does in .NET programming: to scope a type of contract and reduce the overall chance of a collision. You can use the `Namespace` property of the `ServiceContract` attribute to provide a namespace:

```
[ServiceContract(Namespace = "MyNamespace")]
interface IMyContract
{...}
```

Unspecified, the contract namespace defaults to <http://tempuri.org>. For outward-facing services, you typically use your company's URL, and for intranet services, you can use any meaningful unique name, such as `MyApplication`.

By default, the exposed name of the contract will be the name of the interface. However, you can use an alias for a contract to expose a different name to the clients in the metadata, by using the `Name` property of the `ServiceContract` attribute:

```
[ServiceContract(Name = "IMyContract")]
interface IMyOtherContract
{...}
```

Similarly, the name of the publicly exposed operation defaults to the method name, but you can use the `Name` property of the `OperationContract` attribute to alias it to a different publicly exposed name:

```
[ServiceContract]
interface IMyContract
{
    [OperationContract(Name = "SomeOperation")]
    void MyMethod(string text);
}
```

You will see a use for these properties in [Chapter 2](#).

Hosting

The WCF service class cannot exist in a void. Every WCF service must be hosted in a Windows process called the *host process*. A single host process can host multiple services, and the same service type can be hosted in multiple host processes. WCF has no restrictions regarding whether or not the host process is also the client process, although having a separate process promotes fault and security isolation. It is also immaterial who provides the process and what kind of process is involved. The host can be provided by Internet Information Services (IIS), by the Windows Activation Service (WAS) on versions of Windows before Windows Server 2008 R2 and Windows 7, or by the developer as part of the application.



In-process (or *in-proc*) hosting, where the service resides in the same process as the client, is a special case. By definition, the developer provides the host for the in-proc case.

IIS Hosting

The main advantage of hosting a service on the Microsoft IIS web server is that the host process is launched automatically upon the first client request, and IIS manages the lifecycle of the host process. The main disadvantage of IIS hosting is that you can only use HTTP.

Hosting in IIS is very similar to hosting a web service or ASP.NET Web API service. You need to provide a virtual directory under IIS and supply an *.svc* file. IIS uses the *.svc* file to identify the service code behind the file and class. **Example 1-2** shows the syntax for the *.svc* file.

Example 1-2. A .svc file

```
<%@ ServiceHost
    Language      = "C#"
    Debug        = "true"
    CodeBehind   = "~/App_Code/MyService.cs"
    Service      = "MyService"
%>
```



You can even inject the service code inline in the *.svc* file, but that is not advisable.

When you use IIS hosting, the base address used for the service always has to be the same as the address of the *.svc* file.

Using Visual Studio

You can use Visual Studio to generate a boilerplate IIS-hosted service. From the File menu, select New Web Site, then select WCF Service from the New Web Site dialog box. Visual Studio creates a new website, service code, and a matching *.svc* file. You can also use the Add New Item dialog box to add another service later.

The Web.Config file

The website config file (*web.config*) typically lists the types you want to expose as services. You need to use fully qualified type names, including the assembly name if the service type comes from an unreferenced assembly:

```
<system.serviceModel>
    <services>
        <service name = "MyNamespace.MyService">
            ...
        </service>
```

```
</services>
</system.serviceModel>
```

Instead of defining an .svc file, you can provide the service type and its address information directly in the application *web.config* file in the **serviceHostingEnvironment** section. In fact, you can list as many services as you like there:

```
<system.serviceModel>
  <serviceHostingEnvironment>
    <serviceActivations>
      <add relativeAddress = "MyService.svc" service =
        "MyNamespace.MyService"/>
      <add relativeAddress = "MyOtherService.svc" service =
        "MyOtherService"/>
    </serviceActivations>
  </serviceHostingEnvironment>
  <services>
    <service name = "MyNamespace.MyService">
      ...
    </service>
    <service name = "MyOtherService">
      ...
    </service>
  </services>
</system.serviceModel>
```

Self-Hosting

Self-hosting is the technique in which the developer is responsible for providing and managing the lifecycle of the host process. Use self-hosting when you want a process (or machine) boundary between the client and the service and when you are using the service in-proc—that is, in the same process as the client. You can provide any Windows process, such as a Windows Forms application, a WPF application, a Console application, or a Windows Service. Note that the process must be running before the client calls the service, which typically means you have to prelaunch it. This is not an issue for Windows Services or in-proc hosting. You can provide a host with only a few lines of code. Unlike IIS, a self-hosted service can use any WCF transport protocol, and you can take advantage of all the WCF features, including the Service Bus, discovery, and utilizing a singleton service.

As with IIS hosting, the hosting application config file (*app.config*) typically lists the types of the services you wish to host and expose to the world:

```
<system.serviceModel>
  <services>
    <service name = "MyNamespace.MyService">
      ...
    </service>
  </services>
</system.serviceModel>
```

In addition, the host process must explicitly register the service types at runtime and open the host for client calls, which is why the host process must be running before the client calls arrive. Creating the host is typically done in the `Main()` method using the class `ServiceHost`, defined in [Example 1-3](#).

Example 1-3. The ServiceHost class

```
public interface ICommunicationObject
{
    void Open();
    void Close();
    //More members
}
public abstract class CommunicationObject : ICommunicationObject
{...}
public abstract class ServiceHostBase : CommunicationObject, IDisposable, ...
{...}
public class ServiceHost : ServiceHostBase
{
    public ServiceHost(Type serviceType, params Uri[] baseAddresses);
    //More members
}
```

You need to provide the constructor of `ServiceHost` with the service type and optionally with default base addresses. The set of base addresses can be an empty set, and even if you provide base addresses, you can configure the service to use different base addresses. Having a set of base addresses enables the service to accept calls on multiple addresses and protocols and to use only a relative URI.

Note that each `ServiceHost` instance is associated with a particular service type, and if the host process needs to host multiple types of services, you will need a matching number of `ServiceHost` instances. By calling the `Open()` method on the host, you allow calls in, and by calling the `Close()` method, you gracefully exit the host instance, allowing calls in progress to complete while refusing future client calls even if the host process is still running. Closing the service host is typically done when the host process shuts down. For example, to host this service in a Windows Forms application:

```
[ServiceContract]
interface IMyContract
{...}
class MyService : IMyContract
{...}
```

you would write the following hosting code:

```
static void Main()
{
    ServiceHost host = new ServiceHost(typeof(MyService));
```

```

host.Open();

//Can do blocking calls:
Application.Run(new MyForm());

host.Close();
}

```

Opening a host loads the WCF runtime and launches worker threads to monitor incoming requests. The monitoring threads dispatch incoming calls to worker threads from the I/O completion thread pool (where there are up to 1,000 threads by default). Because worker threads are involved, you can perform blocking operations after opening the host.

Because the host is closed gracefully, the amount of time it will take is undetermined. By default, the host will block for 10 seconds waiting for `Close()` to return and will proceed with the shutdown after that timeout has expired. Before opening the host, you can configure a different close timeout with the `CloseTimeout` property of `ServiceHostBase`:

```

public abstract class ServiceHostBase : ...
{
    public TimeSpan CloseTimeout
    {get;set;}
    //More members
}

```

For example, you can use programmatic calls to set the close timeout to 20 seconds:

```

ServiceHost host = new ServiceHost(...);
host.CloseTimeout = TimeSpan.FromSeconds(20);
host.Open();

```

You can do the same in a config file by placing the close timeout in the host section of the service:

```

<system.serviceModel>
    <services>
        <service name = "MyNamespace.MyService">
            <host>
                <timeouts
                    closeTimeout = "00:00:20"
                />
            </host>
            ...
        </service>
    </services>
</system.serviceModel>

```

Using Visual Studio

Visual Studio allows you to add a WCF service to any application project by selecting WCF Service from the Add New Item dialog box. A service added this way is, of course, in-proc toward the host process, but out-of-proc clients can also access it.

Self-hosting and base addresses

You can launch a service host without providing any base address by omitting the base addresses altogether:

```
ServiceHost host = new ServiceHost(typeof(MyService));
```



Do not provide a `null` instead of an empty list, because that will throw an exception:

```
ServiceHost host;
host = new ServiceHost(typeof(MyService),null);
```

You can also register multiple base addresses separated by commas, as in the following snippet, as long as the addresses do not use the same transport scheme (note the use of the `params` qualifier in [Example 1-3](#)):

```
Uri tcpBaseAddress = new Uri("net.tcp://localhost:8001/");
Uri httpBaseAddress = new Uri("http://localhost:8002/");

ServiceHost host = new ServiceHost(typeof(MyService),
    tcpBaseAddress,httpBaseAddress);
```

WCF also lets you list the base addresses in the host config file:

```
<system.serviceModel>
  <services>
    <service name = "MyNamespace.MyService">
      <host>
        <baseAddresses>
          <add baseAddress = "net.tcp://localhost:8001/">
          <add baseAddress = "http://localhost:8002/">
        </baseAddresses>
      </host>
      ...
    </service>
  </services>
</system.serviceModel>
```

When you create the host, it will use whichever base addresses it finds in the config file, plus any base addresses you provide programmatically. Take extra care to ensure the configured base addresses and the programmatic ones do not overlap in the scheme.



On any machine running Windows XP or later, for HTTP addresses other than port 80, you will need to launch the host process (or Visual Studio while testing or debugging) as an administrator. Instead of doing that every time, you can instruct Windows to reserve the port namespace for the user running the host. Do this using the `netsh.exe` command-line utility. For example, to reserve the HTTP port 8002 on the local machine, you will need to run this command at a command prompt launched as an administrator:

```
netsh http add urlacl url=http://+:8002/user=
"MachineOrDomain\UserName"
```

You can even register multiple hosts for the same type, as long as the hosts use different base addresses:

```
Uri baseAddress1 = new Uri("net.tcp://localhost:8001/");
ServiceHost host1 = new ServiceHost(typeof(MyService),baseAddress1);
host1.Open();

Uri baseAddress2 = new Uri("net.tcp://localhost:8002/");
ServiceHost host2 = new ServiceHost(typeof(MyService),baseAddress2);
host2.Open();
```

However, with the exception of some threading issues discussed in [Chapter 8](#), opening multiple hosts this way offers no real advantage. In addition, opening multiple hosts for the same type does not work with base addresses supplied in the config file and requires use of the `ServiceHost` constructor.

Advanced hosting features

The `ICommunicationObject` interface that `ServiceHost` supports offers some advanced features, listed in [Example 1-4](#).

Example 1-4. The ICommunicationObject interface

```
public interface ICommunicationObject
{
    void Open();
    void Close();
    void Abort();

    event EventHandler Closed;
    event EventHandler Closing;
    event EventHandler Faulted;
    event EventHandler Opened;
    event EventHandler Opening;

    IAsyncResult BeginClose(AsyncCallback callback,object state);
    IAsyncResult BeginOpen(AsyncCallback callback,object state);
    void EndClose(IAsyncResult result);
```

```

void EndOpen(IAsyncResult result);

CommunicationState State
{
    get; }
    //More members
}
public enum CommunicationState
{
    Created,
    Opening,
    Opened,
    Closing,
    Closed,
    Faulted
}

```

If opening or closing the host is a lengthy operation, you can do so asynchronously with the `BeginOpen()` and `BeginClose()` methods.

You can subscribe to hosting events such as state changes or faults, and you can use the `State` property to query for the host status. Finally, the `ServiceHost` class also offers the `Abort()` method. `Abort()` is an ungraceful exit—when called, it immediately aborts all service calls in progress and shuts down the host. Active clients will each get an exception.

The `ServiceHost<T>` class

You can improve on the WCF-provided `ServiceHost` class by defining the `ServiceHost<T>` class, as shown in [Example 1-5](#).

Example 1-5. The `ServiceHost<T>` class

```

public class ServiceHost<T> : ServiceHost
{
    public ServiceHost() : base(typeof(T))
    {}

    public ServiceHost(params string[] baseAddresses) : base(typeof(T),
        baseAddresses.Select(address=>new Uri(address)).ToArray())
    {}
    public ServiceHost(params Uri[] baseAddresses) : base(typeof(T),baseAddresses)
    {}
}

```

`ServiceHost<T>` provides simple constructors that do not require the service type as a construction parameter and that can operate on raw strings instead of the cumbersome `Uri`. I'll add quite a few extensions, features, and capabilities to `ServiceHost<T>` throughout this book.

WAS Hosting

The problem with hosting in IIS is that it is a web server, not a hosting engine. It therefore requires you to masquerade your service as a website. While ASP.NET encapsulates this step for you, it causes a significant increase in internal complexity, involving the HTTP modules and the ASP.NET pipeline. The problem is that the more moving parts involved, the higher the likelihood of something going wrong. Moreover, limiting the service to using only HTTP makes IIS ill-suited for intranet applications.

With the next wave of Windows, Microsoft rectified this issue by providing a general-purpose hosting engine called the Windows Activation Service (WAS). WAS is a system service available with Windows versions Windows Server 2008 or later. The WAS is a true general-purpose hosting engine. It can host websites (in fact, IIS 7 or later will host its websites in the WAS by default), but it can just as easily host your services, allowing you to use any transport, such as TCP, IPC, or MSMQ. You can install and configure the WAS separately from IIS. Hosting a WCF service in the WAS is designed to look just like hosting in IIS. You need to either supply an .svc file, just as with IIS, or provide the equivalent information in the config file. All the other development aspects, such as support in Visual Studio, remain exactly the same. Because the WAS is a system service, you do not need to pre-launch your service host process. When the first client call arrives, the WAS will intercept it, launch a worker process to host your service, and forward it the call.

WAS offers many features beyond self-hosting, including application pooling, recycling, idle time management, identity management, and isolation. To use the WAS, you must target a platform that supports it, such as a Windows Server 2008 (or later) machine for scalability, or a Windows 7 (or later) client machine for a handful of clients.

That said, self-hosted processes do offer singular advantages, such as in-proc hosting, dealing well with unknown customer environments, and easy programmatic access to the advanced hosting features described previously.

Custom Hosting in IIS/WAS

It is often the case that you need to interact with the host instance. While this is integral to the use of a self-hosting solution, when using IIS or WAS, you have no direct access to the host. To overcome this hurdle, WCF provides a hook called a *host factory*. Using the `Factory` tag in the .svc file, you can specify a class you provide that creates the host instance:

```
<%@ ServiceHost
    Language = "C#"
    Debug    = "true"
    CodeBehind = "~/App_Code/MyService.cs"
```

```
Service      = "MyService"  
Factory     = "MyServiceFactory"  
%>
```

You can also specify the host factory in the config file when not using an .svc file explicitly:

```
<serviceActivations>  
    <add relativeAddress = "MyService.svc"  
        service = "MyService"  
        factory = "MyServiceFactory"  
    />  
</serviceActivations>
```

The host factory class must derive from the `ServiceHostFactory` class and override the `CreateServiceHost()` virtual method:

```
public class ServiceHostFactory : ...  
{  
    protected virtual ServiceHost CreateServiceHost(Type serviceType,  
                                                    Uri[] baseAddresses);  
    //More members  
}
```

For example:

```
class MyServiceFactory : ServiceHostFactory  
{  
    protected override ServiceHost CreateServiceHost(Type serviceType,  
                                                    Uri[] baseAddresses)  
    {  
        ServiceHost host = new ServiceHost(serviceType,baseAddresses);  
  
        //Custom steps here  
  
        return host;  
    }  
}
```



The `CreateServiceHost()` method is logically the `Main()` method of an IIS or WAS hosted service, and you can use it just for that purpose.

Choosing a Host

Although WCF offers such a variety of options, from IIS, to the WAS, to self-hosting, it is easy to choose the correct host, as shown in [Figure 1-4](#). For an Internet application (i.e., an application that receives calls from clients across the Internet), IIS or the WAS provide the best capabilities to harden your services against the security con-

cerns of access over the Internet. Otherwise prefer self-hosting your services. Self-hosting provides your organization with a hosting model that significantly simplifies the administration and provisioning of intranet services. WCF services do not require the many additional lifecycle management facilities that the WAS provides.

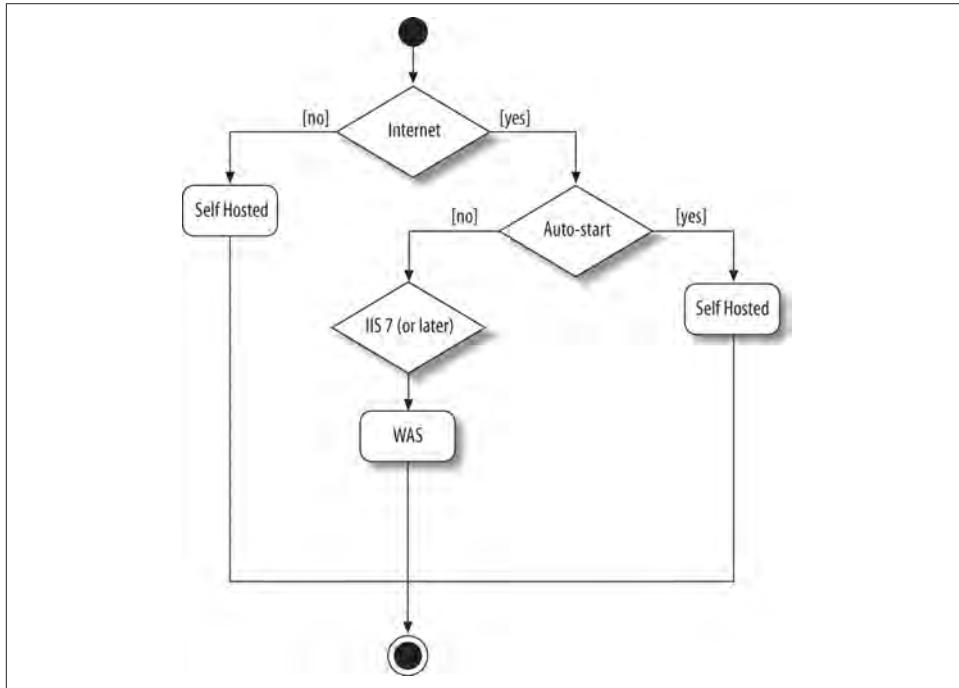


Figure 1-4. Choosing a host for any service



Microsoft will retire Windows Server AppFabric in April 2016. You should not target any of its features, particularly the service auto-start feature, for any of your systems.

Bindings

There are multiple aspects of communication with any given service, and there are many possible communication patterns. Messages can follow a synchronous request-reply or asynchronous fire-and-forget pattern, messages can be bidirectional, messages can be delivered immediately or queued, and the queues can be durable or volatile. As discussed previously, there are many possible transport protocols for the messages, such as HTTP (or HTTPS), TCP, IPC, MSMQ. There are also a few possible message encoding options. You can choose plain text to enable interoperability, binary encoding to optimize performance, or the Message Transport Optimization

Mechanism (MTOM) for large payloads. Finally, there are multiple options for securing messages. You can choose not to secure them at all, to provide transport-level security only, or to provide message-level privacy and security, and of course, there are numerous ways to authenticate and authorize the clients. Message delivery might be unreliable or reliable end-to-end across intermediaries and dropped connections, and the messages might be processed in the order they were sent or in the order they were received. Your service might need to interoperate with other services or clients that are aware of only the basic web service protocol, or with clients and services capable of using the score of WS-* protocols, such as WS-Security and WS-Atomic Transactions. Your service may need to have the capability to interoperate with any client or you may want to restrict your service to interoperate only with another WCF service or client.

If you were to count all the possible communication and interaction options, you'd probably find that the number of permutations is in the tens of thousands. Some of those choices may be mutually exclusive and some may mandate other choices. Clearly, both the client and the service must be aligned on all these options in order to communicate properly. Managing this level of complexity adds no business value to most applications, and yet the productivity and quality implications of making the wrong decisions are severe.

To simplify these choices and make them manageable, WCF groups together sets of communication aspects in *bindings*. A binding is merely a consistent, canned set of choices regarding the transport protocol, message encoding, communication pattern, reliability, security, transaction propagation, and interoperability. All you need to do is determine the target scenario for your service, and WCF makes a correct multidimensional decision for you regarding all the aspects of the communication. Ideally, you can extract all these "plumbing" aspects from your service code and allow the service to focus solely on the implementation of the business logic. Bindings allow you to use the same service logic over drastically different plumbing.

You can use the WCF-provided bindings out of the box, you can tweak their properties, or you can write your own custom bindings from scratch. The service publishes its choice of binding in its metadata, enabling clients to query for the type and specific properties of the binding. This is important because the client must use the exact same binding values as the service. A single service can support multiple bindings on separate addresses.

The Common Bindings

WCF defines five frequently used bindings:

Basic binding

Offered by the `BasicHttpBinding` class, basic binding is designed to expose a WCF service as a legacy ASMX web service so that old clients can work with new

services. The basic binding makes your service look, on the wire, like a legacy web service that communicates over the basic web service profile. When used by clients, this binding enables new WCF clients to work with old ASMX services.

TCP binding

Offered by the `NetTcpBinding` class, TCP binding uses TCP for cross-machine communication on the intranet. It supports a variety of features, including reliability, transactions, and security, and is optimized for WCF-to-WCF communication. As a result, it requires both the client and the service to use WCF.

IPC binding

Offered by the `NetNamedPipeBinding` class, IPC binding uses named pipes as a transport for same-machine communication. It is the most secure binding, as it cannot accept calls from outside the machine. The IPC binding supports a variety of features similar to the TCP binding. It is also the most performant binding, as IPC is a lighter protocol than TCP.



The `NetNamedPipeBinding` class is inconsistently named, because the binding naming convention is to refer to the protocol, not the communication mechanism (thus, we have `NetTcpBinding` rather than `NetSocketBinding`). The correct name for this binding should have been `NetIpcBinding`. Throughout this book, I will refer to the `NetNamedPipeBinding` as the IPC binding.

Web Service (WS) binding

Offered by the `WSHttpBinding` class, the WS binding uses HTTP or HTTPS for transport and offers a variety of features (such as reliability, transactions, and security) over the Internet, all using the WS-* standards. This binding is designed to interoperate with any party that supports the WS-* standards.

MSMQ binding

Offered by the `NetMsmqBinding` class, the MSMQ binding uses MSMQ for transport and offers support for disconnected queued calls. Use of this binding is the subject of [Chapter 9](#).

Format and encoding

Each of the frequently used bindings uses a different transport scheme and encoding, as listed in [Table 1-1](#). Where multiple encodings are possible, the defaults are shown in bold.

Table 1-1. Transport and encoding for common bindings

Name	Transport	Encoding	Interoperable
BasicHttpBinding	HTTP/HTTPS	Text, MTOM	Yes
NetTcpBinding	TCP	Binary	No
NetNamedPipeBinding	IPC	Binary	No
WSHttpBinding	HTTP/HTTPS	Text, MTOM	Yes
NetMsmqBinding	MSMQ	Binary	No

Having text-based encoding typically enables a WCF service (or client) to communicate over HTTP with any other service (or client), regardless of its technology and across firewalls. Binary encoding over TCP, IPC, or MSMQ yields the best performance, but it does so at the expense of interoperability because it mandates WCF-to-WCF communication. That said, with the TCP, IPC, and MSMQ bindings, interoperability is often not required. In the case of IPC, because the call can never leave the client machine, the client can rest assured that the target machine is running Windows and has WCF installed on it. In the case of the TCP binding, while your application may need to interoperate with other applications written in other technologies, applications themselves do tend to be homogeneous internally. As such, as long as your application spans only the local intranet, you can typically assume a homogeneous Windows environment without internal firewalls between machines. Finally, the MSMQ binding requires the use of MSMQ server, which of course is Windows-specific.



The binary encoder for the TCP, IPC, and MSMQ bindings is proprietary to WCF. Do not attempt to write a custom parser for it on other platforms. Microsoft reserves the right to change its format over time in order to keep optimizing and evolving it.

Choosing a Binding

When choosing a binding for your service, you should follow the decision diagram shown in [Figure 1-5](#).

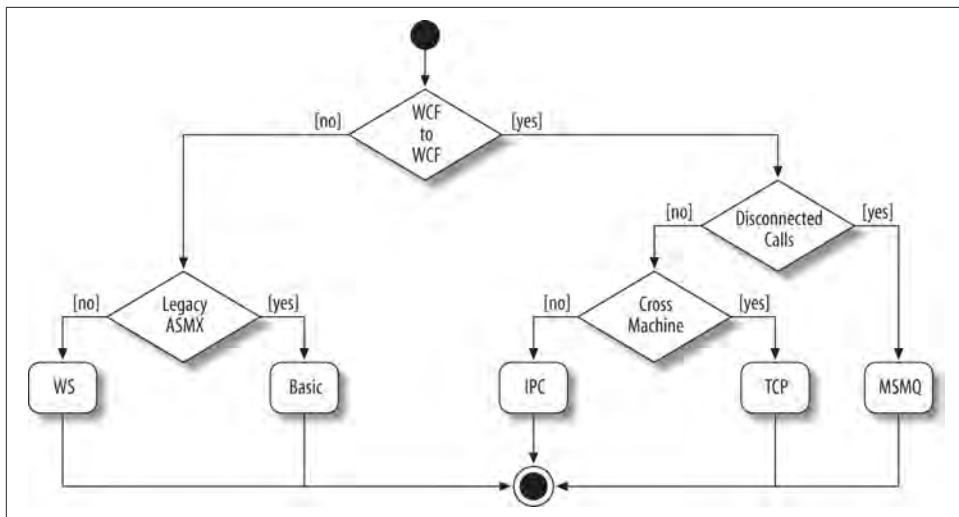


Figure 1-5. Choosing a binding

The first question you should ask yourself is whether your service needs to interact with non-WCF clients. If it does, and those clients expect the basic web service protocol (ASMX web services), choose the `BasicHttpBinding`, which exposes your WCF service to the outside world as if it were an ASMX web service (i.e., a WSI-basic profile). The downside of this choice is that you cannot take advantage of most of the standard WS-* protocols. If, however, the non-WCF client can understand these standards, you can instead choose the `WS` binding. If you can assume the client is a WCF client and requires offline or disconnected interaction, choose the `NetMsmqBinding`, which uses MSMQ for transporting the messages. If the client requires connected communication but could be calling across machine boundaries, choose the `NetTcpBinding`, which communicates over TCP. If the client is on the same machine as the service, choose the `NetNamedPipeBinding`, which uses IPC to maximize performance.



Most bindings work well even outside their target scenarios. For example, you can use the `TCP` binding for same-machine or even in-proc communication, and you can use the `basic` binding for intranet WCF-to-WCF communication. However, do try to choose a binding according to [Figure 1-5](#).

Additional Bindings

In addition to the five frequently used bindings described so far, WCF provides three specializations of these bindings: the `BasicHttpContextBinding`, the `WSHttpContextBinding`, and the `NetTcpHttpContextBinding`. The context bindings (described in [Appendix A](#)) provide a way to access the context information of the message being processed.

dix B) all derive from their respective regular bindings, adding support for a context protocol. The context protocol allows you to pass out-of-band parameters to the service. You can also use the context bindings for durable services support, as described in [Chapter 4](#).

WCF defines eight bindings that we do not think you should use frequently. These bindings (listed next) are each designed for a specific target scenario and you cannot use them easily outside that scenario. This book makes little or no use of these bindings, due to their somewhat esoteric nature and primarily the availability of better design alternatives.

WS dual binding

Offered by the `WSDualHttpBinding` class, this is similar to the WS binding, except it also supports bidirectional duplex communication from the service to the client, as discussed in [Chapter 5](#). While this binding does use industry standards (it is nothing more than two `WSHttpBinding` bindings wired up against each other to support callbacks), there is no industry standard for setting up the callback, and therefore the `WSDualHttpBinding` is not interoperable. This binding is a legacy from the first release of WCF. The availability of the WebSocket binding `NetHttpBinding` deprecates the `WSDualHttpBinding` for callbacks over the Internet.

WebSocket binding

Offered by the `NetHttpBinding` class, this binding is a de facto .NET 4.5 replacement for the .NET 3.0 `WSDualHttpBinding`, as it provides a better model for callbacks over the Internet. As with the `WSDualHttpBinding`, the `NetHttpBinding` offers a variety of additional features such as reliability and security. Unless you absolutely need to use callbacks over the Internet, you should instead select one of the common bindings as prescribed in [“Choosing a Binding” on page 26](#).



The WebSocket protocol and associated WebSocket API are both industry standards but the `NetHttpBinding` represents a proprietary WCF-to-WCF implementation. Therefore, the `NetHttpBinding` is incompatible with the interoperable WebSocket API. We find this lack of standards compliance severely limits the WebSocket binding's reach and applicability. Consequently, we do not consider the WebSocket binding one of the recommended bindings and have reduced the role of the `NetHttpBinding` to advanced scenarios where callbacks over the Internet are appropriate.

Federated WS binding

Offered by the `WSFederationHttpBinding` class, this is a specialization of the WS binding that offers support for federated security. Federated security is beyond the scope of this book, as industry presently lacks good support (both in technol-

ogy and in business models) for true federated scenarios. I do expect federation to become mainstream as time goes by.

Federated WS 2007 binding

Offered by the `WS2007FederationHttpBinding` class, this is an update of `WSFederationHttpBinding`.

MSMQ integration binding

Offered by the `MsmqIntegrationBinding` class, this is the analogous queued-world binding to the basic binding. The integration binding converts WCF messages to and from MSMQ messages and is designed to interoperate with legacy MSMQ clients, which was a fringe scenario at the time WCF first came out (it is even less applicable today).

WS 2007 binding

Offered by the `WS2007HttpBinding` class, this binding derives from the `WSHttpBinding` class; it adds support for the coordination standard and updates for the transaction, security, and reliability standards.

UDP binding

Offered by the `UdpBinding` class, this binding provides support for the User Datagram Protocol (UDP). UDP provides no guarantee of message delivery, ordering, or duplicate message detection. Because of these shortcomings, the `UdpBinding`'s usefulness is limited to the few scenarios that can tolerate message loss, lack of ordering, or duplication. In most business scenarios, using UDP is therefore out of the question. While you could entertain the use of UDP for events broadcasting, a far better management and reliable design solution is to use a pub/sub service, as is shown in [Appendix D](#).

Web binding

Offered by the `WebHttpBinding` class, this binding allows your service to accept simple calls over web protocols such as HTTP-GET using the REST/POX/JSON patterns. The `WebHttpBinding` is now superseded by the new ASP.NET Web API framework initially released in .NET 4.0 and supported in IIS 7 and above. The ASP.NET Web API provides a much cleaner model for implementing HTTP-based services and REST-like interactions.



As of .NET 4.5, `NetPeerTcpBinding` has been marked `Obsolete` and should no longer be used.

Using a Binding

Each binding offers literally dozens of configurable properties. There are three ways of working with bindings: you can use the built-in bindings as they are, if they fit your requirements; you can tweak and configure some of their properties, such as transaction propagation, reliability, and security; or you can write your own custom bindings. The most common scenario is using an existing binding mostly as it is and merely configuring two or three of its aspects. Application developers will hardly ever need to write a custom binding, but framework developers may need to.

Endpoints

Every service is associated with an address that defines where the service is, a binding that defines how to communicate with the service, and a contract that defines what the service does. This triumvirate governing the service is easy to remember as the *ABC* of the service. WCF formalizes this relationship in the form of an *endpoint*. The endpoint is the fusion of the address, contract, and binding (see [Figure 1-6](#)).

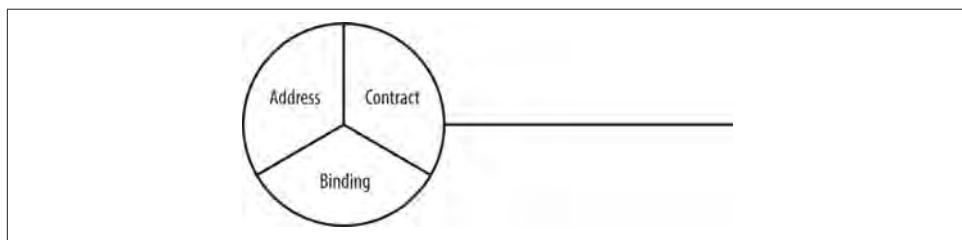


Figure 1-6. The endpoint

Every endpoint must have all three elements, and the host exposes the endpoint. Logically, the endpoint is the service's interface and is analogous to a CLR or COM interface. Note the use of the traditional “lollipop” notation to denote an endpoint in [Figure 1-6](#).



Conceptually, even in C# or VB, there are endpoints: the address is the memory address of the type's virtual table, the binding is CLR, and the contract is the interface itself. Because in classic .NET programming you never deal with addresses or bindings, you take them for granted, and you've probably grown accustomed to equating in your mind's eye the interface (which is merely a programming construct) with all that it takes to interface with an object. The WCF endpoint is a true interface because it contains all the information required to interface with the object. In WCF, the address and the binding are not preordained and you must specify them.

Every service must expose at least one business endpoint, and each endpoint has exactly one contract. All endpoints on a service have unique addresses, and a single service can expose multiple endpoints. These endpoints can use the same or different bindings and can expose the same or different contracts. There is absolutely no relationship between the various endpoints a service provides.

It is important to point out that nothing in the service code pertains to its endpoints, and they are always external to the service code. You can configure endpoints either administratively (using a config file) or programmatically.

Administrative Endpoint Configuration

Configuring an endpoint administratively requires placing the endpoint details in the hosting process config file. For example, given this service definition:

```
namespace MyNamespace
{
    [ServiceContract]
    interface IMyContract
    {...}
    class MyService : IMyContract
    {...}
}
```

Example 1-6 shows the required entries in the config file. Under each service type, you list its endpoints.

Example 1-6. Administrative endpoint configuration

```
<system.serviceModel>
<services>
    <service name = "MyNamespace.MyService">
        <endpoint
            address  = "http://localhost:8000/MyService"
            binding   = "wsHttpBinding"
            contract = "MyNamespace.IMyContract"
        />
    </service>
</services>
</system.serviceModel>
```

When you specify the service and the contract type, you need to use fully qualified type names. I will omit the namespace in the examples throughout the remainder of this book, but you should use a namespace when applicable. Note that if the endpoint provides a base address, that address scheme must be consistent with the binding, such as HTTP with `WSHttpBinding`. A mismatch causes an exception at service load time.

Example 1-7 shows a config file defining a single service that exposes multiple endpoints. You can configure multiple endpoints with the same base address as long as the URI is different.

Example 1-7. Multiple endpoints on the same service

```
<service name = "MyService">
  <endpoint
    address  = "http://localhost:8000/MyService"
    binding  = "wsHttpBinding"
    contract = "IMyContract"
  />
  <endpoint
    address  = "net.tcp://localhost:8001/MyService"
    binding  = "netTcpBinding"
    contract = "IMyContract"
  />
  <endpoint
    address  = "net.tcp://localhost:8002/MyService"
    binding  = "netTcpBinding"
    contract = "IMyOtherContract"
  />
</service>
```

Using base addresses

In **Example 1-7**, each endpoint provided its own base address. When you provide an explicit base address, it overrides any base address the host may have provided.

You can also have multiple endpoints use the same base address, as long as the endpoint addresses differ in their URLs:

```
<service name = "MyService">
  <endpoint
    address  = "net.tcp://localhost:8001/MyService"
    binding  = "netTcpBinding"
    contract = "IMyContract"
  />
  <endpoint
    address  = "net.tcp://localhost:8001/MyOtherService"
    binding  = "netTcpBinding"
    contract = "IMyContract"
  />
</service>
```

Alternatively, if the host provides a base address with a matching transport scheme, you can leave out the address. In this case, the endpoint address will be the same as the base address of the matching transport:

```
<endpoint
  binding  = "wsHttpBinding"
```

```
    contract = "IMyContract"
/>>
```

If the host does not provide a matching base address, loading the service host will fail with an exception.

When you configure the endpoint address, you can add just the relative URI under the base address:

```
<endpoint
    address  = "SubAddress"
    binding   = "wsHttpBinding"
    contract  = "IMyContract"
/>
```

The endpoint address in this case will be the matching base address plus the URI, and again, the host must provide a matching base address.

Binding configuration

You can use the config file to customize the binding used by the endpoint. To that end, add the `bindingConfiguration` tag to the `endpoint` section and name a customized section in the `bindings` section of the config file. [Example 1-8](#) demonstrates using this technique to enable transaction propagation. [Chapter 7](#) explains the function of the `transactionFlow` tag.

Example 1-8. Service-side binding configuration

```
<system.serviceModel>
    <services>
        <service name = "MyService">
            <endpoint
                address  = "net.tcp://localhost:8000/MyService"
                bindingConfiguration = "TransactionalTCP"
                binding   = "netTcpBinding"
                contract  = "IMyContract"
            />
            <endpoint
                address  = "net.tcp://localhost:8001/MyService"
                bindingConfiguration = "TransactionalTCP"
                binding   = "netTcpBinding"
                contract  = "IMyOtherContract"
            />
        </service>
    </services>
    <bindings>
        <netTcpBinding>
            <binding name = "TransactionalTCP"
                transactionFlow = "true"
            />
        </netTcpBinding>
    </bindings>
</system.serviceModel>
```

```
</bindings>
</system.serviceModel>
```

As shown in [Example 1-8](#), you can reuse the named binding configuration in multiple endpoints simply by referring to it.

Default binding

WCF allows you to use a *default binding* that affects all endpoints of all services of the application that uses the config file. A default binding is simply a nameless binding section. For example, in the case of TCP:

```
<netTcpBinding>
  <binding
    transactionFlow = "true"
  />
</netTcpBinding>
```

The default binding implicitly configures all endpoints that do not explicitly reference a binding configuration.

For example, using a default binding, [Example 1-8](#) is reduced to:

```
<system.serviceModel>
  <services>
    <service name = "MyService">
      <endpoint
        address = "net.tcp://localhost:8000/MyService"
        binding = "netTcpBinding"
        contract = "IMyContract"
      />
      <endpoint
        address = "net.tcp://localhost:8001/MyService"
        binding = "netTcpBinding"
        contract = "IMyOtherContract"
      />
    </service>
  </services>
  <bindings>
    <netTcpBinding>
      <binding
        transactionFlow = "true"
      />
    </netTcpBinding>
  </bindings>
</system.serviceModel>
```

You can only have at most one default binding configuration per binding type.

The problem with the default binding is that when you combine default bindings with named binding configurations, as shown in [Figure 1-7](#), the config file may become difficult for humans to parse and understand.

```

<endpoint
    binding = "netTcpBinding"
/>
<endpoint
    binding = "netTcpBinding"
/>
<endpoint
    bindingConfiguration = "MyBinding"
    binding = "netTcpBinding"
/>

<bindings>
    <netTcpBinding>
        <binding
            ...
        />
        <binding name = "MyBinding"
            ...
        />
    </netTcpBinding>
</bindings>

```

The diagram shows a portion of a WCF configuration file. It includes two endpoint definitions, each pointing to a 'netTcpBinding'. Below them is a 'bindings' section containing a 'netTcpBinding' definition. This definition contains two 'binding' elements. The first 'binding' is unnamed, indicated by three dots. The second 'binding' is named 'MyBinding', also indicated by three dots. Arrows point from the 'binding' and 'name' attributes of the second binding back to its respective labels in the caption.

Figure 1-7. Named and default binding configuration

While [Figure 1-8](#) is a perfectly valid configuration, I recommend against mixing named and default bindings. Either have all your binding configurations named or use only the default configuration. Another advantage of a named configuration is that it allows you to weave a bit of documentation via the binding configuration name as to what that configuration is trying to achieve. Most if not all of the binding configurations in this book are named for precisely that reason.

Programmatic Endpoint Configuration

Programmatic endpoint configuration is equivalent to administrative configuration, but instead of resorting to a config file, you rely on programmatic calls to add endpoints to the `ServiceHost` instance. Again, these calls are always outside the scope of the service code. `ServiceHost` provides overloaded versions of the `AddServiceEndpoint()` method:

```

public class ServiceHost : ServiceHostBase
{
    public ServiceEndpoint AddServiceEndpoint(Type implementedContract,
                                                Binding binding,
                                                string address);
    //Additional members
}

```

You can provide `AddServiceEndpoint()` methods with either relative or absolute addresses, just as with a config file. [Example 1-9](#) demonstrates programmatic configuration of the same endpoints as in [Example 1-7](#).

Example 1-9. Service-side programmatic endpoint configuration

```
ServiceHost host = new ServiceHost(typeof(MyService));

Binding wsBinding = new WSHttpBinding();
Binding tcpBinding = new NetTcpBinding();

host.AddServiceEndpoint(typeof(IMyContract),wsBinding,
    "http://localhost:8000/MyService");
host.AddServiceEndpoint(typeof(IMyContract),tcpBinding,
    "net.tcp://localhost:8001/MyService");
host.AddServiceEndpoint(typeof(IMyOtherContract),tcpBinding,
    "net.tcp://localhost:8002/MyService");

host.Open();
```

When you add an endpoint programmatically, the address is given as a string, the contract as a Type, and the binding as one of the subclasses of the abstract class `Binding`, as in:

```
public class NetTcpBinding : Binding,...
```

To rely on the host base address, provide an empty string if you want to use only the base address, or just the URI to use the base address plus that URI:

```
Uri tcpBaseAddress = new Uri("net.tcp://localhost:8000/");

ServiceHost host = new ServiceHost(typeof(MyService),tcpBaseAddress);

Binding tcpBinding = new NetTcpBinding();

//Use base address as address
host.AddServiceEndpoint(typeof(IMyContract),tcpBinding,"");
//Add relative address
host.AddServiceEndpoint(typeof(IMyContract),tcpBinding,"MyService");
//Ignore base address
host.AddServiceEndpoint(typeof(IMyContract),tcpBinding,
    "net.tcp://localhost:8001/MyService");
host.Open();
```

As with administrative configuration using a config file, the host must provide a matching base address; otherwise, an exception occurs. In fact, in terms of capabilities, there is no difference between programmatic and administrative configuration. When you use a config file, all WCF does is parse the file and execute the appropriate programmatic calls in its place.

Binding configuration

You can programmatically set the properties of the binding used. For example, the following is the code required to enable transaction propagation (similar to [Example 1-8](#)):

```
ServiceHost host = new ServiceHost(typeof(MyService));
NetTcpBinding tcpBinding = new NetTcpBinding();

tcpBinding.TransactionFlow = true;

host.AddServiceEndpoint(typeof(IMyContract),tcpBinding,
                      "net.tcp://localhost:8000/MyService");
host.Open();
```

Note that when you're dealing with specific binding properties, you typically interact with a concrete binding subclass, such as `NetTcpBinding`, rather than its abstract base class, `Binding` (as was done in [Example 1-9](#)).

All the binding classes also offer a constructor that takes a string, for example:

```
public class NetTcpBinding : Binding, ...
{
    public NetTcpBinding(string configurationName);
    //More members
}
```

You can use that constructor to programmatically initialize a binding object based on settings found in a particular binding section in the config file. You can also pass an empty string to instruct WCF to use the default (nameless) binding configuration. If the config file does not contain a default binding definition, you will encounter a `KeyNotFoundException`.

Default Endpoints

If the service host does not define any endpoints (neither in config nor programmatically) but does provide at least one base address, WCF will by default add endpoints to the service. These are called the *default endpoints*. WCF will add an endpoint per base address per contract, using the base address as the endpoint's address. WCF will infer the binding from the scheme of the base address. For HTTP, WCF will use the basic binding. Note that the default bindings will affect the default endpoints. WCF will also name the endpoint by concatenating the binding name and the contract name.

For example, given this service definition:

```
[ServiceContract]
interface IMyContract
{...}
```

```

[ServiceContract]
interface IMyOtherContract
{...}

class MyService : IMyContract,IMyOtherContract
{...}

```

for this hosting code:

```

Uri httpBaseAddress = new Uri("http://localhost:8000/");
Uri tcpBaseAddress = new Uri("net.tcp://localhost:9000/");
Uri ipcBaseAddress = new Uri("net.pipe://localhost/");

ServiceHost host = new ServiceHost(typeof(MyService),httpBaseAddress,
                                    tcpBaseAddress,ipcBaseAddress);
host.Open();

```

Assuming no config file is used to define any additional endpoints, WCF will add these endpoints, as if they were defined in config:

```

<service name = "MyService">
    <endpoint name = "BasicHttpBinding_IMyContract"
              address = "http://localhost:8000/"
              binding = "basicHttpBinding"
              contract = "IMyContract"
    />
    <endpoint name = "NetTcpBinding_IMyContract"
              address = "net.tcp://localhost:9000"
              binding = "netTcpBinding"
              contract = "IMyContract"
    />
    <endpoint name = "NetNamedPipeBinding_IMyContract"
              address = "net.pipe://localhost/"
              binding = "netNamedPipeBinding"
              contract = "IMyContract"
    />
    <endpoint name = "BasicHttpBinding_IMyOtherContract"
              address = "http://localhost:8000/"
              binding = "basicHttpBinding"
              contract = "IMyOtherContract"
    />
    <endpoint name = "NetTcpBinding_IMyOtherContract"
              address = "net.tcp://localhost:9000"
              binding = "netTcpBinding"
              contract = "IMyOtherContract"
    />
    <endpoint name = "NetNamedPipeBinding_IMyOtherContract"
              address = "net.pipe://localhost/"
              binding = "netNamedPipeBinding"
              contract = "IMyOtherContract"
    />
</service>

```

Note that WCF will provide the same address multiple times to different endpoints. While this works as far as invocation (because the host monitors the incoming ports or pipes only once and simply dispatches the message internally to the correct endpoint), this configuration will fail metadata publishing due to an internal limitation of WCF.

You can also add the default endpoints explicitly using the `AddDefaultEndpoints()` method of `ServiceHost`:

```
public class ServiceHost : ...
{
    public void AddDefaultEndpoints();
    //More members
}
```

You can add the default endpoints even if you have added other endpoints conventionally using a config file or programmatically. The only thing to watch for is conflict with other endpoints that use the base address as their address.

Protocol mapping

For the default endpoints, WCF will infer the binding to use from the scheme of the base address. This inferring is called *protocol mapping*. In the case of TCP, IPC, and MSMQ, there is only a single mapping option. However, in the case of HTTP (or HTTPS), WCF will default to the basic binding for mapping. If you like to rely on the WS binding instead (as you should in most cases), you need to override the default protocol mapping using the `protocolMapping` section in the config file:

```
<system.serviceModel>
  <protocolMapping>
    <add
      scheme = "http"
      binding = "wsHttpBinding"
    />
  </protocolMapping>
</system.serviceModel>
```

You can also specify a particular binding configuration to use:

```
<protocolMapping>
  <add
    scheme = "http"
    binding = "wsHttpBinding"
    bindingConfiguration = "..."
  />
</protocolMapping>
```

You must do protocol mapping administratively in the config file. There is no equivalent programmatic way.



Protocol mapping is the only configuration option offered by WCF in the service model section that does not have a programmatic equivalent.

The Configure() Method

Prior to .NET 4.5, when using programmatic configuration, your hosting code was often coupled to the hosting process (or Windows Service) with which it was deployed. When using self-hosting, you were required to provide programmatic configuration by interacting directly with a service host instance. When hosting in WAS, you were required to provide a service host factory. And if you had evolved your configuration strategy to leverage a centralized configuration store, you were required to merge your host's stored configuration with its file-based configuration.

As part of the WCF simplification features introduced in .NET 4.5, the `Configure()` method provides you with a hosting process-independent approach to configuring your services programmatically. Regardless of the hosting environment, a `Configure()` method allows you to define the configuration for your service's endpoints directly within its code.

You enable this feature for your service by adding a public static method named `Configure()` with the following signature to your service's implementation:

```
class MyService : IMyContract
{
    public static void Configure(ServiceConfiguration config)
    {...}

    // More members
}
```



The `Configure()` method convention is very specific. If you do not mark your `Configure()` method both as `public` and as `static`, WCF will not detect it.

The `ServiceConfiguration` argument passed to the `Configure()` method provides many of the same methods for programmatic service configuration as offered by the `ServiceHost` class:

```
public class ServiceConfiguration
{
    public void AddServiceEndpoint(ServiceEndpoint endpoint);
    public ServiceEndpoint AddServiceEndpoint(Type contractType,
                                              Binding binding,
```

```

        Uri address);

    public Collection<ServiceEndpoint> EnableProtocol(Binding protocol);

    public void LoadFromConfiguration();
    public void LoadFromConfiguration(Configuration configuration);

    //More members
}

```

This allows you to use the `ServiceConfiguration` class to configure your service with the same techniques described thus far. [Example 1-10](#) shows the same programmatic endpoint configuration as [Example 1-9](#), but instead uses the `ServiceConfiguration` parameter of the `Configure()` method.

Example 1-10. Configure() method programmatic endpoint configuration

```

public static void Configure(ServiceConfiguration config)
{
    Binding wsBinding = new WSHttpBinding();
    Binding tcpBinding = new NetTcpBinding();

    config.AddServiceEndpoint(typeof(IMyContract),wsBinding,
                            "http://localhost:8000/MyService");
    config.AddServiceEndpoint(typeof(IMyContract),tcpBinding,
                            "net.tcp://localhost:8001/MyService");
    config.AddServiceEndpoint(typeof(IMyOtherContract),tcpBinding,
                            "net.tcp://localhost:8002/MyService");
}

```

If present, WCF will call your service's `Configure()` method during service host instantiation before the service host is opened. This gives you the opportunity to configure your service's endpoints programmatically or explicitly read your service's configuration section by calling one of the `LoadFromConfiguration()` methods offered by the `ServiceConfiguration` class.



When you use a `Configure()` method, WCF will ignore the service configuration section for your service in the app or web config file.

You can also add default endpoints to available base addresses by calling the `EnableProtocol()` method with a specific binding instance. `EnableProtocol()` will then add default endpoints using the same conventions as described in [“Default Endpoints” on page 37](#). For example, the following `EnableProtocol()` call will produce

the same endpoint configurations for the `BasicHttpBinding` as the default endpoints added by WCF for each HTTP compatible base address:

```
config.EnableProtocol(new BasicHttpBinding());
```

Of course, because you provide the binding, you can use `EnableProtocol()` to better control the default endpoint configurations WCF produces. WCF will still add default endpoints for any base address for which you have not explicitly provided an endpoint or enabled a protocol. Be aware that if you change the name of the binding you specify when enabling a protocol, you create an endpoint that no longer complies with WCF's default endpoint naming conventions.



When you enable a protocol, WCF will ignore the default binding, default endpoint behavior and any protocol mappings related to the binding type you've specified.

While a `Configure()` method provides you with a host-agnostic approach to configuring your services programmatically, it also couples your service's configuration directly to its implementation and therefore its deployment. In simple scenarios, you may find this approach desirable, even labor saving. But as your service-oriented system grows, you will find it necessary to establish a cleaner separation of concerns.

Using a `Configure()` method carries a price when you have multiple services that all share very similar, if not identical behavior. Obviously, repeating the same programmatic configuration across all of them is hardly advisable. At the same time, resorting to file-based configuration is often undesirable as well and still carries the same duplication and additional long-term maintenance costs across hosting environments.

Although the convention for implementing a `Configure()` method is public and static, you can still mitigate the duplication of configuration code by relegating the `Configure()` method to a base class. This allows you to provide a single expression of configuration for a set of related services:

```
class MyServiceBase
{
    public static void Configure(ServiceConfiguration config)
    {...}
}

class MyService : MyServiceBase, IMyContract
{...}

class MyOtherService : MyServiceBase, IMyOtherContract
{...}
```

Because your `Configure()` method must be static, you cannot mark it as virtual in your base class. Implementing a `Configure()` method in a child class will hide the implementation in your base class. In this case, you should mark as new any customizations added to your child classes and call your base class' implementation:

```
class MyService : MyServiceBase, IMyContract
{
    public new static void Configure(ServiceConfiguration config)
    {
        MyServiceBase.Configure(config);
        //Optional additional processing
    }
}

class MyOtherService : MyServiceBase, IMyOtherContract
{
    public new static void Configure(ServiceConfiguration config)
    {
        MyServiceBase.Configure(config);
        //Optional additional processing
    }
}
```

While alleviating code redundancy, inheritance still does not provide the cleanest separation between your service's implementation and its configuration.

You should always endeavor to keep the many facets of your services encapsulated from one another to promote reuse, flexibility, testability, maintainability, and most of all, agility. Anything from simple helper classes to a central configuration repository or even Inversion of Control (IoC) approaches can help you to encapsulate your configuration infrastructure code from your service's implementation.

In conjunction with these options, as your needs evolve, you can further extend your control of programmatic configuration, as well as many other service capabilities, by creating a custom `ServiceHost`.

Metadata Exchange

By default, the service will not publish its metadata. However, this does not preclude clients that have obtained the metadata via some other mechanism (such as a project reference to a class library containing the contracts) from invoking operations on the service.

Publishing your service's metadata involves significant effort, as you have to convert CLR types and binding information into WSDL or some other low-level representation, and all that effort does not add any business value. Fortunately, the host already knows everything there is to know about your service and its endpoints, so it can publish the metadata for you if explicitly instructed to do so.

WCF offers two options for publishing a service's metadata: you can provide the metadata over HTTP-GET, a simple text-based protocol that most platforms support, or you can use a dedicated endpoint.

Metadata over HTTP-GET

WCF can provide the metadata for your service over HTTP-GET automatically; all you need to do is enable it by adding an explicit service behavior. Behaviors are described fully in subsequent chapters. For now, all you need to know is that a behavior is a local aspect of the service, such as whether or not it wants to have the host publish its metadata over HTTP-GET. You can add this behavior administratively or programmatically.

Enabling metadata exchange administratively

Example 1-11 shows a host application config file where both hosted services reference a custom behavior section that enables metadata publishing over HTTP-GET.

Example 1-11. Enabling metadata exchange behavior using a config file

```
<system.serviceModel>
  <services>
    <service name = "MyService" behaviorConfiguration = "MEXGET">
      <host>
        <baseAddresses>
          <add baseAddress = "http://localhost:8000/" />
        </baseAddresses>
      </host>
      ...
    </service>
    <service name = "MyOtherService" behaviorConfiguration = "MEXGET">
      <host>
        <baseAddresses>
          <add baseAddress = "http://localhost:8001/" />
        </baseAddresses>
      </host>
      ...
    </service>
  </services>
  <behaviors>
    <serviceBehaviors>
      <behavior name = "MEXGET">
        <serviceMetadata httpGetEnabled = "true"/>
      </behavior>
    </serviceBehaviors>
  </behaviors>
</system.serviceModel>
```

By default, the address the clients need to use for HTTP-GET is the registered HTTP base address of the service. If the host is not configured with an HTTP base address, loading the service will throw an exception. You can also specify a different address (or just a URI appended to the HTTP base address) at which to publish the metadata by setting the `httpGetUrl` property of the `serviceMetadata` tag:

```
<behavior name = "MEXGET">
    <serviceMetadata httpGetEnabled = "true" httpGetUrl = "MyMEXAddress"/>
</behavior>
```

Once you have enabled the metadata exchange over HTTP-GET, you can navigate to the address you configured (the HTTP base address, by default, or an explicit address) using a browser. If all is well, you will get a confirmation page like the one shown in [Figure 1-8](#), letting you know that you have successfully hosted a service. The confirmation page is unrelated to IIS hosting, and you can use a browser to navigate to the service address even when self-hosting.

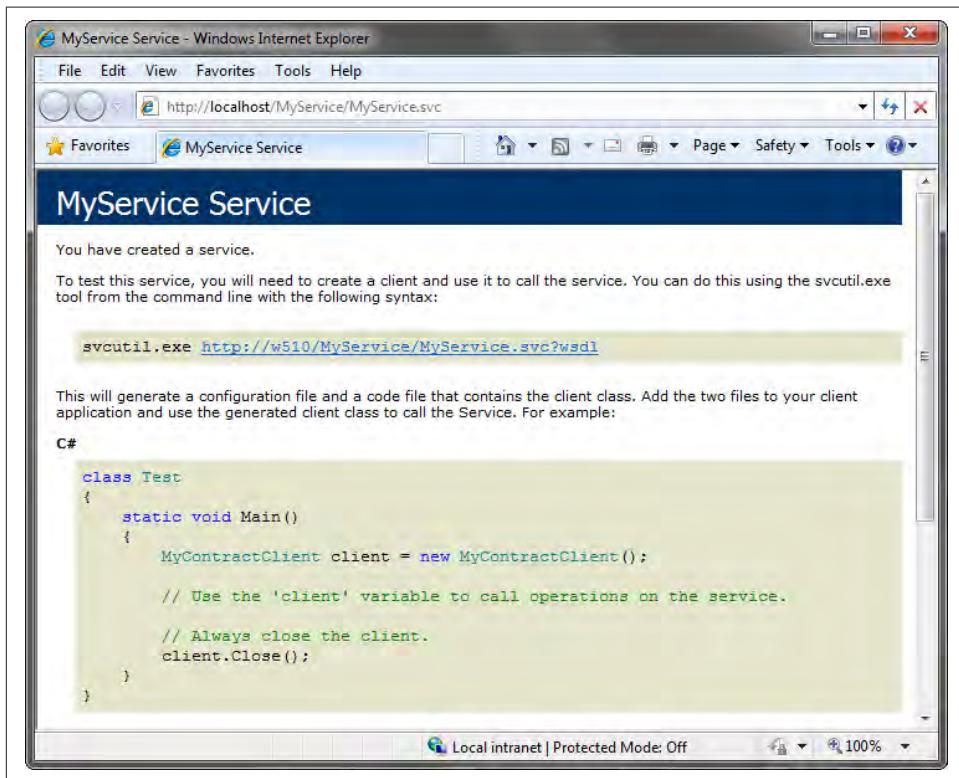


Figure 1-8. A service confirmation page

Enabling metadata exchange programmatically

To enable metadata exchange over HTTP-GET programmatically, you first need to add the behavior to the collection of behaviors the host maintains for the service type. The `ServiceHostBase` class offers the `Description` property of the type `ServiceDescription`:

```
public abstract class ServiceHostBase : ...
{
    public ServiceDescription Description
    {get;}
    //More members
}
```

The service description, as its name implies, is the description of the service with all its aspects and behaviors. `ServiceDescription` contains a property called `Behaviors` of the type `KeyedByTypeCollection<T>`, with `IServiceBehavior` as the generic type parameter:

```
public class KeyedByTypeCollection<T> : KeyedCollection<Type, T>
{
    public U Find<U>();
    public U Remove<U>();
    //More members
}
public class ServiceDescription
{
    public KeyedByTypeCollection<IServiceBehavior> Behaviors
    {get;}

    //More members
}
```

`IServiceBehavior` is the interface that all behavior classes and attributes implement. `KeyedByTypeCollection<T>` offers the generic method `Find<U>()`, which returns the requested behavior if it is in the collection, and `null` otherwise. A given behavior type can be found in the collection at most once.

Example 1-12 shows how to enable the metadata exchange behavior programmatically.

Example 1-12. Enabling the metadata exchange behavior programmatically

```
ServiceHost host = new ServiceHost(typeof(MyService));

ServiceMetadataBehavior metadataBehavior;
metadataBehavior = host.Description.Behaviors.Find<ServiceMetadataBehavior>();
if(metadataBehavior == null)
{
    Debug.Assert(host.BaseAddresses.Any(baseAddress=>baseAddress.Uri.Scheme ==
        "http"));
```

```

        metadataBehavior = new ServiceMetadataBehavior();
        metadataBehavior.HttpGetEnabled = true;
        host.Description.Behaviors.Add(metadataBehavior);
    }

host.Open();

```

Notice the defensive manner in which the hosting code first verifies that no metadata behavior was provided in the config file, by calling the `Find<T>()` method of `KeyedByTypeCollection<I>` and using `ServiceMetadataBehavior` as the type parameter. `ServiceMetadataBehavior` is defined in the `System.ServiceModel.Description` namespace:

```

public class ServiceMetadataBehavior : IServiceBehavior
{
    public bool HttpGetEnabled
    {get;set;}

    public Uri HttpGetUrl
    {get;set;}
    //More members
}

```

If the returned behavior is `null`, it means the config file contains no metadata behavior. In this case, the hosting code creates a new `ServiceMetadataBehavior` instance, sets `HttpGetEnabled` to `true`, and adds it to the behaviors in the service description. By checking defensively for the presence of the behavior first, the hosting code avoids overriding the config file and always allowing the administrator to tweak the behavior or turn it on or off. Note also that the code asserts the presence of an HTTP base address.

The Metadata Exchange Endpoint

Publishing metadata over HTTP-GET is merely a WCF feature; there are no guarantees that other platforms you interact with will support it. There is, however, a standard way of publishing metadata over a special endpoint, called the *metadata exchange endpoint* (sometimes referred to as the *MEX endpoint*). [Figure 1-9](#) shows a service with business endpoints and a metadata exchange endpoint. However, you typically do not show the metadata exchange endpoint in your design diagrams.

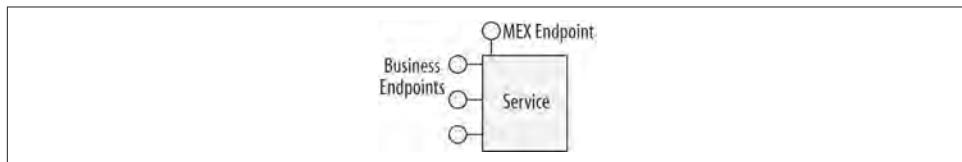


Figure 1-9. The metadata exchange endpoint

The MEX endpoint supports an industry standard for exchanging metadata, represented in WCF by the `IMetadataExchange` interface:

```
[ServiceContract(...)]  
public interface IMetadataExchange  
{  
    [OperationContract(...)]  
    Message Get(Message request);  
    //More members  
}
```

The details of this interface are inconsequential. Like most of these industry standards, it is difficult to implement, but fortunately WCF can have the service host automatically provide the implementation of `IMetadataExchange` and expose the metadata exchange endpoint. All you need to do is designate the address and the binding to use and add the service metadata behavior. For the bindings, WCF provides dedicated binding transport elements for the HTTP, HTTPS, TCP, and IPC protocols. For the address, you can provide a full address or use any of the registered base addresses. There is no need to enable the HTTP-GET option, but there is no harm in doing so. [Example 1-13](#) shows a service that exposes three MEX endpoints, over HTTP, TCP, and IPC. For demonstration purposes, the TCP and IPC MEX endpoints use relative addresses and the HTTP endpoint uses an absolute address.

Example 1-13. Adding MEX endpoints

```
<services>  
    <service name = "MyService" behaviorConfiguration = "MEX">  
        <host>  
            <baseAddresses>  
                <add baseAddress = "net.tcp://localhost:8001/" />  
                <add baseAddress = "net.pipe://localhost/" />  
            </baseAddresses>  
        </host>  
        <endpoint  
            address = "MEX"  
            binding = "mexTcpBinding"  
            contract = "IMetadataExchange"  
        />  
        <endpoint  
            address = "MEX"  
            binding = "mexNamedPipeBinding"  
            contract = "IMetadataExchange"  
        />  
        <endpoint  
            address = "http://localhost:8000/MEX"  
            binding = "mexHttpBinding"  
            contract = "IMetadataExchange"  
        />  
    ...
```

```

</service>
</services>
<behaviors>
  <serviceBehaviors>
    <behavior name = "MEX">
      <serviceMetadata/>
    </behavior>
  </serviceBehaviors>
</behaviors>

```



In [Example 1-13](#), all you need to do to have the host implement the MEX endpoint for your service is include the `serviceMetadata` tag in the behavior. If you do not reference the behavior, the host will expect your service to implement `IMetadataExchange`. While this normally adds no value, it is the only way to provide for custom implementation of `IMetadataExchange` for advanced interoperability needs.

Standard endpoints

In the vast majority of cases, a MEX endpoint always has the same three elements: the contract is always `IMetadataExchange`, the binding is always the reserved binding element, and the only variable is the address (and even that is typically just the base address). Having developers stipulate these endpoint elements time and time again is excessive. To streamline this and similar infrastructure endpoints, WCF provides pre-canned definitions of several endpoint types, called *standard endpoints*. WCF offers standard endpoints for metadata exchange, discovery, announcements, workflow, and web. You can use the standard endpoints both in config and programmatically.

You can reference the desired standard endpoint with the `kind` tag:

```

<endpoint
  kind = "..."
/>

```

Whatever is not specified (usually the address or the binding) always defaults to some predefined value, depending on the other fields of the endpoint. [Appendix C](#) will take advantage of the standard discovery and announcements endpoints. In the context of this section, you can use the `kind` value of `mexEndpoint` to define the MEX endpoint.

For example, suppose you do not specify an address and binding, like so:

```

<service ...
  <host>
    <baseAddresses>
      <add baseAddress = "http://..."/>
      <add baseAddress = "net.tcp://..."/>
    </baseAddresses>
  </host>

```

```
<endpoint  
    kind = "mexEndpoint"  
/>  
...  
</service>
```

WCF will add a MEX endpoint whose address is the HTTP base address. This mandates the presence of an HTTP base address and that no other endpoint is using the base address for its address.

You can also append a URI to the base address:

```
<endpoint  
    kind = "mexEndpoint"  
    address = "MEX"  
/>
```

If you specify the binding, WCF will infer the correct base address to use from the binding type, for example:

```
<service ...  
    <host>  
        <baseAddresses>  
            <add baseAddress = "http://..."/>  
            <add baseAddress = "net.tcp://..."/>  
        </baseAddresses>  
    </host>  
  
    <endpoint  
        kind = "mexEndpoint"  
        binding = "mexTcpBinding"  
    />  
    <endpoint  
        kind = "mexEndpoint"  
        address = "MEX"  
        binding = "mexTcpBinding"  
    />  
    ...  
</service>
```

You can also specify a fully qualified address irrespective of the base address.

Note that WCF is not smart enough to infer the binding to use from the address scheme, meaning the following configuration is invalid:

```
<!-- Invalid configuration -->  
<endpoint  
    kind = "mexEndpoint"  
    address = "net.tcp://..."  
/>
```

Adding MEX endpoints programmatically

Like any other endpoint, you can only add a metadata exchange endpoint programmatically before opening the host. WCF does not offer a dedicated binding type for the metadata exchange endpoint. Instead, you need to construct a custom binding that uses the matching transport binding element and provide that binding element as a construction parameter to an instance of a custom binding. To streamline this process, use the `MetadataExchangeBindings` static helper class defined as:

```
public static class MetadataExchangeBindings
{
    public static Binding CreateMexHttpBinding();
    public static Binding CreateMexNamedPipeBinding();
    public static Binding CreateMexTcpBinding();

    //More members
}
```

Finally, call the `AddServiceEndpoint()` method of the host, providing it with the address, the MEX binding, and the `IMetadataExchange` contract type. [Example 1-14](#) shows the code required to add a MEX endpoint over TCP. Note that before adding the endpoint, you must verify the presence of the metadata behavior.

Example 1-14. Adding a TCP MEX endpoint programmatically

```
Uri tcpBaseAddress = new Uri("net.tcp://localhost:9000/");
ServiceHost host = new ServiceHost(typeof(MyService),tcpBaseAddress);

ServiceMetadataBehavior metadataBehavior;
metadataBehavior = host.Description.Behaviors.Find<ServiceMetadataBehavior>();
if(metadataBehavior == null)
{
    metadataBehavior = new ServiceMetadataBehavior();
    host.Description.Behaviors.Add(metadataBehavior);
}
Binding binding = MetadataExchangeBindings.CreateMexTcpBinding();
host.AddServiceEndpoint(typeof(IMetadataExchange),binding,"MEX");
host.Open();
```

You can also add a MEX endpoint using the standard MEX endpoint. To do so, use the type `ServiceMetadataEndpoint`, defined as:

```
public class ServiceMetadataEndpoint : ServiceEndpoint
{
    public ServiceMetadataEndpoint();
    public ServiceMetadataEndpoint(EndpointAddress address);
    public ServiceMetadataEndpoint(Binding binding,EndpointAddress address);
}
```

The default constructor of `ServiceMetadataEndpoint` defaults to using the HTTP base address and binding. The constructor that takes an endpoint address must receive a fully qualified HTTP or HTTPS address:

```
ServiceHost host = new ServiceHost(typeof(MyService));
host.Description.Behaviors.Add(new ServiceMetadataBehavior());

EndpointAddress address = new EndpointAddress("http://localhost:8000/MEX");

ServiceEndpoint endpoint = new ServiceMetadataEndpoint(address);
host.AddServiceEndpoint(endpoint);
...
host.Open();
```

In addition, `ServiceMetadataEndpoint` will never use the host base addresses.

Streamlining with `ServiceHost<T>`

You can extend `ServiceHost<T>` to automate the code in Examples 1-12 and 1-14. `ServiceHost<T>` offers the `EnableMetadataExchange()` method, which you can call to both publish metadata over HTTP-GET and add the MEX endpoints:

```
public class ServiceHost<T> : ServiceHost
{
    public void EnableMetadataExchange(bool enableHttpGet = true);

    public bool HasMexEndpoint
    {get;}
    public void AddAllMexEndPoints();
    //More members
}
```

The default `EnableMetadataExchange()` publishes metadata over HTTP-GET, and if no MEX endpoint is available, `EnableMetadataExchange()` adds a MEX endpoint for each registered base address scheme. Using `ServiceHost<T>`, 1-12 and 1-14 are reduced to:

```
ServiceHost<MyService> host = new ServiceHost<MyService>();
host.EnableMetadataExchange();
host.Open();
```

`EnableMetadataExchange()` will not override the behavior in the config file if one is present.

`ServiceHost<T>` offers the `HasMexEndpoint` Boolean property, which returns `true` if the service has any MEX endpoint (regardless of transport protocol), and the `AddAllMexEndPoints()` method, which adds a MEX endpoint for each registered base address of the scheme type of HTTP, TCP, or IPC. Example 1-15 shows the implementation of these methods.

Example 1-15. Implementing EnableMetadataExchange and its supporting methods

```
public class ServiceHost<T> : ServiceHost
{
    public void EnableMetadataExchange(bool enableHttpGet = true)
    {
        if(State == CommunicationState.Opened)
        {
            throw new InvalidOperationException("Host is already opened");
        }
        ServiceMetadataBehavior metadataBehavior
            = Description.Behaviors.Find<ServiceMetadataBehavior>();

        if(metadataBehavior == null)
        {
            metadataBehavior = new ServiceMetadataBehavior();
            Description.Behaviors.Add(metadataBehavior);

            if(BaseAddresses.Any(uri=>uri.Scheme == "http"))
            {
                metadataBehavior.HttpGetEnabled = enableHttpGet;
            }
        }
        AddAllMexEndPoints();
    }
    public bool HasMexEndpoint
    {
        get
        {
            return Description.Endpoints.Any(
                endpoint=>endpoint.Contract.ContractType ==
                typeof(IMetadataExchange));
        }
    }
    public void AddAllMexEndPoints()
    {
        Debug.Assert(HasMexEndpoint == false);

        foreach(Uri baseAddress in BaseAddresses)
        {
            Binding binding = null;

            switch(baseAddress.Scheme)
            {
                case "net.tcp":
                {
                    binding = MetadataExchangeBindings.CreateMexTcpBinding();
                    break;
                }
                case "net.pipe":
                {...}
                case "http":
```

```
        {...}
        case "https":
        {...}
    }
    if(binding != null)
    {
        AddServiceEndpoint(typeof(IMetadataExchange),binding,"MEX");
    }
}
```

`EnableMetadataExchange()` verifies that the host has not been opened yet using the `State` property of the `CommunicationObject` base class. The `HasMexEndpoint` property checks whether a given endpoint's contract is indeed `IMetadataExchange`. `Any()` invokes the expression on the endpoints in the collection, returning `true` when any one of the endpoints in the collection satisfies the predicate (that is, if the invocation of the lambda expression method returned `true`) and `false` otherwise. The `AddAllMexEndPoints()` method iterates over the `BaseAddresses` collection. For each base address found, it creates a matching MEX binding and adds the MEX endpoint with a MEX URI under the base address.

The Metadata Explorer

The metadata exchange endpoint provides metadata that describes not just contracts and operations, but also information about data contracts, security, transactions, reliability, and faults. To visualize the metadata of a running service, I developed the Metadata Explorer tool, which is available along with the rest of the source code for this book. [Figure 1-10](#) shows the Metadata Explorer reflecting the endpoints of [Example 1-7](#). To use the Metadata Explorer, simply provide it with the HTTP-GET address or the metadata exchange endpoint of the running service, and it will reflect the returned metadata.

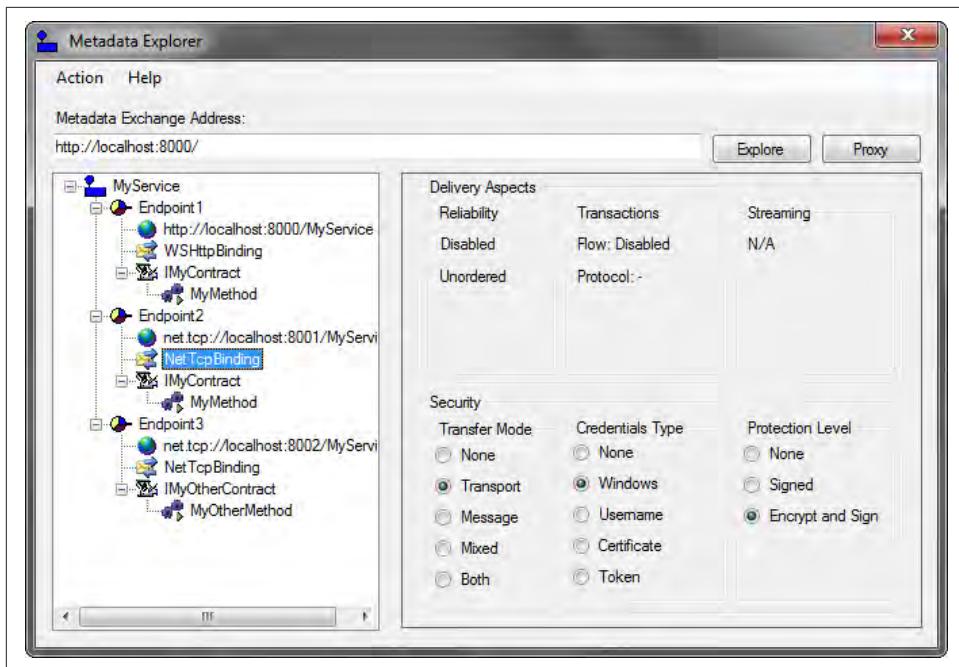


Figure 1-10. The Metadata Explorer

The WCF-Provided Test Host

Visual Studio ships with a ready-made, general-purpose service host. The executable is called *WcfSvcHost.exe* and after a normal installation, it is found in *C:\Program Files\Microsoft Visual Studio <version>\Common7\IDE*.

For ease of use, I recommend adding that location to your system's Path variable. *WcfSvcHost* is a simple command-line utility that accepts two parameters: the file-name (and path) reference to a .NET assembly containing the service class or classes, and a filename (and path) reference to the host config file. For example:

```
WcfSvcHost.exe /service:MyService.dll /config:App.config
```

The specified service assembly can be a class library assembly (DLL) or an application assembly (EXE). *WcfSvcHost* launches a new process that automatically hosts all the service classes listed in the services section of the specified config file. These service classes and their service contracts and data contracts need not be public types—they can be internal. In addition, the autohosted services need not provide any metadata, but they can publish metadata if you configure them to.

WcfSvcHost is a Windows Forms application that resides as a desktop tray icon. If you click the *WcfSvcHost* tray icon, it will bring up a dialog box listing all the hosted services. The dialog box also shows the status of the service and its metadata address,

which you can copy to the clipboard (perhaps for use later, when adding a reference to the service). Closing the WcfSvcHost UI merely collapses it back to the tray. To close the host, simply select Exit from the tray icon's context menu. Terminating the host this way is an ungraceful exit; WcfSvcHost will abort all calls currently in progress and the clients are likely to receive an exception.

WcfSvcHost negates the need during development for a separate host assembly accompanying your service library. Developing such host projects is a repetitive task, as these hosts typically contain the same lines of code over and over again, and they tend to bloat the solution when you have multiple service libraries. For development and testing purposes, you can integrate WcfSvcHost directly into your Visual Studio service library projects. In the Debug pane of the project properties, specify *WcfSvcHost.exe* as the external program to start and your class library name and its config file (the one autogenerated and autocopied to the *bin* folder) as arguments.

With that done, when you launch the class library it will automatically be hosted by WcfSvcHost, with the debugger attached to that process. When you stop debugging, Visual Studio will abort the host ungracefully.

The last feature of WcfSvcHost is its ability to automatically launch a client application and even provide the client with optional parameters specific for that application:

```
WcfSvcHost.exe /service:MyService.dll /config:App.config  
               /client:MyClient.exe   /clientArg:123,ABC
```

This is useful in automated testing and even simple deployment scenarios to launch both the host and the client.

The main advantage of using WcfSvcHost is that during development, you will not need to develop, build, and own a separate host project. Its major drawback is that it is only suitable for simple scenarios in which you do not require programmatic access to the host instance before opening it or programmatic access to its event model once it is open. Unlike hosting with IIS or WAS, there is no equivalent service host factory support. Consequently, there is no ability to dynamically add base addresses, configure endpoints, throttle the calls, configure custom behaviors at the host level, and so on. My experience with WCF is that in all but the simplest cases, eventually you will need programmatic access to the host instance. Thus, I do not view WcfSvcHost as a full-fledged production-worthy host, like WAS or a dedicated self-host.

More on Behavior Configuration

The service metadata behavior demonstrated in the previous section is just one of many such ready-to-use behaviors, and you will see many examples in the subsequent chapters. You can configure behaviors at the service level (as with the metadata behavior) or at the endpoint level:

```

<services>
    <service name = "MyService" behaviorConfiguration = "MyServiceBehavior">
        <endpoint behaviorConfiguration = "MyEndpointBehavior"
            ...
        />
    </service>
</services>
<behaviors>
    <endpointBehaviors>
        <behavior name = "MyEndpointBehavior">
            ...
        </behavior>
    </endpointBehaviors>
    <serviceBehaviors>
        <behavior name = "MyServiceBehavior">
            ...
        </behavior>
    </serviceBehaviors>
</behaviors>

```

Similar to default bindings, WCF allows for the notion of a default behavior. A *default behavior* is a nameless behavior (either a service or an endpoint level) that implicitly affects all services or endpoints that do not explicitly reference a behavior configuration. For example, consider the services MyService1, MyService2, and MyService3. To add the service metadata exchange on all services in the application, you can use this config file:

```

<services>
    <service name = "MyService1">
        ...
    </service>
    <service name = "MyService2">
        ...
    </service>
</services>
<behaviors>
    <serviceBehaviors>
        <behavior>
            <serviceMetadata/>
        </behavior>
    </serviceBehaviors>
</behaviors>

```

Along with this hosting code:

```

ServiceHost host1 = new ServiceHost(typeof(MyService1));
ServiceHost host2 = new ServiceHost(typeof(MyService2));
ServiceHost host3 = new ServiceHost(typeof(MyService3));
host3.AddServiceEndpoint(...);

host1.Open();

```

```
host2.Open();
host3.Open();
```

Note that the default behavior affects all services in the application that do not reference a behavior, even those (like `MyService3`) that do not rely at all on the config file.

You can have at most one default service behavior and one default endpoint behavior.

As with default bindings, the problem with the default behaviors is that a config file may get difficult for humans to parse and understand once you combine default behaviors with named behaviors, as shown in [Figure 1-11](#).

The diagram shows a portion of a WCF configuration file. It includes sections for services and behaviors. Services are defined with attributes like name and behaviorConfiguration. Behaviors are defined with attributes like name. A box highlights the behaviors section, specifically the serviceBehaviors and behavior elements. Arrows point from the highlighted elements to their corresponding definitions further down in the code.

```
<services>
  <service name = "...> behaviorConfiguration = "MyBehavior"> ←
    ...
  </service>

  <service name = "...> ←
    ...
  </service>
  <service name = "...> ←
    ...
  </service>
</services>

<behaviors>
  <serviceBehaviors>
    <behavior> ←
      ...
      </behavior>
      <behavior name = "MyBehavior"> ←
        ...
        </behavior>
    </serviceBehaviors>
  </behaviors>
```

Figure 1-11. Named and default behavior configuration

Because of this difficulty, coupled with the side effect of implicitly affecting even services that do not rely on the config file at all, I recommend exercising caution when utilizing a default behavior. Only use it when you want to affect all services in the application. Never mix and match them with named behaviors, since any service using a named behavior will be exempt from the default behavior. In the interest of readability, most if not all of the behaviors in this book are explicitly named, except in the rare cases in which a default behavior is required.

Client-Side Programming

To invoke operations on a service, a client first needs to import the service contract to the client's native representation. If the client uses WCF, the common way of invoking operations is to use a proxy. The proxy is a CLR class that exposes a single CLR interface representing the service contract. If the service supports several contracts (over

at least as many endpoints), the client needs one proxy per contract type. The proxy provides the same operations as the service's contract, but also has additional methods for managing the proxy lifecycle and the connection to the service. The proxy completely encapsulates every aspect of the service: its location, its implementation technology and runtime platform, and the communication transport.

Generating the Proxy

You can use Visual Studio to import the service metadata and generate a proxy. If the service is self-hosted outside the solution, first launch the service and then select Add Service Reference from the client project's context menu. If the service is self-hosted in the same solution, first launch it without the debugger, and then select the Add Service Reference option from the context menu.

If the service is hosted in IIS or WAS, there is no need to prelaunch the service; simply select Add Service Reference from the client project's context menu, and Visual Studio will bring up the Add Service Reference dialog, shown in [Figure 1-12](#).



For the Add Service Reference option to appear in a project's context menu, the project must be configured to target .NET Framework 3.0 or later.

In the Add Service Reference dialog box, specify the service metadata address (not the service URL, as the dialog box states) and click Go to view the available service endpoints (not Services, as labeled). Specify a namespace (such as `MyService`) to contain the generated proxy, then click OK to generate the proxy and update the config file. Use the Discover button to discover WCF services in your own solution, as long as they are hosted either in a website project or in one of the WCF service library project types. In the case of a website project, Visual Studio will either retrieve the metadata from IIS or launch IIS Express. In the case of a WCF service library, WCF will automatically launch its host (`WcfSvcHost`, described in [“The WCF-Provided Test Host” on page 55](#)) to get the metadata.

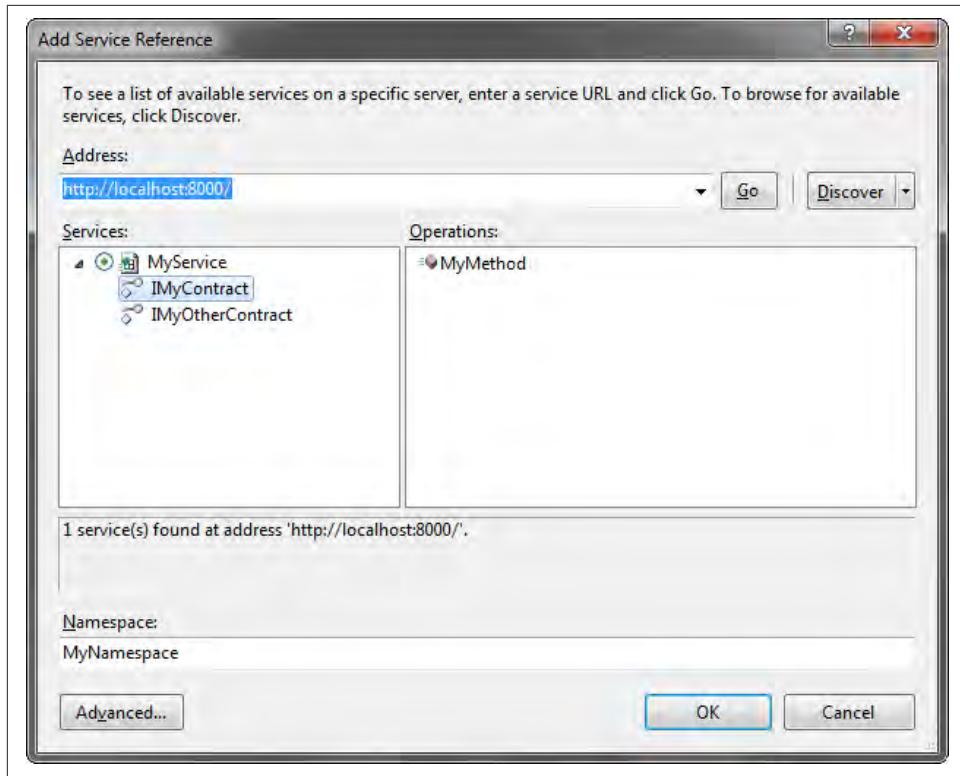


Figure 1-12. Generating a proxy using Visual Studio

Click the Advanced button to bring up the Service Reference Settings dialog box, where you can tweak the proxy generation options (see [Figure 1-13](#)).

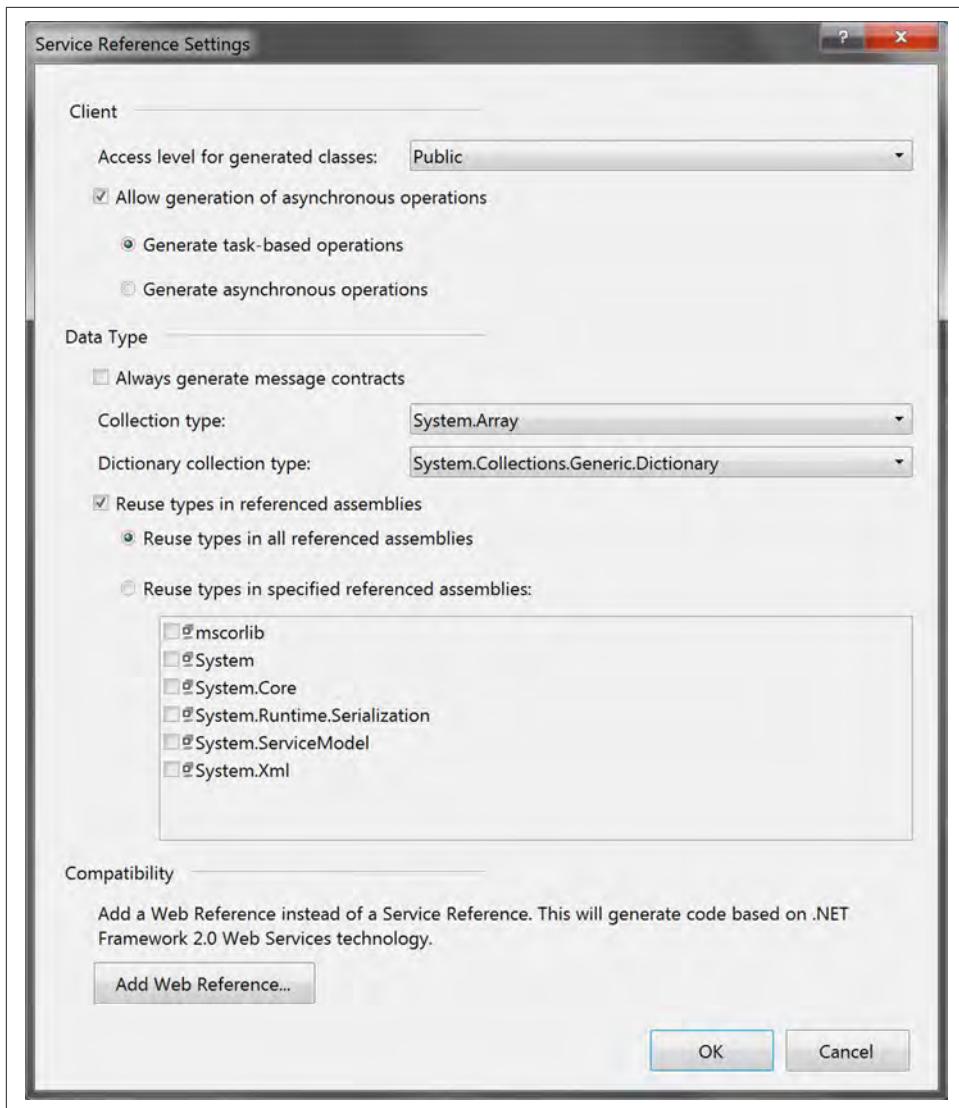


Figure 1-13. Advanced options for the service reference

The more intuitive options let you configure the visibility of the generated proxy and contracts (public or internal), and you can generate message contracts for your data types for advanced interoperability scenarios, where you have to comply with an existing (typically custom) message format. You can also click the Add Web Reference button to convert the reference to an old ASMX web service reference, as long as the service is using the basic binding.

Once you've added a reference, your project will have a new folder called *Service References*. In it, you'll find a service reference item for each referenced service (see Figure 1-14).

At any point, you can right-click a reference and select Update Service Reference to regenerate the proxy. This is possible because each service reference item also contains a file that records the original metadata address used. You can also select Configure Service Reference to bring up a dialog box similar to the advanced settings dialog box used when adding a reference. The Configure Service Reference dialog box lets you change the service metadata address, as well as the rest of the advanced proxy settings.

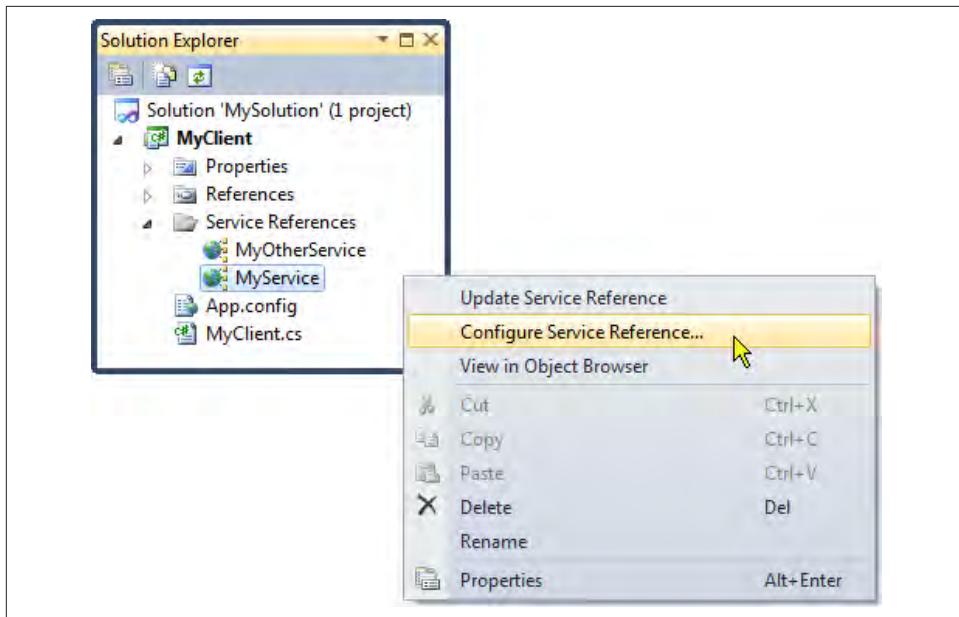


Figure 1-14. The Service References folder

Generating the proxy using SvcUtil

As an alternative to Visual Studio, you can use the *SvcUtil.exe* command-line utility to import the service metadata and generate a proxy. You need to provide SvcUtil with the metadata exchange address and, optionally, with a proxy filename. The default proxy filename is *output.cs*, but you can use the */out* switch to indicate a different name.

For example, if you're hosting the service *MyService* in the WAS and have enabled metadata publishing over HTTP-GET, you can simply run this command line:

```
SvcUtil http://localhost/MyService/MyService.svc /out:Proxy.cs
```

With self-hosting, suppose that the self-hosted service has enabled metadata publishing over HTTP-GET on the address

```
http://localhost:8000/
```

and has exposed MEX endpoints using these addresses:

```
http://localhost:8000/MEX  
http://localhost:8001/MEX  
net.tcp://localhost:8002/MEX  
net.pipe://localhost/MyPipe/MEX
```

After launching the host, you'll be able to use the following commands to generate the proxy:

```
SvcUtil http://localhost:8000 /out:Proxy.cs  
SvcUtil http://localhost:8000/MEX /out:Proxy.cs  
SvcUtil http://localhost:8001/MEX /out:Proxy.cs  
SvcUtil net.tcp://localhost:8002/MEX /out:Proxy.cs  
SvcUtil net.pipe://localhost/MyPipe/MEX /out:Proxy.cs
```



The main advantage of using SvcUtil over Visual Studio is that you can include the command line for generating the proxy as a pre-build event.

SvcUtil offers numerous command-line switches that correspond to the options in the Visual Studio advanced settings dialog box shown in [Figure 1-14](#).

Regardless of whether you use Visual Studio or SvcUtil to generate the proxy, [Example 1-16](#) shows the imported contract and generated proxy for this service definition:

```
[ServiceContract(Namespace = "MyNamespace")]
interface IMyContract
{
    [OperationContract]
    void MyMethod();
}
class MyService : IMyContract
{
    public void MyMethod()
    {...}
}
```

You can safely remove many of the gunk attributes the tools generate, which merely state the defaults, so that you end up with the cleaned-up proxy shown in [Example 1-16](#).

Example 1-16. Client proxy file

```
[ServiceContract(Namespace = "MyNamespace")]
interface IMyContract
{
    [OperationContract]
    void MyMethod();
}

class MyContractClient : ClientBase<IMyContract>, IMyContract
{
    public MyContractClient()
    {}
    public MyContractClient(string endpointName) : base(endpointName)
    {}
    public MyContractClient(Binding binding, EndpointAddress remoteAddress) :
        base(binding, remoteAddress)
    {}

    /* Additional constructors */

    public void MyMethod()
    {
        Channel.MyMethod();
    }
}
```



You can also instruct SvcUtil to forgo generating the proxy and only generate the service and data contracts for your service by specifying the /servicecontract switch.

The most glaring aspect of the proxy class is that it has no reference to the service-implementing class, only to the contract exposed by the service. You can use the proxy in conjunction with a client-side config file that provides the address and the binding, or you can use it without a config file. Note that each proxy instance points at exactly one endpoint. The endpoint to interact with is provided to the proxy at construction time. As mentioned previously, if the service-side contract does not provide a namespace, it will default to the <http://tempuri.org> namespace.

Administrative Client Configuration

The client needs to know where the service is located and must use the same binding as the service, and of course, import the service contract definition. In essence, this is exactly the same information captured in the service's endpoint. To reflect that, the client config file can contain information about the target endpoints and even uses the same endpoint configuration scheme as the host.

Example 1-17 shows the client configuration file required to interact with a service whose host is configured according to **Example 1-6**.

Example 1-17. Client config file

```
<system.serviceModel>
  <client>
    <endpoint name = "MyEndpoint"
      address   = "http://localhost:8000/MyService"
      binding   = "wsHttpBinding"
      contract  = "IMyContract"
    />
  </client>
</system.serviceModel>
```

The client config file may list as many endpoints as the services it deals with support, and the client may use any one of them. **Example 1-18** shows the client config file matching the host config file of **Example 1-7**. There is no relationship between the various endpoints in the client's config file: they could all be pointing at the same endpoint on the service, at different endpoints on the service, at different endpoints on different services, or any mix and match in between. Note that on the client side, you typically name endpoints with unique names (you will see why shortly). Naming the endpoints on the client side is optional, just as it is optional on the service side, yet on the service side you typically do not name the endpoints, while on the client side, you typically do.

Example 1-18. Client config file with multiple target endpoints

```
<system.serviceModel>
  <client>
    <endpoint name = "FirstEndpoint"
      address   = "http://localhost:8000/MyService"
      binding   = "wsHttpBinding"
      contract  = "IMyContract"
    />
    <endpoint name = "SecondEndpoint"
      address   = "net.tcp://localhost:8001/MyService"
      binding   = "netTcpBinding"
      contract  = "IMyContract"
    />
    <endpoint name = "ThirdEndpoint"
      address   = "net.tcp://localhost:8002/MyService"
      binding   = "netTcpBinding"
      contract  = "IMyOtherContract"
    />
  </client>
</system.serviceModel>
```

Binding configuration

You can customize the client-side bindings to match the service binding in a manner identical to the service configuration, as shown in [Example 1-19](#).

Example 1-19. Client-side binding configuration

```
<system.serviceModel>
  <client>
    <endpoint name = "MyEndpoint"
      address   = "net.tcp://localhost:8000/MyService"
      bindingConfiguration = "TransactionalTCP"
      binding   = "netTcpBinding"
      contract  = "IMyContract"
    />
  </client>
  <bindings>
    <netTcpBinding>
      <binding name = "TransactionalTCP"
        transactionFlow = "true"
      />
    </netTcpBinding>
  </bindings>
</system.serviceModel>
```

The client can rely on a default binding configuration, just as the service does. Note that if the config file contains both client and service sections, the default binding will affect both.

Generating the client config file

When you add a service reference in Visual Studio, it will also try to automatically edit the client's config file and insert the required `client` section describing the service's endpoints in it. In versions of Visual Studio prior to 2012, Visual Studio was not smart enough to infer the cleanest binding values and it would therefore butcher the config file by stating all the default values for the bindings, which effectively rendered the file unreadable. Visual Studio 2012 and later solves this problem by only adding values to the config file for binding properties that are set to nondefault values.

Like Visual Studio, SvcUtil also autogenerates a client-side config file called `output.config`. You can specify a different config filename using the `/config` switch:

```
SvcUtil http://localhost:8002/MyService/out:Proxy.cs /config:App.Config
```

As with Visual Studio, the SvcUtil with .NET 4.5 is smart enough to only produce configuration values for binding properties that are set to nondefault values. However, unlike with Visual Studio, with SvcUtil you can suppress generating the config file by using the `/noconfig` switch:

```
SvcUtil http://localhost:8002/MyService/out:Proxy.cs /noconfig
```

Although these tools no longer litter the config file with redundant default values, we still recommend never letting SvcUtil or Visual Studio control the config file.

In-proc configuration

With in-proc hosting, the client config file is also the service host config file, and the same file contains both service and client entries, as shown in [Example 1-20](#).

Example 1-20. In-proc hosting config file

```
<system.serviceModel>
  <services>
    <service name = "MyService">
      <endpoint
        address  = "net.pipe://localhost/MyPipe"
        binding   = "netNamedPipeBinding"
        contract = "IMyContract"
      />
    </service>
  </services>
  <client>
    <endpoint name = "MyEndpoint"
      address  = "net.pipe://localhost/MyPipe"
      binding   = "netNamedPipeBinding"
      contract = "IMyContract"
    />
  </client>
</system.serviceModel>
```

Note the use of the named pipe binding for in-proc hosting.

The SvcConfigEditor

WCF provides a config file editor called *SvcConfigEditor.exe* that can edit both host and client configuration files (see [Figure 1-15](#)). You can launch the editor from within Visual Studio by right-clicking on the configuration file (for either the client or the host) and selecting Edit WCF Configuration.



If you are using a version of Visual Studio earlier than 2012, you will have to launch the editor first from the Tools menu.

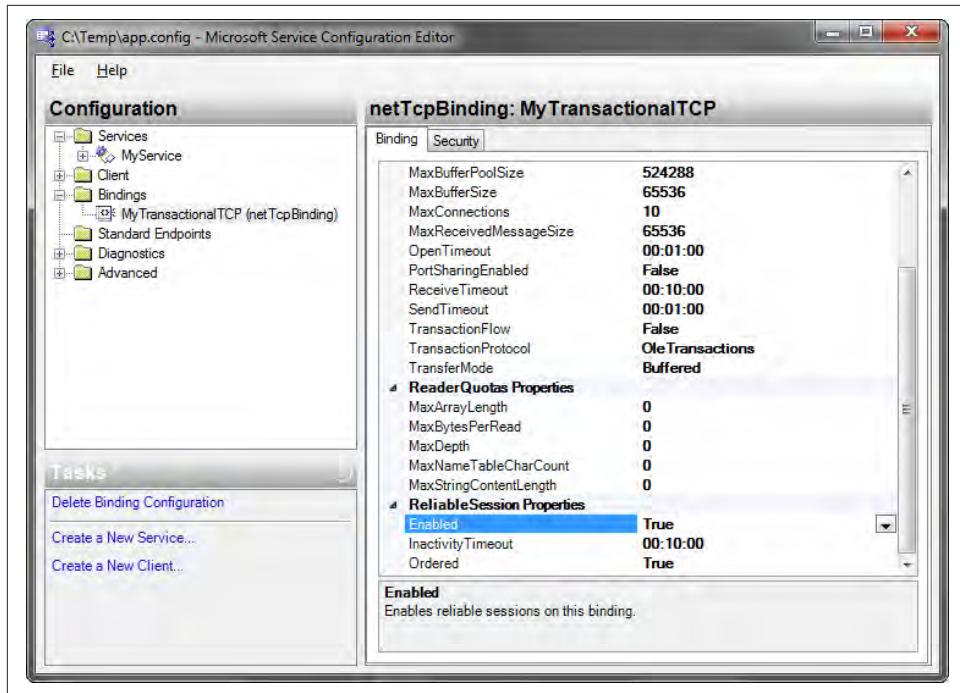


Figure 1-15. SvcConfigEditor is used to edit both host and client config files

I have mixed feelings about SvcConfigEditor. On the one hand, it edits the config files nicely and it saves developers the need to learn the configuration scheme. On the other hand, it does not shield developers from needing to thoroughly understand WCF configuration, and it's often faster to do the light editing that's typically required in a config file by hand than it is using Visual Studio.

Working with the proxy

The proxy class derives from the class `ClientBase<T>`, defined as:

```
public abstract class ClientBase<T> : ICommunicationObject, IDisposable
{
    protected ClientBase(string endpointName);
    protected ClientBase(Binding binding, EndpointAddress remoteAddress);

    public void Open();
    public void Close();

    protected T Channel
    {get; }

    //Additional members
}
```

`ClientBase<T>` accepts a single generic type parameter identifying the service contract that this proxy encapsulates. The `Channel` property of `ClientBase<T>` is of that type parameter. As shown in [Example 1-16](#), the generated subclass of `ClientBase<T>` simply delegates the method call to `Channel`. Calling the method on the `Channel` property sends the appropriate WCF message to the service.

To use the proxy, the client first needs to instantiate a proxy object and to provide the constructor with endpoint information: either the endpoint section name from the config file or the endpoint address and binding objects if you're not using a config file. The client can then use the proxy methods to call the service and when it is done, the client needs to close the proxy instance. For example, given the same definitions as in Examples [1-16](#) and [1-17](#), the client constructs the proxy, identifying the endpoint to use from the config file, and then invokes the method and closes the proxy:

```
MyContractClient proxy = new MyContractClient("MyEndpoint");
proxy.MyMethod();
proxy.Close();
```

When specifying the endpoint name to the proxy, its constructor also verifies that the contract configured for that endpoint matches the proxy's type parameter. Because of this verification ability, if exactly one endpoint is defined in the client config file for the type of contract the proxy is using, the client can omit the endpoint name from the proxy's constructor:

```
MyContractClient proxy = new MyContractClient();
proxy.MyMethod();
proxy.Close();
```

The proxy will simply look up that endpoint (named or not in the config file) and use it. However, if you use this technique when multiple (or zero) endpoints are available for the same contract type, the proxy will throw an exception.

Closing the proxy

It is a recommended best practice to always close the proxy when the client is done using it. Closing the proxy releases the connection held toward the service, which is particularly important to do in the presence of a transport session (as discussed later in this chapter). It also helps ensure the threshold for the maximum number of connections on the client's machine is not reached. Furthermore, as you will see in [Chapter 4](#), closing the proxy terminates the session with the service instance.

Instead of closing the proxy, you can use its `Dispose()` method. Internally, `Dispose()` just calls `Close()`. The advantage of the `Dispose()` method is that you can use the `using` statement to call it even in the face of exceptions:

```
using(MyContractClient proxy = new MyContractClient())
{
```

```
        //Any exception here automatically closes the proxy;  
    }
```

If the client is declaring the contract directly instead of the concrete proxy class, the client can query for the presence of `IDisposable`:

```
IMyContract proxy = new MyContractClient();  
proxy.MyMethod();  
IDisposable disposable = proxy as IDisposable;  
if(disposable != null)  
{  
    disposable.Dispose();  
}
```

Alternatively, the client can collapse the query inside the `using` statement:

```
IMyContract proxy = new MyContractClient();  
using(proxy as IDisposable)  
{  
    proxy.MyMethod();  
}
```



While the results of calling `Dispose()` and `Close()` are identical, you will see in [Chapter 6](#) that it is always better to call `Close()` than to use the `using` statement.

Call timeout

Each call made by a WCF client must complete within a configurable timeout. If, for whatever reason, the call duration exceeds the timeout, the call is aborted and the client gets a `TimeoutException`. This behavior is very handy, since it offers an elegant way to deal with deadlocks on the service side or just poor availability. In traditional .NET, the client has to spin a worker thread and have the worker thread call the class (and potentially hang), and the client then monitors some timed-out event that the worker thread has to signal when done. This is obviously a complicated programming model. The advantage of using a proxy for every call is that the proxy can do all this for you. The exact value of the timeout is a property of the binding, and the default timeout is one minute. To provide a different timeout, set the `SendTimeout` property of the abstract `Binding` base class:

```
public abstract class Binding : ...  
{  
    public TimeSpan SendTimeout  
    {get;set;}  
    //More members  
}
```

For example, here's how to configure the `WSHttpBinding` with a five-minute call timeout:

```
<client>
    <endpoint
        ...
        binding = "wsHttpBinding"
        bindingConfiguration = "LongTimeout"
        ...
    />
</client>
<bindings>
    <wsHttpBinding>
        <binding name = "LongTimeout" sendTimeout = "00:05:00"/>
    </wsHttpBinding>
</bindings>
```

Programmatic Client Configuration

Instead of relying on a config file, the client can programmatically construct address and binding objects matching the service endpoint and provide them to the proxy constructor. There is no need to provide the contract, since that was provided in the form of the generic type parameter of the proxy. To represent the address, the client needs to instantiate an `EndpointAddress` class, defined as:

```
public class EndpointAddress
{
    public EndpointAddress(string uri);
    //More members
}
```

Example 1-21 demonstrates this technique, showing the code equivalent of Example 1-17 targeting the service in Example 1-9.

Example 1-21. Programmatic client configuration

```
Binding wsBinding = new WSHttpBinding();
EndpointAddress endpointAddress = new EndpointAddress(
    "http://localhost:8000/MyService");

MyContractClient proxy = new MyContractClient(wsBinding, endpointAddress);
proxy.MyMethod();
proxy.Close();
```

Similar to using a `binding` section in a config file, the client can programmatically configure the binding properties:

```
WSHttpBinding wsBinding = new WSHttpBinding();
wsBinding.SendTimeout = TimeSpan.FromMinutes(5);
wsBinding.TransactionFlow = true;
```

```

EndpointAddress endpointAddress = new EndpointAddress
    ("http://localhost:8000/MyService");

MyContractClient proxy = new MyContractClient(wsBinding,endpointAddress);
proxy.MyMethod();
proxy.Close();

```

Again, note the use of the concrete subclass of `Binding` in order to access binding-specific properties such as the transaction flow.

The WCF-Provided Test Client

Visual Studio ships with a simple general-purpose test client for rudimentary testing that you can use to invoke operations on most services. The test client is called `WcfTestClient.exe`, and after a normal installation it is found in `C:\Program Files\Microsoft Visual Studio <version>\Common7\IDE`. You can provide `WcfTestClient` with a single command-line argument containing the metadata address of the service to test:

```
WcfTestClient.exe http://localhost:9000/
```

You can specify any metadata address (be it an HTTP-GET address or a metadata exchange endpoint over HTTP, TCP, or IPC). You can also specify multiple metadata addresses:

```
WcfTestClient.exe http://localhost:8000/ net.tcp://localhost:9000/MEX
```

You can also launch the test client without a command-line parameter. Once it's running, you can add a new service by selecting Add Service from the File menu and specifying the metadata address in the Add Service dialog box. You can also remove a service by right-clicking it in the services tree.

`WcfTestClient` is a Windows Forms application. The tree control in the left pane contains the tested services and their endpoints. You can drill into an endpoint's contract and select an operation to display a dedicated tab for that invocation in the pane on the right. For example, for this simple contract and implementation:

```

[ServiceContract]
interface IMyContract
{
    [OperationContract]
    string MyMethod(int someNumber,string someText);
}
class MyService : IMyContract
{
    public string MyMethod(int someNumber,string someText)
    {
        return "Hello";
    }
}

```

The method tab will let you provide an integer and a string as operation parameters in the Request section, as shown in [Figure 1-16](#).

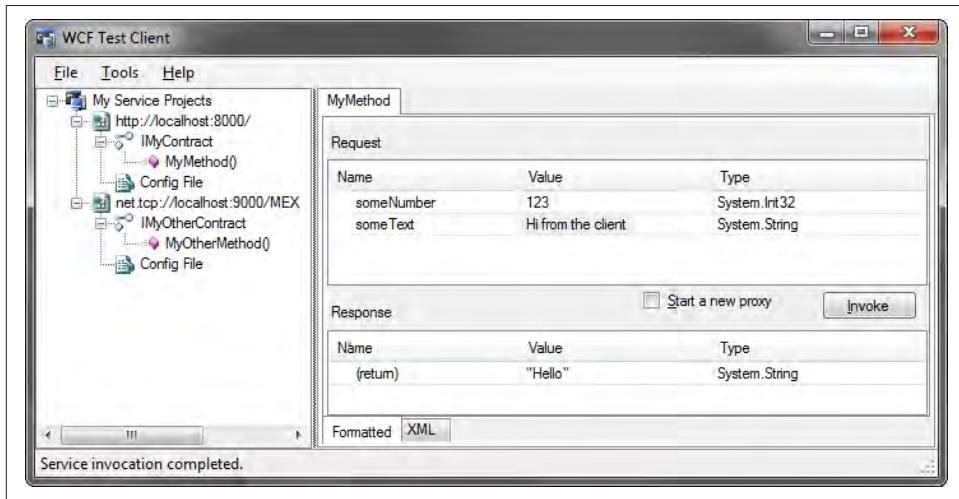


Figure 1-16. Using WcfTestClient

When you click the Invoke button, WcfTestClient will dispatch the call to the service and display the returned value or out-parameters in the Response section. In case of an exception, WcfTestClient will display the exception information in a message box and allow you to issue additional calls. All calls are made on new proxy instances. In addition, all calls are made asynchronously so that the UI is kept responsive. However, while the calls are asynchronous, WcfTestClient will let you dispatch only one call at a time.

WcfTestClient functions by silently creating an assembly from a proxy file, complete with a config file, and then loading it for use from a temporary location. If you click the Config File item in the tree on the left, you can actually grab that config file (the same config file generated when adding a service reference) and display the config file in its own tab. You can even edit the config file using SvcConfigEditor.

WcfTestClient allows you to invoke operations with enumerations, composite parameters such as classes or structures (each of which is a composite of other classes or structures), and even collections and arrays of parameters. Simply expand the items in the Request section, set their values from the drop-down lists (e.g., enum values), and invoke the call. If the operation accepts a collection or an array, you will also need to set the length. Again, the Response pane will contain any composite returned value or out-parameters.

As with WcfSvcHost (see the sidebar “[The WCF-Provided Test Host](#)” on page 55), you can integrate WcfTestClient directly into your Visual Studio solution. First, add a

class library project to the solution and delete from it all references, folders, and source files, since you have no need for those. Next, set *WcfTestClient.exe* as the external start program and provide the metadata address (or addresses) of the tested service (or services). This may be the *.svc* address of an IIS- or WAS-hosted project or, for that matter, any other metadata address of a host project, whether inside or outside your solution.

Of course, you can combine the use of *WcfTestClient* and *WcfSvcHost* in a single step to automatically host a service in a service library and test it:

```
WcfSvcHost.exe /service:MyService.dll    /config:App.config  
/client:WcfTestClient.exe /clientArg:http://localhost:9000/
```

However, with *WcfSvcHost*, specifying the metadata arguments is optional. By default, *WcfSvcHost* will pipe into the specified client application any metadata addresses it finds in the service config file. You should specify a service's metadata address explicitly only if the service (or services) does not provide its own metadata, or if you would like the test client to use different addresses. If the service config file contains multiple metadata endpoints for a given service, they will be provided in this precedence order: HTTP, TCP, IPC, HTTP-GET.

You can incorporate these steps into Visual Studio for a seamless hosting and testing experience. To do this, specify *WcfSvcHost.exe* as the startup program along with the config file and specify *WcfTestClient.exe* as the client. When you invoke *WcfTestClient* using */client*, closing the test client also terminates the host.

WCF Service Libraries

Visual Studio offers several WCF project templates. The WCF Service Library project is merely the prebuilt usage of *WcfSvcHost* and *WcfTestClient*. The main difference when using the WCF Service Library template is that there is no need to specify the config file and specify *WcfSvcHost.exe* as the startup program, because the project file will contain a new *ProjectTypeGuids* element for a WCF service library. The WCF Service Library also provides you with a simple template for a service contract and its implementation, as well as the matching config file.

Using the WCF Service Application is similar to creating a new WCF service from the new website option in Visual Studio.

The Syndication Service Library project allows you to implement an RSS feed over a WCF endpoint. It starts off with a simple service contract that returns a feed, its implementation, and a matching config file. You can host and expose your feed like any other service. The syndicated endpoints make use of the *WebHttpBinding* binding.

The Workflow Service Application template allows you to implement a service as a sequential workflow.

Programmatic Versus Administrative Configuration

The two techniques shown so far for configuring the client and the service complement each other. Administrative configuration gives you the option to change major aspects of the service and the client post-deployment, without the need to rebuild or redeploy. The major downside of administrative configuration is that it is not type-safe; you will discover certain configuration errors only at runtime and WCF configuration files often become unwieldy as a system matures.



Unlike previous versions of Visual Studio, Visual Studio 2012 and later will now validate the majority of the values within the `<serviceModel>` section of a config file. Visual Studio still does not validate configuration values that contain external references, such as the address property.

Programmatic configuration is useful when the configuration decision either is completely dynamic (i.e., when it is taken at runtime based on the current input or conditions) or is static and never changes (in which case, you might as well hardcode it). For example, if you are interested in hosting in-proc calls only, you can hardcode the use of the `NetNamedPipeBinding` and its configuration.

Configuration Policy

Generally, for each application, it's best to select a single configuration approach and stick with it. If you find you must mix different configuration approaches, you should do so consistently or things can become confusing. Defining and conveying a *configuration policy* for your application is a great way to establish a consistent configuration approach across your entire development effort. Your configuration policy should clearly define if developers may use file-based config, rely on default endpoints, hard-code programmatic defaults or dynamically apply configuration based on runtime conditions.

Your configuration policy should also address when developers should leverage various configuration mechanisms for both client and service. For example, your service configuration policy should clearly state when a `Configure()` method, programmatic `ServiceHost` usage, or even a custom `ServiceHost` implementation is appropriate. And your client configuration policy should suggest when direct channel usage, a proxy-based hierarchy, or even a custom `ChannelFactory` is necessary.

Since the options are many but the correct choices are few, you should consider folding your configuration policy into the WCF portion of your infrastructure. Doing so allows you to lower the bar of entry for your developer community, wrap best practice usage, enforce policy, and enable extensibility.

WCF Architecture

So far in this chapter, I've covered all that is required to set up and consume simple WCF services. However, as you'll see in the rest of the book, WCF offers immensely valuable support for reliability, transactions, concurrency management, security, and instance activation, all of which rely on the WCF interception-based architecture. Having the client interact with a proxy means that WCF is always present between the service and the client, intercepting the call and performing pre-call and post-call processing. The interception starts when the proxy serializes the call stack frame to a message and sends the message down a chain of *channels*. The channel is merely an interceptor whose purpose is to perform a specific task. Each client-side channel does pre-call processing of the message. The exact structure and composition of the chain depend mostly on the binding. For example, one of the channels may be responsible for encoding the message (binary, text, or MTOM), another for passing the security call context, another for propagating the client transaction, another for managing the reliable session, another for encrypting the message body (if so configured), and so on. The last channel on the client side is the transport channel, which sends the message over the configured transport to the host.

On the host side, the message goes through another chain of channels that perform host-side pre-call processing of the message. The first channel on the host side is the transport channel, which receives the message from the transport. Subsequent channels perform various tasks, such as decryption of the message body, decoding of the message, joining the propagated transaction, setting the security principal, managing the session, and activating the service instance. The last channel on the host side passes the message to the dispatcher. The dispatcher converts the message to a stack frame and calls the service instance. This sequence is depicted in [Figure 1-17](#).

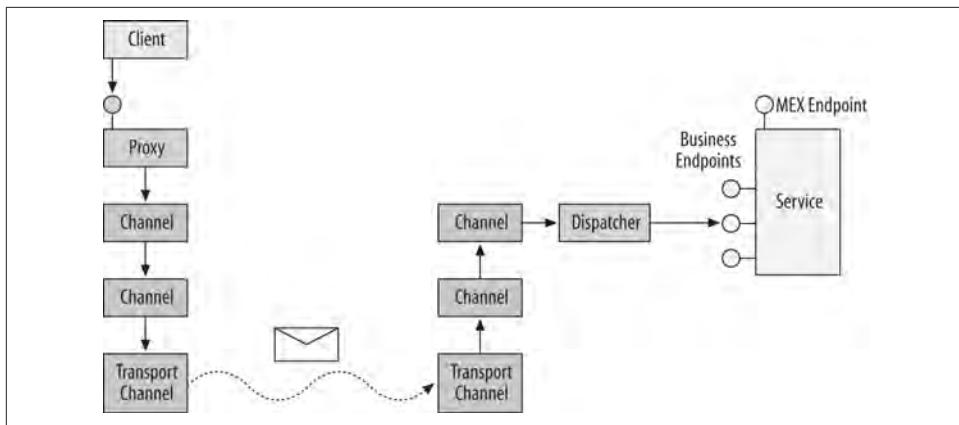


Figure 1-17. The WCF architecture

The service has no way of knowing that it was not called by a local client. In fact, it *was* called by a local client—the dispatcher. The interception both on the client and the service sides ensures that the client and the service get the runtime environments they require to operate properly.

The service instance executes the call and returns control to the dispatcher, which then converts the returned values and error information (if any) into a return message. The process is then reversed: the dispatcher passes the message through the host-side channels to perform post-call processing, such as managing the transaction, deactivating the instance, encoding the reply, encrypting it, and so on. The returned message then goes to the transport channel, which sends it to the client-side channels for client-side post-call processing. This process in turn consists of tasks such as decryption, decoding, committing or aborting the transaction, and so on. The last channel passes the message to the proxy, which converts the returned message to a stack frame and returns control to the client.

Most noteworthy is that almost all the points in the architecture provide hooks for extensibility—you can provide custom channels for proprietary interaction, custom behaviors for instance management, custom security behavior, and so on. In fact, the standard facilities that WCF offers are all implemented using the same extensibility model. You will see many examples and uses of extensibility throughout this book.

Host Architecture

It is also important to explore how the transition is made from a technology-neutral, service-oriented interaction to CLR interfaces and classes. The host performs the bridging. Each .NET host process can have many app domains, and each app domain can have zero or more service host instances. Each service host instance is dedicated to a particular service type. Thus, when you create a host instance, you are in effect

registering that service host instance with all the endpoints for that type on the host machine that correspond to its base addresses. Each service host instance has zero or more *contexts*. The context is the innermost execution scope of the service instance. A context is associated with zero or one service instance, meaning it could also be empty (i.e., not associated with any service instance). This architecture is shown in Figure 1-18.

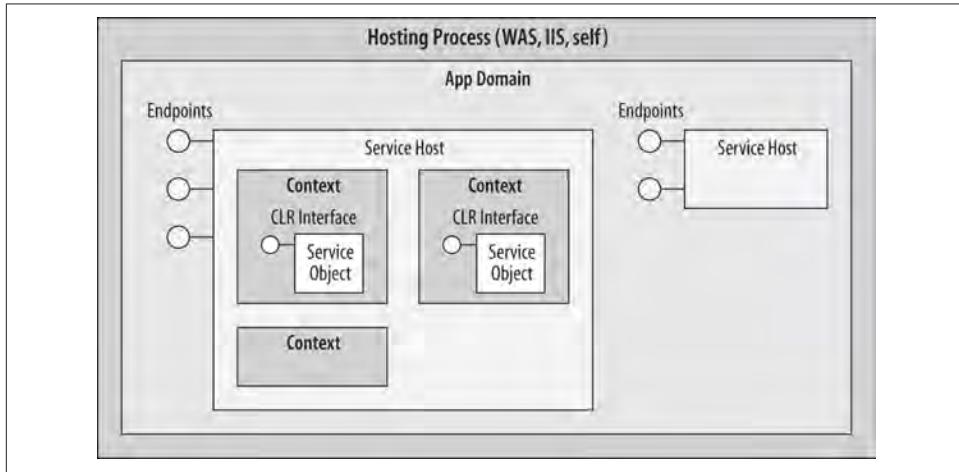


Figure 1-18. The WCF host architecture



The WCF context is conceptually similar to the Enterprise Services context or the .NET context-bound object context.

It is the combined work of the service host and the context that exposes a native CLR type as a service. After the message is passed through the channels, the host maps that message to a new or existing context (and the object instance inside) and lets it process the call.

Working with Channels

You can use channels directly to invoke operations on a service without ever resorting to using a proxy class. The `ChannelFactory<T>` class (and its supporting types), shown in Example 1-22, allows you to create a proxy on the fly.

Example 1-22. The `ChannelFactory<T>` class

```
public class ContractDescription
{
```

```

public Type ContractType
{get;set;}
//More members
}

public class ServiceEndpoint
{
    public ServiceEndpoint(ContractDescription contract,Binding binding,
        EndpointAddress address);
    public EndpointAddress Address
    {get;set;}
    public Binding Binding
    {get;set;}
    public ContractDescription Contract
    {get;}
    //More members
}

public abstract class ChannelFactory : ...
{
    public ServiceEndpoint Endpoint
    {get;}
    //More members
}
public class ChannelFactory<T> : ChannelFactory, ...
{
    public ChannelFactory(ServiceEndpoint endpoint);
    public ChannelFactory(string configurationName);
    public ChannelFactory(Binding binding,EndpointAddress endpointAddress);

    public static T CreateChannel(Binding binding,EndpointAddress endpointAddress);
    public T CreateChannel();

    //More members
}

```

You need to provide the constructor of `ChannelFactory<T>` with the endpoint. This can be the endpoint name from the client config file, the binding and address objects, or a `ServiceEndpoint` object. Next, use the `CreateChannel()` method to obtain a reference to the proxy and use its methods. Finally, close the proxy by either casting it to `IDisposable` and calling the `Dispose()` method or casting it to `ICommunicationObject` and calling the `Close()` method:

```

ChannelFactory<IMyContract> factory = new ChannelFactory<IMyContract>();

IMyContract proxy1 = factory.CreateChannel();
using(proxy1 as IDisposable)
{
    proxy1.MyMethod();
}

```

```

IMyContract proxy2 = factory.CreateChannel();
proxy2.MyMethod();
ICommunicationObject channel = proxy2 as ICommunicationObject;
Debug.Assert(channel != null);
channel.Close();

```

You can also use the shorthand static `CreateChannel()` method to create a proxy given a binding and an address without directly constructing an instance of `ChannelFactory<T>`:

```

Binding binding = new NetTcpBinding();
EndpointAddress address = new EndpointAddress("net.tcp://localhost:8000");

IMyContract proxy = ChannelFactory<IMyContract>.CreateChannel(binding,address);
using(proxy as IDisposable)
{
    proxy.MyMethod();
}

```

The InProcFactory Class

To demonstrate the power of `ChannelFactory<T>`, consider my static helper class `InProcFactory`, defined as:

```

public static class InProcFactory
{
    public static I CreateInstance<S,I>() where I : class
                                                where S : I;
    public static void CloseProxy<I>(I instance) where I : class;
    //More members
}

```

`InProcFactory` is designed to streamline and automate in-proc hosting. The `CreateInstance()` method takes two generic type parameters: the type of the service `S` and the type of the supported contract `I`. `CreateInstance()` constrains `S` to derive from `I`. Using `InProcFactory` is straightforward:

```

IMyContract proxy = InProcFactory.CreateInstance<MyService,IMyContract>();

proxy.MyMethod();

InProcFactory.CloseProxy(proxy);

```

It literally takes a service class and hoists it up as a WCF service. This is very similar to the C# `new` operator, as these two lines are equivalent in their coupling to the service type:

```

IMyContract proxy = InProcFactory.CreateInstance<MyService,IMyContract>();
IMyContract obj = new MyService();

```

In the case of C#, the compiler verifies that the type supports the requested interface and then, in effect, casts the interface into the variable. In the absence of compiler

support, `InProcFactory` requires the interface type so it will know which interface type to return.

Implementing `InProcFactory<T>`

All in-proc calls should use named pipes and should flow all transactions. You can use programmatic configuration to automate the configurations of both the client and the service, and use `ChannelFactory<T>` to avoid the need for a proxy. **Example 1-23** shows the implementation of `InProcFactory` with some of the code removed for brevity.

Example 1-23. The InProcFactory class

```
public static class InProcFactory
{
    static readonly string BaseAddress = "net.pipe://localhost/" + Guid.NewGuid();
    static readonly Binding Binding;
    static Dictionary<Type, Tuple<ServiceHost, EndpointAddress>> m_Hosts =
        new Dictionary<Type, Tuple<ServiceHost, EndpointAddress>>();
    static InProcFactory()
    {
        NetNamedPipeBinding binding = new NetNamedPipeBinding();
        binding.TransactionFlow = true;
        Binding = binding;
        AppDomain.CurrentDomain.ProcessExit += delegate
        {
            foreach(Tuple<ServiceHost, EndpointAddress>
                record in m_Hosts.Values)
            {
                record.Item1.Close();
            }
        };
    }
    public static I CreateInstance<S,I>() where I : class
        where S : I
    {
        EndpointAddress address = GetAddress<S,I>();
        return ChannelFactory<I>.CreateChannel(Binding,address);
    }
    static EndpointAddress GetAddress<S,I>() where I : class
        where S : class,I
    {
        Tuple<ServiceHost,EndpointAddress> record;

        if(m_Hosts.ContainsKey(typeof(S)))
        {
            hostRecord = m_Hosts[typeof(S)];
        }
        else
        {
            ServiceHost host = new ServiceHost(typeof(S));
            hostRecord = new Tuple<ServiceHost, EndpointAddress>(host, address);
            m_Hosts[typeof(S)] = hostRecord;
        }
    }
}
```

```

        string address = BaseAddress + Guid.NewGuid();
        record = new Tuple<ServiceHost,EndpointAddress>(
            host,new EndpointAddress(address));
        m_Hosts[typeof(S)] = record;
        host.AddServiceEndpoint(typeof(I),Binding,address);
        host.Open();
    }
    return hostRecord;
}
public static void CloseProxy<I>(I instance) where I : class
{
    ICommunicationObject proxy = instance as ICommunicationObject;
    Debug.Assert(proxy != null);
    proxy.Close();
}
}

```

`InProcFactory`'s static constructor is called once per app domain, allocating in each case a new unique base address using a GUID. This allows you to use `InProcFactory` multiple times on the same machine, across app domains and processes.

The main challenge facing `InProcFactory` is that `CreateInstance()` can be called to instantiate services of every type. For every service type, there should be a single matching host (an instance of `ServiceHost`). Allocating a host instance for each call is not a good idea. The problem is what `CreateInstance()` should do when it is asked to instantiate a second object of the same type, like so:

```

IMyContract proxy1 = InProcFactory.CreateInstance<MyService,IMyContract>();
IMyContract proxy2 = InProcFactory.CreateInstance<MyService,IMyContract>();

```

The solution is for `InProcFactory` to internally manage a dictionary that maps service types to a particular host instance and the endpoint address using a tuple. When `CreateInstance()` is called to create an instance of a particular type, it looks in the dictionary using a helper method called `GetAddress()`. If the dictionary does not already contain the service type, this helper method creates a host instance for it. If it needs to create a host, `GetAddress()` programmatically adds an endpoint to that host, using a new GUID as the unique pipe name. `GetAddress()` stores the new host and its address in the dictionary. `CreateInstance()` then uses `ChannelFactory<T>` to create the proxy. In its static constructor, which is called upon the first use of the class, `InProcFactory` subscribes to the process exit event using an anonymous method to close all hosts when the process shuts down. Finally, to help the clients close the proxy, `InProcFactory` provides the `CloseProxy()` method, which queries the proxy to `ICommunicationObject` and closes it.

The WcfWrapper

If you wish to completely approximate the C# programming model, you can wrap the in-proc factory (and thus, all of WCF) with my helper base class `WcfWrapper`, shown in [Example 1-24](#).

Example 1-24. The WcfWrapper class

```
public abstract class WcfWrapper<S,I> : IDisposable, ICommunicationObject
    where I : class
    where S : class, I
{
    protected I Proxy
    {get; private set;}

    protected WcfWrapper()
    {
        Proxy = InProcFactory.CreateInstance<S,I>();
    }

    public void Dispose()
    {
        Close();
    }

    public void Close()
    {
        InProcFactory.CloseProxy(Proxy);
    }

    void ICommunicationObject.Close()
    {
        (Proxy as ICommunicationObject).Close();
    }
    //Rest of ICommunicationObject
}
```

Using `WcfWrapper<S,I>` is simple—derive from it and the contract and implement the operations on the contract by delegating to the `Proxy` property. For example, for this service definition:

```
[ServiceContract]
interface IMyContract{
    [OperationContract]
    string MyMethod();
}

class MyService : IMyContract
{
    public string MyMethod()
```

```
    {...}  
}
```

This is the matching wrapper class:

```
class MyClass : WcfWrapper<MyService, IMyContract>, IMyContract  
{  
    public string MyMethod()  
    {  
        return Proxy.MyMethod();  
    }  
}
```

Using the wrapper class is now indistinguishable from regular C# code and yet all the calls are actually WCF calls:

```
MyClass obj = new MyClass();  
string text = obj.MyMethod();  
obj.Close();
```

[Appendix A](#) further discusses the profound implications of this programming model.

Transport-Level Sessions

In traditional programming, an object is indirectly associated with a client by virtue of the call stack. That is, each object is tied to a particular client. But in WCF, since the client sends a message to the service and never invokes the instance directly, such an association is not possible. The analogous concept in WCF is the *transport session*, which ensures that all messages coming from a particular client are sent to the same transport channel on the host. It is as if the client and the channel maintain a logical session at the transport level (hence the name). As in traditional programming, with a transport session, the calls (or rather, the messages) are strictly processed in the order in which they were received. The transport session is unrelated to any application-level session the client may or may not have with the instance itself. Note that using a transport session is optional and is largely an aspect of the binding configuration, so the client and the service may or may not have a transport session. The transport session is one of the key fundamental concepts of WCF, affecting reliability, instance management, error management, synchronization, transactions, and security.

A transport session relies on WCF's ability to identify the client and correlate all its messages to a particular channel. Thus, there has to be something in the transport or in the message that identifies the client.

Transport Session and Binding

The TCP, IPC and WebSocket bindings are connection-full. That is, all calls from the client come on the same connection or pipe, enabling WCF to easily identify the cli-

ent. However, HTTP is, by definition, a connectionless protocol, and every message from the client comes on a new connection. Consequently, when using the basic binding, there is never a transport session. Or, more precisely, there is a transport session, but it lasts for only one call and after the call returns, the channel is destroyed along with the connection. The next call will come on a new connection and will be routed to a new channel. The WS binding can improve on this situation by emulating a transport session. If configured to do so, it will insert a unique ID identifying the client in each message and will keep sending this ID for every call from that client. You will see more about this ID in [Chapter 4](#).

Transport Session Termination

Typically, the transport session will end once the client closes the proxy. However, in case the client terminates ungracefully or in case of a communication problem, each transport session also has an idle-time timeout that defaults to 10 minutes. The transport session will automatically terminate after 10 minutes of inactivity from the client, even if the client still intends to use the proxy. If the client tries to use its proxy after the transport session has been terminated due to the idle timeout, it will get a `CommunicationObjectFaultedException`. You can configure different timeouts on the client and the service by setting different values in the binding. The bindings that support a transport-level session provide the `ReliableSession` property, which can be of the type `ReliableSession` or `OptionalReliableSession`. The `ReliableSession` class offers the `InactivityTimeout` property, which you can use to configure a new idle-time timeout:

```
public class ReliableSession
{
    public TimeSpan InactivityTimeout
    {get;set;}
    //More members
}
public class OptionalReliableSession : ReliableSession
{...}
public class NetTcpBinding : Binding, ...
{
    public OptionalReliableSession ReliableSession
    {get;set;}
    //More members
}
public abstract class WSHttpBindingBase : ...
{
    public OptionalReliableSession ReliableSession
    {get;set;}
    //More members
}
public class WSHttpBinding : WSHttpBindingBase, ...
{...}
```

For example, here is the code required to programmatically configure an idle timeout of 25 minutes for the TCP binding:

```
NetTcpBinding tcpSessionBinding = new NetTcpBinding();
tcpSessionBinding.ReliableSession.InactivityTimeout = TimeSpan.FromMinutes(25);
```

Here is the equivalent configuration setting using a config file:

```
<netTcpBinding>
  <binding name = "TCPsession">
    <reliableSession inactivityTimeout = "00:25:00"/>
  </binding>
</netTcpBinding>
```

If both the client and the service configure a timeout, the shorter timeout prevails.



There is another esoteric service-side configuration for session termination: the `ServiceBehavior` attribute offers an advanced option for managing session shutdown via the `AutomaticSessionShutdown` property. This property is intended for optimizing certain callback scenarios and you can safely ignore it in most cases. In a nutshell, `AutomaticSessionShutdown` defaults to `true` so that when the client closes the proxy, the session is terminated. Setting it to `false` causes the session to continue until the service explicitly closes its sending channel. When this attribute is set to `false`, the client of a duplex session (discussed in [Chapter 5](#)) must manually close the output session on the duplex client channel; otherwise, the client will hang waiting for the session to terminate.

Reliability

WCF and other service-oriented technologies make a distinction between transport reliability and message reliability. *Transport reliability* (such as that offered by TCP/IP) offers point-to-point guaranteed delivery at the network packet level and also guarantees in-order delivery of the packets. Transport reliability is not resilient to dropped network connections or a variety of other communication problems.

Message reliability, as the name implies, deals with reliability at the message level, independent of how many packets are required to deliver the message. Message reliability provides end-to-end guaranteed delivery and order of messages, regardless of how many intermediaries are involved and how many network hops are required to deliver the message from the client to the service. Message reliability is based on an industry standard for reliable message-based communication that maintains a session at the transport level and supports retries in case of transport failures, such as dropping a wireless connection. It automatically deals with congestion, message buffering, and flow control, and can adjust the flow of messages accordingly. Message reliability

also deals with connection management, verifying connections and cleaning them up when they are no longer needed.



Message reliability does not guarantee message delivery. It provides only a guarantee that if the message does not reach its destination, the sender will know about it.

Bindings, Reliability, and Ordered Messages

In WCF, you control and configure reliability in the binding. A particular binding can support or not support reliable messaging and, if it's supported, you can enable or disable it. Whether a binding supports reliability is driven by the target scenario for that particular binding. [Table 1-2](#) summarizes the relationship between binding, reliability, and ordered delivery for the five recommended bindings and lists the respective default values.

Table 1-2. Reliability and ordered delivery

Binding name	Supports reliability	Default reliability	Supports ordered delivery	Default ordered delivery
BasicHttpBinding	No	N/A	No	N/A
NetTcpBinding	Yes	Off	Yes	On
NetNamedPipeBinding	No	N/A (On)	Yes	N/A (On)
WSHttpBinding	Yes	Off	Yes	On
NetMsmqBinding	No	N/A	No	N/A

The `BasicHttpBinding` and the `NetMsmqBinding` do not support reliability. The `BasicHttpBinding` is oriented toward the legacy ASMX web services world, which does not support reliability, while the `NetMsmqBinding` is for disconnected calls and has its own notion of reliability (discussed in [Chapter 9](#)).

Reliability is disabled by default, but you can enable it in the `NetTcpBinding`, and the `WSHttpBinding` bindings. Finally, the `NetNamedPipeBinding` is considered inherently reliable because it always has exactly one hop from the client to the service.

Message reliability also provides ordered delivery assurance, allowing execution of messages in the order in which they were sent, not the order in which they were delivered. In addition, it guarantees that each message is delivered exactly once.

WCF lets you enable reliability but not ordered delivery, in which case messages are executed in the order in which they were received. The default for all bindings that support reliability is that when reliability is enabled, ordered delivery is enabled as

well. Ordered delivery requires reliability. Thus, if ordered delivery is turned on but reliability is turned off, the calls will not be delivered in order.

Configuring Reliability

You can configure reliability (and ordered delivery) both programmatically and administratively. When you enable reliability, you must do so on both the client and the service host sides, or the client will not be able to communicate with the service. You can only configure reliability for the bindings that support it. [Example 1-25](#) shows a service-side config file that uses a `binding` configuration section to enable reliability when using the TCP binding.

Example 1-25. Enabling reliability with the TCP binding

```
<system.serviceModel>
  <services>
    <service name = "MyService">
      <endpoint
        address   = "net.tcp://localhost:8000/MyService"
        binding   = "netTcpBinding"
        bindingConfiguration = "ReliableTCP"
        contract  = "IMyContract"
      />
    </service>
  </services>
  <bindings>
    <netTcpBinding>
      <binding name = "ReliableTCP">
        <reliableSession enabled = "true"/>
      </binding>
    </netTcpBinding>
  </bindings>
</system.serviceModel>
```

When it comes to programmatic configuration, the TCP and WS bindings both offer a construction parameter and a property for configuring reliability. For example, the `NetTcpBinding` binding accepts a Boolean construction parameter for enabling reliability:

```
public class NetTcpBinding : Binding, ...
{
  public NetTcpBinding(...,bool reliableSessionEnabled);
  //More members
}
```

You can also enable reliability post-construction by accessing the `ReliableSession` property:

```

public class ReliableSession
{
    public bool Ordered
    {get;set;}
    //More members
}
public class OptionalReliableSession : ReliableSession
{
    public bool Enabled
    {get;set;}
    //More members
}
public class NetTcpBinding : Binding, ...
{
    public OptionalReliableSession ReliableSession
    {get;}
    //More members
}

```

Requiring Ordered Delivery

In theory, the service code and the contract definition should be independent of the binding used and its properties. The service should not care about the binding, and nothing in the service code pertains to the binding used. The service should be able to work with any aspect of the configured binding. In practice, however, the service implementation or the contract itself may depend on ordered delivery of the messages. To enable the contract or service developer to constrain the allowed bindings, WCF defines the `DeliveryRequirementsAttribute`:

```

[AttributeUsage(AttributeTargets.Class|AttributeTargets.Interface,
               AllowMultiple = true)]
public sealed class DeliveryRequirementsAttribute : Attribute, ...
{
    public Type TargetContract
    {get;set;}
    public bool RequireOrderedDelivery
    {get;set;}
    //More members
}

```

You can apply the `DeliveryRequirements` attribute at the service level, affecting all endpoints of the service, or only at those endpoints that expose a particular contract. When applied at the service level, requiring ordered delivery is an implementation decision. For example, to demand that all endpoints of the service, regardless of contracts, have ordered delivery enabled, apply the attribute directly on the service class:

```

[DeliveryRequirements(RequireOrderedDelivery = true)]
class MyService : IMyContract, IMyOtherContract
{...}

```

By setting the `TargetContract` property, you can demand that only endpoints of the service that support the specified contract be constrained to have reliable ordered delivery:

```
[DeliveryRequirements(TargetContract = typeof(IMyContract),
                      RequireOrderedDelivery = true)]
class MyService : IMyContract, IMyOtherContract
{...}
```

You can also use the attribute at the contract level, affecting all services that support that contract. When applied at the contract level, requiring ordered delivery is a design decision. Enforcing the constraint is done at service load time. If an endpoint has a binding that does not support reliability, supports reliability but has it disabled, or has reliability enabled but ordered delivery disabled, loading the service will fail with an `InvalidOperationException`.

By applying the `DeliveryRequirements` attribute on the contract interface, you place the constraint on all services that support it:

```
[ServiceContract]
[DeliveryRequirements(RequireOrderedDelivery = true)]
interface IMyContract
{...}

class MyService : IMyContract
{...}

class MyOtherService : IMyContract
{...}
```

The default value of `RequireOrderedDelivery` is `false`, so merely applying the attribute has no effect. For example, these statements are equivalent:

```
[ServiceContract]
interface IMyContract
{...}

[ServiceContract]
[DeliveryRequirements]
interface IMyContract
{...}

[ServiceContract]
[DeliveryRequirements(RequireOrderedDelivery = false)]
interface IMyContract
{...}
```



The IPC binding satisfies the ordered delivery constraint.

Service Contracts

The `ServiceContract` attribute presented in the previous chapter exposes a programming construct (the interface) as a service-oriented contract, allowing you to program in languages such as C#, while exposing the construct as WCF contracts and services. This chapter starts by discussing how to better bridge the gap between the two programming models by enabling operation overloading and contract inheritance. Next, it presents a few simple yet powerful service contract design and factoring guidelines and techniques. The chapter ends by showing you how to interact programmatically at runtime with the metadata of the exposed contracts.

Operation Overloading

Programming languages such as C++ and C# support *method overloading*; that is, defining two methods with the same name but with different parameters. For example, this is a valid C# interface definition:

```
interface ICalculator
{
    int Add(int arg1,int arg2);
    double Add(double arg1,double arg2);
}
```

However, operation overloading is invalid in the world of WSDL-based operations, since all operations must have unique names (they are identified by name in the messages). Consequently, while the following contract definition compiles, it will throw an `InvalidOperationException` at the service host load time:

```
//Invalid contract definition:
[ServiceContract]
interface ICalculator
{
    [OperationContract]
```

```

    int Add(int arg1,int arg2);
    [OperationContract]
    double Add(double arg1,double arg2);
}

```

However, you can manually enable operation overloading. The trick is using the `Name` property of the `OperationContract` attribute to alias the operation:

```

[AttributeUsage(AttributeTargets.Method)]
public sealed class OperationContractAttribute : Attribute
{
    public string Name
    {get;set;}
    //More members
}

```

You need to alias the operation both on the service and on the client side. On the service side, provide a unique name for each overloaded operation, as shown in [Example 2-1](#).

Example 2-1. Service-side operation overloading

```

[ServiceContract]
interface ICalculator
{
    [OperationContract(Name = "AddInt")]
    int Add(int arg1,int arg2);

    [OperationContract(Name = "AddDouble")]
    double Add(double arg1,double arg2);
}

```

When the client imports the contract and generates the proxy, the imported operations will have the aliased names:

```

[ServiceContract]
interface ICalculator
{
    [OperationContract]
    int AddInt(int arg1,int arg2);

    [OperationContract]
    double AddDouble(double arg1,double arg2);
}
class CalculatorClient : ClientBase<ICalculator>,ICalculator
{
    public int AddInt(int arg1,int arg2)
    {
        return Channel.AddInt(arg1,arg2);
    }
    public double AddDouble(double arg1,double arg2)
    {

```

```

        return Channel.AddDouble(arg1,arg2);
    }
    //Rest of the proxy
}

```

The client can use the generated proxy and contract as they are, but you can also rework them to provide overloading on the client side. Rename the methods on the imported contract and the proxy to the overloaded names, and make sure the proxy class makes calls on the internal proxy using the overloaded methods, as in:

```

public int Add(int arg1,int arg2)
{
    return Channel.Add(arg1,arg2);
}

```

Finally, use the `Name` property on the imported contract on the client side to alias and overload the methods, matching the imported operation names, as shown in [Example 2-2](#).

Example 2-2. Client-side operation overloading

```

[ServiceContract]
interface ICalculator
{
    [OperationContract(Name = "AddInt")]
    int Add(int arg1,int arg2);

    [OperationContract(Name = "AddDouble")]
    double Add(double arg1,double arg2);
}

class CalculatorClient : ClientBase<ICalculator>,ICalculator
{
    public int Add(int arg1,int arg2)
    {
        return Channel.Add(arg1,arg2);
    }
    public double Add(double arg1,double arg2)
    {
        return Channel.Add(arg1,arg2);
    }
    //Rest of the proxy
}
CalculatorClient proxy = new CalculatorClient();

int result1 = proxy.Add(1,2);
double result2 = proxy.Add(1.0,2.0);

proxy.Close();

```

Since the tools provide little value, it is always better to handcraft your proxy classes to the WCF services you own. However, proxy generation is still the best solution for third-party services where the vendor has not provided its own helper classes or proxies. While you can certainly handcraft or modify proxies for third parties to improve the programming experience, you must in this case also weigh the maintenance concerns and potential confusion that it may create for your developers because any subsequent generation of the proxy will overwrite your changes requiring you to reapply them.

Contract Inheritance

Service contract interfaces can derive from each other, enabling you to define a hierarchy of contracts. However, the `ServiceContract` attribute is not inheritable:

```
[AttributeUsage(Inherited = false,...)]
public sealed class ServiceContractAttribute : Attribute
{...}
```

Consequently, every level in the interface hierarchy must explicitly have the `ServiceContract` attribute, as shown in [Example 2-3](#).

Example 2-3. Service-side contract hierarchy

```
[ServiceContract]
interface ISimpleCalculator
{
    [OperationContract]
    int Add(int arg1,int arg2);
}
[ServiceContract]
interface IScientificCalculator : ISimpleCalculator
{
    [OperationContract]
    int Multiply(int arg1,int arg2);
}
```

When it comes to implementing a contract hierarchy, a single service class can implement the entire hierarchy, just as with classic C# programming:

```
class MyCalculator : IScientificCalculator
{
    public int Add(int arg1,int arg2)
    {
        return arg1 + arg2;
    }
    public int Multiply(int arg1,int arg2)
    {
        return arg1 * arg2;
```

```
}
```

The host can expose a single endpoint for the bottommost interface in the hierarchy:

```
<service name = "MyCalculator">
    <endpoint
        address = "http://localhost:8001/MyCalculator/"
        binding = "basicHttpBinding"
        contract = "IScientificCalculator"
    />
</service>
```

Client-Side Contract Hierarchy

When a client imports the metadata of a service endpoint whose contract is part of an interface hierarchy, the resulting contract on the client side will not maintain the original hierarchy. Instead, it will include a flattened hierarchy in the form of a single contract named after the endpoint's contract. The single contract will have a union of all the operations from all the interfaces leading down to it in the hierarchy, including itself. However, the imported interface definition will maintain, in the `Action` and `ReplyAction` properties of the `OperationContract` attribute, the name of the original contract that defined each operation:

```
[AttributeUsage(AttributeTargets.Method)]
public sealed class OperationContractAttribute : Attribute
{
    public string Action
    {get;set;}
    public string ReplyAction
    {get;set;}
    //More members
}
```

Finally, a single proxy class will implement all methods in the imported contract. Given the definitions in [Example 2-3](#), [Example 2-4](#) shows the imported contract and the generated proxy class.

Example 2-4. Client-side flattened hierarchy

```
[ServiceContract]
interface IScientificCalculator
{
    [OperationContract(Action = ".../ISimpleCalculator/Add",
                      ReplyAction = ".../ISimpleCalculator/AddResponse")]
    int Add(int arg1,int arg2);

    [OperationContract(Action = ".../IScientificCalculator/Multiply",
                      ReplyAction = ".../IScientificCalculator/MultiplyResponse")]
    int Multiply(int arg1,int arg2);
```

```

}

class ScientificCalculatorClient :
    ClientBase<IScientificCalculator>,IScientificCalculator
{
    public int Add(int arg1,int arg2)
    {...}
    public int Multiply(int arg1,int arg2)
    {...}
    //Rest of the proxy
}

```

Restoring the hierarchy on the client

The client can manually rework the proxy and the imported contract definitions to restore the contract hierarchy, as shown in [Example 2-5](#).

Example 2-5. Client-side contract hierarchy

```

[ServiceContract]
interface ISimpleCalculator
{
    [OperationContract]
    int Add(int arg1,int arg2);
}

class SimpleCalculatorClient : ClientBase<ISimpleCalculator>,ISimpleCalculator
{
    public int Add(int arg1,int arg2)
    {
        return Channel.Add(arg1,arg2);
    }
    //Rest of the proxy
}

[ServiceContract]
interface IScientificCalculator : ISimpleCalculator
{
    [OperationContract]
    int Multiply(int arg1,int arg2);
}

class ScientificCalculatorClient :
    ClientBase<IScientificCalculator>,IScientificCalculator
{
    public int Add(int arg1,int arg2)
    {
        return Channel.Add(arg1,arg2);
    }
    public int Multiply(int arg1,int arg2)
    {
        return Channel.Multiply(arg1,arg2);
    }
}

```

```
//Rest of the proxy  
}
```

Using the value of the `Action` property in the various operations, the client can factor out the definitions of the comprising contracts in the service contract hierarchy and provide interface and proxy definitions (for example, `ISimpleCalculator` and `SimpleCalculatorClient` in [Example 2-5](#)). There is no need to set the `Action` and `ReplyAction` properties, and you can safely remove them all. Next, manually add the interface to the inheritance chain as required:

```
[ServiceContract]  
interface IScientificCalculator : ISimpleCalculator  
{...}
```

Although the service may have exposed just a single endpoint for the bottommost interface in the hierarchy, the client can view it as different endpoints with the same address, where each endpoint corresponds to a different level in the contract hierarchy:

```
<client>  
    <endpoint name = "SimpleEndpoint"  
        address  = "http://localhost:8001/MyCalculator/"  
        binding   = "basicHttpBinding"  
        contract  = "ISimpleCalculator"  
    />  
    <endpoint name = "ScientificEndpoint"  
        address  = "http://localhost:8001/MyCalculator/"  
        binding   = "basicHttpBinding"  
        contract  = "IScientificCalculator"  
    />  
</client>
```

The client can now write the following code, taking full advantage of the contract hierarchy:

```
SimpleCalculatorClient proxy1 = new SimpleCalculatorClient();  
proxy1.Add(1,2);  
proxy1.Close();  
  
ScientificCalculatorClient proxy2 = new ScientificCalculatorClient();  
proxy2.Add(3,4);  
proxy2.Multiply(5,6);  
proxy2.Close();
```

The advantage of the proxy refactoring in [Example 2-5](#) is that each level in the contract is kept separately and decoupled from the levels underneath it. Anyone on the client side that expects a reference to `ISimpleCalculator` can now be given a reference to `IScientificCalculator`:

```
void UseCalculator(ISimpleCalculator calculator)  
{...}
```

```

ISimpleCalculator proxy1 = new SimpleCalculatorClient();
ISimpleCalculator proxy2 = new ScientificCalculatorClient();
IScientificCalculator proxy3 = new ScientificCalculatorClient();
SimpleCalculatorClient proxy4 = new SimpleCalculatorClient();
ScientificCalculatorClient proxy5 = new ScientificCalculatorClient();

UseCalculator(proxy1);
UseCalculator(proxy2);
UseCalculator(proxy3);
UseCalculator(proxy4);
UseCalculator(proxy5);

```

However, there is no Is-A relationship between the proxies. The **IScientificCalculator** interface derives from **ISimpleCalculator**, but a **ScientificCalculatorClient** is not a **SimpleCalculatorClient**. In addition, you have to repeat the implementation of the base contract in the proxy for the subcontract. You can rectify that by using a technique I call *proxy chaining*, illustrated in [Example 2-6](#).

Example 2-6. Proxy chaining

```

class SimpleCalculatorClient :
    ClientBase<IScientificCalculator>,ISimpleCalculator
{
    public int Add(int arg1,int arg2)
    {
        return Channel.Add(arg1,arg2);
    }
    //Rest of the proxy
}

class ScientificCalculatorClient :
    SimpleCalculatorClient,IScientificCalculator
{
    public int Multiply(int arg1,int arg2)
    {
        return Channel.Multiply(arg1,arg2);
    }
    //Rest of the proxy
}

```

Only the proxy that implements the topmost base contract derives directly from **ClientBase<T>**, providing it as a type parameter with the bottommost subinterface. All the other proxies derive from the proxy immediately above them and the respective contract.

Proxy chaining gives you an Is-A relationship between the proxies, as well as enabling code reuse. Anyone on the client side that expects a reference to **SimpleCalculatorClient** can now be given a reference to **ScientificCalculatorClient**:

```

void UseCalculator(SimpleCalculatorClient calculator)
{...}

SimpleCalculatorClient proxy1 = new SimpleCalculatorClient();
SimpleCalculatorClient proxy2 = new ScientificCalculatorClient();
ScientificCalculatorClient proxy3 = new ScientificCalculatorClient();

UseCalculator(proxy1);
UseCalculator(proxy2);
UseCalculator(proxy3);

```

Service Contract Factoring and Design

Syntax aside, how do you go about designing service contracts? How do you know which operations to allocate to which service contract? How many operations should each contract have? Answering these questions has little to do with WCF and a lot to do with abstract service-oriented analysis and design. An in-depth discussion of how to decompose a system into services and how to discover contract methods is beyond the scope of this book. Nonetheless, this section offers a few pieces of advice to guide you in your service contracts design effort.

Contract Factoring

A service contract is a grouping of logically related operations. What constitutes “logically related” is usually domain-specific. You can think of service contracts as different facets of some entity. Once you have identified (after requirements analysis) all the operations the entity supports, you need to allocate those operations to contracts. This is called *service contract factoring*. When you factor a service contract, always think in terms of reusable elements. In a service-oriented application, the basic unit of reuse is the service contract. Ask yourself, will this particular contract factoring yield contracts that other entities in the system can reuse? What facets of the entity can you logically factor out and use with other entities?

As a concrete yet simple example, suppose you wish to model a dog service. The requirements are that the dog should be able to bark and fetch, that it should have a veterinary clinic registration number, and that you should be able to vaccinate it. You could define the `IDog` service contract and have different kinds of services, such as the `PoodleService` and the `GermanShepherdService`, implement the `IDog` contract:

```

[ServiceContract]
interface IDog
{
    [OperationContract]
    void Fetch();

    [OperationContract]
    void Bark();
}

```

```

[OperationContract]
long GetVetClinicNumber();

[OperationContract]
void Vaccinate();
}

class PoodleService : IDog
{...}
class GermanShepherdService : IDog
{...}

```

However, this composition of the `IDog` service contract is not well factored. Although all the operations are things a dog should support, `Fetch()` and `Bark()` are more logically related to each other than to `GetVetClinicNumber()` and `Vaccinate()`. `Fetch()` and `Bark()` involve one facet of the dog, as a living, active canine entity, while `GetVetClinicNumber()` and `Vaccinate()` involve a different facet, one that relates it as a record of a pet in a veterinary clinic. A better approach is to factor out the `GetVetClinicNumber()` and `Vaccinate()` operations to a separate contract called `IPet`:

```

[ServiceContract]
interface IPet
{
    [OperationContract]
    long GetVetClinicNumber();

    [OperationContract]
    void Vaccinate();
}

[ServiceContract]
interface IDog
{
    [OperationContract]
    void Fetch();

    [OperationContract]
    void Bark();
}

```

Because the pet facet is independent of the canine facet, other entities (such as cats) can reuse the `IPet` service contract and support it:

```

[ServiceContract]
interface ICat
{
    [OperationContract]
    void Purr();

    [OperationContract]

```

```

        void CatchMouse();
    }

class PoodleService : IDog,IPet
{...}

class SiameseService : ICat,IPet
{...}

```

This factoring, in turn, allows you to decouple the clinic-management aspect of the application from the actual service (be it dogs or cats). Factoring operations into separate interfaces is usually done when there is a weak logical relation between the operations. However, identical operations are sometimes found in several unrelated contracts, and these operations are logically related to their respective contracts. For example, both cats and dogs need to shed fur and nurse their offspring. Logically, shedding is just as much a dog operation as barking, and just as much a cat operation as purring.

In such cases, you can factor the service contracts into a hierarchy of contracts instead of separate contracts:

```

[ServiceContract]
interface IMammal
{
    [OperationContract]
    void Shed();

    [OperationContract]
    void Lactate();
}

[ServiceContract]
interface IDog : IMammal
{...}[ServiceContract]
interface ICat : IMammal
{...}

```

Factoring Metrics

As you can see, proper contract factoring results in more specialized, loosely coupled, fine-tuned, and reusable contracts, and subsequently, those benefits apply to the system as well. In general, contract factoring results in contracts with fewer operations.

When you design a service-based system, however, you need to balance two counteracting forces (see [Figure 2-1](#)). One is the cost of implementing the service contracts, and the other is the cost of putting them together or integrating them into a cohesive application.

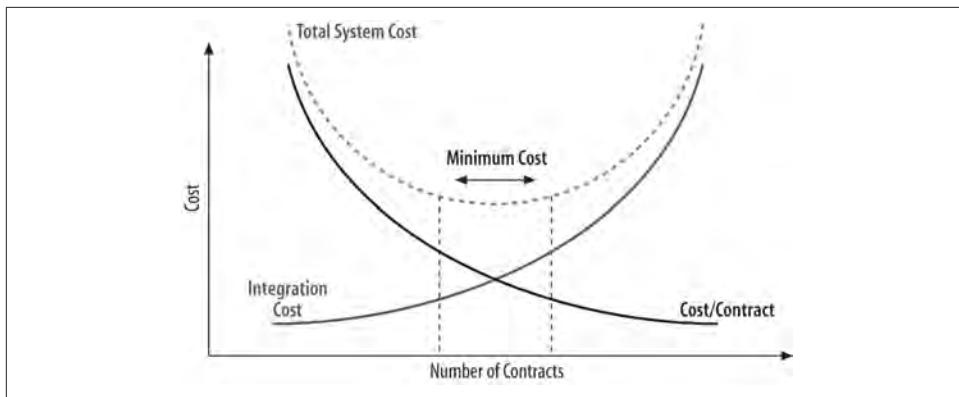


Figure 2-1. Balancing the number of services and their size

If you have too many granular service contracts, it will be easy to implement each contract, but the overall cost of integrating all those service contracts will be prohibitive. On the other hand, if you have only a few complex, large service contracts, the cost of implementing those contracts will be a prohibitive factor, even though the cost of integrating them might be low.

The relationship between size of a service contract and cost of implementation is not linear, because complexity is not linear to size—something that's twice as big may be four or six times as complex. Similarly, the relationship between integration cost and the number of service contracts to integrate is not linear, because the number of possible connections is not linear to the number of participating services.

In any given system, the total effort involved in designing and maintaining the services that implement the contracts is the sum of those two factors (cost of implementation and cost of integration). As you can see from [Figure 2-1](#), there is an area of minimum cost or effort in relation to the size and number of service contracts. A well-designed system has not too many but not too few services, and those services are not too big but not too small.

Because these contract-factoring issues are independent of the service technology used, I can extrapolate from my own and others' experiences of factoring and architecting large-scale applications and share a few rules of thumb and metrics I have collected about service-contract factoring.

Service contracts with just one operation are possible, but you should avoid them. A service contract is a facet of an entity, and that facet must be pretty dull if you can express it with just one operation. Examine that single operation, and ask yourself some questions about it. Is it using too many parameters? Is it too coarse, and therefore should it be factored into several operations? Should you factor this operation into an already existing service contract?

The optimal number of service contract members (in my opinion and experience) is between three and five. If you design a service contract with more operations—say, six to nine—you are still doing relatively well. However, look at the operations and try to determine whether any can be collapsed into each other, since it's quite possible to over-factor operations. If you have a service contract with 12 or more operations, you should definitely look for ways to factor the operations into either separate service contracts or a hierarchy of contracts. Your coding standard should set some upper limit never to be exceeded, regardless of the circumstances (say, 20).

Another rule involves the use of property-like operations, such as this:

```
[OperationContract]
long GetVetClinicNumber();
```

You should avoid such operations. Service contracts allow clients to invoke abstract operations, without caring about actual implementation details. Properties in general are better than public member variables, but methods are better than properties. I say that properties provide *just-enough encapsulation*, and this is why WCF (unlike C#) does not support properties directly—although you can easily circumvent that by defining property-like operations. Such property-like operations would encapsulate the business logic of setting and reading the variables' values on the service side. Ideally, however, you shouldn't bother clients with properties at all. That is, clients should be able to simply invoke operations and let the service worry about how to manage its state. The interaction should be in terms of *DoSomething()*, like *Vaccinate()*. How the service goes about doing that and whether or not a vet clinic number is involved should be of no concern to the client.



A word of caution about factoring metrics: rules of thumb and generic metrics are only tools to help you evaluate your particular design. There is no substitute for domain expertise and experience. Always be practical, apply judgment, and question what you do in light of these guidelines.

Contract Queries

Sometimes the client needs to programmatically verify whether a particular endpoint (identified by its address) supports a particular contract. For example, imagine an application where the application administrator (or even the end user) specifies or configures the application during setup (or at runtime) to consume and interact with a service. If the service does not support the required contracts, the application should alert the user that an invalid address was specified, and ask for an alternative or a correct address. For example, the Credentials Manager application used in [Chapter 10](#) has just such a feature: the user needs to provide it with the address of the security credentials service that manages account membership and roles. Credentials

Manager only allows the user to select a valid address, after verifying that the address supports the required service contracts.

Programmatic Metadata Processing

In order to support such functionality, the application needs to retrieve the service's metadata and see if at least one of the endpoints supports the requested contract. As explained in [Chapter 1](#), the metadata may be available either in special metadata exchange endpoints (if the service supports them), or over the HTTP-GET protocol. When you use HTTP-GET, the address of the metadata exchange is the HTTP-GET address suffixed by ?wsdl. To ease the task of parsing the returned metadata WCF offers a few helper classes, available in the `System.ServiceModel.Description` namespaces, as shown in [Example 2-7](#).

Example 2-7. Metadata processing supporting types

```
public enum MetadataExchangeClientMode
{
    MetadataExchange,
    HttpGet
}
public class MetadataSet : ...
{...}
public class ServiceEndpointCollection : Collection<ServiceEndpoint>
{...}

public class MetadataExchangeClient
{
    public MetadataExchangeClient();
    public MetadataExchangeClient(Binding mexBinding);
    public MetadataSet GetMetadata(Uri address,MetadataExchangeClientMode mode);
    //More members
}
public abstract class MetadataImporter
{
    public abstract ServiceEndpointCollection ImportAllEndpoints();
    //More members
}
public class WsdlImporter : MetadataImporter
{
    public WsdlImporter(MetadataSet metadata);
    //More members
}
public class ServiceEndpoint
{
    public EndpointAddress Address
    {get;set;}
    public Binding Binding
    {get;set;}
```

```

public ContractDescription Contract
{
    get;
    //More members
}
public class ContractDescription
{
    public static ContractDescription GetContract(Type contractType);

    public string Name
    {get;set;}
    public string Namespace
    {get;set;}
    //More members
}

```

You can provide the `MetadataExchangeClient` constructor with an already-initialized binding instance that has some custom values, such as a capacity for larger messages if the metadata returned exceeds the default received message size. The `GetMetadata()` method of `MetadataExchangeClient` accepts an endpoint address instance wrapping the metadata exchange address, as well as an enum specifying the access method. It returns the metadata in an instance of `MetadataSet`, but you should not work with that type directly. Instead, instantiate a subclass of `MetadataImporter`, such as `WsdlImporter`, and provide the raw metadata as a construction parameter; then call the `ImportAllEndpoints()` method to obtain a collection of all endpoints found in the metadata. The endpoints are represented by the `ServiceEndpoint` class.

`ServiceEndpoint` provides the `Contract` property of the type `ContractDescription`. `ContractDescription` provides the name and namespace of the contract. You can convert a CLR type (typically an interface with the `ServiceContract` attribute) to a `ContractDescription` via the `GetContract()` static method of `ContractDescription`.

Using HTTP-GET, to find out if a specified base address supports a particular contract, follow the steps just described. This will yield the collection of endpoints. Then, find out if any of the endpoints in the collection support the contract by comparing the `Name` and `Namespace` properties in the `ContractDescription` with the requested contract, as shown in [Example 2-8](#).

Example 2-8. Querying an address for a contract

```

Uri mexAddress = new Uri("...?wsdl");

MetadataExchangeClient mexClient = new MetadataExchangeClient(mexAddress,
    MetadataExchangeClientMode.HttpGet);

MetadataSet metadata = mexClient.GetMetadata();
MetadataImporter importer = new WsdlImporter(metadata);
ServiceEndpointCollection endpoints = importer.ImportAllEndpoints();

```

```

ContractDescription description =
    ContractDescription.GetContract(typeof(IMyContract));

bool contractSupported = endpoints.Any(endpoint =>
    endpoint.Contract.Namespace == description.Namespace &&
    endpoint.Contract.Name == description.Name);

```



The Metadata Explorer tool presented in [Chapter 1](#) follows steps similar to those in [Example 2-8](#) to retrieve the service endpoints. When given an HTTP-based address, the tool tries both HTTP-GET and an HTTP-based metadata exchange endpoint. The Metadata Explorer can also retrieve the metadata using a TCP- or IPC-based metadata exchange endpoint. The bulk of the implementation of the tool was in processing the metadata and rendering it because the difficult task of retrieving and parsing the metadata is done by the WCF-provided classes.

The MetadataResolver class

[Example 2-8](#) queries the metadata in two phases: `MetadataExchangeClient` is used to obtain the metadata, and `MetadataImporter` is used to parse it and extract the endpoints. WCF combines these two steps for you with the `MetadataResolver` static class:

```

public static class MetadataResolver
{
    public static ServiceEndpointCollection Resolve(Type contract,
                                                    EndpointAddress address);
    public static ServiceEndpointCollection Resolve(Type contract, Uri address,
                                                    MetadataExchangeClientMode mode);
    //Additional members
}

```

Here is the same querying logic as that in [Example 2-8](#), using the `MetadataResolver`:

```

bool contractSupported = false;

Uri mexAddress = new Uri("...?wsdl");

ServiceEndpointCollection endpoints = MetadataResolver.Resolve
    (typeof(IMyContract), mexAddress,
     MetadataExchangeClientMode.HttpGet);

if(endpoints.Count > 0)
{
    contractSupported = true;
}

```

The MetadataHelper Class

While the WCF-provided `MetadataResolver` is a step in the right direction, I wanted to streamline obtaining metadata further and encapsulate advanced yet sometimes necessary steps such as setting the metadata message size. To this end, I created a general-purpose static utility class called `MetadataHelper`, which offers methods such as `QueryContract()`:

```
public static class MetadataHelper
{
    public static bool QueryContract(string mexAddress, Type contractType);
    //More members
}
```

You can provide `MetadataHelper` with either the `Type` of the contract you wish to query for, or the name and namespace of the contract:

```
string address = "...";
bool contractSupported = MetadataHelper.QueryContract(address, typeof
    (IMyContract));
```

For a metadata exchange address, you can provide `MetadataHelper` with an HTTP-GET address, or a MEX endpoint address over HTTP, HTTPS, TCP, or IPC. **Example 2-9** shows the implementation of `MetadataHelper.QueryContract()`, with some of the error-handling code removed.

Example 2-9. Implementing `MetadataHelper.QueryContract()`

```
public static class MetadataHelper
{
    const int MessageSizeMultiplier = 5;

    static ServiceEndpointCollection QueryMexEndpoint(string mexAddress,
                                                    BindingElement bindingElement)
    {
        dynamic element = bindingElement;
        element.MaxReceivedMessageSize *= MessageSizeMultiplier;

        CustomBinding binding = new CustomBinding(element);

        MetadataExchangeClient mexClient = new MetadataExchangeClient(binding);
        MetadataSet metadata = mexClient.
            GetMetadata(new EndpointAddress(mexAddress));
        MetadataImporter importer = new WsdlImporter(metadata);
        return importer.ImportAllEndpoints();
    }

    public static ServiceEndpoint[] GetEndpoints(string mexAddress)
    {
        /* Some error handling */
    }
}
```

```

Uri address = new Uri(mexAddress);
ServiceEndpointCollection endpoints = null;
BindingElement bindingElement = null;

if(address.Scheme == Uri.UriSchemeNetTcp)
{
    bindingElement = new TcpTransportBindingElement();
}
if(address.Scheme == Uri.UriSchemeNetPipe)
{...}
if(address.Scheme == Uri.UriSchemeHttp) //Checks for HTTP-GET as well
{...}
if(address.Scheme == Uri.UriSchemeHttps) //Checks for HTTPS-GET as well
{...}

endpoints = QueryMexEndpoint(mexAddress,bindingElement);

return endpoints.ToArray();
}
public static bool QueryContract(string mexAddress,Type contractType)
{

//Check mex address and contract namespace.
if(string.IsNullOrWhiteSpace(mexAddress))
{
    throw new ArgumentException("Invalid MEX address");
}

//Check contract type.
if(contractType == null)
{
    throw new ArgumentException("Invalid contract type");
}

if(contractType.IsInterface == false)
{
    Debug.Assert(false,contractType + " is not an interface");
    return false;
}
object[] attributes = contractType.GetCustomAttributes(
                           typeof(ServiceContractAttribute),false);
if(attributes.Length == 0)
{
    Debug.Assert(false,"Interface " + contractType +
                  " does not have the ServiceContractAttribute");
    return false;
}
ContractDescription description = ContractDescription.
                                         GetContract(contractType);
return QueryContract(mexAddress,description.Namespace,description.Name);
}
public static bool QueryContract(string mexAddress,string contractNamespace,

```

```

                string contractName)
{
    // Check mex address and contract namespace.
    if(string.IsNullOrWhiteSpace(mexAddress))
    {
        throw new ArgumentException("Invalid MEX address");
    }

    if(string.IsNullOrWhiteSpace(contractNamespace))
    {
        throw new ArgumentException("Invalid contract namespace");
    }

    if(string.IsNullOrWhiteSpace(contractName))
    {
        throw new ArgumentException("Invalid contract name");
    }

    ServiceEndpoint[] endpoints = GetEndpoints(mexAddress);

    return endpoints.Any(endpoint =>
        endpoint.Contract.Namespace == contractNamespace &&
        endpoint.Contract.Name == contractName);
}
}

```

In [Example 2-9](#), the `GetEndpoints()` method parses out the scheme of the metadata exchange address. According to the transport scheme found (e.g., TCP), `GetEndpoints()` constructs a matching binding element. `GetEndpoints()` then uses the `QueryMexEndpoint()` private method to actually retrieve the metadata. `QueryMexEndpoint()` accepts the metadata exchange endpoint address and the binding element to use. `QueryMexEndpoint()` must adjust the received message size because the maximum received message size defaults to 64K. While this is adequate for simple services, services that have many endpoints and use complex types will generate larger messages, causing the call to `MetadataExchangeClient.GetMetadata()` to fail. My experimentations indicate that 5 is an adequate boost factor for most cases. You control the maximum received message size via the `MaxReceivedMessageSize` property of each specific binding element. For example, in the case of the TCP binding element:

```

public abstract class TransportBindingElement : BindingElement
{
    public virtual long MaxReceivedMessageSize
    {get;set;}
}

public class TcpTransportBindingElement :
    ConnectionOrientedTransportBindingElement
{...}

```

```

public abstract class ConnectionOrientedTransportBindingElement :
    TransportBindingElement,...  

{...}

```

Since `QueryMexEndpoint()` uses the abstract base class `BindingElement`, which does not have a `MaxReceivedMessageSize` property, `QueryMexEndpoint()` relies on the `dynamic` directive to instruct the compiler to set the property on the derived element. `QueryMexEndpoint()` then wraps the binding element with a custom binding and provides it to an instance of `MetadataExchangeClient`, which retrieves the metadata and returns the endpoint collection. Instead of returning a `ServiceEndpointCollection`, `GetEndpoints()` uses LINQ to return an array of endpoints.

The `QueryContract()` method that accepts a `Type` first verifies that the type is an interface and that it is decorated with the `ServiceContract` attribute. Because the `ServiceContract` attribute can be used to alias both the name and namespace of the requested type of contract, `QueryContract()` uses `ContractDescription` instead of the raw type and calls the `QueryContract()` method that operates on the name and namespace. That version of `QueryContract()` calls `GetEndpoints()` to obtain the array of endpoints and returns true if it finds at least one endpoint that supports the contract.

Example 2-10 shows additional metadata querying methods offered by `MetadataHelper`.

Example 2-10. The MetadataHelper class

```

public static class MetadataHelper  

{  

    public static ServiceEndpoint[] GetEndpoints(string mexAddress);  

    public static ServiceEndpoint[] GetEndpoints(string mexAddress,  

                                                Type contractType);  

    public static string[] GetAddresses(Type bindingType, string mexAddress,  

                                       Type contractType);  

    public static string[] GetAddresses(string mexAddress, Type contractType);  

    public static string[] GetAddresses(Type bindingType, string mexAddress,  

                                       string contractNamespace,  

                                       string contractName);  

    public static string[] GetAddresses(string mexAddress, string contractNamespace,  

                                       string contractName);  

    public static ContractDescription[] GetContracts(Type bindingType,  

                                                    string mexAddress);  

    public static ContractDescription[] GetContracts(string mexAddress);  

    public static string[] GetOperations(string mexAddress, Type contractType);  

    public static string[] GetOperations(string mexAddress,  

                                       string contractNamespace,  

                                       string contractName);  

    public static bool QueryContract(string mexAddress, Type contractType);  

}

```

```
public static bool QueryContract(string mexAddress,
                                string contractNamespace, string contractName);
//More members
}
```

These powerful and useful features are often required during setup or in administration applications and tools, and yet their implementation is all based on processing the array of endpoints returned from the `GetEndpoints()` method.

The `GetAddresses()` methods return either all the endpoint addresses that support a particular contract, or only the addresses of those endpoints that also use a particular binding. Similarly, `GetContracts()` returns all the contracts supported across all endpoints, or the contracts supported across all endpoints that use a particular binding. Finally, `GetOperations()` returns all the operations on a particular contract.



[Chapter 10](#) uses the `MetadataHelper` class in the Credentials Manager application, [Appendix C](#) uses it in conjunction with discovery, and [Appendix D](#) uses it for administering persistent subscribers.

Data Contracts

WCF provides the ability to host and expose native CLR types (interfaces and classes) as services, as well as the ability to consume services as native CLR interfaces and classes. WCF service operations accept and return CLR types such as integers and strings, and the WCF clients pass in and process returned CLR types. However, such CLR types are specific to .NET. This poses a problem because one of the core tenets of service orientation is that services do not betray their implementation technologies across the service boundary. As a result, any client, regardless of its own technology, can interact with any service. This, of course, means that WCF cannot allow you to expose the CLR data types across the service boundary. What you need is a way of converting CLR types to and from a standard neutral representation. The service needs a formal way of declaring how the conversion is to take place. This formal specification is called a *data contract*, and it is the subject of this chapter. The first part of the chapter shows how data contracts enable type marshaling and conversions, and how the infrastructure deals with class hierarchies and data contract versioning. The second part shows how to use various .NET types, such as enumerations, delegates, data tables, and collections, as data contracts.

Serialization

The data contract is part of the contractual obligation the service supports, just like the service operations are part of that contract. The data contract is published in the service's metadata, allowing clients to convert the neutral, technology-agnostic representation of the data types to their native representations. Because objects and local references are CLR concepts, you cannot pass CLR objects and references to and from a WCF service operation. Allowing you to do so not only would violate the core service-oriented principle discussed previously, but also would be impractical, since an object is comprised of both its state and the code manipulating it. There is no way

of sending the code or the logic as part of a C# or Visual Basic method invocation, let alone marshaling it to another platform and technology. In fact, when passing an object (or a value type) as an operation parameter, all you really need to send is the state of that object, and you let the receiving side convert it back to its own native representation. Such an approach for passing state around is called *marshaling by value*. The easiest way to perform marshaling by value is to rely on the built-in support most platforms (.NET included) offer for serialization. The approach is simple enough, as shown in [Figure 3-1](#).

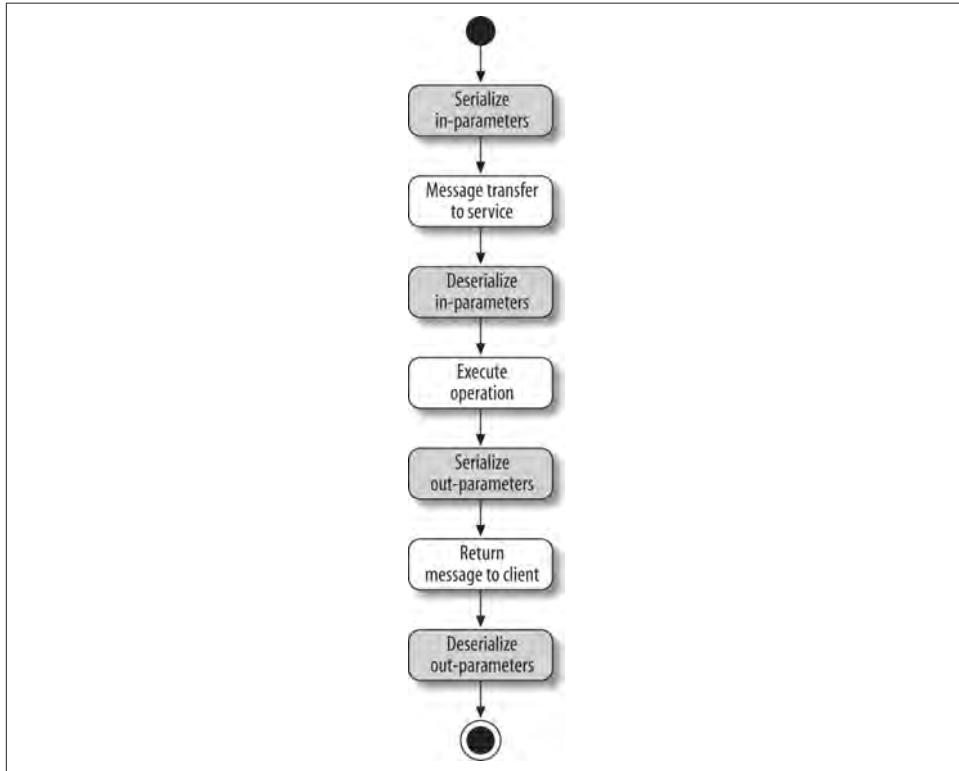


Figure 3-1. Serialization and deserialization during an operation call

On the client side, WCF will serialize the in-parameters from the CLR native representation to a neutral representation of the data contract and bundle them in the outgoing message to the client. Once the message is received on the service side, WCF will deserialize it and convert the neutral representation of the data contract to the corresponding CLR representation before dispatching the call to the service. The service will then process the native CLR parameters. Once the service has finished executing the operation, WCF will serialize the out-parameters and the returned values into a neutral representation of the data contract, package them in the returned mes-

sage, and post the returned message to the client. Finally, on the client side, WCF will deserialize the returned values into native CLR types and return them to the client.



The double dose of serialization and deserialization in every call is the real bottleneck of WCF, performance-wise. The cost of running a message through the interceptors on the client and service sides is minuscule compared with the overhead of serialization.

.NET Serialization

WCF can make use of .NET's ready-made support for serialization. .NET automatically serializes and deserializes objects using reflection. .NET captures the value of each of the object's fields and serializes it to memory, to a file, or to a network connection. For deserializing, .NET creates a new object of the matching type, reads its persisted field values, and sets the values of its fields using reflection. Because reflection can access private fields, .NET takes a complete snapshot of the state of an object during serialization and perfectly reconstructs that state during deserialization. .NET serializes the object state into a *stream*, which is a sequence of bytes, independent of a particular medium such as a file, memory, a communication port, or any other resource.

The Serializable attribute

By default, user-defined types (classes and structs) are not serializable. The reason is that .NET has no way of knowing whether a reflection-based dump of the object state to a stream makes sense. Perhaps the object members have some transient value or state (such as an open database connection or communication port). If .NET simply serialized the state of such an object after constructing a new object by deserializing it from the stream, you could end up with a defective object. Consequently, serialization has to be performed by consent of the class's developer.

To indicate to .NET that instances of your class are serializable, add the `SerializableAttribute` to your class or struct definition:

```
[AttributeUsage(AttributeTargets.Delegate |
                 AttributeTargets.Enum      |
                 AttributeTargets.Struct    |
                 AttributeTargets.Class,
                 Inherited = false)]
public sealed class SerializableAttribute : Attribute
{}
```

For example:

```
[Serializable]
class MyClass
{...}
```

The NonSerialized attribute

When a class is serializable, .NET insists that all its member variables be serializable as well, and if it discovers a non-serializable member, it throws an exception. However, what if the class or a struct that you want to serialize has a member that cannot be serialized? That type will not have the `Serializable` attribute and will preclude the containing type from being serialized. Commonly, that non-serializable member is a reference type requiring some special initialization. The solution to this problem requires marking such a member as non-serializable and taking a custom step to initialize it during deserialization.

To allow a serializable type to contain a non-serializable type as a member variable, you need to mark the member with the `NonSerialized` field attribute. For example:

```
class MyOtherClass
{...}

[Serializable]
class MyClass
{
    [NonSerialized]
    MyOtherClass m_OtherClass;
    /* Methods and properties */
}
```

When .NET serializes a member variable, it first reflects it to see whether it has the `NonSerialized` attribute. If so, .NET ignores that variable and simply skips over it.

You can even use this technique to exclude from serialization normally serializable types, such as `string`:

```
[Serializable]
class MyClass
{
    [NonSerialized]
    string m_Name;
}
```

The .NET formatters

.NET offers two formatters for serializing and deserializing types. The `BinaryFormatter` serializes into a compact binary format, enabling fast serialization and deserialization. The `SoapFormatter` uses a .NET-specific SOAP XML format.

Both formatters support the `IFormatter` interface, defined as:

```
public interface IFormatter
{
    object Deserialize(Stream serializationStream);
    void Serialize(Stream serializationStream, object graph);
    // More members
```

```

}

public sealed class BinaryFormatter : IFormatter, ...
{...}
public sealed class SoapFormatter : IFormatter, ...
{...}

```

In addition to the state of the object, both formatters persist the type's assembly and versioning information to the stream so that they can deserialize it back to the correct type. This renders them inadequate for service-oriented interaction, however, because it requires the other party not only to have the type assembly, but also to be using .NET. The use of the Stream is also an imposition, because it requires the client and the service to somehow share the stream.

The WCF Formatters

Due to the deficiencies of the classic .NET formatters, WCF has to provide its own service-oriented formatter. The formatter, `DataContractSerializer`, is capable of sharing just the data contract, not the underlying type information. `DataContractSerializer` is defined in the `System.Runtime.Serialization` namespace and is partially listed in [Example 3-1](#).

Example 3-1. The DataContractSerializer formatter

```

public abstract class XmlObjectSerializer
{
    public virtual object ReadObject(Stream stream);
    public virtual object ReadObject(XmlReader reader);
    public virtual void WriteObject(XmlWriter writer,object graph);
    public void WriteObject(Stream stream,object graph);
    //More members
}
public sealed class DataContractSerializer : XmlObjectSerializer
{
    public DataContractSerializer(Type type);
    //More members
}

```

`DataContractSerializer` captures only the state of the object according to the serialization or data contract schema. Note that `DataContractSerializer` does not support `IFormatter`.

WCF uses `DataContractSerializer` automatically under the covers, and developers should never need to interact with it directly. However, you can use `DataContractSerializer` to serialize types to and from a .NET stream, similar to using the legacy formatters. Unlike when using the binary or SOAP formatters, however, you need to

supply the `DataContractSerializer` constructor with the type to operate on, because no type information will be present in the stream:

```
MyClass obj1 = new MyClass();
DataContractSerializer formatter = new DataContractSerializer(typeof(MyClass));

using(Stream stream = new MemoryStream())
{
    formatter.WriteObject(stream,obj1);
    stream.Position = 0;
    MyClass obj2 = (MyClass)formatter.ReadObject(stream);
}
```

While you can use `DataContractSerializer` with .NET streams, you can also use it in conjunction with XML readers and writers when the only form of input is the raw XML itself, as opposed to some medium such as a file or memory.

Note the use of the amorphous `object` in the definition of `DataContractSerializer` in [Example 3-1](#). This means that there will be no compile-time-type safety, because the constructor can accept one type, the `WriteObject()` method can accept a second type, and the `ReadObject()` method can cast to yet a third type.

To compensate for that, you can define your own generic wrapper around `DataContractSerializer`, as shown in [Example 3-2](#).

Example 3-2. The generic DataContractSerializer<T>

```
public class DataContractSerializer<T> : XmlObjectSerializer
{
    DataContractSerializer m_DataContractSerializer;

    public DataContractSerializer()
    {
        m_DataContractSerializer = new DataContractSerializer(typeof(T));
    }
    public new T ReadObject(Stream stream)
    {
        return (T)m_DataContractSerializer.ReadObject(stream);
    }
    public new T ReadObject(XmlReader reader)
    {
        return (T)m_DataContractSerializer.ReadObject(reader);
    }
    public void WriteObject(Stream stream,T graph)
    {
        m_DataContractSerializer.WriteObject(stream,graph);
    }
    public void WriteObject(XmlWriter writer,T graph)
    {
        m_DataContractSerializer.WriteObject(writer,graph);
    }
}
```

```
    }
    //More members
}
```

The generic class `DataContractSerializer<T>` is much safer to use than the object-based `DataContractSerializer`:

```
MyClass obj1 = new MyClass();
DataContractSerializer<MyClass>
formatter = new DataContractSerializer<MyClass>();
using(Stream stream = new MemoryStream())
{
    formatter.WriteObject(stream,obj1);
    stream.Position = 0;
    MyClass obj2 = formatter.ReadObject(stream);
}
```

WCF also offers the `NetDataContractSerializer` formatter, which is polymorphic with `IFormatter`:

```
public sealed class NetDataContractSerializer : IFormatter, ...
{...}
```

As its name implies, similar to the legacy .NET formatters, the `NetDataContractSerializer` formatter captures the type information in addition to the state of the object. It is used just like the legacy formatters:

```
MyClass obj1 = new MyClass();
IFormatter formatter = new NetDataContractSerializer();

using(Stream stream = new MemoryStream())
{
    formatter.Serialize(stream,obj1);
    stream.Position = 0;
    MyClass obj2 = (MyClass)formatter.Deserialize(stream);
}
```

`NetDataContractSerializer` is designed to complement `DataContractSerializer`. You can serialize a type using `NetDataContractSerializer` and deserialize it using `DataContractSerializer`:

```
MyClass obj1 = new MyClass();
IFormatter formatter1 = new NetDataContractSerializer();

using(Stream stream = new MemoryStream())
{
    formatter1.Serialize(stream,obj1);

    stream.Position = 0;

    DataContractSerializer formatter2 = new DataContractSerializer
        (typeof(MyClass));
```

```
    MyClass obj2 = (MyClass)formatter2.ReadObject(stream);
}
```

This capability opens the way for versioning tolerance and for migrating legacy code that shares type information into a more service-oriented approach where only the data schema is maintained.

Data Contract via Serialization

When a service operation accepts or returns any type or parameter, WCF uses `DataContractSerializer` to serialize and deserialize that parameter. This means you can pass any serializable type as a parameter or returned value from a contract operation, as long as the other party has the definition of the data schema or the data contract. All the .NET built-in primitive types are serializable. For example, here are the definitions of the `int` and `string` types:

```
[Serializable]
public struct Int32 : ...
{...}

[Serializable]
public sealed class String : ...
{...}
```

This is the only reason why any of the service contracts shown in the previous chapters actually worked: WCF offers *implicit data contracts* for the primitive types because there is an industry standard for the schemas of those types.

To use a custom type as an operation parameter, there are two requirements: first, the type must be serializable, and second, both the client and the service need to have a local definition of that type that results in the same data schema.

Consider the `IContactManager` service contract used to manage a contacts list:

```
[Serializable]
struct Contact
{
    public string FirstName;
    public string LastName;
}

[ServiceContract]
interface IContactManager
{
    [OperationContract]
    void AddContact(Contact contact);

    [OperationContract]
    Contact[] GetContacts();
}
```

If the client uses an equivalent definition of the `Contact` structure, it can pass a contact to the service. An equivalent definition might be anything that results in the same data schema for serialization. For example, the client might use this definition instead:

```
[Serializable]
struct Contact
{
    public string FirstName;
    public string LastName;
    [NonSerialized]
    public string Address;
}
```

Data Contract Attributes

Although using the `Serializable` attribute is workable, it is not ideal for service-oriented interaction between clients and services. Rather than denoting all members in a type as serializable and therefore part of the data schema for that type, it would be preferable to have an opt-in approach, where only members the contract developer wants to explicitly include in the data contract are included. The `Serializable` attribute forces the data type to be serializable in order to be used as a parameter in a contract operation, and it does not offer clean separation between the ability to use the type as a WCF operation parameter (the “serviceness” aspect of the type) and the ability to serialize it. The attribute offers no support for aliasing type names or members, or for mapping a new type to a predefined data contract. The attribute operates directly on member fields and completely bypasses any logical properties used to access those fields. It would be better to allow those properties to add their values when accessing the fields. Finally, there is no direct support for versioning, because the formatter supposedly captures all versioning information. Consequently, it is difficult to deal with versioning over time.

Yet again, the WCF solution is to come up with new, service-oriented opt-in attributes. The first of these attributes is the `DataContractAttribute`, defined in the `System.Runtime.Serialization` namespace:

```
[AttributeUsage(AttributeTargets.Enum |
                AttributeTargets.Struct |
                AttributeTargets.Class,
                Inherited = false,
                AllowMultiple = false)]
public sealed class DataContractAttribute : Attribute
{
    public string Name
    {get;set;}
    public string Namespace
    {get;set;}}
```

```
        //More members  
    }
```

Applying the `DataContract` attribute on a class or struct does not cause WCF to serialize any of its members:

```
[DataContract]  
struct Contact  
{  
    //Will not be part of the data contract  
    public string FirstName;  
    public string LastName;  
}
```

All the `DataContract` attribute does is opt-in the type, indicating that the type can be marshaled by value. To serialize any of its members, you must apply the `DataMember` Attribute, defined as:

```
[AttributeUsage(AttributeTargets.Field|AttributeTargets.Property,  
                 Inherited = false,AllowMultiple = false)]  
public sealed class DataMemberAttribute : Attribute  
{  
    public bool IsRequired  
    {get;set;}  
    public string Name  
    {get;set;}  
    public int Order  
    {get;set;}  
  
    //More members  
}
```

You can apply the `DataMember` attribute on the fields directly:

```
[DataContract]  
struct Contact  
{  
    [DataMember]  
    public string FirstName;  
  
    [DataMember]  
    public string LastName;  
}
```

You can also apply the `DataMember` attribute on properties (either explicit properties, where you provide the property implementation, or automatic properties, where the compiler generates the underlying member and access implementation):

```
[DataContract]  
struct Contact  
{  
    string m_FirstName;
```

```

[DataMember]
public string FirstName
{
    get
    {
        return m_FirstName;
    }
    set
    {
        m_FirstName = value;
    }
}

[DataMember]
public string LastName
{get;set;}
}

```

As with service contracts, the visibility of the data members and the data contract itself is of no consequence to WCF. Thus, you can include internal types with private data members in the data contract:

```

[DataContract]
struct Contact
{
    [DataMember]
    string FirstName
    {get;set;}

    [DataMember]
    string LastName
    {get;set;}
}

```

Some of the code in this chapter applies the `DataMember` attribute directly on public data members for brevity's sake. In real code, you should, of course, use properties instead of public members.



Data contracts are case-sensitive both at the type and member levels.

Importing a Data Contract

When a data contract is used in a contract operation, it is published in the service metadata. When the client uses a tool such as Visual Studio to import the definition of the data contract, the client will end up with an equivalent definition, but not nec-

essarily an identical one. The difference is a function of the tool, not the published metadata. With Visual Studio, the imported definition will maintain the original type designation of a class or a structure as well as the original type namespace, but with SvcUtil, only the data contract will maintain the namespace. Take, for example, the following service-side definition:

```
namespace MyNamespace
{
    [DataContract]
    struct Contact
    {...}

    [ServiceContract]
    interface IContactManager
    {
        [OperationContract]
        void AddContact(Contact contact);

        [OperationContract]
        Contact[] GetContacts();
    }
}
```

The imported definition will be:

```
namespace MyNamespace
{
    [DataContract]
    struct Contact
    {...}
}

[ServiceContract]
interface IContactManager
{
    [OperationContract]
    void AddContact(Contact contact);

    [OperationContract]
    Contact[] GetContacts();
}
```

To override this default and provide an alternative namespace for the data contract, you can assign a value to the `Namespace` property of the `DataContract` attribute. The tools treat the provided namespace differently. Given this service-side definition:

```
namespace MyNamespace
{
    [DataContract(Namespace = "MyOtherNamespace")]
    struct Contact
    {...}
}
```

Visual Studio imports it exactly as defined, while SvcUtil imports it as published:

```
namespace MyOtherNamespace
{
    [DataContract]
    struct Contact
    {...}
}
```

When using Visual Studio, the imported definition will always have properties decorated with the `DataMember` attribute, even if the original type on the service side did not define any properties. For example, for this service-side definition:

```
[DataContract]
struct Contact
{
    [DataMember]
    public string FirstName;

    [DataMember]
    public string LastName;
}
```

The imported client-side definition will be:

```
[DataContract]
public partial struct Contact
{
    string FirstNameField;
    string LastNameField;

    [DataMember]
    public string FirstName
    {
        get
        {
            return FirstNameField;
        }
        set
        {
            FirstNameField = value;
        }
    }

    [DataMember]
    public string LastName
    {
        get
        {
            return LastNameField;
        }
        set
        {
```

```
        LastNameField = value;
    }
}
}
```

The client can, of course, manually rework any imported definition to be just like a service-side definition.



Even if the `DataMember` attribute on the service side is applied on a private field or property, as shown here:

```
[DataContract]
struct Contact
{
    [DataMember]
    string FirstName
    {get;set;}

    [DataMember]
    string LastName;
}
```

the imported definition will have a public property instead.

When the `DataMember` attribute is applied on a property (on either the service or the client side), that property must have `get` and `set` accessors. Without them, you will get an `InvalidOperationException` at call time. The reason is that when the property itself is the data member, WCF uses the property during serialization and deserialization, letting you apply any custom logic in the property.



Do not apply the `DataMember` attribute on both a property and its underlying field—this will result in duplication of the members on the importing side.

It is important to realize that the method just described for utilizing the `DataMember` attribute applies to both the service and the client side. When the client uses the `DataMember` attribute (and its related attributes, described elsewhere in this chapter), it affects the data contract it is using to either serialize and send parameters to the service or deserialize and use the values returned from the service. It is quite possible for the two parties to use equivalent yet not identical data contracts, and as you will see later, even to use nonequivalent data contracts. The client controls and configures its data contract independently of the service.

Data Contracts and the Serializable Attribute

The service can still use a type that is only marked with the `Serializable` attribute:

```
[Serializable]
struct Contact
{
    string m_FirstName;
    public string LastName;
}
```

When importing the metadata of such a type, the imported definition will use the `DataContract` attribute. In addition, since the `Serializable` attribute affects only fields, it will be as if every serializable member (whether public or private) is a data member, resulting in a set of wrapping properties named exactly like the original fields:

```
[DataContract]
public partial struct Contact
{
    string LastNameField;
    string m_FirstNameField;

    [DataMember(...)]
    public string LastName
    {
        ... //Accesses LastNameField
    }
    [DataMember(...)]
    public string m_FirstName
    {
        ... //Accesses m_FirstNameField
    }
}
```

The client can also use the `Serializable` attribute on its data contract to have it marshaled in much the same way.



.NET's legacy formatters cannot serialize a type marked only with the `DataContract` attribute. If you want to serialize such a type, you must apply both the `DataContract` attribute and the `Serializable` attribute on it. In the resulting data contract for the type, the effect will be the same as if only the `DataContract` attribute had been applied, and you will still need to use the `DataMember` attribute on the members you want to serialize.

Data Contracts and XML Serialization

.NET offers yet another serialization mechanism: raw XML serialization, using a dedicated set of attributes. When you're dealing with a data type that requires explicit control over XML serialization, you can use the `XmlSerializerFormatAttribute` on individual operations in the contract definition to instruct WCF to use XML serialization at runtime. If all the operations in the contract require this form of serialization, you can use the `/serializer:XmlSerializer` switch of `SvcUtil` (described in [Chapter 1](#)) to instruct it to automatically apply the `XmlSerializerFormat` attribute on all operations in all imported contracts. Use this switch with caution, though, because it will affect all data contracts, including those that do not require explicit control over XML serialization.

Inferred Data Contracts

WCF provides support for *inferred data contracts*. If the marshaled type is a public type and it is not decorated with the `DataContract` attribute, WCF will automatically infer such an attribute and apply the `DataMember` attribute to all public members (fields or properties) of the type.

For example, given this service contract definition:

```
public struct Contact
{
    public string FirstName
    {get;set;}

    public string LastName;

    internal string PhoneNumber;

    string Address;
}
[ServiceContract]
interface IContactManager
{
    [OperationContract]
    void AddContact(Contact contact);
```

```
    } ...
```

WCF will infer a data contract, as if the service contract developer had defined it as:

```
[DataContract]
public class Contact
{
    [DataMember]
    public string FirstName
    {get;set;}

    [DataMember]
    public string LastName;
}
```

The inferred data contract will be published in the service metadata.

If the type already contains `DataMember` attributes (but not a `DataContract` attribute), these data member contracts will be ignored as if they were not present. If the type does contain a `DataContract` attribute, no data contract is inferred. Likewise, if the type is internal, no data contract is inferred. Furthermore, all subclasses of a class that utilizes an inferred data contract must themselves be inferable; that is, they must be public classes and have no `DataContract` attribute.



Microsoft calls inferred data contracts POCOs, or “plain old CLR objects.”

In my opinion, relying on inferred data contracts is a sloppy hack that goes against the grain of most everything else in WCF. Much as WCF does not infer a service contract from a mere interface definition or enable transactions or reliability by default, it should not infer a data contract. Service orientation (with the exception of security) is heavily biased toward opting out by default, as it should be, to maximize encapsulation and decoupling. Do use the `DataContract` attribute, and be explicit about your data contracts. This will enable you to tap into data contract features such as versioning. The rest of this book does not use or rely on inferred data contracts.

Composite Data Contracts

When you define a data contract, you can apply the `DataMember` attribute on members that are themselves data contracts, as shown in [Example 3-3](#).

Example 3-3. A composite data contract

```
[DataContract]
class Address
{
    [DataMember]
    public string Street;

    [DataMember]
    public string City;

    [DataMember]
    public string State;

    [DataMember]
    public string Zip;
}

[DataContract]
struct Contact
{
    [DataMember]
    public string FirstName;

    [DataMember]
    public string LastName;

    [DataMember]
    public Address Address;
}
```

Being able to aggregate other data contracts in this way illustrates the fact that data contracts are actually recursive in nature. When you serialize a composite data contract, the `DataContractSerializer` will chase all applicable references in the object graph and capture their state as well. When you publish a composite data contract, all its comprising data contracts will be published as well. For example, using the same definitions as those in [Example 3-3](#), the metadata for this service contract:

```
[ServiceContract]
interface IContactManager
{
    [OperationContract]
    void AddContact(Contact contact);

    [OperationContract]
    Contact[] GetContacts();
}
```

will include the definition of the `Address` structure as well.

Data Contract Events

.NET provides support for serialization events for serializable types, and WCF provides the same support for data contracts. WCF calls designated methods on your data contract when serialization and deserialization take place. Four serialization and deserialization events are defined. The *serializing event* is raised just before serialization takes place and the *serialized event* is raised just after serialization. Similarly, the *deserializing event* is raised just before deserialization and the *deserialized event* is raised after deserialization. You designate methods as serialization event handlers using method attributes, as shown in [Example 3-4](#).

Example 3-4. Applying the serialization event attributes

```
[DataContract]
class MyDataContract
{
    [OnSerializing]
    void OnSerializing(StreamingContext context)
    {...}

    [OnSerialized]
    void OnSerialized(StreamingContext context)
    {...}

    [OnDeserializing]
    void OnDeserializing(StreamingContext context)
    {...}

    [OnDeserialized]
    void OnDeserialized(StreamingContext context)
    {...}
    //Data members
}
```

Each serialization event-handling method must have the following signature:

```
void <Method Name>(StreamingContext context);
```

If the serialization event attributes (defined in the `System.Runtime.Serialization` namespace) are applied on methods with incompatible signatures, WCF will throw an exception.

The `StreamingContext` structure informs the type of why it is being serialized, but it can be ignored for WCF data contracts.

As their names imply, the `OnSerializing` attribute designates a method to handle the serializing event and the `OnSerialized` attribute designates a method to handle the serialized event. Similarly, the `OnDeserializing` attribute designates a method to

handle the deserializing event and the `OnDeserialized` attribute designates a method to handle the deserialized event.

Figure 3-2 is an activity diagram depicting the order in which events are raised during serialization.

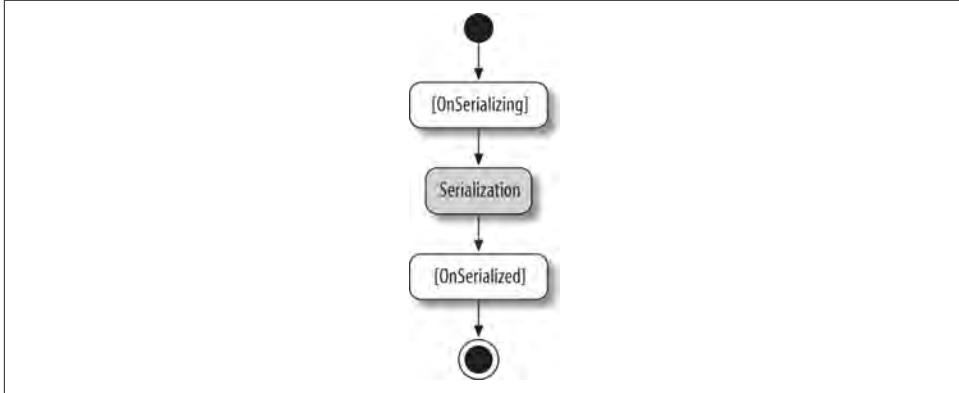


Figure 3-2. Events raised during serialization

WCF first raises the serializing event, thus invoking the corresponding event handler. Next, WCF serializes the object, and finally, the serialized event is raised and its event handler is invoked.

Figure 3-3 is an activity diagram depicting the order in which deserialization events are raised. WCF first raises the deserializing event, thus invoking the corresponding event handler. Next, WCF deserializes the object, and finally the deserialized event is raised and its event handler is invoked.

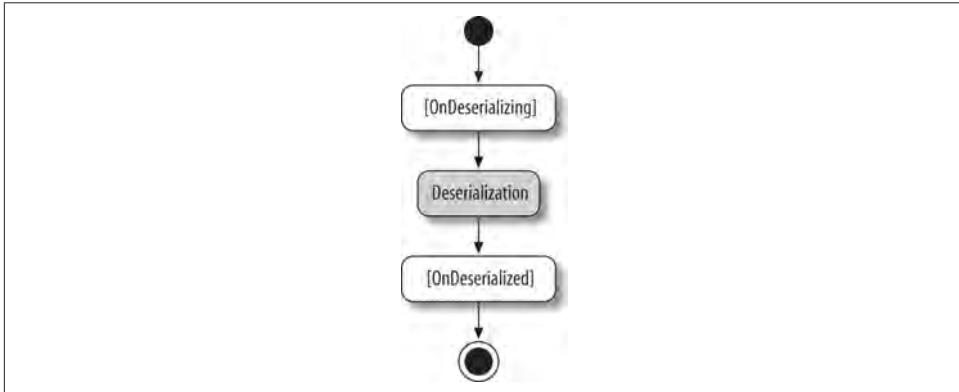


Figure 3-3. Events raised during deserialization

Note that in order to call the deserializing event-handling method, WCF has to first construct an object. However, it does so without ever calling your data contract class's default constructor.



WCF does not allow you to apply the same serialization event attribute on multiple methods of the data contract type. This is somewhat regrettable, because it precludes support for partial classes where each part deals with its own serialization events.

Using the deserializing event

Since no constructor calls are ever made during deserialization, the deserializing event-handling method is logically your deserialization constructor. It is intended for performing some custom pre-deserialization steps, typically, initialization of class members not marked as data members. Any value settings on members marked as data members will be in vain because WCF will set those members again during deserialization using values from the message. Other steps you can take in the deserializing event-handling method are setting specific environment variables (such as thread local storage), performing diagnostics, or signaling some global synchronization events. I would even go as far as to say that if you do provide such a deserializing event-handling method, you should have only a default constructor and have both the default constructor and the event handler call the same helper method so that anyone instantiating the type using regular .NET will perform exactly the same steps that you do and you will have a single place to maintain that code:

```
[DataContract]
class MyClass
{
    public MyClass()
    {
        OnDeserializing();
    }
    [OnDeserializing]
    void OnDeserializing(StreamingContext context)
    {
        OnDeserializing();
    }
    void OnDeserializing()
    {...}
}
```

Using the deserialized event

The deserialized event lets your data contract initialize or reclaim non-data members while utilizing already deserialized values. [Example 3-5](#) demonstrates this point, using the deserialized event to initialize a database connection. Without the event, the

data contract will not be able to function properly—since the constructor is never called, it will have a null for the connection object.

Example 3-5. Initializing non-serializable resources using the deserialized event

```
[DataContract]
class MyDataContract
{
    IDbConnection m_Connection;

    [OnDeserialized]
    void OnDeserialized(StreamingContext context)
    {
        m_Connection = new SqlConnection(...);
    }
    /* Data members */
}
```

Shared Data Contracts

When adding a service reference in Visual Studio, you must provide a unique namespace for each service reference. The imported types will be defined in that new namespace. This presents a problem when adding references for two different services that share the same data contract, since you will get two distinct types in two different namespaces representing the same data contract.

By default, however, if any of the assemblies referenced by the client has a data contract type that matches a data contract type already exposed in the metadata of the referenced service, Visual Studio will not import that type again. It is worth emphasizing again that the existing data contract reference must be in another referenced assembly, not in the client project itself. This limitation may be addressed in a future release of Visual Studio, but for now, the workaround and best practice is obvious: factor all of your shared data contracts to a designated class library, and have all clients reference that assembly. You can then control and configure which referenced assemblies (if any) to consult regarding those shared data contracts via the advanced settings dialog box for the service reference, shown in [Figure 1-10](#). The “Reuse types in referenced assemblies” checkbox is checked by default, but you can turn off this feature if you so desire. Despite its name, it will share only data contracts, not service contracts. Using the radio buttons below it, you can also instruct Visual Studio to reuse data contracts across all referenced assemblies, or restrict the sharing to specific assemblies by selecting them in the list.

Data Contract Hierarchy

Your data contract class may be a subclass of another data contract class. WCF requires that every level in the class hierarchy explicitly opt in for a given data contract, because the `DataContract` attribute is not inheritable:

```
[DataContract]
class Contact
{
    [DataMember]
    public string FirstName;

    [DataMember]
    public string LastName;
}
[DataContract]
class Customer : Contact
{
    [DataMember]
    public int CustomerNumber;
}
```

Failing to designate every level in the class hierarchy as serializable or as a data contract will result in an `InvalidOperationException` at the service load time. WCF lets you mix the `Serializable` and `DataContract` attributes in the class hierarchy:

```
[Serializable]
class Contact
{...}

[DataContract]
class Customer : Contact
{...}
```

However, the `Serializable` attribute will typically be at the root of the class hierarchy, if it appears at all, because new classes should use the `DataContract` attribute. When you export a data contract hierarchy, the metadata maintains the hierarchy, and all levels of the class hierarchy are exported when you make use of the subclass in a service contract:

```
[ServiceContract]
interface IContactManager
{
    [OperationContract]
    void AddCustomer(Customer customer); //Contact is exported as well
    ...
}
```

Known Types

In traditional object-oriented programming, a reference to a subclass is also a reference to its base class, so the subclass maintains an Is-A relationship with its base class. Any method that expects a reference to a base class can also accept a reference to its subclass. This is a direct result of the way the compiler spans the state of the subclass in memory, by appending it right after the base class section.

Although languages such as C# let you substitute a subclass for a base class in this manner, this is not the case with WCF operations. By default, you cannot use a subclass of a data contract class instead of its base class. Consider this service contract:

```
[ServiceContract]
interface IContactManager
{
    //Cannot accept Customer object here:
    [OperationContract]
    void AddContact(Contact contact);

    //Cannot return Customer objects here:
    [OperationContract]
    Contact[] GetContacts();
}
```

Suppose the client defined the `Customer` class as well:

```
[DataContract]
class Customer : Contact
{
    [DataMember]
    public int CustomerNumber;
}
```

While the following code will compile successfully, it will fail at runtime:

```
Contact contact = new Customer();

ContactManagerClient proxy = new ContactManagerClient();

//Service call will fail:
proxy.AddContact(contact);

proxy.Close();
```

The reason is that you are not actually passing an object reference; you are instead passing the object's state. When you pass in a `Customer` instead of a `Contact`, as in the previous example, the service does not know it should deserialize the `Customer` portion of the state.

Likewise, when a `Customer` is returned instead of a `Contact`, the client does not know how to deserialize it, because all it knows about are `Contacts`, not `Customers`:

```

//////////////////////////// Service Side /////////////////////////
[DataContract]
class Customer : Contact
{
    [DataMember]
    public int CustomerNumber;
}
class CustomerManager : IContactManager
{
    List<Customer> m_Customers = new List<Customer>();

    public Contact[] GetContacts()
    {
        return m_Customers.ToArray();
    }
    //Rest of the implementation
}
//////////////////////////// Client Side /////////////////////////
ContactManagerClient proxy = new ContactManagerClient();
//Call will fail if there are items in the list:
Contact[] contacts = proxy.GetContacts();
proxy.Close();

```

The solution is to explicitly tell WCF about the `Customer` class using the `KnownTypeAttribute` attribute, defined as:

```

[AttributeUsage(AttributeTargets.Struct|AttributeTargets.Class,
               AllowMultiple = true)]
public sealed class KnownTypeAttribute : Attribute
{
    public KnownTypeAttribute(Type type);
    //More members
}

```

The `KnownType` attribute allows you to designate acceptable subclasses for the data contract:

```

[DataContract]
[KnownType(typeof(Customer))]
class Contact
{...}

[DataContract]
class Customer : Contact
{...}

```

On the host side, the `KnownType` attribute affects all contracts and operations using the base class, across all services and endpoints, allowing it to accept subclasses instead of base classes. In addition, it includes the subclass in the metadata so that the client will have its own definition of the subclass and will be able to pass the subclass instead of the base class. If the client also applies the `KnownType` attribute on its copy of the base class, it can in turn receive the known subclass back from the service.

Service Known Types

The downside of using the `KnownType` attribute is that it may be too broad in scope. WCF also provides the `ServiceKnownTypeAttribute`, defined as:

```
[AttributeUsage(AttributeTargets.Interface |  
    AttributeTargets.Method |  
    AttributeTargets.Class,  
    AllowMultiple = true)]  
public sealed class ServiceKnownTypeAttribute : Attribute  
{  
    public ServiceKnownTypeAttribute(Type type);  
    //More members  
}
```

Instead of using the `KnownType` attribute on the base data contract, you can apply the `ServiceKnownType` attribute on a specific operation on the service side. Then, only that operation (across all supporting services) can accept the known subclass:

```
[DataContract]  
class Contact  
{...}  
  
[DataContract]  
class Customer : Contact  
{...}  
  
[ServiceContract]  
interface IContactManager  
{  
    [OperationContract]  
    [ServiceKnownType(typeof(Customer))]  
    void AddContact(Contact contact);  
  
    [OperationContract]  
    Contact[] GetContacts();  
}
```

Other operations cannot accept the subclass.

When the `ServiceKnownType` attribute is applied at the contract level, all the operations in that contract can accept the known subclass across all implementing services:

```
[ServiceContract]  
[ServiceKnownType(typeof(Customer))]  
interface IContactManager  
{  
    [OperationContract]  
    void AddContact(Contact contact);  
  
    [OperationContract]  
    Contact[] GetContacts();  
}
```



Do not apply the `ServiceKnownType` attribute on the service class itself. Although the code will compile, this will have an effect only when you don't define the service contract as an interface (something I strongly discourage in any case). If you apply the `ServiceKnownType` attribute on the service class while there is a separate contract definition, it will have no effect.

Whether you apply the `ServiceKnownType` attribute at the operation or the contract level, the exported metadata and the generated proxy will have no trace of it and will include the `KnownType` attribute on the base class only. For example, given this service-side definition:

```
[ServiceContract]
[ServiceKnownType(typeof(Customer))]
interface IContactManager
{...}
```

The imported definition will be:

```
[DataContract]
[KnownType(typeof(Customer))]
class Contact
{...}
[DataContract]
class Customer : Contact
{...}
[ServiceContract]
interface IContactManager
{...}
```

You can manually rework the client-side proxy class to correctly reflect the service-side semantic by removing the `KnownType` attribute from the base class and applying the `ServiceKnownType` attribute to the appropriate level in the contract.

Multiple Known Types

You can apply the `KnownType` and `ServiceKnownType` attributes multiple times to inform WCF about as many known types as required:

```
[DataContract]
class Contact
{...}

[DataContract]
class Customer : Contact
{...}

[DataContract]
class Person : Contact
{...}
```

```

[ServiceContract]
[ServiceKnownType(typeof(Customer))]
[ServiceKnownType(typeof(Person))]
interface IContactManager
{...}

```

The WCF formatter uses reflection to collect all the known types of the data contracts, then examines the provided parameter to see if it is of any of the known types.

Note that you must explicitly add all levels in the data contract class hierarchy. Adding a subclass does not add its base class(es):

```

[DataContract]
class Contact
{...}

[DataContract]
class Customer : Contact
{...}

[DataContract]
class Person : Customer
{...}

[ServiceContract]
[ServiceKnownType(typeof(Customer))]
[ServiceKnownType(typeof(Person))]
interface IContactManager
{...}

```

Configuring Known Types

The main downside of the known types attributes is that they require the service or the client to know in advance about all possible subclasses the other party may want to use. Adding a new subclass necessitates changing the code, recompiling, and redeploying. To alleviate this, WCF lets you configure the known types in the service's or client's config file, as shown in [Example 3-6](#). You need to provide not just the type names, but also the names of their containing assemblies.

Example 3-6. Known types in config file

```

<system.runtime.serialization>
    <dataContractSerializer>
        <declaredTypes>
            <add type = "Contact,MyClassLibrary,Version = 1.0.0.0,Culture = neutral,
                  PublicKeyToken = null">
                <knownType type = "Customer,MyOtherClassLibrary,Version = 1.0.0.0,
                  Culture = neutral,PublicKeyToken = null"/>
            </add>
        </declaredTypes>
    </dataContractSerializer>
</system.runtime.serialization>

```

```
</declaredTypes>
</DataContractSerializer>
</system.runtime.serialization>
```

When not relying on string name or assembly version resolution, you can just use the assembly-friendly name:

```
<add type = "Contact,MyClassLibrary">
    <knownType type = "Customer,MyOtherClassLibrary"/>
</add>
```

Including the known types in the config file has the same effect as applying the Known Type attribute on the data contract, and the published metadata will include the known types definition.



Using a config file to declare a known type is the only way to add a known type that is internal to another assembly.

Data Contract Resolvers

The final technique for addressing known types would be to do so programmatically. This is the most powerful technique, since you can extend it to completely automate dealing with the known type issues. This is possible using a mechanism called data contract resolvers introduced by WCF in .NET 4.0. In essence, you are given a chance to intercept the operation's attempt to serialize and deserialize parameters and resolve the known types at runtime both on the client and service sides.

The first step in implementing a programmatic resolution is to derive from the abstract class `DataContractResolver` defined as:

```
public abstract class DataContractResolver
{
    protected DataContractResolver();

    public abstract bool TryResolveType(Type type, Type declaredType,
                                         DataContractResolver knownTypeResolver,
                                         out XmlDictionaryString typeName,
                                         out XmlDictionaryString typeNamespace);

    public abstract Type ResolveName(string typeName, string typeNamespace,
                                     Type declaredType,
                                     DataContractResolver knownTypeResolver);
}
```

Your implementation of the `TryResolveType()` is called when WCF tries to serialize a type into a message and the type provided (the `type` parameter) is different from the

type declared in the operation contract (the `declaredType` parameter). If you want to serialize the type, you need to provide some unique identifiers to serve as keys into a dictionary that maps identifiers to types. WCF will provide you those keys during deserialization so that you can bind against that type. Note that the namespace key cannot be an empty string or a null. While virtually any unique string value will do for the identifiers, I recommend simply using the CLR type name and namespace. Set the type name and namespace into the `typeName` and `typeNamespace` out parameters.

If you return `true` from `TryResolveType()`, the type is considered resolved as if you had applied the `KnownType` attribute. If you return `false`, WCF fails the call. Note that `TryResolveType()` must resolve all known types, even those types that are decorated with the `KnownType` attribute or are listed in the config file. This presents a potential risk: it requires the resolver to be coupled to all known types in the application and will fail the operation call with other types that may come over time. It is therefore preferable as a fallback contingency to try to resolve the type using the default known types resolver that WCF would have used if your resolver was not in use. This is exactly what the `knownTypeResolver` parameter is for. If your implementation of `TryResolveType()` cannot resolve the type, it should delegate to `knownTypeResolver`.

The `ResolveName()` is called when WCF tries to deserialize a type out of a message, and the type provided (the `type` parameter) is different from the type declared in the operation contract (the `declaredType` parameter). In this case, WCF provides you with the type name and namespace identifiers so that you can map them back to a known type.

For example, consider again these two data contracts:

```
[DataContract]
class Contact
{...}

[DataContract]
class Customer : Contact
{...}
```

Example 3-7 lists a simple resolver for the `Customer` type.

Example 3-7. The CustomerResolver

```
class CustomerResolver : DataContractResolver
{
    string Namespace
    {
        get
        {
            return typeof(Customer).Namespace ?? "global";
        }
    }
}
```

```

        }

        string Name
        {
            get
            {
                return typeof(Customer).Name;
            }
        }

    public override Type ResolveName(string typeName, string typeNamespace,
                                    Type declaredType,
                                    DataContractResolver knownTypeResolver)
    {
        if(typeName == Name && typeNamespace == Namespace)
        {
            return typeof(Customer);
        }
        else
        {
            return knownTypeResolver.
                ResolveName(typeName, typeNamespace, declaredType, null);
        }
    }

    public override bool TryResolveType(Type type, Type declaredType,
                                    DataContractResolver knownTypeResolver,
                                    out XmlDictionaryString typeName,
                                    out XmlDictionaryString typeNamespace)
    {
        if(type == typeof(Customer))
        {
            XmlDictionary dictionary = new XmlDictionary();
            typeName = dictionary.Add(Name);
            typeNamespace = dictionary.Add(Namespace);
            return true;
        }
        else
        {
            return knownTypeResolver.
                TryResolveType(type, declaredType, null, out typeName, out typeNamespace);
        }
    }
}

```

Installing the data contract resolver

The resolver must be attached as a behavior for each operation on the proxy or the service endpoint. As explained in [Chapter 1](#), a behavior is a local aspect of the service that does not affect in any way the message or the communication channel to the service. For example, how you choose to resolve a known type (be it declaratively via the

`KnownType` attribute or programmatically with a resolver) is a local implementation detail, on both the client and the service sides.

In WCF, every endpoint is represented by the type `ServiceEndpoint`. The `ServiceEndpoint` has a property called `Contract` of the type `ContractDescription`:

```
public class ServiceEndpoint
{
    public ContractDescription Contract
    {get;set;}

    //More members
}
```

`ContractDescription` has a collection of operation descriptions, with an instance of `OperationDescription` for every operation on the contract:

```
public class ContractDescription
{
    public OperationDescriptionCollection Operations
    {get;}

    //More members
}
public class OperationDescriptionCollection :
    Collection<OperationDescription>
{...}
```

Each `OperationDescription` has a collection of operation behaviors of the type `IOperationBehavior`:

```
public class OperationDescription
{
    public KeyedByTypeCollection<IOperationBehavior> Behaviors
    {get;}
    //More members
}
```

In its collection of behaviors, every operation always has a behavior called `DataContractSerializerOperationBehavior` with a `DataContractResolver` property:

```
public class DataContractSerializerOperationBehavior : IOperationBehavior, ...
{
    public DataContractResolver DataContractResolver
    {get;set;}
    //More members
}
```

The `DataContractResolver` property defaults to `null`, but you can set it to your custom resolver.

To install a resolver on the host side, you must iterate over the collection of endpoints in the service description maintained by the host:

```

public class ServiceHost : ServiceHostBase
{...}

public abstract class ServiceHostBase : ...
{
    public ServiceDescription Description
    {get;}
    //More members
}

public class ServiceDescription
{
    public ServiceEndpointCollection Endpoints
    {get;}

    //More members
}

public class ServiceEndpointCollection : Collection<ServiceEndpoint>
{...}

```

Suppose you have the following service definition and are using the resolver in [Example 3-7](#):

```

[ServiceContract]
interface IContactManager
{
    [OperationContract]
    void AddContact(Contact contact);
    ...
}
class ContactManager : IContactManager
{...}

```

[Example 3-8](#) shows how to install the resolver on the host for the `ContactManager` service.

Example 3-8. Installing a resolver on the host

```

ServiceHost host = new ServiceHost(typeof(ContactManager));

foreach(ServiceEndpoint endpoint in host.Description.Endpoints)
{
    foreach(OperationDescription operation in endpoint.Contract.Operations)
    {
        DataContractSerializerOperationBehavior behavior =
            operation.Behaviors.Find<DataContractSerializerOperationBehavior>();
        behavior.DataContractResolver = new CustomerResolver();
    }
}
host.Open();

```

On the client side, you follow similar steps, except you need to set the resolver on the single endpoint of the proxy or the channel factory. For example, given this proxy class definition:

```
class ContactManagerClient : ClientBase<IContactManager>, IContactManager
{...}
```

Example 3-9 shows how to install the resolver on the proxy in order to call the service of **Example 3-8** with a known type.

Example 3-9. Installing a resolver on the proxy

```
ContactManagerClient proxy = new ContactManagerClient();

foreach(OperationDescription operation in proxy.Endpoint.Contract.Operations)
{
    DataContractSerializerOperationBehavior behavior =
        operation.Behaviors.Find<DataContractSerializerOperationBehavior>();

    behavior.DataContractResolver = new CustomerResolver();
}

Customer customer = new Customer();
...

proxy.AddContact(customer);
```

The generic resolver

Writing and installing a resolver for each type is obviously a lot of work, requiring you to meticulously track all known types, something that is error-prone and can quickly get out of hand in an evolving system. To automate implementing a resolver, I wrote the class `GenericResolver` defined as:

```
public class GenericResolver : DataContractResolver
{
    public Type[] KnownTypes
    {get; }

    public GenericResolver();
    public GenericResolver(Type[] typesToResolve);

    public static GenericResolver Merge(GenericResolver resolver1,
                                       GenericResolver resolver2);
}
```

`GenericResolver` offers two constructors. One constructor can accept an array of known types to resolve. The types in the array can include *bounded generic types* (that is, generic types for which you have already specified type parameters). The parameterless constructor will automatically add as known types all classes and structs in the

calling assembly and all public classes and structs in assemblies referenced by the calling assemblies. The parameterless constructor will not add types originating in a .NET Framework-referenced assembly. Note that the parameterless constructor will also ignore generic types (since there is no way of inferring the type parameters used in code). In addition, `GenericResolver` offers the `Merge()` static method that you can use to merge the known types of two resolvers, returning a `GenericResolver` that resolves the union of the two resolvers provided. [Example 3-10](#) shows the pertinent portion of `GenericResolver` without reflecting the types in the assemblies, which has nothing to do with WCF.

Example 3-10. Implementing GenericResolver (partial)

```
public class GenericResolver : DataContractResolver
{
    const string DefaultNamespace = "global";

    readonly Dictionary<Type, Tuple<string, string>> m_TypeToNames;
    readonly Dictionary<string, Dictionary<string, Type>> m_NamesToType;

    public Type[] KnownTypes
    {
        get
        {
            return m_TypeToNames.Keys.ToArray();
        }
    }

    //Get all types in calling assembly and referenced assemblies
    static Type[] ReflectTypes()
    {...}

    public GenericResolver() : this(ReflectTypes())
    {}
    public GenericResolver(Type[] typesToResolve)
    {
        m_TypeToNames = new Dictionary<Type, Tuple<string, string>>();
        m_NamesToType = new Dictionary<string, Dictionary<string, Type>>();

        foreach(Type type in typesToResolve)
        {
            string typeNamespace = GetNamespace(type);
            string typeName = GetName(type);

            m_TypeToNames[type] = new Tuple<string, string>(typeNamespace, typeName);

            if(m_NamesToType.ContainsKey(typeNamespace) == false)
            {
                m_NamesToType[typeNamespace] = new Dictionary<string, Type>();
            }
        }
    }
}
```

```

        m_NamesToType[typeNamespace][typeName] = type;
    }
}
static string GetNamespace(Type type)
{
    return type.Namespace ?? DefaultNamespace;
}
static string GetName(Type type)
{
    return type.Name;
}

public static GenericResolver Merge(GenericResolver resolver1,
                                    GenericResolver resolver2)
{
    if(resolver1 == null)
    {
        return resolver2;
    }
    if(resolver2 == null)
    {
        return resolver1;
    }
    List<Type> types = new List<Type>();

    types.AddRange(resolver1.KnownTypes);
    types.AddRange(resolver2.KnownTypes);

    return new GenericResolver(types.ToArray());
}
public override Type ResolveName(string typeName,string typeNamespace,
                                 Type declaredType,
                                 DataContractResolver knownTypeResolver)
{
    if(m_NamesToType.ContainsKey(typeNamespace))
    {
        if(m_NamesToType[typeNamespace].ContainsKey(typeName))
        {
            return m_NamesToType[typeNamespace][typeName];
        }
    }
    return knownTypeResolver.
        ResolveName(typeName,typeNamespace,declaredType,null);
}
public override bool TryResolveType(Type type,Type declaredType,
                                    DataContractResolver knownTypeResolver,
                                    out XmlDictionaryString typeName,
                                    out XmlDictionaryString typeNamespace)
{
    if(m_TypeToNames.ContainsKey(type))
    {
        XmlDictionary dictionary = new XmlDictionary();

```

```

        typeNamespace = dictionary.Add(m_TypeToNames[type].Item1);
        typeName      = dictionary.Add(m_TypeToNames[type].Item2);
        return true;
    }
    else
    {
        return knownTypeResolver.
            TryResolveType(type, declaredType, null, out typeName,
                           out typeNamespace);
    }
}

```

The most important members of `GenericResolver` are the `m_TypeToNames` and the `m_NamesToType` dictionaries. `m_TypeToNames` maps a type to a tuple of its name and namespace. `m_NamesToType` maps a type namespace and name to the actual type. The constructor that takes the array of types initializes those two dictionaries. The `Merge()` method uses the helper constructor to merge the two resolvers. The `TryResolveType()` method uses the provided type as a key into the `m_TypeToNames` dictionary to read the type's name and namespace. The `ResolveName()` method uses the provided namespace and name as keys into the `m_NamesToType` dictionary to return the resolved type.

Installing the generic resolver

While you could use tedious code similar to Examples 3-8 and 3-9 to install `GenericResolver`, it is best to streamline it with extension methods. To that end, use my `AddGenericResolver()` methods of `GenericResolverInstaller` defined as:

```

public static class GenericResolverInstaller
{
    public static void AddGenericResolver(this ServiceHost host,
                                         params Type[] typesToResolve);

    public static void AddGenericResolver<T>(this ClientBase<T> proxy,
                                              params Type[] typesToResolve)
                                              where T : class;

    public static void AddGenericResolver<T>(this ChannelFactory<T> factory,
                                              params Type[] typesToResolve)
                                              where T : class;
}

```

The `AddGenericResolver()` method accepts a `params` array of types, which means an open-ended, comma-separated list of types. If you do not specify types, that will make `AddGenericResolver()` add as known types all classes and structs in the calling assembly plus the public classes and structs in referenced assemblies. For example, given these known types:

```
[DataContract]
class Contact
{...}

[DataContract]
class Customer : Contact
{...}[DataContract]
class Employee : Contact
{...}
```

Example 3-11 shows several examples of using the `AddGenericResolver()` extension method.

Example 3-11. Installing GenericResolver

```
//Host side

ServiceHost host1 = new ServiceHost(typeof(ContactManager));
//Resolve all types in this and referenced assemblies
host1.AddGenericResolver();
host1.Open();

ServiceHost host2 = new ServiceHost(typeof(ContactManager));
//Resolve only Customer and Employee
host2.AddGenericResolver(typeof(Customer),typeof(Employee));
host2.Open();

ServiceHost host3 = new ServiceHost(typeof(ContactManager));
//Can call AddGenericResolver() multiple times
host3.AddGenericResolver(typeof(Customer));
host3.AddGenericResolver(typeof(Employee));
host3.Open();

//Client side

ContactManagerClient proxy = new ContactManagerClient();
//Resolve all types in this and referenced assemblies
proxy.AddGenericResolver();

Customer customer = new Customer();
...
proxy.AddContact(customer);
```

`GenericResolverInstaller` not only installs the `GenericResolver`, it also tries to merge it with the old generic resolver (if present). This means you can call the `AddGenericResolver()` method multiple times. This is handy when adding bounded generic types:

```
[DataContract]
class Customer<T> : Contact
{...}
```

```

ServiceHost host = new ServiceHost(typeof(ContactManager));

//Add all non-generic known types
host.AddGenericResolver(); //Add the generic types
host.AddGenericResolver(typeof(Customer<int>),Customer<string>));

host.Open();

```

Example 3-12 shows the partial implementation of `GenericResolverInstaller`.

Example 3-12. Implementing GenericResolverInstaller

```

public static class GenericResolverInstaller
{
    public static void AddGenericResolver(this ServiceHost host,
                                         params Type[] typesToResolve)
    {
        foreach(ServiceEndpoint endpoint in host.Description.Endpoints)
        {
            AddGenericResolver(endpoint,typesToResolve);
        }
    }

    static void AddGenericResolver(ServiceEndpoint endpoint,Type[] typesToResolve)
    {
        foreach(OperationDescription operation in endpoint.Contract.Operations)
        {
            DataContractSerializerOperationBehavior behavior = operation.
                Behaviors.Find<DataContractSerializerOperationBehavior>();

            GenericResolver newResolver;

            if(typesToResolve == null || typesToResolve.Any() == false)
            {
                newResolver = new GenericResolver();
            }
            else
            {
                newResolver = new GenericResolver(typesToResolve);
            }

            GenericResolver oldResolver = behavior.DataContractResolver
                                         as GenericResolver;
            behavior.DataContractResolver =
                GenericResolver.Merge(oldResolver,newResolver);
        }
    }
}

```

If no types are provided, `AddGenericResolver()` will use the parameterless constructor of `GenericResolver`. Otherwise, it will use only the specified types by calling the other constructor. Note the merging with the old resolver if present.

GenericResolver and ServiceHost<T>

Using a generic resolver is great, and arguably should always be associated with any host since you have no way of knowing in advance all the known types your system will encounter. In that case, instead of explicitly adding `GenericResolver` to the host (as in [Example 3-11](#)), you can also have a custom host type that adds the generic resolver implicitly. For example, `My ServiceHost<T>` (presented in [Chapter 1](#)) does just that:

```
public class ServiceHost<T> : ServiceHost
{
    protected override void OnOpening()
    {
        this.AddGenericResolver();
        ...
    }
}
```

This means that the following code is all you will need when it comes to known types:

```
ServiceHost host = new ServiceHost<ContactManager>();
host.Open();
```



The `InProcFactory` presented in [Chapter 1](#) actually uses `ServiceHost<T>` internally, and as such, it benefits automatically from the generic resolver without any additional steps.

Generic resolver attribute

If your service relies on the generic resolver by design, it is better not to be at the mercy of the host and to declare your need for the generic resolver at design time. To that end, I wrote the `GenericResolverBehaviorAttribute`:

```
[AttributeUsage(AttributeTargets.Class)]
public class GenericResolverBehaviorAttribute : Attribute,IServiceBehavior
{
    void IServiceBehavior.Validate(ServiceDescription serviceDescription,
                                    ServiceHostBase serviceHostBase)
    {
        ServiceHost host = serviceHostBase as ServiceHost;
        host.AddGenericResolver();
    }
    //More members
}
```

This concise attribute makes the service independent of the host and the config file:

```
[GenericResolverBehavior]
class ContactManager : IContactManager
{...}
```

`GenericResolverBehaviorAttribute` derives from `IServiceBehavior`, which is a special WCF interface and the most commonly used extension in WCF. Subsequent chapters will make extensive use of `IServiceBehavior` and will discuss its various methods. Briefly, when the host loads the service, it uses reflection to determine if the service class has an attribute that supports `IServiceBehavior`, and if so, the host calls the `IServiceBehavior` methods, specifically the `Validate()` method, which lets the attribute interact with the host. In the case of `GenericResolverBehaviorAttribute`, it adds the generic resolver to the host.

Objects and Interfaces

The base type of a data contract class or a struct can be an interface:

```
interface IContact
{
    string FirstName
    {get;set;}
    string LastName
    {get;set;}
}
[DataContract]
class Contact : IContact
{...}
```

You can use such a base interface in your service contract or as a data member in a data contract if you use the `ServiceKnownType` attribute to designate the actual data type:

```
[ServiceContract]
[ServiceKnownType(typeof(Contact))]
interface IContactManager
{
    [OperationContract]
    void AddContact(IContact contact);

    [OperationContract]
    IContact[] GetContacts();
}
```

You cannot apply the `KnownType` attribute on the base interface, because the interface itself will not be included in the exported metadata. Instead, the imported service contract will be object-based and it will not include the data contract interface:

```
//Imported definitions:
[DataContract]
```

```

class Contact
{...}

[ServiceContract]
interface IContactManager
{
    [OperationContract]
    [ServiceKnownType(typeof(Contact))]
    [ServiceKnownType(typeof(object[]))]
    void AddContact(object contact);

    [OperationContract]
    [ServiceKnownType(typeof(Contact))]
    [ServiceKnownType(typeof(object[]))]
    object[] GetContacts();
}

```

The imported definition will always have the `ServiceKnownType` attribute applied at the operation level, even if it was originally defined at the scope of the contract. In addition, every operation will include a union of all the `ServiceKnownType` attributes required by all the operations, including a redundant service known type attribute for the array. These are relics from a time when these definitions were required in a beta version of WCF.

You can manually rework the imported definition to have only the required `Service KnownType` attributes:

```

[DataContract]
class Contact
{...}

[ServiceContract]
interface IContactManager
{
    [OperationContract]
    [ServiceKnownType(typeof(Contact))]
    void AddContact(object contact);

    [OperationContract]
    [ServiceKnownType(typeof(Contact))]
    object[] GetContacts();
}

```

Or better yet, if you have the definition of the base interface on the client side or if you refactor that definition, you can use that instead of `object`. This gives you an added degree of type safety as long as you add a derivation from the interface to the data contract:

```

[DataContract]
class Contact : IContact
{...}

```

```

[ServiceContract]
interface IContactManager
{
    [OperationContract]
    [ServiceKnownType(typeof(Contact))]
    void AddContact(IContact contact);

    [OperationContract]
    [ServiceKnownType(typeof(Contact))]
    IContact[] GetContacts();
}

```

However, you cannot replace the `object` in the imported contract with the concrete data contract type, because it is no longer compatible:

```

//Invalid client-side contract
[ServiceContract]
interface IContactManager
{
    [OperationContract]
    void AddContact(Contact contact);

    [OperationContract]
    Contact[] GetContacts();
}

```

Data Contract Equivalence

Two data contracts are considered *equivalent* if they have the same wire representation—that is, if they have the same schema. This can be the case if they define the same type (but not necessarily the same version of the type) or if the two data contracts refer to two different types with the same data contract and data member names. Equivalent data contracts are interchangeable: WCF will let any service that was defined with one data contract operate with an equivalent data contract.

The most common way of defining an equivalent data contract is to use the `DataContract` and `DataMember` attributes' `Name` properties to map one data contract to another. In the case of the `DataContract` attribute, the `Name` property defaults to the type's name, so these two definitions are identical:

```

[DataContract]
struct Contact
{...}

[DataContract(Name = "Contact")]
struct Contact
{...}

```

In fact, the full name of the data contract always includes its namespace as well, but as you have seen, you can assign a different namespace.

In the case of the `DataMember` attribute, the `Name` property defaults to the member name, so these two definitions are identical:

```
[DataMember]  
string FirstName;  
  
[DataMember(Name = "FirstName")]  
string FirstName;
```

By assigning different names to the data contract and data members, you can generate an equivalent data contract from a different type. For example, these two data contracts are equivalent:

```
[DataContract]  
struct Contact  
{  
    [DataMember]  
    public string FirstName;  
  
    [DataMember]  
    public string LastName;  
}  
[DataContract(Name = "Contact")]  
struct Person  
{  
    [DataMember(Name = "FirstName")]  
    public string Name;  
  
    [DataMember(Name = "LastName")]  
    public string Surname;  
}
```

In addition to having identical names, the types of the data members have to match.



A class and a structure that support the same data contract are interchangeable.

Serialization Order

In .NET, a subclass can define a member of the same name and type as a private member of its base class, and in turn, its own subclass can do the same:

```
class A  
{  
    string Name;
```

```

}
class B : A
{
    string Name;
}
class C : B
{
    string Name;
}

```

If the class hierarchy is also a data contract, this presents a problem when serializing an instance of the subclass into a message, since the message will contain multiple copies of a data member with the same name and type. To distinguish between them, WCF places the data members in the message in a particular order.

The default serialization order inside a type is simply alphabetical, and across a class hierarchy the order is top-down. In case of a mismatch in the serialization order, the members will be initialized to their default values. For example, when serializing a `Customer` instance defined as:

```

[DataContract]
class Contact
{
    [DataMember]
    public string FirstName;
    [DataMember]
    public string LastName;
}
[DataContract]
class Customer : Contact
{
    [DataMember]
    public int CustomerNumber;
}

```

the members will be serialized in the following order: `FirstName`, `LastName`, `CustomerNumber`.

Of course, equivalent data contracts must serialize and deserialize their members in the same order. The problem now is that combining a data contract hierarchy with aliasing contracts and members might break the serialization order. For example, the following data contract is not equivalent to the `Customer` data contract:

```

[DataContract(Name = "Customer")]
public class Person
{
    [DataMember(Name = "FirstName")]
    public string Name;

    [DataMember(Name = "LastName")]
    public string Surname;
}

```

```
[DataMember]
public int CustomerNumber;
}
```

because the serialization order is `CustomerNumber`, `FirstName`, `LastName`.

To resolve this conflict, you need to provide WCF with the order of serialization by setting the `Order` property of the `DataMember` attribute. The value of the `Order` property defaults to `-1`, meaning the default WCF ordering, but you can assign it values indicating the required order:

```
[DataContract(Name = "Customer")]
public class Person
{
    [DataMember(Name = "FirstName", Order = 1)]
    public string Name;

    [DataMember(Name = "LastName", Order = 2)]
    public string Surname;

    [DataMember(Order = 3)]
    public int CustomerNumber;
}
```

When renaming data members, you must take care to manually change their order. Even without renaming, the sorting can quickly get out of hand with a large number of data members. Fortunately, if another member has the same value for its `Order` property, WCF will order them alphabetically. You can take advantage of this behavior by assigning the same number to all members coming from the same level in the original class hierarchy, or better yet, simply assigning them their levels in that hierarchy:

```
[DataContract(Name = "Customer")]
public class Person
{
    [DataMember(Name = "FirstName", Order = 1)]
    public string Name;

    [DataMember(Name = "LastName", Order = 1)]
    public string Surname;

    [DataMember(Order = 2)]
    public int CustomerNumber;
}
```

Versioning

Services should be decoupled from their clients as much as possible, especially when it comes to versioning and technologies. Any version of the client should be able to

consume any version of the service and should do so without resorting to version numbers (such as those in assemblies), because those are .NET-specific. When a service and a client share a data contract, an important objective is to allow the service and client to evolve their versions of the data contract separately. To allow such decoupling, WCF needs to enable both backward and forward compatibility, without even sharing types or version information. There are three main versioning scenarios:

- New members
- Missing members
- Round-tripping, in which a new data contract version is passed to and from a client or service with an older version, requiring both backward and forward compatibility

By default, data contracts are version-tolerant and will silently ignore incompatibilities.

New Members

The most common change made to a data contract is adding new members on one side and sending the new contract to an old client or service. When deserializing the type, `DataContractSerializer` will simply ignore the new members. As a result, both the service and the client can accept data with new members that were not part of the original contract. For example, suppose the service is built against this data contract:

```
[DataContract]
struct Contact
{
    [DataMember]
    public string FirstName;

    [DataMember]
    public string LastName;
}
```

yet the client sends it this data contract instead:

```
[DataContract]
struct Contact
{
    [DataMember]
    public string FirstName;

    [DataMember]
    public string LastName;

    [DataMember]
}
```

```
    public string Address;  
}
```

Note that adding new members and having them ignored in this way breaks the data contract schema compatibility, because a service (or a client) that is compatible with one schema is all of a sudden compatible with a new schema.

Missing Members

By default, WCF lets either party remove members from the data contract. That is, you can serialize a type without certain members and send it to another party that expects the missing members. Although normally you probably won't intentionally remove members, the more likely scenario is when a client that is written against an old definition of the data contract interacts with a service written against a newer definition of that contract that expects new members. When, on the receiving side, `DataContractSerializer` does not find in the message the information required to deserialize those members, it will silently deserialize them to their default values; that is, `null` for reference types and a zero whitewash for value types. In effect, it will be as if the sending party never initialized those members. This default policy enables a service to accept data with missing members or return data with missing members to the client. [Example 3-13](#) demonstrates this point.

Example 3-13. Missing members are initialized to their default values

```
//////////////////////////////////////////////////////////////// Service Side ///////////////////////////////  
[DataContract]  
struct Contact  
{  
    [DataMember]  
    public string FirstName;  
  
    [DataMember]  
    public string LastName;  
  
    [DataMember]  
    public string Address;  
}  
  
[ServiceContract]  
interface IContactManager  
{  
    [OperationContract]  
    void AddContact(Contact contact);  
    ...  
}  
  
class ContactManager : IContactManager  
{
```

```

public void AddContact(Contact contact)
{
    Trace.WriteLine("First name = " + contact.FirstName);
    Trace.WriteLine("Last name = " + contact.LastName);
    Trace.WriteLine("Address = " + (contact.Address ?? "Missing"));
    ...
}
...
}
//////////////////////////// Client Side //////////////////////////////
[DataContract]
struct Contact
{
    [DataMember]
    public string FirstName;

    [DataMember]
    public string LastName;
}

Contact contact = new Contact()
{
    FirstName = "Juval",
    LastName = "Lowy"
};

ContactManagerClient proxy = new ContactManagerClient();
proxy.AddContact(contact);

proxy.Close();

```

The output of [Example 3-13](#) will be:

```

First name = Juval
Last name = Lowy
Address = Missing

```

because the service received `null` for the `Address` data member and coalesced the trace to `Missing`. The problem with [Example 3-13](#) is that you will have to manually compensate this way at every place the service (or any other service or client) uses this data contract.

Using the `OnDeserializing` event

When you do want to share your compensation logic across all parties using the data contract, it's better to use the `OnDeserializing` event to initialize potentially missing data members based on some local heuristic. If the message contains values for those members, they will override your settings in the `OnDeserializing` event. If it doesn't, the event-handling method provides some nondefault values. Using the technique shown here:

```

[DataContract]
struct Contact
{
    [DataMember]
    public string FirstName;

    [DataMember]
    public string LastName;

    [DataMember]
    public string Address;

    [OnDeserializing]
    void OnDeserializing(StreamingContext context)
    {
        Address = "Some default address";
    }
}

```

the output of [Example 3-13](#) will be:

```

First name = Juval
Last name = Lowy
Address = Some default address

```

Required members

Unlike ignoring new members, which for the most part is benign, the default handling of missing members may very likely cause the receiving side to fail further down the call chain, because the missing members may be essential for correct operation. This may have disastrous results. You can instruct WCF to avoid invoking the operation and to fail the call if a data member is missing by setting the `IsRequired` property of the `DataMember` attribute to `true`:

```

[DataContract]
struct Contact
{
    [DataMember]
    public string FirstName;

    [DataMember]
    public string LastName;

    [DataMember(IsRequired = true)]
    public string Address;
}

```

The default value of `IsRequired` is `false`; that is, to ignore the missing member. When, on the receiving side, `DataContractSerializer` does not find the information required to deserialize a member marked as required in the message, it will abort the call, resulting in a `NetDispatcherFaultException` on the sending side. For instance,

if the data contract on the service side in [Example 3-13](#) were to mark the `Address` member as required, the call would not reach the service. The fact that a particular member is required is published in the service metadata, and when it is imported to the client, the generated proxy definition will also have that member as required.

Both the client and the service can mark some or all of the data members in their data contracts as required, completely independently of each other. The more members that are marked as required, the safer the interaction with a service or a client will be, but at the expense of flexibility and versioning tolerance.

When a data contract that has a required new member is sent to a receiving party that is not even aware of that member, such a call is actually valid and will be allowed to go through. In other words, if Version 2 (V2) of a data contract has a new member for which `IsRequired` is set to `true`, you can send V2 to a party expecting Version 1 (V1) that does not even have the member in the contract, and the new member will simply be ignored. `IsRequired` has an effect only when the V2-aware party is missing the member. Assuming that V1 does not know about a new member added by V2, [Table 3-1](#) lists the possible permutations of allowed or disallowed interactions as a product of the versions involved and the value of the `IsRequired` property.

Table 3-1. Versioning tolerance with required members

<code>IsRequired</code>	V1 to V2	V2 to V1
False	Yes	Yes
True	No	Yes

An interesting situation relying on required members has to do with serializable types. Since serializable types have no tolerance for missing members by default, the resulting data contract will have all data members as required when they are exported. For example, this `Contact` definition:

```
[Serializable]
struct Contact
{
    public string FirstName;
    public string LastName;
}
```

will have the metadata representation:

```
[DataContract]
struct Contact
{
    [DataMember(IsRequired = true)]
    public string FirstName
    {get;set;}

    [DataMember(IsRequired = true)]
```

```

    public string LastName
    {get;set;}
}

```

To set the same versioning tolerance regarding missing members as the `DataContract` attribute, apply the `OptionalField` attribute on the optional member. For example, this `Contact` definition:

```

[Serializable]
struct Contact
{
    public string FirstName;

    [OptionalField]
    public string LastName;
}

```

will have the metadata representation:

```

[DataContract]
struct Contact
{
    [DataMember(IsRequired = true)]
    public string FirstName
    {get;set;}

    [DataMember]
    public string LastName
    {get;set;}
}

```

Versioning Round-Trip

The versioning tolerance techniques discussed so far for ignoring new members and defaulting missing ones are suboptimal: they enable a point-to-point client-to-service call, but have no support for a wider-scope pass-through scenario. Consider the two interactions shown in [Figure 3-4](#).

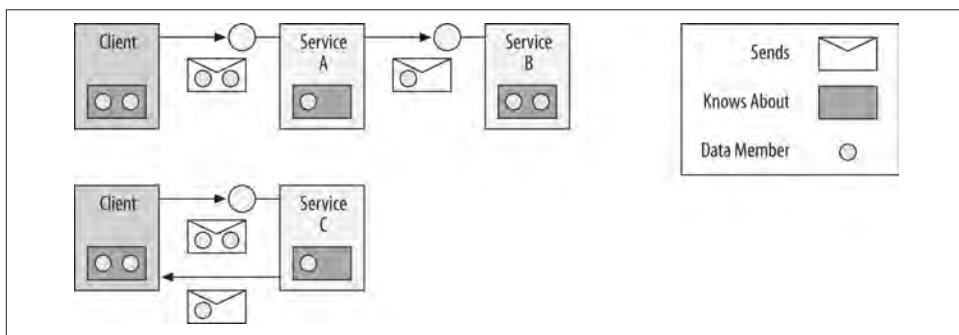


Figure 3-4. Versioning round-trip may degrade overall interaction

In the first interaction, a client that is built against a new data contract with new members passes that data contract to Service A, which does not know about the new members. Service A then passes the data to Service B, which is aware of the new data contract. However, the data passed from Service A to Service B does not contain the new members—they were silently dropped during deserialization from the client because they were not part of the data contract for Service A. A similar situation occurs in the second interaction, where a client that is aware of the new data contract with new members passes the data to Service C, which is aware only of the old contract that does not have the new members. The data Service C returns to the client will not have the new members.

This situation of new-old-new interaction is called a *versioning round-trip*. WCF supports handling of versioning round-trips by allowing a service (or client) with knowledge of only the old contract to simply pass through the state of the members defined in the new contract without dropping them. The problem is how to enable services/clients that are not aware of the new members to serialize and deserialize those unknown members without their schemas, and where to store them between calls. WCF's solution is to have the data contract type implement the **IExtensibleDataObject** interface, defined as:

```
public interface IExtensibleDataObject
{
    ExtensionDataObject ExtensionData
    {get;set;}
}
```

IExtensibleDataObject defines a single property of the type **ExtensionDataObject**. The exact definition of **ExtensionDataObject** is irrelevant, since developers never have to interact with it directly. **ExtensionDataObject** has an internal linked list of object references and type information, and that is where the unknown data members are stored. In other words, if the data contract type supports **IExtensibleDataObject**, when unrecognized new members are available in the message, they are deserialized and stored in that list. When the service (or client) calls out—passing the old data contract type, which now includes the unknown data members inside **ExtensionDataObject**—the unknown members are serialized out into the message in order. If the receiving side knows about the new data contract, it will get a valid new data contract without any missing members. [Example 3-14](#) demonstrates implementing and relying on **IExtensibleDataObject**. As you can see, the implementation is straightforward: just add an **ExtensionDataObject** automatic property with explicit interface implementation.

Example 3-14. Implementing IExtensibleDataObject

```
[DataContract]
class Contact : IExtensibleDataObject
```

```

{
    ExtensionDataObject IExtensibleDataObject.ExtensionData
    {get;set;}

    [DataMember]
    public string FirstName;

    [DataMember]
    public string LastName;
}

```

Schema compatibility

While implementing `IExtensibleDataObject` enables round-tripping, it has the downside of enabling a service that is compatible with one data contract schema to interact successfully with another service that expects another data contract schema. In some esoteric cases, the service may decide to disallow round-tripping and enforce its own version of the data contract on downstream services. Using the `ServiceBehavior` attribute (discussed at length in [Chapter 4](#)), services can instruct WCF to override the handling of unknown members by `IExtensibleDataObject` and ignore them even if the data contract supports `IExtensibleDataObject`. The `ServiceBehavior` attribute offers the Boolean property `IgnoreExtensionDataObject`, defined as:

```

[AttributeUsage(AttributeTargets.Class)]
public sealed class ServiceBehaviorAttribute : Attribute, ...
{
    public bool IgnoreExtensionDataObject
    {get;set;}
    //More members
}

```

The default value of `IgnoreExtensionDataObject` is `false`. Setting it to `true` ensures that all unknown data members across all data contracts used by the service will always be ignored:

```

[ServiceBehavior(IgnoreExtensionDataObject = true)]
class ContactManager : IContactManager
{...}

```

When you import a data contract using Visual Studio, the generated data contract type always supports `IExtensibleDataObject`, even if the original data contract did not. I believe that the best practice is to always have your data contracts implement `IExtensibleDataObject` and to avoid setting `IgnoreExtensionDataObject` to `true`. `IExtensibleDataObject` decouples the service from its downstream services, allowing them to evolve separately.

Data Contract Factoring and Design

Chapter 2 revealed many techniques, metrics, guidelines and even a mindset for service contract factoring. These tips and techniques help you to produce service contracts that are more specialized, loosely coupled, and reusable. The result is improving the design and maintainability of your service-oriented system. Yet at the same time an integral part of a comprehensive contract factoring strategy is carefully considering the impact data contract factoring will have on your overall system design.

Data Contract Design

You can think of a data contract as declaring the information payload for a service interaction, but not the interaction itself. As explained earlier, WCF uses data contracts for all the input parameters and return values of a service operation, even those that are primitive types. For these, WCF provides implicit data contracts. You may use custom data contracts to express the input parameters required by a service operation (the request), the values returned by a service operation (the response), or both. You can use data contracts alongside other primitive types within a service operation's inputs. And as you have seen, you can design your data contracts as hierarchies and then compose them into larger composite data contracts.

In WCF, when you use data contracts in a service operation, they become part of the service's contract definition. This means that when you change a data contract, you also change the contractual declaration of all the services that use them. The notable consequence is that when you change a service's contract you disrupt the ecosystem in which the service interacts. This disruption ultimately produces a cascade of additional maintenance, testing, and deployment for you and your team.

As you have learned throughout this chapter, WCF provides a number of techniques, such as data contract equivalence and data contract versioning that you can leverage when designing your service interactions to minimize or contain the impact of data contract changes. While these techniques help you to insulate the layers of your system from data contract changes, they do not inform you on how to promote decoupling between the service facets of your system or the operations within them.

Data Contract Factoring

As with service contract factoring, you should take a holistic, domain-informed approach to data contract factoring. This includes analyzing not only the associated service contracts, but the use cases they represent. Later in this section, you will see the important role use cases play in creating a consistent, pragmatic approach to data contract factoring. But first, you must ensure your approach to data contract factoring supports the intent of your service contract factoring strategy.

Consider the dog service example presented in [Chapter 2](#). The original IDog contract before service contract factoring was:

```
[ServiceContract]
interface IDog
{
    [OperationContract]
    void Fetch();

    [OperationContract]
    void Bark();

    [OperationContract]
    long GetVetClinicNumber();

    [OperationContract]
    void Vaccinate();
}
class PoodleService : IDog
{...}
```

As explained in [Chapter 2](#), the original operation `GetVetClinicNumber()` is a property-like operation, and does not comply with service-oriented design best practices. For the purpose of this chapter, suppose the contract has instead the behavioral `ScheduleAppointment()` operation:

```
[ServiceContract]
interface IDog
{
    [OperationContract]
    void Fetch();

    [OperationContract]
    void Bark();

    [OperationContract]
    void ScheduleAppointment();

    [OperationContract]
    void Vaccinate();
}
```

Suppose that you wished to introduce a single, canonical Dog data contract that would support all the required behaviors of a dog service such as barking, fetching, scheduling an appointment with the veterinarian and vaccinating the dog (see [Example 3-15](#)).

Example 3-15. A possible dog data contract

```
[DataContract]
class Dog
```

```

{
    [DataMember]
    public string ThingToFetch;

    [DataMember]
    public int InterestLevel;

    [DataMember]
    public ushort FetchTimeout;

    [DataMember]
    public BarkStyle BarkStyle;

    [DataMember]
    public Volume BarkVolume;

    [DataMember]
    public uint BarkLength;

    [DataMember]
    public string Name;

    [DataMember]
    public string OwnerName;

    [DataMember]
    public ulong LicenseId;

    [DataMember]
    public ushort Age;

    [DataMember]
    public ulong VetClinicNumber;

    [DataMember]
    public IEnumerable<Appointment> Appointments;

    [DataMember]
    public IEnumerable<Vaccination> Vaccinations;
}

enum BarkStyle
{
    Growl,
    Bark,
    Yelp
}

enum Volume
{
    Low,
    Normal,
}

```

```

    High
}

[DataContract]
struct Vaccination
{
    [DataMember]
    public DateTime Date;

    [DataMember]
    public string Description;
}

[DataContract]
struct Appointment
{
    [DataMember]
    public DateTime Date;

    [DataMember]
    public string ReasonforVisit;
}

```

With the updated IDog service contract:

```

[ServiceContract]
interface IDog
{
    [OperationContract]
    void Fetch(Dog dog);

    [OperationContract]
    void Bark(Dog dog);

    [OperationContract]
    void ScheduleAppointment(Dog dog);

    [OperationContract]
    void Vaccinate(Dog dog);
}

```

The dog service contract factoring example of [Chapter 2](#) determined that IDog was poorly factored due to the weak logical cohesion between its aspects as a pet and a dog. To solve this problem the example showed contract segregation as a factoring strategy. Applying the same approach to your updated IDog service contract results in the following contracts:

```

[ServiceContract]
interface IDog
{
    [OperationContract]
    void Fetch(Dog dog);
}

```

```

[OperationContract]
void Bark(Dog dog);
}

[ServiceContract]
interface IPet
{
    [OperationContract]
    void ScheduleAppointment(Dog dog);

    [OperationContract]
    void Vaccinate(Dog dog);
}

```

The implementation of the service is still as with [Chapter 2](#):

```

class PoodleService : IDog,IPet
{...}

```

This naïve use of the Dog data contract in both service contracts at first glance looks satisfactory, as it addresses the needs of any dog service by the original requirements. But when you analyze these new, more succinct service contracts, the use cases their operations represent and your monolithic Dog data contract, you see clearly something is wrong. IPet is now poorly factored. Its Dog data contract has members that have nothing to do with IPet. Similarly, IDog is also poorly factored as it has members required only by IPet. You have now coupled IDog and IPet to each other throughout their lifetimes via the Dog data contract. It would also be wrong to add a cat to the system, and expect it to support IPet with a Dog contract.

During your initial analysis, one thing should have stood out immediately. You have misplaced the use of the canine dog in relation to the clinical activities for a pet. Simply by name, it does not look right. The coupling force that allows you to do this in your design is the low cohesion of the Dog data contract.

To solve the problem, your first inclination might be to factor out a Pet data contract from Dog and create a hierarchical relationship between them. You may also decide to make Pet abstract, to ensure you only have concrete instances of pets running around. On the surface, this approach looks as if it will allow you to solve multiple problems. You can broaden the usage of IPet to accept all pets while allowing you to still pass dogs around.

Such a hierarchical refactoring of the Dog data contract would look like:

```

[DataContract]
[KnownType(typeof(Dog))]
abstract class Pet
{...}

[DataContract]

```

```

class Dog : Pet
{...}

[ServiceContract]
interface IDog
{...}

[ServiceContract]
interface IPet
{
    [OperationContract]
    void ScheduleAppointment(Pet pet);

    [OperationContract]
    void Vaccinate(Pet pet);
}

```

But when you analyze the result of your refactoring, you notice you have not solved the problem at all. The inheritance of `Pet` and the required known type association of `Dog` still couple together the canine and clinical facets of your system through their data contracts.

The problem becomes even more acute as your system evolves. Attempting to add cats to your menagerie reveals that service facet coupling will actually increase with this design as you introduce new concepts.

```

[DataContract]
[KnownType(typeof(Dog))]
[KnownType(typeof(Cat))]
abstract class Pet
{...}

[DataContract]
class Cat : Pet
{...}

```

There is yet another subtle, but significant problem that will arise from this approach; the unintentional smearing of business logic across your system. Your data contract design does not support the original intent of your service contract refactoring. The intent was to promote the autonomy of the canine, feline and clinical facets of your system by decoupling them. In this design, a `Dog` or `Cat` is also a `Pet`. By design, you are prescribing that developers must pass dogs from your canine facet and cats from your feline facet into your system's clinical facet.

Think about a possible use case for vaccination. If the use case suggests that the act of vaccination operate against a clinical `Pet`, but your design stipulates that developers may only pass a `Dog` or `Cat` to `Vaccinate()`, then there is nothing stopping a developer (other than code review) from adding logic to `Vaccinate()` that is `Dog` or `Cat` specific. Once this occurs, you may find that developers have irrevocably changed the

logic within the `Vaccinate()` operation. You have allowed developers to intimately couple the clinical facet of your system to forms of `Dog` and `Cat` more geared toward canine and feline behavior and not at all related to clinical activity.

Removing the abstract modifier from `Pet` will not improve your design. Nothing is stopping wayward developers from passing dogs and cats into your clinical facet. And the potential for coupling is even worse within the canine behaviors of `Bark()` and `Fetch()`. Clinical information is readily available through the inheritance of `Pet`, allowing developers to errantly introduce clinical logic into canine behaviors.

Through poor data contract design, you have undone all the value of your original service contract factoring and caused your overall system design to suffer. Changes to `Dog`, `Cat` or any other specialization you introduce will now resonate across the clinical facet of your system. Likewise, changes to `Pet` will resonate across your canine and feline service facets, not just during the initial refactor, but throughout their lifetimes.

Instead of a hierarchical approach, in this scenario you should apply the same contract factoring strategy to your data contracts as you did to your service contracts. The final service contract refactoring in [Chapter 2](#) explicitly did not use a hierarchical strategy between `IDog` and `IPet`, because not all dogs are pets. By recognizing this important distinction, the contract design ([Example 3-16](#)) sought to improve the logical relation between contract operations through segregation.

Example 3-16. The complete dog contract

```
[DataContract]
class Dog
{
    [DataMember]
    public string ThingToFetch;

    [DataMember]
    public int InterestLevel;

    [DataMember]
    public ushort FetchTimeout;

    [DataMember]
    public BarkStyle BarkStyle;

    [DataMember]
    public Volume BarkVolume;

    [DataMember]
    public uint BarkLength;
}
```

```

[DataContract]
class Pet
{
    [DataMember]
    public string Name;

    [DataMember]
    public string OwnerName;

    [DataMember]
    public ulong LicenseId;

    [DataMember]
    public ushort Age;

    [DataMember]
    public ulong VetClinicNumber;

    [DataMember]
    public IEnumerable<Appointment> Appointments;

    [DataMember]
    public List<Vaccination> Vaccinations;
}

enum BarkStyle
{
    Growl,
    Bark,
    Yelp
}

enum Volume
{
    Low,
    Normal,
    High
}

[DataContract]
struct Vaccination
{
    [DataMember]
    public DateTime Date;

    [DataMember]
    public string Description;
}

[DataContract]
struct Appointment
{

```

```

[DataMember]
public DateTime Date;

[DataMember]
public string ReasonforVisit;
}

[ServiceContract]
interface IDog
{
    [OperationContract]
    void Fetch(Dog dog);

    [OperationContract]
    void Bark(Dog dog);
}

[ServiceContract]
interface IPet
{
    [OperationContract]
    void ScheduleAppointment(Pet pet);

    [OperationContract]
    void Vaccinate(Pet pet);
}

class PoodleService : IDog,IPet
{
    public void Fetch(Dog dog)
    {...}

    public void Bark(Dog dog)
    {...}

    public void ScheduleAppointment(Pet pet)
    {...}

    public void Vaccinate(Pet pet)
    {...}
}

```

Because you have promoted the autonomy of the clinical facet of your system and protected against the errant coupling of service facets through misplaced business logic, your new design would more easily support the possibility of a much more dramatic change in requirements, a clinical service:

```

class ClinicService : IPet
{
    public void ScheduleAppointment(Pet pet)
    {...}
}

```

```
    public void Vaccinate(Pet pet)
    {...}
}
```

In this case, by using the same contract factoring strategy for both service and data contracts, you preserve and extend the value of your original intent and ensure through your design that the service facets of your system remain decoupled and autonomous.

Service Facet Coupling

As your service-oriented system matures, you must continually guard against coupling together the service facets of your system. Once you have coupled one service facet to another, untangling them can be very costly and impractical, if not impossible. You have now coupled them throughout their lifetimes, something you might call *lifetime coupling*.

As explained in [Appendix A](#), a tenet of service orientation is that services are autonomous. Put another way, system service facets are autonomous. Changes to one facet should not affect another. You should endeavor to insulate your system's service facets from one another as much as possible. Doing so will afford your systems greater flexibility, resiliency, and agility throughout their lifecycle.

The observation that a hierarchical factoring strategy was not appropriate in the previous example does not mean it is not valuable. [Chapter 2](#) presents a second service contract factoring example involving mammalian behaviors where the requirements warrant a hierarchical technique. This example shows the value of grouping operations common to multiple, unrelated service contracts into a base contract reusable by all.

The original mammal hierarchical service contract factoring:

```
[ServiceContract]
interface IMammal
{
    [OperationContract]
    void Shed();

    [OperationContract]
    void Lactate();
}

[ServiceContract]
interface IDog : IMammal
{...}

[ServiceContract]
```

```
interface ICat : IMammal
{...}
```

A possible mammal data contract design might be:

```
[DataContract]
[KnownType(typeof(Dog))]
[KnownType(typeof(Cat))]
class Mammal
{
    [DataMember]
    public int Coat;

    [DataMember]
    public int Gender;
}

[DataContract]
class Dog : Mammal
{...}

[DataContract]
class Cat : Mammal
{...}
```

With the updated service contract:

```
[ServiceContract]
interface IMammal
{
    [OperationContract]
    void Shed(Mammal mammal);

    [OperationContract]
    void Lactate(Mammal mammal);
}

[ServiceContract]
interface IDog : IMammal
{...}

[ServiceContract]
interface ICat : IMammal
{...}
```



Instead of the multiple KnownType attributes, you should prefer one of the approaches based on the GenericResolver.

The level of coupling inherent in this design is acceptable, even necessary. Unlike the canine/clinic example from before, mammal operations are logically related to their respective contracts. All dogs and cats are mammals. This makes mammal behaviors common and pervasive. A change in the behavior of a mammal necessitates a required change to all.

Another scenario that necessitates a hierarchical data contract factoring strategy is the need for dog-specific information within new or evolving clinical use cases. You must be able to accommodate these new requirements in your data contract design without reintroducing coupling between the canine and clinical facets of your system. In this case, you should not introduce these new properties into your canine Dog data contract, but instead create a new clinical Dog data contract that inherits from Pet:

```
[DataContract]
[KnownType(typeof(Dog))]
class Pet
{...}

[DataContract]
class Dog : Pet
{
    [DataMember]
    public bool RequiresMuzzle;

    [DataMember]
    public bool KennelCough;
}
```

Here a hierarchical factoring strategy is appropriate. This Dog data contract is specific to clinical facet needs and preserves the autonomy of your system's service facets.

Diverging versus Converging Service Facets

If two service facets are indeed autonomous, you will notice that as they evolve the behavioral and informational needs of their use cases will diverge and become more differentiated over time. This observation confirms your original decision to use a strategy of segregation to factor out these facets and to also keep their data needs separate.

On the other hand, if you notice that their behavioral and informational needs are becoming more and more similar over time, they are beginning to converge, probably because the separate factoring was questionable in the first place. You should endeavor to collapse and merge the converging service facets of your system whenever possible. You are no longer gaining value by keeping them segregated.

Operation Factoring

So far this section has presented a few data contract design techniques focused on helping you to promote the decoupling and preserve the autonomy of your system's service facets. You can further promote the decoupling of the service facets and layers of your system by analyzing the use cases that service contract operations represent.

As discussed in [Chapter 2](#), well-factored service contracts often possess a relatively small number of operations that focus on business use cases. In service-oriented systems, you can observe that not all use cases evolve at the same rate. Some change much more frequently than others. You may also notice that logic changes within a use case often precipitate new data needs. If your contract design prescribes that all the operations within your service contract must share a common data contract, you potentially allow frequently changing use cases to affect other use cases as well.

Consider again the canine facet contracts of [Example 3-16](#). It is clear in this simple example that the Dog data contract is not cohesive. It contains properties related to multiple use cases. This may not present a problem if system needs require that the barking and fetching use cases change in unison. But if these use cases change separately, you will find yourself constantly managing these changes across all the use cases of your service facet.

An additional problem is that this canine Dog data contract does not clearly convey intent. It is not evident to the consumers of IDog which properties are relevant in any given use case. Applying a prefix somewhat improves the clarity of your design, such as the Bark prefixes in the Dog data contract of [Example 3-16](#), but it will be very difficult for you to enforce this policy across your team and the approach is not type safe. A prefix also does nothing to help you improve the decoupling of your service use cases.



You should consider prefixes or suffixes on service contract operations, data contract properties, or naming in general as a design *smell*, a sign that there is a problem in your contract design.

To improve the cohesion of your data contracts, whenever possible you should constrain your data contracts to the specific needs of each use case.

Data contracts for IDog factored by use case:

```
[DataContract]
class BarkRequest
{
    [DataMember]
    public BarkStyle Style;
```

```

[DataMember]
public Volume Volume;

[DataMember]
public uint Length;
}

[DataContract]
class FetchRequest
{
    [DataMember]
    public string ThingToRetrieve;

    [DataMember]
    public int InterestLevel;

    [DataMember]
    public ushort Timeout;
}

[ServiceContract]
interface IDog
{
    [OperationContract]
    void Fetch(FetchRequest request);

    [OperationContract]
    void Bark(BarkRequest request);
}

```

As an additional option, if your operations return data, you can promote the greatest degree of use case decoupling by providing response data contracts for them as well.

Chunky, not Chatty

In service-oriented design, we recommend making your service interactions *chunky*, not *chatty*. This means you should endeavor to design your service operations in such a way as to require the fewest number of service calls to satisfy the needs of a given use case. Ideally, executing a given use case should only require a single service call from the initiating client

To support this approach, for use cases that require data retrieval or submission, you should design your service operations to gather or receive in a single service call as much of the data needed by the use case as possible. When following this rule of thumb, you will often find that the data contracts associated with these more coarse-grained service operations will be larger (chunkier) and that the number of service interactions your contract design requires will be fewer (less chatty).

The Data Transfer Object Pattern

Service orientation provides no mechanism for you to share logic with another platform or technology. If your data contract design relies on logic to manipulate its state and that logic is internal to the classes implementing your data contracts, it may present a problem to developers who wish to consume your services from other platforms or code bases. Those developers will not have access to your original types and therefore no idea that these additional logic requirements exist. At best, you will have to convey these implementation details each time manually. At worst, developers from other platforms will be unable to create data contracts that comply with the specific needs of your system.

If the layers of your system do share a common platform and technology, you can certainly share the classes that implement your data contracts throughout your system. But if you choose to also include logic in these classes, you run the risk of once again errantly coupling the layers of your system through logic that may need to change independently at each layer over time.

It is therefore a best practice to constrain your data contracts to simple classes that contain only primitive types or other data contracts as members, but no logic. This guidance is compatible with the well-known object-oriented design pattern, the Data Transfer Object (DTO). The original intent of the DTO pattern was to use simple data structures to improve the efficiency of remote calls and to insulate the layers or tiers of an application from each other. When employing the pattern, developers promoted decoupling between the layers of their application by ensuring the DTOs they created were unique and separate from the data structures used internally by the business logic. This allowed them to change the logic within each layer without the changes always resonating from layer to layer or even throughout the entire application.

By now, this should all sound very familiar to you. True to its original intent of isolating application layers and optimizing data transfer, the DTO pattern is equally valuable in service-oriented design as it was in object-oriented design, perhaps even more so. The DTO pattern provides you with a solid starting point from which you can craft your own approach to data contract design.

Pragmatic pattern usage

As with all good things, you should use judgment when applying design patterns. Taking the DTO pattern to its ultimate conclusion would require you to produce a unique data contract for the request and response of every single service operation of every service facet in your entire system. In theory, this would afford your design the greatest degree of decoupling. In practice, the effort required to create and maintain all these data contracts is often excessive.

The highly valuable graph presented in [Figure 2-1](#), which compares count versus cost, applies to data contract factoring as well. For a given set of use cases, you should aspire to have a balanced number of data contracts (i.e., neither extreme of the possible number of data contracts for your system is good). Often, providing a single, canonical data contract or data contract hierarchy for your entire system will quickly lead to a suffocating level of coupling. Likewise, it makes little sense for you to produce hundreds of single property data contracts just to abide by the guidance of a pattern. To help you converge on the proper level of granularity, here are a few observations from large-scale systems that you can use to guide you in creating a consistent, pragmatic approach to data contract factoring.

If your design calls for many small data contracts or a few very large ones, this indicates that there may be a problem in your service contract factoring. Your service contract operations may be too small or too coarse. Or there are additional service facets within your service contracts that you have not yet recognized to factor out. Refer to [Chapter 2](#) for practical insights on improving your service contract design. You may also find it useful to first design and factor your service contracts. Then design and factor your data contracts so that they maximize the original intent of your service contract factoring strategy.

Rarely do you need DTOs for both the request and response of a service operation. When deciding whether or not to use the full expression of the DTO pattern, you should always consider the needs of the use case encapsulated by the operation. For that, you must rely on your expertise within your given domain and your observations from past experience.

If use case logic or data will change rarely, the risk of coupling your service facets or system layers together is low. By recognizing this, you can improve the overall maintainability of your system and reduce the number of DTOs you need by sharing existing data contracts between use cases and system layers. In many cases, simply using primitive types for operation parameters or return values will suffice.

On the other hand, if use case logic will change frequently, you should always use unique DTOs for both the request and response of the operations that encapsulate them. This approach will give you the best opportunity to insulate the other use cases of your service facet, as well as the layers of your system, from these frequently changing use cases. By using a pragmatic approach to the DTO pattern, you can lower the effort required to maintain your system, while still promoting the decoupling and preserving the autonomy of your system's service facets.

Enumerations

Enumerations are always, by definition, serializable. When you define a new enum, there is no need to apply the `DataContract` attribute to it, and you can use it freely in

a data contract, as shown in [Example 3-17](#). All the values in the enum will implicitly be included in the data contract.

Example 3-17. Using an enum in a data contract

```
enum ContactType
{
    Customer,
    Vendor,
    Partner
}

[DataContract]
struct Contact
{
    [DataMember]
    public ContactType ContactType;

    [DataMember]
    public string FirstName;

    [DataMember]
    public string LastName;
}
```

If you want to exclude certain enum values from the data contract, you need to first decorate the enum with the `DataContract` attribute, and then explicitly apply the `EnumMemberAttribute` to all enum values you want to include in the enum data contract. The `EnumMember` attribute is defined as:

```
[AttributeUsage(AttributeTargets.Field, Inherited = false)]
public sealed class EnumMemberAttribute : Attribute
{
    public string Value
    {get;set;}
}
```

Any enum value not decorated with the `EnumMember` attribute will not be part of the data contract for that enum. For example, this enum:

```
[DataContract]
enum ContactType
{
    [EnumMember]
    Customer,

    [EnumMember]
    Vendor,

    //Will not be part of data contract
}
```

```
    Partner  
}
```

will result in this wire representation:

```
enum ContactType  
{  
    Customer,  
    Vendor  
}
```

The other use for the `EnumMember` attribute is to alias certain enum values to an existing enum data contract using the `Value` property. For example, this enum:

```
[DataContract]  
enum ContactType  
{  
    [EnumMember(Value = "MyCustomer")]  
    Customer,  
  
    [EnumMember]  
    Vendor,  
  
    [EnumMember]  
    Partner  
}
```

will result in this wire representation:

```
enum ContactType  
{  
    MyCustomer,  
    Vendor,  
    Partner  
}
```

The effect the `EnumMember` attribute has is local to the party using it. When publishing the metadata (or when defining it on the client side), the resulting data contract has no trace of it, and only the final product is used.

Delegates and Data Contracts

All delegate definitions are compiled into serializable classes, so in theory your data contract types can contain delegates as member variables:

```
[DataContract]  
class MyDataContract  
{  
    [DataMember]  
    public EventHandler MyEvent;  
}
```

Or even as events (note the use of the `field` qualifier):

```
[DataContract]
class MyDataContract
{
    [field:DataMember]
    public event EventHandler MyEvent;
}
```

In practice, however, when the data contract refers to a custom delegate, the imported data contract will contain an invalid delegate definition. While you could manually fix that definition, the bigger problem is that when you serialize an object that has a delegate member variable, the internal invocation list of the delegates is serialized, too. In most cases, this is not the desired effect with services and clients, because the exact structure of the list is local to the client or the service and should not be shared across the service boundary. In addition, there are no guarantees that the target objects in the internal list are serializable or are valid data contracts. Consequently, sometimes the serialization will work, and sometimes it will fail.

The simplest way to avoid this pitfall is not to apply the `DataMember` attribute on delegates. If the data contract is a serializable type, you need to explicitly exclude the delegate from the data contract:

```
[Serializable]
public class MyDataContract
{
    [NonSerialized]
    public EventHandler MyEvent;
}
```

Generics

You cannot define WCF contracts that rely on generic type parameters. Generics are specific to .NET, and using them would violate the service-oriented nature of WCF. However, you can use bounded generic types in your data contracts, as long as you specify the type parameters in the service contract and as long as the specified type parameters have valid data contracts, as shown in [Example 3-18](#).

Example 3-18. Using bounded generic types

```
[DataContract]
class MyClass<T>
{
    [DataMember]
    public T MyMember;
}

[ServiceContract]
interface IMyContract
{
```

```
[OperationContract]
void MyMethod(MyClass<int> obj);
}
```

When you import the metadata of a data contract such as the one in [Example 3-18](#), the imported types have all type parameters replaced with specific types, and the data contract itself is renamed to this format:

```
<Original name>Of<Type parameter names><hash>
```

Using the same definitions as in [Example 3-18](#), the imported data contract and service contract will look like this:

```
[DataContract]
class MyClassOfint
{
    int MyMemberField;

    [DataMember]
    public int MyMember
    {
        get
        {
            return MyMemberField;
        }
        set
        {
            MyMemberField = value;
        }
    }
}

[ServiceContract]
interface IMyContract
{
    [OperationContract]
    void MyMethod(MyClassOfint obj);
}
```

If, however, the service contract were to use a custom type such as `SomeClass` instead of `int`:

```
[DataContract]
class SomeClass
{...}

[DataContract]
class MyClass<T>
{...}

[OperationContract]
void MyMethod(MyClass<SomeClass> obj);
```

the exported data contract might look like this:

```
[DataContract]
class SomeClass
{...}

[DataContract]
class MyClassOfSomeClassMTRdqN6P
{...}

[OperationContract(...)]
void MyMethod(MyClassOfSomeClassMTRdqN6P obj);
```

where MTRdqN6P is some quasi-unique hash of the generic type parameter and the containing namespace. Different data contracts and namespaces will generate different hashes. The hash is in place to reduce the overall potential for a conflict with another data contract that might use another type parameter with the same name. No hash is created for the implicit data contracts of the primitives when they are used as generic type parameters, since the type `int` is a reserved word and the definition of `MyClassOfint` is unique.

In most cases, the hash is a cumbersome over-precaution. You can specify a different name for the exported data contract by simply assigning it to the data contract's `Name` property. For example, given this service-side data contract:

```
[DataContract]
class SomeClass
{...}

[DataContract(Name = "MyClass")]
class MyClass<T>
{...}

[OperationContract]
void MyMethod(MyClass<SomeClass> obj);
```

the exported data contract will be:

```
[DataContract]
class SomeClass
{...}

[DataContract]
class MyClass
{...}

[OperationContract]
void MyMethod(MyClass obj);
```

However, by doing this, you run the risk of some ambiguity, since two different custom generic types will result in the same type name.

If you still want to combine the name of the generic type parameter with that of the data contract, use the {<number>} directive, where the number is the ordinal number of the type parameter. For example, given this service-side definition:

```
[DataContract]
class SomeClass
{...}

[DataContract(Name = "MyClassOf{0}{1}")]
class MyClass<T,U>
{...}

[OperationContract]
void MyMethod(MyClass<SomeClass,int> obj);
```

the exported definition will be:

```
[DataContract]
class SomeClass
{...}

[DataContract]
class MyClassOfSomeClassint
{...}

[OperationContract(...)]
void MyMethod(MyClassOfSomeClassint obj);
```



The number of type parameters specified is not verified at compile time. Any mismatch will yield a runtime exception.

Finally, you can append # after the number to generate the unique hash. For example, given this data contract definition:

```
[DataContract]
class SomeClass
{...}

[DataContract(Name = "MyClassOf{0}{#}")]
class MyClass<T>
{...}

[OperationContract]
void MyMethod(MyClass<SomeClass> obj);
```

the exported definition will be:

```
[DataContract]
class SomeClass
```

```

{...}

[DataContract]
class MyClassOfSomeClassMTRdqN6P
{...}

[OperationContract]
void MyMethod(MyClassOfSomeClassMTRdqN6P obj);

```

Collections

In .NET, a *collection* is any type that supports the `IEnumerable` or `IEnumerable<T>` interface. All of the built-in collections in .NET, such as the array, the list, and the stack, support these interfaces. A data contract can include a collection as a data member, and a service contract can define operations that interact with a collection directly. Because .NET collections are .NET-specific, WCF cannot expose them in the service metadata, yet because they are so useful, WCF offers dedicated marshaling rules for collections.

Whenever you define a service operation that uses the collection interfaces `IEnumerable<T>`, `IList<T>`, or `ICollection<T>`, the resulting metadata always uses an array. For example, this service contract definition and implementation:

```

[ServiceContract]
interface IContactManager
{
    [OperationContract]
    IEnumerable<Contact> GetContacts();
    ...
}
class ContactManager : IContactManager
{
    List<Contact> m_Contacts = new List<Contact>();

    public IEnumerable<Contact> GetContacts()
    {
        return m_Contacts;
    }
    ...
}

```

will still be exported as:

```

[ServiceContract]
interface IContactManager
{
    [OperationContract]
    Contact[] GetContacts();
}

```

Concrete Collections

If the collection in the contract is a concrete collection (not an interface) and is a serializable collection—that is, it is marked with the `Serializable` attribute but not with the `DataContract` attribute—WCF can normalize the collection automatically to an array of the collection’s type, provided the collection contains an `Add()` method with either one of these signatures:

```
public void Add(object obj); //Collection uses IEnumerable
public void Add(T item);    //Collection uses IEnumerable<T>
```

For example, consider this contract definition:

```
[ServiceContract]
interface IContactManager
{
    [OperationContract]
    void AddContact(Contact contact);

    [OperationContract]
    List<Contact> GetContacts();
}
```

The list class is defined as:

```
public interface ICollection<T> : IEnumerable<T>
{
    ...
}
public interface IList<T> : ICollection<T>
{
    ...
}
[Serializable]
public class List<T> : IList<T>
{
    public void Add(T item);
    //More members
}
```

Because it is a valid collection and it has an `Add()` method, the resulting representation of the contract will be:

```
[ServiceContract]
interface IContactManager
{
    [OperationContract]
    void AddContact(Contact contact);

    [OperationContract]
    Contact[] GetContacts();
}
```

That is, a `List<Contact>` is marshaled as a `Contact[]`. The service may still return a `List<Contact>`, and yet the client will interact with an array, as shown in [Example 3-19](#).

Example 3-19. Marshaling a list as an array

```
//////////////////////////// Service Side //////////////////////////////
[ServiceContract]
interface IContactManager
{
    [OperationContract]
    void AddContact(Contact contact);

    [OperationContract]
    List<Contact> GetContacts();
}
//Service implementation
class ContactManager : IContactManager
{
    List<Contact> m_Contacts = new List<Contact>();

    public void AddContact(Contact contact)
    {
        m_Contacts.Add(contact);
    }

    public List<Contact> GetContacts()
    {
        return m_Contacts;
    }
}
//////////////////////////// Client Side //////////////////////////////
[ServiceContract]
interface IContactManager
{
    [OperationContract]
    void AddContact(Contact contact);

    [OperationContract]
    Contact[] GetContacts();
}
class ContactManagerClient : ClientBase<IContactManager>, IContactManager
{
    public Contact[] GetContacts()
    {
        return Channel.GetContacts();
    }
    //More members
}
//Client code
ContactManagerClient proxy = new ContactManagerClient();
Contact[] contacts = proxy.GetContacts();
proxy.Close();
```

Note that while the collection must have the Add() method to be marshaled as an array, the collection need not implement the Add() method at all.

Custom Collections

It's not just the built-in collections that can be marshaled automatically as arrays—any custom collection can abide by the same prerequisites and be marshaled as an array, as shown in [Example 3-20](#). In this example, the collection `MyCollection<string>` is marshaled as a `string[]`.

Example 3-20. Marshaling a custom collection as an array

```
//////////////////////////// Service Side //////////////////////////////
[Serializable]
public class MyCollection<T> : IEnumerable<T>
{
    public void Add(T item)
    {}

    IEnumerator<T> IEnumerable<T>.GetEnumerator()
    {...}
    //Rest of the implementation
}

[ServiceContract]
interface IMyContract
{
    [OperationContract]
    MyCollection<string> GetCollection();
}

//////////////////////////// Client Side //////////////////////////////
[ServiceContract]
interface IMyContract
{
    [OperationContract]
    string[] GetCollection();
}
```

The CollectionDataContract Attribute

The mechanism shown so far for marshaling a concrete collection is suboptimal. For one thing, it requires the collection to be serializable and does not work with the service-oriented `DataContract` attribute. Also, while one party is dealing with a collection, the other is dealing with an array, and the two are not semantically equivalent: the collection is likely to offer some advantages, or it would not have been chosen in the first place. Furthermore, there is no compile-time or runtime verification of the presence of the `Add()` method or the `IEnumerable` and `IEnumerable<T>` interfaces, resulting in an unworkable data contract if they are missing. WCF's solution is yet another dedicated attribute called `CollectionDataContractAttribute`, defined as:

```

[AttributeUsage(AttributeTargets.Struct|AttributeTargets.Class,Inherited =
    false)]
public sealed class CollectionDataContractAttribute : Attribute
{
    public string Name
    {get;set;}

    public string Namespace
    {get;set;}
    //More members
}

```

The `CollectionDataContract` attribute is analogous to the `DataContract` attribute, and similarly, it does not make the collection serializable. When applied on a collection, the `CollectionDataContract` attribute exposes the collection to the client as a generic linked list. While the linked list may have nothing to do with the original collection, it does offer a more collection-like interface than an array.

For example, given this collection definition:

```

[CollectionDataContract(Name = "MyCollectionOf{0}")]
public class MyCollection<T> : IEnumerable<T>
{
    public void Add(T item)
    {}

    IEnumerator<T> IEnumerable<T>.GetEnumerator()
    {...}
    //Rest of the implementation
}

```

and this service-side contract definition:

```

[ServiceContract]
interface IContactManager
{
    [OperationContract]
    void AddContact(Contact contact);

    [OperationContract]
    MyCollection<Contact> GetContacts();
}

```

the definitions the client ends up with after importing the metadata will be:

```

[CollectionDataContract]
public class MyCollectionOfContact : List<Contact>
{}

[ServiceContract]
interface IContactManager
{
    [OperationContract]

```

```
void AddContact(Contact contact);

[OperationContract]
MyCollectionOfContact GetContacts();
}
```

In addition, at service load time the `CollectionDataContract` attribute verifies the presence of the `Add()` method as well as either `IEnumerable` or `IEnumerable<T>`. Failing to have these on the collection will result in an `InvalidOperationException`.

Note that you cannot apply both the `DataContract` attribute and the `CollectionDataContract` attribute on a collection. Again, this is verified at service load time.

Referencing a Collection

WCF even lets you preserve the same collection on the client side as on the service side. The advanced settings dialog box for the service reference (see [Figure 1-10](#)) contains a Collection type combo box that lets you specify how to represent to the client certain kinds of collections and arrays found in the service metadata. For example, if the service operation returns one of the collections `IEnumerable<T>`, `IList<T>`, or `ICollection<T>`, by default the proxy will present it as an array (the default item in the combo box). However, you can request Visual Studio to use another collection, such as `BindingList` for data binding, a `List<T>`, `Collection`, or `LinkedList<T>`, and so on. If a conversion is possible, the proxy will use the requested collection type instead of an array, for example:

```
[OperationContract]
List<int> GetNumbers();
```

When you then define the collection in another assembly referenced by the client's project, as I've discussed, that collection will be imported as-is. This feature is very useful when interacting with one of the built-in .NET collections, such as the `Stack<T>` collection defined in the `System.dll`, which is referenced by practically all .NET projects.

Dictionaries

A *dictionary* is a special type of collection that maps one data instance to another. As such, dictionaries do not marshal well either as arrays or as lists. Not surprisingly, dictionaries get their own representation in WCF.

If the dictionary is a serializable collection that supports the `IDictionary` interface, it will be exposed as a `Dictionary<object,object>`. For example, this service contract definition:

```
[Serializable]
public class MyDictionary : IDictionary
{...}
```

```

[ServiceContract]
interface IContactManager
{
    ...
    [OperationContract]
    MyDictionary GetContacts();
}

```

will be exposed as this definition:

```

[ServiceContract]
interface IContactManager
{
    ...
    [OperationContract]
    Dictionary<object,object> GetContacts();
}

```

This, by the way, includes using the `HashTable` collection.

If the serializable collection supports the `IDictionary<K,T>` interface, as in:

```

[Serializable]
public class MyDictionary<K,T> : IDictionary<K,T>
{...}

[ServiceContract]
interface IContactManager
{
    ...
    [OperationContract]
    MyDictionary<int,Contact> GetContacts();
}

```

the exported representation will be as a `Dictionary<K,T>`:

```

[ServiceContract]
interface IContactManager
{
    ...
    [OperationContract]
    Dictionary<int,Contact> GetContacts();
}

```

This includes making direct use of `Dictionary<K,T>` in the original definition, instead of `MyDictionary<K,T>`.

If instead of merely being a serializable collection, the dictionary is decorated with the `CollectionDataContract` attribute, it will be marshaled as a subclass of the respective representation. For example, this service contract definition:

```

[CollectionDataContract]
public class MyDictionary : IDictionary

```

```

    {...}

    [ServiceContract]
    interface IContactManager
    {
        ...
        [OperationContract]
        MyDictionary GetContacts();
    }

```

will have this representation:

```

    [CollectionDataContract]
    public class MyDictionary : Dictionary<object,object>
    {}

    [ServiceContract]
    interface IContactManager
    {
        ...
        [OperationContract]
        MyDictionary GetContacts();
    }

```

while this generic collection:

```

    [CollectionDataContract(Name = "MyDictionary")]
    public class MyDictionary<K,T> : IDictionary<K,T>
    {...}

    [ServiceContract]
    interface IContactManager
    {
        ...
        [OperationContract]
        MyDictionary<int,Contact> GetContacts();
    }

```

will be published in the metadata as:

```

    [CollectionDataContract]
    public class MyDictionary : Dictionary<int,Contact>
    {}

    [ServiceContract]
    interface IContactManager
    {
        ...
        [OperationContract]
        MyDictionary GetContacts();
    }

```

As for a collection, in the advanced settings dialog box for a service reference (see [Figure 1-10](#)), you can request other dictionary types, such as the `SortedDictionary`

`ary<T,K>`, `HashTable`, or `ListDictionary` type, and the proxy will use that dictionary instead if possible.

Instance Management

Instance management is my name for the set of techniques WCF uses to bind client requests to service instances, governing which service instance handles which client request, and when. Instance management is necessary because of the extent to which applications differ in their needs for scalability, performance, throughput, durability, transactions, and queued calls—there simply isn’t a one-size-fits-all solution. However, there are a few canonical instance management techniques that are applicable across the range of applications, thus enabling a wide variety of scenarios and programming models. These techniques are the subject of this chapter, and understanding them is essential to developing scalable and consistent applications. WCF supports three types of instance activation: *per-call services* allocate (and destroy) a new service instance for each client request; *sessionful services* allocate a service instance for each client connection; and *singleton services* share the same service instance for all clients, across all connections and activations. This chapter provides the rationale for each of these instance management modes, and offers guidelines on when and how to best use them. It also addresses some related topics, such as behaviors, contexts, demarcating operations, instance deactivation, durable services, and throttling.¹

Behaviors

By and large, the service instance mode is strictly a service-side implementation detail that should not manifest itself on the client side in any way. To support that and a few other local service-side aspects, WCF defines the notion of *behaviors*. A behavior is a

¹ This chapter contains excerpts from my articles “WCF Essentials: Discover Mighty Instance Management Techniques for Developing WCF Apps” (*MSDN Magazine*, June 2006) and “Managing State with Durable Services” (*MSDN Magazine*, October 2008).

local attribute of the service or the client that does not affect its communication patterns. Clients should be unaware of service behaviors, and behaviors do not manifest themselves in the service's binding or published metadata. You have already seen two service behaviors in the previous chapters: [Chapter 1](#) uses the service metadata behavior to instruct the host to publish the service's metadata over HTTP-GET or to implement the MEX endpoint, and [Chapter 3](#) uses the service behavior to ignore the data object extension. No client can ever tell simply by examining the communication and the exchanged messages if the service is ignoring the data object extension or who published its metadata.

WCF defines two types of declarative service-side behaviors, governed by two corresponding attributes. The `ServiceBehaviorAttribute` is used to configure *service behaviors*; that is, behaviors that affect all endpoints (all contracts and operations) of the service. Apply the `ServiceBehavior` attribute directly on the service implementation class.

Use the `OperationBehaviorAttribute` to configure *operation behaviors*; that is, behaviors that affect only the implementation of a particular operation. The `OperationBehavior` attribute can be applied only on a method that implements a contract operation, never on the operation definition in the contract itself. You will see the use of `OperationBehavior` attribute later in this chapter and in subsequent chapters as well.

In the context of this chapter, the `ServiceBehavior` attribute is used to configure the service instance mode. As shown in [Example 4-1](#), the attribute defines the `InstanceContextMode` property of the enum type `InstanceContextMode`. The value of the `InstanceContextMode` enum controls which instance mode is used for the service.

Example 4-1. ServiceBehaviorAttribute used to configure the instance context mode

```
public enum InstanceContextMode
{
    PerCall,
    PerSession,
    Single
}
[AttributeUsage(AttributeTargets.Class)]
public sealed class ServiceBehaviorAttribute : Attribute, ...
{
    public InstanceContextMode InstanceContextMode
    {get;set;}
    //More members
}
```

The enum is correctly called `InstanceContextMode` rather than `InstanceMode` because it actually controls the instantiation mode of the context hosting the instance,

rather than that of the instance itself (recall from [Chapter 1](#) that the instance context is the innermost execution scope of the service). By default, however, the instance and its context are treated as a single unit, so the enum does control the life of the instance as well. You will see later in this chapter and in subsequent chapters how (and when) you can disengage the two, and for what purposes.

Per-Call Services

When the service type is configured for *per-call activation*, a service instance (the CLR object) exists only while a client call is in progress. Every client request (that is, a method call on the WCF contract) gets a new dedicated service instance. The following list explains how per-call activation works, and the steps are illustrated in [Figure 4-1](#):

1. The client calls the proxy and the proxy forwards the call to the service.
2. WCF creates a new context with a new service instance and calls the method on it.
3. When the method call returns, if the object implements `IDisposable`, WCF calls `IDisposable.Dispose()` on it. WCF then destroys the context.
4. The client calls the proxy and the proxy forwards the call to the service.
5. WCF creates an object and calls the method on it.

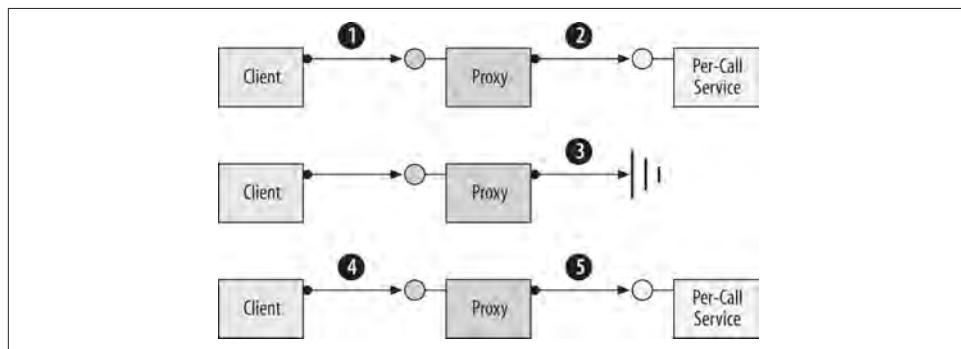


Figure 4-1. Per-call instantiation mode

Disposing of the service instance is an interesting point. As I just mentioned, if the service supports the `IDisposable` interface, WCF will automatically call the `Dispose()` method, allowing the service to perform any required cleanup. Note that `Dispose()` is called on the same thread that dispatched the original method call, and that `Dispose()` has an operation context (presented later). Once `Dispose()` is called, WCF disconnects the instance from the rest of the WCF infrastructure, making it a candidate for garbage collection.

Benefits of Per-Call Services

In the classic client/server programming model, using languages such as C++ or C#, every client gets its own dedicated server object. The fundamental problem with this approach is that it doesn't scale well. Imagine an application that has to serve many clients. Typically, these clients create the objects they need when the client application starts and dispose of them when the client application shuts down. What impedes scalability with the client/server model is that the client applications can hold onto objects for long periods of time, while actually using them for only a fraction of that time. Those objects may hold expensive or scarce resources, such as database connections, communication ports, or files. If you allocate an object for each client, you will tie up such crucial and/or limited resources for long periods, and you will eventually run out of resources.

A better activation model is to allocate an object for a client only while a call is in progress from the client to the service. That way, you have to create and maintain in memory only as many objects as there are concurrent calls, not as many objects as there are outstanding clients. My personal experience indicates that in a typical Enterprise system, especially one that involves users, at most 1% of all clients make concurrent calls (in a high-load Enterprise system, that figure rises to 3%). Thus, if your system can concurrently sustain 100 expensive service instances, it can still typically serve as many as 10,000 outstanding clients. This is precisely the benefit the per-call instance activation mode offers. In between calls, the client holds a reference on a proxy that doesn't have an actual object at the end of the wire. This means that you can dispose of the expensive resources the service instance occupies long before the client closes the proxy. By that same token, acquiring the resources is postponed until they are actually needed by a client.

Keep in mind that creating and destroying a service instance repeatedly on the service side without tearing down the connection to the client (with its client-side proxy) is a lot cheaper than repeatedly creating an instance and a connection. The second benefit is that forcing the service instance to reallocate or connect to its resources on every call caters very well to transactional resources and transactional programming (discussed in [Chapter 7](#)), because it eases the task of enforcing consistency with the instance state. The third benefit of per-call services is that they can be used in conjunction with queued disconnected calls (described in [Chapter 9](#)), because they allow easy mapping of service instances to discrete queued messages.

Configuring Per-Call Services

To configure a service type as a per-call service, you apply the `ServiceBehavior` attribute with the `InstanceContextMode` property set to `InstanceContextMode.PerCall`:

```

[ServiceContract]
interface IMyContract
{...}

[ServiceBehavior(InstanceContextMode = InstanceContextMode.PerCall)]
class MyService : IMyContract
{...}

```

Example 4-2 lists a simple per-call service and its client. As you can see from the program output, for each client method call a new service instance is constructed.

Example 4-2. Per-call service and client

```

/////////// Service Code ///////////
[ServiceContract]
interface IMyContract
{
    [OperationContract]
    void MyMethod();
}
[ServiceBehavior(InstanceContextMode = InstanceContextMode.PerCall)]
class MyService : IMyContract, IDisposable
{
    int m_Counter = 0;

    MyService()
    {
        Trace.WriteLine("MyService.MyService()");
    }
    public void MyMethod()
    {
        m_Counter++;
        Trace.WriteLine("Counter = " + m_Counter);
    }
    public void Dispose()
    {
        Trace.WriteLine("MyService.Dispose()");
    }
}
/////////// Client Code ///////////
MyContractClient proxy = new MyContractClient();

proxy.MyMethod();
proxy.MyMethod();

proxy.Close();

//Possible output
MyService.MyService()
Counter = 1
MyService.Dispose()
MyService.MyService()

```

```
Counter = 1
MyService.Dispose()
```

Per-Call Services and Transport Sessions

The use of a per-call service is independent from the presence of a transport session (described in [Chapter 1](#)). A transport session correlates all messages from a particular client to a particular channel. If the service is configured for per-call instantiation, there can still be a transport session, but for every call WCF will create a new context used just for that call. If transport-level sessions are not used, as you will see later, the service always behaves as a per-call service, regardless of its configuration.

If the per-call service has a transport session, communication from the client is subjected to the inactivity timeout of the transport session (which defaults to 10 minutes). Once the timeout has expired, the client can no longer use the proxy to invoke operations on the per-call service, since the transport session has ended.

The biggest effect transport sessions have on per-call services is that when the service is configured for single-threaded access (the WCF default, explained in [Chapter 8](#)), the transport session enforces a lock-step execution, where calls to the per-call service from the same proxy are serialized. That is, even if the client issues the calls concurrently, they are executed against different instances, one at a time, in order. This has particular implications for disposing of the instance. WCF does not block the client while it disposes of the service instance. However, if during the call to `Dispose()` the client has issued a second call, that call will be allowed to access a new instance only after `Dispose()` has returned. For example, the output at the end of [Example 4-2](#) represents a case where there is a transport session, since the second call can only execute once `Dispose()` has returned. If [Example 4-2](#) had no transport session, you might end up with the same output but also an out-of-order invocation where `Dispose()` is nonblocking, such as:

```
MyService.MyService()
Counter = 1
MyService.MyService()
Counter = 1
MyService.Dispose()
MyService.Dispose()
```

Designing Per-Call Services

Although in theory you can use the per-call instance activation mode on any service type, in practice you need to design the service and its contracts to support this mode from the ground up. The main problem is that the client doesn't know it's getting a new instance each time it makes a call. Per-call services must be *state-aware*; that is, they must proactively manage their state, giving the client the illusion of a continuous session. A state-aware service isn't the same as a stateless service. In fact, if the per-call

service were truly stateless, there would be no need for per-call activation in the first place. It is precisely because it has state, and an expensive state at that, that you need the per-call mode. An instance of a per-call service is created just before every method call and is destroyed immediately after each call. Therefore, at the beginning of each call, the object should initialize its state from values saved in some storage, and at the end of the call it should return its state to the storage. Such storage is typically either a database or the file system, but volatile storage (e.g., static variables) may be used instead.

Not all of the object's state can be saved as-is, however. For example, if the state contains a database connection, the object must reacquire the connection at construction or at the beginning of every call and dispose of the connection at the end of the call or in its implementation of `IDisposable.Dispose()`.

Using the per-call instance mode has one important implication for operation design: every operation must include a parameter to identify the service instance whose state needs to be retrieved. The instance uses that parameter to get *its* state from the storage, and not the state of another instance of the same type. Consequently, state storage is typically keyed (for example, as a static dictionary in memory or a database table). Examples of such state parameters are the account number for a bank account service, the order number for an order-processing service, and so on.

Example 4-3 shows a template for implementing a per-call service.

Example 4-3. Implementing a per-call service

```
[DataContract]
class Param
{...}

[ServiceContract]
interface IMyContract
{
    [OperationContract]
    void MyMethod(Param stateIdentifier);
}

[ServiceBehavior(InstanceContextMode = InstanceContextMode.PerCall)]
class MyPerCallService : IMyContract, IDisposable
{
    public void MyMethod(Param stateIdentifier)
    {
        GetState(stateIdentifier);
        DoWork();
        SaveState(stateIdentifier);
    }
    void GetState(Param stateIdentifier)
    {...}
```

```

void DoWork()
{...}
void SaveState(Param stateIdentifier)
{...}
public void Dispose()
{...}
}

```

The class implements the `MyMethod()` operation, which accepts a parameter of type `Param` (a pseudotype invented for this example) that identifies the instance:

```
public void MyMethod(Param stateIdentifier);
```

The instance then uses the identifier to retrieve its state and to save the state back at the end of the method call. Any piece of state that is common to all clients can be allocated at the constructor and disposed of in `Dispose()`.

The per-call activation mode works best when the amount of work to be done in each method call is finite, and there are no more activities to complete in the background once a method returns. Because the object will be discarded once the method returns, you should not spin off background threads or dispatch asynchronous calls back into the instance.

Since the per-call service retrieves its state from some storage in every method call, per-call services work very well in conjunction with a load-balancing machine, as long as the state repository is some global resource accessible to all machines. The load balancer can redirect calls to different machines at will, knowing that each per-call service can execute the call after retrieving its state.

Per-call services and performance

Per-call services clearly offer a trade-off between performance (the overhead of retrieving and saving the instance state on each method call) and scalability (holding onto the state and the resources it ties in). There are no hard-and-fast rules as to when and to what extent you should trade some performance for a lot of scalability. You may need to profile your system and ultimately design some services to use per-call activation and others not to use it.

Cleanup operations

Whether or not the service type supports `IDisposable` is an implementation detail and is of no relevance to the client. In fact, the client has no way of calling the `Dispose()` method anyway. When you design a contract for a per-call service, avoid defining operations that are dedicated for state or resource cleanup, like this:

```
//Avoid
[ServiceContract]
interface IMyContract
```

```

{
    void DoSomething();
    void Cleanup();
}
[ServiceBehavior(InstanceContextMode = InstanceContextMode.PerCall)]
class MyPerCallService : IMyContract, IDisposable
{
    public void DoSomething()
    {...}
    public void Cleanup()
    {...}
    public void Dispose()
    {
        Cleanup();
    }
}

```

The folly of such a design is obvious: if the client does call the cleanup method, it has the detrimental effect of creating an object just so the client can call `Cleanup()` on it, followed by a call to `IDisposable.Dispose()` (if present) by WCF to do the cleanup again.

Choosing Per-Call Services

While the programming model of per-call services may look somewhat alien to application developers transitioning to service developers, per-call services are actually the preferred instance management mode for many WCF services. This is simply because per-call services scale better, or at least are scale-invariant. When designing a service, my golden rule for scalability is 10X. That is, every service should be designed to handle a load at least an order of magnitude greater than what its requirements call for. In every other engineering discipline, engineers never design a system to handle its exact nominal specified load. You would not want to enter a building whose beams could support only the exact load they were required to handle, ride in an elevator whose cable could handle only the exact number of passengers it's rated for, and so on. Software systems are no different—why design a system for the specific current load while every other person in the company is working to increase business and the implied load? You should design software systems to last years and to sustain current and future loads. As a result, when using the 10X golden rule, you very quickly end up needing the scalability of the per-call service.

Per-Session Services

WCF can maintain a logical session between a client and a particular service instance. When the client creates a new proxy to a service configured as a *sessionful service*, the client gets a new dedicated service instance that is independent of all other instances of the same service. That instance will typically remain in service until the client no

longer needs it. This activation mode (sometimes also referred to as the *private-session mode*) is very much like the classic client/server model: each private session uniquely binds a proxy and its set of client- and service-side channels to a particular service instance, or more specifically, to its context. It follows that a transport session is required for the private-session instantiation mode, as discussed later in this section.

Because the service instance remains in memory throughout the session, it can maintain state in memory, and the programming model is very much like that of the classic client/server model. Consequently, it suffers from the same scalability and transaction issues as the classic client/server model. A service configured for private sessions cannot typically support more than a few dozen (or perhaps up to one or two hundred) outstanding clients, due to the cost associated with each such dedicated service instance.



The client session is per service endpoint per proxy. If the client creates another proxy to the same or a different endpoint, that second proxy will be associated with a new instance and session.

Configuring Private Sessions

There are three elements to supporting a session: behavior, binding, and contract. The behavior part is required so that WCF will keep the service instance context alive throughout the session, and to direct the client messages to it. This local behavior facet is achieved by setting the `InstanceContextMode` property of the `ServiceBehavior` attribute to `InstanceContextMode.PerSession`:

```
[ServiceBehavior(InstanceContextMode = InstanceContextMode.PerSession)]
class MyService : IMyContract
{...}
```

Since `InstanceContextMode.PerSession` is the default value of the `InstanceContextMode` property, these definitions are equivalent:

```
class MyService : IMyContract
{...}

[ServiceBehavior]
class MyService : IMyContract
{...}

[ServiceBehavior(InstanceContextMode = InstanceContextMode.PerSession)]
class MyService : IMyContract
{...}
```

The session typically terminates when the client closes the proxy, which causes the proxy to notify the service that the session has ended. If the service supports `IDisposable`, the `Dispose()` method will be called asynchronously to the client. However, `Disposed()` will be called on a worker thread without an operation context.

In order to correlate all messages from a particular client to a particular instance, WCF needs to be able to identify the client. As explained in [Chapter 1](#), this is exactly what the transport session achieves. If your service is designed to be used as a sessionful service, there has to be some contract-level way for you to express that expectation. The contractual element is required across the service boundary, because the client-side WCF runtime needs to know it should use a session. To that end, the `ServiceContract` attribute offers the property `SessionMode`, of the enum type `SessionMode`:

```
public enum SessionMode
{
    Allowed,
    Required,
    NotAllowed
}
[AttributeUsage(AttributeTargets.Interface|AttributeTargets.Class,
    Inherited=false)]
public sealed class ServiceContractAttribute : Attribute
{
    public SessionMode SessionMode
    {get;set;}
    //More members
}
```

`SessionMode` defaults to `SessionMode.Allowed`. The configured `SessionMode` value is included in the service metadata and is reflected correctly when the client imports the contract metadata. The enum value of `SessionMode` has nothing to do with the service session; in fact, its proper name should have been `TransportSessionMode` since it pertains to the transport session, not to the logical session maintained between the client and the instance.

SessionMode.Allowed

`SessionMode.Allowed` is the default value of the `SessionMode` property, so these definitions are equivalent:

```
[ServiceContract]
interface IMyContract
{...}

[ServiceContract(SessionMode = SessionMode.Allowed)]
interface IMyContract
{...}
```

All bindings support configuring the contract on the endpoint with `SessionMode.Allowed`. When the `SessionMode` property is configured with this value, transport sessions are allowed, but not enforced. The exact resulting behavior is a product of the service configuration and the binding used. If the service is configured for per-call activation, it still behaves as per-call service, as is the case in [Example 4-2](#). When the service is configured for per-session activation, it will behave as a per-session service only if the binding used maintains a transport-level session. For example, the `BasicHttpBinding` can never have a transport-level session, due to the connectionless nature of the HTTP protocol. The `WSHttpBinding` without Message security and without reliable messaging will also not maintain a transport-level session. In both of these cases, although the service is configured with `InstanceContextMode.PerSession` and the contract with `SessionMode.Allowed`, the service will behave as a per-call service.

However, if you use the `WSHttpBinding` with Message security (its default configuration) or with reliable messaging, or if you use the `NetTcpBinding` or the `NetNamedPipeBinding`, the service will behave as a per-session service. For example, assuming use of the `NetTcpBinding`, this service behaves as sessionful:

```
[ServiceContract]
interface IMyContract
{...}

class MyService : IMyContract
{...}
```

Note that the previous code snippet simply takes the default of both the `SessionMode` and the `InstanceContextMode` properties.

SessionMode.Required

The `SessionMode.Required` value mandates the use of a transport-level session, but not necessarily an application-level session. You cannot have a contract configured with `SessionMode.Required` with a service endpoint whose binding does not maintain a transport-level session, and this constraint is verified at the service load time. However, you can still configure the service to be a per-call service, and the service instance will be created and destroyed on each client call. Only if the service is configured as a sessionful service will the service instance persist throughout the client's session:

```
[ServiceContract(SessionMode = SessionMode.Required)]
interface IMyContract
{...}

class MyService : IMyContract
{...}
```



When designing a sessionful contract, I recommend explicitly using `SessionMode.Required` and not relying on the default of `SessionMode.Allowed`. The rest of the code samples in this book actively apply `SessionMode.Required` when sessionful interaction is by design.

Example 4-4 lists the same service and client as in [Example 4-2](#), except the contract and service are configured to require a private session. As you can see from the output, the client got a dedicated instance.

Example 4-4. Per-session service and client

```
/////////////// Service Code ///////////
[ServiceContract(SessionMode = SessionMode.Required)]
interface IMyContract
{
    [OperationContract]
    void MyMethod();
}
class MyService : IMyContract, IDisposable
{
    int m_Counter = 0;

    MyService()
    {
        Trace.WriteLine("MyService.MyService()");
    }
    public void MyMethod()
    {
        m_Counter++;
        Trace.WriteLine("Counter = " + m_Counter);
    }
    public void Dispose()
    {
        Trace.WriteLine("MyService.Dispose()");
    }
}
/////////////// Client Code ///////////
MyContractClient proxy = new MyContractClient();

proxy.MyMethod();
proxy.MyMethod();

proxy.Close();

//Output
MyService.MyService()
Counter = 1
Counter = 2
MyService.Dispose()
```

SessionMode.NotAllowed

SessionMode.NotAllowed disallows the use of a transport-level session, which precludes an application-level session. Regardless of the service configuration, when this value is used, the service will always behave as a per-call service.

Since both the TCP and IPC protocols maintain a session at the transport level, you cannot configure a service endpoint that uses the NetTcpBinding or the NetNamedPipeBinding to expose a contract marked with SessionMode.NotAllowed, and this is verified at the service load time. However, the use of the WSHttpBinding with an emulated transport session is still allowed. In the interest of readability, I recommend that when selecting SessionMode.NotAllowed, you always also configure the service as per-call:

```
[ServiceContract(SessionMode = SessionMode.NotAllowed)]
interface IMyContract
{...}

[ServiceBehavior(InstanceContextMode = InstanceContextMode.PerCall)]
class MyService : IMyContract
{...}
```

Since the BasicHttpBinding cannot have a transport-level session, endpoints that use it behave as if the contract is always configured with SessionMode.NotAllowed. I view SessionMode.NotAllowed as a setting available for the sake of completeness more than anything else, and I would not explicitly choose it.

Bindings, contracts, and service behavior

Table 4-1 summarizes the resulting instance mode as a product of the binding being used, the session mode in the contract, and the configured instance context mode in the service behavior. The table does not list invalid configurations, such as SessionMode.Required with the BasicHttpBinding.

Table 4-1. Instance mode as a product of the binding, contract configuration, and service behavior

Binding	Session mode	Context mode	Instance mode
Basic	Allowed/NotAllowed	PerCall/PerSession	PerCall
TCP, IPC	Allowed/Required	PerCall	PerCall
TCP, IPC	Allowed/Required	PerSession	PerSession
WS (no Message security, no reliability)	NotAllowed/Allowed	PerCall/PerSession	PerCall
WS (with Message security or reliability)	Allowed/Required	PerSession	PerSession
WS (with Message security or reliability)	NotAllowed	PerCall/PerSession	PerCall

Consistent configuration

I strongly recommend that if one contract the service implements is a sessionful contract, then all contracts should be sessionful, and that you should avoid mixing per-call and sessionful contracts on the same per-session service type (even though WCF allows it):

```
[ServiceContract(SessionMode = SessionMode.Required)]
interface IMyContract
{...}

[ServiceContract(SessionMode = SessionMode.NotAllowed)]
interface IMyOtherContract
{...}

//Avoid
class MyService : IMyContract,IMyOtherContract
{...}
```

The reason is obvious: per-call services need to proactively manage their state, while per-session services do not. While the two contracts will be exposed on two different endpoints and can be consumed independently by two different clients, this duality requires cumbersome implementation for the underlying service class.

Sessions and Reliability

The session between the client and the service instance is only as reliable as the underlying transport session. Consequently, a service that implements a sessionful contract should have all of its endpoints that expose that contract use bindings that enable reliable transport sessions. Make sure to always use a binding that supports reliability and to explicitly enable it at both the client and the service, either programmatically or administratively, as shown in [Example 4-5](#).

Example 4-5. Enabling reliability for per-session services

```
<!--Host configuration:-->
<system.serviceModel>
  <services>
    <service name = "MyPerSessionService">
      <endpoint
        address   = "net.tcp://localhost:8000/MyPerSessionService"
        binding   = "netTcpBinding"
        bindingConfiguration = "TCPSession"
        contract  = "IMyContract"
      />
    </service>
  </services>
  <bindings>
    <netTcpBinding>
```

```

<binding name = "TCPsession">
    <reliableSession enabled = "true"/>
</binding>
</netTcpBinding>
</bindings>
</system.serviceModel>

<!--Client configuration:-->
<system.serviceModel>
    <client>
        <endpoint
            address   = "net.tcp://localhost:8000/MyPerSessionService/"
            binding   = "netTcpBinding"
            bindingConfiguration = "TCPsession"
            contract  = "IMyContract"
        />
    </client>
    <bindings>
        <netTcpBinding>
            <binding name = "TCPsession">
                <reliableSession enabled = "true"/>
            </binding>
        </netTcpBinding>
    </bindings>
</system.serviceModel>

```

The one exception to this rule is the IPC binding. This binding has no need for the reliable messaging protocol (all calls will be on the same machine anyway), and it is considered an inherently reliable transport.

Just as a reliable transport session is optional, so is ordered delivery of messages, and WCF will provide for a session even when ordered delivery is disabled. However, by the very nature of an application session, a client that interacts with a sessionful service expects all messages to be delivered in the order they are sent. Luckily, ordered delivery is enabled by default when reliable transport sessions are enabled, so no additional setting is required.

The Session ID

Every session has a unique ID that both the client and the service can obtain. The session ID is largely in the form of a GUID, and it can be used for logging and diagnostics. The service can access the session ID via the *operation call context*, which is a set of properties (including the session ID) that are used for callbacks, message headers, transaction management, security, host access, and access to the object representing the execution context itself. Every service operation has an operation call context, accessible via the `OperationContext` class. A service can obtain a reference to the operation context of the current method via the `Current` static method of the `OperationContext` class:

```

public sealed class OperationContext : ...
{
    public static OperationContext Current
    {get;set;}
    public string SessionId
    {get;}
}

```

To access the session ID, the service needs to read the value of the `SessionId` property, which returns (almost) a GUID in the form of a string. In the case of the TCP binding without reliability, it will be followed by the ordinal number of the session from that host:

```

string sessionID = OperationContext.Current.SessionId;
Trace.WriteLine(sessionID);
//Traces:
//uuid:8a0480da-7ac0-423e-9f3e-b2131bcbad8d:id=1

```

If a per-call service without a transport session accesses the `SessionId` property, the session ID will be `null`, since there is no session and therefore no ID.

The client can access the session ID via the proxy. As introduced in [Chapter 1](#), the class `ClientBase<T>` is the base class of the proxy. `ClientBase<T>` provides the read-only property `InnerChannel` of the type `IClientChannel`. `IClientChannel` derives from the interface `IContextChannel`, which provides a `SessionId` property that returns the session ID in the form of a string:

```

public interface IContextChannel : ...
{
    string SessionId
    {get;}
    //More members
}
public interface IClientChannel : IContextChannel, ...
{...}
public abstract class ClientBase<T> : ...
{
    public IClientChannel InnerChannel
    {get;}
    //More members
}

```

Given the definitions in [Example 4-4](#), the client might obtain the session ID like this:

```

MyContractClient proxy = new MyContractClient();
proxy.MyMethod();

string sessionID = proxy.InnerChannel.SessionId;
Trace.WriteLine(sessionID);

```

However, the degree to which the client-side session ID matches that of the service (and even when the client is allowed to access the `SessionId` property) is a product of

the binding used and its configuration. What correlates the client-side and service-side session IDs is the reliable session at the transport level. If the TCP binding is used, when a reliable session is enabled (as it should be) the client can obtain a valid session ID only after issuing the first method call to the service to establish the session, or after explicitly opening the proxy. In this case, the session ID obtained by the client will match that of the service. (If the client accesses the session ID before the first call, the `SessionId` property will be set to `null`.) If the TCP binding is used but reliable sessions are disabled, the client can access the session ID before making the first call, but the ID obtained will be different from that obtained by the service. With the WS binding, if reliable messaging is enabled, the session ID will be `null` until after the first call (or after the client opens the proxy), but after that the client and the service will always have the same session ID. Without reliable messaging, the client must first use the proxy (or just open it) before accessing the session ID, or risk an `InvalidOperationException`. After opening the proxy, the client and the service will have a correlated session ID. With the IPC binding, the client can access the `SessionId` property before making the first call, but the client will always get a session ID different from that of the service. When using this binding, it is therefore better to ignore the session ID altogether.

Session Termination

Typically, the session will end once the client closes the proxy. However, in case the client neglects to close the proxy, or when the client terminates ungracefully or there is a communication problem, the session will also terminate once the inactivity timeout of the transport session is exceeded.

Singleton Service

The *singleton service* is the ultimate shareable service. When you configure a service as a singleton, all clients are independently connected to the same single well-known instance context and implicitly to the same instance inside, regardless of which endpoint of the service they connect to. The singleton is created exactly once, when the host is created, and lives forever: it is disposed of only when the host shuts down.



IIS or the WAS only create a singleton when the first request to any service in that process is made. To maintain the semantic of the singleton, you must use self-hosting.

Using a singleton does not require clients to maintain a logical session with the singleton instance, or to use a binding that supports a transport-level session. If the contract the client consumes has a session, during the call the singleton will have the

same session ID as the client (binding permitting), but closing the client proxy will terminate only the transport session, not the singleton context and the instance inside. If the singleton service supports contracts without a session, those contracts will not be per-call: they too will be connected to the same instance. By its very nature, the singleton is shared, and each client should simply create its own proxy or proxies to it.

You configure a singleton service by setting the `InstanceContextMode` property to `InstanceContextMode.Single`:

```
[ServiceBehavior(InstanceContextMode = InstanceContextMode.Single)]
class MySingleton : ...
{...}
```

Example 4-6 demonstrates a singleton service with two contracts, one that requires a session and one that does not. As you can see from the client call, the calls on the two endpoints were routed to the same instance, and closing the proxies did not terminate the singleton.

Example 4-6. A singleton service and client

```
/////////////// Service Code ///////////
[ServiceContract(SessionMode = SessionMode.Required)]
interface IMyContract
{
    [OperationContract]
    void MyMethod();
}

[ServiceContract(SessionMode = SessionMode.NotAllowed)]
interface IMyOtherContract
{
    [OperationContract]
    void MyOtherMethod();
}

[ServiceBehavior(InstanceContextMode=InstanceContextMode.Single)]
class MySingleton : IMyContract,IMyOtherContract, IDisposable
{
    int m_Counter = 0;

    public MySingleton()
    {
        Trace.WriteLine("MySingleton.MySingleton()");
    }
    public void MyMethod()
    {
        m_Counter++;
        Trace.WriteLine("Counter = " + m_Counter);
    }
    public void MyOtherMethod()
    {
```

```

        m_Counter++;
        Trace.WriteLine("Counter = " + m_Counter);
    }
    public void Dispose()
    {
        Trace.WriteLine("Singleton.Dispose()");
    }
}
//////////////////////////// Client Code /////////////////////
MyContractClient proxy1 = new MyContractClient();
proxy1.MyMethod();
proxy1.Close();

MyOtherContractClient proxy2 = new MyOtherContractClient();
proxy2.MyOtherMethod();
proxy2.Close();

//Output
MySingleton.MySingleton()
Counter = 1
Counter = 2

```

Initializing a Singleton

Sometimes, you may not want to create and initialize the singleton using just the default constructor. Perhaps initializing the state requires some custom steps or specific knowledge that the clients should not be bothered with, or that is not available to the clients. Whatever the reason, you may want to create the singleton using some other mechanism besides the WCF service host. To support such scenarios, WCF allows you to directly create the singleton instance beforehand using normal CLR instantiation, initialize it, and then open the host with that instance in mind as the singleton service. The `ServiceHost` class offers a dedicated constructor that accepts an object:

```

public class ServiceHost : ServiceHostBase, ...
{
    public ServiceHost(object singletonInstance, params Uri[] baseAddresses);
    public object SingletonInstance
    {get;}
    //More members
}

```

Note that the `object` must be configured as a singleton. For instance, consider the code in [Example 4-7](#). The class `MySingleton` will be first initialized and then hosted as a singleton.

Example 4-7. Initializing and hosting a singleton

```
//Service code
[ServiceContract]
interface IMyContract
{
    [OperationContract]
    void MyMethod();
}

[ServiceBehavior(InstanceContextMode = InstanceContextMode.Single)]
class MySingleton : IMyContract
{
    public int Counter
    {get;set;}

    public void MyMethod()
    {
        Counter++;
        Trace.WriteLine("Counter = " + Counter);
    }
}

//Host code
MySingleton singleton = new MySingleton();
singleton.Counter = 287;

ServiceHost host = new ServiceHost(singleton);
host.Open();

//Client code
MyContractClient proxy = new MyContractClient();
proxy.MyMethod();
proxy.Close();

//Output:
Counter = 288
```

If you do initialize and host a singleton this way, you may also want to be able to access it directly on the host side. WCF enables downstream objects to reach back into the singleton directly using the `SingletonInstance` property of `ServiceHost`. Any party on the call chain leading down from an operation call on the singleton can always access the host via the operation context's read-only `Host` property:

```
public sealed class OperationContext : ...
{
    public ServiceHostBase Host
    {get;}
    //More members
}
```

Once you have the singleton reference, you can interact with it directly:

```
ServiceHost host = OperationContext.Current.Host as ServiceHost;
Debug.Assert(host != null);
MySingleton singleton = host.SingletonInstance as MySingleton;
Debug.Assert(singleton != null);
singleton.Counter = 388;
```

If no singleton instance was provided to the host, `SingletonInstance` returns `null`.

Streamlining with `ServiceHost<T>`

The `ServiceHost<T>` class presented in [Chapter 1](#) can be extended to offer type-safe singleton initialization and access:

```
public class ServiceHost<T> : ServiceHost
{
    public ServiceHost(T singleton,params Uri[] baseAddresses) :
        base(singleton,baseAddresses)
    {}
    public virtual T Singleton
    {
        get
        {
            if(SingletonInstance == null)
            {
                return default(T);
            }
            return (T)SingletonInstance;
        }
    }
    //More members
}
```

The type parameter provides type-safe binding for the object used for construction:

```
MySingleton singleton = new MySingleton();
singleton.Counter = 287;

ServiceHost<MySingleton> host = new ServiceHost<MySingleton>(singleton);
host.Open();
```

and the object returned from the `Singleton` property:

```
ServiceHost<MySingleton> host = OperationContext.Current.Host as
    ServiceHost<MySingleton>;
Debug.Assert(host != null);
host.Singleton.Counter = 388;
```



The `InProcFactory<T>` (presented in [Chapter 1](#)) is similarly extended to initialize a singleton instance.

Choosing a Singleton

The singleton service is the sworn enemy of scalability. The reason has to do with singleton state synchronization, rather than the cost of that single instance. Having a singleton implies that the singleton has some valuable state that you wish to share across multiple clients. The problem is that if the singleton's state is mutable and multiple clients connect to the singleton, they may all do so concurrently, and the incoming client calls will be on multiple worker threads. The singleton must therefore synchronize access to its state to avoid state corruption. This, in turn, means that only one client at a time can access the singleton. This constraint may degrade throughput, responsiveness, and availability to the point that the singleton is unusable in a decent-sized system. For example, if an operation on a singleton takes one-tenth of a second, the singleton can service only 10 clients per second. If there are many more clients (say 20 or 100), the system's performance will be inadequate.

In general, you should use a singleton only if it maps well to a natural singleton in the application domain. A *natural singleton* is a resource that is, by its very nature, single and unique. Examples of natural singletons are a global logbook to which all services should log their activities, a single communication port, or a single mechanical motor. Avoid using a singleton if there is even the slightest chance that the business logic will allow more than one such service in the future (for example, adding another motor or a second communication port). The reason is clear: if your clients all depend on implicitly being connected to the well-known instance and more than one service instance is available, the clients will suddenly need to have a way to bind to the correct instance. This can have severe implications for the application's programming model. Because of these limitations, I recommend that you avoid singletons in the general case and find ways to share the state of the singleton instead of the singleton instance itself. That said, there are cases when using a singleton is acceptable, as mentioned earlier.

Demarcating Operations

Sometimes, a sessionful contract has an implied order of operation invocations. Some operations cannot be called first, while other operations must be called last. For example, consider this contract, used to manage customer orders:

```
[ServiceContract(SessionMode = SessionMode.Required)]
interface IOrderManager
{
    [OperationContract]
    void SetCustomerId(int customerId);

    [OperationContract]
    void AddItem(int itemId);

    [OperationContract]
```

```

        decimal GetTotal();

        [OperationContract]
        bool ProcessOrders();
    }

```

The contract has the following constraints: the client must provide the customer ID as the first operation in the session, or else no other operations can take place; items may be added, and the total calculated, as often as the client wishes; processing the order terminates the session, and therefore must come last. In classic .NET, such requirements often forced the developers to support some state machine or state flags and to verify the state on every operation.

WCF, however, allows contract designers to designate contract operations as operations that can or cannot start or terminate the session, using the `IsInitiating` and `IsTerminating` properties of the `OperationContract` attribute:

```

[AttributeUsage(AttributeTargets.Method)]
public sealed class OperationContractAttribute : Attribute
{
    public bool IsInitiating
    {get;set;}
    public bool IsTerminating
    {get;set;}
    //More members
}

```

These properties can be used to demarcate the boundary of the session; hence, I call this technique *demarcating operations*. At service load time (or during the proxy use time on the client side), if these properties are set to their nondefault values, WCF verifies that the demarcating operations are part of a contract that mandates sessions (i.e., that `SessionMode` is set to `SessionMode.Required`) and throws an `InvalidOperationException` otherwise. Both a sessionful service and a singleton can implement contracts that use demarcating operations to manage their client sessions.

The default values of these properties are `true` for `IsInitiating` and `false` for `IsTerminating`. Consequently, these two definitions are equivalent:

```

[OperationContract]
void MyMethod();

[OperationContract(IsInitiating = true,IsTerminating = false)]
void MyMethod();

```

As you can see, you can set both properties on the same method. In addition, operations do not demarcate the session boundary by default—operations can be called first, last, or in between any other operations in the session. Using nondefault values enables you to dictate that a method is not called first, or that it is called last, or both:

```

[ServiceContract(SessionMode = SessionMode.Required)]
interface IMyContract
{
    [OperationContract]
    void StartSession();

    [OperationContract(IsInitiating = false)]
    void CannotStart();

    [OperationContract(IsTerminating = true)]
    void EndSession();

    [OperationContract(IsInitiating = false,IsTerminating = true)]
    void CannotStartCanEndSession();
}

```

Going back to the order-management contract, you can use demarcating operations to enforce the interaction constraints:

```

[ServiceContract(SessionMode = SessionMode.Required)]
interface IOrderManager
{
    [OperationContract]
    void SetCustomerId(int customerId);

    [OperationContract(IsInitiating = false)]
    void AddItem(int itemId);

    [OperationContract(IsInitiating = false)]
    decimal GetTotal();

    [OperationContract(IsInitiating = false,IsTerminating = true)]
    bool ProcessOrders();
}

//Client code
OrderManagerClient proxy = new OrderManagerClient();

proxy.SetCustomerId(123);
proxy.AddItem(4);
proxy.AddItem(5);
proxy.AddItem(6);
proxy.ProcessOrders();

proxy.Close();

```

When `IsInitiating` is set to `true` (its default), it means the operation will start a new session if it is the first method the client calls but will be part of the ongoing session if another operation is called first. When `IsInitiating` is set to `false`, it means that a client can never call that operation as the first operation in a new session, and that the method can only be part of an ongoing session.

When `IsTerminating` is set to `false` (its default), it means the session continues after the operation returns. When `IsTerminating` is set to `true`, it means the session terminates once the method returns, and WCF disposes of the service instance asynchronously. The client will not be able to issue additional calls on the proxy. Note that the client should still close the proxy.



When you generate a proxy to a service that uses demarcating operations, the imported contract definition contains the property settings. In addition, WCF enforces the demarcation separately on the client and service sides, so you could actually employ them independently.

Instance Deactivation

Conceptually, the sessionful service instance management technique as described so far connects a client (or clients) to a service instance. Yet, the real picture is more complex. Recall from [Chapter 1](#) that each service instance is hosted in a context, as shown in [Figure 4-2](#).

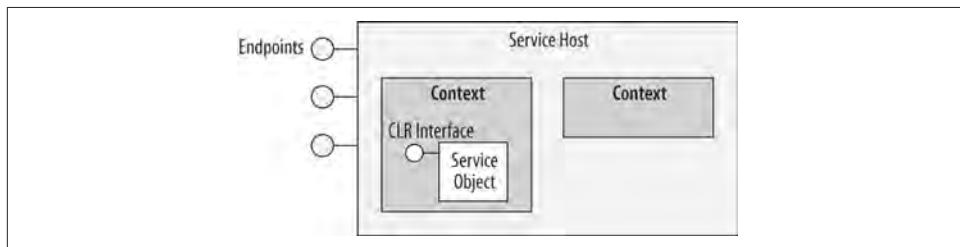


Figure 4-2. Contexts and instances

What sessions actually do is correlate the client messages not to the instance, but to the context that hosts it. When the session starts, the host creates a new context. When the session ends, the context is terminated. By default, the lifetime of the context is the same as that of the instance it hosts. However, for optimization and extensibility purposes, WCF provides the service designer with the option of separating the two lifetimes and deactivating the instance separately from its context. In fact, WCF also allows a context to exist without an associated instance at all, as shown in [Figure 4-2](#). I call this instance management technique *context deactivation*. The common way of controlling context deactivation is via the `ReleaseInstanceMode` property of the `OperationBehavior` attribute:

```
public enum ReleaseInstanceMode
{
    None,
    BeforeCall,
    AfterCall,
```

```

        BeforeAndAfterCall,
    }
    [AttributeUsage(AttributeTargets.Method)]
    public sealed class OperationBehaviorAttribute : Attribute, ...
    {
        public ReleaseInstanceMode ReleaseInstanceMode
        {get;set;}
        //More members
    }
}

```

`ReleaseInstanceMode` is of the enum type `ReleaseInstanceMode`. The various values of `ReleaseInstanceMode` control when to release the instance in relation to the method call: before, after, before and after, or not at all. When releasing the instance, if the service supports `IDisposable`, the `Dispose()` method is called and `Dispose()` has an operation context.

You typically apply instance deactivation on some but not all service methods, or with different values on different methods:

```

[ServiceContract(SessionMode = SessionMode.Required)]
interface IMyContract
{
    [OperationContract]
    void MyMethod();

    [OperationContract]
    void MyOtherMethod();
}
class MyService : IMyContract, IDisposable
{
    [OperationBehavior(ReleaseInstanceMode = ReleaseInstanceMode.AfterCall)]
    public void MyMethod()
    {...}
    public void MyOtherMethod()
    {...}
    public void Dispose()
    {...}
}

```

The reason you typically apply it sporadically is that if you were to apply it uniformly, you would end up with a per-call-like service, in which case you might as well have configured the service as per-call.

If relying on instance deactivation assumes a certain call order, you can try to enforce that order using demarcating operations.

Configuring with `ReleaseInstanceMode.None`

The default value for the `ReleaseInstanceMode` property is `ReleaseInstanceMode.None`, so these two definitions are equivalent:

```
[OperationContract(ReleaseInstanceMode = ReleaseInstanceMode.None)]
public void MyMethod()
{...}

public void MyMethod()
{...}
```

`ReleaseInstanceMode.None` means that the instance lifetime is not affected by the call, as shown in [Figure 4-3](#).

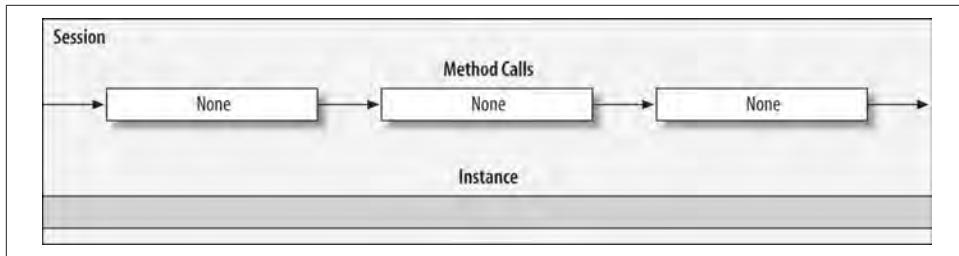


Figure 4-3. Instance lifetime with methods configured with `ReleaseInstanceMode.None`

Configuring with `ReleaseInstanceMode.BeforeCall`

When a method is configured with `ReleaseInstanceMode.BeforeCall`, if there is already an instance in the session, before forwarding the call WCF will deactivate it, create a new instance in its place, and let that new instance service the call, as shown in [Figure 4-4](#).

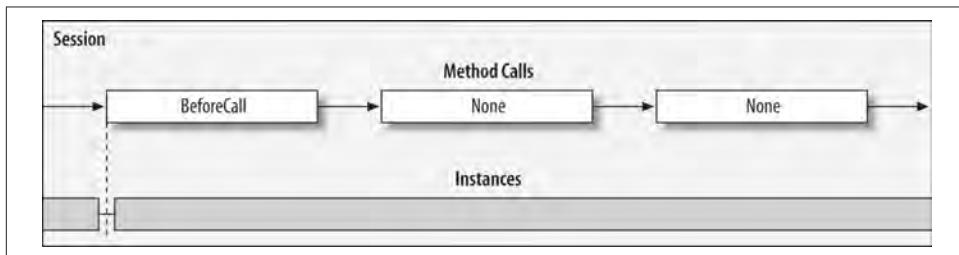


Figure 4-4. Instance lifetime with methods configured with `ReleaseInstanceMode.BeforeCall`

WCF deactivates the instance and calls `Dispose()` before the call is done on the incoming call thread, while the client blocks. This ensures that the deactivation is indeed done before the call, not concurrently with it. `ReleaseInstanceMode.BeforeCall` is designed to optimize methods such as `Create()` that acquire some valuable resources, yet wish to release the previously allocated resources. Instead of acquiring the resources when the session starts, you wait until the call to the `Create()` method and then both release the previously allocated resources and allocate new ones. After

Create() is called, you are ready to start calling other methods on the instance, which are typically configured with ReleaseInstanceMode.None.

Configuring with ReleaseInstanceMode.AfterCall

When a method is configured with ReleaseInstanceMode.AfterCall, WCF deactivates the instance after the call, as shown in [Figure 4-5](#).

This is designed to optimize a method such as Cleanup() that cleans up valuable resources held by the instance, without waiting for the session to terminate. ReleaseInstanceMode.AfterCall is typically applied on methods called after methods configured with ReleaseInstanceMode.None.

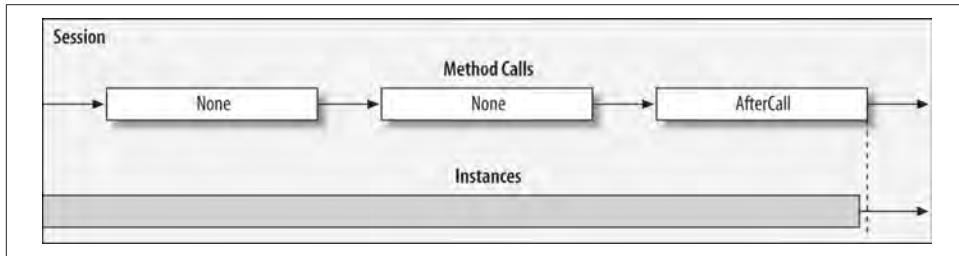


Figure 4-5. Instance lifetime with methods configured with ReleaseInstanceMode.AfterCall

Configuring with ReleaseInstanceMode.BeforeAndAfterCall

As its name implies, configuring a method with ReleaseInstanceMode.BeforeAndAfterCall has the combined effect of using ReleaseInstanceMode.BeforeCall and ReleaseInstanceMode.AfterCall. If the context has an instance before the call is made, just before the call WCF deactivates that instance and creates a new instance to service the call. It then deactivates the new instance after the call, as shown in [Figure 4-6](#).

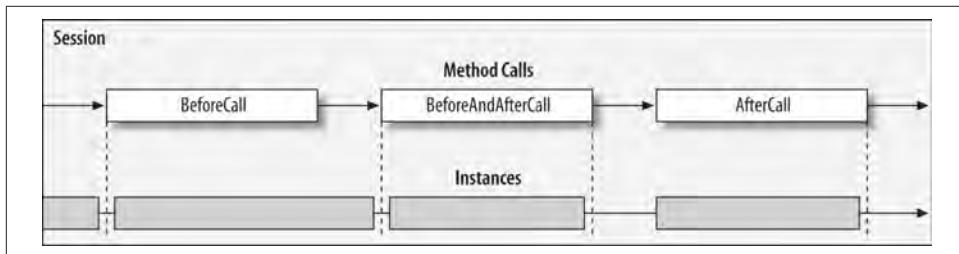


Figure 4-6. Instance lifetime with methods configured with ReleaseInstanceMode.BeforeAndAfterCall

`ReleaseInstanceMode.BeforeAndAfterCall` may look superfluous at first glance, but it actually complements the other values. It is designed to be applied on methods called after methods marked with `ReleaseInstanceMode.BeforeCall` or `None`, or before methods marked with `ReleaseInstanceMode.AfterCall` or `None`. Consider a situation where the sessionful service wants to benefit from state-aware behavior (like a per-call service), while holding onto resources only when needed to optimize resource allocation and security lookup. If `ReleaseInstanceMode.BeforeCall` were the only available option, there would be a period of time after the call when the resources would still be allocated to the object, but would not be in use. A similar situation would occur if `ReleaseInstanceMode.AfterCall` were the only available option, because there would be a period of time before the call when the resources would be wasted.

Explicit Deactivation

Instead of making a design-time decision on which methods to use to deactivate an instance, you can make a runtime decision to deactivate the instance after the method returns. You do that by calling the `ReleaseServiceInstance()` method on the instance context. You obtain the instance context via the `InstanceContext` property of the operation context:

```
public sealed class InstanceContext : ...
{
    public void ReleaseServiceInstance();
    //More members
}
public sealed class OperationContext : ...
{
    public InstanceContext InstanceContext
    {get;}
    //More members
}
```

Example 4-8 demonstrates using explicit deactivation to implement a custom instance management technique that is dependent on the value of a counter.

Example 4-8. Using `ReleaseServiceInstance()`

```
[ServiceContract(SessionMode = SessionMode.Required)]
interface IMyContract
{
    [OperationContract]
    void MyMethod();
}
class MyService : IMyContract, IDisposable
{
    int m_Counter = 0;
```

```

public void MyMethod()
{
    m_Counter++;

    if(m_Counter > 4)
    {
        OperationContext.Current.InstanceContext.ReleaseServiceInstance();
    }
}
public void Dispose()
{...}
}

```

Calling `ReleaseServiceInstance()` has a similar effect to using `ReleaseInstanceMode.AfterCall`. When used in a method decorated with `ReleaseInstanceMode.BeforeCall`, it has a similar effect to using `ReleaseInstanceMode.BeforeAndAfterCall`.



Instance deactivation affects a singleton as well, although combining the two makes little sense—by its very definition, it is permissible and even desirable to never deactivate the singleton.

Using Instance Deactivation

Instance deactivation is an optimization technique, and like all optimization techniques, you should avoid it in the general case. It adds complexity to the application and makes the code less approachable and maintainable to all but the WCF expert. Consider using instance deactivation only after failing to meet both your performance and scalability goals and when careful examination and profiling has proven beyond a doubt that using instance deactivation will improve the situation. If scalability and throughput are your concerns, you should take advantage of the simplicity of the per-call instancing mode and avoid instance deactivation. The main reason I share this technique with you is that WCF itself makes extensive use of instance deactivation; thus, knowledge of it is instrumental in demystifying other aspects of WCF, such as durable services and transactions.

Durable Services

Consider the case of a long-running business process or workflow, comprised of multiple execution sequences, that lasts many days or even weeks.



I use the term *workflow* to denote a business workflow in general, not one that is necessarily supported by or related to the product called Windows Workflow.

Such long-running processes may involve clients (or even end users) that connect to the application, perform a finite amount of work, transition the workflow to a new state, and then disconnect for an indeterminate amount of time before connecting again and continuing to execute the workflow. The clients may at any point also decide to terminate the workflow and start a new one, or the backend service supporting the workflow may end it. Obviously, there is little point in keeping proxies and services in memory waiting for the clients to call. Such an approach will not robustly withstand the test of time; at the very least, timeout issues will inevitably terminate the connection, and there is no easy way to allow machines on both sides to reboot or log off. The need to allow the clients and the services to have independent lifecycles is an important one in a long-running business process, because without it there is no way to enable the clients to connect, perform some work against the workflow, and disconnect. On the host side, over time you may even want to redirect calls between machines.

The solution for long-running services is to avoid keeping the service state in memory, and to handle each call on a new instance with its own temporary in-memory state. For every operation, the service should retrieve its state from some durable storage (such as a file or a database), perform the requested unit of work for that operation, and then save the state back to the durable storage at the end of the call. Services that follow this model are called *durable services*. Since the durable storage can be shared across machines, using durable services also gives you the ability to route calls to different machines at different times, be it for scalability, redundancy, or maintenance purposes.

Durable Services and Instance Management Modes

This approach to state management for durable services is very much like the one proposed previously for per-call services, which proactively manage their state. Using per-call services makes additional sense because there is no point in keeping the instance around between calls if its state is coming from durable storage. The only distinguishing aspect of a durable service compared with a classic per-call service is that the state repository needs to be durable.

While in theory nothing prevents you from basing a durable service on a sessionful or even a singleton service and having that service manage its state in and out of the durable storage, in practice this would be counterproductive. In the case of a sessionful service, you would have to keep the proxy open on the client side for long periods of time, thus excluding clients that terminate their connections and then reconnect.

In the case of a singleton service, the very notion of a singleton suggests an infinite lifetime with clients that come and go, so there is no need for durability. Consequently, the per-call instantiation mode offers the best choice all around. Note that with durable per-call services, because the primary concern is long-running workflows rather than scalability or resource management, supporting `IDisposable` is optional. It is also worth pointing out that the presence of a transport session is optional for a durable service, since there is no need to maintain a logical session between the client and the service. The transport session will be a facet of the transport channel used and will not be used to dictate the lifetime of the instance.

Initiating and terminating

When the long-running workflow starts, the service must first write its state to the durable storage, so that subsequent operations will find the state in the storage. When the workflow ends, the service must remove its state from the storage; otherwise, over time, the storage will become bloated with instance states not required by anyone.

Instance IDs and Durable Storage

Since a new service instance is created for every operation, an instance must have a way of looking up and loading its state from the durable storage. The client must therefore provide some state identifier for the instance. That identifier is called the *instance ID*. To support clients that connect to the service only occasionally, and client applications or even machines that recycle between calls, as long as the workflow is in progress the client will typically save the instance ID in some durable storage on the client side (such as a file) and provide that ID for every call. When the workflow ends, the client can discard that ID. For an instance ID, it is important to select a type that is serializable and equatable. Having a serializable ID is important because the service will need to save the ID along with its state into the durable storage. Having an equatable ID is required in order to allow the service to obtain the state from the storage. All the .NET primitives (such as `int`, `string`) and `Guid` qualify as instance IDs.

The durable storage is usually some kind of dictionary that pairs the instance ID with the instance state. The service typically will use a single ID to represent all of its state, although more complex relationships involving multiple keys and even hierarchies of keys are possible. For simplicity's sake, I will limit the discussion here to a single ID. In addition, the service often uses a dedicated helper class or a structure to aggregate all its member variables, and stores that type in and retrieves it from the durable storage. Finally, access to the durable storage itself must be thread-safe and synchronized. This is required because multiple instances may try to access and modify the store concurrently.

To help you implement and support simple durable services, I wrote the `FileInstanceStore<ID,T>` class:

```
public interface IInstanceStore<ID,T> where ID : IEquatable<ID>
{
    void RemoveInstance(ID instanceId);
    bool ContainsInstance(ID instanceId);
    T this[ID instanceId]
    {get;set;}
}

public class FileInstanceStore<ID,T> : IInstanceStore<ID,T> where ID :
    IEquatable<ID>
{
    protected readonly string Filename;

    public FileInstanceStore(string fileName);

    //Rest of the implementation
}
```

`FileInstanceStore<ID,T>` is a general-purpose file-based instance store. `FileInstanceStore<ID,T>` takes two type parameters: the `ID` type parameter is constrained to be an equatable type, and the `T` type parameter represents the instance state. `FileInstanceStore<ID,T>` verifies at runtime in a static constructor that both `T` and `ID` are serializable types.

`FileInstanceStore<ID,T>` provides a simple indexer allowing you to read and write the instance state to the file. You can also remove an instance state from the file, and check whether the file contains the instance state. These operations are defined in the `IInstanceStore<ID,T>` interface. The implementation of `FileInstanceStore<ID,T>` encapsulates a dictionary, and on every access it serializes and deserializes the dictionary to and from the file. When `FileInstanceStore<ID,T>` is used for the first time, if the file is empty, `FileInstanceStore<ID,T>` will initialize it with an empty dictionary.

Explicit Instance IDs

The simplest way a client can provide the instance ID to the service is as an explicit parameter for every operation designed to access the state. [Example 4-9](#) demonstrates such a client and service, along with the supporting type definitions.

Example 4-9. Passing explicit instance IDs

```
[DataContract]
class SomeKey : IEquatable<SomeKey>
{...}
```

```

[ServiceContract]
interface IMyContract
{
    [OperationContract]
    void MyMethod(SomeKey instanceId);
}

//Helper type used by the service to capture its state
[Serializable]
struct MyState
{...}

[ServiceBehavior(InstanceContextMode = InstanceContextMode.PerCall)]
class MyService : IMyContract
{
    public void MyMethod(SomeKey instanceId)
    {
        GetState(instanceId);
        DoWork();
        SaveState(instanceId);
    }
    void DoWork()
    {...}

    //Get and set MyState from durable storage
    void GetState(SomeKey instanceId)
    {...}

    void SaveState(SomeKey instanceId)
    {...}
}

```

To make [Example 4-9](#) more concrete, consider [Example 4-10](#), which supports a pocket calculator with durable memory stored in a file.

Example 4-10. Calculator with explicit instance ID

```

[ServiceContract]
interface ICalculator
{
    [OperationContract]
    double Add(double number1,double number2);

    /* More arithmetic operations */

    //Memory management operations

    [OperationContract]
    void MemoryStore(string instanceId,double number);

    [OperationContract]

```

```

    void MemoryClear(string instanceId);
}

[ServiceBehavior(InstanceContextMode = InstanceContextMode.PerCall)]
class MyCalculator : ICalculator
{
    static IInstanceStore<string,double> Memory =
        new FileInstanceStore<string,double>(Settings.Default.MemoryFileName);

    public double Add(double number1,double number2)
    {
        return number1 + number2;
    }
    public void MemoryStore(string instanceId,double number)
    {
        lock(typeof(MyCalculator))
        {
            Memory[instanceId] = number;
        }
    }
    public void MemoryClear(string instanceId)
    {
        lock(typeof(MyCalculator))
        {
            Memory.RemoveInstance(instanceId);
        }
    }
    //Rest of the implementation
}

```

In [Example 4-10](#), the filename is available in the properties of the project in the `Settings` class. All instances of the calculator use the same static memory, in the form of a `FileInstanceStore<string,double>`. The calculator synchronizes access to the memory in every operation across all instances by locking on the service type. Clearing the memory signals to the calculator the end of the workflow, so it purges its state from the storage.

Instance IDs in Headers

Instead of explicitly passing the instance ID, the client can provide the instance ID in the message headers. Using message headers as a technique for passing out-of-band parameters used for custom contexts is described in detail in [Appendix B](#). In this case, the client can use my `HeaderClientBase<T,H>` proxy class, and the service can read the ID in the relevant operations using my `GenericContext<H>` helper class. The service can use `GenericContext<H>` as-is or wrap it in a dedicated context.

The general pattern for this technique is shown in [Example 4-11](#).

Example 4-11. Passing instance IDs in message headers

```
[ServiceContract]
interface IMyContract
{
    [OperationContract]
    void MyMethod();
}
//Client-side
class MyContractClient : HeaderClientBase<IMyContract,SomeKey>,IMyContract
{
    public MyContractClient(SomeKey instanceId)
    {}
    public MyContractClient(SomeKey instanceId,string endpointName) :
        base(instanceId,endpointName)
    {}

    //More constructors

    public void MyMethod()
    {
        Channel.MyMethod();
    }
}
//Service-side
[ServiceBehavior(InstanceContextMode = InstanceContextMode.PerCall)]
class MyService : IMyContract
{
    public void MyMethod()
    {
        SomeKey instanceId = GenericContext<SomeKey>.Current.Value;
        ...
    }
    //Rest same as Example 4-9
}
```

Again, to make [Example 4-11](#) less abstract, [Example 4-12](#) shows the calculator using the message headers technique.

Example 4-12. Calculator with instance ID in headers

```
[ServiceContract]
interface ICalculator
{
    [OperationContract]
    double Add(double number1,double number2);

    /* More arithmetic operations */

    //Memory management operations
```

```

[OperationContract]
void MemoryStore(double number);

[OperationContract]
void MemoryClear();
}
//Client-side
class MyCalculatorClient : HeaderClientBase<ICalculator, string>, ICalculator
{
    public MyCalculatorClient(string instanceId)
    {}

    public MyCalculatorClient(string instanceId, string endpointName) :
        base(instanceId, endpointName)
    {}

    //More constructors

    public double Add(double number1, double number2)
    {
        return Channel.Add(number1, number2);
    }

    public void MemoryStore(double number)
    {
        Channel.MemoryStore(number);
    }

    //Rest of the implementation
}
//Service-side
//If using GenericContext<T> is too raw, can encapsulate:
static class CalculatorContext
{
    public static string Id
    {
        get
        {
            return GenericContext<string>.Current.Value ?? String.Empty;
        }
    }
}

[ServiceBehavior(InstanceContextMode = InstanceContextMode.PerCall)]
class MyCalculator : ICalculator
{
    static IIInstanceStore<string, double> Memory =
        new FileInstanceStore<string, double>(Settings.Default.MemoryFileName);

    public double Add(double number1, double number2)
    {
        return number1 + number2;
    }
}

```

```

    }
    public void MemoryStore(double number)
    {
        lock(typeof(MyCalculator))
        {
            Memory[CalculatorContext.Id] = number;
        }
    }
    public void MemoryClear()
    {
        lock(typeof(MyCalculator))
        {
            Memory.RemoveInstance(CalculatorContext.Id);
        }
    }
    //Rest of the implementation
}

```

Context Bindings for Instance IDs

WCF provides dedicated bindings for passing custom context parameters. These bindings, called *context bindings*, are also explained in [Appendix B](#). Clients can use the `ContextClientBase<T>` class to pass the instance ID over the context binding protocol. Since the context bindings require a key and a value for every contextual parameter, the clients will need to provide both to the proxy. Using the same `IMyContract` as in [Example 4-11](#), such a proxy will look like this:

```

class MyContractClient : ContextClientBase<IMyContract>, IMyContract
{
    public MyContractClient(string key, string instanceId) : base(key,instanceId)
    {}
    public MyContractClient(string key, string instanceId, string endpointName) :
        base(key,instanceId,endpointName)
    {}

    //More constructors

    public void MyMethod()
    {
        Channel.MyMethod();
    }
}

```

Note that the context protocol only supports strings for keys and values. Because the value of the key must be known to the service in advance, the client might as well hardcode the same key in the proxy itself. The service can then retrieve the instance ID using my `ContextManager` helper class (described in [Appendix B](#)). As with message headers, the service can also encapsulate the interaction with `ContextManager` in a dedicated context class.

Example 4-13 shows the general pattern for passing an instance ID over the context bindings. Note that the proxy hardcodes the key for the instance ID, and that the same ID is known to the service.

Example 4-13. Passing the instance ID over a context binding

```
//Client-side
class MyContractClient : ContextClientBase<IMyContract>, IMyContract
{
    public MyContractClient(string instanceId) : base("MyKey",instanceId)
    {}

    public MyContractClient(string instanceId,string endpointName) :
        base( "MyKey",instanceId,endpointName)
    {}

    //More constructors

    public void MyMethod()
    {
        Channel.MyMethod();
    }
}

//Service-side
[ServiceBehavior(InstanceContextMode = InstanceContextMode.PerCall)]
class MyService : IMyContract
{
    public void MyMethod()
    {
        string instanceId = ContextManager.GetContext("MyKey");

        GetState(instanceId);
        DoWork();
        SaveState(instanceId);
    }

    void DoWork()
    {...}

    //Get and set state from durable storage
    void GetState(string instanceId)
    {...}

    void SaveState(string instanceId)
    {...}
}
```

Example 4-14 shows the matching concrete calculator example.

Example 4-14. Calculator with instance ID over context binding

```
//Client-side
class MyCalculatorClient : ContextClientBase<ICalculator>, ICalculator
{
    public MyCalculatorClient(string instanceId) : base("CalculatorId",instanceId)
    {}
    public MyCalculatorClient(string instanceId, string endpointName) :
        base("CalculatorId",instanceId,endpointName)
    {}

    //More constructors

    public double Add(double number1, double number2)
    {
        return Channel.Add(number1, number2);
    }
    public void MemoryStore(double number)
    {
        Channel.MemoryStore(number);
    }

    //Rest of the implementation
}
//Service-side
static class CalculatorContext
{
    public static string Id
    {
        get
        {
            return ContextManager.GetContext("CalculatorId") ?? String.Empty;
        }
    }
}

[ServiceBehavior(InstanceContextMode = InstanceContextMode.PerCall)]
class MyCalculator : ICalculator
{
    //Same as Example 4-12
}
```

Using the standard ID for context binding

The need to hardcode and know in advance the key used for the instance ID is a liability. The context bindings were designed with durable services in mind, so every context binding always contains an autogenerated instance ID in the form of a Guid (in string format), accessible via the reserved key of `instanceId`. The client and the service will see the same value for the instance ID. The value is initialized once the first call on the proxy returns, after the binding has had the chance to correlate it

between the client and the service. Like any other parameter passed over a context binding, the value of the instance ID is immutable throughout the life of the proxy.

To streamline interacting with the standard instance ID, I extended `ContextManager` with ID management methods, properties, and proxy extension methods, as shown in [Example 4-15](#).

Example 4-15. Standard instance ID management with ContextManager

```
public static class ContextManager
{
    public const string InstanceIdKey = "instanceId";

    public static Guid InstanceId
    {
        get
        {
            string id = GetContext(InstanceIdKey) ?? Guid.Empty.ToString();
            return new Guid(id);
        }
    }

    public static Guid GetInstanceId(IClientChannel innerChannel)
    {
        try
        {
            string instanceId =
                innerChannel.GetProperty<IContextManager>().GetContext()[InstanceIdKey];
            return new Guid(instanceId);
        }
        catch(KeyNotFoundException)
        {
            return Guid.Empty;
        }
    }

    public static void SetInstanceId(IClientChannel innerChannel,Guid instanceId)
    {
        SetContext(innerChannel,InstanceIdKey,instanceId.ToString());
    }

    public static void SaveInstanceId(Guid instanceId,string fileName)
    {
        using(Stream stream =
              new FileStream(fileName,FileMode.OpenOrCreate,FileAccess.Write))
        {
            IFormatter formatter = new BinaryFormatter();
            formatter.Serialize(stream,instanceId);
        }
    }

    public static Guid LoadInstanceId(string fileName)
    {
        try
```

```

    {
        using(Stream stream = new FileStream(fileName, FileMode.Open,
                                              FileAccess.Read))
        {
            IFormatter formatter = new BinaryFormatter();
            return (Guid)formatter.Deserialize(stream);
        }
    }
    catch
    {
        return Guid.Empty;
    }
}
//More members
}

```

`ContextManager` offers the `GetInstanceId()` and `SetInstanceId()` methods to enable the client to read an instance ID from and write it to the context. The service uses the `InstanceId` read-only property to obtain the ID. `ContextManager` adds type safety by treating the instance ID as a `Guid` and not as a `string`. It also adds error handling.

Finally, `ContextManager` provides the `LoadInstanceId()` and `SaveInstanceId()` methods to read the instance ID from and write it to a file. These methods are handy on the client side to store the ID between client application sessions against the service.

While the client can use `ContextClientBase<T>` (as in [Example 4-13](#)) to pass the standard ID, it is better to tighten it and provide built-in support for the standard instance ID, as shown in [Example 4-16](#).

Example 4-16. Extending ContextClientBase<T> to support standard IDs

```

public abstract class ContextClientBase<T> : ClientBase<T> where T : class
{
    public Guid InstanceId
    {
        get
        {
            return ContextManager.GetInstanceId(InnerChannel);
        }
    }
    protected ContextClientBase(Guid instanceId) :
        this(ContextManager.InstanceIdKey,instanceId.ToString())
    {}

    public ContextClientBase(Guid instanceId,string endpointName) :
        this(ContextManager.InstanceIdKey,instanceId.ToString(),endpointName)
    {}
}

```

```
//More constructors  
}
```

Example 4-17 shows the calculator client and service using the standard ID.

Example 4-17. Calculator using standard ID

```
//Client-side  
class MyCalculatorClient : ContextClientBase<ICalculator>, ICalculator  
{  
    public MyCalculatorClient()  
    {}  
    public MyCalculatorClient(Guid instanceId) : base(instanceId)  
    {}  
    public MyCalculatorClient(Guid instanceId, string endpointName) :  
        base(instanceId, endpointName)  
    {}  
  
    //Rest same as Example 4-14  
}  
//Service-side  
[ServiceBehavior(InstanceContextMode = InstanceContextMode.PerCall)]  
class MyCalculator : ICalculator  
{  
    static IIInstanceStore<Guid, double> Memory =  
        new FileInstanceStore<Guid, double>(Settings.Default.MemoryFileName);  
  
    public double Add(double number1, double number2)  
    {  
        return number1 + number2;  
    }  
    public void MemoryStore(double number)  
    {  
        lock(typeof(MyCalculator))  
        {  
            Memory[ContextManager.InstanceId] = number;  
        }  
    }  
    public void MemoryClear()  
    {  
        lock(typeof(MyCalculator))  
        {  
            Memory.RemoveInstance(ContextManager.InstanceId);  
        }  
    }  
    //Rest of the implementation  
}
```

Automatic Durable Behavior

All the techniques shown so far for durable services require a nontrivial amount of work by the service—in particular, providing a durable state storage and explicitly managing the instance state against it in every operation. Given the repetitive nature of this work, WCF can automate it for you, and serialize and deserialize the service state on every operation from an indicated state store, using the standard instance ID.

When you let WCF manage your instance state, it follows these rules:

- If the client does not provide an ID, WCF will create a new service instance by exercising its constructor. After the call, WCF will serialize the instance to the state store.
- If the client provides an ID to the proxy and the store already contains state matching that ID, WCF will not call the instance constructor. Instead, the call will be serviced on a new instance deserialized out of the state store.
- When the client provides a valid ID, for every operation WCF will deserialize an instance out of the store, call the operation, and serialize the new state modified by the operation back to the store.
- If the client provides an ID not found in the state store, WCF will throw an exception.

The durable service behavior attribute

To enable this automatic durable behavior, WCF provides the `DurableService` behavior attribute, defined as:

```
public sealed class DurableServiceAttribute : Attribute, IServiceBehavior, ...  
{ ... }
```

You apply this attribute directly on the service class. Most importantly, the service class must be marked either as serializable or as a data contract with the `DataMember` attribute on all members requiring durable state management:

```
[Serializable]  
[DurableService]  
class MyService : IMyContract  
{  
    /* Serializable member variables only */  
  
    public void MyMethod()  
    {  
        //Do work  
    }  
}
```

The instance can now manage its state in member variables, just as if it were a regular instance, trusting WCF to manage those members for it. If the service is not marked as serializable (or a data contract), the first call to it will fail once WCF tries to serialize it to the store. Any service relying on automatic durable state management must be configured as per-session, yet it will always behave as a per-call service (WCF uses context deactivation after every call). In addition, the service must use one of the context bindings with every endpoint to enable the standard instance ID, and the contract must allow or require a transport session, but cannot disallow it. These two constraints are verified at service load time.

The durable operation behavior attribute

A service can optionally use the `DurableOperation` behavior attribute to instruct WCF to purge its state from the store at the end of the workflow:

```
[AttributeUsage(AttributeTargets.Method)]
public sealed class DurableOperationAttribute : Attribute, ...
{
    public bool CanCreateInstance
    {get;set;}

    public bool CompletesInstance
    {get;set;}
}
```

Setting the `CompletesInstance` property to `true` instructs WCF to remove the instance ID from the store once the operation call returns. The default value of the `CompletesInstance` property is `false`. In case the client does not provide an instance ID, you can also prevent an operation from creating a new instance by setting the `CanCreateInstance` property to `false`. [Example 4-18](#) demonstrates the use of the `CompletesInstance` property on the `MemoryClear()` operation of the calculator.

Example 4-18. Using `CompletesInstance` to remove the state

```
[Serializable]
[DurableService]
class MyCalculator : ICalculator
{
    double Memory
    {get;set;}

    public double Add(double number1,double number2)
    {
        return number1 + number2;
    }
    public void MemoryStore(double number)
    {
        Memory = number;
```

```

}
[DurableOperation(CompletesInstance = true)]
public void MemoryClear()
{
    Memory = 0;
}
//Rest of the implementation
}

```

The problem with relying on `CompletesInstance` is that the context ID is immutable. This means that if the client tries to make additional calls on the proxy after calling an operation for which `CompletesInstance` is set to `true`, all of those calls will fail, since the store will no longer contain the instance ID. The client must be aware, therefore, that it cannot continue to use the same proxy: if the client wants to make further calls against the service, it must do so on a new proxy that does not have an instance ID yet, and by doing so, the client will start a new workflow. One way of enforcing this is to simply close the client program after completing the workflow (or create a new proxy reference). Using the proxy definition of [Example 4-17](#), [Example 4-19](#) shows how to manage the calculator proxy after clearing the memory while seamlessly continuing to use the proxy.

Example 4-19. Resetting the proxy after completing a workflow

```

class CalculatorProgram
{
    MyCalculatorClient m_Proxy;

    public CalculatorProgram()
    {
        Guid calculatorId =
            ContextManager.LoadInstanceId(Settings.Default.CalculatorIdFileName);

        m_Proxy = new MyCalculatorClient(calculatorId);
    }

    public void Add()
    {
        m_Proxy.Add(2,3);
    }

    public void MemoryClear()
    {
        m_Proxy.MemoryClear();

        ResetDurableSession(ref m_Proxy);
    }

    public void Close()
    {
        ContextManager.SaveInstanceId(m_Proxy.InstanceId,
                                      Settings.Default.CalculatorIdFileName);
        m_Proxy.Close();
    }
}

```

```

}
void ResetDurableSession(ref MyCalculatorClient proxy)
{
    ContextManager.SaveInstanceId(Guid.Empty,
        Settings.Default.CalculatorIdFileName);
    Binding binding = proxy.Endpoint.Binding;
    EndpointAddress address = proxy.Endpoint.Address;

    proxy.Close();

    proxy = new MyCalculatorClient(binding, address);
}
}

```

Example 4-19 uses my `ContextManager` helper class to load an instance ID and save it to a file. The constructor of the client program creates a new proxy using the ID found in the file. As shown in [Example 4-15](#), if the file does not contain an instance ID, `LoadInstanceId()` returns `Guid.Empty`. My `ContextClientBase<T>` is designed to expect an empty GUID for the context ID: if an empty GUID is provided, `ContextClientBase<T>` constructs itself without an instance ID, thus ensuring a new workflow. After clearing the memory of the calculator, the client calls the `ResetDurableSession()` helper method. `ResetDurableSession()` first saves an empty GUID to the file, and then duplicates the existing proxy. It copies the old proxy's address and binding, closes the old proxy, and sets the proxy reference to a new proxy constructed using the same address and binding as the old one and with an implicit empty GUID for the instance ID.

Programmatic instance management

WCF offers a simple helper class for durable services called `DurableOperationContext`:

```

public static class DurableOperationContext
{
    public static void AbortInstance();
    public static void CompleteInstance();
    public static Guid InstanceId
    {get;}
}

```

The `CompleteInstance()` method lets the service programmatically (instead of declaratively via the `DurableOperation` attribute) complete the instance and remove the state from the store once the call returns. `AbortInstance()`, on the other hand, cancels any changes made to the store during the call, as if the operation was never called. The `InstanceId` property is similar to `ContextManager.InstanceId`.

Persistence providers

While the `DurableService` attribute instructs WCF when to serialize and deserialize the instance, it does not say anything about where to do so, or for that matter, provide any information about the state storage. WCF actually uses a bridge pattern in the form of a provider model, which lets you specify the state store separately from the attribute. The attribute is thus decoupled from the store, allowing you to rely on the automatic durable behavior for any compatible storage.

If a service is configured with the `DurableService` attribute, you must configure its host with a persistence provider factory. The factory derives from the abstract class `PersistenceProviderFactory`, and it creates a subclass of the abstract class `PersistenceProvider`:

```
public abstract class PersistenceProviderFactory : CommunicationObject
{
    protected PersistenceProviderFactory();
    public abstract PersistenceProvider CreateProvider(Guid id);
}

public abstract class PersistenceProvider : CommunicationObject
{
    protected PersistenceProvider(Guid id);

    public Guid Id
    {get; }

    public abstract object Create(object instance,TimeSpan timeout);
    public abstract void Delete(object instance,TimeSpan timeout);
    public abstract object Load(TimeSpan timeout);
    public abstract object Update(object instance,TimeSpan timeout);

    //Additional members
}
```

The most common way of specifying the persistence provider factory is to include it in the host config file as a service behavior, and to reference that behavior in the service definition:

```
<behaviors>
    <serviceBehaviors>
        <behavior name = "DurableService">
            <persistenceProvider
                type = "...type...,...assembly ..."
                <!-- Provider-specific parameters -->
            />
        </behavior>
    </serviceBehaviors>
</behaviors>
```

Once the host is configured with the persistence provider factory, WCF uses the created `PersistenceProvider` for every call to serialize and deserialize the instance. If no persistence provider factory is specified, WCF aborts creating the service host.

Custom persistence providers

A nice way to demonstrate how to write a simple custom persistence provider is my `FilePersistenceProviderFactory`, defined as:

```
public class FilePersistenceProviderFactory : PersistenceProviderFactory
{
    public FilePersistenceProviderFactory();
    public FilePersistenceProviderFactory(string fileName);
    public FilePersistenceProviderFactory(NameValueCollection parameters);
}
public class FilePersistenceProvider : PersistenceProvider
{
    public FilePersistenceProvider(Guid id, string fileName);
}
```

`FilePersistenceProvider` wraps my `FileInstanceStore<ID,T>` class. The constructor of `FilePersistenceProviderFactory` requires you to specify the desired filename. If no filename is specified, `FilePersistenceProviderFactory` defaults the filename to *Instances.bin*.

The key for using a custom persistence factory in a config file is to define a constructor that takes a `NameValueCollection` of parameters. These parameters are simple text-formatted pairs of the keys and values specified in the provider factory behavior section in the config file. Virtually any free-formed keys and values will work. For example, here's how to specify the filename:

```
<behaviors>
    <serviceBehaviors>
        <behavior name = "Durable">
            <persistenceProvider
                type = "FilePersistenceProviderFactory,ServiceModelEx"
                fileName = "MyService.bin"
            />
        </behavior>
    </serviceBehaviors>
</behaviors>
```

The constructor can then use the `parameters` collection to access these parameters:

```
string fileName = parameters["fileName"];
```

The SQL Server persistence provider

WCF ships with a persistence provider, which stores the instance state in a dedicated SQL Server table. After a default installation, the installation scripts for the database

are found in `C:\Windows\Microsoft.NET\Framework\<.NET version>\SQL\EN`. Note that with the WCF-provided SQL persistence provider you can only use SQL Server 2005 or later for state storage. The SQL provider comes in the form of `SqlPersistenceProviderFactory` and `SqlPersistenceProvider`, found in the `System.WorkflowServices` assembly under the `System.ServiceModel.Persistence` namespace.

All you need to do is specify the SQL provider factory and the connection string name:

```
<connectionStrings>
    <add name = "DurableServices"
        connectionString = "..."
        providerName = "System.Data.SqlClient"
    />
</connectionStrings>

<behaviors>
    <serviceBehaviors>
        <behavior name = "Durable">
            <persistenceProvider type =
                "System.ServiceModel.Persistence.SqlPersistenceProviderFactory,
                System.WorkflowServices,Version=4.0.0.0,Culture=neutral,
                PublicKeyToken=31bf3856ad364e35"
                connectionStringName = "DurableServices"
            />
        </behavior>
    </serviceBehaviors>
</behaviors>
```

You can also instruct WCF to serialize the instances as text (instead of the default binary serialization), perhaps for diagnostics or analysis purposes:

```
<persistenceProvider
    type = "System.ServiceModel.Persistence.SqlPersistenceProviderFactory,
        System.WorkflowServices,Version=4.0.0.0,Culture=neutral,
        PublicKeyToken=31bf3856ad364e35"
    connectionStringName = "DurableServices"
    serializeAsText = "true"
/>
```



With .NET 4.5, Microsoft relocated the various Workflow Foundation V3.0 types formally maintained in the `System.WorkflowServices` assembly. The move also included a deprecation warning stating that the Workflow Foundation `System.Activities` assembly in .NET 4.5 should provide replacement types. However, the assembly does not contain replacement types.

Normally it is poor practice to use compiler directives to cover up compilation warnings in your code, unless you are compensating for an oversight such as in this case. You can clean up your build and remove the warning messages by applying the following compiler directive: [add the restore pragma after your code]

```
#pragma warning disable 618
```

Throttling

While it is not a direct instance management technique, *throttling* enables you to restrain client connections and the load they place on your service. You need throttling because software systems are not elastic, as shown in [Figure 4-7](#).

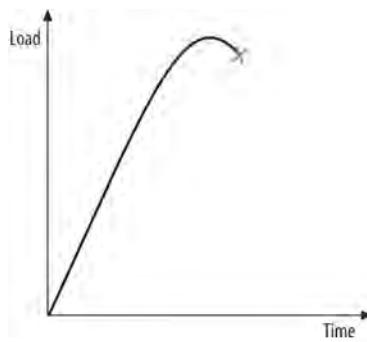


Figure 4-7. The inelastic nature of all software systems

In other words, you cannot keep increasing the load on the system and expect an infinite, gradual decline in its performance, as if stretching chewing gum. Most systems will initially handle the increase in load well, but then begin to yield and abruptly snap and break. All software systems behave this way, for reasons that are beyond the scope of this book and are related to queuing theory and the overhead inherent in managing resources. This snapping, inelastic behavior is of particular concern when there are spikes in load, as shown in [Figure 4-8](#).

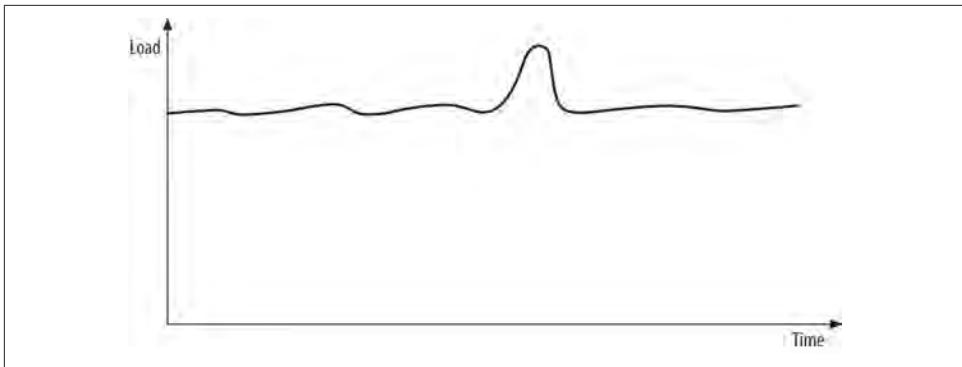


Figure 4-8. A spike in load may push the system beyond its design limit

Even if a system is handling a nominal load well (the horizontal line in [Figure 4-8](#)), a spike may push it beyond its design limit, causing it to snap and resulting in the clients experiencing a significant degradation in their level of service. Spikes can also pose a challenge in terms of the rate at which the load grows, even if the absolute level reached would not otherwise cause the system problems.

Throttling enables you to avoid maxing out your service and the underlying resources it allocates and uses. When throttling is engaged, if the settings you configure are exceeded, WCF will automatically place the pending callers in a queue and serve them out of the queue in order. If a client's call timeout expires while its call is pending in the queue, the client will get a `TimeoutException`. Throttling is inherently an unfair technique, because those clients whose requests are buffered will see a degradation in their level of service. However, in this case, it is better to be smart than just: if all the callers in the spike are allowed in, that will be fair, but all callers will then see a significant drop in the level of service as the system snaps. Throttling therefore makes sense when the area under the spike is relatively small compared with the area under the entire load graph, implying that the probability of the same caller being queued successively is very low. Every once in a while, in response to a spike, some callers will get buffered, but the system as a whole will still function well. Throttling does *not* work well when the load increases to a new level and remains constant at that level for a long time (as shown in [Figure 4-9](#)). In that case, all it does is defer the problems a bit, eventually causing all callers to time out. Such a system should be designed from the ground up to handle the higher level of load.

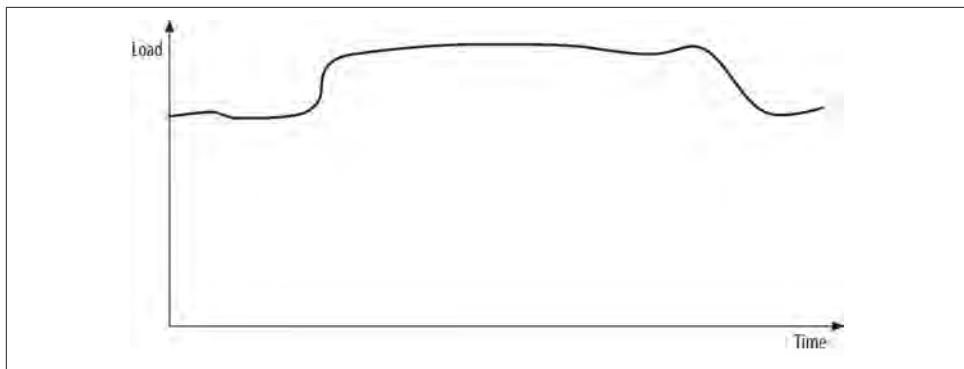


Figure 4-9. Inadequate justification for throttling

Throttling is done per service type; that is, it affects all instances of the service and all its endpoints. This is done by associating the throttle with every channel dispatcher the service uses.

WCF lets you control some or all of the following service consumption parameters:

Maximum number of concurrent sessions

Indicates the overall number of outstanding clients that can have a transport session with the service. In plain terms, this represents the maximum overall number of outstanding clients using TCP, IPC, or either of the WS bindings (with reliability, security, or both). Because the connectionless nature of a basic HTTP connection implies a very short transport session that exists only for the duration of the call, this number usually has no effect on clients using the basic binding or a WS binding without a transport session; such clients are instead limited by the maximum allowed number of concurrent calls. The default value is 100 times the processor (or cores) count.

Maximum number of concurrent calls

Limits the total number of calls that can currently be in progress across all service instances. This number should usually be kept at 1%–3% of the maximum number of concurrent sessions. The default value is 16 times the processor (or cores) count.

Maximum number of concurrent instances

Controls the total number of concurrently alive contexts. Unless you set this value explicitly, it will implicitly equate to the sum of the maximum concurrent calls and maximum of the concurrent sessions (116 times the processor count). Once set, it will retain its value regardless of the other two properties. How instances map to contexts is a product of the instance context management mode, as well as context and instance deactivation. With a per-session service, the maximum number of instances is both the total number of concurrently

active instances and the total number of concurrent sessions. When instance deactivation is employed, there may be far fewer instances than contexts, and yet clients will be blocked if the number of contexts has reached the maximum number of concurrent instances. With a per-call service, the number of instances is actually the same as the number of concurrent calls. Consequently, the maximum number of instances with a per-call service is the lesser of the configured maximum concurrent instances and maximum concurrent calls. The value of this parameter is ignored with a singleton service, since it can only have a single instance anyway.



Throttling is an aspect of hosting and deployment. When you design a service, you should make no assumptions about throttling configuration—always assume your service will bear the full brunt of the client's load. This is why, although it is fairly easy to write a throttling behavior attribute, WCF does not offer one.

Configuring Throttling

Administrators typically configure throttling in the config file. This enables you to throttle the same service code differently over time or across deployment sites. The host can also programmatically configure throttling based on some runtime decisions.

Administrative throttling

Example 4-20 shows how to configure throttling in the host config file. Using the `behaviorConfiguration` tag, you add to your service a custom behavior that sets throttled values.

Example 4-20. Administrative throttling

```
<system.serviceModel>
  <services>
    <service name = "MyService" behaviorConfiguration = "ThrottledBehavior">
      ...
    </service>
  </services>
  <behaviors>
    <serviceBehaviors>
      <behavior name = "ThrottledBehavior">
        <serviceThrottling
          maxConcurrentCalls      = "500"
          maxConcurrentSessions   = "10000"
          maxConcurrentInstances  = "100"
        />
      </behavior>
    </serviceBehaviors>
  </behaviors>
</system.serviceModel>
```

```
</serviceBehaviors>
</behaviors>
</system.serviceModel>
```

Programmatic throttling

The host process can programmatically throttle the service based on some runtime parameters. You can only configure the throttle programmatically before the host is opened. Although the host can override the throttling behavior found in the config file by removing it and adding its own, you typically should provide a programmatic throttling behavior only when there is no throttling behavior in the config file.

The `ServiceHostBase` class offers the `Description` property of the type `ServiceDescription`:

```
public abstract class ServiceHostBase : ...
{
    public ServiceDescription Description
    {get;}
    //More members
}
```

The service description, as its name implies, is a description of the service, with all its aspects and behaviors. `ServiceDescription` contains a property called `Behaviors` of the type `KeyedByTypeCollection<T>`, with `IServiceBehavior` as the generic parameter.

[Example 4-21](#) shows how to set the throttled behavior programmatically.

Example 4-21. Programmatic throttling

```
ServiceHost host = new ServiceHost(typeof(MyService));

ServiceThrottlingBehavior throttle;
throttle = host.Description.Behaviors.Find<ServiceThrottlingBehavior>();
if(throttle == null)
{
    throttle = new ServiceThrottlingBehavior();
    throttle.MaxConcurrentCalls      = 500;
    throttle.MaxConcurrentSessions  = 10000;
    throttle.MaxConcurrentInstances = 100;
    host.Description.Behaviors.Add(throttle);
}

host.Open();
```

First, the hosting code verifies that no service throttling behavior was provided in the config file. This is done by calling the `Find<T>()` method of `KeyedByTypeCollection<T>`, using `ServiceThrottlingBehavior` as the type parameter:

```

public class ServiceThrottlingBehavior : IServiceBehavior
{
    public int MaxConcurrentCalls
    {get;set;}
    public int MaxConcurrentSessions
    {get;set;}
    public int MaxConcurrentInstances
    {get;set;}
    //More members
}

```

If the returned throttle is null, then the hosting code creates a new `ServiceThrottlingBehavior`, sets its values, and adds it to the behaviors in the service description.

Streamlining with `ServiceHost<T>`

Using C# 3.0 extensions, you can extend `ServiceHost` (or any subclass of it, such as `ServiceHost<T>`) to automate the code in [Example 4-21](#), as shown in [Example 4-22](#).

Example 4-22. Extending `ServiceHost` to handle throttling

```

public static class ServiceThrottleHelper
{
    public static void SetThrottle(this ServiceHost host,
                                    int maxCalls,int maxSessions,int maxInstances)
    {
        ServiceThrottlingBehavior throttle = new ServiceThrottlingBehavior();
        throttle.MaxConcurrentCalls      = maxCalls;
        throttle.MaxConcurrentSessions  = maxSessions;
        throttle.MaxConcurrentInstances = maxInstances;
        host.SetThrottle(throttle);
    }
    public static void SetThrottle(this ServiceHost host,
                                  ServiceThrottlingBehavior serviceThrottle,
                                  bool overrideConfig)
    {
        if(host.State == CommunicationState.Opened)
        {
            throw new InvalidOperationException("Host is already opened");
        }
        ServiceThrottlingBehavior throttle =
            host.Description.Behaviors.Find<ServiceThrottlingBehavior>();
        if(throttle == null)
        {
            host.Description.Behaviors.Add(serviceThrottle);
            return;
        }
        if(overrideConfig == false)
        {
            return;
        }
    }
}

```

```

        host.Description.Behaviors.Remove(throttle);
        host.Description.Behaviors.Add(serviceThrottle);
    }
    public static void SetThrottle(this ServiceHost host,
                                  ServiceThrottlingBehavior serviceThrottle)
    {
        host.SetThrottle(serviceThrottle, false);
    }
}

```

`ServiceThrottleHelper` offers the `SetThrottle()` method, which accepts the throttle to use, and a Boolean flag indicating whether or not to override the configured values, if present. The default value (using an overloaded version of `SetThrottle()`) is `false`. `SetThrottle()` verifies that the host hasn't been opened yet using the `State` property of the `CommunicationObject` base class. If it is required to override the configured throttle, `SetThrottle()` removes it from the description. The rest of [Example 4-22](#) is similar to [Example 4-21](#). Here is how to use `ServiceHost<T>` to set a throttle programmatically:

```

ServiceHost<MyService> host = new ServiceHost<MyService>();
host.SetThrottle(12,34,56);
host.Open();

```



The `InProcFactory<T>` class presented in [Chapter 1](#) was similarly extended to streamline throttling.

Reading throttle values

Service developers can read the throttle values at runtime, for diagnostic and analytical purposes. For a service instance to access its throttle properties from its dispatcher at runtime, it must first obtain a reference to the host from the operation context.

The host base class `ServiceHostBase` offers the read-only `ChannelDispatchers` property:

```

public abstract class ServiceHostBase : CommunicationObject, ...
{
    public ChannelDispatcherCollection ChannelDispatchers
    {get;}
    //More members
}

```

`ChannelDispatchers` is a strongly typed collection of `ChannelDispatcherBase` objects:

```
public class ChannelDispatcherCollection :  
    SynchronizedCollection<ChannelDispatcherBase>  
{...}
```

Each item in the collection is of the type `ChannelDispatcher`. `ChannelDispatcher` offers the property `ServiceThrottle`:

```
public class ChannelDispatcher : ChannelDispatcherBase  
{  
    public ServiceThrottle ServiceThrottle  
    {get;set;}  
    //More members  
}  
public sealed class ServiceThrottle  
{  
    public int MaxConcurrentCalls  
    {get;set;}  
    public int MaxConcurrentSessions  
    {get;set;}  
    public int MaxConcurrentInstances  
    {get;set;}  
}
```

`ServiceThrottle` contains the configured throttle values:

```
class MyService : ...  
{  
    public void MyMethod() //Contract operation  
    {  
        ChannelDispatcher dispatcher = OperationContext.Current.Host.  
            ChannelDispatchers[0] as ChannelDispatcher;  
  
        ServiceThrottle serviceThrottle = dispatcher.ServiceThrottle;  
  
        Trace.WriteLine("Max Calls = " + serviceThrottle.MaxConcurrentCalls);  
        Trace.WriteLine("Max Sessions = " +  
            serviceThrottle.MaxConcurrentSessions);  
        Trace.WriteLine("Max Instances = " +  
            serviceThrottle.MaxConcurrentInstances);  
    }  
}
```

Note that the service can only read the throttle values and has no way of affecting them. If the service tries to set the throttle values, it will get an `InvalidOperationException`.

Again, you can streamline the throttle lookup via `ServiceHost<T>`. First, add a `ServiceThrottle` property:

```
public class ServiceHost<T> : ServiceHost  
{  
    public ServiceThrottle Throttle  
    {
```

```

    get
    {
        if(State == CommunicationState.Created)
        {
            throw new InvalidOperationException("Host is not opened");
        }

        ChannelDispatcher dispatcher = OperationContext.Current.Host.
                                         ChannelDispatchers[0] as ChannelDispatcher;
        return dispatcher.ServiceThrottle;
    }
}
//More members
}

```

Then, use `ServiceHost<T>` to host the service and use the `ServiceThrottle` property to access the configured throttle:

```

//Hosting code
ServiceHost<MyService> host = new ServiceHost<MyService>();
host.Open();

class MyService : ...
{
    public void MyMethod()
    {
        ServiceHost<MyService> host = OperationContext.Current.Host as
                                         ServiceHost<MyService>;

        ServiceThrottle serviceThrottle = host.Throttle;
        ...
    }
}

```



You can only access the `Throttle` property of `ServiceHost<T>` after the host is opened, because the dispatcher collection is initialized only after that point.

CHAPTER 5

Operations

The classic object- or component-oriented programming models offered only a single way for clients to call a method: the client would issue a call, block while the call was in progress, and continue executing once the method returned. Any other calling model had to be handcrafted, often incurring productivity and quality penalties. While WCF supports this classic invocation model, it also provides built-in support for additional operation types: one-way calls for fire-and-forget operations, duplex callbacks for allowing the service to call back to the client, and streaming to allow the client or the service to handle large payloads. In general, the type of operation used is part of the service contract and is an intrinsic part of the service design. The operation type may even place some constraints on the allowed bindings. Consequently, clients and services should be designed from the ground up with the operation type in mind, and you cannot switch easily between the various operation types. This chapter is dedicated to the various ways of invoking WCF operations and the related design guidelines.¹ Two other ways of invoking operations—asynchronously and queued—are addressed in subsequent chapters.

Request-Reply Operations

All the samples in the previous chapters included contracts whose operations were of the type known as *request-reply*. As the name implies, in these operations, the client issues a request in the form of a message and blocks until it gets the reply message. If the service does not respond within a default timeout of one minute, the client will get a `TimeoutException`. Request-reply is the default operation mode. Programming against request-reply operations is simple enough and resembles programming using

¹ This chapter contains excerpts from my article “WCF Essentials: What You Need to Know About One-Way Calls, Callbacks, and Events” (*MSDN Magazine*, October 2006).

the classic client/server model. The returned response message containing the results or returned values is converted to normal method return values. In addition, the proxy will throw an exception on the client side if there are any communication or service-side exceptions. Apart from the `NetPeerTcpBinding` and the `NetMsmqBinding`, all bindings support request-reply operations.

One-Way Operations

There are cases when an operation has no return value, and the client does not care about the success or failure of the invocation. To support this sort of fire-and-forget invocation, WCF offers *one-way* operations: once the client issues the call, WCF generates a request message, but no correlated reply message will ever return to the client. As a result, one-way operations cannot return values, and any exceptions thrown on the service side will not make their way to the client.

Ideally, when the client calls a one-way method, it should be blocked only for the briefest moment required to dispatch the call. However, in reality, one-way calls do not equate to asynchronous calls. When one-way calls reach the service, they may not be dispatched all at once but may instead be buffered on the service side to be dispatched one at a time, according to the service's configured concurrency mode behavior. (Chapter 8 discusses concurrency management and one-way calls in depth.) The number of messages the service can buffer (be they one-way or request-reply operations) is a product of the configured channel and reliability mode. If the number of messages exceeds the buffer's capacity, the client will be blocked even if it has issued a one-way call. However, once the call is deposited in the buffer, the client will be unblocked and can continue executing while the service processes the operation in the background.

It's also wrong to equate one-way calls with concurrent calls. If the client uses the same proxy yet utilizes multiple threads to invoke one-way calls, the calls may or may not execute concurrently on the service, and the exact nature of the interaction will be determined by the service concurrency management mode and the transport session (see Chapter 8 for more on this subject).

All of the WCF bindings support one-way operations.

Configuring One-Way Operations

The `OperationContract` attribute offers the Boolean `IsOneWay` property:

```
[AttributeUsage(AttributeTargets.Method)]
public sealed class OperationContractAttribute : Attribute
{
    public bool IsOneWay
    {get;set;}
```

```
    //More members  
}
```

`IsOneWay` defaults to `false`, which means a request-reply operation (hence the WCF default). However, setting `IsOneWay` to `true` configures the method as a one-way operation:

```
[ServiceContract]  
interface IMyContract  
{  
    [OperationContract(IsOneWay = true)]  
    void MyMethod();  
}
```

The client doesn't have to do anything special or different when invoking a one-way operation. The value of the `IsOneWay` property is reflected in the service metadata. Note that both the service contract definition and the definition imported by the client must have the same value for `IsOneWay`.

Because there is no reply associated with a one-way operation, there is no point in having any returned values or results. For example, here is an invalid definition of a one-way operation that returns a value:

```
//Invalid contract  
[ServiceContract]  
interface IMyContract  
{  
    [OperationContract(IsOneWay = true)]  
    int MyMethod();  
}
```

In fact, WCF enforces this by verifying the method signature when loading the host or opening the proxy and throwing an `InvalidOperationException` in the case of a mismatch.

One-Way Operations and Reliability

The fact that the client does not care about the result of the invocation does not mean that the client does not care whether the invocation took place at all. In general, you should turn on reliability for your services, even for one-way calls. This will ensure delivery of the requests to the service. However, the client may or may not care about the invocation order of the one-way operations. This is one of the main reasons why WCF allows you to separate enabling reliable delivery from enabling ordered delivery and execution of messages. Obviously, both the client and the service have to agree beforehand on these details, or the binding configuration will not match.

One-Way Operations and Sessionful Services

WCF will let you design a sessionful contract with one-way operations:

```

[ServiceContract(SessionMode = SessionMode.Required)]
interface IMyContract
{
    [OperationContract(IsOneWay = true)]
    void MyMethod();
}

```

With this configuration, if the client issues a one-way call and then closes the proxy while the method executes, the client will still be blocked until the operation completes.

Although they are technically possible, I believe that in general one-way operations in a sessionful contract (and per-session instantiation) indicate bad design. The reason is that having a session usually implies that the service manages state on behalf of the client. Any exception that may happen will be likely to fault that state, and yet the client may be unaware of it. In addition, typically the client (or the service) will choose a sessionful interaction because the contract used requires some lock-step execution advancing through some state machine. One-way calls do not fit this model very well. Consequently, I recommend that one-way operations should be applied on per-call or singleton services only.

If you do employ one-way operations on a sessionful contract, strive to make only the last operation terminating the session a one-way operation (and make sure it complies with one-way rules, such as having a `void` return type). You can use demarcating operations to enforce that:

```

[ServiceContract(SessionMode = SessionMode.Required)]
interface IOOrderManager
{
    [OperationContract]
    void SetCustomerId(int customerId);

    [OperationContract(IsInitiating = false)]
    void AddItem(int itemId);

    [OperationContract(IsInitiating = false)]
    decimal GetTotal();

    [OperationContract(IsOneWay = true, IsInitiating = false,
                      IsTerminating = true)]
    void ProcessOrders();
}

```

One-Way Operations and Exceptions

Although one-way operations do not return values or exceptions from the service itself, it's wrong to perceive them as a one-way street or a "black hole" from which nothing can come out. The client should still expect exceptions from a one-way call, and can even deduce that the call failed on the service. When dispatching a one-way

operation, any error because of communication problems (such as a wrong address or the host being unavailable) will throw an exception on the side of the client trying to invoke the operation. Furthermore, depending on the service instance mode and the binding used, the client may be affected by service-side exceptions. (The following discussion assumes that the service does not throw a `FaultException` or a derived exception, as discussed in [Chapter 6](#).)

When there is no transport session (for example, when using the `BasicHttpBinding` or the `WSHttpBinding` without reliable messaging and Message security), if an exception takes place during the invocation of a one-way operation, the client is unaffected and can continue to issue calls on the same proxy instance:

```
[ServiceContract]
interface IMyContract
{
    [OperationContract(IsOneWay = true)]
    void MethodWithError();

    [OperationContract]
    void MethodWithoutError();
}
class MyService : IMyContract
{
    public void MethodWithError()
    {
        throw new Exception();
    }
    public void MethodWithoutError()
    {}
}
//Client side without transport session:
MyContractClient proxy = new MyContractClient();
proxy.MethodWithError();
proxy.MethodWithoutError();
proxy.Close();
```

However, in the presence of a transport session, a service-side exception—including one thrown by a one-way operation—will fault the channel, and the client will not be able to issue any new calls using the same proxy instance:

```
[ServiceContract]
interface IMyContract
{
    [OperationContract(IsOneWay = true)]
    void MethodWithError();

    [OperationContract]
    void MethodWithoutError();
}
class MyService : IMyContract
```

```

{
    public void MethodWithError()
    {
        throw new Exception();
    }
    public void MethodWithoutError()
    {}
}
//Client side with transport session
MyContractClient proxy = new MyContractClient();
proxy.MethodWithError();
try
{
    proxy.MethodWithoutError(); //Will throw because channel faulted
    proxy.Close();
}
catch
{}
```

The client will not even be able to safely close the proxy.

A one-way call is therefore not fire-and-forget in nature, since the client can discover that something went wrong on the service during a one-way invocation. [Chapter 8](#) shows you how to use one-way operations for true asynchronous fire-and-forget operations.

Callback Operations

WCF supports allowing a service to call back to its clients. During a callback, in many respects the tables are turned: the service is the client, and the client becomes the service (see [Figure 5-1](#)). Callback operations can be used in a variety of scenarios and applications, but they are especially useful when it comes to events, or notifying the client(s) that some event has happened on the service side.

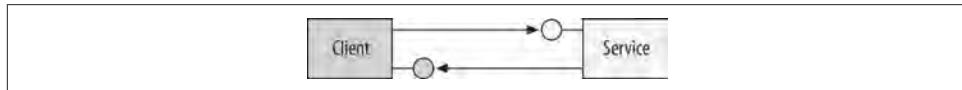


Figure 5-1. A callback allows the service to call back to the client

Callbacks are also commonly referred to as *duplex operations*. There are two immediate challenges to supporting duplex communication. First, how does the service know where the callback endpoint is? Second, how does the client facilitate hosting the callback object?

Not all bindings support callback operations. Only bidirectional-capable bindings support callback operations. For example, because of its connectionless nature, HTTP cannot be used for callbacks, and therefore you cannot use callbacks over the `BasicHttpBinding` or the `WSHttpBinding`. The only two commonly used bindings

that offer callbacks are the `NetTcpBinding` and the `NetNamedPipeBinding`, because by their very nature, the TCP and the IPC protocols support duplex communication.

The WSDualHttpBinding

To support callbacks over HTTP, WCF offers the `WSDualHttpBinding`, which actually sets up two WS channels: one for the calls from the client to the service and one for the calls from the service to the client. Duplex callbacks are nonstandard, as there is no industry reference that states how the client endpoint address is passed to the service or how the service publishes the callback contract in the first place. Duplex callbacks (including those over the WS dual binding) are a pure Microsoft feature. However, the `WSDualHttpBinding` is mostly unusable, since it is practically impossible to tunnel through various communication barriers separating the service from the client and the need to find a specific web-server machine makes this impractical. These connectivity problems are addressed with the Azure Service Bus, which supports duplex callbacks across the cloud using the `NetTcpRelayBinding`. The `NetTcpRelayBinding` by and large deprecates the `WSDualHttpBinding` in the vast majority of callback cases.

The Callback Contract

Callback operations are part of the service contract, and it is up to the service contract to define its own callback contract. A service contract can have at most one callback contract. Once defined, the clients are required to support the callback and provide the callback endpoint to the service in every call. To define a callback contract, the `ServiceContract` attribute offers the `CallbackContract` property of the type `Type`:

```
[AttributeUsage(AttributeTargets.Interface|AttributeTargets.Class)]
public sealed class ServiceContractAttribute : Attribute
{
    public Type CallbackContract
    {get;set;}
    //More members
}
```

When you define a service contract with a callback contract, you need to provide the `ServiceContract` attribute with the type of the callback contract and the definition of the callback contract, as shown in [Example 5-1](#).

Example 5-1. Defining and configuring a callback contract

```
interface ISomeCallbackContract
{
    [OperationContract]
```

```

    void OnCallback();
}

[ServiceContract(CallbackContract = typeof(ISomeCallbackContract))]
interface IMyContract
{
    [OperationContract]
    void DoSomething();
}

```

Note that the callback contract need not be marked with the `ServiceContract` attribute—the `ServiceContract` attribute is implied because it is defined as a callback contract and will be included in the service metadata. Of course, you still need to mark all the callback interface methods with the `OperationContract` attribute.

When the client imports the metadata of the callback contract, the imported callback interface will not have the same name on the client as in the original service-side definition. The name on the client will be the name of the service contract interface, suffixed with the word `Callback`. For example, a client that imports the definitions of [Example 5-1](#) will end up with these definitions instead:

```

interface IMyContractCallback
{
    [OperationContract]
    void OnCallback();
}
[ServiceContract(CallbackContract = typeof(IMyContractCallback))]
interface IMyContract
{
    [OperationContract]
    void DoSomething();
}

```



For simplicity's sake, I recommend using this naming convention also on the service side (i.e., naming the callback contract with the service contract interface name suffixed by `Callback`).

Client Callback Setup

It is up to the client to host the callback object and expose a callback endpoint. Recall from [Chapter 1](#) that the innermost execution scope of the service instance is the instance context. The `InstanceContext` class provides a constructor that takes the service instance to the host:

```

public sealed class InstanceContext : CommunicationObject, ...
{
    public InstanceContext(object implementation);
}

```

```

    public object GetServiceInstance();
    //More members
}

```

All the client needs to do to host a callback object is instantiate the callback object and construct a context around it:

```

class MyCallback : IMyContractCallback
{
    public void OnCallback()
    {...}
}
IMyContractCallback callback = new MyCallback();
InstanceContext context = new InstanceContext(callback);

```

It is also worth mentioning that although the callback methods are on the client side, they are WCF operations in every respect and therefore have an operation call context, accessible via `OperationContext.Current`.

Duplex proxies

Whenever a client is interacting with a service endpoint whose contract defines a callback contract, the client must use a proxy that will set up the bidirectional communication and pass the callback endpoint reference to the service. To that end, the proxy the client uses must derive from the specialized proxy class `DuplexClientBase<T>`, shown in [Example 5-2](#).

Example 5-2. The DuplexClientBase<T> class

```

public interface IDuplexContextChannel : IContextChannel
{
    InstanceContext CallbackInstance
    {get;set;}
    //More members
}
public abstract class DuplexClientBase<T> : ClientBase<T> where T : class
{
    protected DuplexClientBase(InstanceContext callbackContext);
    protected DuplexClientBase(InstanceContext callbackContext, string endpointName);
    protected DuplexClientBase(InstanceContext callbackContext, Binding binding,
                               EndpointAddress remoteAddress);
    protected DuplexClientBase(object callbackInstance);
    protected DuplexClientBase(object callbackInstance, string endpointName);
    protected DuplexClientBase(object callbackInstance, Binding binding,
                               EndpointAddress remoteAddress);

    public IDuplexContextChannel InnerDuplexChannel
    {get;}
    //More members
}

```

The client needs to provide the constructor of `DuplexClientBase<T>` with the instance context hosting the callback object (as well as the service endpoint information, as with a regular proxy). The proxy will construct an endpoint around the callback context, while inferring the details of the callback endpoint from the service endpoint configuration: the callback endpoint contract is the one defined by the service contract callback type. The callback endpoint will use the same binding (and transport) as the outgoing call. For the address, WCF will use the client's machine name. Simply passing the instance context to the duplex proxy and using the proxy to call the service will expose that client-side callback endpoint. To streamline the process, `DuplexClientBase<T>` also offers constructors that accept the callback object directly and wrap it internally with a context. If for any reason the client needs to access that context, `DuplexClientBase<T>` additionally offers the `InnerDuplexChannel` property of the type `IDuplexContextChannel`, which provides access to the context via the `CallbackInstance` property.

When you use Visual Studio to generate a proxy class targeting a service with a callback contract, the tool will generate a class that derives from `DuplexClientBase<T>`, as shown in [Example 5-3](#).

Example 5-3. VS 2010-generated duplex proxy

```
class MyContractClient : DuplexClientBase<IMyContract>, IMyContract
{
    public MyContractClient(InstanceContext callbackContext) : base(callbackContext)
    {}
    public MyContractClient(InstanceContext callbackContext, string endpointName) :
        base(callbackContext, endpointName)
    {}
    public MyContractClient(InstanceContext callbackContext, Binding binding,
        EndpointAddress remoteAddress) :
        base(callbackContext, binding, remoteAddress)
    {}
    //More constructors

    public void DoSomething()
    {
        Channel.DoSomething();
    }
}
```

Using that derived proxy class, the client can construct a callback instance, host it in a context, create a proxy, and call the service, thus passing the callback endpoint reference:

```
class MyCallback : IMyContractCallback
{
    public void OnCallback()
    {...}
```

```

}

IMyContractCallback callback = new MyCallback();
InstanceContext context = new InstanceContext(callback);

MyContractClient proxy = new MyContractClient(context);
proxy.DoSomething();

```

Note that as long as the client is expecting callbacks, the client cannot close the proxy. Doing so will close the callback endpoint and cause an error on the service side when the service tries to call back.

It is often the case that the client itself implements the callback contract, in which case the client will typically use a member variable for the proxy and close it when the client is disposed of, as shown in [Example 5-4](#).

Example 5-4. Client implementing the callback contract

```

class MyClient : IMyContractCallback, IDisposable
{
    MyContractClient m_Proxy;

    public void CallService()
    {
        InstanceContext context = new InstanceContext(this);
        m_Proxy = new MyContractClient(context);
        m_Proxy.DoSomething();
    }
    public void OnCallback()
    {...}

    public void Dispose()
    {
        m_Proxy.Close();
    }
}

```

The generated proxy does not take advantage of the streamlined constructors of `DuplexClientBase<T>` that accept the callback object directly, but you can rework the proxy manually to add that support, as shown in [Example 5-5](#).

Example 5-5. Using a reworked object-based proxy

```

class MyContractClient : DuplexClientBase<IMyContract>, IMyContract
{
    public MyContractClient(object callbackInstance) : base(callbackInstance)
    {}
    //More constructors
    public void DoSomething()
    {
        Channel.DoSomething();
    }
}

```

```

        }
    }
    class MyClient : IMyContractCallback, IDisposable
    {
        MyContractClient m_Proxy;

        public void CallService()
        {
            m_Proxy = new MyContractClient(this);
            m_Proxy.DoSomething();
        }
        public void OnCallback()
        {...}
        public void Dispose()
        {
            m_Proxy.Close();
        }
    }
}

```

Service-Side Callback Invocation

The client-side callback endpoint reference is passed along with every call the client makes to the service, and it is part of the incoming message. The `OperationContext` class provides the service with easy access to the callback reference via the generic method `GetCallbackChannel<T>()`:

```

public sealed class OperationContext : ...
{
    public T GetCallbackChannel<T>();
    //More members
}

```

Exactly what the service does with the callback reference and when it decides to use it is completely at the discretion of the service. The service can extract the callback reference from the operation context and store it for later use, or it can use it during the service operation to call back to the client. [Example 5-6](#) demonstrates the first option.

Example 5-6. Storing the callback references for later use

```

[ServiceBehavior(InstanceContextMode = InstanceContextMode.PerCall)]
class MyService : IMyContract
{
    static List<IMyContractCallback> m_Callbacks = new List<IMyContractCallback>();
    public void DoSomething()
    {
        IMyContractCallback callback = OperationContext.Current.
            GetCallbackChannel<IMyContractCallback>();

        if(m_Callbacks.Contains(callback) == false)
        {

```

```

        m_Callbacks.Add(callback);
    }
}
public static void CallClients()
{
    Action<IMyContractCallback> invoke = callback => callback.OnCallback();
    m_Callbacks.ForEach (invoke);
}
}

```

The service uses a static, generic list to store references to interfaces of the type `IMyContractCallback`. Because the service is not aware of which client is calling it and whether or not the client has called it already, in every call the service checks to see whether the list already contains the passed-in callback reference. If the list does not contain the reference, the service adds the callback to the list.

The service class also offers the static method `CallClients()`, which any party on the host side can use to call back to the clients:

```
MyService.CallClients();
```

Here, the invoking party is using some host-side thread for the callback invocation. That thread is unrelated to any thread executing the incoming service call.



Example 5-6 (and similar examples in this chapter) does not synchronize access to the callback list. Obviously, real application code will need to do that. Concurrency management (and, in particular, synchronizing access to shared resources) is discussed in [Chapter 8](#).

Callback reentrancy

The service may also want to invoke the callback reference that's passed in (or a saved copy of it) during the execution of a contract operation. However, such invocations are disallowed by default. The reason is the default service concurrency management. By default, the service class is configured for single-threaded access: the service instance context is associated with a lock, and only one thread at a time can own the lock and access the service instance inside that context. Calling out to the client during an operation call requires blocking the service thread and invoking the callback. The problem is that processing the reply message from the client on the same channel once the callback returns requires reentering the same context and negotiating ownership of the same lock, which will result in a deadlock. Note that the service may still invoke callbacks to other clients or call other services; it is the callback to its calling client that will cause the deadlock.

To prevent such a deadlock, if the single-threaded service instance tries to call back to its client, WCF will throw an `InvalidOperationException`. There are three possible

solutions. The first is to configure the service for multithreaded access. Callbacks to the calling client will then be allowed because the service instance will not be associated with a lock; however, this will increase the burden on the service developer, because of the need to provide synchronization for the service. The second solution is to configure the service for reentrancy. When configured for reentrancy, the service instance context is still associated with a lock, and only single-threaded access is allowed. However, if the service is calling back to its client, WCF will silently release the lock first. [Chapter 8](#) is dedicated to the synchronization modes and their implications on the programming model. For now, all you need to know is that you can set the concurrency behavior to either multithreaded or reentrant using the `ConcurrencyMode` property of the `ServiceBehavior` attribute:

```
public enum ConcurrencyMode
{
    Single, //Default
    Reentrant,
    Multiple
}

[AttributeUsage(AttributeTargets.Class)]
public sealed class ServiceBehaviorAttribute : ...
{
    public ConcurrencyMode ConcurrencyMode
    {get;set;}
    //More members
}
```

[Example 5-7](#) demonstrates a service configured for reentrancy. During the operation execution, the service accesses the operation context, grabs the callback reference to its calling client, and invokes it. Control will only return to the service once the callback returns, and the service's own thread will need to reacquire the lock.

Example 5-7. Configuring for reentrancy to allow callbacks

```
[ServiceContract(CallbackContract = typeof(IMyContractCallback))]
interface IMyContract
{
    [OperationContract]
    void DoSomething();
}

interface IMyContractCallback
{
    [OperationContract]
    void OnCallback();
}

[ServiceBehaviorConcurrencyMode = ConcurrencyMode.Reentrant]
class MyService : IMyContract
{
    public void DoSomething()
```

```

    {
        IMyContractCallback callback = OperationContext.Current.
            GetCallbackChannel<IMyContractCallback>();
        callback.OnCallback();
    }
}

```

The third solution that allows the service to safely call back to the calling client is to configure the callback contract operations as one-way operations. Doing so will enable the service to call back even when the concurrency mode is set to single-threaded, because there will not be any reply message to contend for the lock. **Example 5-8** demonstrates this configuration. Note that the service defaults to single-threaded concurrency mode.

Example 5-8. One-way callbacks are allowed by default

```

[ServiceContract(CallbackContract = typeof(IMyContractCallback))]
interface IMyContract
{
    [OperationContract]
    void DoSomething();
}

interface IMyContractCallback
{
    [OperationContract(IsOneWay = true)]
    void OnCallback();
}

class MyService : IMyContract
{
    public void DoSomething()
    {
        IMyContractCallback callback = OperationContext.Current.
            GetCallbackChannel<IMyContractCallback>();
        callback.OnCallback();
    }
}

```

Callback Connection Management

The callback mechanism supplies nothing like a higher-level protocol for managing the connection between the service and the callback endpoint. It is up to the developer to come up with some application-level protocol or a consistent pattern for managing the lifecycle of the connection. As mentioned previously, the service can only call back to the client if the client-side channel is still open, which is typically achieved by not closing the proxy. Keeping the proxy open will also prevent the callback object from being garbage-collected. If the service maintains a reference on a callback endpoint and the client-side proxy is closed or the client application itself is gone, when the service invokes the callback it will get an `ObjectDisposedException`

from the service channel. It is therefore preferable for the client to inform the service when it no longer wishes to receive callbacks or when the client application is shutting down. To that end, you can add an explicit `Disconnect()` method to the service contract. Since every method call carries the callback reference with it, in the `Disconnect()` method the service can remove the callback reference from its internal store.

In addition, for symmetry's sake, I recommend adding an explicit `Connect()` method. Having a `Connect()` method will enable the client to connect or disconnect multiple times, as well as provide a clearly delineated point in time as to when to expect a callback (only after a call to `Connect()`). [Example 5-9](#) demonstrates this technique. In both the `Connect()` and `Disconnect()` methods, the service needs to obtain the callback reference. In `Connect()`, the service verifies that the callback list does not already contain the callback reference before adding it to the list (this makes multiple calls to `Connect()` benign). In `Disconnect()`, the service verifies that the list contains the callback reference, and it throws an exception otherwise.

Example 5-9. Explicit callback connection management

```
[ServiceContract(CallbackContract = typeof(IMyContractCallback))]
interface IMyContract
{
    [OperationContract]
    void DoSomething();

    [OperationContract]
    void Connect();

    [OperationContract]
    void Disconnect();
}
interface IMyContractCallback
{
    [OperationContract]
    void OnCallback();
}
[ServiceBehavior(InstanceContextMode = InstanceContextMode.PerCall)]
class MyService : IMyContract
{
    static List<IMyContractCallback> m_Callbacks = new List<IMyContractCallback>();

    public void Connect()
    {
        IMyContractCallback callback = OperationContext.Current.
            GetCallbackChannel<IMyContractCallback>();
        if(m_Callbacks.Contains(callback) == false)
        {
            m_Callbacks.Add(callback);
        }
    }
}
```

```

    }
    public void Disconnect()
    {
        IMyContractCallback callback = OperationContext.Current.
            GetCallbackChannel<IMyContractCallback>();
        if(m_Callbacks.Contains(callback))
        {
            m_Callbacks.Remove(callback);
        }
        else
        {
            throw new InvalidOperationException("Cannot find callback");
        }
    }
    public static void CallClients()
    {
        Action<IMyContractCallback> invoke = callback => callback.OnCallback();
        m_Callbacks.ForEach(invoke);
    }
    public void DoSomething()
    {...}
}

```

Connection management and instance mode

A per-call service can use a callback reference during the operation call itself, or store it in some kind of a global repository such as a static variable, as you have seen in the examples so far. The per-call service must use some static variable to store the reference, since any instance state the service may use to store the reference will be gone when the operation returns. Using a `Disconnect()`-like method is therefore especially required for per-call services, as without it, the shared store will become bloated over time with dead callback references. A similar need exists with a singleton service. The singleton lifetime has no end, so it will accumulate an unlimited number of callback references, and as time goes by most of them will become stale because the callback clients will no longer be running. Having a `Disconnect()` method will keep the singleton connected only to the relevant alive clients.

Interestingly enough, a per-session service may get by without a `Disconnect()` method, as long as it maintains the callback reference in some instance member variable. The reason is that the service instance will automatically be disposed of when the session ends (when the client closes the proxy or times out), and there is no danger in keeping the reference throughout the session, as it is guaranteed to always be valid. However, if the sessionful service stores its callback reference in some global repository for the use of other host-side parties or across sessions, adding a `Disconnect()` method is required in order to remove the callback reference explicitly, because the callback reference is not available during the call to `Dispose()`.

You may also want to add the `Connect()` and `Disconnect()` pair on a sessionful service simply as a feature, because it enables the client to decide when to start or stop receiving callbacks during the session.

The Duplex Proxy and Type Safety

The WCF-provided `DuplexClientBase<T>` is not strongly typed to the callback interface used. The compiler will let you pass in any object, including an invalid callback interface. The compiler will even let you use for `T` a service contract type that has no callback contract defined at all. At runtime, you can successfully instantiate the proxy. The incompatibility will be discovered only when you try to use the proxy, yielding an `InvalidOperationException`. Much the same way, `InstanceContext` is object-based and is not verified at compile time to actually have a valid callback contract instance. When it's passed as a constructor parameter to the duplex proxy, there is no compile-time check to correlate the `InstanceContext` with the callback instance the duplex proxy expects, and the error will be discovered only when you try to use the proxy. You can use generics to compensate to some degree for these oversights and discover the error at runtime, as soon as you declare the proxy.

First, define the type-safe, generic `InstanceContext<T>` class, shown in [Example 5-10](#).

Example 5-10. The InstanceContext<T> class

```
public class InstanceContext<T>
{
    public InstanceContext Context
    {get;private set;}

    public InstanceContext(T callbackInstance)
    {
        Context = new InstanceContext(callbackInstance);
    }
    public void ReleaseServiceInstance()
    {
        Context.ReleaseServiceInstance();
    }
    public T ServiceInstance
    {
        get
        {
            return (T)Context.GetServiceInstance();
        }
    }
}
```

By using generics, you also provide type-safe access to the hosted callback object and capture the desired callback type.

Next, define a new type-safe, generic subclass of `DuplexClientBase<T>`, as shown in [Example 5-11](#).

Example 5-11. The `DuplexClientBase<T,C>` class

```
//T is the service contract and C is the callback contract
public abstract class DuplexClientBase<T,C> : DuplexClientBase<T> where T : class
{
    protected DuplexClientBase(InstanceContext<C> context) : base(context.Context)
    {}
    protected DuplexClientBase(InstanceContext<C> context, string endpointName) :
        base(context.Context, endpointName)
    {}
    protected DuplexClientBase(InstanceContext<C> context, Binding binding,
        EndpointAddress remoteAddress) :
        base(context.Context, binding, remoteAddress)
    {}
    protected DuplexClientBase(C callback) : base(callback)
    {}
    protected DuplexClientBase(C callback, string endpointName) :
        base(callback, endpointName)
    {}
    protected DuplexClientBase(C callback, Binding binding,
        EndpointAddress remoteAddress) :
        base(callback, binding, remoteAddress)
}

/* More constructors */

static DuplexClientBase()
{
    VerifyCallback();
}
internal static void VerifyCallback()
{
    Type contractType = typeof(T);
    Type callbackType = typeof(C);

    object[] attributes = contractType.GetCustomAttributes(
        typeof(ServiceContractAttribute), false);
    if(attributes.Length == 0)
    {
        throw new InvalidOperationException("Type of " + contractType +
            " is not a service contract");
    }
    ServiceContractAttribute serviceContractAttribute;
    serviceContractAttribute = attributes[0] as ServiceContractAttribute;
    if(callbackType != serviceContractAttribute.CallbackContract)
```

```
        {
            throw new InvalidOperationException("Type of " + callbackType +
                " is not configured as callback contract for " + contractType);
        }
    }
}
```

The `DuplexClientBase<T,C>` class uses two type parameters: `T` is used for the service contract type parameter and `C` is used for the callback contract type parameter. The constructors of `DuplexClientBase<T,C>` can accept either a raw `C` instance or an instance of `InstanceContext<C>` wrapping a `C` instance. These enable the compiler to ensure that only compatible contexts are used. However, .NET does not support a way to constrain a declarative relationship between `T` and `C`. The workaround is to perform a single runtime check before any use of `DuplexClientBase<T,C>`, and abort the use of the wrong type immediately, before any damage can be done. The trick is to place the runtime verification in the C# static constructor. The static constructor of `DuplexClientBase<T,C>` calls the static helper method `VerifyCallback()`. `VerifyCallback()` uses reflection to first verify that `T` is decorated with the `ServiceContract` attribute. Then it verifies that it has a type set for the callback contract that is the type parameter `C`. If not, an exception will be thrown in the static constructor, enabling you to discover the error as soon as possible at runtime.



Performing the callback contract verification in the static constructor is a technique applicable to any constraint that you cannot enforce at compile time, yet you have some programmatic way of determining and enforcing it at runtime.

Next, if you used Visual Studio to generate the proxy, you need to rework the proxy to derive from the type-safe `DuplexClientBase<T, C>` class:

```
class MyContractClient : DuplexClientBase<IMyContract, IMyContractCallback>,  
    IMyContract  
{  
    public MyContractClient(InstanceContext<IMyContractCallback> context) :  
        base(context)  
    {}  
    public MyContractClient(IMyContractCallback callback) : base(callback)  
    {}  
  
    /* Rest of the constructors */  
  
    public void DoSomething()  
    {  
        Channel.DoSomething();  
    }  
}
```

You can provide the reworked proxy either with a type-safe instance context, or with the callback instance directly:

```
//Client code
class MyCallback : IMyContractCallback
{...}

IMyContractCallback callback = new MyCallback();
MyContractClient proxy1 = new MyContractClient(callback);

InstanceContext<IMyContractCallback> context = new InstanceContext
                                                <IMyContractCallback>(callback);
MyContractClient proxy2 = new MyContractClient(context);
```

Either way, the compiler will verify that the type parameters provided to the proxy match the context type parameter or the callback instance, and the static constructor will verify the relationship between the service contract and the callback instance upon instantiation.

The Duplex Factory

Similar to the `ChannelFactory<T>` class, WCF also offers `DuplexChannelFactory<T>`, which can be used for setting up duplex proxies programmatically:

```
public class DuplexChannelFactory<T> : ChannelFactory<T>
{
    public DuplexChannelFactory(object callback);
    public DuplexChannelFactory(object callback, string endpointName);
    public DuplexChannelFactory(InstanceContext context, string endpointName);

    public T CreateChannel(InstanceContext context);
    public static T CreateChannel(object callback, string endpointName);
    public static T CreateChannel(InstanceContext context, string endpointName);
    public static T CreateChannel(object callback, Binding binding,
                                EndpointAddress endpointAddress);
    public static T CreateChannel(InstanceContext context, Binding binding,
                                EndpointAddress endpointAddress);
    //More members
}
```

`DuplexChannelFactory<T>` is used just like its base class, `ChannelFactory<T>`, except its constructors expect either a callback instance or a callback context. Note again the use of `object` for the callback instance and the lack of type safety. [Example 5-12](#) shows my reworked `DuplexChannelFactory<T,C>` class, which provides both compile-time and runtime type safety (similar to the fixed-up `DuplexClientBase<T>` class presented in [Example 5-11](#)).

Example 5-12. The DuplexChannelFactory<T,C> class

```
public class DuplexChannelFactory<T,C> : DuplexChannelFactory<T> where T : class
{
    static DuplexChannelFactory()
    {
        DuplexClientBase<T,C>.VerifyCallback();
    }

    public static T CreateChannel(C callback, string endpointName)
    {
        return DuplexChannelFactory<T>.CreateChannel(callback, endpointName);
    }
    public static T CreateChannel(InstanceContext<C> context, string endpointName)
    {
        return DuplexChannelFactory<T>.CreateChannel(context.Context, endpointName);
    }
    public static T CreateChannel(C callback, Binding binding,
                                EndpointAddress endpointAddress)
    {
        return DuplexChannelFactory<T>.CreateChannel(callback, binding,
                                                       endpointAddress);
    }
    public static T CreateChannel(InstanceContext<C> context, Binding binding,
                                EndpointAddress endpointAddress)
    {
        return DuplexChannelFactory<T>.CreateChannel(context, binding,
                                                       endpointAddress);
    }
    public DuplexChannelFactory(C callback) : base(callback)
    {}
    public DuplexChannelFactory(C callback, string endpointName) :
                                base(callback, endpointName)
    {}
    public DuplexChannelFactory(InstanceContext<C> context, string endpointName) :
                                base(context.Context, endpointName)
    {}
    //More constructors
}
```

As an example of utilizing the duplex channel factory, consider [Example 5-13](#), which adds callback ability to the InProcFactory static helper class presented in [Chapter 1](#).

Example 5-13. Adding duplex support to InProcFactory

```
public static class InProcFactory
{
    public static I CreateInstance<S,I,C>(C callback) where I : class
                                                where S : class, I
    {
        InstanceContext<C> context = new InstanceContext<C>(callback);
```

```

        return CreateInstance<S,I,C>(context);
    }
    public static I CreateInstance<S,I,C>(InstanceContext<C> context)
        where I : class
        where S : class,I
    {
        HostRecord hostRecord = GetHostRecord<S,I>();
        return DuplexChannelFactory<I,C>.CreateChannel(context,Binding,
                                                        hostRecord.Address);
    }
    //More members
}
//Sample client
IMyContractCallback callback = new MyCallback();

IMyContract proxy = InProcFactory.CreateInstance<MyService,IMyContract,
                                                IMyContractCallback>(callback);
proxy.DoSomething();
InProcFactory.CloseProxy(proxy);

```

Callback Contract Hierarchy

An interesting constraint on the design of callback contracts is that a service contract can designate a callback contract only if that contract is a sub-interface of all callback contracts defined by the contract's own base contracts. For example, here is an invalid callback contract definition:

```

interface ICallbackContract1
{...}

interface ICallbackContract2
{...}

[ServiceContract(CallbackContract = typeof(ICallbackContract1))]
interface IMyBaseContract
{...}

//Invalid
[ServiceContract(CallbackContract = typeof(ICallbackContract2))]
interface IMySubContract : IMyBaseContract
{...}

```

`IMySubContract` cannot designate `ICallbackContract2` as a callback contract because `ICallbackContract2` is not a sub-interface of `ICallbackContract1`, which `IMyBaseContract` (the base of `IMySubContract`) defines as its own callback contract.

The reason for this constraint is obvious: if a client passes an endpoint reference to a service implementation of `IMySubContract`, that callback reference must satisfy the callback type expected by `IMyBaseContract`. WCF verifies the callback contract hier-

archy at service load time and throws an `InvalidOperationException` in the case of a violation.

The straightforward way to satisfy the constraint is to reflect the service contract hierarchy in the callback contract hierarchy:

```
interface ICallbackContract1
{...}

interface ICallbackContract2 : ICallbackContract1
{...}

[ServiceContract(CallbackContract = typeof(ICallbackContract1))]
interface IMyBaseContract
{...}

[ServiceContract(CallbackContract = typeof(ICallbackContract2))]
interface IMSubContract : IMyBaseContract
{...}
```

However, you can also use multiple interface inheritance by a single callback contract and avoid mimicking the service contract hierarchy:

```
interface ICallbackContract1
{...}
interface ICallbackContract2
{...}
interface ICallbackContract3 : ICallbackContract2, ICallbackContract1
{...}

[ServiceContract(CallbackContract = typeof(ICallbackContract1))]
interface IMyBaseContract1
{...}
[ServiceContract(CallbackContract = typeof(ICallbackContract2))]
interface IMyBaseContract2
{...}
[ServiceContract(CallbackContract = typeof(ICallbackContract3))]
interface IMSubContract : IMyBaseContract1, IMyBaseContract2
{...}
```

Note, also, that a service can implement its own callback contract:

```
[ServiceContract(CallbackContract = typeof(IMyContractCallback))]
interface IMyContract
{...}
[ServiceContract]
interface IMyContractCallback
{...}
class MyService : IMyContract, IMyContractCallback
{...}
```

The service can even store a reference to itself in some callback store (if it wishes to be called back as if it were a client).

Events

The basic WCF callback mechanism does not indicate anything about the nature of the interaction between the client and the service. They may be equal peers in a commutative interaction, each calling and receiving calls from the other. However, the canonical use for duplex callbacks is with *events*. Events allow the client or clients to be notified about something that has occurred on the service side. An event may result from a direct client call, or it may be the result of something the service monitors. The service firing the event is called the *publisher*, and the client receiving the event is called the *subscriber*. Events are a required feature in almost any type of application, as shown in [Figure 5-2](#).

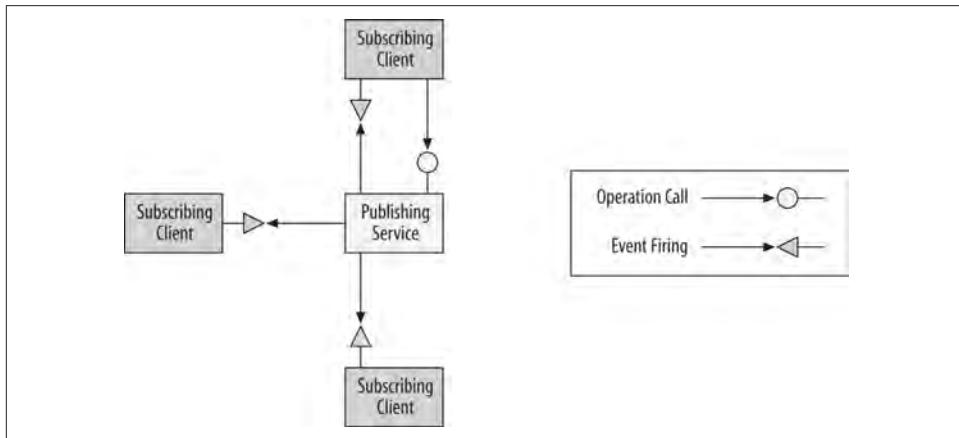


Figure 5-2. A publishing service can fire events at multiple subscribing clients

While events in WCF are nothing more than callback operations, by their very nature events usually imply a looser relationship between the publisher and the subscriber than the typical relationship between a client and a service. When dealing with events, the service typically publishes the same event to multiple subscribing clients. The publisher often does not care about the order of invocation of the subscribers, or any errors the subscribers might have while processing the events. All the publisher knows is that it should deliver the event to the subscribers. If they have a problem with the event, there is nothing the service can do about it anyway. In addition, the service does not care about returned results from the subscribers. Consequently, event-handling operations should have a `void` return type, should not have any outgoing parameters, and should be marked as one-way. I also recommend factoring the events to a separate callback contract, and not mixing events with regular callbacks in the same contract:

```
interface IMyEvents
{
    [OperationContract(IsOneWay = true)]
```

```

void OnEvent1();

[OperationContract(IsOneWay = true)]
void OnEvent2(int number);

[OperationContract(IsOneWay = true)]
void OnEvent3(int number,string text);
}

```

On the subscriber side, even when using one-way callback operations, the implementation of the event-handling methods should be of short duration. There are two reasons for this. First, if there is a large volume of events to publish, the publisher may get blocked if a subscriber has maxed out its ability to queue up callbacks because it is still processing the previous events. Blocking the publisher may prevent the event from reaching other subscribers in a timely manner. Second, if there are a large number of subscribers to the event, the accumulated processing time of each subscriber could exceed the publisher's timeout.

The publisher may add dedicated operations to its contract, allowing clients to explicitly subscribe to or unsubscribe from the events. If the publisher supports multiple event types, it may want to allow the subscribers to choose exactly which event(s) they want to subscribe to or unsubscribe from.

How the service internally goes about managing the list of subscribers and their preferences is a completely service-side implementation detail that should not affect the clients. The publisher can even use .NET delegates to manage the list of subscribers and the publishing act itself. [Example 5-14](#) demonstrates this technique, as well as the other design considerations discussed so far.

Example 5-14. Events management using delegates

```

enum EventType
{
    Event1 = 1,
    Event2 = 2,
    Event3 = 4,
    AllEvents = Event1|Event2|Event3
}
[ServiceContract(CallbackContract = typeof(IMyEvents))]
interface IMyContract
{
    [OperationContract]
    void DoSomething();

    [OperationContract]
    void Subscribe(EventType mask);

    [OperationContract]
    void Unsubscribe(EventType mask);
}

```

```

}

[ServiceBehavior(InstanceContextMode = InstanceContextMode.PerCall)]
class MyPublisher : IMyContract
{
    static Action m_Event1             = delegate{};
    static Action <int> m_Event2       = delegate{};
    static Action <int,string> m_Event3 = delegate{};

    public void Subscribe(EventType mask)
    {
        IMyEvents subscriber = OperationContext.Current.
            GetCallbackChannel<IMyEvents>();

        if((mask & EventType.Event1) == EventType.Event1)
        {
            m_Event1 += subscriber.OnEvent1;
        }
        if((mask & EventType.Event2) == EventType.Event2)
        {
            m_Event2 += subscriber.OnEvent2;
        }
        if((mask & EventType.Event3) == EventType.Event3)
        {
            m_Event3 += subscriber.OnEvent3;
        }
    }

    public void Unsubscribe(EventType mask)
    {
        //Similar to Subscribe() but uses -=
    }

    public static void FireEvent(EventType eventType)
    {
        switch(eventType)
        {
            case EventType.Event1:
            {
                m_Event1();
                return;
            }
            case EventType.Event2:
            {
                m_Event2(42);
                return;
            }
            case EventType.Event3:
            {
                m_Event3(42,"Hello");
                return;
            }
            default:
            {
                throw new InvalidOperationException("Unknown event type");
            }
        }
    }
}

```

```

        }
    }
}
public void DoSomething()
{...}
}

```

The service contract `IMyContract` defines the `Subscribe()` and `Unsubscribe()` methods. These methods each take an enum of the type `EventType`, whose individual fields are set to integer powers of 2. This enables the subscribing client to combine the values into a mask indicating the types of events it wants to subscribe to or unsubscribe from. For example, to subscribe to `Event1` and `Event3` but not `Event2`, the subscriber would call `Subscribe()` like this:

```

class MySubscriber : IMyEvents
{
    void OnEvent1()
    {...}
    void OnEvent2(int number)
    {...}
    void OnEvent3(int number,string text)
    {...}
}
IMyEvents subscriber = new MySubscriber();
InstanceContext context = new InstanceContext(subscriber);
MyContractClient proxy = new MyContractClient(context);
proxy.Subscribe(EventType.Event1|EventType.Event3);

```

Internally, `MyPublisher` maintains three static delegates, each corresponding to an event type.

Both the `Subscribe()` and `Unsubscribe()` methods check the supplied `EventType` value and either add the subscriber's callback to or remove it from the corresponding delegate. To fire an event, `MyPublisher` offers the static `FireEvent()` method. `FireEvent()` accepts the event to fire and invokes the corresponding delegate.

Again, the fact that the `MyPublisher` service uses delegates is purely an implementation detail simplifying event lookup. The service could have used a linked list, although that would require more complex code.



[Appendix D](#) presents a number of approaches for supporting a better design of events, called the publish-subscribe pattern.

Streaming

By default, when the client and the service exchange messages, these messages are buffered on the receiving end and delivered only once the entire message has been received. This is true whether it is the client sending a message to the service or the service returning a message to the client. As a result, when the client calls the service, the service is invoked only after the client's message has been received in its entirety; likewise, the client is unblocked only once the returned message with the results of the invocation has been received in its entirety.

For sufficiently small messages, this exchange pattern provides for a simple programming model because the latency caused by receiving the message is usually negligible compared with the message processing itself. However, when it comes to much larger messages—such as ones involving multimedia content, large files, or batches of data—blocking until the entire message has been received may be impractical. To handle such cases, WCF enables the receiving side (be it the client or the service) to start processing the data in the message while the message is still being received by the channel. This type of processing is known as *streaming transfer mode*. With large payloads, streaming provides improved throughput and responsiveness because neither the receiving nor the sending side is blocked while the message is being sent or received.

I/O Streams

For message streaming, WCF requires the use of the .NET Stream class. In fact, the contract operations used for streaming look just like conventional I/O methods. The Stream class is the base class of all the I/O streams in .NET (such as the FileStream, NetworkStream, and MemoryStream classes), allowing you to stream content from any of these I/O sources. All you need to do is return or receive a Stream as an operation parameter, as shown in [Example 5-15](#).

Example 5-15. Streaming operations

```
[ServiceContract]
interface IMyContract
{
    [OperationContract]
    Stream StreamReply1();

    [OperationContract]
    void StreamReply2(out Stream stream);

    [OperationContract]
    void StreamRequest(Stream stream);

    [OperationContract(IsOneWay = true)]
```

```
    void OneWayStream(Stream stream);  
}
```

Note that you can only define as an operation parameter the abstract class `Stream` or a specific serializable subclass such as `MemoryStream`. Subclasses such as `FileStream` are not serializable; you will have to use the base `Stream` instead.

WCF lets services stream the reply, the request, or both the request and the reply.

Streaming and Binding

Only the TCP, IPC, and basic HTTP bindings support streaming. With all of these bindings, streaming is disabled by default, and the binding will buffer the message in its entirety even when a `Stream` is used. You have to enable streaming by setting the `TransferMode` property according to the desired streaming mode. For example, when using the `BasicHttpBinding`:

```
public enum TransferMode  
{  
    Buffered, //Default  
    Streamed,  
    StreamedRequest,  
    StreamedResponse  
}  
public class BasicHttpBinding : Binding,...  
{  
    public TransferMode TransferMode  
    {get;set;}  
    //More members  
}
```

`TransferMode.Streamed` supports all streaming modes, and this is the only transfer mode that can support all the operations in [Example 5-15](#). However, if the contract contains only a specific type of streaming, such as streamed reply:

```
[ServiceContract]  
interface IMyContract  
{  
    //Stream reply  
    [OperationContract]  
    Stream GetStream1();  
  
    [OperationContract]  
    int MyMethod();  
}
```

you can have a buffered request and streamed reply by selecting `TransferMode.StreamedResponse`.

You will need to configure the binding on the client or service side (or both) per the required stream mode:

```
<configuration>
  <system.serviceModel>
    <client>
      <endpoint
        binding = "basicHttpBinding"
        bindingConfiguration = "StreamedHTTP"
        ...
      />
    </client>
    <bindings>
      <basicHttpBinding>
        <binding name = "StreamedHTTP"
          transferMode = "Streamed"
        />
      </basicHttpBinding>
    </bindings>
  </system.serviceModel>
</configuration>
```

Streaming and Transport

It is important to realize that WCF streaming is merely a programming model nicety. The underlying transport itself is not streamed, and the default maximum message size is set to 64K. This may be a problem with the sort of data you are likely to use streaming with, because streamed messages tend to be very large (hence the motivation for streaming in the first place). If the default limit proves insufficient, you can increase the maximum message size on the receiving side by setting the `MaxReceivedMessageSize` property to the expected maximum message size:

```
public class BasicHttpBinding : Binding, ...
{
  public long MaxReceivedMessageSize
  {get;set;}
  //More members
}
```

You can even use trial-and-error at runtime to find out what the streamed message size is and set the binding accordingly.

Typically, though, you would place that piece of configuration in the config file and avoid doing it programmatically, as message size tends to be deployment-specific:

```
<bindings>
  <basicHttpBinding>
    <binding name = "StreamedHTTP"
      transferMode = "Streamed"
      maxReceivedMessageSize = "120000"
    />
```

```
</basicHttpBinding>  
</bindings>
```



When you use streaming, you cannot use message-level transfer security. This is the main reason why only the TCP, IPC, and basic bindings (and their subclasses) support streaming: with these bindings, you typically do not (and, in the case of IPC, cannot) use message security. When streaming with the TCP binding, you also cannot enable reliable messaging. You'll see more on security in [Chapter 10](#).

Using streamed messages has a few additional implications. First, you need to synchronize access to the streamed content; for example, by opening the file stream in a read-only mode to allow other parties to access the file, or opening the stream in an exclusive mode to prevent others from accessing it if so required. In addition, you cannot use streaming when the contract is configured with `SessionMode.Required`.

Any service operation can, at any moment, encounter an unexpected error. The question is how (if at all) that error should be reported back to the client. Concepts such as exceptions and exception handling are technology-specific and should not transcend the service boundary. In addition, attempts by clients to handle errors invariably lead to increased coupling. Typically, error handling is a local implementation detail that should not affect the client, partly because the client may not care about the details of the errors (other than the fact that something went wrong), but mostly because in a well-designed application, the service is encapsulated so that the client does not have to do anything meaningful about the error anyway. A well-designed service should be as autonomous as possible, and should not depend on its clients to handle and recover errors. Anything beyond a blank error notification should in fact be part of the contractual interaction between the client and the service. This chapter describes just how the service and the client should handle these declared faults, and how you can extend and improve on the basic mechanism.

Error Isolation and Decoupling

In traditional .NET programming, any unhandled exception (except `ThreadAbortException`) immediately terminated the app domain (and thus, often the process) in which it occurred. While this is a very conservative behavior, it does not provide for proper fault isolation, which would enable the client to keep functioning even after the object blew up. Much the same way, after any unhandled error on the client side, the object would go down with the ship. Developers that did not like this had to provide for process (or app domain) isolation between the client and the object, which greatly complicated the programming model. That is not the WCF behavior, however. If a service call on behalf of one client causes an exception, it must not be allowed to take down the hosting process. Other clients accessing the service, or other services

hosted by the same process, should not be affected. As a result, when an unhandled exception leaves the service scope, the dispatcher silently catches and handles it by serializing it in the returned message to the client. When the returned message reaches the proxy, the proxy throws an exception on the client side. This behavior provides every WCF service with process-level isolation. The client and service can share a process, and yet be completely isolated as far as errors. The only exceptions that will take down the host process are critical errors that blow up .NET itself, such as stack overflows. Fault isolation, however, is only one of three key error-decoupling features of WCF. The second is error masking, and the third is faulting the channel.

Error Masking

The client can actually encounter three types of errors when trying to invoke a service. The first type of error is a communication error, which may occur because of network unavailability, an incorrect address, the host process not running, and so on. Communication exceptions are manifested on the client side by a `CommunicationException` or a `CommunicationException`-derived class such as `EndpointNotFoundException`.

The second type of error the client might encounter is related to the state of the proxy and the channels. There are many such possible exceptions. For example, these errors may occur when the client is trying to access an already closed proxy, resulting in an `ObjectDisposedException`; when there is a mismatch between the contract and the binding security protection level, resulting in an `InvalidOperationException`; when the client's credentials are denied by the service, resulting in a `SecurityNegotiationException` in case of authentication failure, or `SecurityAccessDeniedException` in case of authorization failure; or when the transport session times out, resulting in a `TimeoutException`.

The third type of error is an error that originates in the execution of the service call itself, as a result of either the service throwing an exception, or the service calling another object or resource and having that internal call throw an exception.

As stated at the beginning of this chapter, it is a common illusion that clients care about errors or have anything meaningful to do when they occur. Any attempt to bake such capabilities into the client creates an inordinate degree of coupling between the client and the object, raising serious design questions. How could the client possibly know more about the error than the service, unless it is tightly coupled to it? What if the error originated several layers below the service—should the client be coupled to those low-level layers? Should the client try the call again? How often and how frequently? Should the client inform the user of the error? Is there a user?

All that the client cares about is that something went wrong. The best practice for most clients is to simply let the exception go up the call chain. The topmost client

typically will catch the exception, not in order to handle it, but simply to prevent the application from shutting down abruptly. A well-designed client should never care about the actual error; WCF enforces this. In the interest of encapsulation and decoupling, by default all exceptions thrown on the service side always reach the client as `FaultExceptions`:

```
[...]
public class FaultException : CommunicationException
{...}
```

By having all service exceptions be indistinguishable from one another, WCF decouples the client from the service. The less the client knows about what happened on the service side, the more decoupled the interaction will be.

Channel Faulting

In traditional .NET programming, the client can catch the exception and keep calling the object. Consider this definition of a class and an interface:

```
interface IMyContract
{
    void MyMethod();
}
class MyClass : IMyContract
{...}
```

If the client snuffs out the exception thrown by the object, it can call it again:

```
IMyContract obj = new MyClass();
try
{
    obj.MyMethod();
}
catch
{}
obj.MyMethod();
```

This is a fundamental flaw of .NET as a platform. Exceptions, by their very nature, are for exceptional cases. Here, something totally unexpected and horrible has happened. How could the client possibly pretend otherwise? The object may be hopelessly broken, and yet the client keeps using it. In classic .NET, developers that did not approve of this behavior had to maintain a flag in each object, set the flag before throwing an exception (or after catching any downstream exceptions), and check the flag inside any public method, refusing to use the object if it was called after an exception had been thrown. This, of course, is cumbersome and tedious. WCF automates this best practice. If the service has a transport session, any unhandled exceptions (save those derived from `FaultException`, as described next) fault the channel (the proxy's state is changed to `CommunicationState.Faulted`), thus preventing the client from using

the proxy, and the object behind it, after an exception. In other words, for this service and proxy definition:

```
[ServiceContract]
interface IMyContract
{
    [OperationContract]
    void MyMethod();
}

class MyClass : IMyContract
{...}

class MyContractClient : ClientBase<IMyContract>, IMyContract
{...}
```

the following client code results in a `CommunicationObjectFaultedException`:

```
IMyContract proxy = new MyContractClient();
try
{
    proxy.MyMethod();
}
catch
{}
//Throws CommunicationObjectFaultedException
proxy.MyMethod();
```

The obvious conclusion is that the client should never try to use a WCF proxy after an exception. If there was a transport session, the client cannot even close the proxy.



If there is no transport-level session, the client can technically keep using the proxy after an exception, except again, it should not.

The only thing a client might safely do after an exception is to abort the proxy, perhaps to trigger tracing, or to raise events for state changes in the proxy, or to prevent others from using the proxy (even if there was no transport session):

```
MyContractClient proxy = new MyContractClient();
try
{
    proxy.MyMethod();
}
catch
{
    proxy.Abort();
}
```

The problem now is that you have to repeat this over and over for every method invocation. It is better to encapsulate this in the proxy itself:

```

class MyContractClient : ClientBase<IMyContract>, IMyContract
{
    public void MyMethod()
    {
        try
        {
            Channel.MyMethod();
        }
        catch
        {
            Abort();
            throw;
        }
    }
}

```

Closing the proxy and the using statement

I recommend against relying on the `using` statement to close the proxy. The reason is that in the presence of a transport session, any service-side exception will fault the channel. Trying to dispose of the proxy when the channel is faulted throws a `CommunicationObjectFaultedException`, so code after the `using` statement will never get called, even if you catch all exceptions inside the `using` statement:

```

using(MyContractClient proxy = new MyContractClient())
{
    try
    {
        proxy.MyMethod();
    }
    catch
    {}
}
Trace.WriteLine("This trace may never get called");

```

This reduces the readability of the code and may introduce defects, since the code will behave differently than most developers will expect. The only remedy is to encase the `using` statement itself in a `try/catch` statement:

```

try
{
    using(MyContractClient proxy = new MyContractClient())
    {
        try
        {
            proxy.MyMethod();
        }
        catch
        {}
    }
}

```

```

    catch
    {}
    Trace.WriteLine("This trace always gets called");

```

It is therefore far better to call `Close()`. In the case of an exception, the exception will skip over the call to `Close()`:

```

MyContractClient proxy = new MyContractClient();
proxy.MyMethod();
proxy.Close();

```

You can, of course, catch the exception, but now the code is readable:

```

MyContractClient proxy = new MyContractClient();

try
{
    proxy.MyMethod();
    proxy.Close();
}
catch
{
    proxy.Abort();
}
Trace.WriteLine("This trace always gets called");

```

Exceptions and instance management

When the service is configured as per-call or as sessionful (which mandates the use of a transport session), the client can never access the same instance after an exception occurs. With a per-call service this is, of course, always true, but with a sessionful service this is the result of faulting the channel and terminating the transport session. The one exception to the rule here is a singleton. When the client calls a singleton service and encounters an exception, the singleton instance is not terminated and continues running. If there was no transport session (or if the exception was a `FaultException`-derived class, as described next), the client can keep using the proxy to connect to the singleton object. Even if the channel is faulted, the client can create a new proxy instance and reconnect to the singleton.

In the case of a durable service, the `DurableService` attribute offers the `UnknownExceptionAction` property, defined as:

```

public enum UnknownExceptionAction
{
    TerminateInstance,
    AbortInstance
}

[AttributeUsage(AttributeTargets.Class)]
public sealed class DurableServiceAttribute : ...
{

```

```

public UnknownExceptionAction UnknownExceptionAction
{get;set;}
//More members
}

```

`UnknownExceptionAction` defaults to `UnknownExceptionAction.TerminateInstance`, meaning that any unhandled exception will not only fault the channel but also remove the instance state from the store, thus terminating the workflow. This behavior is analogous to simply faulting the channel with a regular service, preventing future use of the object. The value `UnknownExceptionAction.AbortInstance`, on the other hand, terminates the channel to the client but keeps the state in the store. While any changes made to the instance are not persisted, this value is analogous to not faulting the channel in the case of a regular service.

Fault Propagation

Although the default error-masking policy of WCF is a best practice, there are times when you should refrain from relying on it. This is typically the case when there is an existing application (or communication pattern) in place, and the service is required to throw particular exceptions as it processes inputs, reaches certain states, or encounters errors. The client is required to respond to these exceptions in a prescribed way. Obviously, controlling the flow of the application using exceptions is hardly a good idea, as it leads to nonstructured programming and couples the client to the service. And yet, the underlying requirements remain: the service is required to report specific errors to the client, and the default masking of the errors by WCF precludes that. Another fundamental problem pertaining to propagating the error to the client is that exceptions are technology-specific, and as such should not be shared across the service boundary. For seamless interoperability, you need a way to map technology-specific exceptions to some neutral error information. This representation is called a *SOAP fault*. SOAP faults are based on an industry standard that is independent of any technology-specific exceptions, such as CLR, Java, or C++ exceptions. To return a SOAP fault (or just a fault, for short), the service cannot throw a raw CLR exception. Instead, the service could throw an instance of the `FaultException<T>` class, defined in [Example 6-1](#).

Example 6-1. The FaultException<T> class

```

[Serializable]
[...]
public class FaultException : CommunicationException
{
    public FaultException();
    public FaultException(string reason);
    public FaultException(FaultReason reason);
    public virtual MessageFault CreateMessageFault();
}

```

```

        //More members
    }

[Serializable]
public class FaultException<T> : FaultException
{
    public FaultException(T detail);
    public FaultException(T detail, string reason);
    public FaultException(T detail, FaultReason reason);
    //More members
}

```

`FaultException<T>` is a specialization of `FaultException`, so any client that programs against `FaultException` will be able to handle `FaultException<T>` as well. Deriving `FaultException` from `CommunicationException` also enables clients to handle all communication and service-side exceptions in a single catch.

The type parameter `T` for `FaultException<T>` conveys the error details. The detailing type can be any type, and doesn't necessarily have to be an `Exception`-derived class. The only constraint is that the type can be serialized by WCF.

Example 6-2 demonstrates a simple calculator service that throws a `FaultException<DivideByZeroException>` in its implementation of the `Divide()` operation when asked to divide by zero.

Example 6-2. Throwing a `FaultException<T>`

```

[ServiceContract]
interface ICalculator
{
    [OperationContract]
    double Divide(double number1, double number2);
    //More methods
}

class Calculator : ICalculator
{
    public double Divide(double number1, double number2)
    {
        if(number2 == 0)
        {
            DivideByZeroException exception = new DivideByZeroException();
            throw new FaultException<DivideByZeroException>(exception);
        }
        return number1 / number2;
    }
    //Rest of the implementation
}

```

Instead of `FaultException<DivideByZeroException>`, the service could also have thrown a non-Exception-derived class:

```
throw new FaultException<double>(number2);
```

However, I find that using an Exception-derived detailing type is more in line with conventional .NET programming practices and results in more readable code. In addition, it allows for exception promotion, discussed later in this chapter.

The `reason` parameter passed to the constructor of `FaultException<T>` is used as the exception message. You can pass a `string` for the reason:

```
DivideByZeroException exception = new DivideByZeroException("number2 is 0");
throw new FaultException<DivideByZeroException>(exception, "Reason: " +
    exception.Message);
```

or you can pass a `FaultReason`, which is useful when localization is required.

Fault Contracts

By default, any exception thrown by a service reaches the client as a `FaultException`. This is the case even if the service throws a `FaultException<T>`. The reason is that anything beyond communication errors that the service wishes to share with the client must be part of the service contract in order for the service to inform WCF that it wishes to pierce the error mask. To that end, WCF provides *fault contracts*, which are a way for the service to list the types of errors it can throw. The idea is that these types of errors should be the same as the type parameters used with `FaultException<T>`, and by listing them in fault contracts, the service enables its WCF clients to distinguish between contracted faults and other errors.

The service defines its fault contracts using the `FaultContractAttribute`:

```
[AttributeUsage(AttributeTargets.Method, AllowMultiple = true, Inherited = false)]
public sealed class FaultContractAttribute : Attribute
{
    public FaultContractAttribute(Type detailType);
    //More members
}
```

You apply the `FaultContract` attribute directly on a contract operation, specifying the error detailing type, as shown in [Example 6-3](#).

Example 6-3. Defining a fault contract

```
[ServiceContract]
interface ICalculator
{
    [OperationContract]
    double Add(double number1, double number2);
```

```

[OperationContract]
[FaultContract(typeof(DivideByZeroException))]
double Divide(double number1,double number2);
//More methods
}

```

The effect of the `FaultContract` attribute is limited to the method it decorates. That is, only that method can throw that fault and have it propagated to the client.

In addition, if the operation throws an exception that is not in the contract, it will reach the client as a plain `FaultException`. To propagate the exception, the service must throw exactly the same detailing type listed in the fault contract. For example, to satisfy this fault contract definition:

```
[FaultContract(typeof(DivideByZeroException))]
```

The service must throw a `FaultException<DivideByZeroException>`. The service cannot even throw a subclass of the fault contract's detailing type and have it satisfy the contract:

```

[ServiceContract]
interface IMyContract
{
    [OperationContract]
    [FaultContract(typeof(Exception))]
    void MyMethod();
}

class MyService : IMyContract
{
    public void MyMethod()
    {
        //Will not satisfy contract
        throw new FaultException<DivideByZeroException>(new DivideByZeroException());
    }
}

```

The `FaultContract` attribute is configured to allow multiple usages, so you can list multiple fault contracts in a single operation:

```

[ServiceContract]
interface ICalculator
{
    [OperationContract]
    [FaultContract(typeof(InvalidOperationException))]
    [FaultContract(typeof(string))]
    double Add(double number1,double number2);

    [OperationContract]
    [FaultContract(typeof(DivideByZeroException))]
    double Divide(double number1,double number2);
}

```

```
    //More methods  
}
```

This enables the service to throw any of the exceptions in the contracts and have them propagate to the client.



You cannot provide a fault contract on a one-way operation, because there is no reply message returned from a one-way operation:

```
//Invalid definition  
[ServiceContract]  
interface IMyContract  
{  
    [OperationContract(IsOneWay = true)]  
    [FaultContract(...)]  
    void MyMethod();  
}
```

Trying to do so will result in an `InvalidOperationException` at service load time (or when the proxy is created).

Fault handling

The fault contracts are published along with the rest of the service metadata. When a WCF client imports that metadata, the contract definitions contain the fault contracts as well as the fault detailing type definition, including the relevant data contracts. This last point is important if the detailing type is some custom exception type with various dedicated fields.

The client can expect to catch and handle the imported fault types. For example, when you write a client against the contract shown in [Example 6-3](#), the client can catch `FaultException<DivideByZeroException>`:

```
CalculatorClient proxy = new CalculatorClient();  
try  
{  
    proxy.Divide(2,0);  
    proxy.Close();  
}  
  
catch(FaultException<DivideByZeroException> exception)  
{...}  
  
catch(FaultException exception)  
{...}  
  
catch(CommunicationException exception)  
{...}  
  
catch(TimeoutException exception)
```

```
{...}  
  
catch(Exception exception)  
{...}
```

Note that the client can still encounter communication exceptions, or any other exception thrown by the service.

The client can choose to treat all non-communication service-side exceptions uniformly by simply handling only the `FaultException` base exception:

```
CalculatorClient proxy = new CalculatorClient();  
try  
{  
    proxy.Divide(2,0);  
    proxy.Close();  
}  
  
catch(FaultException exception)  
{...}  
  
catch(CommunicationException exception)  
{...}
```



You can improve on the tools-generated proxy: manually change the definition of the imported contract by removing the fault contract on the client side. In this case, even when the service throws an exception listed in a service-side fault contract, the exception will manifest itself on the client as a plain `FaultException`, not as the contracted fault. This will decouple the client from the specific errors the service throws.

Faults and channels

Listing an expected error in a fault contract hardly makes it an exceptional unexpected case. As a result, when the service throws an exception listed in a service-side fault contract, the exception will not fault the communication channel. The client can catch that exception and continue using the proxy, or safely close the proxy. This enables the service class to treat the errors listed in the fault contracts differently from regular exceptions, knowing that they will not fault the channel. This ability is not limited to the service class, though. If any downstream .NET class the service invokes throws such an error, it will not fault the channel to the client. The problem is, how can downstream classes know about the fault contracts of the upstream services that call them? Clearly, having this knowledge present downstream introduces undesirable coupling into the system.

To support the ability of a downstream class to throw an exception without faulting the channel, WCF treats any `FaultException` (or `FaultException`-derived class) as a

special case that does not fault the channel. In fact, WCF does not actually treat `FaultException<T>` itself as a special case at all—the reason a fault listed in the contract does not fault the channel is because it is derived from `FaultException`, not because it is listed in the contract.

Even without any fault contracts, the service (or any downstream object it uses) can throw an instance of `FaultException` directly:

```
throw new FaultException("Some reason");
```

The `Message` property of the exception object on the client side will be set to the `reason` construction parameter of `FaultException`. I call this throwing an *unknown fault* that will not fault the communication channel, so the client can keep using the proxy as if the exception was part of a fault contract. Throwing an unknown fault also allows the client to handle the exception separately from any other communication error.



Any `FaultException<T>` thrown by the service will reach the client as either a `FaultException<T>` or a `FaultException`. If no fault contract is in place (or if `T` is not in the contract), both a `FaultException` and a `FaultException<T>` thrown by the service will reach the client as `FaultException`.

Fault Debugging

A deployed service should be decoupled as much as possible from its clients, declaring in the service fault contracts only the absolute bare minimum and providing as little information as possible about any errors that occur. However, during testing and debugging, it is very useful to include all exceptions in the information sent back to the client. In fact, for a test client, it is instrumental to know exactly which error was thrown as a result of a particular input or use case, to see if the test cases break the service as they should. In such a case, dealing with the all-encompassing yet opaque `FaultException` is simply inadequate. For that purpose, you should use the `ExceptionDetail` class, defined as:

```
[DataContract]
public class ExceptionDetail
{
    public ExceptionDetail(Exception exception);

    [DataMember]
    public string HelpLink
    {get;}

    [DataMember]
    public ExceptionDetail InnerException
    {get;}
```

```

[DataMember]
public string Message
{get;}

[DataMember]
public string StackTrace
{get;}

[DataMember]
public string Type
{get;}
}

```

You need to create an instance of `ExceptionDetail` and initialize it with the exception you want to propagate to the client. Next, instead of throwing the intended exception, throw a `FaultException<ExceptionDetail>` with the instance of `ExceptionDetail` as a construction parameter, and also provide the original exception's message as the fault reason. This sequence is shown in [Example 6-4](#).

Example 6-4. Including the service exception in the fault message

```

[ServiceContract]
interface IMyContract
{
    [OperationContract]
    void MethodWithError();
}
class MyService : IMyContract
{
    public void MethodWithError()
    {
        InvalidOperationException exception =
            new InvalidOperationException("Some error");
        ExceptionDetail detail = new ExceptionDetail(exception);
        throw new FaultException<ExceptionDetail>(detail,exception.Message);
    }
}

```

Doing so will enable the client to discover the original exception type and message. The client-side fault object will have a `Detail.Type` property containing the name of the original service exception, and the `Message` property will contain the original exception message. [Example 6-5](#) shows the client code processing the exception thrown in [Example 6-4](#).

Example 6-5. Processing the included exception

```

MyContractClient proxy = new MyContractClient();
try

```

```

{
    proxy.MethodWithError();
}
catch(FaultException<ExceptionDetail> exception)
{
    Debug.Assert(exception.Detail.Type ==
                    typeof(InvalidOperationException).ToString());
    Debug.Assert(exception.Message == "Some error");
}

```

Since `FaultException<ExceptionDetail>` derives from `FaultException`, throwing it will not fault the channel. I do not consider this the desired behavior.

Including exceptions declaratively

The `ServiceBehavior` attribute offers `IncludeExceptionDetailInFaults`, a Boolean property defined as:

```

[AttributeUsage(AttributeTargets.Class)]
public sealed class ServiceBehaviorAttribute : Attribute, ...
{
    [DefaultValue(false)]
    public bool IncludeExceptionDetailInFaults
    {get;set;}
    //More members
}

```

`IncludeExceptionDetailInFaults` defaults to `false`. Setting it to `true`, as in this snippet:

```

[ServiceBehavior(IncludeExceptionDetailInFaults = true)]
class MyService : IMyContract
{...}

```

has a similar effect as the code in [Example 6-4](#), only automated. All non-contractual faults and exceptions thrown by the service or any of its downstream objects are propagated to the client and included in the returned fault message for the client program to process them, as in [Example 6-5](#):

```

[ServiceBehavior(IncludeExceptionDetailInFaults = true)]
class MyService : IMyContract
{
    public void MethodWithError()
    {
        throw new InvalidOperationException("Some error");
    }
}

```

Any fault thrown by the service (or its downstream objects) that is listed in the fault contracts is unaffected and is propagated as-is to the client.

Another important difference between using the declarative support for including the fault details compared with manually throwing `FaultException<ExceptionDetail>` is that it will correctly fault the channel, preventing the client from reusing the proxy (if a transport session was present).

While including all exceptions is beneficial for debugging, great care should be taken to avoid shipping and deploying the service with `IncludeExceptionDetailInFaults` set to `true`. To avoid this potential pitfall automatically, you can use conditional compilation, as shown in [Example 6-6](#).

Example 6-6. Setting `IncludeExceptionDetailInFaults` to true in debug only

```
public static class DebugHelper
{
    public const bool IncludeExceptionDetailInFaults =
#if DEBUG
    true;
#else
    false;
#endif
}
[ServiceBehavior(IncludeExceptionDetailInFaults =
    DebugHelper.IncludeExceptionDetailInFaults)]
class MyService : IMyContract
{...}
```

Host and exception diagnostics

Obviously, including all exceptions in the fault message contributes greatly in debugging, but it's also useful when you're trying to analyze a problem in an already deployed service. Fortunately, you can set `IncludeExceptionDetailInFaults` to `true` both programmatically and administratively in the host config file. To set this behavior programmatically, before opening the host you need to find the service behavior in the service description and set the `IncludeExceptionDetailInFaults` property:

```
ServiceHost host = new ServiceHost(typeof(MyService));

ServiceBehaviorAttribute debuggingBehavior = host.Description.Behaviors.Find
    <ServiceBehaviorAttribute>();

debuggingBehavior.IncludeExceptionDetailInFaults = true;

host.Open();
```

You can streamline this procedure by encapsulating it in `ServiceHost<T>`, as shown in [Example 6-7](#).

Example 6-7. ServiceHost<T> and returning unknown exceptions

```
public class ServiceHost<T> : ServiceHost
{
    public bool IncludeExceptionDetailInFaults
    {
        set
        {
            if(State == CommunicationState.Opened)
            {
                throw new InvalidOperationException("Host is already opened");
            }
            Description.Behaviors.Find<ServiceBehaviorAttribute>().DebuggingBehavior.
                IncludeExceptionDetailInFaults = value;
        }
        get
        {
            return Description.Behaviors.Find<ServiceBehaviorAttribute>().
                DebuggingBehavior.IncludeExceptionDetailInFaults;
        }
    }
    //More members
}
```

Using ServiceHost<T> is trivial and readable:

```
ServiceHost<MyService> host = new ServiceHost<MyService>();
host.IncludeExceptionDetailInFaults = true;
host.Open();
```

To apply this behavior administratively, add a custom behavior section in the host config file and reference it in the service definition, as shown in [Example 6-8](#).

Example 6-8. Administratively including exceptions in the fault message

```
<system.serviceModel>
    <services>
        <service name = "MyService" behaviorConfiguration = "Debugging">
            ...
        </service>
    </services>
    <behaviors>
        <serviceBehaviors>
            <behavior name = "Debugging">
                <serviceDebug includeExceptionDetailInFaults = "true"/>
            </behavior>
        </serviceBehaviors>
    </behaviors>
</system.serviceModel>
```

The advantage of administrative configuration in this case is the ability to toggle the behavior in production post-deployment without affecting the service code.

Exception extraction

While including the exception details in the fault is a useful diagnostic technique, it is a cumbersome programming model: the client has to take extra steps to extract the error information out of the `ExceptionDetail` object. More deterring is the fact that the client must use a single massive `catch` statement (that catches a single `FaultException<ExceptionDetail>`) to catch all possible exceptions, and sort them all inside the `catch` statement. In the world of .NET, this is akin to always catching a mere `Exception`, and avoiding cascading `catch` statements.

In addition, when writing a test client, you want to know as much as possible about the original exception that happened on the service side, since your test cases are predicated on producing specific errors. The test client could extract the original exception from the `ExceptionDetail` object and recursively build the inner exception chain. However, that would be tedious and redundant, and it would require repeated code on every use of the service by the test client. It is therefore better to encapsulate these steps in the proxy using C# extensions. To that end, I wrote the `ExtractException()` extension method to `FaultException<ExceptionDetail>`, defined as:

```
public static class DebugHelper
{
    public static Exception ExtractException(this FaultException<ExceptionDetail>
                                              fault);
    //More members
}
```

The implementation of `FaultException<ExceptionDetail>` has nothing to do with WCF, so I won't show it here (but it is available with `ServiceModelEx`). The best way of using the extension is to encapsulate it within the proxy, as shown in [Example 6-9](#).

Example 6-9. Automatically extracting the exception

```
[ServiceContract]
interface IMyContract
{
    [OperationContract]
    void MethodWithError();
}
class MyContractClient : ClientBase<IMyContract>, IMyContract
{
    public MyContractClient()
    {}
    /* More constructors */
    public void MethodWithError()
```

```

    {
        try
        {
            Channel.MethodWithError();
        }
        catch(FaultException<ExceptionDetail> exception)
        {
            Abort();
            throw exception.ExtractException();
        }
    }
}

```

In [Example 6-9](#), in the case of a `FaultException<ExceptionDetail>`, the proxy aborts itself (to prevent the proxy from being used again) regardless of whether a transport session is present or how exactly the service threw the exception. The proxy uses the extension method to throw the extracted exception, allowing the client to catch the raw CLR exception. For example, for this service definition in debug mode:

```

[ServiceContract]
interface IMyContract
{
    [OperationContract]
    void MyMethod();
}

[ServiceBehavior(IncludeExceptionDetailInFaults =
    DebugHelper.IncludeExceptionDetailInFaults)]
class MyService : IMyContract
{
    public void MyMethod()
    {
        throw new InvalidOperationException();
    }
}

```

when using the proxy from [Example 6-9](#), the client can expect to catch an `InvalidOperationException`:

```

MyContractClient proxy = new MyContractClient();
try
{
    proxy.MyMethod();
}
catch(InvalidOperationException exception)
{...}

```



Exception extraction should be used judiciously, only in specific diagnostic and testing cases, since it negates the core benefit of fault masking and decoupling from the nature of the error and the technology.

Faults and Callbacks

Callbacks to the client can, of course, fail due to communication exceptions, or because the callback itself threw an exception. Similar to service contract operations, callback contract operations can define fault contracts, as shown in [Example 6-10](#).

Example 6-10. Callback contract with fault contract

```
[ServiceContract(CallbackContract = typeof(IMyContractCallback))]
interface IMyContract
{
    [OperationContract]
    void DoSomething();
}

interface IMyContractCallback
{
    [OperationContract]
    [FaultContract(typeof(InvalidOperationException))]
    void OnCallBack();
}
```



Callbacks in WCF are usually configured as one-way calls, and as such cannot define their own fault contracts.

However, unlike with a normal service invocation, what is propagated to the service and how the error manifests itself also depend upon the following:

- When the callback is being invoked (i.e., whether the callback is invoked during a service call to its calling client or is invoked out-of-band by some other party on the host side)
- The type of the exception thrown

If the callback is invoked out-of-band—that is, by some party other than the service during a service operation—the callback behaves like a normal WCF operation invocation. [Example 6-11](#) demonstrates out-of-band invocation of the callback contract defined in [Example 6-10](#).

Example 6-11. Fault handling in out-of-band invocation

```
[ServiceBehavior(InstanceContextMode = InstanceContextMode.PerCall)]
class MyService : IMyContract
{
    static List<IMyContractCallback> m_Callbacks = new List<IMyContractCallback>();

    public void DoSomething()
    {
        IMyContractCallback callback =
            OperationContext.Current.GetCallbackChannel<IMyContractCallback>();

        if(m_Callbacks.Contains(callback) == false)
        {
            m_Callbacks.Add(callback);
        }
    }
    public static void CallClients()
    {
        Action<IMyContractCallback> invoke = (callback)=>
        {
            try
            {
                callback.OnCallBack();
            }
            catch(FaultException<...> exception)
            {...}
            catch(FaultException exception)
            {...}
            catch(CommunicationException exception)
            {...}
        };
        m_Callbacks.ForEach(invoke);
    }
}
```

As you can see, it is valid to expect to handle the callback fault contract, because faults are propagated to the host side according to it. If the client callback throws a `FaultException` or any of its subclasses, it will not fault the callback channel, and you can catch the exception and continue using the callback channel. However, as with service calls, after an exception that is not part of the fault contract occurs, you should avoid using the callback channel.

Likewise, when the service calls back to its calling client, if the callback throws a `FaultException` or any of its subclasses, it will not fault the callback channel, and the service can catch the exception and continue using the callback channel (just as with the out-of-band invocation):

```
[ServiceBehaviorConcurrencyMode = ConcurrencyMode.Reentrant]
class MyService : IMyContract
{
```

```

public void DoSomething()
{
    IMyContractCallback callback = OperationContext.Current.GetCallbackChannel
        <IMyContractCallback>();

    try
    {
        callback.OnCallBack();
    }
    catch(FaultException exception)
    {...}
}
}

```

Note that the service must be configured for reentrancy to avoid a deadlock, as explained in [Chapter 5](#).

The scenario gets considerably more complex when the service invokes the callback during a service operation, calling back to its calling client, and the exception is not `FaultException` and does not derive from `FaultException`. Recall that all bindings capable of duplex communication maintain a transport-level session. The exception during the callback terminates the transport session from the client to the service.

Since both the TCP and IPC bindings use the same transport for calls from the client to the service and callbacks from the service to the client, when the callback throws such an exception the client that called the service in the first place immediately receives a `CommunicationException`, even if the service catches the exception. This is a direct result of reusing the same transport for both directions, and having faulted the callback transport (which is tantamount to faulting the client-to-service transport as well). The service can catch and handle the exception, but the client still gets its exception:

```

[ServiceBehaviorConcurrencyMode = ConcurrencyMode.Reentrant]
class MyService : IMyContract
{
    public void DoSomething()
    {
        IMyContractCallback callback = OperationContext.Current.GetCallbackChannel
            <IMyContractCallback>();

        try
        {
            callback.OnCallBack();
        }
        catch(FaultException exception) //Client still gets CommunicationException
        {...}
    }
}

```

Callback debugging

While the callback can use the same technique shown in [Example 6-4](#) to manually include the exception in the fault message, the `CallbackBehavior` attribute provides the Boolean property `IncludeExceptionDetailInFaults`, which can be used to include all non-contract exceptions in the message:

```
[AttributeUsage(AttributeTargets.Class)]
public sealed class CallbackBehaviorAttribute : Attribute, ...
{
    public bool IncludeExceptionDetailInFaults
    {get;set;}
    //More members
}
```

As for the service, including the exceptions is instrumental in debugging:

```
[CallbackBehavior(IncludeExceptionDetailInFaults =
                  DebugHelper.IncludeExceptionDetailInFaults)]
class MyClient : IMyContractCallback
{
    public void OnCallBack()
    {
        ...
        throw new InvalidOperationException();
    }
}
```

You can also configure this behavior administratively in the client config file:

```
<client>
    <endpoint ... behaviorConfiguration = "Debugging"
    ...
    />
</client>
<behaviors>
    <endpointBehaviors>
        <behavior name = "Debugging">
            <callbackDebug includeExceptionDetailInFaults = "true"/>
        </behavior>
    </endpointBehaviors>
</behaviors>
```

Note the use of the `endpointBehaviors` tag to affect the client's callback endpoint.

Error-Handling Extensions

WCF enables developers to customize the default exception reporting and propagation behavior, and even to provide for a hook for custom logging. This extensibility is applied per channel dispatcher (that is, per endpoint), although you are most likely to simply utilize it across all dispatchers.

To install your own error-handling extension, you need to provide the dispatchers with an implementation of the `IErrorHandler` interface, defined as:

```
public interface IErrorHandler
{
    bool HandleError(Exception error);
    void ProvideFault(Exception error,MessageVersion version,ref Message fault);
}
```

Any party can provide this implementation, but typically it will be provided either by the service itself or by the host. In fact, you can have multiple error-handling extensions chained together. You will see how to install the extensions later in this section.

Providing a Fault

The `ProvideFault()` method of the extension object is called immediately after any unhandled exception is thrown by the service or any object on the call chain downstream from a service operation. WCF calls `ProvideFault()` before returning control to the client, and before terminating the session (if present) and disposing of the service instance (if required). Because `ProvideFault()` is called on the incoming call thread while the client is still blocked waiting for the operation to complete, you should avoid lengthy execution inside `ProvideFault()`.

Using `ProvideFault()`

`ProvideFault()` is called regardless of the type of exception thrown, be it a regular CLR exception, an unlisted fault, or a fault listed in the fault contract. The `error` parameter is a reference to the exception just thrown. If `ProvideFault()` does nothing, the exception the client gets will be determined by the fault contract (if any) and the exception type being thrown, as discussed previously in this chapter:

```
class MyErrorHandler : IErrorHandler
{
    public bool HandleError(Exception error)
    {...}

    public void ProvideFault(Exception error,MessageVersion version,
                           ref Message fault)
    {
        //Nothing here- exception will go up as usual
    }
}
```

However, `ProvideFault()` can examine the `error` parameter and either return it to the client as-is, or provide an alternative fault. This alternative behavior will affect even exceptions that are in the fault contracts. To provide an alternative fault, you need to use the `CreateMessageFault()` method of `FaultException` to create an alternative fault message. If you are providing a new fault contract message, you must cre-

ate a new detailing object, and you cannot reuse the original `error` reference. You then provide the created fault message to the static `CreateMessage()` method of the `Message` class:

```
public abstract class Message : ...
{
    public static Message CreateMessage(MessageVersion version,
                                         MessageFault fault, string action);
    //More members
}
```

Note that you need to provide `CreateMessage()` with the action of the fault message used. This intricate sequence is demonstrated in [Example 6-12](#).

Example 6-12. Creating an alternative fault

```
class MyErrorHandler : IErrorHandler
{
    public bool HandleError(Exception error)
    {...}
    public void ProvideFault(Exception error, MessageVersion version,
                           ref Message fault)
    {
        FaultException<int> faultException = new FaultException<int>(3);
        MessageFault messageFault = faultException.CreateMessageFault();
        fault = Message.CreateMessage(version, messageFault, faultException.Action);
    }
}
```

In [Example 6-12](#), the `ProvideFault()` method provides `FaultException<int>` with a value of 3 as the fault thrown by the service, irrespective of the actual exception that was thrown.

The implementation of `ProvideFault()` can also set the `fault` parameter to `null`:

```
class MyErrorHandler : IErrorHandler
{
    public bool HandleError(Exception error)
    {...}
    public void ProvideFault(Exception error, MessageVersion version,
                           ref Message fault)
    {
        fault = null; //Suppress any faults in contract
    }
}
```

Doing so will result in all exceptions being propagated to the client as `FaultExceptions`, even if the exceptions were listed in the fault contracts. Setting `fault` to `null` is therefore an effective way of suppressing any fault contracts that may be in place.

Exception promotion

One possible use for `ProvideFault()` is a technique I call *exception promotion*. A service may use downstream objects, which could be called by a variety of services. In the interest of decoupling, these objects may very well be unaware of the particular fault contracts of the service calling them. In case of errors, the objects simply throw regular CLR exceptions. If a downstream object throws an exception of type `T`, where `FaultException<T>` is part of the operation fault contract, by default the service will report that exception to the client as an opaque `FaultException`. What the service could do instead is use an error-handling extension to examine the exception thrown. If that exception is of the type `T`, where `FaultException<T>` is part of the operation fault contract, the service could then promote that exception to a full-fledged `FaultException<T>`. For example, given this service contract:

```
[ServiceContract]
interface IMyContract
{
    [OperationContract]
    [FaultContract(typeof(InvalidOperationException))]
    void MyMethod();
}
```

if the downstream object throws an `InvalidOperationException`, `ProvideFault()` will promote it to `FaultException<InvalidOperationException>`, as shown in [Example 6-13](#).

Example 6-13. Exception promotion

```
class MyErrorHandler : IErrorHandler
{
    public bool HandleError(Exception error)
    {...}
    public void ProvideFault(Exception error, MessageVersion version,
                           ref Message fault)
    {
        if(error is InvalidOperationException)
        {
            FaultException<InvalidOperationException> faultException =
                new FaultException<InvalidOperationException>(
                    new InvalidOperationException(error.Message));
            MessageFault messageFault = faultException.CreateMessageFault();
            fault = Message.CreateMessage(version,messageFault,
                                          faultException.Action);
        }
    }
}
```

The problem with such explicit exception promotion is that the code is coupled to a specific fault contract, and implementing it across all services requires a lot of tedious

work—not to mention that any change to the fault contract will necessitate a change to the error extension.

Fortunately, you can automate exception promotion using my `ErrorHandlerHelper` static class:

```
public static class ErrorHandlerHelper
{
    public static void PromoteException(Type serviceType,
                                         Exception error,
                                         MessageVersion version,
                                         ref Message fault);
    //More members
}
```

The `ErrorHandlerHelper.PromoteException()` method requires the service type as a parameter. It uses reflection to examine all the interfaces and operations on that service type, looking for fault contracts for the particular operation (it gets the faulted operation by parsing the `error` object). `PromoteException()` lets exceptions in the contract go up the call stack unaffected, but it will promote a CLR exception to a contracted fault if the exception type matches any one of the detailing types defined in the fault contracts for that operation.

Using `ErrorHandlerHelper`, [Example 6-13](#) can be reduced to one or two lines of code:

```
class MyErrorHandler : IErrorHandler
{
    public bool HandleError(Exception error)
    {...}
    public void ProvideFault(Exception error, MessageVersion version,
                             ref Message fault)
    {
        Type serviceType = ...;
        ErrorHandlerHelper.PromoteException(serviceType,error,version,ref fault);
    }
}
```

The implementation of `PromoteException()` has little to do with WCF, so it is not listed in this chapter. However, you can examine it as part of the source code available with *ServiceModelEx*. The implementation makes use of some advanced C# programming techniques, such as generics late binding.

Handling a Fault

The `HandleError()` method of `IErrorHandler` is defined as:

```
bool HandleError(Exception error);
```

`HandleError()` is called by WCF after control returns to the client. `HandleError()` is strictly for service-side use, and nothing it does affects the client in any way. Calling in the background enables you to perform lengthy processing, such as logging to a database without impeding the client.

Because you can have multiple error-handling extensions installed in a list, WCF also enables you to control whether extensions down the list should be used. If `HandleError()` returns `false`, WCF will continue to call `HandleError()` on the rest of the installed extensions. If `HandleError()` returns `true`, WCF stops invoking the error-handling extensions. Obviously, most extensions should return `false`.

The `error` parameter of `HandleError()` is the original exception thrown. The classic use for `HandleError()` is for logging and tracing, as shown in [Example 6-14](#).

Example 6-14. Logging the error log to a logbook service

```
class MyErrorHandler : IErrorHandler
{
    public bool HandleError(Exception error)
    {
        try
        {
            LogbookServiceClient proxy = new LogbookServiceClient();
            proxy.Log(...);
            proxy.Close();
        }
        catch
        {}
        return false;
    }
    public void ProvideFault(Exception error, MessageVersion version,
                           ref Message fault)
    {...}
}
```

The logbook service

The source code available with this book in *ServiceModelEx* contains a standalone service called `LogbookManager` that is dedicated to error logging. `LogbookManager` logs the errors into a SQL Server database. The service contract also provides operations for retrieving the entries in the logbook and clearing the logbook. *ServiceModelEx* also contains a simple logbook viewer and management tool. In addition to error logging, `LogbookManager` allows you to log entries explicitly into the logbook, independently of exceptions. The architecture of this framework is depicted in [Figure 6-1](#).

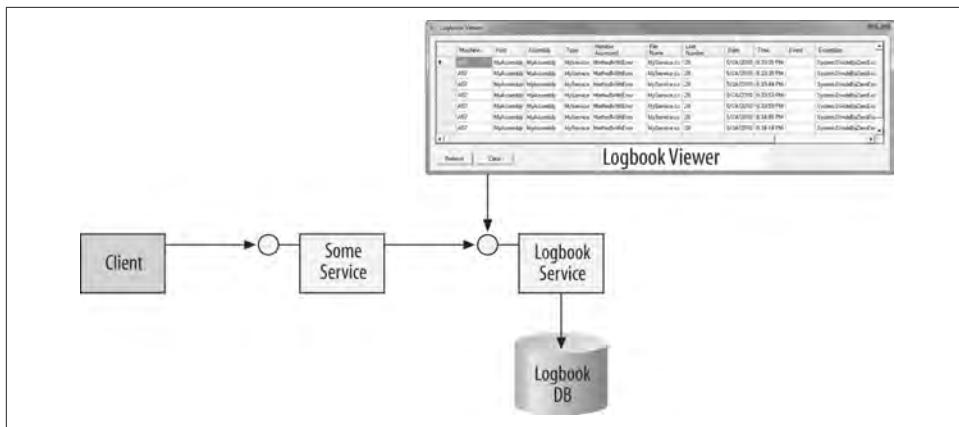


Figure 6-1. The logbook service and viewer

You can automate error logging to `LogbookManager` using the `.LogError()` method of my `ErrorHandlerHelper` static class:

```
public static class ErrorHandlerHelper
{
    public static void LogError(Exception error);
    //More members
}
```

The `error` parameter is simply the exception you wish to log. `.LogError()` encapsulates the call to `LogbookManager`. For example, instead of the code in [Example 6-14](#), you can simply write a single line:

```
class MyErrorHandler : IErrorHandler
{
    public bool HandleError(Exception error)
    {
        ErrorHandlerHelper.LogError(error);
        return false;
    }
    public void ProvideFault(Exception error, MessageVersion version,
                           ref Message fault)
    {...}
}
```

In addition to capturing the raw exception information, `.LogError()` performs extensive parsing of the exception and other environment variables for a comprehensive record of the error and its related information.

Specifically, `.LogError()` captures the following information:

- Where the exception occurred (machine name and host process name)

- The code where the exception took place (the assembly name, the filename, and the line number if debug symbols are provided)
- The type where the exception took place and the member being accessed
- The date and time when the exception occurred
- The exception name and message

Implementing `LogError()` has little to do with WCF, so this method is not shown in this chapter. The code, however, makes extensive use of interesting .NET programming techniques such as string and exception parsing, along with obtaining the environment information. The error information is passed to `LogbookManager` in a dedicated data contract.

Installing Error-Handling Extensions

Every channel dispatcher in WCF offers a collection of error extensions:

```
public class ChannelDispatcher : ChannelDispatcherBase
{
    public Collection<IErrorHandler> ErrorHandlers
    {get;}
    //More members
}
```

Installing your own custom implementation of `IErrorHandler` requires merely adding it to the desired dispatcher (usually all of them).

You must add the error extensions before the first call arrives to the service, but after the host constructs the collection of dispatchers. This narrow window of opportunity exists after the host is initialized, but before it is opened. To act in that window, the best solution is to treat error extensions as custom service behaviors, because the behaviors are given the opportunity to interact with the dispatchers at just the right time. As mentioned in [Chapter 4](#), all service behaviors implement the `IServiceBehavior` interface, defined as:

```
public interface IServiceBehavior
{
    void AddBindingParameters(ServiceDescription description,
                             ServiceHostBase host,
                             Collection<ServiceEndpoint> endpoints,
                             BindingParameterCollection bindingParameters);

    void ApplyDispatchBehavior(ServiceDescription description,
                             ServiceHostBase host);

    void Validate(ServiceDescription description, ServiceHostBase host);
}
```

The `ApplyDispatchBehavior()` method is your cue to add the error extensions to the dispatchers. You can safely ignore all other methods of `IServiceBehavior` and provide empty implementations for them.

In `ApplyDispatchBehavior()`, you need to access the collection of dispatchers available in the `ChannelDispatchers` property of `ServiceHostBase`:

```
public class ChannelDispatcherCollection : SynchronizedCollection<ChannelDispatcherBase>
{
    public abstract class ServiceHostBase : ...
    {
        public ChannelDispatcherCollection ChannelDispatchers
        {get;}
        //More members
    }
}
```

Each item in `ChannelDispatchers` is of the type `ChannelDispatcher`. You can add the implementation of `IErrorHandler` to all dispatchers, or just add it to specific dispatchers associated with a particular binding. [Example 6-15](#) demonstrates adding an implementation of `IErrorHandler` to all of a service's dispatchers.

Example 6-15. Adding an error extension object

```
class MyErrorHandler : IErrorHandler
{...}

class MyService : IMyContract,IServiceBehavior
{
    public void ApplyDispatchBehavior(ServiceDescription description,
                                      ServiceHostBase host)
    {
        IErrorHandler handler = new MyErrorHandler();
        foreach(ChannelDispatcher dispatcher in host.ChannelDispatchers)
        {
            dispatcher.ErrorHandlers.Add(handler);
        }
    }
    public void Validate(...)
    {}
    public void AddBindingParameters(...)
    {}
    //More members
}
```

In [Example 6-15](#), the service itself implements `IServiceBehavior`. In `ApplyDispatchBehavior()`, the service obtains the dispatchers collection and adds an instance of the `MyErrorHandler` class to each dispatcher.

Instead of relying on an external class to implement `IErrorHandler`, the service class itself can support `IErrorHandler` directly, as shown in [Example 6-16](#).

Example 6-16. Service class supporting `IErrorHandler`

```
class MyService : IMyContract,IServiceBehavior,IErrorHandler
{
    public void ApplyDispatchBehavior(ServiceDescription description,
                                      ServiceHostBase host)
    {
        foreach(ChannelDispatcher dispatcher in host.ChannelDispatchers)
        {
            dispatcher.ErrorHandlers.Add(this);
        }
    }
    public bool HandleError(Exception error)
    {...}

    public void ProvideFault(Exception error,MessageVersion version,
                           ref Message fault)
    {...}
    //More members
}
```

The `ErrorHandlerBehavior`

The problem with Examples [6-15](#) and [6-16](#) is that they pollute the service class code with WCF plumbing; instead of focusing exclusively on the business logic, the service also has to wire up error extensions. Fortunately, you can provide the same plumbing declaratively using my `ErrorHandlerBehaviorAttribute`, defined as:

```
public class ErrorHandlerBehaviorAttribute : Attribute,IErrorHandler,
                                            IServiceBehavior
{
    protected Type ServiceType
    {get;set;}
}
```

Applying the `ErrorHandlerBehavior` attribute is straightforward:

```
[ErrorHandlerBehavior]
class MyService : IMyContract
{...}
```

The attribute installs itself as an error-handling extension. Its implementation uses `ErrorHandlerHelper` both to automatically promote exceptions to fault contracts, if required, and to automatically log the exceptions to `LogbookManager`. [Example 6-17](#) lists the implementation of the `ErrorHandlerBehavior` attribute.

Example 6-17. The ErrorHandlerBehavior attribute

```
[AttributeUsage(AttributeTargets.Class)]
public class ErrorHandlerBehaviorAttribute : Attribute,IServiceBehavior,
                                         IErrorHandler
{
    protected Type ServiceType
    {get;set;}

    void IServiceBehavior.ApplyDispatchBehavior(ServiceDescription description,
                                                ServiceHostBase host)
    {
        ServiceType = description.ServiceType;
        foreach(ChannelDispatcher dispatcher in host.ChannelDispatchers)
        {
            dispatcher.ErrorHandlers.Add(this);
        }
    }
    bool IErrorHandler.HandleError(Exception error)
    {
        ErrorHandlerHelper.LogError(error);
        return false;
    }
    void IErrorHandler.ProvideFault(Exception error,MessageVersion version,
                                    ref Message fault)
    {
        ErrorHandlerHelper.PromoteException(ServiceType,error,version,ref fault);
    }
    void IServiceBehavior.Validate(...)
    {}
    void IServiceBehavior.AddBindingParameters(...)
    {}
}
```

Note in [Example 6-17](#) that `ApplyDispatchBehavior()` saves the service type in a protected property. The reason is that the call to `ErrorHandlerHelper.PromoteException()` in `ProvideFault()` requires the service type.

The Host and Error Extensions

While the `ErrorHandlerBehavior` attribute greatly simplifies the act of installing an error extension, the attribute does require the service developer to apply the attribute. It would be nice if the host could add error extensions independently of whether or not the service provides any. However, due to the narrow timing window available for installing extensions, having the host add such an extension requires multiple steps. First, you need to provide an error-handling extension type that supports both `IServiceBehavior` and `IErrorHandler`. The implementation of `IServiceBehavior` will add the error extension to the dispatchers, as shown previously. Next, you must

derive a custom host class from `ServiceHost` and override the `OnOpening()` method defined by the `CommunicationObject` base class:

```
public abstract class CommunicationObject : ICommunicationObject
{
    protected virtual void OnOpening();
    //More members
}
public abstract class ServiceHostBase : CommunicationObject ,...
{...}
public class ServiceHost : ServiceHostBase, ...
{...}
```

In `OnOpening()`, you need to add the custom error-handling type to the collection of service behaviors in the service description. That behaviors collection was described in Chapters 1 and 4:

```
public class Collection<T> : IList<T>, ...
{
    public void Add(T item);
    //More members
}
public abstract class KeyedCollection<K,T> : Collection<T>
{...}
public class KeyedByTypeCollection<I> : KeyedCollection<Type,I>
{...}
public class ServiceDescription
{
    public KeyedByTypeCollection<IServiceBehavior> Behaviors
    {get;}
}
public abstract class ServiceHostBase : ...
{
    public ServiceDescription Description
    {get;}
    //More members
}
```

This sequence of steps is already encapsulated and automated in `ServiceHost<T>`:

```
public class ServiceHost<T> : ServiceHost
{
    public void AddErrorHandler(IErrorHandler errorHandler);
    public void AddErrorHandler();
    //More members
}
```

`ServiceHost<T>` offers two overloaded versions of the `AddErrorHandler()` method. The one that takes an `IErrorHandler` object will internally associate it with a behavior, so you can provide it with any class that supports just `IErrorHandler`, not `IServiceBehavior`:

```

class MyService : IMyContract
{...}

class MyErrorHandler : IErrorHandler
{...}

ServiceHost<MyService> host = new ServiceHost<MyService>();
host.AddErrorHandler(new MyErrorHandler());
host.Open();

```

The `AddErrorHandler()` method that takes no parameters will install an error-handling extension that uses `ErrorHandlerHelper`, just as if the service class was decorated with the `ErrorHandlerBehavior` attribute:

```

class MyService : IMyContract
{...}

ServiceHost<MyService> host = new ServiceHost<MyService>();
host.AddErrorHandler();
host.Open();

```

Actually, for this last example, `ServiceHost<T>` does internally use an instance of the `ErrorHandlerBehavior` attribute.

Example 6-18 shows the implementation of the `AddErrorHandler()` method.

Example 6-18. Implementing AddErrorHandler()

```

public class ServiceHost<T> : ServiceHost
{
    class ErrorHandlerBehavior : IServiceBehavior, IErrorHandler
    {
        IErrorHandler m_ErrorHandler;

        public ErrorHandlerBehavior(IErrorHandler errorHandler)
        {
            m_ErrorHandler = errorHandler;
        }

        void IServiceBehavior.ApplyDispatchBehavior(ServiceDescription description,
                                                    ServiceHostBase host)
        {
            foreach(ChannelDispatcher dispatcher in host.ChannelDispatchers)
            {
                dispatcher.ErrorHandlers.Add(this);
            }
        }

        bool IErrorHandler.HandleError(Exception error)
        {
            return m_ErrorHandler.HandleError(error);
        }

        void IErrorHandler.ProvideFault(Exception error, MessageVersion version,
                                         ref Message fault)
    }
}

```

```

    {
        m_ErrorHandler.ProvideFault(error,version,ref fault);
    }
    //Rest of the implementation
}

List<IServiceBehavior> m_ErrorHandlers = new List<IServiceBehavior>();

public void AddErrorHandler(IErrorHandler errorHandler)
{
    if(State == CommunicationState.Opened)
    {
        throw new InvalidOperationException("Host is already opened");
    }
    IServiceBehavior errorHandlerBehavior =
                    new ErrorHandlerBehavior(errorHandler);
    m_ErrorHandlers.Add(errorHandlerBehavior);
}
public void AddErrorHandler()
{
    AddErrorHandler(new ErrorHandlerBehaviorAttribute());
}
protected override void OnOpening()
{
    foreach(IServiceBehavior behavior in m_ErrorHandlers)
    {
        Description.Behaviors.Add(behavior);
    }
    base.OnOpening();
}
//Rest of the implementation
}

```

To avoid forcing the provided `IErrorHandler` reference to also support `IServiceBehavior`, `ServiceHost<T>` defines a private nested class called `ErrorHandlerBehavior`. `ErrorHandlerBehavior` implements both `IErrorHandler` and `IServiceBehavior`. To construct `ErrorHandlerBehavior`, you need to provide it with an implementation of `IErrorHandler`. That implementation is saved for later use. The implementation of `IServiceBehavior` adds the instance itself to the error-handler collection of all dispatchers. The implementation of `IErrorHandler` simply delegates to the saved construction parameter. `ServiceHost<T>` defines a list of `IServiceBehavior` references in the `m_ErrorHandlers` member variable. The `AddErrorHandler()` method that accepts an `IErrorHandler` reference uses it to construct an instance of `ErrorHandlerBehavior` and then adds it to `m_ErrorHandlers`. The `AddErrorHandler()` method that takes no parameter uses an instance of the `ErrorHandlerBehavior` attribute, because the attribute is merely a class that supports `IErrorHandler`. Finally, the `OnOpening()` method iterates over `m_ErrorHandlers`, adding each behavior to the behaviors collection.

Callbacks and Error Extensions

The client-side callback object can also provide an implementation of `IErrorHandler` for error handling. Compared with the service-error extensions, the main difference is that to install the callback extension you need to use the `IEndpointBehavior` interface, defined as:

```
public interface IEndpointBehavior
{
    void AddBindingParameters(ServiceEndpoint endpoint,
                             BindingParameterCollection bindingParameters);
    void ApplyClientBehavior(ServiceEndpoint endpoint,
                            ClientRuntime clientRuntime);
    void ApplyDispatchBehavior(ServiceEndpoint endpoint,
                             EndpointDispatcher endpointDispatcher);
    void Validate(ServiceEndpoint endpoint);
}
```

`IEndpointBehavior` is the interface all callback behaviors support. The only relevant method for the purpose of installing an error extension is the `ApplyClientBehavior()` method, which lets you associate the error extension with the single dispatcher of the callback endpoint. The `clientRuntime` parameter is of the type `ClientRuntime`, which offers the `CallbackDispatchRuntime` property of the type `DispatchRuntime`. The `DispatchRuntime` class offers the `ChannelDispatcher` property, with its collection of error handlers:

```
public sealed class ClientRuntime
{
    public DispatchRuntime CallbackDispatchRuntime
    {get;}
    //More members
}
public sealed class DispatchRuntime
{
    public ChannelDispatcher ChannelDispatcher
    {get;}
    //More members
}
```

As with a service-side error-handling extension, you need to add to that collection your custom error-handling implementation of `IErrorHandler`.

The callback object itself can implement `IEndpointBehavior`, as shown in [Example 6-19](#).

Example 6-19. Implementing IEndpointBehavior

```
class MyErrorHandler : IErrorHandler
{...}
```

```

class MyClient : IMyContractCallback, IEndpointBehavior
{
    public void OnCallBack()
    {...}

    void IEndpointBehavior.ApplyClientBehavior(ServiceEndpoint serviceEndpoint,
                                                ClientRuntime clientRuntime)
    {
        IErrorHandler handler = new MyErrorHandler();

        clientRuntime.CallbackDispatchRuntime.ChannelDispatcher.ErrorHandlers.
            Add(handler);
    }

    void IEndpointBehavior.AddBindingParameters(...)
    {}
    void IEndpointBehavior.ApplyDispatchBehavior(...)
    {}
    void IEndpointBehavior.Validate(...)
    {}
    //More members
}

```

Instead of using an external class for implementing `IErrorHandler`, the callback class itself can implement `IErrorHandler` directly:

```

class MyClient : IMyContractCallback, IEndpointBehavior, IErrorHandler
{
    public void OnCallBack()
    {...}

    void IEndpointBehavior.ApplyClientBehavior(ServiceEndpoint serviceEndpoint,
                                                ClientRuntime clientRuntime)
    {
        clientRuntime.CallbackDispatchRuntime.ChannelDispatcher.ErrorHandlers.
            Add(this);
    }
    public bool HandleError(Exception error)
    {...}
    public void ProvideFault(Exception error, MessageVersion version,
                            ref Message fault)
    {...}
    //More members
}

```

The `CallbackErrorHandlerBehavior` attribute

Code such as that shown in [Example 6-19](#) can be automated with the `CallbackErrorHandlerBehaviorAttribute`, defined as:

```

public class CallbackErrorHandlerBehaviorAttribute :
    ErrorHandlerBehaviorAttribute, IEndpointBehavior

```

```

{
    public CallbackErrorHandlerBehaviorAttribute(Type clientType);
}

```

The `CallbackErrorHandlerBehavior` attribute derives from the service-side `ErrorHandlerBehavior` attribute and adds explicit implementation of `IEndpointBehavior`. The attribute uses `ErrorHandlerHelper` to promote and log the exception.

In addition, the attribute requires as a construction parameter the type of the callback on which it is applied:

```

[CallbackErrorHandlerBehavior(typeof(MyClient))]
class MyClient : IMyContractCallback
{
    public void OnCallBack()
    {...}
}

```

The type is required because there is no other way to get hold of the callback type, which is required by `ErrorHandlerHelper.PromoteException()`.

The implementation of the `CallbackErrorHandlerBehavior` attribute is shown in [Example 6-20](#).

Example 6-20. Implementing the `CallbackErrorHandlerBehavior` attribute

```

public class CallbackErrorHandlerBehaviorAttribute :
    ErrorHandlerBehaviorAttribute, IEndpointBehavior
{
    public CallbackErrorHandlerBehaviorAttribute(Type clientType)
    {
        ServiceType = clientType;
    }
    void IEndpointBehavior.ApplyClientBehavior(ServiceEndpoint serviceEndpoint,
                                                ClientRuntime clientRuntime)
    {
        clientRuntime.CallbackDispatchRuntime.ChannelDispatcher.ErrorHandlers.
            Add( this );
    }
    void IEndpointBehavior.AddBindingParameters(...)
    {}
    void IEndpointBehavior.ApplyDispatchBehavior(...)
    {}
    void IEndpointBehavior.Validate(...)
    {}
}

```

Note in [Example 6-20](#) how the provided callback client type is stored in the `ServiceType` property, defined as protected in [Example 6-17](#).

CHAPTER 7

Transactions

Transactions are the key to building robust, high-quality applications. WCF provides simple, declarative transaction support for service developers, enabling you to configure parameters such as enlistment and voting, all outside the scope of your service. In addition, WCF allows client applications to create transactions and to propagate transactions across service boundaries. This chapter starts by introducing the problem space transactions address and the basic transactions terminology, and then discusses the support for transactions and transaction management offered by WCF and .NET. The rest of the chapter is dedicated to transactional programming models, both for services and clients, and to how transactions relate to other aspects of WCF, such as instance management and callbacks.

The Recovery Challenge

Proper error handling and recovery is the Achilles' heel of many applications. When an application fails to perform a particular operation, you should recover from it and restore the system—that is, the collection of interacting services and clients—to a consistent state (usually, the state the system was at before the operation that caused the error took place). Operations that can fail typically consist of multiple potentially concurrent smaller steps. Some of those steps can fail while others succeed. The problem with recovery is the sheer number of partial success and partial failure permutations that you have to code against. For example, an operation comprising 10 smaller concurrent steps has some three million recovery scenarios, because for the recovery logic, the order in which the suboperations fail matters as well, and the factorial of 10 is roughly three million.

Trying to handcraft recovery code in a decent-sized application is often a futile attempt, resulting in fragile code that is very susceptible to any changes in the application execution or the business use case, incurring both productivity and perfor-

mance penalties. The productivity penalty results from all the effort required for handcrafting the recovery logic. The performance penalty is inherited with such an approach because you need to execute huge amounts of code after every operation to verify that all is well. In reality, developers tend to deal only with the easy recovery cases (that is, the cases that they are both aware of and know how to handle). More insidious error scenarios, such as intermediate network failures or disk crashes, go unaddressed. In addition, because recovery is all about restoring the system to a consistent state (typically the state before the operations), the real problem has to do with the steps that succeeded, rather than those that failed. The failed steps failed to affect the system; the challenge is actually the need to undo successful steps, such as deleting a row from a table, or a node from a linked list, or a call to a remote service. The scenarios involved could be very complex, and your manual recovery logic is almost certain to miss a few successful suboperations.

The more complex the recovery logic becomes, the more error-prone the recovery itself becomes. If you have an error in the recovery, how would you recover the recovery? How do developers go about designing, testing, and debugging complex recovery logic? How do they simulate the endless number of errors and failures that are possible? Not only that, but what if before the operation failed, as it was progressing along executing its suboperations successfully, some other party accessed your application and acted upon the state of the system—the state that you are going to roll back during the recovery? That other party is now acting on inconsistent information, and by definition, is in error too. Moreover, your operation may be just one step in some other, much wider operation that spans multiple services from multiple vendors on multiple machines. How would you recover the system as a whole in such a case? Even if you have a miraculous way of recovering your service, how would that recovery logic plug into the cross-service recovery? As you can see, it is practically impossible to write error-recovery code by hand.

Transactions

The best (and perhaps only) way of maintaining system consistency and dealing properly with the error-recovery challenge is to use *transactions*. A transaction is a set of potentially complex operations, in which the failure of any single operation causes the entire set to fail, as one atomic operation. As illustrated in [Figure 7-1](#), while the transaction is in progress the system is allowed to be in a temporary inconsistent state, but once the transaction is complete it is guaranteed to be in a consistent state. That state may be either a new consistent state (B), or the original consistent state the system was in before the transaction started (A).

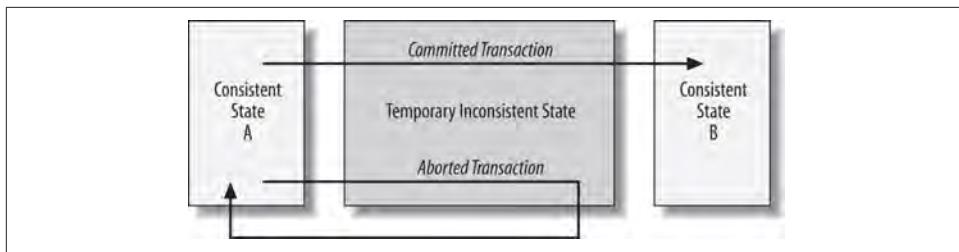


Figure 7-1. A transaction transfers the system between consistent states

A transaction that executes successfully and manages to transfer the system from the consistent state A to the consistent state B is called a *committed transaction*. If the transaction encounters an error during its execution and rolls back all the intermediate steps that have already succeeded, it is called an *aborted transaction*. If the transaction failed to either commit or abort, it is called an *in-doubt transaction*. In-doubt transactions usually require administrator or user assistance to resolve and are beyond the scope of this book.

Transactional Resources

Transactional programming requires working with a resource (such as a database or a message queue) that is capable of participating in a transaction, and being able to commit or roll back the changes made during the transaction. Such resources have been around in one form or another for decades. Traditionally, you had to inform a resource that you would like to perform transactional work against it. This act is called *enlisting* the resource in the transaction. Some resources support *autoenlistment*; that is, they can detect that they are being accessed by a transaction and automatically enlist in it. Once the resource is enlisted, you can then perform work against the resource. If no error occurs, the resource is asked to commit the changes made to its state; if any error is encountered, the resource is asked to roll back the changes. During a transaction, it is vital that you do not access any nontransactional resources (such as the file system on Windows XP), because changes made to those resources will not roll back if the transaction is aborted.

Transaction Properties

When you make use of transactions in your service-oriented applications, you must abide by four core properties, known as *ACID* (atomic, consistent, isolated, and durable). When you design transactional services, you must adhere to the ACID requirements—they are not optional. As you will see throughout this chapter, WCF enforces them rigorously.

The atomic property

In order for it to be *atomic*,¹ when a transaction completes, all the individual changes it has made to the resource state must be made as if they were all one atomic, indivisible operation. The changes made to the resource are made as if everything else in the universe stops, the changes are made, and then everything resumes. It must not be possible for a party outside the transaction to observe the resources involved with only some of the changes made, but not all of them. A transaction should not leave anything to be done in the background once it is done, as those operations would violate atomicity. Every operation resulting from the transaction must be included in the transaction itself.

Atomic transactions make client applications a lot easier to develop. The client does not have to manage partial failures of its requests, or have complex recovery logic. The client knows that the transaction will either succeed or fail as a whole. In the case of failure, the client can choose to issue a new request (start a new transaction), or do something else, such as alerting the user. The important thing is that the client does not have to recover the system.

The consistent property

Consistent means the transaction must leave the system in a consistent state. Note that consistency is different from atomicity. Even if all the changes are committed as one atomic operation, the transaction is required to guarantee that all those changes are consistent (i.e., that they “make sense”). Usually, it is up to the developer to ensure that the semantics of the operations are consistent. All the transaction is required to do is to transfer the system from one consistent state to another.

The isolated property

Isolated means no other entity (transactional or not) is able to see the intermediate state of the resources during the transaction, because that state may be inconsistent. In fact, even if it is consistent, the transaction could still abort, and the changes could be rolled back. Isolation is crucial to overall system consistency. Suppose transaction A allows transaction B access to its intermediate state, and then transaction A aborts while transaction B decides to commit. The problem is that transaction B based its execution on a system state that was rolled back, and therefore transaction B is left unknowingly inconsistent.

¹ The word “atom” comes from the Greek word “atomos,” meaning *indivisible*. The ancient Greeks thought that if you started dividing matter, and continued dividing it, eventually you would get to indivisible pieces, which they called “atomos.” The ancient Greeks were, of course, wrong, as atoms can be divided into subatomic particles such as electrons, protons, and neutrons. Transactions, however, are truly atomic.

Managing isolation is not trivial. The resources participating in a transaction must lock the data accessed by the transaction from all other parties, and must unlock access to that data when the transaction commits or aborts.

The durable property

Traditionally, transactional support by a resource implies not just a transaction-aware resource, but also a *durable* one. This is because at any moment the application could crash, and the memory it was using could be erased. If the changes to the system's state were in-memory changes, they would be lost, and the system would be left in an inconsistent state. However, durability is really a range of options. How resilient to such catastrophes the resource should be is an open question that depends on the nature and sensitivity of the data, your budget, available time and available system administration staff, and so on.

If durability is a range that actually means various degrees of persistence, then you can also consider the far end of the spectrum: volatile, in-memory resources. The advantage of volatile resources is that they offer better performance than durable resources, and more importantly, they allow you to approximate much better conventional programming models while using transaction support for error recovery. You will see later in this chapter how and when your services can benefit from volatile resource managers (VRMs).

Transaction Management

WCF services can work directly against a transactional resource and manage the transaction explicitly using a programming model such as that offered by ADO.NET. As shown in [Example 7-1](#), using this model, you are responsible for explicitly starting and managing the transaction.

Example 7-1. Explicit transaction management

```
[ServiceContract]
interface IMyContract
{
    [OperationContract]
    void MyMethod();
}

class MyService : IMyContract
{
    public void MyMethod()
    {
        //Avoid this programming model:

        string connectionString = "...";

        using(IDbConnection connection = new SqlConnection(connectionString))
```

```

{
    connection.Open();

    using(IDbCommand command = new SqlCommand())
    {
        command.Connection = connection;

        using(IDbTransaction transaction =
                connection.BeginTransaction())//Enlisting
        {
            command.Transaction = transaction;

            try
            {
                /* Interact with database here, then commit the transaction */
                transaction.Commit();
            }
            catch
            {
                transaction.Rollback(); //Abort transaction

                throw;
            }
        }
    }
}

```

You obtain an object representing the underlying database transaction by calling `BeginTransaction()` on the connection object. `BeginTransaction()` returns an implementation of the interface `IDbTransaction`, used to manage the transaction. When the database is enlisted, it does not really execute any of the requests that are made. Instead, it merely logs the requests against the transaction. If at the end all updates and other changes made to the database are consistent and no error has taken place, you simply call `Commit()` on the transaction object. This instructs the database to commit the changes as one atomic operation. If any exception occurred, the call to `Commit()` is skipped over, and the `catch` block aborts the transaction by calling `Rollback()`. Aborting the transaction instructs the database to discard all the changes logged so far.

The transaction management challenge

While the explicit programming model is straightforward, requiring nothing of the service performing the transaction, it is most suitable for a client calling a single service interacting with a single database (or a single transactional resource), where the service starts and manages the transaction, as shown in [Figure 7-2](#).

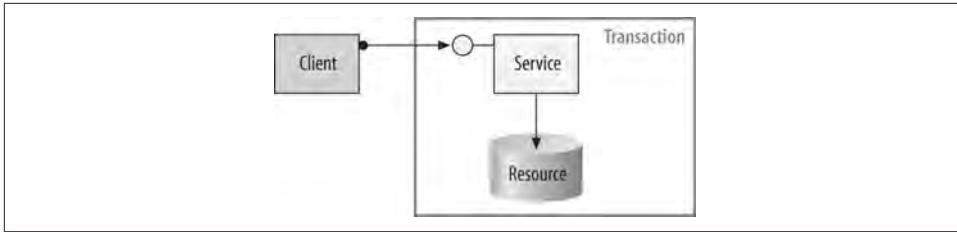


Figure 7-2. Single service/single resource transaction

This is due to the transaction coordination problem. Consider, for example, a service-oriented application where the client interacts with multiple services that in turn interact with each other and with multiple resources, as shown in [Figure 7-3](#).

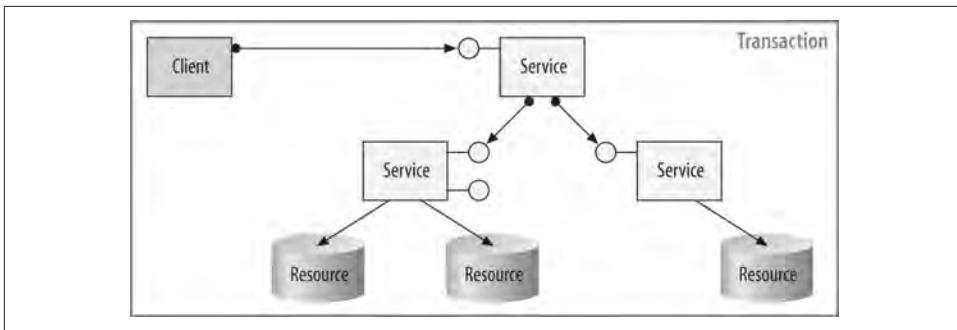


Figure 7-3. Distributed transactional service-oriented application

The question now is, which one of the participating services is responsible for beginning the transaction and enlisting each resource? If all of them will do that, you will end up with multiple transactions. Putting the enlistment logic in the service code will create a great deal of coupling between the services and the resources. Furthermore, which one of the services is responsible for committing or rolling back the transaction? How would one service know what the rest of the services feel about the transaction? How would the service managing the transaction inform the other services about the transaction's ultimate outcome? Trying to pass the transaction object or some identifier as an operation parameter is not service-oriented, because the clients and the services could all be using different implementation platforms and technologies. The services could also be deployed in different processes or even across different machines or sites. In this case, issues such as network failures or machine crashes introduce additional complexity for managing the transaction, because one service might crash while others continue processing the transaction.

One possible solution is to couple the clients and the services by adding logic for coordinating the transaction, but such an approach is very fragile and would not withstand even minor changes to the business flow or the number of participating

services. In addition, if different vendors developed the services, this will preclude any such coordination. And even if you find a way of solving the coordination problem at the service level, when multiple resources are involved you have multiple independent points of failure, because each of the resources could fail independently of the services.

Distributed transactions

The type of transaction just described is called a *distributed transaction*. A distributed transaction contains two or more independent services (often in different execution contexts), or even just a single service with two or more transactional resources. It is impractical to try to explicitly manage the potential error cases of a distributed transaction. For a distributed transaction, you need to rely on the two-phase commit protocol, and a dedicated transaction manager. A *transaction manager* is a third party that will manage the transaction for you, because the last thing you want is to place the transaction management logic in your service code.

The two-phase commit protocol

To overcome the complexity of a distributed transaction, the transaction manager uses a transaction management protocol called the *two-phase commit* protocol to decide on the outcome of the transaction as well as to commit or roll back the changes to the system state. The two-phase commit protocol is what enforces atomicity and consistency in a distributed system. The protocol enables WCF to support transactions that involve multiple clients, services, and resources. You will see later in this chapter just how transactions start and how they flow across service boundaries. For now, the important thing to note is that while a transaction is in progress, the transaction manager stays largely out of the way. New services may join the transaction, and every resource accessed is enlisted with that transaction. The services execute business logic, and the resources record the changes made under the scope of the transaction. During the transaction, all the services (and the clients participating in the transaction) must *vote* on whether they want to commit the changes they've performed or abort the transaction for whatever reason.

When the transaction ends (you will see when transactions end later in this chapter), the transaction manager checks the combined vote of the participating services. If any service or client voted to abort, the transaction is doomed: all the participating resources are instructed to discard the changes made during the transaction. If, however, all the services in the transaction voted to commit, the two-phase commit protocol starts. In the first phase, the transaction manager asks all the resources that took part in the transaction if they have any reservations about committing the changes recorded during the transaction. That is, if they were asked to commit, would they? Note that the transaction manager is not instructing the resources to commit the changes at this point; it is merely asking for their votes on the matter. At the end of

the first phase, the transaction manager has the combined vote of the resources. The second phase of the protocol is acting upon that combined vote. If all the resources voted to commit the transaction in the first phase, the transaction manager instructs all of them to commit the changes. But if even one of the resources said in phase one that it would not commit the changes, then in phase two the transaction manager instructs all the resources to roll back the changes made, thus aborting the transaction and restoring the system to its pre-transaction state.

It is important to emphasize that a resource voting that it would commit if asked to constitutes an unbreakable promise. If a resource votes to commit a transaction, it means that it cannot fail if subsequently, in the second phase, it is instructed to commit. The resource should verify before voting to commit that all the changes are consistent and legitimate. A resource can never go back on its vote. This is the basis for enabling distributed transactions, and the various resource vendors have gone to great lengths to implement this behavior exactly.

Resource Managers

A *resource manager* (RM) is any resource that supports both automatic enlistment and the two-phase commit protocol managed by one of the transaction managers. The resource must be able to detect that it is being accessed by a transaction and automatically enlist in it exactly once. The RM can be either a durable resource or a volatile resource, such as a transactional integer, string, or collection. While the RM must support the two-phase commit protocol, it can optionally also implement an optimized protocol used when it is the only RM in the transaction. That optimized protocol is called the *single-phase commit* protocol, when the RM is the one informing the transaction manager in one step about the success or failure of an attempt to commit.

Transaction Propagation

WCF can propagate transactions across the service boundary. This enables a service to participate in a client's transaction, and the client to include operations on multiple services in the same transaction. The client itself may or may not be a WCF service. Both the binding and the operation contract configuration control the decision as to whether or not the client's transaction is propagated to the service. I call any binding that is capable of propagating the client's transaction to the service if configured to do so a *transaction-aware* binding. Only the TCP, IPC, and WS bindings are transaction-aware.

Transaction Flow and Bindings

By default, transaction-aware bindings do not propagate transactions. The reason is that, like most everything else in WCF, this is an opt-in setting: the service host or

administrator has to explicitly give its consent to accepting incoming transactions, potentially from across the organization or the business boundaries. To propagate a transaction, you must explicitly enable it in the binding on both the service host and client sides. All transaction-aware bindings offer the Boolean property `TransactionFlow`, such as:

```
public class NetTcpBinding : Binding, ...
{
    public bool TransactionFlow
    {get;set;}
    //More members
}
```

`TransactionFlow` defaults to `false`. To enable propagation, simply set this property to `true`, either programmatically or in the host config file. For example, in the case of the TCP binding:

```
NetTcpBinding tcpBinding = new NetTcpBinding();
tcpBinding.TransactionFlow = true;
```

or when using a config file:

```
<bindings>
<netTcpBinding>
    <binding name = "TransactionalTCP"
        transactionFlow = "true"
    />
</netTcpBinding>
</bindings>
```

Transactions and Reliability

Strictly speaking, transactions do not require reliable messaging. The reason is that when reliability is disabled, if WCF messages are dropped or the client or service becomes disconnected, the transaction will abort. Because the client is guaranteed complete success or complete failure of the transactional operation, transactions are reliable in their own way. However, enabling reliability will decrease the likelihood of aborted transactions, because it will make the communication reliable; this means the transaction will be less likely to abort due to communication problems. I therefore recommend as a best practice also enabling reliability when enabling transactions with the `NetTcpBinding` and `WSHttpBinding`:

```
<netTcpBinding>
    <binding name = "TransactionalTCP" transactionFlow = "true">
        <reliableSession enabled = "true"/>
    </binding>
</netTcpBinding>
```

There is no need to enable reliability for the `NetNamedPipeBinding` because, as discussed in [Chapter 1](#), this binding is always reliable.

Transaction Flow and the Operation Contract

Using a transaction-aware binding and even enabling transaction flow does not mean that the service wants to use the client's transaction in every operation, or that the client necessarily has a transaction to propagate in the first place. Such service-level decisions should be part of the contractual agreement between the client and the service. To that end, WCF provides the `TransactionFlowAttribute` method attribute, which controls if and when the client's transaction flows into the service:

```
public enum TransactionFlowOption
{
    Allowed,
    NotAllowed,
    Mandatory
}

[AttributeUsage(AttributeTargets.Method)]
public sealed class TransactionFlowAttribute : Attribute, IOperationBehavior
{
    public TransactionFlowAttribute(TransactionFlowOption flowOption);
}
```

Note that the `TransactionFlow` attribute is a method-level attribute because WCF insists that the decision on transaction flow be made on a per-operation level, not at the service level:

```
[ServiceContract]
interface IMyContract
{
    [OperationContract]
    [TransactionFlow(TransactionFlowOption.Allowed)]
    void MyMethod();
}
```

This is deliberate, to enable the granularity of having some methods that use the client's transaction and some that do not.

The value of the `TransactionFlow` attribute is included in the published metadata of the service, so when you import a contract definition, the imported definition will contain the configured value. WCF will also let you apply the `TransactionFlow` attribute directly on the service class implementing the operation:

```
[ServiceContract]
interface IMyContract
{
    [OperationContract]
```

```

        void MyMethod();
    }
    class MyService : IMyContract
    {
        [TransactionFlow(TransactionFlowOption.Allowed)]
        public void MyMethod()
        {...}
    }

```

However, such use is discouraged because it splits the definition of the logical service contract that will be published.

TransactionFlowOption.NotAllowed

When the operation is configured to disallow transaction flow, the client cannot propagate its transaction to the service. Even if transaction flow is enabled at the binding and the client has a transaction, it will be silently ignored and will not propagate to the service. As a result, the service will never use the client's transaction, and the service and the client can select any binding with any configuration. `TransactionFlowOption.NotAllowed` is the default value of the `TransactionFlowOption` attribute, so these two definitions are equivalent:

```

[ServiceContract]
interface IMyContract
{
    [OperationContract]
    void MyMethod();
}

[ServiceContract]
interface IMyContract
{
    [OperationContract]
    [TransactionFlow(TransactionFlowOption.NotAllowed)]
    void MyMethod();
}

```

TransactionFlowOption.Allowed

When the operation is configured to allow transaction flow, if the client has a transaction, the service will allow the client's transaction to flow across the service boundary. However, just because the client propagates its transaction doesn't mean the service will necessarily use it. When you choose `TransactionFlowOption.Allowed`, the service can be configured to use any binding, be it transaction-aware or not, but the client and the service must be compatible in their binding configuration. In the context of transaction flow, "compatible" means that when the service operation allows transaction flow but the binding disallows it, the client should also disallow it in the binding on its side: trying to flow the client's transaction will cause an error, because the service will not understand the transaction information in the message. However,

when the service-side binding configuration is set to allow transaction flow, the client may or may not want to enable propagation on its side, so it may elect to set `TransactionFlow` to `false` in the binding even if the service has it set to `true`.

`TransactionFlowOption.Mandatory`

When the operation is configured with `TransactionFlowOption.Mandatory`, the service and the client must use a transaction-aware binding with transaction flow enabled. WCF verifies this requirement at the service load time and throws an `InvalidOperationException` if the service has at least one incompatible endpoint. `TransactionFlowOption.Mandatory` means the client must have a transaction to propagate to the service. Trying to call a service operation without a transaction results in a `FaultException` on the client side stating that the service requires a transaction. With mandatory flow, the client's transaction always propagates to the service, but again, the service may or may not use the client's transaction.



The test client `WcfTestClient.exe` discussed in [Chapter 1](#) does not support mandatory transaction flow. It does not create a transaction on the client side, and therefore will fail all calls to an operation that mandates transaction flow.

One-Way Calls

Propagating the client's transaction to the service requires, by its very nature, allowing the service to abort the client's transaction if so desired. This implies that you cannot flow the client's transaction to a service over a one-way operation, because that call does not have a reply message. WCF validates this at the service load time and will throw an exception when a one-way operation is configured with anything but `TransactionFlowOption.NotAllowed`:

```
//Invalid definition:  
[ServiceContract]  
interface IMyContract  
{  
    [OperationContract(IsOneWay = true)]  
    [TransactionFlow(TransactionFlowOption.Allowed)]  
    void MyMethod();  
}
```

Transaction Protocols and Managers

By and large, WCF developers need never concern themselves with transaction protocols and transaction managers. You should rely on WCF (and .NET) to select the correct protocol and transaction manager, and focus instead on implementing your logic. However, developers typically care a great deal about this issue, so I've chosen

to include a discussion of transaction protocols and managers (mostly to demonstrate that there is no practical reason to actually deal with them in the first place).

The transaction management protocol WCF chooses depends on the execution scope of the participating parties in the transaction. The word *protocol* may be misleading here, because in the abstract the protocol being used is the two-phase commit protocol. The differences between the transaction management protocols have to do with the type of remote calls and communication protocol used, and the kind of boundaries it can cross. The options are:

The lightweight protocol

This protocol is used to manage transactions in a local context only, inside the same app domain. It cannot propagate the transaction across the app domain boundary (let alone the process or machine boundary), nor can it flow the transaction across any service boundary (that is, from a client to a service). The lightweight protocol is used only inside a service or between two parties in the same app domain, never between services. The lightweight protocol yields the best performance compared with the other protocols.

The OleTx protocol

This protocol is used to propagate transactions across app domain, process, and machine boundaries, and to manage the two-phase commit protocol. The protocol uses RPC calls, and the exact binary format of the calls is Windows-specific. As a result of the use of both the RPC and the Windows-specific format, it cannot be used across firewalls or to interoperate with non-Windows parties. This is usually not a problem, because the primary use for the OleTx protocol is for managing transactions in an intranet, in a homogeneous Windows environment.

The WS-Atomic Transaction (WSAT) protocol

This protocol is similar to the OleTx protocol in that it too can propagate transactions across app domain, process, and machine boundaries and can be used to manage the two-phase commit protocol. However, unlike the OleTx protocol, the WSAT protocol is based on an industry standard and can typically be used across firewalls. Although you can use the WSAT protocol in an intranet in a heterogeneous environment, its primary use is for transaction management across the Internet, where multiple transaction managers are involved.

Protocols and Bindings

No binding supports the lightweight protocol, because the protocol cannot propagate transactions across the service boundary anyway. However, the various transaction-aware bindings differ in their support for the two other transaction-management protocols. The TCP and IPC bindings can be configured to work with both the OleTx and WSAT protocols, or with just one of them. Both bindings default to the OleTx protocol and will switch to the WSAT protocol if required. In addition, these two

intranet bindings let you configure the protocol either in a config file or programmatically, like any other binding property.

WCF provides the `TransactionProtocol` abstract class, defined as:

```
public abstract class TransactionProtocol
{
    public static TransactionProtocol Default
    {get;}
    public static TransactionProtocol OleTransactions
    {get;}
    public static TransactionProtocol WSAtomicTransactionOctober2004
    {get;}
    public static TransactionProtocol WSAtomicTransaction11
    {get;}
}
```

Both the TCP and IPC bindings offer the `TransactionProtocol` property of the matching type. For example:

```
public class NetTcpBinding : Binding, ...
{
    TransactionProtocol TransactionProtocol
    {get;set;}
    //More members
}
```

To set the protocol programmatically, first construct the specific binding type, then set the property using one of the static properties:

```
NetTcpBinding tcpBinding = new NetTcpBinding();
//Protocol only matters with propagation
tcpBinding.TransactionFlow = true;
tcpBinding.TransactionProtocol = TransactionProtocol.
    WSAtomicTransactionOctober2004;
```

Note that the transaction protocol configuration is meaningful only when transaction propagation is enabled as well.

To configure the protocol in a config file, define a `binding` section as usual:

```
<bindings>
    <netTcpBinding>
        <binding name = "TransactionalTCP"
            transactionFlow = "true"
            transactionProtocol = "WSAtomicTransactionOctober2004">
        />
    </netTcpBinding>
</bindings>
```

When you configure a protocol for the TCP or IPC binding, the service and the client must use the same protocol.

Since the TCP and IPC bindings can be used only in an intranet, there is really no practical value to configuring them for the WSAT protocol, and this ability is available largely for the sake of completeness.

The WS bindings (`WSHttpBinding`, `WS2007HttpBinding`, `WSDualHttpBinding`, `WSFederationHttpBinding`, and `WS2007FederationHttpBinding`) are designed for use across the Internet, when multiple transaction managers are involved, using the WSAT protocol. However, in an Internet scenario where only a single transaction manager is involved, these bindings will default to the OleTx protocol. There is no need or ability to configure a particular protocol.

Transaction Managers

Recall from the discussion at the beginning of this chapter that the last thing you should do is manage a transaction yourself. The best solution is to have a third party, called the transaction manager, manage the two-phase commit protocol for your clients and services. WCF can work with two different transaction managers in a provider model, as shown in [Figure 7-4](#).

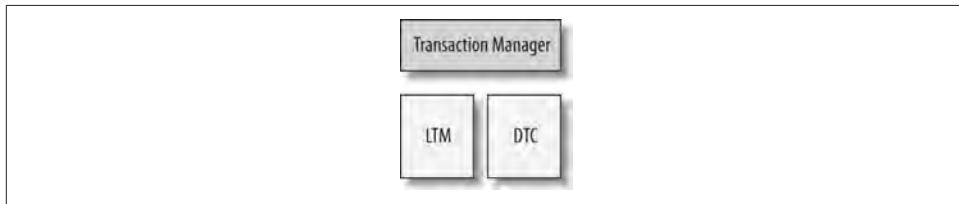


Figure 7-4. WCF transaction managers

These transaction managers are the *Lightweight Transaction Manager* (LTM) and the *Distributed Transaction Coordinator* (DTC). As a function of what the application does, the services it calls, and the resources it consumes, .NET (and Windows) will assign the appropriate transaction manager. Because the transaction manager is assigned automatically, your code is decoupled from the transaction management and from the transaction protocol used. Again, developers need never bother themselves with the transaction managers, and the following discussion is intended only to alleviate some common concerns regarding performance and efficiency.

The LTM

The LTM can manage only a local transaction; that is, a transaction inside a single app domain. The LTM uses the lightweight transaction protocol to manage the two-phase commit protocol. It can only manage a transaction that involves at most a single durable resource manager. The LTM can also manage as many volatile resource managers as are present. If only a single resource manager is present, and that resource supports single-phase commit, the LTM will use that optimized protocol.

Most importantly, the LTM can only manage a transaction inside a single service, and only when that service does not flow the transaction to other services. The LTM is the most performant transaction manager, on a par performance-wise with performing direct transactions against the resource.

The DTC

The DTC is capable of managing transactions across any execution boundary, from the most local scope (a transaction within the same app domain) to the most remote scope (a transaction that crosses process, machine, or site boundaries). The DTC can use either the OleTx or the WSAT protocol. The DTC is the transaction manager used when transactions flow across the service boundary. The DTC can easily manage a transaction that involves any number of services and resource managers.

The DTC is a system service available by default on every machine running WCF, and WCF (and .NET) is tightly integrated with the DTC. The DTC creates new transactions, collects the votes of the resource managers, and instructs the resource managers to abort or commit the transaction. For example, consider the application shown in [Figure 7-5](#), where a nontransactional client calls to a service on Machine A. The service on Machine A is configured to use a transaction. That service becomes the *root* of the transaction, and it will get the opportunity not just to start the transaction but also to indicate when the transaction is done.

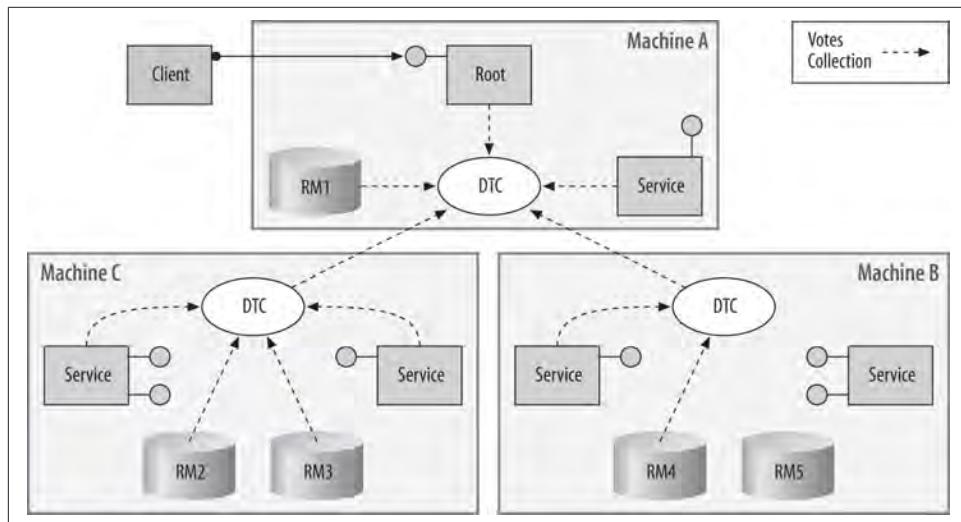


Figure 7-5. A DTC-managed transaction



Every transaction in WCF has at most one root service, because a non-service client can also be the root of the transaction.

When a service that is part of a transaction on Machine A tries to access another service or a resource on Machine B, it actually has a proxy to the remote service or resource. That proxy propagates the transaction ID to Machine B. The call interception on Machine B contacts the local DTC on Machine B, passing it the transaction ID and informing it to start managing that transaction on Machine B. Because the transaction ID gets propagated to Machine B, resource managers on Machine B can now autoenlist in it. Similarly, the transaction ID is propagated to Machine C.

When the transaction is done, if the combined services' vote was to try to commit the transaction, it is time to start the two-phase commit protocol. The DTC on the root machine collects the resource managers' votes on that machine and contacts the DTC on every other machine that took part in the transaction, instructing them to conduct the first phase on their machines. The DTCs on the remote machines collect the resource managers' votes on their machines and forward the results back to the DTC on the root machine. After the DTC on the root machine receives the results from all the remote DTCs, it has the combined resource managers' vote. If all of them vote to commit, the DTC on the root machine again contacts all the DTCs on the remote machines, instructing them to conduct phase two on their respective machines and to commit the transaction. If even one resource manager voted to abort the transaction, however, the DTC on the root machine informs all the DTCs on the remote machines to conduct phase two on their respective machines and to abort the transaction. Note that only the DTC on the root machine has the combined vote of phase one, and only it can instruct the final abort or commit.

Transaction Manager Promotion

The WCF transaction managers' architecture shown in [Figure 7-4](#) is that of a provider model. While presently there are only two such providers, in the past there was a third (the Kernel Transaction Manager, or KTM, now defunct). There is every reason to assume that in the future Microsoft will introduce additional transaction managers, likely as part of new development platforms or operating systems. The provider architecture is designed to dynamically assign the appropriate transaction manager for the transaction. If one transaction manager is inadequate, .NET will *promote* the transaction; that is, ask the next-level-up transaction manager to handle the transaction. A single transaction can be promoted multiple times. Once promoted, the transaction stays elevated and cannot be demoted. The previous transaction manager used to manage the transaction is relegated to a pass-through mode. Because of this dynamic promotion, developers are precluded from interacting with the transaction

managers directly (doing so would bypass promotion). Promotion is yet another reason why you should not write code such as that in [Example 7-1](#): it eliminates any chance of promotion.

LTM promotion

Every transaction in .NET always starts out as a transaction managed by the LTM. As long as the transaction interacts with a single durable resource and as long as there is no attempt to flow the transaction to a WCF service, the LTM can manage the transaction and yield the best throughput and performance. The LTM can manage as many volatile resource managers as required. However, if the transaction tries to enlist a second durable resource or the transaction is propagated to a service, .NET will promote the transaction from the LTM to the DTC.

Resources and promotion

The only resources that can participate in an LTM transaction and promote from it are volatile resource managers and the various flavors of SQL Server, from SQL Server 2005 and later. Legacy resource managers such as SQL Server 2000, Oracle, DB2, and MSMQ can only participate in DTC transactions. Consequently, when a legacy resource is accessed by an LTM transaction, even if it is the single resource in the transaction, the transaction is automatically promoted to the DTC.

The relationship between resources and transaction managers is summarized in [Table 7-1](#).

Table 7-1. Resources and transaction managers

Resource	LTM	DTC
Volatile	Yes	Yes
SQL Server 2005/2008	Yes	Yes
Kernel	No	Yes
Any other RM	No	Yes



With an Oracle database, ODP.NET Release 10.2.0.3 can support a local LTM transaction only or a distributed DTC transaction only, but cannot promote from the LTM to the DTC. An LTM transaction encountering a promotion event is not promoted to the DTC, but rather aborts with an exception. A transaction can also be configured to always start as a DTC transaction.

The Transaction Class

The `Transaction` class from the `System.Transactions` namespace represents the transaction that all .NET transaction managers work with:

```
[Serializable]
public class Transaction : IDisposable, ISerializable
{
    public static Transaction Current
    {get; set;}

    public void Rollback(); //Abort the transaction
    public void Dispose();

    //More members
}
```

Developers rarely need to interact with the `Transaction` class directly. The main use of the `Transaction` class is to manually abort a transaction by calling the `Rollback()` method. Additional features of the `Transaction` class include enlisting resource managers, setting the isolation level, subscribing to transaction events, cloning the transaction for concurrent threads, and obtaining the transaction status and other information.

The Ambient Transaction

.NET 2.0 defined a concept called the *ambient transaction*, which is the transaction in which your code executes. To obtain a reference to the ambient transaction, call the static `Current` property of `Transaction`:

```
Transaction ambientTransaction = Transaction.Current;
```

If there is no ambient transaction, `Current` will return `null`. Every piece of code, be it client or service, can always reach out for its ambient transaction. The ambient transaction object is stored in the thread local storage (TLS). As a result, when the thread winds its way across multiple objects and methods on the same call chain, all objects and methods can access their ambient transactions.

In the context of WCF, the ambient transaction is paramount. When present, any resource manager will automatically enlist in the ambient transaction. When a client calls a WCF service, if the client has an ambient transaction and the binding and the contract are configured to allow transaction flow, the ambient transaction will propagate to the service.



The client cannot propagate an already aborted transaction to the service. Doing so will yield an exception.

Local Versus Distributed Transactions

The `Transaction` class is used both for local and distributed transactions. Each transaction object has two identifiers used to identify the local and the distributed transaction. You obtain the transaction identifiers by accessing the `TransactionInformation` property of the `Transaction` class:

```
[Serializable]
public class Transaction : IDisposable, ISerializable
{
    public TransactionInformation TransactionInformation
    {get;}
    //More members
}
```

The `TransactionInformation` property is of the type `TransactionInformation`, defined as:

```
public class TransactionInformation
{
    public Guid DistributedIdentifier
    {get;}
    public string LocalIdentifier
    {get;}
    //More members
}
```

`TransactionInformation` offers access to the two identifiers. The main use of these identifiers is for logging, tracing, and analysis. In this chapter, I will use the identifiers as a convenient way to demonstrate transaction flow in code as a result of configuration.

The local transaction identifier

The *local transaction identifier* (local ID) contains both an identifier for the LTM in the current app domain as well as an ordinal number enumerating the transaction. You access the local ID via the `LocalIdentifier` property of `TransactionInformation`. The local ID is always available with the ambient transaction, and as such is never null: as long as there is an ambient transaction, it will have a valid local ID.

The value of the local ID has two parts: a constant GUID that is unique for each app domain and represents the assigned LTM for that app domain, and an incremented integer enumerating the transactions managed so far by that LTM.

For example, if a service traces three consecutive transactions, starting with the first call, it will get something like this:

```
8947aec9-1fac-42bb-8de7-60df836e00d6:1  
8947aec9-1fac-42bb-8de7-60df836e00d6:2  
8947aec9-1fac-42bb-8de7-60df836e00d6:3
```

The GUID is constant per app domain. If the service is hosted in the same app domain as the client, they will have the same GUID. If the client makes a cross-app domain call, the client will have its own unique GUID identifying its own local LTM.

The distributed transaction identifier

The *distributed transaction identifier* (distributed ID) is generated automatically whenever an LTM-managed transaction is promoted to a DTC-managed transaction (for example, when the ambient transaction flows to another service). You access the distributed ID via the `DistributedIdentifier` property of `TransactionInformation`. The distributed ID is unique per transaction, and no two transactions will ever have the same distributed ID. Most importantly, the distributed ID will be uniform across the service boundaries and across the entire call chain, from the topmost client through every service and object down the call chain. As such, it is useful in logging and tracing. Note that for a transaction that has not yet been promoted, the value of the distributed ID will be `Guid.Empty`. The distributed ID is usually `Guid.Empty` on the client side when the client is the root of the transaction and it has not yet called a service, and on the service side it will be empty if the service does not use the client's transaction and instead starts its own local transaction.

Transactional Service Programming

For services, WCF offers a simple and elegant declarative programming model. This model is, however, unavailable for non-service code called by services and for non-service WCF clients.

Setting the Ambient Transaction

By default, the service class and all its operations have no ambient transaction. This is the case even when the client's transaction is propagated to the service. Consider the following service:

```
[ServiceContract]  
interface IMyContract  
{  
    [OperationContract]  
    [TransactionFlow(TransactionFlowOption.Mandatory)]  
    void MyMethod();  
}  
class MyService : IMyContract
```

```

{
    public void MyMethod()
    {
        Transaction transaction = Transaction.Current;
        Debug.Assert(transaction == null);
    }
}

```

The ambient transaction of the service will be `null`, although the mandatory transaction flow guarantees the client's transaction propagation. To have an ambient transaction, for each contract method the service must indicate that it wants WCF to scope the body of the method with a transaction. For that purpose, WCF provides the `TransactionScopeRequired` property of the `OperationBehaviorAttribute`:

```

[AttributeUsage(AttributeTargets.Method)]
public sealed class OperationBehaviorAttribute : Attribute, ...
{
    public bool TransactionScopeRequired
    {get;set;}
    //More members
}

```

The default value of `TransactionScopeRequired` is `false`, which is why by default the service has no ambient transaction. Setting `TransactionScopeRequired` to `true` provides the operation with an ambient transaction:

```

class MyService : IMyContract
{
    [OperationBehavior(TransactionScopeRequired = true)]
    public void MyMethod()
    {
        Transaction transaction = Transaction.Current;
        Debug.Assert(transaction != null);
    }
}

```

If the client's transaction is propagated to the service, WCF will set the client's transaction as the operation's ambient transaction. If not, WCF will create a new transaction for that operation and set the new transaction as the ambient transaction.



The service class constructor does not have a transaction: it can never participate in the client's transaction, and you cannot ask WCF to scope it with a transaction. Unless you manually create a new ambient transaction (as shown later), do not perform transactional work in the service constructor and never expect to participate in the transaction of the client that created the instance inside the constructor.

Figure 7-6 demonstrates which transaction a WCF service uses as a product of the binding configuration, the contract operation, and the local operation behavior attribute.

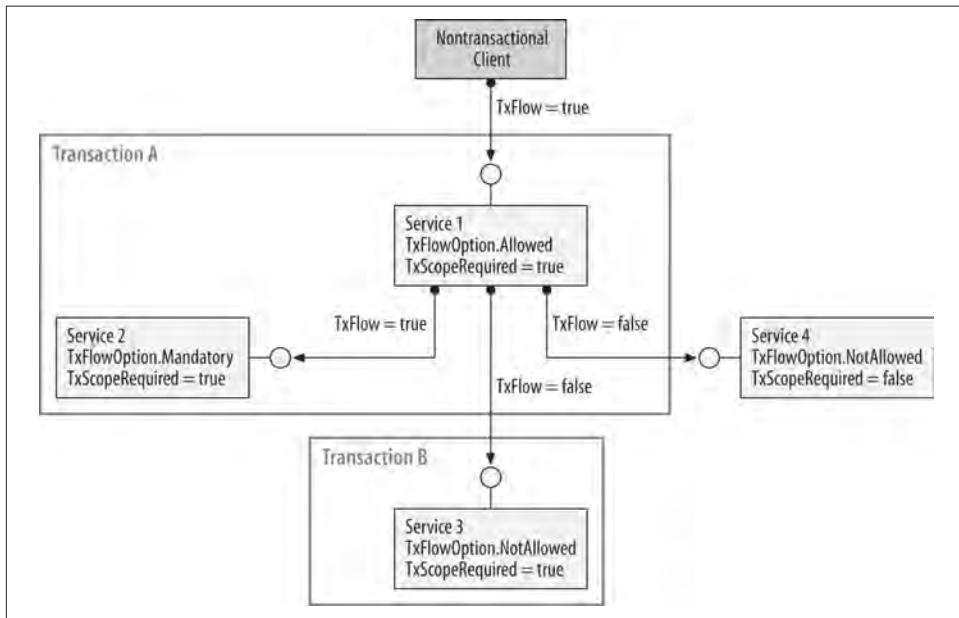


Figure 7-6. Transaction propagation as the product of contract, binding, and operation behavior

In the figure, a nontransactional client calls Service 1. The operation contract is configured with `TransactionFlowOption.Allowed`. Even though transaction flow is enabled in the binding, since the client has no transaction, no transaction is propagated. The operation behavior on Service 1 is configured to require a transaction scope. As a result, WCF creates a new transaction for Service 1 (Transaction A in **Figure 7-6**). Service 1 then calls three other services, each configured differently. The binding used for Service 2 has transaction flow enabled, and the operation contract mandates the flow of the client transaction. Since the operation behavior is configured to require transaction scope, WCF sets Transaction A as the ambient transaction for Service 2. The call to Service 3 has the binding and the operation contract disallow transaction flow. However, since Service 3 has its operation behavior require a transaction scope, WCF creates a new transaction for Service 3 (Transaction B) and sets it as the ambient transaction for Service 3. Similar to Service 3, the call to Service 4 has the binding and the operation contract disallow transaction flow. But since Service 4 does not require a transaction scope, it has no ambient transaction.

Transaction Propagation Modes

Which transaction the service uses is determined by the flow property of the binding (two values), the flow option in the operation contract (three values), and the value of the transaction scope property in the operation behavior (two values). There are therefore 12 possible configuration settings. Out of these 12, 4 are inconsistent and are precluded by WCF (such as flow disabled in the binding, yet mandatory flow in the operation contract) or are just plain impractical or inconsistent. [Table 7-2](#) lists the remaining eight permutations.²

Table 7-2. Transaction modes as the product of binding, contract, and behavior

Binding transaction flow	TransactionFlowOption	TransactionScopeRequired	Transaction mode
False	Allowed	False	None
False	Allowed	True	Service
False	NotAllowed	False	None
False	NotAllowed	True	Service
True	Allowed	False	None
True	Allowed	True	Client/Service
True	Mandatory	False	None
True	Mandatory	True	Client

Those eight permutations actually result in only four transaction propagation modes. I call these four modes *Client/Service*, *Client*, *Service*, and *None*. [Table 7-2](#) uses bold type to indicate the recommended way to configure each mode. Each of these modes has its place in designing your application, and understanding how to select the correct mode is not only a key to sound design, but also greatly simplifies thinking about and configuring transaction support.

Client/Service transaction mode

The Client/Service mode, as its name implies, ensures the service uses the client's transaction if possible, or a service-side transaction when the client does not have a transaction. To configure this mode:

1. Select a transactional binding and enable transaction flow by setting `TransactionFlow` to `true`.
2. Set the transaction flow option in the operation contract to `TransactionFlowOption.Allowed`.

² I first presented my designation of transaction propagation modes in the article "WCF Transaction Propagation" (*MSDN Magazine*, May 2007).

- Set the `TransactionScopeRequired` property of the operation behavior to `true`.

The Client/Service mode is the most decoupled configuration, because in this mode the service minimizes its assumptions about what the client is doing. The service will join the client's transaction if the client has a transaction to flow, which is always good for overall system consistency: if the service has a transaction separate from that of the client, one of those transactions could commit while the other aborts, leaving the system in an inconsistent state. However, if the service joins the client's transaction, all the work done by the client and the service (and potentially other services the client calls) will be committed or aborted as one atomic operation. If the client does not have a transaction, the service still requires the protection of a transaction, so this mode provides a contingent transaction to the service by making it the root of a new transaction.

[Example 7-2](#) shows a service configured for the Client/Service transaction mode.

Example 7-2. Configuring for the Client/Service transaction mode

```
[ServiceContract]
interface IMyContract
{
    [OperationContract]
    [TransactionFlow(TransactionFlowOption.Allowed)]
    void MyMethod();
}

class MyService : IMyContract
{
    [OperationBehavior(TransactionScopeRequired = true)]
    public void MyMethod()
    {
        Transaction transaction = Transaction.Current;
        Debug.Assert(transaction != null);
    }
}
```

Note in [Example 7-2](#) that the service can assert that it always has a transaction, but it cannot assume or assert whether it is the client's transaction or a locally created one. The Client/Service mode is applicable when the service can be used standalone or as part of a bigger transaction. When you select this mode, you should be mindful of potential deadlocks—if the resulting transaction is a service-side transaction, it may deadlock with other transactions trying to access the same resources, because the resources will isolate access per transaction and the service-side transaction will be a new transaction. When you use the Client/Service mode, the service may or may not be the root of the transaction, and the service must not behave differently when it is the root and when it is joining the client's transaction.

Requiring transaction flow

The Client/Service mode requires the use of a transaction-aware binding with transaction flow enabled, but this is not enforced by WCF at service load time. To tighten this loose screw, you can use my `BindingRequirementAttribute`:

```
[AttributeUsage(AttributeTargets.Class)]
public class BindingRequirementAttribute : Attribute,IServiceBehavior
{
    public bool TransactionFlowEnabled //Default is false
    {get;set;}
    //More members
}
```

You apply the attribute directly on the service class. The default of `TransactionFlowEnabled` is `false`. However, when you set it to `true`, per endpoint, if the contract of the endpoint has at least one operation with the `TransactionFlow` attribute configured with `TransactionFlowOption.Allowed`, the `BindingRequirement` attribute will enforce that the endpoint uses a transaction-aware binding with the `TransactionFlowEnabled` property set to `true`:

```
[ServiceContract]
interface IMyContract
{
    [OperationContract]
    [TransactionFlow(TransactionFlowOption.Allowed)]
    void MyMethod();
}

[BindingRequirement(TransactionFlowEnabled = true)]
class MyService : IMyContract
{...}
```

To enforce the binding requirement, in the case of a mismatch an `InvalidOperationException` is thrown when the host is launched. [Example 7-3](#) shows a somewhat simplified implementation of the `BindingRequirement` attribute.

Example 7-3. BindingRequirement attribute implementation

```
[AttributeUsage(AttributeTargets.Class)]
public class BindingRequirementAttribute : Attribute,IServiceBehavior
{
    public bool TransactionFlowEnabled
    {get;set;}

    void IServiceBehavior.Validate(ServiceDescription description,
                                  ServiceHostBase host)
    {
        if(TransactionFlowEnabled == false)
        {
```

```

        return;
    }
    foreach(ServiceEndpoint endpoint in description.Endpoints)
    {
        foreach(OperationDescription operation in endpoint.Contract.Operations)
        {
            TransactionFlowAttribute attribute =
                operation.Behaviors.Find<TransactionFlowAttribute>();
            if(attribute != null)
            {
                if(attribute.Transactions == TransactionFlowOption.Allowed)
                {
                    try
                    {
                        dynamic binding = endpoint.Binding;
                        if(binding.TransactionFlow == false)
                        {
                            throw new InvalidOperationException(...);
                        }
                        continue;
                    }
                    catch(RuntimeBinderException)//Not transaction aware binding
                    {
                        throw new InvalidOperationException(...);
                    }
                }
            }
        }
    }
}
void IServiceBehavior.AddBindingParameters(...)
{}
void IServiceBehavior.ApplyDispatchBehavior(...)
{}
}

```

The `BindingRequirementAttribute` class is a service behavior, so it supports the `IServiceBehavior` interface introduced in [Chapter 6](#). The `Validate()` method of `IServiceBehavior` is called during the host launch time, enabling you to abort the service load sequence. The first thing `Validate()` does is to check whether the `TransactionFlowEnabled` property is set to `false`. If so, `Validate()` does nothing and returns. If `TransactionFlowEnabled` is `true`, `Validate()` iterates over the collection of service endpoints available in the service description. For each endpoint, it obtains the collection of operations, and for each operation, it accesses its collection of operation behaviors. All operation behaviors implement the `IOperationBehavior` interface, including the `TransactionFlowAttribute`. If the `TransactionFlowAttribute` behavior is found, `Validate()` checks whether the attribute is configured with `TransactionFlowOption.Allowed`. If so, `Validate()` checks the binding. For each transaction-aware binding, it verifies that the `TransactionFlow` property is set to

true, and if not, it throws an `InvalidOperationException`. `Validate()` also throws an `InvalidOperationException` if a nontransactional binding is used for the endpoint.



The technique shown in [Example 7-3](#) for implementing the `BindingRequirement` attribute is a general-purpose technique you can use to enforce any binding requirement or custom validation rule. For example, the `BindingRequirement` attribute has another property, called `WCFOnly`, that enforces the use of WCF-to-WCF bindings only, and a `ReliabilityRequired` property that insists on the use of a reliable binding with reliability enabled:

```
[AttributeUsage(AttributeTargets.Class)]
public class BindingRequirementAttribute :
    Attribute, IServiceBehavior
{
    public bool ReliabilityRequired
    {get;set;}
    public bool TransactionFlowEnabled
    {get;set;}
    public bool WCFOnly
    {get;set;}
}
```

Client transaction mode

The Client mode ensures the service uses only the client's transaction. To configure this mode:

1. Select a transactional binding and enable transaction flow by setting `TransactionFlow` to `true`.
2. Set the transaction flow option in the operation contract to `TransactionFlowOption.Mandatory`.
3. Set the `TransactionScopeRequired` property of the operation behavior to `true`.

You should select the Client transaction mode when the service must use its client's transactions and can never be used standalone, by design. The main motivation for this is to maximize overall system consistency, since the work of the client and the service is always treated as one atomic operation. Another motivation is that by having the service share the client's transaction you reduce the potential for a deadlock, because all resources accessed will enlist in the same transaction. This means no other transactions will compete for access to the same resources and underlying locks.

[Example 7-4](#) shows a service configured for the Client transaction mode.

Example 7-4. Configuring for the Client transaction mode

```
[ServiceContract]
interface IMyContract
{
    [OperationContract]
    [TransactionFlow(TransactionFlowOption.Mandatory)]
    void MyMethod();
}
class MyService : IMyContract
{
    [OperationBehavior(TransactionScopeRequired = true)]
    public void MyMethod()
    {
        Transaction transaction = Transaction.Current;
        Debug.Assert(transaction.TransactionInformation.
            DistributedIdentifier != Guid.Empty);
    }
}
```

Note in [Example 7-4](#) that `MyMethod()` asserts the fact that the ambient transaction is a distributed one, meaning it originated with the client.

Service transaction mode

The Service mode ensures that the service always has a transaction, separate from any transaction its clients may or may not have. The service will always be the root of a new transaction. To configure this mode:

1. You can select any binding. If you select a transaction-aware binding, leave its default value for the `TransactionFlow` property, or explicitly set it to `false`.
2. Do not apply the `TransactionFlow` attribute, or configure it with `TransactionFlowOption.NotAllowed`.
3. Set the `TransactionScopeRequired` property of the operation behavior to `true`.

You should select the Service transaction mode when the service needs to perform transactional work outside the scope of the client's transaction (e.g., when you want to perform some logging or audit operations, or when you want to publish events to subscribers regardless of whether the client's transaction commits or aborts). As an example, consider a logbook service that performs error logging into a database. When an error occurs on the client side, the client will use the logbook service to log it or some other entries. But after it's logged, the error on the client side aborts the client's transaction. If the service were to use the client's transaction, once the client's transaction aborts, the logged error would be discarded from the database, and you would have no trace of it (defeating the purpose of the logging in the first place).

Configuring the service to have its own transaction, on the other hand, ensures that the log of the error is committed even when the client's transaction aborts.

The downside, of course, is the potential for jeopardizing the consistency of the system, because the service's transaction could abort while the client's commits. To avoid this pitfall, if the service-side transaction aborts, WCF throws an exception on the calling client side, even if the client was not using transactions or if the binding did not propagate any transaction. I therefore recommend that you only choose the Service mode if you have a supporting heuristic. The heuristic must be that the service's transaction is much more likely to succeed and commit than the client's transaction. In the example of the logging service, this is often the case, because once deterministic logging is in place it will usually work (unlike business transactions, which may fail for a variety of reasons).

In general, you should be extremely careful when using the Service transaction mode, and verify that the two transactions (the client's and the service's) do not jeopardize consistency if one aborts and the other commits. Logging and auditing services are the classic candidates for this mode.

Example 7-5 shows a service configured for the Service transaction mode.

Example 7-5. Configuring for the Service transaction mode

```
[ServiceContract]
interface IMyContract
{
    [OperationContract]
    void MyMethod();
}

class MyService : IMyContract
{
    [OperationBehavior(TransactionScopeRequired = true)]
    public void MyMethod()
    {
        Transaction transaction = Transaction.Current;
        Debug.Assert(transaction.TransactionInformation.
            DistributedIdentifier == Guid.Empty);
    }
}
```

Note in **Example 7-5** that the service can assert that it actually has a local transaction.

None transaction mode

If the None transaction mode is configured, the service never has a transaction. To configure this mode:

1. You can select any binding. If you select a transaction-aware binding, leave its default value for the `TransactionFlow` property, or explicitly set it to `false`.
2. Do not apply the `TransactionFlow` attribute, or configure it with `TransactionFlowOption.NotAllowed`.
3. You do not need to set the `TransactionScopeRequired` property of the operation behavior, but if you do, you should set it to `false`.

The `None` transaction mode is useful when the operations performed by the service are nice to have but not essential, and should not abort the client's transaction if they fail. For example, a service that prints a receipt for a money transfer should not be able to abort the client transaction if it fails because the printer is out of paper. Another example where the `None` mode is useful is when you want to provide some custom behavior, and you need to perform your own programmatic transaction support or manually enlist resources (for example, when calling legacy code, as in [Example 7-1](#)). Obviously, there is danger when using the `None` mode because it can jeopardize the system's consistency. Say the calling client has a transaction and it calls a service configured for the `None` transaction mode. If the client aborts its transaction, changes made to the system state by the service will not roll back. Another pitfall of this mode is that if a service configured for the `None` mode calls another service configured for the `Client` mode, the call will fail because the calling service has no transaction to propagate.

[Example 7-6](#) shows a service configured for the `None` transaction mode.

Example 7-6. Configuring for the None transaction mode

```
[ServiceContract]
interface IMyContract
{
    [OperationContract]
    void MyMethod();
}

class MyService : IMyContract
{
    public void MyMethod()
    {
        Transaction transaction = Transaction.Current;
        Debug.Assert(transaction == null);
    }
}
```

Note that the service in [Example 7-6](#) can assert that it has no ambient transaction.

The `None` mode allows you to have a nontransactional service be called by a transactional client. As stated previously, the `None` mode is typically used for services that

perform nice-to-have operations. The problem with this usage is that any exception thrown by the None service will abort the calling client's transaction, which should be avoided for mere nice-to-have operations. The solution is to have the client catch all exceptions from the None service to avoid contaminating the client's transaction. For example, here's how a client could call the service from [Example 7-6](#):

```
MyContractClient proxy = new MyContractClient();
try
{
    proxy.MyMethod();
    proxy.Close();
}
catch
{}
```



You need to encase the call to the None service in a `try/catch` statement even when configuring that service's operations as one-way operations, because one-way operations can still throw delivery exceptions.

Choosing a service transaction mode

The Service and None transaction modes are somewhat esoteric. They are useful in the context of the particular scenarios I've mentioned, but in other scenarios they harbor the danger of jeopardizing the system's consistency. You should typically use the Client/Service or Client transaction mode. Choose between these two based on the ability of the service to be used standalone (that is, based on the consistency consequences of using the service in its own transaction, and the potential for a deadlock). Avoid the Service and None modes.

Voting and Completion

Although WCF (and .NET underneath) is responsible for every aspect of transaction propagation and the overall management of the two-phase commit protocol across the resource managers, it does not itself know whether a transaction should commit or abort. WCF simply has no way of knowing whether the changes made to the system state are consistent (that is, if they make sense). Every participating service must vote on the outcome of the transaction and voice an opinion about whether the transaction should commit or abort. In addition, WCF does not know when to start the two-phase commit protocol; that is, when the transaction ends and when all the services are done with their work. That too is something the services (actually, just the root service) need to indicate to WCF. WCF offers two programming models for services to vote on the outcome of the transaction: a declarative model and an explicit model. As you will see, voting is strongly related to completing and ending the transaction.

Declarative voting

WCF can automatically vote on behalf of a service to commit or abort the transaction. Automatic voting is controlled via the Boolean `TransactionAutoComplete` property of the `OperationBehavior` attribute:

```
[AttributeUsage(AttributeTargets.Method)]
public sealed class OperationBehaviorAttribute : Attribute, ...
{
    public bool TransactionAutoComplete
    {get;set;}
    //More members
}
```

The `TransactionAutoComplete` property defaults to `true`, so these two definitions are equivalent:

```
[OperationBehavior(TransactionScopeRequired = true,
    TransactionAutoComplete = true)]
public void MyMethod()
{...}

[OperationBehavior(TransactionScopeRequired = true)]
public void MyMethod()
{...}
```

When this property is set to `true`, if there were no unhandled exceptions in the operation, WCF will automatically vote to commit the transaction. If there was an unhandled exception, WCF will vote to abort the transaction. Note that although WCF has to catch the exception in order to abort the transaction, it then rethrows it, allowing it to go up the call chain.

In order to rely on automatic voting, the service method must have `TransactionScopeRequired` set to `true`, because automatic voting only works when it was WCF that set the ambient transaction for the service.

It is very important when `TransactionAutoComplete` is set to `true` to avoid catching and handling exceptions and explicitly voting to abort:

```
//Avoid
[OperationBehavior(TransactionScopeRequired = true)]
public void MyMethod()
{
    try
    {
        ...
    }
    catch
    {
        Transaction.Current.Rollback();
    }
}
```

Although your service catches the exception, the operation will still result in an exception since WCF will throw an exception such as `TransactionAbortedException` on the client side. WCF does that because your service could be part of a much larger transaction that spans multiple services, machines, and sites. All other parties involved in this transaction are working hard, consuming system resources and locking out other parties, yet it is all in vain because your service has voted to abort, and nobody knows about it. By returning an exception to the client WCF ensures that the exception will abort all objects in its path, eventually reaching the root service or client and terminating the transaction. This will improve throughput and performance. If you want to catch the exception for some local handling such as logging, make sure to rethrow it:

```
[OperationBehavior(TransactionScopeRequired = true)]
public void MyMethod()
{
    try
    {
        ...
    }
    catch
    {
        /* Some local handling here */
        throw;
    }
}
```

Explicit voting

Explicit voting is required when `TransactionAutoComplete` is set to `false`. You can only set `TransactionAutoComplete` to `false` when `TransactionScopeRequired` is set to `true`.

When declarative voting is disabled, WCF will vote to abort all transactions by default, regardless of exceptions or a lack thereof. You must explicitly vote to commit using the `SetTransactionComplete()` method of the operation context:

```
public sealed class OperationContext : ...
{
    public void SetTransactionComplete();
    //More members
}
```

Make sure you do not perform any work, especially transactional work, after the call to `SetTransactionComplete()`. Calling `SetTransactionComplete()` should be the last line of code in the operation just before it returns:

```
[OperationBehavior(TransactionScopeRequired = true,
                    TransactionAutoComplete = false)]
public void MyMethod()
{
```

```

    /* Do transactional work here, then: */
    OperationContext.Current.SetTransactionComplete();
}

```

If you try to perform any transactional work (including accessing `Transaction.Current`) after the call to `SetTransactionComplete()`, WCF will throw an `InvalidOperationException` and abort the transaction.

Not performing any work after `SetTransactionComplete()` ensures that any exception raised before the call to `SetTransactionComplete()` will cause `SetTransactionComplete()` to be skipped, so WCF will default to aborting the transaction. As a result, there is no need to catch the exception, unless you want to do some local handling. As with declarative voting, since the method aborts, WCF will return a `TransactionAbortedException` to the client. In the interest of readability, if you do catch the exception, make sure to rethrow it:

```

[OperationBehavior(TransactionScopeRequired = true,
                  TransactionAutoComplete = false)]
public void MyMethod()
{
    try
    {
        /* Do transactional work here, then: */
        OperationContext.Current.SetTransactionComplete();
    }    catch
    {
        /* Do some error handling then */
        throw;
    }
}

```

Explicit voting is designed for the case when the vote depends on other information obtained throughout the transaction (besides exceptions and errors). However, for the vast majority of applications and services, you should prefer the simplicity of declarative voting.



Setting `TransactionAutoComplete` to `false` should not be done lightly. In fact, it is only allowed for per-session services with required session mode, because it has drastic effects on the service instance's affinity to a transaction. (In order to obtain information for the vote throughout a transaction, it must be the same transaction and the same instance.) You will see later why, when, and how you can set `TransactionAutoComplete` to `false`.

Terminating a transaction

When the transaction ends is determined by who starts it. Consider a client that either does not have a transaction or just does not propagate its transaction to the

service. If that client calls a service operation configured with `TransactionScopeRequired` set to `true`, that service operation becomes the root of the transaction. The root service can call other services and propagate the transaction to them. The transaction will end once the root operation completes the transaction, which it can do either declaratively by setting `TransactionAutoComplete` to `true`, or explicitly by setting it to `false` and calling `SetTransactionComplete()`. This is partly why both `TransactionAutoComplete` and `SetTransactionComplete()` are named the way they are; they are used for more than just voting; they complete and terminate the transaction for a root service. Note, however, that any of the downstream services called by the root operation can only use them to vote on the transaction, not to complete it. Only the root both votes on and completes the transaction.

When a non-service client starts the transaction, the transaction ends when the client disposes of the transaction object. You will see more on that in the section on explicit transaction programming.

Transaction Isolation

In general, the more isolated transactions are, the more consistent their results will be. The highest degree of isolation is called **Serializable**. At this level, the results obtained from a set of concurrent transactions are identical to the results that would be obtained by running each transaction serially. To achieve this goal, all the resources a transaction touches must be locked from any other transaction. If other transactions try to access those resources, they are blocked and cannot continue executing until the original transaction commits or aborts.

The isolation level is defined using the `IsolationLevel` enumeration, found in the `System.Transactions` namespace:

```
public enum IsolationLevel
{
    Serializable,
    RepeatableRead,
    ReadCommitted,
    ReadUncommitted,
    Snapshot, //Special form of ReadCommitted supported by SQL 2005 and later
    Chaos,    //No isolation whatsoever
    Unspecified
}
```

The difference between the four isolation levels (`ReadUncommitted`, `ReadCommitted`, `RepeatableRead`, and `Serializable`) is in the way the different levels use read and write locks. A lock can be held only while the transaction is accessing the data in the resource manager, or it can be held until the transaction is committed or aborted: the former is better for throughput; the latter for consistency. The two kinds of locks and the two kinds of operations (read/write) give four basic isolation levels. However, not

all resource managers support all levels of isolation, and they may elect to take part in the transaction at a higher level than the one configured for it. Every isolation level apart from `Serializable` is susceptible to some sort of inconsistency resulting from more than one transaction accessing the same information.

Selecting an isolation level other than `Serializable` is commonly used for read-intensive systems, and it requires a solid understanding of transaction processing theory and of the semantics of the transaction itself, the concurrency issues involved, and the consequences for system consistency. The reason other isolation levels are available is that a high degree of isolation comes at the expense of overall system throughput, because the resource managers involved have to hold on to both read and write locks for as long as a transaction is in progress, and all other transactions are blocked. However, there are some situations where you may be willing to trade system consistency for throughput by lowering the isolation level. Imagine, for example, a banking system where one of the requirements is to retrieve the total amount of money in all customer accounts combined. Although it would be possible to execute that transaction with the `Serializable` isolation level, if the bank has hundreds of thousands of accounts, it might take quite a while to complete. The transaction might also time out and abort, because some accounts could be accessed by other transactions at the same time. However, the number of accounts may be a blessing in disguise. On average (statistically speaking), if the transaction is allowed to run at a lower transaction level, it may get the wrong balance for some accounts, but those incorrect balances will tend to cancel each other out. The actual resulting error may be acceptable for the bank's needs.

In WCF, the isolation level is a service behavior, so all service operations use the same configured isolation level. Isolation is configured via the `TransactionIsolationLevel` property of the `ServiceBehavior` attribute:

```
[AttributeUsage(AttributeTargets.Class)]
public sealed class ServiceBehaviorAttribute : Attribute, ...
{
    public IsolationLevel TransactionIsolationLevel
    {get;set;}
    //More members
}
```

You can only set the `TransactionIsolationLevel` property if the service has at least one operation configured with `TransactionScopeRequired` set to `true`. There is no way to configure the isolation level in the host configuration file.

Isolation and transaction flow

The default value of `TransactionIsolationLevel` is `IsolationLevel.Unspecified`, so these two statements are equivalent:

```

class MyService : IMyContract
{...}

[ServiceBehavior(TransactionIsolationLevel = IsolationLevel.Unspecified)]
class MyService : IMyContract
{...}

```

When a service configured with `IsolationLevel.Unspecified` joins the client transaction, the service will use the client's isolation level. However, if the service specifies an isolation level other than `IsolationLevel.Unspecified`, the client must match that level, and a mismatch will throw an exception.

When the service is the root of the transaction and the service is configured with `IsolationLevel.Unspecified`, WCF will set the isolation level to `IsolationLevel.Serializable`. If the root service provides a level other than `IsolationLevel.Unspecified`, WCF will use that specified level.

Transaction Timeout

The introduction of isolation locks raises the possibility of a deadlock when one transaction tries to access a resource manager owned by another. If a transaction takes a long time to complete, it may be indicative of a transactional deadlock. To address that possibility, the transaction will automatically abort if it takes longer than a predetermined timeout (60 seconds, by default) to complete, even if no exceptions took place. Once it's aborted, any attempt to flow that transaction to a service will result in an exception.

The timeout is a service behavior property, and all operations across all endpoints of the service use the same timeout. You configure the timeout by setting the `TransactionTimeout` time-span string property of `ServiceBehaviorAttribute`:

```

[AttributeUsage(AttributeTargets.Class)]
public sealed class ServiceBehaviorAttribute : Attribute, ...
{
    public string TransactionTimeout
    {get;set;}
    //More members
}

```

For example, you would use the following to configure a 30-second timeout:

```

[ServiceBehavior(TransactionTimeout = "00:00:30")]
class MyService : ...
{...}

```

You can also configure the transaction timeout in the host config file by creating a custom `behavior` section and referencing it in the `service` section:

```

<services>
    <service name = "MyService" behaviorConfiguration =

```

```

        "ShortTransactionBehavior">
    ...
  </service>
</services>
<behaviors>
  <serviceBehaviors>
    <behavior name = "ShortTransactionBehavior"
      transactionTimeout = "00:00:30"
    />
  </serviceBehaviors>
</behaviors>

```

The maximum allowed transaction timeout is 10 minutes, and this value will be used even when a larger value is specified. If you want to override the default maximum timeout of 10 minutes and specify, say, 40 minutes, add (or modify) the following in *machine.config*:

```

<configuration>
  <system.transactions>
    <machineSettings maxTimeout = "00:40:00"/>
  </system.transactions>
</configuration>

```



Setting any value in *machine.config* will affect all applications on the machine.

Configuring such a long timeout is useful mostly for debugging, when you want to try to isolate a problem in your business logic by stepping through your code and you do not want the transaction you're debugging to time out while you figure out the problem. Be extremely careful with using a long timeout in all other cases, because it means there are no safeguards against transaction deadlocks.

You may also want to set the timeout to a value less than the default 60 seconds. You typically do this in two cases. The first is during development, when you want to test the way your application handles aborted transactions. By setting the timeout to a small value (such as one millisecond), you can cause your transactions to fail so you can observe your error-handling code.

The second case where it can be useful to set the transaction timeout to less than the default value is when you have reason to believe that a service is involved in more than its fair share of resource contention, resulting in deadlocks. If you are unable to redesign and redeploy the service, you want to abort the transaction as soon as possible and not wait for the default timeout to expire.

Transaction flow and timeout

When a transaction flows into a service that is configured with a shorter timeout than that of the incoming transaction, the transaction adopts the service's timeout and the service gets to enforce the shorter timeout. This behavior is designed to support resolving deadlocks in problematic services, as just discussed. When a transaction flows into a service that is configured with a longer timeout than the incoming transaction, the service configuration has no effect.

Explicit Transaction Programming

The transactional programming model described so far can only be used declaratively by transactional services. Non-service clients, nontransactional services, and plain .NET objects called downstream by a service cannot take advantage of it. For all these cases, WCF relies on the transactional infrastructure available with .NET in the `System.Transactions` namespace. You may also rely on `System.Transactions` even in transactional services, when exploiting some advanced features such as transaction events, cloning, asynchronous commits, and manual transactions. I described the capabilities of `System.Transactions` in my MSDN whitepaper “Introducing `System.Transactions` in the .NET Framework 2.0” (published April 2005; updated December 2005). The following sections contain excerpts from that article describing how to use the core aspects of `System.Transactions` in the context of WCF. Refer to the whitepaper for detailed discussions of the rest of the features.

The TransactionScope Class

The most common way of using transactions explicitly is via the `TransactionScope` class:

```
public sealed class TransactionScope : IDisposable
{
    public TransactionScope();
    //Additional constructors

    public void Complete();
    public void Dispose();
}
```

As its name implies, the `TransactionScope` class is used to scope a code section with a transaction, as demonstrated in [Example 7-7](#).

Example 7-7. Using TransactionScope

```
using(TransactionScope scope = new TransactionScope())
{
    /* Perform transactional work here */
```

```
//No errors- commit transaction  
scope.Complete();  
}
```

The scope constructor can create a new LTM transaction and make it the ambient transaction by setting `Transaction.Current`, or can join an existing ambient transaction. `TransactionScope` is a disposable object—if the scope creates a new transaction, the transaction will end once the `Dispose()` method is called (this is done at the end of the `using` statement in [Example 7-7](#)). The `Dispose()` method also restores the original ambient transaction (`null`, in the case of [Example 7-7](#)).

Finally, if the `TransactionScope` object is not used inside a `using` statement, it will become eligible for garbage collection once the transaction timeout has expired and the transaction is aborted.

TransactionScope voting

The `TransactionScope` object has no way of knowing whether the transaction should commit or abort. To address this, every `TransactionScope` object has a consistency bit, which by default is set to `false`. You can set the consistency bit to `true` by calling the `Complete()` method. Note that you can only call `Complete()` once; subsequent calls to `Complete()` will raise an `InvalidOperationException`. This is deliberate, to encourage developers to have no transactional code after the call to `Complete()`.

If the transaction ends (due to calling `Dispose()` or garbage collection) and the consistency bit is set to `false`, the transaction will abort. For example, the following scope object will abort its transaction, because the consistency bit is never changed from its default value:

```
using(TransactionScope scope = new TransactionScope())  
{}
```

By having the call to `Complete()` as the last action in the scope, you have an automated way for voting to abort in the case of an error: any exception thrown inside the scope will skip over the call to `Complete()`, the `finally` block in the `using` statement will dispose of the `TransactionScope` object, and the transaction will abort. On the other hand, if you do call `Complete()` and the transaction ends with the consistency bit set to `true`, as in [Example 7-7](#), the transaction will try to commit. Note that after calling `Complete()`, you cannot access the ambient transaction, and an attempt to do so will result in an `InvalidOperationException`. You can access the ambient transaction again (via `Transaction.Current`) once the scope object has been disposed of.

The fact that the code in the scope called `Complete()` does not guarantee that the transaction will be committed. Even if you call `Complete()` and the scope is disposed of, all that will do is try to commit the transaction. The ultimate success or failure of that attempt depends upon the outcome of the two-phase commit protocol, which

may involve multiple resources and services of which your code is unaware. As a result, `Dispose()` will throw a `TransactionAbortedException` if it fails to commit the transaction. You can catch and handle that exception, perhaps by alerting the user, as shown in [Example 7-8](#).

Example 7-8. TransactionScope and error handling

```
try
{
    using(TransactionScope scope = new TransactionScope())
    {
        /* Perform transactional work here */
        //No errors- commit transaction
        scope.Complete();
    }
}
catch(TransactionAbortedException exception)
{
    Trace.WriteLine(exception.Message);
}
catch //Any other exception took place
{
    Trace.WriteLine("Cannot complete transaction");
    throw;
}
```

Transaction Flow Management

Transaction scopes can nest both directly and indirectly. In [Example 7-9](#), `scope2` simply nests directly inside `scope1`.

Example 7-9. Direct scope nesting

```
using(TransactionScope scope1 = new TransactionScope())
{
    using(TransactionScope scope2 = new TransactionScope())
    {
        scope2.Complete();
    }
    scope1.Complete();
}
```

The scope can also nest indirectly. This happens when you call a method that uses `TransactionScope` from within a method that uses its own scope, as is the case with the `RootMethod()` in [Example 7-10](#).

Example 7-10. Indirect scope nesting

```
void RootMethod()
{
    using(TransactionScope scope = new TransactionScope())
    {
        /* Perform transactional work here */
        SomeMethod();
        scope.Complete();
    }
}
void SomeMethod()
{
    using(TransactionScope scope = new TransactionScope())
    {
        /* Perform transactional work here */
        scope.Complete();
    }
}
```

A transaction scope can also nest in a service method, as in [Example 7-11](#). The service method may or may not be transactional.

Example 7-11. Scope nesting inside a service method

```
class MyService : IMyContract
{
    [OperationBehavior(TransactionScopeRequired = true)]
    public void MyMethod()
    {
        using(TransactionScope scope = new TransactionScope())
        {
            scope.Complete();
        }
    }
}
```

If the scope creates a new transaction for its use, it is called the *root scope*. Whether or not a scope becomes a root scope depends on the scope configuration and the presence of an ambient transaction. Once a root scope is established, there is an implicit relationship between it and all its nested scopes and any downstream services called.

The `TransactionScope` class provides several overloaded constructors that accept an enum of the type `TransactionScopeOption`:

```
public enum TransactionScopeOption
{
    Required,
    RequiresNew,
    Suppress
}
```

```

public sealed class TransactionScope : IDisposable
{
    public TransactionScope(TransactionScopeOption scopeOption);
    public TransactionScope(TransactionScopeOption scopeOption,
                           TransactionOptions transactionOptions);
    public TransactionScope(TransactionScopeOption scopeOption,
                           TimeSpan scopeTimeout);
    //Additional constructors and members
}

```

The value of `TransactionScopeOption` lets you control whether the scope takes part in a transaction and, if so, whether it will join the ambient transaction or will be the root scope of a new transaction.

For example, here is how you specify the value of `TransactionScopeOption` in the scope's constructor:

```

using(TransactionScope scope = new TransactionScope
      ( TransactionScopeOption.Required))
{...}

```

The default value for the scope option is `TransactionScopeOption.Required`, meaning this is the value used when you call one of the constructors that does not accept a `TransactionScopeOption` parameter. So, these two definitions are equivalent:

```

using(TransactionScope scope = new TransactionScope())
{...}

using(TransactionScope scope = new TransactionScope
      (TransactionScopeOption.Required))
{...}

```

The `TransactionScope` object determines which transaction to belong to when it is constructed. Once that's been determined, the scope will always belong to that transaction. `TransactionScope` bases its decision on two factors: whether an ambient transaction is present, and the value of the `TransactionScopeOption` parameter.

A `TransactionScope` object has three options:

- Join the ambient transaction
- Be a new scope root (that is, start a new transaction and have that transaction be the new ambient transaction inside its own scope)
- Not take part in a transaction at all

If the scope is configured with `TransactionScopeOption.Required` and an ambient transaction is present, the scope will join that transaction. If, on the other hand, there is no ambient transaction, the scope will create a new transaction and become the root scope.

If the scope is configured with `TransactionScopeOption.RequiresNew`, it will always be a root scope. It will start a new transaction, and its transaction will be the new ambient transaction inside the scope.

If the scope is configured with `TransactionScopeOption.Suppress`, it will never be part of a transaction, regardless of whether an ambient transaction is present. A scope configured with `TransactionScopeOption.Suppress` will always have `null` as its ambient transaction.

Voting inside a nested scope

It is important to realize that although a nested scope can join the ambient transaction of its parent scope, the two scope objects will have two distinct consistency bits. Calling `Complete()` in the nested scope has no effect on the parent scope:

```
using(TransactionScope scope1 = new TransactionScope())
{
    using(TransactionScope scope2 = new TransactionScope())
    {
        scope2.Complete();
    }
    //scope1's consistency bit is still false
}
```

Only if all the scopes, from the root scope down to the last nested scope, vote to commit the transaction will the transaction commit. In addition, only the root scope dictates the life span of the transaction. When a `TransactionScope` object joins an ambient transaction, disposing of that scope does not end the transaction. The transaction ends only when the root scope is disposed of, or when the service method that started the transaction returns.

`TransactionScopeOption.Required`

`TransactionScopeOption.Required` is not just the most commonly used value; it is also the most decoupled value. If your scope has an ambient transaction, the scope will join that ambient transaction to improve consistency. If there is no ambient transaction to join, the scope will provide the code with a new ambient transaction. When `TransactionScopeOption.Required` is used, the code inside the `TransactionScope` must not behave differently when it is the root and when it is just joining the ambient transaction. It should operate identically in both cases. On the service side, the most common use for `TransactionScopeOption.Required` is by non-service downstream classes called by the service, as shown in [Example 7-12](#).

Example 7-12. Using `TransactionScopeOption.Required` in a downstream class

```
class MyService : IMyContract
{
```

```

[OperationBehavior(TransactionScopeRequired = true)]
public void MyMethod()
{
    MyClass obj = new MyClass();
    obj.SomeMethod();
}
class MyClass
{
    public void SomeMethod()
    {
        using(TransactionScope scope = new TransactionScope())
        {
            //Do some work then
            scope.Complete();
        }
    }
}

```

While the service itself can use `TransactionScopeOption.Required` directly, such a practice adds no value:

```

class MyService : IMyContract
{
    [OperationBehavior(TransactionScopeRequired = true)]
    public void MyMethod()
    {
        //One transaction only
        using(TransactionScope scope = new TransactionScope())
        {
            //Do some work then
            scope.Complete();
        }
    }
}

```

The reason is obvious: the service can simply ask WCF to scope the operation with a transaction scope by setting `TransactionScopeRequired` to `true` (this is also the origin of that property's name). Note that although the service may use declarative voting, any downstream (or directly nested) scope must still explicitly call `Complete()` in order for the transaction to commit. The same is true when the service method uses explicit voting:

```

[OperationBehavior(TransactionScopeRequired = true,
                   TransactionAutoComplete = false)]
public void MyMethod()
{
    using(TransactionScope scope = new TransactionScope())
    {
        //Do some work then
        scope.Complete();
    }
}

```

```

    /* Do transactional work here, then: */
    OperationContext.Current.SetTransactionComplete();
}

```

In short, voting to abort in a scope with `TransactionScopeRequired` nested in a service call will abort the service's transaction regardless of exceptions or the use of declarative voting (via `TransactionAutoComplete`) or explicit voting by the service (via `SetTransactionComplete()`).

`TransactionScopeOption.RequiresNew`

Configuring the scope with `TransactionScopeOption.RequiresNew` is useful when you want to perform transactional work outside the scope of the ambient transaction (for example, when you want to perform some logging or audit operations, or when you want to publish events to subscribers, regardless of whether your ambient transaction commits or aborts):

```

class MyService : IMyContract
{
    [OperationBehavior(TransactionScopeRequired = true)]
    public void MyMethod()
    {
        //Two distinct transactions
        using(TransactionScope scope = new TransactionScope(
            TransactionScopeOption.RequiresNew))
        {
            //Do some work then
            scope.Complete();
        }
    }
}

```

Note that you must complete the scope in order for the new transaction to commit. You might consider encasing a scope that uses `TransactionScopeOption.RequiresNew` in a `try` and `catch` block to isolate it from the service's ambient transaction.

You should be extremely careful when using `TransactionScopeOption.RequiresNew`, and verify that the two transactions (the ambient transaction and the one created for your scope) do not jeopardize consistency if one aborts and the other commits.

`TransactionScopeOption.Suppress`

`TransactionScopeOption.Suppress` is useful for both the client and the service when the operations performed by the code section are nice to have but should not cause the ambient transaction to abort in the event that they fail. `TransactionScopeOption.Suppress` allows you to have a nontransactional code section inside a transactional scope or service operation, as shown in [Example 7-13](#).

Example 7-13. Using TransactionScopeOption.Suppress

```
[OperationBehavior(TransactionScopeRequired = true)]
public void MyMethod()
{
    try
    {
        //Start of nontransactional section
        using(TransactionScope scope =
               new TransactionScope(TransactionScopeOption.Suppress))
        {
            //Do nontransactional work here
        } //Restores ambient transaction here
    }
    catch
    {}
}
```

Note in [Example 7-13](#) that there is no need to call `Complete()` on the suppressed scope. Another example where `TransactionScopeOption.Suppress` is useful is when you want to provide some custom behavior and you need to perform your own programmatic transaction support or manually enlist resources.

That said, you should be careful when mixing transactional scopes or service methods with nontransactional scopes. That can jeopardize isolation and consistency, because changes made to the system state inside the suppressed scope will not roll back along with the containing ambient transaction. In addition, the nontransactional scope may have errors, but those errors should not affect the outcome of the ambient transaction. This is why, in [Example 7-13](#), the suppressed scope is encased in a `try` and `catch` statement that also suppresses any exception coming out of it.



Do not call a service configured for the Client transaction mode (basically, with mandatory transaction flow) inside a suppressed scope, because that call is guaranteed to fail.

TransactionScope timeout

If the code inside the transactional scope takes a long time to complete, it may be indicative of a transactional deadlock. To address that possibility, the transaction will automatically abort if it takes longer than a predetermined timeout to complete (60 seconds, by default). You can configure the default timeout in the application config file. For example, to configure a default timeout of 30 seconds, add this to the config file:

```
<system.transactions>
    <defaultSettings timeout = "00:00:30"/>
</system.transactions>
```

Placing the new default in the application config file affects all scopes used by all clients and services in that application. You can also configure a timeout for a specific transaction scope. A few of the overloaded constructors of `TransactionScope` accept a value of type `TimeSpan`, used to control the timeout of the transaction. For example:

```
public TransactionScope(TransactionScopeOption scopeOption,
                        TimeSpan scopeTimeout);
```

To specify a timeout different from the default of 60 seconds, simply pass in the desired value:

```
TimeSpan timeout = TimeSpan.FromSeconds(30);
using(TransactionScope scope = new TransactionScope(
    TransactionScopeOption.Required,
    timeout))
{...}
```

When a `TransactionScope` joins the ambient transaction, if it specifies a shorter timeout than the one the ambient transaction is set to, it has the effect of enforcing the new, shorter timeout on the ambient transaction; the transaction must end within the nested time specified, or it is automatically aborted. If the scope's timeout is greater than that of the ambient transaction, it has no effect.

TransactionScope isolation level

If the scope is a root scope, by default the transaction will execute with the isolation level set to `Serializable`. Some of the overloaded constructors of `TransactionScope` accept a structure of the type `TransactionOptions`, defined as:

```
public struct TransactionOptions
{
    public IsolationLevel IsolationLevel
    {get;set;}
    public TimeSpan Timeout
    {get;set;}
    //Other members
}
```

Although you can use the `TransactionOptions` `Timeout` property to specify a timeout, the main use for `TransactionOptions` is for specifying the isolation level. You can assign into the `TransactionOptions` `IsolationLevel` property a value of the enum type `IsolationLevel` presented earlier:

```
TransactionOptions options = new TransactionOptions();
options.IsolationLevel = IsolationLevel.ReadCommitted;
options.Timeout = TransactionManager.DefaultTimeout;

using(TransactionScope scope = new TransactionScope(
    TransactionScopeOption.Required,options))
{...}
```

When a scope joins an ambient transaction, it must be configured to use exactly the same isolation level as the ambient transaction; otherwise, an `ArgumentException` is thrown.

Non-Service Clients

Although services can take advantage of `TransactionScope`, by far its primary use is by non-service clients. Using a transaction scope is practically the only way a non-service client can group multiple service calls into single transaction, as shown in [Figure 7-7](#).

Having the option to create a root transaction scope enables the client to flow its transaction to services and to manage and commit the transaction based on the aggregated result of the services, as shown in [Example 7-14](#).

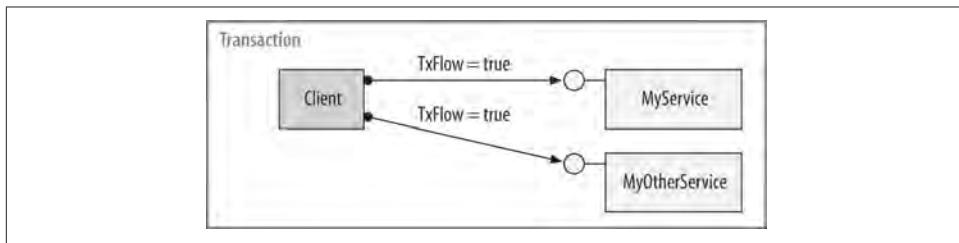


Figure 7-7. A non-service client using a single transaction to call multiple services

Example 7-14. Using TransactionScope to call services in a single transaction

```
//////////////////////////// Service Side //////////////////////////////
[ServiceContract]
interface IMyContract
{
    [OperationContract]
    [TransactionFlow(TransactionFlowOption.Allowed)]
    void MyMethod();
}

[ServiceContract]
interface IMyOtherContract
{
    [OperationContract]
    [TransactionFlow(TransactionFlowOption.Mandatory)]
    void MyOtherMethod();
}

class MyService : IMyContract
{
    [OperationBehavior(TransactionScopeRequired = true)]
    public void MyMethod()
    {...}
}

class MyOtherService : IMyOtherContract
```

```

{
    [OperationBehavior(TransactionScopeRequired = true)]
    public void MyOtherMethod()
    {...}
}
//////////////////////////// Client Side //////////////////////////////
using(TransactionScope scope = new TransactionScope())
{
    MyContractClient proxy1 = new MyContractClient();
    proxy1.MyMethod();
    proxy1.Close();

    MyOtherContractClient proxy2 = new MyOtherContractClient();
    proxy2.MyOtherMethod();
    proxy2.Close();

    scope.Complete();
}

//Can combine in single using block:
using(MyContractClient proxy3 = new MyContractClient())
using(MyOtherContractClient proxy4 = new MyOtherContractClient())
using(TransactionScope scope = new TransactionScope())
{
    proxy3.MyMethod();
    proxy4.MyOtherMethod();
    scope.Complete();
}

```

Service State Management

The sole purpose of transactional programming is to address the recovery challenge by always leaving the system in a consistent state. The state of the system consists of all the resources that were involved in the transaction, plus the in-memory clients and service instances. Besides benefits such as autoenlistment and participation in the two-phase commit protocol, the basic and obvious advantage of using a resource manager is that any change made to its state during a transaction will automatically roll back if the transaction aborts. This, however, is not true when it comes to the in-memory instance members and static members of the participating services, which means that if the transaction aborts, the system will not be left in a consistent state. The problem is compounded by the fact that any transaction a service participates in may span multiple services, machines, and sites. Even if that service instance encounters no errors and votes to commit the transaction, other parties across the service boundary may eventually abort the transaction. If the service were to simply store its state in memory, how would it know about the outcome of the transaction so that it would somehow manually roll back the changes it had made to its state?

The solution to this instance state management problem is to develop the service as a state-aware service and proactively manage its state. As explained in [Chapter 4](#), a state-aware service is not the same as a stateless service. If the service were truly stateless, there would not be any problem with instance state rollback. As long as a transaction is in progress, the service instance is allowed to maintain state in memory. Between transactions, the service should store its state in a resource manager. That state resource manager may be unrelated to any other business-logic-specific resource accessed during the transaction, or it may be one and the same. At the beginning of the transaction, the service should retrieve its state from the resource and, by doing so, enlist the resource in the transaction. At the end of the transaction, the service should save its state back to the resource manager.

The elegant thing about this technique is that it provides for state autorecovery. Any changes made to the instance state will commit or roll back as part of the transaction. If the transaction commits, the next time the service gets its state it will have the new state. If the transaction aborts, the next time it will have its pre-transaction state. Either way, the service will have a consistent state ready to be accessed by a new transaction. To force the service instance to purge all its in-memory state this way, by default, once the transaction completes, WCF destroys the service instance, ensuring there are no leftovers in memory that might jeopardize consistency.

The Transaction Boundary

There are two remaining problems when writing transactional state-aware services. The first is how a service can know when transactions start and end, so that it can get and save its state. The service may be part of a much larger transaction that spans multiple services and machines. At any moment between service calls, the transaction might end. Who will call the service, letting it know to save its state? The second problem has to do with isolation. Different clients might call the service concurrently, on different transactions. How can the service isolate changes made to its state by one transaction from another? The service cannot allow cross-transactional calls, because doing so would jeopardize isolation. If the other transaction were to access its state and operate based on its values, that transaction would be contaminated with foul state if the original transaction aborted and the changes rolled back.

The solution to both problems is for the service to equate method boundaries with transaction boundaries. At the beginning of every method call, the service should read its state, and at the end of each method call the service should save its state to the resource manager. Doing so ensures that if a transaction ends between method calls, the service's state will either persist or roll back with it. Because the service equates method boundaries with transaction boundaries, the service instance must also vote on the transaction's outcome at the end of every method call. From the service perspective, the transaction completes once the method returns. This is really why the `TransactionAutoComplete` property is named that instead of something like `Transac`

`tionAutoVote`: the service states that, as far as it is concerned, the transaction is complete. If the service is also the root of the transaction, completing it will indeed terminate the transaction.

In addition, reading and storing the state in the resource manager in each method call addresses the isolation challenge, because the service simply lets the resource manager isolate access to the state between concurrent transactions.

Instance Management and Transactions

As hinted previously, the transactional configuration of the service is intimately related to the service instance lifecycle, and it drastically changes the programming model. All transactional services must store their state in a resource manager or managers. Those resource managers could be volatile or durable, shared between the instances or per instance, and could support multiple services, all according to your design of both the service and its resources.

Volatile Resource Managers

In the article “Volatile Resource Managers in .NET Bring Transactions to the Common Type” (*MSDN Magazine*, May 2005), I presented my technique for implementing a general-purpose volatile resource manager called `Transactional<T>`:

```
public class Transactional<T> : ...
{
    public Transactional(T value);
    public Transactional();
    public T Value
    {get;set;}

    /* Conversion operators to and from T */
}
```

By specifying any serializable type parameter (such as an `int` or a `string`) to `Transactional<T>`, you turn that type into a full-blown volatile resource manager that auto-enlists in the ambient transaction, participates in the two-phase commit protocol, and isolates the current changes from all other transactions using my original transaction-based lock.

For example, in the following code snippet, the scope is not completed. As a result, the transaction aborts, and the values of `number` and `city` revert to their pre-transaction state:

```
Transactional<int> number = new Transactional<int>(3);
Transactional<string> city =
    new Transactional<string>("New York");

using(TransactionScope scope = new TransactionScope())
```

```

{
    city.Value = "London";
    number.Value = 4;
    number.Value++;
    Debug.Assert(number.Value == 5);
    Debug.Assert(number == 5);
}
Debug.Assert(number == 3);
Debug.Assert(city == "New York");

```

In addition to `Transactional<T>`, I have provided a transactional array called `TransactionalArray<T>`, and transactional versions for all of the collections in `System.Collections.Generic`, such as `TransactionalDictionary<K,T>` and `TransactionalList<T>`. The volatile resource managers (or VRMs, for short) are available with `ServiceModelEx`. Their implementation has nothing to do with WCF, so I chose not to include it in this book. That implementation, however, makes intensive use of some of the more advanced features of C#, `System.Transactions`, and .NET system programming, so you may find it of interest on its own merit.

Per-Call Transactional Services

With a per-call service, once the call returns, the instance is destroyed. Therefore, the resource manager used to store the state between calls must be outside the scope of the instance. Because there could be many instances of the same service type accessing the same resource manager, every operation must contain some parameters that allow the service instance to find its state in the resource manager and bind against it. The best approach is to have each operation contain some key as a parameter identifying the state. I call that parameter the *state identifier*. The client must provide the state identifier with every call to the per-call service. Typical state identifiers are account numbers, order numbers, and so on. For example, the client creates a new transactional order-processing object, and on every method call the client must provide the order number as a parameter, in addition to other parameters.

Example 7-15 shows a template for implementing a transactional per-call service.

Example 7-15. Implementing a transactional service

```

[DataContract]
struct Param
{...}

[ServiceContract]
interface IMyContract
{
    [OperationContract]
    [TransactionFlow(...)]
    void MyMethod(Param stateIdentifier);

```

```

}
[ServiceBehavior(InstanceContextMode = InstanceContextMode.PerCall)]
class MyService : IMyContract, IDisposable
{
    [OperationBehavior(TransactionScopeRequired = true)]
    public void MyMethod(Param stateIdentifier)
    {
        GetState(stateIdentifier);
        DoWork();
        SaveState(stateIdentifier);
    }
    void GetState(Param stateIdentifier)
    {...}
    void DoWork()
    {...}
    void SaveState(Param stateIdentifier)
    {...}
    public void Dispose()
    {...}
}

```

The `MyMethod()` signature contains a state identifier parameter of the type `Param` (a pseudotype invented for this example), used to get the state from a resource manager with the `GetState()` helper method. The service instance then performs its work using the `DoWork()` helper method and saves its state back to the resource manager using the `SaveState()` method, specifying its identifier.

Note that not all of the service instance's state can be saved by value to the resource manager. If the state contains references to other objects, `GetState()` should create those objects, and `SaveState()` (or `Dispose()`) should dispose of them.

Because the service instance goes through the trouble of retrieving its state and saving it on every method call, transactional programming is natural for per-call services. The behavioral requirements for a state-aware transactional object and a per-call object are the same: both retrieve and save their state at the method boundaries. Compare [Example 7-15](#) with [Example 4-3](#). The only difference is that the state store used by the service in [Example 7-15](#) should be transactional.

As far as a per-call service call is concerned, transactional programming is almost incidental. Every call on the service gets a new instance, and that call may or may not be in the same transaction as the previous call (see [Figure 7-8](#)).

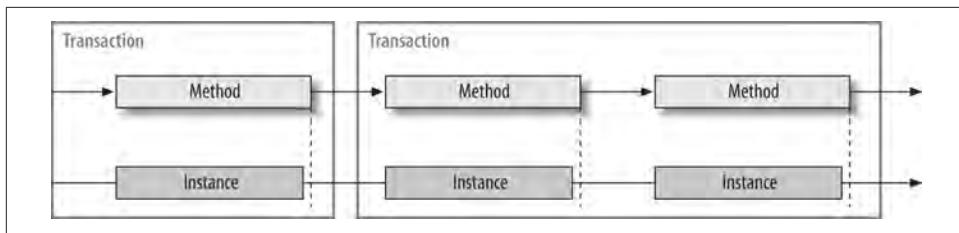


Figure 7-8. Per-call service and transactions

Regardless of transactions, in every call the service gets its state from a resource manager and then saves it back, so the methods are always guaranteed to operate either on consistent state from the previous transaction or on the temporary yet well-isolated state of the current transaction in progress. A per-call service must vote and complete its transaction in every method call. In fact, a per-call service must always use auto-completion (i.e., have `TransactionAutoComplete` set to its default value, `true`).

From the client's perspective, the same service proxy can participate in multiple transactions or in the same transaction. For example, in the following code snippet, every call will be in a different transaction:

```
MyContractClient proxy = new MyContractClient();

using(TransactionScope scope = new TransactionScope())
{
    proxy.MyMethod(...);
    scope.Complete();
}

using(TransactionScope scope = new TransactionScope())
{
    proxy.MyMethod(...);
    scope.Complete();
}

proxy.Close();
```

Or, the client can use the same proxy multiple times in the same transaction, and even close the proxy independently of any transactions:

```
MyContractClient proxy = new MyContractClient();
using(TransactionScope scope = new TransactionScope())
{
    proxy.MyMethod(...);
    proxy.MyMethod(...);
    scope.Complete();
}
proxy.Close();
```



The call to `Dispose()` on a per-call service has no ambient transaction.

With a per-call service, any resource manager can be used to store the service state. For example, you might use a database, or you might use volatile resource managers accessed as static member variables, as shown in [Example 7-16](#).

Example 7-16. Per-call service using a VRM

```
[ServiceContract]
interface ICounterManager
{
    [OperationContract]
    [TransactionFlow(...)]
    void Increment(string stateIdentifier);
}

[ServiceBehavior(InstanceContextMode = InstanceContextMode.PerCall)]
class MyService : ICounterManager
{
    static TransactionalDictionary<string,int> m_StateStore =
        new TransactionalDictionary<string,int>();

    [OperationBehavior(TransactionScopeRequired = true)]
    public void Increment(string stateIdentifier)
    {
        if(m_StateStore.ContainsKey(stateIdentifier) == false)
        {
            m_StateStore[stateIdentifier] = 0;
        }
        m_StateStore[stateIdentifier]++;
    }
}
```

The transaction lifecycle

When the per-call service is the root of a transaction (that is, when it is configured for the Client/Service transaction mode and there is no client transaction, or when it is configured for the Service transaction mode), the transaction ends once the service instance is deactivated. WCF completes and ends the transaction as soon as the method returns, even before `Dispose()` is called. When the client is the root of the transaction (or whenever the client's transaction flows to the service and the service joins it), the transaction ends when the client's transaction ends.

Per-Session Transactional Services

While it is possible to develop transactional sessionful services with great ease using my volatile resource managers, WCF was designed without them in mind, simply because these technologies evolved more or less concurrently. Consequently, the WCF architects did not trust developers to properly manage the state of their sessionful service in the face of transactions—something that is rather cumbersome and difficult, as you will see, if all you have at your disposal is raw .NET and WCF. The WCF architects made the extremely conservative decision to treat a sessionful transactional service as a per-call service by default in order to enforce a proper state-aware programming model. In fact, the default transaction configuration of WCF will turn any service, regardless of its instancing mode, into a per-call service. This, of course, negates the very need for a per-session service in the first place. That said, WCF does allow you to maintain the session semantic with a transactional service, using several distinct programming models. A per-session transactional service instance can be accessed by multiple transactions, or the instance can establish an affinity to a particular transaction, in which case, until it completes, only that transaction is allowed to access it. However, as you will see, unless you use volatile resource managers this support harbors a disproportional cost in programming model complexity and constraints.

Releasing the service instance

The lifecycle of any transactional service is controlled by the `ServiceBehavior` attribute's Boolean property, `ReleaseServiceInstanceOnTransactionComplete`:

```
[AttributeUsage(AttributeTargets.Class)]
public sealed class ServiceBehaviorAttribute : Attribute, ...
{
    public bool ReleaseServiceInstanceOnTransactionComplete
    {get;set;}
    //More members
}
```

When `ReleaseServiceInstanceOnTransactionComplete` is set to `true` (the default value), it disposes of the service instance once the instance completes the transaction. WCF uses context deactivation (discussed in [Chapter 4](#)) to terminate the sessionful service instance and its in-memory state, while maintaining the transport session and the instance context.

Note that the release takes place once the instance completes the transaction, not necessarily when the transaction really completes (which could be much later). When `ReleaseServiceInstanceOnTransactionComplete` is `true`, the instance has two ways of completing the transaction and being released: at the method boundary if the method has `TransactionAutoComplete` set to `true`, or when any method that has `TransactionAutoComplete` set to `false` calls `SetTransactionComplete()`.

`ReleaseServiceInstanceOnTransactionComplete` has two interesting interactions with other service and operation behavior properties. First, it cannot be set (to either true or false) unless at least one operation on the service has `TransactionScopeRequired` set to true. This is validated at the service load time by the `set` accessor of the `ReleaseServiceInstanceOnTransactionComplete` property.

For example, this is a valid configuration:

```
[ServiceBehavior(ReleaseServiceInstanceOnTransactionComplete = true)]
class MyService : IMyContract
{
    [OperationBehavior(TransactionScopeRequired = true)]
    public void MyMethod()
    {...}

    [OperationBehavior(...)]
    public void MyOtherMethod()
    {...}
}
```

What this constraint means is that although the default of `ReleaseServiceInstanceOnTransactionComplete` is true, the following two definitions are not semantically equivalent, because the second one will throw an exception at the service load time:

```
class MyService : IMyContract
{
    public void MyMethod()
    {...}
}

//Invalid definition:
[ServiceBehavior(ReleaseServiceInstanceOnTransactionComplete = true)]
class MyService : IMyContract
{
    public void MyMethod()
    {...}
}
```

The second constraint involved in using `ReleaseService InstanceOnTransactionComplete` relates to concurrent multithreaded access to the service instance.

Concurrency management is the subject of the next chapter. For now, all you need to know is that the `ConcurrencyMode` property of the `ServiceBehavior` attribute controls concurrent access to the service instance:

```
public enum ConcurrencyMode
{
    Single,
    Reentrant,
    Multiple
}
```

```
[AttributeUsage(AttributeTargets.Class)]
public sealed class ServiceBehaviorAttribute : ...
{
    public ConcurrencyMode ConcurrencyMode
    {get;set;}
    //More members
}
```

The default value of `ConcurrencyMode` is `ConcurrencyMode.Single`.

At the service load time, WCF will verify that, if `TransactionScopeRequired` is set to `true` for at least one operation on the service when `ReleaseServiceInstanceOnTransactionComplete` is `true` (by default or explicitly), the service concurrency mode is `ConcurrencyMode.Single`.

For example, given this contract:

```
[ServiceContract]
interface IMyContract
{
    [OperationContract]
    [TransactionFlow(...)]
    void MyMethod();

    [OperationContract]
    [TransactionFlow(...)]
    void MyOtherMethod();
}
```

the following two definitions are equivalent and valid:

```
class MyService : IMyContract
{
    [OperationBehavior(TransactionScopeRequired = true)]
    public void MyMethod()
    {...}

    public void MyOtherMethod()
    {...}
}

[ServiceBehaviorConcurrencyMode = ConcurrencyMode.Single,
 ReleaseServiceInstanceOnTransactionComplete = true]
class MyService : IMyContract
{
    [OperationBehavior(TransactionScopeRequired = true)]
    public void MyMethod()
    {...}

    public void MyOtherMethod()
    {...}
}
```

The following definition is also valid, since no method requires a transaction scope even though `ReleaseServiceInstanceOnTransactionComplete` is `true`:

```
[ServiceBehaviorConcurrencyMode = ConcurrencyMode.Multiple)]
class MyService : IMyContract
{
    public void MyMethod()
    {...}

    public void MyOtherMethod()
    {...}
}
```

In contrast, the following definition is invalid, because at least one method requires a transaction scope, `ReleaseServiceInstanceOnTransactionComplete` is `true`, and yet the concurrency mode is not `ConcurrencyMode.Single`:

```
//Invalid configuration:
[ServiceBehaviorConcurrencyMode = ConcurrencyMode.Multiple]
class MyService : IMyContract
{
    [OperationBehavior(TransactionScopeRequired = true)]
    public void MyMethod()
    {...}

    public void MyOtherMethod()
    {...}
}
```



The concurrency constraint applies to all instancing modes.

The `ReleaseServiceInstanceOnTransactionComplete` property can enable a transactional sessionful interaction between the client and the service. With its default value of `true`, once the service instance completes the transaction (either declaratively or explicitly), the return of the method will deactivate the service instance as if it were a per-call service.

For example, the service in [Example 7-17](#) behaves just like a per-call service.

Example 7-17. Per-session yet per-call transactional service

```
[ServiceContractSessionMode = SessionMode.Required]
interface IMyContract
{
    [OperationContract]
    [TransactionFlow(...)]
```

```

        void MyMethod();
    }
    class MyService : IMyContract
    {
        [OperationContract(TransactionScopeRequired = true)]
        public void MyMethod()
        {...}
    }

```

Every time the client calls `MyMethod()`, the client will get a new service instance. The new client call may come in on a new transaction as well, and the service instance has no affinity to any transaction. The relationship between the service instances and the transactions is just as in [Figure 7-8](#). The service needs to proactively manage its state just as it did in [Example 7-15](#), as demonstrated in [Example 7-18](#).

Example 7-18. Proactive state management by default with a per-session transactional service

```

[DataContract]
struct Param
{...}

[ServiceContract(SessionMode = SessionMode.Required)]
interface IMyContract
{
    [OperationContract]
    [TransactionFlow(...)]
    void MyMethod(Param stateIdentifier);
}
class MyService : IMyContract
{
    [OperationContract(TransactionScopeRequired = true)]
    public void MyMethod(Param stateIdentifier)
    {
        GetState(stateIdentifier);
        DoWork();
        SaveState(stateIdentifier);
    }
    void GetState(Param stateIdentifier)
    {...}
    void DoWork()
    {...}
    void SaveState(Param stateIdentifier)
    {...}
}

```

The transactional per-session service can also, of course, use VRMs, as was done in [Example 7-16](#).

Disabling releasing the service instance

Obviously, a configuration such as that in [Example 7-17](#) or [Example 7-18](#) adds no value to configuring the service as sessionful. The client must still pass a state identifier, and the service is de facto a per-call service. To behave as a per-session service, the service can set `ReleaseServiceInstanceOnTransactionComplete` to `false`, as in [Example 7-19](#).

Example 7-19. Per-session transactional service

```
[ServiceContract(SessionMode = SessionMode.Required)]
interface IMyContract
{
    [OperationContract]
    [TransactionFlow(...)]
    void MyMethod();
}

[ServiceBehavior(ReleaseServiceInstanceOnTransactionComplete = false)]
class MyService : IMyContract
{
    [OperationBehavior(TransactionScopeRequired = true)]
    public void MyMethod()
    {...}
}
```

When `ReleaseServiceInstanceOnTransactionComplete` is `false`, the instance will not be disposed of when transactions complete, as shown in [Figure 7-9](#).

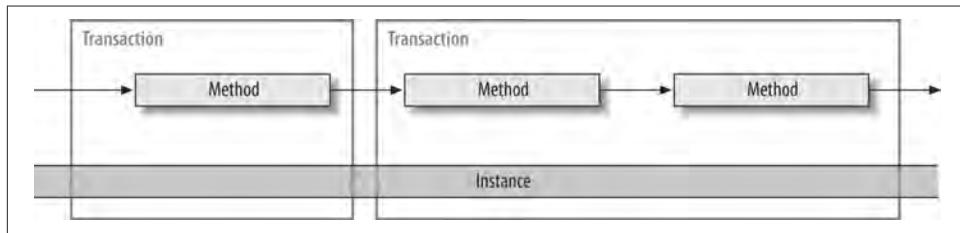


Figure 7-9. Sessionful transactional instance and transactions

The interaction in [Figure 7-9](#) might, for example, be the result of the following client code, where all calls went to the same service instance:

```
MyContractClient proxy = new MyContractClient();
using(TransactionScope scope = new TransactionScope())
{
    proxy.MyMethod();
    scope.Complete();
}

using(TransactionScope scope = new TransactionScope())
{
```

```

        proxy.MyMethod();
        proxy.MyMethod();
        scope.Complete();
    }
    proxy.Close();
}

```

State-aware per-session services

When `ReleaseServiceInstanceOnTransactionComplete` is `false`, WCF will stay out of the way and will let the service developer worry about managing the state of the service instance. Obviously, you have to somehow monitor transactions and roll back any changes made to the state of the instance if a transaction aborts. The per-session service still must equate method boundaries with transaction boundaries, because every method may be in a different transaction, and a transaction may end between method calls in the same session. There are two possible programming models. The first is to be state-aware, but use the session ID as a state identifier. With this model, at the beginning of every method the service gets its state from a resource manager using the session ID as a key, and at the end of every method the service instance saves the state back to the resource manager, as shown in [Example 7-20](#).

Example 7-20. State-aware, transactional per-session service

```

[ServiceBehavior(ReleaseServiceInstanceOnTransactionComplete = false)]
class MyService : IMyContract, IDisposable
{
    readonly string m_StateIdentifier;

    public MyService()
    {
        InitializeState();
        m_StateIdentifier = OperationContext.Current.SessionId;
        SaveState();
    }

    [OperationBehavior(TransactionScopeRequired = true)]
    public void MyMethod()
    {
        GetState();
        DoWork();
        SaveState();
    }

    public void Dispose()
    {
        RemoveState();
    }

    //Helper methods

    void InitializeState()
    {...}
}

```

```

void GetState()
{
    //Use m_StateIdentifier to get state
    ...
}
void DoWork()
{...}
void SaveState()
{
    //Use m_StateIdentifier to save state
    ...
}
void RemoveState()
{
    //Use m_StateIdentifier to remove the state from the RM
    ...
}
}

```

In [Example 7-20](#), the constructor first initializes the state of the object and then saves the state to a resource manager, so that any method can retrieve it. Note that the per-session object maintains the illusion of a stateful, sessionful interaction with its client. The client does not need to pass an explicit state identifier, but the service must be disciplined and retrieve and save the state in every operation call. When the session ends, the service purges its state from the resource manager in the `Dispose()` method.

Stateful per-session services

The second, more modern programming model is to use volatile resource managers for the service members, as shown in [Example 7-21](#).

Example 7-21. Using volatile resource managers to achieve a stateful per-session transactional service

```

[ServiceBehavior(ReleaseServiceInstanceOnTransactionComplete = false)]
class MyService : IMyContract
{
    Transactional<string> m_Text = new Transactional<string>("Some initial value");

    TransactionalArray<int> m_Numbers = new TransactionalArray<int>(3);

    [OperationBehavior(TransactionScopeRequired = true)]
    public void MyMethod()
    {
        m_Text.Value = "This value will roll back if the transaction aborts";

        //These will roll back if the transaction aborts
        m_Numbers[0] = 11;
        m_Numbers[1] = 22;
    }
}

```

```
    m_Numbers[2] = 33;
}
}
```

Example 7-21 uses my `Transactional<T>` and `TransactionalArray<T>` volatile resource managers. The per-session service can safely set `ReleaseServiceInstanceOnTransactionComplete` to `false` and yet freely access its members. The use of the volatile resource managers enables a stateful programming model, and the service instance simply accesses its state as if no transactions were involved. The volatile resource managers autoenlist in the transaction and isolate that transaction from all other transactions. Any changes made to the state will commit or roll back with the transaction.

Transaction lifecycle

When the per-session service is the root of the transaction, the transaction ends once the service completes the transaction, which is when the method returns. When the client is the root of the transaction (or when a transaction flows to the service), the transaction ends when the client's transaction ends. If the per-session service provides an `IDisposable` implementation, the `Dispose()` method will not have any transaction, regardless of the root.

Concurrent transactions

Because a per-session service can engage the same service instance in multiple client calls, it can also sustain multiple concurrent transactions. Given the service definition of **Example 7-19**, **Example 7-22** shows some client code that launches concurrent transactions on the same instance. `scope2` will use a new transaction separate from that of `scope1`, and yet access the same service instance in the same session.

Example 7-22. Launching concurrent transactions

```
using(TransactionScope scope1 = new TransactionScope())
{
    MyContractClient proxy = new MyContractClient();
    proxy.MyMethod();

    using(TransactionScope scope2 =
        new TransactionScope(TransactionScopeOption.RequiresNew))
    {
        proxy.MyMethod();
        scope2.Complete();
    }
    proxy.MyMethod();

    proxy.Close();
```

```
    scope1.Complete();  
}
```

The resulting transactions of [Example 7-22](#) are depicted in [Figure 7-10](#).

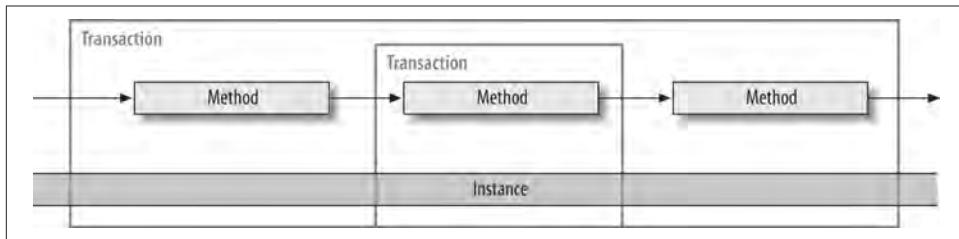


Figure 7-10. Concurrent transactions



Code such as that in [Example 7-22](#) will almost certainly result in a transactional deadlock over the underlying resources the service accesses. The first transaction will obtain the resource lock, and the second transaction will wait to own that lock while the first transaction waits for the second to complete.

Completing on session end

WCF offers yet another programming model for transactional per-session services, which is completely independent of `ReleaseServiceInstanceOnTransactionComplete`. This model is available for the case when the entire session fits into a single transaction, and the service equates session boundaries with transaction boundaries. The idea is that the service should not complete the transaction inside the session, because that is what causes WCF to release the service instance. To avoid completing the transaction, a per-session service can set `TransactionAutoComplete` to `false`, as shown in [Example 7-23](#).

Example 7-23. Setting `TransactionAutoComplete` to `false`

```
[ServiceContract(SessionMode = SessionMode.Required)]  
interface IMyContract  
{  
    [OperationContract]  
    [TransactionFlow(...)]  
    void MyMethod1();  
  
    [OperationContract]  
    [TransactionFlow(...)]  
    void MyMethod2();  
  
    [OperationContract]  
    [TransactionFlow(...)]
```

```

        void MyMethod3();
    }
    class MyService : IMyContract
    {
        [OperationBehavior(TransactionScopeRequired = true,
                           TransactionAutoComplete = false)]
        public void MyMethod1()
        {...}

        [OperationBehavior(TransactionScopeRequired = true,
                           TransactionAutoComplete = false)]
        public void MyMethod2()
        {...}

        [OperationBehavior(TransactionScopeRequired = true,
                           TransactionAutoComplete = false)]
        public void MyMethod3()
        {...}
    }

```

Note that only a per-session service with a contract set to `SessionMode.Required` can set `TransactionAutoComplete` to `false`, and that is verified at the service load time. The problem with [Example 7-23](#) is that the transaction the service participates in will always abort because the service does not vote to commit it by completing it. If the service equates sessions with transactions, the service should vote once the session ends. For that purpose, the `ServiceBehavior` attribute provides the Boolean property `TransactionAutoCompleteOnSessionClose`, defined as:

```

[AttributeUsage(AttributeTargets.Class)]
public sealed class ServiceBehaviorAttribute : Attribute, ...
{
    public bool TransactionAutoCompleteOnSessionClose
    {get;set;}
    //More members
}

```

The default of `TransactionAutoCompleteOnSessionClose` is `false`. However, when set to `true`, it will autocomplete all uncompleted methods in the session. If no exceptions occurred during the session, when `TransactionAutoCompleteOnSessionClose` is `true` the service will vote to commit. For example, here is how to retrofit [Example 7-23](#):

```

[ServiceBehavior(TransactionAutoCompleteOnSessionClose = true)]
class MyService : IMyContract
{...}

```

[Figure 7-11](#) shows the resulting instance and its session.

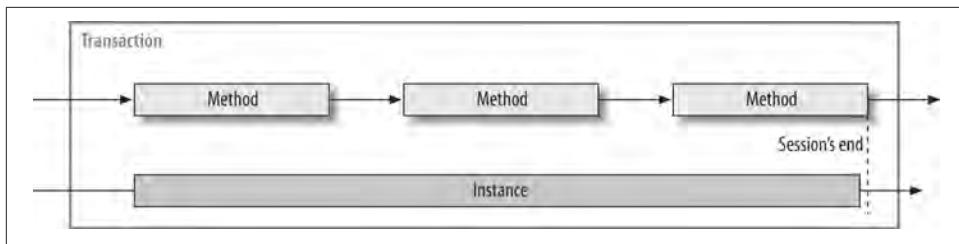


Figure 7-11. Setting `TransactionAutoCompleteOnSessionClose` to true

During the session, the instance can maintain and access its state in normal member variables, and there is no need for state awareness or volatile resource managers.



When joining the client's transaction and relying on autocompletion on session close, the service must avoid lengthy processing in `Dispose()` or, in practical terms, avoid implementing `IDisposable` altogether. The reason is the race condition described here. Recall from [Chapter 4](#) that `Dispose()` is called asynchronously at the end of the session. Autocompletion at session end takes place once the instance has been disposed of. If the client has control before the instance is disposed, the transaction will abort because the service has not yet completed it.

Note that using `TransactionAutoCompleteOnSessionClose` is risky, because it is always subjected to the transaction timeout. Sessions are by their very nature long-lived entities, while well-designed transactions are short-lived. This programming model is available for the case when the vote decision requires information that will be obtained by future calls throughout the session.

Because having `TransactionAutoCompleteOnSessionClose` set to `true` equates the session's end with the transaction's end, it is required that when the client's transaction is used, the client terminates the session within that transaction:

```
using(TransactionScope scope = new TransactionScope())
{
    MyContractClient proxy = new MyContractClient();
    proxy.MyMethod();
    proxy.MyMethod();
    proxy.Close();

    scope.Complete();
}
```

Failing to do so will abort the transaction. As a side effect, the client cannot easily stack the `using` statements of the transaction scope and the proxy, because that may cause the proxy to be disposed of after the transaction:

```

//This always aborts:using(MyContractClient proxy = new MyContractClient())
using(TransactionScope scope = new TransactionScope())
{
    proxy.MyMethod();
    proxy.MyMethod();

    scope.Complete();
}

```

In addition, because the proxy is basically good for only one-time use, there is little point in storing the proxy in member variables.

Transactional affinity

Setting `TransactionAutoComplete` to `false` has a unique effect that nothing else in WCF provides: it creates an affinity between the service instance context and the transaction, so that only that single transaction can ever access that service instance context. Unless context deactivation is used, this affinity is therefore to the instance as well. The affinity is established once the first transaction accesses the service instance, and once established it is fixed for the life of the instance (until the session ends). Transactional affinity is available only for per-session services, because only a per-session service can set `TransactionAutoComplete` to `false`. Affinity is crucial because the service is not state-aware—it uses normal members, and it must isolate access to them from any other transaction, in case the transaction to which it has an affinity aborts. Affinity thus offers a crude form of transaction-based locking. With transaction affinity, code such as that in [Example 7-22](#) is guaranteed to deadlock (and eventually abort due to timing out) because the second transaction is blocked (independently of any resources the service accesses) waiting for the first transaction to finish, while the first transaction is blocked waiting for the second.

Hybrid state management

WCF also supports a hybrid of two of the sessionful programming models discussed earlier, combining both a state-aware and a regular sessionful transactional per-session service. The hybrid mode is designed to allow the service instance to maintain state in memory until it can complete the transaction, and then recycle that state using `ReleaseServiceInstanceOnTransactionComplete` as soon as possible, instead of delaying completing the transaction until the end of the session. Consider the service in [Example 7-24](#), which implements the contract from [Example 7-23](#).

Example 7-24. Hybrid per-session service

```

[ServiceBehavior(TransactionAutoCompleteOnSessionClose = true)]
class MyService : IMyContract
{
    [OperationBehavior(TransactionScopeRequired = true,
                      TransactionAutoComplete = false)]

```

```

public void MyMethod1()
{
}
[OperationBehavior(TransactionScopeRequired = true,
    TransactionAutoComplete = false)]
public void MyMethod2()
{
}
[OperationBehavior(TransactionScopeRequired = true)]
public void MyMethod3()
{
}
}

```

The service uses the default of `ReleaseServiceInstanceOnTransactionComplete` (true), yet it has two methods (`MyMethod1()` and `MyMethod2()`) that do not complete the transaction with `TransactionAutoComplete` set to false, resulting in an affinity to a particular transaction. The affinity isolates its members from any other transaction. The problem now is that the transaction will always abort, because the service does not complete it. To compensate for that, the service offers `MyMethod3()`, which does complete the transaction. Because the service uses the default of `ReleaseServiceInstanceOnTransactionComplete` (true), after `MyMethod3()` is called, the transaction is completed and the instance is disposed of, as shown in [Figure 7-12](#). Note that `MyMethod3()` could have instead used explicit voting via `SetTransactionComplete()`. The important thing is that it completes the transaction. If the client does not call `MyMethod3()`, purely as a contingency, the service in [Example 7-24](#) relies on `TransactionAutoCompleteOnSessionClose` being set to true to complete and commit the transaction.

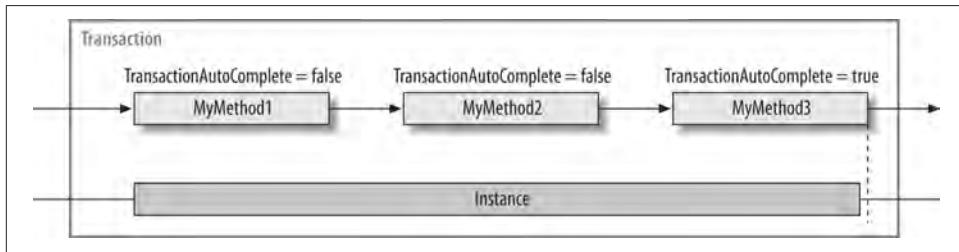


Figure 7-12. Hybrid state management

The hybrid mode is inherently a brittle proposition. The first problem is that the service instance must complete the transaction before it times out, but since there is no telling when the client will call the completing method, you risk timing out before that. In addition, the service holds onto any locks on resource managers it may access for the duration of the session, and the longer the locks are held, the higher the likelihood is of other transactions timing out or deadlocking with this service's transaction. Finally, the service is at the mercy of the client, because the client must call the completing method to end the session. You can and should clearly document the need to call that operation at the end of the transaction:

```

[ServiceContract(SessionMode = SessionMode.Required)]
interface IMyContract
{
    [OperationContract]
    [TransactionFlow(...)]
    void MyMethod1();

    [OperationContract]
    [TransactionFlow(...)]
    void MyMethod2();

    [OperationContract]
    [TransactionFlow(...)]
    void CompleteTransaction();
}

```

Both equating sessions with transactions (while relying solely on `TransactionAutoCompleteOnSessionClose`) and using the hybrid mode are potential solutions for situations when the transaction execution and subsequent voting decision require information obtained throughout the session. Consider, for example, the following contract used for order processing:

```

[ServiceContract(SessionMode = SessionMode.Required)]
interface IOrderManager
{
    [OperationContract]
    [TransactionFlow(...)]
    void SetCustomerId(int customerId);

    [OperationContract]
    [TransactionFlow(...)]
    void AddItem(int itemId);    [OperationContract]
    [TransactionFlow(...)]
    bool ProcessOrders();
}

```

The implementing service can only process the order once it has the customer ID and all of the ordered items. However, relying on transactional sessions in this way usually indicates poor design, because of the inferior throughput and scalability implications. Good transactions are inherently short while sessions are inherently long (up to 10 minutes by default), so they are inherently incompatible. The disproportional complexity of prolonging a single transaction across a session outweighs the perceived benefit of using a session. It is usually better to factor the contract so that it provides every operation with all the information it needs to complete and vote:

```

[ServiceContract(SessionMode = ...)]
interface IOrderManager
{
    [OperationContract]
    [TransactionFlow(...)]

```

```
        bool ProcessOrders(int customerId,int[] itemIds);
    }
```

Done this way, you can either implement the service as per-call or maintain a sessionful programming model, avoid placing operation order constraints on the client, and use any VRMs as member variables and access other transactional resources. You clearly separate the contract from its implementation, both on the client and the service side.

Transactional Durable Services

Recall from [Chapter 4](#) that a durable service retrieves its state from the configured store and then saves its state back into that store on every operation. The state store may or may not be a transactional resource manager. If the service is transactional, it should of course use only a transactional durable storage and enlist it in each operation's transaction. That way, if a transaction aborts, the state store will be rolled back to its pre-transaction state. However, WCF does not know whether a service is designed to propagate its transactions to the state store, and by default it will not enlist the storage in the transaction even if the storage is a transactional resource manager, such as SQL Server. To instruct WCF to propagate the transaction and enlist the underlying storage, set the `SaveStateInOperationTransaction` property of the `DurableService` attribute to `true`:

```
public sealed class DurableServiceAttribute : ...
{
    public bool SaveStateInOperationTransaction
    {get;set;}
}
```

`SaveStateInOperationTransaction` defaults to `false`, which means the state storage will not participate in the transaction. It is therefore important to always set `SaveStateInOperationTransaction` to `true` to ensure consistent state management in the presence of transactions. Since only a transactional service could benefit from having `SaveStateInOperationTransaction` set to `true`, if it is `true` then WCF will insist that all operations on the service either have `TransactionScopeRequired` set to `true` or have mandatory transaction flow. If the operation is configured with `TransactionScopeRequired` set to `true`, the ambient transaction of the operation will be the one used to enlist the storage. If the operation is configured for mandatory transaction flow, the client's transaction will be used to enlist the storage (regardless of whether the operation does or does not have an ambient transaction).

Instance ID management

As explained in [Chapter 4](#), the `DurableService` behavior attribute enforces strict management of the instance ID passed over the context binding. The first operation to start the workflow will have no instance ID, in which case, WCF will create a new

instance ID, use it to save the newly created instance state to the storage, and then send the instance ID back to the client. From that point on, until the end of the workflow, the client must pass the same instance ID to the service. If the client provides an instance ID that is not present in the storage, WCF will throw an exception. This presents a potential pitfall with transactional durable services: suppose the client starts a workflow and propagates its transaction to the service. The first operation creates the instance ID, executes successfully, and stores the state in the storage. However, what would happen if the transaction were then to abort, due to some other party (such as the client or another service involved in the transaction) voting to abort? The state storage would roll back the changes made to it, including the newly created instance state and the corresponding ID. The next call coming from the client will present the same ID created by the first call, except now the state storage will not have any record of that ID, so WCF will reject the call, throw an exception, and prevent any other call to the service with that ID from ever executing.

To avoid this pitfall, you need to add to the service contract an explicit first operation whose sole purpose is to guarantee that the first call successfully commits the instance ID to the state storage. For example, in the case of a calculator service, this would be your `PowerOn()` operation. You should explicitly block the client's transaction (by using the default value of `TransactionFlowOption.NotAllowed`), and avoid placing any code in that method body, thus precluding anything that could go wrong from aborting the transaction. You can enforce having the client call the initiating operation first using demarcating operations (discussed in [Chapter 4](#)).

A similar pitfall exists at the end of the workflow. By setting the `CompletesInstance` property of the `DurableOperation` attribute to `true`, you indicate to WCF that the workflow has ended and that WCF should purge the instance state from the storage. However, if the client's transaction aborts after the last operation in the service has executed successfully, the storage will roll back and keep the orphaned state indefinitely. To avoid bloating the state storage with zombie instances (the product of aborted transactions of the completing instance operations), you need to add to the service contract an explicit operation whose sole purpose is to complete the instance and to commit successfully, irrespective of whether the client's transaction commits. For example, in the case of a calculator service, this would be your `PowerOff()` operation. Again, block any client transaction from propagating to the service, and avoid placing any code in the completing method.

[Example 7-25](#) shows a template for defining and implementing a transactional durable service, adhering to these guidelines.

Example 7-25. Transactional durable service

```
[ServiceContract]  
interface IMyContract
```

```

{
    [OperationContract]
    void SaveState();

    [OperationContract(IsInitiating = false)]
    void ClearState();

    [OperationContract(IsInitiating = false)]
    [TransactionFlow(...)]
    void MyMethod1();

    [OperationContract(IsInitiating = false)]
    [TransactionFlow(...)]
    void MyMethod2();
}

[Serializable]
[DurableService(SaveStateInOperationTransaction = true)]
class MyService: IMyContract
{
    [OperationBehavior(TransactionScopeRequired = true)]
    public void SaveState()
    {}

    [DurableOperation(CompletesInstance = true)]
    [OperationBehavior(TransactionScopeRequired = true)]
    public void ClearState()
    {}

    [OperationBehavior(TransactionScopeRequired = true)]
    public void MyMethod1()
    {...}
    [OperationBehavior(TransactionScopeRequired = true)]
    public void MyMethod2()
    {...}
}

```

Transactional Behavior

As far as the `DurableService` attribute is concerned, the word *durable* in its name is a misnomer, since it does not necessarily indicate a durable behavior. All it means is that WCF will automatically deserialize the service state from a configured storage and then serialize it back again on every operation. Similarly, the `persistence provider` behavior (see [Chapter 4](#)) does not necessarily mean persistence, since any provider that derives from the prescribed abstract provider class will comply with WCF's expectation of the behavior.³

³ I presented my approach for transactional behavior in the January 2009 issue of *MSDN Magazine*.

The fact that the WCF durable service infrastructure is, in reality, a serialization infrastructure enabled me to leverage it into yet another technique for managing service state in the face of transactions, while relying underneath on a volatile resource manager, without having the service instance do anything about it. This further streamlines the transactional programming model of WCF and yields the benefit of the superior programming model of transactions for mere objects.

The first step was to define two transactional in-memory provider factories:

```
public abstract class MemoryProviderFactory : PersistenceProviderFactory
{...}

public class TransactionalMemoryProviderFactory : MemoryProviderFactory
{...}
public class TransactionalInstanceProviderFactory : MemoryProviderFactory
{...}
```

The `TransactionalMemoryProviderFactory` uses my `TransactionalDictionary<ID,T>` to store the service instances.



Unrelated to this section and transactions, you can configure the service to use the `TransactionalMemoryProviderFactory` with or without transactions by simply listing it in the `persistence providers` section of the service behaviors:

```
<behavior name = "TransactionalMemory">
    <persistenceProvider
        type = "ServiceModelEx.
            TransactionalMemoryProviderFactory,
            ServiceModelEx"
    />
</behavior>
```

This will enable you to store the instances in memory, instead of in a file or SQL Server database. This is useful for quick testing and for stress testing, since it avoids the inherent I/O latency of a durable persistent storage.

The in-memory dictionary is shared among all clients and transport sessions, and as long as the host is running, `TransactionalMemoryProviderFactory` allows clients to connect and disconnect from the service. When using `TransactionalMemoryProviderFactory` you should designate a completing operation that removes the instance state from the store as discussed in [Chapter 4](#), using the `CompletesInstance` property of the `DurableOperation` attribute.

`TransactionalInstanceProviderFactory`, on the other hand, matches each transport session with a dedicated instance of `Transactional<T>`. There is no need to call

any completing operation since the service state will be cleaned up with garbage collection after the session is closed.

Next, I defined the `TransactionalBehaviorAttribute`, shown in [Example 7-26](#).

Example 7-26. The TransactionalBehavior attribute

```
[AttributeUsage(AttributeTargets.Class)]
public class TransactionalBehaviorAttribute : Attribute,IServiceBehavior
{
    public bool TransactionRequiredAllOperations
    {get;set;}

    public bool AutoCompleteInstance
    {get;set;}

    public TransactionalBehaviorAttribute()
    {
        TransactionRequiredAllOperations = true;
        AutoCompleteInstance = true;
    }
    void IServiceBehavior.Validate(ServiceDescription description,
                                    ServiceHostBase host)
    {
        DurableServiceAttribute durable = new DurableServiceAttribute();
        durable.SaveStateInOperationTransaction = true;
        description.Behaviors.Add(durable);

        PersistenceProviderFactory factory;
        if(AutoCompleteInstance)
        {
            factory = new TransactionalInstanceProviderFactory();
        }
        else
        {
            factory = new TransactionalMemoryProviderFactory();
        }

        PersistenceProviderBehavior persistenceBehavior =
            new PersistenceProviderBehavior(factory);
        description.Behaviors.Add(persistenceBehavior);

        if(TransactionRequiredAllOperations)
        {
            foreach(ServiceEndpoint endpoint in description.Endpoints)
            {
                foreach(OperationDescription operation in endpoint.Contract.Operations)
                {
                    OperationBehaviorAttribute operationBehavior =
                        operation.Behaviors.Find<OperationBehaviorAttribute>();
                    operationBehavior.TransactionScopeRequired = true;
                }
            }
        }
    }
}
```

```

        }
    }
}
void IServiceBehavior.AddBindingParameters(...)
{}
void IServiceBehavior.ApplyDispatchBehavior(...)
{}
}

```

`TransactionalBehavior` is a service behavior attribute. It always performs these configurations for the service. First, it injects into the service description a `DurableService` attribute with `SaveStateInOperationTransaction` set to `true`. Second, it adds the use of either `TransactionalMemoryProviderFactory` or `TransactionalInstanceProviderFactory` for the persistent behavior according to the value of the `AutoCompleteInstance` property. If `AutoCompleteInstance` is set to `true` (the default) then `TransactionalBehavior` will use `TransactionalInstanceProviderFactory`. Finally, `TransactionalBehavior` provides the `TransactionRequiredAllOperations` property. When the property is set to `true` (the default) `TransactionalBehavior` will set `TransactionScopeRequired` to `true` on all the service operation behaviors, thus providing all operations with an ambient transaction. When it is explicitly set to `false`, the service developer can choose which operations will be transactional.

As a result, using the attribute like so:

```

[Serializable]
[TransactionalBehavior]
class MyService : IMyContract
{
    public void MyMethod()
    {...}
}

```

is equivalent to this service declaration and configuration:

```

[Serializable]
[DurableService(SaveStateInOperationTransaction = true)]
class MyService : IMyContract
{
    [OperationBehavior(TransactionScopeRequired = true)]
    public void MyMethod()
    {...}
}

<services>
    <service name = "MyService" behaviorConfiguration = "TransactionalBehavior">
        ...
    </service>
</services>
<behaviors>

```

```

<serviceBehaviors>
    <behavior name = "TransactionalBehavior">
        <persistenceProvider
            type = "ServiceModelEx.TransactionalInstanceProviderFactory,
                ServiceModelEx"
        />
    </behavior>
</serviceBehaviors>
</behaviors>

```

When using the `TransactionalBehavior` attribute with the default values, the client need not manage or interact in any way with the instance ID as shown in [Chapter 4](#). All the client needs to do is use the proxy over one of the context bindings, and let the binding manage the instance ID. For example, for this service definition:

```

[ServiceContract]
interface IMyContract
{
    [OperationContract]
    [TransactionFlow(TransactionFlowOption.Allowed)]
    void IncrementCounter();
}

[Serializable]
[TransactionalBehavior]
class MyService : IMyContract
{
    int m_Counter = 0;

    public void IncrementCounter()
    {
        m_Counter++;
        Trace.WriteLine("Counter = " + m_Counter);
    }
}

```

the following client code:

```

MyContractClient proxy = new MyContractClient();

using(TransactionScope scope = new TransactionScope())
{
    proxy.IncrementCounter();
    scope.Complete();
}

//This transaction will abort since the scope is not completed
using(TransactionScope scope = new TransactionScope())
{
    proxy.IncrementCounter();
}

using(TransactionScope scope = new TransactionScope())

```

```

{
    proxy.IncrementCounter();
    scope.Complete();
}

proxy.Close();

```

yields this output:

```

Counter = 1
Counter = 2
Counter = 2

```

Note that the service was interacting with a normal integer as its member variable.

In-proc transactions

The `TransactionalBehavior` attribute substantially simplifies transactional programming and is a fundamental breakthrough in software engineering, where memory itself can be transactional and it is possible for every object to be transactional, as if it is allocated off a transactional heap (for more on my vision for the future of the platform, see [Appendix A](#)). `TransactionalBehavior` maintains the programming model of conventional, plain .NET, yet it provides the full benefits of transactions.

To allow the efficient use of `TransactionalBehavior` even in the most intimate execution scopes, `ServiceModelEx` contains the `NetNamedPipeContextBinding` class. As the binding's name implies, it is the IPC binding plus the context protocol (required by the `DurableService` attribute). [Appendix B](#) walks through implementing the `NetNamedPipeContextBinding` class.



Supporting `TransactionalBehavior` over IPC was my main motivation for developing the `NetNamedPipeContextBinding`.

To make the programming model of `TransactionalBehavior` even more accessible, the `InProcFactory` class from [Chapter 1](#) actually uses `NetNamedPipeContextBinding` instead of the built-in `NetNamedPipeBinding`. `InProcFactory` also flows transactions over the binding. Combined with my `WcfWrapper` class, this enables the programming model of [Example 7-27](#), without ever resorting to host management or client or service config files.

Example 7-27. Combining TransactionalBehavior with the InProcFactory

```

[ServiceContract]
interface IMyContract
{

```

```

[OperationContract]
[TransactionFlow(TransactionFlowOption.Allowed)]
void IncrementCounter();
}

[Serializable][TransactionalBehavior]
class MyService : IMyContract
{
    int m_Counter = 0;

    public void IncrementCounter()
    {
        m_Counter++;
        Trace.WriteLine("Counter = " + m_Counter);
    }
}

class MyClass : WcfWrapper<MyService,IMyContract>,IMyContract
{
    public void IncrementCounter()
    {
        Proxy.IncrementCounter();
    }
}

//Client-code

MyClass proxy = new MyClass();

using(TransactionScope scope = new TransactionScope())
{
    proxy.IncrementCounter();
    scope.Complete();
}

//This transaction will abort since the scope is not completed
using(TransactionScope scope = new TransactionScope())
{
    proxy.IncrementCounter();
}

using(TransactionScope scope = new TransactionScope())
{
    proxy.IncrementCounter();
    scope.Complete();
}

proxy.Close();

//Traces:
Counter = 1

```

```
Counter = 2
Counter = 2
```

Transactional Singleton Service

By default, a transactional singleton behaves like a per-call service. The reason is that by default `ReleaseServiceInstanceOnTransactionComplete` is set to `true`, so after the singleton autocompletes a transaction, WCF disposes of the singleton, in the interest of state management and consistency. This, in turn, implies that the singleton must be state-aware and must proactively manage its state in every method call, in and out of a resource manager. The big difference compared to a per-call service is that WCF will enforce the semantic of the single instance, so at any point there will be at most a single instance running. WCF uses concurrency management and instance deactivation to enforce this rule. Recall that when `ReleaseServiceInstanceOnTransactionComplete` is `true`, the concurrency mode must be `ConcurrencyMode.Single` to disallow concurrent calls. WCF keeps the singleton context and merely deactivates the instance hosted in the context, as discussed in [Chapter 4](#). What this means is that although the singleton needs to be state-aware, it does not need the client to provide an explicit state identifier in every call. The singleton can use any type-level constant to identify its state in the state resource manager, as shown in [Example 7-28](#).

Example 7-28. State-aware singleton

```
[ServiceBehavior(InstanceContextMode = InstanceContextMode.Single)]
class MySingleton : IMyContract
{
    readonly static string m_StateIdentifier = typeof(MySingleton).GUID.ToString();

    [OperationBehavior(TransactionScopeRequired = true)]
    public void MyMethod()
    {
        GetState();
        DoWork();
        SaveState();
    }

    //Helper methods
    void GetState()
    {
        //Use m_StateIdentifier to get state
    }
    void DoWork()
    {}
    public void SaveState()
    {
        //Use m_StateIdentifier to save state
    }
    public void RemoveState()
```

```

    {
        //Use m_StateIdentifier to remove the state from the resource manager
    }
}

//Hosting code
MySingleton singleton = new MySingleton();
singleton.SaveState(); //Create the initial state in the resource manager

ServiceHost host = new ServiceHost(singleton);
host.Open();

/* Some blocking calls */

host.Close();
singleton.RemoveState();

```

In this example, the singleton uses the unique GUID associated with every type as a state identifier. At the beginning of every method call the singleton reads its state, and at the end of each method call it saves the state back to the resource manager. However, the first call on the first instance must also be able to bind to the state, so you must prime the resource manager with the state before the first call ever arrives. To that end, before launching the host, you need to create the singleton, save its state to the resource manager, and then provide the singleton instance to `ServiceHost` (as explained in [Chapter 4](#)). After the host shuts down, make sure to remove the singleton state from the resource manager, as shown in [Example 7-28](#). Note that you cannot create the initial state in the singleton constructor, because the constructor will be called for each operation on the singleton and will override the previous saved state.

While a state-aware singleton is certainly possible (as demonstrated in [Example 7-28](#)), the overall complexity involved makes it a technique to avoid. It is better to use a stateful transactional singleton, as presented next.

Stateful singleton service

By setting `ReleaseServiceInstanceOnTransactionComplete` to `false`, you regain the singleton semantic. The singleton will be created just once, when the host is launched, and the same single instance will be shared across all clients and transactions. The problem is, of course, how to manage the state of the singleton. The singleton has to have state; otherwise, there is no point in using a singleton in the first place. The best solution (as before, with the stateful per-session service) is to use volatile resource managers as member variables, as shown in [Example 7-29](#).

Example 7-29. Achieving a stateful singleton transactional service

```
/////////////// Service Side /////////////////////////////////
[ServiceBehavior(InstanceContextMode = InstanceContextMode.Single,
    ReleaseServiceInstanceOnTransactionComplete = false)]
```

```

class MySingleton : IMyContract
{
    Transactional<int> m_Counter = new Transactional<int>();

    [OperationContract(TransactionScopeRequired = true)]
    public void MyMethod()
    {
        m_Counter.Value++;
        Trace.WriteLine("Counter: " + m_Counter.Value);
    }
}

////////// Client Side ///////////
using(TransactionScope scope1 = new TransactionScope())
{
    MyContractClient proxy = new MyContractClient();
    proxy.MyMethod();
    proxy.Close();
    scope1.Complete();
}

using(TransactionScope scope2 = new TransactionScope())
{
    MyContractClient proxy = new MyContractClient();
    proxy.MyMethod();
    proxy.Close();
}
using(TransactionScope scope3 = new TransactionScope())
{
    MyContractClient proxy = new MyContractClient();
    proxy.MyMethod();
    proxy.Close();
    scope3.Complete();
}

////////// Output ///////////
Counter: 1
Counter: 2
Counter: 2

```

In [Example 7-29](#), a client creates three transactional scopes, each with its own new proxy to the singleton. In each call, the singleton increments a counter it maintains as a `Transactional<int>` volatile resource manager. `scope1` completes the transaction and commits the new value of the counter (1). In `scope2`, the client calls the singleton and temporarily increments the counter to 2. However, `scope2` does not complete its transaction. The volatile resource manager therefore rejects the increment and reverts to its previous value of 1. The call in `scope3` then increments the counter again from 1 to 2, as shown in the trace output.

Note that when setting `ReleaseServiceInstanceOnTransactionComplete`, the singleton must have at least one method with `TransactionScopeRequired` set to `true`.

In addition, the singleton must have `TransactionAutoComplete` set to true on every method, which of course precludes any transactional affinity and allows concurrent transactions. All calls and all transactions are routed to the same instance. For example, the following client code will result in the transaction diagram shown in Figure 7-13:

```
using (MyContractClient proxy = new MyContractClient())
using(TransactionScope scope = new TransactionScope())
{
    proxy.MyMethod();
    scope.Complete();
}

using(MyContractClient proxy = new MyContractClient())
using(TransactionScope scope = new TransactionScope())
{
    proxy.MyMethod();
    proxy.MyMethod();
    scope.Complete();
}
```

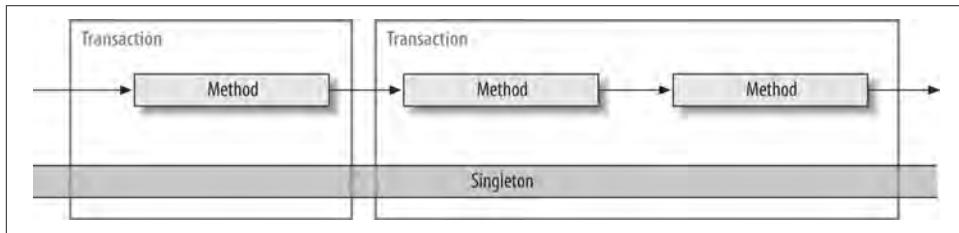


Figure 7-13. Stateful transactional singleton

Instancing Modes and Transactions

To summarize the topic of instance management modes and transactions, Table 7-3 lists the possible configurations discussed so far and their resulting effects. Other combinations may technically be allowed, but I've omitted them because they are either nonsensical or plainly disallowed by WCF.

With so many options, which mode should you choose? I find that the complexity of an explicit state-aware programming model with sessionful and singleton services outweighs any potential benefits, and this is certainly the case with the hybrid mode as well. Equating sessions with transactions is often impractical and indicates a bad design. For both sessionful and singleton services, I prefer the simplicity and elegance of volatile resource managers as member variables. You can also use a durable service on top of a transactional durable storage or the `TransactionalBehavior` attribute for the ultimate programming fusion of regular .NET programming with transactions.

Table 7-3. Possible instancing modes, configurations, and transactions

Configured instancing mode	Autocomplete	Release on complete	Complete on session close	Resulting instancing mode	State mgmt.	Trans. affinity
Per-call	True	True/False	True/False	Per-call	State-aware	Call
Session	True	True	True/False	Per-call	State-aware	Call
Session	True	False	True/False	Session	VRM members	Call
Session	False	True/False	True	Session	Stateful	Instance context
Session	Hybrid	True	True/False	Hybrid	Hybrid	Instance context
Durable service and transactional behavior	True	True/False	True/False	Per-call	Stateful	Call
Singleton	True	True	True/False	Per-call	State-aware	Call
Singleton	True	False	True/False	Singleton	VRM members	Call

Table 7-4 lists these recommended configurations. None of the recommended options relies on transactional affinity or autocompletion on session close, but they all use autocompletion.

Table 7-4. Recommended instancing modes, configurations, and transactions

Configured instancing mode	Release on complete	Resulting instancing mode	State management
Per-call	True/False	Per-call	State-aware
Session	False	Session	VRM members
Durable service and transactional behavior	True/False	Per-call	Stateful
Singleton	False	Singleton	VRM members

Callbacks

Callback contracts, just like service contracts, can propagate the service transaction to the callback client. To enable this, you apply the `TransactionFlow` attribute, as with a service contract. For example:

```
interface IMyContractCallback
{
    [OperationContract]
    [TransactionFlow(TransactionFlowOption.Allowed)]
    void OnCallback();
}
[ServiceContract(CallbackContract = typeof(IMyContractCallback))]
interface IMyContract
{...}
```

The callback method implementation can use the `OperationBehavior` attribute (just like a service operation) and specify whether to require a transaction scope and auto-completion:

```
class MyClient : IMyContractCallback
{
    [OperationBehavior(TransactionScopeRequired = true)]
    public void OnCallback()
    {
        Transaction transaction = Transaction.Current;
        Debug.Assert(transaction != null);
    }
}
```

Callback Transaction Modes

The callback client can have four modes of configuration: Service, Service/Callback, Callback, and None. These are analogous to the service transaction modes, except the service now plays the client role and the callback plays the service role. For example, to configure the callback for the Service transaction mode (that is, to always use the service's transaction), follow these steps:

1. Use a transaction-aware duplex binding with transaction flow enabled.
2. Set transaction flow to mandatory on the callback operation.
3. Configure the callback operation to require a transaction scope.

Example 7-30 shows a callback client configured for the Service transaction mode.

Example 7-30. Configuring the callback for the Service transaction mode

```
interface IMyContractCallback
{
    [OperationContract]
    [TransactionFlow(TransactionFlowOption.Mandatory)]
    void OnCallback();
}

class MyClient : IMyContractCallback
{
    [OperationBehavior(TransactionScopeRequired = true)]
    public void OnCallback()
    {
        Transaction transaction = Transaction.Current;
        Debug.Assert(transaction.TransactionInformation.
            DistributedIdentifier != Guid.Empty);
    }
}
```

When the callback operation is configured for mandatory transaction flow, WCF will enforce the use of a transaction-aware binding with transaction flow enabled.

When you configure the callback for the Service/Callback transaction propagation mode, WCF does not enforce enabling of transaction flow in the binding. You can use my `BindingRequirement` attribute to enforce this:

```
interface IMyContractCallback
{
    [OperationContract]
    [TransactionFlow(TransactionFlowOption.Allowed)]
    void OnCallback();
}[[BindingRequirement(TransactionFlowEnabled = true)]]
class MyClient : IMyContractCallback
{
    [OperationBehavior(TransactionScopeRequired = true)]
    public void OnCallback()
    {...}
}
```

I extended my `BindingRequirement` attribute to verify the callback binding by implementing the `IEndpointBehavior` interface:

```
public interface IEndpointBehavior
{
    void AddBindingParameters(ServiceEndpoint endpoint,
                             BindingParameterCollection bindingParameters);
    void ApplyClientBehavior(ServiceEndpoint endpoint,
                            ClientRuntime clientRuntime);
    void ApplyDispatchBehavior(ServiceEndpoint endpoint,
                             EndpointDispatcher endpointDispatcher);
    void Validate(ServiceEndpoint serviceEndpoint);
}
```

As explained in [Chapter 6](#), the `IEndpointBehavior` interface lets you configure the client-side endpoint used for the callback by the service. In the case of the `BindingRequirement` attribute, it uses the `IEndpointBehavior.Validate()` method, and the implementation is almost identical to that of [Example 7-3](#).

Isolation and timeouts

As with a service, the `CallbackBehavior` attribute enables a callback type to control its transaction's timeout and isolation level:

```
[AttributeUsage(AttributeTargets.Class)]
public sealed class CallbackBehaviorAttribute: Attribute, IEndpointBehavior
{
    public IsolationLevel TransactionIsolationLevel
    {get;set;}
    public string TransactionTimeout
    {get;set;}}
```

```
    //More members  
}
```

These properties accept the same values as in the service case, and the same reasoning can be used to choose a particular value.

Callback Voting

By default, WCF will use automatic voting for the callback operation, just as with a service operation. Any exception in the callback will result in a vote to abort the transaction, and without an error WCF will vote to commit the transaction, as is the case in [Example 7-30](#). However, unlike with a service instance, the callback instance lifecycle is managed by the client, and it has no instancing mode. Any callback instance can be configured for explicit voting by setting `TransactionAutoComplete` to `false`. Voting can then be done explicitly using `SetTransactionComplete()`:

```
class MyClient : IMyContractCallback  
{  
    [OperationBehavior(TransactionScopeRequired = true,  
                      TransactionAutoComplete = false)]  
    public void OnCallback()  
    {  
        /* Do some transactional work then */  
  
        OperationContext.Current.SetTransactionComplete();  
    }  
}
```

As with a per-session service, explicit voting is for the case when the vote depends on other things besides exceptions. Do not perform any work—especially transactional work—after the call to `SetTransactionComplete()`. Calling `SetTransactionComplete()` should be the last line of code in the callback operation, just before returning. If you try to perform any transactional work (including accessing `Transaction.Current`) after the call to `SetTransactionComplete()`, WCF will throw an `InvalidOperationException` and abort the transaction.

Using Transactional Callbacks

While WCF provides the infrastructure for propagating the service's transaction to the callback, in reality callbacks and service transactions do not mix well. First, callbacks are usually one-way operations, and as such cannot propagate transactions. Second, to be able to invoke the callback to its calling client, the service cannot be configured with `ConcurrencyMode.Single`; otherwise, WCF will abort the call to avoid a deadlock. Typically, services are configured for either the Client/Service or the Client transaction propagation mode. Ideally, a service should be able to propagate its original calling client's transaction to all callbacks it invokes, even if the callback is to the calling client. Yet, for the service to use the client's transaction, `TransactionSco`

`peRequired` must be set to `true`. Since `ReleaseServiceInstanceOnTransactionComplete` is `true` by default, it requires `ConcurrencyMode.Single`, thus precluding the callback to the calling client.

Out-of-band transactional callbacks

There are two types of transactional callbacks. The first is out-of-band callbacks made by non-service parties on the host side using callback references stored by the service. Such parties can easily propagate their transactions to the callback (usually in a `TransactionScope`) because there is no risk of a deadlock, as shown in [Example 7-31](#).

Example 7-31. Out-of-band callbacks

```
[ServiceBehavior(InstanceContextMode = InstanceContextMode.PerCall)]
class MyService : IMyContract
{
    static List<IMyContractCallback> m_Callbacks = new List<IMyContractCallback>();

    public void MyMethod()
    {
        IMyContractCallback callback = OperationContext.Current.
            GetCallbackChannel<IMyContractCallback>();

        if(m_Callbacks.Contains(callback) == false)
        {
            m_Callbacks.Add(callback);
        }
    }

    public static void CallClients()
    {
        Action<IMyContractCallback> invoke = (callback)=>
        {
            using(TransactionScope scope =
                new TransactionScope())
            {
                callback.OnCallback();
                scope.Complete();
            };
        };

        m_Callbacks.ForEach(invoke);
    }
}

//Out-of-band callbacks:
MyService.CallClients();
```

Service transactional callbacks

The second option is to carefully configure the transactional service so that it is able to call back to its calling client. To that end, configure the service with `Concurrency`

Mode.Reentrant, set ReleaseServiceInstanceOnTransactionComplete to false, and make sure at least one operation has TransactionScopeRequired set to true, as shown in [Example 7-32](#).

Example 7-32. Configuring for transactional callbacks

```
[ServiceContract(CallbackContract = typeof(IMyContractCallback))]
interface IMyContract
{
    [OperationContract]
    [TransactionFlow(TransactionFlowOption.Allowed)]
    void MyMethod();
}
interface IMyContractCallback
{
    [OperationContract]
    [TransactionFlow(TransactionFlowOption.Allowed)]
    void OnCallback();
}
[ServiceBehavior(InstanceContextMode = InstanceContextMode.PerCall,
                 ConcurrencyMode = ConcurrencyMode.Reentrant,
                 ReleaseServiceInstanceOnTransactionComplete = false)]
class MyService : IMyContract
{
    [OperationBehavior(TransactionScopeRequired = true)]
    public void MyMethod()
    {
        Trace.WriteLine("Service ID: " +
                        Transaction.Current.TransactionInformation.DistributedIdentifier);

        IMyContractCallback callback =
            OperationContext.Current.GetCallbackChannel<IMyContractCallback>();
        callback.OnCallback();
    }
}
```

The rationale behind this constraint is explained in the next chapter.

Given the definitions of [Example 7-32](#), if transaction flow is enabled in the binding, the following client code:

```
class MyCallback : IMyContractCallback
{
    [OperationBehavior(TransactionScopeRequired = true)]
    public void OnCallback()
    {
        Trace.WriteLine("OnCallback ID: " + Transaction.Current.
                        TransactionInformation.
                        DistributedIdentifier);
    }
}
```

```
MyCallback callback = new MyCallback();
InstanceContext context = new InstanceContext(callback);
MyContractClient proxy = new MyContractClient(context);

using(TransactionScope scope = new TransactionScope())
{
    proxy.MyMethod();

    Trace.WriteLine("Client ID:      " + Transaction.Current.
                    TransactionInformation.
                    DistributedIdentifier);
    scope.Complete();
}
proxy.Close();
```

yields output similar to this:

```
Service ID: 23627e82-507a-45d5-933c-05e5e5a1ae78
OnCallback ID: 23627e82-507a-45d5-933c-05e5e5a1ae78
Client ID: 23627e82-507a-45d5-933c-05e5e5a1ae78
```

This indicates that the client's transaction was propagated to the service and into the callback.

Concurrency Management

Incoming client calls are dispatched to the service on threads from the Windows I/O completion thread pool (the pool has 1,000 threads by default). Multiple clients can make multiple concurrent calls, and the service itself can sustain those calls on multiple threads. If the calls are dispatched to the same service instance, you must provide thread-safe access to the service's in-memory state or risk state corruption and errors. The same is true for the client's in-memory state during callbacks, since callbacks too are dispatched on threads from the I/O completion thread pool. In addition to synchronizing access to the instance state when applicable, all services also need to synchronize access to resources shared between instances, such as static variables. Another dimension altogether for concurrency management is ensuring that, if required, the service (or the resources it accesses) executes on particular threads.

WCF offers two modes for synchronization. *Automatic synchronization* instructs WCF to synchronize access to the service instance. Automatic synchronization is simple to use, but it is available only for service and callback classes. *Manual synchronization*, on the other hand, puts the full burden of synchronization on the developer and requires application-specific integration. The developer needs to employ .NET synchronization locks, which is by far an expert discipline. The advantages of manual synchronization are that it is available for service and non-service classes alike, and it allows developers to optimize throughput and scalability. This chapter starts by describing the basic concurrency modes available and then presents more advanced aspects of concurrency management, such as dealing with resource safety and synchronization, thread affinity and custom synchronization contexts, callbacks, and asynchronous calls. Throughout, the chapter shares best practices, concurrency management design guidelines, and custom techniques.

Instance Management and Concurrency

Service-instance thread safety is closely related to the service instancing mode. A per-call service instance is thread-safe by definition, because each call gets its own dedicated instance. That instance is accessible only by its assigned worker thread, and because no other threads will be accessing it, it has no need for synchronization. However, a per-call service is typically state-aware. The state store can be an in-memory resource such as static dictionary, and it can be subject to multithreaded access because the service can sustain concurrent calls, whether from the same client or from multiple clients. Consequently, you must synchronize access to the state store.

A per-session service always requires concurrency management and synchronization, because the client may use the same proxy and yet dispatch calls to the service on multiple client-side threads. A singleton service is even more susceptible to concurrent access, and must support synchronized access. The singleton has some in-memory state that all clients implicitly share. On top of the possibility of the client dispatching calls on multiple threads, as with a per-session service, a singleton may simply have multiple clients in different execution contexts, each using its own thread to call the service. All of these calls will enter the singleton on different threads from the I/O completion thread pool—hence the need for synchronization.

Service Concurrency Modes

Concurrent access to the service instance is governed by the `ConcurrencyMode` property of the `ServiceBehavior` attribute:

```
public enum ConcurrencyMode
{
    Single,
    Reentrant,
    Multiple
}

[AttributeUsage(AttributeTargets.Class)]
public sealed class ServiceBehaviorAttribute : ...
{
    public ConcurrencyMode ConcurrencyMode
    {get;set;}
    //More members
}
```

The value of the `ConcurrencyMode` enum controls if and when concurrent calls are allowed. The name `ConcurrencyMode` is actually incorrect; the proper name for this property would have been `ConcurrencyContextMode`, since it synchronizes access not to the instance, but rather to the context containing the instance (much the same way

`InstanceContextMode` controls the instantiation of the context, not the instance). The significance of this distinction—i.e., that the synchronization is related to the context and not to the instance—will become evident later.

ConcurrencyMode.Single

When the service is configured with `ConcurrencyMode.Single`, WCF will provide automatic synchronization to the service context and disallow concurrent calls by associating the context containing the service instance with a synchronization lock. Every call coming into the service must first try to acquire the lock. If the lock is unowned, the caller will be allowed in. Once the operation returns, WCF will unlock the lock, thus allowing in another caller.

The important thing is that only one caller at a time is ever allowed. If there are multiple concurrent callers while the lock is locked, the callers are all placed in a queue and are served out of the queue in order. If a call times out while blocked, WCF will remove the caller from the queue and the client will get a `TimeoutException`. `ConcurrencyMode.Single` is the WCF default setting, so these definitions are equivalent:

```
class MyService : IMyContract
{...}

[ServiceBehavior(ConcurrencyMode = ConcurrencyMode.Single)]
class MyService : IMyContract
{...}
```

Because the default concurrency mode is synchronized access, the susceptible instancing modes of per-session and singleton are also synchronized by default. Note that even calls to a per-call service instance are synchronized by default.

Synchronized access and transactions

As explained in [Chapter 7](#), WCF will verify at service load time whether at least one operation on the service has `TransactionScopeRequired` set to `true` and that `ReleaseServiceInstanceOnTransactionComplete` is `true`. In this case, the service concurrency mode must be `ConcurrencyMode.Single`. This is done deliberately to ensure that the service instance can be recycled at the end of the transaction without any danger of there being another thread accessing the disposed instance.

ConcurrencyMode.Multiple

When the service is configured with `ConcurrencyMode.Multiple`, WCF will stay out of the way and will not synchronize access to the service instance in any way. `ConcurrencyMode.Multiple` simply means that the service instance is not associated with any synchronization lock, so concurrent calls are allowed on the service instance. Put differently, when a service instance is configured with `ConcurrencyMode.Multiple`,

WCF will not queue up the client messages and dispatch them to the service instance as soon as they arrive.



A large number of concurrent client calls will not result in a matching number of concurrently executing calls on the service. The maximum number of concurrent calls dispatched to the service is determined by the configured maximum concurrent calls throttle value.

Obviously, this is of great concern to sessionful and singleton services, which must manually synchronize access to their instance state. The common way of doing that is to use .NET locks such as `Monitor` or a `WaitHandle`-derived class. Manual synchronization is not for the faint of heart, but it does enable the service developer to optimize the throughput of client calls on the service instance: you can lock the service instance just when and where synchronization is required, thus allowing other client calls on the same service instance in between the synchronized sections. [Example 8-1](#) shows a manually synchronized sessionful service whose client performs concurrent calls.

Example 8-1. Manual synchronization using fragmented locking

```
[ServiceContract(SessionMode = SessionMode.Required)]
interface IMyContract
{
    void MyMethod();
}

[ServiceBehaviorConcurrencyMode = ConcurrencyMode.Multiple)]
class MyService : IMyContract
{
    int[] m_Numbers;
    List<string> m_Names;

    public void MyMethod()
    {
        lock(m_Numbers)
        {
            ...
        }

        /* Don't access members here */

        lock(m_Names)
        {
            ...
        }
    }
}
```

The service in [Example 8-1](#) is configured for concurrent access. Since the critical sections of the operations that require synchronization are any member variable accesses, the service uses a `Monitor` (encapsulated in the `lock` statement) to lock the member variable before accessing it. I call this synchronization technique *fragmented locking*, since it locks only when needed and only what is being accessed. Local variables require no synchronization, because they are visible only to the thread that created them on its own call stack.

There are two problems with fragmented locking: it is both error- and deadlock-prone. Fragmented locking only provides for thread-safe access if every other operation on the service is as disciplined about always locking the members before accessing them. But even if all operations lock all members, you still risk deadlocks: if one operation on thread A locks member M1 while trying to access member M2, while another operation executing concurrently on thread B locks member M2 while trying to access member M1, you will end up with a deadlock.



WCF resolves service call deadlocks by eventually timing out the call and throwing a `TimeoutException`. Avoid using a long send timeout, as it decreases WCF's ability to resolve deadlocks in a timely manner.

It is better to reduce the fragmentation by locking the entire service instance instead:

```
public void MyMethod()
{
    lock(this)
    {
        ...
    }

    /* Don't access members here */

    lock(this)
    {
        ...
    }
}
```

This approach, however, is still fragmented and thus error-prone—if at some point in the future someone adds a method call in the unsynchronized code section that does access the members, it will not be a synchronized access. It is better still to lock the entire body of the method:

```
public void MyMethod()
{
    lock(this)
    {
        ...
    }
}
```

```
    }
}
```

The problem with this approach is that in the future someone maintaining this code may err and place some code before or after the lock statement. Your best option therefore is to instruct the compiler to automate injecting the call to lock the instance using the `MethodImpl` attribute with the `MethodImplOptions.Synchronized` flag:

```
[ServiceBehaviorConcurrencyMode = ConcurrencyMode.Multiple)]
class MyService : IMyContract
{
    int[] m_Numbers;
    List<string> m_Names;

    [MethodImpl(MethodImplOptions.Synchronized)]
    public void MyMethod()
    {
        ...
    }
}
```

You will need to repeat the assignment of the `MethodImpl` attribute on all the service operation implementations.

While this code is thread-safe, you actually gain little from the use of `ConcurrencyMode.Multiple`: the net effect in terms of synchronization is similar to using `ConcurrencyMode.Single`, yet you have increased the overall code complexity and reliance on developers' discipline. In general, you should avoid `ConcurrencyMode.Multiple`. However, there are cases where `ConcurrencyMode.Multiple` is useful, as you will see later in this chapter.

Unsynchronized access and transactions

When the service is configured for `ConcurrencyMode.Multiple`, if at least one operation has `TransactionScopeRequired` set to `true`, then `ReleaseServiceInstanceOnTransactionComplete` must be set to `false`. For example, this is a valid definition, even though `ReleaseServiceInstanceOnTransactionComplete` defaults to `true`, because no method has `TransactionScopeRequired` set to `true`:

```
[ServiceBehaviorConcurrencyMode = ConcurrencyMode.Multiple)]
class MyService : IMyContract
{
    public void MyMethod()
    {...}
    public void MyOtherMethod()
    {...}
}
```

The following, on the other hand, is an invalid definition because at least one method has `TransactionScopeRequired` set to `true`:

```
//Invalid configuration:
[ServiceBehaviorConcurrencyMode = ConcurrencyMode.Multiple]
class MyService : IMyContract
{
    [OperationBehavior(TransactionScopeRequired = true)]
    public void MyMethod()
    {...}
    public void MyOtherMethod()
    {...}
}
```

A transactional unsynchronized service must explicitly set `ReleaseServiceInstanceOnTransactionComplete` to `false`:

```
[ServiceBehaviorConcurrencyMode = ConcurrencyMode.Multiple,
 ReleaseServiceInstanceOnTransactionComplete = false]
class MyService : IMyContract
{
    [OperationBehavior(TransactionScopeRequired = true)]
    public void MyMethod()
    {...}
    public void MyOtherMethod()
    {...}
}
```

The rationale behind this constraint is that only a sessionful or a singleton service could possibly benefit from unsynchronized access, so in the case of transactional access, WCF wants to enforce the semantic of the configured instancing mode. In addition, this will avoid having one caller access the instance, complete the transaction, and release the instance, all while another caller is using the instance.

ConcurrencyMode.Reentrant

The `ConcurrencyMode.Reentrant` value is a refinement of `ConcurrencyMode.Single`. Similar to `ConcurrencyMode.Single`, `ConcurrencyMode.Reentrant` associates the service context with a synchronization lock, so concurrent calls on the same instance are never allowed. However, if the reentrant service calls out to another service or a callback, and that call chain (or *causality*) somehow winds its way back to the service instance, as shown in [Figure 8-1](#), that call is allowed to reenter the service instance.

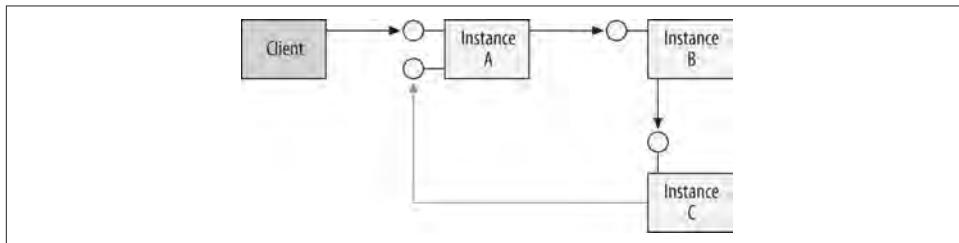


Figure 8-1. Call reentrancy

The implementation of `ConcurrencyMode.Reentrant` is very simple—when the reentrant service calls out over WCF, WCF silently releases the synchronization lock associated with the instance context. `ConcurrencyMode.Reentrant` is designed to avoid the potential deadlock of reentrancy, although it will release the lock in case of a callout. If the service were to maintain the lock while calling out, if the causality tried to enter the same context, a deadlock would occur.

Reentrancy support is instrumental in a number of cases:

- A singleton service calling out risks a deadlock if any of the downstream services it calls tries to call back into the singleton.
- In the same app domain, if the client stores a proxy reference in some globally available variable, then some of the downstream objects called by the service use the proxy reference to call back to the original service.
- Callbacks on non-one-way operations must be allowed to reenter the calling service.
- If the callout the service performs is of long duration, even without reentrancy, you may want to optimize throughput by allowing other clients to use the same service instance while the callout is in progress.



A service configured with `ConcurrencyMode.Multiple` is by definition also reentrant, because no lock is held during the callout. However, unlike a reentrant service, which is inherently thread-safe, a service configured with `ConcurrencyMode.Multiple` must provide for its own synchronization (for example, by locking the instance during every call, as explained previously). It is up to the developer of such a service to decide if it should release the lock before calling out to avoid a reentrancy deadlock.

Designing for reentrancy

It is very important to recognize the liability associated with reentrancy. When a reentrant service calls out, it must leave the service in a workable, consistent state, because others could be allowed into the service instance while the service is calling out. A consistent state means that the reentrant service must have no more interactions with its own members or any other local object or static variable, and that when the callout returns, the reentrant service should simply be able to return control to its client. For example, suppose the reentrant service modifies the state of some linked list and leaves it in an inconsistent state—say, missing a head node—because it needs to get the value of the new head from another service. If the reentrant service then calls out to the other service, it leaves other clients vulnerable, because if they call into the reentrant service and access the linked list they will encounter an error.

Moreover, when the reentrant service returns from its callout, it must refresh all local method state. For example, if the service has a local variable that contains a copy of the state of a member variable, that local variable may now have the wrong value, because during the callout another party could have entered the reentrant service and modified the member variable.

Reentrancy and transactions

A reentrant service faces exactly the same design constraints regarding transactions as a service configured with `ConcurrencyMode.Multiple`; namely, if at least one operation has `TransactionScopeRequired` set to `true`, then `ReleaseServiceInstanceOnTransactionComplete` must be set to `false`. This is done to maintain the instance context mode semantics.

Callbacks and reentrancy

Consider now the case of a service designed for single-threaded access with `ConcurrencyMode.Single` and with duplex callbacks. When a call from the client enters the context, it acquires the synchronization lock. If that service obtains the callback reference and calls back to the calling client, that callout will block the thread used to issue the call from the client while still maintaining the lock on the context. The callback will reach the client, execute there, and return with a reply message from the client. Unfortunately, when the reply message is sent to the same service instance context, it will first try to acquire the lock—the same lock already owned by the original call from the client, which is still blocked waiting for the callback to return—and a deadlock will ensue. To avoid this deadlock, during the operation execution, WCF disallows callbacks from the service to its calling client as long as the service is configured for single-threaded access.

There are three ways of safely allowing the callback. The first is to configure the service for reentrancy. When the service invokes the proxy to the callback object, WCF will silently release the lock, thus allowing the reply message from the callback to acquire the lock when it returns, as shown in [Example 8-2](#).

Example 8-2. Configure for reentrancy to allow callbacks

```
interface IMyContractCallback
{
    [OperationContract]
    void OnCallback();
}

[ServiceContract(CallbackContract = typeof(IMyContractCallback))]
interface IMyContract
{
    [OperationContract]
    void MyMethod();
```

```

}

[ServiceBehaviorConcurrencyMode = ConcurrencyMode.Reentrant)]
class MyService : IMyContract
{
    public void MyMethod()
    {
        IMyContractCallback callback = OperationContext.Current.
            GetCallbackChannel<IMyContractCallback>();
        callback.OnCallback();
    }
}

```

Control will only return to the service once the callback returns, and the service's own thread will need to reacquire the lock. Configuring for reentrancy is required even of a per-call service, which otherwise has no need for anything but `ConcurrencyMode.Single`. Note that the service may still invoke callbacks to other clients or call other services; it is the callback to the calling client that is disallowed.

You can, of course, configure the service for concurrent access with `ConcurrencyMode.Multiple` to avoid having any lock.

The third option (as mentioned in [Chapter 5](#)), and the only case where a service configured with `ConcurrencyMode.Single` can call back to its clients, is when the callback contract operation is configured as one-way because there will not be any reply message to contend for the lock.

Instances and Concurrent Access

Using the same proxy, a single client can issue multiple concurrent calls to a service. The client can use multiple threads to invoke calls on the service, or it can issue one-way calls in rapid succession on the same thread. In both of these cases, whether the calls from the same client are processed concurrently is the product of the service's configured instancing mode, the service's concurrency mode, and the configured delivery mode (that is, the transport session). The following discussion applies equally to request-reply and one-way calls.

Per-Call Services

In the case of a per-call service, if there is no transport-level session, concurrent processing of calls is allowed. Calls are dispatched as they arrive, each to a new instance, and execute concurrently. This is the case regardless of the service concurrency mode. I consider this to be the correct behavior.

If the per-call service has a transport-level session, whether concurrent processing of calls is allowed is a product of the service concurrency mode. If the service is config-

ured with `ConcurrencyMode.Single`, concurrent processing of the pending calls is not allowed, and the calls are dispatched one at a time. The reason is that with `ConcurrencyMode.Single` WCF tries to maintain the guarantee of the transport session that messages are processed strictly in the order in which they were received in that session by having exactly one outstanding instance per channel. You should avoid lengthy processing of calls, because it may risk call timeouts.

While this is a direct result of the channel's architecture, I consider this to be a flawed design. If the service is configured with `ConcurrencyMode.Multiple`, concurrent processing is allowed. Calls are dispatched as they arrive, each to a new instance, and execute concurrently. An interesting observation here is that in the interest of throughput, it is a good idea to configure a per-call service with `ConcurrencyMode.Multiple`—the instance itself will still be thread-safe (so you will not incur the synchronization liability), yet you will allow concurrent calls from the same client.



Two clients using two different proxies will have two distinct channels and will have no issue with concurrent calls. It is only concurrent calls on the same transport session that are serialized one at a time to the per-call service.

When the service is configured with `ConcurrencyMode.Reentrant`, if the service does not call out, it behaves similarly to a service configured with `ConcurrencyMode.Single`. If the service does call out, the next call is allowed in, and the returning call has to negotiate the lock like all other pending calls.

Sessionful and Singleton Services

In the case of a sessionful or a singleton service, the configured concurrency mode alone governs the concurrent execution of pending calls. If the service is configured with `ConcurrencyMode.Single`, calls will be dispatched to the service instance one at a time, and pending calls will be placed in a queue. You should avoid lengthy processing of calls, because it may risk call timeouts.

If the service instance is configured with `ConcurrencyMode.Multiple`, concurrent processing of calls from the same client is allowed. Calls will be executed by the service instance as fast as they come off the channel (up to the throttle limit). Of course, as is always the case with a stateful unsynchronized service instance, you must synchronize access to the service instance or risk state corruption.

If the service instance is configured with `ConcurrencyMode.Reentrant`, it behaves just as it would with `ConcurrencyMode.Single`. However, if the service calls out, the next call is allowed to execute. You must follow the guidelines discussed previously regarding programming in a reentrant environment.



For a service configured with `ConcurrencyMode.Multiple` to experience concurrent calls, the client must use multiple worker threads to access the same proxy instance. However, if the client threads rely on the auto-open feature of the proxy (that is, just invoking a method and having that call open the proxy if the proxy is not yet open) and call the proxy concurrently, then the calls will actually be serialized until the proxy is opened, and will be concurrent after that. If you want to dispatch concurrent calls regardless of the state of the proxy, the client needs to explicitly open the proxy (by calling the `Open()` method) before issuing any calls on the worker threads.

Resources and Services

Synchronizing access to the service instance using `ConcurrencyMode.Single` or an explicit synchronization lock only manages concurrent access to the service instance state itself. It does not provide safe access to the underlying resources the service may be using. These resources must also be thread-safe. For example, consider the application shown in [Figure 8-2](#).

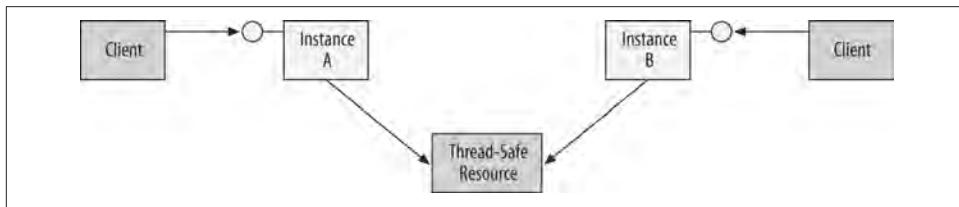


Figure 8-2. Applications must synchronize access to resources

Although the service instances are thread-safe, the two instances try to concurrently access the same resource (such as a static variable, a helper static class, or a file), and therefore the resource itself must have synchronized access. This is true regardless of the service instancing mode. Even a per-call service could run into the situation shown in [Figure 8-2](#).

Deadlocked Access

The naïve solution to providing thread-safe access to resources is to provide each resource with its own lock, potentially encapsulating that lock in the resource itself, and ask the resource to lock the lock when it's accessed and unlock the lock when the service is done with the resource. The problem with this approach is that it is deadlock-prone. Consider the situation depicted in [Figure 8-3](#).

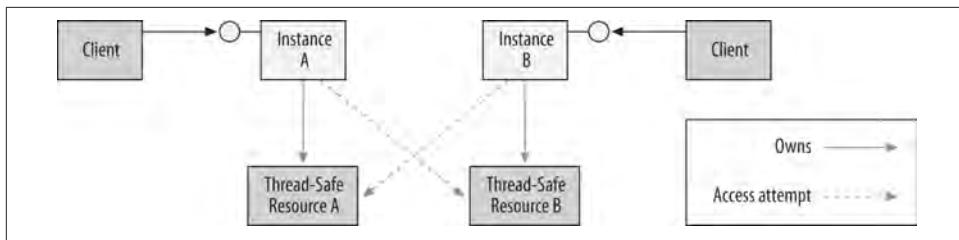


Figure 8-3. Deadlock over resources access

In the figure, Instance A of the service accesses the thread-safe Resource A. Resource A has its own synchronization lock, and Instance A acquires that lock. Similarly, Instance B accesses Resource B and acquires its lock. A deadlock then occurs when Instance A tries to access Resource B while Instance B tries to access Resource A, since each instance will be waiting for the other to release its lock.

The concurrency and instancing modes of the service are almost irrelevant to avoiding this deadlock. The only case that avoids it is if the service is configured both with `InstanceContextMode.Single` and `ConcurrencyMode.Single`, because a synchronized singleton by definition can only have one client at a time and there will be no other instance to deadlock with over access to resources. All other combinations are still susceptible to this kind of deadlock. For example, a per-session synchronized service may have two separate thread-safe instances associated with two different clients, yet the two instances can deadlock when accessing the resources.

Deadlock Avoidance

There are a few possible ways to avoid the deadlock. If all instances of the service meticulously access all resources in the same order (e.g., always trying to acquire the lock of Resource A first, and then the lock of Resource B), there will be no deadlock. The problem with this approach is that it is difficult to enforce, and over time, during code maintenance, someone may deviate from this strict guideline (even inadvertently, by calling methods on helper classes) and trigger the deadlock.

Another solution is to have all resources use the same shared lock. In order to minimize the chances of a deadlock, you'll also want to minimize the number of locks in the system and have the service itself use the same lock. To that end, you can configure the service with `ConcurrencyMode.Multiple` (even with a per-call service) to avoid using the WCF-provided lock. The first service instance to acquire the shared lock will lock out all other instances and own all underlying resources. A simple technique for using such a shared lock is locking on the service type, as shown in Example 8-3.

Example 8-3. Using the service type as a shared lock

```
[ServiceBehavior(InstanceContextMode = InstanceContextMode.PerCall,
                 ConcurrencyMode = ConcurrencyMode.Multiple)]
class MyService : IMyContract
{
    public void MyMethod()
    {
        lock(typeof(MyService))
        {
            ...
            MyResource.DoWork();
            ...
        }
    }
}
static class MyResource
{
    public static void DoWork()
    {
        lock(typeof(MyService))
        {
            ...
        }
    }
}
```

The resources themselves must also lock on the service type (or some other shared type agreed upon in advance). There are two problems with the approach of using a shared lock. First, it introduces coupling between the resources and the service, because the resource developer has to know about the type of the service or the type used for synchronization. While you could get around that by providing the type as a resource construction parameter, it will likely not be applicable with third-party-provided resources. The second problem is that while your service instance is executing, all other instances (and their respective clients) will be blocked. Therefore, in the interest of throughput and responsiveness, you should avoid lengthy operations when using a shared lock.

If you think the situation in [Example 8-3](#), where the two instances are of the same service, is problematic, imagine what happens if the two instances are of different services. The observation to make here is that services should never share resources. Regardless of concurrency management, resources are local implementation details and therefore should not be shared across services. Most importantly, sharing resources across the service boundary is also deadlock-prone. Such shared resources have no easy way to share locks across technologies and organizations, and the services need to somehow coordinate the locking order. This necessitates a high degree of coupling between the services, violating the best practices and tenets of service-orientation.

Resource Synchronization Context

Incoming service calls execute on worker threads from the I/O completion thread pool and are unrelated to any service or resource threads. This means that by default the service cannot rely on any kind of *thread affinity* (that is, always being accessed by the same thread). Much the same way, the service cannot by default rely on executing on any host-side custom threads created by the host or service developers. The problem with this situation is that some resources may rely on thread affinity. For example, user interface resources updated by the service must execute and be accessed only by the user interface (UI) thread. Other examples are a resource (or a service) that makes use of the thread local storage (TLS) to store out-of-band information shared globally by all parties on the same thread (using the TLS mandates use of the same thread), or accessing legacy COM or ActiveX components, which also require thread affinity (due to their own use of the TLS). In addition, for scalability and throughput purposes, some resources or frameworks may require access by their own pool of threads.

Whenever an affinity to a particular thread or threads is expected, the service cannot simply execute the call on the incoming WCF worker thread. Instead, the service must marshal the call to the correct thread(s) required by the resource it accesses.

.NET Synchronization Contexts

.NET 2.0 introduced the concept of a *synchronization context*. The idea is that any party can provide an execution context and have other parties marshal calls to that context. The synchronization context can be a single thread or any number of designated threads, although typically it will be just a single, yet particular thread. All the synchronization context does is assure that the call executes on the correct thread or threads.

Note that the word *context* is overloaded. Synchronization contexts have absolutely nothing to do with the service instance context or the operation context described so far in this book. They are simply the synchronizational context of the call.

While synchronization contexts are a simple enough design pattern to use conceptually, implementing a synchronization context is a complex programming task that is not normally intended for developers to attempt.

The `SynchronizationContext` class

The `SynchronizationContext` class from the `System.Threading` namespace represents a synchronization context:

```
public delegate void SendOrPostCallback(object state);  
  
public class SynchronizationContext
```

```

{
    public virtual void Post(SendOrPostCallback callback,object state);
    public virtual void Send(SendOrPostCallback callback,object state);
    public static void SetSynchronizationContext(SynchronizationContext context);
    public static SynchronizationContext Current
    {get;}
    //More members
}

```

Every thread in .NET may have a synchronization context associated with it. You can obtain a thread's synchronization context by accessing the static `Current` property of `SynchronizationContext`. If the thread does not have a synchronization context, `Current` will return `null`. You can also pass the reference to the synchronization context between threads, so that one thread can marshal a call to another thread.

To represent the call to invoke in the synchronization context, you wrap a method with a delegate of the type `SendOrPostCallback`. Note that the signature of the delegate uses an `object`. If you want to pass multiple parameters, pack those in a structure and pass the structure as an `object`.



Synchronization contexts use an amorphous `object`. Exercise caution when using synchronization contexts, due to the lack of compile-time type safety. Instead of an `object`, you can use anonymous methods and outer variables (closures) that are type-safe.

Working with the synchronization context

There are two ways of marshaling a call to the synchronization context: synchronously and asynchronously, by sending or posting a work item, respectively. The `Send()` method will block the caller until the call has completed in the other synchronization context, while `Post()` will merely dispatch it to the synchronization context and then return control to its caller.

For example, to synchronously marshal a call to a particular synchronization context, you first somehow obtain a reference to that synchronization context, and then use the `Send()` method:

```

//Obtain synchronization context
SynchronizationContext context = ...

SendOrPostCallback doWork = (arg)=>
{
    //The code here is guaranteed to
    //execute on the correct thread(s)
};
context.Send(doWork,"Some argument");

```

Example 8-4 shows a less abstract example.

Example 8-4. Calling a resource on the correct synchronization context

```
class MyResource
{
    public int DoWork()
    {...}
    public SynchronizationContext MySynchronizationContext
    {get;}
}
class MyService : IMyContract
{
    MyResource GetResource()
    {...}

    public void MyMethod()
    {
        MyResource resource = GetResource();
        SynchronizationContext context = resource.MySynchronizationContext;
        int result = 0;
        SendOrPostCallback doWork = _=>
        {
            result = resource.DoWork();
        };
        context.Send(doWork,null);
    }
}
```

In [Example 8-4](#), the service `MyService` needs to interact with the resource `MyResource` and have it perform some work by executing the `DoWork()` method and returning a result. However, `MyResource` requires that all calls to it execute on its particular synchronization context. `MyResource` makes that execution context available via the `MySynchronizationContext` property. The service operation `MyMethod()` executes on a WCF worker thread. `MyMethod()` first obtains the resource and its synchronization context, then defines a lambda expression that wraps the call to `DoWork()` and assigns that expression to the `doWork` delegate of the type `SendOrPostCallback`. Finally, `MyMethod()` calls `Send()` and passes `null` for the argument, since the `DoWork()` method on the resource requires no parameters. Note the technique used in [Example 8-4](#) to retrieve a returned value from the invocation. Since `Send()` returns `void`, the lambda expression assigns the returned value of `DoWork()` into an outer variable.

The problem with [Example 8-4](#) is the excessive degree of coupling between the service and the resource. The service needs to know that the resource is sensitive to its synchronization context, obtain the context, and manage the execution. You must also duplicate such code in any service using the resource. It is much better to encapsulate the need in the resource itself, as shown in [Example 8-5](#).

Example 8-5. Encapsulating the synchronization context

```
class MyResource
{
    public int DoWork()
    {
        int result = 0;
        SendOrPostCallback doWork = _=>
        {
            result = DoWorkInternal();
        };
        MySynchronizationContext.Send(doWork,null);
        return result;
    }
    SynchronizationContext MySynchronizationContext
    {get;}
    int DoWorkInternal()
    {...}
}
class MyService : IMyContract
{
    MyResource GetResource()
    {...}
    public void MyMethod()
    {
        MyResource resource = GetResource();
        int result = resource.DoWork();
    }
}
```

Compare Example 8-5 to Example 8-4. All the service in Example 8-5 has to do is access the resource: it is up to the service internally to marshal the call to its synchronization context.



The most straightforward way to illustrate the usefulness and power of synchronization contexts, along with concrete examples of the abstract patterns discussed so far is when utilizing synchronization contexts along with one of the Windows user interface frameworks.

Note that while the discussions that follow focus on the UI case, the patterns, design guidelines, and consequences of those design decisions and even the best practices apply to most other cases of a synchronization context. For simplicity's sake, the rest of the discussion in this chapter will refer only to Windows Forms, although it applies equally to all of Microsoft's desktop UI frameworks and to large degree Microsoft's web technologies as well.

The UI Synchronization Context

A Windows UI application relies on the underlying Windows messages and a message-processing loop (the *message pump*) to process them. The message loop must have thread affinity, because messages to a window are delivered only to the thread that created it. In general, you must always marshal to the UI thread any attempt to access a Windows control or form, or risk errors and failures. This becomes an issue if your services need to update some user interface as a result of client calls or some other event. Fortunately, Windows Forms supports the synchronization context pattern. Every thread that pumps Windows messages has a synchronization context, which is the `WindowsFormsSynchronizationContext` class:

```
public sealed class WindowsFormsSynchronizationContext : SynchronizationContext,  
    ...  
{...}
```

Whenever you create any Windows Forms control or form, that control or form ultimately derives from the class `Control`. The constructor of `Control` checks whether the current thread that creates it already has a synchronization context, and if it does not, `Control` installs `WindowsFormsSynchronizationContext` as the current thread's synchronization context.

`WindowsFormsSynchronizationContext` converts the call to `Send()` or `Post()` to a custom Windows message and posts that Windows message to the UI thread's message queue. Every Windows Forms UI class that derives from `Control` has a dedicated method that handles this custom message by invoking the supplied `SendOrPostCallback` delegate. At some point, the UI thread processes the custom Windows message and the delegate is invoked.

Because the window or control can also be called already in the correct synchronization context, to avoid a deadlock when calling `Send()`, the implementation of the Windows Forms synchronization context verifies that marshaling the call is indeed required. If marshaling is not required, it uses direct invocation on the calling thread.

UI access and updates

When a service needs to update a user interface, it must have some proprietary mechanisms to find the window to update in the first place. And once the service has the correct window, it must somehow get hold of that window's synchronization context and marshal the call to it. Such a possible interaction is shown in [Example 8-6](#).

Example 8-6. Using the form synchronization context

```
partial class MyForm : Form  
{  
    Label m_CounterLabel;  
    public readonly SynchronizationContext MySynchronizationContext;
```

```

public MyForm()
{
    InitializeComponent();
    MySynchronizationContext = SynchronizationContext.Current;
}
void InitializeComponent()
{
    ...
    m_CounterLabel = new Label();
    ...
}

public int Counter
{
    get
    {
        return Convert.ToInt32(m_CounterLabel.Text);
    }
    set
    {
        m_CounterLabel.Text = value.ToString();
    }
}
}

[ServiceContract]
interface IFormManager
{
    [OperationContract]
    void IncrementLabel();
}

class MyService : IFormManager
{
    public void IncrementLabel()
    {
        MyForm form = Application.OpenForms[0] as MyForm;
        Debug.Assert(form != null);

        SendOrPostCallback callback = _=>
        {
            form.Counter++;
        };
        form.MySynchronizationContext.Send(callback,null);
    }
}

static class Program
{
    static void Main()
    {
        ServiceHost host = new ServiceHost(typeof(MyService));
        host.Open();
    }
}

```

```

        Application.Run(new MyForm());

        host.Close();
    }
}

```

Example 8-6 shows the form `MyForm`, which provides the `MySynchronizationContext` property that allows its clients to obtain its synchronization context. `MyForm` initializes `MySynchronizationContext` in its constructor by obtaining the synchronization context of the current thread. The thread has a synchronization context because the constructor of `MyForm` is called after the constructor of its topmost base class, `Control`, was called, and `Control` has already attached the Windows Forms synchronization context to the thread in its constructor.

`MyForm` also offers a `Counter` property that updates the value of a counting Windows Forms label. Only the thread that owns the form can access that label. `MyService` implements the `IncrementLabel()` operation. In that operation, the service obtains a reference to the form via the static `OpenForms` collection of the `Application` class:

```

public class FormCollection : ReadOnlyCollectionBase
{
    public virtual Form this[int index]
    {get;}
    public virtual Form this[string name]
    {get;}
}

public sealed class Application
{
    public static FormCollection OpenForms
    {get;}
    //Rest of the members
}

```

Once `IncrementLabel()` has the form to update, it accesses the synchronization context via the `MySynchronizationContext` property and calls the `Send()` method. `Send()` is provided with an anonymous method that accesses the `Counter` property. **Example 8-6** is a concrete example of the programming model shown in **Example 8-4**, and it suffers from the same deficiency: namely, tight coupling between all service operations and the form. If the service needs to update multiple controls, that also results in a cumbersome programming model. Any change to the user interface lay-

out, the controls on the forms, and the required behavior is likely to cause major changes to the service code.

The Task Parallel Library and Synchronization Context

.NET 4.0 introduced the parallel computing library, with programming models and helper types designed to streamline writing, executing, and synchronizing concurrent or parallel programs. The parallel task library does offer a wrapper around synchronization context in the form of a dedicated task scheduler, available with the static method `FromCurrentSynchronizationContext()` of `TaskScheduler`:

```
public abstract class TaskScheduler
{
    public static TaskScheduler FromCurrentSynchronizationContext();
    //More members
}
```

`FromCurrentSynchronizationContext()` obtains the synchronization context of the current thread and returns a task scheduler that marshals all its tasks to that synchronization context. Using the task parallel library, the pertinent elements of [Example 8-6](#) can be written as:

```
partial class MyForm : Form
{
    public int m_Counter
    {get;set;}

    public readonly TaskScheduler Scheduler;

    public MyForm()
    {
        InitializeComponent();
        Scheduler = TaskScheduler.FromCurrentSynchronizationContext();
    }
    ...
}

class MyService : IFormManager
{
    public void IncrementLabel()
    {
        MyForm form = Application.OpenForms[0] as MyForm;
        Debug.Assert(form != null);

        TaskScheduler scheduler = form.Scheduler;
        TaskFactory factory = Task.Factory;

        factory.StartNew(()=>form.Counter++,
                        CancellationToken.None,TaskCreationOptions.None,scheduler);
    }
}
```

Since the two programming models are equivalent in capabilities and complexity, I see no clear advantage of one over the other in simple cases. The advantage of the parallel task library is that it offers a clean way to combine accessing a resource that is sensitive to the affinity of its accessing thread (or threads) as part of a larger set of parallel tasks.

Safe controls

A better approach is to encapsulate the interaction with the Windows Forms synchronization context in safe controls or safe methods on the form, to decouple them from the service and to simplify the overall programming model. [Example 8-7](#) lists the code for `SafeLabel`, a `Label`-derived class that provides thread-safe access to its `Text` property. Because `SafeLabel` derives from `Label`, you still have full design-time visual experience and integration with Visual Studio, yet you can surgically affect just the property that requires the safe access.

Example 8-7. Encapsulating the synchronization context

```
public class SafeLabel : Label
{
    SynchronizationContext m_SynchronizationContext =
        SynchronizationContext.Current;
    override public string Text
    {
        set
        {
            m_SynchronizationContext.Send(_=> base.Text = value,null);
        }
        get
        {
            string text = String.Empty;
            m_SynchronizationContext.Send(_=> text = base.Text,null);
            return text;
        }
    }
}
```

Upon construction, `SafeLabel` caches its synchronization context. `SafeLabel` overrides its base class's `Text` property and uses the lambda expression method in the `get` and `set` accessors to send the call to the correct UI thread. Note in the `get` accessor the use of an outer variable to return a value from `Send()`, as discussed previously. Using `SafeLabel`, the code in [Example 8-6](#) is reduced to the code shown in [Example 8-8](#).

Example 8-8. Using a safe control

```
class MyForm : Form
{
    Label m_CounterLabel;

    public MyForm()
    {
        InitializeComponent();
    }
    void InitializeComponent()
    {
        ...
        m_CounterLabel = new SafeLabel();
        ...
    }
    public int Counter
    {
        get
        {
            return Convert.ToInt32(m_CounterLabel.Text);
        }
        set
        {
            m_CounterLabel.Text = value.ToString();
        }
    }
}
class MyService : IFormManager
{
    public void IncrementLabel()
    {
        MyForm form = Application.OpenForms[0] as MyForm;
        Debug.Assert(form != null);

        form.Counter++;
    }
}
```

Note in [Example 8-8](#) that the service simply accesses the form directly:

```
form.Counter++;
```

and that the form is written as a normal form. [Example 8-8](#) is a concrete example of the programming model shown in [Example 8-5](#).

ServiceModelEx contains not only *SafeLabel*, but also other controls you are likely to update at runtime such as *SafeButton*, *SafeListBox*, *SafeProgressBar*, *SafeStatus Bar*, *SafeTrackBar*, and *SafeTextBox*.



Exercise caution when using the safe controls (or in the general case, safe resources that encapsulate their own synchronization context). While having safe resources does simplify accessing each individual resource, if you have to access multiple resources, you will pay the penalty of marshaling to the synchronization context with every one of them. With multiple resources, it is better to lump all the accesses into a single method and marshal just the call to that method to the target synchronization context.

Service Synchronization Context

The programming techniques shown so far put the onus of accessing the resource on the correct thread squarely on the service or resource developer. It would be preferable if the service had a way of associating itself with a particular synchronization context, and could have WCF detect that context and automatically marshal the call from the worker thread to the associated service synchronization context. In fact, WCF lets you do just that. You can instruct WCF to maintain an affinity between all service instances from a particular host and a specific synchronization context. The **Service Behavior** attribute offers the **UseSynchronizationContext** Boolean property, defined as:

```
[AttributeUsage(AttributeTargets.Class)]
public sealed class ServiceBehaviorAttribute : ...
{
    public bool UseSynchronizationContext
    {get;set;}
    //More members
}
```

The affinity between the service type, its host, and a synchronization context is locked in when the host is opened. If the thread opening the host has a synchronization context and **UseSynchronizationContext** is **true**, WCF will establish an affinity between that synchronization context and all instances of the service hosted by that host. WCF will automatically marshal all incoming calls to the service's synchronization context. All the thread-specific information stored in the TLS, such as the client's transaction or the security information (discussed in Chapter 10), will be marshaled correctly to the synchronization context.

If **UseSynchronizationContext** is **false**, regardless of any synchronization context the opening thread might have, the service will have no affinity to any synchronization context. Likewise, even if **UseSynchronizationContext** is **true**, if the opening thread has no synchronization context the service will not have one either.

The default value of **UseSynchronizationContext** is **true**, so these definitions are equivalent:

```
[ServiceContract]
interface IMyContract
{...}

class MyService : IMyContract
{...}
[ServiceBehavior(UseSynchronizationContext = true)]
class MyService : IMyContract
{...}
```

Hosting on the UI Thread

Again, I will use the UI thread affinity as a model for demonstrating the way WCF is integrated with synchronization context, but as you will see later on, the discussion here is just as relevant with more powerful examples, especially those involving custom synchronization contexts.

You can use the `UseSynchronizationContext` property to enable the service to update user interface controls and windows directly, without resorting to techniques such as those illustrated in Examples 8-6 and 8-7. WCF greatly simplifies UI updates by providing an affinity between all service instances from a particular host and a specific UI thread. To achieve that end, host the service on the UI thread that also creates the windows or controls with which the service needs to interact. Since the Windows Forms synchronization context is established during the instantiation of the base window, you need to open the host before that. For example, this sequence from Example 8-6:

```
ServiceHost host = new ServiceHost(typeof(MyService));
host.Open();

Application.Run(new MyForm());
```

will not have the host associate itself with the form synchronization context, since the host is opened before the form is created.

However, this minute change in the order of the lines of instantiation will achieve the desired effect:

```
Form form = new MyForm();

ServiceHost host = new ServiceHost(typeof(MyService));
host.Open();

Application.Run(form);
```

Although this change has no apparent effect in classic .NET, it is actually monumental for WCF, since now the thread that opened the host does have a synchronization context, and the host will use it for all calls to the service. The problem with this approach is that it is fragile—most developers maintaining your code will not be

aware that simply rearranging the same independent lines of code will have this effect. It is also wrong to design the form and the service that needs to update it so that they are both at the mercy of the `Main()` method and the hosting code to such a degree.

The simple solution is to have the window (or the thread-sensitive resource) that the service needs to interact with be the one that opens the host, as shown in [Example 8-9](#).

Example 8-9. The form hosting the service

```
class MyService : IMyContract
{...}

partial class HostForm : Form
{
    ServiceHost m_Host;
    Label m_CounterLabel;

    public HostForm()
    {
        InitializeComponent();

        m_Host = new ServiceHost(typeof(MyService));

        m_Host.Open();
    }
    void OnFormClosed(object sender,EventArgs e)
    {
        m_Host.Close();
    }

    public int Counter
    {
        get
        {
            return Convert.ToInt32(m_CounterLabel.Text);
        }
        set
        {
            m_CounterLabel.Text = value.ToString();
        }
    }
}
static class Program
{
    static void Main()
    {
        Application.Run(new HostForm());
    }
}
```

The service in [Example 8-9](#) defaults to using whichever synchronization context its host encounters. The form `HostForm` stores the service host in a member variable so that the form can close the service when the form is closed. The constructor of `HostForm` already has a synchronization context, so when it opens the host, an affinity to that synchronization context is established.

Accessing the form

EvenAlthough the form hosts the service in [Example 8-9](#), the service instances must have some proprietary application-specific mechanism to reach into the form. If a service instance needs to update multiple forms, you can use the `Application.OpenForms` collections (as in [Example 8-6](#)) to find the correct form. Once the service has the form, it can freely access it directly, as opposed to the code in [Example 8-6](#), which required marshaling:

```
class MyService : IFormManager
{
    public void IncrementLabel()
    {
        HostForm form = Application.OpenForms[0] as HostForm;
        Debug.Assert(form != null);
        form.Counter++;
    }
}
```

You could also store references to the forms to use in static variables, but the problem with such global variables is that if multiple UI threads are used to pump messages to different instances of the same form type, you cannot use a single static variable for each form type—you need a static variable for each thread used, which complicates things significantly.

Instead, the form (or forms) can store a reference to itself in the TLS, and have the service instance access that store and obtain the reference. However, using the TLS is a cumbersome and non-type-safe programming model. An improvement on this approach is to use thread-relative static variables. By default, static variables are visible to all threads in an app domain. With thread-relative static variables, each thread in the app domain gets its own copy of the static variable. You use the `ThreadStaticAttribute` to mark a static variable as thread-relative. Thread-relative static variables are always thread-safe because they can be accessed only by a single thread and because each thread gets its own copy of the static variable. Thread-relative static variables are stored in the TLS, yet they provide a type-safe, simplified programming model. [Example 8-10](#) demonstrates this technique.

Example 8-10. Storing form reference in a thread-relative static variable

```
partial class HostForm : Form
{
    Label m_CounterLabel;
    ServiceHost m_Host;

    [ThreadStatic]
    static HostForm m_CurrentForm;

    public static HostForm CurrentForm
    {
        get
        {
            return m_CurrentForm;
        }
        private set
        {
            m_CurrentForm = value;
        }
    }
    public int Counter
    {
        get
        {
            return Convert.ToInt32(m_CounterLabel.Text);
        }
        set
        {
            m_CounterLabel.Text = value.ToString();
        }
    }
    public HostForm()
    {
        InitializeComponent();

        CurrentForm = this;

        m_Host = new ServiceHost(typeof(MyService));
        m_Host.Open();
    }
    void OnFormClosed(object sender,EventArgs e)
    {
        m_Host.Close();
    }
}
[ServiceContract]
interface IFormManager
{
    [OperationContract]
    void IncrementLabel();
}
```

```

class MyService : IFormManager
{
    public void IncrementLabel()
    {
        HostForm form = HostForm.CurrentForm;
        form.Counter++;
    }
}
static class Program
{
    static void Main()
    {
        Application.Run(new HostForm());
    }
}

```

The form `HostForm` stores a reference to itself in a thread-relative static variable called `m_CurrentForm`. The service accesses the static property `CurrentForm` and obtains a reference to the instance of `HostForm` on that UI thread.

Multiple UI threads

Your service host process can actually have multiple UI threads, each pumping messages to its own set of windows. Such a setup is usually required with UI-intensive applications that want to avoid having multiple windows sharing a single UI thread and hosting the services, because while the UI thread is processing a service call (or a complicated UI update), not all of the windows will be responsive. Since the service synchronization context is established per host, if you have multiple UI threads you will need to open a service host instance for the same service type on each UI thread. Each service host will therefore have a different synchronization context for its service instances. As mentioned in [Chapter 1](#), in order to have multiple hosts for the same service type, you must provide each host with a different base address. The easiest way of doing that is to provide the form constructor with the base address to use as a construction parameter. I also recommend in such a case to use base address-relative addresses for the service endpoints. The clients will still invoke calls on the various service endpoints, yet each endpoint will now correspond to a different host, according to the base address schema and the binding used. [Example 8-11](#) demonstrates this configuration.

Example 8-11. Hosting on multiple UI threads

```

partial class HostForm : Form
{
    public HostForm(string baseAddress)
    {
        InitializeComponent();
    }
}

```

```

        CurrentForm = this;

        m_Host = new ServiceHost(typeof(MyService), new Uri(baseAddress));
        m_Host.Open();
    }
    //Rest same as Example 8-10
}
static class Program
{
    static void Main()
    {
        ParameterizedThreadStart threadMethod = (baseAddress)=>
        {
            string address = baseAddress as string;
            Application.Run(new HostForm(address));
        };
        Thread thread1 = new Thread(threadMethod);
        thread1.Start("http://localhost:8001/");

        Thread thread2 = new Thread(threadMethod);
        thread2.Start("http://localhost:8002/");
    }
}
/* MyService same as Example 8-10 */

/////////////////////// Host Config File ///////////////////
<services>
    <service name = "MyService">
        <endpoint
            address = "MyService"
            binding = "basicHttpBinding"
            contract = "IFormManager"
        />
    </service>
</services>
/////////////////////// Client Config File ///////////////////
<client>
    <endpoint name = "Form A"
        address = "http://localhost:8001/MyService/"
        binding = "basicHttpBinding"
        contract = "IFormManager"
    />
    <endpoint name = "Form B"
        Address = "http://localhost:8002/MyService/"
        binding = "basicHttpBinding"
        contract = "IFormManager"
    />
</client>

```

In [Example 8-11](#), the `Main()` method launches two UI threads, each with its own instance of `HostForm`. Each form instance accepts as a construction parameter a base address that it in turn provides for its own host instance. Once the host is opened, it

establishes an affinity to that UI thread's synchronization context. Calls from the client to the corresponding base address are now routed to the respective UI thread.

A Form as a Service

The main motivation for hosting a WCF service on a UI thread is when the service needs to update the UI or the form. The problem is, how does the service reach out and obtain a reference to the form? While the techniques and ideas shown in the examples so far certainly work, the separation between the service and the form is artificial. It would be simpler if the form were the service and hosted itself. For this to work, the form (or any window) must be a singleton service. The reason is that singleton is the only instancing mode that enables you to provide WCF with a live instance to host. In addition, it wouldn't be desirable to use a per-call form that exists only during a client call (which is usually very brief), or a sessionful form that only a single client can establish a session with and update. When a form is also a service, having that form as a singleton is the best instancing mode all around. [Example 8-12](#) lists just such a service.

Example 8-12. Form as a singleton service

```
[ServiceContract]
interface IFormManager
{
    [OperationContract]
    void IncrementLabel();
}

[ServiceBehavior(InstanceContextMode = InstanceContextMode.Single)]
partial class MyForm : Form, IFormManager
{
    Label m_CounterLabel;
    ServiceHost m_Host;

    public MyForm()
    {
        InitializeComponent();
        m_Host = new ServiceHost(this);
        m_Host.Open();
    }

    void OnFormClosed(object sender, EventArgs args)
    {
        m_Host.Close();
    }

    public void IncrementLabel()
    {
        Counter++;
    }

    public int Counter
    {
```

```

    get
    {
        return Convert.ToInt32(m_CounterLabel.Text);
    }
    set
    {
        m_CounterLabel.Text = value.ToString();
    }
}
}

```

`MyForm` implements the `IFormManager` contract and is configured as a WCF singleton service. `MyForm` has a `ServiceHost` as a member variable, as before. When `MyForm` constructs the host, it uses the host constructor that accepts an object reference, as shown in [Chapter 4](#). `MyForm` passes itself as the object. `MyForm` opens the host when the form is created and closes the host when the form is closed. Updating the form's controls as a result of client calls is done by accessing them directly, because the form, of course, runs in its own synchronization context.

The `FormHost<F>` class

You can streamline and automate the code in [Example 8-12](#) using my `FormHost<F>` class, defined as:

```

[ServiceBehavior(InstanceContextMode = InstanceContextMode.Single)]
public abstract class FormHost<F> : Form where F : Form
{
    public FormHost(params string[] baseAddresses);

    protected ServiceHost<F> Host
    {get;}
}

```

Using `FormHost<F>`, [Example 8-12](#) is reduced to:

```

partial class MyForm : FormHost<MyForm>, IFormManager
{
    Label m_CounterLabel;

    public MyForm()
    {
        InitializeComponent();
    }
    public void IncrementLabel()
    {
        Counter++;
    }
    public int Counter
    {
        get
        {

```

```

        return Convert.ToInt32(m_CounterLabel.Text);
    }
    set
    {
        m_CounterLabel.Text = value.ToString();
    }
}
}

```



The Windows Forms designer is incapable of rendering a form that has an abstract base class, let alone one that uses generics. You will have to change the base class to `Form` for visual editing, then revert to `FormHost<F>` for debugging. To compensate, copy the Debug configuration into a new solution configuration called Design, then add the `DESIGN` symbol to the Design configuration. Finally, define the form to render properly in design mode and to execute properly in debug and release modes:

```

#if DESIGN
public partial class MyForm : Form, IFormManager
#else
public partial class MyForm : FormHost<MyForm>, IFormManager
#endif
{...}

```

Example 8-13 shows the implementation of `FormHost<F>`.

Example 8-13. Implementing FormHost<F>

```

[ServiceBehavior(InstanceContextMode = InstanceContextMode.Single)]
public abstract class FormHost<F> : Form where F : Form
{
    protected ServiceHost<F> Host
    {get;private set;}

    public FormHost(params string[] baseAddresses)
    {
        Host = new ServiceHost<F>(this as F,baseAddresses);

        Load += delegate
        {
            if(Host.State == CommunicationState.Created)
            {
                Host.Open();
            }
        };
        FormClosed += delegate
        {
            if(Host.State == CommunicationState.Opened)
            {

```

```

        Host.Close();
    }
}
}

```

`FormHost<F>` is an abstract generic class configured as a singleton service. It takes a single type parameter, `F`, which is constrained to be a Windows Forms `Form` class. `FormHost<F>` uses my `ServiceHost<T>` as a member variable, specifying `F` for the type parameter for the host. `FormHost<F>` offers the host access to the derived forms, mostly for advanced configuration, so the `Host` property is marked as `protected`. The constructor of `FormHost<F>` creates the host, but does not open it. The reason is that the subform may want to perform some host initialization, such as configuring a throttle, and this initialization can only be done before opening the host. The subclass should place that initialization in its own constructor:

```

public MyForm()
{
    InitializeComponent();
    Host.SetThrottle(10,20,1);
}

```

To allow for this, the constructor uses an anonymous method to subscribe to the form's `Load` event, where it first verifies that the subform has not yet opened the host and then opens it. In a similar manner, the constructor subscribes to the form's `FormClosed` event, where it closes the host.

The UI Thread and Concurrency Management

Whenever you use hosting on the UI thread (or in any other case of a single-thread affinity synchronization context), deadlocks are possible. For example, the following setup is guaranteed to result in a deadlock: a Windows Forms application is hosting a service with `UseSynchronizationContext` set to `true`, and UI thread affinity is established; the Windows Forms application then calls the service in-proc over one of its endpoints. The call to the service blocks the UI thread, while WCF posts a message to the UI thread to invoke the service. That message is never processed due to the blocking UI thread—hence the deadlock.

Another possible case for a deadlock occurs when a Windows Forms application is hosting a service with `UseSynchronizationContext` set to `true` and UI thread affinity established. The service receives a call from a remote client, which is marshaled to the UI thread and eventually executed on that thread. If the service is allowed to call out to another service, that may result in a deadlock if the callout causality tries somehow to update the UI or call back to the service's endpoint, since all service instances associated with any endpoint (regardless of the service instancing mode) share the same UI thread. Similarly, you risk a deadlock if the service is configured for reentrancy

and it calls back to its client: a deadlock will occur if the callback causality tries to update the UI or enter the service, since that reentrance must be marshaled to the blocked UI thread.

UI responsiveness

Every client call to a service hosted on the UI thread is converted to a Windows message and is eventually executed on the UI thread—the same thread that is responsible for updating the UI and for continuing to respond to user input, as well as updating the user about the state of the application. While the UI thread is processing the service call, it does not process UI messages. Consequently, you should avoid lengthy processing in the service operation, because that can severely degrade the UI's responsiveness. You can alleviate this somewhat by pumping Windows messages in the service operation, either by explicitly calling the static method `Application.DoEvents()` to process all the queued-up Windows messages or by using a method such as `MessageBox.Show()` that pumps some but not all of the queued messages. The downside of trying to refresh the UI this way is that it may dispatch queued client calls to the service instance and may cause unwanted reentrancy or a deadlock.

To make things even worse, what if clients dispatch a number of calls to the service all at once? Depending on the service concurrency mode (discussed next) even if those service calls are of short duration, the calls will all be queued back-to-back in the Windows message queue, and processing them in order might take time—and all the while, the UI will not be updated.

Whenever you're hosting on a UI thread, carefully examine the calls' duration and frequency to see whether the resulting degradation in UI responsiveness is acceptable. What is acceptable may be application-specific, but as a rule of thumb, most users will not mind a UI latency of less than half a second, will notice a delay of more than three quarters of a second, and will be annoyed if the delay is more than a second. If that is the case, consider hosting parts of the UI (and the associated services) on multiple UI threads, as explained previously. By having multiple UI threads you maximize responsiveness, because while one thread is busy servicing a client call, the rest can still update their windows and controls. If using multiple UI threads is impossible in your application and processing service calls introduces unacceptable UI responsiveness, examine what the service operations do and what is causing the latency. Typically, the latency would be caused not by the UI updates but rather by performing lengthy operations, such as calling other services, or computational-intensive operations, such as image processing. Because the service is hosted on the UI thread, WCF performs all of that work on the UI thread, not just the critical part that interacts with the UI directly. If that is indeed your situation, disallow the affinity to the UI thread altogether by setting `UseSynchronizationContext` to `false`:

```
[ServiceBehavior(UseSynchronizationContext = false)]
class MyService : IMyContract
```

```

{
    public void MyMethod()
    {
        Debug.Assert(Application.MessageLoop == false);
        //Rest of the implementation
    }
}

```

(You can even assert that the thread executing the service call does not have a message loop.) Perform the lengthy operations on the incoming worker thread, and use safe controls (such as `SafeLabel`) to marshal the calls to the UI thread only when required, as opposed to all the time. The downside of this approach is that it is an expert programming model: the service cannot be the window or form itself (by relying on the simplicity of `FormHost<F>`), so you need a way of binding to the form, and the service developer has to work together with the UI developers to ensure they use the safe controls or provide access to the form's synchronization context.

The UI thread and concurrency modes

A service with a UI thread affinity is inherently thread-safe because only that UI thread can ever call its instances. Since only a single thread (and the same thread, at that) can ever access an instance, that instance is by definition thread-safe. The service is single-threaded anyway, so configuring the service with `ConcurrencyMode.Single` adds no safety. When you configure with `ConcurrencyMode.Single`, concurrent client calls are first queued up by the instance lock and then dispatched to the service's message loop one at a time, in order. These client calls are therefore given the opportunity of being interleaved with other UI Windows messages. `ConcurrencyMode.Single` thus yields the best responsiveness, because the UI thread will alternate between processing client calls and user interactions. When you configure the service with `ConcurrencyMode.Multiple`, client calls are dispatched to the service message loop as soon as they arrive off the channel and are invoked in order. The problem is that this mode allows the possibility of a batch of client calls arriving either back-to-back or in close proximity to one another in the Windows message queue, and while the UI thread processes that batch, the UI will be unresponsive. Consequently, `ConcurrencyMode.Multiple` is the worst option for UI responsiveness. When configured with `ConcurrencyMode.Reentrant`, the service is not reentrant at all, and deadlocks are still possible, as explained at the beginning of this section. Clearly, the best practice with UI thread affinity is to configure the service with `ConcurrencyMode.Single`. Avoid `ConcurrencyMode.Multiple` due to its detrimental effect on responsiveness and `ConcurrencyMode.Reentrant` due to its unfulfilled ability.

Custom Service Synchronization Contexts

While a synchronization context is a general-purpose pattern out of the box, .NET only implements a handful of them, with the two useful ones being the Windows Forms synchronization context and the WPF synchronization context (there is also the default implementation that uses the .NET thread pool). As it turns out, the ability to automatically marshal calls to a custom synchronization context is one of the most powerful extensibility mechanisms in WCF.

The Thread Pool Synchronizer

There are two aspects to developing a custom service synchronization context: the first is implementing a custom synchronization context, and the second is installing it or even applying it declaratively on the service. *ServiceModelEx* contains the `ThreadPoolSynchronizer` class, defined as:

```
public class ThreadPoolSynchronizer : SynchronizationContext, IDisposable
{
    public ThreadPoolSynchronizer(uint poolSize);
    public ThreadPoolSynchronizer(uint poolSize, string poolName);

    public void Dispose();
    public void Close();
    public void Abort();

    protected Semaphore CallQueued
    {get;}
}
```

Implementing a custom synchronization context has nothing to do with WCF and is therefore not discussed in this book, although the implementation code is available with *ServiceModelEx*.

`ThreadPoolSynchronizer` marshals all calls to a custom thread pool, where the calls are first queued up, then multiplexed on the available threads. The size of the pool is provided as a construction parameter. If the pool is maxed out, any calls that come in will remain pending in the queue until a thread is available.

You can also provide a pool name (which will be the prefix of the name of each of the threads in the pool). Disposing of or closing the `ThreadPoolSynchronizer` kills all threads in the pool gracefully; that is, the `ThreadPoolSynchronizer` waits for the engaged threads to complete their tasks. The `Abort()` method is an ungraceful shutdown, as it terminates all threads abruptly.

The classic use for a custom thread pool is with a server application (such as a web server or an email server) that needs to maximize its throughput by controlling the underlying worker threads and their assignment. However, such usage is rare, since

most application developers do not write servers anymore. The real use of `ThreadPoolSynchronizer` is as a stepping-stone to implement other synchronization contexts, which are useful in their own right.

To associate your service with the custom thread pool, you can manually attach `ThreadPoolSynchronizer` to the thread opening the host using the static `SetSynchronizationContext()` method of `SynchronizationContext`, as shown in [Example 8-14](#).

Example 8-14. Using ThreadPoolSynchronizer

```
SynchronizationContext syncContext = new ThreadPoolSynchronizer(3);

SynchronizationContext.SetSynchronizationContext(syncContext);

using(syncContext as IDisposable)
{
    ServiceHost host = new ServiceHost(typeof(MyService));
    host.Open();
    /* Some blocking operations */

    host.Close();
}
```

In [Example 8-14](#), the thread pool will have three threads. The service `MyService` will have an affinity to those three threads, and all calls to the service will be channeled to them, regardless of the service concurrency mode or instancing mode, and across all endpoints and contracts supported by the service. After closing the host, the example disposes of `ThreadPoolSynchronizer` to shut down the threads in the pool.

Note that a service executing in a custom thread pool is not thread-safe (unless the pool size is 1), so the preceding discussion of concurrency management still applies. The only difference is that now you control the threads.

Declaratively attaching a custom synchronization context

The problem with [Example 8-14](#) is that the service is at the mercy of the hosting code. If by design the service is required to execute in the pool, it would be better to apply the thread pool declaratively, as part of the service definition.

To that end, I wrote the `ThreadPoolBehaviorAttribute`:

```
[AttributeUsage(AttributeTargets.Class)]
public class ThreadPoolBehaviorAttribute : Attribute,
                                         IContractBehavior, IServiceBehavior
{
    public ThreadPoolBehaviorAttribute(uint poolSize, Type serviceType);
    public ThreadPoolBehaviorAttribute(uint poolSize, Type serviceType,
                                      string poolName);
}
```

You apply this attribute directly on the service, while providing the service type as a constructor parameter:

```
[ThreadPoolBehavior(3,typeof(MyService))]  
class MyService : IMyContract  
{...}
```

The attribute provides an instance of `ThreadPoolSynchronizer` to the dispatchers of the service's endpoints. The key in implementing the `ThreadPoolBehavior` attribute is knowing how and when to hook up the dispatchers with the synchronization context. The `ThreadPoolBehavior` attribute supports the special WCF extensibility interface `IContractBehavior`, introduced in [Chapter 5](#):

```
public interface IContractBehavior  
{  
    void ApplyDispatchBehavior(ContractDescription description,  
                               ServiceEndpoint endpoint,  
                               DispatchRuntime dispatchRuntime);  
    //More members  
}
```

When a service is decorated with an attribute that supports `IContractBehavior`, after opening the host (but before forwarding calls to the service), for each service endpoint WCF calls the `ApplyDispatchBehavior()` method and provides it with the `DispatchRuntime` parameter, allowing you to affect an individual endpoint dispatcher's runtime and set its synchronization context. Each endpoint has its own dispatcher, and each dispatcher has its own synchronization context, so the attribute is instantiated and `ApplyDispatchBehavior()` is called for each endpoint.

[Example 8-15](#) lists most of the implementation of `ThreadPoolBehaviorAttribute`.

Example 8-15. Implementing ThreadPoolBehaviorAttribute

```
[AttributeUsage(AttributeTargets.Class)]  
public class ThreadPoolBehaviorAttribute : Attribute,IContractBehavior,  
                                         IServiceBehavior  
{  
    protected string PoolName  
    {get;set;}  
    protected uint PoolSize  
    {get;set;}  
    protected Type ServiceType  
    {get;set;}  
  
    public ThreadPoolBehaviorAttribute(uint poolSize,Type serviceType) :  
                                         this(poolSize,serviceType,null)  
    {}  
    public ThreadPoolBehaviorAttribute(uint poolSize,Type serviceType,  
                                         string poolName)  
    {
```

```

        PoolName    = poolName;
        ServiceType = serviceType;
        PoolSize    = poolSize;
    }
    protected virtual ThreadPoolSynchronizer ProvideSynchronizer()
    {
        if(ThreadPoolHelper.HasSynchronizer(ServiceType) == false)
        {
            return new ThreadPoolSynchronizer(PoolSize,PoolName);
        }
        else
        {
            return ThreadPoolHelper.GetSynchronizer(ServiceType);
        }
    }

    void IContractBehavior.ApplyDispatchBehavior(ContractDescription description,
                                                ServiceEndpoint endpoint,
                                                DispatchRuntime dispatchRuntime)
    {
        PoolName = PoolName ?? "Pool executing endpoints of " + ServiceType;

        lock(typeof(ThreadPoolHelper))
        {
            ThreadPoolHelper.ApplyDispatchBehavior(ProvideSynchronizer(),
                                                    PoolSize,ServiceType,PoolName,dispatchRuntime);
        }
    }
    void IServiceBehavior.Validate(ServiceDescription description,
                                   ServiceHostBase serviceHostBase)
    {
        serviceHostBase.Closed += delegate
        {
            ThreadPoolHelper.CloseThreads(ServiceType);
        };
    }
    //Rest of the implementation
}
public static class ThreadPoolHelper
{
    static Dictionary<Type,ThreadPoolSynchronizer> m_Synchronizers =
        new Dictionary<Type,ThreadPoolSynchronizer>();

    [MethodImpl(MethodImplOptions.Synchronized)]
    internal static bool HasSynchronizer(Type type)
    {
        return m_Synchronizers.ContainsKey(type);
    }

    [MethodImpl(MethodImplOptions.Synchronized)]
    internal static ThreadPoolSynchronizer GetSynchronizer(Type type)
    {

```

```

        return m_Synchronizers[type];
    }
    [MethodImpl(MethodImplOptions.Synchronized)]
    internal static void ApplyDispatchBehavior(ThreadPoolSynchronizer synchronizer,
                                                uint poolSize, Type type,
                                                string poolName,
                                                DispatchRuntime dispatchRuntime)
    {
        if(HasSynchronizer(type) == false)
        {
            m_Synchronizers[type] = synchronizer;
        }
        dispatchRuntime.SynchronizationContext = m_Synchronizers[type];
    }
    [MethodImpl(MethodImplOptions.Synchronized)]
    public static void CloseThreads(Type type)
    {
        if(HasSynchronizer(type))
        {
            m_Synchronizers[type].Dispose();
            m_Synchronizers.Remove(type);
        }
    }
}

```

The constructors of the `ThreadPoolBehavior` attribute save the provided service type and pool name. The name is simply passed to the constructor of `ThreadPoolSynchronizer`.

It is a best practice to separate the implementation of a WCF custom behavior attribute from the actual behavior: let the attribute merely decide on the sequence of events, and have a helper class provide the actual behavior. Doing so enables the behavior to be used separately (for example, by a custom host). This is why the `ThreadPoolBehavior` attribute does not do much. It delegates most of its work to a static helper class called `ThreadPoolHelper`. `ThreadPoolHelper` provides the `HasSynchronizer()` method, which indicates whether the specified service type already has a synchronization context, and the `GetSynchronizer()` method, which returns the synchronization context associated with the type. The `ThreadPoolBehavior` attribute uses these two methods in the virtual `ProvideSynchronizer()` method to ensure that it creates the pool exactly once per service type. This check is required because `ApplyDispatchBehavior()` may be called multiple times (once per endpoint). The `ThreadPoolBehavior` attribute is also a custom service behavior, because it implements `IServiceBehavior`. The `Validate()` method of `IServiceBehavior` provides the service host instance the `ThreadPoolBehavior` attribute uses to subscribe to the host's `Closed` event, where it asks `ThreadPoolHelper` to terminate all the threads in the pool by calling `ThreadPoolHelper.CloseThreads()`.

`ThreadPoolHelper` associates all dispatchers of all endpoints of that service type with the same instance of `ThreadPoolSynchronizer`. This ensures that all calls are routed to the same pool. `ThreadPoolHelper` has to be able to map a service type to a particular `ThreadPoolSynchronizer`, so it declares `m_Synchronizers`, a static dictionary that uses service types as keys and `ThreadPoolSynchronizer` instances as values.

In `ApplyDispatchBehavior()`, `ThreadPoolHelper` checks to see whether `m_Synchronizers` already contains the provided service type. If the type is not found, `ThreadPoolHelper` adds the provided `ThreadPoolSynchronizer` to `m_Synchronizers`, associating it with the service type.

The `DispatchRuntime` class provides the `SynchronizationContext` property `ThreadPoolHelper` uses to assign a synchronization context for the dispatcher:

```
public sealed class DispatchRuntime
{
    public SynchronizationContext SynchronizationContext
    {get;set;}
    //More members
}
```

Before making the assignment, `ThreadPoolHelper` verifies that the dispatcher has no other synchronization context, since that would indicate some unresolved conflict. After that, it simply assigns the `ThreadPoolSynchronizer` instance to the dispatcher:

```
dispatchRuntime.SynchronizationContext = m_Synchronizers[type];
```

This single line is all that is required to have WCF use the custom synchronization context from now on. In the `CloseThreads()` method, `ThreadPoolHelper` looks up the `ThreadPoolSynchronizer` instance in the dictionary and disposes of it (thus gracefully terminating all the worker threads in the pool). `ThreadPoolHelper` also verifies that the provided pool size value does not exceed the maximum concurrent calls value of the dispatcher's throttle (this is not shown in [Example 8-15](#)).

Replacing the I/O Completion Thread Pool

As explained in [Chapter 4](#), WCF provides you the option to throttle client connections to avoid maxing out server resources. .NET also provides management of thread consumption within .NET processes of both its worker and I/O completion thread pools. Normally, you do not need to worry about configuring these thresholds. .NET takes care of setting them appropriately when your process starts based on system factors such as the number of available CPUs. However, there are scenarios involving the I/O completion thread pool where a spike in client requests may cause a significant degradation in the responsiveness of your services, even though the number of requests has not yet reached your WCF throttling values. Here, you can apply a

custom synchronization context to move your services to a custom thread pool to improve their responsiveness.

Both the Windows I/O completion thread pool and the .NET thread pool offer a minimum thread threshold. For both thread pools, Windows will create new threads on demand in response to incoming requests until the number of active threads within the pool reaches the minimum thread threshold. Once a thread pool reaches its minimum thread threshold it no longer creates new threads on demand. Instead, each thread pool uses its own algorithm to determine how it will process new requests. The .NET thread pool uses a balanced approach between creating new threads and waiting for existing threads to become available. The I/O completion thread pool uses a fixed approach for new thread creation. For this pool, once reaching its minimum thread threshold the system will only add a new thread to the pool every 500 milliseconds.

Under nominal load conditions, you may never notice this delay in thread creation impacting the overall responsiveness of your service. If the rate of new requests rises moderately, the I/O completion thread pool can keep pace with the rise. But if your service receives a sudden spike in new requests (though still well below your WCF throttling values) after the I/O completion thread pool has reached its minimum thread threshold, the delay in creating new threads will cause a slow, linear rise in operation execution times until the demand is finally met. As a result, if operation execution times become too long clients will begin to experience a very low level of service and even incur call timeout errors.

Compounding the problem, there is a well-documented bug within the I/O completion thread pool that causes it to enforce its growth algorithm well before reaching its minimum thread threshold. The problem becomes most acute if your service is already performing a number of lengthy operations utilizing as many threads as the pool's minimum thread threshold. In this case, your system will exhibit the problem almost immediately. Worst of all, once the I/O completion thread pool enters this state a process reset is your only option to correct the problem.

To address the problem, you need to create and install a custom synchronization context that moves your service's workload off of the I/O completion thread pool to another thread pool. If you want to discriminately apply the fix only to services that are experiencing the problem, you can use the `ThreadPoolSynchronizer` and apply the `ThreadPoolBehaviorAttribute` to those services directly:

```
[ThreadPoolBehavior(123,typeof(MyService))]  
class MyService : IMyContract  
{...}
```

You will need to do some performance tuning to decide on the appropriate pool size.

Alternatively, you can use the technique from [Example 8-14](#) and install `ThreadPoolSynchronizer` at host process startup as the current synchronization context before opening all the service hosts for your services. In this case, all services will have an affinity to the thread pool provided by a single instance of `ThreadPoolSynchronizer`. If you need finer control over the worker threads used by WCF, you can derive a custom synchronization context from `ThreadPoolSynchronizer` and modify its thread pool management.

If you do not need control over the threads WCF uses you can provide your service with a thread pool that is elastic and responds appropriately to both increases and decreases in request demand. Luckily, .NET already provides such a thread pool in the form of the .NET thread pool. You can use the existing .NET `SynchronizationContext` class to switch your WCF service to the .NET thread pool. You can then install `SynchronizationContext` for all your services at host startup using the technique from [Example 8-14](#) as mentioned earlier.

Thread Affinity

A pool size of 1 will in effect create an affinity between a particular thread and all service calls, regardless of the service's concurrency and instancing modes. This is particularly useful if the service is required not merely to update some UI but to also create a UI (for example, creating a pop-up window and then periodically showing, hiding, and updating it). Having created the window, the service must ensure that the creating thread is used to access and update it. Thread affinity is also required for a service that accesses or creates resources that use the TLS. To formalize such requirements, I created the specialized `AffinitySynchronizer` class, implemented as:

```
public class AffinitySynchronizer : ThreadPoolSynchronizer
{
    public AffinitySynchronizer() : this("AffinitySynchronizer Worker Thread")
    {}
    public AffinitySynchronizer(string threadName) : base(1,threadName)
    {}
}
```

While you can install `AffinitySynchronizer`, as shown in [Example 8-14](#), if by design the service is required to always execute on the same thread it is better not to be at the mercy of the host and the thread that happens to open it. Instead, use my `ThreadAffinityBehaviorAttribute`:

```
[ThreadAffinityBehavior(typeof(MyService))]
class MyService : IMyContract
{...}
```

`ThreadAffinityBehaviorAttribute` is a specialization of `ThreadPoolBehaviorAttribute` that hardcodes the pool size as 1, as shown in [Example 8-16](#).

Example 8-16. Implementing ThreadAffinityBehaviorAttribute

```
[AttributeUsage(AttributeTargets.Class)]
public class ThreadAffinityBehaviorAttribute : ThreadPoolBehaviorAttribute
{
    public ThreadAffinityBehaviorAttribute(Type serviceType) :
        this(serviceType, "Affinity Worker Thread")
    {}

    public ThreadAffinityBehaviorAttribute(Type serviceType, string threadName) :
        base(1, serviceType, threadName)
    {}
}
```

When relying on thread affinity, all service instances are always thread-safe, since only a single thread (and the same thread, at that) can access them.

When the service is configured with `ConcurrencyMode.Single`, it gains no additional thread safety because the service instance is single-threaded anyway. You do get double queuing of concurrent calls, though: all concurrent calls to the service are first queued in the lock's queue and then dispatched to the single thread in the pool one at a time. With `ConcurrencyMode.Multiple`, calls are dispatched to the single thread as fast as they arrive and are then queued up to be invoked later, in order and never concurrently. That may reduce the code readability. Finally, with `ConcurrencyMode.Renentrant`, the service is, of course, not reentrant, because the incoming reentering call will be queued up and a deadlock will occur while the single thread is blocked on the callout. It is therefore best to use the default of `ConcurrencyMode.Single` when relying on thread affinity.

The host-installed synchronization context

If the affinity to a particular synchronization context is a host decision, you can streamline the code in [Example 8-14](#) by encapsulating the installation of the synchronization context with extension methods. For example, the use of thread affinity is such a socialized case; you could define the following extension methods:

```
public static class HostThreadAffinity
{
    public static void SetThreadAffinity(this ServiceHost host, string threadName);
    public static void SetThreadAffinity(this ServiceHost host);
}
```

`SetThreadAffinity()` works equally well on `ServiceHost` and my `ServiceHost<T>`:

```
ServiceHost<MyService> host = new ServiceHost<MyService>();
host.SetThreadAffinity();

host.Open();
```

[Example 8-17](#) lists the implementation of the `SetThreadAffinity()` methods.

Example 8-17. Adding thread affinity to the host

```
public static class HostThreadAffinity
{
    public static void SetThreadAffinity(this ServiceHost host, string threadName)
    {
        if(host.State == CommunicationState.Opened)      {
            throw new InvalidOperationException("Host is already opened");
        }

        AffinitySynchronizer affinitySynchronizer =
            new AffinitySynchronizer(threadName);

        SynchronizationContext.SetSynchronizationContext(affinitySynchronizer);

        host.Closing += delegate
        {
            using(affinitySynchronizer);
        };
    }
    public static void SetThreadAffinity(this ServiceHost host)
    {
        SetThreadAffinity(host, "Executing all endpoints of " +
                           host.Description.ServiceType);
    }
}
```

`HostThreadAffinity` offers two versions of `SetThreadAffinity()`: the parameterized version takes the thread name to provide for `AffinitySynchronizer`'s worker thread, while the parameterless version calls the other `SetThreadAffinity()` method, specifying a thread name inferred from the hosted service type (such as “Executing all endpoints of `MyService`”). `SetThreadAffinity()` first checks that the host has not yet been opened, because you can only attach a synchronization context before the host is opened. If the host has not been opened, `SetThreadAffinity()` constructs a new `AffinitySynchronizer`, providing it with the thread name to use, and attaches it to the current thread. Finally, `SetThreadAffinity()` subscribes to the host's `Closing` event in order to call `Dispose()` on the `AffinitySynchronizer`, to shut down its worker thread. Since the `AffinitySynchronizer` member can be `null` if no one calls `SetThreadAffinity()`, `OnClosing()` uses the `using` statement, which internally checks for `null` assignment before calling `Dispose()`.

Priority Processing

By default, all calls to your WCF service will be processed in the order in which they arrive. This is true both if you use the I/O completion thread pool or a custom thread pool. Normally, this is exactly what you want. But what if some calls have higher priority and you want to process them as soon as they arrive, rather than in order? Even

worse, when such calls arrive, what if the load on your service is such that the underlying service resources are exhausted? What if the throttle is maxed out? In these cases, your higher-priority calls will be queued just like all the other calls, waiting for the service or its resources to become available.¹

Synchronization contexts offer an elegant solution to this problem: you can assign a priority to each call and have the synchronization context sort the calls as they arrive before dispatching them to the thread pool for execution. This is exactly what my `PrioritySynchronizer` class does:

```
public enum CallPriority
{
    Low,
    Normal,
    High
}
public class PrioritySynchronizer : ThreadPoolSynchronizer
{
    public PrioritySynchronizer(uint poolSize);
    public PrioritySynchronizer(uint poolSize, string poolName);

    public static CallPriority Priority
    {get;set;}
}
```

`PrioritySynchronizer` derives from `ThreadPoolSynchronizer` and adds the sorting just mentioned. Since the `Send()` and `Post()` methods of `SynchronizationContext` do not take a priority parameter, the client of `PrioritySynchronizer` has two ways of passing the priority of the call: via the `Priority` property, which stores the priority (a value of the enum type `CallPriority`) in the TLS of the calling thread, or via the message headers. If unspecified, `Priority` defaults to `CallPriority.Normal`.

In addition to the `PrioritySynchronizer` class, I also provide the matching `PriorityCallsBehaviorAttribute`, shown in [Example 8-18](#).

Example 8-18. Implementing PriorityCallsBehaviorAttribute

```
[AttributeUsage(AttributeTargets.Class)]
public class PriorityCallsBehaviorAttribute : ThreadPoolBehaviorAttribute
{
    public PriorityCallsBehaviorAttribute(uint poolSize, Type serviceType) :
        this(poolSize, serviceType, null)
    {}
    public PriorityCallsBehaviorAttribute(uint poolSize, Type serviceType,
        string poolName) : base(poolSize, serviceType, poolName)
```

¹ I first presented my technique for priority processing of WCF calls in my article “Synchronization Contexts in WCF” (*MSDN Magazine*, November 2007).

```

    {}
protected override ThreadPoolSynchronizer ProvideSynchronizer()
{
    if(ThreadPoolHelper.HasSynchronizer(ServiceType) == false)
    {
        return new PrioritySynchronizer(PoolSize,PoolName);
    }
    else
    {
        return ThreadPoolHelper.GetSynchronizer(ServiceType);
    }
}
}

```

Using the `PriorityCallsBehavior` attribute is straightforward:

```

[PriorityCallsBehavior(3,typeof(MyService))]
class MyService : IMyContract
{...}

```

`PriorityCallsBehaviorAttribute` overrides `ProvideSynchronizer()` and provides an instance of `PrioritySynchronizer` instead of `ThreadPoolSynchronizer`. Because `PrioritySynchronizer` derives from `ThreadPoolSynchronizer`, this is transparent as far as `ThreadPoolHelper` is concerned.

The real challenge in implementing and supporting priority processing is providing the call priority from the client to the service, and ultimately to `PrioritySynchronizer`. Using the `Priority` property of `PrioritySynchronizer` is useful only for non-WCF clients that interact directly with the synchronization context; it is of no use for a WCF client, whose thread is never used to access the service. While you could provide the priority as an explicit parameter in every method, I wanted a generic mechanism that can be applied on any contract and service. To achieve that goal you have to pass the priority of the call out-of-band, via the message headers, using the techniques described in [Appendix B](#) (which explains in detail the use of the incoming and outgoing headers, including augmenting WCF with general-purpose management of extraneous information sent from the client to the service). In effect, I provide a generic yet type-safe and application-specific custom context via my `GenericContext<T>` class, available in `ServiceModelEx`:

```

[DataContract]
public class GenericContext<T>
{
    [DataMember]
    public readonly T Value;

    public GenericContext();
    public GenericContext(T value);
    public static GenericContext<T> Current
}

```

```
    {get;set;}  
}
```

Literally any data contract (or serializable) type can be used for the type parameter in the custom context, including of course the `CallPriority` enum.

On the service side, any party can read the value out of the custom headers:

```
CallPriority priority = GenericContext<CallPriority>.Current.Value;
```

This is exactly what `PrioritySynchronizer` does when looking for the call priority. It expects the client to provide the priority either in the TLS (via the `Priority` property) or in the form of a custom context that stores the priority in the message headers.

The client can use my `HeaderClientBase<T,H>` proxy class (also discussed in [Appendix B](#)) to pass the priority to the service in the message headers, or, even better, define a general-purpose priority-enabled proxy class, `PriorityClientBase<T>`, shown in [Example 8-19](#).

Example 8-19. Defining PriorityClientBase<T>

```
public abstract partial class PriorityClientBase<T> :  
    HeaderClientBase<T,CallPriority> where T : class  
{  
    public PriorityClientBase() : this(PrioritySynchronizer.Priority)  
    {}  
  
    public PriorityClientBase(string endpointName) :  
        this(PrioritySynchronizer.Priority,endpointName)  
    {}  
  
    public PriorityClientBase(Binding binding,EndpointAddress remoteAddress) :  
        this(PrioritySynchronizer.Priority,binding,remoteAddress)  
    {}  
  
    public PriorityClientBase(CallPriority priority) : base(priority)  
    {}  
  
    public PriorityClientBase(CallPriority priority,string endpointName) :  
        base(priority,endpointConfigurationName)  
    {}  
  
    public PriorityClientBase(CallPriority priority,Binding binding,  
        EndpointAddress remoteAddress) : base(priority,binding,remoteAddress)  
    {}  
    /* More constructors */  
}
```

`PriorityClientBase<T>` hardcodes the use of `CallPriority` for the type parameter `H`. `PriorityClientBase<T>` defaults to reading the priority from the TLS (yielding

`CallPriority.Normal` when no priority is found), so it can be used like any other proxy class. With very minor changes to your existing proxy classes, you can now add priority-processing support:

```
class MyContractClient : PriorityClientBase<IMyContract>, IMyContract
{
    //Reads priority from TLS
    public MyContractClient()
    {}

    public MyContractClient(CallPriority priority) : base(priority)
    {}
    public void MyMethod()
    {
        Channel.MyMethod();
    }
}

MyContractClient proxy = new MyContractClient(CallPriority.High);
proxy.MyMethod();
proxy.Close();
```

Callbacks and Client Safety

There are quite a few cases when a client might receive concurrent callbacks. For instance, if the client has provided a callback reference to multiple services, those services could call back to the client concurrently. Even if it has only provided a single callback reference, the service might launch multiple threads and use all of them to call on that single reference. Duplex callbacks enter the client on worker threads, and if they are processed concurrently without synchronization they might corrupt the client's state. The client must therefore synchronize access to its own in-memory state, as well as to any resources the callback thread might access. Similar to a service, a callback client can use either manual or declarative synchronization. The `CallbackBehavior` attribute introduced in [Chapter 6](#) offers the `ConcurrencyMode` and the `UseSynchronizationContext` properties:

```
[AttributeUsage(AttributeTargets.Class)]
public sealed class CallbackBehaviorAttribute : Attribute, ...
{
    public ConcurrencyMode ConcurrencyMode
    {get;set;}
    public bool UseSynchronizationContext
    {get;set;}
}
```

Both of these properties default to the same values as with the `ServiceBehavior` attribute and behave in a similar manner. For example, the default of the `ConcurrencyMode` property is `ConcurrencyMode.Single`, so these two definitions are equivalent:

```

class MyClient : IMyContractCallback
{...}

[CallbackBehavior(ConcurrencyMode = ConcurrencyMode.Single)]
class MyClient : IMyContractCallback
{...}

```

Callbacks with ConcurrencyMode.Single

When the callback class is configured with `ConcurrencyMode.Single` (the default), only one callback at a time is allowed to enter the callback object. The big difference, compared with a service, is that callback objects often have an existence independent of WCF. While the service instance is owned by WCF and only ever accessed by worker threads dispatched by WCF, a callback object may also interact with local client-side threads. It fact, it always interacts with at least one additional thread: the thread that called the service and provided the callback object. These client threads are unaware of the synchronization lock associated with the callback object when it is configured with `ConcurrencyMode.Single`. All that `ConcurrencyMode.Single` does for a callback object is serialize the access by WCF threads. You must therefore manually synchronize access to the callback state and any other resource accessed by the callback method, as shown in [Example 8-20](#).

Example 8-20. Manually synchronizing the callback with ConcurrencyMode.Single

```

interface IMyContractCallback
{
    [OperationContract]
    void OnCallback();
}

[ServiceContract(CallbackContract = typeof(IMyContractCallback))]
interface IMyContract
{
    [OperationContract]
    void MyMethod();
}

class MyClient : IMyContractCallback, IDisposable
{
    MyContractClient m_Proxy;

    public void CallService()
    {
        m_Proxy = new MyContractClient(new InstanceContext(this));
        m_Proxy.MyMethod();
    }
    //This method invoked by one callback at a time, plus client threads
    public void OnCallback()
    {
        //Access state and resources, synchronize manually
    }
}

```

```

        lock(this)
        {...}
    }
    public void Dispose()
    {
        m_Proxy.Close();
    }
}

```

Callbacks with ConcurrencyMode.Multiple

When you configure the callback class with `ConcurrencyMode.Multiple`, WCF will allow concurrent calls on the callback instance. This means you need to synchronize access in the callback operations, as shown in [Example 8-21](#), because they could be invoked concurrently both by WCF worker threads and by client-side threads.

Example 8-21. Manually synchronizing the callback with ConcurrencyMode.Multiple

```

[CallbackBehavior(ConcurrencyMode = ConcurrencyMode.Multiple)]
class MyClient : IMyContractCallback, IDisposable
{
    MyContractClient m_Proxy;

    public void CallService()
    {
        m_Proxy = new MyContractClient(new InstanceContext(this));
        m_Proxy.MyMethod();
    }
    //This method can be invoked concurrently by callbacks,
    //plus client threads
    public void OnCallback()
    {
        //Access state and resources, synchronize manually
        lock(this)
        {...}
    }
    public void Dispose()
    {
        m_Proxy.Close();
    }
}

```

Callbacks with ConcurrencyMode.Reentrant

The callback object can perform outgoing calls over WCF, and those calls may eventually try to reenter the callback object. To avoid the deadlock that would occur when using `ConcurrencyMode.Single`, you can configure the callback class with `ConcurrencyMode.Reentrant` as needed:

```
[CallbackBehavior(ConcurrencyMode = ConcurrencyMode.Reentrant)]
class MyClient : IMyContractCallback
{...}
```

Configuring the callback for reentrancy also enables other services to call it when the callback object itself is engaged in WCF callouts.

Callbacks and Synchronization Contexts

Like a service invocation, a callback may need to access resources that rely on some kind of thread(s) affinity. In addition, the callback instance itself may require thread affinity for its own use of the TLS, or for interacting with a UI thread. While the callback can use techniques such as those in Examples 8-4 and 8-5 to marshal the interaction to the resource synchronization context, you can also have WCF associate the callback with a particular synchronization context by setting the `UseSynchronizationContext` property to `true`. However, unlike the service, the client does not use any host to expose the endpoint. If the `UseSynchronizationContext` property is `true`, the synchronization context to use is locked in when the proxy is opened (or, more commonly, when the client makes the first call to the service using the proxy, if `Open()` is not explicitly called). If the client is using the channel factory, the synchronization context to use is locked in when the client calls `CreateChannel()`. If the calling client thread has a synchronization context, this will be the synchronization context used by WCF for all callbacks to the client's endpoint associated with that proxy. Note that only the first call made on the proxy (or the call to `Open()` or `CreateChannel()`) is given the opportunity to determine the synchronization context. Subsequent calls have no say in the matter. If the calling client thread has no synchronization context, even if `UseSynchronizationContext` is `true`, no synchronization context will be used for the callbacks.

Callbacks and the UI Synchronization Context

If the callback object is running in a Windows Forms synchronization context, or if it needs to update some UI, you must marshal the callbacks or the updates to the UI thread. You can use techniques such as those in Example 8-6 or Example 8-8. However, the more common use for UI updates over callbacks is to have the form itself implement the callback contract and update the UI, as in Example 8-22.

Example 8-22. Relying on the UI synchronization context for callbacks

```
partial class MyForm : Form,IMyContractCallback
{
    MyContractClient m_Proxy;

    public MyForm()
    {
```

```

        InitializeComponent();
        m_Proxy = new MyContractClient(new InstanceContext(this));
    }
    //Called as a result of a UI event
    public void OnCallService(object sender,EventArgs args)
    {
        m_Proxy.MyMethod(); //Affinity established here
    }
    //This method always runs on the UI thread
    public void OnCallback()
    {
        //No need for synchronization and marshaling
        Text = "Some Callback";
    }
    public void OnClose(object sender,EventArgs args)
    {
        m_Proxy.Close();
    }
}

```

In [Example 8-22](#) the proxy is first used in the `OnCallService()` method, which is called by the UI thread as a result of some UI event. Calling the proxy on the UI synchronization context establishes the affinity to it, so the callback can directly access and update the UI without marshaling any calls. In addition, since only one thread (and the same thread, at that) will ever execute in the synchronization context, the callback is guaranteed to be synchronized.

You can also explicitly establish the affinity to the UI synchronization context by opening the proxy in the form's constructor without invoking an operation. This is especially useful if you want to dispatch calls to the service on worker threads (or perhaps even asynchronously as discussed at the end of this chapter) and yet have the callbacks enter on the UI synchronization context, as shown in [Example 8-23](#).

Example 8-23. Explicitly opening a proxy to establish a synchronization context

```

partial class MyForm : Form,IMyContractCallback
{
    MyContractClient m_Proxy;

    public MyForm()
    {
        InitializeComponent();

        m_Proxy = new MyContractClient(new InstanceContext(this));

        //Establish affinity to UI synchronization context here:
        m_Proxy.Open();
    }
    //Called as a result of a UI event
    public void CallService(object sender,EventArgs args)

```

```

    {
        Thread thread = new Thread(()=>m_Proxy.MyMethod());
        thread.Start();
    }
    //This method always runs on the UI thread
    public void OnCallback()
    {
        //No need for synchronization and marshaling
        Text = "Some Callback";
    }
    public void OnClose(object sender,EventArgs args)
    {
        m_Proxy.Close();
    }
}

```

UI thread callbacks and responsiveness

When callbacks are being processed on the UI thread, the UI itself is not responsive. Even if you perform relatively short callbacks, you must bear in mind that if the callback class is configured with `ConcurrencyMode.Multiple` there may be multiple callbacks back-to-back in the UI message queue, and processing them all at once will degrade responsiveness. You should avoid lengthy callback processing on the UI thread, and opt for configuring the callback class with `ConcurrencyMode.Single` so that the callback lock will queue up the callbacks. They can then be dispatched to the callback object one at a time, giving them the chance of being interleaved among the UI messages.

UI thread callbacks and concurrency management

Configuring the callback for affinity to the UI thread may trigger a deadlock. Suppose a Windows Forms client establishes an affinity between a callback object (or even itself) and the UI synchronization context, and then calls a service, passing the callback reference. The service is configured for reentrancy, and it calls back to the client. A deadlock now occurs because the callback to the client needs to execute on the UI thread, and that thread is blocked waiting for the service call to return. For example, [Example 8-22](#) has the potential for this deadlock. Configuring the callback as a one-way operation will not resolve the problem here, because the one-way call still needs to be marshaled first to the UI thread. The only way to resolve the deadlock in this case is to turn off using the UI synchronization context by the callback, and to manually and asynchronously marshal the update to the form using its synchronization context. [Example 8-24](#) demonstrates using this technique.

Example 8-24. Avoiding a callback deadlock on the UI thread

```
//////////////////////////// Client Side /////////////////////
[CallbackBehavior(UseSynchronizationContext = false)]
```

```

partial class MyForm : Form,IMyContractCallback
{
    SynchronizationContext m_Context;
    MyContractClient m_Proxy;
    public MyForm()
    {
        InitializeComponent();
        m_Context = SynchronizationContext.Current;
        m_Proxy = new MyContractClient(new InstanceContext(this));
    }

    public void CallService(object sender,EventArgs args)
    {
        m_Proxy.MyMethod();
    }
    //Callback runs on worker threads
    public void OnCallback()
    {
        SendOrPostCallback setText = _=>
        {
            Text = "Manually marshaling to UI thread";
        };
        m_Context.Post(setText,null);
    }
    public void OnClose(object sender,EventArgs args)
    {
        m_Proxy.Close();
    }
}
//////////////////////////// Service Side /////////////////////
[ServiceContract(CallbackContract = typeof(IMyContractCallback))]
interface IMyContract
{
    [OperationContract]
    void MyMethod();
}

interface IMyContractCallback{
    [OperationContract]
    void OnCallback();
}

[ServiceBehaviorConcurrencyMode = ConcurrencyMode.Reentrant]
class MyService : IMyContract
{
    public void MyMethod()
    {
        IMyContractCallback callback = OperationContext.Current.
            GetCallbackChannel<IMyContractCallback>();
        callback.OnCallback();
    }
}

```

As shown in [Example 8-24](#), you must use the `Post()` method of the synchronization context. Under no circumstances should you use the `Send()` method—even though the callback is executing on the worker thread, the UI thread is still blocked on the outbound call. Calling `Send()` would trigger the deadlock you are trying to avoid because `Send()` will block until the UI thread can process the request. The callback in [Example 8-24](#) cannot use any of the safe controls (such as `SafeLabel`) either, because those too use the `Send()` method.

Callback Custom Synchronization Contexts

As with a service, you can install a custom synchronization context for the use of the callback. All that is required is that the thread that opens the proxy (or calls it for the first time) has the custom synchronization context attached to it. [Example 8-25](#) shows how to attach my `ThreadPoolSynchronizer` to the callback object by setting it before using the proxy.

Example 8-25. Setting custom synchronization context for the callback

```
interface IMyContractCallback
{
    [OperationContract]
    void OnCallback();
}

[ServiceContract(CallbackContract = typeof(IMyContractCallback))]
interface IMyContract
{
    [OperationContract]
    void MyMethod();
}

class MyClient : IMyContractCallback
{
    //This method is always invoked by the same thread
    public void OnCallback()
    {...}
}

MyClient client = new MyClient();
InstanceContext callbackContext = new InstanceContext(client);
MyContractClient proxy = new MyContractClient(callbackContext);

SynchronizationContext synchronizationContext = new ThreadPoolSynchronizer(3);
SynchronizationContext.SetSynchronizationContext(synchronizationContext);

using(synchronizationContext as IDisposable)
{
    proxy.MyMethod();
    /*Some blocking operations until after the callback*/
```

```

        proxy.Close();
    }
}

```

While you could manually install a custom synchronization context (as in [Example 8-25](#)) by explicitly setting it before opening the proxy, it is better to do so declaratively, using an attribute. To affect the callback endpoint dispatcher, the attribute needs to implement the `IEndpointBehavior` interface presented in [Chapter 6](#):

```

public interface IEndpointBehavior
{
    void ApplyClientBehavior(ServiceEndpoint endpoint, ClientRuntime
                           clientRuntime);
    //More members
}

```

In the `ApplyClientBehavior` method, the `ClientRuntime` parameter contains a reference to the endpoint dispatcher with the `CallbackDispatchRuntime` property:

```

public sealed class ClientRuntime
{
    public DispatchRuntime CallbackDispatchRuntime
    {get;}
    //More members
}

```

The rest is identical to the service-side attribute, as demonstrated by my `CallbackThreadPoolBehaviorAttribute`, whose implementation is shown in [Example 8-26](#).

Example 8-26. Implementing `CallbackThreadPoolBehaviorAttribute`

```

[AttributeUsage(AttributeTargets.Class)]
public class CallbackThreadPoolBehaviorAttribute : ThreadPoolBehaviorAttribute,
                                                IEndpointBehavior
{
    public CallbackThreadPoolBehaviorAttribute(uint poolSize, Type clientType) :
                                                this(poolSize, clientType, null)
    {}
    public CallbackThreadPoolBehaviorAttribute(uint poolSize, Type clientType,
                                              string poolName) : base(poolSize, clientType, poolName)
    {
        AppDomain.CurrentDomain.ProcessExit += delegate
        {
            ThreadPoolHelper.CloseThreads(ServiceType);
        };
    }
    void IEndpointBehavior.ApplyClientBehavior(ServiceEndpoint serviceEndpoint,
                                                ClientRuntime clientRuntime)
    {
        IContractBehavior contractBehavior = this;
        contractBehavior.ApplyDispatchBehavior(null, serviceEndpoint,

```

```

        clientRuntime.CallbackDispatchRuntime);
}
//Rest of the implementation
}

```

In fact, I wanted to reuse as much of the service attribute as possible in the callback attribute. To that end, `CallbackThreadPoolBehaviorAttribute` derives from `ThreadPoolBehaviorAttribute`. Its constructors pass the client type as the service type to the base constructors. The `CallbackThreadPoolBehavior` attribute's implementation of `ApplyClientBehavior()` queries its base class for `IContractBehavior` (this is how a subclass uses an explicit private interface implementation of its base class) and delegates the implementation to `ApplyDispatchBehavior()`.

The big difference between a client callback attribute and a service attribute is that the callback scenario has no host object to subscribe to its `Closed` event. To compensate, the `CallbackThreadPoolBehavior` attribute monitors the process exit event to close all the threads in the pool.

If the client wants to expedite closing those threads, it can use `ThreadPoolBehavior.CloseThreads()`, as shown in [Example 8-27](#).

Example 8-27. Using the `CallbackThreadPoolBehavior` attribute

```

interface IMyContractCallback
{
    [OperationContract]
    void OnCallback();
}

[ServiceContract(CallbackContract = typeof(IMyContractCallback))]
interface IMyContract
{
    [OperationContract]
    void MyMethod();
}

[CallbackThreadPoolBehavior(3,typeof(MyClient))]
class MyClient : IMyContractCallback, IDisposable
{
    MyContractClient m_Proxy;

    public MyClient()
    {
        m_Proxy = new MyContractClient(new InstanceContext(this));
    }

    public void CallService()
    {
        m_Proxy.MyMethod();
    }
}

```

```

}

//Called by threads from the custom pool
public void OnCallback()
{...}

public void Dispose()
{
    m_Proxy.Close();
    ThreadPoolHelper.CloseThreads(typeof(MyClient));
}
}

```

Callback thread affinity

Just like on the service side, if you want all the callbacks to execute on the same thread (perhaps to create some UI on the callback side), you can configure the callback class to have a pool size of 1. Or, better yet, you can define a dedicated callback attribute such as my `CallbackThreadAffinityBehaviorAttribute`:

```

[AttributeUsage(AttributeTargets.Class)]
public class CallbackThreadAffinityBehaviorAttribute :
    CallbackThreadPoolBehaviorAttribute
{
    public CallbackThreadAffinityBehaviorAttribute(Type clientType) : this(
        clientType,"Callback Worker Thread")
    {}
    public CallbackThreadAffinityBehaviorAttribute(Type clientType,string
        threadName) : base( 1,
        clientType,threadName)
    {}
}

```

The `CallbackThreadAffinityBehavior` attribute makes all callbacks across all callback contracts the client supports execute on the same thread, as shown in [Example 8-28](#).

Example 8-28. Applying the `CallbackThreadAffinityBehavior` attribute

```

[CallbackThreadAffinityBehavior(typeof(MyClient))]
class MyClient : IMyContractCallback, IDisposable
{
    MyContractClient m_Proxy;

    public void CallService()
    {
        m_Proxy = new MyContractClient(new InstanceContext(this));
        m_Proxy.MyMethod();
    }
    //This method invoked by same callback thread, plus client threads
    public void OnCallback()

```

```

{
    //Access state and resources, synchronize manually
}
public void Dispose()
{
    m_Proxy.Close();
}
}

```

Note that although WCF always invokes the callback on the same thread, you still may need to synchronize access to it if other client-side threads access the method as well.

Asynchronous Calls

When a client calls a service, usually the client is blocked while the service executes the call, and control returns to the client only when the operation completes its execution and returns. However, there are quite a few cases in which you will want to call operations asynchronously; that is, you'll want control to return immediately to the client while the service executes the operation in the background and then somehow let the client know that the method has completed execution and provide the client with the results of the invocation. Such an execution mode is called *asynchronous operation invocation*, and the action is known as an *asynchronous call*. Asynchronous calls allow you to improve client responsiveness and availability.

Requirements for an Asynchronous Mechanism

To make the most of the various options available with WCF asynchronous calls, you should be aware of the generic requirements set for any service-oriented asynchronous call support. These requirements include the following:

- The same service code should be used for both synchronous and asynchronous invocation. This allows service developers to focus on business logic and cater to both synchronous and asynchronous clients.
- A corollary of the first requirement is that the client should be the one to decide whether to call a service synchronously or asynchronously. That, in turn, implies that the client will have different code for each case (whether to invoke the call synchronously or asynchronously).
- The client should be able to issue multiple asynchronous calls and have multiple asynchronous calls in progress, and it should be able to distinguish between multiple methods' completions.
- Since a service operation's output parameters and return values are not available when control returns to the client, the client should have a way to harvest the results when the operation completes.

- Similarly, communication errors or errors on the service side should be communicated back to the client side. Any exception thrown during operation execution should be played back to the client later.
- The implementation of the mechanism should be independent of the binding and transfer technology used. Any binding should support asynchronous calls.
- The mechanism should not use technology-specific constructs such as .NET exceptions or delegates.
- The asynchronous calls mechanism should be straightforward and simple to use (this is less of a requirement and more of a design guideline). For example, the mechanism should, as much as possible, hide its implementation details, such as the worker threads used to dispatch the call.

The client has a variety of options for handling operation completion. After it issues an asynchronous call, it can choose to:

- Perform some work while the call is in progress and then block until completion.
- Perform some work while the call is in progress and then poll for completion.
- Receive notification when the method has completed. The notification will be in the form of a callback on a client-provided method. The callback should contain information identifying which operation has just completed and its return values.
- Perform some work while the call is in progress, wait for a predetermined amount of time, and then stop waiting, even if the operation execution has not yet completed.
- Wait simultaneously for completion of multiple operations. The client can also choose to wait for all or any of the pending calls to complete.

WCF offers all of these options to clients. The WCF support is strictly a client-side facility, and in fact the service is unaware it is being invoked asynchronously. This means that intrinsically any service supports asynchronous calls, and that you can call the same service both synchronously and asynchronously. In addition, because all of the asynchronous invocation support happens on the client side regardless of the service, you can use any binding for the asynchronous invocation.

Proxy-Based Asynchronous Calls

Because the client decides if the call should be synchronous or asynchronous, you need to create a different proxy for the asynchronous case. In Visual Studio 2013 and later, when adding a service reference, you can click the Advanced button in the Add Service Reference dialog to bring up the Service Reference Settings dialog that lets you tweak the proxy generation. By default, the “Allow Generation of Asynchronous Operations” option is checked and the “Generate task-based operations” radio button

is selected. These settings will cause the tool to generate a proxy that contains task-based asynchronous methods in addition to the synchronous ones. For each operation in the original contract, the asynchronous proxy and contract will contain one additional method. If the original operation returns a value, the additional method will be of the form:

```
Task<returned type> <Operation>Async(<in arguments>);
```

If the original operation does not return a value, the additional method will be of the form:

```
Task <Operation>Async(<in arguments>);
```

The Service Reference Settings dialog also offers the “Generate asynchronous operations” option which allows you to create a proxy with legacy asynchronous programming model (APM) methods. For each operation in the original contract, the asynchronous proxy and contract will contain two additional methods of this form:

```
[OperationContract(AsyncPattern = true)]
IAsyncResult Begin<Operation>(<in arguments>,
                               AsyncCallback callback, object asyncState);
<returned type> End<Operation>(<out arguments>, IAsyncResult result);
```

The `OperationContract` attribute offers the `AsyncPattern` Boolean property, defined as:

```
[AttributeUsage(AttributeTargets.Method)]
public sealed class OperationContractAttribute : Attribute
{
    public bool AsyncPattern
    {get;set;}
    //More members
}
```

The `AsyncPattern` property defaults to false. `AsyncPattern` has meaning only on the client side; it is merely a validation flag indicating to the proxy to verify that the method on which this flag is set to true has a `Begin<Operation>()`-compatible signature and that the defining contract has a matching method with an `End<Operation>()`-compatible signature. These requirements are verified at the proxy load time. `AsyncPattern` binds the underlying synchronous method with the `Begin/End` pair and correlates the synchronous execution with the asynchronous one. While both the APM and task-based programming models produce nearly identical asynchronous behavior as described below, always prefer to use the more succinct task-based programming model which supports a fluid, compositional style of asynchronous programming.

When you make a proxy call using the task-based version of an operation, WCF will detect that the method returns a task and execute the operation asynchronously. Briefly, when the client invokes a proxy method of the form `Task<returned type>`

<Operation>Async() or Task <Operation>Async(), this tells WCF not to try to directly invoke a method with that name on the service. Instead, WCF should use a thread from the I/O completion thread pool to initiate the call to the underlying service method. A small group of threads from the I/O completion thread pool will then monitor the status of the service's response. Upon receiving a response from the service, Windows will assign a thread from the I/O completion thread pool to handle the response and then forward processing of the response to WCF. The client will be blocked for only the slightest moment it takes to dispatch the call request to the thread pool. WCF correlates the response of the asynchronous invocation with conveying results and exceptions to the associated task.



It is the return value of Task not the Async suffix that WCF uses to detect task-based asynchronous methods. A task-based operation will still operate asynchronously without the suffix. The proxy generation tool simply adorns asynchronous operation definitions with the suffix to avoid naming collisions.

It is also best practice to use the Async suffix with your task-based operations to clearly convey to the users of your proxies that the operation will execute asynchronously.

Example 8-29 shows a calculator contract, its implementing service, and the generated asynchronous proxy.

Example 8-29. Asynchronous contract and proxy

```
/////////////// Service Side ///////////
[ServiceContract]
interface ICalculator
{
    [OperationContract]
    int Add(int number1,int number2);
    //More operations
}
class Calculator : ICalculator
{
    public int Add(int number1,int number2)
    {
        return number1 + number2;
    }
    //Rest of the implementation
}
////////////// Client Side ///////////
[ServiceContract]
interface ICalculator
{
    [OperationContract]
```

```

        int Add(int number1,int number2);

        [OperationContract]
        Task<int> AddAsync(int number1,int number2);
        //Rest of the methods
    }
    class CalculatorClient : ClientBase<ICalculator>,ICalculator
    {
        public int Add(int number1,int number2)
        {
            return Channel.Add(number1,number2);
        }
        public Task<int> AddAsync(int number1,int number2)
        {
            return Channel.AddAsync(number1,number2);
        }
        //Rest of the methods and constructors
    }
}

```

Asynchronous Invocation

`<Operation>Async()` accepts the input parameters of the original synchronous operation, which may include primitive types or data contracts passed by value only. The original method's return value becomes the generic type parameter of a task of type `Task<T>`. It is important to note that in .NET you cannot use explicit output parameters designated using the `out` and `ref` modifiers with task-based methods. This is because a caller would expect these values to be set upon the return of the task-based method, but the asynchronous implementation of the method may not set them until some time later during its completion.

To maintain parity with .NET, WCF also abides by this constraint. While your contract and proxy code containing such an asynchronous operation definition will compile, you will receive an `InvalidOperationException` when you attempt to invoke *any* method on the asynchronous contract. For example, manually placing the following operation definition on your client-side contract and calling any asynchronous method from your proxy will produce the error.

```

[OperationContract]
Task<string> MyMethodAsync(int number1,out int number2,ref int number3);

```

If your service-side contract does include output parameters, the proxy generation tool will compensate for the task-based restriction by creating additional request and response message contracts to wrap any output parameters that exist within your operation. This in turn will add unnecessary complexity to your client-side programming model. Instead, prefer to use the techniques discussed in [Chapter 3](#) to create well-factored data contracts that contain these values instead of expressing them as output parameters.

When WCF receives the response of a task-based asynchronous operation it transfers the return value to the returned task. If you have no interest in the results or the errors of the operation, you could use code like the following to asynchronously invoke the `AddAsync()` method of the `Calculator` service from [Example 8-29](#) using the asynchronous proxy:

```
CalculatorClient proxy = new CalculatorClient();
proxy.AddAsync(2,3); //Dispatched asynchronously
proxy.Close();
```

As long as the client has the definition of the asynchronous contract, you can also invoke the operation asynchronously using a channel factory:

```
ChannelFactory<ICalculator> factory = new ChannelFactory<ICalculator>(...);
ICalculator proxy = factory.CreateChannel();
proxy.AddAsync(2,3);
(proxy as ICommunicationObject).Close();
```

The problem with such an invocation is that the client has no way of getting its results or errors.

The Task class

Every task-based asynchronous operation returns either a task of type `Task` or `Task<T>` from the `System.Threading.Tasks` namespace:

```
public class Task : ...
{...}

public class Task<T> : Task
{
    public T Result
    {get;}
    //More members
}
```

The returned task object is uniquely associated to a single invocation of a task-based asynchronous operation. You can access the `Result` property of the task to retrieve the results of a specific asynchronous method execution. `Result` will block its caller until the operation it's waiting for (encapsulated by the associated task) completes and it can return the results or errors. If the method is already complete by the time you access `Result`, `Result` will not block the caller and will just return the results. [Example 8-30](#) shows the entire sequence.

Example 8-30. Simple asynchronous execution sequence

```
CalculatorClient proxy = new CalculatorClient();
Task<int> task1 = proxy.AddAsync(2,3);
Task<int> task2 = proxy.AddAsync(4,5);
```

```

/* Do some work */

int sum;

sum = task1.Result; //This may block
Debug.Assert(sum == 5);

sum = task2.Result; //This may block
Debug.Assert(sum == 9);

proxy.Close();

```

As simple as [Example 8-30](#) is, it does demonstrate a few key points. The first point is that the same proxy instance can invoke multiple asynchronous calls. The caller can distinguish among the different pending calls using each unique Task object returned from a task-based operation. In fact, when the caller makes asynchronous calls, as in [Example 8-30](#), it must save the Task objects. In addition, the caller should make no assumptions about the order in which the pending calls will complete. It is quite possible that the second call will complete before the first one. You can also call Result multiple times to retrieve results. The first call may or may not block the caller, while all subsequent calls will return results immediately.

Asynchronous calls and transport sessions

If the proxy is not using a transport session, the client can close the proxy immediately after the call to a task-based operation and still be able to call Result later:

```

CalculatorClient proxy = new CalculatorClient();
Task<int> task = proxy.AddAsync(2,3);
proxy.Close();

/* Do some work */

//Sometime later:
int sum = task.Result;
Debug.Assert(sum == 5);

```

Polling or Waiting for Completion

When a client calls Result, the client is blocked until the asynchronous method returns. This may be fine if the client has a finite amount of work to do while the call is in progress, and if after completing that work the client cannot continue its execution without the returned value. However, what if the client only wants to check that the operation has completed? What if the client wants to wait for completion for a fixed timeout and then, if the operation has not completed, do some additional finite processing and wait again? .NET supports these alternative programming models to calling Result.

The Task object returned from a task-based proxy operation provides helper methods that allow you to wait for an asynchronous operation to complete.

```
public class Task : ...
{
    public void Wait();
    public bool Wait(int millisecondsTimeout);
    public static void WaitAll(params Task[] tasks);
    public static int WaitAny(params Task[] tasks);
    //More members
}
```

The `Wait()` method of the `Task` class returns only when the task is completed.

Example 8-31 demonstrates using `Wait()`.

Example 8-31. Using Task.Wait() to block until completion

```
CalculatorClient proxy = new CalculatorClient();
Task<int> task = proxy.AddAsync(2,3);

/* Do some work */

task.Wait(); //This may block
int sum = task.Result; //This will not block
Debug.Assert(sum == 5);

proxy.Close();
```

Logically, **Example 8-31** is identical to **Example 8-30**, which called only `Result`. If the operation is still executing when `Wait()` is called, `Wait()` will block. But if by the time `Wait()` is called the method execution is complete, `Wait()` will not block, and the client will proceed to call `Result` for the returned value. The important difference between Examples **8-30** and **8-31** is that the call to `Result` in **Example 8-31** is guaranteed not to block its caller.

Example 8-32 demonstrates a more practical way of using `Wait()`, by specifying a timeout (10 milliseconds in this example). When you specify a timeout, `Wait()` returns when the method execution is completed or when the timeout has elapsed, whichever condition is met first.

Example 8-32. Using Wait() to specify wait timeout

```
CalculatorClient proxy = new CalculatorClient();
Task<int> task = proxy.AddAsync(2,3);
while(task.IsCompleted == false)
{
    task.Wait(10); //This may block
    /* Do some optional work */
```

```
}
```

```
int sum = task.Result; //This will not block
```

Example 8-32 uses another handy property of Task, called `IsCompleted`, which lets you check the status of the call without waiting or blocking. You can even use `IsCompleted` in a strict polling mode:

```
CalculatorClient proxy = new CalculatorClient();
Task<int> task = proxy.AddAsync(2,3);

//Sometime later:
if(task.IsCompleted)
{
    int sum = task.Result; //This will not block
    Debug.Assert(sum == 5);
}
else
{
    //Do some optional work
}
proxy.Close();
```

The task-based helpers really shine when you use them to manage multiple concurrent asynchronous methods in progress. You can use Task's static `WaitAll()` method to wait for completion of multiple asynchronous methods, as shown in **Example 8-33**.

Example 8-33. Waiting for completion of multiple methods

```
CalculatorClient proxy = new CalculatorClient();
Task<int> task1 = proxy.AddAsync(2,3);
Task<int> task2 = proxy.AddAsync(4,5);

Task.WaitAll(task1,task2);

int sum;

//These access of Result will not block

sum = task1.Result;
Debug.Assert(sum == 5);

sum = task2.Result;
Debug.Assert(sum == 9);

proxy.Close();
```

Note that you still need to call `Result` to access the returned values. Instead of waiting for all of the methods to return, you can choose to wait for any of them to return, using the `WaitAny()` static method of the `Task` class. Like `Wait()`, both `WaitAll()`

and `WaitAny()` have overloaded versions that let you specify a timeout to wait instead of waiting indefinitely.



The `Task` class also provides overloads for `Wait()`, `WaitAll()` and `WaitAny()` that take as a parameter an instance of the `CancellationToken` structure from the `System.Threading` namespace. These overloads are not applicable to the execution of task-based WCF proxy operations. Once the proxy issues a service request onto the network the system cannot retract it, nor can a client remotely cancel a service operation's execution because WCF processes the service operation synchronously.

Continuations

Instead of blocking, waiting, and polling for asynchronous call completion, .NET offers another programming model altogether—*continuations*. With this model, you can enlist a task to continue execution once its *antecedent* task has completed. The method encapsulated within the continuation task may not run on the thread that executed the antecedent task. Unless specified otherwise, by default the continuation task's method will execute on a thread from the .NET thread pool. To manually enlist a continuation task, you can use the `ContinueWith()` helper method provided by the `Task<T>` class:

```
public class Task<T> : Task
{
    public Task ContinueWith(Action<Task<T>> continuationAction);
    //More members
}
```

Example 8-34 demonstrates asynchronous call management using a continuation.

Example 8-34. Managing an asynchronous call with a continuation

```
class MyClient : IDisposable
{
    CalculatorClient m_Proxy = new CalculatorClient();

    public void CallAsync()
    {
        Task<int> task = m_Proxy.AddAsync(2,3);
        task.ContinueWith(OnCompletion);
    }
    void OnCompletion(Task<int> task)
    {
        int sum = task.Result;
        Debug.Assert(sum == 5);
    }
    public void Dispose()
```

```

    {
        m_Proxy.Close();
    }
}

```

Unlike in the programming models described so far, when you use a continuation, there's no need to save the returned Task object. This is because when .NET calls the continuation task, it provides the antecedent Task object as a parameter. Because .NET provides a unique Task object for each asynchronous method, you can channel multiple asynchronous method continuations to the same completion method:

```

class MyClient : IDisposable
{
    CalculatorClient m_Proxy = new CalculatorClient();

    public void CallAsync()
    {
        m_Proxy.AddAsync(2,3).ContinueWith(OnCompletion);
        m_Proxy.AddAsync(4,5).ContinueWith(OnCompletion);
    }

    void OnCompletion(Task<int> task)
    {
        int sum = task.Result;
    }

    public void Dispose()
    {
        m_Proxy.Close();
    }
}

```

Instead of using a class method as a completion method, you can just as easily use a local anonymous method or a lambda expression, as shown in [Example 8-35](#).

Example 8-35. Lambda completion method

```

CalculatorClient proxy = new CalculatorClient();
int sum;
Action<Task<int>> onCompletion = (task) =>
{
    sum = task.Result;
    Debug.Assert(sum == 5);
    proxy.Close();
};

proxy.AddAsync(2,3).ContinueWith(onCompletion);

```

Note in [Example 8-35](#) that the anonymous method assigns to an outer variable (`sum`) to provide the result of the `AddAsync()` operation.

Language-based continuations

Instead of manually providing continuations through `ContinueWith()`, you can use the `async` and `await` keywords introduced in .NET 4.5 to create continuation tasks automatically. You apply the `async` keyword to a .NET method to indicate to the compiler that the method may contain one or more task-based asynchronous method calls. You then apply the `await` keyword to each call. The location of the `await` keyword informs the compiler where it should automatically enlist a continuation, similar to using the anonymous method in [Example 8-35](#). The compiler will consider all of the code preceding `await` as part of an antecedent task's initialization. All the code following `await` the compiler will consider a continuation of the antecedent task. The task returned from the call to which you've applied `await` becomes the antecedent task. As with `ContinueWith()`, by default the automatic continuation will execute on a thread from the .NET thread pool. [Example 8-36](#) shows the equivalent task-based call and continuation of [Example 8-34](#) using `async` and `await`.

Example 8-36. Managing an asynchronous call with async and await

```
class MyClient : IDisposable
{
    CalculatorClient m_Proxy = new CalculatorClient();

    public async Task CallAsync()
    {
        int sum = await m_Proxy.AddAsync(2,3);
        Debug.Assert(sum == 5);
    }
    public void Dispose()
    {
        m_Proxy.Close();
    }
}
```



The `async` and `await` pattern is also called the Task Asynchronous Pattern or TAP.

As in [Example 8-34](#), the calling thread will perform the call to `AddAsync()` and WCF will dispatch the call asynchronously. The significant difference in [Example 8-36](#) is that once WCF is finished dispatching the `AddAsync()` call, the `CallAsync()` method will exit. When WCF eventually receives the reply of the `AddAsync()` call, it will convey the results to the returned task. A thread from the thread pool will then execute the remainder of the `CallAsync()` method as a continuation, including the assignment of `sum`. Notice that `await` does not return a task. As a programming convenience,

nience, the compiler-generated support for `await` calls `Result` for you returning the results of the task associated with the `AddAsync()` call.

Note that in [Example 8-36](#) the return value of the `CallAsync()` method changed from `void` to `Task`. You should always return `Task` or `Task<T>` from your task-based methods. This will ensure that your task-based methods conform to the .NET programming model for consuming tasks by supporting task composition, conveying task results to callers and propagating exceptions that occur within task operations.

Continuations are by far the preferred model in any event-driven application. An event-driven application has methods that trigger events (or requests) and methods that handle those events and fire their own events as a result. Writing an application as event-driven makes it easier to manage multiple threads, events, and callbacks and allows for scalability, responsiveness, and performance.

The last thing you want in an event-driven application is to block, since then your application does not process events. Continuation tasks allow you to treat the completion of the asynchronous operation as yet another event in your system. The other options (waiting, blocking, and polling) are available for applications that are strict, predictable, and deterministic in their execution flow. Prefer to use continuations whenever possible.

Continuations and thread safety

Because the completion method is executed on a thread from the thread pool, you must provide for thread safety in the completion method and in the object that provides it. This means that you must use synchronization locks to access the member variables of the client, even outer variables to anonymous completion methods. You need to provide for synchronization between client-side threads and the worker thread from the pool, and potentially synchronizing between multiple worker threads all calling concurrently into the completion method to handle their respective asynchronous call completion. Therefore, you need to make sure the completion method is reentrant and thread-safe.

Passing state information

The `Task` and `Task<T>` classes provide overloads of `ContinueWith()` that allow you to pass into your continuation an amorphous object. This object, known as a *state object*, is provided as an optional container for whatever need you deem fit. The party handling the continuation can access such a container object as a parameter passed into the continuation method. Although you can certainly use state objects with any of the other asynchronous call programming models (blocking, waiting, or polling), they are most useful in conjunction with continuations. The reason is simple: when you are using a continuation, the state object offers the only way to pass in additional parameters to the continuation method, whose signature is predetermined.

Example 8-37 demonstrates how you might use a state object to pass an integer value as an additional parameter to the completion method. Note that the completion method must downcast the `state` parameter to the actual type.

Example 8-37. Passing an additional parameter using a state object

```
class MyClient : IDisposable
{
    CalculatorClient m_Proxy = new CalculatorClient();

    public void CallAsync()
    {
        int asyncState = 4; //int, for example
        m_Proxy.AddAsync(2,3).ContinueWith(OnCompletion,asyncState);
    }
    void OnCompletion(Task<int> task,object state)
    {
        int asyncState = (int)state;
        Debug.Assert(asyncState == 4);

        int sum = task.Result;
        Debug.Assert(sum == 5);
    }
    public void Dispose()
    {
        m_Proxy.Close();
    }
}
```

A common use for the state object is to pass the proxy used for the task-based asynchronous call instead of saving it as a member variable:

```
class MyClient
{
    public void CallAsync()
    {
        CalculatorClient proxy = new CalculatorClient();
        proxy.AddAsync(2,3).ContinueWith(OnCompletion,proxy);
    }
    void OnCompletion(Task<int> task,object state)
    {
        CalculatorClient proxy = state as CalculatorClient;
        Debug.Assert(proxy != null);

        int sum = task.Result;
        Debug.Assert(sum == 5);

        proxy.Close();
    }
}
```

Continuation synchronization context

The completion method, by default, is called on a thread from the thread pool. This presents a serious problem if the completion method is to access some resources that have an affinity to a particular thread or threads and are required to run in a particular synchronization context. The classic example is a Windows Forms application that dispatches a lengthy service call asynchronously (to avoid blocking the UI), and then wishes to update the UI with the result of the invocation. Using the raw `ContinueWith()` is disallowed, since only the UI thread is allowed to update the UI. You must marshal the call from the completion method to the correct synchronization context, using any of the techniques described previously (such as safe controls). [Example 8-38](#) demonstrates such a completion method that interacts directly with its containing form, ensuring that the UI update will be in the UI synchronization context.

Example 8-38. Relying on continuation synchronization context

```
partial class CalculatorForm : Form
{
    CalculatorClient m_Proxy;
    SynchronizationContext m_SynchronizationContext;

    public CalculatorForm()
    {
        InitializeComponent();
        m_Proxy = new CalculatorClient();
        m_SynchronizationContext = SynchronizationContext.Current;
    }

    public void CallAsync(object sender, EventArgs args)
    {
        m_Proxy.AddAsync(2,3).ContinueWith(OnCompletion);
    }

    void OnCompletion(Task<int> task)
    {
        SendOrPostCallback callback = delegate
        {
            Text = "Sum = " + task.Result;
        };
        m_SynchronizationContext.Send(callback,null);
    }

    public void OnClose(object sender, EventArgs args)
    {
        m_Proxy.Close();
    }
}
```

To better handle this situation, the `Task` and `Task<T>` classes provide overloads of `ContinueWith()` that allow you to pass as a parameter a `TaskScheduler`. `ContinueWith()` will use this task scheduler to invoke the completion method:

```
class Task<T> : Task
{
    public Task ContinueWith(Action<Task<T>> continuationAction,
                           TaskScheduler scheduler);
    //More methods
}
```

As described previously, you can use `FromCurrentSynchronizationContext()` to capture the calling thread's synchronization context and obtain a `TaskScheduler` that marshals all of its calls to that synchronization context:

```
public abstract class TaskScheduler
{
    public static TaskScheduler FromCurrentSynchronizationContext();
    //More members
}
```

You can then assign this task scheduler to your continuation, as shown in [Example 8-39](#).

Example 8-39. Synchronization-context-friendly continuation

```
partial class CalculatorForm : Form
{
    CalculatorClient m_Proxy;

    public CalculatorForm()
    {
        InitializeComponent();

        m_Proxy = new CalculatorClient();
    }
    void CallAsync(object sender,EventArgs args)
    {
        //Sync context picked up here
        m_Proxy.AddAsync(2,3).ContinueWith(OnCompletion,
                                         TaskScheduler.FromCurrentSynchronizationContext());
    }
    //Called on the UI thread
    void OnCompletion(Task<int> task)
    {
        Text = "Sum = " + task.Result;
    }
    public void OnClose(object sender,EventArgs args)
    {
        m_Proxy.Close();
    }
}
```

```
    }
}
```

Alternatively, you can use `async` and `await` to achieve the same result. The big difference with using `async` and `await` as opposed to `ContinueWith()` is that by default `await` will capture the current synchronization context and marshal the continuation to it. The compiler generated support for `await` will pick up the synchronization context of the client, create a task scheduler from it, and then assign this task scheduler to the continuation, as shown in [Example 8-40](#).

Example 8-40. Synchronization-context-friendly continuation using await

```
partial class CalculatorForm : Form
{
    CalculatorClient m_Proxy;

    public CalculatorForm()
    {
        InitializeComponent();

        m_Proxy = new CalculatorClient();
    }

    async void CallAsync(object sender,EventArgs args)
    {
        int sum = await m_Proxy.AddAsync(2,3); //Sync context picked up here

        //Called on the UI thread
        Text = "Sum = " + sum;
    }

    public void OnClose(object sender,EventArgs args)
    {
        m_Proxy.Close();
    }
}
```

Note that if you were to add a continuation to [Example 8-40](#) such that:

```
int OnCompletion(Task<int> task)
{
    int result = task.Result;
    // No sync context here!
    Text = "Sum = " + result;
    return result;
}

async void CallAsync(object sender,EventArgs args)
{
    int sum = await m_Proxy.AddAsync(2,3).ContinueWith<int>(OnCompletion);

    //Called on the UI thread
}
```

```
    Text = "Sum = " + sum;
}
```

Your completion method would not have a synchronization context. `await` only applies to the overall statement and its continuation in the context of the caller, not to any continuations within the statement. You will still need to use `FromCurrentSyncrhonizationContext()` to designate that your inline continuation should also run on the current synchronization context.

Async/await and deadlocks

If you were to change the definition of the `CalculatorClient.AddAsync()` method from [Example 8-29](#) to use the `async` and `await` pattern as shown:

```
class CalculatorClient : ClientBase<ICalculator>, ICalculator
{
    public async Task<int> AddAsync(int number1,int number2)
    {
        int sum = await Channel.AddAsync(number1,number2);
        return sum;
    }
    //Rest of the methods and constructors
}
```

A problem will arise if you make a call to `AddAsync()` and then wait for task completion in an environment that possesses single-threaded affinity, such as a Windows user interface. Thread deadlock will result when the compiler generated support for `await` attempts to marshal the continuation back to the ambient synchronization context it captured at the time of the call. A deadlock occurs because you have already blocked the UI thread waiting for the task to complete. For example, using the definition of `AddAsync()` above the following code will always deadlock:

```
partial class CalculatorForm : Form
{
    CalculatorClient m_Proxy;

    public CalculatorForm()
    {
        InitializeComponent();

        m_Proxy = new CalculatorClient();
    }
    async void CallAsync(object sender,EventArgs args)
    {
        Task<int> task = m_Proxy.AddAsync(2,3);

        int sum = task.Result; //This blocks the UI thread

        //Called on the UI thread
        Text = "Sum = " + sum;
    }
}
```

```

        }
        public void OnClose(object sender,EventArgs args)
        {
            m_Proxy.Close();
        }
    }
}

```

The problem is not unique to asynchronous calls in WCF using the `async` and `await` pattern. Thread deadlock will occur in .NET if you apply the same approach within a thread-affinity context such as user interface. For example, the following .NET code will produce thread deadlock:

```

class MyCalculator
{
    public async Task<int> AddAsync(int number1,int number2)
    {
        //Deadlock occurs when await attempts to return
        int sum = await Task.Run(()=> number1 + number2);
        return sum;
    }
}
partial class CalculatorForm : Form
{
    MyCalculator m_MyCalculator;

    public CalculatorForm()
    {
        InitializeComponent();

        MyCalculator m_MyCalculator = new MyCalculator();
    }
    async void CallAsync(object sender,EventArgs args)
    {
        Task<int> task = m_MyCalculator.AddAsync(2,3);

        int sum = task.Result; //This blocks the UI thread

        Text = "Sum = " + sum;
    }
    public void OnClose(object sender,EventArgs args)
    {
        m_Proxy.Close();
    }
}

```

The root cause of the deadlock stems from a design flaw in the code above of mixing asynchronous and synchronous techniques. When you use the `async` and `await` pattern, you introduce non-blocking execution flow into your application. While the non-blocking flow may initiate many asynchronous tasks throughout its scope, a single thread executes the flow making the flow itself serial in time. Because of this, the non-blocking flow works best when you allow it to continue throughout your applica-

cation so that your code eventually relinquishes the associated thread back to the thread pool to perform other tasks. When you initiate non-blocking execution flow on a thread and then choose to explicitly block that thread by waiting on a task, you prevent the non-blocking flow from completing. The associated thread cannot continue the non-blocking flow until the task upon which you are waiting completes. If your application is multi-threaded, when you chose to explicitly block you are needlessly consuming system resources which may ultimately hinder your application's overall responsiveness. If your application is single-threaded (or has thread affinity), when you choose to explicitly block after initiating a non-blocking execution flow with `async` and `await`, by default you will always produce a deadlock. This is because you have blocked the only thread on which your continuation can run.

The simplest way to avoid the possibility of thread deadlock when using asynchronous proxy calls in your UI is to not apply the `async` and `await` pattern to your proxy methods. Instead, simply return the task as shown in [Example 8-29](#). This allows the consumers of your proxies to choose how they wish to handle asynchronous interactions in their environment. You should also avoid mixing asynchronous and synchronous techniques. Instead, apply a consistent technique across all your asynchronous interactions so that you allow the non-blocking execution flow to complete. For example, if the inner scope of your interaction applies the `async` and `await` pattern, use `async` and `await` throughout:

```
partial class CalculatorForm : Form
{
    CalculatorClient m_Proxy;

    public CalculatorForm()
    {
        InitializeComponent();

        m_Proxy = new CalculatorClient();
    }

    async void CallAsync(object sender,EventArgs args)
    {
        int sum = await m_Proxy.AddAsync(2,3);

        Text = "Sum = " + sum;
    }

    public void OnClose(object sender,EventArgs args)
    {
        m_Proxy.Close();
    }
}
```

You can also use the `Task.ConfigureAwait()` method to control whether or not the compiler generated support for `await` attempts to capture the ambient synchronization context, if one exists.

```

public struct ConfiguredTaskAwaitable<T>
{
}

public class Task<T> : Task
{
    public ConfiguredTaskAwaitable<T> ConfigureAwait(bool onCapturedContext);

    //More members
}

class MyCalculator
{
    public async Task<int> AddAsync(int number1,int number2)
    {
        int sum = await Task.Run(()=> number1 + number2).ConfigurAwait(false);
        return sum;
    }
}

```

Setting `ConfigureAwait()` to `false` instructs the system not to capture the ambient synchronization context at the point of `await` and therefore it will not marshal the continuation to it. This guards the eventual execution of the continuation from any synchronization or single-threaded deadlock concerns:

One-Way Asynchronous Operations

In some cases, such as firing an event with a large payload or broadcasting a larger number of one-way operations, you may want to invoke a one-way operation asynchronously. The problem is that one-way calls are not well aligned with asynchronous calls: one of the main features of asynchronous calls is their ability to retrieve and correlate a reply message; no such message is available with a one-way call. If you do invoke a task-based one-way operation, the call will return control immediately to the calling thread and the task associated with the call will complete as soon as the worker thread has finished dispatching the call. Aside from communication errors, the task-based call will not encounter any exceptions. If you provide a continuation task for a task-based one-way operation, the completion method will execute on a thread from the thread pool immediately after WCF dispatches the call. The only justification for invoking a one-way operation asynchronously is to avoid the potential blocking of the one-way call, in which case you should not enlist a continuation task or `await` the returned task, as shown in [Example 8-41](#).

Example 8-41. Invoking a one-way operation asynchronously

```

[ServiceContract]
interface IMyContract
{
    [OperationContract(IsOneWay = true)]
    void MyMethod(string text);
}

```

```

[OperationContract(IsOneWay = true)]
Task MyMethodAsync(string text);
}
MyContractClient proxy = MyContractClient();
proxy.MyMethodAsync("Async one way");

//Sometime later:
proxy.Close();

```

Note that because there is no return value, your task-based one-way operations may only return a Task. Returning a Task<T> would imply a return value and will result in runtime errors just as with synchronous one-way operations.

The problem with [Example 8-41](#) is the potential race condition of closing the proxy. It is possible to push the asynchronous call and then close the proxy before the worker thread used has had a chance to invoke the call. If you want to close the proxy immediately after asynchronously invoking the one-way call, you need to provide a continuation for closing the proxy:

```

MyContractClient proxy = MyContractClient();

proxy.MyMethodAsync("Async one way").ContinueWith(_=>proxy.Close());

```

However, this is still flawed, because it complicates issuing a number of asynchronous, one-way calls, since there is no single place to wait for all the dispatching to complete. One approach you might consider is to apply `await` to each call to create a chain of continuations as shown in [Example 8-42](#).

Example 8-42. Multiple asynchronous continuations with await

```

async Task CallMultipleOneWayAsync()
{
    MyContractClient proxy = MyContractClient();
    await proxy.MyMethodAsync("Async one way 1");
    await proxy.MyMethodAsync("Async one way 2");

    //Sometime later:
    proxy.Close();
}

```

With the approach shown in [Example 8-42](#), when the client now calls multiple task-based one-way operations, each call is always dispatched asynchronously but each continuation task will not execute until its associated antecedent task has completed. The multiple continuations in [Example 8-42](#) execute in turn until the last one executes and calls `Close()`. This is [Example 8-42](#) expressed as multiple, manual continuations:

```

void CallMultipleAsync()
{
    MyContractClient proxy = MyContractClient();
    proxy.MyMethodAsync("Async one way 1").ContinueWith(_=>proxy.MyMethodAsync(
        "Async one way 2")).ContinueWith(_=>proxy.Close());
}

```

While the technique used in [Example 8-42](#) solves the problem of providing a single place to close the proxy after issuing multiple asynchronous one-way calls, it does not avoid the problem of the one-way calls blocking and executing asynchronously but serially. To handle that, you can use the `OneWayClientBaseAsync<T>` helper class from `ServiceModelEx`, defined as:

```

public class OneWayClientBaseAsync<T> : ClientBase<T> where T : class
{
    protected async Task Invoke(Task task);
}

```

Using the contract of [Example 8-41](#), you need to derive your proxy from `OneWayClientBaseAsync<T>` and implement the methods as follows:

```

class MyContractClient : OneWayClientBaseAsync<IMyContract>, IMyContract
{
    public void MyMethod(string text)
    {
        Invoke(Channel.MyMethodAsync(text));
    }
    Task IMyContract.MyMethodAsync()
    {
        throw new InvalidOperationException("Do not call directly,
                                         use MyMethod()");
    }
}

```

When the client now calls the original one-way operation:

```

MyContractClient proxy = MyContractClient();
proxy.MyMethod("Async one way");

//Sometime later:
proxy.Close();

```

the proxy always dispatches the call asynchronously and closing the proxy waits for all the pending asynchronous one-way calls to complete dispatch. You never call `MyMethodAsync()` directly so you should throw an exception in it. Note that the deliberate decision not to mark `MyMethod()` as `async` so that you can `await` the call to `Invoke()`. This pattern follows the rationale for asynchronous one-way calls stated at the beginning of this section. By not awaiting the call, the design explicitly chooses to suppress exception propagation to the caller. [Example 8-43](#) shows the implementation of `OneWayClientBaseAsync<T>`.

Example 8-43. Implementing OneWayClientBaseAsync<T>

```
public class OneWayClientBaseAsync<T> : ClientBase<T> where T : class
{
    List<Task> m_PendingOperations = new List<Task>();

    static OneWayClientBaseAsync()
    {
        Type type = typeof(T);
        Debug.Assert(type.IsInterface);

        MethodInfo[] methods = type.GetMethods();

        foreach(MethodInfo method in methods)
        {
            object[] attributes = method.GetCustomAttributes(
                typeof(OperationContractAttribute), true);

            if(attributes.Length > 0)
            {
                OperationContractAttribute operationContract = attributes[0] as
                    OperationContractAttribute;
                if(operationContract.IsOneWay == false)
                {
                    throw new InvalidOperationException(method.Name + " not one-way");
                }
                else
                {
                    if(method.Name.EndsWith("Async"))
                    {
                        Debug.Assert(method.ReturnType == typeof(Task));
                    }
                }
            }
        }
    }

    protected async Task Invoke(Task task)
    {
        lock(m_PendingOperations)
        {
            m_PendingOperations.Add(task);
        }
        try
        {
            await task.ConfigureAwait(false);
        }
        finally
        {
            lock(m_PendingOperations)
            {
                m_PendingOperations.Remove(task);
            }
        }
    }
}
```

```

        task.AsyncWaitHandle.Close();
    }
}
}
public new void Close()
{
    lock(m_PendingOperations)
    {
        Task[] tasks = m_PendingOperations.ToArray();
        Task.WaitAll(tasks);
    }

    base.Close();
}
public void Dispose()
{
    Close();
}
}

```

The static constructor of `OneWayClientBaseAsync<T>` verifies that all operations on the contract are one-way. It also verifies that if the operation name has the suffix `Async` then that the operation returns a `Task`. You will need to omit these checks if you want to mix in the contract non-one-way asynchronous calls. `OneWayClientBaseAsync<T>` maintains a list of tasks. Each task represents a pending asynchronous one-way invocation that has yet to be dispatched by the thread from the thread pool. When closing or disposing of `OneWayClientBaseAsync<T>`, it waits for all the tasks to complete before proceeding with the closing. The `Invoke()` method captures the task associated with the asynchronous operation, adds it to the list, and enlists a continuation that removes the task from the list. The `Invoke()` method immediately awaits the returned task, so as soon as the worker thread from the thread pool dispatches the call, it signals the continuation. `Invoke()` also sets `ConfigureAwait()` to `false` to ensure the system will not marshal the continuation to the ambient synchronization context, if one exists.

Asynchronous Error Handling

Return values are not the only elements unavailable at the time an asynchronous task-based call is dispatched: exceptions are missing as well. After making a task-based call, control returns to the client, but it may be some time before the asynchronous method encounters an error and throws an exception, and sometime after that before the client actually evaluates the call's associated task. WCF must therefore provide some way for the client to know that an exception was thrown and allow the client to handle it. When the asynchronous method throws an exception, the proxy catches it and conveys the exception to the call's associated task. The client can then

use the following members of the `Task` class to evaluate the task's status and exceptions:

```
public enum TaskStatus
{
    RanToCompletion,
    Canceled,
    Faulted,
    //More members
}

public class AggregateException : Exception
{
    public ReadOnlyCollection<Exception> InnerExceptions
    {get;}
    //More members
}

public class Task : ...
{
    public AggregateException Exception
    {get;}
    public bool IsCompleted
    {get;}
    public bool IsFaulted
    {get;}
    public TaskStatus Status
    {get;}
    //More members
}
```

If you have provided a continuation for the call, .NET calls the completion method immediately after receiving the exception. The antecedent task will carry with it the exception information into the continuation like so:

```
Action<Exception> processException = (exception)=>
{
    if(exception.InnerException != null)
    {
        processException(exception.
                           InnerException);
    }
    //process exception here
};

Action<Task<int>> onCompletion = (task)=>
{
    if(task.Status == TaskStatus.RanToCompletion)
    {
        sum = task.Result;
        Debug.Assert(sum == 5);
    }
}
```

```

        else
        {
            if(task.IsFaulted)
            {
                IEnumerable<Exception> exceptions =
                    task.Exception.InnerExceptions;

                foreach(Exception exception in
                    exceptions)
                {
                    processException(exception);
                }
            }
        };
    };
}

```

When the client finally evaluates the task, its `Result` property will not be set, its `Exception` property will contain the exceptions caught by the proxy and its `Status` property will be set to `TaskStatus.Faulted`. The `IsCompleted` and `IsFaulted` properties will both return `true`. Accessing the `Result` property of a faulted task will cause the task to re-throw its `AggregateException`.

If you have awaited the call, the compiler-generated support for `await` will automatically re-throw the task's exception into the caller's context via the continuation as shown:

```

public async Task CallAsync()
{
    try
    {
        int sum = await m_Proxy.AddAsync(2,3);
        Debug.Assert(sum == 5);
    }
    catch(AggregateException exception)
    {...}
}

```

The exact exception thrown is compliant with the fault contract and the exception type, as explained in [Chapter 6](#).



If fault contracts are defined on the service operation contract, the `FaultContract` attribute should be applied only on the synchronous operations.

Asynchronous calls and timeouts

Since the task-based asynchronous invocation mechanism is nothing but a convenient programming model on top of the actual synchronous service operation, the

underlying synchronous call can still time out. This will result in a `TimeoutException` when the client evaluates the task. It is therefore wrong to equate asynchronous calls with lengthy operations. By default, asynchronous calls are still relatively short, but unlike synchronous calls, they are non-blocking. For lengthy asynchronous calls you will need to provide an adequately long send timeout.

Cleaning up after task completion

To support compatibility with the existing asynchronous programming models in .NET, `Task` implements `IAsyncResult`. This allows you to obtain a `WaitHandle` via the `AsyncWaitHandle` property:

```
public class Task : IAsyncResult, ...
{
    WaitHandle AsyncWaitHandle
    {get;}
    //More members
}
```

A task only creates the handle for you when you access the property. It is important to note that the completion of your task will not close the handle. Instead, the task's finalizer will close the handle when .NET garbage-collects your task. As with any other case of using an unmanaged resource, you have to be mindful about your application-deterministic finalization needs. It is possible (in theory, at least) for the application to dispatch asynchronous calls faster than .NET can collect the handles, resulting in a resource leak. To compensate, you can explicitly close that handle after accessing `Result`. For example, using the same definitions as those in [Example 8-34](#):

```
void OnCompletion(Task<int> task)
{
    using(task.AsyncWaitHandle)
    {
        int sum = task.Result;
        Debug.Assert(sum == 5);
    }
}
```

Asynchronous Calls and Transactions

Transactions do not mix well with asynchronous calls, for a few reasons. First, well-designed transactions are of short duration, yet the main motivation for using asynchronous calls is because of the latency of the operations. Second, the client's ambient transaction will not by default flow to the service, because the asynchronous operation is invoked on a worker thread, not the client's thread. While it is possible to develop a proprietary mechanism that uses cloned transactions, this is esoteric at best and should be avoided. Finally, when a transaction completes, it should have no leftover activities to do in the background that could commit or abort independently of

the transaction; however, this will be the result of spawning an asynchronous operation call from within a transaction. In short, do not mix transactions with asynchronous calls.

Synchronous Versus Asynchronous Calls

Although it is technically possible to call the same service synchronously and asynchronously, the likelihood that a service will be accessed both ways is low.

The reason is that using a service asynchronously necessitates drastic changes to the workflow of the client, and consequently the client cannot simply use the same execution sequence logic as with synchronous access. Consider, for example, an online store application. Suppose the client (a server-side object executing a customer request) accesses a **Store** service, where it places the customer's order details. The **Store** service uses three well-factored helper services to process the order: **Order**, **Shipment**, and **Billing**. In a synchronous scenario, the **Store** service first calls the **Order** service to place the order. Only if the **Order** service succeeds in processing the order (i.e., if the item is available in the inventory) does the **Store** service then call the **Shipment** service, and only if the **Shipment** service succeeds does the **Store** service access the **Billing** service to bill the customer. This sequence is shown in [Figure 8-4](#).

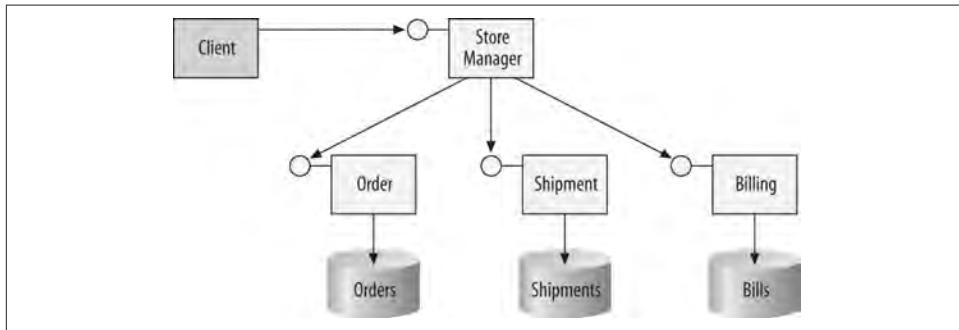


Figure 8-4. Synchronous processing of an order

The downside to the workflow shown in [Figure 8-4](#) is that the store must process orders synchronously and serially. On the surface, it might seem that if the **Store** service invoked its helper objects asynchronously, it would increase throughput, because it could process incoming orders as fast as the client submitted them. The problem in doing so is that it is possible for the calls to the **Order**, **Shipment**, and **Billing** services to fail independently, and if they do, all hell will break loose. For example, the **Order** service might discover that there were no items in the inventory matching the customer request, while the **Shipment** service tried to ship the nonexistent item and the **Billing** service had already billed the customer for it.

Using asynchronous calls on a set of interacting services requires that you change your code and your workflow. As illustrated in [Figure 8-5](#), to call the helper services asynchronously, you need to string them together. The **Store** service should call only the **Order** service, which in turn should call the **Shipment** service only if the order processing was successful, to avoid the potential inconsistencies just mentioned. Similarly, only in the case of successful shipment should the **Shipment** service asynchronously call the **Billing** service.

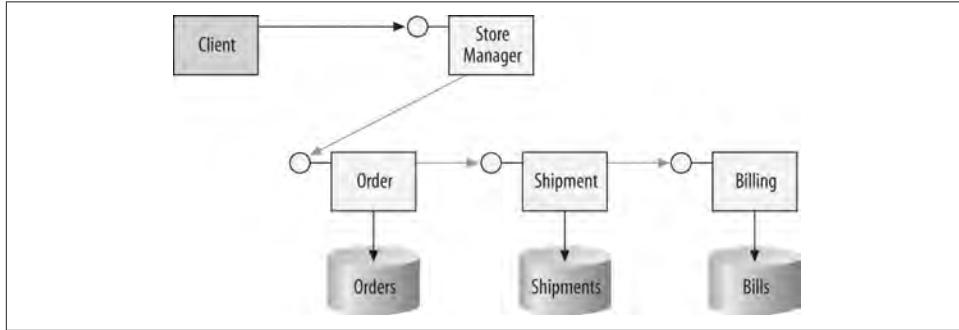


Figure 8-5. Revised workflow for asynchronous processing of an order

In general, if you have more than one service in your asynchronous workflow, you should have each service invoke the next one in the logical execution sequence. Needless to say, such a programming model introduces tight coupling between services (they have to know about one another) and changes to their interfaces (you have to pass in additional parameters, which are required for the desired invocation of services downstream).

The conclusion is that using asynchronous instead of synchronous invocation introduces major changes to the service interfaces and the client workflow. Asynchronous invocation on a service that was built for synchronous execution works only in isolated cases. When dealing with a set of interacting services, it is better to simply spin off a worker thread to call them and use the worker thread to provide asynchronous execution. This will preserve the service interfaces and the original client execution sequence.

Asynchronous Service Operations

The service-side also supports the task-based programming model. This allows you to apply the `async` and `await` pattern to a service operation to make its internal *execution* asynchronous. This means that while the internal flow of your service operation will be asynchronous at the points where you apply `await`, the operation type of the underlying call will not change. WCF does this by detecting that your service

operation is task-based and managing an `await` on the operation invocation internally.

You apply the tasked-based programming model to a service operation as shown in [Example 8-44](#). It is important to note that if the operation type is request-reply, even though your service operation will suspend execution on the awaited proxy call, WCF will not return a reply to the caller until the final continuation within your service operation completes.

Example 8-44. Asynchronous service operation

```
[ServiceContract]
interface IMyOtherContract
{
    [OperationContract]
    Task<int> MyOtherMethodAsync();
}

class MyOtherClient : ClientBase<IMyOtherContract>, IMyOtherContract
{...}

[ServiceContract]
interface IMyContract
{
    [OperationContract]
    Task<int> MyMethodAsync();
}

class MyService : IMyContract
{
    public async Task<int> MyMethodAsync()
    {
        MyOtherClient proxy = new MyOtherClient();

        int result = await proxy.MyOtherMethodAsync();

        //Do work in continuation

        //Eventually reply
        return result;
    }
}
```

In [Example 8-44](#), `MyMethodAsync()` will not return a reply to the caller until the task returned from `MyOtherMethodAsync()` completes and the compiler-generated support for `await` executes the remainder of `MyMethodAsync()` as a continuation.

Asynchronous service operations and timeouts

As with proxy-based asynchronous calls, an asynchronous service operation is nothing more than a convenient non-blocking programming model on top of an actual synchronous operation. The underlying synchronous operation can still timeout. You should not equate asynchronous service operations with long running operations. This behavior can mislead developers who may infer that by making a service operation asynchronous they somehow change the semantics of the underlying operation type so that the operation returns a reply to the client at the point of the first `await`. Developers then incorrectly assume that by changing the operation type their service operation is no longer subjected to timeouts. Another common pitfall is making subsequent asynchronous proxy calls within an asynchronous service operation, as shown in [Example 8-44](#). Although the service operation and the proxy call are both asynchronous, all WCF timeouts are still in play across all the parties involved.

Asynchronous service operation execution

It is important to note in [Example 8-44](#) that the part of `MyMethodAsync()` leading up to the point of `await` WCF will execute on a thread from the I/O completion thread pool. While the awaited task and its continuation .NET will execute on separate threads from the .NET thread pool. This means that the part of `MyMethodAsync()` leading up to the point of `await` executes under the context of WCF, while the remainder does not.



For asynchronous service operations on services with `PerCall` instancing that implement `IDisposable`, .NET will also call `IDisposable.Dispose()` on a separate thread from the .NET thread pool.

Asynchronous service operations and context

Within your asynchronous service operation upon crossing the threshold of the first `await` you lose nearly all WCF context. This includes among other ambient aspects `OperationContext.Current` and `Transaction.Current`. For example, the following code has no operation context after the `await`:

```
class MyService : IMyContract
{
    public async Task<int> MyMethodAsync()
    {
        MyOtherClient proxy = new MyOtherClient();

        int result = await proxy.MyOtherMethodAsync();
        Debug.Assert(OperationContext.Current == null);
        ...
    }
}
```

```
    }  
}
```

The implications on the WCF programming model are severe because you lose the simplified programming model afforded by the ambient `OperationContext`. Of particular importance is the loss of access to the incoming and outgoing message headers for the operation which WCF makes available to you through `OperationContext` (a technique used throughout this book and detailed in [Appendix B](#)). The end result is that out of the box, asynchronous service operations no longer support foundational, modern service-oriented programming techniques such as continuing the flow of out-of-band parameters across all calls and service interactions.

For example, you can easily build upon the priority processing techniques described previously in this chapter to convey a call's priority via message headers from one service to another. In the following example, the `ServicePriorityClientBase` class simplifies the programming model and enforces this behavior by implicitly propagating the current call's priority in its own outgoing calls:

```
public abstract class ServicePriorityClientBase<T> : PriorityClientBase<T>  
    where T : class  
{  
    public ServicePriorityClientBase() : base(GenericContext<CallPriority>.  
        Current.Value)  
    {}  
}
```

Unfortunately, these powerful techniques will not work out of the box within an asynchronous service operation when you try and construct a priority-aware call chain:

```
//Client side  
class MyContractClient : PriorityClientBase<IMyContract>, IMyContract  
{  
    public MyContractClient(CallPriority priority) : base(priority)  
    {}  
    public void MyMethod()  
    {  
        Channel.MyMethod();  
    }  
}  
  
MyContractClient proxy = new MyContractClient(CallPriority.High);  
proxy.MyMethod();  
proxy.Close();  
  
//Service-side  
class MyOtherContractClient : ServicePriorityClientBase<IMyOtherContract>,  
    IMyOtherContract  
{  
    public Task<int> MyOtherMethodAsync()
```

```

        {
            return Channel.MyOtherMethodAsync();
        }
    }
[PriorityCallsBehavior(...)]
class MyService : IMyContract
{
    public async Task<int> MyMethodAsync()
    {
        Debug.Assert(GenericContext<CallPriority>.Current != null);

        //Implicitly passes priority in call
        MyOtherContractClient proxy1 = new MyServicePriorityClient();

        int result = await proxy1.MyOtherMethodAsync();

        //Call priority no longer available to include in subsequent calls
        Debug.Assert(GenericContext<CallPriority>.Current == null);

        //This call will fail
        MyOtherContractClient proxy2 = new MyServicePriorityClient();

        result = await proxy2.MyOtherMethodAsync();

        return result;
    }
}
[PriorityCallsBehavior(...)]
class MyOtherService : IMyOtherContract
{
    public async Task<int> MyOtherMethodAsync()
    {
        int result = 0;

        Debug.Assert(GenericContext<CallPriority>.Current != null);

        return result;
    }
}

```

Gone also is the fluid declarative programming model of WCF's support for transaction flow. For example, the following code will lose the ambient transaction after the `await` call:

```

class MyService : IMyContract
{
    [OperationBehavior(TransactionScopeRequired = true)]
    public async Task<int> MyMethodAsync()
    {
        MyOtherClient proxy = new MyOtherClient();

```

```

        int result = await proxy.MyOtherMethodAsync();
        Debug.Assert(Transaction.Current == null);
        ...
    }
}

```

And if you attempt to use the Client transaction mode by setting `TransactionFlow` to `true` in the binding, `TransactionFlowOption.Mandatory` in the operation contract and `TransactionScopeRequired` to `true` in the service operation behavior, any subsequent proxy calls you make will fail:

```

[ServiceContract]
interface IMyOtherContract
{
    [OperationContract]
    [TransactionFlow(TransactionFlowOption.Mandatory)]
    Task<int> MyOtherMethodAsync();
}

class MyService : IMyContract
{
    [OperationBehavior(TransactionScopeRequired = true)]
    public async Task<int> MyMethodAsync()
    {
        MyOtherClient proxy = new MyOtherClient();

        int result = await proxy.MyOtherMethodAsync();

        //This call will fail
        result = await proxy.MyOtherMethodAsync();

        return result;
    }
}

```

A similar problem will occur if you attempt to encapsulate an awaited proxy call within a `TransactionScope`. Upon the return of the awaited call, the transaction scope cannot complete because its associated transaction is no longer available. In this case, `Complete()` will always throw an exception:

```

class MyService : IMyContract
{
    public async Task<int> MyMethodAsync()
    {
        MyOtherClient proxy = new MyOtherClient();

        int result = 0;

        using(TransactionScope scope = new TransactionScope())
        {
            result = await proxy.MyOtherMethodAsync();
            //Complete() will fail
        }
    }
}

```

```

        scope.Complete();
    }

    return result;
}
}

```

.NET 4.5.1 and above improve upon this problem by introducing the `TransactionScopeAsyncFlowOption` enumeration:

```

enum TransactionScopeAsyncFlowOption
{
    Suppress,
    Enabled
}

```

And a supporting set of new `TransactionScope` constructors:

```

class TransactionScope : ...
{
    public TransactionScope(TransactionScopeAsyncFlowOption asyncFlowOption);
    //More constructors
}

```

To use this new feature, you set `TransactionScopeAsyncFlowOption.Enabled` when you create your `TransactionScope`:

```

class MyService : IMyContract
{
    public async Task<int> MyMethodAsync()
    {
        MyOtherClient proxy = new MyOtherClient();

        int result = 0;

        using(TransactionScope scope = new TransactionScope(
            TransactionScopeAsyncFlowOption.
            Enabled))
        {
            result = await proxy.MyOtherMethodAsync();
            //Complete() will not fail
            scope.Complete();
        }

        return result;
    }
}

```

The problem with this approach is that you must explicitly opt in to flow a transaction across each awaited call you make, making the programming model somewhat error prone.

In an attempt to solve the loss of the `OperationContext`, you could choose to use `OperationContextScope` (detailed in [Appendix B](#)) to spin a new context for your proxy call and have it automatically reinstate the original context when you dispose of it:

```
public sealed class OperationContextScope : IDisposable
{
    public OperationContextScope(IContextChannel channel);
    public OperationContextScope(OperationContext context);
    public void Dispose();
}

using(OperationContextScope scope = new OperationContextScope(...))
{
    //Do work with new context
    ...
}//Restores previous context here
```

This technique would allow you in theory to flow headers in your call by copying them from the original operation context to the new one before making your call. Unfortunately, this technique will not work. An operation context scope cannot span threads. Meaning, you cannot create an operation context scope on one thread and dispose of it on another. This is exactly what happens when you try to encapsulate an awaited proxy call within an operation context scope:

```
class MyService : IMyContract
{
    public async Task<int> MyMethodAsync()
    {
        int result = 0;
        MyOtherClient proxy = new MyOtherClient();

        using(OperationContextScope scope = new OperationContextScope
            (proxy.InnerChannel))
        {
            result = await proxy.MyOtherMethodAsync();
        }
        //Exception will happen here
        return result;
    }
}
```

To circumvent this issue, you may choose to capture the current operation context before making your proxy call and then use this context after your call returns to spin a new operation context scope. You could then work against the original operation context within your continuation:

```
class MyService : IMyContract
{
    public async Task<int> MyMethodAsync()
```

```

{
    MyOtherClient proxy = new MyOtherClient();

    OperationContext context = OperationContext.Current;

    int result = await proxy.MyOtherMethodAsync();
    using(OperationContextScope scope = new OperationContextScope(context))
    {
        //OperationContext available here
        Debug.Assert(OperationContext.Current != null);
    }

    //OperationContext no longer available here
    Debug.Assert(OperationContext.Current == null);

    return result;
}
}

```

The problem with this approach is that each subsequent call you need to make will require you to spin a new operation context scope. This results in a very unwieldy programming model. It also does nothing to provide access to the original operation context for use within each subsequent proxy call. Nor does it reinstate the original operation context before your operation returns so that you may as an option pass headers back to the original caller:

```

class MyService : IMyContract
{
    public async Task<int> MyMethodAsync()
    {
        MyOtherClient proxy = new MyOtherClient();

        OperationContext context = OperationContext.Current;

        int result = await proxy.MyOtherMethodAsync();
        using(OperationContextScope scope = new OperationContextScope(context))
        {
            //OperationContext available here
            Debug.Assert(OperationContext.Current != null);
        }

        //OperationContext no longer available here
        Debug.Assert(OperationContext.Current == null);

        result = await proxy.MyOtherMethodAsync();
        using(OperationContextScope scope = new OperationContextScope(context))
        {
            //OperationContext available here
            Debug.Assert(OperationContext.Current != null);
        }

        //OperationContext no longer available here
    }
}

```

```

        Debug.Assert(OperationContext.Current == null);

        return result;
    }
}

```

Instead of using `OperationContextScope`, you might simply attempt to capture the ambient `OperationContext` at the start of your operation and access this reference in the continuation of each `await`:

```

class MyService : IMyContract
{
    public async Task<int> MyMethodAsync()
    {
        ServiceHost host1 = OperationContext.Current.Host;

        MyOtherClient proxy = new MyOtherClient();

        OperationContext operationContext = OperationContext.Current;

        int result = await proxy.MyOtherMethodAsync();

        //Access the original context, such as:
        ServiceHost host2 = operationContext.Host;

        Debug.Assert(host1 == host2);

        result = await proxy.MyOtherMethodAsync();

        ServiceHost host3 = operationContext.Host;
        Debug.Assert(host2 == host3);

        return result;
    }
}

```

This technique works well within the local scope of your operation. But it does nothing to reinstate the original context on any new threads the system may use so that you can support other objects that expect the previous context. It also does not support any code you may have within WCF extensibility points that expects `OperationContext` to exist during subsequent proxy calls.

To solve these problems, you might attempt to capture the ambient `OperationContext`, `Transaction`, and `SynchronizationContext` at the start of your operation and reinstate them in the continuation of each `await`:

```

class MyService : IMyContract
{
    [OperationBehavior(TransactionScopeRequired = true)]
    public async Task<int> MyMethodAsync()
    {

```

```

MyOtherClient proxy = new MyOtherClient();

OperationContract operationContext = OperationContext.Current;
Transaction transaction = Transaction.Current;
SynchronizationContext syncContext = SynchronizationContext.Current;

int result = await proxy.MyOtherMethodAsync();
OperationContext.Current = operationContext;
Transaction.Current = transaction;
SynchronizationContext.SetSynchronizationContext(syncContext);

//Access contexts and Transaction here
result = await proxy.MyOtherMethodAsync();

OperationContract.Current = operationContext;
Transaction.Current = transaction;
SynchronizationContext.SetSynchronizationContext(syncContext);

//Access contexts and Transaction here
return result;
}

```

While it does work, any manual technique such as this will produce a very brittle programming model for your developers which will ultimately degrade the quality of your codebase over its lifetime. Every time a developer has to make a change to your code, they will have to remember to apply the manual pattern you've established. This produces a very error prone approach to programming. And if they forget, they will only detect the failure at runtime.

Restoring service operation context

We consider the lack of parity with standard WCF behavior introduced by asynchronous service operations a design flaw of WCF or at the very least a regretful oversight. The manually intensive solutions shown thus far of reinstating each WCF relevant context for each asynchronous proxy call are far from ideal. To address this concern, the sections that follow present several techniques increasing in sophistication you can use to simplify your programming model for reestablishing WCF relevant context flow in your asynchronous service operations.



It is important to note that the techniques that follow are only applicable to asynchronous service operations that maintain a sequential non-blocking execution flow using the `async` and `await` pattern. Using these techniques in task-based composition scenarios such as `WaitAll()` or `WhenAll()` will produce indeterminate results. This is because, unlike the `async` and `await` pattern, asynchronous combinatory scenarios do not provide you any control over the order of operation of your tasks. A lack of control over the order of operation of your tasks will inevitably lead to state corruption within the ambient contexts of your asynchronous service operation.

Instead of manually reinstating each WCF relevant context for every asynchronous proxy call, you can create a helper to encapsulate these manual steps and simplify your programming model. Your helper would be easiest to use if it shared a programming model common to other WCF or .NET concepts. Since your design goal is to restore all contexts relevant to WCF after the scope of an asynchronous proxy call, the .NET concept of a *scope* seems the most appropriate programming model.

To this end, *ServiceModelEx* provides the `AsyncContextScope` helper class. `AsyncContextScope` manages the capture and restoration of `OperationContext.Current`, `Transaction.Current` and `SynchronousContext.Current` for each asynchronous proxy call you make. Upon construction, `AsyncContextScope` captures the relevant ambient contexts and then restores them back to the ambient context accessors once you dispose of or close the scope.

`AsyncContextScope` is defined as:

```
public class AsyncContextScope : IDisposable
{
    public AsyncContextScope();
    public AsyncContextScope(AsyncContextScope asyncContextScope);
    public void Dispose();
    public void Close();
    //More members
}
```

You use `AsyncContextScope` much like you would `OperationContextScope` or `TransactionScope`:

```
class MyService : IMyContract
{
    [OperationBehavior(TransactionScopeRequired = true)]
    public async Task<int> MyMethodAsync()
    {
        MyOtherClient proxy = new MyOtherClient();

        int result;
```

```

        using(AsyncContextScope asyncScope = new AsyncContextScope())
        {
            result = await proxy.MyOtherMethodAsync();
        }
        Debug.Assert(OperationContext.Current != null);

        return result;
    }
}

```

You can also nest multiple `AsyncContextScopes` by forwarding the outer scope into the inner scope through the appropriate constructor:

```

class MyService : IMyContract
{
    [OperationBehavior(TransactionScopeRequired = true)]
    public async Task<int> MyMethodAsync()
    {
        MyOtherClient proxy = new MyOtherClient();

        int result;
        using(AsyncContextScope asyncScope1 = new AsyncContextScope())
        {
            result = await proxy.MyOtherMethodAsync();
            using(AsyncContextScope asyncScope2 = new AsyncContextScope
                ( asyncScope1))
            {
                result = await proxy.MyOtherMethodAsync();
            }
            Debug.Assert(OperationContext.Current != null);
        }
        Debug.Assert(OperationContext.Current != null);

        return result;
    }
}

```

`AsyncContextScope` uses the helper class `AsyncContext` that encapsulates all three contexts. `AsyncContext` internally captures and reinstates all contexts relevant to WCF. Through its various constructors you can even provide to `AsyncContext` an existing `AsyncContext` or specific instances of each context. `AsyncContext`'s default constructor captures and stores the current ambient value for each context, if one exists:

```

public class AsyncContext
{
    public readonly OperationContext OperationContext;
    public readonly Transaction Transaction;
    public readonly SynchronizationContext SynchronizationContext;

    public AsyncContext() : this(OperationContext.Current, Transaction.Current,
                                SynchronizationContext.Current)

```

```

    {}
    public AsyncContext(AsyncContext asyncContext) : this(asyncContext.
        OperationContext,
        asyncContext.Transaction,
        asyncContext.Synchronization
        Context)
    {}
    public AsyncContext(OperationContext operationContext, Transaction transaction,
                        SynchronizationContext syncContext)
    {
        OperationContext = operationContext;
        Transaction = transaction;
        SynchronizationContext = syncContext;
    }
    public void Restore()
    {
        Transaction.Current = Transaction;
        SynchronizationContext.SetSynchronizationContext(SynchronizationContext);
        OperationContext.Current = OperationContext;
    }
}

```

Example 8-45 shows the implementation of `AsyncContextScope`. You can use it as a model of how to create your own asynchronous context scope. `AsyncContextScope` also provides direct access within the scope to all contexts through the `AsyncContext` property.

Example 8-45. The AsyncContextScope class

```

public class AsyncContextScope : IDisposable
{
    public AsyncContext AsyncContext
    {get;private set;}

    public AsyncContextScope()
    {
        AsyncContext = new AsyncContext();
    }
    public AsyncContextScope(AsyncContextScope asyncContextScope)
    {
        AsyncContext = new AsyncContext(asyncContextScope.AsyncContext);
        AsyncContext.Restore();
    }
    public void Dispose()
    {
        Close();
    }
    public void Close()
    {
        AsyncContext.Restore();
    }
}

```

```
    }  
}
```

While `AsyncContextScope` certainly helps simplify your asynchronous service operation programming model, the scope technique still represents a manual programming effort that is less than ideal. Nor does it coincide with the fluid programming model that the `async` and `await` pattern promotes.

Alternatively, you can choose to create a custom `TaskAwaiter`. Briefly, an *awaiter* is a programming construct .NET uses within its compiler-generated support for `await` that defines pre and post await events. Creating a custom awaierter allows you to achieve the same result as `AsyncContextScope` but with a simpler, more fluid task-based programming model.

In general, creating a custom awaierter is an advanced task-based programming technique that has nothing to do with WCF. However in this case, you do not need to write your custom awaierter from the ground up. Instead, you simply need to follow the .NET awaierter pattern and wrap the .NET's existing `TaskAwaiter` in your custom awaierter. This approach provides you access to `await`'s pre and post execution events without having to concern yourself with the complex task-based security and synchronization details that accompany writing a custom awaierter from scratch.

.NET's awaierter pattern prescribes only that an object instance provides a `GetAwaiter()` method, which returns a type that implements the `INotifyCompletion` interface from the `System.Runtime.CompilerServices` namespace:

```
public interface INotifyCompletion  
{  
    void OnCompleted(Action continuation);  
}
```

In addition, the awaierter must support the following two members:

```
bool IsCompleted  
{get;}  
  
void GetResult();
```

If the task associated with the awaierter returns a value of the type `T`, then `GetResult()` must also return that value:

```
T GetResult();
```

The `INotifyCompletion` method of `OnCompleted()` represents the awaierter's pre-await event and `GetResult()` represents the awaierter's post-await event.

To flow context across your asynchronous proxy call, you would capture the ambient contexts within `OnCompleted()` and restore them in `GetResult()`. You can then leverage the convenience of `AsyncContext` in conjunction with the standard

`TaskAwaiter` to achieve context flow. By providing a unique awainer per proxy call you also greatly simplify the design of your awainer.

You can use `WcfContextAwaiter` and `WcfContextAwaiter<T>` defined in `ServiceModelEx` in writing your awainer:

```
public class WcfContextAwaiter : INotifyCompletion
{
    public virtual WcfContextAwaiter GetAwaiter();
    public virtual bool IsCompleted
    {get;}
    public virtual void OnCompleted(Action continuation);
    public void GetResult();
}

public class WcfContextAwaiter<T> : WcfContextAwaiter
{
    public new WcfContextAwaiter<T> GetAwaiter();
    public override bool IsCompleted
    {get;}
    public override void OnCompleted(Action continuation);
    public new T GetResult();
}
```

`ServiceModelEx` also provides the `FlowWcfContext()` extension methods for the `Task` and `Task<T>` that complement the `async` and `await` pattern:

```
public static class AsyncExtensions
{
    public static WcfContextAwaiter FlowWcfContext(this Task task)
    {
        return new WcfContextAwaiter(task);
    }
    public static WcfContextAwaiter<T> FlowWcfContext<T>(this Task<T> task)
    {
        return new WcfContextAwaiter<T>(task);
    }
}
```

You apply `FlowWcfContext()` to each asynchronous proxy call. This means you create a unique awainer for each proxy call thus ensuring that the contexts the awainer captures are unique to each call:

```
class MyService : IMyContract
{
    [OperationBehavior(TransactionScopeRequired = true)]
    public async Task<int> MyMethodAsync()
    {
        int result = 0;
        MyOtherClient proxy = new MyOtherClient();

        result = await proxy.MyOtherMethodAsync().FlowWcfContext();
```

```

        Debug.Assert(OperationContext.Current != null);
        Debug.Assert(Transaction.Current != null);

        result = await proxy.MyOtherMethodAsync().FlowWcfContext();
        Debug.Assert(OperationContext.Current != null);
        Debug.Assert(Transaction.Current != null);

        return result;
    }
}

```

The Async Context Synchronizer

The refinements for restoring context within your asynchronous service operations presented thus far will certainly work. But their manual programming models burden developers with boilerplate code and still allow them to make mistakes. Perhaps the most appropriate solution to this problem would be one that returns to WCF the implicit ambient programming model and continuity of context that existed within service operations before .NET introduced the `async` and `await` pattern. You can renew this continuity by creating and installing a custom synchronization context that captures all the relevant ambient contexts of an operation instance and manages their restoration upon each continuation within your service operation. This capability will be especially valuable with pass-through edge servers and third-party integration services that possess solely I/O-bound workloads (as detailed later in this chapter) where asynchronous service operations shine. To this end, `ServiceModelEx` provides the `AsyncContextSynchronizer` class. `AsyncContextSynchronizer` derives from `ThreadPoolSynchronizer`, defined as:

```

public class AsyncContextSynchronizer : ThreadPoolSynchronizer
{
}

```

`AsyncContextSynchronizer` derives from `ThreadPoolSynchronizer` because it requires a dedicated thread pool that installs on each thread within the pool an instance of `AsyncContextSynchronizer` as the current synchronization context. Once installed for your service, `AsyncContextSynchronizer` receives method invocation notifications for both WCF service operations and .NET task continuations.

Upon each method invocation, `AsyncContextSynchronizer` captures the WCF relevant ambient contexts, if they exist. A thread from the `AsyncContextSynchronizer` thread pool then processes the method invocation. The details of `AsyncContextSynchronizer`'s implementation have nothing to do with WCF and are therefore beyond the scope of this book. Briefly, `AsyncContextSynchronizer` uses .NET's `ExecutionContext` concept to store each service operation's associated `AsyncContext` so that it may access it across multiple threads. `ExecutionContext` allows .NET to maintain a *logical call context* in asynchronous environments without the use of thread local storage (TLS) by supporting the flow of ambient contexts and logical call state from

thread to thread. Every asynchronous aspect of the .NET framework (except those marked as `Unsafe`) uses `ExecutionContext` to flow ambient context and state from thread to thread.

When processing a method invocation, `AsyncContextSynchronizer` first determines whether a method invocation is a service operation or a task continuation. When processing a service operation, `AsyncContextSynchronizer` saves the associated `AsyncContext` to the active `ExecutionContext` before invoking the service operation method. When processing a task continuation, `AsyncContextSynchronizer` first retrieves the execution context associated with the continuation. It then temporarily restores the task continuation's execution context so that it may retrieve the saved `AsyncContext`. Through `AsyncContext` the `AsyncContextSynchronizer` then restores to the thread that will execute the task continuation the service operation's WCF relevant ambient contexts. Upon completion, `AsyncContextSynchronizer` restores to the current thread the execution context that existed on that thread before the call.

In this way, `AsyncContextSynchronizer` returns to WCF the implicit ambient programming model and continuity of context that existed previously within service operations without the manual programming concerns of other approaches such as `AsyncContextScope` or `WcfContextAwaiter<T>`.

You can install `AsyncContextSynchronizer` using any of the techniques described previously for installing a custom synchronization context at the service, endpoint or host level. Within `ServiceModelEx`, you will find the `AsyncContextBehaviorAttribute` and the `SetAsyncContext()` extension method that you can use to install `AsyncContextSynchronizer` for your services.

For example, if by design only some of your services will contain asynchronous service operations, you can apply `AsyncContextBehaviorAttribute` to each instead of being at the mercy of the host and the thread that happens to open it:

```
[AsyncContextBehavior(5,typeof(MyService))]
class MyService : IMyContract
{...}
```

`AsyncContextBehaviorAttribute` is a specialization of `ThreadPoolBehaviorAttribute` that overrides `ProvideSynchronizer()` as shown:

```
public class AsyncContextBehaviorAttribute : ThreadPoolBehaviorAttribute
{
    public AsyncContextBehaviorAttribute(uint poolSize,Type serviceType) : this(
        poolSize,serviceType,"Async Context Pool Thread: ")
    {}
    public AsyncContextBehaviorAttribute(uint poolSize,Type serviceType,
                                         string threadName) : base(poolSize,serviceType,
                                         threadName)
    {}
    protected override ThreadPoolSynchronizer ProvideSynchronizer()
```

```

    {
        if(ThreadPoolHelper.HasSynchronizer(ServiceType) == false)
        {
            return new AsyncContextSynchronizer(PoolSize,PoolName);
        }
        else
        {
            return ThreadPoolHelper.GetSynchronizer(ServiceType);
        }
    }
}

```

CPU-bound versus I/O-bound operations

Even if the `async` and `await` pattern did work out of the box with standard WCF behavior on the service-side it would still do little to improve the throughput or performance of your business-related services. Service operations that perform common business activities, such as calculations and data transformation, represent CPU-intensive or *CPU-bound* activities. Regardless of whether you initiate them synchronously or asynchronously, these activities consume CPU resources and require a dedicated thread throughout their execution. In contrast, operations that interact with input and output devices, such as network or communication ports, represent I/O intensive or *I/O-bound* activities. These activities behave differently when you initiate them asynchronously. Asynchronous I/O-bound activities do not consume CPU resources nor do they require a dedicated thread during the entirety of their execution as they wait for a response from the device. Instead, Windows uses a small pool of threads to efficiently wait for device responses, signaling the appropriate application when a given I/O response arrives.

The majority of operations in a service-oriented application perform some type of business-related activity. This makes sense since you commonly use services to encapsulate and centralize your system's business logic. This means that most service-oriented applications perform more CPU-bound, rather than I/O-bound, activity. Here, asynchronous service operations will not help you. Asynchronous or not, all this CPU-bound activity will consume the same amount of CPU resources and require the same number of threads to execute. In fact, you often incur higher CPU utilization when you introduce asynchronous activity into your service operations due to the additional thread context switching you impose upon the system. The net result is that under load, your system's throughput and performance will be much the same if you use asynchronous service operations as it would be if your operations were synchronous. In the case of CPU-bound activity, by introducing asynchronous service operations you have increased the overall complexity of your system and your programming model without gaining any appreciable benefit in throughput, performance or responsiveness. For this reason, it is best practice in general not to wrap your CPU-bound activities with the `async` and `await` pattern. Instead, you should

leave CPU-bound activity synchronous and let the consumers of your API or services decide whether they need to invoke your code asynchronously or not.

Due to the loss of parity out of the box with standard WCF, the many programming concerns it presents and most of all the unwarranted complexity it introduces, we recommend that you avoid using asynchronous service operations in WCF with your business-related services.

Service operation parallelism

There are scenarios where leveraging tasks to introduce asynchronous or parallel activities into your service operations will in fact provide additional throughput and performance benefit without the need to resort to asynchronous service operations. The patterns that follow are most useful in service-oriented infrastructure scenarios where you can leverage calls that are inherently asynchronous and one-way, such as publish/subscribe, discovery, and announcements. Even so, you may occasionally find them useful in the context of your business-related services as well. As per the previous discussion on CPU-bound versus I/O-bound operations, these patterns are best suited to parallelize your I/O-bound activities. Since your business-related operations are often already parallelized per session or per call based on your service's concurrency mode, attempting to further parallelize your CPU-bound service operations using these patterns will only serve to degrade your system's throughput and performance under load.

For example, in scenarios where you need true asynchronous fire-and-forget such as events broadcast, you can efficiently broadcast events asynchronously and in parallel using tasks. To simplify this pattern, *ServiceModelEx* provides the helper method `ForEachAsync<T>()`, which is an extension method of `IEnumerable<T>`. The `ForEachAsync<T>()` method iterates over the associated `IEnumerable<T>` creating a task for each element within the collection. These tasks each perform a unit of work defined by an `Action<T>` delegate operating on a given item in the collection. `ForEachAsync<T>()` will invoke this delegate on each item in the collection, but do it all asynchronously and in parallel. `ForEachAsync<T>()` will not block its caller or wait for the actions to complete. It therefore will return almost instantly, having blocked its caller for the brief time it takes to dispatch all the parallel calls. `ForEachAsync<T>()` will also return the list of tasks it has dispatched:

```
public static class CollectionExtensions
{
    public static IEnumerable<Task> ForEachAsync<T>(this IEnumerable<T>
        collection, Action<T> action)
    {
        List<Task> tasks = new List<Task>();

        foreach(T item in collection)
        {
```

```

        tasks.Add(Task.Run(()=> action(item)));
    }
    return tasks;
}
}

```

For example, to ensure efficient broadcast of callbacks to all clients in scenarios where you are not concerned about callback exceptions, you can easily update [Example 5-6](#) to use `ForEachAsync<T>()`:

```

class MyService : IMyContract
{
    static List<IMyContractCallback> m_Callbacks = new List
        <IMyContractCallback>();

    public void DoSomething()
    {
        IMyContractCallback callback = OperationContext.Current.GetCallbackChannel
            <IMyContractCallback>();

        if(m_Callbacks.Contains(callback) == false)
        {
            m_Callbacks.Add(callback);
        }
    }
    public static void CallClients()
    {
        Action<IMyContractCallback> invoke = callback=>callback.OnCallback();
        m_Callbacks.ForEachAsync(invoke);
    }
}

```

`ForEachAsync<T>()` allows you to keep your service operation timeouts short to avoid the poor service responsiveness that often accompanies lengthy timeouts. It is also important to note that this example does not exhibit the problems described with asynchronous service operations. All WCF context is still available once `ForEachA sync<T>()` returns.

There are also scenarios where you may need to have your tasks execute in parallel, but block synchronously for all tasks to complete so that you can perform further related work in your operation such as closing and cleaning up shared resources. To achieve this, you can use the existing `Parallel.ForEach<T>()` helper in the `System.Threading.Tasks` namespace:

```

public struct ParallelLoopResult
{...}

public static class Parallel
{
    public static ParallelLoopResult ForEach<T>(IEnumerable<T> source,
                                                Action<T> body);
}

```

```

}

class MyService : IMyContract
{
    ...

    public static void CallClients()
    {
        Action<IMyContractCallback> invoke = callback=>callback.OnCallback();
        try
        {
            Parallel.ForEach(m_Callbacks, invoke);

            //Do further work once completed
        }
        catch(AggregateException exception)
        {
            ...
        }
    }
}

```

When you use `Parallel.ForEach<T>()` within a service operation, all the parallel tasks it creates must complete either successfully or in error before `Parallel.ForEach<T>()` returns control to you. Because your call to `Parallel.ForEach<T>()` itself is fully synchronous, all WCF context is available after `Parallel.ForEach<T>()` completes. `Parallel.ForEach<T>()` will also aggregate exceptions from all tasks that complete in error and provides them to you in an instance of `AggregateException`. The tasks `Parallel.ForEach<T>()` creates are autonomous by default, so failing tasks do not prevent other tasks from completing successfully.

To complement this pattern, `ForEachAsync<T>()` returns an `IEnumerable<Task>` allowing you to wait on all tasks. This produces an equivalent programming model to `Parallel.ForEach<T>()`:

```

class MyService : IMyContract
{
    ...

    public static void CallClients()
    {
        Action<IMyContractCallback> invoke = callback=>callback.OnCallback();
        try
        {
            Task.WaitAll(m_Callbacks.ForEachAsync(invoke).ToArray());

            //Do further work once completed
        }
        catch(AggregateException exception)
        {

```

```
        }
    }
}
```

The patterns shown in the previous two examples still subject your service operation to possible timeout concerns. This makes them no better than initiating a non-blocking execution flow using `async` and `await` and then explicitly waiting on a subsequent task. To preserve your original design goal of efficient asynchronous I/O processing and to insulate your service operation from timeout concerns, you can spin another task to encapsulate the wait:

```
class MyService : IMyContract
{
    ...

    public static void CallClients()
    {
        Action<IMyContractCallback> invoke = callback=>callback.OnCallback();

        Action process = delegate
        {
            try
            {
                Task.WaitAll(m_Callbacks.ForEachAsync(invoke).
                    ToArray());
            }
            catch(AggregateException exception)
            {
                ...
            }
        };
        Task.Run(process);
    }
}
```

You can also pair `ForEachAsync<T>()` with `Task.WhenAll()` or `Task.WhenAny()` so that you may provide asynchronous completion when all tasks complete or when any of the tasks complete respectively. Unlike `WaitAll()`, `WhenAll()` and `WhenAny()` return a task upon which you can wait. If you instead do nothing with the returned task, they will immediately return control to the caller. In the following example, `MyMethod()` will exit immediately, all actions will execute asynchronously and in parallel and the system will only call the `dowork` continuation delegate some time later when all tasks complete:

```
class MyService : IMyContract
{
    ...

    public static void CallClients()
```

```

    {
        Action<IMyContractCallback> invoke = callback=>callback.OnCallback();

        Action<Task> doWork = (task)=>
        {
            //Do work asynchronously once completed
        };
        Task.WhenAll(m_Callbacks.ForEachAsync(invoke)).ContinueWith(doWork);
    }
}

```

Note that when you do not wait upon the task `WhenAll()` returns, `AggregateException` is not thrown into the caller's context. This is acceptable because the system assigns the `AggregateException` to the antecedent task which you will receive in your `doWork` completion method.

To round out a consistent and simplified programming model, `ServiceModelEx` also provides the `ParallelForEach<T>()` helper method that wraps `Parallel.ForEach<T>()`. The result is a fluid, almost natural-language-like code. `ParallelForEach<T>()` spares developers from having to interact with the `Parallel` class directly. You can also use it as an example of how to control which of the many `Parallel.ForEach<T>()` overloads you wish to support:

```

public static class CollectionExtensions
{
    public static void ParallelForEach<T>(this IEnumerable<T> collection,
                                            Action<T> action)
    {
        Parallel.ForEach(collection,action);
    }
}

class MyService : IMyContract
{
    ...

    public static void CallClients()
    {
        Action<IMyContractCallback> invoke = callback=>callback.OnCallback();

        Action process = delegate
        {
            try
            {
                m_Callbacks.ParallelForEach(invoke);
            }
            catch(AggregateException exception)
            {
                ...
            };
        };
    }
}

```

```
        Task.Run(process);  
    }  
}
```

Queued Services

WCF enables disconnected work: the client posts messages to a queue, and the service processes them. Such interaction enables different possibilities from those presented so far, and in turn, a different programming model. This chapter starts by showing you how to set up and configure simple queued services, and then focuses on aspects such as transactions, instance management, and failures, and their impact on both the business model of the service and its implementation. The chapter ends with my solution for a response service and a discussion of using the HTTP bridge for queued calls over the Internet.

Disconnected Services and Clients

The previous chapters were all predicated on a connected interaction between the client and the service, where both sides must be up and running to be able to interact with each other. However, there are quite a few cases (as well as the overall business model justification) for wanting to have disconnected interaction in a service-oriented application:

Availability

The client may need to work against the service even when the client is disconnected—for example, when using a mobile device. The solution is to queue up requests against a local queue and send them to the service when the client is connected. Similarly, if the service is offline (perhaps because of network problems or machine crashes), you want clients to be able to continue working against the service. When the service is connected again, it can retrieve the pending calls from a queue. Even when both the client and the service are alive and running, network connectivity may be unavailable, and yet both the client and the service may want to continue with their work. Using queues at both ends will facilitate that.

Disjoint work

Whenever it is possible to decompose a business workflow into several operations that are separated in time—that is, where each operation must take place, but not necessarily immediately or in a particular order—it is usually a good idea to use queuing, because it will improve availability and throughput. You can queue up the operations and have them execute independently of one another.

Compensating work

When your business transaction may take hours or days to complete, you typically split it into at least two transactions. The first queues up the work to be completed immediately by initiating an external sequence, and the second verifies the success of the first and compensates for its failure if necessary.

Load leveling

Most systems do not have a constant level of load, as shown in [Figure 9-1](#). If you design the system for the peak load, you will be wasting system resources through most of the load cycle, and if you design the system to handle the average load, you will not be able to handle the peak. But with queued calls, the service can simply queue up the excess load and process it at leisure. This enables you to design a system for a nominal average of the desired throughput, as opposed to the maximum load.

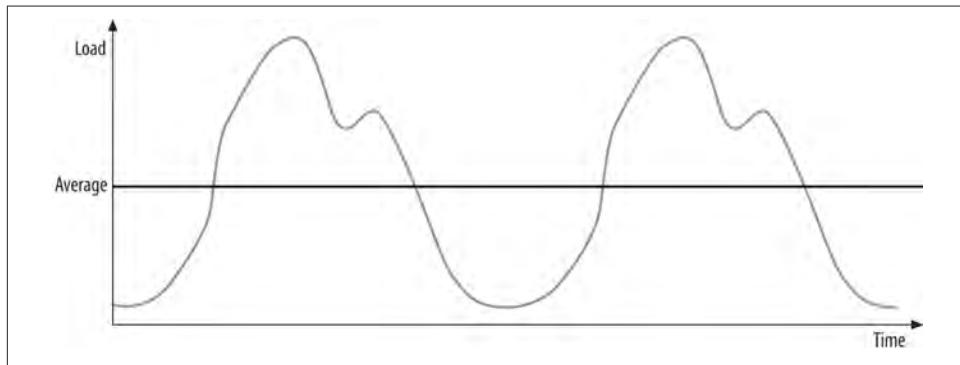


Figure 9-1. Fluctuating load

Queued Calls

WCF provides support for queued calls using the `NetMsmqBinding`. With this binding, instead of transporting the messages over TCP, HTTP, or IPC, WCF transports the messages over MSMQ. WCF packages each SOAP message into an MSMQ message and posts it to a designated queue. Note that there is no direct mapping of WCF messages to MSMQ messages, just like there is no direct mapping of WCF messages to TCP packets. A single MSMQ message can contain multiple WCF messages, or just a single one, according to the contract session mode (as discussed at length later). In

effect, instead of sending the WCF message to a live service, the client posts the message to an MSMQ queue. All that the client sees and interacts with is the queue, not a service endpoint. As a result, the calls are inherently asynchronous (because they will execute later, when the service processes the messages) and disconnected (because the service or client may interact with local queues).

Queued Calls Architecture

As with every WCF service, in the case of a queued service, the client interacts with a proxy, as shown in [Figure 9-2](#).

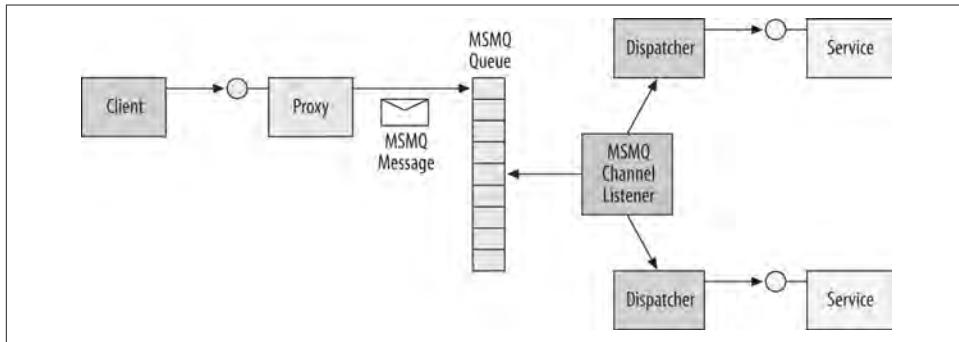


Figure 9-2. Queued calls architecture

However, since the proxy is configured to use the MSMQ binding, it does not send the WCF message to any particular service. Instead, it converts the call (or calls) to an MSMQ message (or messages) and posts it to the queue specified in the endpoint's address. On the service side, when a service host with a queued endpoint is launched, the host installs a queue listener, which is similar conceptually to the listener associated with a port when using TCP or HTTP. The queue's listener detects that there is a message in the queue, de-queues the message, and then creates the host side's chain of interceptors, ending with a dispatcher. The dispatcher calls the service instance as usual. If multiple messages are posted to the queue, the listener can create new instances as fast as the messages come off the queue, resulting in asynchronous, disconnected, and concurrent calls.

If the host is offline, messages destined for the service will simply remain pending in the queue. The next time the host is connected, the messages will be played to the service. Obviously, if both the client and the host are alive and running and are connected, the host will process the calls immediately.

Queued Contracts

A potentially disconnected call made against a queue cannot possibly return any values, because no service logic is invoked at the time the message is dispatched to the

queue. Not only that, but the call may be dispatched to the service and processed after the client application has shut down, when there is no client available to process the returned values. In much the same way, the call cannot return to the client any service-side exceptions, and there may not be a client around to catch and handle the exceptions anyway. In fact, WCF disallows using fault contracts on queued operations. Since the client cannot be blocked by invoking the operation—or rather, the client is blocked, but only for the briefest moment it takes to queue up the message—the queued calls are inherently asynchronous from the client's perspective. All of these are the classic characteristics of one-way calls. Consequently, any contract exposed by an endpoint that uses the `NetMsmqBinding` can have only one-way operations, and WCF verifies this at service (and proxy) load time:

```
//Only one-way calls allowed on queued contracts
[ServiceContract]
interface IMyContract
{
    [OperationContract(IsOneWay = true)]
    void MyMethod();
}
```

Because the interaction with MSMQ is encapsulated in the binding, there is nothing in the service or client invocation code pertaining to the fact that the calls are queued. The queued service and client code look like any other WCF service and client code, as shown in [Example 9-1](#).

Example 9-1. Implementing and consuming a queued service

```
//////////////////////////// Service Side ///////////////////////
[ServiceContract]
interface IMyContract
{
    [OperationContract(IsOneWay = true)]
    void MyMethod();
}
class MyService : IMyContract
{
    public void MyMethod()
    {...}
}
//////////////////////////// Client Side //////////////////////
MyContractClient proxy = new MyContractClient();
proxy.MyMethod();
proxy.Close();
```

Configuration and Setup

When you define an endpoint for a queued service, the endpoint address must contain the queue's name and designation (that is, the type of the queue). MSMQ defines

two types of queues: public and private. *Public queues* require an MSMQ domain controller installation or Active Directory integration and can be accessed across machine boundaries. Applications in production often require public queues due to the secure and disconnected nature of such queues. *Private queues* are local to the machine on which they reside and do not require a domain controller. Such a deployment of MSMQ is called a *workgroup installation*. During development, and for private queues they set up and administer, developers usually resort to a workgroup installation.

You designate the queue type (private or public) as part of the queued endpoint address:

```
<endpoint
    address = "net.msmq://localhost/private/MyServiceQueue"
    binding = "netMsmqBinding"
    ...
/>
```

In the case of a public queue, you can omit the `public` designator and have WCF infer the queue type. With private queues, you must include the designator. Also note that there is no \$ sign in the queue's type.

Workgroup installation and security

When you're using private queues in a workgroup installation, you typically disable MSMQ security on the client and service sides. [Chapter 10](#) discusses in detail how to secure WCF calls, including queued calls. Briefly, the default MSMQ security configuration expects users to present certificates for authentication, and MSMQ certificate-based security requires an MSMQ domain controller. Alternatively, selecting Windows security for transport security over MSMQ requires Active Directory integration, which is not possible with an MSMQ workgroup installation. For now, [Example 9-2](#) shows how to disable MSMQ security.

Example 9-2. Disabling MSMQ security

```
<system.serviceModel>
    ...
    <endpoint
        address = "net.msmq://localhost/private/MyServiceQueue"
        binding = "netMsmqBinding"
        bindingConfiguration = "NoMSMQSecurity"
        contract = "..." 
    />
    ...
<bindings>
    <netMsmqBinding>
        <binding name = "NoMSMQSecurity">
            <security mode = "None"/>
        </binding>
    </netMsmqBinding>
</bindings>
```

```
</binding>
</netMsmqBinding>
</bindings>
</system.serviceModel>
```



If you must for some reason enable security for development in a workgroup installation, you can configure the service to use message security with username credentials.

Creating the queue

On both the service and the client side, the queue must exist before client calls are queued up against it. There are several options for creating the queue. The administrator (or the developer, during development) can use the MSMQ control panel applet to create the queue, but that is a manual step that should be automated. The host process can use the API of `System.Messaging` to verify that the queue exists before opening the host. The class `MessageQueue` offers the `Exists()` method for verifying that a queue is created, and the `Create()` methods for creating a queue:

```
public class MessageQueue : ...
{
    public static MessageQueue Create(string path); //Nontransactional
    public static MessageQueue Create(string path, bool transactional);
    public static bool Exists(string path);
    public void Purge();
    //More members
}
```

If the queue is not present, the host process can first create it and then proceed to open the host. [Example 9-3](#) demonstrates this sequence.

Example 9-3. Verifying a queue on the host

```
ServiceHost host = new ServiceHost(typeof(MyService));

if(MessageQueue.Exists(@".\private$\MyServiceQueue") == false)
{
    MessageQueue.Create(@".\private$\MyServiceQueue", true);
}
host.Open();
```

In this example, the host verifies against the MSMQ installation on its own machine that the queue is present before opening the host. If it needs to, the hosting code creates the queue. Note the use of the `true` value for the transactional queue, as discussed later. Note also the use of the `$` sign in the queue designation.

The obvious problem with [Example 9-3](#) is that it hardcodes the queue name, not once, but twice. It is preferable to read the queue name from the application config file by storing it in an application setting, although there are problems even with that approach. First, you have to constantly synchronize the queue name in the application settings and in the endpoint's address. Second, you still have to repeat this code every time you host a queued service. Fortunately, it is possible to encapsulate and automate the code in [Example 9-3](#) in my `ServiceHost<T>`, as shown in [Example 9-4](#).

Example 9-4. Creating the queue in ServiceHost<T>

```
public class ServiceHost<T> : ServiceHost
{
    protected override void OnOpening()
    {
        foreach(ServiceEndpoint endpoint in Description.Endpoints)
        {
            endpoint.VerifyQueue();
        }
        base.OnOpening();
    }
    //More members
}
public static class QueuedServiceHelper
{
    public static void VerifyQueue(this ServiceEndpoint endpoint)
    {
        if(endpoint.Binding is NetMsmqBinding)
        {
            string queue = GetQueueFromUri(endpoint.Address.Uri);
            if(MessageQueue.Exists(queue) == false)
            {
                MessageQueue.Create(queue,true);
            }
        }
    }
    //Parses the queue name out of the address
    static string GetQueueFromUri(Uri uri)
    {...}
}
```

In [Example 9-4](#), `ServiceHost<T>` overrides the `OnOpening()` method of its base class. This method is called before opening the host, but after calling the `Open()` method. `ServiceHost<T>` iterates over the collection of configured endpoints. For each endpoint, if the binding used is `NetMsmqBinding`—that is, if queued calls are expected—`ServiceHost<T>` calls the extension method `VerifyQueue()` of the `ServiceEndpoint` type and asks it to verify the presence of the queue. The static extension `VerifyQueue()` method of `QueuedServiceHelper` parses the queue's name out of the end-

point's address and uses code similar to that in [Example 9-3](#) to create the queue if needed.

Using `ServiceHost<T>`, [Example 9-3](#) is reduced to:

```
ServiceHost<MyService> host = new ServiceHost<MyService>();
host.Open();
```

The client must also verify that the queue exists before dispatching calls to it. [Example 9-5](#) shows the required steps on the client side.

Example 9-5. Verifying the queue by the client

```
if(MessageQueue.Exists(@".\private$\MyServiceQueue") == false)
{
    MessageQueue.Create(@".\private$\MyServiceQueue",true);
}
MyContractClient proxy = new MyContractClient();
proxy.MyMethod();
proxy.Close();
```

Again, you should not hardcode the queue name and should instead read the queue name from the application config file by storing it in an application setting. And again, you will face the challenges of keeping the queue name synchronized in the application settings and in the endpoint's address, and of writing queue verification logic everywhere your clients use the queued service. You can use `QueuedServiceHelper` directly on the endpoint behind the proxy, but that forces you to create the proxy (or a `ServiceEndpoint` instance) just to verify the queue. You can, however, extend my `QueuedServiceHelper` to streamline and support client-side queue verification, as shown in [Example 9-6](#).

Example 9-6. Extending QueuedServiceHelper to verify the queue on the client side

```
public static class QueuedServiceHelper
{
    public static void VerifyQueues()
    {
        Configuration config = ConfigurationManager.OpenExeConfiguration(
            ConfigurationUserLevel.None);
        ServiceModelSectionGroup sectionGroup =
            ServiceModelSectionGroup.GetSectionGroup(config);

        foreach(ChannelEndpointElement endpointElement in
            sectionGroup.Client.Endpoints)
        {
            if(endpointElement.Binding == "netMsmqBinding")
            {
                string queue = GetQueueFromUri(endpointElement.Address);
```

```

        if(MessageQueue.Exists(queue) == false)
        {
            MessageQueue.Create(queue,true);
        }
    }
}
//More members
}

```

Example 9-6 uses the type-safe programming model that is offered by the `ConfigurationManager` class to parse a configuration file. It loads the WCF section (the `ServiceModelSectionGroup`) and iterates over all the endpoints defined in the client config file. For each endpoint that is configured with the MSMQ binding, `VerifyQueues()` creates the queue if required.

Using `QueuedServiceHelper`, **Example 9-5** is reduced to:

```

QueuedServiceHelper.VerifyQueues();

MyContractClient proxy = new MyContractClient();
proxy.MyMethod();
proxy.Close();

```

Note that the client application needs to call `QueuedServiceHelper.VerifyQueues()` just once anywhere in the application, before issuing the queued calls.

If the client is not using a config file to create the proxy (or is using a channel factory), the client can still use the extrusion method `VerifyQueue()` of the `ServiceEndpoint` class:

```

EndpointAddress address = new EndpointAddress(...);
Binding binding = new NetMsmqBinding(...); //Can still read binding from config

MyContractClient proxy = new MyContractClient(binding,address);
proxy.Endpoint.VerifyQueue();

proxy.MyMethod();

proxy.Close();

```

Queue purging

When a host is launched, it may already have messages in queues, received by MSMQ while the host was offline, and the host will immediately start processing these messages. Dealing with this very scenario is one of the core features of queued services, as it enables you to have disconnected services. While this is, therefore, exactly the sort of behavior you would like when deploying a queued service, it is typically a hindrance in debugging. Imagine a debug session of a queued service. The client issues a few calls and the service begins processing the first call, but while stepping through

the code you notice a defect. You stop debugging, change the service code, and relaunch the host, only to have it process the remaining messages in the queue from the previous debug session, even if those messages break the new service code. Usually, messages from one debug session should not seed the next one. The solution is to programmatically purge the queues when the host shuts down, in debug mode only. You can streamline this with my `ServiceHost<T>`, as shown in [Example 9-7](#).

Example 9-7. Purging the queues on host shutdown during debugging

```
public static class QueuedServiceHelper
{
    public static void PurgeQueue(ServiceEndpoint endpoint)
    {
        if(endpoint.Binding is NetMsmqBinding)
        {
            string queueName = GetQueueFromUri(endpoint.Address.Uri);
            if(MessageQueue.Exists(queueName) == true)
            {
                MessageQueue queue = new MessageQueue(queueName);
                queue.Purge();
            }
        }
    }
    //More members
}
public class ServiceHost<T> : ServiceHost
{
    protected override void OnClosing()
    {
        PurgeQueues();
        //More cleanup if necessary
        base.OnClosing();
    }
    [Conditional("DEBUG")]
    void PurgeQueues()
    {
        foreach(ServiceEndpoint endpoint in Description.Endpoints)
        {
            QueuedServiceHelper.PurgeQueue(endpoint);
        }
    }
    //More members
}
```

In this example, the `QueuedServiceHelper` class offers the static method `PurgeQueue()`. As its name implies, `PurgeQueue()` accepts a service endpoint. If the binding used by that endpoint is the `NetMsmqBinding`, `PurgeQueue()` extracts the queue name out of the endpoint's address, creates a new `MessageQueue` object, and purges it. `ServiceHost<T>` overrides the `OnClosing()` method, which is called when the host

shuts down gracefully. It then calls the private `PurgeQueues()` method. `PurgeQueues()` is marked with the `Conditional` attribute, using `DEBUG` as a condition. This means that while the body of `PurgeQueues()` always compiles, its call sites are conditioned on the `DEBUG` symbol. In debug mode only, `OnClosing()` will actually call `PurgeQueues()`. `PurgeQueues()` iterates over all endpoints of the host, calling `QueuedServiceHelper.PurgeQueue()` on each.



The `Conditional` attribute is the preferred way in .NET for using conditional compilation and avoiding the pitfalls of explicit conditional compilation with `#if`.

Queues, services, and endpoints

WCF requires you to always dedicate a queue per endpoint for each service. This means a service with two contracts needs two queues for the two corresponding endpoints:

```
<service name = "MyService">
  <endpoint
    address = "net.msmq://localhost/private/MyServiceQueue1"
    binding = "netMsmqBinding"
    contract = "IMyContract"
  />
  <endpoint
    address = "net.msmq://localhost/private/MyServiceQueue2"
    binding = "netMsmqBinding"
    contract = "IMyOtherContract"
  />
</service>
```

The reason is that the client actually interacts with a queue, not a service endpoint. In fact, there may not even be a service at all; there may only be a queue. Two distinct endpoints cannot share queues because they will get each other's messages. Since the WCF messages in the MSMQ messages will not match, WCF will silently discard those messages it deems invalid, and you will lose the calls. Much the same way, two polymorphic endpoints on two services cannot share a queue, because they will eat each other's messages.

Metadata

WCF cannot exchange metadata over MSMQ. Consequently, it is customary for even a service that will always have only queued calls to also expose a MEX endpoint or to enable metadata exchange over HTTP-GET, because the service's clients still need a way to retrieve the service description and bind against it.

WAS hosting

When hosting a queued service in the WAS, the name of the queue must be the same as the name of the .svc file leading to the virtual directory. For example:

```
<endpoint  
    address = "net.msmq://localhost/private/WASService.svc"  
    binding = "netMsmqBinding"  
    contract = "IMyQueuedContract"  
/>
```

Transactions

MSMQ is a WCF transactional resource manager. When you create a queue (either programmatically or administratively), you can create the queue as a transactional queue. If the queue is transactional, it is durable, and messages always persist to disk. More importantly, posting messages to and removing messages from the queue will always be done under a transaction. If the code that tries to interact with the queue has an ambient transaction, the queue will silently join that transaction. If no ambient transaction is present, MSMQ will start a new transaction for that interaction. It is as if the queue is encased in a `TransactionScope` constructed with `TransactionScopeOption.Required`. Once in a transaction, the queue will commit or roll back along with the accessing transaction. For example, if the accessing transaction posts a message to the queue and then aborts, the queue will reject the message.

Delivery and Playback

When a nontransactional client calls a queued service, client-side failures after the call will not roll back posting the message to the queue, and the queued call will be dispatched to the service. However, a client calling a queued service may call under a transaction, as shown in [Figure 9-3](#).

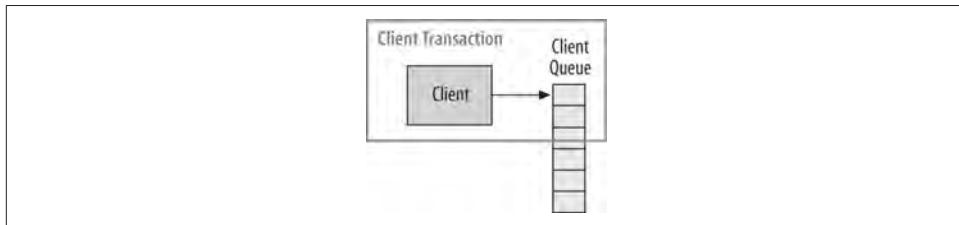


Figure 9-3. Posting to a client-side queue

The client calls are converted to WCF messages and then packaged in an MSMQ message (or messages). If the client's transaction commits, these MSMQ messages are posted to the queue and persist there. If the client's transaction aborts, the queue discards these MSMQ messages. In effect, WCF provides clients of a queued service with

an auto-cancellation mechanism for their asynchronous, potentially disconnected calls. Normal connected asynchronous calls cannot be combined easily, if at all, with transactions, because once the call is dispatched there is no way to recall it in case the original transaction aborts. Unlike connected asynchronous calls, queued service calls are designed for this very transactional scenario. In addition, the client may interact with multiple queued services in the same transaction. Aborting the client transaction for whatever reason will automatically cancel all calls to those queued services.

The delivery transaction

Since the client may not be on the same machine as the service, and since the client, the service, or both could be disconnected, MSMQ maintains a client-side queue as well. The client-side queue serves as a “proxy” to the service-side queue. In the case of a remote queued call, the client first posts the message to the client-side queue. When (or if) the client is connected, MSMQ will deliver the queued messages from the client-side queue to the service-side queue, as shown in [Figure 9-4](#).

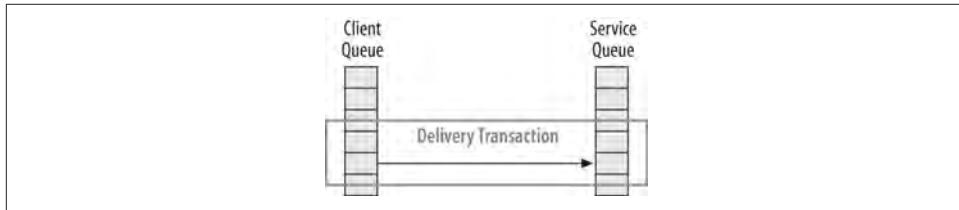


Figure 9-4. The delivery transaction

Since MSMQ is a resource manager, removing the message from the client-side queue will create a transaction (if indeed the queue is transactional). If MSMQ fails to deliver the message to the service-side queue for whatever reason (such as a network fault or service machine crash), the delivery transaction will abort, the message removal from the client-side queue will be rolled back, and the message posting to the service-side queue will also be canceled, resulting in the message being back in the client-side queue. At this point, MSMQ will try again to deliver the message. Thus, while you can configure and control failure handling (as you will see later), excluding fatal errors that can never be resolved, queued services actually enjoy a guaranteed delivery mechanism; if it is technically possible to deliver the message (within the confines of the failure-handling modes), the message will get from the client to the service. In effect, this is WCF's way of providing reliable messaging for queued services. Of course, there is no direct support for the reliable messaging protocol, as there is with connected calls; this is just the analogous mechanism.

The playback transaction

When WCF removes a message from the queue for playback to the service, this kick-starts a new transaction (assuming the queue is transactional), as shown in Figure 9-5.

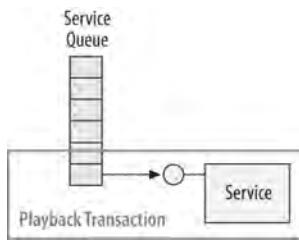


Figure 9-5. Playback transaction

The service is usually configured to participate in the playback transaction. If the playback transaction aborts (usually due to service-side exceptions), the message rolls back to the queue, where WCF detects it and dispatches it again to the service. This, in effect, yields an auto-retry mechanism. Consequently, you should keep the service's processing of the queued call relatively short, or risk aborting the playback transaction. An important observation here is that it is wrong to equate queued calls with lengthy asynchronous calls.

Service Transaction Configuration

As just demonstrated, assuming transactional queues, there are actually three transactions involved in every queued call: client, delivery, and playback, as shown in Figure 9-6.

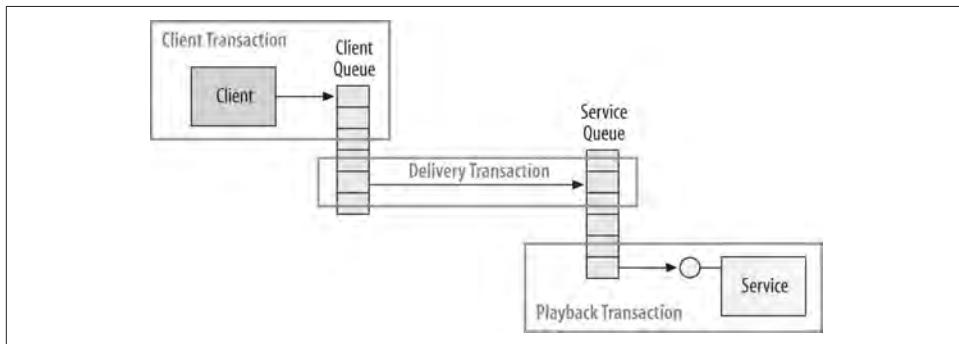


Figure 9-6. Queued calls and transactions

From a design perspective, you rarely, if ever, depict the delivery transaction in your design diagrams and you simply take it for granted. In addition, the service will never

participate in the client's transaction, so in effect my four logical transactional modes from [Chapter 7](#) (Client, Client/Service, Service, None) do not apply with queued services. Configuring the service contract operation with `TransactionFlowOption.Allowed` or `TransactionFlowOption.NotAllowed` leads to the same result—the client transaction is never provided to the service. Not only that, but `TransactionFlowOption.Mandatory` is disallowed for configuration on a queued contract, and this constraint is verified at the service load time. The real issue is the relation between the playback transaction and the service transactional configuration.

Participating in the playback transaction

From a WCF perspective, the playback transaction is treated as the incoming transaction to the service. To participate in the playback transaction, the service needs to have the operation behavior configured with `TransactionScopeRequired` set to `true`, as shown in [Example 9-8](#) and graphically in [Figure 9-5](#).

Example 9-8. Participating in the playback transaction

```
[ServiceContract]
interface IMyContract
{
    [OperationContract(IsOneWay = true)]
    void MyMethod();
}
class MyService : IMyContract
{
    [OperationBehavior(TransactionScopeRequired = true)]
    public void MyMethod()
    {
        Transaction transaction = Transaction.Current;
        Debug.Assert(transaction.TransactionInformation.
            DistributedIdentifier != Guid.Empty);
    }
}
```

An interesting point made in [Example 9-8](#) is that with MSMQ, every transaction always uses the DTC for transaction management, even in the case of a single service and a single playback.

Ignoring the playback transaction

If the service is configured for not having any transactions (like the service shown in [Example 9-9](#)), WCF will still use a transaction to read the message from the queue, except that transaction will always commit (barring an unforeseen failure in MSMQ itself). Exceptions and failures at the service itself will not abort the playback transaction.

Example 9-9. Ignoring the playback transaction

```
[ServiceContract]
interface IMyContract
{
    [OperationContract(IsOneWay = true)]
    void MyMethod();
}
class MyService : IMyContract
{
    public void MyMethod()
    {
        Transaction transaction = Transaction.Current;
        Debug.Assert(transaction == null);
    }
}
```

This scenario is depicted graphically in [Figure 9-7](#).

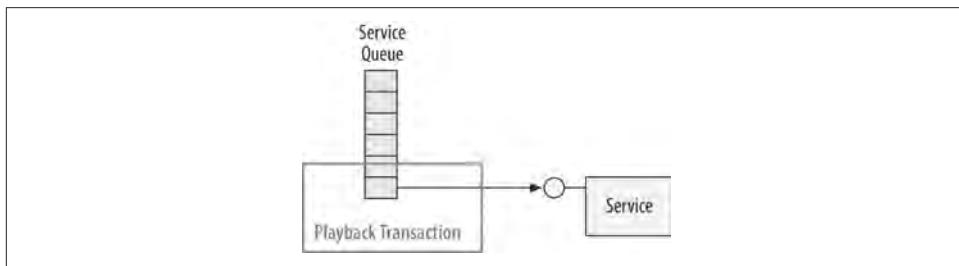


Figure 9-7. Ignoring the playback transaction

Services that do not participate in the playback transaction will not have the benefit of automated retries by WCF in the case of a playback failure, and it is possible for the played-back call to fail while the de-queued transaction commits. The main motivation for configuring queued services this way is to accommodate lengthy processing. If the service does not participate in the playback transaction, the call can take any amount of time to complete.

Using a separate transaction

You can also write a service so that it manually requires a new transaction, as shown in [Example 9-10](#).

Example 9-10. Using a new transaction

```
class MyService : IMyContract
{
    public void MyMethod()
    {
        using(TransactionScope scope = new TransactionScope())
```

```

{
    ...
    scope.Complete();
}
}
}

```

This scenario is depicted in [Figure 9-8](#).

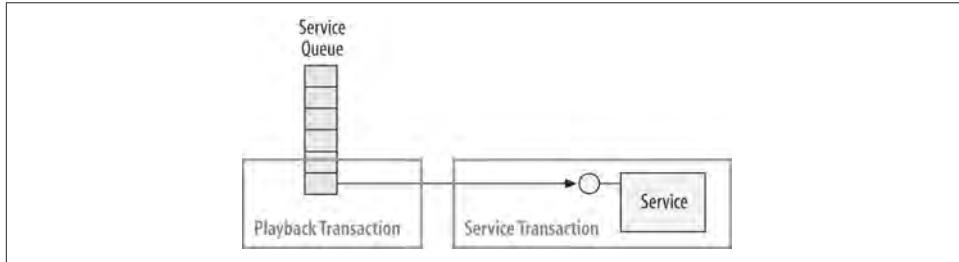


Figure 9-8. Using a new transaction

When the service uses its own new transaction for each message, it should also prevent participating in the playback transaction (by defaulting to the `TransactionScopeRequired` value of `false`) so as not to affect the playback transaction in any way. Again, this negates the benefit of the auto-retry mechanism. However, having a new transaction separate from the playback transaction gives the service the opportunity to perform its own transactional work. You would typically configure a service to use its own transaction when the queued operation being called is nice to have and should be performed under the protection of a transaction, yet does not need to be retried in case of a failure.

Nontransactional Queues

The MSMQ queues described so far were both durable and transactional. The messages persisted to the disk, and posting a message to and reading it from the queue was transactional. However, MSMQ also supports nontransactional queues. Such queues can be durable and persist on the disk or can be volatile (stored in memory). If the queue is volatile, the messages in the queue will not persist across a machine shutdown or a machine crash or just recycling of the MSMQ service.

When you create a queue (either using the MSMQ administration tool or programmatically), you can configure it to be transactional or not, and that selection is fixed for the life of the queue. Nontransactional queues do not offer any of the benefits of transactional messaging systems, such as auto-cancellation, guaranteed delivery, and auto-retries. When using a nontransactional queue, if the client transaction aborts, the message or messages will stay in the queue and be delivered to the service. If the playback transaction aborts, the messages will be lost.

As inadvisable as it is, WCF can work with nontransactional queues. `MsmqBindingBase` (the base class of `NetMsmqBinding`) offers the two Boolean properties `Durable` and `ExactlyOnce`, and these properties default to `true`:

```
public abstract class MsmqBindingBase : Binding, ...
{
    public bool Durable
    {get;set;}
    public bool ExactlyOnce
    {get;set;}
    //More members
}
public class NetMsmqBinding : MsmqBindingBase
{...}
```

To work with a nontransactional queue, the `ExactlyOnce` property must be set to `false`. This will enable you to work both with volatile and durable queues. However, because of the lack of guaranteed delivery, when using a volatile queue WCF requires that you set the `ExactlyOnce` property of the binding to `false`; otherwise, WCF will throw an `InvalidOperationException` at the service load time. Consequently, here is a consistent configuration for a volatile nontransactional queue:

```
<netMsmqBinding>
    <binding name = "VolatileQueue"
        durable      = "false"
        exactlyOnce = "false"
    />
</netMsmqBinding>
```

Instance Management

The contract session mode and the service instance mode have a paramount effect on the behavior of the queued calls, the way the calls are played back to the service, and the overall program workflow and allowed assumptions. The MSMQ binding cannot maintain a transport session in the connected sense, since the client is inherently disconnected. Instead, the equivalent MSMQ concept is called a *sessionogram*. If the contract is configured with `SessionMode.Allowed` (the default) or `SessionMode.NotAllowed`, there will be no sessionogram. Every call the client makes on the proxy will be converted to a single WCF message, and those WCF messages will be placed in individual MSMQ messages and posted to the queue. A client making two calls on the proxy will result in two MSMQ messages. If the contract is configured with `SessionMode.Required`, all the calls made by the client against the same proxy will be packaged in a single MSMQ message, in the order in which they were made and posted to the queue. On the service side, WCF will play the calls from the MSMQ message in the order they were made (like a recording) to the same service instance. This mode is therefore analogous to a transport session and a sessionful service.

Per-Call Queued Services

In the case of a per-call service, the client has no way of knowing whether its calls will eventually end up being played to a queued per-call service. All the client sees is the session mode of the contract. If the session mode is either `SessionMode.Allowed` or `SessionMode.NotAllowed`, there will be no sessiongram. In this case, regardless of whether the service is configured as per-call or sessionful it will amount to the same result: per-call processing and instantiation.

Nontransactional clients

When a client without an ambient transaction calls a sessiongram-less queued endpoint (as in [Example 9-11](#)), the MSMQ messages generated for each call are posted to the queue immediately after each call. If the client has an exception, the messages posted up to that point are not rejected and are delivered to the service.

Example 9-11. Nontransactional client of a sessionless queued endpoint

```
[ServiceContract]
interface IMyContract
{
    [OperationContract(IsOneWay = true)]
    void MyMethod();
}

//Client code
using(TransactionScope scope =
    new TransactionScope(TransactionScopeOption.Suppress))
{
    MyContractClient proxy = new MyContractClient();

    proxy.MyMethod(); //Message posts to queue here
    proxy.MyMethod(); //Message posts to queue here

    proxy.Close();
}
```

Transactional clients

With a transactional client (that is, client code with an ambient transaction) of a sessiongram-less queued endpoint (as in [Example 9-12](#)), the messages corresponding to each call are posted to the queue only when the client's transaction commits. If the client transaction aborts, all of those messages are rejected from the queue and all calls are canceled.

Example 9-12. Transactional client of a sessionless queued endpoint

```
[ServiceContract]
interface IMyContract
```

```

{
    [OperationContract(IsOneWay = true)]
    void MyMethod();
}
//Client code
using(TransactionScope scope = new TransactionScope())
{
    MyContractClient proxy = new MyContractClient();

    proxy.MyMethod(); //Message written to queue
    proxy.MyMethod(); //Message written to queue

    proxy.Close();
    scope.Complete();
} //Messages committed to queue here

```

There is no relationship between the proxy and the ambient transaction. If the client uses a transaction scope (as in [Example 9-12](#)), the client can close the proxy inside or outside the scope and may continue to use the proxy even after the transaction ends, or in a new transaction. The client may also close the proxy before or after the call to `Complete()`.

Per-call processing

On the host side, the queued calls are dispatched separately to the service, and each call is played to a separate service instance. This is the case even if the service instance mode is per-session. I therefore recommend that when using a sessiongram-less queued contract, you should always explicitly configure the service as per-call and configure the contract for disallowing sessions, to increase the readability of the code and clearly convey your design decision:

```

[ServiceContract(SessionMode = SessionMode.NotAllowed)]
interface IMyContract
{...}

[ServiceBehavior(InstanceContextMode = InstanceContextMode.PerCall)]
class MyService : IMyContract
{...}

```

After each call, the service instance is disposed of, just as with a connected per-call service. The per-call service may or may not be transactional. If it is transactional and the playback transaction is aborted, only that particular call is rolled back to the queue for a retry. As you will see later, due to concurrent playback and WCF's failure-handling behavior, calls to a per-call queued service can execute and complete in any order, and the client cannot make any assumptions about call ordering. Note that even calls dispatched by a transactional client may fail or succeed independently. Never assume order of calls with a sessiongram-less queued service.

Sessionful Queued Services

For sessionful queued services, the service contract must be configured with `SessionMode.Required`:

```
[ServiceContract(SessionMode = SessionMode.Required)]
interface IMyContract
{...}

class MyService : IMyContract
{...}
```

As mentioned previously, when the client queues up calls against a sessionful queued endpoint, all calls made throughout the session are grouped into a single MSMQ message. Once that single message is dispatched and played to the service, WCF creates a new dedicated service instance to handle all the calls in the message. All calls in the message are played back to that instance in their original order. After the last call, the instance is disposed of automatically.

WCF will provide both the client and the service with a unique session ID. However, the client session ID will be uncorrelated to that of the service. To approximate the session semantic, all calls on the same instance on the host side will share the same session ID.

Clients and transactions

In the case of a sessionful queued endpoint, the client must have an ambient transaction in order to call the proxy. Nontransactional clients are disallowed and will result in an `InvalidOperationException`:

```
[ServiceContract(SessionMode = SessionMode.Required)]
interface IMyContract
{
    [OperationContract(IsOneWay = true)]
    void MyMethod();
}

using(TransactionScope scope = new TransactionScope
        (TransactionScopeOption.Suppress))
{
    MyContractClient proxy = new MyContractClient();

    proxy.MyMethod(); //Throws InvalidOperationException
    proxy.MyMethod();

    proxy.Close();
}
```

For a transactional client, WCF posts a single message to the queue when the transaction commits, and that single message is rejected from the queue if the transaction aborts:

```
using(TransactionScope scope = new TransactionScope())
{
    MyContractClient proxy = new MyContractClient();

    proxy.MyMethod();
    proxy.MyMethod();

    proxy.Close(); //Finish composing message, writes to queue

    scope.Complete();
} //Single message committed to queue here
```

It is important to note that the single message prepared by the proxy must be posted to the queue within the same client transaction—that is, the client must end the session inside the transaction. If the client does not close the proxy before the transaction is complete, the transaction will always abort:

```
MyContractClient proxy = new MyContractClient();
using(TransactionScope scope = new TransactionScope())
{
    proxy.MyMethod();
    proxy.MyMethod();

    scope.Complete();
}//Transaction aborts
proxy.Close();
```

This is required to enforce the atomicity of the sessiongram. All the calls in the session should either be posted to or rejected from the queue. If the client were to use the proxy in a second transaction that could commit or abort independently of the first, the results could be ambiguous or even dangerous.

An interesting side effect of this edict is that there is no point in storing a proxy to a queued sessionful endpoint in a member variable, because that proxy can only be used once in a single transaction and cannot be reused across client transactions.

Not only does the client have to close the proxy before the transaction ends, but when using a transaction scope, the client must close the proxy before completing the transaction. The reason is that closing the proxy to a queue's sessionful endpoint requires accessing the current ambient transaction, which is not possible after calling `Complete()`. Trying to do so results in an `InvalidOperationException`:

```
MyContractClient proxy = new MyContractClient();
using(TransactionScope scope = new TransactionScope())
{
    proxy.MyMethod();
    proxy.MyMethod();
```

```
    scope.Complete();
    proxy.Close(); //Transaction aborts
}
```

A corollary of this requirement is that you cannot stack using statements in any order, because doing so may result in calling `Dispose()` in the wrong order (first on the scope, and then on the proxy):

```
using(MyContractClient proxy = new MyContractClient())
using(TransactionScope scope = new TransactionScope())
{
    proxy.MyMethod();
    proxy.MyMethod();

    scope.Complete();

} //Transaction aborts
```

Services and transactions

A sessionful queued service must be configured to use transactions in all operations by setting `TransactionScopeRequired` to true. Failing to do so will abort all playback transactions. The service is required to have a transaction in every operation so that all the calls in the session fail or succeed as one atomic operation (i.e., so that a failure in one of the operations causes the entire queued session to fail). In addition, the transaction must be the same transaction for all operations in the session. Partial success is impossible here, because WCF cannot return only a portion of the MSMQ message back to the queue after a failure of one of the operations but not the others. The service must equate the session boundary with the transaction boundary. Do this by setting `TransactionAutoComplete` to false on all operations and relying on `TransactionAutoCompleteOnSessionClose` to true. This will also have the added benefit of creating the affinity to the same transaction in all operations.



Only a sessionful service can support a sessiongram contract, since only a service configured with `InstanceContextMode.PerSession` can set `TransactionAutoComplete` to false.

To further enforce this constraint, the service cannot rely on setting `ReleaseServiceInstanceOnTransactionComplete` to false in order to restore the instance semantics while completing in each operation. Trying to do so will cause all queued sessions to always abort.

Example 9-13 is a template for implementing a queued sessionful service.

Example 9-13. Implementing a sessionful queued service

```
[ServiceContract(SessionMode = SessionMode.Required)]
interface IMyContract
{
    [OperationContract(IsOneWay = true)]
    void MyMethod1();

    [OperationContract(IsOneWay = true)]
    void MyMethod2();

    [OperationContract(IsOneWay = true)]
    void MyMethod3();
}

[ServiceBehavior(TransactionAutoCompleteOnSessionClose = true)]
class MyService : IMyContract
{
    [OperationBehavior(TransactionScopeRequired = true,
                      TransactionAutoComplete = false)]
    public void MyMethod1()
    {...}

    [OperationBehavior(TransactionScopeRequired = true,
                      TransactionAutoComplete = false)]
    public void MyMethod2()
    {...}

    [OperationBehavior(TransactionScopeRequired = true,
                      TransactionAutoComplete = false)]
    public void MyMethod3()
    {...}
}
```

Singleton Service

A queued singleton service can never have a session and can only implement sessionless contracts. Configuring the `SessionMode` as either `SessionMode.Allowed` or `SessionMode.NotAllowed` has the same result: a sessionless interaction. Consequently, I recommend always explicitly configuring the contracts of a queued singleton service as sessionless:

```
[ServiceContract(SessionMode = SessionMode.NotAllowed)]
interface IMyContract
{...}
[ServiceBehavior(InstanceContextMode=InstanceContextMode.Single)]
class MyService : IMyContract
{...}
```

A nontransactional queued singleton service behaves like a regular WCF singleton as far as instancing. Regardless of the way the clients use their proxies, individual calls

on the proxies are packaged into separate MSMQ messages and dispatched separately to the singleton, as with a per-call service. However, unlike with a per-call service, all these calls will be played back to the same single instance.

A transactional queued singleton, on the other hand, behaves by default like a per-call service, because after every call that completes the transaction WCF will release the singleton instance. The only difference between a true per-call service and a singleton is that WCF will allow at most a single instance of the singleton, regardless of the number of queued messages. While you could apply the techniques described in [Chapter 7](#) to create a state-aware transactional singleton, you can also restore the singleton semantic by setting the `ReleaseServiceInstanceOnTransactionComplete` property to `false` and use volatile resource managers.

[Example 9-14](#) shows a template for implementing a transactional queued singleton.

Example 9-14. Transactional queued singleton

```
[ServiceContract(SessionMode = SessionMode.NotAllowed)]
interface IMyContract
{
    [OperationContract(IsOneWay = true)]
    void MyMethod();
}

[ServiceBehavior(InstanceContextMode=InstanceContextMode.Single,
                ReleaseServiceInstanceOnTransactionComplete = false)]
class MySingleton : IMyContract, IDisposable
{
    [OperationBehavior(TransactionScopeRequired = true)]
    public void MyMethod()
    {...}
    //More members
}
```

Calls and order

Because the calls are packaged into individual MSMQ messages, they may be played to the singleton in any order (due to retries and transactions). In addition, calls may complete in any order, and even calls dispatched by a transactional client may fail or succeed independently. Never assume order of calls with a singleton.

Concurrency Management

As with a connected service, the `ConcurrencyMode` property governs concurrent playback of queued messages. With a per-call service, all queued messages are played at once to different instances as fast as they come off the queue, up to the limit of the configured throttle. There is no need to configure for reentrancy to support callbacks,

because the operation contexts can never have callback references. There is also no need to configure for multiple concurrent access, because no two messages will ever share an instance. In short, with a queued per-call service, the concurrency mode is ignored.

When it comes to a sessionful queued service, you are required to configure the service with `ConcurrencyMode.Single`. The reason is that it is the only concurrency mode that allows you to turn off auto-completion, which is essential to maintain the session semantic. The calls in the message are always played to the same service instance, one at a time.

A queued singleton is really the only instancing mode that has any leeway with its concurrency mode. If the singleton is configured with `ConcurrencyMode.Single`, WCF will retrieve the messages all at once from the queue (up to the thread pool and throttling limits) and then queue up the calls in the internal queue the context lock maintains. Calls will be dispatched to the singleton one at a time. If the singleton is configured with `ConcurrencyMode.Multiple`, WCF will retrieve the messages all at once from the queue (up to the thread pool and throttling limits) and play them concurrently to the singleton. Obviously, in that case the singleton must provide for synchronized access to its state. If the singleton is also transactional, it is prone to transactional deadlocks over prolonged isolation maintained throughout each transaction.

Throttling

Queued calls have a nasty side effect of excelling at turning a low level of load into a high level of stress. Imagine an offline queued service that sustained relatively low load, such as a call per minute for one day. Once the host is launched, WCF flushes the queued calls (all 1,440 of them) to the service all at once, subjecting it to high stress. The fact that there are over 1,000 messages in the queue does not mean that your design supports 1,000 concurrent instances and calls.

Throttling a queued service is your way of controlling the stress on the service and avoiding turning load into stress. The important value to throttle is the number of concurrent playbacks. This is an effective way of throttling the number of played messages, because if the maximum number of concurrent calls is exceeded (overall stress), the excess messages will stay in the queue. With a per-call service, the throttle controls the overall number of allowed concurrent instances (and their implied resource consumption). With a per-session service, the throttle controls the number of allowed sessions. In the case of a queued singleton, you can combine a throttle value with `ConcurrencyMode.Multiple` to control just how many concurrent players are allowed (stress) and how many messages to keep in the queue (buffered load).

Delivery Failures

As discussed in [Chapter 6](#), a connected call may fail due to either communication failures or service-side errors. Similarly, a queued call can fail due to delivery failures or service-side playback errors. WCF provides dedicated error-handling mechanisms for both types of errors, and understanding them and integrating your error-handling logic with them is an intrinsic part of using queued services.

While MSMQ can guarantee delivery of a message if it is technically possible to do so, there are multiple examples of when it is not possible to deliver the message. These include but are not limited to:

Timeouts and expiration

As you will see shortly, each message has a timestamp, and the message has to be delivered and processed within the configured timeout. Failure to do so will cause the delivery to fail.

Security mismatches

If the security credentials in the message (or the chosen authentication mechanism itself) do not match up with what the service expects, the service will reject the message.

Transactional mismatches

The client cannot use a local nontransactional queue while posting a message to a transactional service-side queue.

Network problems

If the underlying network fails or is simply unreliable, the message may never reach the service.

Machine crashes

The service machine may crash due to software or hardware failures and will not be able to accept the message to its queue.

Purges

Even if the message is delivered successfully, the administrator (or any application, programmatically) can purge messages out of the queue and avoid having the service process them.

Quota breaches

Each queue has a quota controlling the maximum amount of data it can hold. If the quota is exceeded, future messages are rejected.

After every delivery failure, the message goes back to the client's queue, where MSMQ will continuously retry to deliver it. While in some cases, such as intermittent network failures or quota issues, the retries may eventually succeed, there are many

cases where MSMQ will never succeed in delivering the message. In practical terms, a large enough number of retry attempts may be unacceptable and may create a dangerous amount of thrashing. Delivery-failure handling deals with how to let MSMQ know that it should not retry forever, how many attempts it should make before giving up, how much time can elapse before it gives up, and what it should do with the failed messages.

`MsmqBindingBase` offers a number of properties governing handling of delivery failures:

```
public abstract class MsmqBindingBase : Binding, ...
{
    public TimeSpan TimeToLive
    {get;set;}

    //DLQ settings
    public Uri CustomDeadLetterQueue
    {get;set;}
    public DeadLetterQueue DeadLetterQueue
    {get;set;}

    //More members
}
```

The Dead-Letter Queue

In messaging systems, after an evident failure to deliver a message, that message goes to a special queue called the *dead-letter queue* (DLQ). The DLQ is somewhat analogous to a classic dead-letter mailbox at the main post office. In the context of this discussion, failure to deliver constitutes not only failure to reach the service-side queue, but also failure to commit the playback transaction. MSMQ on the client and on the service side constantly acknowledge to each other receipt and processing of messages. If the service-side MSMQ successfully receives and retrieves the message from the service-side queue (that is, if the playback transaction committed), it sends a positive acknowledgment (ACK) to the client-side MSMQ. The service-side MSMQ can also send a negative acknowledgment (NACK) to the client. When the client-side MSMQ receives a NACK, it posts the message to the DLQ. If the client-side MSMQ receives neither an ACK nor a NACK, the message is considered in-doubt.

With MSMQ, you can configure a service-specific DLQ where only messages destined to that specific service go. Application-specific dead-letter queues grossly simplify both the administrator's and the developer's work.



When dealing with a nondurable queue, failed nontransactional messages go to a special system-wide DLQ.

Time to Live

With MSMQ, each message carries a timestamp initialized when the message is first posted to the client-side queue. In addition, every queued WCF message has a timeout, controlled by the `TimeToLive` property of `MsmqBindingBase`. After posting a message to the client-side queue, WCF mandates that the message must be delivered and processed within the configured timeout. Note that successful delivery to the service-side queue is not good enough—the call must be processed as well. The `TimeToLive` property is therefore somewhat analogous to the `SendTimeout` property of the connected bindings.

The `TimeToLive` property is relevant only to the posting client; it has no effect on the service side, nor can the service change it. `TimeToLive` defaults to one day. After continuously trying and failing to deliver (and process) a message for as long as `TimeToLive` allows, MSMQ stops trying and moves the message to the configured DLQ.

You can configure the time-to-live value either programmatically or administratively. For example, using a config file, here is how to configure a time to live of five minutes:

```
<bindings>
  <netMsmqBinding>
    <binding name = "ShortTimeout"
      timeToLive = "00:05:00"
    />
  </netMsmqBinding>
</bindings>
```

The main motivation for configuring a short timeout is when dealing with time-sensitive calls that must be processed in a timely manner. However, time-sensitive queued calls go against the grain of disconnected queued calls in general: the more time-sensitive the calls are, the more questionable the use of queued services is in the first place. The correct way of viewing time to live is as a last-resort heuristic used to eventually bring to the attention of the administrator the fact that the message was not delivered, not as a way to enforce business-level interpretation of the message's sensitivity.

Configuring the Dead-Letter Queue

`MsmqBindingBase` offers the `DeadLetterQueue` property, of the enum type `DeadLetterQueue`:

```
public enum DeadLetterQueue
{
    None,
    System,
    Custom
}
```

When `DeadLetterQueue` is set to `DeadLetterQueue.None`, WCF makes no use of a dead-letter queue. After a failure to deliver, WCF silently discards the message as if the call never happened. `DeadLetterQueue.System` is the default value of the property. As its name implies, it uses the system-wide DLQ: after a delivery failure, WCF moves the message from the client-side queue to the system-wide DLQ.



The system-wide DLQ is a transactional queue, so you must have the `ExactlyOnce` binding property set to its default value of `true` and the `Durable` property set to its default value of `true`.

When `DeadLetterQueue` is set to `DeadLetterQueue.Custom`, the application can take advantage of a dedicated DLQ. `DeadLetterQueue.Custom` requires that the application specify the custom DLQ address in the `CustomDeadLetterQueue` property of the binding. The default value of `CustomDeadLetterQueue` is `null`, but when `DeadLetterQueue.Custom` is employed, `CustomDeadLetterQueue` cannot be `null`:

```
<netMsmqBinding>
    <binding name = "CustomDLQ"
        deadLetterQueue = "Custom"
        customDeadLetterQueue = "net.msmq://localhost/private/MyCustomDLQ">
    </binding>
</netMsmqBinding>
```

Conversely, when the `DeadLetterQueue` property is set to any other value besides `DeadLetterQueue.Custom`, then `CustomDeadLetterQueue` must be `null`.

It is important to realize that the custom DLQ is just another MSMQ queue. It is up to the client-side developer to also deploy a DLQ service that processes its messages. All WCF does is automate the act of moving the message to the DLQ once a failure is detected.

Custom DLQ verification

If a custom DLQ is required, as with any other queue, the client should verify at run-time (before issuing queued calls) that the custom DLQ exists, and if necessary, create it. Following the pattern presented previously, you can automate and encapsulate this with the `ServiceEndpoint` extension method `VerifyQueue()` of `QueuedServiceHelper`, shown in [Example 9-15](#).

Example 9-15. Verifying a custom DLQ

```
public static class QueuedServiceHelper
{
    public static void VerifyQueue(this ServiceEndpoint endpoint)
    {
        if(endpoint.Binding is NetMsmqBinding)
        {
            string queue = GetQueueFromUri(endpoint.Address.Uri);
            if(MessageQueue.Exists(queue) == false)
            {
                MessageQueue.Create(queue,true);
            }
            NetMsmqBinding binding = endpoint.Binding as NetMsmqBinding;
            if(binding.DeadLetterQueue == DeadLetterQueue.Custom)
            {
                Debug.Assert(binding.CustomDeadLetterQueue != null);
                string DLQ = GetQueueFromUri(binding.CustomDeadLetterQueue);
                if(MessageQueue.Exists(DLQ) == false)
                {
                    MessageQueue.Create(DLQ,true);
                }
            }
        }
    }
    //More members
}
```



By default, queue permissions for new queues do not grant everyone send access. This includes the MSMQ service itself. To enable your queue to send to a custom DLQ, you must create the registry setting `HKEY_LOCAL_MACHINE\SOFTWARE\Microsoft\MSMQ\Parameters\Security\PermitAnonEveryoneSend` and set it to 1.

Processing the Dead-Letter Queue

The client needs to somehow process the accumulated messages in the DLQ. In the case of the system-wide DLQ, the client can provide a mega-service that supports all contracts of all queued endpoints on the system to enable it to process all failed messages. This is clearly an impractical idea, though, because that service could not possibly know about all queued contracts, let alone provide meaningful processing for all applications. The only feasible way to make this solution work would be to restrict the client side to at most a single queued service per system. Alternatively, you can write a custom application for direct administration and manipulation of the system DLQ using the types in `System.Messaging`. That application will parse and extract the relevant messages and process them. The problem with that approach (besides the inordinate amount of work involved) is that if the messages are protected and encrypted (as they should be), the application will have a hard time dealing with and

distinguishing between them. In practical terms, the only possible solution for a general client-side environment is a custom DLQ. When using a custom DLQ, you also provide a client-side service whose queue is the application's custom DLQ. That service will process the failed messages according to the application-specific requirements.

Defining the DLQ service

Implementing the DLQ service is done like any other queued service. The only requirement is that the DLQ service be polymorphic with the original service's contract. If multiple queued endpoints are involved, you will need a DLQ per contract per endpoint. [Example 9-16](#) shows a possible setup.

Example 9-16. DLQ service config file

```
<!-- Client side -->
<system.serviceModel>
    <client>
        <endpoint
            address = "net.msmq://localhost/private/MyServiceQueue"
            binding = "netMsmqBinding"
            bindingConfiguration = "MyCustomDLQ"
            contract = "IMyContract"
        />
    </client>
    <bindings>
        <netMsmqBinding>
            <binding name = "MyCustomDLQ"
                deadLetterQueue = "Custom"
                customDeadLetterQueue = "net.msmq://localhost/private/MyCustomDLQ">
            </binding>
        </netMsmqBinding>
    </bindings>
</system.serviceModel>

<!-- DLQ service side -->
<system.serviceModel>
    <services>
        <service name = "MyDLQService">
            <endpoint
                address = "net.msmq://localhost/private/MyCustomDLQ"
                binding = "netMsmqBinding"
                contract = "IMyContract"
            />
        </service>
    </services>
</system.serviceModel>
```

The client config file defines a queued endpoint with the `IMyContract` contract. The client uses a custom binding section to define the address of the custom DLQ. A separate queued service (potentially on a separate machine) also supports the `IMyContract` contract. The DLQ service uses as its address the DLQ defined by the client.

Failure properties

The DLQ service typically needs to know why the queued call delivery failed. WCF therefore offers the `MsmqMessageProperty` class, used to find out the cause of the failure and the current status of the message. `MsmqMessageProperty` is defined in the `System.ServiceModel.Channels` namespace:

```
public sealed class MsmqMessageProperty
{
    public const string Name = "MsmqMessageProperty";

    public int AbortCount
    {get;}
    public DeliveryFailure? DeliveryFailure
    {get;}
    public DeliveryStatus? DeliveryStatus
    {get;}
    public int MoveCount
    {get;}
    //More members
}
```

The DLQ service needs to obtain the `MsmqMessageProperty` from the operation context's incoming message properties:

```
public sealed class OperationContext : ...
{
    public MessageProperties IncomingMessageProperties
    {get;}
    //More members
}
public sealed class MessageProperties : IDictionary<string,object>,... {
    public object this[string name]
    {get;set;}
    //More members
}
```

When a message is passed to the DLQ, WCF will add to its properties an instance of `MsmqMessageProperty` detailing the failure. `MessageProperties` is merely a collection of message properties that you can access using a string as a key. To obtain the `MsmqMessageProperty`, use the constant `MsmqMessageProperty.Name`, as shown in [Example 9-17](#).

Example 9-17. Obtaining the MsmqMessageProperty

```
[ServiceContract(SessionMode = SessionMode.NotAllowed)]
interface IMyContract
{
    [OperationContract(IsOneWay = true)]
    void MyMethod(string someValue);
}
[ServiceBehavior(InstanceContextMode = InstanceContextMode.PerCall)]
class MyDLQService : IMyContract
{
    [OperationBehavior(TransactionScopeRequired = true)]
    public void MyMethod()
    {
        MsmqMessageProperty msmqProperty = OperationContext.Current.
            IncomingMessageProperties[MsmqMessageProperty.Name] as MsmqMessageProperty;

        Debug.Assert(msmqProperty != null);
        //Process msmqProperty
    }
}
```

Note in [Example 9-17](#) the use of the practices discussed so far for configuring the session mode, instance management, and transactions—the DLQ service is, after all, just another queued service.

The properties of `MsmqMessageProperty` detail the reasons for failure and offer some contextual information. `MoveCount` is the number of attempts made to play the message to the service, and `AbortCount` is the number of attempts made to read the message from the queue. `AbortCount` is less relevant to recovery attempts, because it falls under the responsibility of MSMQ and usually is of no concern. `DeliveryStatus` is a nullable enum of the type `DeliveryStatus`, defined as:

```
public enum DeliveryStatus
{
    InDoubt,
    NotDelivered
}
```

When a regular WCF queued service processes a delivered call, `DeliveryStatus` is set to `null`. With a DLQ service, `DeliveryStatus` will be set to `DeliveryStatus.InDoubt` unless the message was positively not delivered (i.e., a NACK was received). For example, expired messages are considered in-doubt because their time to live elapsed before the service could acknowledge them one way or the other.

The `DeliveryFailure` property is a nullable enum of the type `DeliveryFailure`, defined as follows (without the specific numerical values):

```
public enum DeliveryFailure
{
```

```

        AccessDenied,
        NotTransactionalMessage,
        Purged,
        QueueExceedMaximumSize,
        ReachQueueTimeout,
        ReceiveTimeout,
        Unknown
        //More members
    }
}

```



When a regular WCF queued service processes a queued call and access `MsmqMessageProperty`, both `DeliveryStatus` and `DeliveryFailure` are set to null.

Implementing a DLQ service

The DLQ service cannot affect a message's properties (for example, extending its time to live). Handling of delivery failures typically involves some kind of compensating workflow: notifying the administrator; trying to resend a new message, or resending a new request with extended timeout; logging the error; or perhaps doing nothing (i.e., merely processing the failed call and returning, thus discarding the message).

Example 9-18 demonstrates a possible DLQ service implementation.

Example 9-18. Implementing a DLQ service

```

[ServiceBehavior(InstanceContextMode = InstanceContextMode.PerCall)]
class MyDLQService : IMyContract
{
    [OperationBehavior(TransactionScopeRequired = true)]
    public void MyMethod(string someValue)
    {
        MsmqMessageProperty msmqProperty = OperationContext.Current.
            IncomingMessageProperties[MsmqMessageProperty.Name] as MsmqMessageProperty;
        //If tried more than 25 times: discard message
        if(msmqProperty.MoveCount >= 25)
        {
            return;
        }
        //If timed out: try again
        if(msmqProperty.DeliveryStatus == DeliveryStatus.InDoubt)
        {
            if(msmqProperty.DeliveryFailure == DeliveryFailure.ReceiveTimeout)
            {
                MyContractClient proxy = new MyContractClient();
                proxy.MyMethod(someValue);
                proxy.Close();
            }
        }
    }
}

```

```

        return;
    }
    if(msmqProperty.DeliveryStatus == DeliveryStatus.InDoubt ||
       msmqProperty.DeliveryFailure == DeliveryFailure.Unknown)
    {
        NotifyAdmin();
    }
}
void NotifyAdmin()
{...}
}

```

The DLQ service in [Example 9-18](#) examines the cause of the failure. If WCF has tried more than 25 times to deliver the message, the DLQ service simply gives up and drops the message. If the cause for the failure was a timeout, the DLQ service tries again by creating a proxy to the queued service and calling it, passing the same arguments from the original call (the in-parameters to the DLQ service operation). If the message is in-doubt or an unknown failure took place, the service notifies the application administrator.

Playback Failures

Even after successful delivery, a message may still fail during playback to the service. Such failures typically abort the playback transaction, which causes the message to return to the service queue. WCF will then detect the message in the queue and retry. If the next call fails too, the message will go back to the queue again, and so on. Continuously retrying this way is often unacceptable. If the initial motivation for the queued service was load leveling, WCF's auto-retry behavior will generate considerable stress on the service. You need a smart failure-handling schema that deals with the case when the call never succeeds (and, of course, defines "never" in practical terms). The failure handling will determine after how many attempts to give up, after how long to give up, and even the interval at which to try. Different systems need different retry strategies and have different sensitivity to the additional thrashing and probability of success. For example, retrying 10 times with a single retry once every hour is not the same strategy as retrying 10 times at 1-minute intervals, or the same as retrying 5 times, with each attempt consisting of a batch of 2 successive retries separated by a day. In general, it is better to hedge your bets on the causes for the failure and the probability of future success by retrying in a series of batches, to deal with sporadic and intermediate infrastructure issues as well as fluctuating application state. A series of batches, each batch comprised of a set number of retries in rapid succession, may just be able to catch the system in a state that will allow the call to succeed. If it doesn't, deferring some of the retries to a future batch allows the system some time to recuperate. Additionally, once you have given up on retries, what should you do with the failed message, and what should you acknowledge to its sender?

Poison Messages

Transactional messaging systems are inherently susceptible to repeated failure, because the retries thrashing can bring the system to its knees. Messages that continuously fail playbacks are referred to as *poison messages*, because they literally poison the system with futile retries. Transactional messaging systems must actively detect and eliminate poison messages. Since there is no telling whether just one more retry might actually succeed, you can use the following simple heuristic: all things being equal, the more the message fails, the higher the likelihood is of it failing again. For example, if the message has failed just once, retrying seems reasonable. But if the message has already failed 1,000 times, it is very likely it will fail again the 1,001st time, so it is pointless to try again. In this case, the message should be deemed a poison message. What exactly constitutes “pointless” (or just wasteful) is obviously application-specific, but it is a configurable decision. `MsmqBindingBase` offers a number of properties governing the handling of playback failures:

```
public abstract class MsmqBindingBase : Binding, ...
{
    //Poison message handling
    public int ReceiveRetryCount
    {get;set;}

    public int MaxRetryCycles
    {get;set;}

    public TimeSpan RetryCycleDelay
    {get;set;}

    public ReceiveErrorHandling ReceiveErrorHandling
    {get;set;}

    //More members
}
```

Poison Message Handling

WCF retries playing back a failed message in series of batches, for the reasoning just presented. WCF provides each queued endpoint with a retry queue and an optional poison messages queue. After all the calls in the batch have failed, the message does not return to the endpoint queue. Instead, it goes to the retry queue (WCF will create that queue on the fly). Once the message is deemed poisonous, you may have WCF move that message to the poison queue.

Retry batches

In each batch, WCF will immediately retry for `ReceiveRetryCount` times after the first call failure. `ReceiveRetryCount` defaults to five retries, or a total of six attempts,

including the first attempt. After a batch has failed, the message goes to the retry queue. After a delay of `RetryCycleDelay` minutes, the message is moved from the retry queue to the endpoint queue for another retry batch. The retry delay defaults to 30 minutes. Once that batch fails, the message goes back to the retry queue, where it will be tried again after the delay has expired. Obviously, this cannot go on indefinitely. The `MaxRetryCycles` property controls how many batches at the most to try. The default of `MaxRetryCycles` is two cycles only, resulting in three batches in total. After `MaxRetryCycles` number of retry batches, the message is considered a poison message.

When configuring non-default values for `MaxRetryCycles`, I recommend setting its value in direct proportion to `RetryCycleDelay`. The reason is that the longer the delay is, the more tolerant your system will be of additional retry batches, because the overall stress will be somewhat mitigated (having been spread over a longer period of time). With a short `RetryCycleDelay`, you should minimize the number of allowed batches, because you are trying to avoid approximating continuous thrashing.

Finally, the `ReceiveErrorHandling` property governs what to do after the last retry fails and the message is deemed poisonous. The property is of the enum type `ReceiveErrorHandling`, defined as:

```
public enum ReceiveErrorHandling
{
    Fault,
    Drop,
    Reject,
    Move
}
```

ReceiveErrorHandling.Fault

The `Fault` value considers the poison message as a catastrophic failure and actively faults the MSMQ channel and the service host. Doing so prevents the service from processing any other messages, be they from a queued client or a regular connected client. The poison message will remain in the endpoint queue and must be removed from it explicitly by the administrator or by some compensating logic, since WCF will refuse to process it again if you merely restart the host. In order to continue processing client calls of any sort, you must open a new host (after you have removed the poison message from the queue). While you could install an error-handling extension (as discussed in [Chapter 6](#)) to do some of that work, in practice there is no avoiding involving the application administrator.

`ReceiveErrorHandling.Fault` is the default value of the `ReceiveErrorHandling` property. With this setting, no acknowledgment of any sort is sent to the sender of the poison message. `ReceiveErrorHandling.Fault` is both the most conservative poison

message strategy and the least useful from the system perspective, since it amounts to a stalemate.

ReceiveErrorHandling.Drop

The `Drop` value, as its name implies, silently ignores the poison message by dropping it and having the service keep processing other messages. You should configure for `ReceiveErrorHandling.Drop` if you have high tolerance for both errors and retries. If the message is not crucial (i.e., it is used to invoke a nice-to-have operation), dropping and continuing is acceptable. In addition, while `ReceiveErrorHandling.Drop` does allow for retries, conceptually you should not have too many retries—if you care that much about the message succeeding, you should not just drop it after the last failure.

Configuring for `ReceiveErrorHandling.Drop` also sends an ACK to the sender, so from the sender's perspective, the message was delivered and processed successfully. For many applications, `ReceiveErrorHandling.Drop` is an adequate choice.

ReceiveErrorHandling.Reject

The `ReceiveErrorHandling.Reject` value actively rejects the poison message and refuses to have anything to do with it. Similar to `ReceiveErrorHandling.Drop`, it drops the message, but it also sends a NACK to the sender, thus signaling ultimate delivery and processing failure. The sender responds by moving the message to the sender's dead-letter queue. `ReceiveErrorHandling.Reject` is a consistent, defensive, and adequate option for the vast majority of applications, yet be aware it is not the default.

ReceiveErrorHandling.Move

The `ReceiveErrorHandling.Move` value is the advanced option for services that wish to defer judgment on the failed message to a dedicated third party. `ReceiveErrorHandling.Move` moves the message to the dedicated poison messages queue, and it does not send back an ACK or a NACK. Acknowledging processing of the message will be done after it is processed from the poison messages queue. While `ReceiveErrorHandling.Move` is a great choice if indeed you have some additional error recovery or compensation workflow to execute in case of a poison message, a relatively smaller set of applications will find it useful, due to its increased complexity and intimate integration with the system.

Configuration sample

Example 9-19 shows a configuration section from a host config file, configuring poison message handling.

Example 9-19. Poison message handling

```
<bindings>
<netMsmqBinding>
<binding name = "PoisonMessageHandling"
    receiveRetryCount      = "2"
    retryCycleDelay        = "00:05:00"
    maxRetryCycles         = "2"
    receiveErrorHandling   = "Move"
/>
</netMsmqBinding>
</bindings>
```

Figure 9-9 illustrates graphically the resulting behavior in the case of a poison message from Example 9-19.

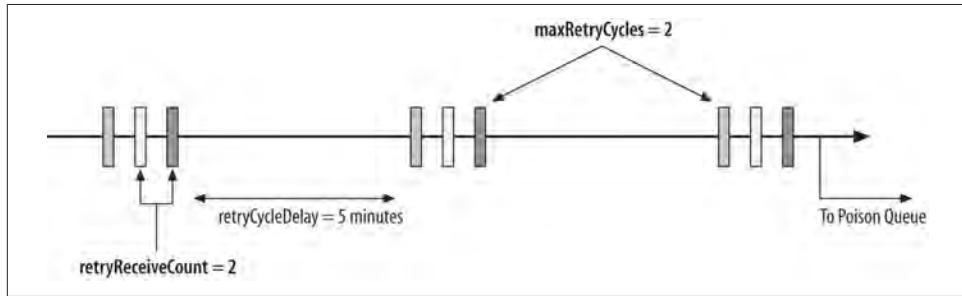


Figure 9-9. Poison message handling

Poison message service

Your service can provide a dedicated poison message-handling service to handle messages posted to its poison messages queue when the binding is configured with `ReceiveErrorHandling.Move`. The poison message service must be polymorphic with the service's queued endpoint contract. WCF will retrieve the poison message from the poison queue and play it to the poison service. It is therefore important that the poison service does not throw unhandled exceptions or abort the playback transaction (configuring it to ignore the playback transaction, as in Example 9-9, or to use a new transaction, as in Example 9-10, is a good idea). Such a poison message service typically engages in some kind of compensating work associated with the failed message, such as refunding a customer for a missing item in the inventory. Alternatively, a poison service could do any number of things, including notifying the administrator, logging the error, or just ignoring the message altogether by simply returning.

The poison message service is developed and configured like any other queued service. The only difference is that the endpoint address must be the same as the original endpoint address, suffixed by `;poison`. Example 9-20 demonstrates the required configuration of a service and its poison message service. In Example 9-20, the service

and its poison message service share the same host process, but that is certainly optional.

Example 9-20. Configuring a poison message service

```
<system.serviceModel>
    <services>
        <service name = "MyService">
            <endpoint
                address = "net.msmq://localhost/private/MyServiceQueue"
                binding = "netMsmqBinding"
                bindingConfiguration = "PoisonMessageSettings"
                contract = "IMyContract"
            />
        </service>
        <service name = "MyPoisonServiceMessageHandler">
            <endpoint
                address = "net.msmq://localhost/private/MyServiceQueue;poison"
                binding = "netMsmqBinding"
                contract = "IMyContract"
            />
        </service>
    </services>
    <bindings>
        <netMsmqBinding>
            <binding name = "PoisonMessageSettings"
                receiveRetryCount      = "..."
                retryCycleDelay       = "..."
                maxRetryCycles        = "..."
                receiveErrorHandling = "Move"
            />
        </netMsmqBinding>
    </bindings>
</system.serviceModel>
```

Receive Context

As mentioned at the beginning of this chapter, you should avoid having two or more service endpoints monitoring the same queue, since this will result in them processing each other's messages. You may be tempted, however, to leverage this behavior as a load balancing of sort: you can deploy your service on multiple machines while having all the machines share the same queue. The problem in this scenario is poison message handling. It is possible for one service to return the message to the queue after an error for a retry, and then have a second service start processing that message, not knowing it was already tried or how many times. I believe the fundamental problem here is not with load balancing queued calls and playback errors; rather, it is with the need to load balance the queued calls in the first place. Load balancing is done in the interest of scalability and throughput. Both scalability and throughput have a tem-

poral quality—they imply a time constraint on the level of performance of the service, and yet queued calls, by their very nature, indicate that the client does not care exactly when the calls execute.

Nonetheless, to enable multiple services to share a queue and manage playback errors, WCF defines the helper class `ReceiveContext`, defined as:

```
public abstract class ReceiveContext
{
    public virtual void Abandon(TimeSpan timeout);
    public virtual void Complete(TimeSpan timeout);

    public static bool TryGet(Message message,
                             out ReceiveContext property);
    public static bool TryGet(MessageProperties properties,
                             out ReceiveContext property);
    //More members
}
```

You enable the use of `ReceiveContext` with the `ReceiveContextEnabled` attribute:

```
public sealed class ReceiveContextEnabledAttribute : Attribute,
                                                IOperationBehavior
{
    public bool ManualControl
    {get;set;}

    //More members
}
```

After a failure, you can use `ReceiveContext` to lock the message in the queue and prevent other services from processing it. This, however, results in a cumbersome programming model that is nowhere near as elegant as the transaction-driven queued services. I recommend you design the system so that you do not need to load balance your queued services and that you avoid `ReceiveContext` altogether.

Queued Versus Connected Calls

Although it is technically possible to use the same service code both connected and queued (with simple changes such as configuring operations as one-way, or adding another contract for the one-way operations), in reality it is unlikely that you will actually use the same service both ways. The reasons are similar to the arguments made in the context of asynchronous calls, discussed in [Chapter 8](#). Synchronous calls and asynchronous calls addressing the same business scenario often have to use different workflows, and these differences will necessitate changes to the service code to adapt it for each case. The use of queued calls adds yet another barrier for using the same service code (both connected and disconnected): changes to the transactional

semantics of the service. Consider, for example, [Figure 9-10](#), which depicts an online store application that uses connected calls only.

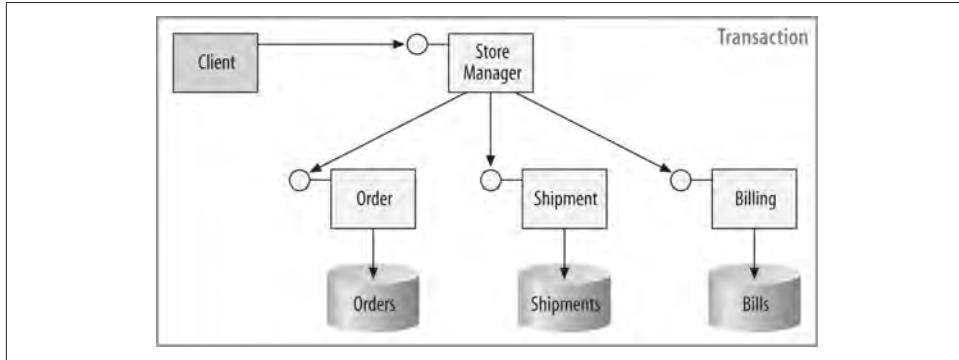


Figure 9-10. A connected application relies on a single transaction

The Store service uses three well-factored helper services to process the order: Order, Shipment, and Billing. In the connected scenario, the Store service calls the Order service to place the order. Only if the Order service succeeds in processing the order (that is, if the item is available in the inventory) does the Store service call the Shipment service, and only if the Shipment service succeeds does the Store service access the Billing service to bill the customer. The connected case involves exactly one transaction created by the client, and all operations commit or abort as one atomic operation. Now, suppose the Billing service also exposes a queued endpoint for the use of the Store service, as shown in [Figure 9-11](#).

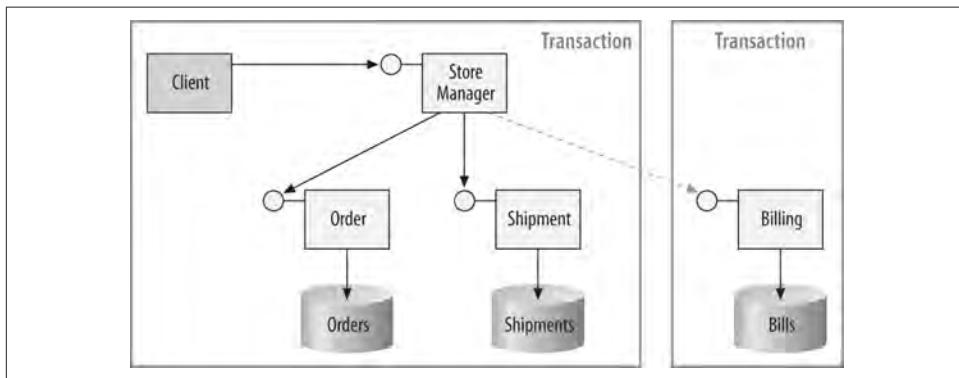


Figure 9-11. A disconnected application relies on multiple transactions

The queued call to the Billing service will be played to the service in a separate transaction from that of the rest of the store, and it could commit or abort separately from the transaction that groups Order and Shipment. This, in turn, could jeopardize

the system's consistency, so you must include some logic in the Billing service to detect the failure of the other service and to initiate some compensating logic in the event that it fails to do its work. As a result, the Billing service will no longer be the same service used in the connected case.

Requiring Queuing

Since not every service can be connected and queued, and since some services may be designed for a particular option and only that option, WCF lets you constrain a service's communication pattern. The `DeliveryRequirements` attribute presented in [Chapter 1](#) also lets you insist on queued or connected delivery of messages to the service:

```
public enum QueuedDeliveryRequirementsMode
{
    Allowed,
    Required,
    NotAllowed
}
[AttributeUsage(AttributeTargets.Interface|AttributeTargets.Class,
    AllowMultiple = true)]
public sealed class DeliveryRequirementsAttribute : Attribute, ...
{
    public QueuedDeliveryRequirementsMode QueuedDeliveryRequirements
    {get;set;}
    public bool RequireOrderedDelivery
    {get;set;}
    public Type TargetContract
    {get;set;}
}
```

This attribute can be used to constrain a contract (and all its supporting endpoints) or a particular service type. The default value of the `QueuedDeliveryRequirements` property is `QueuedDeliveryRequirementsMode.Allowed`, so these definitions are equivalent:

```
[ServiceContract]
interface IMyContract
{...}

[ServiceContract]
[DeliveryRequirements]
interface IMyContract
{...}

[ServiceContract]
[DeliveryRequirements(QueuedDeliveryRequirements =
    QueuedDeliveryRequirementsMode.Allowed)]
interface IMyContract
{...}
```

`QueuedDeliveryRequirementsMode.Allowed` grants permission for using the contract or the service with either connected or queued calls. `QueuedDeliveryRequirementsMode.NotAllowed` explicitly disallows the use of the MSMQ binding, so all calls on the endpoint must be connected calls. Use this value when the contract or the service is explicitly designed to be used in a connected fashion only. `QueuedDeliveryRequirementsMode.Required` is the opposite: it mandates the use of the MSMQ binding on the endpoint, and it should be used when the contract or the service is designed from the ground up to be queued.

Even though the `DeliveryRequirements` attribute offers the `RequireOrderedDelivery` property (discussed in [Chapter 1](#)), if `QueuedDeliveryRequirementsMode.Required` is used, then `RequireOrderedDelivery` must be `false`, because queued calls inherently are unordered and messages may be played back in any order.

When the `DeliveryRequirements` attribute is applied on an interface, it affects all services that expose endpoints with that contract:

```
[ServiceContract]
[DeliveryRequirements(QueuedDeliveryRequirements =
                      QueuedDeliveryRequirementsMode.Required)]
interface IMyQueuedContract
{...}
```

The client as well can apply the `DeliveryRequirements` attribute on its copy of the service contract.

When the `DeliveryRequirements` attribute is applied on a service class, it affects all endpoints of that service:

```
[DeliveryRequirements(QueuedDeliveryRequirements =
                      QueuedDeliveryRequirementsMode.Required)]
class MyQueuedService : IMyQueuedContract,IMyOtherContract
{...}
```

When applied on a service class while using the `TargetContract` property, the attribute affects all endpoints of the service that expose the specified contract:

```
[DeliveryRequirements(TargetContract = typeof(IMyQueuedContract),
                      QueuedDeliveryRequirements =
                      QueuedDeliveryRequirementsMode.Required)]
class MyService : IMyQueuedContract,IMyOtherContract
{...}
```

The Response Service

The programming model of queued calls described so far was one-sided: the client posted a one-way message to a queue, and the service processed that message. This model is sufficient when the queued operations are one-way calls by nature. However, the queued service may need to report back to its client on the result of the invoca-

tion, or return results or even errors. By default, this is not possible: WCF equates queued calls with one-way calls, which inherently forbids any such response. In addition, queued services (and their clients) are potentially disconnected. If a client posts a queued call to a disconnected service, by the time the service finally gets the message and processes it, there may no longer be a client to return the values to. The solution is to have the service report back to a client-provided queued service. I call such a service a *response service*.¹ **Figure 9-12** shows the architecture of such a solution.

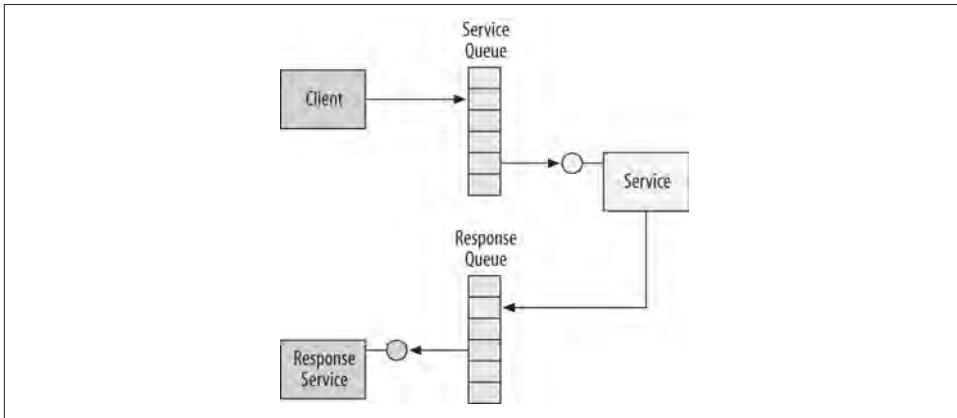


Figure 9-12. A response service

The response service is just another queued service in the system. The response service may be disconnected toward the client as well, or it may share the client's process, or it may be hosted in a separate process or even on a separate machine. If the response service shares the client's process, when the client is launched the response service will start processing the queued responses. Having the response service in a separate process (or even on a separate machine) from the client's helps to further decouple lifeline-wise the response service from the client or clients that use it.



Not all queued services require a response service. Be pragmatic, and use a response service only where appropriate; that is, where it adds the most value.

Designing a Response Service Contract

As with any WCF service, the client and the service need to agree beforehand on the response contract and what it will be used for; that is, whether it will be used for

¹ I first published my initial technique for a response service in the February 2007 issue of *MSDN Magazine*.

returned values and error information, or just returned values. Note that you can also split the response service into two services, and have one response service for results and another for faults and errors. As an example, consider the `ICalculator` contract implemented by the queued `MyCalculator` service:

```
[ServiceContract]
interface ICalculator
{
    [OperationContract(IsOneWay = true)]
    void Add(int number1,int number2);
    //More operations
}
class MyCalculator : ICalculator
{...}
```

The `MyCalculator` service is required to respond to its client with the result of the calculation and report on any errors. The result of the calculation is an integer, and the error is in the form of the `ExceptionDetail` data contract presented in [Chapter 6](#). The `ICalculatorResponse` contract could be defined as:

```
[ServiceContract]
interface ICalculatorResponse
{
    [OperationContract(IsOneWay = true)]
    void OnAddCompleted(int result,ExceptionDetail error);
    //More operations
}
```

The response service supporting `ICalculatorResponse` needs to examine the returned error information; notify the client application, the user, or the application administrator on the method completion; and make the results available to the interested parties. [Example 9-21](#) shows a simple response service that supports `ICalculatorResponse`.

Example 9-21. A simple response service

```
class MyCalculatorResponse : ICalculatorResponse
{
    public void OnAddCompleted(int result,ExceptionDetail error)
    {
        if(error != null)
        {
            //Handle error
        }
        else
        {
            MessageBox.Show("Result = " + result,"MyCalculatorResponse");
        }
    }
}
```

```
//More operations  
}
```

As demonstrated by [Example 9-21](#), the response service is just that—a simple service. There is nothing special about it other than its designation as a response service.

Response address and method ID

There are two immediate problems with the implementation of both `MyCalculator` and `MyCalculatorResponse`. The first is that the same response service could be used to handle the response (or completion) of multiple calls on multiple queued services, and yet, as listed in [Example 9-21](#), `MyCalculatorResponse` (and more importantly, the clients it serves) has no way of distinguishing between responses. The solution for that is to have the client that issued the original queued call tag the call by associating it with some unique ID, or at least an ID that is unique enough across that client's application. The queued service `MyCalculator` needs to pass that ID to the response service `MyCalculatorResponse`, so that it can apply its custom logic regarding that ID. Note that the service typically has no direct use for the ID; all it needs to do is pass it along.

The second problem is how to enable the queued service to discover the address of the response service. Unlike with duplex callbacks, there is no built-in support in WCF for passing the response service's reference to the queued service, so the queued service needs to manually construct a proxy to the response service and invoke the operations of the response contract. While the response contract is decided upon at design time, and the binding is always `NetMsmqBinding`, the queued service lacks the address of the response service to be able to respond. You could place that address in the service host config file (in a `client` section) but such a course of action is to be avoided. The main reason is that the same queued service could be called by multiple clients, each with its own dedicated response service and address.

One possible solution is to explicitly pass both the client-managed ID and the desired response service address as parameters to every operation on the queued service contract:

```
[ServiceContract]  
interface ICalculator  
{  
    [OperationContract(IsOneWay = true)]  
    void Add(int number1,int number2,string responseAddress, string methodId);  
}
```

Much the same way, the queued service could explicitly pass the method ID to the response service as a parameter to every operation on the queued response contract:

```
[ServiceContract]  
interface ICalculatorResponse  
{
```

```

    [OperationContract(IsOneWay = true)]
    void OnAddCompleted(int result,ExceptionDetail error,string methodId);
}

```

The ResponseContext class

While passing the address and the ID as explicit parameters would work, it does distort the original contract, and it introduces plumbing-level parameters alongside business-level parameters in the same operation. A better solution is to have the client store the response address and operation ID in the outgoing message headers of the call. Using the message headers this way is a general-purpose technique for passing out-of-band information to the service (information that is otherwise not present in the service contract). [Appendix B](#) explains in detail the use of the incoming and outgoing headers, including the related techniques and supporting classes in *ServiceModelEx*.

Since the client needs to pass both the address and the method ID in the message headers, a single primitive type parameter will not do. Instead, use my ResponseContext class, defined in [Example 9-22](#).

Example 9-22. The ResponseContext class

```

[DataContract]
public class ResponseContext
{
    [DataMember]
    public readonly string ResponseAddress;

    [DataMember]
    public readonly string FaultAddress;

    [DataMember]
    public readonly string MethodId;

    public ResponseContext(string responseAddress,string methodId) :
        this(responseAddress,methodId,null)
    {}
    public ResponseContext(string responseAddress) : this(responseAddress,
        Guid.NewGuid().ToString())
    {}
    public ResponseContext(string responseAddress,string methodId,
        string faultAddress)
    {
        ResponseAddress = responseAddress;
        MethodId = methodId;
        FaultAddress = faultAddress;
    }
    public static ResponseContext Current

```

```

{
    get
    {
        return GenericContext<ResponseContext>.Current.Value;
    }
    set
    {
        GenericContext<ResponseContext>.Current =
            new GenericContext<ResponseContext>(value);
    }
}
//More members
}

```

ResponseContext provides a place to store both the response address and the ID. In addition, if the client wants to use a separate response service for faults, ResponseContext provides a field for the fault response service address. (This chapter makes no use of that feature.) The client is responsible for constructing an instance of ResponseContext with a unique ID. While the client can supply that ID as a construction parameter, the client can also use the constructor of ResponseContext, which takes just the response address, and have that constructor generate a GUID for the ID. To streamline the act of storing a ResponseContext instance in and retrieving it from the headers, ResponseContext provides the Current property, which merely encapsulates my GenericContext<T>. The client can provide an ID for each method call (even when dealing with a sessionful queued service) by using a different instance of ResponseContext for each call.

Client-Side Programming

My HeaderClientBase<T,H> proxy base class (defined in [Appendix B](#)) is designed to automate passing information in the headers from the client to the service:

```

public abstract class InterceptorClientBase<T> : ClientBase<T> where T : class
{
    protected virtual void PreInvoke(ref Message request)
    {}
    //More members
}public abstract partial class HeaderClientBase<T,H> : InterceptorClientBase<T>
    where T : class
{
    public H Header
    {get;protected set;}

    public HeaderClientBase(H header);
    public HeaderClientBase(H header,string endpointName);

    //More members
}

```

However, when it comes to the response context, there are differences compared with the raw headers management discussed in [Appendix B](#): specifically, changing the context (that is, the headers) on each call as opposed to setting it only at construction time, generating method IDs and providing them to the client, and enqueueing rather than merely invoking the service call. While the client can easily use my `HeaderClientBase<T,H>` to do all that, all clients will have to repeat such code for every contract and proxy they have. It is better to automate and encapsulate these steps in a dedicated proxy base class such as my `ClientResponseBase<T>`, shown in [Example 9-23](#).

Example 9-23. The ClientResponseBase<T> class

```
public abstract class ClientResponseBase<T> :  
    HeaderClientBase<T,ResponseContext> where T : class  
{  
    protected readonly string ResponseAddress;  
  
    public ClientResponseBase(string responseAddress)  
    {  
        ResponseAddress = responseAddress;  
        Endpoint.VerifyQueue();  
    }  
    public ClientResponseBase(string responseAddress, string endpointName)  
    {...}  
    public ClientResponseBase(string responseAddress,  
        NetMsmqBinding binding, EndpointAddress address)  
    {...}  
    /* More constructors */  
  
    protected override void PreInvoke(ref Message request)  
    {  
        string methodId = GenerateMethodId();  
        Header = new ResponseContext(ResponseAddress, methodId);  
        base.PreInvoke(ref request);  
    }  
    protected virtual string GenerateMethodId()  
    {  
        return Guid.NewGuid().ToString();  
    }  
}
```

The constructors of `ClientResponseBase<T>` accept the response address and the regular proxy parameters, such as the endpoint name, address, and binding. The constructors store the response address in the read-only public field `ResponseAddress`. In addition, the constructors use the `VerifyQueue()` endpoint extension method to verify that the service queue (and the DLQ) exists and to create it if necessary.

`ClientResponseBase<T>` provides the virtual `GenerateMethodId()` method, which by default uses a GUID for the method ID. However, your subclasses of `ClientResponseBase<T>` can override it and provide their own unique strings, such as an incremented integer.

The heart of `ClientResponseBase<T>` is the overridden `PreInvoke()` method. `PreInvoke()` is defined as virtual in the `InterceptorClientBase<T>` base class of `HeaderClientBase<T,H>`. `InterceptorClientBase<T>` is part of a generic interception framework I wrote (defined in [Appendix E](#)) that enables you to perform custom pre-call and post-call interception steps. For every operation invoked by the client, `PreInvoke()` generates a new method ID, provides it to a new `ResponseContext` object (along with the response address supplied to the constructor), and assigns the new `ResponseContext` object to the `Header` property of `HeaderClientBase<T,H>`. Thanks to generics, `Header` is of the type `ResponseContext`.

Using `ClientResponseBase<T>`

You use `ClientResponseBase<T>` like a regular proxy; for example, given this calculator contract:

```
[ServiceContract]
interface ICalculator
{
    [OperationContract(IsOneWay = true)]
    void Add(int number1,int number2);
    //More operations
}
```

[Example 9-24](#) shows the matching service proxy.

Example 9-24. Deriving from ClientResponseBase<T>

```
class CalculatorClient : ClientResponseBase<ICalculator>,ICalculator
{
    public CalculatorClient(string responseAddress) : base(responseAddress)
    {}
    public CalculatorClient(string responseAddress,string endpointName) :
        base(responseAddress,endpointName)
    {}
    public CalculatorClient(string responseAddress,
                           NetMsmqBinding binding,EndpointAddress address) :
        base(responseAddress,binding,address)
    {}
    //More constructors

    public void Add(int number1,int number2)
    {
        Channel.Add(number1,number2);
    }
}
```

```
//More operations  
}
```

Using the proxy in [Example 9-24](#) yields this straightforward client code:

```
string responseAddress =  
    "net.msmq://localhost/private/MyCalculatorResponseQueue";  
  
CalculatorClient proxy = new CalculatorClient(responseAddress);  
proxy.Add(2,3);  
proxy.Close();
```

Note how closely the client that provides the response address to the proxy corresponds to a client that provides a duplex callback object to a proxy (as demonstrated in [Chapter 5](#)). In the queued services world, the response service address is the equivalent callback reference.



A queued response service is not limited to being used only with a queued service. You can use the same technique to pass the address and method ID to a connected service and have that service respond to a client-provided queued response service. You will need to rework `ClientResponseBase<T>` so that it uses only `Binding`.

When managing the responses on the client side using a `ClientResponseBase<T>`-derived proxy, it is often very handy to have the invoking client obtain the method ID used to dispatch the call. You can do this easily with the `Header` property:

```
CalculatorClient proxy = new CalculatorClient(responseAddress);  
proxy.Add(2,3);  
string methodId = proxy.Header.MethodId;  
proxy.Close();
```

Queued Service-Side Programming

The service needs to construct a proxy that will dispatch messages to the client-side response service. To simplify this, use my `ServiceResponseBase<T>`, defined as:

```
public abstract class ServiceResponseBase<T> : HeaderClientBase  
    <T,ResponseContext> where T : class  
{  
    public ServiceResponseBase();  
    public ServiceResponseBase(string bindingName);  
    public ServiceResponseBase(NetMsmqBinding binding);  
}
```

`ServiceResponseBase<T>` automates reading the response context from the message headers and writing the response itself to the outgoing headers. Other than that, you

can use it like any other proxy base class. [Example 9-25](#) demonstrates the usage of `ServiceResponseBase<T>`.

Example 9-25. Using ServiceResponseBase<T>

```
class CalculatorResponseClient : ServiceResponseBase<ICalculatorResponse>,
                                ICalculatorResponse
{
    public void OnAddCompleted(int result,ExceptionDetail error)
    {
        Channel.OnAddCompleted(result,error);
    }
}
class MyCalculator : ICalculator
{
    [OperationBehavior(TransactionScopeRequired = true)]
    public void Add(int number1,int number2)
    {
        int result = 0;
        ExceptionDetail error = null;

        try
        {
            result = number1 + number2;
        }
        //Don't rethrow
        catch(Exception exception)
        {
            error = new ExceptionDetail(exception);
        }
        finally
        {
            CalculatorResponseClient proxy = new CalculatorResponseClient();

            proxy.OnAddCompleted(result,error);

            proxy.Close();
        }
    }
}
```

In [Example 9-25](#), the `MyCalculator` service catches any exception thrown by the business logic operation and wraps that exception with an `ExceptionDetail` object. The service does not rethrow the exception. As you will see later, in the context of transactions and response services, rethrowing the exception would also cancel the response. Moreover, when using a response service, being able to respond in case of an error is a much better strategy than relying on WCF's playback error handling.

In the `finally` statement, regardless of exceptions, the service responds. It creates a new proxy to the response service to enqueue the response. The proxy in [Example 9-25](#) will default to using the same MSMQ binding as the host.

[Example 9-26](#) shows the implementation of `ServiceResponseBase<T>`.

Example 9-26. Implementing ServiceResponseBase<T>

```
public abstract class ServiceResponseBase<T> : HeaderClientBase<T,ResponseContext>
{
    public ServiceResponseBase() : this(OperationContext.Current.Host,
                                         Description.Endpoints[0].Binding as NetMsmqBinding)
    {}
    public ServiceResponseBase(string bindingName) :
        this(new NetMsmqBinding(bindingName))
    {}

    public ServiceResponseBase(NetMsmqBinding binding) :
        base(ResponseContext.Current,binding,
             new EndpointAddress(ResponseContext.Current.ResponseAddress))
    {
        Endpoint.VerifyQueue();
    }
}
```

The default constructor of `ServiceResponseBase<T>` uses the same queued binding the host was using to dequeue the client's call. You can also specify an MSMQ binding section in the config file or provide the constructor with the binding instance to use. Both of these constructors delegate the work to the third constructor, which accepts the MSMQ binding to use. That constructor reads the response address out of the response context and provides those two along with the response context to the base constructor of `HeaderClientBase<T,H>`. It also verifies the presence of the response queue.

Note that `ServiceResponseBase<T>` sends the response service the entire response context (not just the ID). This is done both for simplicity's sake and because it may be beneficial for the response service to have access to the fault and response address used.

Response Service-Side Programming

The response service accesses its response context, reads from it the method ID, and responds accordingly. [Example 9-27](#) demonstrates a possible implementation of such a response service.

Example 9-27. Implementing a response service

```
class MyCalculatorResponse : ICalculatorResponse
{
    [OperationBehavior(TransactionScopeRequired = true)]
    public void OnAddCompleted(int result, ExceptionDetail error)
    {
        string methodId = ResponseContext.Current.MethodId;
        ...
    }
}
```



It is common for the response service to update the application's user interfaces with the queued results (or errors). [Chapter 8](#) introduced my `FormHost<F>` class, which you can certainly leverage to support the queued response contract. For example:

```
class CalculatorResponse :
    FormHost<CalculatorResponse>,
    ICalculatorResponse
{
    [OperationBehavior(TransactionScopeRequired=true)]
    public void OnAddCompleted(int result,
                               ExceptionDetail error)
    {
        Text = "Add returned: " + result;
        ...
    }
}
```



In fact, nothing prevents you from having the client itself be the response service as well.

Transactions

A queued service typically queues up the response as part of the incoming playback transaction. Given the queued service definition of [Example 9-28](#), [Figure 9-13](#) depicts the resulting transaction and the participating actions.

Example 9-28. Queuing up a response as part of the playback transaction

```
class MyCalculator : ICalculator
{
    [OperationBehavior(TransactionScopeRequired = true)]
    public void Add(int number1, int number2)
```

```

{
    ...
    ...
    try
    {
        ...
    }
    catch //Do not rethrow
    {
        ...
    }
    finally
    {
        CalculatorResponseClient proxy = new CalculatorResponseClient();

        proxy.OnAddCompleted(result,error);

        proxy.Close();
    }
}
}

```

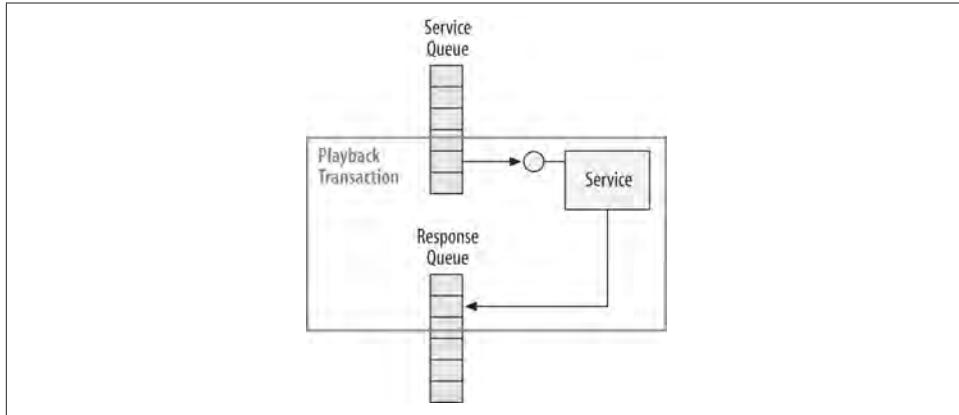


Figure 9-13. Queuing up in the playback transaction

Design-wise, the nice thing about having the queued call playback and the queued response in the same transaction is that if the playback transaction is aborted for whatever reason (including due to other services in the transaction aborting), the response is canceled automatically. This is by far the most common choice for most applications.

Note in [Example 9-28](#) that the service catches all exceptions and does not rethrow them. This is important, because any unhandled exception (or rethrown exception) will abort the response, so there won't be any point in the service bothering to respond. Using a response service intrinsically means that the service does not rely on

the automatic retry mechanism of WCF, and it handles its own business logic failures because the clients expect it to respond in a prescribed manner.

Using a new transaction

As an alternative to always having the response be part of the playback transaction, the service can respond in a new transaction by encasing the response in a new transaction scope, as shown in [Example 9-29](#) and illustrated in [Figure 9-14](#).

Example 9-29. Responding in a new transaction

```
class MyCalculator : ICalculator
{
    [OperationBehavior(TransactionScopeRequired = true)]
    public void Add(int number1,int number2)
    {
        ...
        finally
        {
            using(TransactionScope transactionScope =
                new TransactionScope(TransactionScopeOption.RequiresNew))
            {
                CalculatorResponseClient proxy = new CalculatorResponseClient();

                proxy.OnAddCompleted(result,error);

                proxy.Close();
            }
        }
    }
}
```

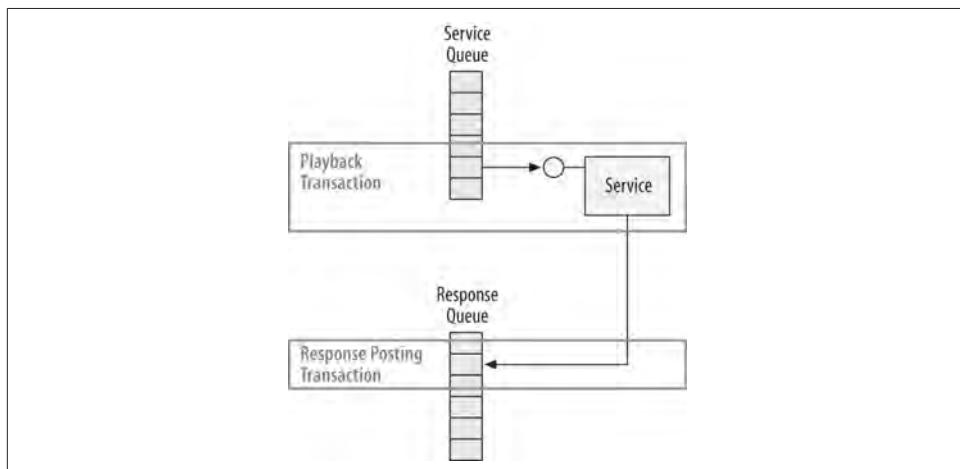


Figure 9-14. Responding in a new transaction

Responding in a new transaction is required in two cases. The first is when the service wants to respond regardless of the outcome of the playback transaction (which could be aborted by other downstream services). The second case is when the response is nice to have, and the service does not mind if the playback transaction commits but the response aborts.

Response service and transactions

Since a response service is just another queued service, the mechanics of managing and participating in a transaction are just like those of any other queued service. However, there are a few points worth mentioning in this particular context. The response service can process the response as part of the incoming response playback transaction:

```
class MyCalculatorResponse : ICalculatorResponse
{
    [OperationContract(TransactionScopeRequired = true)]
    public void OnAddCompleted(...)
    {...}
}
```

This is by far the most common option, because it allows for retries. That said, the response service should avoid lengthy processing of the queued response, because it may risk aborting the playback transaction. The response service can process the response in a separate transaction if the response is nice to have (as far as the provider of the response service is concerned):

```
class MyCalculatorResponse : ICalculatorResponse
{
    public void OnAddCompleted(int result,ExceptionDetail error)
    {
        using(TransactionScope scope = new TransactionScope())
        {...}
    }
}
```

When the response is processed in a new transaction, if that transaction aborts, WCF will not retry the response out of the response service's queue. Finally, for response processing of long duration, you can configure the response service not to use a transaction at all (including the playback transaction):

```
class MyCalculatorResponse : ICalculatorResponse
{
    public void OnAddCompleted(...)
    {...}
}
```

The HTTP Bridge

The MSMQ binding is designed to be employed in the intranet. It cannot go through firewalls by default, and more importantly, it uses a Microsoft-specific encoding and message format. Even if you could tunnel through the firewall, you would need the other party to use WCF as well. While requiring WCF at both ends is a reasonable assumption in the intranet, it is unrealistic to demand that from Internet-facing clients and services, and it violates the core service-oriented principles that service boundaries should be explicit and that the implementation technology used by a service should be immaterial to its clients. That said, Internet services may benefit from queued calls just like intranet clients and services, and yet the lack of an industry standard for such queued interoperability (and the lack of support in WCF) prevents such interaction. The solution to that problem is a technique I call the *HTTP bridge*. Unlike most of my other techniques shown in this book, the HTTP bridge is a configuration pattern rather than a set of helper classes. The HTTP bridge, as its name implies, is designed to provide queued calls support for clients and services connected over the Internet. The bridge requires the use of the `WSHttpBinding` (rather than the basic binding) because it is a transactional binding. There are two parts to the HTTP bridge. The bridge enables WCF clients to queue up calls to an Internet service that uses the WS binding, and it enables a WCF service that exposes an HTTP endpoint over the WS binding to queue up calls from its Internet clients. You can use each part of the bridge separately, or you can use them in conjunction. The bridge can only be used if the remote service contract can be queued (that is, if the contract has only one-way operations), but that is usually the case; otherwise, the client would not have been interested in the bridge in the first place.

Designing the Bridge

Since you cannot really queue up calls with the WS binding, you can facilitate that instead using an intermediary bridging client and service. When the client wishes to queue up a call against an Internet-based service, the client will in fact queue up the call against a local (that is, intranet-based) queued service called `MyClientHttpBridge`. In its processing of the queued call, the client-side queued bridge service will use the WS binding to call the remote Internet-based service. When an Internet-based service wishes to receive queued calls, it will use a queue. But because non-WCF clients cannot access that queue over the Internet, the service will use a façade: a dedicated connected service called `MyServiceHttpBridge` that exposes a WS-binding endpoint. In its processing of the Internet call, `MyServiceHttpBridge` simply makes a queued call against the local service. [Figure 9-15](#) shows the HTTP bridge architecture.

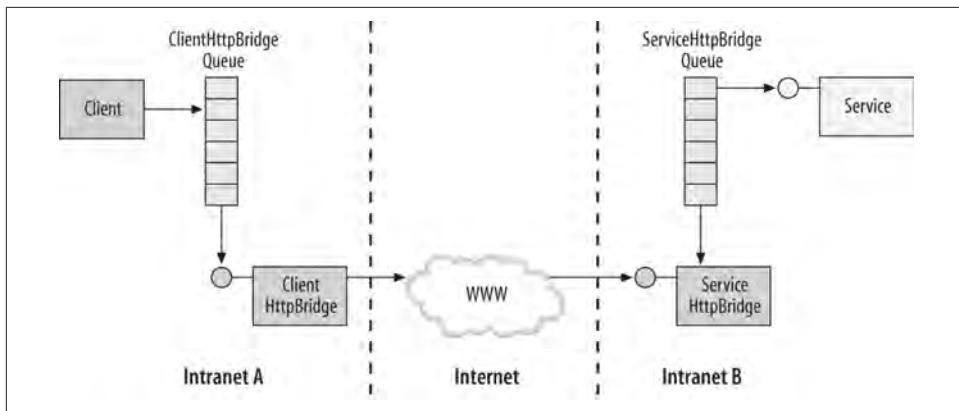


Figure 9-15. The HTTP bridge

Transaction Configuration

It is important to use transactions between `MyClientHttpBridge`, the client side of the bridge, and the remote service, and it is important to configure the service-side bridge (`MyServiceHttpBridge`) to use the Client transaction mode discussed in [Chapter 7](#). The rationale is that by using a single transaction from the playback of the client call to the `MyClientHttpBridge` to the `MyServiceHttpBridge` (if present) you will approximate the transactional delivery semantic of a normal queued call, as shown in [Figure 9-16](#).

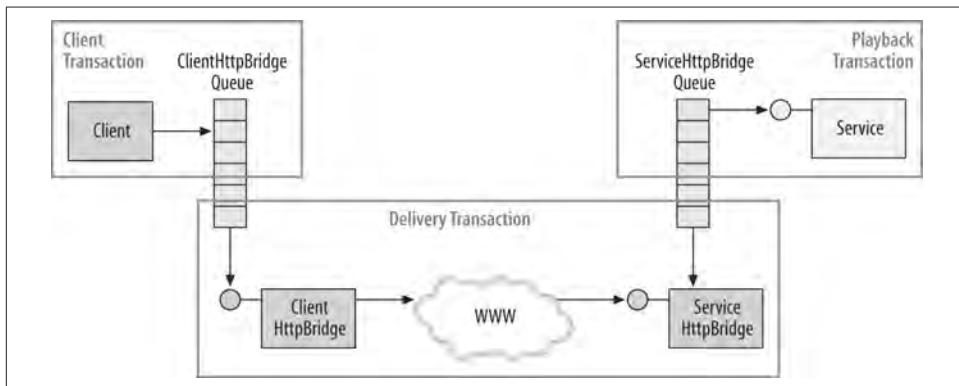


Figure 9-16. The HTTP bridge and transactions

Compare [Figure 9-16](#) with [Figure 9-6](#). If the delivery transaction in the bridge aborts for any reason, the message will roll back to the `MyClientHttpBridge` queue for another retry. To maximize the chances of successful delivery, you should also turn on reliability for the call between `MyClientHttpBridge` and the remote service.

Service-Side Configuration

MyServiceHttpBridge converts a regular connected call over the WS binding into a queued call and posts it to the service queue. MyServiceHttpBridge implements a contract that is similar, but not identical, to that of the queued service. The reason is that the service-side bridge should be able to participate in the incoming transaction, but transactions cannot flow over one-way operations. The solution is to modify the contract to support (indeed, mandate) transactions. For example, if this is the original service contract:

```
[ServiceContract]
public interface IMyContract
{
    [OperationContract(IsOneWay = true)]
    void MyMethod();
}
```

then MyServiceHttpBridge should expose this contract instead:

```
[ServiceContract]
public interface IMyContractHttpBridge
{
    [OperationContract]
    [TransactionFlow(TransactionFlowOption.Mandatory)]
    void MyMethod();
}
```

In essence, you need to set `IsOneWay` to `false` and use `TransactionFlowOption.Mandatory`. For readability's sake, I recommend that you also rename the interface by suffixing it with `HttpBridge`. MyServiceHttpBridge can be hosted anywhere in the service's intranet, including in the service's own process. [Example 9-30](#) shows the required configuration of the service and its HTTP bridge.

Example 9-30. Service-side configuration of the HTTP bridge

```
<!-- MyService Config File -->
<services>
    <service name = "MyService">
        <endpoint
            address   = "net.msmq://localhost/private/MyServiceQueue"
            binding   = "netMsmqBinding"
            contract  = "IMyContract"
        />
    </service>
</services>

<!-- MyServiceHttpBridge Config File -->
<services>
    <service name  = "MyServiceHttpBridge">
        <endpoint
```

```

address  = "http://localhost:8001/MyServiceHttpBridge"
binding  = "wsHttpBinding"
bindingConfiguration = "ReliableTransactedHTTP"
contract = "IMyContractHttpBridge"
/>
</service>
</services>

<client>
<endpoint
    address  = "net.msmq://localhost/private/MyServiceQueue"
    binding  = "netMsmqBinding"
    contract = "IMyContract"
/>
</client>

<bindings>
<wsHttpBinding>
    <binding name = "ReliableTransactedHTTP" transactionFlow = "true">
        <reliableSession enabled = "true"/>
    </binding>
</wsHttpBinding>
</bindings>

```

The service `MyService` exposes a simple queued endpoint with `IMyContract`. The service `MyServiceHttpBridge` exposes an endpoint with `WSHttpBinding` and the `IMyContractHttpBridge` contract. `MyServiceHttpBridge` is also a client of the queued endpoint defined by the service. [Example 9-31](#) shows the corresponding implementation. Note that `MyServiceHttpBridge` is configured for the Client transaction mode.

Example 9-31. Service-side implementation of the HTTP bridge

```

class MyService : IMyContract
{
    //This call comes in over MSMQ
    [OperationBehavior(TransactionScopeRequired = true)]
    public void MyMethod()
    {...}
}
class MyServiceHttpBridge : IMyContractHttpBridge
{
    //This call comes in over HTTP
    [OperationBehavior(TransactionScopeRequired = true)]
    public void MyMethod()
    {
        MyContractClient proxy = new MyContractClient();

        //This call goes out over MSMQ
        proxy.MyMethod();
    }
}

```

```

        proxy.Close();
    }
}

```

Client-Side Configuration

The client uses queued calls against the local `MyClientHttpBridge` service. `MyClientHttpBridge` can be hosted in the same process as the client, in a different process, or even on a separate machine on the client's intranet. The local `MyClientHttpBridge` service uses the `WSHttpBinding` to call the remote service. The client needs to retrieve the metadata of the remote Internet service (such as the definition of `IMyContractHttpBridge`) and convert it to a queued contract (such as `IMyContract`). [Example 9-32](#) shows the required configuration of the client and its HTTP bridge.

Example 9-32. Client-side configuration of the HTTP bridge

```

<!-- Client Config File -->
<client>
    <endpoint
        address = "net.msmq://localhost/private/MyClientHttpBridgeQueue"
        binding = "netMsmqBinding"
        contract = "IMyContract"
    />
</client>

<!-- MyClientHttpBridge Config File -->
<services>
    <service name = "MyClientHttpBridge">
        <endpoint
            address = "net.msmq://localhost/private/MyClientHttpBridgeQueue"
            binding = "netMsmqBinding"
            contract = "IMyContract"
        />
    </service>
</services>
<client>
    <endpoint
        address = "http://localhost:8001/MyServiceHttpBridge"
        binding = "wsHttpBinding"
        bindingConfiguration = "ReliableTransactedHTTP"
        contract = "IMyContractHttpBridge"
    />
</client>
<bindings>
    <wsHttpBinding>
        <binding name = "ReliableTransactedHTTP" transactionFlow = "true">
            <reliableSession enabled = "true"/>
        </binding>
    </wsHttpBinding>

```

```
</wsHttpBinding>  
</bindings>
```

MyClientHttpBridge exposes a simple queued endpoint with IMyContract. MyClientHttpBridge is also a client of the connected WS-binding endpoint defined by the service. [Example 9-33](#) shows the corresponding implementation.

Example 9-33. Client-side implementation of the HTTP bridge

```
MyContractClient proxy = new MyContractClient();  
  
//This call goes out over MSMQ  
proxy.MyMethod();  
  
proxy.Close();  
  
///////////////// Client-Side Bridge Implementation ///////////////////  
class MyClientHttpBridge : IMyContract  
{  
    //This call comes in over MSMQ  
    [OperationBehavior(TransactionScopeRequired = true)]  
    public void MyMethod()  
    {  
        MyContractHttpBridgeClient proxy = new MyContractHttpBridgeClient();  
  
        //This call goes out over HTTP  
        proxy.MyMethod();  
  
        proxy.Close();  
    }  
}
```


There are several aspects pertaining to secure interaction between a client and a service. As in traditional client/server and component-oriented applications, the service needs to authenticate its callers and often also authorize the callers before executing sensitive operations. In addition, regardless of the technology, when securing a service (and its clients) as in any distributed system, you need to secure the messages while they are en route from the client to the service. Once the messages arrive securely and are authenticated and authorized, the service has a number of options regarding the identity it uses to execute the operation. This chapter will explore these classic security aspects—authentication, authorization, transfer security, and identity management—as well as something more abstract, which I call *overall security policy*: that is, your own personal and your company’s (or customer’s) approach to and mindset regarding security. This chapter starts by defining the various aspects of security in the context of WCF and the options available to developers when it comes to utilizing WCF and .NET security. Then, it explains how to secure the canonical and prevailing types of applications. Finally, I will present my declarative security framework, which vastly reduces the complexity of the WCF security programming model by eliminating the need to understand and tweak the many details of WCF security.

Authentication

Authentication is the act of verifying that the caller of a service is indeed who that caller claims to be. While authentication is typically referred to in the context of verification of the caller, from the client perspective there is also a need for service authentication; that is, assuring the client that the service it calls really is the service it intends to call. This is especially important with clients who call over the Internet,

because if a malicious party subverts the client's DNS service, it could hijack the client's calls. WCF offers various authentication mechanisms:

No authentication

The service does not authenticate its callers, and virtually all callers are allowed.

Windows authentication

The service typically uses Kerberos when a Windows Domain Server is available, or NTLM when deployed in a workgroup configuration. The caller provides the service with its Windows credentials (such as a ticket or a token) and the service authenticates that against Windows.

Username and password

The caller provides the service with a username and a password. The service then uses these credentials against some kind of credentials store, such as Windows accounts or a custom credentials store (such as a dedicated database).

X509 certificate

The client identifies itself using a certificate. Typically, that certificate is known in advance to the service. The service looks up the certificate on the host side and validates it, thus authenticating the client. Alternatively, the service may implicitly trust the issuer of the certificate and hence the client presenting it.

Custom mechanism

WCF allows developers to replace the built-in authentication mechanisms with any protocol and credential type, such as using biometrics. These custom solutions are beyond the scope of this book.

Issued token

The caller and the service can both rely on a secure token service to issue the client a token that the service recognizes and trusts. Such a service is typically federated and encapsulates the act of authenticating and securing the call. Azure's Access Control Service (ACS) is an example of such a secure token service. However, federated security and ACS are beyond the scope of this book.

Authorization

Authorization is concerned with what the caller is allowed to do: typically, which operations the client is allowed to invoke on the service. Authorizing of the caller is done under the assumption that the caller is indeed who the caller claims to be—in other words, authorization is meaningless without authentication. For authorization, the service typically relies on some kind of credentials store, where callers are mapped to logical roles. When authorizing an operation, the operation declares or explicitly demands that only certain roles can access it, and the service needs to look up the caller's role or roles from the store and verify that the caller is a member of the

requested roles. Out of the box, WCF supports two credentials stores: the service can use Windows groups (and accounts) for authorization, or it can use an ASP.NET provider (such as the SQL Server provider) to store user accounts and roles. WCF also supports custom role repositories, but I have found that the easiest option by far for implementing a custom store is to implement a custom ASP.NET provider. This chapter will address the ASP.NET providers at length later.



WCF offers an elaborate and extensible infrastructure for authenticating and authorizing the caller based on a set of claims contained in the message. However, discussion of this mechanism is beyond the scope of this book.

Transfer Security

Both authentication and authorization deal with two local aspects of security—if (and to what extent) to grant access to the caller once the service has received the message. In this respect, WCF services are not much different from traditional client/server classes. However, both authentication and authorization are predicated on secure delivery of the message itself. The *transfer* of the message from the client to the service has to be secure, or both authentication and authorization are moot. There are three essential aspects to transfer security, and all three aspects must be enforced to provide for secure services. Message *integrity* deals with how to ensure that the message itself is not tampered with en route from the client to the service. A malicious party or intermediary could, in practice, intercept the message and modify its content; for example, altering the account numbers in the case of a transfer operation in a banking service. Message *privacy* deals with ensuring the confidentiality of the message, so that no third party can even read the contents of the message. Privacy complements integrity. Without it, even if the malicious party does not tamper with the message, that party can still cause harm by gleaning sensitive information (again, such as account numbers) from the message content. Finally, transfer security must provide for *mutual authentication*, which deals with assuring the client that only the proper service is able to read the content of its message—in other words, that the client connects to the correct service. Once the credentials in the message are received, the service must authenticate those credentials locally. The mutual authentication mechanism also needs to detect and eliminate replay attacks and denial-of-service (DOS) attacks. In a replay attack, a malicious party records a valid message from the wire and later sends that valid message back to the service. With a DOS attack, a malicious party floods the service with messages (either valid messages or bogus invalid messages) at such a frequency as to degrade the service's availability.

Transfer Security Modes

WCF supports five different ways of accomplishing the three aspects of transfer security. Choosing the correct transfer security mode is perhaps the prime decision to be made in the context of securing a service. The five transfer security modes are None, Transport security, Message security, Mixed, and Both.

None transfer security mode

As its name implies, the *None* transfer security mode has transfer security completely turned off—in fact, all aspects of WCF security are turned off. No client credentials are provided to the service, and the message itself is wide open to any malicious party to do with it as it pleases. Obviously, setting transfer security to None is highly inadvisable.

Transport transfer security mode

When configured for *Transport* security, WCF uses a secure communication protocol. The available secure transports are HTTPS, TCP, IPC, and MSMQ. Transport security encrypts all communication on the channel and thus provides for integrity, privacy, and mutual authentication. Integrity is provided because without knowing the encryption key, any attempt to modify the message will corrupt it so that it will become useless. Privacy is provided because no party other than the recipient can see the content of the message. Mutual authentication is supported because only the intended recipient of the message can read it; the client need not be concerned with message rerouting to malicious endpoints, as those will not be able to use the message. Once the message is decrypted, the service can read the client's credentials and authenticate the client.

Transport security requires the client and the service to negotiate the details of the encryption, but that is done automatically as part of the communication protocol in the respective binding. Transport security can benefit from hardware acceleration done on the network card so as to avoid burdening the host machine's CPU with the encryption and decryption of the messages. Hardware acceleration obviously caters to high throughput, and it may even make the security overhead unnoticeable. Transport security is the simplest way of achieving transfer security, and the most performant option. Its main downside is that it can only guarantee transfer security point-to-point, meaning when the client connects directly to the service. Having multiple intermediaries between the client and the service renders Transport security questionable, as those intermediaries may not be secure. Consequently, Transport security is typically used only by intranet applications, where you can ensure a single hop between the client and the service in a controlled environment.



When configuring any of the HTTP bindings for Transport security, WCF verifies at the service load time that the corresponding address on the endpoint uses HTTPS rather than mere HTTP.

Message transfer security mode

The *Message* transfer security mode simply encrypts the message itself. By encrypting the message, you gain integrity and privacy and enable mutual authentication, for the same reason that Transport security provides these features when the communication channel is encrypted. However, encrypting the message rather than the transport enables the service to communicate securely over nonsecure transports, such as HTTP. Because of that, Message security provides for end-to-end security, regardless of the number of intermediaries involved in transferring the message and regardless of whether or not the transport is secure. In addition, Message security is based on a set of industry standards designed both for interoperability and for thwarting common attacks such as replay and DOS attacks, and the support WCF offers for it is both rich and extensible. The downside of Message security is that it may introduce call latency due to its inherent overhead. Message security is typically used by Internet applications, where the call patterns are less chatty and the transport is not necessarily secure.

Mixed transfer security mode

The *Mixed* transfer security mode uses Transport security for message integrity and privacy as well as service authentication, and it uses Message security for securing the client's credentials. The Mixed mode tries to combine the advantages of both Transport and Message security by benefiting from the secure transport and even hardware acceleration offered by Transport security to cater to high throughput, and from the extensibility and richer types of client credentials offered by Message security. The downside of the Mixed mode is that it is only secure point-to-point, as a result of the use of Transport security. Application developers rarely need to use the Mixed mode, but it is available for advanced cases.

Both transfer security mode

As its name implies, the *Both* transfer security mode uses both Transport security and Message security. The message itself is secured using Message security, and then it is transferred to the service over a secure transport. The Both mode maximizes security, yet it may be overkill for most applications (with the exception perhaps of disconnected applications, where the additional latency it introduces will go unnoticed).

Transfer Security Mode Configuration

Configuring the transfer security mode is done in the binding, and both the client and the service must use the same transfer security mode and, of course, comply with its requirements. Like any other binding configuration, you can configure transfer security either programmatically or administratively, in a config file. All the common bindings offer a construction parameter indicating the transfer security mode, and all bindings offer a `Security` property with a `Mode` property identifying the configured mode using a dedicated enumeration. As shown in [Table 10-1](#), not all bindings support all transfer security modes: the supported modes are driven by the target scenarios for the binding.

Table 10-1. Bindings and transfer security modes

Name	None	Transport	Message	Mixed	Both
BasicHttpBinding	Yes (default)	Yes	Yes	Yes	No
NetTcpBinding	Yes	Yes (default)	Yes	Yes	No
NetNamedPipeBinding	Yes	Yes (default)	No	No	No
WSHttpBinding	Yes	Yes	Yes (default)	Yes	No
NetMsmqBinding	Yes	Yes (default)	Yes	No	Yes

The intranet bindings (`NetTcpBinding`, `NetNamedPipeBinding`, and `NetMsmqBinding`) all default to Transport security. Thus, no special programming is required on behalf of the service or client developer. The reason is that on the intranet calls are typically point-to-point, and Transport security yields the best performance. However, the intranet bindings can also be configured for the None transfer mode; that is, they can be used on the same transport protocol, only without security. The `NetNamedPipeBinding` supports only None and Transport security—there is no sense in using Message security over IPC, since with IPC there is always exactly one hop from the client to the service. Also note that only the `NetMsmqBinding` supports the Both mode.

The Internet bindings all default to Message security, to enable them to be used over nonsecure transports (that is, HTTP) and to accommodate multiple hops and intermediaries.

With one noticeable exception, all of the WCF bindings are configured with some kind of transfer security and are therefore secure by default. Only the `BasicHttpBinding` defaults to having no security. The reason is that the basic binding is designed to make a WCF service look like a legacy ASMX service, and ASMX is unsecured by default. That said, you can and should configure the `BasicHttpBinding` to use a different transfer security mode, such as Message security.

Specific binding configurations

The `BasicHttpBinding` uses the `BasicHttpSecurityMode` enum for transfer mode configuration. The enum is available via the `Mode` property of the `Security` property of the binding:

```
public enum BasicHttpSecurityMode
{
    None,
    Transport,
    Message,
    TransportWithMessageCredential,
    TransportCredentialOnly
}
public sealed class BasicHttpSecurity
{
    public BasicHttpSecurityMode Mode
    {get;set;}
    //More members
}
public class BasicHttpBinding : Binding, ...
{
    public BasicHttpBinding();
    public BasicHttpBinding(BasicHttpSecurityMode securityMode);
    public BasicHttpSecurity Security
    {get;}
    //More members
}
```

`Security` is of the type `BasicHttpSecurity`. One of the constructors of `BasicHttpBinding` takes the `BasicHttpSecurityMode` enum as a parameter. To secure the basic binding for Message security, you can either construct it secured or set the security mode post-construction. Consequently, in [Example 10-1](#), `binding1` and `binding2` are equivalent.

Example 10-1. Programmatically securing the basic binding

```
BasicHttpBinding binding1 = new BasicHttpBinding(BasicHttpSecurityMode.Message);

BasicHttpBinding binding2 = new BasicHttpBinding();
binding2.Security.Mode = BasicHttpSecurityMode.Message;
```

Instead of programmatic settings, you can use a config file, as in [Example 10-2](#).

Example 10-2. Administratively securing the basic binding

```
<bindings>
    <basicHttpBinding>
        <binding name = "SecuredBasic">
            <security mode = "Message"/>
```

```
</binding>
</basicHttpBinding>
</bindings>
```

The rest of the bindings all use their own enumerations and dedicated security classes, yet they are configured just as in Examples 10-1 and 10-2. For example, the `NetTcpBinding` and the `WSHttpBinding` use the `SecurityMode` enum, defined as:

```
public enum SecurityMode
{
    None,
    Transport,
    Message,
    TransportWithMessageCredential //Mixed
}
```

These bindings offer a matching construction parameter and a matching `Security` property.

The `NetNamedPipeBinding` uses the `NetNamedPipeSecurityMode` enum, which supports only the `None` and `Transport` security modes:

```
public enum NetNamedPipeSecurityMode
{
    None,
    Transport
}
```

The `NetMsmqBinding` uses the `NetMsmqSecurityMode` enum:

```
public enum NetMsmqSecurityMode
{
    None,
    Transport,
    Message,
    Both
}
```

`NetMsmqSecurityMode` is the only enum that offers the `Both` transfer mode.

The reason that almost every common binding has its own dedicated enum for the security mode is that the designers of WCF security opted for increased safety at the expense of overall complexity. They could have defined just a single all-inclusive enum with values corresponding to the five possible transfer security modes, but then it would have been possible at compile time to assign invalid values, such as `Message` security for the `NetNamedPipeBinding`. Opting for specialized enums makes configuring security less error-prone, yet there are more moving parts to come to terms with.

Transport Security and Credentials

WCF lets you select from a number of possible client credential types. For example, the client can identify itself using a classic username and password, or a Windows security token. Windows credentials can then be authenticated using NTLM or Kerberos, when available. Alternatively, the client can use an X509 certificate, or choose to provide no credentials at all and be anonymous. When configuring transfer security for Transport security, however, not all bindings support all client credential types, as shown in [Table 10-2](#).

Table 10-2. Bindings and Transport security client credentials

Name	None	Windows	Username	Certificate
BasicHttpBinding	Yes (default)	Yes	Yes	Yes
NetTcpBinding	Yes	Yes (default)	No	Yes
NetNamedPipeBinding	No	Yes (default)	No	No
WSHttpBinding	Yes	Yes (default)	Yes	Yes
NetMsmqBinding	Yes	Yes (default)	No	Yes

Which types of credentials a binding supports is largely a product of the target scenario for which the binding is designed. For example, all of the intranet bindings default to Windows credentials since they are used in a Windows environment, and the `BasicHttpBinding` defaults to no credentials, just like a classic ASMX web service. The odd default is that of the `WSHttpBinding`, which defaults to Windows credentials to enable the binding to be used over Transport security with minimum effort out of the box.

Message Security and Credentials

When it comes to using Message transfer security, WCF lets applications use the same types of credentials as with Transport security, with the addition of the issued token credential type. Again, when configured for Message security, not all bindings support all client credential types, as shown in [Table 10-3](#).

Table 10-3. Bindings and Message security client credentials

Name	None	Windows	Username	Certificate	Issued token
BasicHttpBinding	No	No	No	Yes	No
NetTcpBinding	Yes	Yes (default)	Yes	Yes	Yes
NetNamedPipeBinding	N/A	N/A	N/A	N/A	N/A
WSHttpBinding	Yes	Yes (default)	Yes	Yes	Yes
NetMsmqBinding	Yes	Yes (default)	Yes	Yes	Yes

While it makes sense that all intranet bindings that support Message security default to Windows credentials, it is interesting to note that the `WSHttpBinding` as Internet binding also defaults to Windows credentials, even though (as discussed later) Internet applications rarely use Windows credentials over HTTP. The reason for this default is to enable developers to securely use the WS binding out of the box, in its correct transfer security mode without resorting first to custom credentials stores.



The `BasicHttpBinding` supports username client credentials for Message security only when configured for Mixed mode. This may be a source of runtime validation errors, since the `BasicHttpMessageCredentialType` enum contains the `BasicHttpMessageCredentialType.UserName` value.

Identity Management

Identity management is the security aspect that deals with which security identity the client sends to the service, and in turn, what the service can do with the client's identity. Not only that, but when designing a service, you need to decide in advance which identity the service will execute under. The service can execute under its own identity; it can impersonate the client's identity (when applicable); or it can use a mixture of identities, alternating in a single operation between its own identity, the client's identity, or even a third identity altogether. Selecting the correct identity has drastic implications for the application's scalability and administration cost. In WCF, when enabled, the security identity flows down the call chain, and each service can find out who its caller is, regardless of the identity of the service.

Overall Policy

To the traditional commonplace security aspects of authentication, authorization, transfer security, and identity management, I would like to add one that is less technical and conventional, but to me just as important: what is your business's approach, or even your personal approach, to security? That is, what is your *security policy*? I believe that in the vast majority of cases, applications simply cannot afford not to be secured. And while security carries with it performance and throughput penalties, these should be of no concern. Simply put, it costs to live. Paying the security penalty is an unavoidable part of designing and administering modern connected applications. Gone are the days when developers could afford not to care about security and deploy applications that relied on the ambient security of the target environment, such as physical security provided by employee access cards or firewalls.

Since most developers cannot afford to become full-time security experts (nor should they), the approach I advocate for overall security policy is simple: crank security all the way up until someone complains. If the resulting application performance and

throughput are still adequate with the maximum security level, leave it at that level. Only if the resulting performance is inadequate should you engage in detailed threat analysis to find out what you can trade in security in exchange for performance. In my experience, you will rarely need to actually go this route; most developers should never need to compromise security this way.

The security strategies described in this chapter follow my overall security policy. WCF's overall approach to security is very much aligned with my own, and I will explicitly point out the few places it is not (and how to rectify it). With the noticeable exception of the `BasicHttpBinding`, WCF is secured by default, and even the `BasicHttpBinding` can easily be secured. All other WCF bindings by default authenticate all callers to the service and rely on transfer security.

Scenario-Driven Approach

Security is by far the most intricate area of WCF. The following list shows the elements that govern security in every WCF operation call:

- Service contract
- Operation contract
- Fault contract
- Service behavior
- Operation behavior
- Host configuration
- Method configuration and code
- Client-side behavior
- Proxy configuration
- Binding configuration

Each of the items in the list may have a dozen or more security-related properties. Obviously, there are an overwhelming number of possible combinations and permutations. In addition, not all combinations are allowed or supported, and not all allowed combinations make sense or are consistent; for example, while technically possible, it does not make sense to use a certificate for client credentials in a homogeneous Windows intranet, much as it makes little sense to use Windows accounts in an Internet application. The solution I chose for this book is to focus on a few key scenarios (and slight variations of them) that address the security needs of the majority of applications today.

The scenarios are:

- Intranet application
- Internet application
- Business-to-business application
- Anonymous application
- No security

I will demonstrate how to make each of these scenarios consistent and secure. In each scenario I will discuss how to support the security aspects of transfer security, authentication, authorization, and identity management. If you need an additional scenario, you can follow my analysis approach to derive the required security aspects and settings.

Intranet Application Scenario

The characteristics of the intranet application are that both the clients and the service use WCF, and that they are deployed in the same intranet. The clients reside behind the firewall, and you can use Windows-based security for transfer security, authentication, and authorization. You can rely on Windows accounts and groups to store the client's credentials. The intranet scenario addresses a wide range of business applications, from finance to manufacturing to in-house IT applications. The intranet scenario is also the richest scenario of all in the options it offers developers for configuring security.

This section on the intranet scenario will define the terminology, techniques, and types used in the other scenarios.

Securing the Intranet Bindings

For the intranet scenario, you should use the intranet bindings: namely, `NetTcpBinding`, `NetNamedPipeBinding`, and `NetMsmqBinding`. You can rely on Transport mode for transfer security because the calls are invariably point-to-point. Conveniently, Transport security is the default transfer mode of the intranet bindings (see [Table 10-1](#)). You can also use the default for the client credentials type, which is Windows (see [Table 10-2](#)). You need to configure this on both the client and the service.

Transport security protection level

Each of the three intranet bindings has a configurable protection level, which is the master switch for Transport protection. The three protection levels are:

None

When configured for this protection level, WCF does not protect the message on transfer from the client to the service. Any malicious party can read the content of the message, or even alter it.

Signed

When configured for this protection level, WCF ensures that the message could have come only from an authenticated sender and that the message integrity was not compromised during transfer. To accomplish this, WCF appends an encrypted checksum to the message. Upon receiving the message, the service calculates the checksum and compares it to the original. If the two do not match, the message is rejected. As a result, the message is impervious to tampering. However, the message content is still visible during the transfer.

Encrypted and Signed

When configured for this protection level, WCF both signs the message and encrypts its content. The Encrypted and Signed protection level provides integrity, privacy, and authenticity.

The Signed protection level offers a clear trade-off between a measured degree of security and performance. However, I consider this to be a trade-off to avoid, and I recommend that you always opt instead for the Encrypted and Signed protection level. WCF represents the protection level with the `ProtectionLevel` enum, defined as:

```
public enum ProtectionLevel
{
    None,
    Sign,
    EncryptAndSign
}
```

Not all Internet bindings default to the same protection level. Both the `NetTcpBinding` and the `NetNamedPipeBinding` default to Encrypted and Signed, yet the `NetMsmqBinding` defaults to Signed.

NetTcpBinding configuration

`NetTcpBinding` takes a construction parameter indicating the desired transfer security mode:

```
public class NetTcpBinding : ...
{
    public NetTcpBinding(SecurityMode securityMode);
    public NetTcpSecurity Security
    {get;}
    //More members
}
```

The `Security` property of the type `NetTcpSecurity` contains the transfer mode (Transport or Message) and two respective properties with their specific settings:

```
public sealed class NetTcpSecurity
{
    public SecurityMode Mode
    {get;set;}
    public MessageSecurityOverTcp Message
    {get;}
    public TcpTransportSecurity Transport
    {get;}
}
```

In the intranet security scenario, you should select Transport security for the transfer security mode and set the values of the `Transport` property of the type `TcpTransportSecurity`:

```
public sealed class TcpTransportSecurity
{
    public TcpClientCredentialType ClientCredentialType
    {get;set;}

    public ProtectionLevel ProtectionLevel
    {get;set;}
}
```

The `Transfer` property should be initialized with the client credential type set to Windows using the `TcpClientCredentialType` enum, defined as:

```
public enum TcpClientCredentialType
{
    None,
    Windows,
    Certificate
}
```

The `Transport` property should also have the protection level set to `ProtectionLevel.EncryptAndSign`. Since both of those settings are the defaults for this binding, these two declarations are equivalent:

```
NetTcpBinding binding1 = new NetTcpBinding();

NetTcpBinding binding2 = new NetTcpBinding(SecurityMode.Transport);
binding2.Security.Transport.ClientCredentialType =
    TcpClientCredentialType.Windows;
binding2.Security.Transport.ProtectionLevel = ProtectionLevel.EncryptAndSign;
```

Alternatively, you can configure the binding using a config file:

```
<bindings>
    <netTcpBinding>
        <binding name = "TCPWindowsSecurity">
            <security mode = "Transport">
```

```

<transport
    clientCredentialType = "Windows"
    protectionLevel = "EncryptAndSign"
/>
</security>
</binding>
</netTcpBinding>
</bindings>

```



The `NetTcpContextBinding` and the `WSHttpContextBinding` also offer the `ContextProtectionLevel` property of the type `ProtectionLevel`, used to indicate the desired protection level for the custom context. `ContextProtectionLevel` defaults to `ProtectionLevel.Sign`. When using Transport security, the value of `ContextProtectionLevel` is ignored (since the transport protects the whole message during transfer). Due to a bug in WCF, it is also ignored when using Message security. The same is true for the `NetNamedPipeContextBinding` defined in [Appendix B](#) (since it uses the same WCF facility).

For the sake of completeness, although it's not required by the intranet scenario, here is how to configure `NetTcpBinding` for Message security with username client credentials:

```

public enum MessageCredentialType
{
    None,
    Windows,
    UserName,
    Certificate,
    IssuedToken
}
public sealed class MessageSecurityOverTcp
{
    public MessageCredentialType ClientCredentialType
    {get;set;}
    //More members
}
NetTcpBinding binding = new NetTcpBinding(SecurityMode.Message);
binding.Security.Message.ClientCredentialType = MessageCredentialType.UserName;

```

`NetTcpSecurity` offers the `Message` property of the type `MessageSecurityOverTcp`. You'll need to set the credentials type using the `MessageCredentialType` enum. Most bindings use the `MessageCredentialType` enum for representing Message security client credentials.

[Figure 10-1](#) shows the security-related elements of the `NetTcpBinding`.

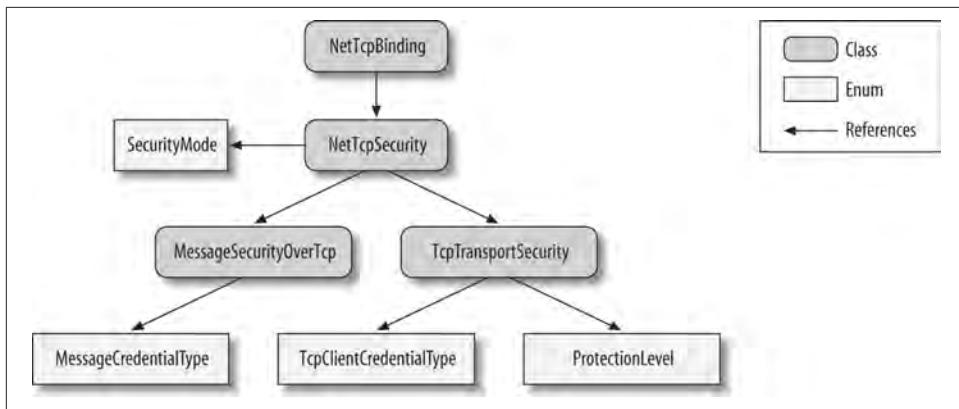


Figure 10-1. NetTcpBinding and security

`NetTcpBinding` has a reference to `NetTcpSecurity`, which uses the `SecurityMode` enum to indicate the transfer security mode. When Transport security is used, `NetTcpSecurity` will use an instance of `TcpTransportSecurity` containing the client credentials type via the `TcpClientCredentialType` enum and the configured protection level via the `ProtectionLevel` enum. When Message security is used, `NetTcpSecurity` will use an instance of `MessageSecurityOverTcp` containing the client credentials type via the `MessageCredentialType` enum.

NetNamedPipeBinding configuration

`NetNamedPipeBinding` takes a construction parameter indicating the desired transfer security mode:

```

public class NetNamedPipeBinding : Binding, ...
{
    public NetNamedPipeBinding(NetNamedPipeSecurityMode securityMode);

    public NetNamedPipeSecurity Security
    {get;}
    //More members
}
  
```

The `Security` property of the type `NetNamedPipeSecurity` contains the transfer mode (Transport or None) and a single property with the specific Transport settings:

```

public sealed class NetNamedPipeSecurity
{
    public NetNamedPipeSecurityMode Mode
    {get;set;}
    public NamedPipeTransportSecurity Transport
    {get;}
}
  
```

For the intranet security scenario, select Transport security for the transfer security mode and set the values of the `Transport` property of the type `NamedPipeTransportSecurity`:

```
public sealed class NamedPipeTransportSecurity
{
    public ProtectionLevel ProtectionLevel
    {get;set;}
}
```

The `Transfer` property should be initialized with the protection level set to `ProtectionLevel.EncryptAndSign`. Because this is the default for the binding, these two declarations are equivalent:

```
NetNamedPipeBinding binding1 = new NetNamedPipeBinding();

NetNamedPipeBinding binding2 = new NetNamedPipeBinding(NetNamedPipeSecurityMode.
                                                       Transport);
binding2.Security.Transport.ProtectionLevel = ProtectionLevel.EncryptAndSign;
```

You can also configure the binding administratively, using a config file:

```
<bindings>
<netNamedPipeBinding>
    <binding name = "IPCWindowsSecurity">
        <security mode = "Transport">
            <transport protectionLevel = "EncryptAndSign"/>
        </security>
    </binding>
</netNamedPipeBinding>
</bindings>
```

There is no need (or option) to set the client credentials type, since only Windows credentials are supported (see [Table 10-2](#)). [Figure 10-2](#) shows the security-related elements of the `NetNamedPipeBinding`.

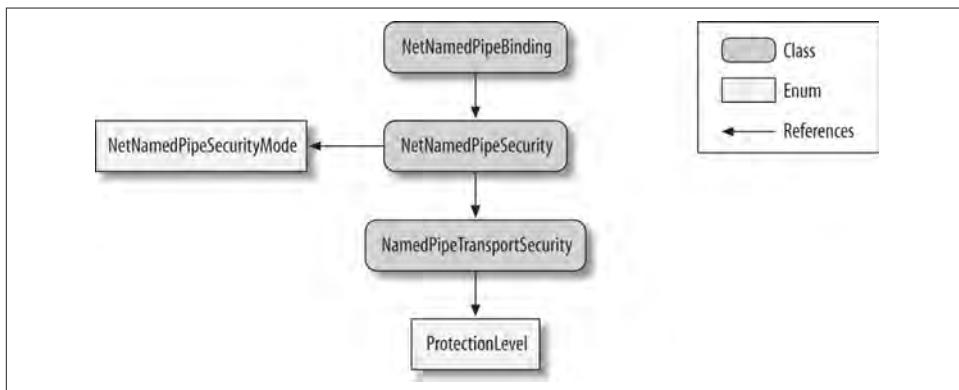


Figure 10-2. NetNamedPipeBinding and security

`NetNamedPipeBinding` has a reference to `NetNamedPipeSecurity`, which uses the `NetNamedPipeSecurityMode` enum to indicate the transfer security mode. When Transport security is used, `NetTcpSecurity` will use an instance of `NamedPipeTransportSecurity` containing the configured protection level via the `ProtectionLevel` enum.

NetMsmqBinding configuration

`NetMsmqBinding` offers a construction parameter for the transfer security mode and a `Security` property:

```
public class NetMsmqBinding : MsmqBindingBase
{
    public NetMsmqBinding(NetMsmqSecurityMode securityMode);
    public NetMsmqSecurity Security
    {get;}
    //More members
}
```

The `Security` property of the type `NetMsmqSecurity` contains the transfer mode (Transport or Message) and two respective properties with their specific settings:

```
public sealed class NetMsmqSecurity
{
    public NetMsmqSecurityMode Mode
    {get;set;}
    public MsmqTransportSecurity Transport
    {get;}
    public MessageSecurityOverMsmq Message
    {get;}
}
```

For the intranet security scenario, select Transport security for the transfer security mode and set the values of the `Transport` property of the type `MsmqTransportSecurity`:

```
public sealed class MsmqTransportSecurity
{
    public MsmqAuthenticationMode MsmqAuthenticationMode
    {get;set;}
    public ProtectionLevel MsmqProtectionLevel
    {get;set;}
    //More members
}
```

The `Transfer` property should be initialized with the client credential type set to Windows domain using the `MsmqAuthenticationMode` enum, defined as:

```
public enum MsmqAuthenticationMode
{
    None,
```

```

        WindowsDomain,
        Certificate
    }
}

```

Windows domain is the default credentials type. In addition, you need to set the protection level to `ProtectionLevel.EncryptAndSign` because the MSMQ binding defaults to `ProtectionLevel.Signed`. The following two definitions are equivalent:

```

NetMsmqBinding binding1 = new NetMsmqBinding();
binding1.Security.Transport.MsmqProtectionLevel = ProtectionLevel.EncryptAndSign;

NetMsmqBinding binding2 = new NetMsmqBinding();
binding2.Security.Mode = NetMsmqSecurityMode.Transport;
binding2.Security.Transport.MsmqAuthenticationMode =
    MsmqAuthenticationMode.WindowsDomain;
binding2.Security.Transport.MsmqProtectionLevel = ProtectionLevel.EncryptAndSign;

```

Alternatively, you can configure the binding using a config file:

```

<bindings>
    <netMsmqBinding>
        <binding name = "MSMQWindowsSecurity">
            <security mode = "Transport">
                <transport
                    msmqAuthenticationMode = "WindowsDomain"
                    msmqProtectionLevel = "EncryptAndSign"
                />
            </security>
        </binding>
    </netMsmqBinding>
</bindings>

```

Figure 10-3 shows the security-related elements of the `NetMsmqBinding`.

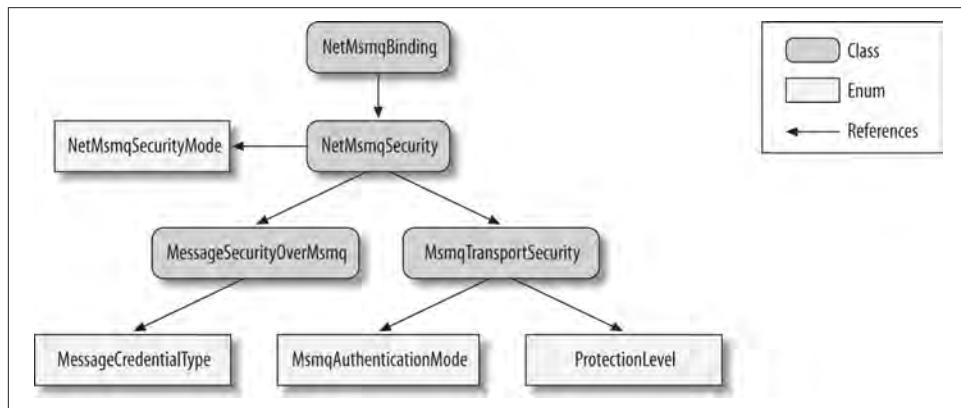


Figure 10-3. `NetMsmqBinding` and security

`NetMsmqBinding` has a reference to `NetMsmqSecurity`, which uses the `NetMsmqSecurityMode` enum to indicate the transfer security mode. When Transport security is used, `NetMsmqSecurity` will use an instance of `MsmqTransportSecurity` containing the client credentials type via the `MsmqAuthenticationMode` enum, and the configured protection level via the `ProtectionLevel` enum. There are similar references to types controlling Message security.

Constraining Message Protection

While a service should ideally use the highest possible level of security, it is actually at the mercy of its host, because the host is the one configuring the binding. This is especially problematic if the service is to be deployed in an unknown environment with an arbitrary host. To compensate, WCF lets service developers insist on a protection level, or rather, constrain the minimum protection level at which their service is willing to operate. Both the service and the client can constrain the protection level, independently of each other. You can constrain the protection level in three places. When constrained at the service contract, all operations on the contract are considered sensitive and protected. When constrained at the operation contract, only that operation is protected; other operations on the same contract are not. Finally, you can constrain the protection level for an individual fault contract. This can be required because sometimes the error information returned to the client is sensitive, containing parameter values, exception messages, and the call stack. The respective contract attributes offer the `ProtectionLevel` property of the enum type `ProtectionLevel`:

```
[AttributeUsage(AttributeTargets.Interface|AttributeTargets.Class,
               Inherited = false)]
public sealed class ServiceContractAttribute : Attribute
{
    public ProtectionLevel ProtectionLevel
    {get;set;}
    //More members
}
[AttributeUsage(AttributeTargets.Method)]
public sealed class OperationContractAttribute : Attribute
{
    public ProtectionLevel ProtectionLevel
    {get;set;}
    //More members
}
[AttributeUsage(AttributeTargets.Method,AllowMultiple = true,
               Inherited = false)]
public sealed class FaultContractAttribute : Attribute
{
    public ProtectionLevel ProtectionLevel
    {get;set;}
    //More members
}
```

As an example, here is how to set the protection level on a service contract:

```
[ServiceContract(ProtectionLevel = ProtectionLevel.EncryptAndSign)]
interface IMyContract
{...}
```

Setting the `ProtectionLevel` property on the contract attributes merely indicates the low-water mark; that is, the minimum protection level accepted by this contract. If the binding is configured for a lower protection level, it will result in an `InvalidOperationException` at the service load time or the time the proxy is opened. If the binding is configured for a higher level, the contract will use that level. The `ProtectionLevel` property on the contract attributes defaults to `ProtectionLevel.None`, meaning it has no effect.

The desired protection constraint is considered a local implementation detail of the service, so the required protection level is not exported with the service metadata. Consequently, the client may require a different level and enforce it separately from the service.



Even though the Internet bindings do not offer a protection level property, the protection level constraint at the service-, operation-, or fault-contract level is satisfied when using Transport or Message security. The constraint is not satisfied when security is turned off by using the `None` security mode.

Authentication

By default, when a client calls a proxy that targets an endpoint whose binding is configured for using Windows credentials with Transport security, there is nothing explicit the client needs to do to pass its credentials. WCF will automatically pass the Windows identity of the client's process to the service:

```
class MyContractClient : ClientBase<IMyContract>, IMyContract
{...}

MyContractClient proxy = new MyContractClient();
proxy.MyMethod(); //Client identity passed here
proxy.Close();
```

When the service receives the call, WCF will authenticate the caller on the service side. If the client's credentials represent a valid Windows account, the caller will be allowed to access the requested operation on the service.

Providing alternative Windows credentials

Instead of using the identity of the process in which it happens to be running, the client can pass alternative Windows credentials. The `ClientBase<T>` base class offers the `ClientCredentials` property of the type `ClientCredentials`:

```
public abstract class ClientBase<T> : ...
{
    public ClientCredentials ClientCredentials
    {get;}
}
public class ClientCredentials : ..., IEndpointBehavior
{
    public WindowsClientCredential Windows
    {get;}
    //More members
}
```

`ClientCredentials` contains the property `Windows` of the type `WindowsClientCredential`, defined as:

```
public sealed class WindowsClientCredential
{
    public NetworkCredential ClientCredential
    {get;set;}
    //More members
}
```

`WindowsClientCredential` has the property `ClientCredential` of the type `NetworkCredential`, which is where the client needs to set the alternative credentials:

```
public class NetworkCredential : ...
{
    public NetworkCredential();
    public NetworkCredential(string userName, string password);
    public NetworkCredential(string userName, string password, string domain);

    public string Domain
    {get;set;}
    public string UserName
    {get;set;}
    public string Password
    {get;set;}
}
```

Example 10-3 demonstrates how to use these classes and properties to provide alternative Windows credentials.

Example 10-3. Providing alternative Windows credentials

```
MyContractClient proxy = new MyContractClient();
proxy.ClientCredentials.Windows.ClientCredential.Domain = "MyDomain";
```

```
proxy.ClientCredentials.Windows.ClientCredential.UserName = "MyUsername";
proxy.ClientCredentials.Windows.ClientCredential.Password = "MyPassword";

proxy.MyMethod();
proxy.Close();
```

Once you specify an alternative identity and open the proxy, the proxy cannot use any other identity later.



If you do try specifying alternative credentials after opening the proxy, those credentials will be silently ignored.

Clients can use the technique demonstrated in [Example 10-3](#) when the credentials provided are collected dynamically at runtime, perhaps using a login dialog box.

When working with a channel factory instead of a proxy class, the `ChannelFactory` base class offers the `Credentials` property of the type `ClientCredentials`:

```
public abstract class ChannelFactory : ...
{
    public ClientCredentials Credentials
    {get;}
    //More members
}
public class ChannelFactory<T> : ChannelFactory, ...
{
    public T CreateChannel();
    //More members
}
```

In this case, simply set the alternative credentials in the `Credentials` property, as was done in [Example 10-3](#):

```
ChannelFactory<IMyContract> factory = new ChannelFactory<IMyContract>(...);

factory.Credentials.Windows.ClientCredential.Domain    = "MyDomain";
factory.Credentials.Windows.ClientCredential.UserName = "MyUsername";
factory.Credentials.Windows.ClientCredential.Password = "MyPassword";

IMyContract proxy = factory.CreateChannel();
```

Note that you cannot use the static `CreateChannel()` methods of `ChannelFactory<T>`, since you have to first instantiate a factory in order to access the `Credentials` property.

Identities

All Windows processes run with an authenticated security identity, and the process hosting a WCF service is no different. The identity is actually a Windows account whose security token is attached to the process (and to every thread in the process). However, it is up to the application administrator to decide which identity to use. One option is to have the host run with an *interactive* user identity; that is, the identity of the user who launched the host process. An interactive identity is typically used when self-hosting and is ideal for debugging, because the debugger will automatically attach itself to the host process when launched from within Visual Studio. However, relying on an interactive identity is impractical for deployment on a server machine, where there will not necessarily be a logged-on user, and if there is a logged-on user that user may not have the necessary credentials to perform the requested work. For production deployment, you typically rely on a *designated account*, which is a preset Windows account used primarily by your service or services. To launch the service under a designated account, you can use the “Run as” shell option. However, “Run as” is useful only for simple testing. You can also have an Windows service as your host and use the Control Panel Services applet to assign a designated identity to the host. If you’re hosting in IIS or the WAS, you can use those environments’ configuration tools to assign a designated identity to the process from the pool.

The `IIdentity` interface

In .NET, the `IIdentity` interface (from the `System.Security.Principal` namespace) represents a security identity:

```
public interface IIdentity
{
    string AuthenticationType
    {get;}
    bool IsAuthenticated
    {get;}
    string Name
    {get;}
}
```

The interface lets you know whether the identity behind the interface is authenticated (and, if so, which authentication mechanism was used) and allows you to obtain the name of the identity. Out of the box, WCF takes advantage of three implementations of `IIdentity` offered by .NET: `WindowsIdentity`, `GenericIdentity`, and `X509Identity`. The `WindowsIdentity` class represents a Windows account. The `GenericIdentity` class is a general-purpose class whose main use is to wrap an identity name with an `IIdentity`. With both `GenericIdentity` and `WindowsIdentity`, if the identity name is an empty string, that identity is considered unauthenticated, and any other non-zero-length name is considered authenticated. Finally, `X509Identity` is an inter-

nal class that represents an identity that was authenticated using an X509 certificate. The identity behind an `X509Identity` is always authenticated.

Working with `WindowsIdentity`

The `WindowsIdentity` class offers a few useful methods above and beyond the mere implementation of `IIdentity`:

```
public class WindowsIdentity : IIdentity, ...
{
    public WindowsIdentity(string sUserPrincipalName);
    public static WindowsIdentity GetAnonymous();
    public static WindowsIdentity GetCurrent();
    public virtual bool IsAnonymous
    {get;}
    public virtual bool IsAuthenticated
    {get;}
    public virtual string Name
    {get;}
    //More members
}
```

The `IsAnonymous` Boolean property indicates whether the underlying identity is anonymous and the `GetAnonymous()` method returns an anonymous Windows identity, typically used for impersonation to mask the real identity:

```
WindowsIdentity identity = WindowsIdentity.GetAnonymous();
Debug.Assert(identity.Name == "");
Debug.Assert(identity.IsAuthenticated == false);
Debug.Assert(identity.IsAnonymous == true);
```

The `GetCurrent()` static method returns the identity of the process where it is called. That identity is always non-anonymous and authenticated:

```
WindowsIdentity currentIdentity = WindowsIdentity.GetCurrent();
Debug.Assert(currentIdentity.Name != "");
Debug.Assert(currentIdentity.IsAuthenticated == true);
Debug.Assert(currentIdentity.IsAnonymous == false);
```

The Security Call Context

Every operation on a secured WCF service has a security call context. The security call context is represented by the class `ServiceSecurityContext`, defined as:

```
public class ServiceSecurityContext
{
    public static ServiceSecurityContext Current
    {get;}
    public bool IsAnonymous
    {get;}
    public IIdentity PrimaryIdentity
    {get;}
```

```
public WindowsIdentity WindowsIdentity  
{get;}  
//More members  
}
```

The main use for the security call context is for custom security mechanisms, as well as analysis and auditing. While it is presented here in the context of the intranet scenario, all other secured scenarios have use for the security call context as well.

Note that in spite of its name, this is the security context of the call, not the service. Two operations on the same service can definitely have different security call contexts.

The security call context is stored in the TLS, so every method on every object down the call chain from the service can access the security call context, including your service constructor. To obtain your current security call context, simply access the `Current` static property. Another way of accessing the security call context is via the `ServiceSecurityContext` property of the `OperationContext`:

```
public sealed class OperationContext : ...  
{  
    public ServiceSecurityContext ServiceSecurityContext  
    {get;}  
    //More members  
}
```

Regardless of which mechanism you use, you will get the same object:

```
ServiceSecurityContext context1 = ServiceSecurityContext.Current;  
ServiceSecurityContext context2 =  
    OperationContext.Current.ServiceSecurityContext;  
Debug.Assert(context1 == context2);
```



Your service has a security call context only if security is enabled. When security is disabled, `ServiceSecurityContext.Current` returns null.

The `PrimaryIdentity` property of `ServiceSecurityContext` contains the identity of the immediate client up the call chain. If the client is unauthenticated, `PrimaryIdentity` will reference an implementation of `IIdentity` with a blank identity. When Windows authentication is used, the `PrimaryIdentity` property will be set to an instance of `WindowsIdentity`.

The `WindowsIdentity` property is meaningful only when using Windows authentication, and it will always be of the type `WindowsIdentity`. When valid Windows credentials are provided, the `WindowsIdentity` property will contain the corresponding client identity and will match the value of `PrimaryIdentity`.



The constructor of a singleton service does not have a security call context, since it is called when the host is launched, not as a result of a client call.

Impersonation

Some resources, such as the file system, SQL Server, sockets, and even DCOM objects, grant access to themselves based on the caller's security token. Typically, the host process is assigned an identity with elevated permissions that are required to access such resources, so that it can function properly. Clients, however, typically have restricted credentials compared with those of the service. Legacy technologies such as unmanaged Visual Basic or C++ did not offer role-based security support, so developers used *impersonation* to address this credentials gap. Impersonation lets the service assume the client's identity, primarily in order to verify whether the client is authorized to perform the work it's asking the service to do. Impersonation has a number of key detrimental effects on your application, which will be discussed at the end of this section. Instead of impersonation, you should apply role-based security to authorize the callers, coupled with a trusted subsystem pattern across layers. That said, many developers are accustomed to designing systems using impersonation, so both .NET and WCF support this technique.

Manual impersonation

The service can impersonate its calling client by calling the `Impersonate()` method of the `WindowsIdentity` class:

```
public class WindowsIdentity : IIdentity, ...
{
    public virtual WindowsImpersonationContext Impersonate();
    //More members
}
public class WindowsImpersonationContext : IDisposable
{
    public void Dispose();
    public void Undo();
}
```

`Impersonate()` returns an instance of `WindowsImpersonationContext` containing the service's previous identity. To revert back to that identity, the service calls the `Undo()` method. To impersonate a client, the service needs to call `Impersonate()` on the identity of the caller, which is available via the `WindowsIdentity` property of its security call context, as shown in [Example 10-4](#).

Example 10-4. Explicit impersonation and reversion

```
class MyService : IMyContract
{
    public void MyMethod()
    {
        WindowsImpersonationContext impersonationContext =
            ServiceSecurityContext.Current.WindowsIdentity.Impersonate();
        try
        {
            /* Do work as client */
        }
        finally
        {
            impersonationContext.Undo();
        }
    }
}
```

Note in [Example 10-4](#) that the call to `Undo()` is in the `finally` statement, so the service will revert to its old identity even if exceptions occur. To somewhat simplify reverting, the `WindowsImpersonationContext` implementation of `Dispose()` also reverts, which enables you to use it in a `using` statement:

```
public void MyMethod()
{
    using(ServiceSecurityContext.Current.WindowsIdentity.Impersonate())
    {
        /* Do work as client */
    }
}
```

Declarative impersonation

Instead of impersonating manually, you can instruct WCF to automatically impersonate the caller of the method. The `OperationBehavior` attribute offers the `Impersonation` property of the enum type `ImpersonationOption`:

```
public enum ImpersonationOption
{
    NotAllowed,
    Allowed,
    Required
}
[AttributeUsage(AttributeTargets.Method)]
public sealed class OperationBehaviorAttribute : Attribute, IOperationBehavior
{
    public ImpersonationOption Impersonation
    {get;set;}
    //More members
}
```

The default value is `ImpersonationOption.NotAllowed`. This value indicates that WCF should not auto-impersonate, but you can write code (as in [Example 10-4](#)) that explicitly impersonates.

`ImpersonationOption.Allowed` instructs WCF to automatically impersonate the caller whenever Windows authentication is used, but it has no effect with other authentication mechanisms. When WCF auto-impersonates, it will also auto-revert to the previous service identity once the method returns.

The `ImpersonationOption.Required` value mandates the use of Windows authentication and will throw an exception if any other authentication mechanism is used. As its name implies, with this setting WCF will always auto-impersonate (and revert) in every call to the operation:

```
class MyService : IMyContract
{
    [OperationBehavior(Impersonation = ImpersonationOption.Required)]
    public void MyMethod()
    {
        /* Do work as client */
    }
}
```

Note that there is no way to use declarative impersonation with the service constructor because you cannot apply the `OperationBehavior` attribute on a constructor. Constructors can only use manual impersonation. If you do impersonate in the constructor, always revert as well in the constructor, to avoid side effects on the operations of the service (and even other services in the same host).

Impersonating all operations

In the event that you need to enable impersonation in all the service operations, the `ServiceHostBase` class has the `Authorization` property of the type `ServiceAuthorizationBehavior`:

```
public abstract class ServiceHostBase : ...
{
    public ServiceAuthorizationBehavior Authorization
    {get;}
    //More members
}
public sealed class ServiceAuthorizationBehavior : IServiceBehavior
{
    public bool ImpersonateCallerForAllOperations
    {get;set;}
    //More members
}
```

`ServiceAuthorizationBehavior` provides the Boolean property `ImpersonateCallerForAllOperations`, which is `false` by default. Contrary to what its name implies, when set to `true`, this property merely verifies that the service does not have any operations configured with `ImpersonationOption.NotAllowed`. This constraint is verified at service load time, yielding an `InvalidOperationException` when violated.

In effect, when Windows authentication is used, this will amount to the service automatically impersonating the client in all operations, but all the operations must be explicitly decorated with `ImpersonationOption.Allowed` or `ImpersonationOption.Required`. `ImpersonateCallerForAllOperations` has no effect on constructors.

You can set the `ImpersonateCallerForAllOperations` property programmatically or in the config file. If you set it programmatically, you can do so only before opening the host:

```
ServiceHost host = new ServiceHost(typeof(MyService));
host.Authorization.ImpersonateCallerForAllOperations = true;
host.Open();
```

If you set it using a config file, you need to reference the matching service behavior in the service declaration:

```
<services>
    <service name = "MyService" behaviorConfiguration= "ImpersonateAll">
        ...
    </service>
</services>
<behaviors>
    <serviceBehaviors>
        <behavior name = "ImpersonateAll">
            <serviceAuthorization impersonateCallerForAllOperations = "true"/>
        </behavior>
    </serviceBehaviors>
</behaviors>
```

To automate impersonating in all operations without the need to apply the `OperationBehavior` attribute on every method, I wrote the `SecurityHelper` static class, with the `ImpersonateAll()` extension methods:

```
public static class SecurityHelper
{
    public static void ImpersonateAll(this ServiceHostBase host);
    public static void ImpersonateAll(this ServiceDescription description);
    //More members
}
```

The extension methods work on both `ServiceHost` and `ServiceHost<T>`.

You can only call `ImpersonateAll()` before opening the host:

```

//Will impersonate in all operations
class MyService : IMyContract
{
    public void MyMethod()
    {...}
}
ServiceHost host = new ServiceHost(typeof(MyService));
host.ImpersonateAll();
host.Open();

```

Example 10-5 shows the implementation of `ImpersonateAll()`.

Example 10-5. Implementing SecurityHelper.ImpersonateAll()

```

public static class SecurityHelper
{
    public static void ImpersonateAll(this ServiceHostBase host)
    {
        host.Authorization.ImpersonateCallerForAllOperations = true;
        host.Description.ImpersonateAll();
    }
    public static void ImpersonateAll(this ServiceDescription description)
    {
        foreach(ServiceEndpoint endpoint in description.Endpoints)
        {
            if(endpoint.Contract.Name == "IMetadataExchange")
            {
                continue;
            }
            foreach(OperationDescription operation in endpoint.Contract.Operations)
            {
                OperationBehaviorAttribute attribute = operation.Behaviors.
                    Find<OperationBehaviorAttribute>();
                attribute.Impersonation = ImpersonationOption.Required;
            }
        }
    }
    //More members
}

```

In **Example 10-5**, `ImpersonateAll()` (for the sake of good manners) first sets the `ImpersonateCallerForAllOperations` property of the provided host to `true`, then obtains the service description from the host and calls the other overloaded extension method of `ServiceDescription`. This version explicitly configures all operations with `ImpersonationOption.Required`, by iterating over the endpoints collection of the service description. For each endpoint (except the metadata exchange endpoints), `ImpersonateAll()` accesses the operations collection of the contract. For each operation, there is always exactly one `OperationBehaviorAttribute` in the collection of

operation behaviors, even if you did not provide one explicitly. The method then simply sets the `Impersonation` property to `ImpersonationOption.Required`.

Restricting impersonation

Authorization and authentication protect the service from being accessed by unauthorized, unauthenticated, potentially malicious clients. However, how should the client be protected from malicious services? One of the ways an adversarial service could abuse the client is by assuming the client's identity and credentials and causing harm while masquerading as the client. This tactic enables the malicious service both to leave an identity trail pointing back to the client and to elevate its own potentially demoted, less-privileged credentials to the client's level.

In some cases, the client may not want to allow the service to obtain its identity at all. WCF therefore lets the client indicate the degree to which the service can obtain the client's identity and how it can use it. Impersonation is actually a range of options indicating the level of trust between the client and the service. The `WindowsClientCredential` class provides the `AllowedImpersonationLevel` enum of the type `TokenImpersonationLevel`, found in the `System.Security.Principal` namespace:

```
public enum TokenImpersonationLevel
{
    None,
    Anonymous,
    Identification,
    Impersonation,
    Delegation
}
public sealed class WindowsClientCredential
{
    public TokenImpersonationLevel AllowedImpersonationLevel
    {get;set;}
    //More members
}
```

The client can use `AllowedImpersonationLevel` to restrict the allowed impersonation level both programmatically and administratively. For example, to programmatically restrict the impersonation level to `TokenImpersonationLevel.Identification`, before opening the proxy the client would write:

```
MyContractClient proxy = new MyContractClient();
proxy.ClientCredentials.Windows.AllowedImpersonationLevel =
    TokenImpersonationLevel.
    Identification;
proxy.MyMethod();
proxy.Close();
```

When using a config file, the administrator should define the allowed impersonation level as a custom endpoint behavior and reference it from the relevant endpoint section:

```
<client>
  <endpoint behaviorConfiguration = "ImpersonationBehavior"
    ...
  />
</client>
<behaviors>
  <endpointBehaviors>
    <behavior name = "ImpersonationBehavior">
      <clientCredentials>
        <windows allowedImpersonationLevel = "Identification"/>
      </clientCredentials>
    </behavior>
  </endpointBehaviors>
</behaviors>
```

`TokenImpersonationLevel.None` simply means that no impersonation level is assigned, so the client provides no identity information. This setting therefore amounts to the same thing as `TokenImpersonationLevel.Anonymous`, where the client provides no credentials at all. These two values are, of course, the safest from the client's perspective, but they are the least useful options from the application's perspective, since the service cannot perform any authentication or authorization. Not sharing credentials is possible only if the service is configured for anonymous access or for having no security, which is not the case with the intranet scenario. If the service is configured for Windows security, these two values yield an `ArgumentOutOfRangeException` on the client side.

With `TokenImpersonationLevel.Identification`, the service can identify the client (i.e., obtain the security identity of the calling client). The service, however, is not allowed to impersonate the client—everything the service does must be done under the service's own identity. Trying to impersonate will throw an `ArgumentOutOfRangeException` on the service side. Note, however, that if the service and the client are on the same machine, the service will still be able to impersonate the client, even when `TokenImpersonationLevel.Identification` is used. `TokenImpersonationLevel.Identification` is the default value used with Windows security and is the recommended value for the intranet scenario.

`TokenImpersonationLevel.Impersonation` grants the service permission both to obtain the client's identity and to impersonate the client. Impersonation indicates a great deal of trust between the client and the service, since the service can do anything the client can do, even if the service host is configured to use a less privileged identity. The only difference between the real client and the impersonating service is that if the service is on a separate machine from the client, it cannot access resources or objects on other machines as the client, because the service machine does not

really have the client's password. In the case where the service and the client are on the same machine, the service impersonating the client can make one network hop to another machine, since the machine it resides on can still authenticate the impersonated client identity.

Finally, `TokenImpersonationLevel.Delegation` provides the service with the client's Kerberos ticket. In this case, the service can freely access resources on any machine as the client. If service is also configured for delegation, when it calls other downstream services the client's identity could be propagated further and further down the call chain. Delegation-required Kerberos authentication is not possible on Windows workgroup installations. Both the client and server user accounts must be properly configured in Active Directory to support delegation, due to the enormous trust (and hence security risk) involved. Delegation uses by default another security service called *cloaking*, which propagates the caller's identity along the call chain.

Delegation is extremely dangerous from the client's perspective, since the client has no control over who ends up using its identity, or where. When the impersonation level is set to `TokenImpersonationLevel.Impersonation`, the client takes a calculated risk: it knows which services it is accessing, and if those services are on a different machine, the client identity cannot propagate across the network. I consider delegation something that enables the service not just to impersonate the client, but to act as an imposter; security-wise, as far as the client is concerned, this is tantamount to waiving security.

Avoiding impersonation

You should design your services so that they do not rely on impersonation, and your clients should use `TokenImpersonationLevel.Identification`. Impersonation is a relic of the '90s, typically used in classic two-tier systems in the absence of role-based security support, where scalability was not a concern and managing a small number of identities across resources was doable.

As a general design guideline, the further down the call chain from the client, the less relevant the client's identity is. If you use some kind of layered approach in your system design, each layer should run under its own identity, authenticate its immediate callers, and implicitly trust its calling layer to authenticate its callers, thereby maintaining a chain of trusted, authenticated callers. This is called the *trusted subsystem* pattern. Impersonation, on the other hand, requires you to keep propagating the identity further and further down the call chain, all the way to the underlying resources. Doing so impedes scalability, because many resources (such as SQL Server connections) are allocated per identity. With impersonation, you will need as many resources as clients, and you will not be able to benefit from resource pooling (such as connection pooling). Impersonation also complicates resource administration, because you need to grant access to the resources to all of the original client identities,

and there could be numerous such identities to manage. A service that always runs under its own identity poses no such problems, regardless of how many identities access that service. To control access to the resources, you should use authorization, as discussed next.

Multitier systems that do use impersonation typically gravitate toward delegation, since that is the only way to propagate the client identities across tiers and machines. In fact, the main reason developers today use impersonation has little to do with resource access authorization (which can easily be accomplished with role-based security); instead, it is used as a mechanism for auditing and identity propagation. If the application is required to provide at lower layers the identity of the topmost client or all clients up the chain, impersonation (if not full-fledged delegation) may look like a viable option. There are three good solutions for these requirements. First, if the business use cases require you to provide the top-level identity to downstream parties, there is nothing wrong with providing it as explicit method arguments since they are part of the required behavior of the system. The second solution is to use security audits (discussed later) and leave a trail across the call chain. At any point, you can reconstruct that chain of identities from the local audits. The third option is to propagate the identity of the original caller (or the entire stack of callers) in the message headers. Doing that transparently across the call chain requires passing the identities out-of-band in the headers and using the elegant generic interception technique described in [Appendix E](#). *ServiceModelEx* contains those helper classes (look for `SecurityCallStackClientBase<T>`, `OperationSecurityCallStackAttribute`, and `SecurityCallStackBehaviorAttribute`).

Finally, relying on impersonation precludes non-Windows authentication mechanisms. If you do decide to use impersonation, use it judiciously and only as a last resort, when there is no other, better design approach.



Impersonation is not possible with queued services.

Authorization

While authentication deals with verifying that the client is indeed who the client claims to be, most applications also need to verify that the client (or more precisely, the identity it presents) has permission to perform the operation. Since it would be impractical to program access permissions for each individual identity, it is better to grant permissions to the roles clients play in the application domain. A *role* is a symbolic category of identities that share the same security privileges. When you assign a role to an application resource, you are granting access to that resource to anyone

who is a member of that role. Discovering the roles clients play in your business domain is part of your application-requirements analysis and design, just like factoring services and interfaces. By interacting with roles instead of particular identities, you isolate your application from changes made in real life, such as adding new users, moving existing users between positions, promoting users, or users leaving their jobs. .NET allows you to apply role-based security both declaratively and programmatically, if the need to verify role membership is based on a dynamic decision.

The security principal

For security purposes, it is convenient to lump together an identity and the information about its role membership. This representation is called the *security principal*.

The principal in .NET is any object that implements the `IPrincipal` interface, defined in the `System.Security.Principal` namespace:

```
public interface IPrincipal
{
    IIdentity Identity
    {get;}
    bool IsInRole(string role);
}
```

The `IsInRole()` method simply returns `true` if the identity associated with this principal is a member of the specified role, and `false` otherwise. The `Identity` read-only property provides access to read-only information about the identity, in the form of an object implementing the `IIdentity` interface. Out of the box, .NET offers several implementations of `IPrincipal`. `GenericPrincipal` is a general-purpose principal that has to be preconfigured with the role information. It is typically used when no authorization is required, in which case `GenericPrincipal` wraps a blank identity. The `WindowsPrincipal` class looks up role membership information inside Windows groups.

Every .NET thread has a principal object associated with it, obtained via the `CurrentPrincipal` static property of the `Thread` class:

```
public sealed class Thread
{
    public static IPrincipal CurrentPrincipal
    {get;set;}
    //More members
}
```

For example, here is how to discover the username as well as whether or not the caller was authenticated:

```
IPrincipal principal = Thread.CurrentPrincipal;
string userName = principal.Identity.Name;
bool isAuthenticated = principal.Identity.isAuthenticated;
```

Selecting an authorization mode

As presented earlier, the `ServiceHostBase` class provides the `Authorization` property of the type `ServiceAuthorizationBehavior`. `ServiceAuthorizationBehavior` has the `PrincipalPermissionMode` property of the enum type `PrincipalPermissionMode`, defined as:

```
public enum PrincipalPermissionMode
{
    None,
    UseWindowsGroups,
    UseAspNetRoles,
    Custom
}
public sealed class ServiceAuthorizationBehavior : IServiceBehavior
{
    public PrincipalPermissionMode PrincipalPermissionMode
    {get;set;}
    //More members
}
```

Before opening the host, you can use the `PrincipalPermissionMode` property to select the principal mode; that is, which type of principal to install to authorize the caller.

If `PrincipalPermissionMode` is set to `PrincipalPermissionMode.None`, principal-based authorization is impossible. After authenticating the caller (if authentication is required at all), WCF installs `GenericPrincipal` with a blank identity and attaches it to the thread that invokes the service operation. That principal will be available via `Thread.CurrentPrincipal`.

When `PrincipalPermissionMode` is set to `PrincipalPermissionMode.UseWindowsGroups`, WCF installs a `WindowsPrincipal` with an identity matching the provided credentials. If no Windows authentication took place (because the service did not require it), WCF will install a `WindowsPrincipal` with a blank identity.

`PrincipalPermissionMode.UseWindowsGroups` is the default value of the `PrincipalPermissionMode` property, so these two definitions are equivalent:

```
ServiceHost host1 = new ServiceHost(typeof(MyService));

ServiceHost host2 = new ServiceHost(typeof(MyService));
host2.Authorization.PrincipalPermissionMode = PrincipalPermissionMode.
    UseWindowsGroups;
```

When using a config file, you need to reference a custom behavior section assigning the principal mode:

```
<services>
    <service name = "MyService" behaviorConfiguration = "WindowsGroups">
```

```

    ...
  </service>
</services>
<behaviors>
  <serviceBehaviors>
    <behavior name = "WindowsGroups">
      <serviceAuthorization principalPermissionMode = "UseWindowsGroups"/>
    </behavior>
  </serviceBehaviors>
</behaviors>

```

Declarative role-based security

You apply service-side declarative role-based security using the attribute `PrincipalPermissionAttribute`, defined in the `System.Security.Permissions` namespace:

```

public enum SecurityAction
{
  Demand,
  //More members
}

[AttributeUsage(AttributeTargets.Class | AttributeTargets.Method)]
public sealed class PrincipalPermissionAttribute : CodeAccessSecurityAttribute
{
  public PrincipalPermissionAttribute(SecurityAction action);

  public bool Authenticated
  {get;set; }
  public string Name
  {get;set;}
  public string Role
  {get;set;}
  //More members
}

```

The `PrincipalPermission` attribute lets you declare the required role membership. For the intranet scenario, when you specify a Windows group as a role, you don't have to prefix the role name with your domain or machine name (if you wish to authorize against its roles). You can also explicitly specify another domain, if you have a trust relationship with it.

In [Example 10-6](#), the declaration of the `PrincipalPermission` attribute grants access to `MyMethod()` only to callers whose identities belong to the Managers group.

Example 10-6. Declarative role-based security on the intranet

```

[ServiceContract]
interface IMyContract
{
  [OperationContract]

```

```

    void MyMethod();
}
class MyService : IMyContract
{
    [PrincipalPermission(SecurityAction.Demand,Role = "Manager")]
    public void MyMethod()
    {...}
}

```

If the caller is not a member of that role, .NET throws an exception of type `SecurityException`.



When experimenting with Windows role-based security, you often add users to or remove users from user groups. Because Windows caches user-group information at login time, the changes you make are not reflected until the next login.

If multiple roles are allowed to access the method, you can apply the attribute multiple times:

```

[PrincipalPermission(SecurityAction.Demand,Role = "Manager")]
[PrincipalPermission(SecurityAction.Demand,Role = "Customer")]
public void MyMethod()
{...}

```

When multiple `PrincipalPermission` attributes are used, .NET verifies that the caller is a member of at least one of the demanded roles. If you want to verify that the caller is a member of both roles, you need to use programmatic role membership checks, discussed later.

While the `PrincipalPermission` attribute by its very definition can be applied on methods and classes, in a WCF service class you can apply it only on methods. The reason is that in WCF, unlike with normal classes, the service class constructor always executes under a `GenericPrincipal` with a blank identity, regardless of the authentication mechanisms used. As a result, the identity under which the constructor is running is unauthenticated and will always fail any kind of authorization attempt (even if the client is a member of the role and even when not using Windows groups):

```

//Will always fail
[PrincipalPermission(SecurityAction.Demand,Role = "...")]
class MyService : IMyContract
{...}

```



Avoid sensitive work that requires authorization in the service constructor. With a per-call service, perform such work in the operations themselves, and with a sessionful service, provide a dedicated `Initialize()` operation where you can initialize the instance and authorize the callers.

By setting the `Name` property of the `PrincipalPermission` attribute, you can even insist on granting access only to a particular user:

```
[PrincipalPermission(SecurityAction.Demand,Name = "John")]
```

or to a particular user that is a member of a particular role:

```
[PrincipalPermission(SecurityAction.Demand,Name = "John",
Role = "Manager")]
```

These practices are inadvisable, however, because it is best to avoid hardcoding user-names.



Declarative role-based security hardcodes the role name. If your application looks up role names dynamically you have to use programmatic role verification, as presented next.

Programmatic role-based security

Sometimes you need to programmatically verify role membership. Usually, you need to do that when the decision as to whether to grant access depends both on role membership and on some other values known only at call time, such as parameter values, time of day, and location. Another case in which programmatic role membership verification is needed is when you're dealing with localized user groups. To demonstrate the first category, imagine a banking service that lets clients transfer sums of money between two specified accounts. Only customers and tellers are allowed to call the `TransferMoney()` operation, with the following business rule: if the amount transferred is greater than 50,000, only tellers are allowed to do the transfer. Declarative role-based security can verify that the caller is either a teller or a customer, but it cannot enforce the additional business rule. For that, you need to use the `IsInRole()` method of `IPrincipal`, as shown in [Example 10-7](#).

Example 10-7. Programmatic role-based security

```
[ServiceContract]
interface IBankAccounts
{
    [OperationContract]
    void TransferMoney(double sum, long sourceAccount, long destinationAccount);
}
```

```

static class AppRoles
{
    public const string Customer = "Customer";
    public const string Teller   = "Teller";
}
class BankService : IBankAccounts
{

    [PrincipalPermission(SecurityAction.Demand,Role = AppRoles.Customer)]
    [PrincipalPermission(SecurityAction.Demand,Role = AppRoles.Teller)]
    public void TransferMoney(double sum,long sourceAccount,long destinationAccount)
    {
        IPrincipal principal = Thread.CurrentPrincipal;
        Debug.Assert(principal.Identity.IsAuthenticated);

        bool isCustomer = principal.IsInRole(AppRoles.Customer);
        bool isTeller   = principal.IsInRole(AppRoles.Teller);

        if(isCustomer && ! isTeller)
        {
            if(sum > 50000)      {
                string message = "Caller does not have sufficient authority to" +
                                  "transfer this sum";
                throw new SecurityException(message);
            }
        }
        DoTransfer(sum,sourceAccount,destinationAccount);
    }
    //Helper method
    void DoTransfer(double sum,long sourceAccount,long destinationAccount)
    {...}
}

```

Example 10-7 also demonstrates a number of other points. First, even though it uses programmatic role membership verification with the value of the `sum` argument, it still uses declarative role-based security as the first line of defense, allowing access only to clients who are members of the Customer or Teller roles. Second, you can programmatically assert that the caller is authenticated using the `IsAuthenticated` property of `IIdentity`. Finally, note the use of the `AppRoles` static class to encapsulate the actual string used for the role to avoid hardcoding the roles in multiple places.



There is a complete disconnect between role-based security and the actual principal type. When the `PrincipalPermission` attribute is asked to verify role membership, it simply gets hold of its thread's current principal in the form of `IPrincipal`, and calls its `IsInRole()` method. This is also true of programmatic role membership verification that uses only `IPrincipal`, as shown in [Example 10-7](#). The separation of the `IPrincipal` interface from its implementation is the key to providing other role-based security mechanisms besides Windows groups, as you will see in the other scenarios.

Identity Management

In the intranet scenario, after successful authentication, WCF will attach to the operation thread a principal identity of the type `WindowsIdentity`, which will have the value of its `Name` property set to the username (or Windows account) provided by the client. Since valid credentials are provided, the security call context's two identities—the primary identity and the Windows identity—will be set to the same identity as the principal identity. All three identities will be considered authenticated. The identities and their values are shown in [Table 10-4](#).

Table 10-4. Identity management in the intranet scenario

Identity	Type	Value	Authenticated
Thread principal	<code>WindowsIdentity</code>	Username	Yes
Security context primary	<code>WindowsIdentity</code>	Username	Yes
Security context Windows	<code>WindowsIdentity</code>	Username	Yes

Windows Roles Localization

If your application is deployed in international markets and you use Windows groups as roles, it's likely the role names will not match. In the intranet scenario, the principal object attached to the thread accessing the service is of the type `WindowsPrincipal`:

```
public class WindowsPrincipal : IPrincipal
{
    public WindowsPrincipal(WindowsIdentity ntIdentity);

    //IPrincipal implementation
    public virtual IIIdentity Identity
    {get;}
    public virtual bool IsInRole(string role);

    //Additional methods:
    public virtual bool IsInRole(int rid);
```

```
        public virtual bool IsInRole(WindowsBuiltInRole role);
    }
```

WindowsPrincipal provides two additional IsInRole() methods that are intended to ease the task of localizing Windows groups. You can provide IsInRole() with an enum of the type WindowsBuiltInRole matching the built-in Windows roles, such as WindowsBuiltInRole.Administrator or WindowsBuiltInRole.User. The other version of IsInRole() accepts an integer indexing specific roles. For example, a role index of 512 maps to the Administrators group. The MSDN Library contains a list of both the predefined indexes and ways to provide your own aliases and indexes to user groups.

Note that while the host processes retain their designated identities, the principal identity will be that of the caller. I call this behavior *soft impersonation*. When it is used in conjunction with role-based security, it largely negates the need to ever perform real impersonation and replace the security token with that of the client.

Callbacks

When it comes to security on the intranet, there are several key differences between normal service operations and callbacks. First, with a callback contract you can only assign a protection level at the operation level, not the callback contract level. For example, this protection-level constraint will be ignored:

```
[ServiceContract(CallbackContract = typeof(IMyContractCallback))]
interface IMyContract
{...}

//Demand for protection level will be ignored
[ServiceContract(ProtectionLevel = ProtectionLevel.EncryptAndSign)]
interface IMyContractCallback
{...}
```

Only the service contract designating the callback contract can set a contract-level protection constraint. WCF deliberately ignores the service contract attribute on the callback contract (as explained in [Chapter 5](#)) to avoid a potential conflict between two contract attributes that apply to the same channel.

You can take advantage of operation-level demand for a protection level as follows:

```
[ServiceContract(CallbackContract = typeof(IMyContractCallback))]
interface IMyContract
{...}

interface IMyContractCallback
{
    [OperationContract(ProtectionLevel = ProtectionLevel.EncryptAndSign)]
```

```
        void OnCallback();
    }
```

All calls into the callback object come in with an unauthenticated principal, even if Windows security was used across the board to invoke the service. As a result, the principal identity will be set to a Windows identity with a blank identity, which will preclude authorization and role-based security.

While the callback does have a security call context, the Windows identity will be set to a `WindowsIdentity` instance with a blank identity, which will preclude impersonation. The only meaningful information will be in the primary identity, which will be set to the service host's process identity and machine name:

```
class MyClient : IMyContractCallback
{
    public void OnCallback()
    {
        IPrincipal principal = Thread.CurrentPrincipal;
        Debug.Assert(principal.Identity.IsAuthenticated == false);

        ServiceSecurityContext context = ServiceSecurityContext.Current;
        Debug.Assert(context.PrimaryIdentity.Name == "MyHost/localhost");

        Debug.Assert(context.IsAnonymous == false);
    }
}
```

I recommend avoiding any sensitive work in the callback, since you cannot easily use role-based security.

Internet Application Scenario

In the Internet scenario, the clients or services may not be using WCF, or even Windows. If you are writing an Internet service or client, you cannot assume the use of WCF on the other end. In addition, an Internet application typically has a relatively large number of clients calling the service. These client calls originate from outside the firewall. You need to rely on HTTP for transport, and multiple intermediaries are possible. In an Internet application, you typically do not want to use Windows accounts and groups for credentials; instead, the application needs to access some custom credentials store. That said, you could still be using Windows security, as demonstrated later.

Securing the Internet Bindings

In an Internet application, you must use Message security for the transfer security mode to provide for end-to-end security across all intermediaries. The client should provide credentials in the form of a username and password, as this is a safe, low common denominator that all platforms support. For the Internet scenario, you

should use the `WSHttpBinding`. You cannot use the basic binding because it does not provide for username credentials over Message security. In addition, if you have an intranet application that uses the `NetTcpBinding` but you do not wish to use Windows security for user accounts and groups, you should follow the same configuration as with the WS-based binding. This is done uniformly across these bindings by selecting `MessageCredentialType.Username` for the client credentials type used with Message security. You need to configure the bindings this way both at the client and at the service.

WSHttpBinding configuration

`WSHttpBinding` offers the `Security` property of the type `WSHttpSecurity`:

```
public class WSHttpBinding : WSHttpBindingBase
{
    public WSHttpBinding();
    public WSHttpBinding(SecurityMode securityMode);
    public WSHttpSecurity Security
    {get;}
    //More members
}
```

With `WSHttpSecurity`, you need to set the `Mode` property of the type `SecurityMode` to `SecurityMode.Message`. The `Message` property of `WSHttpSecurity` will then take effect:

```
public sealed class WSHttpSecurity
{
    public SecurityMode Mode
    {get;set;}
    public NonDualMessageSecurityOverHttp Message
    {get;}
    public HttpTransportSecurity Transport
    {get;}
}
```

`Message` is of the type `NonDualMessageSecurityOverHttp`, which derives from `MessageSecurityOverHttp`:

```
public class MessageSecurityOverHttp
{
    public MessageCredentialType ClientCredentialType
    {get;set;}
    //More members
}
public sealed class NonDualMessageSecurityOverHttp : MessageSecurityOverHttp
{...}
```

You need to set the `ClientCredentialType` property of `MessageSecurityOverHttp` to `MessageCredentialType.Username`. Recall that the default Message security credentials type of the `WSHttpBinding` is Windows (see [Table 10-3](#)).

Because Message security is the default security mode of the `WSHttpBinding` (see [Table 10-1](#)), these three definitions are equivalent:

```
WSHttpBinding binding1 = new WSHttpBinding();
binding1.Security.Message.ClientCredentialType = MessageCredentialType.UserName;

WSHttpBinding binding2 = new WSHttpBinding(SecurityMode.Message);
binding2.Security.Message.ClientCredentialType = MessageCredentialType.UserName;

WSHttpBinding binding3 = new WSHttpBinding();
binding3.Security.Mode = SecurityMode.Message;
binding3.Security.Message.ClientCredentialType = MessageCredentialType.UserName;
```

You can achieve the same configuration using a config file as follows:

```
<bindings>
    <wsHttpBinding>
        <binding name = "UserNameWS">
            <security mode = "Message">
                <message clientCredentialType = "UserName"/>
            </security>
        </binding>
    </wsHttpBinding>
</bindings>
```

Or, since Message security is the default, you can omit explicitly setting the mode in the config file:

```
<bindings>
    <wsHttpBinding>
        <binding name = "UserNameWS">
            <security>
                <message clientCredentialType = "UserName"/>
            </security>
        </binding>
    </wsHttpBinding>
</bindings>
```

[Figure 10-4](#) shows the security-related elements of the `WSHttpBinding`.

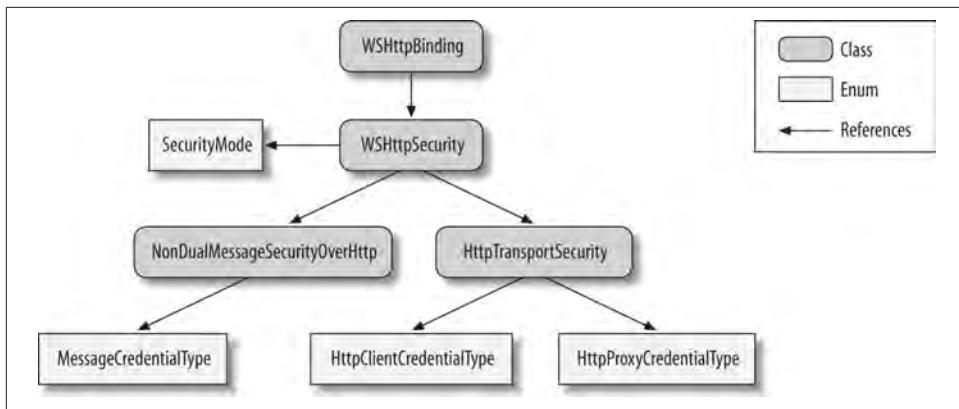


Figure 10-4. *WSHttpBinding and security*

`WSHttpBinding` has a reference to `WSHttpSecurity`, which uses the `SecurityMode` enum to indicate the transfer security mode. When Transport security is used, `WSHttpSecurity` will use an instance of `HttpTransportSecurity`. When Message security is used, `WSHttpSecurity` will use an instance of `NonDualMessageSecurityOverHttp` containing the client credentials type via the `MessageCredentialType` enum.

Message Protection

Since in the Internet scenario the client's message is transferred to the service over plain HTTP, it is vital to protect its content (both the client's credentials and the body of the message) by encrypting it. Encryption will provide for message integrity and privacy. One technical option for encryption is to use the client's password. However, WCF never uses this option, for a number of reasons. First, there are no guarantees that the password is strong enough, so anyone monitoring the communication could potentially break the encryption using a dictionary attack. Second, this approach forces the service (or more precisely, its host) to have access to the password, thus coupling the host to the credentials store. Finally, while the password may protect the message, it will not authenticate the service to the client.

Instead, to protect the message, WCF uses an X509 certificate. The certificate provides strong protection, and it authenticates the service to the client. A certificate works by using two keys, called the *public* and *private keys*, as well as a *common name* (CN) such as "MyCompanyCert." What is important about those keys is that anything encrypted with the public key can only be decrypted with the matching private one. The certificate contains the public key and the common name, and the private key is kept in some secure storage on the host machine to which the host has access. The host makes the certificate (and its public key) publicly available, so any client can access the host's endpoints and obtain the public key.

In a nutshell, what happens during a call is that WCF on the client's side uses the public key to encrypt all messages to the service. Upon receiving the encrypted message, WCF decrypts the message on the host side using the private key. Once the message is decrypted, WCF will read the client's credentials from the message, authenticate the client, and allow it to access the service. The real picture is a bit more complex, because WCF also needs to secure the reply messages and callbacks from the service to the client. One of the standards WCF supports deals with setting up such a secure conversation. In fact, several calls are made before the first request message from the client to the service, where WCF on the client's side generates a temporary shared secret it passes encrypted (using the service certificate) to the service. The client and the service will use that shared secret to protect all subsequent communication between them.

Configuring the host certificate

The `ServiceHostBase` class offers the `Credentials` property of the type `ServiceCredentials`. `ServiceCredentials` is a service behavior:

```
public abstract class ServiceHostBase : ...
{
    public ServiceCredentials Credentials
    {get;}
    //More members
}
public class ServiceCredentials : ...,IServiceBehavior
{
    public X509CertificateRecipientServiceCredential ServiceCertificate
    {get;}
    //More members
}
```

`ServiceCredentials` provides the `ServiceCertificate` property of the type `X509CertificateRecipientServiceCredential`:

```
public sealed class X509CertificateRecipientServiceCredential
{
    public void SetCertificate(StoreLocation storeLocation,
                               StoreName storeName,
                               X509FindType findType,
                               object findValue);
    //More members
}
```

You can use the `SetCertificate()` method to instruct WCF where and how to load the service certificate. You typically provide this information in the host config file as a custom behavior under the `serviceCredentials` section, as shown in [Example 10-8](#).

Example 10-8. Configuring the service certificate

```
<services>
    <service name = "MyService" behaviorConfiguration = "Internet">
        ...
    </service>
</services>
<behaviors>
    <serviceBehaviors>
        <behavior name = "Internet">
            <serviceCredentials>
                <serviceCertificate
                    findValue      = "MyServiceCert"
                    storeLocation = "LocalMachine"
                    storeName     = "My"
                    x509FindType  = "FindBySubjectName"
                />
            </serviceCredentials>
        </behavior>
    </serviceBehaviors>
</behaviors>
```

Using the host certificate

The client developer can obtain the service certificate using any out-of-band mechanism (such as email, or via a public web page). The client can then include in its config file in the endpoint behavior section detailed information about the service certificate, such as where it is stored on the client side and how to find it. This is by far the most secure option from the client's perspective, because any attempt to subvert the client's address resolving and redirect the call to a malicious service will fail since the other service will not have the correct certificate. This is the least flexible option as well, however, because every time the client needs to interact with a different service, the client administrator will need to rework the client's config file.

A reasonable alternative to explicitly referencing the certificates of all services the client may interact with is to store those certificates in the client's Trusted People certificate folder. The administrator can then instruct WCF to allow calls only to services whose certificates are in that folder. In that case, the client will need to obtain the service certificate at runtime as part of the initial pre-call negotiation, check to see whether it is in the Trusted People store, and if so, proceed to use it to protect the message. This certificate negotiation behavior is the default for the WS bindings. You can disable it and use a hard-configured certificate instead, but for the Internet scenario I strongly recommend using certificate negotiation and storing the certificates in the Trusted People store.

Service certificate validation

To instruct WCF as to what degree to validate and trust the service certificate, add a custom endpoint behavior to the client's config file. The behavior should use the `clientCredentials` section. `ClientCredentials` is an endpoint behavior that offers the `ServiceCertificate` property of the type `X509CertificateRecipientClientCredential`:

```
public class ClientCredentials : ..., IEndpointBehavior
{
    public X509CertificateRecipientClientCredential ServiceCertificate
    {get;}
    //More members
}
```

`X509CertificateRecipientClientCredential` offers the `Authentication` property of the type `X509CertificateRecipientClientCredential`:

```
public sealed class X509CertificateRecipientClientCredential
{
    public X509ServiceCertificateAuthentication Authentication
    {get;}
    //More members
}
```

`X509CertificateRecipientClientCredential` provides the `CertificateValidationMode` property of the enum type `X509CertificateValidationMode`:

```
public enum X509CertificateValidationMode
{
    None,
    PeerTrust,
    ChainTrust,
    PeerOrChainTrust,
    Custom
}

public class X509ServiceCertificateAuthentication
{
    public X509CertificateValidationMode CertificateValidationMode
    {get;set;}
    //More members
}
```

[Example 10-9](#) demonstrates setting the service certificate validation mode in the client's config file.

Example 10-9. Validating the service certificate

```
<client>
<endpoint behaviorConfiguration = "ServiceCertificate"
...
```

```

        </endpoint>
    </client>
    <behaviors>
        <endpointBehaviors>
            <behavior name = "ServiceCertificate">
                <clientCredentials>
                    <serviceCertificate>
                        <authentication certificateValidationMode = "PeerTrust"/>
                    </serviceCertificate>
                </clientCredentials>
            </behavior>
        </endpointBehaviors>
    </behaviors>

```

`X509CertificateValidationMode.PeerTrust` instructs WCF to trust the negotiated service certificate if it is present in the client's Trusted People store. `X509CertificateValidationMode.ChainTrust` instructs WCF to trust the certificate if it was issued by a root authority (such as VeriSign or Thwart) whose certificate is found in the client's Trusted Root Authority folder. `X509CertificateValidationMode.ChainTrust` is the default value used by WCF. `X509CertificateValidationMode.PeerOrChainTrust` allows either of those options. Since there are a number of illicit ways of obtaining a valid certificate from a public root authority, I do not recommend using this value. `X509CertificateValidationMode.PeerOrChainTrust` is available for tightly controlled environments that purge all public root authorities and install their own root certificates, which are used to sign other certificates.

Working with a test certificate

Developers often do not have access to their organizations' certificates, and therefore resort to using test certificates such as the ones generated by the *MakeCert.exe* command-line utility. There are two problems with test certificates. The first is that they will fail the default certificate validation on the client side, since the client uses `X509CertificateValidationMode.ChainTrust` by default. You can easily overcome this by installing the test certificate in the client's Trusted People store and using `X509CertificateValidationMode.PeerTrust`. The second problem is that WCF by default expects the service certificate name to match the service host's domain (or machine) name. This provides yet another line of defense, since typically with an Internet-facing service, the host domain name will match its certificate common name. To compensate, the client must explicitly specify the test certificate name in the endpoint identity's `dns` section:

```

<client>
    <endpoint
        address  = "http://localhost:8001/MyService"
        binding  = "wsHttpBinding"
        contract = "IMyContract">
        <identity>

```

```
<dns value = "MyServiceCert"/>
</identity>
</endpoint>
</client>
```

Authentication

The client needs to provide its credentials to the proxy. The `ClientCredentials` property (presented earlier) of the `ClientBase<T>` base class has the `UserName` property of the type `UserNamePasswordClientCredential`:

```
public class ClientCredentials : ..., IEndpointBehavior
{
    public UserNamePasswordClientCredential UserName
    {get;} //More members
}

public sealed class UserNamePasswordClientCredential
{
    public string UserName
    {get;set;}
    public string Password
    {get;set;}
}
```

The client uses `UserNamePasswordClientCredential` to pass its username and password to the service, as demonstrated in [Example 10-10](#).

Example 10-10. Providing username and password credentials

```
MyContractClient proxy = new MyContractClient();

proxy.ClientCredentials.UserName.UserName = "MyUsername";
proxy.ClientCredentials.UserName.Password = "MyPassword";

proxy.MyMethod();
proxy.Close();
```



The client need not provide a domain name (if Windows security is used) or application name (if the ASP.NET providers are used). The host will use its service domain or a configured application name, as appropriate.

When working with a channel factory instead of a proxy class, you must set the `Credentials` property of the factory with the credentials:

```
ChannelFactory<IMyContract> factory = new ChannelFactory<IMyContract>("");
```

```

factory.Credentials.UserName.UserName = "MyUsername";
factory.Credentials.UserName.Password = "MyPassword";

IMyContract proxy = factory.CreateChannel();
using(proxy as IDisposable)
{
    proxy.MyMethod();
}

```

Note that you cannot use the static `CreateChannel()` methods of `ChannelFactory<T>`, since you have to instantiate a factory in order to access the `Credentials` property.

Once the username and password credentials are received by the WCF on the service side, the host can choose to authenticate them as Windows credentials, ASP.NET membership provider's credentials, or even custom credentials. Whichever option you choose, make sure it matches your role-based policy configuration.

The `ServiceCredentials` class (available via the `Credentials` property of `ServiceHostBase`) provides the `UserNameAuthentication` property of the type `UserNamePasswordServiceCredential`:

```

public class ServiceCredentials : ...,IServiceBehavior
{
    public UserNamePasswordServiceCredential UserNameAuthentication
    {get;}
    //More members
}

```

`UserNamePasswordServiceCredential` has the `UserNamePasswordValidationMode` property of a matching enum type:

```

public enum UserNamePasswordValidationMode
{
    Windows,
    MembershipProvider,
    Custom
}
public sealed class UserNamePasswordServiceCredential
{
    public MembershipProvider MembershipProvider
    {get;set;}
    public UserNamePasswordValidationMode UserNamePasswordValidationMode
    {get; set;}
    //More members
}

```

By setting the `UserNamePasswordValidationMode` property, the host chooses how to authenticate the incoming username and password credentials.

Using Windows Credentials

While not necessarily common, WCF lets the Internet-facing service authenticate the incoming credentials as Windows credentials. To authenticate the client's username and password as Windows credentials, you need to set `UserNamePasswordValidationMode` to `UserNamePasswordValidationMode.Windows`. Because `UserNamePasswordValidationMode.Windows` is the default value of the `UserNamePasswordValidationMode` property, these two definitions are equivalent:

```
ServiceHost host1 = new ServiceHost(typeof(MyService));  
  
ServiceHost host2 = new ServiceHost(typeof(MyService));  
host2.Credentials.UserNameAuthentication.UserNamePasswordValidationMode =  
    UserNamePasswordValidationMode.Windows;
```

When using a config file, add a custom behavior that assigns the username and password authentication mode along with the service certificate information, as shown in [Example 10-11](#).

Example 10-11. Internet security with Windows credentials

```
<services>  
    <service name = "MyService" behaviorConfiguration = "UsernameWindows">  
        ...  
    </service>  
</services>  
<behaviors>  
    <serviceBehaviors>  
        <behavior name = "UsernameWindows">  
            <serviceCredentials>  
                <userNameAuthentication userNamePasswordValidationMode = "Windows"/>  
                <serviceCertificate  
                    ...  
                />  
            </serviceCredentials>  
        </behavior>  
    </serviceBehaviors>  
</behaviors>
```

As with the programmatic case, adding this line to the config file:

```
<userNameAuthentication userNamePasswordValidationMode = "Windows"/>
```

is optional because it is the default setting.

Authorization

If the `PrincipalPermissionMode` property of `ServiceAuthorizationBehavior` is set to its default value of `PrincipalPermissionMode.UseWindowsGroups`, once the username and password are authenticated against Windows, WCF installs a Windows

principal object and attaches it to the thread. This enables the service to freely use Windows groups for authorization, just as with the intranet case, both declaratively and programmatically.

Identity management

As long as the principal permission mode is set to `PrincipalPermissionMode.UseWindowsGroups`, the identity management aspect of the Internet scenario is just as with the intranet scenario, including the identities of the security call context, as shown in [Table 10-4](#). The main difference between an intranet application and an Internet application that both use Windows credentials is that with the latter the client cannot dictate the allowed impersonation level, and the host can impersonate at will. This is because WCF will assign `TokenImpersonationLevel.Impersonation` to the Windows identity of the security call context.

Using the ASP.NET Providers

By default, role-based security in WCF uses Windows user groups for roles and Windows accounts for security identities. There are several drawbacks to this default policy. First, you may not want to assign a Windows account for every client of your Internet application. Second, the security policy is only as granular as the user groups in the hosting domain. Often you do not have control over your end customers' IT departments, and if you deploy your application in an environment in which the user groups are coarse or don't map well to the actual roles users play in your application, or if the group names are slightly different, Windows role-based security will be of little use to you. Role localization presents yet another set of challenges, because role names will likely differ between customer sites in different locales. Consequently, Internet applications hardly ever use Windows accounts and groups. Out of the box, .NET 2.0 (and later) provides a custom credential management infrastructure called the ASP.NET providers. Despite its name, non-ASP.NET applications (such as WCF applications) can easily use it to authenticate users and authorize them, without ever resorting to Windows accounts.

One of the concrete implementations of the ASP.NET providers includes a SQL Server store. SQL Server is often the repository of choice for Internet applications, so I will use it in this scenario. To use the SQL Server provider, run the setup file `aspnet_regsql.exe`, found under `%Windir%\Microsoft.NET\Framework\v4.0.30319`. The setup program will create a new database called `aspnetdb`, containing the tables and stored procedures required to manage the credentials.

The SQL Server credentials store is well designed and uses the latest best practices for credential management, such as password salting, stored procedures, normalized tables, and so on. In addition to providing a high-quality, secure solution, this infrastructure aids productivity, saving developers valuable time and effort. That said, the

credential management architecture is that of a provider model, and you can easily add other storage options if required, such as an Access database.

The credentials providers

Figure 10-5 shows the architecture of the ASP.NET credentials providers.

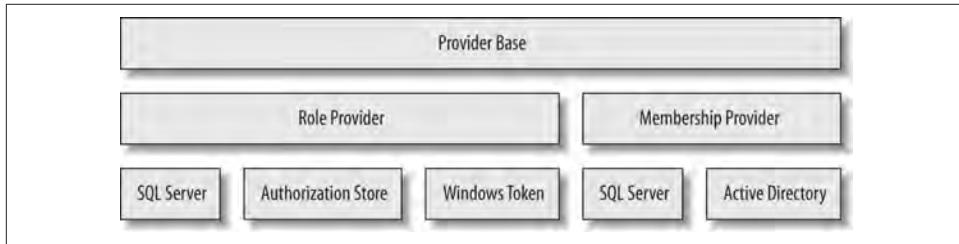


Figure 10-5. The ASP.NET provider model

Membership providers are responsible for managing users (usernames and passwords), and role providers are responsible for managing roles. Out of the box, ASP.NET offers support for membership stores in SQL Server or Active Directory, and roles can be stored in SQL Server, a file (the authorization store provider), or Windows groups (the Windows token provider).

Username and password authentication is done using a class called `MembershipProvider` from the `System.Web.Security` namespace, defined as:

```
public abstract class MembershipProvider : ProviderBase
{
    public abstract string ApplicationName
    {get;set;}
    public abstract bool ValidateUser(string username, string password);
    //Additional members
}
```

`MembershipProvider`'s goal is to encapsulate the actual provider used and the details of the actual data access, as well as to enable changing the membership provider without affecting the application itself. Depending on the configured security provider in the host config file, WCF will use a concrete data access class such as `SqlMembershipProvider`, targeting SQL Server or SQL Server Express:

```
public class SqlMembershipProvider : MembershipProvider
{...}
```

However, WCF interacts only with the `MembershipProvider` base functionality. WCF obtains the required membership provider by accessing the `Provider` static property of the `Membership` class, defined as:

```
public static class Membership
{
```

```
public static string ApplicationName  
{get;set;}  
public static MembershipProvider Provider  
{get;}  
public static bool ValidateUser(string username,string password);  
//Additional members  
}
```

Membership offers many members, which support the many aspects of user management. `Membership.Provider` retrieves the type of the configured provider from the `System.Web` section in the host config file. Unspecified, the role provider defaults to `SqlMembershipProvider`.



Because all membership providers derive from the abstract class `MembershipProvider`, if you write your own custom credential provider it needs to derive from `MembershipProvider` as well.

A single credentials store can serve many applications, and those applications may define the same usernames. To allow for that, every record in the credentials store is scoped by an application name (similar to the way usernames in Windows are scoped by a domain or machine name).

The `ApplicationName` property of `Membership` is used to set and retrieve the application name, and the `ValidateUser()` method is used to authenticate the specified credentials against the store, returning `true` if they match and `false` otherwise. `Membership.ValidateUser()` is shorthand for retrieving and using the configured provider.

If you have configured your application to use the ASP.NET credentials store for authorization and if you enabled roles support, after authentication WCF will install an instance of the internal class `RoleProviderPrincipal` and attach it to the thread invoking the operation:

```
sealed class RoleProviderPrincipal : IPrincipal  
{...}
```

`RoleProviderPrincipal` uses the abstract class `RoleProvider` for authorization:

```
public abstract class RoleProvider : ProviderBase  
{  
    public abstract string ApplicationName  
    {get;set;}  
    public abstract bool IsUserInRole(string username,string roleName);  
    //Additional members  
}
```

The `ApplicationName` property of `RoleProvider` binds the role provider to the particular application. The `IsUserInRole()` method verifies the user's role membership. Just as all membership providers must derive from `MembershipProvider`, all role providers (including custom role providers) must derive from `RoleProvider`.

`RoleProvider` encapsulates the actual provider used, and the role provider to use is specified in the host config file. Depending on the configured role provider, `RoleProviderPrincipal` uses a corresponding data access class such as `SqlRoleProvider` to authorize the caller:

```
public class SqlRoleProvider : RoleProvider  
{...}
```

You can obtain the required role provider by accessing the `Provider` static property of the `Roles` class, defined as:

```
public static class Roles  
{  
    public static string ApplicationName  
    {get;set;}  
    public static bool IsUserInRole(string username, string roleName);  
    public static RoleProvider Provider  
    {get;}  
    //Additional members  
}
```

`Roles.IsUserInRole()` is shorthand for first accessing `Roles.Provider` and then calling `IsUserInRole()` on it. `Roles.Provider` retrieves the type of the configured provider from the host config file. If unspecified, the role provider defaults to `SqlRoleProvider`.

Credentials administration

If you use SQL Server, .NET installs website administration pages under `\Inetpub\wwwroot\aspnet_webadmin\<version number>`. Developers can configure the application directly from within Visual Studio. When you select ASP.NET Configuration from the Web Site menu, Visual Studio will launch the ASP.NET development server used for the administration pages, browse to the ASP.NET administration pages, and allow you to configure various parameters, including security. You can configure the following aspects for your application:

- Create new users and delete existing ones
- Create new roles and delete existing ones
- Allocate users to roles
- Retrieve a user's details
- Set a user's status

- Use additional features not relevant to this chapter

Shortcomings of Visual Studio

There are a number of significant shortcomings to using the Visual Studio–driven administration pages. First, you need Visual Studio. It is unlikely that application or system administrators will have Visual Studio, let alone know how to use it. The administration pages use “/” by default for the application name, and do not offer any visual way to modify that. Also, you must create a web application to activate the administration pages and there is no remote access: the application and Visual Studio must be co-located in order for Visual Studio to be able to access the application’s configuration file, and the ASP.NET development server used for the administration pages cannot accept remote calls. The browser-based user interface is somewhat annoying (you need to frequently click the Back button) and rather dull. Furthermore, many features that administrators are likely to want to use are not available via the administration pages, despite the fact that the underlying provider classes support those features. Some of the things missing from the Visual Studio–driven administration pages include the ability to:

- Update most if not all of the details in a user account
- Retrieve a user’s password
- Change a user’s password
- Reset a user’s password
- Retrieve information about the number of current online users
- Remove all users from a role in one operation
- Retrieve information about the password management policy (such as length, reset policy, type of passwords, etc.)
- Test user credentials
- Verify user role membership

There are additional features that administrators are likely to want, yet they are not supported even by the provider classes. These features include the ability to retrieve a list of all of the applications in the store, the ability to remove all users from an application, the ability to remove all roles from an application, the ability to delete an application (and all its associated users and roles), and the ability to delete all applications.



The IIS control panel applet also offers some administrative support for managing the roles and membership providers. However, this support is on a par with that of Visual Studio.

Credentials Manager

This tools disparity motivated me to develop the Credentials Manager application, a smart client application that compensates for all of the shortcomings just listed. [Figure 10-6](#) shows a screenshot of Credentials Manager.¹

In Credentials Manager, which is available with *ServiceModelEx*, I wrapped the ASP.NET providers with a WCF service (which can be self-hosted or IIS- or WAS-hosted) and added the missing features, such as the ability to delete an application.

Credentials Manager uses the dedicated WCF service to administer the credentials store. In addition, it lets administrators select the address of the credentials service at runtime, and using the `MetadataHelper` class presented in [Chapter 2](#), it verifies that the address provided does indeed support the required contracts.

¹ I first published an earlier version of Credentials Manager in my article “Manage Custom Security Credentials the Smart (Client) Way” (*CoDe Magazine*, November 2005).

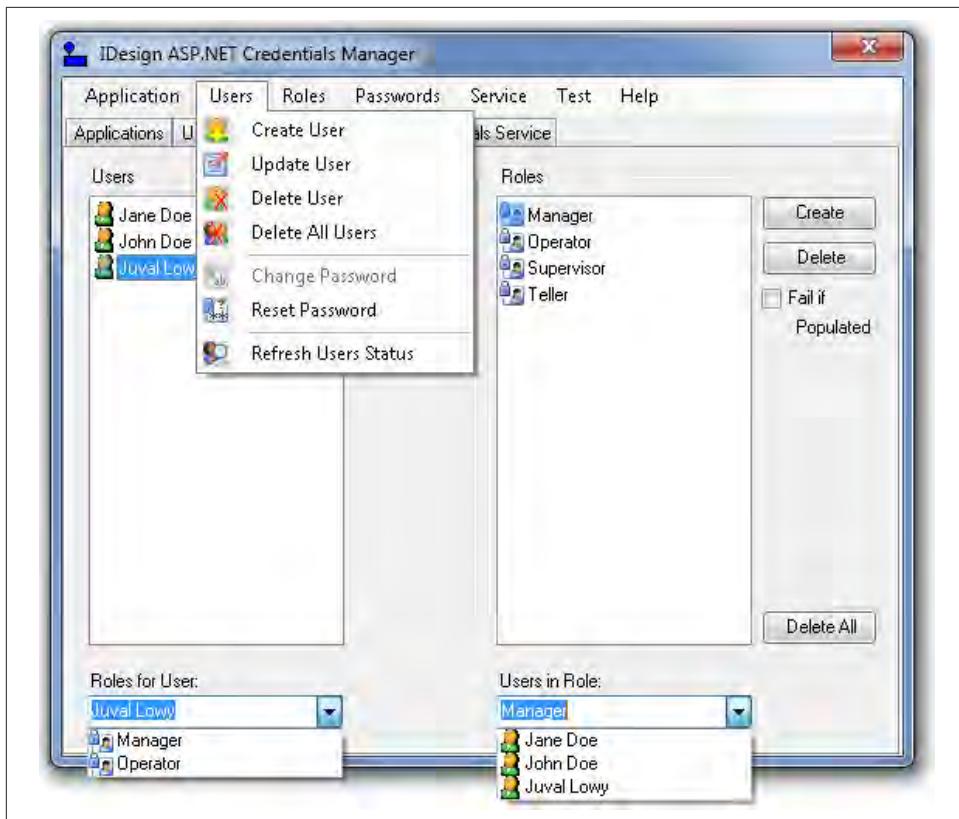


Figure 10-6. The Credentials Manager utility

Authentication

To authenticate the client's username and password using an ASP.NET provider, set the `UserNamePasswordValidationMode` property to `UserNamePasswordValidationMode.MembershipProvider`:

```
ServiceHost host = new ServiceHost(typeof(MyService));
host.Credentials.UserNameAuthentication.UserNamePasswordValidationMode =
    UserNamePasswordValidationMode.MembershipProvider;
```

Which provider is used depends on the host config file. In addition, the host config file must contain any provider-specific settings such as a SQL Server connection string, as shown in [Example 10-12](#).

Example 10-12. Internet security using an ASP.NET SQL Server provider

```
<connectionStrings>
<add name= "AspNetDb" connectionString = "data source=(local);
Integrated Security=SSPI;Initial Catalog=aspnetdb"/>
```

```

</connectionStrings>

<system.serviceModel>
    <services>
        <service name = "MyService" behaviorConfiguration = "ASPNETProviders">
            <endpoint
                ...
            />
        </service>
    </services>
    <behaviors>
        <serviceBehaviors>
            <behavior name = "ASPNETProviders">
                <serviceCredentials>
                    <userNameAuthentication
                        userNamePasswordValidationMode = "MembershipProvider"/>
                    <serviceCertificate
                        ...
                    />
                </serviceCredentials>
            </behavior>
        </serviceBehaviors>
    </behaviors>
</system.serviceModel>

```

The default application name will be a useless /, so you must assign your application's name. Once the ASP.NET providers are configured, WCF initializes the `Membership.Provider` property of `UserNamePasswordServiceCredential` with an instance of the configured membership provider. You can programmatically access that membership provider and set its application name:

```

ServiceHost host = new ServiceHost(typeof(MyService));
Debug.Assert(host.Credentials.UserNameAuthentication.MembershipProvider != null);
Membership.ApplicationName = "MyApplication";
host.Open();

```

You can also configure the application name in the config file, but for that you need to define a custom ASP.NET membership provider, as shown in [Example 10-13](#).

Example 10-13. Configuring the application name for the membership provider

```

<system.web>
    <membership defaultProvider = " MySqlMembershipProvider ">
        <providers>
            <add name = " MySqlMembershipProvider "
                type = " System.Web.Security.SqlMembershipProvider "
                connectionStringName = " AspNetDb "
                applicationName = " MyApplication "
            />
        </providers>
    </membership>

```

```
</system.web>
<connectionStrings>
    <add name = "AspNetDb"
        ...
    />
</connectionStrings>
```

First, you add a `system.Web` section with a `providers` section, where you add a custom membership provider and set that to be the new default membership provider. Next, you need to list the fully qualified type name of the new provider. Nothing prevents you from referencing an existing implementation of a membership provider (such as `SqlMembershipProvider`, as in [Example 10-13](#)). When using the SQL provider, you must also list the connection string to use, and you cannot rely on the default connection string from `machine.config`. Most importantly, you must set the `ApplicationName` tag to the desired application name.

Authorization

To support authorizing the users, the host must enable role-based security by adding this to the config file:

```
<system.web>
    <roleManager enabled = "true"/>
</system.web>
```



To enable the role manager programmatically, you have to use reflection.

Enabling roles this way will initialize the `Roles` class and have its `Provider` property set to the configured provider. To use the ASP.NET role provider, set the `PrincipalPermissionMode` property to `PrincipalPermissionMode.UseAspNetRoles`:

```
ServiceHost host = new ServiceHost(typeof(MyService));
host.Authorization.PrincipalPermissionMode =
    PrincipalPermissionMode.UseAspNetRoles;
host.Open();
```

Alternatively, when using a config file, you can add a custom behavior to that effect:

```
<services>
    <service name = "MyService" behaviorConfiguration = "ASPNETProviders">
        ...
    </service>
</services>
<behaviors>
    <serviceBehaviors>
        <behavior name = "ASPNETProviders">
```

```

<serviceAuthorization principalPermissionMode = "UseAspNetRoles"/>
...
</behavior>
</serviceBehaviors>
</behaviors>

```

After authenticating the client, the `RoleProvider` property of `ServiceAuthorizationBehavior` will be set to the configured role provider:

```

public sealed class ServiceAuthorizationBehavior : IServiceBehavior
{
    public RoleProvider RoleProvider
    {get;set;}
    //More members
}

```

The default application name will be a useless `/`, so you must assign your application's name using the static helper class `Roles`:

```

ServiceHost host = new ServiceHost(typeof(MyService));
Debug.Assert(host.Credentials.UserNameAuthentication.MembershipProvider != null);
Roles.ApplicationName = "MyApplication";

```

You can also configure the application name in the config file, but for that you need to define a custom ASP.NET role provider, as shown in [Example 10-14](#).

Example 10-14. Configuring the application name for the role provider

```

<system.web>
    <roleManager enabled = "true" defaultProvider = " MySqlRoleManager">
        <providers>
            <add name = " MySqlRoleManager"
                type = "System.Web.Security.SqlRoleProvider"
                connectionStringName = "AspNetDb"
                applicationName = "MyApplication"
            />
        </providers>
    </roleManager>
</system.web>
<connectionStrings>
    <add name = "AspNetDb"
        ...
    />
</connectionStrings>

```

As with the membership provider, you add a `system.Web` section with a `providers` section, where you add a custom role provider and set that to be the new default role provider. Next you need to list the fully qualified type name of the new provider. As with the membership provider, you can reference any existing implementation of a role provider, such as `SqlRoleProvider`, in which case you must also list the connec-

tion string to use. Finally, you must set the `ApplicationName` tag to the desired application name.

Declarative role-based security

You can use the `PrincipalPermission` attribute to verify role membership just as in the intranet scenario, because all the attribute does is access the principal object attached to the thread, which WCF has already set to `RoleProviderPrincipal`. **Example 10-15** demonstrates declarative role-based security using the ASP.NET providers.

Example 10-15. ASP.NET role provider declarative role-based security

```
class MyService : IMyContract
{
    [PrincipalPermission(SecurityAction.Demand, Role = "Manager")]
    public void MyMethod()
    {...}
}
```

Identity Management

In the Internet scenario, when you use the ASP.NET providers, the identity associated with the principal object is a `GenericIdentity` that wraps the username provided by the client. That identity is considered authenticated. The security call context's primary identity will match the principal identity. The Windows identity, on the other hand, will be set to a Windows identity with a blank username; that is, it is unauthenticated. **Table 10-5** shows the identities in this scenario.

Table 10-5. Identity management in the Internet scenario with providers

Identity	Type	Value	Authenticated
Thread principal	GenericIdentity	Username	Yes
Security context primary	GenericIdentity	Username	Yes
Security context Windows	WindowsIdentity	-	No

Impersonation

Since no valid Windows credentials are provided, the service cannot impersonate any of its clients.

Business-to-Business Application Scenario

In the business-to-business scenario, the service and its clients are disparate business entities. They do not share credentials or accounts, and the communication between

them is typically closed to the public. There are relatively few clients interacting with the service, and the client can only interact with the service after an elaborate business agreement has been established and other conditions have been met. Instead of Windows accounts or usernames, the clients identify themselves to the service using X509 certificates. These certificates are usually known a priori to the service. The client or service may not necessarily be using WCF, or even Windows. Therefore, if you are writing a service or a client, you cannot assume the use of WCF at the other end. The client calls originate from outside the firewall, and you need to rely on HTTP for transport. Also, multiple intermediaries are possible.

Securing the Business-to-Business Bindings

For the business-to-business scenario, you should use the Internet bindings; namely, `BasicHttpBinding` and `WSHttpBinding`. You must use Message security for the transfer security mode, to provide for end-to-end security across all intermediaries. The message will be protected using a service-side certificate, just as with the Internet scenario. However, unlike with the Internet scenario, here the clients provide credentials in the form of a certificate. This is done uniformly across these bindings by selecting `MessageCredentialType.Certificate` for the client credentials type to be used with the Message security mode. You need to configure this on both the client and the service. For example, to configure the `WSHttpBinding` programmatically, you would write:

```
WSHttpBinding binding = new WSHttpBinding();
binding.Security.Message.ClientCredentialType = MessageCredentialType.
    Certificate;
```

Or with a config file:

```
<bindings>
    <wsHttpBinding>
        <binding name = "WSCertificateSecurity">
            <security mode = "Message">
                <message clientCredentialType = "Certificate"/>
            </security>
        </binding>
    </wsHttpBinding>
</bindings>
```

Authentication

The service administrator has a number of options as to how to authenticate the certificates sent by the clients. If its certificate is validated, the client is considered authenticated. If no validation is done on the service side, merely sending a certificate will do. If the validation mode is set to use a chain of trust and a trusted root authority issued the certificate, the client will be considered authenticated. However, the best way of validating the client's certificate is to use *peer trust*. With this approach, the

service administrator installs the certificates of all the clients allowed to interact with the service in the Trusted People store on the service's local machine. When the service receives the client's certificate, it verifies that the certificate is in the trusted store, and if so, the client is considered authenticated. I recommend using peer trust in the business-to-business scenario.

The `ServiceCredentials` class offers the `ClientCertificate` property of the type `X509CertificateInitiatorServiceCredential`:

```
public class ServiceCredentials : ...,IServiceBehavior
{
    public X509CertificateInitiatorServiceCredential ClientCertificate
    {get;}
    //More members
}
```

`X509CertificateInitiatorServiceCredential` provides the `Authentication` property of the type `X509ClientCertificateAuthentication`, which lets you configure the certificate validation mode:

```
public sealed class X509CertificateInitiatorServiceCredential
{
    public X509ClientCertificateAuthentication Authentication
    {get;}
    //More members
}
public class X509ClientCertificateAuthentication
{
    public X509CertificateValidationMode CertificateValidationMode
    {get;set;} //More members
}
```

Example 10-16 demonstrates the settings required in the host config file for the business-to-business scenario. Note in **Example 10-16** that the host still needs to provide its own certificate for Message security.

Example 10-16. Configuring the host for business-to-business security

```
<services>
    <service name = "MyService" behaviorConfiguration = "BusinessToBusiness">
        ...
    </service>
</services>
<behaviors>
    <serviceBehaviors>
        <behavior name = "BusinessToBusiness">
            <serviceCredentials>
                <serviceCertificate
                    ...
                />
                <clientCertificate>
```

```

        <authentication certificateValidationMode = "PeerTrust"/>
    </clientCertificate>
</serviceCredentials>
</behavior>
</serviceBehaviors>
</behaviors>

```

The client needs to reference the certificate to use by including its location, name, and lookup method. This is done by accessing the `ClientCredentials` property of the proxy, which offers the `ClientCertificate` property of the type `X509CertificateInitiatorClientCredential`:

```

public class ClientCredentials : ..., IEndpointBehavior
{
    public X509CertificateInitiatorClientCredential ClientCertificate
    {get;}
    //More members
}
public sealed class X509CertificateInitiatorClientCredential
{
    public void SetCertificate(StoreLocation storeLocation,
                               StoreName storeName,
                               X509FindType findType,
                               object findValue);
    //More members
}

```

However, the client will typically set these values in its config file, as shown in [Example 10-17](#).

Example 10-17. Setting the client's certificate

```

<client>
    <endpoint behaviorConfiguration = "BusinessToBusiness"
              ...
              />
</client>
...
<behaviors>
    <endpointBehaviors>
        <behavior name = "BusinessToBusiness">
            <clientCredentials>
                <clientCertificate
                    minValue      = "MyClientCert"
                    storeLocation = "LocalMachine"
                    storeName     = "My"
                    x509FindType  = "FindBySubjectName"
                />
                ...
            </clientCredentials>
        </behavior>
    </endpointBehaviors>
</behaviors>

```

```
</endpointBehaviors>
</behaviors>
```

The config file must also indicate the service certificate validation mode. When using the `BasicHttpBinding`, since that binding cannot negotiate the service certificate, the client's config file needs to contain in the service certificate section of the endpoint behavior the location of the service certificate to use. Note that when using a service test certificate, as with the Internet scenario, the client's config file must still include the information regarding the endpoint's identity.

If the client is required to always provide the same certificate, the client developer can encapsulate setting the certificate in the proxy constructors:

```
class MyContractClient: ClientBase<...>, ...
{
    public MyContractClient()
    {
        SetCertificate();
    }
    /* More constructors */

    void SetCertificate()
    {
        ClientCredentials.ClientCertificate.SetCertificate(
            StoreLocation.LocalMachine,
            StoreName.My,
            X509FindType.FindBySubjectName,
            "MyClientCert");
    }
    //Rest of the proxy
}
```

Once the client certificate is configured, there is no need to do anything special with the proxy class:

```
MyContractClient proxy = new MyContractClient();
proxy.MyMethod();
proxy.Close();
```

Authorization

By default, the service cannot employ principal-based, role-based security. The reason is that the credentials provided—namely, the client's certificate—do not map to either Windows or ASP.NET user accounts. Because business-to-business endpoints and services are often dedicated to a small set of clients or even a particular client, this lack of authorization support may not pose a problem. If that is indeed your case, you should set the `PrincipalPermissionMode` property to `PrincipalPermissionMode.None`, so that WCF will attach a generic principal with a blank identity as opposed to a `WindowsIdentity` instance with a blank identity.

If, on the other hand, you would still like to authorize the clients, you can actually achieve just that. In essence, all you need to do is deploy some credentials store, add each client's certificate name—that is, its common name and its thumbprint—to that repository, and then perform access checks against that store as needed.

In fact, nothing prevents you from taking advantage of the ASP.NET role provider for authorization, even if you didn't use the membership provider for authentication. This ability to use the providers separately was a core design goal for the ASP.NET provider model.

First, you need to enable the role provider in the host config file and configure the application name as in [Example 10-14](#) (or provide the application name programmatically).

Next, add the client certificate and thumbprint to the membership store as a user, and assign roles to it. For example, when using a certificate whose common name is MyClientCert, you need to add a user by that name (such as “CN=MyClientCert; 12A06153D25E94902F50971F68D86DCDE2A00756”) to the membership store, and provide a password. The password, of course, is irrelevant and will not be used. Once you have created the user, assign it to the appropriate roles in the application.

Most importantly, set the `PrincipalPermissionMode` property to `PrincipalPermissionMode.UseAspNetRoles`. [Example 10-18](#) lists the required settings in the host config file.

Example 10-18. ASP.NET role-based security for the business-to-business scenario

```
<system.web>
    <roleManager enabled = "true" defaultProvider = "...">
        ...
    </roleManager>
</system.web>

<system.serviceModel>
    <services>
        <service name = "MyService" behaviorConfiguration = "BusinessToBusiness">
            ...
        </service>
    </services>
    <behaviors>
        <serviceBehaviors>
            <behavior name = "BusinessToBusiness">
                <serviceCredentials>
                    <serviceCertificate
                        ...
                    />
                    <clientCertificate>
                        <authentication certificateValidationMode = "PeerTrust"/>
                    </clientCertificate>
                </serviceCredentials>
            </behavior>
        </serviceBehaviors>
    </behaviors>
</system.serviceModel>
```

```

        </serviceCredentials>
        <serviceAuthorization principalPermissionMode = "UseAspNetRoles"/>
    </behavior>
</serviceBehaviors>
</behaviors>
<bindings>
    ...
</bindings>
</system.serviceModel>

```

Now you can use role-based security, just as in [Example 10-15](#).

Identity Management

If the `PrincipalPermissionMode` property is set to `PrincipalPermissionMode.None`, then the principal identity will be a `GenericIdentity` with a blank username. The security call context's primary identity will be of the type `X509Identity` and will contain the client certificate's common name and its thumbprint. The security call context's Windows identity will have a blank username, since no valid Windows credentials were provided. If the `PrincipalPermissionMode` property is set to `PrincipalPermissionMode.UseAspNetRoles`, then both the principal identity and the security call context's primary identity will be set to an instance of `X509Identity` containing the client certificate and thumbprint. The security call context's Windows identity will have a blank username, as before. [Table 10-6](#) details this setup.

Table 10-6. Identity management in the business-to-business scenario with role providers

Identity	Type	Value	Authenticated
Thread principal	<code>X509Identity</code>	Client cert name	Yes
Security context primary	<code>X509Identity</code>	Client cert name	Yes
Security context Windows	<code>WindowsIdentity</code>	-	No

Impersonation

Since no valid Windows credentials are provided, the service cannot impersonate any of its clients.

Host Security Configuration

While [Figure 10-7](#) is not specific to the business-to-business scenario, having covered this scenario, this is the first point in this chapter where I can show all the pieces of the service host pertaining to security.

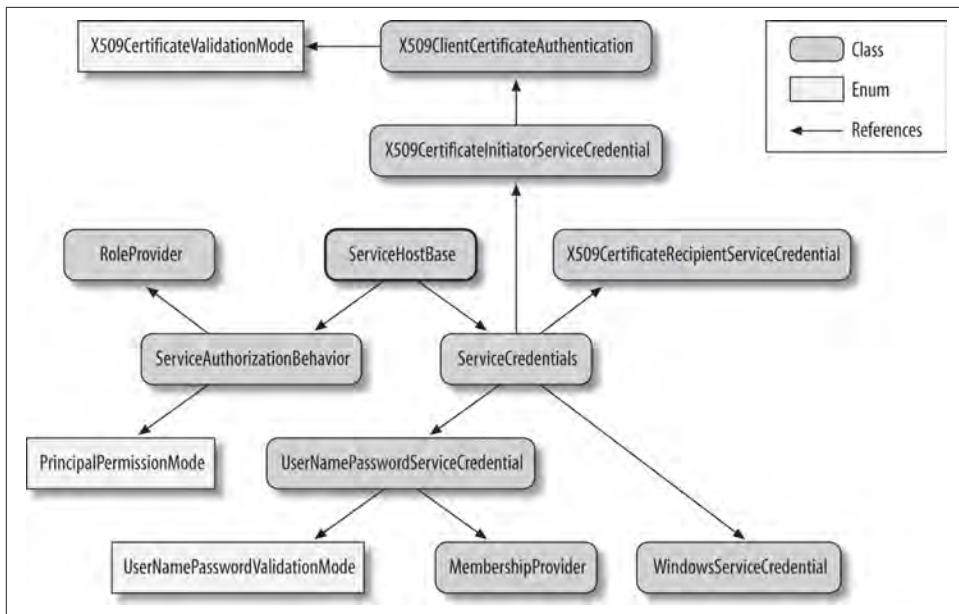


Figure 10-7. The security elements of `ServiceHostBase`

Anonymous Application Scenario

In the anonymous scenario, the clients access the service without presenting any credentials—they are anonymous. Nevertheless, the clients and the service do require secure message transfer, impervious to tampering and sniffing. Both Internet-facing and intranet-based applications may need to provide for anonymous yet end-to-end secure access. The anonymous scenario can have any number of clients, small or large. The clients may connect over HTTP, TCP, or MSMQ.

Securing the Anonymous Bindings

The need to secure the message and the fact that the clients may be calling over the Internet with multiple intermediaries mean that in the anonymous scenario you should use Message security. With Message security, you can easily satisfy both requirements by setting the `ClientCredentialType` property to `MessageCredentialType.None`. The service needs to be configured with a certificate to secure the message itself. For the anonymous scenario, you can use only the `WSHttpBinding`, `NetTcpBinding`, and `NetMsmqBinding`—a mixture of both Internet and intranet bindings, as is required in this scenario. You cannot use the `NetNamedPipeBinding` or the `BasicHttpBinding`, as the former does not support Message security and the latter does not support having no credentials in the message (see Tables 10-1 and 10-3).

Configuring the allowed bindings is done similarly to the previous scenarios. The noticeable difference is in configuring for no client credentials. For example, here's how to configure the `WSHttpBinding`:

```
WSHttpBinding binding = new WSHttpBinding();
binding.Security.Message.ClientCredentialType = MessageCredentialType.None;
```

You can also do this using a config file:

```
<bindings>
  <wsHttpBinding>
    <binding name = "WSAnonymous">
      <security>
        <message clientCredentialType = "None"/>
      </security>
    </binding>
  </wsHttpBinding>
</bindings>
```

Authentication

No client authentication is done in the anonymous scenario, of course, and the client need not provide any credentials to the proxy. For the purposes of service authentication to the client and message protection, the service needs to provide its certificate, as in [Example 10-8](#).

Authorization

Since the clients are anonymous (and unauthenticated), authorization and role-based security are precluded. The service host should set the `PrincipalPermissionMode` property to `PrincipalPermissionMode.None` to have WCF install a generic principal with a blank identity, instead of a Windows principal with a blank identity.

Identity Management

Assuming the use of `PrincipalPermissionMode.None`, the identity associated with the principal object is a `GenericIdentity` with a blank username. That identity is considered unauthenticated. The security call context's primary identity will match the principal identity. The Windows identity, on the other hand, will be set to a Windows identity with a blank username—that is, it will be unauthenticated. [Table 10-7](#) shows the identities in this scenario.

Table 10-7. Identity management in the anonymous scenario

Identity	Type	Value	Authenticated
Thread principal	GenericIdentity	-	No
Security context primary	GenericIdentity	-	No

Identity	Type	Value	Authenticated
Security context Windows	WindowsIdentity	-	No

Impersonation

Since the clients are anonymous, the service cannot impersonate any of its clients.

Callbacks

While the call from the client to the service is anonymous, the service does reveal its identity to the client. The primary identity of the security call context will be set to an instance of the `X509Identity` class, with the name set to the common name of the service host certificate suffixed by the certificate's thumbprint. The rest of the information is masked out. The principal identity will be set to a Windows identity with a blank username, which will preclude authorization and role-based security, as it is considered anonymous. The security call context's Windows identity will be set to a `WindowsIdentity` instance with a blank identity, which will preclude impersonation. Avoid sensitive work in the callback, since you cannot use role-based security.

No Security Scenario

In this last scenario, your application turns off security completely. The service does not rely on any transfer security, and it does not authenticate or authorize its callers. Obviously, such a service is completely exposed, and you generally need a very good business justification for relinquishing security. Both Internet and intranet services can be configured for no security, and they can accept any number of clients.

Unsecuring the Bindings

To turn off security, you need to set the transfer security mode to `None`. This will also avoid storing any client credentials in the message. All bindings support no transfer security (see [Table 10-1](#)).

Configuring the allowed bindings is done similarly to the previous scenarios, except the transfer security mode is set to `None`. For example, here's how to configure the `NetTcpBinding` programmatically:

```
NetTcpBinding binding = new NetTcpBinding(SecurityMode.None);
```

And here's how to do this using a config file:

```
<bindings>
  <netTcpBinding>
    <binding name = "NoSecurity">
      <security mode = "None"/>
    </binding>
```

```
</netTcpBinding>  
</bindings>
```

Authentication

No client authentication is done in this scenario, and the client does not need to provide any credentials to the proxy. Nor does the client ever authenticate the service.

Authorization

Since the clients are anonymous (and unauthenticated), authorization and role-based security are precluded. WCF will automatically set the `PrincipalPermissionMode` property to `PrincipalPermissionMode.None` to install a generic principal with a blank identity.

Identity Management

The identity associated with the principal object is a `GenericIdentity` with a blank username. That identity is considered unauthenticated. Unlike all the previous scenarios, in the no security scenario, the operation has no security call context, and the `ServiceSecurityContext.Current` returns `null`. [Table 10-8](#) shows the identities in this scenario.

Table 10-8. Identity management in the no security scenario

Identity	Type	Value	Authenticated
Thread principal	<code>GenericIdentity</code>	-	No
Security context primary	-	-	-
Security context Windows	-	-	-

Impersonation

Because the clients are anonymous, the service cannot impersonate any of its clients.

Callbacks

Unlike the intranet or the anonymous scenarios, in the absence of transfer security, callbacks come in under the client's own identity. The principal identity will be set to an instance of `WindowsIdentity` with the client's username. The callback will be authenticated, but there is no point in either impersonation or using role-based security since the client will only be authorizing itself. In addition, the security call context of the callback will be set to `null`.

Scenarios Summary

Now that you have seen the making of the five key scenarios, Tables 10-9 and 10-10 serve as a summary of their key elements. Table 10-9 lists the bindings used in each scenario. Note again that while technically you could use other bindings in almost all of the scenarios, my binding selections are aligned with the contexts in which the scenarios are used.

Table 10-9. Bindings and security scenarios

Binding	Intranet	Internet	B2B	Anonymous	None
BasicHttpBinding	No	No	Yes	No	Yes
NetTcpBinding	Yes	Yes	No	Yes	Yes
NetNamedPipeBinding	Yes	No	No	No	Yes
WSHttpBinding	No	Yes	Yes	Yes	Yes
NetMsmqBinding	Yes	No	No	Yes	Yes

Table 10-10 shows how each of the security aspects defined at the beginning of this chapter (transfer security, service and client authentication, authorization, and impersonation) relates to each scenario.

Table 10-10. The security aspects of the various scenarios

Aspect	Intranet	Internet	B2B	Anonymous	None
Transport security	Yes	No	No	No	No
Message security	No	Yes	Yes	Yes	No
Service authentication	Windows	Certificate	Certificate	Certificate	No
Client authentication	Windows	ASP.NET	Certificate	No	No
Authorization	Windows	ASP.NET	No/ASP.NET	No	No
Impersonation	Yes	No	No	No	No

Declarative Security Framework

WCF security is truly a vast topic. The number of details to master is daunting, and intricate relationships exist between the various aspects. The programming model is very complex, and at first you're likely to have an inescapable feeling of navigating a maze. To make things even worse, getting it wrong has severe implications both at the application and the business level. To simplify things, I came up with a declarative security framework for WCF. For the service, I have provided a security attribute (with matching support for the host), and for the client I have provided a few helper classes and secure proxy classes. My declarative framework grossly simplifies WCF security configuration, placing it on a par with other aspects of WCF configuration

such as transactions and synchronization. My goal was to provide a declarative model that would be simple to use and would minimize the need to understand the many details of security. As a developer, all you need to do is select the correct scenario (out of the five common scenarios discussed in this chapter), and my framework will automate the configuration. Not only that, but my framework mandates the correct options and enforces my recommendations. At the same time, my model maintains granularity and allows developers to control the underlying configuration if the need arises.

The SecurityBehaviorAttribute

Example 10-19 lists the definition of the `SecurityBehaviorAttribute` and the `ServiceSecurity` enum. `ServiceSecurity` defines the five scenarios supported by my framework.

Example 10-19. The SecurityBehaviorAttribute

```
public enum ServiceSecurity
{
    None,
    Anonymous,
    BusinessToBusiness,
    Internet,
    Intranet
}
[AttributeUsage(AttributeTargets.Class)]
public class SecurityBehaviorAttribute : Attribute, IServiceBehavior
{
    public SecurityBehaviorAttribute(ServiceSecurity mode);
    public SecurityBehaviorAttribute(ServiceSecurity mode,
                                    string serviceCertificateName);
    public SecurityBehaviorAttribute(ServiceSecurity mode,
                                    StoreLocation storeLocation,
                                    StoreName storeName,
                                    X509FindType findType,
                                    string serviceCertificateName);

    public bool ImpersonateAll
    {get;set;}
    public string ApplicationName
    {get;set;}
    public bool UseAspNetProviders
    {get;set;}
}
```

When applying the `SecurityBehavior` attribute, you need to provide it with the target scenario in the form of a `ServiceSecurity` value. You can use just the constructors of the `SecurityBehavior` attribute, or you can set the properties. Unset, the properties all default to reasonable values in the context of the target scenario. When

selecting a scenario, the configured behavior follows to the letter my previous descriptions of the individual scenarios. The `SecurityBehavior` attribute yields a composable security model, allowing quite a few permutations and sub-scenarios. When using the attribute, you can even have a security-free host config file, or you can combine settings from the config file with values driven by the attribute. Similarly, your hosting code can be free of security, or you can combine programmatic host security with the attribute.

Configuring an intranet service

To configure a service for the intranet security scenario, apply `SecurityBehavior` with `ServiceSecurity.Intranet`:

```
[ServiceContract]
interface IMyContract
{
    [OperationContract]
    void MyMethod();
}
[SecurityBehavior(ServiceSecurity.Intranet)]
class MyService : IMyContract
{
    public void MyMethod()
    {...}
}
```

Even though the service contract used may not constrain the protection level, the attribute programmatically adds that demand to enforce message protection. You can use Windows groups for role-based security:

```
[SecurityBehavior(ServiceSecurity.Intranet)]
class MyService : IMyContract
{
    [PrincipalPermission(SecurityAction.Demand,Role = "Customer")]
    public void MyMethod()
    {...}
}
```

The service can programmatically impersonate the callers, or use the operation behavior attribute for individual method impersonation. You can also configure the service to automatically impersonate all callers in all methods via the `ImpersonateAll` property. `ImpersonateAll` defaults to `false`, but when it's set to `true` the attribute will impersonate all callers in all operations without your needing to apply any operation behavior attributes or do any host configuration:

```
[SecurityBehavior(ServiceSecurity.Intranet,ImpersonateAll = true)]
class MyService : IMyContract
{...}
```

Configuring an Internet service

With the Internet scenario, you need to both configure the service for this scenario and select the service certificate to use. Note in [Example 10-19](#) that the `ServiceBehavior` attribute constructor may take the service certificate name. If it's unspecified, the service certificate is loaded from the host config file as with [Example 10-8](#):

```
[SecurityBehavior(ServiceSecurity.Internet)]
class MyService : IMyContract
{...}
```

You can also specify the service certificate name, in which case the specified certificate is loaded from the `LocalMachine` store from the `My` folder by name:

```
[SecurityBehavior(ServiceSecurity.Internet,"MyServiceCert")]
class MyService : IMyContract
{...}
```

If the certificate name is set to an empty string, the `SecurityBehavior` attribute will infer the certificate name by using the hosting machine name (or domain) for the certificate name and load such a certificate from the `LocalMachine` store from the `My` folder by name:

```
[SecurityBehavior(ServiceSecurity.Internet,"")]
class MyService : IMyContract
{...}
```

Finally, the attribute lets you explicitly specify the store location, the store name, and the lookup method:

```
[SecurityBehavior(ServiceSecurity.Internet,
                  StoreLocation.LocalMachine,StoreName.My,
                  X509FindType.FindBySubjectName,"MyServiceCert")]
class MyService : IMyContract
{...}
```

Note that you can combine an explicit location with an inferred certificate name:

```
[SecurityBehavior(ServiceSecurity.Internet,
                  StoreLocation.LocalMachine,StoreName.My,
                  X509FindType.FindBySubjectName,"")]
class MyService : IMyContract
{...}
```

Which credentials store to authenticate the client against is indicated by the `UseAspNetProviders` property. `UseAspNetProviders` defaults to `false`, meaning that the default is to authenticate the client's username and password as Windows credentials (as in [Example 10-11](#)). Because of that, when `UseAspNetProviders` is `false` you can by default use Windows groups for authorization and even impersonate all callers:

```
[SecurityBehavior(ServiceSecurity.Internet,"MyServiceCert",ImpersonateAll =
true )]
```

```
class MyService : IMyContract
{...}
```

If `UseAspNetProviders` is set to `true`, instead of Windows credentials the `SecurityBehavior` attribute will use the ASP.NET membership and role providers, as prescribed for the Internet scenario:

```
[SecurityBehavior(ServiceSecurity.Internet,"MyServiceCert",
                  UseAspNetProviders = true)]
class MyService : IMyContract
{
    [PrincipalPermission(SecurityAction.Demand,Role = "Manager")]
    public void MyMethod()
    {...}
}
```



The attribute will programmatically enable the role manager section in the config file.

The `SecurityBehavior` attribute allows the use of the `NetTcpBinding` with `ServiceSecurity.Internet` along with ASP.NET providers to allow intranet applications to avoid using Windows accounts and groups, as explained previously.

Next is the issue of supplying the application name for the ASP.NET providers. That is governed by the `ApplicationName` property. If no value is assigned, the `SecurityBehavior` attribute will look up the application name from the config file, as in Examples 10-13 and 10-14. If no value is found in the host config file, the attribute will not default to using the meaningless / from `machine.config`; instead, it will by default use the host assembly name for the application name. If the `ApplicationName` property is assigned a value, that value will override whatever application name is present in the host config file:

```
[SecurityBehavior(ServiceSecurity.Internet,"MyServiceCert",
                  UseAspNetProviders = true, ApplicationName = "MyApplication")]
class MyService : IMyContract
{...}
```

Configuring a business-to-business service

To configure a service for the business-to-business scenario, you must set `ServiceSecurity` to `ServiceSecurity.BusinessToBusiness`. The `SecurityBehavior` attribute will use peer trust for validating the client's certificate. Configuring the service certificate is done just as with `ServiceSecurity.Internet`. For example:

```
[SecurityBehavior(ServiceSecurity.BusinessToBusiness)]
class MyService : IMyContract
```

```

[...]
[SecurityBehavior(ServiceSecurity.BusinessToBusiness,"")]
class MyService : IMyContract
{...}

[SecurityBehavior(ServiceSecurity.BusinessToBusiness,"MyServiceCert")]
class MyService : IMyContract
{...}

```

By default, with `ServiceSecurity.BusinessToBusiness`, the attribute will set the `PrincipalPermissionMode` property of the host to `PrincipalPermissionMode.None`, and the service will not be able to authorize its callers. However, setting the `UseAspNetProviders` property to `true` will enable use of the ASP.NET role providers, as in [Example 10-18](#):

```

[SecurityBehavior(ServiceSecurity.BusinessToBusiness,UseAspNetProviders = true)]
class MyService : IMyContract
{...}

```

When using the ASP.NET role providers, the application name is looked up and decided upon just as with `ServiceSecurity.Internet`:

```

[SecurityBehavior(ServiceSecurity.BusinessToBusiness,"MyServiceCert",
                  UseAspNetProviders = true,ApplicationName = "MyApplication")]
class MyService : IMyContract
{...}

```

Configuring an anonymous service

To allow anonymous callers, you need to configure the attribute with `ServiceSecurity.Anonymous`. Configuring the service certificate is done just as with `ServiceSecurity.Internet`. For example:

```

[SecurityBehavior(ServiceSecurity.Anonymous)]
class MyService : IMyContract
{...}

[SecurityBehavior(ServiceSecurity.Anonymous,"")]
class MyService : IMyContract
{...}

[SecurityBehavior(ServiceSecurity.Anonymous,"MyServiceCert")]
class MyService : IMyContract
{...}

```

Configuring a no-security service

To turn off security completely, provide the attribute with `ServiceSecurity.None`:

```
[SecurityBehavior(ServiceSecurity.None)]
class MyService : IMyContract
{...}
```

Implementing the SecurityBehavior attribute

Example 10-20 is a partial listing of the implementation of `SecurityBehaviorAttribute`.

Example 10-20. Implementing SecurityBehaviorAttribute

```
[AttributeUsage(AttributeTargets.Class)]
class SecurityBehaviorAttribute : Attribute,IServiceBehavior
{
    readonly SecurityBehavior m_SecurityBehavior;

    public bool ImpersonateAll
    {get;set;}
    public string ApplicationName
    {get;set;}
    public bool UseAspNetProviders
    {get;set;}

    public SecurityBehaviorAttribute(ServiceSecurity mode)
    {
        m_SecurityBehavior = new SecurityBehavior(mode);
    }
    public SecurityBehaviorAttribute(ServiceSecurity mode,
                                    string serviceCertificateName)
    {
        m_SecurityBehavior = new SecurityBehavior(mode,serviceCertificateName);
    }

    void IServiceBehavior.AddBindingParameters(ServiceDescription description,
                                                ServiceHostBase serviceHostBase,
                                                Collection<ServiceEndpoint> endpoints,
                                                BindingParameterCollection parameters)
    {
        m_SecurityBehavior.AddBindingParameters(description,serviceHostBase,
                                                endpoints,parameters);
    }
    void IServiceBehavior.Validate(ServiceDescription description,
                                ServiceHostBase serviceHostBase)
    {
        m_SecurityBehavior.UseAspNetProviders = UseAspNetProviders;
        m_SecurityBehavior.ApplicationName = ApplicationName;
        m_SecurityBehavior.ImpersonateAll = ImpersonateAll;
        m_SecurityBehavior.Validate(description,serviceHostBase);
    }
    //Rest of the implementation
}
```

The `SecurityBehavior` attribute is a service behavior attribute, so you can apply it directly on the service class. When the `AddBindingParameters()` method of `IServiceBehavior` is called, the `SecurityBehavior` attribute enforces the binding configuration that matches the requested scenario. The `Validate()` method of `IServiceBehavior` is where the `SecurityBehavior` attribute configures the host. Other than that, all the attribute really does is sequence the overall order of configuration. The actual configuration is accomplished using a helper class called `SecurityBehavior`. Recall from other examples that it is always best to separate the attribute from its behavior, so you can reuse the behavior elsewhere. The `SecurityBehavior` attribute constructs an instance of `SecurityBehavior`, providing it with the scenario (the `mode` parameter) as well as the certificate name in the matching constructor. `SecurityBehavior` provides systematic, meticulous setting of all security scenarios using programmatic calls, encapsulating all the explicit steps described previously for each scenario. `SecurityBehavior` is a service behavior in its own right, and it is designed to even be used standalone, independent of the attribute. [Example 10-21](#) contains a partial listing of `SecurityBehavior`, demonstrating how it operates.

Example 10-21. Implementing SecurityBehavior (partial)

```
class SecurityBehavior : IServiceBehavior
{
    readonly ServiceSecurity m_Mode;
    readonly StoreLocation m_StoreLocation;
    readonly StoreName m_StoreName;
    readonly X509FindType m_FindType;
    readonly string m_SubjectName;

    public bool ImpersonateAll
    {get;set;}
    public bool UseAspNetProviders
    {get;set;}
    public string ApplicationName
    {get;set;}

    public SecurityBehavior(ServiceSecurity mode) :
        this(mode,StoreLocation.LocalMachine,X509FindType.FindBySubjectName,null)
    {}
    public SecurityBehavior(ServiceSecurity mode,StoreLocation storeLocation,
                           StoreName storeName,X509FindType findType,
                           string subjectName)
    {...} //Sets the corresponding members

    public void Validate(ServiceDescription description,
                        ServiceHostBase serviceHostBase)
    {
        if(m_SubjectName != null)
        {
```

```

switch(m_Mode)
{
    case ServiceSecurity.Anonymous:
    case ServiceSecurity.BusinessToBusiness:
    case ServiceSecurity.Internet:
    {
        string subjectName;
        if(m_SubjectName != String.Empty)
        {
            subjectName = m_SubjectName;
        }
        else
        {
            subjectName = description.Endpoints[0].Address.Uri.Host;
        }
        serviceHostBase.Credentials.ServiceCertificate.
            SetCertificate(m_StoreLocation,m_StoreName,m_FindType,subjectName);
        break;
    }
}
.
.
.

}

public void AddBindingParameters(ServiceDescription description,
                                ServiceHostBase serviceHostBase,
                                Collection<ServiceEndpoint> endpoints,
                                BindingParameterCollection parameters)
{
.
.
.

switch(m_Mode)
{
    case ServiceSecurity.Intranet:
    {
        ConfigureIntranet(endpoints);
        break;
    }
    case ServiceSecurity.Internet:
    {
        ConfigureInternet(endpoints,UseAspNetProviders);
        break;
    }
.
.
.

}
internal static void ConfigureInternet(IEnumerable<ServiceEndpoint> endpoints)
{

```

```

foreach(ServiceEndpoint endpoint in endpoints)
{
    Binding binding = endpoint.Binding;
    if(binding is WSHttpBinding)
    {

        WSHttpBinding wsBinding = (WSHttpBinding)binding;
        wsBinding.Security.Mode = SecurityMode.Message;
        wsBinding.Security.Message.ClientCredentialType =
            MessageCredentialType.UserName;
        continue;
    }
    .
    .
    .
    throw new InvalidOperationException(binding.GetType() +
        " is unsupported with ServiceSecurity.Internet");
}
//Rest of the implementation
}

```

The constructors of `SecurityBehavior` store in member variables the construction parameters, such as the security mode and the details of the certificate. The `Validate()` method is a decision tree that configures the host according to the scenario and the provided information, supporting the behavior of the `SecurityBehavior` attribute. `AddBindingParameters()` calls a dedicated helper method for each scenario to configure the collection of endpoints the host exposes. Each helper method (such as `ConfigureInternet()`) iterates over the collection of service endpoints. For each endpoint, it verifies whether the binding used matches the scenario and then configures the binding according to the scenario.

Host-Side Declarative Security

While configuring declarative security via the `SecurityBehavior` attribute is easy and handy, often it is up to the host to configure security, and the service just focuses on the business logic. In addition, you may be required to host services you do not develop, and those services may not happen to use my declarative security framework. The natural next step is to add declarative security support to the service host class as a set of `SetSecurityBehavior()` extension methods:

```

public static class SecurityHelper
{
    public static void SetSecurityBehavior(this ServiceHost host,
                                          ServiceSecurity mode,
                                          bool useAspNetProviders,
                                          string applicationName,
                                          bool impersonateAll = false);

```

```

public static void SetSecurityBehavior(this ServiceHost host,
                                      ServiceSecurity mode,
                                      string serviceCertificateName,
                                      bool useAspNetProviders,
                                      string applicationName,
                                      bool impersonateAll = false);

public static void SetSecurityBehavior(this ServiceHost host,
                                      ServiceSecurity mode,
                                      StoreLocation storeLocation,
                                      StoreName storeName,
                                      X509FindType findType,
                                      string serviceCertificateName,
                                      bool useAspNetProviders,
                                      string applicationName,
                                      bool impersonateAll = false);
//More members
}

```

Using declarative security via the host follows the same consistent guidelines as with the `SecurityBehavior` attribute. For example, here is how to configure the host (and the service) for Internet security with ASP.NET providers:

```

ServiceHost<MyService> host = new ServiceHost<MyService>();
host.SetSecurityBehavior(ServiceSecurity.Internet,
                        "MyServiceCert",true,"MyApplication");
host.Open();

```

Example 10-22 shows a partial listing of the declarative security support in the extension methods for the host.

Example 10-22. Adding declarative security extensions for the host

```

public static class SecurityHelper
{
    public static void SetSecurityBehavior(this ServiceHost host,
                                          ServiceSecurity mode,
                                          StoreLocation storeLocation,
                                          StoreName storeName,
                                          X509FindType findType,
                                          string serviceCertificateName,
                                          bool useAspNetProviders,
                                          string applicationName,
                                          bool impersonateAll = false)
    {
        if(host.State == CommunicationState.Opened)
        {
            throw new InvalidOperationException("Host is already opened");
        }
        SecurityBehavior securityBehavior = new SecurityBehavior(mode,storeLocation,
                                                               storeName,findType,
                                                               serviceCertificateName);
    }
}

```

```

        securityBehavior.UseAspNetProviders = useAspNetProviders;
        securityBehavior.ApplicationName = applicationName;
        securityBehavior.ImpersonateAll = impersonateAll;

        host.Description.Behaviors.Add(securityBehavior);
    }
    //More members
}

```

The implementation of `SetSecurityBehavior()` relies on the fact that the `SecurityBehavior` class supports `IServiceBehavior`. `SetSecurityBehavior()` initializes an instance of `SecurityBehavior` with the supplied parameters and then adds it to the collection of behaviors in the service description, as if the service were decorated with the `SecurityBehavior` attribute.

Client-Side Declarative Security

WCF does not allow attributes to be applied on the proxy class, and while a contract-level attribute is possible, the client may need to provide its credentials and other settings at runtime. The first step in supporting declarative security on the client side is my `SecurityHelper` static helper class with its set of extension methods for the proxy, defined in [Example 10-23](#).

Example 10-23. Adding declarative security extensions for the proxy

```

public static class SecurityHelper
{
    public static void UnsecuredProxy<T>(this ClientBase<T> proxy) where T : class;
    public static void AnonymousProxy<T>(this ClientBase<T> proxy) where T : class;
    public static void SecureProxy<T>(this ClientBase<T> proxy,
                                    string userName, string password) where T : class;
    public static void SecureProxy<T>(this ClientBase<T> proxy,
                                    string domain, string userName, string password) where T : class;
    public static void SecureProxy<T>(this ClientBase<T> proxy, string domain,
                                    string userName, string password, TokenImpersonationLevel impersonationLevel)
                                    where T : class;
    public static void SecureProxy<T>(this ClientBase<T> proxy,
                                    string clientCertificateName) where T : class;
    public static void SecureProxy<T>(this ClientBase<T> proxy,
                                    StoreLocation storeLocation, StoreName storeName,
                                    X509FindType findType, string clientCertificateName) where T : class;
    //More members
}

```

You can use `SecurityHelper` to configure a plain proxy according to the desired security scenario and behavior, using the dedicated static extension methods `SecurityHelper` offers. You can configure the proxy only before opening it. There is no need for any security settings in the client's config file or elsewhere in the client's code.

`SecurityHelper` is smart, and it will select the correct security behavior based on provided parameters and the method invoked. There is no need to explicitly use the `ServiceSecurity` enum.

For example, here is how to secure a proxy for the intranet scenario and provide it with the client's Windows credentials:

```
MyContractClient proxy = new MyContractClient();
proxy.SecureProxy("MyDomain", "MyUsername", "MyPassword");
proxy.MyMethod();
proxy.Close();
```

For the Internet scenario, the client only needs to provide the username and the password (remember that the decision as to whether those are Windows or ASP.NET provider credentials is a service-side decision):

```
MyContractClient proxy = new MyContractClient();
proxy.SecureProxy("MyUsername", "MyPassword");
proxy.MyMethod();
proxy.Close();
```

For the business-to-business scenario, the client can specify a `null` or an empty string for the client certificate name if it wants to use the certificate in its config file, or it can list the certificate name explicitly:

```
MyContractClient proxy = new MyContractClient();
proxy.SecureProxy("MyClientCert");
proxy.MyMethod();
proxy.Close();
```

`SecurityHelper` will load the certificate from the client's `LocalMachine` store from the `My` folder by name. The client can also specify all the information required to find and load the certificate. To keep the design of `SecurityHelper` simple, when using the `BasicHttpBinding` in the business-to-business scenario, the client must explicitly specify the service certificate location, either in the config file or programmatically.

For an anonymous client, use the `AnonymousProxy()` method:

```
MyContractClient proxy = new MyContractClient();
proxy.AnonymousProxy();
proxy.MyMethod();
proxy.Close();
```

and for no security at all, use the `UnsecuredProxy()` method:

```
MyContractClient proxy = new MyContractClient();
proxy.UnsecuredProxy();
proxy.MyMethod();
proxy.Close();
```

Implementing SecurityHelper

Internally, `SecurityHelper` uses `SecurityBehavior` to configure the proxy's endpoint and set the credentials, as shown in [Example 10-24](#).

Example 10-24. Implementing SecurityHelper (partial)

```
public static class SecurityHelper
{
    public static void SecureProxy<T>(this ClientBase<T> proxy,
                                      string userName, string password) where T : class
    {
        if(proxy.State == CommunicationState.Opened)
        {
            throw new InvalidOperationException("Proxy channel is already opened");
        }

        ServiceEndpoint[] endpoints = {factory.Endpoint};

        SecurityBehavior.ConfigureInternet(endpoints, true);

        proxy.ClientCredentials.UserName.UserName = userName;
        proxy.ClientCredentials.UserName.Password = password;
        proxy.ClientCredentials.ServiceCertificate.Authentication.
            CertificateValidationMode = X509CertificateValidationMode.PeerTrust;
    }
    //Rest of the implementation
}
```

The `SecureClientBase<T>` class

The advantage of using the `SecurityHelper` extensions is that they can operate on any proxy—even a proxy the client developer is not responsible for creating. The disadvantage is that it is an extra step the client has to take. If you are responsible for generating the proxy, you can take advantage of my `SecureClientBase<T>` class, defined in [Example 10-25](#).

Example 10-25. The `SecureClientBase<T>` class

```
public abstract class SecureClientBase<T> : ClientBase<T> where T : class
{
    //These constructors target the default endpoint
    protected SecureClientBase();
    protected SecureClientBase(ServiceSecurity mode);
    protected SecureClientBase(string userName, string password);
    protected SecureClientBase(string domain, string userName, string password,
                              TokenImpersonationLevel impersonationLevel);
    protected SecureClientBase(string domain, string userName, string password);
    protected SecureClientBase(string clientCertificateName);
    protected SecureClientBase(StoreLocation storeLocation, StoreName storeName,
```

```

        X509FindType findType, string clientCertificateName);
    //More constructors for other types of endpoints
}

```

`SecureClientBase<T>` derives from the conventional `ClientBase<T>` and adds declarative security support. You need to derive your proxy from `SecureClientBase<T>` instead of `ClientBase<T>`, provide constructors that match your security scenario, and call the base constructors of `SecureClientBase<T>` with the supplied credentials and endpoint information:

```

class MyContractClient : SecureClientBase<IMyContract>, IMyContract
{
    public MyContractClient(ServiceSecurity mode) : base(mode)
    {}
    public MyContractClient(string userName, string password) : base
        (userName, password)
    {}

    /* More constructors */

    public void MyMethod()
    {
        Channel.MyMethod();
    }
}

```

Using the derived proxy is straightforward. For example, for the Internet scenario:

```

MyContractClient proxy = new MyContractClient("MyUsername", "MyPassword");
proxy.MyMethod();
proxy.Close();

```

or for the Anonymous scenario:

```

MyContractClient proxy = new MyContractClient(ServiceSecurity.Anonymous);
proxy.MyMethod();
proxy.Close();

```

The implementation of `SecureClientBase<T>` simply uses the extensions of `SecurityHelper` (as shown in [Example 10-26](#)), so `SecureClientBase<T>` follows the same behaviors regarding, for example, the client certificate.

Example 10-26. Implementing `SecureClientBase<T>` (partial)

```

public class SecureClientBase<T> : ClientBase<T> where T : class
{
    protected SecureClientBase(ServiceSecurity mode)
    {
        switch(mode)
        {
            case ServiceSecurity.None:
            {

```

```

        this.UnsecuredProxy();
        break;
    }
    case ServiceSecurity.Anonymous:
    {
        this.AnonymousProxy();
        break;
    }
    ...
}
protected SecureClientBase(string userName,string password)
{
    this.SecureProxy(userName,password);
}
//More constructors
}

```

Secure channel factory

If you are not using a proxy at all, `SecurityHelper` and `SecureClientBase<T>` will be of little use to you. For that case, I added to `SecurityHelper` a set of extension methods to the `ChannelFactory<T>` class, defined in [Example 10-27](#).

Example 10-27. Adding declarative security extensions for ChannelFactory<T>

```

public static class SecurityHelper
{
    public static void SetSecurityMode<T>(this ChannelFactory<T> factory,
                                         ServiceSecurity mode);
    public static void SetCredentials<T>(this ChannelFactory<T> factory,
                                         string userName,string password);
    public static void SetCredentials<T>(this ChannelFactory<T> factory,
                                         string domain,string userName,string password);
    public static void SetCredentials<T>(this ChannelFactory<T> factory,
                                         string clientCertificateName);
    //More members
}

```

You need to call the `SetSecurityMode()` method or one of the `SetCredentials()` methods that fits your target scenario before opening the channel. For example, with a proxy to an Internet security-based service:

```

ChannelFactory<IMyContract> factory = new ChannelFactory<IMyContract>( . . . );
factory.SetCredentials("MyUsername","MyPassword");

IMyContract proxy = factory.CreateChannel();

using(proxy as IDisposable)
{

```

```
        proxy.MyMethod();
    }
```

Implementing the `ChannelFactory<T>` extensions was very similar to implementing the extensions for `ClientBase<T>`, so I have omitted that code.

Duplex clients and declarative security

I also provide the `SecureDuplexClientBase<T,C>` class (similar to `SecureClientBase<T>`), which is defined in [Example 10-28](#).

Example 10-28. The `SecureDuplexClientBase<T,C>` class

```
public abstract class SecureDuplexClientBase<T,C> : DuplexClientBase<T,C>
{
    protected SecureDuplexClientBase(C callback);
    protected SecureDuplexClientBase(ServiceSecurity mode,C callback);
    protected SecureDuplexClientBase(string userName,string password,C callback);
    protected SecureDuplexClientBase(string domain,string userName,string password,
                                    TokenImpersonationLevel impersonationLevel,C callback);
    protected SecureDuplexClientBase(string domain,string userName,string password,
                                    C callback);
    protected SecureDuplexClientBase(string clientCertificateName,C callback);
    protected SecureDuplexClientBase(StoreLocation storeLocation,
                                    StoreName storeName,X509FindType findType,
                                    string clientCertificateName,C callback);

    /* More constructors with InstanceContext<C> and constructors that
       target the configured endpoint and a programmatic endpoint */
}
```

`SecureDuplexClientBase<T,C>` derives from my type-safe `DuplexClientBase<T,C>` class, presented in [Chapter 5](#), and it adds declarative scenario-based security support. As when using the `DuplexClientBase<T,C>` class, you need to derive your proxy class from it and take advantage of either the callback parameter or the type-safe context `InstanceContext<C>`. For example, given this service contract and callback contract definition:

```
[ServiceContract(CallbackContract = typeof(IMyContractCallback))]
interface IMyContract
{
    [OperationContract]
    void MyMethod();
}
interface IMyContractCallback
{
    [OperationContract]
    void OnCallback();
}
```

your derived proxy class will look like this:

```
class MyContractClient : SecureDuplexClientBase<IMyContract, IMyContractCallback>,  
    IMyContract  
{  
    public MyContractClient(IMyContractCallback callback) : base(callback)  
    {}  
    public MyContractClient(ServiceSecurity mode, IMyContractCallback callback) :  
        base(mode, callback)  
    {}  
    /* More constructors */  
  
    public void MyMethod()  
    {  
        Channel.MyMethod();  
    }  
}
```

When using `SecureDuplexClientBase<T,C>`, provide the security scenario or credentials, the callback object, and the endpoint information. For example, when targeting the anonymous scenario:

```
class MyClient : IMyContractCallback  
{...}  
  
IMyContractCallback callback = new MyClient();  
  
MyContractClient proxy = new MyContractClient(ServiceSecurity.Anonymous,  
                                              callback);  
proxy.MyMethod();  
  
proxy.Close();
```

The implementation of `SecureDuplexClientBase<T,C>` is almost identical to that of `SecureClientBase<T>`, with the main difference being a different base class. Note that there was no point in defining declarative extensions for `DuplexClientBase<T>`, since you should not use it in the first place due to its lack of type safety.

Extensions for the duplex factory

When you're not using a `SecureDuplexClientBase<T,C>`-derived proxy to set up the bidirectional communication, you can use my declarative extensions for the `DuplexChannelFactory<T,C>` channel factory, defined in [Example 10-29](#).

Example 10-29. Adding declarative security extensions for DuplexChannelFactory<T,C>

```
public static class SecurityHelper  
{  
    public static void SetSecurityMode<T,C>(this DuplexChannelFactory<T,C> factory,  
                                             ServiceSecurity mode);
```

```

public static void SetCredentials<T,C>(this DuplexChannelFactory<T,C> factory,
                                         string userName, string password);
public static void SetCredentials<T,C>(this
                                         DuplexChannelFactory<T,C> factory,
                                         string domain, string userName, string password);
public static void SetCredentials<T,C>(this DuplexChannelFactory<T,C> factory,
                                         string clientCertificateName);
//More members
}

```

You need to call the `SetSecurityMode()` method or one of the `SetCredentials()` methods that fits your target scenario before opening the channel. For example, when targeting the Internet scenario:

```

class MyClient : IMyContractCallback
{...}

IMyContractCallback callback = new MyClient();

DuplexChannelFactory<IMyContract, IMyContractCallback> factory =
    new DuplexChannelFactory<IMyContract, IMyContractCallback>
        (callback, "");
factory.SetCredentials("MyUsername", "MyPassword");

IMyContract proxy = factory.CreateChannel();
using(proxy as IDisposable)
{
    proxy.MyMethod();
}

```

Implementing the extensions for `DuplexChannelFactory<T,C>` was very similar to implementing those for `ChannelFactory<T>`.

Security Auditing

I will end this chapter by presenting a useful feature WCF supports called *security audits*. As its name implies, a security audit is a logbook of the security-related events in your services. WCF can log authentication and authorization attempts, their times and locations, and the calling clients' identities. The class `ServiceSecurityAuditBehavior` governs auditing; it is listed in [Example 10-30](#) along with its supporting enumerations.

Example 10-30. The ServiceSecurityAuditBehavior class

```

public enum AuditLogLocation
{
    Default, //Decided by the operating system
    Application,
    Security
}

```

```

}
public enum AuditLevel
{
    None,
    Success,
    Failure,
    SuccessOrFailure
}
public sealed class ServiceSecurityAuditBehavior : IServiceBehavior
{
    public AuditLogLocation AuditLogLocation
    {get;set;}
    public AuditLevel MessageAuthenticationAuditLevel
    {get;set;}
    public AuditLevel ServiceAuthorizationAuditLevel
    {get;set;}
    //More members
}

```

`ServiceSecurityAuditBehavior` is a service behavior. The `AuditLogLocation` property specifies where to store the log entries: in the application log or in the security log, *both of which* are in the event log on the host computer. The `MessageAuthenticationAuditLevel` property governs the authentication audit verbosity. Its default value is `AuditLevel.None`. For performance's sake, you may want to audit only failures. For diagnostic purposes, you can also audit successful authentications. Similarly, you use the `ServiceAuthorizationAuditLevel` property to control authorization audit verbosity. It is also disabled by default.

Configuring Security Audits

The typical way of enabling a security audit is in the host config file, by adding a custom behavior section and referencing it in the service declaration, as shown in [Example 10-31](#).

Example 10-31. Configuring a security audit administratively

```

<system.serviceModel>
    <services>
        <service name = "MyService" behaviorConfiguration = "MySecurityAudit">
            ...
        </service>
    </services>
    <behaviors>
        <serviceBehaviors>
            <behavior name = "MySecurityAudit">
                <serviceSecurityAudit
                    auditLogLocation = "Default"
                    serviceAuthorizationAuditLevel = "SuccessOrFailure"
                    messageAuthenticationAuditLevel = "SuccessOrFailure"

```

```
        />
    </behavior>
</serviceBehaviors>
</behaviors>
</system.serviceModel>
```

You can also configure security auditing programmatically, by adding the behavior to the host at runtime before opening it. As you do when adding other behaviors programmatically, you can check that the host does not already have an audit behavior to avoid overriding the config file, as shown in [Example 10-32](#).

Example 10-32. Enabling a security audit programmatically

```
ServiceHost host = new ServiceHost(typeof(MyService));

ServiceSecurityAuditBehavior securityAudit =
    host.Description.Behaviors.Find<ServiceSecurityAuditBehavior>();
if(securityAudit == null)
{
    securityAudit = new ServiceSecurityAuditBehavior();

    securityAudit.MessageAuthenticationAuditLevel = AuditLevel.SuccessOrFailure;
    securityAudit.ServiceAuthorizationAuditLevel = AuditLevel.SuccessOrFailure;
    host.Description.Behaviors.Add(securityAudit);
}
host.Open();
```

You can streamline the code in [Example 10-32](#) by adding the `SecurityAuditEnabled` Boolean property to `ServiceHost<T>`:

```
public class ServiceHost<T> : ServiceHost
{
    public bool SecurityAuditEnabled
    {get;set;}
    //More members
}
```

Using `ServiceHost<T>`, [Example 10-32](#) is reduced to:

```
ServiceHost<MyService> host = new ServiceHost<MyService>();
host.SecurityAuditEnabled = true;
host.Open();
```

[Example 10-33](#) shows the implementation of the `SecurityAuditEnabled` property.

Example 10-33. Implementing the `SecurityAuditEnabled` property

```
public class ServiceHost<T> : ServiceHost
{
    public bool SecurityAuditEnabled
    {
```

```

get
{
    ServiceSecurityAuditBehavior securityAudit =
        Description.Behaviors.Find<ServiceSecurityAuditBehavior>();
    if(securityAudit != null)
    {
        return securityAudit.MessageAuthenticationAuditLevel ==
            AuditLevel.SuccessOrFailure
            &&
            securityAudit.ServiceAuthorizationAuditLevel ==
            AuditLevel.SuccessOrFailure;
    }
    else
    {
        return false;
    }
}
set
{
    if(State == CommunicationState.Opened)
    {
        throw new InvalidOperationException("Host is already opened");
    }
    ServiceSecurityAuditBehavior securityAudit =
        Description.Behaviors.Find<ServiceSecurityAuditBehavior>();
    if(securityAudit == null && value == true)
    {
        securityAudit = new ServiceSecurityAuditBehavior();
        securityAudit.MessageAuthenticationAuditLevel =
            AuditLevel.SuccessOrFailure;
        securityAudit.ServiceAuthorizationAuditLevel =
            AuditLevel.SuccessOrFailure;
        Description.Behaviors.Add(securityAudit);
    }
}
}
//More members
}

```

In the `get` accessor, the `SecurityAuditEnabled` property accesses the description of the service and looks for an instance of `ServiceSecurityAuditBehavior`. If one is found, and if both the authentication and the authorization audits are set to `AuditLevel.SuccessOrFailure`, `SecurityAuditEnabled` returns `true`; otherwise, it returns `false`. In the `set` accessor, the property enables the security audit only if the description does not contain a previous value (because the config file does not contain the audit behavior). If no prior behavior is found, `SecurityAuditEnabled` sets both the authentication and authorization audits to `AuditLevel.SuccessOrFailure`.

Declarative Security Auditing

You can also write an attribute that surfaces the security audit options at the service level. I chose to add that support in the form of a single Boolean property of the `SecurityBehavior` attribute called `SecurityAuditEnabled`:

```
[AttributeUsage(AttributeTargets.Class)]
public class SecurityBehaviorAttribute : Attribute,IServiceBehavior
{
    public bool SecurityAuditEnabled
    {get;set;}
    //More members
}
```

The default of `SecurityAuditEnabled` is `false` (i.e., no security audit). Using this property complements the rest of the declarative security model. For example:

```
[SecurityBehavior(ServiceSecurity.Internet,UseAspNetProviders = true,
                  SecurityAuditEnabled = true)]
class MyService : IMyContract
{...}
```

Example 10-34 shows how that support was added to the `SecurityBehavior` attribute.

Example 10-34. Implementing a declarative security audit

```
[AttributeUsage(AttributeTargets.Class)]
public class SecurityBehaviorAttribute : Attribute,IServiceBehavior
{
    public bool SecurityAuditEnabled
    {get;set;}

    void IServiceBehavior.Validate(ServiceDescription description,
                                    ServiceHostBase serviceHostBase)
    {
        if(SecurityAuditEnabled)
        {
            ServiceSecurityAuditBehavior securityAudit =
                serviceHostBase.Description.Behaviors.Find<ServiceSecurityAuditBehavior>();
            if(securityAudit == null)
            {
                securityAudit = new ServiceSecurityAuditBehavior();
                securityAudit.AuditLogLocation = AuditLogLocation.Application;
                securityAudit.MessageAuthenticationAuditLevel =
                    AuditLevel.SuccessOrFailure;
                securityAudit.ServiceAuthorizationAuditLevel =
                    AuditLevel.SuccessOrFailure;
                serviceHostBase.Description.Behaviors.Add(securityAudit);
            }
            //Rest same as Example 10-20
        }
    }
}
```

```
    }
    //Rest of the implementation
}
```

The `Validate()` method of `IServiceBehavior` enables auditing using the same verbosity level as `ServiceHost<T>`, again avoiding overriding the config file.

The Azure Service Fabric

The *Windows Azure Service Fabric* (or just the *Service Fabric*) represents nothing less than the next evolutionary breakthrough in software development. With the advent of the Service Fabric, Microsoft makes public the long anticipated service-oriented platform upon which it has built and refined significant pieces of Microsoft Azure. As you will see, Service Fabric also delivers on nearly all of the rationale in the *What's Next* section of [Appendix A](#).

Most significantly, the Service Fabric is a uniquely integrated platform. The Service Fabric defines a prescribed approach to service-oriented system programming and deployment with models for both development and operations that are in line with current industry best practice. This consistent, end-to-end approach to system management reduces the many programming, connectivity, and administration headaches common to service-oriented systems in the wild. This makes Azure Service Fabric *the killer app* for modern DevOps and modern software systems both in the Cloud and on-premises. This chapter starts with an overview of what the Service Fabric has to offer and the rationale behind it, and then moves on to discuss the Service Fabric's service-oriented programming model. The chapter concludes by presenting how to best prepare your systems for it and even how to start leveraging its programming model today, in a forward-compatible manner to affect a smooth transition.

Why Azure Service Fabric

Azure Service Fabric offers many benefits to modern software developers and software development organizations alike. Its integrated platform provides continuity to your overall system delivery experience by streamlining system lifecycle management across both development and operations. This ensures that deployment-related activities and decisions do not fall between the cracks and that the overall experience of development and deployment has a high degree of design integrity.

The Service Fabric's integrated technology stack and simple programming model bring consistency to a varied and ever-changing technology landscape, allowing you to design, build, test, and consume a variety of different service types in a single environment. The Service Fabric also helps you address the demise of Moore's Law, which demands that you now use scalable, parallel programming techniques to improve the responsiveness and performance of your software systems. To do this effectively, you need an environment built to support such patterns and a programming paradigm that simplifies their expression. The Service Fabric is what Microsoft used internally to build Windows Azure itself, and is the single most important software technology factor behind the success of that platform. The Service Fabric allowed Microsoft to build Azure in a fraction of the time and cost it would have taken using the prevailing tools at the time, and it allowed for levels of scalability, elasticity, and performance that were practically out of reach before it. But most importantly, the Service Fabric is what unleashed Microsoft's creativity, because it made extending and innovating so much faster than before.

An Integrated Platform

As an integrated platform, Azure Service Fabric helps you address the problems that often plague service-oriented system delivery. The process of service-oriented system deployment often lacks clarity in the shared responsibilities between the system development and operation. Developers may design and implement the system in such a way that makes it difficult for operations to scale the system or ensure its availability or manage it efficiently. The Service Fabric allows you to introduce consistency, repeatability, predictability, and standardization across your system's lifecycle for all the roles involved, from development through deployment to administration and eventual decommissioning.

Design Consistency

Often due to insufficient technical leadership, service-oriented systems suffer from inconsistent or even competing design and implementation approaches. Each team may define a different structure for their services or implement them using different approaches for hosting, security, concurrency, instancing, transactions or even defining service behavior. These inconsistencies eventually introduce internal integration problems that in turn lengthen development cycles and ultimately erode system quality. In contrast, the Service Fabric introduces a single system structure. This inherently promotes design and implementation consistency and integrity across all teams, even operations.

Programming Model Consistency

In the current technology landscape there are numerous ways to establish service interaction such as Uri-based, contract-based, or raw message-based interactions. There are also numerous ways to establish service connectivity such as HTTP, TCP, IPC, or queuing, each with its own unique programming model and implications. Again, this often introduces inconsistent or even competing approaches to service connectivity within a given system requiring developers to learn and master multiple disparate service programming models. Instead, the Service Fabric establishes a consistent, constrained, and repeatable service programming model for each service type that is informed by Microsoft's experience in building their Windows Azure Cloud platform using it. This helps you to remove much of the friction developers experience in their day-to-day work by reducing the number of connectivity models they must master as they design, build, test, and consume services. These more clearly defined connectivity choices also guide them in selecting the appropriate mode of interaction for a given service based on its role in the system.

Future Compatible

Most of all, Azure Service Fabric allows you to centralize under a single platform, technology stack, and programming model your most valuable software assets: the hard-won domain expertise that you have captured in your business logic. The Service Fabric enables you to insulate these key software assets from disruptive industry pressures and changes in technology by enforcing a structure on your systems and your code that can transition directly to the new world of actor-based architectures that is just now becoming viable.

You must first recognize the value behind the Actor Model. It represents a straightforward programming model that alleviates the challenges common to the most difficult programming task—that is, writing concurrent programs. The Actor Model is the only way to reason about highly parallel programming problems. You must then prepare your systems and your code for the age of pervasive compute, where there will be an actor for every *thing* on the Internet of Things (IoT). This will require new parallel programming techniques for the ingestion, aggregation, processing, and interpretation of continual, unsolicited data streams. This brave new world also signals “the end of the cylinder,” your database, where you will replace it with meshes of smart networks and behavior-less actors, as shown in [Figure 11-1](#). [Figure 11-1](#) shows a simple sorting network that will sort eight values in five steps. The network is built of simple stateful actors that contain but a single behavior; a `CompareAndExchange()` operation. You can build complex sorting meshes using basic, simple building blocks such as these allowing for sublinear sorting performance.

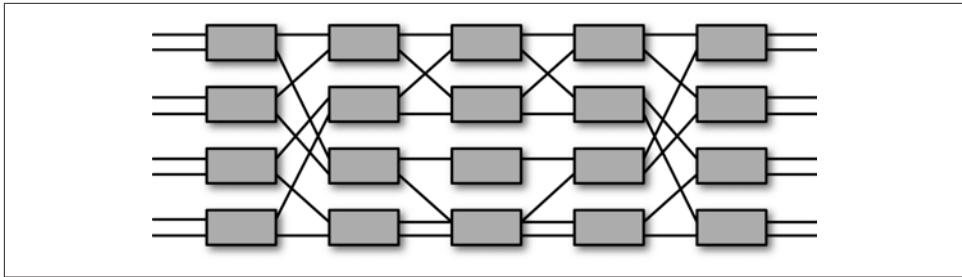


Figure 11-1. Mesh of actors

Or you can model your database as a graph of actors, as shown in [Figure 11-2](#). [Figure 11-2](#) shows a graph of stateful actors with arbitrary, directed relationships shown as *vectors* or *edges*. Each edge holds *properties* about that specific relationship between the associated actors. You query the graph by creating a *visitor* that possesses a specific filter or query pattern. The visitor then traverses the network, selecting relationships based on its filter that matches against the edge's properties. Although relationships are directional, the visitor can still navigate them regardless of direction. For each edge that satisfies its filter, the visitor visits the associated actor node and emits an event with that actor's data. A stateful aggregator actor then listens for visitor events and records the data. You can then obtain the result of your query directly from the aggregator actor. What a graph of actors allows you to do is have a concurrent set of visitors querying your database in parallel supporting mass scalability and high throughput.

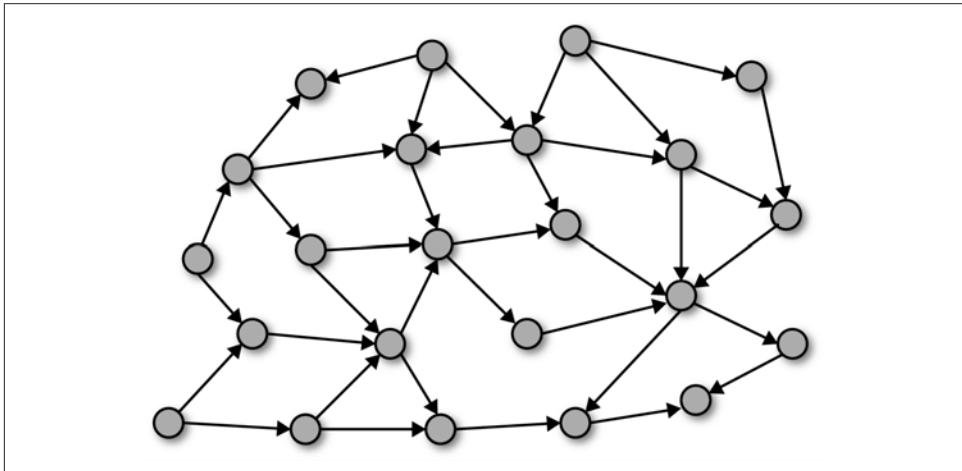


Figure 11-2. Graph of actors

As an industry we will have to adopt these mesh-based parallel computing models because it will be the only way to process in a reasonable way (ideally a way that is

invariant to the size of the data set) the torrents of data coming out of the IoT. But the same holds true for any future business that will need to process the volumes of data, users, devices, and programs our much more connected and integrated world will produce. The Actor Model enables you to model in software (over many cores) processing models that were reserved in the past for high-end, dedicated switching hardware. While you may not process billions of devices or users tomorrow, the support for the Actor Model in a present-day framework is the gateway to the future. You can start today to prepare your code and your business for what is coming.

Preparing for the Service Fabric

To seize the opportunity that this powerful new platform represents, you must first prepare your existing systems for a smooth transition to the Service Fabric environment. While readying your systems, it will be important that you avoid big bang approaches and instead perform your preparation one service or system area at a time. It will also be highly advantageous to test the services you have prepared in your current production environment before transitioning them to the Azure Service Fabric. It will be some time before Azure Service Fabric matures for full production use, particularly for on-premises scenarios. And unless you have developed your service-oriented systems using `InProcFactory` and the helpers from `ServiceModelEx`, Service Fabric's system structure and programming model represent a significant departure from the way you structure your code and your systems today. To prepare, you must start writing forward-compatible code today.

ServiceModelEx.ServiceFabric

To help you start preparing now, this chapter presents our `ServiceModelEx.ServiceFabric`. `ServiceModelEx.ServiceFabric` is a ready-made framework we developed that is compatible in both form and function with the essential aspects of Azure Service Fabric's programming model. `ServiceModelEx.ServiceFabric` is production ready. Its design uses the same proven service-oriented techniques presented throughout this book. It is built atop the mature code of `ServiceModelEx`, which numerous service-oriented systems use in production every day. To facilitate a smooth transition, `ServiceModelEx.ServiceFabric` provides a simple transition strategy: create a solution targeting the Azure Service Fabric, transfer your Service Fabric ready code to the solution, recompile and deploy.

`ServiceModelEx.ServiceFabric` also enables you to broaden your investment in mastering the Azure Service Fabric. `ServiceModelEx.ServiceFabric` extends the Service Fabric's programming model beyond the Windows Server environment to all Windows desktops and devices. You can consider `ServiceModelEx.ServiceFabric` a *microservice fabric*, enabling the programming model anywhere WCF runs. Through `ServiceModelEx.ServiceFabric` you can even contemplate creating a *nano service*.

fabric for use on small device platforms by using `ServiceModelEx.ServiceFabric` as a model for development on Windows 10 IoT Core or even on a sensor itself using the .NET Micro Framework. Creating continuity in your programming experience across all environments, nodes, and devices will become an essential necessity to building quality-oriented, cost-effective systems in the emerging world of the *Internet of Things* (IoT). This is the only development model that will scale with the expanding needs of the IoT. Developers must be able to learn a single service-oriented programming model and apply their skills across all system layers. Nothing else will work or will make economic sense.

Finally, you can use `ServiceModelEx.ServiceFabric` to lower the bar of entry for developers to learn the Service Fabric by reducing the amount of effort involved in working with it. With `ServiceModelEx.ServiceFabric`, there is nothing to install or maintain. `ServiceModelEx.ServiceFabric` represents a lightweight, fully testable Service Fabric perfect for development desktops, builder servers, and marketing laptops.

Configuring the Service Fabric

A full exploration of the Azure Service Fabric is well beyond the scope of this book. Nonetheless, to understand how to leverage the Service Fabric, you must first have a basic understanding of its capabilities and system structure. This involves both operations and development.

IT Operations

For IT operations, Azure Service Fabric represents a modern service-oriented infrastructure with prescribed topology, operations flow, lifecycle management, tooling, and system structure. The Service Fabric automates system deployment, hosting, upgrading, scaling, decommissioning, failover, self-healing, workload balancing, address resolution and discovery, security, state replication, and state persistence.

The Service Fabric provides a modern, cluster-based topology. Each node is a separate machine running Windows Server onto which you have installed the Service Fabric. All nodes are aware of each other through automatic endpoint address configuration and connect to each other upon the startup of each Service Fabric node, as shown in [Figure 11-3](#). It is important to note that the Service Fabric does not provide its many capabilities (as listed before) through a single fabric service process. Instead, each fabric capability is a separate fabric service and process deployed to all nodes that is also built using the Service Fabric. This extends the same benefits listed earlier to all the services that comprise the Service Fabric itself, including state management.

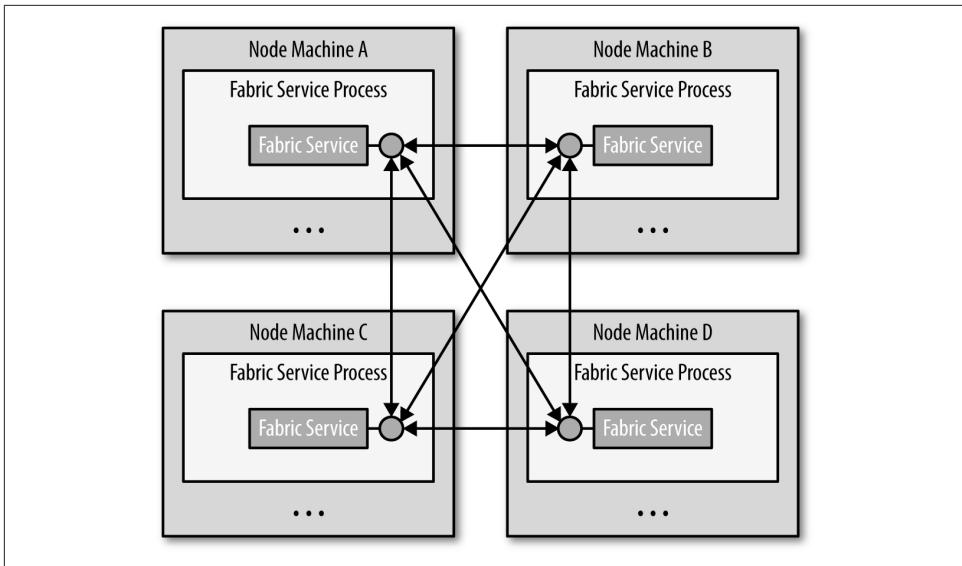


Figure 11-3. Service Fabric cluster topology



To maximize the Service Fabric's reliability, in a production cluster each node represents either a physical, or more likely, a virtual machine. The Service Fabric's single machine development deployment is the only configuration that allows multiple nodes to exist on a single machine.

The Service Fabric's cluster of nodes ensures high availability, density, and scalability. It achieves availability by providing resiliency to node failures through service redundancy and state replication. This allows the Service Fabric to self-heal by replacing or rebuilding failed nodes with copies of Service Fabric code, your service code and state from other nodes. The Service Fabric achieves scalability by supporting near transparent state management to manage the state of expensive instances, which as [Chapter 4](#) explains is paramount in supporting a large number of greedy clients. This allows the Service Fabric to automatically redistribute service code and state as you add or remove nodes.

It promotes high density, which is the efficient use of cluster resources by dynamically balancing its workload across all available nodes. Service Fabric's topology also sets the stage for it to eventually evolve into an expansive computing network, a mesh of nodes all working on a single program that the entire mesh executes.

Service Fabric is fully instrumented, providing transparency into every aspect of operations and enabling automated health monitoring and corrective action. You can query all health metrics through one of Service Fabric's APIs or view them directly

through the Service Fabric Explorer. As shown in **Figure 11-4**, Service Fabric Explorer is currently a somewhat crude tool that will evolve significantly as Service Fabric matures.

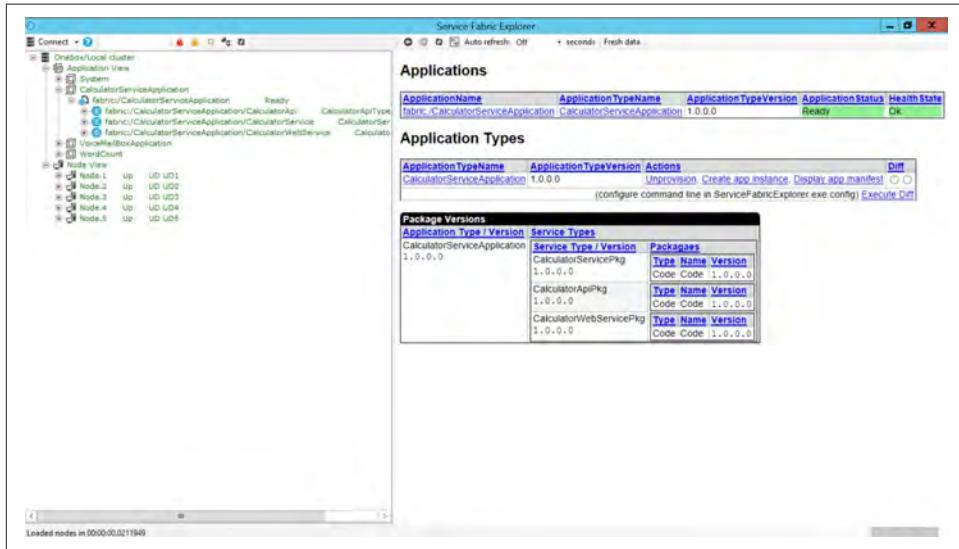


Figure 11-4. The Service Fabric Explorer

Service Fabric presents a common, Cloud-based flow to the day-to-day IT operations of packaging, uploading, registering, provisioning, validating, and managing system elements. It also supports the modern DevOps approach to system management by empowering multiple roles to easily participate across the full spectrum of lifecycle management including the deployment, testing, upgrading, maintenance, and decommissioning of systems. Finally, Service Fabric provides a full array of tooling options suitable to any role participating in the system lifecycle. From command line PowerShell scripts for automation to programmatic control through fully capable REST and C# APIs to automated operations via the Azure Management Portal or Visual Studio.

Development

The Service Fabric's programming model clarifies many of the common service-oriented design decisions you must make by defining the available modes of service interaction and enforcing where, when, and how you use them in service and system composition. In line with modern software architecture, Service Fabric provides three modes of service interaction. It provides an *API* interaction mode through ASP.NET Web API that offers interoperable interaction and a programming model compatible with web and mobile clients. It provides an *Intranet* interaction mode through WCF that offers efficient, internal interaction and a programming model compatible with

backend .NET server and client environments. The Service Fabric also provides a *Component* interaction mode through a new *actor* framework. Analogous to InProc Factory services, Service Fabric's actor framework completely hides actor connectivity concerns and enables you to create more granular services. Though this may change in the future, it is relevant to note that Service Fabric's actor framework is actually built atop WCF and uses a prescribed NetTcpBinding configuration internally.

What Is an Actor?

An actor is a specialization of a service, preferably granular in intent that abides by a specific set of design rules:

1. Actors may only interact with other actors asynchronously through immutable message exchange.
2. Actors must provide an explicit, cohesive, behavioral boundary through an interface-based programming construct.
3. Actors must always internalize their state, if any. Only by receiving and processing messages may an actor initiate a change in its state.
4. Actors must process messages one at a time to ensure the consistency of their state. This means actors are single-threaded, and every actor operation runs to completion before another operation may start. Actor frameworks call this approach *turn-based concurrency*.
5. When performing an operation, an actor may send asynchronously a finite number of messages to other actors, create a finite number of new actors, and optionally designate the behavior it will use for the next message it processes.

This constrained programming model makes actors ideal for crafting network-based parallel programming patterns. While the Actor Model is not a novel or modern idea (it dates back to 1973), it remained out of reach until the emergence of modern platforms such as Service Fabric or Akka.

In Service Fabric, service integration testing is a first-class citizen. Service Fabric provides a fully capable framework for service testing and system integration testing named *Testability*. Testability contains an API of predetermined actions that inject faults into the Service Fabric while your services are active. This allows you to test numerous fault, failover, and self-healing concerns for your services. You can also compose individual test actions into broader scenarios that Service Fabric will orchestrate for you. You can define, run, and automate Testability tests through PowerShell scripts or a fully capable C# API.

Both

For both development and operations, Service Fabric establishes a common system structure and unit of deployment for system management. Development uses this common structure to describe the elements of a system's composition. Development then shares this system description with IT operations through simple XML files called *manifests*. A manifest file contains basic descriptive information like name and version associated with one or more packages. A *package* is nothing more than a collection of one or more related files in a specific folder layout. Service Fabric manages a system's elements as a collection of packages based on the information in the associated manifest files. Service Fabric tooling automates the process of validating package layout, reading and validating manifest files, assembling related files into packages, and deploying the resulting packages into a cluster. Ultimately, Service Fabric hosts system element packages in one or more Windows processes.

System structure

Azure Service Fabric's unit of deployment is the *application*. Each application you deploy to a cluster is a collection of one or more services, as shown in [Figure 11-5](#).

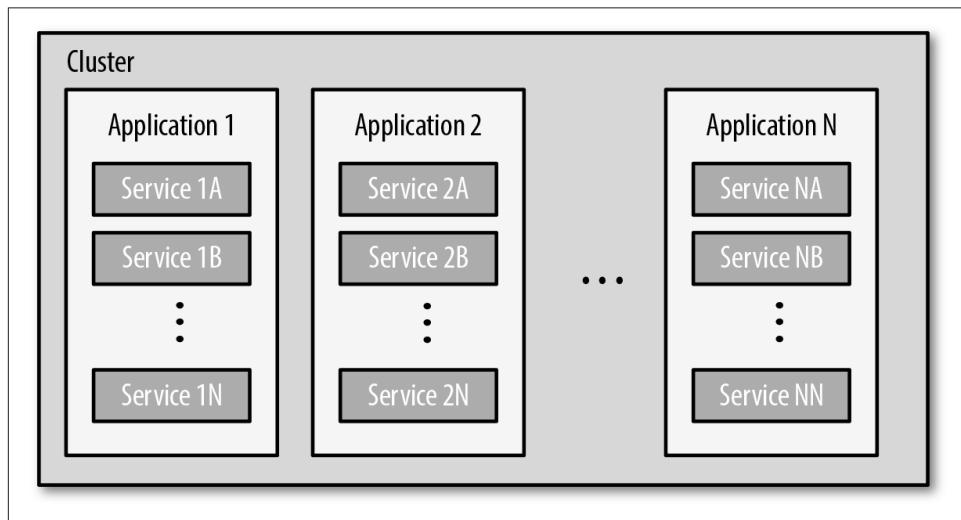


Figure 11-5. Multiple logical applications within a cluster

The Service Fabric automatically deploys your services across the nodes in your cluster. It determines the appropriate placement of services to nodes based on your application configuration. As shown in [Figure 11-6](#), the Service Fabric deployed Service 1A to all nodes, Service 2B to only nodes 1 and 2, and Service NN to only node N.

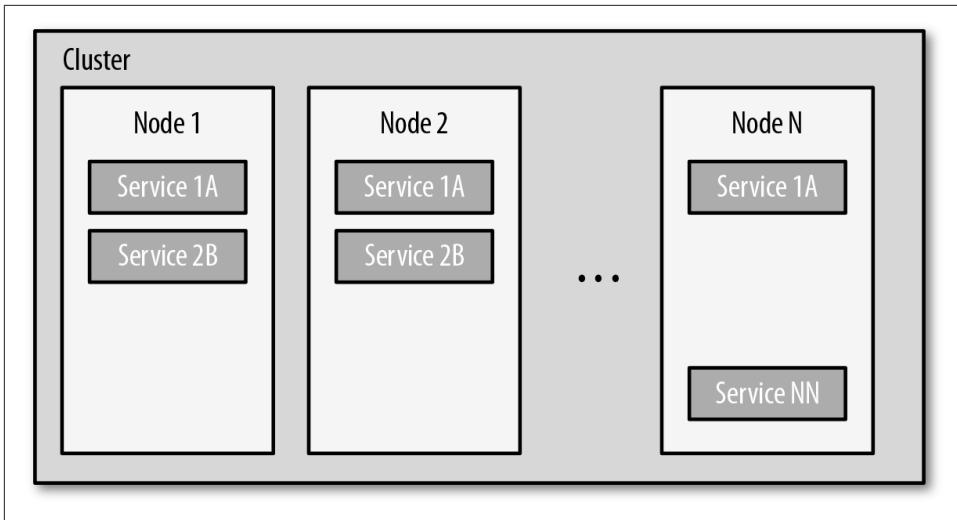


Figure 11-6. Service deployment across a cluster

Development defines an application's configuration and shares it with IT operations through an *ApplicationManifest.xml* file. Unlike the `system.servicemodel` section in a WCF configuration file, your application manifest file contains only limited configuration pertaining to deployment that the Service Fabric needs to properly deploy each service, as shown in [Example 11-1](#).

Example 11-1. The ApplicationManifest.xml file

```
<ApplicationManifest ApplicationTypeName = "MyApplication"
                     ApplicationTypeVersion = "...">
  <ServiceManifestImport>
    <ServiceManifestRef ServiceManifestName = "MyServicePkg"
                       ServiceManifestVersion = "..."/>
    <ServiceManifestRef ServiceManifestName = "MyWebServicePkg"
                       ServiceManifestVersion = "..."/>
    <ServiceManifestRef ServiceManifestName = "MyStatefulServicePkg"
                       ServiceManifestVersion = "..."/>
  </ServiceManifestImport>
  ...
  <DefaultServices>
    <Service Name = "MyService">
      <StatelessService ServiceTypeName = "MyServiceType">
        ...
      </StatelessService>
    </Service>
    <Service Name = "MyWebService">
      <StatelessService ServiceTypeName = "MyWebServiceType">
        ...
      </StatelessService>
    </Service>
  </DefaultServices>
</ApplicationManifest>
```

```

</Service>
<Service Name = "MyStatefulService">
    <StatefulService ServiceTypeName = "MyStatefulServiceType">
        ...
        </StatefulService>
    </Service>
</DefaultServices>
</ApplicationManifest>

```



The Service Fabric enforces during compilation as a convention the full Visual Studio project name and the *Pkg* suffix on all Reliable Actor package names. To maintain consistent package naming across all of your services, you should consider applying the same convention to your Reliable Service package names as well.

A service in the Service Fabric is a collection of code, configuration, and optional static data packages defined in a *ServiceManifest.xml* file. Again unlike WCF configuration files, a *ServiceManifest.xml* file contains only limited information about your service, as shown in [Example 11-2](#). The Service Fabric needs only a limited description of your service to properly host it because it prescribes nearly all of your service's addressing, connectivity, and behavioral configuration for you. For each service in your application, you create a service manifest file that defines the name of the service type, the name of the process that will host your service code, and basic configuration for your service's endpoint. As shown in [Example 11-1](#), the application manifest file can import a service manifest by referencing its package name, and you can provide application-specific deployment configuration for your service by referencing its type.

Example 11-2. The ServiceManifest.xml file

```

<ServiceManifest Name = "MyServicePkg" Version = "...">
    <ServiceTypes>
        <StatelessServiceType ServiceTypeName = "MyServiceType"/>
    </ServiceTypes>
    <CodePackage Name = "Code" Version = "...">
        <EntryPoint>
            <ExeHost>
                <Program>
                    MyService.exe
                </Program>
            </ExeHost>
        </EntryPoint>
    </CodePackage>
    <ConfigPackage Name = "Config" Version = "..." />
    <DataPackage Name = "Data" Version = "..." />
    <Resources>
        <Endpoints>
            ...
        </Endpoints>
    </Resources>
</ServiceManifest>

```

```
</Resources>  
</ServiceManifest>
```

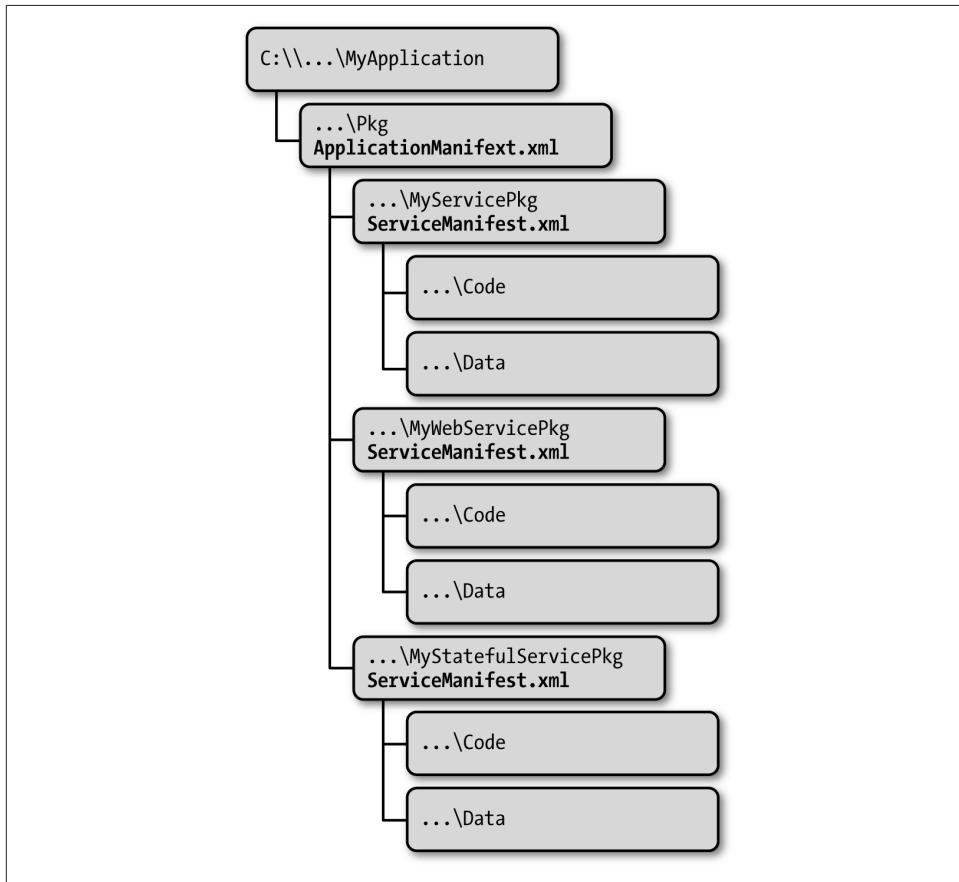


Figure 11-7. The application package file and folder layout

Naming

For both applications and services in Service Fabric you must specify the type and name for each element. You can consider the type of an application or service its *class* of which you can create one or more instances that you make unique by name. This allows you to deploy multiple instances of the same service type into a single application and multiple instances of the same application type into a single cluster, as shown in [Figure 11-8](#).

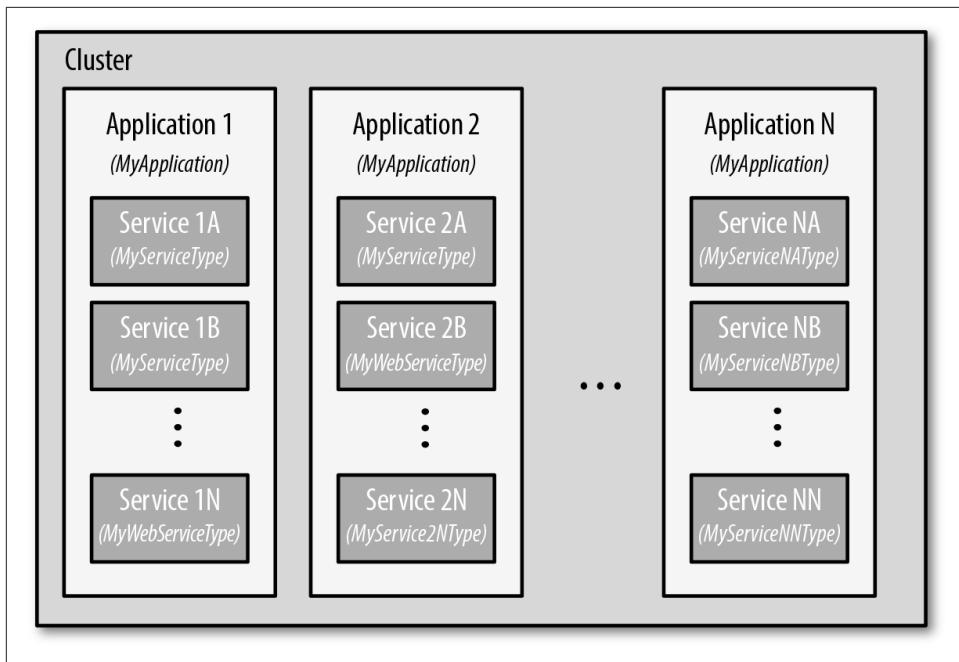


Figure 11-8. Multiple logical application and service types within a cluster

For application naming, Service Fabric uses the application type name in the application manifest as the application name verbatim. To deploy multiple applications of the same type, you must copy the existing application manifest file and change the application type name within.

For service naming, you specify both the service type and the service instance name in your service manifest file. For example, you could define two services for the application *MyApplication* that are separate instances of the *MyServiceType*, as shown in [Example 11-3](#). You would then access each by name in the service address.

Example 11-3. Multiple service instances of the same type

```
<ApplicationManifest ApplicationTypeName = "MyApplication"
                     ApplicationTypeVersion = "...">
  <ServiceManifestImport>
    <ServiceManifestRef ServiceManifestName = "MyServicePackage"
                       ServiceManifestVersion = "..." />
  </ServiceManifestImport>
  ...
  <DefaultServices>
    <Service Name = "MyService">
      <StatelessService ServiceTypeName = "MyServiceType">
        ...
      </StatelessService>
    </Service>
  </DefaultServices>
</ApplicationManifest>
```

```

        </StatelessService>
    </Service>
    <Service Name = "MyOtherService">
        <StatelessService ServiceTypeName = "MyServiceType">
            ...
        </StatelessService>
    </Service>
</DefaultServices>
</ApplicationManifest>

```

The Service Fabric might then deploy these services as shown in [Figure 11-9](#) using a deployment plan similar to the one depicted in [Figure 11-6](#). The Service Fabric deployed MyService to all nodes and MyOtherService to only nodes 1 and 2.

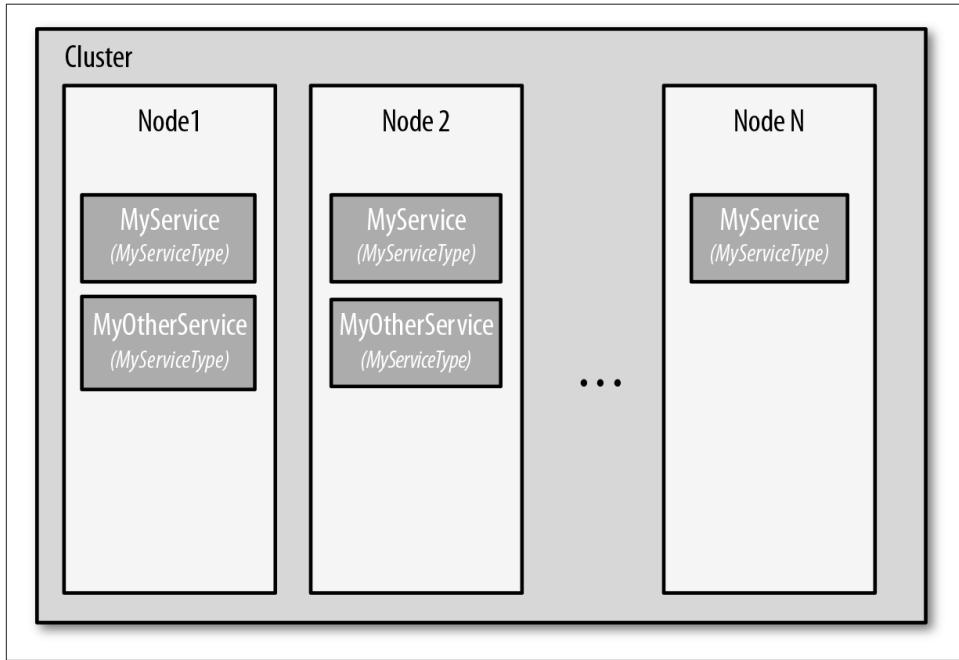


Figure 11-9. Service type deployment across a cluster

Programming the Service Fabric

Azure Service Fabric's programming model provides the appropriate concepts for each mode of interaction your system will need, making it easier for you to design, build, test, and consume different service types in a single environment. Service Fabric's system structure also informs its programming model. This relationship between system structure and programming model is important because it helps you to ensure that your services and your system architecture will match well with Service Fabric's capabilities and deployment model.

The Service Programming Model

To assist you in creating services that conform to the interaction modes it supports, Azure Service Fabric establishes two distinct service programming models: a Service Model and an Actor Model. Service Fabric's Service Model, called *Reliable Services*, provides the necessary programming constructs for creating services that support the common API and Intranet interaction modes.

Service Fabric's Actor Model, called *Reliable Actors*, is a novel service programming construct that supports the further shrinking of the service boundary. Reliable Actors effectively provide platform-level support for the notion of *every class as a service*. The platform can even support actors of finer granularity. The Actor Model allows you to create a variety of component-like services. Due to this, it is incorrect to consider actors simply as object equivalents. For example, actors can perform activities such as aggregation, computation, transformation, and generation. They can provide access to external data or represent the data itself as storage or state. Actors can also provide access to resources within a Service Fabric cluster or those external to it, such as file stores and even other third-party systems. This broad expressive capability makes actors ideal for service orientation, as you can easily compose them into use case sequences, workflow orchestrations, or other compositional structures. Their compositional expressiveness also makes actors ideal for representing new parallel programming patterns such as mesh-based topologies, sorting networks, and graph-based data forms. You can also use actors for more mundane tasks such as scheduling, data I/O, and data flow. The simple Reliable Actor programming model and direct platform support lower the bar of entry for developers, making it easier to implement these advanced programming concepts.

Reliable Services

The `Microsoft.ServiceFabric.Services` namespace contains support for the Reliable Services programming model. Service Fabric provides reliability for these services by enabling redundancy for service instances and replication for service state. It is important to note that service reliability in this context is about providing resiliency in the face of service and node failures, not connectivity reliability as is the case with WCF. Through the Reliable Services programming model you can create both background services for performing long-running tasks and interactive services handling client requests.

Background Reliable Services do not accept requests from clients, so they do not need to provide a listener for the service. Instead, these services perform long-running background tasks via a `RunAsync()` method that Service Fabric calls when it first activates the service. Both Reliable Service base classes, the `StatelessServiceBase` and the `StatefulServiceBase`, define the same `RunAsync()` signature, as shown in [Example 11-4](#). To keep your background service instance active and responsive

within the Service Fabric, it is best practice that your RunAsync() method enters a loop and explicitly cancels the loop by calling the ThrowIfCancellationRequested() method, which throws an OperationCanceledException. This approach ensures that .NET will also cancel any tasks waiting on the loop.

Example 11-4. The RunAsync() method

```
public interface IStatelessServiceInstance
{...}

public abstract class StatelessServiceBase : IStatelessServiceInstance
{
    protected virtual async Task RunAsync(CancellationToken cancellationToken)
    {...}
    //More methods
}

public class StatelessService : StatelessServiceBase
{...}
class MyStatelessService : StatelessService
{
    protected override async Task RunAsync(CancellationToken cancellationToken)
    {
        while(true)
        {
            cancellationToken.ThrowIfCancellationRequested();
            ...
        }
    }
}

public interface IStatefulServiceReplica
{...}

public abstract class StatefulServiceBase : IStatefulServiceReplica
{
    protected virtual async Task RunAsync(CancellationToken cancellationToken)
    {...}
    //More methods
}

public class StatefulService : StatefulServiceBase
{...}

class MyStatefulService : StatefulService
{
    protected override async Task RunAsync(CancellationToken cancellationToken)
    {
        while(true)
        {
            cancellationToken.ThrowIfCancellationRequested();
        }
    }
}
```

```
    ...
}
```

As the name implies, interactive Reliable Services accept client requests by providing a listener for the service. These services support two modes of interaction: API and Intranet. The service, not its configuration, explicitly defines the connectivity it will use via the `CreateCommunicationListener()` method, as shown in [Example 11-5](#).

Example 11-5. The CreateCommunicationListener() method

```
public interface ICommunicationListener
{...}

public abstract class StatelessServiceBase : IStatelessServiceInstance
{
    public ICommunicationListener CommunicationListener
    {get;}
    protected abstract ICommunicationListener CreateCommunicationListener()
    {...}
    //More methods
}

public abstract class StatefulServiceBase : IStatefulServiceReplica
{
    public ICommunicationListener CommunicationListener
    {get;}
    protected abstract ICommunicationListener CreateCommunicationListener()
    {...}
    //More methods
}
```

Your service's `CreateCommunicationListener()` method should return an instance of `ICommunicationListener` specific to the interaction mode you require. For example, the `ServiceCommunicationListener<T>` class is the listener for Service Fabric's *default* Reliable Service programming model. `T` in this case is an interface of service operations. Unlike background services, default interactive services must implement an interface to conform to the default Reliable Service programming model. [Example 11-6](#) shows the Reliable Service `MyStatelessService` explicitly defining its connectivity through its `CreateCommunicationListener()` method.

Example 11-6. Using the ServiceCommunicationListener<T> class

```
public class ServiceCommunicationListener<T> : ICommunicationListener
    where T : class
{
    public ServiceCommunicationListener(T serviceImplementationType)
```

```

    {...}
    public ServiceCommunicationListener(T serviceImplementationType,
                                         string endpointResourceName)
    {...}
    //More methods
}

public interface IService
{}
interface IMyService : IService
{...}

class MyStatelessService : StatelessService,IMyService
{
    protected override ICommunicationListener CreateCommunicationListener()
    {
        ICommunicationListener communicationListener =
            new ServiceCommunicationListener<IMyService>(this);
        return communicationListener;
    }
}

```

Similar code to [Example 11-6](#) applies to the stateful service.

You can also specify an endpoint configuration for your service listener by endpoint resource name. You must set this name on the listener either through the listener's constructor or an associated property on the listener class before Service Fabric opens the listener. You define the endpoint's configuration in the Resources section of your service's manifest, as shown in [Example 11-7](#).

Example 11-7. Configuring a service's endpoint

```

<ServiceManifest Name = "MyServicePackage" Version = "...">
    ...
    <Resources>
        <Endpoints>
            <Endpoint
                Name = "ServiceEndpoint"
                Protocol = "tcp"
                Port = "8000"
            />
        </Endpoints>
    </Resources>
</ServiceManifest>

```

Since Service Fabric must manage the actual physical details of your service's endpoint, it limits the options you have to control it. Because of this, you should consider the limited endpoint configuration you can provide a *virtual* endpoint. You can only provide values for the endpoint's protocol, port, and certificate reference. The available protocols are TCP, HTTP, and HTTPS. If you do not provide an endpoint con-

figuration, Service Fabric will use TCP as the default protocol and it will auto-assign the endpoint's port. In fact, the only reason to provide an endpoint configuration would be to provide a specific port or certificate for TCP, or indicate a different protocol altogether. Providing a specific port value disables Service Fabric's port auto-assignment feature, though it does not disable Service Fabric's management of the endpoint's physical address. Be aware that Service Fabric relies on auto-assigning ports to ensure it can create unique endpoints for every service instance as defined in your application manifest. Specifying a port value limits your service deployment to a single instance. You must ensure this configuration is compatible with your application's deployment plan. Otherwise, your service will fail to activate.



The type of communication listener your service provides must support the endpoint protocol setting in your service's manifest or the service will fail to activate.

API services for interoperability

API services are Reliable Services that are interoperable and potentially accessible over the Internet. They provide broad reach to services that extend beyond the firewall of the environment in which the Service Fabric cluster resides. You should consider API services your system's *public* API. This makes API services ideal for supporting the *multiplicity of clients*, as this interaction mode is compatible with most common mobile platforms, browser-based languages, and tooling, and in line with web developer capabilities.

To enable this broad reach, API services use HTTP connectivity and Uri-based service interaction. Service Fabric uses ASP.NET Web API as the supportive programming framework for API services. API services are the only service type within Service Fabric meant to be publicly accessible over the internet. This means web clients can access your API service by formalizing an HTTP Uri using a well-known, fully qualified domain name (FQDN) and its service name such as `http://mycompanydomain/myapiservice/....` To enable public access, you will most often associate your FQDN with an HTTP gateway or NAT in the DMZ of your environment that will forward HTTP requests to your cluster. So that your API service may receive requests over the standard HTTP port 80, you must also disable Service Fabric's port auto-assignment feature by providing a port setting in the service endpoint's configuration. Additionally, only single instance stateless services can be publicly accessible API services.

You define the API interaction mode for your service by creating an `OwinCommunicationListener`, as shown in [Example 11-8](#). This listener enables and controls the self-hosting of ASP.NET Web API.

Example 11-8. Using the OwinCommunicationListener class

```
public class OwinCommunicationListener : ICommunicationListener
{
    public class MyAPIService : StatelessService
    {
        protected override ICommunicationListener CreateCommunicationListener()
        {
            ICommunicationListener communicationListener =
                new OwinCommunicationListener(...);
            return communicationListener;
        }
        //No other methods
    }
}
```

To provide the best programming experience and conform to best practice, the API service class is only responsible for creating the listener. As shown in [Example 11-8](#), it has no interface. Instead, you define a Web API controller class to contain all of the web operations for your service just as you would for an ASP.NET Web API service hosted in IIS. In fact, the ASP.NET Web API in Service Fabric and the one in IIS are one and the same. This means you can transfer your Web API service code directly to a solution targeting the Service Fabric. The only difference is the way in which Service Fabric hosts your API service. Since a publicly accessible API service in the Service Fabric must only be stateless, you must ensure your existing ASP.NET Web API service is also stateless.



As of this writing, the `OwinCommunicationListener` is not part of the Service Fabric programming framework. It is available as community sample code. The production version of the Service Fabric will contain these framework additions directly as part of full support for ASP.NET 5.

Service Fabric self-hosts your API service using Katana, an *OWIN*-based framework (Open Web Interface for .NET) created by Microsoft. Katana provides the capability for you to self-host an ASP.NET Web API in Service Fabric. The OWIN compatible listener controls the actual web server startup. Service Fabric then manages the hosting process as it would any other self-hosted service.



OWIN, the Open Web Interface for .NET, is an open source specification for creating lightweight web servers. OWIN achieves this by allowing the decoupling of web application specifics from the underlying web server infrastructure itself. This allows you to host the resulting lightweight web server in a simple Windows EXE or Windows Service. Katana is an open source OWIN implementation built and released by Microsoft.

Intranet services

Intranet services are Reliable Services that should primarily play a compositional role within your system where they orchestrate calls to other services and actors. Intranet services are only accessible from within the environment where the Service Fabric cluster resides. You should consider Intranet services your system's *private API*. This makes Intranet services ideal for efficient, connected interactions compatible with server platforms, backend languages, and tools and in line with service and system developer capabilities.

Intranet services provide a simple interface-based programming model that is easy to test, mock, and simulate. This makes Intranet services perfect for consumption by all backend clients such as websites, API services, other Service Fabric services, as well as simple application services hosted in Azure App Service. Intranet services are also the preferred interaction mode for client applications behind the environment's firewall such as Windows applications and Windows services.

When you use them in their compositional role, Intranet services can also act as *microservices*. You can think of microservices as compositional services that orchestrate calls to one or more component services. Compositional Intranet services also support the best practice design principle of encapsulating system orchestrations such as use cases and workflows within a service, instead of pushing them to the client tier. Pushing system orchestrations to your clients can become very problematic and costly, especially with the multiplicity of clients your systems will need to support. If you push service orchestrations to your clients, as these orchestrations change and evolve you will need to make frequent, redundant updates to your clients on every client platform, instead of making changes once within a service operation. Pushing service orchestrations to your clients also prevents reuse of this important type of business logic. Orchestrating services from your clients can produce very *chatty* designs that create an unnecessary volume of calls over the network. Chatty service designs become a worst case scenario with mobile clients because you are now putting more calls over the slowest part of the network: the cell tower.

To ensure the broadest reuse of your service orchestrations, you should consider API services as clients to your compositional Intranet services. This approach brings consistency to your client programming model by providing the consumers of your services a single, simple, interface-based programming model that they can easily mock and test. It also promotes single point of entry to your system, always a desired design goal. In general, whenever possible, you should prefer an interface-based programming model over any alternative.

Service Fabric provides two programming models for Intranet services: the default service programming model and the WCF programming model.

Default Intranet services

As you have already seen, default Intranet services provide a straightforward interface-based programming model. In this case, the service interface is not a service contract, though request and response parameters must still be data contracts. A default Intranet service interface must also inherit from the empty marker interface, `IService`. Default Intranet services use a fully prescribed connectivity configuration that completely hides service connectivity concerns. Very much like Reliable Actors, Default Intranet services are actually built atop WCF and use `NetTcpBinding` internally. They also carry an additional constraint of only being accessible from within a Service Fabric cluster.

To define a Reliable Service as a Default Intranet service, your service must specify `ServiceCommunicationListener<T>` as its listener, as shown in [Example 11-9](#). It is important to note that although this type of Reliable Service is built atop WCF, the programming model does not afford you access to any of WCF's extensibility points.

Example 11-9. Defining a default Intranet service

```
[DataContract]
class MyDataContract
{
    [DataMember]
    public string MyValue
    {get;set;}
}

//Not a service contract
interface IMyService : IService
{
    Task MyMethodAsync(MyDataContract request);
}

//No service behaviors
class MyService : StatelessService,IMyService
{
    public Task MyMethodAsync(MyDataContract request)
    {...}
    protected override ICommunicationListener CreateCommunicationListener()
    {
        ICommunicationListener communicationListener =
            new ServiceCommunicationListener<IMyService>(this);
        return communicationListener;
    }
}
```

The behavioral configuration for Default Intranet services uses `ConcurrencyMode.Multiple`, `InstanceContextMode.Single`, no reliable messaging and no transaction flow. The prescribed `NetTcpBinding` sets `MaxReceiveMessageSize` to 4MB,

which for most common service scenarios is too high. It sets the `OpenTimeout` and `CloseTimeouts` to 5 seconds, which allows connectivity failures to fail fast. This is the best practice configuration for service infrastructure in dynamically managed environments. Service Fabric proxies provide automatic retry on these transient connectivity faults. The Service Fabric will attempt to retry until it can achieve a successful connection. For Reliable Services, you can control the overall length of the retry period by providing a `CancellationToken` in your proxy call. The binding also sets the `SendTimeout` and `ReceiveTimeout` to an effectively infinite value. These settings allow successful connections to persist indefinitely within the cluster as long as the endpoints are available. The Service Fabric promotes all service exceptions to faults on the service side and performs exception extraction on the client side, rethrowing exceptions within the proxy as described in [Chapter 6](#). This means that regardless of the type of failure whether it is due to connectivity, timeout, security or service fault you will receive only the .NET exception from your Reliable Service and Reliable Actor proxy calls.

WCF Intranet services

Unlike Default Intranet services, WCF Intranet services provide an explicit, but constrained WCF programming model. For example, you cannot control a WCF Intranet service's `ServiceBehaviorAttribute`. Service Fabric will override it with the required best practice settings compatible for connectivity within the cluster. You can provide a custom `IServiceBehavior` that affords you access to WCF's extensibility points. You can also provide a binding configuration at the time of listener creation of `NetTcpBinding`, `BasicHttpBinding`, or `WSHttpBinding`.

To define a Reliable Service as a WCF Intranet service, your service must specify the `WcfCommunicationListener` as its listener, as shown in [Example 11-10](#). Note that unlike the listener for a Default Intranet service, the listener for a WCF Intranet service does not take a service contract as a generic parameter. Instead, it is weakly typed and takes the service contract type as a parameter and either a reference to a service implementation instance or the type of the service implementation class as a parameter. This makes its construction similar to that of a WCF `ServiceHost`. [Example 11-10](#) also shows that the WCF listener supports setting an endpoint resource name as a property instead of providing it through a constructor.

Example 11-10. Defining a WCF Intranet service

```
public class WcfCommunicationListener : ICommunicationListener
{
    public WcfCommunicationListener(Type communicationInterfaceType, object service);
    public WcfCommunicationListener(Type communicationInterfaceType,
                                    Type communicationImplementationType);
    public Binding Binding
    {get;set;}
```

```

public string EndpointResourceName
{get;set;}
//More methods
}

[ServiceContract]
interface IMyWcfContract
{
    [OperationContract]
    Task MyMethodAsync(MyDataContract request);
}

[AttributeUsage(AttributeTargets.Class)]
class MyServiceBehavior : Attribute,IServiceBehavior
{...}

//Optional
[MyServiceBehavior]
class MyWcfService : StatelessService,IMyWcfContract
{
    public Task MyMethodAsync(MyDataContract request)
    {...}
    protected override ICommunicationListener CreateCommunicationListener()
    {
        WcfCommunicationListener communicationListener
            = new WcfCommunicationListener(typeof(IMyWcfContract),this);

        communicationListener.EndpointResourceName = "ServiceEndpoint";
        communicationListener.Binding = new NetTcpBinding(...);

        return communicationListener;
    }
}

```

Intranet service hosting

Service Fabric hosts both types of Intranet service using self-hosted WCF. The listener controls the `ServiceHost`, which allows subclasses of `WcfCommunicationListener` to explicitly control and configure the service's `ServiceHost`. `WcfCommunicationListener` also allows subclasses to explicitly control the service's endpoint configuration. Both options enable you to access additional WCF extensibility points as described throughout the book. However, the likelihood of you doing that is somewhat diminished over the Service Fabric. The reason is that most of the extensibility and power of WCF demonstrated in this book (and codified in *Service-ModelEx*) is for the purpose of providing your own platform for the services to execute in, your own superior runtime. The main reason to adopt the Service Fabric is precisely because it provides such a ready-made platform already.

Reliable Actors

The `Microsoft.ServiceFabric.Actors` namespace contains support for the Reliable Actors programming model. As with Reliable Services, the Service Fabric provides reliability for Reliable Actors by enabling redundancy for actor types and replication for actor state.

Reliable Actors (or just actors) primarily play a component role in your system that you should only access from within the Service Fabric cluster. They allow you to create services of fine granularity. You should consider actors the components your Intranet services compose. Think of Internet services as directors of a play, and the Reliable Actors as the actors in it. Actors should perform specific activities, and the Intranet services sequence them into a desired workflow, executing a required use case. This design technique ensures that you encapsulate actor orchestration instead of pushing it out into your clients. Actors are also the essential programming construct for creating advanced parallel programming patterns over a mesh or a sorting network.

Reliable Actors present a strict but straightforward interface-based programming model. Actors inherit behavior from a base class, either `Actor` for stateless actors or `Actor<S>` for stateful actors. Actors may only support a single interface that must inherit from the empty marker interface `IActor`. Actor operations can only be request/response. This constraint ensures that actors can always promote exceptions to faults and return them back to callers. All operation parameters must be data contracts serializable by `DataContractSerializer`. Finally, all actor operations must be Task-based, a constraint Service Fabric validates at runtime. One reason Reliable Actor operations must be Task-based is to allow the underlying actor framework to call your actor operations asynchronously without blocking. This allows the Service Fabric's actor framework to handle the high volume of actor requests necessary to support Cloud-scale deployments. Task-based operations also promote efficient thread utilization within each Service Fabric host process. Another reason that Reliable Actor operations are Task-based is to allow you to take what would otherwise be a sequence of synchronous actor-to-actor calls and transform them into a series of concurrent and parallel one-way calls. This technique of using non-blocking calls to other actors is more in line with the intent of the Actor Model pattern. Ideally, actors should not block when calling other actors. This allows them to remain responsive and to continue to efficiently process messages when under load. If a request requires a response, the actor called should send it as a separate message.

While all actor operations must be Task-based, the Service Fabric does not require that you apply the `async/await` pattern (TAP) to them. Instead, it should become your best practice to reserve applying `async/await` to only those actor operations that encapsulate asynchronous activity. This will more clearly convey your intent to others who might consume or extend your actor. As per [Chapter 8](#), the majority of actor

operations will not encapsulate asynchronous activity. Most often, you will not apply `async/await` to these operations to avoid the nagging compiler warning stating your operation is missing an `await`. Instead, you will most often use `Task.FromResult<T>()` to return your actor operation's result synchronously, as shown in [Example 11-11](#).

Example 11-11. Defining a stateless Reliable Actor

```
public abstract class ActorBase
{...}
public abstract class Actor : ActorBase
{...}

public interface IActor
{}
interface IMyActor : IActor
{
    Task MyMethodAsync(MyDataContract someParam);
    Task<MyDataContract> MyOtherMethodAsync(MyDataContract someParam);
}

class MyActor : Actor,IMyActor
{
    public Task MyMethodAsync(MyDataContract someParam)
    {
        //Do synchronous work here
        return Task.FromResult(true);
    }
    public Task<MyDataContract> MyOtherMethodAsync(MyDataContract someParam)
    {
        //Do synchronous work here
        return Task.FromResult(new MyDataContract());
    }
}
```

As shown in [Example 11-11](#), the Reliable Actors programming model completely hides connectivity concerns. Using WCF as a foundation, Reliable Actors offer the same fully prescribed connectivity configuration as Default Reliable Services.

Reliable Actor garbage collection

The Service Fabric provides control over Reliable Actor garbage collection using a declarative model through the `ActorGarbageCollectionAttribute`, as shown in [Example 11-12](#). After the actor is idle for more than the allotted timeout, Service Fabric garbage collection first calls the actor's `OnDeactivateAsync()` method, then removes the actor instance from memory. The Service Fabric determines the actor idle time by monitoring various actor activity. Actor garbage collection has a default idle timeout of 1 hour and a default scan interval of 1 minute. It is essential to note

that actor garbage collection does not remove an actor's state from the backing store. The only way for you to remove actor state is to do it manually. This chapter presents a number of techniques you can use to simplify this manual process.

Example 11-12. The ActorGarbageCollectionAttribute

```
[AttributeUsage(AttributeTargets.Class)]
public class ActorGarbageCollectionAttribute : Attribute
{
    public long IdleTimeoutInSeconds
    {get;set;}
    public long ScanIntervalInSeconds
    {get;set;}
    public ActorGarbageCollectionAttribute()
    {...}
}
```

Reliable Actor lifecycle events

The Service Fabric provides lifecycle events for Reliable Actors. It calls the `OnActivateAsync()` method upon the first activation of an actor and calls the `OnDeactivateAsync()` method upon the expiration of the garbage collection idle period. Inheritors can override these lifecycle events to perform actor lifecycle related actions. Normally, you do not need to worry about handling lifecycle events for your actors. One scenario where you might consider using `OnDeactivateAsync()` would be to clean up an actor's state.

```
public abstract class ActorBase
{
    public virtual Task OnActivateAsync();
    public virtual Task OnDeactivateAsync();
    //More methods
}

class MyActor : Actor, IMyActor
{
    public override Task OnActivateAsync()
    {
        //Do first time activation work here.
        ...
    }
    public override Task OnDeactivateAsync()
    {
        //Do additional cleanup before instance disposal here.
        ...
    }
    //More Methods
}
```



The Service Fabric also provides asynchronous client callback events, timers, and reminders for Reliable Actors. These are advanced Reliable Actor programming constructs that are beyond the scope of this chapter. In general, you should use these constructs with great care because they do not promote loose coupling between the elements of your system and they cripple much of the simplicity the Actor Model introduces to parallel processing.

Reliable Actor hosting

The Service Fabric currently hosts Reliable Actors using self-hosted WCF. The Service Fabric controls the `ServiceHost` responsible for actor connectivity and manages the hosting process.

State Management

State and the reliable management of state are essential aspects of Service Fabric's service programming model. The options for state management that Service Fabric provides are rich and varied. You can create reliable queues for your services to improve their throughput. Or you can create reliable caches and even use reliable collections to model distributed database concepts. And with stateful Reliable Actors you can parallelize your state to enable powerful new parallel data patterns such as sorting, graphs, and visitors.

There are two modes of service state: stateless and stateful. Reliable Services and Reliable Actors support both state programming models. You define a service's state behavior programmatically by inheriting from the appropriate base class.

Stateless services and actors

The term *state* when it comes to the Service Fabric refers to state the Service Fabric has to manage, not to the instance state as in [Chapter 4](#). Stateless services and actors possess no explicit state that Service Fabric manages. To define a stateless Reliable Service, your service must inherit from the `StatelessService` class found in the `Microsoft.ServiceFabric.Services` namespace. To define a stateless Reliable Actor, your actor must inherit from the `Actor` class found in the `Microsoft.ServiceFabric.Actors` namespace.



Once activated, stateless Reliable Service instances remain in memory indefinitely. This means any class level variables they may possess will retain their state across service calls, giving the illusion that these services also maintain state. Unfortunately, you will lose this volatile, in-memory state in the event of an instance fault or if the Service Fabric moves your service to another node within the cluster. You should never rely on this *implicit state* in your services and actors.

Stateful services and actors

By definition, stateful services and actors possess explicit state that Service Fabric manages. State in Service Fabric is transactional, reliable, distributed, and optionally persistent. Service Fabric uses a best practice *replica* model to ensure state reliability. Service Fabric maintains multiple, distributed instances of service or actor state along with its associated service code packages. Each replica set has one primary replica and one or more secondary replicas. You define the minimum and target sizes of a replica set through the service section of your application manifest, as shown in [Example 11-13](#). If Service Fabric cannot satisfy the replica set's minimum size, your service will fail to activate.

Example 11-13. Defining stateful service replica set size

```
<ApplicationManifest ...>
  ...
  <DefaultServices>
    <Service Name = "MyStatefulService">
      <StatefulService ServiceTypeName = "MyStatefulServiceType"
                      MinReplicaSetSize = "2"
                      TargetReplicaSetSize = "3">
        ...
        </StatefulService>
      </Service>
    <Service Name = "MyStatefulActor">
      <StatefulService ServiceTypeName = "MyStatefulActorType"
                      MinReplicaSetSize = "4"
                      TargetReplicaSetSize = "5">
        ...
        </StatefulService>
      </Service>
    </DefaultServices>
  </ApplicationManifest>
```

Each replica represents a full instance of your service or actor including all packages hosted by a Windows process. By default, only the endpoint of the primary replica is accessible. During deployment, the Service Fabric automatically places replicas across your cluster ensuring proper distribution to maximize reliability. This means each

replica instance must be on a different cluster node, as shown in [Figure 11-10](#). If the count of active, healthy nodes in your cluster cannot satisfy your replica set minimum requirement, your service or actor will fault. When persisting state, the Service Fabric replicates the state across the replica set within a transaction. The Service Fabric also uses other advanced features to ensure state consistency. Upon primary replica failure, the Service Fabric automatically promotes a secondary replica to the primary position, activates its endpoint, and updates the endpoint's address within the cluster. The Service Fabric will also self-heal the replica set by adding secondary replicas up to the target size when new nodes become available. [Figure 11-10](#) shows the possible replica set placement for three services all with a `TargetReplicaSetSize` of 3.

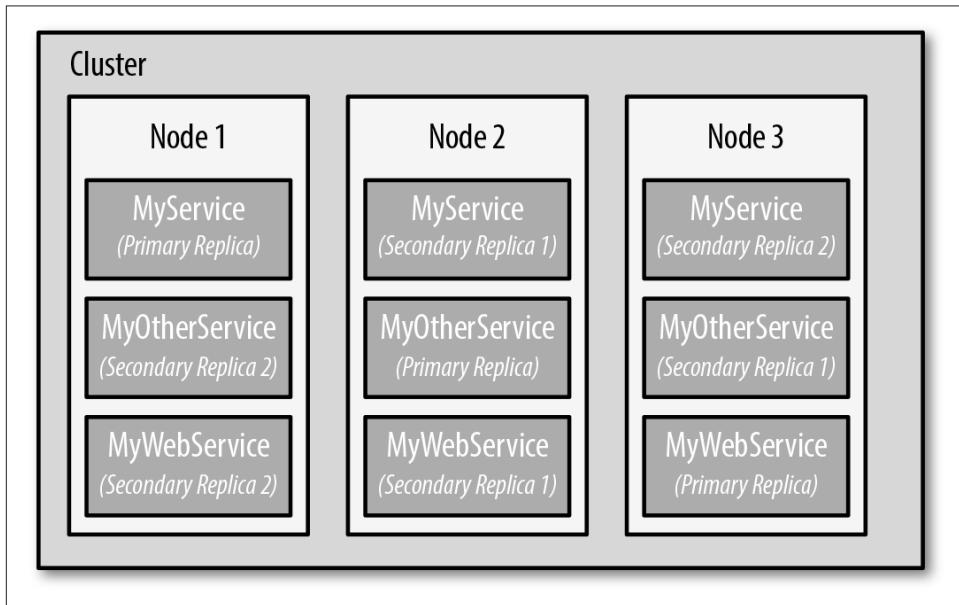


Figure 11-10. Replica placement across a cluster

As you might expect, you must set a stateful service's `MinReplicaSetSize` to a value less than or equal to its `TargetReplicaSetSize` or the service will fail to activate. Additionally, if you want to ensure that your replica set will survive the loss of a primary replica while also being able to heal itself without the primary replica incurring the necessary replication overhead, you should always use a `TargetReplicaSetSize` of no less than 3. While you can set your service's replica set size to 1, you should avoid doing so, as this defeats the purpose of replication and therefore disables reliability for your stateful service. Because a `TargetReplicaSetSize` of 1 is the default, you must always set `TargetReplicaSetSize` to ensure your stateful service will be reliable.

The nature of your system and its sensitivity to failure are important factors when determining the level of reliability that your stateful service requires. There is a cost in both service performance and hosting associated with making your services highly reliable. The larger your replica set, the longer it will take to replicate state across it. And you may also need to add additional nodes to your cluster to support a large replica set size. For example, if your system represents mission-critical software responsible for anything from real-time machine control to monitoring the vital signs of patients, the reliability of your service state will be very important to you. In this case, the higher cost of a greater degree of reliability may be well worth it. If, on the other hand, scalability and response time are more important than reliability, choose the other way around. This is clearly the trade-off demonstrated in today's cellular phone networks.

All state representations within the Service Fabric must be serializable and by default a data contract as `DataContractSerializer` is the default serializer. Once you have deployed a stateful service or actor, you cannot change the type of state management for it without deploying a new application instance. As with stateless services, stateful service and actor instances may accrue volatile implicit state that you will lose upon primary replica fault or Service Fabric cluster management. What you can do, however, is interact with other reliable services or actors that the Service Fabric manages, and even have those other components store your state.

Stateful Reliable Services

To define a stateful Reliable Service, your service must inherit from the `StatefulService` class. All background and interactive Reliable Services may possess state:

```
//Background or API service
class MyStatefulService : StatefulService
{...}

class MyStatefulDefaultService : StatefulService, IMyService
{...}

class MyStatefulWcfService : StatefulService, IMyWcfContract
{...}
```

Stateful Reliable Services maintain state within one of the *Reliable Collection* types found in `Microsoft.ServiceFabric.Data.Collections`. All reliable collection types are distributed, transactional resources that replicate state across all the replicas for a service and persist their state to the local disk of each node on which a replica resides.

There are two types of reliable collection; a reliable dictionary and a reliable queue. Both are instances of `IReliableState`:

```
public interface IReliableState
{...}
```

```

public interface IReliableDictionary<K,V> : IReliableState, ...
    where K : IComparable<K>,IEquatable<K>
{...}

public interface IReliableQueue<T> : IReliableState, ...
{...}

```

For stateful Reliable Services you must manage service state manually through a *state manager*. You use the `StatefulService.StateManager` property to access the state manager and to manage service state. `StatefulService.StateManager` returns an instance of `IReliableStateManager`. As shown in [Example 11-14](#), you cannot directly create an instance of a reliable collection. Instead, `IReliableStateManager` acts as a class factory manufacturing reliable collection instances for you by the reliable collection interface type `T`. This allows Service Fabric to own and manage the instance of a reliable collection internally.

Example 11-14. Defining Reliable Service state

```

public interface IReliableStateManager : ...
{
    Task<T> GetOrAddAsync<T>(string name) where T : IReliableState;
    Task RemoveAsync(string name);
    //More methods
}

public class StatefulService : StatefulServiceBase
{
    public IReliableStateManager StateManager
    {get;protected set;}
    //More methods
}

class MyStatefulService : StatefulService,IMyService
{
    IReliableDictionary<string,string> State
    {
        get
        {
            IReliableDictionary<string,string> state =
                StateManager.GetOrAddAsync<IReliableDictionary<string,string>>(
                    "MyReliableDictionary");
            return state;
        }
    }

    public Task MyMethodAsync(...)
    {
        IReliableDictionary<string,string> state = State;
        ...
    }
}

```

```
//More methods  
}
```

As shown in [Example 11-14](#), the Service Fabric maintains state per service type instance and you access an instance of service state by name. Using the same state name in a different service will reference a different instance of state. Once you add state for the service to Service Fabric, it exists until you remove it.

In addition to manually managing state for stateful Reliable Services, you must also manually manage the scope of the state transaction. As shown in [Example 11-15](#), you manage stateful Reliable Service transactions through the `StateManager` property. Service Fabric persists the state across all replica instances within the transaction upon successful commit.

Example 11-15. Managing the Reliable Service state transaction

```
public interface IReliableStateManager : ...  
{  
    ITransaction CreateTransaction();  
    Task<T> GetOrAddAsync<T>(ITransaction tx, string name) where T : IReliableState;  
    Task RemoveAsync(ITransaction tx, string name);  
    //More methods  
}  
  
public interface IReliableDictionary<K,V> : ...  
{  
    Task<V> AddOrUpdateAsync(ITransaction tx, K key, V addValue,  
                             Func<K,V,V> updateValueFactory);  
    //More methods  
}  
class MyStatefulService : StatefulService, IMyService  
{  
    IReliableDictionary<string, long> State  
    {...}  
    public async Task MyMethodAsync(MyDataContract someParam)  
    {  
        IReliableDictionary<string, long> state = State;  
  
        using(ITransaction transaction = stateManager.CreateTransaction())  
        {  
            Func<string, long, long> updateValue = (key, oldValue)=>oldValue + 1;  
            long currentCount = await State.AddOrUpdateAsync(transaction,  
                                                someParam.MyValue, 1, updateValue);  
  
            await transaction.CommitAsync();  
        }  
    }  
    //More methods  
}
```

As shown in [Example 11-15](#), you should take a pessimistic approach to state transaction committal by explicitly waiting for the result of the commit operation. This ensures your service operation will not return until Service Fabric has either successfully committed your state across all replicas or a fault has occurred.



You should always wait on the result of the transaction commit. Taking an optimistic approach to state transaction committal by not waiting will inevitably lead to state inconsistencies.

Stateful Reliable Actors

To define a stateful Reliable Actor, your actor must inherit from the `Actor<S>` class, as shown in [Example 11-16](#).

Example 11-16. Defining a stateful Reliable Actor

```
public abstract class Actor<S> : ActorBase, ... where S : class
{
    public S State
    {get;set;}
}

[DataContract]
public class MyState
{
    [DataMember]
    public string MyValue
    {get;set;}
}

class MyStatefulActor : Actor<MyState>, IMyActor
{...}
```

Service Fabric automatically manages Reliable Actor state through an *actor state provider*. You define which actor state provider Service Fabric will use to manage your actor's state by annotating it with an `actor state provider` attribute. There are three types of actor state providers: KVS (key-value store), Reliable Dictionary, and Volatile. You designate the KVS actor state provider by using the `KVSActorStateProviderAttribute`. KVS is the default actor state provider that Service Fabric uses if you do not provide one. You designate the reliable dictionary actor state provider by using the `ReliableDictionaryActorStateProviderAttribute`. And you designate the volatile actor state provider by using the `VolatileActorStateProviderAttribute`, as shown in [Example 11-17](#).

Example 11-17. The actor state providers

```
[AttributeUsage(AttributeTargets.Class)]
public abstract class ActorStateProviderAttribute : Attribute
{...}

public class KvsActorStateProviderAttribute : ActorStateProviderAttribute
{...}

public class ReliableDictionaryActorStateProviderAttribute :
    ActorStateProviderAttribute
{...}

public class VolatileActorStateProviderAttribute : ActorStateProviderAttribute
{...}
```

Both the KVS and reliable dictionary actor state providers are transactional, reliable, distributed, and persistent. The volatile actor state provider is transactional, reliable, and distributed, but as its name implies, it is not persistent. This actor state provider only holds its state in the volatile memory of each replica's Windows process host. This means that if all of the Windows processes hosting the replica instances that hold your actor's state fail, you will lose all of your state.

Unlike Stateful Reliable Services, Service Fabric automatically manages your stateful Reliable Actor state for you. The mechanism behind Service Fabric's actor state management is much like WCF's durable services. Service Fabric loads your actor's state upon first activation and updates its state in a transaction at the end of each operation. It is important to remember that the Service Fabric never removes actor state. You can only clear it manually from within an actor operation or lifecycle method. You access state within your actor's operations through the `Actor<S>.State` property, as shown in [Example 11-18](#).

Example 11-18. Accessing stateful Reliable Actor state

```
class MyStatefulActor : Actor<MyState>, IMyActor
{
    public Task MyMethodAsync(MyDataContract someParam)
    {
        State.MyValue = someParam.MyValue;
        return Task.FromResult(true);
    }
    //More methods
}
```

Instead of using `OnDeactivateAsync()` to manually cleanup your state during Reliable Actor garbage collection, you can provide for explicit state cleanup by simply creating a `Cleanup()` operation on your actor and calling it at the appropriate time. Within `Cleanup()` you can then set your actor state to `null`:

```

interface IMyActor : IActor
{
    void Cleanup();
}
class MyActor : Actor<MyState>, IMyActor
{
    public void Cleanup()
    {
        State = null;
    }
    //More methods
}

```

The client, which should be a Reliable Service, will have to call `Cleanup()` once the workflow it performs is complete. If your workflow already possesses an explicit completion operation such as `SubmitInvoice()`, you should encapsulate the call to `Cleanup()` within your workflow's completion operation.

Hosting

Before Service Fabric can activate your Reliable Services or Reliable Actors you must first register them with the Service Fabric runtime. This makes Service Fabric aware of your service or actor instances and enables the Service Fabric runtime to manage and diagnose them. You register your services or actors during the hosting Windows process' startup in its `Main()` method. You first create an instance of the `FabricRuntime` class. You then register each service or actor with the fabric runtime. `FabricRuntime` is the service administration and diagnostics interface for the Service Fabric. After registration, you must keep the instance of `FabricRuntime` in memory by blocking the main thread. This ensures the host process will maintain an active interaction with the Service Fabric, as shown in [Example 11-19](#). Note that the `RegisterActor` extension method is currently part of the Service Fabric's programming model. Future versions may absorb this unnecessary extension into the `FabricRuntime` class.

Example 11-19. Registering services and actors with the Service Fabric

```

public sealed class FabricRuntime : IDisposable
{
    public static FabricRuntime Create();
    public void RegisterServiceType(string serviceTypeName,
                                    Type serviceTypeImplementation);
    //More methods
}

public static class ActorRegistration
{
    public static void RegisterActor(this FabricRuntime runtime, Type actorType);
    //More methods
}

```

```

class Program
{
    static void Main()
    {
        using(FabricRuntime fabricRuntime = FabricRuntime.Create())
        {
            fabricRuntime.RegisterServiceType("MyServiceType",typeof(MyService));
            fabricRuntime.RegisterActor(typeof(MyActor));
            fabricRuntime.RegisterActor(typeof(MyStatefulActor));

            Thread.Sleep(Timeout.Infinite);
        }
    }
}

```

The Client Programming Model

Service Fabric provides a client programming model to match each different service programming model it defines. For Reliable Services, there is a prescribed default client programming model that you use with default services and a configurable client programming model that you use with both API and WCF services. There is a straightforward, prescribed client programming model for Reliable Actors as well. It is interesting to note that both the Reliable Service default and the Reliable Actor client programming models are very similar. This reveals that there are effectively two client models: a prescribed model and configurable model.

Addressing

By default, all Service Fabric addressing is *virtual*. This means the address a client uses to access a given service or actor instance within the Service Fabric does not represent the physical address of the endpoint and that the address requires further resolution. This enables the Service Fabric to support location transparency along with easy and consistent deployment. Virtual addressing is also essential in allowing the Service Fabric to dynamically manage the cluster and balance the workload across it.

Service virtual addresses are unique by application and of the form:

`fabric:[/application name]/[service name]`

Note that a Service Fabric virtual address format explicit does not use `fabric://<...>`. This is to ensure that you do not mistakenly try to use a Service Fabric address as a normal endpoint address.

A client uses the service virtual address to create a proxy to a given service or actor instance. Internally, a client factory *resolves* against the Service Fabric to obtain the service endpoint's physical address within the cluster. Service Fabric endpoint resolution is similar in intent to WCF address discovery. The client factory sends a resolu-

tion request to the Service Fabric in the form of a query using the service's virtual address and in response receives the resolved service's endpoint information. This information contains the endpoint's physical address within the cluster. The client factory then uses this address to manufacture a proxy instance.

Default Reliable Service proxies

To establish connectivity to a default Reliable Service, you use the `ServiceProxy` class. `ServiceProxy` provides the `Create<T>()` factory method that creates the appropriate proxy instance for a service interface of type `T`, as shown in [Example 11-20](#). `ServiceProxy` will use the same prescribed configuration as that of the default Intranet service. Note that the `serviceName` parameter in `Create<T>()` is really the service virtual address, not the service name value from your application manifest file. Since you can only access default Intranet services from within a Service Fabric cluster, there is no way for you to specify the Service Fabric cluster's FQDN.

Note that [Example 11-20](#) does not close the proxy. Azure Service Fabric manages channels internally and cleans them up automatically. There is no need to call `Close` or `Dispose` on the proxy. In fact, `ServiceProxy` does not even provide these methods. This is true of all proxies that the Service Fabric provides.

Example 11-20. Calling a default Reliable Service

```
public interface IServiceProxy
{...}

public abstract class ServiceProxy : IServiceProxy
{
    public static T Create<T>(Uri serviceName);
    //More methods
}

class MyClient
{
    public async Task CallService()
    {
        Uri serviceAddress = new Uri("fabric:/myapplication/myservice");
        IMyService proxy = ServiceProxy.Create<IMyService>(serviceAddress);
        await proxy.MyMethodAsync(...);
        //No proxy.Close() necessary or possible
    }
}
```

API and WCF Reliable Service proxies

To access API and WCF Reliable Services, .NET clients use the `ServicePartitionClient<T>` class. `ServicePartitionClient<T>` provides you greater control over the

selection of connectivity, binding configuration, as well as the configuration for Service Fabric resolution by allowing you to provide an instance of `ICommunicationClientFactory<T>` during client construction:

```
public interface ICommunicationClient
{...}

public class ServicePartitionClient<T> : ... where T : ICommunicationClient
{
    public ServicePartitionClient(ICommunicationClientFactory<T> factory,
                                  Uri serviceName)
    {...}
    public Task InvokeWithRetryAsync(Func<TCommunicationClient, Task> func,
                                     CancellationToken cancellationToken, params Type[] doNotRetryExceptionTypes)
    {...}
    //More methods
}
```

As you will see later, you can use the additional control `ServicePartitionClient<T>` provides to specify a cluster endpoint base address, customize your WCF binding configuration to support legacy service scenarios, and control the way your client interacts with the Service Fabric.

You select the interaction mode for your proxy by providing the appropriate client factory. All client factories inherit from the `CommunicationClientFactoryBase<T>` class, which manufactures the appropriate `ICommunicationClient` for use with `ServicePartitionClient<T>`. For example, to create a client to a WCF Intranet service, you would use `WcfCommunicationClientFactory<T>` to manufacture instances of `WcfCommunicationClient<T>`, as shown in [Example 11-21](#). In the case of WCF services, `T` is a service contract.

Example 11-21. Calling a WCF Reliable Service

```
public abstract class CommunicationClientFactoryBase<T> : ...
    where T : ICommunicationClient
{...}

public class WcfCommunicationClient<T> : ICommunicationClient where T : class
{...}

public class WcfCommunicationClientFactory<T> :
    CommunicationClientFactoryBase<WcfCommunicationClient<T>> where T : class
{...}

class MyClient
{
    ServicePartitionClient<WcfCommunicationClient<IMyWcfContract>> GetProxy(
        string serviceAddress)
{}
```

```

        WcfCommunicationClientFactory<IMyWcfContract> factory =
            new WcfCommunicationClientFactory<IMyWcfContract>();

        ServicePartitionClient<WcfCommunicationClient<IMyWcfContract>> client =
            new ServicePartitionClient<WcfCommunicationClient<IMyWcfContract>>(factory,
                serviceAddress);

        return client;
    }

    public async Task CallWcfService()
    {
        Uri serviceAddress = new Uri("fabric:/myapplication/mywcfservice");

        ServicePartitionClient<WcfCommunicationClient<IMyWcfContract>> proxy =
            GetProxy(serviceAddress);

        await proxy.InvokeWithRetryAsync(client=>client.Channel.MyMethodAsync(...));
    }
}

```

ServicePartitionClient<T> does not present an interface-based programming model:

```

public class ServicePartitionClient<T> : ...
{
    public Task InvokeWithRetryAsync(Func<T,Task> func,
                                    params Type[] doNotRetryExceptionTypes);
    //More overloads
}

```

For WCF, it is better to encapsulate its usage in a wrapper class. [Example 11-22](#) shows the ServiceFabricClientBase<T> from *ServiceModelEx*.

Example 11-22. The ServiceFabricClientBase<T> wrapper class

```

public class ServiceFabricClientBase<T> :
    ServicePartitionClient<WcfCommunicationClient<T>> where T : class
{
    public ServiceFabricClientBase(Uri address) : base(
        new WcfCommunicationClientFactory<T>(),address)
    {}

    //More methods
}

class MyWcfProxy : ServiceFabricClientBase<IMyWcfContract>,IMyWcfContract
{
    public MyWcfProxy(Uri address) : base(address)
    {}

    public Task MyMethodAsync(MyDataContract someParam)
    {
        return InvokeWithRetryAsync(client=>client.Channel.MyMethodAsync(someParam));
    }
}

```

```

        }
    }

    class MyClient
    {
        public async Task CallWcfService()
        {
            Uri serviceAddress = new Uri("fabric:/myapplication/mywcfservice");

            IMyWcfContract proxy = new MyWcfProxy(serviceAddress);
            MyDataContract someParam = ...;
            await proxy.MyMethodAsync(someParam);
        }
    }
}

```

If you do not provide the WCF client factory with a binding configuration, as shown in [Example 11-22](#), the factory will use the default Reliable Service configuration based on `NetTcpBinding`. Consequently, these are not equivalent:

```

//Uses Default connectivity binding configuration
WcfCommunicationClientFactory<IMyWcfContract> factory =
    new WcfCommunicationClientFactory<IMyWcfContract>();

//Uses standard NetTcpBinding defaults
NetTcpBinding binding = new NetTcpBinding();
WcfCommunicationClientFactory<IMyWcfContract> factory =
    new WcfCommunicationClientFactory<IMyWcfContract>(binding);

```

A client factory resolves against the Service Fabric to obtain the service's endpoint physical address using an instance of the `ServicePartitionResolver` class. You can explicitly configure the resolver or use the default that resolves to `localhost`:

```

public class ServicePartitionResolver : ...
{
    public static ServicePartitionResolver GetDefault()
    {...}
    //More methods
}

```

These are equivalent:

```

WcfCommunicationClientFactory<IMyWcfContract> factory =
    new WcfCommunicationClientFactory<IMyWcfContract>();

WcfCommunicationClientFactory<IMyWcfContract> factory =
    new WcfCommunicationClientFactory<IMyWcfContract>(
        ServicePartitionResolver.GetDefault());

```

`ServicePartitionResolver` uses the `FabricClient` class to perform service resolution. `FabricClient` is the client administration and diagnostic interface for the Service Fabric. It represents the C# Service Fabric management API providing full administrative control over all Service Fabric subsystems. `FabricClient` provides

many constructor overloads that allow you to specify among other things the Service Fabric cluster FQDN and to control `FabricClient` settings, as shown in [Example 11-23](#).

Example 11-23. The FabricClient class

```
public sealed class FabricClientSettings
{
    public TimeSpan ConnectionInitializationTimeout
    {get;set;}
    public TimeSpan HealthOperationTimeout
    {get;set;}
    public TimeSpan KeepAliveInterval
    {get;set;}
    //More methods
}

public sealed class FabricClient : IDisposable
{
    public FabricClient(params string[] hostEndpoints)
    {...}
    public FabricClient(FabricClientSettings settings);
    {...}
    //More constructors and methods
}
```

`FabricClient` uses the public nested class `ServiceManagementClient` to obtain a service endpoint's physical address. Note again the `serviceName` parameter is the service virtual address, not service name value from your application manifest:

```
public sealed class ResolvedServiceEndpoint
{
    public string Address
    {get;}
    //More methods
}

public sealed class ResolvedServicePartition
{
    public ResolvedServiceEndpoint GetEndpoint()
    {...}
    //More methods
}

public sealed class FabricClient : ...
{
    public ServiceManagementClient ServiceManager
    {get;}
    public sealed class ServiceManagementClient : ...
    {
        public Task<ResolvedServicePartition> ResolveServicePartitionAsync(
```

```

        Uri serviceName);
    }
    //More methods
}
//More methods
}

```

`ServicePartitionResolver` provides many constructor overloads that allow you to specify the Service Fabric cluster FQDN, to control `FabricClient` settings and to provide an instance of `FabricClient` that you have already configured, as shown in [Example 11-24](#).

Example 11-24. The ServicePartitionResolver class

```

public delegate FabricClient CreateFabricClientDelegate();

public class ServicePartitionResolver : ...
{
    public ServicePartitionResolver(params string[] connectionEndpoints)
    {...}
    public ServicePartitionResolver(CreateFabricClientDelegate createFabricClient)
    {...}
    public ServicePartitionResolver(FabricClientSettings settings,
                                    params string[] connectionEndpoints)
    {...}
    //More constructors
}

```

These are equivalent:

```

ServicePartitionResolver serviceResolver = ServicePartitionResolver.GetDefault();

ServicePartitionResolver serviceResolver = new
    ServicePartitionResolver(()=>new FabricClient());

```

As are these, even though the parameter names differ between `hostEndpoints` and `connectionEndpoints`:

```

ServicePartitionResolver serviceResolver =
    new ServicePartitionResolver("localhost:19000");

ServicePartitionResolver serviceResolver =
    new ServicePartitionResolver(()=>new FabricClient("localhost:19000"));

```

Reliable Actor proxies

Much like the Default Reliable Service proxy, to establish connectivity to an instance of a Reliable Actor, you use the `ActorProxy` class. `ActorProxy` provides the `Create<T>()` factory method that creates an appropriate proxy instance for the actor interface of type `T`, as shown in [Example 11-25](#). `ActorProxy` will use the actor framework's prescribed connectivity configuration for the proxy. Note that the `serviceUri`

parameter in `Create<T>()` is again the service virtual address. You designate which actor instance the proxy should resolve to through a new `ActorId` construct. An `ActorId` is also a data contract. This allows Service Fabric to serialize it as internal state and to pass it as WCF context. Since you can only access Reliable Actors from within a Service Fabric cluster, there is no way for you to specify the Service Fabric cluster's FQDN.

Example 11-25. Calling a Reliable Actor

```
public interface IActorProxy
{...}

[DataContract]
public class ActorId : IEquatable<ActorId>
{...}

public abstract class ActorProxy : IActorProxy
{
    public static T Create<T>(ActorId actorId, Uri serviceUri) where T : IActor
    {...}
}

class MyClient
{
    public async Task CallActor()
    {
        Uri actorAddress = new Uri("fabric:/myapplication/myactor");

        IMyActor proxy = ActorProxy.Create<IMyActor>(ActorId.NewGuid(), actorAddress);
        MyDataContract someParam = ...;
        await proxy.MyMethodAsync(someParam);
    }
}
```

Reliable Actor proxies resolve internally against the Service Fabric to obtain an actor endpoint's physical address. They use the same approach as Reliable Services using internally a `FabricClient` instance, a `ServicePartitionResolver` instance, a `ServicePartitionClient<T>` instance, and a special internal WCF-based actor client factory.

Reliable Actor instancing

Unlike WCF, service behavior does not control Reliable Actor instancing. Instead, an actor's client actually controls it through explicit `ActorId` management. The Service Fabric conveys the `ActorId` you provide as context in each message it sends. The Service Fabric then uses the value within the `ActorId` and the service virtual address you provide to resolve to the correct actor instance in the cluster. This means many differ-

ent clients can use the same `ActorId` value and service address to resolve to the same actor instance and therefore its state. This type of `ActorId` usage allows Service Fabric to provide the concept of actor *shared sessions*.

As shown in [Example 11-26](#), the Service Fabric provides three types of `ActorId`: `Guid`, `long`, and `string`.

Example 11-26. The ActorId class

```
public enum ActorIdKind
{
    Long,
    Guid,
    String
}

[DataContract(Name = "ActorId")]
public class ActorId : IEquatable<ActorId>
{
    public ActorId(Guid id);
    public ActorId(string id);
    public ActorId(long id);

    public ActorIdKind Kind
    {get;}

    public static ActorId NewId();
    public Guid GetGuidId();
    public long GetLongId();
    public string GetStringId();
    //More methods
}
```

Through managing an actor's `ActorId`, you can model WCF-like instancing modes. You can model the `PerCall` instancing mode by providing a new `ActorId` value and new proxy instance for every call. This is easiest to achieve using the `Guid` `ActorId` type. Since the `Guid` type is the default `ActorId` type, `ActorId.NewId()` creates a new `ActorId` of type `Guid` with a new `Guid` value for every proxy instance:

```
Uri actorAddress = new Uri("fabric:/myapplication/myactor");

//Always resolve to different actor instances
IMyActor proxy = ActorProxy.Create<IMyActor>(ActorId.NewId(),actorAddress);
proxy.MyMethodAsync(...);

IMyActor proxy = ActorProxy.Create<IMyActor>(ActorId.NewId(),actorAddress);
proxy.MyMethodAsync(...);
```

You should prefer `PerCall` actor instancing for your stateless actors because it allows the Service Fabric to balance your actor's load evenly across all available instances.

While the Service Fabric will not prevent it, using a `PerCall` instancing mode with a stateful actor makes little sense, since your actor will never be able to accrue state.

You can model WCF's `Single` instancing mode by providing the same `ActorId` value for every call. Again, the singleton actor instance to which a specific `ActorId` value resolves is unique only at the application level, not globally within the Service Fabric cluster. A singleton actor instance is easiest to achieve by using the `string ActorId` type. `string`-based `ActorIds` allow you to create the notion of *named actors*. This gives the `ActorId` a key-like usage. Named actors allow you to create unique compound names from data identifiers that are relevant to your application. For example, you could define a named stateful actor that stores user state using an `ActorId` format such as `User[name][id]` with values like `UserMonty1111`. Service Fabric will always resolve this `ActorId` to the same instance of `MyActor` within `MyApplication` as shown:

```
Uri actorAddress = new Uri("fabric:/myapplication/myactor");
ActorId actorId = new ActorId("UserMonty1111");

//Always resolves to same actor instance
IMyActor proxy = ActorProxy.Create<IMyActor>(actorId,actorAddress);
```

You can use the same `ActorId` and service address with multiple proxy instances, all of which will resolve to same actor instance:

```
//Resolve to same actor instance
IMyActor proxy1 = ActorProxy.Create<IMyActor>(actorId,actorAddress);
IMyActor proxy2 = ActorProxy.Create<IMyActor>(actorId,actorAddress);
```

Whenever possible, you should use a named actors approach for your stateful actors. Named actors provide the most straightforward mechanism for accessing actor state throughout a system's lifetime. They also save you from having to devise additional mechanisms for storing and sharing `ActorId` instances as state.

You can also model WCF's `PerSession` instancing mode by using a unique `ActorId` for each proxy instance and making multiple calls against it as shown:

```
Uri actorAddress = new Uri("fabric:/myapplication/myactor");

IMyActor proxy = ActorProxy.Create<IMyActor>(ActorId.NewGuid(),actorAddress);
//Resolve to same actor instance
proxy.MyMethodAsync(...);
proxy.MyOtherMethodAsync(...);
```

Or you can use an incrementing long value to provide a shared but unique per-session value:

```
long id = 0;
Uri actorAddress = new Uri("fabric:/myapplication/myactor");

IMyActor proxy1 = ActorProxy.Create<IMyActor>(new ActorId(id++),actorAddress);
```

```
//Resolve to same actor instance
proxy1.MyMethodAsync(...);
proxy1.MyOtherMethodAsync(...);

IMyActor proxy2 = ActorProxy.Create<IMyActor>(new ActorId(id++),actorAddress);
//Resolve to same actor instance
proxy2.MyMethodAsync(...);
proxy2.MyOtherMethodAsync(...);
```

You would use this sort of ID to typically keep track of some kind of resource you consume during the session.

Modeling the Service Fabric

The most difficult aspect of any technology shift is the break with legacy code. This will be especially hard with the new programming models of the Service Fabric which represent such a departure from the past. Developers the world over will need to either migrate their legacy code or find clever ways of interoperating with it. For the reasons discussed earlier, we view it as vitally important that developers who are not presently using the Service Fabric can prepare their code for it. Ideally, you should be able to write code that is ready for the Service Fabric today while not using the Service Fabric. This is exactly what our `ServiceModelEx.ServiceFabric` framework enables. In fact, it provides parity with the programming model to the point that all you need to do is change the namespace at the top, compile to the Service Fabric, and deploy the old/new code.

`ServiceModelEx.ServiceFabric` not only achieves parity with the Service Fabric's programming model—in many respects, it also matches its behavior. This makes `ServiceModelEx.ServiceFabric` an operational Service Fabric with a very small footprint that allows you to deploy Service Fabric code to environments other than Windows Server, such as Windows desktops, Windows devices, or even smaller IoT devices that run Windows. Because it is operational, you can also deploy your `ServiceModelEx.ServiceFabric`-based code into production in your current environment to ensure it is operating correctly and to gain insights into how it might behave when you deploy it into the Azure Service Fabric.

`ServiceModelEx.ServiceFabric` allows you to learn how to program the Service Fabric immediately in .NET in a lightweight form with nothing to install. This lightweight footprint makes `ServiceModelEx.ServiceFabric` ideal for fine-grained testing of your services and actors without having to deploy every time to the Service Fabric to test these small incremental changes. It also enables the many intermediate integration tests you must perform to properly test your system as you build it.

While `ServiceModelEx.ServiceFabric` models many aspects of the Service Fabric, it stops short of modeling aspects of Service Fabric's infrastructure as these are not aspects of Service Fabric's programming model but instead purely aspects of its

operational capabilities. In particular, `ServiceModelEx.ServiceFabric` does not model nor provide support for Service Fabric's notion of *partitioning*. The partitioning of service state in the Service Fabric is an advanced topic and primarily related to providing Cloud environment scalability for stateful services and actors. Future versions of `ServiceModelEx.ServiceFabric` may provide basic support for service state partitioning through TCP-based discovery.

Modeling Approach

`ServiceModelEx.ServiceFabric` models aspects of both the Reliable Service and Reliable Actor programming models using central themes from this book. It also relies heavily on concepts from *ServiceModelEx*.

For state management, `ServiceModelEx.ServiceFabric` uses WCF's durable service framework, as well as the Volatile Resource Managers found in *ServiceModelEx*. To model Service Fabric's `ActorId` concept it uses WCF headers and its extensibility support for flowing headers as contexts. `ServiceModelEx.ServiceFabric` also uses WCF extensibility points to provide exception promotion, generic invocation, and message interception. Finally, `ServiceModelEx.ServiceFabric` introduces a number of new concepts such as a channel invoker, an instance context provider, and the advanced design principle of shared sessions.

`ServiceModelEx.ServiceFabric` bases its programming model on `InProcFactory` which provides a prescribed approach to hosting, connectivity, addressing, proxy creation, channel management, and WCF extensibility. All these are very similar in spirit to that of the Service Fabric. `ServiceModelEx.ServiceFabric` improves testability by enabling support for type decoupling.

The `ServiceModelEx` compiler switch enables seamless code transfer between environments without code modification:

```
#if ServiceModelEx
    using ServiceModelEx.ServiceFabric;
    using ServiceModelEx.ServiceFabric.Services;
    using ServiceModelEx.ServiceFabric.Actors;
#else
    using Microsoft.ServiceFabric;
    using Microsoft.ServiceFabric.Services;
    using Microsoft.ServiceFabric.Actors;
#endif

#if ServiceModelEx
    [Serializable]
    [ApplicationManifest("MyApplication", "MyActor")]
#endif
```

Note that you will only need this compilation flag to conditionally shield the different namespaces and a few class attributes. Nothing in the `ServiceModelEx.ServiceFabric`'s programming model requires this compilation switch.

The ApplicationManifestAttribute

The `ApplicationManifestAttribute` class shown in [Example 11-27](#) is found in the `ServiceModelEx.ServiceFabric` namespace. It is the only aspect of `ServiceModelEx.ServiceFabric`'s programming model that is not part of the Azure Service Fabric. `ApplicationManifestAttribute` is a stand-in for the Service Fabric's application manifest configuration file. In `ServiceModelEx.ServiceFabric` there is no need to also model the config file. `ServiceModelEx.ServiceFabric` uses the `ApplicationManifestAttribute` to enforce compatibility between your proxies, services, and actors. It uses the values of `ApplicationManifestAttribute` to ensure that your virtual addressing will be correct in Azure Service Fabric. It also enables `ServiceModelEx.ServiceFabric` to perform the same manifest validations as Azure Service Fabric, such as validating the address format, validating interface requirements, and ensuring the service or actor exists and that it possesses the proper base class. As shown in [Example 11-27](#), you can apply multiple `ApplicationManifestAttributes` to a given service or actor. This allows you to model the deployment of the same service or actor to different applications.

Example 11-27. The ApplicationManifestAttribute class

```
[AttributeUsage(AttributeTargets.Class,AllowMultiple = true)]
public class ApplicationManifestAttribute : Attribute
{
    public string ApplicationName
    {get;}
    public string ServiceName
    {get;}

    public ApplicationManifestAttribute(string applicationName,string serviceName)
    {...}
}
```

Modeling Reliable Services

The namespace `ServiceModelEx.ServiceFabric.Services` provides support for modeling the most common Reliable Service scenarios that promote best practice service-oriented design. With this in mind, `ServiceModelEx.ServiceFabric` only provides modeling support for compositional services that are stateless Intranet services. For compositional services, `ServiceModelEx.ServiceFabric` supports both the Default and WCF programming models.

`ServiceModelEx.ServiceFabric` does not provide modeling support for API services because until Azure Service Fabric's on-premises cluster matures, you are better off serving on-premises public API endpoints through IIS. And since your ASP.NET Web API code is already 100% compatible with Azure Service Fabric API services, there is no real need to model it. The only difference will be a new hosting scheme. And finally, as explained previously, your API Reliable Services should just call your Intranet compositional services anyway.

Modeling Reliable Service proxies and services

`ServiceModelEx.ServiceFabric`'s Reliable Service proxies and services are the same as those within Azure Service Fabric both in programming model and behavior, as shown in [Example 11-28](#) and [Example 11-29](#), including not having to close the proxy. To model them, you must use the same service virtual address in your proxy as you define in the service's `ApplicationManifestAttribute`. And, as shown in [Example 11-28](#), you must also make your default service interfaces service contracts. This last modeling requirement does not in any way affect service behavior in Azure Service Fabric. You can leave the service contract attributes, remove them later after your transition is complete, or use the compiler switch to conditionally shield them from compilation in your Azure Service Fabric code.

Note that when using `ServiceModelEx.ServiceFabric` you still provide the service virtual address to resolve to your Reliable Services, as shown in [Example 11-28](#) and [Example 11-29](#). This ensures that your client code will also transfer to the Azure Service Fabric directly without change to even your service addresses. Internally, `ServiceModelEx.ServiceFabric`'s proxies and service hosting use a simple addressing convention based on your service's virtual address to define the address for each service's endpoint.

Example 11-28. Modeling the default Reliable Service

```
public static class ServiceProxy
{
    public static I Create<I>(Uri serviceAddress)
    {...}
}

public interface IService
{...}

[ServiceContract]
interface IMyService : IService
{
    [OperationContract]
    Task MyMethodAsync(MyDataContract request);
}
```

```

class MyClient
{
    public async Task CallDefaultService()
    {
        Uri serviceAddress = new Uri("fabric:/myapplication/myservice");
        IMyService proxy = ServiceProxy.Create<IMyService>(serviceAddress);
        ...
    }
}

#if ServiceModelEx
[ApplicationManifest("MyApplication", "MyService")]
#endif
class MyService : StatelessService, IMyService
{
    //IMyService methods
    protected override ICommunicationListener CreateCommunicationListener()
    {
        ServiceCommunicationListener<IMyService> communicationListener =
            new ServiceCommunicationListener<IMyService>(this);

        return communicationListener;
    }
}

```

As shown in [Example 11-29](#), ServiceModelEx.ServiceFabric extends its WCF Reliable Service support to include the ServicePartitionResolver class. While Service ModelEx.ServiceFabric's ServicePartitionResolver does not yet perform address resolution like Azure Service Fabric, it does allow you to provide programmatically the base address of your Intranet Reliable Services.

Example 11-29. Modeling the WCF Reliable Service

```

public class ServicePartitionResolver
{
    public ServicePartitionResolver(string fabricEndpoint)
    {...}
}

public interface ICommunicationClient
{...}
public class ServicePartitionClient<T> : ... where T : ICommunicationClient
{...}
public class WcfCommunicationClient<T> : ICommunicationClient where T : class
{...}
public class WcfCommunicationClientFactory<T> : ... where T : class
{...}

public class ServiceFabricClientBase<T> :

```

```

    ServicePartitionClient
    <WcfCommunicationClient<T>> where T : class
{...}

class MyWcfProxy : ServiceFabricClientBase<IMyWcfContract>, IMyWcfContract
{...}

class MyClient
{
    public async Task CallWcfService()
    {
        Uri serviceAddress = new Uri("fabric:/myapplication/mywcfservice");
        IMyWcfContract proxy = new MyWcfProxy(serviceAddress);
        await proxy.MyMethodAsync(...);
        //No close necessary
    }
}

#if ServiceModelEx
[ApplicationManifest("MyApplication", "MyWcfService")]
#endif
class MyWcfService : StatelessService, IMyWcfContract
{
    //IMyWcfContract methods
    public Task MyMethodAsync(MyDataContract request)
    {...}
    protected override ICommunicationListener CreateCommunicationListener()
    {
        WcfCommunicationListener communicationListener
            = new WcfCommunicationListener(typeof(IMyWcfContract), this);
        ...
        return communicationListener;
    }
}

```

Modeling Reliable Actors

The namespace `ServiceModelEx.ServiceFabric.Actors` provides full programming model support for Reliable Actors including the actor concurrency model, `ActorId`, proxy, lifecycle events, state management, named actors and actor garbage collection. `ServiceModelEx.ServiceFabric` also provides additional support for completing an actor instance which also removes its state from the backing state store.

Modeling actor concurrency

Reliable Actors in Azure Service Fabric use *turn-based concurrency*. This means only a single thread may be operating within the actor at a time and that each actor operation runs to completion before the next may begin. Unlike synchronizing calls using manual locks in a multithreaded scenario, turn-based concurrency buffers all requests as they arrive and only presents a single request to the actor at a time.

By default, Reliable Actors in Azure Service Fabric allow reentrancy. In `ServiceModeEx.ServiceFabric` Reliable Actors use WCF's `ConcurrencyMode.Single` to enforce turn-based concurrency. `ServiceModelEx.ServiceFabric` Reliable Actors do not allow reentrancy since it is a poor design practice to rely on it.

Modeling actor lifecycle events

Reliable Actor lifecycle events are the same in `ServiceModelEx.ServiceFabric` as they are in Azure Service Fabric. `ServiceModelEx.ServiceFabric` calls the `OnActivateAsync()` method upon the first activation of an actor and calls the `OnDeactivateAsync()` method upon the expiration of the garbage collection idle period. Inheritors can override these lifecycle events to perform actor lifecycle related actions. Unlike Azure Service Fabric, there is no actor instance to remove upon actor garbage collection because all actors in `ServiceModelEx.ServiceFabric` effectively model `PerCall` instancing.

Modeling proxies and actors

`ServiceModelEx.ServiceFabric`'s Reliable Actor proxies and services are the same as those within Azure Service Fabric both in programming model and behavior, as shown in [Example 11-30](#). To model them, you must use the same service virtual address in your proxy as you define in the actor's `ApplicationManifestAttribute`. And, as shown in [Example 11-30](#), you must also make your actor interfaces service contracts. Additionally, to support the demarcation of session boundaries, `ServiceModelEx.ServiceFabric` necessitates that your actors require a session. As with Default Reliable Services, this last modeling requirement does not in any way affect actor behavior in Azure Service Fabric. You can leave the service contract attributes, remove them later after your transition is complete, or use the compiler switch to conditionally shield them from compilation in your Azure Service Fabric code.

Example 11-30. Modeling Reliable Actors

```
public static class ActorProxy
{
    public static I Create<I>(ActorId actorId, Uri serviceAddress)
        where I : class, IActor
    //More methods

}
public interface IActor
{...}

[ServiceContract(SessionMode = SessionMode.Required)]
interface IMyActor : IActor
{
    [OperationContract]
```

```

Task MyMethodAsync(MyDataContract someParam);

[OperationContract]
Task<MyDataContract> MyOtherMethodAsync(MyDataContract someParam);
}

class MyClient
{
    public async Task CallActor()
    {
        Uri actorAddress = new Uri("fabric:/myapplication/myactor");
        IMyActor proxy = ActorProxy.Create<IMyActor>(ActorId.NewId(),actorAddress);
        ...
    }
}

#if ServiceModelEx
[ApplicationManifest("MyApplication", "MyActor")]
#endif
class MyActor : Actor,IMyActor
{...}

```

Modeling actor instancing

The `ActorId` concept in `ServiceModelEx.ServiceFabric` is the same as in Azure Service Fabric both in programming model and behavior. `ServiceModelEx.ServiceFabric` also passes it as context so it is also a data contract, as shown. `ServiceModelEx.ServiceFabric` designates the `ActorId` as serializable to support storage to disk, and as you will see later, `ServiceModelEx.ServiceFabric`'s `ActorId` contains additional internal properties to support shared sessions and durable service session association:

```

[Serializable]
[DataContract]
public class ActorId : IEquatable<ActorId>
{...}

```

`ServiceModelEx.ServiceFabric`'s modeling of Reliable Actor instancing is the same as in Azure Service Fabric. `ActorProxy` instance lifetime does not dictate actor instance lifetime. An `ActorProxy` instance does not designate a unique session. `ActorProxy` instances associate with an actor instance solely through an `ActorId`. And multiple `ActorProxy` instances with same the `ActorId` resolve to the same actor instance.

`ServiceModelEx.ServiceFabric`'s proxy instance and channel cleanup is behaviorally the same as in Azure Service Fabric. There is no need to call `Close` or `Dispose` on the proxy. In fact, `ActorProxy` does not even provide these methods. Azure Service Fabric manages channels internally and cleans them up automatically. To provide the same capability, `ServiceModelEx.ServiceFabric` uses a novel `ActorChannelInvoker` concept. `ServiceModelEx.ServiceFabric`'s `ActorChannelInvoker` receives the

appropriate `ChannelFactory<I>` through its constructor and uses the factory to automatically create and close the channel upon each proxy call. This approach does not appreciably affect actor responsiveness because in the case of the `NetNamedPipeBinding` there is no connection to establish. This chapter does not show `ActorChannelInvoker` here because its design and implementation have nothing to do with WCF (but it's available with `ServiceModelEx`). Its implementation does make use of some advanced .NET concepts internally such as `RealProxy`. You may also find the channel invocation concept offers value in its own right, as a channel invoker can still provide all the extensible capabilities of proxy base classes while freeing you from having to create concrete proxy classes for each service interface in your system.

`ServiceModelEx.ServiceFabric`'s `ActorProxy` provides the same automatic retry semantics as you will find in Azure Service Fabric. `ActorProxy` implicitly retries on transient connection errors. `ActorChannelInvoker` performs the retry protocol.

Modeling actor state

`ServiceModelEx.ServiceFabric` Reliable Actors support both the stateless and stateful programming models. As stated previously, stateless actors in Azure Service Fabric can still accrue implicit state. `ServiceModelEx.ServiceFabric` stateless actors instead use the best practice `PerCall` instancing so they cannot accrue implicit instance state. Aside from the `ApplicationManifestAttribute`, the programming model is identical:

```
#if ServiceModelEx
[ApplicationManifest("MyApplication", "MyActor")]
#endif
class MyActor : Actor, IMyActor
{...}
```

In Azure Service Fabric, while actor state must be serializable as a data contract, the actor class itself does not need to be serializable. In `ServiceModelEx.ServiceFabric`, you must also mark the stateful actor class as serializable to support durable service state management:

```
#if ServiceModelEx
[Serializable]
[ApplicationManifest("MyApplication", "MyActor")]
#endif
class MyStatefulActor : Actor<MyState>, IMyActor
{...}
```

In your `ServiceModelEx.ServiceFabric` modeling it is also best practice to ensure only the framework may activate an actor. To enforce this, you should mark all actor operations with `IsInitiating = false`:

```
[ServiceContract(SessionMode = SessionMode.Required)]
interface IMyActor : IActor
```

```

{
    [OperationContract(IsInitiating = false)]
    Task MyMethodAsync(MyDataContract someParam);

    [OperationContract(IsInitiating = false)]
    Task<MyDataContract> MyOtherMethodAsync(MyDataContract someParam);
}

```

To provide robust, production ready state management, `ServiceModelEx.ServiceFabric` supports all three actor state providers, as shown in [Example 11-31](#). `ServiceModelEx.ServiceFabric` leverages WCF's durable service framework to model Azure Service Fabric's Reliable Actor state management. For the persistent providers KVS and reliable dictionary, `ServiceModelEx.ServiceFabric` uses durable services backed by SQL Server for the state store. In Azure Service Fabric, the KVS actor state provider is the default. In `ServiceModelEx.ServiceFabric`, the volatile actor state provider is the default. `ServiceModelEx.ServiceFabric` makes the volatile state provider the default to avoid requiring a dependency on SQL Server out of the box. The volatile actor state provider uses the `TransactionalInstanceStore` VRM to provide transactional, in-memory state storage.

Example 11-31. Modeling the Reliable Actor state providers

```

[AttributeUsage(AttributeTargets.Class)]
public abstract class ActorStateProviderAttribute : Attribute
{}

public class VolatileActorStateProviderAttribute : ActorStateProviderAttribute,
    IServiceProviderBehavior
{...}

public class KvsActorStateProviderAttribute : ActorStateProviderAttribute,
    IServiceProviderBehavior
{...}

public class ReliableDictionaryActorStateProviderAttribute :
    ActorStateProviderAttribute, IServiceProviderBehavior
{...}

#if ServiceModelEx
[Serializable]
[ApplicationManifest("MyApplication", "MyStatefulActor")]
#endif
[ReliableDictionaryActorStateProvider]
class MyStatefulActor : Actor<MyState>, IMyActor
{...}

```

Modeling named actors

ServiceModelEx.ServiceFabric's support for named actors is the same as in Azure Service Fabric with the added benefit explained previously that stateless named actors cannot accrue implicit state. In ServiceModelEx.ServiceFabric, a stateful named actor implies that it may share session within its application scope with many clients. To enable the shared session capability, ServiceModelEx.ServiceFabric associates the named ActorId to a durable service instance Id within an application scope. This allows the same ActorId name and service address to resolve to the same durable service instance state for a given actor. The ActorInstanceContextProvider class enables shared session to durable service instances by controlling the durable service instance Id within each actor request. ServiceModelEx.ServiceFabric must control instance context provisioning because it is the InstanceContext concept within WCF that manages durable service instance sessions internally.

As shown in [Example 11-32](#), ActorInstanceContextProvider installs itself as an IEndpointBehavior in IEndpointBehavior.ApplyDispatchBehavior() through the DispatchRuntime property of the EndpointDispatcher class. The StatefulActorServiceBehavior class then adds ActorInstanceContextProvider into the endpoint behavior collection of each endpoint for the actor. To ensure consistent concurrency management across all of the actor's endpoints, StatefulActorServiceBehavior uses the same instance of ActorInstanceContextProvider for all of them.

Example 11-32. Installing the ActorInstanceContextProvider class

```
public interface IInstanceContextProvider
{
    InstanceContext GetExistingInstanceContext(Message message,
                                              IContextChannel channel);
    void InitializeInstanceContext(InstanceContext instanceContext, Message message,
                                  IContextChannel channel);
    //More methods
}

public sealed class DispatchRuntime
{
    public IInstanceContextProvider InstanceContextProvider
    {get;set;}
    //More methods
}

class ActorInstanceContextProvider : IInstanceContextProvider, IEndpointBehavior
{
    IInstanceContextProvider m_PreviousProvider = null;
    public void ApplyDispatchBehavior(ServiceEndpoint endpoint,
                                      EndpointDispatcher endpointDispatcher)
    {

```

```

        _PreviousProvider =
            endpointDispatcher.DispatchRuntime.InstanceContextProvider;
        endpointDispatcher.DispatchRuntime.InstanceContextProvider = this;
    }
    //More methods
}

class StatefulActorServiceBehavior : IServiceBehavior
{
    public void Validate(ServiceDescription serviceDescription,
                        ServiceHostBase serviceHostBase)
    {
        EnforceStatefulActorEndpointBehaviorPolicy(serviceDescription);
        //More methods
    }

    void EnforceStatefulActorEndpointBehaviorPolicy(
        ServiceDescription serviceDescription)
    {
        ActorInstanceContextProvider contextProvider =
            new ActorInstanceContextProvider();

        foreach(ServiceEndpoint endpoint in serviceDescription.Endpoints)
        {
            //All endpoints for a given actor must share the same
            //ActorInstanceContextProvider
            endpoint.EndpointBehaviors.Add(contextProvider);
            ...
        }
        ...
    }
    //More methods
}

```

`ActorInstanceContextProvider` uses a single `Mutex` to coordinate consistency across `InstanceContext` initialization and retrieval, as shown in [Example 11-33](#).

Example 11-33. Managing Reliable Actor instance access

```

class ActorInstanceContextProvider : IInstanceContextProvider, ...
{
    readonly Mutex m_InstanceAccess = new Mutex(false);
    void InitializeInstanceContext(InstanceContext instanceContext, Message message,
                                  IContextChannel channel)
    {
        try
        {...}
        finally
        {
            m_InstanceAccess.ReleaseMutex();
        }
    }
}

```

```

}

InstanceContext GetExistingInstanceContext(Message message,
                                         IContextChannel channel)
{
    try
    {
        m_InstanceAccess.WaitOne();
        ...
        InstanceContext context =
            _PreviousProvider.GetExistingInstanceContext(message,channel);

        if(context != null)
        {
            m_InstanceAccess.ReleaseMutex();
        }
        return context;
    }
    catch(Exception exception)
    {
        ...
        m_InstanceAccess.ReleaseMutex();
    }
}
//More methods
}

```

To manage ActorIds, `ServiceModelEx.ServiceFabric` provides an internal Actor Manager concept. `ActorManager` supports managing ActorIds for both volatile and persistent stateful actors, as shown in [Example 11-34](#). For actors that use the volatile actor state provider, `ActorManager` keeps ActorIds in memory only. For actors that use either of the persistent actor state providers, `ActorManager` keeps ActorIds in memory and also persists them to disk. This ensures the durable service instance associations for your stateful actors will survive system shutdown and even many types of system failure. `ActorInstanceContextProvider` then uses `ActorManager` to resolve ActorIds with durable service instances.

Example 11-34. The ActorManager class

```

public class ActorId : ...
{
    public Guid InstanceId
    {get;internal set;}
    //More methods
}

[Serializable]
class ActorState
{

```

```

public ActorId ActorId
{get;set;}
public bool StatefulActor
{get;set;}
public DateTime IdleStartTime
{get;set;}
//More methods
}

static class ActorManager
{
    static Dictionary<ActorId,ActorState> m_VolatileInstances =
        new Dictionary<ActorId,ActorState>(...);
    static Dictionary<ActorId,ActorState> m_PersistentInstances =
        new Dictionary<ActorId,ActorState>(...);

    public static Guid GetInstance(ActorId actorId)
    {...}
    public static void SaveInstance(ActorId actorId,ActorState state,
                                    bool persistent)
    {...}
    internal static void RemoveInstance(ActorId actorId)
    {...}
    //More methods
}

```

Modeling actor garbage collection

ServiceModelEx.ServiceFabric provides Reliable Actor garbage collection using the same declarative model as Azure Service Fabric through the `ActorGarbageCollectionAttribute`, as shown in [Example 11-35](#). In `ServiceModelEx.ServiceFabric`, stateful actor instances are never kept in memory so there is nothing to cleanup. Instead, you can use the `OnDeactivateAsync()` method to perform periodic actor cleanup. `ServiceModelEx.ServiceFabric`'s actor garbage collection uses the same default idle timeout of 1 hour and a default scan interval of 1 minute as Azure Service Fabric. In `ServiceModelEx.ServiceFabric`, `ActorManager` performs actor idle timeout monitoring.

Example 11-35. Modeling the ActorGarbageCollectionAttribute

```

[AttributeUsage(AttributeTargets.Class)]
public class ActorGarbageCollectionAttribute : Attribute
{
    public long IdleTimeoutInSeconds
    {get;set;}
    public long ScanIntervalInSeconds
    {get;set;}
    public ActorGarbageCollectionAttribute()

```

```
    {...}  
}
```

Actor instance completion

Unlike Azure Service Fabric, `ServiceModelEx.ServiceFabric` allows you to specify an actor operation that completes the actor instance by removing its state from the backing state store after the operation completes. You must be able to systematically remove actor state. Depending on the design of your system and its state usage profile, you may need to occasionally clean up stale actors or enforce an actor cleanup policy. You use `CompletesActorInstanceAttribute` to designate the actor operation clients should call to complete an actor instance. `ServiceModelEx.ServiceFabric` will first call your actor's `OnCompleteAsync()` method before completing the actor instance. A subsequent call using the same `ActorId` value will activate a new actor instance:

```
[AttributeUsage(AttributeTargets.Method)]  
public class CompletesActorInstanceAttribute : Attribute  
{}
```

Modeling Hosting

`ServiceModelEx.ServiceFabric`'s programming model for Reliable Service and Reliable Actor hosting is the same as Azure Service Fabric in both programming model and behavior. As shown in [Example 11-36](#), upon host process startup in the host's `Main()` method, you must register your service and actor types with the `FabricRuntime`. You must also keep `FabricRuntime` in memory because `ServiceModelEx.ServiceFabric` uses it internally to resolve service and actor types. `ServiceModelEx.ServiceFabric` hosting performs the same validations as Azure Service Fabric such as that the manifest attribute exists, that the manifest names are correct, that there is no type or name duplication within manifests, and confirms that actors have a single interface.

Example 11-36. Modeling Service Fabric hosting

```
public sealed class FabricRuntime : IDisposable  
{  
    public static FabricRuntime Create();  
    public void RegisterServiceType(string serviceTypeName,  
                                    Type serviceTypeImplementation);  
    //More methods  
}  
  
public static class ActorRegistration  
{  
    public static void RegisterActor(this FabricRuntime runtime, Type actorType);  
    //More methods  
}
```

```

}

class Program
{
    static void Main()
    {
        using(FabricRuntime fabricRuntime = FabricRuntime.Create())
        {
            fabricRuntime.RegisterServiceType("MyServiceType",typeof(MyService));
            fabricRuntime.RegisterServiceType("MyWcfServiceType",
                typeof(MyWcfService));

            fabricRuntime.RegisterActor(typeof(MyActor));
            fabricRuntime.RegisterActor(typeof(MyStatefulActor));

            Thread.Sleep(Timeout.Infinite);
        }
    }
}

```

Testing

While it has nothing to do with WCF directly, `ServiceModelEx.ServiceFabric` would not be complete if it lacked a simple, effective and most of all efficient way for you to test your Reliable Services and Reliable Actors. To that end, `ServiceModelEx.ServiceFabric` includes `ServiceModelEx.ServiceFabric.Test`, a lightweight, service-oriented test framework compatible with the *Spiral of Test*. The spiral of test gears its testing approach toward systems that are built using a service-oriented programming model. In these systems, the spiral of test allows you to do *bottom-up testing* with greater control and precision. To support the spiral of test, `ServiceModelEx.ServiceFabric.Test` allows you to go well beyond the simple unit testing of your services by allowing you to incrementally increase the scope of the *unit under test* (UUT). You can increase the scope from plain old C# object (POCO), to actor, to service, through multiple levels of intermediate integration testing all the way to your top-level use cases. This approach allows you to focus your testing on the most problematic aspect of service-oriented systems: service integration.

Performing the many levels of intermediate integration testing necessary to properly test a service-oriented system often involves the cumbersome burden of standing up and tearing down all of the system machinery many times during test iterations. In addition, current testing tools focus primarily on simple unit testing. They lack the capabilities to efficiently orchestrate numerous intermediate integration tests.

`ServiceModelEx.ServiceFabric.Test` solves the many problems involved with service integration testing by using `InProcFactory` to *mock the service environment*, as shown in [Example 11-37](#). To mock the service environment, `ServiceTestBase` uses reflection and late binding techniques to create a `CreateInstance<S,I>()` method

definition and dynamically invoke it. `ServiceTestBase` also provides support for dynamically registering and unregistering test mocks each time you test a service or actor. This technique avoids the unwanted side effect common to mocking frameworks of accumulating test mocks as tests run.

Example 11-37. Mocking the service environment using InProcFactory

```
public abstract class ServiceTestBase
{
    static MethodInfo m_CreateInstanceDefinition = null;
    static ServiceTestBase()
    {
        m_CreateInstanceDefinition = typeof(InProcFactory).GetMethod(
            "CreateInstance", new Type[0]).GetGenericMethodDefinition();
    }

    //Resolve service type by contract type 'I'.
    Type GetServiceType<I>()
    {...}
    //Resolve actor type by contract type 'I'.
    Type GetActorType<I>()
    {...}

    //Manage mocks within a container.
    void RegisterMocks(params object[] mocks)
    {...}
    void UnregisterMocks(params object[] mocks)
    {...}

    void MockEnvironment<I>(I callee, Action<I> callerMock, params object[] mocks)
    {
        RegisterMocks(mocks);
        using (callee as IDisposable)
        {
            callerMock(callee);
        }
        UnregisterMocks();
    }

    void MockServiceEnvironment<I>(Type targetType, Action<I> callerMock,
                                    params object[] mocks)
    {
        MethodInfo createInstance = m_CreateInstanceDefinition.MakeGenericMethod(
            targetType, typeof(I));
        I proxy = (I)createInstance.Invoke(null, new object[0]);
        MockEnvironment<I>(proxy, callerMock, mocks);
    }
    //More members
}
```

By mocking the service environment, `ServiceModelEx.ServiceFabric.Test` provides a simple testing programming model that supports multiple levels of granular service testing:

```
abstract class ServiceTestBase
{
    public void TestActor<I>(Action<I> callerMock,params object[] mocks)
        where I : class
    {
        MockServiceEnvironment<I>(GetActorType<I>(),callerMock,Mocks);
    }

    public void TestService<I>(Action<I> callerMock,params object[] mocks)
        where I : class
    {
        MockServiceEnvironment<I>(GetServiceType<I>(),callerMock,Mocks);
    }
    //More members
}
```

`ServiceModelEx.ServiceFabric.Test` extends the testing programming model to allow you to test your services with or without the many service-oriented aspects in play:

```
abstract class ServiceTestBase
{
    void MockPocoEnvironment<I>(Type targetType,Action<I> callerMock,
        params object[] mocks)
    {
        I poco = (I)Activator.CreateInstance(targetType);
        MockEnvironment<I>(poco,callerMock,Mocks);
    }

    public void TestActorPoco<I>(Action<I> callerMock,params object[] mocks)
        where I : class
    {
        MockPocoEnvironment<I>(GetActorType<I>(),callerMock,Mocks);
    }

    public void TestServicePoco<I>(Action<I> callerMock,params object[] mocks)
        where I : class
    {
        MockPocoEnvironment<I>(GetServiceType<I>(),callerMock,Mocks);
    }
    //More members
}
```

The resulting testing programming model is compatible with existing unit test tools, such as Visual Studio or NUnit and common mocking frameworks such as Moq, as shown in [Example 11-38](#).

Example 11-38. Testing with common tools

```
class IntegrationTestHarness : ServiceTestBase
{}

[TestClass]
class SystemTests
{
    IntegrationTestHarness harness = null;

    [TestInitialize]
    public void Setup()
    {
        harness = new IntegrationTestHarness();
    }

    [TestMethod]
    public async void Test_IMyActor_MyMethodAsync_NoMocks()
    {
        Action<IMyActor> callerMock = (proxy)=>
        {
            await proxy.MyMethodAsync();
        };
        harness.TestActor<IMyActor>(callerMock);
    }
    [TestMethod]
    public async void Test_IMyService_MyMethodAsync_NoMocks()
    {
        Action<IMyService> callerMock = (proxy)=>
        {
            await proxy.MyMethodAsync();
        };
        harness.TestService<IMyService>(callerMock);
    }
    [TestMethod]
    public async void Test_IMyService_MyMethodASync_AllMocks()
    {
        IMyActor mockedActor = new Mock<IMyActor>();
        mockedActor.Setup(...);

        Action<IMyService> callerMock = (proxy)=>
        {
            await proxy.MyMethod();
        };
        harness.TestService<IMyService>(callerMock, mockedActor);
    }
}
```

Most importantly, the testing programming model supports multiple levels of intermediate integration testing. To enable this style of granular, compositional testing, `ServiceModelEx.ServiceFabric`'s `ActorProxy` uses late binding techniques similar

to those shown in [Example 11-37](#) to resolve for the type I in its `Create<I>()` method. For example, for the service and actor interactions defined as:

```
[ServiceContract]
interface IMyOtherActor : IActor
{
    [OperationContract]
    Task MyMethodAsync();
}

class MyOtherActor : Actor,IMyOtherActor
{
    public Task MyMethodAsync()
    {...}
}

class MyActor : Actor,IMyActor
{
    public async Task MyMethodAsync(...)
    {
        IMyOtherActor otherProxy = ActorProxy.Create<IMyOtherActor>(...);
        await otherProxy.MyMethodAsync(...);
    }
}

class MyService : StatelessServiceBase,IMyService
{
    public async Task MyMethodAsync(...)
    {
        IMyActor proxy = ActorProxy.Create<IMyActor>(...);
        await proxy.MyMethodAsync(...);
    }
}
```

You can mix and match mocked actors with real ones. This testing technique allows you to focus in on specific interactions with high iteration testing. In this case, `MyActor` will be a real service, while `MyOtherActor` will be a mock:

```
[TestClass]
class SystemTests
{
    [TestMethod]
    public void Test_IMyService_MyMethodAsync_MyOtherActorMock()
    {
        IMyOtherActor mockedOtherActor = new Mock<IMyOtherActor>();
        mockedOtherActor.Setup(...);

        Action<IMyService> callerMock = (proxy)=>
        {
            proxy.MyMethod();
        };
        harness.TestService<IMyService>(callerMock, mockedOtherActor);
    }
}
```

```
    }
    //More members
}
```

It is the combination of these techniques that allows you to do bottom-up testing with greater control and precision and full support of the Spiral of Test. As shown in [Example 11-39](#), each group of test methods represents a turn of the spiral by increasing the scope of the UUT by one level.

Example 11-39. The Spiral of Test

```
[TestClass]
class ServiceTests
{
    [TestMethod]
    public void Test_IMyOtherActor_MethodAsync_AsPoco()
    {...}
    [TestMethod]
    public void Test_IMyOtherActor_MethodAsync_AsService()
    {...}

    [TestMethod]
    public void Test_IMyActor_MethodAsync_AsPoco_AllMocks()
    {...}
    [TestMethod]
    public void Test_IMyActor_MethodAsync_AsPoco_NoMocks()
    {...}
    [TestMethod]
    public void Test_IMyActor_MethodAsync_AsService_AllMocks()
    {...}
    [TestMethod]
    public void Test_IMyActor_MethodAsync_AsService_NoMocks()
    {...}

    [TestMethod]
    public void Test_IMyService_MethodAsync_AsPoco_AllMocks()
    {...}
    [TestMethod]
    public void Test_IMyService_MethodAsync_AsPoco_MyActorMock()
    {...}
    [TestMethod]
    public void Test_IMyService_MethodAsync_AsPoco_MyOtherActorMock()
    {...}
    [TestMethod]
    public void Test_IMyService_MethodAsync_AsPoco_NoMocks()
    {...}
    [TestMethod]
    public void Test_IMyService_MethodAsync_AsService_AllMocks()
    {...}
    [TestMethod]
    public void Test_IMyService_MethodAsync_AsService_MyActorMock()
```

```
{...}  
[TestMethod]  
public void Test_IMyService_MyMethodAsync_AsService_MyOtherActorMock()  
{...}  
[TestMethod]  
public void Test_IMyService_MyMethodAsync_AsService_NoMocks()  
{...}  
//Etc.  
}
```



Currently, you cannot use these valuable testing techniques in Azure Service Fabric because the current development version of the Service Fabric does not provide for type decoupling (dependency inversion) in any of its proxies or during service instantiation. The production version, or soon thereafter, will provide these capabilities which will unlock the full potential of granular intermediate integration testing.

APPENDIX A

Introduction to Service Orientation

This book is all about designing and developing service-oriented applications using WCF, yet there is considerable confusion and hype concerning what service orientation is and what it means. To make matters worse, most of the vendors in this space equate their definition of service orientation with their products and services. The vendors (Microsoft included) add to the confusion by equating service orientation with high-end Enterprise applications, where handling high scalability and throughput is a must (mostly because they all contend for that market, where the business margins are made).

This appendix presents my understanding of what service orientation is all about and attempts to put it in a concrete context. My take is different from that of the large vendors, but I believe it is more down-to-earth, rooted as it is in trends and the natural evolution of our industry. As you will see, I believe that service orientation is not a breakthrough or a quantum leap of thought, but rather the next gradual step (and probably not the last step) in a long journey that spans decades.

To understand where the software industry is heading with service orientation, you should first appreciate where it came from. After a brief discussion of the history of software engineering and its overarching trend, this appendix defines service-oriented applications (as opposed to mere architecture), explains what services themselves are, and examines the benefits of the methodology. It then presents the main principles of service orientation and augments the abstract tenets with a few more practical and concrete points to which most applications should adhere. Finally, the appendix concludes with a look to the future.

A Brief History of Software Engineering

The first modern computer was an electromechanical, typewriter-sized device developed in Poland in the late 1920s for enciphering messages. The device was later sold to the German Commerce Ministry, and in the 1930s the German military adopted it for enciphering all wireless communication. Today we know it as the Enigma.

Enigma used mechanical rotors to change the route of electrical current flow to a light board in response to a letter key being pressed, resulting in a different letter being output (the ciphered letter). Enigma was not a general-purpose computer: it could only do enciphering and deciphering (which today we call encryption and decryption when done automatically using a mathematical instead of a manual mechanical algorithm). If the operator wanted to change the encryption algorithm, he had to physically alter the mechanical structure of the machine by changing the rotors, their order, their initial positions, and the wired plugs that connected the keyboard to the light board. The “program” was therefore coupled in the extreme to the problem it was designed to solve (encryption), and to the mechanical design of the computer.

The late 1940s and the 1950s saw the introduction of the first general-purpose electronic computers for defense purposes. These machines could run code that addressed any problem, not just a single predetermined task. The downside was that the code executed on these computers was in a machine-specific “language” with the program coupled to the hardware itself. Code developed for one machine could not run on another. In fact, at the time there was no distinction between the software and the hardware (indeed, the word “software” was coined only in 1958). Initially this was not a cause for concern, since there were only a handful of computers in the world anyway. As machines became more prolific, this did turn into a problem. In the early 1960s, the emergence of assembly language decoupled the code from specific machines, enabling it to run on multiple computers. That code, however, was now coupled to the machine architecture: code written for an 8-bit machine could not run on a 16-bit machine, let alone withstand differences in the registers or available memory and memory layout. As a result, the cost of owning and maintaining a program began to escalate. This coincided more or less with the widespread adoption of computers in the civilian and government sectors, where the more limited resources and budgets necessitated a better solution.

In the 1960s, higher-level languages such as COBOL and FORTRAN introduced the notion of a *compiler*: the developer would write in an abstraction of machine programming (the language), and the compiler would translate that into actual assembly code. Compilers for the first time decoupled the code from the hardware and its architecture. The problem with those first-generation languages was that the code resulted in non-structured programming, where the code was internally coupled to

its own structure via the use of jump or go-to statements. Minute changes to the code structure often had devastating effects in multiple places in the program.

The 1970s saw the emergence of structured programming via languages such as C and Pascal, which decoupled the code from its internal layout and structure using functions and structures. The 1970s was also the first time developers and researchers started to examine software as an engineered entity. To drive down the cost of ownership, companies had to start thinking about reuse—that is, what would make a piece of code able to be reused in other contexts. With languages like C, the basic unit of reuse is the *function*. But the problem with function-based reuse is that the function is coupled to the data it manipulates, and if the data is global, a change to benefit one function in one reuse context is likely to damage another function used somewhere else.

Object Orientation

The solution to these problems that emerged in the 1980s, with languages such as Smalltalk and later C++, was object orientation. With object orientation, the functions and the data they manipulated were packaged together in an object. The functions (now called *methods*) encapsulated the logic, and the object encapsulated the data. Object orientation enabled domain modeling in the form of a class hierarchy. The mechanism of reuse was class-based, enabling both direct reuse and specialization via inheritance. But object orientation was not without its own acute problems. First, the generated application (or code artifact) was a single, monolithic application. Languages like C++ have nothing to say about the binary representation of the generated code. Developers had to deploy huge code bases every time they needed to make a change, however minute, and this had a detrimental effect on the development process and on application quality, time to market, and cost. While the basic unit of reuse was a class, it was a class in source format. Consequently, the application was coupled to the language used—you could not have a Smalltalk client consuming a C++ class or deriving from it. Language-based reuse implied uniformity of skill (all developers in the organization had to be skilled enough to use C++), which led to staffing problems. Language-based reuse also inhibited economy of scale, because if the organization was using multiple languages it necessitated duplication of investments in framework and common utilities. Finally, having to access the source files in order to reuse an object coupled developers to each other, complicated source control, and coupled teams together, since it made independent builds difficult. Moreover, inheritance turned out to be a poor mechanism for reuse, often harboring more harm than good because the developer of the derived class needed to be intimately aware of the implementation of the base class (which introduced vertical coupling across the class hierarchy).

Object orientation was oblivious to real-life challenges, such as deployment and versioning issues. Serialization and persistence posed yet another set of problems. Most

applications did not start by plucking objects out of thin air; they had some persistent state that needed to be hydrated into objects. However, there was no way of enforcing compatibility between the persisted state and the potentially new object code. Object orientation assumed the entire application was always in one big process. This prevented fault isolation between the client and the object, and if the object blew up, it took the client (and all other objects in the process) with it. Having a single process implies a single uniform identity for the clients and the objects, without any security isolation. This makes it impossible to authenticate and authorize clients, since they have the same identity as the object. A single process also impedes scalability, availability, responsiveness, throughput, and robustness. Developers could manually place objects in separate processes, yet if the objects were distributed across multiple processes or machines there was no way of using raw C++ for the invocations, since C++ required direct memory references and did not support distribution. Developers had to write host processes and use some remote call technology (such as TCP sockets) to remote the calls, but such invocations looked nothing like native C++ calls and did not benefit from object orientation.

Component Orientation

The solution for the problems of object orientation evolved over time, involving technologies such as the static library (*.lib*) and the dynamic library (*.dll*), culminating in 1994 with the first component-oriented technology, called COM (Component Object Model). Component orientation provided interchangeable, interoperable binary components. With this approach, instead of sharing source files, the client and the server agree on a binary type system (such as IDL) and a way of representing the metadata inside the opaque binary components. The components are discovered and loaded at runtime, enabling scenarios such as dropping a control on a form and having that control be automatically loaded at runtime on the client's machine. The client only programs against an abstraction of the service: a contract called the *interface*. As long as the interface is immutable, the service is free to evolve at will. A proxy can implement the same interface and thus enable seamless remote calls by encapsulating the low-level mechanics of the remote call. The availability of a common binary type system enables cross-language interoperability, so a Visual Basic client can consume a C++ COM component. The basic unit of reuse is the interface, not the component, and polymorphic implementations are interchangeable. Versioning is controlled by assigning a unique identifier for every interface, COM object, and type library.

While COM was a fundamental breakthrough in modern software engineering, most developers found it unpalatable. COM was unnecessarily ugly because it was bolted on top of an operating system that was unaware of it, and the languages used for writing COM components (such as C++ and Visual Basic) were at best object-oriented but not component-oriented. This greatly complicated the programming model, requiring frameworks such as ATL to partially bridge the two worlds. Recognizing

these issues, Microsoft released .NET 1.0 in 2002. .NET is (in the abstract) nothing more than cleaned-up COM, MFC, C++, and Windows, all working seamlessly together under a single new component-oriented runtime. .NET supports all the advantages of COM and mandates and standardizes many of its ingredients, such as type metadata sharing, dynamic component loading, serialization, and versioning.

While .NET is at least an order of magnitude easier to work with than COM, both COM and .NET suffer from a similar set of problems:

Technology and platform

The application and the code are coupled to the technology and the platform. Both COM and .NET predominantly target Windows. Both also expect the client and the service to be either COM or .NET and cannot interoperate natively with other technologies, be they Windows or not. While bridging technologies such as web services make interoperability possible, they force the developers to let go of almost all of the benefits of working with the native framework, and they introduce their own complexities and coupling with regard to the nature of the interoperability mechanism. This, in turn, breaks economy of scale.

Concurrency management

When a vendor ships a component, it cannot assume that its clients will not access it with multiple threads concurrently. In fact, the only safe assumption the vendor can make is that the component will be accessed by multiple threads. As a result, the components must be thread-safe and must be equipped with synchronization locks. However, if an application developer is building an application by aggregating multiple components from multiple vendors, the introduction of multiple locks renders the application deadlock-prone. Avoiding the deadlocks couples the application and the components.

Transactions

If multiple components are to participate in a single transaction, the application that hosts them must coordinate the transaction and flow the transaction from one component to the next, which is a serious programming endeavor. This also introduces coupling between the application and the components regarding the nature of the transaction coordination.

Communication protocols

If components are deployed across process or machine boundaries, they are coupled to the details of the remote calls, the transport protocol used, and its implications for the programming model (e.g., in terms of reliability and security).

Communication patterns

The components may be invoked synchronously or asynchronously, and they may be connected or disconnected. A component may or may not be able to be invoked in either one of these modes, and the application must be aware of its

exact preference. With COM and .NET, developing asynchronous or even queued solutions was still the responsibility of the developer, and any such custom solutions were not only difficult to implement but also introduced coupling between the solution and the components.

Versioning

Applications may be written against one version of a component and yet encounter another in production. Both COM and .NET bear the scars of DLL Hell (which occurs when the client at runtime is trying to use a different, incompatible version of the component than the one against which it was compiled), so both provide a guarantee to the client: that the client would get at runtime exactly the same component versions it was compiled against. This conservative approach stifled innovation and the introduction of new components. Both COM and .NET provided for custom version-resolution policies, but doing so risked DLL Hell-like symptoms. There was no built-in versioning tolerance, and dealing robustly with versioning issues coupled the application to the components it used.

Security

Components may need to authenticate and authorize their callers, but how does a component know which security authority it should use, or which user is a member of which role? Not only that, but a component may want to ensure that the communication from its clients is secure. That, of course, imposes certain restrictions on the clients and in turn couples them to the security needs of the component.

Off-the-shelf plumbing

In the abstract, interoperability, concurrency, transactions, protocols, versioning, and security are the glue—the plumbing—that holds any application together.

In a decent-sized application, the bulk of the development effort and debugging time is spent on addressing such plumbing issues, as opposed to focusing on business logic and features. To make things even worse, since the end customer (or the development manager) rarely cares about plumbing (as opposed to features), the developers typically are not given adequate time to develop robust plumbing. Instead, most hand-crafted plumbing solutions are proprietary (which hinders reuse, migration, and hiring) and are of low quality, because most developers are not security or synchronization experts and because they were not given the time and resources to develop the plumbing properly.

The solution was to use ready-made plumbing that offered such services to components. The first attempt at providing decent off-the-shelf plumbing was MTS (Microsoft Transactions Server), released in 1996. MTS offered support for much more than transactions, including security, hosting, activation, instance management, and syn-

chronization. MTS was followed by J2EE (1998), COM+ (2000), and .NET Enterprise Services (2002). All of these application platforms provided adequate, decent plumbing (albeit with varying degrees of ease of use), and applications that used them had a far better ratio of business logic to plumbing. However, by and large these technologies were not adopted on a large scale, due to what I term the *boundary problem*. Few systems are an island; most have to interact and interoperate with other systems. If the other system doesn't use the same plumbing, you cannot interoperate smoothly. For example, there is no way of propagating a COM+ transaction to a J2EE component. As a result, when crossing the system boundary, a component (say, component A) had to dumb down its interaction to the (not so large) common denominator between the two platforms. But what about component B, next to component A? As far as B was concerned, the component it interacted with (A) did not understand its variety of the plumbing, so B also had to be dumbed down. As a result, system boundaries tended to creep from the outside inward, preventing the ubiquitous use of off-the-shelf plumbing. Technologies like Enterprise Services and J2EE were useful, but they were useful in isolation.

Service Orientation

If you examine the brief history of software engineering just outlined, you'll notice a pattern: every new methodology and technology incorporates the benefits of its preceding technology and improves on the deficiencies of the preceding technology. However, every new generation also introduces new challenges. Therefore, I say that *modern software engineering is the ongoing refinement of the ever-increasing degrees of decoupling*.

Put differently, coupling is bad, but coupling is unavoidable. An absolutely decoupled application would be useless, because it would add no value. Developers can only add value by coupling things together. Indeed, the very act of writing code is coupling one thing to another. The real question is how to wisely choose what to be coupled to. I believe there are two types of coupling, good and bad. Good coupling is business-level coupling. Developers add value by implementing a system use case or a feature, by coupling software functionality together. Bad coupling is anything to do with writing plumbing. What was wrong with .NET and COM was not the concept; it was the fact that developers could not rely on off-the-shelf plumbing and still had to write so much of it themselves. The real solution is not just off-the-shelf plumbing, but rather *standard* off-the-shelf plumbing. If the plumbing is standard, the boundary problem goes away, and applications can utilize ready-made plumbing. However, all technologies (.NET, Java, etc.) use the client thread to jump into the object. How can you possibly take a .NET thread and give it to a Java object? The solution is to avoid call-stack invocation and instead to use message exchange. The technology vendors can standardize the format of the message and agree on ways to represent transactions, security credentials, and so on. When the message is received by the other side, the imple-

mentation of the plumbing there will convert the message to a native call (on a .NET or a Java thread) and proceed to call the object. Consequently, any attempt to standardize the plumbing has to be message-based.

And so, recognizing the problems of the past, in the late 2000s the service-oriented methodology emerged as the answer to the shortcomings of component orientation. In a service-oriented application, developers focus on writing business logic and expose that logic via interchangeable, interoperable service endpoints. Clients consume those endpoints (not the service code, or its packaging). The interaction between the clients and the service endpoint is based on a standard message exchange, and the service publishes some standard metadata describing what exactly it can do and how clients should invoke operations on it. The metadata is the service equivalent of the C++ header file, the COM type library, or the .NET assembly metadata, yet it contains not just operation metadata (such as methods and parameters) but also plumbing metadata. Incompatible clients—that is, clients that are incompatible with the plumbing expectations of the object—cannot call it, since the call will be denied by the platform. This is an extension of the object- and component-oriented compile-time notion that a client that is incompatible with an object's metadata cannot call it. Demanding compatibility with the plumbing (on top of the operations) is paramount. Otherwise, the object must always check on every call that the client meets its expectations in terms of security, transactions, reliability and so on, and thus the object invariably ends up infused with plumbing. Not only that, but the service's endpoint is reusable by any client compatible with its interaction constraints (such as synchronous, transacted, and secure communication), regardless of the client's implementation technology.

In many respects, a service is the natural evolution of the component, just as the component was the natural evolution of the object, which was the natural evolution of the function. Service orientation is, to the best of our knowledge as an industry, the correct way to build maintainable, robust, and secure applications.

The result of improving on the deficiencies of component orientation (i.e., classic .NET) is that when developing a service-oriented application, you decouple the service code from the technology and platform used by the client from many of the concurrency management issues, from transaction propagation and management, and from communication reliability, protocols, and patterns. By and large, securing the transfer of the message itself from the client to the service is also outside the scope of the service, and so is authenticating the caller. The service may still do its own local authorization, however, if the requirements so dictate. Similarly, as long as the endpoint supports the contract the client expects, the client does not care about the version of the service. There are also tolerances built into the standards to deal with versioning tolerance of the data passed between the client and the service.

Benefits of Service Orientation

Service orientation yields maintainable applications because the applications are decoupled on the correct aspects. As the plumbing evolves, the application remains unaffected. A service-oriented application is robust because the developers can use available, proven, and tested plumbing, and the developers are more productive because they get to spend more of the cycle time on the features rather than the plumbing. This is the true value proposition of service orientation: enabling developers to extract the plumbing from their code and invest more in the business logic and the required features.

The many other hailed benefits, such as cross-technology interoperability, are merely a manifestation of the core benefit. You can certainly interoperate without resorting to services, as was the practice until service orientation. The difference is that with ready-made plumbing you rely on the plumbing to provide the interoperability for you.

When you write a service, you usually do not care which platform the client executes on—that is immaterial, which is the whole point of seamless interoperability. However, a service-oriented application caters to much more than interoperability. It enables developers to cross boundaries. One type of boundary is the technology and platform, and crossing that boundary is what interoperability is all about. But other boundaries may exist between the client and the service, such as security and trust boundaries, geographical boundaries, organizational boundaries, timeline boundaries, transaction boundaries, and even business model boundaries. Seamlessly crossing each of these boundaries is possible because of the standard message-based interaction. For example, there are standards for how to secure messages and establish a secure interaction between the client and the service, even though both may reside in domains (or sites) that have no direct trust relationship. There is also a standard that enables the transaction manager on the client side to flow the transaction to the transaction manager on the service side, and have the service participate in that transaction, even though the two transaction managers never enlist in each other's transactions directly.

I believe that every application should be service-oriented, not just Enterprise applications that require interoperability and scalability. Writing plumbing in any type of application is wrong, constituting a waste of your time, effort, and budget, resulting in degradation of quality. Just as with .NET, every application was component-oriented (which was not so easy to do with COM alone) and with C++ every application was object-oriented (which was not so easy to do with C alone), when using WCF, every application should be service-oriented.

Service-Oriented Applications

A *service* is a unit of functionality exposed to the world over standard plumbing. A *service-oriented application* is simply the aggregation of services into a single logical, cohesive application (see [Figure A-1](#)), much as an object-oriented application is the aggregation of objects.

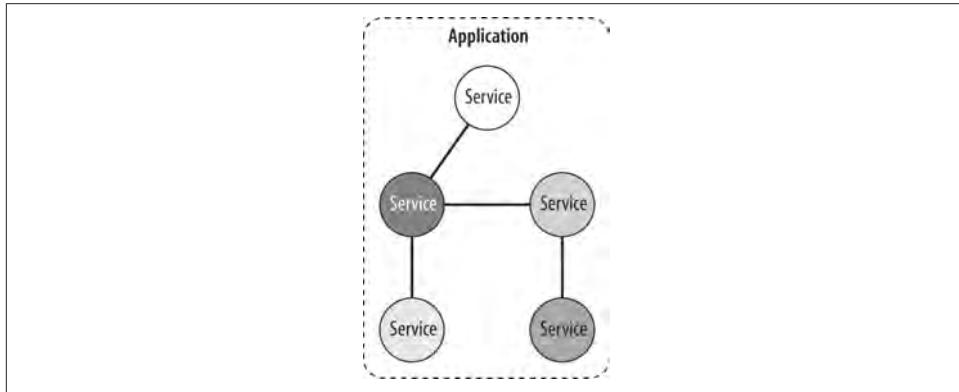


Figure A-1. A service-oriented application

The application itself may expose the aggregate as a new service, just as an object can be composed of smaller objects.

Inside services, developers still use concepts such as specific programming languages, versions, technologies and frameworks, operating systems, APIs, and so on. However, between services you have the standard messages and protocols, contracts, and metadata exchange.

The various services in an application can be all in the same location or be distributed across an intranet or the Internet, and they may come from multiple vendors and be developed across a range of platforms and technologies, versioned independently, and even execute on different timelines. All of those plumbing aspects are hidden from the clients in the application interacting with the services. The clients send the standard messages to the services, and the plumbing at both ends marshals away the differences between the clients and the services by converting the messages to and from the neutral wire representation.

Tenets and Principles

The service-oriented methodology governs what happens in the space between services (see [Figure A-1](#)). There is a small set of principles and best practices for building service-oriented applications, referred to as the *tenets of service-oriented architecture*:

Service boundaries are explicit

Any service is always confined behind boundaries, such as technology and location. The service should not make the nature of these boundaries known to its clients by exposing contracts and data types that betray such details. Adhering to this tenet will make aspects such as location and technology irrelevant. A different way of thinking about this tenet is that the more the client knows about the implementation of the service, the more the client is coupled to the service. To minimize the potential for coupling, the service has to explicitly expose functionality, and only operations (or data contracts) that are explicitly exposed will be shared with the client. Everything else is encapsulated. Service-oriented technologies should adopt an “opt-out by default” programming model, and expose only those things explicitly opted-in. This tenet is the modern incarnation of the old object-oriented adage that the application should maximize encapsulation and information hiding.

Services are autonomous

A service should need nothing from its clients or other services. The service should be operated and versioned independently from the clients, enabling it to evolve separately from them. The service should also be secured independently, so it can protect itself and the messages sent to it regardless of the degree to which the client uses security. Doing this (besides being common sense) further decouples the client and the service.

Services share operational contracts and data schema, not type-specific metadata

What the service decides to expose across its boundary should be type-neutral. The service must be able to convert its native data types to and from some neutral representation and should not share indigenous, technology-specific things such as its assembly version number or its type. In addition, the service should not let its client know about local implementation details such as its instance management mode or its concurrency management mode. The service should only expose logical operations. How the service goes about implementing those operations and how it behaves should not be disclosed to the client.

Services are compatible based on policy

The service should publish a policy indicating what it can do and how clients can interact with it. Any access constraints expressed in the policy (such as the need for reliable communication) should be separate from the service implementation details. Put differently, the service must be able to express, in a standard representation of policy, what it does and how clients should communicate with it. Being unable to express such a policy indicates poor service design. Note that a non-public service may not actually publish any such policy. This tenet simply implies that the service should be able to publish a policy if necessary.

Practical Principles

Well-designed applications should try to maximize adherence to the tenets just listed. However, those tenets are very abstract, and how they are supported is largely a product of the technology used to develop and consume the services, and of the design of the services. Consequently, just as not all code written in C++ is fully object-oriented, not all WCF applications may fully comply with the basic tenets just described. I therefore supplement those tenets with a set of more down-to-earth practical principles:

Services are secure

A service and its clients must use secure communication. At the very least, the transfer of messages from the clients to the service must be secured, and the clients must have a way of authenticating the service. The clients may also provide their credentials in the message so that the service can authenticate and authorize them.

Services leave the system in a consistent state

Conditions such as partially succeeding in executing the client's request are forbidden. All resources the service accesses must be in a consistent state after the client's call. If an error occurs the system state should not be only partially affected, and the service should not require the help of its clients to recover the system back to a consistent state after an error.

Services are thread-safe

The service must be designed so that it can sustain concurrent access from multiple clients. The service should also be able to handle causality and logical thread reentrancy.

Services are reliable

If the client calls a service, the client will always know in a deterministic manner whether the service received the message. In-order processing of messages is optional.

Services are robust

The service should isolate its faults, preventing them from taking it down (or taking down any other services). The service should not require clients to alter their behavior according to the type of error the service has encountered. This helps to decouple the clients from the service on the error-handling dimension.

Optional Principles

While I view the practical principles as mandatory, there is also a set of optional principles that may not be required by all applications (although adhering to them as well is usually a good idea):

Services are interoperable

The service should be designed so that any client, regardless of its technology, can call it.

Services are scale-invariant

It should be possible to use the same service code regardless of the number of clients and the load on the service. This will grossly simplify the cost of ownership of the service as the system grows and allow different deployment scenarios.

Services are available

The service should always be able to accept clients' requests and should have no downtime. Otherwise, if the service has periods of unavailability, the client needs to accommodate them, which in turn introduces coupling.

Services are responsive

The client should not have to wait long for the service to start processing its request. If the service is unresponsive, the client needs to plan for that, which in turn introduces coupling.

Services are disciplined

The service should not block the client for long. The service may perform lengthy processing, but only as long as it does not block the client. Otherwise, the client will need to accommodate that, which in turn introduces coupling.

What's Next?

Since service-oriented frameworks provide off-the-shelf plumbing for connecting services together, the more granular those services are, the more use the application can make of this infrastructure, and the less plumbing the developers have to write. Taken to the ultimate conclusion, every class and primitive should be a service, to maximize the use of the ready-made plumbing and to avoid handcrafting plumbing. This, in theory, will enable effortlessly transactional integers, secure strings, and reliable classes. But in practice, is that viable? Can .NET support it? Will future platforms offer this option?

I believe that as time goes by and service-oriented technologies evolve, the industry will see the service boundary pushed further and further inward, making services more and more granular, until the most primitive building blocks will be services. This would be in line with the historical trend of trading performance for productivity via methodology and abstraction. As an industry, we have always traded performance for productivity. .NET, where every class is treated as a binary component, is slower than COM, but the productivity benefit justifies this. COM itself is orders of magnitude slower than C++, yet developers opted for COM to address the problems of object orientation. C++ is likewise slower than C, but it did offer the crucial

abstractions of objects over functions. C in turn is a lot slower than raw assembly language, but the productivity gains it offered more than made up for that.

My benchmarks show that WCF can easily sustain hundreds of calls per second per class, making it adequate for the vast majority of business applications. While of course there is a performance hit for doing so, the productivity gains more than compensate, and historically, it is evident that this is a trade-off you should make. WCF does have detrimental overhead, but it's to do with ownership, not performance (which is adequate). Imagine a decent-sized application with a few hundred classes, each of which you wish to treat as a service. What would the `Main()` method of such an application look like, with hundreds of service host instances to be instantiated, opened, and closed? Such a `Main()` method would be unmaintainable. Similarly, would a config file with many hundreds of service and client endpoint declarations be workable?

The truth is that in practical terms, WCF cannot support (out of the box) such large-scale granular use. It is designed to be used between applications and across layers in the same application, not in every class. Just as COM had to use C++ and Windows, WCF is bolted on top of .NET. The language used (C# or Visual Basic) is merely component-oriented, not service-oriented, and the platform (.NET) is component-oriented, not service-oriented. What is required is a service-oriented platform, where the basic constructs are not classes but services. The syntax may still define a class, but it will be a service, just as every class in .NET is a binary component, very different from a C++ class. The service-oriented platform will support a config-less metadata repository, much like .NET generalized the type library and IDL concepts of COM. In this regard, WCF is merely a stopgap, a bridging technology between the world of components and the world of service (much like ATL once bridged the world of objects and C++ with the world of components, until .NET stepped in to provide native support for components at the class and primitive level).

A Service-Oriented Platform

If you take a wider view, every new idea in software engineering is implemented in three waves: first there is the methodology, then the technology, then the platform.

For example, object orientation as a methodology originated in the late '70s. The top C developers at the time did develop object-oriented applications, but this required manually passing state handles between functions and managing tables of function pointers for inheritance. Clearly, such practices required a level of conviction and skills that only very few had. With the advent of C++ in the early '80s came the technology, allowing every developer to write object-oriented applications. But C++ on its own was sterile, and required class libraries. Many developers wrote their own, which of course was not productive or scalable. The development of frameworks such as

MFC as an object-oriented platform, with types ranging from strings to windows, is what liberated C++ and enabled it to take off.

Similarly, take component orientation: in the first half of the '90s, developers who wanted to use COM had to write class factories and implement `IUnknown`, and concoct registry scripts and DLL entries. As a methodology, COM was just inaccessible. Then ATL came along, and this technology enabled developers to expose mere C++ classes as binary components. But the programming model was still too complex, since Windows knew nothing about COM, and the language was still object-oriented, lacking support for basic constructs such as interfaces. .NET as a component-oriented runtime provided the missing platform support for components at the class, primitive, language, and class library level.

Service orientation emerged as a methodology in the early 2000s, but at the time it was practically impossible to execute. With WCF, developers can expose mere classes as services, but the ownership overhead prevents widespread and granular use. I do not have a crystal ball, but I see no reason why the waves of methodology/technology/platform should stop now. Extrapolating from the last 30 to 40 years of software engineering, we are clearly missing a service-oriented platform. I believe the next generation of technologies from Microsoft will provide just that.

Every class as a service

Until we have a service-oriented platform, must we suffer the consequences of either unacceptable ownership overhead (granular use of WCF) or productivity and quality penalties (handcrafted custom plumbing)?

[Chapter 1](#) introduces my `InProcFactory` class, which lets you instantiate a service class over WCF:

```
public static class InProcFactory
{
    public static I CreateInstance<S,I>() where I : class
        where S : I;
    public static void CloseProxy<I>(I instance) where I : class;
    //More members
}
```

When using `InProcFactory`, you utilize WCF at the class level without ever resorting to explicitly managing the host or having client or service config files:

```
[ServiceContract]
interface IMyContract
{
    [OperationContract]
    string MyMethod();
}

class MyService : IMyContract
```

```
{...}

IMyContract proxy = InProcFactory.CreateInstance<MyService, IMyContract>();
proxy.MyMethod();
InProcFactory.CloseProxy(proxy);
```

This line:

```
IMyContract proxy = InProcFactory.CreateInstance<MyService, IMyContract>();
```

is syntactically equivalent to the C# way of instantiating a class type:

```
IMyContract proxy = new MyService();
```

The difference syntax-wise is that with C#, there is no need to specify the queried interfaces, since the compiler will examine the class, see if it supports the interface, and implicitly cast the class to the assigned interface variable. As it lacks compiler support for services, `InProcFactory` requires you to specify the required contract.



[Chapter 1](#) also shows my `WcfWrapper` helper class, which can eliminate even that difference, although at the price of defining a wrapper class.

However, the big difference between instantiating the class over WCF rather than C# is that when you do this all the powerful WCF features described in the rest of this book kick in: call timeout, encrypted calls, authentication, identity propagation, transaction propagation, transaction voting, instance management, error masking, channel faulting, fault isolation, buffering and throttling, data versioning tolerance, synchronization, synchronization context affinity, and more. With very little effort, you can also add tracing and logging, authorization, security audits, profiling and instrumentation, and durability, or intercept the calls and add many degrees of extensibility and customization.

`InProcFactory` lets you enjoy the benefits of WCF without suffering the ownership overhead. To me, `InProcFactory` is more than a useful utility—it is a glimpse of the future.

Headers and Contexts

In every method for every call on a .NET object, there are explicit arguments passed in as method parameters and implicit parameters available for the method body. Such implicit parameters include the thread the call executes on (available via `Thread.Current`); the call's app domain (available via `AppDomain.Current`); the call's transactions, if any (available via `Transaction.Current`); the call synchronization context (available via `SynchronizationContext.Current`); the security principal of the call (available via `Thread.CurrentPrincipal`); and even the little-known execution context of the call itself (available via `Thread.CurrentContext`). These *out-of-band* parameters provide in effect the logical execution context of the call, essential for the functioning of the object and .NET itself. But how can you pass additional contextual parameters to your WCF service, forming your own custom context? Such a need is surprisingly common and useful: for example, [Chapter 4](#) uses a custom context to manage the instance IDs of durable services, [Chapter 8](#) uses a custom context to provide the call priority, and [Chapter 9](#) uses a custom context to pass the address of a queued response service. This appendix presents and contrasts two distinct techniques for passing and managing custom contexts, using the message headers or the dedicated context binding. For both options, I will share dedicated helper classes designed to streamline and automate the interaction. You will also see some advanced WCF programming techniques.

Message Headers

Every WCF message contains a collection of outgoing and incoming message headers. When the client wishes to send out-of-band parameters to the service, it does so by adding those parameters to the outgoing headers. The service then reads those parameters from the incoming headers.

The operation context offers collections of incoming and outgoing headers, available via the `IncomingMessageHeaders` and `OutgoingMessageHeaders` properties:

```
public sealed class OperationContext : ...
{
    public MessageHeaders IncomingMessageHeaders
    {get;}

    public MessageHeaders OutgoingMessageHeaders
    {get;}

    //More members
}
```

Each collection is of the type `MessageHeaders` (that is, a collection of `MessageHeader` objects):

```
public sealed class MessageHeaders : ...
{
    public void Add(MessageHeader header);
    public T GetHeader<T>(int index);
    public T GetHeader<T>(string name, string ns);
    //More members
}
```

The class `MessageHeader` is not intended for application developers to interact with directly. Instead, use the `MessageHeader<T>` class, which provides for type-safe and easy conversion from a CLR type parameter to a message header:

```
public abstract class MessageHeader : ...
{...}

public class MessageHeader<T>
{
    public MessageHeader();
    public MessageHeader(T content);
    public T Content
    {get;set;}
    public MessageHeader GetUntypedHeader(string name, string ns);
    //More members
}
```

You can use any serializable or data contract type as the type parameter for `MessageHeader<T>`. You construct a `MessageHeader<T>` around a CLR type, and then use the `GetUntypedHeader()` method to convert it to a `MessageHeader` and store it in the outgoing headers. `GetUntypedHeader()` requires you to provide it with the generic type parameter name and namespace, which will be used later to look up the header from the headers collection. (Actually, using the name and namespace is just a suggestion; any unique value will do for this purpose. Since the type name and namespace combination tends to be unique, it is commonly used.) You perform the lookup

via the `GetHeader<T>()` method of `MessageHeaders`. Calling `GetHeader<T>()` obtains the value of the type parameter of the `MessageHeader<T>` used.

Client-Side Header Interaction

As mentioned previously, the client needs to add the parameter to the outgoing headers collection. However, what if the client is not a WCF service, so it does not have an operation context? As it turns out, this doesn't matter if the client is a service, since once the call enters a service, the operation context becomes immutable, so the client cannot write to its outgoing headers even if it has an operation context. The solution for all clients (services and non-services alike) is to create a new operation context and write to its outgoing headers. WCF enables a client to adopt a new operation context via the `OperationContextScope` class, defined as:

```
public sealed class OperationContextScope : IDisposable
{
    public OperationContextScope(IContextChannel channel);
    public OperationContextScope(OperationContext context);
    public void Dispose();
}
```

Using `OperationContextScope` is a general technique for spinning a new context when the one you have is inadequate. The constructor of `OperationContextScope` replaces the current thread's operation context with the new operation context. Calling `Dispose()` on the `OperationContextScope` instance restores the old context (even if it was `null`). If you do not call `Dispose()`, that may damage other objects on the same thread that expect the previous context. As a result, `OperationContextScope` is designed to be used inside a `using` statement and provide only a scope of code with a new operation context, even in the face of exceptions (hence its name):

```
using(OperationContextScope scope = new OperationContextScope(...))
{
    //Do work with new context
    ...
} //Restores previous context here
```

When constructing a new `OperationContextScope` instance, you provide its constructor with the inner channel of the proxy used for the call (and thus affect the message). **Example B-1** shows the steps required to send an integer to a service in the message headers.

Example B-1. Passing integer in headers by the client

```
[ServiceContract]
interface IMyContract
{
    [OperationContract]
```

```

        void MyMethod();
    }
    class MyContractClient : ClientBase<IMyContract>, IMyContract
    {...}

    //Client code:
    MessageHeader<int> numberHeader = new MessageHeader<int>(123);

    MyContractClient proxy = new MyContractClient();
    using(OperationContractScope contextScope =
            new OperationContextScope(proxy.InnerChannel)){
        OperationContext.Current.OutgoingMessageHeaders.Add(
            numberHeader.GetUntypedHeader("Int32", "System"));

        proxy.MyMethod();
    }
    proxy.Close();
}

```

The client first constructs an instance of `MessageHeader<int>`, initializing it with the value 123. The client then uses the `GetUntypedHeader()` method to convert the type-safe integer header to a non-type-safe representation, using the integer name and namespace as keys, and add that to the outgoing headers inside a new operation context scope. The call to the service is also inside the scope. After exiting the operation context scope, the client closes the proxy in the original operation context scope (if any).

Service-Side Header Interaction

Example B-2 shows the matching service code required to read the integer from the incoming message headers. Note that the service must know in advance the keys required to look up the number from the headers.

Example B-2. Reading integer from headers by the service

```

class MyService : IMyContract
{
    public void MyMethod()
    {
        int number = OperationContext.Current.IncomingMessageHeaders.
            GetHeader<int>("Int32", "System");
        Debug.Assert(number == 123);
    }
}

```

Logically, the service treats the out-of-band integer passed to it as a number context. Any party down the call chain from the service can also read the number context from the operation context.

Encapsulating the Headers

Both the client and the service will benefit greatly from encapsulating the interaction with the message headers by defining a `NumberContext` helper class, as shown in [Example B-3](#).

Example B-3. The NumberContext helper class

```
class NumberContext
{
    public static int Current
    {
        get
        {
            OperationContext context = OperationContext.Current;
            if(context == null)
            {
                return 0;
            }
            return context.IncomingMessageHeaders.GetHeader<int>("Int32", "System");
        }
        set
        {
            OperationContext context = OperationContext.Current;
            MessageHeader<int> numberHeader = new MessageHeader<int>(value);
            context.OutgoingMessageHeaders.Add(
                numberHeader.GetUntypedHeader("Int32", "System"));
        }
    }
}
```

Using `NumberContext` mimics the use of any built-in .NET context, since it offers the `Current` static property, which gets and sets the appropriate headers collection. Using `NumberContext`, [Example B-1](#) is reduced to the code shown in [Example B-4](#).

Example B-4. Using NumberContext by the client

```
MyContractClient proxy = new MyContractClient();
using(OperationContextScope contextScope =
    new OperationContextScope(proxy.InnerChannel))
{
    NumberContext.Current = 123;

    proxy.MyMethod();
}
proxy.Close();
```

Likewise, [Example B-2](#) is reduced to [Example B-5](#).

Example B-5. Using NumberContext by the service

```
class MyService : IMyContract
{
    public void MyMethod()
    {
        int number = NumberContext.Current;
        Debug.Assert(number == 123);
    }
}
```

The GenericContext<T> helper class

While Examples B-4 and B-5 are both a marked improvement over direct interaction with the headers, such use of the headers is still problematic, since it involves defining helper classes with repeated, complicated code for every use of message headers as a logical context. The solution is to generalize the technique shown in Example B-3 using a generic type parameter. This is exactly what my GenericContext<T> class, shown in Example B-6, does. GenericContext<T> and the rest of the helper classes in this chapter are available with *ServiceModelEx*.

Example B-6. The GenericContext<T> class

```
[DataContract]
public class GenericContext<T>
{
    [DataMember]
    public readonly T Value;

    internal static string TypeName;
    internal static string TypeNamespace;

    static GenericContext()
    {
        //Verify [DataContract] or [Serializable] on T
        Debug.Assert(IsDataContract(typeof(T)) || typeof(T).IsSerializable);

        TypeNamespace = "net.clr:" + typeof(T).FullName;
        TypeName = "GenericContext";
    }
    static bool IsDataContract(Type type)
    {
        object[] attributes =
            type.GetCustomAttributes(typeof(DataContractAttribute), false);
        return attributes.Length == 1;
    }

    public GenericContext(T value)
    {
        Value = value;
    }
}
```

```

}

public GenericContext() : this(default(T))
{}
public static GenericContext<T> Current
{
    get
    {
        OperationContext context = OperationContext.Current;
        if(context == null)
        {
            return null;
        }
        try
        {
            return context.IncomingMessageHeaders.
                GetHeader<GenericContext<T>>(TypeName,TypeNamespace);
        }
        catch
        {
            return null;
        }
    }
    set
    {
        OperationContext context = OperationContext.Current;
        Debug.Assert(context != null);

        //Having multiple GenericContext<T> headers is an error
        bool headerExists = false;
        try
        {
            context.OutgoingMessageHeaders.
                GetHeader<GenericContext<T>>(TypeName,TypeNamespace);
            headerExists = true;
        }
        catch(MessageHeaderException exception)
        {
            Debug.Assert(exception.Message == "There is not a header with name " +
                TypeName + " and namespace " +
                TypeNamespace + " in the message.");
        }
        if(headerExists)
        {
            throw new InvalidOperationException("A header with name " + TypeName +
                " and namespace " + TypeNamespace +
                " already exists in the message.");
        }
        MessageHeader<GenericContext<T>> genericHeader =
            new MessageHeader<GenericContext<T>>(value);
        context.OutgoingMessageHeaders.Add(
            genericHeader.GetUntypedHeader(TypeName,TypeNamespace));
    }
}

```

```
        }
    }
}
```

`GenericContext<T>` lets you treat any serializable or data contract type parameter as a logical context, and its static constructor validates that. The type parameter used is a generic yet type-safe and application-specific custom context. `GenericContext<T>` uses “`GenericContext`” for the type name and the full name of `T` for the namespace to reduce the chance of a conflict. `GenericContext<T>` also validates that the outgoing headers do not already contain such a type parameter. Both the client and the service can use `GenericContext<T>` as-is. All a client has to do to pass some custom context to the service is set the static `Current` property inside a new `OperationContextScope`:

```
GenericContext<int>.Current = new GenericContext<int>(123);
```

On the service side, to read the value out of the headers, any downstream party can write:

```
int number = GenericContext<int>.Current.Value;
```

Alternatively, you can wrap `GenericContext<T>` with a dedicated context. Using `GenericContext<T>`, the `NumberContext` of [Example B-3](#) is reduced to the code shown in [Example B-7](#).

Example B-7. NumberContext using GenericContext<T>

```
class NumberContext
{
    public static int Current
    {
        get
        {
            return GenericContext<int>.Current.Value;
        }
        set
        {
            GenericContext<int>.Current = new GenericContext<int>(value);
        }
    }
}
```

Streamlining the Client

Even when using `GenericContext<T>`, the client code (as in [Example B-4](#)) is far too raw and exposed, requiring every invocation of the proxy to use an operation context scope. It is better to encapsulate these steps in the proxy itself. The constructors of the proxy should all take additional parameters for the value to pass in the headers. Inside every method, the proxy will create a new operation context and add the value

to the outgoing headers collection. This will avoid on every invocation polluting the client code with the interaction with the logical context and the operation context.

Using the same contract definition as in [Example B-1](#), [Example B-8](#) shows such a proxy used to pass an out-of-band number in the message headers.

Example B-8. Encapsulating the headers and the operation context scope

```
class MyContractClient : ClientBase<IMyContract>, IMyContract
{
    readonly int Number;

    public MyContractClient(int number)
    {
        Number = number;
    }
    public MyContractClient(int number, string endpointName) : base(endpointName)
    {
        Number = number;
    }

    //More constructors

    public void MyMethod()
    {
        using(OperationContextScope contextScope =
              new OperationContextScope(InnerChannel))
        {
            NumberContext.Current = Number;

            Channel.MyMethod();
        }
    }
}
```

All the constructors of the proxy in [Example B-8](#) accept the number to pass to the service and save it in a read-only variable. The proxy uses the `NumberContext` class of [Example B-7](#) to encapsulate the interaction with the headers.

Using the proxy from [Example B-8](#), the client code from [Example B-4](#) is reduced to:

```
MyContractClient proxy = new MyContractClient(123);
proxy.MyMethod();
proxy.Close();
```

The `HeaderClientBase<T,H>` proxy class

The problem with the technique demonstrated in [Example B-8](#) is that you would have to repeat such code in every method in the proxy, and for every other proxy that wishes to pass out-of-band parameters. It is therefore preferable to encapsulate these

steps further in a dedicated proxy class, and even avoid the interaction with the operation context altogether using message interception. To that end, I wrote `HeaderClientBase<T,H>`, defined in [Example B-9](#).

Example B-9. The HeaderClientBase<T,H> proxy base class

```
public abstract partial class HeaderClientBase<T,H> : InterceptorClientBase<T>
    where T : class
{
    public H Header
    {get;protected set;}

    public HeaderClientBase() : this(default(H))
    {}
    public HeaderClientBase(string endpointName) : this(default(H),endpointName)
    {}

    public HeaderClientBase(H header)
    {
        Header = header;
    }
    public HeaderClientBase(H header,string endpointName) : base(endpointName)
    {
        Header = header;
    }

    //More constructors

    protected override void PreInvoke(ref Message request)
    {
        GenericContext<H> context = new GenericContext<H>(Header);
        MessageHeader<GenericContext<H>> genericHeader =
            new MessageHeader<GenericContext<H>>(context);
        request.Headers.Add(genericHeader.GetUntypedHeader(
            GenericContext<H>.TypeName,GenericContext<H>.TypeNamespace));
    }
}
```

The type parameter `H` can be any serializable or data contract type. In order for you to use it with or without header information, `HeaderClientBase<T,H>` offers two sets of constructors—one set that accepts a header and one set that does not. The constructors that accept the header store it in the protected `Header` property. `HeaderClientBase<T,H>` derives from the `InterceptorClientBase<T>` class defined in [Appendix E](#) as part of a generic interception framework. `InterceptorClientBase<T>` provides the `PreInvoke()` virtual method where its subclasses can hook the outgoing message and interact with it. The overridden version of `PreInvoke()` creates a new instance of `GenericContext<H>`, and manually adds it to the request message headers. Note that the value of the header is read every time by accessing the `Header` property.

Context Bindings

WCF provides three additional bindings dedicated to managing custom contexts. These bindings are the `BasicHttpContextBinding`, the `NetTcpContextBinding`, and the `WSHttpContextBinding`. The context bindings all derive from their respective regular bindings:

```
public class BasicHttpContextBinding : BasicHttpBinding
{
    /* Same constructors as BasicHttpBinding */

    public ProtectionLevel ContextProtectionLevel
    {get;set;}
}

public class NetTcpContextBinding : NetTcpBinding
{
    /* Same constructors as NetTcpBinding */

    public ProtectionLevel ContextProtectionLevel
    {get;set;}
}

public class WSHttpContextBinding : WSHttpBinding
{
    /* Same constructors as WSHttpBinding */

    public ProtectionLevel ContextProtectionLevel
    {get;set;}
}
```

In the case of the `NetTcpContextBinding` and the `WSHttpContextBinding`, the `ContextProtectionLevel` indicates how to protect the context while in transfer, as discussed in [Chapter 10](#).

The context bindings are used exactly the same way as their base bindings, yet they add support for a dedicated context management protocol. These bindings can be used with or without a context. The context protocol lets you pass as a custom context a collection of strings in the form of pairs of keys and values, stored implicitly in the message headers. There are several important differences between using a context binding and using the direct message headers for passing out-of-band parameters to a custom context:

- With a context binding, you can only set the information to pass to the service once, before opening the proxy (or using it for the first time). After that, the custom context is cached, and any attempt to modify it results in an error. With the message headers, every call to the services on the same proxy may contain different headers.
- With the context binding, you can only pass as parameters simple strings in the form of a keys/values dictionary. This is a liability when trying to pass composite

types that go beyond simple values. With message headers, any serializable or data contract type will do.

- The use of strings means there is inherently no type safety with the context parameters. While this is also true with message headers, my `GenericContext<T>` does restore the missing type safety.
- Out of the box, only a limited set of bindings support the context protocol. Glaringly missing are the IPC and MSMQ bindings. The message headers technique works over any binding.

Client-Side Context Binding Interaction

The client sets the context to send to the service using the `IContextManager` interface:

```
public interface IContextManager
{
    IDictionary<string, string> GetContext();
    void SetContext(IDictionary<string, string> context);

    bool Enabled
    {get; set;}
}
```

The client obtains the reference to the `IContextManager` interface by accessing the proxy's inner channel properties:

```
public abstract class ClientBase<T> : ICommunicationObject where T : class
{
    public IClientChannel InnerChannel
    {get;}
    //More members
}
public interface IClientChannel : IContextChannel, ...
{...}
public interface IContextChannel : IChannel, ...
{...}
public interface IChannel : ICommunicationObject
{
    T GetProperty<T>() where T : class;
}
```

The `InnerChannel` property supports the `IChannel` interface, which offers the `GetProperty<T>()` method:

```
MyContractClient proxy = new MyContractClient();
IContextManager contextManager = proxy.InnerChannel.GetProperty
    < IContextManager>();
```

Once the client obtains `IContextManager`, it can copy the current context by calling the `GetContext()` method. The context is merely a dictionary of strings as keys and

values. Since the dictionary returned from `GetContext()` is a copy of the actual context, the client cannot use it to change the context. Instead, the client needs to call the `SetContext()` method, providing the new context. The client can override the old context or just add values to the old context and then set it back in, as shown in [Example B-10](#).

Example B-10. Setting the context on the proxy

```
MyContractClient proxy = new MyContractClient();
IContextManager contextManager = proxy.InnerChannel.GetProperty<IContextManager>();

//Just add in, not overwriting dictionary
IDictionary<string,string> context = contextManager.GetContext();
context["NumberContext"] = "123";
contextManager.SetContext(context);

proxy.MyMethod();

proxy.Close();
```

Service-Side Context Binding Interaction

The service reads the context values from the incoming message properties, accessed via the operation context:

```
public sealed class OperationContext : ...
{
    public MessageProperties IncomingMessageProperties
    {
        get;
    }
    //More members
}
```

`MessageProperties` is a non-type-safe dictionary that accepts a string key and returns the matching object value:

```
public sealed class MessageProperties : IDictionary<string,object>
{...}
```

To obtain the context property, the service uses the static string `ContextMessageProperty.Name`. This returns an object of the type `ContextMessageProperty`, defined as:

```
[Serializable]
public class ContextMessageProperty : IMessageProperty
{
    public IDictionary<string,string> Context
    {get;}
    public static string Name
    {get;}
```

```
        //More members  
    }
```

The Context property of ContextMessageProperty is the same dictionary of parameters passed by the client. [Example B-11](#) shows the required service-side steps to read the number context passed in [Example B-10](#).

Example B-11. Reading the context by the service

```
class MyService : IMyContract  
{  
    public void MyMethod()  
    {  
        ContextMessageProperty contextProperty = OperationContext.Current.  
            IncomingMessageProperties[ContextMessageProperty.Name]  
            as ContextMessageProperty;  
  
        Debug.Assert(contextProperty.Context.ContainsKey("NumberContext"));  
  
        string number = contextProperty.Context["NumberContext"];  
  
        Debug.Assert(number == "123");  
    }  
}
```

Streamlining the Client

You can streamline the steps required of the client to read or write to the context using my ContextManager static helper class, shown in [Example B-12](#).

Example B-12. Client-side methods of ContextManager

```
public static class ContextManager  
{  
    public static void SetContext(IClientChannel innerChannel,  
                                string key, string value)  
    {  
        SetContext(innerChannel, CreateContext(key, value));  
    }  
  
    public static void SetContext(IClientChannel innerChannel,  
                                IDictionary<string, string> context)  
    {  
        IContextManager contextManager = innerChannel.GetProperty<IContextManager>();  
        contextManager.SetContext(context);  
    }  
  
    public static IDictionary<string, string> CreateContext(string key, string value)  
    {  
        IDictionary<string, string> context = new Dictionary<string, string>();  
        context[key] = value;  
        return context;  
    }  
}
```

```

        context[key] = value;
        return context;
    }

    public static IDictionary<string,string> UpdateContext(
                                                IClientChannel innerChannel,
                                                string key,string value)
    {
        IContextManager contextManager = innerChannel.GetProperty<IContextManager>();

        IDictionary<string,string> context =
            new Dictionary<string,string>(contextManager.GetContext());
        context[key] = value;
        return context;
    }

    //Proxy extensions
    public static void SetContext<T>(this ClientBase<T> proxy,
                                    string key,string value) where T : class
    {
        SetContext(proxy.InnerChannel,key,value);
    }

    public static void SetContext<T>(this ClientBase<T> proxy,
                                    IDictionary<string,string> context) where T : class
    {
        SetContext(proxy.InnerChannel,context);
    }
    public static IDictionary<string,string> UpdateContext<T>(
                                                this ClientBase<T> proxy,string key,string value) where T : class
    {
        return UpdateContext(proxy.InnerChannel,key,value);
    }
}

```

ContextManager offers overloaded versions of the SetContext() method that allow the client to set a new context on a proxy's inner channel, using a single key/value pair or a collection of such pairs in a dictionary. These methods are useful both with a proxy class and with a channel factory. ContextManager also exposes setting the context as an extension method on the proxy class. You can use the CreateContext() method to create a new dictionary or the UpdateContext() method to add a key/value pair to an existing context. Using ContextManager, [Example B-10](#) is reduced to:

```

MyContractClient proxy = new MyContractClient();
proxy.SetContext("NumberContext","123");
proxy.MyMethod();
proxy.Close();

```

However, relying on `SetContext()` this way requires you to explicitly use it upon every instantiation of the proxy. It is better to encapsulate `ContextManager` in a dedicated proxy class, such as my `ContextClientBase<T>`:

```
public abstract class ContextClientBase<T> : ClientBase<T> where T : class
{
    public ContextClientBase();
    public ContextClientBase(string endpointName);
    public ContextClientBase(string key, string value);
    public ContextClientBase(IDictionary<string, string> context);
    public ContextClientBase(string key, string value, string endpointName);
    public ContextClientBase(IDictionary<string, string> context,
                           string endpointName);
    //More constructors
}
```

The constructors of `ContextClientBase<T>` accept the usual proxy parameters, such as the endpoint name or binding and address, as well as the contextual parameters to send the service (either a single key/value pair, or a collection of keys and values using a dictionary). Your proxy can derive directly from `ContextClientBase<T>`:

```
class MyContractClient : ContextClientBase<IMyContract>, IMyContract
{
    public MyContractClient(string key, string value) : base(key, value)
    {}
    /* More constructors */
    public void MyMethod()
    {
        Channel.MyMethod();
    }
}
```

Using `ContextClientBase<T>`, [Example B-10](#) is reduced to:

```
MyContractClient proxy = new MyContractClient("NumberContext", "123");
proxy.MyMethod();
proxy.Close();
```

[Example B-13](#) shows the implementation of `ContextClientBase<T>`.

Example B-13. Implementing ContextClientBase<T>

```
public abstract class ContextClientBase<T> : ClientBase<T> where T : class
{
    public ContextClientBase(string key, string value, string endpointName)
        : this(ContextManager.CreateContext(key, value), endpointName)
    {}
    public ContextClientBase(IDictionary<string, string> context, string endpointName)
        : base(endpointName)
    {
        SetContext(context);
    }
}
```

```

/* More constructors */

void SetContext(IDictionary<string,string> context)
{
    VerifyContextBinding();
    ContextManager.SetContext(InnerChannel,context);
}
void VerifyContextBinding()
{
    BindingElementCollection elements = Endpoint.Binding.CreateBindingElements();

    if(elements.Contains(typeof(ContextBindingElement)))
    {
        return;
    }

    throw new InvalidOperationException("Can only use context binding");
}
}

```

A few of the constructors of `ContextClientBase<T>` use `ContextManager` to create a new context and pass it to another constructor, which calls the `SetContext()` helper method. `SetContext()` first verifies that the binding used is indeed a context binding and then uses `ContextManager` to set the context. Verifying that the binding indeed supports the context protocol is done by searching for the `ContextBindingElement` in the collection of binding elements. This way of verifying is better than looking at the binding type, since it also works automatically with a custom context binding.

Streamlining the Service

For the service, the `ContextManager` helper class encapsulates the interaction with operation context and message properties. `ContextManager` provides the `GetContext()` method:

```

public static class ContextManager
{
    public static string GetContext(string key);

    //More members
}

```

Using `GetContext()`, the service code in [Example B-11](#) is reduced to:

```

class MyService : IMyContract
{
    public void MyMethod()
    {
        string number = ContextManager.GetContext("NumberContext");
    }
}

```

```

        Debug.Assert(number == "123");
    }
}

```

Example B-14 shows the implementation of `GetContext()`.

Example B-14. Implementing GetContext()

```

public static class ContextManager
{
    public static string GetContext(string key)
    {
        if(OperationContext.Current == null)
        {
            return null;
        }
        if(OperationContext.Current.IncomingMessageProperties.
            ContainsKey(ContextMessageProperty.Name))
        {
            ContextMessageProperty contextProperty =
                OperationContext.Current.IncomingMessageProperties[ContextMessageProperty.Name]
                    as ContextMessageProperty;
            if(contextProperty.Context.ContainsKey(key) == false)
            {
                return null;
            }
            return contextProperty.Context[key];
        }
        else
        {
            return null;
        }
    }
}

```

`GetContext()` is similar to the explicit steps taken in [Example B-11](#), except it adds state and error management. If the context does not contain the request key (or if no context was found), `GetContext()` returns `null`.

Creating a Custom Context Binding

WCF provides context support for the basic, WS, and TCP bindings. Missing from that list is the IPC binding. It would be valuable to have that support for the IPC binding for custom context support on the same machine. Creating such a custom binding is a worthy exercise, and it serves as a good demonstration of how to write a custom binding.

`ServiceModelEx` contains the `NetNamedPipeContextBinding` class, defined as:

```

public class NetNamedPipeContextBinding : NetNamedPipeBinding
{
    /* Same constructors as NetNamedPipeBinding */

    public ProtectionLevel ContextProtectionLevel
    {get;set;}
}

```

`NetNamedPipeContextBinding` is used exactly like its base class, and you can use it with or without a context. Both the client and the host can use this binding programmatically as-is, by instantiating it like any other built-in binding. However, when using a custom binding in conjunction with a config file, you need to inform WCF where the custom binding is defined.

To that end, `ServiceModelEx` also defines the `NetNamedPipeContextBindingElement` and `NetNamedPipeContextBindingCollectionElement` helper classes:

```

public class NetNamedPipeContextBindingElement : NetNamedPipeBindingElement
{
    public NetNamedPipeContextBindingElement();
    public NetNamedPipeContextBindingElement(string name);
    public ProtectionLevel ContextProtectionLevel
    {get;set;}
}
public class NetNamedPipeContextBindingCollectionElement :
    StandardBindingCollectionElement<NetNamedPipeContextBinding,
                                         NetNamedPipeContextBindingElement>
{
}

```

You need to add the type of `NetNamedPipeContextBindingCollectionElement` and its assembly to the list of binding extensions, naming `NetNamedPipeContextBinding` as a custom binding. You can do this on a per-application basis by adding it to the application config file.

Example B-15 shows such an application-specific config file for the host side, but you have to enter the same directives in the client's config file as well.

Example B-15. Adding per-application administrative custom binding support

```

<system.serviceModel>
    <extensions>
        <bindingExtensions>
            <add name = "netNamedPipeContextBinding"
                 type = "ServiceModelEx.NetNamedPipeContextBindingCollectionElement,
                                         ServiceModelEx"
            />
        </bindingExtensions>
    </extensions>

    <services>
        <service name = "...">

```

```

<endpoint
    address  = "net.pipe://..."
    binding  = "netNamedPipeContextBinding"
    contract = "..."
  />
</service>
</services>
</system.serviceModel>

```

Alternatively, you can add `NetNamedPipeContextBindingCollectionElement` to `machine.config` to affect every application on the machine. In that case, there is no need to list the binding extensions in the client or service config file. [Example B-16](#) shows such a configuration.

Example B-16. Adding machine-wide administrative custom binding support

```

<!--In machine.config-->
<bindingExtensions>
  <add name = "wsHttpContextBinding" type = "..."/>
  <add name = "netTcpContextBinding" type = "..."/>
  <add name = "netNamedPipeContextBinding"
    type = "ServiceModelEx.NetNamedPipeContextBindingCollectionElement,
              ServiceModelEx"/>
  <!--Additional bindings-->
</bindingExtensions>

<!--In app.config-->
<system.serviceModel>
  <services>
    <service name = "...">
      <endpoint
        address  = "net.pipe://..."
        binding  = "netNamedPipeContextBinding"
        contract = "...""
      />
    </service>
  </services>
</system.serviceModel>

```

Of course, you can configure a binding section to customize any property of `NetNamedPipeContextBinding`, whether it comes from `NetNamedPipeBinding` or from `NetNamedPipeContextBinding`:

```

<bindings>
  <netNamedPipeContextBinding>
    <binding name = "TransactionalSecureContextIPC"
      contextProtectionLevel = "EncryptAndSign"
      transactionFlow = "True"
    />

```

```
</netNamedPipeContextBinding>  
</bindings>
```

Implementing NetNamedPipeContextBinding

Example B-17 lists the implementation of NetNamedPipeContextBinding and its supporting classes.

Example B-17. Implementing NetNamedPipeContextBinding

```
public class NetNamedPipeContextBinding : NetNamedPipeBinding  
{  
    internal const string SectionName = "netNamedPipeContextBinding";  
  
    public ProtectionLevel ContextProtectionLevel  
    {get;set;}  
  
    public NetNamedPipeContextBinding()  
    {  
        ContextProtectionLevel = ProtectionLevel.EncryptAndSign;  
    }  
    public NetNamedPipeContextBinding(NetNamedPipeSecurityMode securityMode) :  
        base(securityMode)  
    {  
        ContextProtectionLevel = ProtectionLevel.EncryptAndSign;  
    }  
    public NetNamedPipeContextBinding(string configurationName)  
    {  
        ContextProtectionLevel = ProtectionLevel.EncryptAndSign;  
        ApplyConfiguration(configurationName);  
    }  
    public override BindingElementCollection CreateBindingElements()  
    {  
        BindingElement element = new ContextBindingElement(ContextProtectionLevel,  
            ContextExchangeMechanism.ContextSoapHeader);  
  
        BindingElementCollection elements = base.CreateBindingElements();  
        elements.Insert(0,element);  
  
        return elements;  
    }  
  
    void ApplyConfiguration(string configurationName)  
    {  
        Configuration config =  
            ConfigurationManager.OpenExeConfiguration(ConfigurationUserLevel.None);  
        ServiceModelSectionGroup sectionGroup =  
            ServiceModelSectionGroup.GetSectionGroup(config);  
  
        BindingsSection bindings = sectionGroup.Bindings;
```

```

        NetNamedPipeContextBindingCollectionElement section =
            (NetNamedPipeContextBindingCollectionElement)bindings[SectionName];

        NetNamedPipeContextBindingElement element =
            section.Bindings[configurationName];
        if(element == null)
        {
            throw new ConfigurationErrorsException();
        }
        else
        {
            element.ApplyConfiguration(this);
        }
    }
}

public class NetNamedPipeContextBindingElement : NetNamedPipeBindingElement
{
    const string ContextProtectionLevelName = "contextProtectionLevel";

    public NetNamedPipeContextBindingElement()
    {
        Initialize();
    }

    public NetNamedPipeContextBindingElement(string name) : base(name)
    {
        Initialize();
    }

    void Initialize()
    {
        ConfigurationProperty property =
            new ConfigurationProperty(ContextProtectionLevelName,
                typeof(ProtectionLevel),
                ProtectionLevel.EncryptAndSign);

        Properties.Add(property);

        ContextProtectionLevel = ProtectionLevel.EncryptAndSign;
    }

    protected override void OnApplyConfiguration(Binding binding)
    {
        base.OnApplyConfiguration(binding);

        NetNamedPipeContextBinding netNamedPipeContextBinding =
            binding as NetNamedPipeContextBinding;
        Debug.Assert(netNamedPipeContextBinding != null);

        netNamedPipeContextBinding.ContextProtectionLevel = ContextProtectionLevel;
    }

    protected override Type BindingElementType
    {
        get
        {

```

```

        return typeof(NetNamedPipeContextBinding);
    }
}

public ProtectionLevel ContextProtectionLevel
{
    get
    {
        return (ProtectionLevel)base[ContextProtectionLevelName];
    }
    set
    {
        base[ContextProtectionLevelName] = value;
    }
}

public class NetNamedPipeContextBindingCollectionElement :
    StandardBindingCollectionElement
    <NetNamedPipeContextBinding,NetNamedPipeContextBindingElement>
{}

```

The constructors of `NetNamedPipeContextBinding` all delegate the actual construction to the base constructors of `NetNamedPipeBinding`, and the only initialization they do is setting the context protection level to default to `ProtectionLevel.EncryptAndSign`.

The heart of any binding class is the `CreateBindingElements()` method. `NetNamedPipeContextBinding` accesses its base binding collection of binding elements and adds to it the `ContextBindingElement`. Inserting this element into the collection adds support for the context protocol. The rest of [Example B-17](#) is mere bookkeeping to enable administrative configuration. The `ApplyConfiguration()` method is called by the constructor, which takes the binding section configuration name. `ApplyConfiguration()` uses the `ConfigurationManager` class (discussed in [Chapter 9](#)) to parse out of the config file the `netNamedPipeContextBinding` section, and from it an instance of `NetNamedPipeContextBindingElement`. That binding element is then used to configure the binding instance by calling its `ApplyConfiguration()` method. The constructors of `NetNamedPipeContextBindingElement` add to its base class `Properties` collection of configuration properties a single property for the context protection level. In `OnApplyConfiguration()` (which is called as a result of calling `ApplyConfiguration()` on `NetNamedPipeBindingElement` by `NetNamedPipeContextBinding.ApplyConfiguration()`), the method first configures its base element and then sets the context protection level according to the configured level.

The `NetNamedPipeContextBindingCollectionElement` type is used to bind `NetNamedPipeContextBinding` with the `NetNamedPipeContextBindingElement`. This way, when adding `NetNamedPipeContextBindingCollectionElement` as a binding exten-

sion, the configuration manager knows which type to instantiate and provide with the binding parameters.



Since you can use `NetNamedPipeContextBinding` with or without a context, the `InProcFactory` class presented in [Chapter 1](#) actually uses the `NetNamedPipeContextBinding` to enable transparent support for custom contexts if required.

APPENDIX C

Discovery

All the WCF calls demonstrated throughout this book share two constraints. First, the port or the pipe assigned to the service must be available. The application developer or the administrator literally has to guess or have some way of knowing the port or the pipe is not used by any other application on that machine. TCP port sharing can alleviate this problem, but does not help with the second constraint: the client must *a priori* know the address of the service endpoints—not just the port number or the pipe name, but also the name of the service machine (in the case of TCP).

It would be great if the service could use any available address, decided upon at runtime. In turn, the client would need to discover that address at runtime. In fact, there is even an industry standard-based solution that stipulates how that discovery takes place. That solution, called simply *discovery* (and its supporting mechanisms), is the subject of this appendix.



Discovery is not a new idea, and veteran developers have often opted for not embedding the type of the class to instantiate in their code and instead resorted to a class factory to resolve or discover the actual type supporting the desired interface:

```
IMyInterface obj = MyClassFactory.CreateInstance<IMyInterface>();
```

In WCF, since the client always programs against the interface or the proxy, the address of the service is analogous to the type of the implementation class in regular .NET programming. Address discovery is therefore a modern incarnation of the class factory of old.

Address Discovery

Discovery relies on *UDP* (User Datagram Protocol). Unlike *TCP*, *UDP* is a connectionless protocol, and no direct connection is required between the packet's sender and the receiver. The client uses *UDP* to broadcast discovery requests for any endpoint supporting a specified contract type. Dedicated *UDP* discovery endpoints that the services support will receive these requests (*WCF* offers a standard implementation of the discovery endpoint). The implementation of the discovery endpoint, in turn, responds back to the client with the address of the service endpoints that support the specified contract. Once the client discovers the services, it continues to invoke them as with regular *WCF* calls. This sequence is illustrated in [Figure C-1](#).

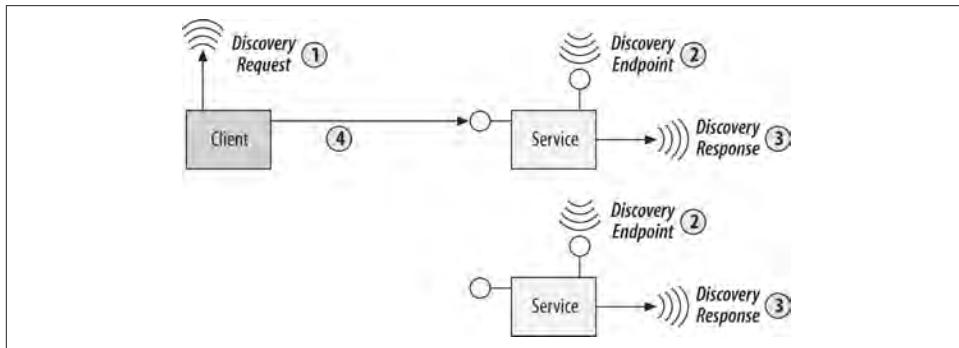


Figure C-1. Address discovery over UDP



Since discovery is predicated on the service responding to the client request, the service process must be running before the client issues the request. This mandates the use of self-hosting. As mentioned in [Chapter 1](#), you should no longer rely on Windows Server AppFabric for its Auto-start feature. Microsoft will be retiring it in April 2016.

Service Configuration

To receive the discovery request, the service must support a discovery endpoint. Much like the MEX endpoint, *WCF* offers a standard discovery endpoint with the type `UdpDiscoveryEndpoint`:

```
public class DiscoveryEndpoint : ServiceEndpoint
{...}
public class UdpDiscoveryEndpoint : DiscoveryEndpoint
{...}
```

The service can have the host implement that endpoint by adding the ServiceDiscoveryBehavior to the collections of behaviors the service supports. You can do that programmatically like this:

```
ServiceHost host = new ServiceHost(...);

host.AddServiceEndpoint(new UdpDiscoveryEndpoint());
ServiceDiscoveryBehavior discovery = new ServiceDiscoveryBehavior();
host.Description.Behaviors.Add(discovery);
host.Open();
```

Example C-1 shows how to add the discovery endpoint and the discovery behavior using the service config file.

Example C-1. Adding discovery endpoint in the config file

```
<services>
  <service name = "MyService">
    <endpoint
      kind = "udpDiscoveryEndpoint"
    />
    ...
  </service>
</services>
<behaviors>
  <serviceBehaviors>
    <behavior>
      <serviceDiscovery/>
    </behavior>
  </serviceBehaviors>
</behaviors>
```



When using discovery for base addresses or business endpoint addresses, it is vital to avoid relative machine names or addresses, such as localhost or 127.0.0.1, and instead use only explicit machine names. The reason is that when the service responds to the client's discovery request, if the response contains a relative machine name, the client will try to invoke the service on the client machine rather than on the service machine.

Dynamic endpoint addresses

Discovery is independent of how the service host defines its endpoints. The service host can respond to discovery requests with the addresses of preconfigured endpoints, where the addresses are either listed in the config file or hardcoded. I call such addresses *static addresses*. Static addresses cater to both clients that use discovery to find the endpoint addresses and clients that use static addresses themselves. The

downsides of static address are the need to deal with configuration and the fact the ports or pipes have to be available on the target machine.

However, what if the client is expected to use discovery to find the service address? In that case, the service is at liberty to configure its endpoint addresses on the fly, dynamically, based on any available port or pipe. I call such addresses *dynamic addresses*.

Note that when using dynamic addresses, the host cannot rely on having WCF implicitly add the default endpoints. This is because WCF will only add the default endpoints if no other endpoint is already defined for the service, and yet the use of dynamic addresses mandates the presence of the discovery endpoint. The host must add the endpoint explicitly by listing them in the config file (using relative addresses only, since the base address is unknown at configuration time), or adding the endpoint programmatically, or explicitly adding the default endpoints.

To automate the use of dynamic addresses, I wrote the `DiscoveryHelper` static helper class with the two properties `AvailableIpcBaseAddress` and `AvailableTcpBaseAddress` (shown in [Example C-2](#)).

Example C-2. Implementing dynamic addresses

```
public static class DiscoveryHelper
{
    public static Uri AvailableIpcBaseAddress
    {
        get
        {
            string machineName = Environment.MachineName;
            return new Uri("net.pipe://" + machineName + "/" + Guid.NewGuid() + "/");
        }
    }

    public static Uri AvailableTcpBaseAddress
    {
        get
        {
            string machineName = Environment.MachineName;
            return new Uri("net.tcp://" + machineName + ":" + FindAvailablePort() + "/");
        }
    }

    static int FindAvailablePort()
    {
        Mutex mutex = new Mutex(false,
                               "ServiceModelEx.DiscoveryHelper.FindAvailablePort");
        try
        {
            mutex.WaitOne();
            IPPEndPoint endPoint = new IPPEndPoint(IPAddress.Any, 0);
            using(Socket socket = new Socket(AddressFamily.InterNetwork,
```

```
        SocketType.Stream,ProtocolType.Tcp))  
    {  
        socket.Bind(endPoint);  
        IPEndPoint local = (IPEndPoint)socket.LocalEndPoint;  
        return local.Port;  
    }  
}  
finally  
{  
    mutex.ReleaseMutex();  
}  
}  
}
```

Implementing `AvailableIpcBaseAddress` is straightforward. Since any uniquely named pipe will do, the property uses a new GUID to name the pipe. Implementing `AvailableTcpBaseAddress` is more challenging because it requires finding an available TCP port (port zero in TCP parlance). To avoid a race condition with other concurrent invocations of `AvailableTcpBaseAddress` on the same machine, `AvailableTcpBaseAddress` uses a named mutex. However, a race condition is still possible with someone else running similar code on the same machine.

As for the machine name, `AvailableTcpBaseAddress` simply uses the current machine name. Example C-3 shows how to use `AvailableTcpBaseAddress`.

Example C-3. Using dynamic addresses

```
Uri baseAddress = DiscoveryHelper.AvailableTcpBaseAddress;

ServiceHost host = new ServiceHost(typeof(MyService),baseAddress);
host.AddDefaultEndpoints();

host.Open();

<service name = "MyService">
    <endpoint
        kind = "udpDiscoveryEndpoint"
    />
</service>

<serviceBehaviors>
    <behavior>
        <serviceDiscovery/>
    </behavior>
</serviceBehaviors>
```

The host in [Example C-3](#) uses the config file to add the discovery endpoint and the discovery behavior. The host uses `AvailableTcpBaseAddress` and explicitly adds the default endpoints.

Enabling discovery

If all you want is a dynamic base address for your service, then the code in [Example C-3](#) is less than perfect, since it still requires you to add discovery, either in the config file or programmatically. You can streamline these steps with my `EnableDiscovery()` host extension, defined as:

```
public static class DiscoveryHelper
{
    public static void EnableDiscovery(this ServiceHost host,
                                       bool enableMEX = true);
}
```

When using `EnableDiscovery()` there is no need for programmatic steps or a config file:

```
Uri baseAddress = DiscoveryHelper.AvailableTcpBaseAddress;

ServiceHost host = new ServiceHost(typeof(MyService),baseAddress);
host.EnableDiscovery();
host.Open();
```

[Example C-4](#) lists the implementation of `EnableDiscovery()`.

Example C-4. Implementing EnableDiscovery()

```
public static class DiscoveryHelper
{
    public static void EnableDiscovery(this ServiceHost host,bool enableMEX = true)
    {
        if(host.Description.Endpoints.Count == 0)
        {
            host.AddDefaultEndpoints();
        }

        host.AddServiceEndpoint(new UdpDiscoveryEndpoint());
        ServiceDiscoveryBehavior discovery = new ServiceDiscoveryBehavior();
        host.Description.Behaviors.Add(discovery);

        if(enableMEX == true)
        {
            host.Description.Behaviors.Add(new ServiceMetadataBehavior());

            foreach(Uri baseAddress in host.BaseAddresses)
            {
                Binding binding = null;
                if(baseAddress.Scheme == "net.tcp")

```

```
        binding = MetadataExchangeBindings.CreateMexTcpBinding();
    }
    if(baseAddress.Scheme == "net.pipe")
    {
        binding = MetadataExchangeBindings.CreateMexNamedPipeBinding();
    }
    if(binding != null)
    {
        host.AddServiceEndpoint(typeof(IMetadataExchange),binding,"MEX");
    }
}
}
```

If the host has not already defined endpoints for the service, `EnableDiscovery()` will add the default endpoints. `EnableDiscovery()` will also default to adding the MEX endpoint to the service on its base addresses.

Client-Side Steps

The client uses the `DiscoveryClient` class to discover all endpoint addresses of all services that support a specified contract:

```
public sealed class DiscoveryClient : ICommunicationObject
{
    public DiscoveryClient();
    public DiscoveryClient(string endpointName);
    public DiscoveryClient(DiscoveryEndpoint discoveryEndpoint);

    public FindResponse Find(FindCriteria criteria);

    //More members
}
```

Logically, `DiscoveryClient` is a proxy to the discovery endpoint. Like all proxies, the client must provide the proxy's constructor with the information about the target endpoint. The client can use a config file to specify the endpoint or programmatically provide the standard UDP discovery endpoint for that purpose, since no further details (such as address or binding) are required. The client then calls the `Find()` method, providing it with the contract type to discover via an instance of `FindCriteria`:

```
public class FindCriteria
{
    public FindCriteria(Type contractType);
    //More members
}
```

`Find()` returns an instance of `FindResponse` which contains a collection of all the discovered endpoints:

```
public class FindResponse
{
    public Collection<EndpointDiscoveryMetadata> Endpoints
    {get;}
    //More members
}
```

Each endpoint is represented by the `EndpointDiscoveryMetadata` class:

```
public class EndpointDiscoveryMetadata
{
    public static EndpointDiscoveryMetadata FromServiceEndpoint(ServiceEndpoint
        endpoint);
    public EndpointAddress Address
    {get;set;}
    //More members
}
```

The main property of the `EndpointDiscoveryMetadata` is `Address`, which finally contains the discovered endpoint address. [Example C-5](#) shows how a client can use these types in conjunction to discover the endpoint address and invoke the service.

Example C-5. Discovering and invoking an endpoint

```
DiscoveryClient discoveryClient = new DiscoveryClient(new UdpDiscoveryEndpoint());

FindCriteria criteria = new FindCriteria(typeof(IMyContract));

FindResponse discovered = discoveryClient.Find(criteria);

discoveryClient.Close();

Debug.Assert(discovered.Endpoints.Count > 0);

//Just grab the first found
EndpointAddress address = discovered.Endpoints[0].Address;
Binding binding = new NetTcpBinding();

IMyContract proxy = ChannelFactory<IMyContract>.CreateChannel(binding,address);
proxy.MyMethod();

(proxy as ICommunicationObject).Close();
```

There are several noteworthy problems with [Example C-5](#):

- While the client may discover multiple endpoints supporting the desired contract, it has no logic to resolve which one to invoke. It simply invokes the first one in the returned collection.

- Discovery is geared toward addresses only. There is no information about which binding to use to invoke the service. [Example C-5](#) simply hardcodes the use of the TCP binding.
- The client will have to repeat these minute steps over and over every time it needs to discover the service address.
- Discovery takes time. By default, `Find()` will wait 20 seconds for the services to respond to the UDP discovery request. Such a delay makes discovery inadequate for use in many applications, certainly when the application performs a high volume of tight calls. While you could shorten that timeout, doing so poses the risk of not discovering any or all of the services. `DiscoveryClient` does offer an asynchronous discovery, but that is of no use for a client that needs to invoke the service before continuing with its execution.

You will see several approaches to addressing these problems in this appendix.



When combining IPC with discovery, the client will discover IPC endpoints on remote machines as well as on the local machine. Since only the IPC endpoints on the local machine are relevant, the client must ignore the remote endpoints by filtering on the local machine name.

Scopes

The use of discovery implies a somewhat loose relationship between the client and the service or services it discovers. This presents another set of problems—how can the client know it has discovered the right endpoint? When multiple compatible endpoints are discovered, which one should the client invoke? Clearly, there is a need for some mechanism that will help the client filter the results of discovery. This is exactly what scopes are about. A *scope* is merely a valid URL (not a mere URI) associated with the endpoint. The service can associate a scope or even multiple scopes with each of its endpoints. The scopes are bundled along with the addresses in the response to the discovery request. In turn, the client can filter the discovered addresses based on the scopes found or, better yet, try to find only relevant scopes in the first place.

Scopes are immensely useful in customizing discovery and in adding sophisticated behavior to your application, especially when writing a framework or administration tools. The classic use for scopes is enabling the client to distinguish between polymorphic services from different applications—that is, when multiple applications share the same intranet and contract type but not the services implementation. However, this is somewhat of a rare occurrence. I find scopes handy when it comes to distinguishing between endpoint types in the same application. For example, suppose you have several implementations for a given contract. You have the operational

mode used in production, the simulation mode used in testing or diagnostics, and an emulation mode used in stress testing. Using scopes, the client can pick and choose the correct implementation type for its needs, and different clients never conflict with each other by consuming one another's services. You can also have the same client pick up a different endpoint based on the context of the invocation. You can have endpoints for profiling, debugging, diagnostics, testing, instrumentation, and so on.

Assigning scopes

The host assigns scopes on a per-endpoint basis using the `EndpointDiscoveryBehavior` class. For example, to apply across all endpoints use a default endpoint behavior:

```
<services>
    <service name = "MyService">
        <endpoint
            ...
        />
        ...
    </service>
</services>
<behaviors>
    ...
    <endpointBehaviors>
        <behavior>
            <endpointDiscovery>
                <scopes>
                    <add scope = "net.tcp://MyApplication"/>
                </scopes>
            </endpointDiscovery>
        </behavior>
    </endpointBehaviors>
</behaviors>
```

You can, of course, apply scopes discretely, based on the type of service, by assigning the behaviors explicitly:

```
<services>
    <service name = "MyService">
        <endpoint behaviorConfiguration = "OperationalScope"
            ...
        />
        ...
    </service>

    <service name = "MySimulator">
        <endpoint behaviorConfiguration = "SimulationScope"
            ...
        />
        ...
    </service>
</services>
```

```

<behaviors>
    ...
    <endpointBehaviors>
        <behavior name = "OperationalScope">
            <endpointDiscovery>
                <scopes>
                    <add scope = "net.tcp://Operational"/>
                </scopes>
            </endpointDiscovery>
        </behavior>
        <behavior name = "SimulationScope">
            <endpointDiscovery>
                <scopes>
                    <add scope = "net.tcp://Simulation"/>
                </scopes>
            </endpointDiscovery>
        </behavior>
    </endpointBehaviors>
</behaviors>

```

A single discovery behavior can list multiple scopes:

```

<endpointDiscovery>
    <scopes>
        <add scope = "net.tcp://MyScope1"/>
        <add scope = "net.tcp://MyScope2"/>
    </scopes>
</endpointDiscovery>

```

If an endpoint has multiple associated scopes, when the client tries to discover the endpoint based on scope matching, the client needs at least one of the scopes to match, but not all of them.

Using scopes

The client has two ways of using scopes. The first is to add the scope to the finding criteria:

```

public class FindCriteria
{
    public Collection<Uri> Scopes
    {get;}
    //More members
}

```

Now, the `Find()` method will return only compatible endpoints that also list that scope.

If the client adds multiple scopes, `Find()` will return only endpoints that support all of the listed scopes. Note that the endpoint may support additional scopes not provided to `Find()`.

The second way of using scopes is to examine the scopes retuned in `FindResponse`:

```
public class EndpointDiscoveryMetadata
{
    public Collection<Uri> Scopes
    {get;}

    //More members
}
```

These are all the scopes supported by the endpoint, and they are useful for additional filtering.



Another use for scopes is in a team environment. When developers develop and test their services on the same intranet, if they work on the same service, or if the services are polymorphic, you run the risk of one developer wanting to test a client against a service on the local machine and unintentionally discovering another developer's service on another machine and invoking that instead. During development, as a precaution, you can use a scope matching the current machine name to avoid the conflict.

Streamlining Discovery

You should encapsulate and automate the manual steps required by the client. To that end, I wrote the helper method `DiscoverAddresses<T>()`, defined as:

```
public static class DiscoveryHelper
{
    public static EndpointAddress[] DiscoverAddresses<T>(Uri scope = null);
    //More members
}
```

Using `DiscoverAddresses<T>()`, [Example C-5](#) is reduced to [Example C-6](#).

Example C-6. Using `DiscoverAddresses<T>()`

```
EndpointAddress[] addresses = DiscoveryHelper.DiscoverAddresses<IMyContract>();

//Just grab the first found
EndpointAddress address = addresses[0];
Binding binding = new NetTcpBinding();

IMyContract proxy = ChannelFactory<IMyContract>.CreateChannel(binding,address);
proxy.MyMethod();
```

```
(proxy as ICommunicationObject).Close();
```

Example C-7 shows the implementation of `DiscoverAddresses<T>()`.

Example C-7. Implementing DiscoverAddresses<T>()

```
public static class DiscoveryHelper
{
    public static EndpointAddress[] DiscoverAddresses<T>(Uri scope = null)
    {
        DiscoveryClient discoveryClient =
            new DiscoveryClient(new UdpDiscoveryEndpoint());

        FindCriteria criteria = new FindCriteria(typeof(T));

        if(scope != null)
        {
            criteria.Scopes.Add(scope);
        }

        FindResponse discovered = discoveryClient.Find(criteria);
        discoveryClient.Close();

        return discovered.Endpoints.Select(endpoint=>endpoint.Address).ToArray();
    }
}
```

However, Example C-6 still has the rest of the problems of Example C-5, namely, which endpoint to invoke in case of discovering multiple endpoints, the discovery latency, and hardcoding the binding.

Discovery cardinality

The most acute of these problems is what I call *discovery cardinality*, that is, how many endpoints are discovered and which one, if any, to invoke. There are several cases of cardinality:

- No endpoint is discovered. In this case, the client needs to deal with the absence of the service. This is no different from any other WCF client whose service is unavailable.
- Exactly one compatible endpoint is discovered. This is by far both the most common and straightforward case—the client simply proceeds to invoke the service.
- Multiple endpoints are discovered. Here, the client (in theory) has two options. The first is to invoke all of the discovered endpoints. This is the case with a publisher firing an event at subscribers (discussed in Appendix D), and is a valid scenario. The second option is to invoke some (including only one) but not all of the

discovered endpoints. I find this scenario to be moot—any attempt to place logic in the client that resolves which endpoint to invoke creates too much coupling across the system and negates the very notion of runtime discovery, namely, that any discovered endpoint will do. If it is possible to discover undesirable endpoints, then using discovery is a poor design choice, and you should instead provide static addresses to the client.

Single endpoint

If the client expects to discover exactly one endpoint (cardinality of one), the client should instruct `Find()` to return as soon as it finds that endpoint. Doing so will drastically reduce the discovery latency and make it adequate for the majority of cases.

The client can configure the cardinality using the `MaxResults` property of `FindCriteria`:

```
public class FindCriteria
{
    public int MaxResults
    {get;set;}
    //More members
}

FindCriteria criteria = new FindCriteria(typeof(IMyContract));
criteria.MaxResults = 1;
```



An interesting situation occurs when the client sets `MaxResults` to 1 and yet multiple compatible endpoints are available. In this case, each invocation of `Find()` is liable to retrieve a different endpoint in a nondeterministic manner. It may never discover a particular endpoint or it may discover the same endpoint repeatedly. Do not attempt to use `MaxResults` set to 1 as a crude load balancer. My testing indicates a disparity of as much as 600:1 between the available endpoints.

You can streamline the case of cardinality of one using my `DiscoveryHelper.DiscoverAddress<T>()` helper method, defined as:

```
public static class DiscoveryHelper
{
    public static EndpointAddress DiscoverAddress<T>(Uri scope = null);
    //More members
}
```

Using `DiscoverAddress<T>()`, [Example C-6](#) is reduced to [Example C-8](#).

Example C-8. Using DiscoverAddress<T>()

```
EndpointAddress address = DiscoveryHelper.DiscoverAddress<IMyContract>();  
  
Binding binding = new NetTcpBinding();  
IMyContract proxy = ChannelFactory<IMyContract>.CreateChannel(binding, address);  
proxy.MyMethod();  
  
(proxy as ICommunicationObject).Close();
```

Example C-9 shows the implementation of `DiscoverAddress<T>()` along with a refactored `DiscoverAddresses<T>()`.

Example C-9. Implementing DiscoverAddress<T>()

```
public static class DiscoveryHelper  
{  
    static EndpointAddress[] Discover<T>(int maxResults, Uri scope)  
    {  
        DiscoveryClient discoveryClient =  
            new DiscoveryClient(new UdpDiscoveryEndpoint());  
        FindCriteria criteria = new FindCriteria(typeof(T));  
        criteria.MaxResults = maxResults;  
        if(scope != null)  
        {  
            criteria.Scopes.Add(scope);  
        }  
        FindResponse discovered = discoveryClient.Find(criteria);  
        discoveryClient.Close();  
  
        return discovered.Endpoints.Select((endpoint) => endpoint.Address).ToArray();  
    }  
    public static EndpointAddress DiscoverAddress<T>(Uri scope = null)  
    {  
        EndpointAddress[] addresses = Discover<T>(1, scope);  
        Debug.Assert(addresses.Length == 1);  
  
        return addresses[0];  
    }  
    public static EndpointAddress[] DiscoverAddresses<T>(Uri scope = null)  
    {  
        return Discover<T>(int.MaxValue, scope);  
    }  
}
```

Binding discovery

So far, when it comes to determining which binding to use, the client has had two options. First, the client can hardcode the binding type and its properties. Second, the client can improve on that approach by inferring the binding type from the discov-

ered address scheme and only hardcode the binding properties per binding type rather than the binding type itself.

However, a third option is available: if the service supports a MEX endpoint—the client can first discover the MEX endpoint address and then proceed to retrieve and process the metadata to obtain the binding type to use, along with its properties. I call this technique *binding discovery*. To help with MEX endpoint discovery, the `FindCriteria` class offers the static method `CreateMetadataExchangeEndpointCriteria()`, which returns a matching criteria for endpoints that support `IMetadataExchange`:

```
public class FindCriteria
{
    public static FindCriteria CreateMetadataExchangeEndpointCriteria();

    //More members
}
```

You can call `CreateMetadataExchangeEndpointCriteria()` and further refine the criteria by adding scopes or setting `MaxResults`.

To streamline binding discovery, use my `DiscoveryHelper.DiscoverBinding<T>()` method, defined as:

```
public static class DiscoveryHelper
{
    public static Binding DiscoverBinding<T>(Uri scope = null);
    //More members
}
```

Using `DiscoveryHelper.DiscoverBinding<T>()`, you can convert [Example C-8](#) to [Example C-10](#).

Example C-10. Using DiscoverBinding<T>()

```
EndpointAddress address = DiscoveryHelper.DiscoverAddress<IMyContract>();
Binding binding = DiscoveryHelper.DiscoverBinding<IMyContract>();

IMyContract proxy = ChannelFactory<IMyContract>.CreateChannel(
    binding, address);
proxy.MyMethod();

(proxy as ICommunicationObject).Close();
```

[Example C-11](#) shows the implementation of `DiscoverBinding<T>()`.

Example C-11. Implementing DiscoverBinding<T>()

```
public static class DiscoveryHelper
{
```

```

public static Binding DiscoverBinding<T>(Uri scope = null)
{
    DiscoveryClient discoveryClient = new DiscoveryClient(
        new UdpDiscoveryEndpoint());
    FindCriteria criteria =
        FindCriteria.CreateMetadataExchangeEndpointCriteria();
    criteria.MaxResults = 1;
    if(scope != null)
    {
        criteria.Scopes.Add(scope);
    }
    FindResponse discovered = discoveryClient.Find(criteria);
    discoveryClient.Close();

    Debug.Assert(discovered.Endpoints.Count == 1);

    Uri mexAddress = discovered.Endpoints[0].Address.Uri;

    ServiceEndpoint[] endpoints = MetadataHelper.GetEndpoints(
        mexAddress.AbsoluteUri, typeof(T));

    Debug.Assert(endpoints.Length == 1);

    return endpoints[0].Binding;
}
}

```

`DiscoverBinding<T>()` assumes there is exactly one MEX endpoint to discover. Once that MEX endpoint is discovered, `DiscoverBinding<T>()` uses my `MetadataHelper` class (discussed in [Chapter 2](#)) to obtain the service metadata. `DiscoverBinding<T>()` assumes there is exactly one endpoint on the service supporting the desired contract and it returns that endpoint's binding.

Discovery factory

There is an obvious problem with [Example C-10](#). The client has to pay for two discovery requests. It would be preferable to combine the two discovery attempts into a single call using my `DiscoveryFactory.CreateChannel<T>()` method:

```

public static class DiscoveryFactory
{
    public static T CreateChannel<T>(Uri scope = null) where T : class;
    //More members
}

```

Using `CreateChannel<T>()`, [Example C-10](#) is reduced to:

```

IMyContract proxy = DiscoveryFactory.CreateChannel<IMyContract>();
proxy.MyMethod();
(proxy as ICommunicationObject).Close();

```

Example C-12 shows the implementation of `CreateChannel<T>()`.

Example C-12. Implementing CreateChannel<T>()

```
public static class DiscoveryFactory
{
    public static T CreateChannel<T>(Uri scope = null) where T : class
    {
        DiscoveryClient discoveryClient =
            new DiscoveryClient(new UdpDiscoveryEndpoint());

        FindCriteria criteria =
            FindCriteria.CreateMetadataExchangeEndpointCriteria();

        criteria.MaxResults = 1;

        if(scope != null)
        {
            criteria.Scopes.Add(scope);
        }
        FindResponse discovered = discoveryClient.Find(criteria);
        discoveryClient.Close();

        Debug.Assert(discovered.Endpoints.Count == 1);

        Uri mexAddress = discovered.Endpoints[0].Address.Uri;

        ServiceEndpoint[] endpoints = MetadataHelper.GetEndpoints(
            mexAddress.AbsoluteUri, typeof(T));
        Debug.Assert(endpoints.Length == 1);

        Binding binding = endpoints[0].Binding;
        EndpointAddress address = endpoints[0].Address;

        return ChannelFactory<T>.CreateChannel(binding, address);
    }
}
```

`CreateChannel<T>()` assumes cardinality of one with the MEX endpoint (that is, only a single discoverable MEX endpoint is found in the local network). It also assumes the metadata contains exactly one endpoint whose contract is the specified type parameter `T`.

Note that `CreateChannel<T>()` uses the MEX endpoint for both the endpoint binding and address. The service is expected to support both a MEX endpoint and a discovery endpoint (although the client never uses the discovery endpoint to find the actual endpoint). This is why the `EnableDiscovery()` extension method of **Example C-4** adds the MEX endpoint by default.

In case there are multiple services supporting the desired service contract or there are multiple MEX endpoints, `DiscoveryFactory` also offers the `CreateChannels<T>()` method:

```
public static class DiscoveryFactory
{
    public static T[] CreateChannels<T>(bool inferBinding = true) where T : class;
    //More members
}
```

By default, `CreateChannels<T>()` will infer which binding to use from the scheme of the service endpoint. If `inferBinding` is `false`, it will discover the binding from the MEX endpoints.

`CreateChannels<T>()` does not assume a cardinality of 1 on the compatible service endpoints or the MEX endpoints, and will return an array of all compatible endpoints.

Example C-13 shows the implementation of `CreateChannels<T>()`.

Example C-13. Implementing CreateChannels<T>()

```
public static class DiscoveryFactory
{
    public static T[] CreateChannels<T>(bool inferBinding = true) where T : class
    {
        if(inferBinding)
        {
            return CreateInferredChannels<T>();
        }
        else
        {
            return CreateChannelsFromMex<T>();
        }
    }
    static T[] CreateChannelsFromMex<T>() where T : class
    {
        DiscoveryClient discoveryClient =
            new DiscoveryClient(new UdpDiscoveryEndpoint());
        FindCriteria criteria =
            FindCriteria.CreateMetadataExchangeEndpointCriteria();

        FindResponse discovered = discoveryClient.Find(criteria);
        discoveryClient.Close();

        List<T> list = new List<T>();

        foreach(EndpointDiscoveryMetadata mexEndpoint in discovered.Endpoints)
        {
            ServiceEndpoint[] endpoints = MetadataHelper.GetEndpoints(
                mexEndpoint.Address.Uri.AbsoluteUri,typeof(T));
        }
    }
}
```

```

        foreach(ServiceEndpoint endpoint in endpoints)
        {
            T proxy = ChannelFactory<T>.CreateChannel(endpoint.Binding,
                                            endpoint.Address);
            list.Add(proxy);
        }
    }
    return list.ToArray();
}
static T[] CreateInferredChannels<T>() where T : class
{
    DiscoveryClient discoveryClient =
        new DiscoveryClient(new UdpDiscoveryEndpoint());
    FindCriteria criteria = new FindCriteria(typeof(T));
    FindResponse discovered = discoveryClient.Find(criteria);
    discoveryClient.Close();

    List<T> list = new List<T>();
    foreach(EndpointDiscoveryMetadata endpoint in discovered.Endpoints)
    {
        Binding binding = InferBindingFromUri(endpoint.Address.Uri);
        T proxy = ChannelFactory<T>.CreateChannel(binding,endpoint.Address);
        list.Add(proxy);
    }
    return list.ToArray();
}
static Binding InferBindingFromUri(Uri address)
{
    switch(address.Scheme)
    {
        case "net.tcp":
        {
            NetTcpBinding tcpBinding =
                new NetTcpBinding(SecurityMode.Transport,true);
            tcpBinding.TransactionFlow = true;
            return tcpBinding;
        }
        case "net.pipe":
        {
            NetNamedPipeBinding ipcBinding = new NetNamedPipeBinding();
            ipcBinding.TransactionFlow = true;
            return ipcBinding;
        }
        case "net.msmq":
        {
            NetMsmqBinding msmqBinding = new NetMsmqBinding();
            msmqBinding.Security.Transport.MsmqProtectionLevel =
                ProtectionLevel.EncryptAndSign;
            return msmqBinding;
        }
        default:

```

```
        {
            throw new InvalidOperationException("Can only create a channel
                                              over TCP/IP/MSMQ bindings");
        }
    }
}
```

`CreateChannels<T>()` uses two private methods, `CreateChannelsFromMex<T>()` and `CreateInferredChannels<T>()`, depending on the value of `inferBinding`. `CreateChannelsFromMex<T>()` is similar to `CreateChannel<T>()`, except it does not assume any cardinality on the MEX or service endpoint. It will return an array of proxies targeting all the compatible endpoints described in the all the MEX endpoints it discovers. `CreateInferredChannels<T>()` also does not assume any cardinality. It will use discovery to find and return proxies to all compatible service endpoints. It uses the `InferBindingFromUri()` helper method to infer the binding type from the endpoint address scheme and hardcodes the binding properties by using transactions, reliability, and security.

Creating discoverable host

My `DiscoveryFactory` class also provides the `CreateDiscoverableHost<T>()` method, defined as:

```
public static class DiscoveryHelper
{
    public static ServiceHost<T> CreateDiscoverableHost<T>(bool supportIpc = false)
    {
        ServiceHost<T> host;
        if(supportIpc == true)
        {
            host = new ServiceHost<T>(DiscoveryHelper.AvailableIpcBaseAddress,
                                         DiscoveryHelper.AvailableTcpBaseAddress);
        }
        else
        {
            host = new ServiceHost<T>(DiscoveryHelper.AvailableTcpBaseAddress);
        }
        host.EnableDiscovery();

        return host;
    }
    //More members
}
```

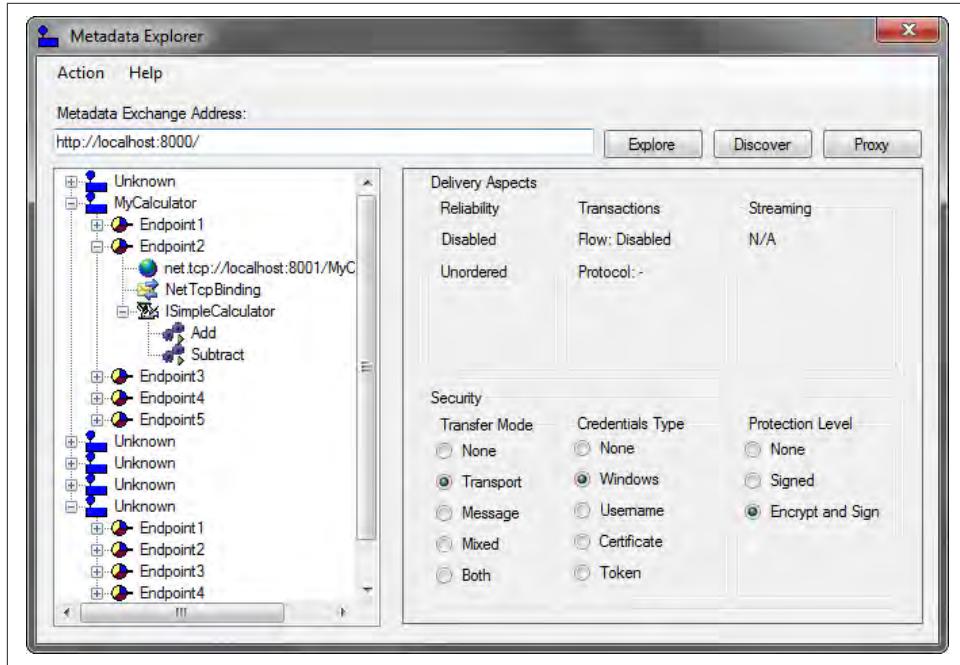
`CreateDiscoverableHost<T>()` returns a host instance with a dynamic TCP base address by default, which supports discovery. The returned host instance also adds

the default endpoints. Consequently, there is no need for any config file or additional programmatic setting other than opening the host:

```
ServiceHost host = DiscoveryHelper.CreateDiscoverableHost<MyService>();  
host.Open();
```

The Metadata Explorer

I have revamped the Metadata Explorer tool presented in [Chapter 1](#) to take advantage of discovery (see [Figure C-2](#)).



Delivery Aspects	Transactions	Streaming
Reliability	Flow: Disabled	N/A
Unordered	Protocol: -	

Security		
Transfer Mode	Credentials Type	Protection Level
<input type="radio"/> None	<input type="radio"/> None	<input type="radio"/> None
<input checked="" type="radio"/> Transport	<input checked="" type="radio"/> Windows	<input type="radio"/> Signed
<input type="radio"/> Message	<input type="radio"/> Username	<input checked="" type="radio"/> Encrypt and Sign
<input type="radio"/> Mixed	<input type="radio"/> Certificate	
<input type="radio"/> Both	<input type="radio"/> Token	

Figure C-2. The MEX Explorer with discovery

Clicking the Discover button triggers a discovery request for all MEX endpoints without any limit on cardinality. The tool then displays all discovered endpoints in the tree.

Ongoing Discovery

The discovery mechanism as described thus far assumes a static services environment. The services are always running and always connected, and it is up to the clients to discover them. Under this assumption, the clients need to make only a single discovery effort. However, what if the environment is more dynamic? What if services come and go or the network is shaky, so that services are intermittently disconnected? For such an environment, it is up to the client to continuously try to discover

the services and maintain a list of available services. To that end, I wrote the class `DiscoveredServices<T>`, defined as:

```
public class DiscoveredServices<T> : AddressesContainer<T> where T: class
{
    public DiscoveredServices(DiscoveredServices<T> container = null);

    public void Abort();
}
```

`DiscoveredServices<T>` is designed to maintain as much as possible an up-to-date list of all discovered services, and it stores the addresses it discovers in its base class, `AddressesContainer<T>`. `AddressesContainer<T>` is a rich address management helper collection that you can use whenever you need to manipulate multiple addresses. `AddressesContainer<T>` supports several iterators, indexers, conversion methods, and queries such as creating a union of two `AddressesContainer<T>` and a complement (finding all the addresses that do not have a specified scope). [Example C-14](#) lists `AddressesContainer<T>`.

Example C-14. The `AddressesContainer<T>` base class

```
public abstract class AddressesContainer<T> : IEnumerable<EndpointAddress>,
    IEnumerable<KeyValuePair<EndpointAddress,Collection<Uri>>>,
    IDisposable where T : class
{
    public AddressesContainer();
    public AddressesContainer(AddressesContainer<T> container);

    public abstract void Close();
    public abstract void Open();

    public void Dispose();

    public EndpointAddress[] Find(Uri scope);

    public EndpointAddress[] FindComplement(Uri scopeToExclude);

    public static Dictionary<EndpointAddress, Collection<Uri>> FindUnion(
        AddressesContainer<T> container1, AddressesContainer<T> container2);

    protected readonly Dictionary<EndpointAddress,Uri[]> Dictionary;

    public void Remove(EndpointAddress address);

    public int Count
    {get;}

    public EndpointAddress[] Addresses
    {get;}
```

```

    public Collection<Uri> this[EndpointAddress address]
    {get;}
    public EndpointAddress this[int index]
    {get;}
    protected string Namespace
    {get;}
}

```

`DiscoveredServices<T>` spins off an ongoing discovery as background activity, shown in [Example C-15](#), without some of the error handling and synchronization code.

Example C-15. Implementing DiscoveredServices<T>

```

public class DiscoveredServices<T> : AddressesContainer<T> where T : class
{
    Thread m_WorkerThread;

    bool Terminate
    {get;set;}

    public override void Open()
    {
        m_WorkerThread = new Thread(Discover);
        m_WorkerThread.Start();
    }
    public override void Close()
    {
        Terminate = true;
        m_WorkerThread.Join();
    }
    public void Abort()
    {
        Terminate = true;
        Thread.Sleep(0);
        m_WorkerThread.Abort();
        m_WorkerThread.Join();
    }
    void Discover()
    {
        while(Terminate == false)
        {
            DiscoveryClient discoveryClient =
                new DiscoveryClient(new UdpDiscoveryEndpoint());
            FindCriteria criteria = new FindCriteria(typeof(T));
            FindResponse discovered = discoveryClient.Find(criteria);
            discoveryClient.Close();

            lock(this)
            {
                Dictionary.Clear();
            }
        }
    }
}

```

```
        foreach(EndpointDiscoveryMetadata endpoint in discovered.Endpoints)
        {
            Dictionary[endpoint.Address] = endpoint.Scopes;
        }
    }
}
```

When you open `DiscoveredServices<T>`, it creates a worker thread whose thread's method repeatedly discovers compatible service endpoints for the service contract `T` and updates the base dictionary. When you close `DiscoveredServices<T>`, it terminates that worker thread gracefully. You can also abort the worker thread.

Announcements

The discovery mechanism presented thus far is passive from the perspective of the service. The client queries the discovery endpoint and the service responds. As an alternative to this passive address discovery, WCF offers an active model in which the service broadcasts its status to all clients and provides its address. The service host broadcasts a “hello” announcement when the host is opened and a “bye” announcement when the host shuts down gracefully. If the host is aborted ungracefully, no “bye” announcement is sent. These announcements are received on a special announcements endpoint hosted by the client (see [Figure C-3](#)).

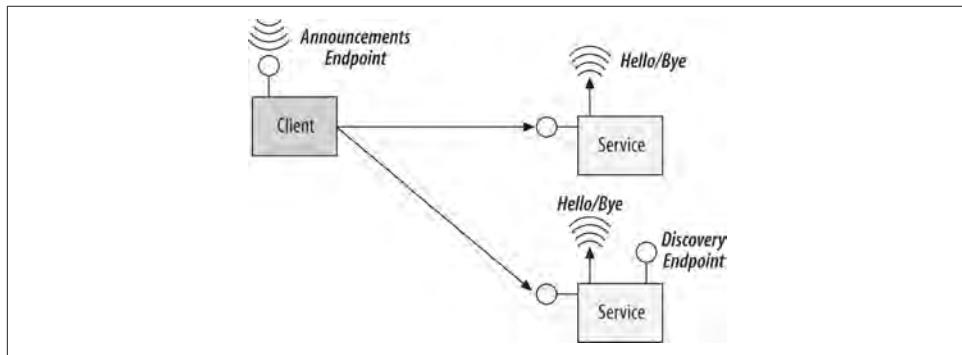


Figure C-3. The announcement architecture

Announcements are an individual endpoint-level mechanism, not a host-level one. The host can choose which endpoint to announce. Each announcement contains the endpoint address, its scopes, and its contract.

Note that announcements are unrelated to address discovery. The host may not support a discovery endpoint at all, and there is no need for the discovery behavior. On

the other hand, the host may both support the discovery endpoint and announce its endpoints, as shown in [Figure C-3](#).

Announcing Endpoints

The host can manually announce its endpoints using the `AnnouncementClient` class, defined as:

```
public sealed class AnnouncementClient : ICommunicationObject
{
    public AnnouncementClient();
    public AnnouncementClient(AnnouncementEndpoint announcementEndpoint);
    public AnnouncementClient(string endpointName);

    public void AnnounceOffline(EndpointDiscoveryMetadata metadata);
    public void AnnounceOnline(EndpointDiscoveryMetadata metadata);

    //More members
}
```

`AnnouncementClient` is actually a proxy to the clients' announcement endpoints. As with every proxy, you need to provide its constructor with information about the target endpoint. As with the discovery endpoints, WCF provides a standard endpoint, the `UdpAnnouncementEndpoint`, defined as:

```
public class AnnouncementEndpoint : ServiceEndpoint
{
}
public class UdpAnnouncementEndpoint : AnnouncementEndpoint
{
}
```

[Example C-16](#) shows a host manually announcing all its endpoints.

Example C-16. Using AnnouncementClient

```
ServiceHost host = ...;
host.Open();

AnnouncementClient announcementClient = new AnnouncementClient(
    new UdpAnnouncementEndpoint());

foreach(ServiceEndpoint endpoint in host.Description.Endpoints)
{
    EndpointDiscoveryMetadata metadata = EndpointDiscoveryMetadata.
        FromServiceEndpoint(endpoint);
    announcementClient.AnnounceOnline(metadata);
}

//Sometimes later:
host.Close();

foreach(ServiceEndpoint endpoint in host.Description.Endpoints)
```

```

{
    EndpointDiscoveryMetadata metadata = EndpointDiscoveryMetadata.
        FromServiceEndpoint(endpoint);
    announcementClient.AnnounceOffline(metadata);
}

announcementClient.Close();

```

The host can provide the information about the announcement endpoint in its config file:

```

<system.serviceModel>
    <client>
        <endpoint
            kind = "udpAnnouncementEndpoint"
        />
    </client>
</system.serviceModel>

```

In this case, there is no need to provide the endpoint instance to the constructor of `AnnouncementClient`:

```
AnnouncementClient announcementProxy = new AnnouncementClient();
```

Automatic announcements

Instead of the tedious manual announcements used in [Example C-16](#), the host can automatically announce its endpoints. All you need to do is provide the information about the client announcement endpoint for the discovery behavior. For example, when using a config file:

```

<behavior>
    <serviceDiscovery>
        <announcementEndpoints>
            <endpoint
                kind = "udpAnnouncementEndpoint"
            />
        </announcementEndpoints>
    </serviceDiscovery>
</behavior>

```

Note again that while the config file contains the discovery behavior, there is no need for a discovery endpoint, but you may add it if you want to.

Receiving Announcements

WCF provides a pre-canned implementation of an announcements endpoint with the `AnnouncementService` class:

```

public class AnnouncementEventArgs : EventArgs
{
    public EndpointDiscoveryMetadata EndpointDiscoveryMetadata

```

```

    {get;}
    //More members
}

[ServiceBehavior(InstanceContextMode = InstanceContextMode.Single,
                 ConcurrencyMode = ConcurrencyMode.Multiple)]
public class AnnouncementService : ...
{
    public event EventHandler<AnnouncementEventArgs> OfflineAnnouncementReceived;
    public event EventHandler<AnnouncementEventArgs> OnlineAnnouncementReceived;
    //More members
}

```

AnnouncementService is a singleton service configured for concurrent access. AnnouncementService provides two event delegates that the client can subscribe to in order to receive the announcements. The client should host the AnnouncementService using the constructor of ServiceHost, which accepts a singleton instance. This is required so the client can interact with the instance and subscribe to the events. In addition, the client must add the UDP announcement endpoint to the host. [Example C-17](#) demonstrates these steps.

Example C-17. Hosting AnnouncementService

```

AnnouncementService announcementService = new AnnouncementService();
announcementService.OnlineAnnouncementReceived += OnHello;
announcementService.OfflineAnnouncementReceived += OnBye;

ServiceHost host = new ServiceHost(announcementService);

host.AddServiceEndpoint(new UdpAnnouncementEndpoint());

host.Open();

void OnHello(object sender, AnnouncementEventArgs args)
{...}

void OnBye(object sender, AnnouncementEventArgs args)
{...}

```

There is one important detail related to receiving announcements. The client receives all notifications of all services on the intranet, regardless of contract type or, for that matter, applications or scopes. The client must filter out the relevant announcements.



The MEX Explorer utilizes announcements of MEX endpoints. In responding to these events, the MEX Explorer refreshes itself and presents the new endpoints or removes those that are no longer running from the tree.

Streamlining Announcements

You can easily add automatic announcement to the `EnableDiscovery()` method of [Example C-4](#) by adding the announcement endpoint to the discovery behavior:

```
public static class DiscoveryHelper
{
    public static void EnableDiscovery(this ServiceHost host,
                                       bool enableMEX = true)
    {
        ...
        host.AddServiceEndpoint(new UdpDiscoveryEndpoint());
        ServiceDiscoveryBehavior discovery = new ServiceDiscoveryBehavior();
        discovery.AnnouncementEndpoints.Add(new UdpAnnouncementEndpoint());
        host.Description.Behaviors.Add(discovery);
        ...
    }
}
```



Since my `CreateDiscoverableHost<T>()` method of `DiscoveryFactory` uses `EnableDiscovery()`, the returned host instance is not just discoverable, it also supports announcements.

The announcements sink

You can greatly simplify and improve on the raw steps required of the client to utilize announcements using my `AnnouncementSink<T>` class, defined as:

```
public class AnnouncementSink<T> : AddressesContainer<T> where T: class
{
    public AnnouncementSink(AnnouncementSink<T> container = null);

    public event Action<string,Uri[]> OnlineAnnouncementReceived;
    public event Action<string,Uri[]> OfflineAnnouncementReceived;
}
```

`AnnouncementSink<T>` automates hosting the announcements endpoint by encapsulating the steps in [Example C-17](#). While `AnnouncementSink<T>` hosts an instance of `AnnouncementService` internally, it improves on its deficiencies. First, `AnnouncementSink<T>` offers two event delegates for notifications. Unlike the raw `AnnouncementService`, `AnnouncementSink<T>` fires these delegates concurrently. In addition, `AnnouncementSink<T>` disables the synchronization context affinity of `AnnouncementService` so that it can accept the announcements on any incoming thread, making it truly concurrent. `AnnouncementSink<T>` filters the contract types and fires its events only when compatible endpoints announce themselves. The only thing the client needs to do is open and close `AnnouncementSink<T>` to indicate when to start and stop receiving notifications.

`AnnouncementSink<T>` derives my general-purpose address container. As long as the service announces itself before the client tries to obtain its address, the client can use `AnnouncementSink<T>` without ever subscribing to its events, simply by using the base `AddressesContainer<T>`:

```
class MyClient
{
    AddressesContainer<IMyContract> m_Addresses;

    public MyClient()
    {
        m_Addresses = new AnnouncementSink<IMyContract>();
        m_Addresses.Open();

        ...
    }
    void CallService()
    {
        EndpointAddress address = m_Addresses[0];

        IMyContract proxy = ChannelFactory<IMyContract>.CreateChannel(
            new NetTcpBinding(),address);
        proxy.MyMethod();
        (proxy as ICommunicationObject).Close();
    }
    void Close()
    {
        m_Addresses.Close();
    }
}
```

The following shows how the client can use the events of `AnnouncementSink<T>`:

```
class MyClient
{
    AnnouncementSink<IMyContract> m_AnnouncementSink;

    public MyClient()
    {
        m_AnnouncementSink = new AnnouncementSink<IMyContract>();
        m_AnnouncementSink.OnlineAnnouncementReceived += OnHello;
        m_AnnouncementSink.Open();

        ...
    }
    void Close()
    {
        m_AnnouncementSink.Close();
    }
    void OnHello(string address,Uri[] scopes)
    {
```

```

        EndpointAddress endpointAddress = new EndpointAddress(address);

        IMyContract proxy = ChannelFactory<IMyContract>.CreateChannel(
            new NetTcpBinding(), endpointAddress);
        proxy.MyMethod();
        (proxy as ICommunicationObject).Close();
    }
}

```

Example C-18 lists the implementation of `AnnouncementSink<T>`.

Example C-18. Implementing AnnouncementSink<T>

```

public abstract class AddressesContainer<T> : ...
{
    protected string Namespace
    {
        get
        {
            ContractDescription description =
                ContractDescription.GetContract(typeof(T));
            return description.Namespace;
        }
    }
    //Rest of the implementation
}

public class AnnouncementSink<T> : AddressesContainer<T> where T : class
{
    readonly ServiceHost m_Host;

    public event Action<string, Uri[]> OnlineAnnouncementReceived = delegate{};
    public event Action<string, Uri[]> OfflineAnnouncementReceived = delegate{};

    public AnnouncementSink(AnnouncementSink<T> container = null) : base(container)
    {
        AnnouncementService announcementService = new AnnouncementService();
        m_Host = new ServiceHost(announcementService);
        m_Host.Description.Behaviors.Find<ServiceBehaviorAttribute>().
            UseSynchronizationContext = false;

        m_Host.AddServiceEndpoint(new UdpAnnouncementEndpoint());

        announcementService.OfflineAnnouncementReceived += OnBye;
        announcementService.OnlineAnnouncementReceived += OnHello;
    }
    public override void Open()
    {
        m_Host.Open();
    }
    public override void Close()
    {

```

```

        m_Host.Close();
    }
    protected void OnHello(object sender, AnnouncementEventArgs args)
    {
        foreach(XmlQualifiedName contract in
            args.EndpointDiscoveryMetadata.ContractTypeNames)
        {
            if(contract.Name == typeof(T).Name && contract.Namespace == Namespace)
            {
                PublishAvailabilityEvent(OnlineAnnouncementReceived,
                    args.EndpointDiscoveryMetadata.Address.Uri.AbsoluteUri,
                    args.EndpointDiscoveryMetadata.Scopes.ToArray());
                Dictionary[args.EndpointDiscoveryMetadata.Address] =
                    args.EndpointDiscoveryMetadata.Scopes.ToArray();
            }
        }
    }
    protected void OnBye(object sender, AnnouncementEventArgs args)
    {
        //Same as OnHello but fires OfflineAnnouncementReceived and
        //removes from dictionary
    }
    void PublishAvailabilityEvent(Action<string, Uri[]> notification,
                                string address, Uri[] scopes)
    {
        Delegate[] subscribers = notification.GetInvocationList();
        Action<Delegate> publish = (subscriber =>
            subscriber.DynamicInvoke(address, scopes));
        subscribers.ForEachAsync(publish);
    }
}

```

The constructor of `AnnouncementSink<T>` is similar to the code in [Example C-17](#). The heart of the implementation is the `OnHello` and `OnBye` methods. These methods examine each announcement, looking for the contract type and namespace. If a match is found, they fire their respective events asynchronously and concurrently using threads from the .NET thread pool. The methods also update the base address container with the address and scopes.

TCP-Based Discovery

WCF's built-in discovery mechanism relies on UDP to broadcast requests of discovery and availability announcements throughout your network. But as a communication protocol, UDP is noisy and unreliable. Because of this, many Enterprise and commercial on-premises environments by policy or by accident do not allow UDP multicast to transit their networks. This means in those environments you cannot use WCF's standard-based mechanism for service discovery. Ironically, it is in these very same environments that you most often need discovery. Enterprise and customer IT

operational policy often requires complete control over where they deploy your services. Without discovery, these volatile addressing scenarios require you to either constantly update all of your client configurations manually or devise a centralized configuration facility.

The reality is that as useful as UDP-based discovery is, it is often unavailable in Enterprise environments. This is especially challenging for independent software vendors who have no control over the end customer's IT department, and where some of the customers support discovery while others don't, forcing drastic branches in the code base, the deployment model or both.

The Discovery Proxy Pattern

To address the lack of discovery you can utilize a simple pattern called the *discovery proxy*. The discovery proxy allows you to provide service discovery to your clients through a dedicated discovery service. That discovery service resides on a well-known address through which clients can query a service registry or repository for services that may exist within the system environment. [Figure C-4](#) shows this pattern. The discovery proxy requires the services to first register with the discovery proxy the full details of their current location and available endpoints (step 1 in [Figure C-4](#)). The discovery proxy then stores each service's configuration in a repository that may be either in-memory or persistent (step 2 in [Figure C-4](#)). Once registered, clients can then query the discovery proxy for services that support a given contract (step 3 in [Figure C-4](#)). In response to the query, the discovery proxy will return the address of each endpoint registered by a service that matched the query parameters. The client can then use the endpoint address to call the discovered service directly (step 4 in [Figure C-4](#)).

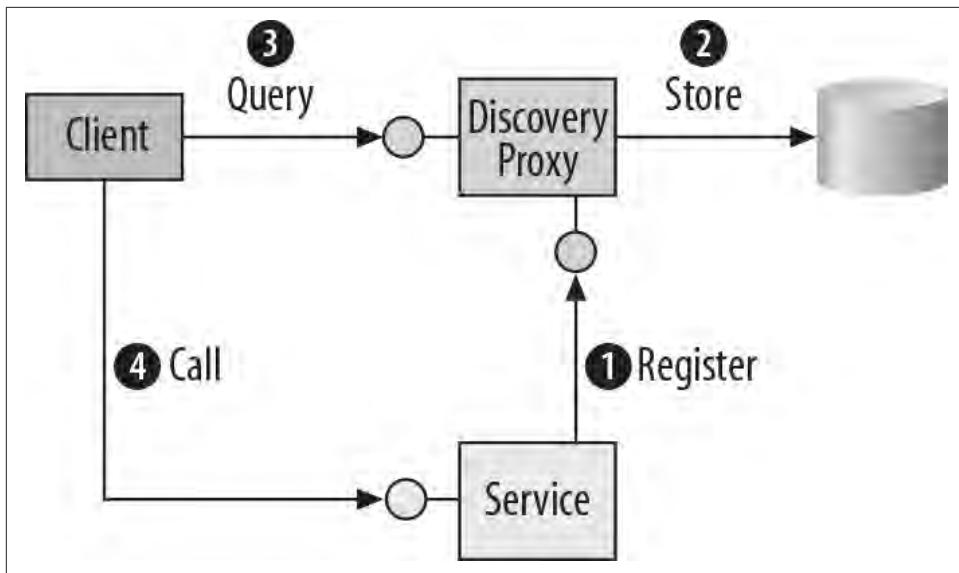


Figure C-4. The Discovery Proxy Pattern

While this discovery proxy design does solve the problem of discovery within intranet environments that do not permit UDP multicast, it does present a number of deficiencies:

- The pattern does not support the real-time detection of changes in the environment. If the service address has changed after it has registered with the discovery proxy then the information in the repository becomes stale and useless for the clients.
- There is no available implementation of the discovery proxy pattern. Everyone has to construct their own version.
- The programming model the discovery pattern imposes on the clients and services does not resemble at all WCF discovery. This lack of consistency makes it difficult to maintain the same code bases across environments with different discovery policies.
- There is no support for service metadata discovery. Clients must have a priori knowledge of the binding requirements of service endpoints.
- The design does not support service availability announcements.

Without the capability of dynamic, runtime discovery of services or service availability, the discovery proxy pattern is a poor substitute for WCF discovery. In essence, it's simply a centralized service metadata repository. In fact, the discovery proxy pattern does not represent a true discovery mechanism at all.

Solution Architecture

At the heart of all true discovery mechanisms is the broadcast of a request for discovery. Broadcasting the request to all available parties avoids all the problems with the discovery proxy pattern by enabling dynamic discovery. While WCF uses UDP for that broadcast, you can approximate WCF discovery by using a publish-subscribe discovery service dedicated to broadcasting discovery requests. Since you are substituting TCP for UDP, we call this approach *TCP-based discovery*. TCP-based discovery maintains the overall communication pattern of broadcast, discovery and invocation making the programming model identical to WCF discovery. Furthermore, the TCP-based discovery service negates the need for a centralized metadata repository.

The discovery service requires only a single well-known base address for both clients issuing discovery requests and subscribing services. To simplify deployment and configuration management, all parties can manage this base address either through their configuration file, programmatically with helper methods or by relying on the framework's default address conventions. The rest of this appendix walks through a small framework that supports TCP-based discovery within the intranet, bringing it on par with the built-in support for discovery in WCF, along with a set of helper classes. It also serves as an example for rolling out your own discovery mechanism.

TCP Discovery Contracts

The discovery service provides the `IDiscoverySubscription` endpoint:

```
[ServiceContract]
public interface IDiscoverySubscription
{
    [OperationContract]
    void Subscribe(string address);

    [OperationContract]
    void Unsubscribe(string address);
}
```

Discoverable services use this endpoint to subscribe to events of discovery request by registering the address of their discovery endpoint. Note that unlike the basic discovery proxy pattern shown in [Figure C-4](#), which requires services to register the full details of their business endpoints for storage in a repository, `IDiscoverySubscription` requires that services register only a single piece of information, the address of their discovery endpoint.

The first step in creating a custom discovery mechanism is to define the contracts for discovery request and responses. The `IDiscovery` contract defines the discovery requests as follows:

```
[ServiceContract(CallbackContract = typeof(IDiscoveryCallback))]
public interface IDiscovery
```

```

{
    [OperationContract(IsOneWay = true)]
    void OnDiscoveryRequest(string contractName, string contractNamespace,
                           Uri[] scopesToMatch);
}

```

The discovery endpoint provides the single operation `IDiscovery.OnDiscoveryRequest()` which allows clients to discover service endpoints that support a particular contract, as with regular WCF. Clients can also pass in an optional set of scopes to match.

The discovery service also supports the `IDiscovery` endpoint, acting as a meta-subscriber. Note that just as with UDP discovery, the discovery information itself is provided via a callback. By matching the lifetime of the request with the callback, clients get the most up-to-date snapshot of discoverable services upon each request.

The discovery service calls back to clients using the `IDiscoveryCallback` contract, defined as:

```

[ServiceContract]
public interface IDiscoveryCallback
{
    [OperationContract(IsOneWay = true)]
    void OnDiscoveryResponse(Uri address, string contractName,
                           string contractNamespace, Uri[] scopes);
}

```

Note that `IDiscovery` defines `IDiscoveryCallback` as its own callback contract, and every discovering client must provide an `IDiscoveryCallback` endpoint. The actual sequence of discovery is that the client fires a discovery request at the discovery service which then publishes that request to all subscribed services that support the discovery endpoint. In response, discoverable services call back to the client with their endpoint addresses through the discovery service.

Using a duplex callback to provide the discovery results allows calling back to clients even on locked-down client machines. The client only needs to keep the `IDiscoveryCallback` endpoint open during a discovery request. [Figure C-5](#) depicts the discovery sequence (compare it to [Figure C-1](#)).

The first step in [Figure C-5](#) is when services subscribe to receive discovery requests. In step 2 a client fires an event of discovery request at the discovery service endpoint supporting `IDiscovery`. The discovery service acts as a publish-subscribe service broadcasting this event at the subscribing discoverable services (step 3 in [Figure C-5](#)). If a service supports the requested contract, it calls back to the discovery service (step 4 in [Figure C-5](#)), which forwards the callback on to the client (step 5 in [Figure C-4](#)). Once the client receives the service endpoint (or endpoints) addresses, it proceeds to call the service as with a regular service call (step 6 in [Figure C-5](#)).

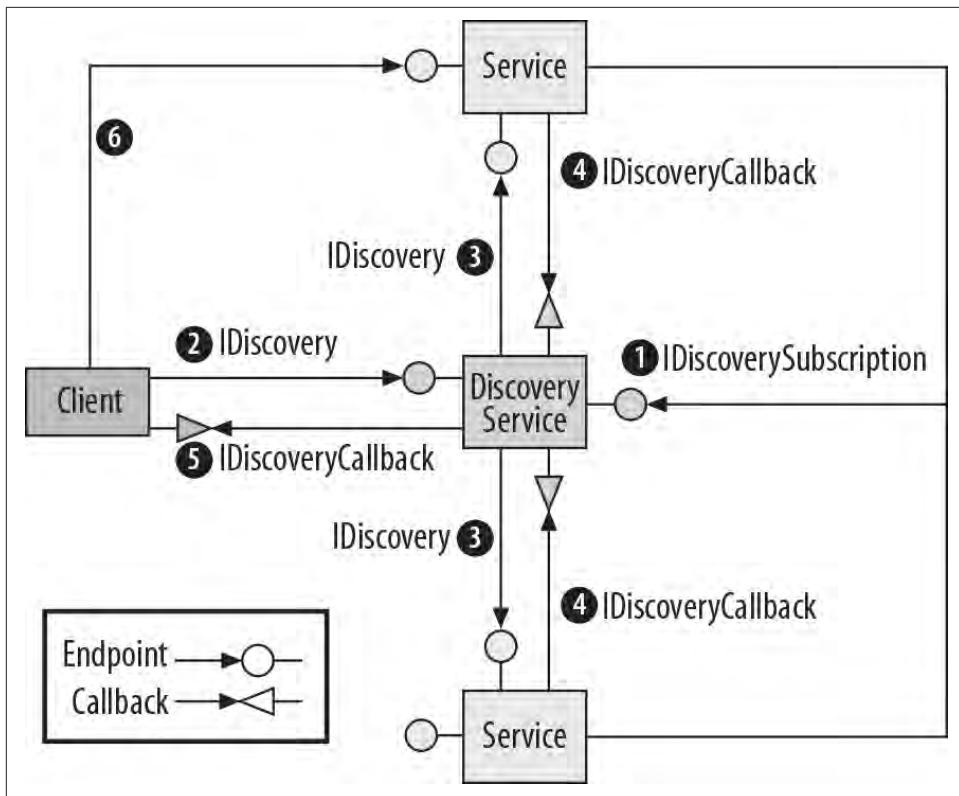


Figure C-5. Discovery via the discovery service

Note that the discovered services do not call the client directly—they call back to the discovery service which has to support the **IDiscoveryCallback** endpoint (as any client of **IDiscovery** endpoint must). The discovery service then forwards the callbacks from the discovered services to the discovering client.

The Discovery Service

There is obviously a lot of work involved in supporting such a discovery mechanism especially at the discovery service. The discovery service must maintain the subscriptions of discoverable services, facilitate the efficient broadcast of discovery requests and manage callbacks to all clients during the discovery process. The TCP-based framework within *ServiceModelEx* provides a ready-made implementation of the discovery service, defined as:

```
[ServiceBehavior(InstanceContextMode = InstanceContextMode.Single,
    UseSynchronizationContext = false,
    ConcurrencyMode = ConcurrencyMode.Multiple)]
```

```
public class DiscoveryService : IDiscovery, IDiscoverySubscription, ...
{...}
```

Subscription management

To manage service subscriptions, `DiscoveryService` maintains a list of discovery endpoint addresses for all subscribing services. As services subscribe and unsubscribe to discovery requests, `DiscoveryService` adds or removes them from an internal list it maintains, as shown in [Example C-19](#).

Example C-19. Discovery service subscription management

```
[ServiceBehavior(...)]
public class DiscoveryService : IDiscoverySubscription, ...
{
    List<string> m_Addresses;

    public DiscoveryService()
    {
        m_Addresses = new List<string>();
    }

    [MethodImpl(MethodImplOptions.Synchronized)]
    void IDiscoverySubscription.Subscribe(string address)
    {
        if(m_Addresses.Contains(address) == false)
        {
            m_Addresses.Add(address);
        }
    }

    [MethodImpl(MethodImplOptions.Synchronized)]
    void IDiscoverySubscription.Unsubscribe(string address)
    {
        if(m_Addresses.Contains(address) == true)
        {
            m_Addresses.Remove(address);
        }
    }
    //Rest of the implementation
}
```

Callback implementation

When clients use `IDiscovery` they must provide a callback reference to receive discovery responses by using a duplex channel. This allows `DiscoveryService` to capture the discovering client's callback reference. Since callback references are not serializable, `DiscoveryService` cannot pass a client's callback reference to subscribing services so that they may call back the client directly. Instead, `DiscoveryService` creates a local callback context in which to maintain the callback to the client and to

control the nature of its concurrency. To achieve this, `DiscoveryService` uses the private nested class `DiscoveryCallback`. `DiscoveryCallback`'s constructor takes as a parameter the discovering client's callback reference. `DiscoveryCallback` also implements the `IDiscoveryCallback` interface. The implementation of `IDiscoveryCallback.OnDiscoveryResponse()` simply forwards the callback from the discovered service to the discovering client's callback, as shown in [Example C-20](#).

Example C-20. The DiscoveryCallback class

```
[ServiceBehavior(...)]
public class DiscoveryService : ...
{
    [CallbackBehavior(UseSynchronizationContext = false,
                      ConcurrencyMode = ConcurrencyMode.Multiple)]
    class DiscoveryCallback : IDiscoveryCallback
    {
        IDiscoveryCallback m_DiscoveryCallback;

        public DiscoveryCallback(IDiscoveryCallback discoveryCallback)
        {
            m_DiscoveryCallback = discoveryCallback;
        }
        public void OnDiscoveryResponse(Uri address, string contractName,
                                         string contractNamespace, Uri[] scopes)
        {
            m_DiscoveryCallback.OnDiscoveryResponse(address, contractName,
                                                     contractNamespace, scopes);
        }
    }
}
```

Note that `DiscoveryCallback` is configured for concurrent access and no use of any synchronization context. This ensures it processes the callbacks as fast as they come in parallel, so that no two responding services have to wait on each other.

Discovery requests

The actual processing of discovery requests takes place within the implementation of `IDiscovery.OnDiscoveryRequest()` in `DiscoveryService`. To make the broadcast of discovery requests as efficient as possible, `DiscoveryService` publishes requests to all subscribing services by iterating over the list of discovery endpoint addresses and firing each discovery request asynchronously and in parallel. During each iteration, the `discover` delegate receives the address of a service's discovery endpoint through a task scheduled on the .NET thread pool. The `discover` delegate uses the instance of `DiscoveryCallback` to create a duplex channel to the service's discovery endpoint. It then sends the discovery request to the service and returns. [Example C-21](#) shows the

publishing mechanism for discovery within `OnDiscoveryRequest()`. For brevity, the sample omits some synchronization, error handling and initialization.

Example C-21. Discovery service request publishing

```
[ServiceBehavior(...)]
public class DiscoveryService : IDiscovery, ...
{
    List<string> m_Addresses;

    [MethodImpl(MethodImplOptions.Synchronized)]
    void IDiscovery.OnDiscoveryRequest(string contractName, string contractNamespace,
                                         Uri[] scopesToMatch)
    {
        //Callback to the client wanting to discover
        IDiscoveryCallback clientCallback = OperationContext.Current.
            GetCallbackChannel<IDiscoveryCallback>();
        DiscoveryCallback serviceCallback = new DiscoveryCallback(clientCallback);

        Action<string> discover = (address)=>
        {
            IDiscovery serviceProxy =
                DuplexChannelFactory<IDiscovery, IDiscoveryCallback>.CreateChannel(
                    serviceCallback, DiscoveryFactory.Binding,
                    new EndpointAddress(address as string));

            EventHandler cleanup = delegate
            {
                (serviceProxy as ICommunicationObject).Close();
            };

            (clientCallback as IClientChannel).Closed += cleanup;
            (clientCallback as IClientChannel).Faulted += cleanup;

            Task.Delay(DiscoveryFactory.Binding.SendTimeout).
                ContinueWith(_=>cleanup(null, EventArgs.Empty));

            serviceProxy.OnDiscoveryRequest(
                contractName, contractNamespace, scopesToMatch);
        };

        m_Addresses.ForEachAsync(discover);
    }
}
```

Because the channel is duplex, `DiscoveryService` cannot close the channel to a subscribing service until after the discovered service has called back to the discovering client. `DiscoveryService` can only safely close the channel to each discovered service once the discovering client has closed its own channel to `DiscoveryService`, or once the publishing timeout has expired. To associate `DiscoveryService`'s channel man-

agement with the duration of the discovery period as controlled by the client, `discover` subscribes the `cleanup` event handler to both the `Closed` and `Faulted` events of the client callback. If the discovering client for some reason fails to close its proxy to the discovery service, `DiscoveryService` must still close all the channels to discoverable services it created during the discovery process. In the eventuality that these events do not take place (if the client simply keeps the channel open indefinitely), the `discover` delegate creates an asynchronous task-based timeout event using the default discovery binding's `SendTimeout`. Upon the expiration of the timeout, the timeout task's continuation calls `cleanup` and closes the channels to the services, even if the channel to the client is still open.

Concurrency management

`DiscoveryService` is fully concurrent. The design configures `DiscoveryService` for multiple concurrency mode and no synchronization context, so multiple clients can initiate discovery requests, all independent of each other. `DiscoveryService` locks all points of access to the shared discovery endpoint address list. `DiscoveryService` is also fully asynchronous. Discovery request publishing is asynchronous and in parallel as are the request timeouts and channel cleanup. As mentioned already, `DiscoveryCallback` is fully concurrent so responding services are not serialized and `DiscoveryCallback` forwards the callbacks to the discovering clients concurrently.

The Discovery Service Host

To simplify the hosting and deployment of `DiscoveryService`, the discovery framework also includes a dedicated `DiscoveryFactory` helper class. `DiscoveryFactory` provides the `CreateDiscoveryService()` method which configures and initializes `DiscoveryService` in a single line of code, as shown in [Example C-22](#).

Example C-22. Hosting the discovery service

```
public static partial class DiscoveryFactory
{
    public static ServiceHost CreateDiscoveryService(uint port = 0);
}

ServiceHost host = DiscoveryFactory.CreateDiscoveryService(8001);
host.Open();
```

Address management

`DiscoveryFactory.CreateDiscoveryService()` will add endpoints for `IDiscovery` and `IDiscoverySubscription`. The `CreateDiscoveryService()` method will use a config file if present or allow for programmatic calls. In the absence of both, `CreateDiscoveryService()` will default to the use of `localhost` for the discovery server

name and 808 for the port. The endpoint addresses will default to the following conventions:

```
net.tcp://<DiscoveryServer>:<Port>/DiscoveryService/Discovery  
net.tcp://<DiscoveryServer>:<Port>/DiscoveryService/DiscoverySubscription
```

To centrally manage `DiscoveryService` endpoint addresses, `DiscoveryFactory` provides the nested public class `Address`. The `Address` class enforces the framework's configuration policy for all clients and services. It also provides the convention-driven default addresses to all parties. `Address` allows you to override the programmatic default conventions either manually or through configuration file settings. For example, the `DiscoveryService` host needs only to set the port if it differs from the default, as shown in [Example C-22](#). Clients and services need only to set the discovery server name when not collocated with the discovery service. And all parties can override the default address conventions either by providing an explicit endpoint address or referencing an endpoint by name within the configuration file or the default endpoint settings, as shown in [Example C-23](#).

Example C-23. The Address class

```
public static partial class DiscoveryFactory  
{  
    public static class Address  
    {  
        public static string DiscoveryServer;  
        {get;set;}  
  
        internal static uint Port  
        {get;set;}  
  
        static Address()  
        {  
            Port = 808;  
            DiscoveryServer = "localhost";  
        }  
  
        //Discovery endpoint address management  
  
        static EndpointAddress m_DiscoveryServiceAddress;  
  
        public static EndpointAddress DiscoveryService  
        {  
            get  
            {  
                if(m_DiscoveryServiceAddress == null)  
                {  
                    m_DiscoveryServiceAddress = GetDiscoveryServiceFromConfig();  
                    if(m_DiscoveryServiceAddress == null)  
                    {  
                        throw new InvalidOperationException("Discovery service endpoint address not found in configuration.");  
                    }  
                }  
            }  
        }  
    }  
}
```

```

        m_DiscoveryServiceAddress = new EndpointAddress(
            "net.tcp://" + DiscoveryServer + ":" + Port + "/DiscoveryService/Discovery");
    }
}
return m_DiscoveryServiceAddress;
}
}
public static void SetDiscoveryService(EndpointAddress address)
{
    m_DiscoveryServiceAddress = address;
}
public static void SetDiscoveryService(string endpointName = null)
{
    m_DiscoveryServiceAddress = GetDiscoveryServiceFromConfig(endpointName);
}
internal static EndpointAddress GetDiscoveryServiceFromConfig(
    string endpointName = null)
{
    return GetAddressFromClientConfig(typeof(IDiscovery), endpointName);
}

//Similar methods for managing discovery subscription address
static EndpointAddress m_DiscoverySubscriptionAddress;
public static EndpointAddress DiscoverySubscription
{get;}
public static void SetDiscoverySubscriptionService(EndpointAddress address);
public static void SetDiscoverySubscriptionService(
    string endpointName = null);

//Generic helper method
static EndpointAddress GetAddressFromClientConfig(Type type,
    string endpointName = null)
{...}
}
}

```

As shown in [Example C-23](#), `Address` manages the default conventions for both the discovery server name and its port. The details of address management for the discovery service endpoint shown in [Example C-23](#) reveal the general approach `Address` uses to enforce configuration policy for all discovery service addresses. The various `Address.SetDiscoveryService()` overloads allow you to set the discovery endpoint address either explicitly or from the configuration file. To enforce configuration policy, `Address.DiscoveryService()` first checks if the discovery endpoint address is already set. If not, it then checks the configuration file through a call to `GetAddressFromClientConfig()` for a compatible client endpoint either by endpoint name or by the `IDiscovery` contract type. If it finds no compatible client endpoint within the configuration file, `Address.DiscoveryService()` applies the default address convention. `Address` enforces configuration policy for the discovery subscription endpoint address in the same manner.

Creating the host

`DiscoveryFactory.CreateDiscoveryService()` enforces the configuration policy for `DiscoveryService`. For each `DiscoveryService` endpoint, it first checks if someone else has already added the endpoint (either programmatically or using a config file). If it finds no such endpoints, `CreateDiscoveryService()` then adds the endpoint programmatically. It uses the `Address` class to apply the current setting of the address appropriate for the given endpoint. When creating the service host for `DiscoveryService`, `CreateDiscoveryService()` also adds a compatible base address using the default conventions for port and discovery server name. This allows for possible extensibility of `DiscoveryService` while still enforcing configuration policy, as shown in [Example C-24](#).

Example C-24. DiscoveryService discovery configuration policy

```
public static partial class DiscoveryFactory
{
    public static ServiceHost CreateDiscoveryService(uint port = 0)
    {
        if(port > 0)
        {
            Address.Port = port;
        }
        Uri baseAddress = new Uri("net.tcp://"+Address.DiscoveryServer+":"+
            Address.port+"/");

        ServiceHost host = new ServiceHost(typeof(DiscoveryService),baseAddress);

        //Adding IDiscovery endpoint
        if(host.Description.Endpoints.Any(endpoint =>
            endpoint.Contract.ContractType == typeof(IDiscovery)) == false)
        {
            host.AddServiceEndpoint(typeof(IDiscovery),Binding,
                Address.DiscoveryService.Uri.AbsoluteUri);
        }

        //Adding IDiscoverySubscription is similar

        return host;
    }
}
```

`DiscoveryFactory` also provides the internal `Binding` property. The framework uses `DiscoveryFactory.Binding` throughout, enforcing a consistent binding policy for all interactions between clients, services, and the discovery service. `DiscoveryFactory.Binding` defines the best practice usage for `NetTcpBinding` within the intranet environment by declaring transport security and enabling reliability and transaction flow. The binding configuration policy also relaxes the reliable session's `InactivityTi`

`meout` (described in [Chapter 1](#)). While it is not best practice in general to set an infinite timeout, the discovery framework must do so in this case to support service availability announcements, as you will see later. Finally, to ensure fast failure and efficient discovery request broadcast, the binding configuration policy sets `SendTimeout` to a very low value of 5 seconds. Unlike business operations, which take an indeterminate amount of time to complete, infrastructure operations such as discovery should always be very short in nature, and any delay indicates a failure that no amount of additional time will remedy. [Example C-25](#) shows the default binding's configuration policy.

Example C-25. The default NetTcpBinding configuration

```
public static partial class DiscoveryFactory
{
    internal static Binding Binding
    {
        get
        {
            NetTcpBinding binding = new NetTcpBinding(SecurityMode.Transport, true);

            binding.TransactionFlow = true;
            binding.ReliableSession.InactivityTimeout = TimeSpan.MaxValue;
            binding.SendTimeout = TimeSpan.FromSeconds(5);

            return binding;
        }
    }
    ...
}
```

Discoverable Services

To simplify the hosting and deployment of discoverable services, `DiscoveryFactory` provides the `CreateDiscoverableHost<T>()` helper method, as shown in [Example C-26](#). `CreateDiscoverableHost<T>()` returns a host that responds to both TCP-based and regular UDP discovery requests. It also adds a MEX endpoint and optionally adds default endpoints for the service.

Example C-26. The CreateDiscoverableHost<T>() method

```
public static partial class DiscoveryFactory
{
    public static ServiceHost<T> CreateDiscoverableHost<T>(Uri scope = null)
        where T : class
    {...}
    //More methods
}
```

You use `CreateDiscoverableHost<T>()` much like a regular service host:

```
class MyService : ...
{...}

ServiceHost host = DiscoveryFactory.CreateDiscoverableHost<MyService>();
host.Open();
```

If needed, you can also easily override the default configuration policy using the `Address` class before creating the discoverable service host by providing a complete endpoint address:

```
//Optional:
EndpointAddress subscriptionAddress = new EndpointAddress(
    "net.tcp://localhost:8001/DiscoveryService/DiscoverySubscription");
DiscoveryFactory.Address.SetDiscoverySubscriptionService(subscriptionAddress);

ServiceHost host = DiscoveryFactory.CreateDiscoverableHost<MyService>();
host.Open();
```

Or by just setting the discovery server name and leaving the remainder of the default configuration policy unchanged:

```
//Optional:
DiscoveryFactory.Address.DiscoveryServer = "MyDiscoveryServer";

ServiceHost host = DiscoveryFactory.CreateDiscoverableHost<MyService>();
host.Open();
```

Implementing IDiscovery

To provide support for TCP-based discovery, `CreateDiscoverableHost<T>()` uses the private nested class `DiscoveryRequestService` which implements `IDiscovery`. `DiscoveryRequestService`'s constructor stores the service's endpoints for processing during discovery requests. To process a discovery request, `IDiscovery.OnDiscoveryRequest()` first obtains a reference to the discovery request callback provided by `DiscoveryService`. It then iterates through the service's endpoints looking for contracts that match the discovery request. If the contract and scopes match, `OnDiscoveryRequest()` responds by calling back to `DiscoveryService` with the endpoint's address and the original discovery request information, as shown in [Example C-27](#).

Example C-27. The DiscoveryRequestService class

```
public static partial class DiscoveryFactory
{
    [ServiceBehavior(InstanceContextMode = InstanceContextMode.Single,
                    ConcurrencyMode = ConcurrencyMode.Multiple,
                    UseSynchronizationContext = false)]
    class DiscoveryRequestService : IDiscovery
    {
```

```

readonly ServiceEndpoint[] Endpoints;

public DiscoveryRequestService(ServiceEndpoint[] endpoints)
{
    Endpoints = endpoints;
}
public void OnDiscoveryRequest(string contractName, string contractNamespace,
                               Uri[] scopesToMatch)
{
    IDiscoveryCallback callback = OperationContext.Current.
        GetCallbackChannel<IDiscoveryCallback>();
    foreach (ServiceEndpoint endpoint in Endpoints)
    {
        if (endpoint.Contract.Name == contractName &&
            endpoint.Contract.Namespace == contractNamespace)
        {
            Uri[] scopes = DiscoveryHelper.LookupScopes(endpoint);
            if (scopesToMatch != null)
            {
                bool scopesMatched = true;
                foreach (Uri scope in scopesToMatch)
                {
                    if (scopes.Any(uri =>
                        uri.AbsoluteUri == scope.AbsoluteUri) == false)
                    {
                        scopesMatched = false;
                        break;
                    }
                }
                if (scopesMatched == false)
                {
                    continue;
                }
            }
            callback.OnDiscoveryResponse(endpoint.Address.Uri,
                                         contractName, contractNamespace, scopes);
        }
    }
}
//Rest of the implementation
}

```

Creating hosts

To facilitate TCP-based discovery, `CreateDiscoverableHost<T>()` actually creates two service hosts. The first service host is for the discoverable service. Since clients can discover the service, there is no need to use a static address so TCP-based discovery uses a dynamic base address for the service. TCP-based discovery enables regular UDP discovery and adds a MEX endpoint. And if the configuration file does not contain endpoints for the service, it also adds the default endpoints for the service. This

makes the discoverable service completely configuration free. The second host is for `DiscoveryRequestService`. Since the service subscribes with its discovery endpoint address to `DiscoveryService` to receive discovery requests, TCP-based discovery can use dynamic addressing here as well. As shown in [Example C-28](#), `CreateDiscoverableHost<T>()` hosts `DiscoveryRequestService` as a singleton. This enables the use of a constructor that takes the endpoints of the discoverable service. TCP-based discovery associates the lifetime of `DiscoveryRequestService` with the discoverable service host through the `Opened` and `Closing` events of the discoverable service host. These events are also where it conveniently subscribes and unsubscribes `DiscoveryRequestService` to discovery requests, as shown in [Example C-28](#).

Example C-28. Discoverable host creation

```
public static partial class DiscoveryFactory
{
    public static ServiceHost<T> CreateDiscoverableHost<T>(Uri scope = null)
        where T : class
    {
        /* Creates host on dynamic address and with support for
           regular UDP discovery, MEX and optionally default service endpoints */
        ServiceHost<T> host = ...;

        DiscoveryRequestService requestService =
            new DiscoveryRequestService(host.Description.Endpoints);

        ServiceHost requestHost = new ServiceHost(requestService);

        Uri requestServiceAddress = new Uri(DiscoveryHelper.
            AvailableTcpBaseAddress.AbsoluteUri +
            typeof(T) + "/" + "DiscoveryEndpoint");

        requestHost.AddServiceEndpoint(typeof(IDiscovery), DiscoveryFactory.Binding,
            requestServiceAddress);
        host.Opened += delegate
        {
            requestHost.Open();

            IDiscoverySubscription pubsubProxy =
                ChannelFactory<IDiscoverySubscription>.CreateChannel(
                    DiscoveryFactory.Binding, Address.DiscoverySubscription);

            pubsubProxy.Subscribe(requestServiceAddress.AbsoluteUri);
            (pubsubProxy as ICommunicationObject).Close();
        };

        host.Closing += delegate
        {
            IDiscoverySubscription pubsubProxy =
                ChannelFactory<IDiscoverySubscription>.CreateChannel(
```

```

        DiscoveryFactory.Binding,Address.DiscoverySubscription);

    pubsubProxy.Unsubscribe(requestServiceAddress,
                           AbsoluteUri);
    (pubsubProxy as ICommunicationObject).Close();

    requestHost.Close();
}
return host;
}
}

```

Discovery Client

For the client, the TCP-based discovery framework provides the helper class `DuplexDiscoveryClient`. `DuplexDiscoveryClient` fires discovery requests at `DiscoveryService` and receives service response callbacks from it. `DuplexDiscoveryClient` then processes these discovery callbacks internally. `DuplexDiscoveryClient`'s design matches closely that of WCF's `DiscoveryClient` for UDP discovery:

```

public class DuplexDiscoveryClient
{
    public DuplexDiscoveryClient()
    public DuplexDiscoveryClient(string endpointName);
    public DuplexDiscoveryClient(EndpointAddress address);

    public FindResponse Find(FindCriteria criteria);
}

```

Setting discovery service addresses

To publish discovery requests, `DuplexDiscoveryClient` must know the address of `DiscoveryService`'s discovery endpoint. As with the other components of this framework, it follows the same approach to configuration policy management. While clients can still use the `Address` class directly, to simplify the client-side programming model, `DuplexDiscoveryClient` provides constructor overloads that encapsulate the use of the `Address` class:

```

public class DuplexDiscoveryClient
{
    //Uses default address convention
    public DuplexDiscoveryClient()
    {}
    public DuplexDiscoveryClient(string endpointName)
    {
        DiscoveryFactory.Address.SetDiscoveryService(endpointName);
    }
    public DuplexDiscoveryClient(EndpointAddress address)
    {
        DiscoveryFactory.Address.SetDiscoveryService(address);
    }
}

```

```

        }
    //Rest of the implementation
}

```

Handling discovery callbacks

To capture service response callbacks from `DiscoveryService`, `DuplexDiscoveryClient` uses the private nested class `ClientDiscoveryResponseCallback`, which delegates the processing of responses back to `DuplexDiscoveryClient`. Its main purpose is to set the attributes for concurrent response processing:

```

public class DuplexDiscoveryClient
{
    [CallbackBehavior(ConcurrencyMode = ConcurrencyMode.Multiple,
                      UseSynchronizationContext = false)]
    class ClientDiscoveryResponseCallback : IDiscoveryCallback
    {
        readonly string ContractNamespace;
        readonly string ContractName;
        readonly Action<Uri, Uri[]> Action;

        public ClientDiscoveryResponseCallback(string contractName,
                                               string contractNamespace,
                                               Action<Uri, Uri[]> action)
        {
            ContractName = contractName;
            ContractNamespace = contractNamespace;
            Action = action;
        }
        public void OnDiscoveryResponse(Uri address, string contractName,
                                       string contractNamespace, Uri[] scopes)
        {
            //Some filtering then:
            Action(address, scopes);
        }
    }
    //Rest of the implementation
}

```

Issuing discovery requests

As with `DiscoveryClient`, `DuplexDiscoveryClient` uses a `Find()` method to manage the process of discovery and handle service responses. [Example C-29](#) shows the relevant aspects of `Find()`. Once again the example removes for brevity synchronization, error handling and some initialization.

Example C-29. Implementing DuplexDiscoveryClient.Find()

```

public class DuplexDiscoveryClient
{

```

```

public FindResponse Find(FindCriteria criteria)
{
    string contractName = criteria.ContractTypeNames[0].Name;
    string contractNamespace = criteria.ContractTypeNames[0].Namespace;
    FindResponse response = DiscoveryHelper.CreateFindResponse();

    ManualResetEvent handle = new ManualResetEvent(false);

    Action<Uri,Uri[]> addEndpoint = (address,scopes)=>
    {
        EndpointDiscoveryMetadata metadata =
            new EndpointDiscoveryMetadata();
        metadata.Address =
            new EndpointAddress(address);
        foreach(Uri scope in scopes)
        {
            metadata.Scopes.Add(scope);
        }

        response.Endpoints.Add(metadata);

        if(response.Endpoints.Count >=
            criteria.MaxResults)
        {
            handle.Set();
        }
    };
}

//To receive the callbacks from the services
ClientDiscoveryResponseCallback callback =
    new ClientDiscoveryResponseCallback(
        contractName,contractNamespace, addEndpoint);

IDiscovery discoveryServiceProxy =
    DuplexChannelFactory<IDiscovery, IDiscoveryCallback>.
        CreateChannel(callback,DiscoveryFactory.Binding,
        DiscoveryFactory.Address.DiscoveryService);

discoveryServiceProxy.OnDiscoveryRequest(criteria.ContractTypeNames[0].Name,
    criteria.ContractTypeNames[0].Namespace,criteria.Scopes.ToArray());

handle.WaitOne(criteria.Duration);
handle.Close();

(discoveryServiceProxy as ICommunicationObject).Close();

return response;
}
//Rest of the implementation
}

```

`Find()` needs a way of receiving callbacks from the discovered services. To that end, every time a client calls it, `Find()` opens and closes a duplex channel using an instance of `ClientDiscoveryResponseCallback`. The constructor of `ClientDiscoveryResponseCallback` accepts a delegate of the type `Action<Uri, Uri[]>`. For every service response, the implementation of `OnDiscoveryResponse()` invokes that delegate, providing it with the discovered address and scope. The `Find()` method uses the `addEndpoint` delegate to aggregate the responses in an instance of `FindResponse`. Unfortunately, there is no public constructor for `FindResponse`, so `Find()` uses the `CreateFindResponse()` method of `DiscoveryHelper`, which in turn uses reflection to instantiate it. `Find()` also creates a waitable event handle. `addEndpoint` signals this handle once the discovery request's cardinality requirement is met. After calling `OnDiscoveryRequest()`, `Find()` waits for the delegate to signal the handle or for the discovery duration to expire, and then closes the channel to stop processing any pending discovery responses that may be en route. This closing is what the Discovery Service implementation of `IDiscovery.OnDiscoveryRequest()` in [Example C-21](#) is monitoring in the `Closed` event of the `IClientChannel` of the duplex callback.

More client-side discovery helpers

Since `DuplexDiscoveryClient` is polymorphic and functionally identical to WCF's `DiscoveryClient`, it also would benefit from a similar streamlined discovery experience. The TCP-based discovery framework's `DiscoveryFactory` offers methods similar to `ServiceModelEx`'s UDP `DiscoveryFactory`. The create channel helpers use MEX-based discovery and `DuplexDiscoveryClient` internally, as shown in [Example C-30](#).

Example C-30. The `CreateChannel<T>()` helper methods

```
public static partial class DiscoveryFactory
{
    public static T CreateChannel<T>(Uri scope = null) where T : class;
    public static T[] CreateChannels<T>(bool inferBinding = true) where T : class;

    //More methods
}
```

`CreateChannel<T>()` assumes cardinality of one and it uses the metadata endpoint to obtain the service's address and binding used to create the proxy. `CreateChannels<T>()` creates proxies to all discovered services, using all discovered metadata endpoints.

You use the create channel helpers similarly to regular WCF:

```
IMyContract proxy = DiscoveryFactory.CreateChannel<IMyContract>();
proxy.MyMethod();
```

And like all the components in the framework, you can easily override the default configuration policy using the `Address` class before creating the proxy by providing a different endpoint address:

```
//Optional:  
EndpointAddress discoveryServiceAddress = new EndpointAddress(  
    "net.tcp://localhost:8001/DiscoveryService/Discovery");  
DiscoveryFactory.Address.SetDiscoveryService(discoveryServiceAddress);  
  
IMyContract proxy = DiscoveryFactory.CreateChannel<IMyContract>();  
proxy.MyMethod();
```

Or by just setting the discovery server name and leaving the remainder of the default configuration policy unchanged:

```
//Optional:  
DiscoveryFactory.Address.DiscoveryServer = "MyDiscoveryServer";  
  
IMyContract proxy = DiscoveryFactory.CreateChannel<IMyContract>();  
proxy.MyMethod();
```

With the aide of `DiscoveryFactory`, the framework's overall programming model for TCP-based discovery becomes simple and succinct:

```
[ServiceContract]  
interface IMyContract  
{...}  
  
class MyService : IMyContract  
{...}  
  
//Client:  
IMyContract proxy = DiscoveryFactory.CreateChannel<IMyContract>();  
  
//Service:  
ServiceHost host = DiscoveryFactory.CreateDiscoverableHost<MyService>();  
host.Open();  
  
//Discovery Service:  
ServiceHost host = DiscoveryFactory.CreateDiscoveryService();  
host.Open();
```

TCP-Based Announcements

All announcements are based on the broadcast of service availability events. Real-time broadcast of availability events ensures clients can respond to dynamic, runtime changes in the environment without polling a centralized service metadata repository for updates. The runtime broadcast of service endpoint addresses also allows the solution to support the dynamic addressing of service endpoints which further simplifies configuration management.

You can approximate WCF announcements without relying on UDP by substituting the UDP broadcast with a publish-subscribe service dedicated to announcements. You can add that functionality to the discovery service discussed thus far. This allows the clients and services to maintain the overall communication pattern of announcement broadcast and invocation making the programming model identical to WCF announcements. In the case of announcements, services are the active publishers and clients are the passive subscribers. For announcements, the discovery service requires only a single well-known base address for services issuing announcement notifications and subscribing clients. To simplify deployment and configuration management, all parties can manage this base address through their configuration file, programmatically with helper methods, or by relying on the framework's default address conventions.

TCP Announcements Contracts

The first step in creating a custom announcements mechanism is to define the `IAnnouncementsSubscription` contract that allows clients to subscribe to service availability events:

```
[ServiceContract(CallbackContract = typeof(IAnnouncements))]
public interface IAnnouncementsSubscription
{
    [OperationContract]
    void Subscribe();

    [OperationContract]
    void Unsubscribe();
}
```

There is no need for the clients to provide their address because the contract uses a duplex callback to transfer the availability events.

The actual announcements are published via a callback of the type `IAnnouncements` defined as:

```
[ServiceContract]
public interface IAnnouncements
{
    [OperationContract(IsOneWay = true)]
    void OnHello(Uri address, string contractName, string contractNamespace,
                Uri[] scopes);

    [OperationContract(IsOneWay = true)]
    void OnBye(Uri address, string contractName, string contractNamespace,
               Uri[] scopes);
}
```

The clients (the actual subscribers) provide their implementation of `IAnnouncements` to the discovery service. The discovery service in turn supports `IAnnouncements` as

well, thus acting as a meta-subscriber toward the publishing services. The announcing services use the discovery service's announcement endpoint to publish events of their availability. The announcements endpoint provides two operations: `OnHello()` and `OnBye()`. Services announce their availability using the `OnHello()` event and retract their availability using the `OnBye()` event. Each availability event contains the announcing service's address, contract, and the scopes it supports.

Since service availability requests are unsolicited, the announcement callback lifetime should match that of the client. This means that the duplex session must be long running, requiring an infinite timeout, as shown in [Example C-25](#). The use of duplex channels allows calling back to clients even on locked-down client machines. [Figure C-6](#) depicts the announcement sequence (compare it to [Figure C-3](#)).

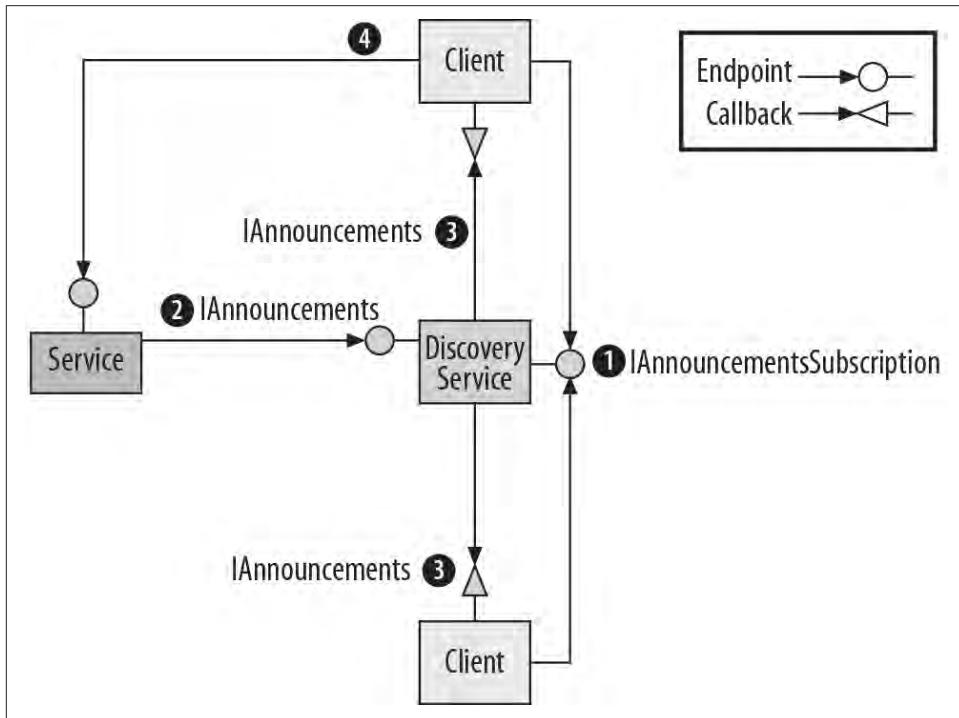


Figure C-6. Announcements via the discovery service

The first step in [Figure C-6](#) is a client subscribing to receive service announcements (step 1 in [Figure C-6](#)). A service then fires an availability event at the discovery service endpoint supporting `IAnnouncements` (step 2 in [Figure C-6](#)). Again, the publish-subscribe mechanism within the discovery service broadcasts the event to all subscribing clients (step 3 in [Figure C-6](#)). Clients filter these events looking for a match to a contract they require and then use the provided service endpoint address to call the service as with a regular service call (step 4 in [Figure C-6](#)).

Discovery Service

Instead of providing a separate publish-subscribe service implementation for announcements that you must deploy, manage, and maintain, it makes sense to add this support directly to your existing discovery service. This would make your discovery service a complete discovery solution on par with WCF's built-in discovery support. As it is the case with discovery, `DiscoverService` encapsulates completely the mechanism of announcements broadcast and subscription management:

```
[ServiceBehavior(...)]
public class DiscoveryService : IAnnouncementsSubscription,IAnnouncements,...
{...}
```

Subscription management

To manage client subscriptions, `DiscoveryService` maintains a list of duplex subscribers. As clients subscribe and unsubscribe to service announcements, `Discovery Service` adds or removes them from the list as shown:

```
[ServiceBehavior(...)]
public class DiscoveryService : IAnnouncementsSubscription, ...
{
    List<IAnnouncements> m_NotifyClients;

    public DiscoveryService()
    {
        m_NotifyClients = new List<IAnnouncements>();
        ...
    }

    [MethodImpl(MethodImplOptions.Synchronized)]
    void IAnnouncementsSubscription.Subscribe()
    {
        IAnnouncements client = OperationContext.Current.GetCallbackChannel
            <IAnnouncements>();

        if(m_NotifyClients.Contains(client) == false)
        {
            m_NotifyClients.Add(client);
        }
    }

    [MethodImpl(MethodImplOptions.Synchronized)]
    void IAnnouncementsSubscription.Unsubscribe()
    {
        IAnnouncements client = OperationContext.Current.GetCallbackChannel
            <IAnnouncements>();

        if(m_NotifyClients.Contains(client))
        {
            m_NotifyClients.Remove(client);
        }
    }
}
```

```

        }
    }
    //Rest of implementation
}

```

Publishing announcements

When clients use `IAnnouncementsSubscription` they must provide a callback reference by using a duplex channel. This allows `DiscoveryService` to capture the subscribing client's callback reference at the point they subscribe and store it for later use.

`DiscoveryService` provides the broadcast of service availability events through the publishing mechanism within `IAnnouncements.OnHello()` and `IAnnouncements.OnBye()`. When announcing services call `OnHello()` or `OnBye()`, `DiscoveryService` iterates over the list of callbacks, forwarding these events on to subscribing clients.

To make the broadcast of service availability events efficient, `DiscoveryService` publishes events to all subscribing clients by iterating in parallel and asynchronously over the list of client callback references. `DiscoveryService` created a delegate of the type `Action<IAnnouncements>` wrapping the appropriate callback to the client (hello or bye). The delegate receives a client callback reference through a task scheduled on the .NET thread pool. It then simply calls the announcements method to publish the event to the client via the callback channel. [Example C-31](#) shows this publishing sequence.

Example C-31. Discovery service announcement publishing

```

[ServiceBehavior(...)]
public class DiscoveryService : IAnnouncements, ...
{
    List<IAnnouncements> m_NotifyClients;

    [MethodImpl(MethodImplOptions.Synchronized)]
    void IAnnouncements.OnHello(Uri address, string contractName,
                                 string contractNamespace, Uri[] scopes)
    {
        Action<IAnnouncements> hello = (client) =>
        {
            client.OnHello(address, contractName,
                           contractNamespace, scopes);
        };
        m_NotifyClients.ForEachAsync(hello);
    }

    [MethodImpl(MethodImplOptions.Synchronized)]
    void IAnnouncements.OnBye(Uri address, string contractName,
                               string contractNamespace, Uri[] scopes)

```

```
{  
    //Similar to OnHello()  
}  
}
```

Concurrency management

As with discovery, `DiscoveryService`'s announcement support is fully concurrent. Clients can subscribe concurrently and services can announce concurrently. `DiscoveryService` locks all points of access to the shared client callback list. `DiscoveryService`'s announcements support is also fully asynchronous. All calls to clients are asynchronous and in parallel. No two announcing services are ever serialized in accessing the discovery service.

The Discovery Service Host

`DiscoveryService`'s announcement support requires no direct interaction other than hosting it, as shown in [Example C-22](#). Hosting code such as [Example C-22](#) will deploy a discovery service that supports discovery and announcements. `DiscoveryFactory.CreateDiscoveryService()` programmatically adds the `IAnnouncementsSubscription` and `IAnnouncements` endpoints using the same default binding, port, and server name configuration policy as with discovery.

Address management

To simplify deployment and management, `DiscoveryFactory`'s helpers support the same configuration policy management and address conventions for announcements as they do for discovery. The announcements endpoint addresses will default to the following conventions:

```
net.tcp://<DiscoveryServer>:<Port>/DiscoveryService/Announcements  
net.tcp://<DiscoveryServer>:<Port>/DiscoveryService/AnnouncementsSubscription
```

`DiscoveryFactory.Address` includes the equivalent convention-driven support for the management of announcement addresses, as shown in [Example C-32](#).

Example C-32. Announcement address conventions

```
public static partial class DiscoveryFactory  
{  
    public static class Address  
    {  
        //Announcements address management  
        public static EndpointAddress Announcements  
        {get;}  
        public static void SetAnnouncements(EndpointAddress address);  
        public static void SetAnnouncements(string endpointName = null);  
    }  
}
```

```

//Announcements subscription address management
public static EndpointAddress AnnouncementsSubscription
{get;}
public static void SetAnnouncementsSubscription(EndpointAddress address);
public static void SetAnnouncementsSubscription(string endpointName = null);
}

```

`DiscoveryFactory.CreateDiscoveryService()` also enforces the configuration policy for `DiscoveryService`'s announcements endpoints, as shown in [Example C-33](#).

Example C-33. Discovery service announcement configuration policy

```

public static partial class DiscoveryFactory
{
    ...
    public static ServiceHost CreateDiscoveryService(uint port = 0)
    {
        ...
        ServiceHost<T> host = ...;

        //Adding IAnnouncements endpoint
        if(host.Description.Endpoints.Any(endpoint =>
            endpoint.Contract.ContractType == typeof(IAnnouncements)) == false)
        {
            host.AddServiceEndpoint(typeof(IAnnouncements),Binding,
                Address.Annotations.Uri.AbsoluteUri);
        }

        //Adding IAnnouncementsSubscription is similar

        return host;
    }
}

```

Announcing Services

The `DiscoveryFactory.CreateDiscoverableHost<T>()` helper method also supports announcements. There is no additional work required when hosting an announcing service. If needed, you can easily override the default configuration policy for the announcements endpoint address by setting the `Address` class before creating the discoverable host:

```

//Optional:
EndpointAddress announcements = new EndpointAddress(
    "net.tcp://localhost:8001/DiscoveryService/Announcements");
DiscoveryFactory.Address.SetAnnouncements(announcements);

ServiceHost host = DiscoveryFactory.CreateDiscoverableHost<MyService>();
host.Open();

```

or by just setting the discovery server name and leaving the remainder of the default configuration policy unchanged.

Firing announcements

To automate TCP-based announcements, the framework associates a service's availability with the lifetime of its service host through the service host's `Opened` and `Closing` events. Within `Opened`, the framework creates a proxy to `DiscoveryService`'s `IAnnouncements` endpoint and fires an `OnHello()` event to announce the availability of each of the service's endpoints. Within `Closing`, it fires the `OnBye()` event to retract each endpoint's availability, as shown in [Example C-34](#).

Example C-34. Publishing availability events

```
public static partial class DiscoveryFactory
{
    public static ServiceHost<T> CreateDiscoverableHost<T>(Uri scope,...)
        where T : class
    {
        ServiceHost<T> host = ...
        ...
        host.Opened += delegate
        {
            ...
            IAnnouncements announcementsProxy =
                ChannelFactory<IAnnouncements>.CreateChannel(
                    DiscoveryFactory.Binding,Address.Annotations);

            PublishAvailabilityEvent(
                host.Description.Endpoints.ToArray(),
                announcementsProxy. OnHello);
        };

        host.Closing += delegate
        {
            ...
            IAnnouncements announcementsProxy =
                ChannelFactory<IAnnouncements>.CreateChannel(
                    DiscoveryFactory.Binding,Address.Annotations);

            PublishAvailabilityEvent(
                host.Description.Endpoints.ToArray(),
                announcementsProxy. OnBye);
        };
        return host;
    }
}
```

The framework uses the `PublishAvailabilityEvent()` helper method to fire endpoint announcements, shown in [Example C-35](#).

Example C-35. The PublishAvailabilityEvent() helper

```
public static partial class DiscoveryFactory
{
    static void PublishAvailabilityEvent(ServiceEndpoint[] endpoints,
                                         Action<Uri, string, string, Uri[]> notification)
    {
        Action<ServiceEndpoint> notify = (endpoint)=>
        {
            Uri[] scopes =
                DiscoveryHelper.LookupScopes(endpoint);
            notification(endpoint.Address.Uri,
                         endpoint.Contract.Name,
                         endpoint.Contract.Namespace, scopes);
        };
        Action publish = ()=>
        {
            endpoints.ParallelForEach(notify);
            (notification.Target as ICommunicationObject).Close();
        };
        Task.Run(publish);
    }
    //Rest of the implementation
}
```

To make the broadcast of availability events efficient, `PublishAvailabilityEvent()` iterates over the service's endpoints in parallel using the `ParallelForEach()` extension method. During each iteration, the `notify` delegate receives a service endpoint on a separate thread from the .NET thread pool. It then looks up any scopes for the endpoint and fires the `notification` event. The `notify` delegate only closes the proxy to `DiscoveryService` once all calls have completed, because `ParallelForEach()` only returns once all iterations are done.

Announcements Client

For the client, the TCP-based framework simplifies the receiving of the service announcements with the helper class `DuplexAnnouncementsSink<T>`. `DuplexAnnouncementsSink<T>` inherits from `AnnouncementSink<T>` to leverage all of its features and provide a programming model on par with the one `ServiceModelEx` provides for UDP-based announcements. `DuplexAnnouncementsSink<T>` implements the `IAnnouncements` callback contract directly. Its configuration allows it to process announcements asynchronously and concurrently as shown:

```
[CallbackBehavior(UseSynchronizationContext = false,
                  ConcurrencyMode = ConcurrencyMode.Multiple)]
```

```

public class DuplexAnnouncementSink<T> : AnnouncementSink<T>, IAnnouncements
    where T : class
{}

```

As with all the components in the TCP-based discovery framework, `DuplexAnnouncementsSink<T>` supports the same configuration policy and address conventions which you can easily override through the `Address` class before creating an instance of `DuplexAnnouncementsSink<T>`. The clients of `DuplexAnnouncementsSink<T>` can subscribe to the delegates of `AnnouncementSink<T>` to receive the announcements, or they can just access the address in the base address container. For example:

```

//Optional:
EndpointAddress announcementsSubscription = new EndpointAddress(
    "net.tcp://localhost:8001/DiscoveryService/AnnouncementsSubscription");
DiscoveryFactory.Address.SetAnnouncementsSubscription(announcementsSubscription);

AnnouncementSink<IMyContract> announcementSink = new DuplexAnnouncementSink
    <IMyContract>();
announcementSink.Open();

EndpointAddress address = announcementSink[0];
Binding binding = ...;

IMyContract proxy = ChannelFactory<IMyContract>.CreateChannel(binding, address);
proxy.MyMethod();
(proxy as ICommunicationObject).Close();

announcementSink.Close();

```

Implementing announcements sink

`DuplexAnnouncementsSink<T>` supports the `IAnnouncements` callback contract through a duplex channel assigning itself as the callback context. `DuplexAnnouncementsSink<T>` overrides the base class's `Open()` method to create the duplex channel and subscribe to availability events. It also overrides the base class's `Close()` method to unsubscribe and close the channel. In between these calls, `DuplexAnnouncementsSink<T>` maintains the channel instance in a local variable, as shown in [Example C-36](#).

Example C-36. The `DuplexAnnouncementSink<T>` class

```

public class DuplexAnnouncementSink<T> : AnnouncementSink<T>, IAnnouncements
    where T : class
{
    IAnnouncementsSubscription m_SubscriptionProxy;

    public override void Open()

```

```

{
    base.Open();

    if(m_SubscriptionProxy == null)
    {
        m_SubscriptionProxy =
            DuplexChannelFactory<IAnnouncementsSubscription,IAnnouncements>.
                CreateChannel( this,DiscoveryFactory.Binding,
                DiscoveryFactory.Address.AnnouncementsSubscription);
        m_SubscriptionProxy.Subscribe();
    }
}
public override void Close()
{
    if((m_SubscriptionProxy as ICommunicationObject).State ==
                                               CommunicationState.Opened)
    {
        m_SubscriptionProxy.Unsubscribe();
        (m_SubscriptionProxy as ICommunicationObject).Close();

        m_SubscriptionProxy = null;
    }

    base.Close();
}

void IAnnouncements.OnHello(Uri address,string contractName,
                           string contractNamespace,Uri[] scopes)
{
    AnnouncementEventArgs args = DiscoveryHelper.CreateAnnouncementArgs(address,
                           contractName,contractNamespace,scopes);
    OnHello(this,args);
}

void IAnnouncements.OnBye(Uri address,string contractName,
                           string contractNamespace,Uri[] scopes)
{
    AnnouncementEventArgs args = DiscoveryHelper.CreateAnnouncementArgs(address,
                           contractName,contractNamespace,scopes);
    OnBye(this,args);
}
}

```

The client then calls `Open()` and `Close()` on an instance of `DuplexAnnouncementsSink<T>` to control its subscription lifetime and control when to receive announcements.

The implementation of the event-handling methods of `IAnnouncements` creates an `AnnouncementEventArgs` instance, populating it with the announced service address, contract, and scopes, and then calls the base class implementation of the respective announcement methods as if they were called using regular WCF discovery. This

both populates the base class of the `AddressesContainer<T>` and fires the appropriate events of `AnnouncementSink<T>`. Note that to create an instance of `AnnouncementEventArgs`, you must use reflection due to the lack of a public constructor.

Again, with the aide of `DiscoveryFactory`, the framework's overall programming model for TCP-based announcements becomes simple and succinct:

```
[ServiceContract]
interface IMyContract
{...}

class MyService : IMyContract
{...}

//Client:
AnnouncementSink<IMyContract> sink = new DuplexAnnouncementSink<IMyContract>();
sink.Open();

IMyContract proxy = ChannelFactory<IMyContract>.CreateChannel(...,sink[0]);

//Service:
ServiceHost host = DiscoveryFactory.CreateDiscoverableHost<MyService>();
host.Open();

//Discovery Service:
ServiceHost host = DiscoveryFactory.CreateDiscoveryService();
host.Open();
```

APPENDIX D

Publish-Subscribe Service

Using raw duplex callbacks for events has numerous drawbacks. The primary concern is that it often introduces too much coupling between the publisher and the subscribers. The subscriber has to know where all the publishing services are in the application and connect to them. Any publisher that the subscriber is unaware of will not be able to notify the subscriber of events. This in turn makes adding new subscribers (or removing existing ones) difficult in an already deployed application. There is no way for a subscriber to ask to be notified whenever anyone in the application raises a particular type of event. In addition, the subscriber must make multiple, potentially expensive calls to each publisher, both to subscribe and to unsubscribe. Different publishers may fire the same event but offer slightly different ways to subscribe and unsubscribe, which, of course, couples the subscribers to those methods.

Much the same way, the publisher can only send notifications to subscribers it knows about. There is no way for the publisher to deliver an event to whomever wishes to receive it, or to broadcast an event. In addition, all publishers must have the necessary code to manage the list of subscribers and the publishing act itself. This code has almost nothing to do with the business problem the service is designed to solve, and it can get fairly complex if advanced features, such as concurrent publishing, are employed. Since the publishers cannot assume the subscribers are all disciplined (i.e., do not take long to process the event), they must publish concurrently on multiple threads. To avoid maxing out the system when there are many subscribers, the publisher needs to multiplex the events on threads from a thread pool. Such publishing logic is not trivial, yet all publishers will have to have it.

Furthermore, duplex-based callbacks introduce coupling between the lifetime of the publisher and the subscribers. The subscribers have to be up and running in order to subscribe to and receive events.

There is no way for a subscriber to ask that if an event is fired, the application create an instance of the subscriber and let it handle the event.

Since proxy references are not serializable, they must reside in memory in some form of list. If the publisher's process (or host machine) crashes, the subscriptions will be lost, yet none of the subscribers will be aware of it.

Security represents yet another dimension of coupling: the subscribers need to be able to authenticate themselves against all publishers, across all security modes and credentials used. The publisher must also have sufficient security credentials to fire the event, and different subscribers may have different role membership mechanisms.

Finally, you must set up subscriptions programmatically. There is no easy administrative way to configure subscriptions in the application or to administratively change the subscriber's preferences when the system is running.

These problems are not actually specific to WCF duplex calls. They also characterize past technologies, such as COM connection points and .NET delegates—all are tightly coupled event-management mechanisms that rely on object references.

The Publish-Subscribe Design Pattern

The solution to the problems just described is to design around them using what is known as the *publish-subscribe* design pattern. The idea behind the pattern is a simple one: decouple the publishers from the subscribers by introducing a dedicated subscription service and a dedicated publishing service in between, as shown in [Figure D-1](#).

Figure D-1. A publish-subscribe system

Subscribers that want to subscribe to events register with the subscription service, which manages the lists of subscribers (and provides a similar ability to unsubscribe). Similarly, all publishers use the publisher service to fire their events and avoid delivering the events directly to the subscribers. The subscription and publishing services provide a layer of indirection that decouples your system. No longer do the subscribers have any knowledge about the identity of the publishers—they can subscribe to a type of event and will receive that event from any publisher. The subscription mechanism is uniform across all publishers. In fact, no publisher has to manage any subscription list, and the publishers have no idea who the subscribers are. They simply deliver the events to the publishing service to be delivered to any interested subscriber.

Subscriber Types

You can even define two types of subscribers: *transient subscribers* are in memory running subscribers and *persistent subscribers* are subscribers that persist on the disk, representing services to invoke when the event takes place. For transient subscribers, you can use the duplex callback mechanism as a handy way of passing the callback reference to the running service. For the persistent subscribers, all you really need to record is the subscriber address as a reference. When the event is raised, the publishing service will call to the persistent subscriber address and deliver the event to it. Another important distinction between the two types of subscriptions is that you can store a persistent subscription on the disk or in a database. Doing so will persist the subscription across application shutdowns or machine crashes and reboots, thus enabling administrative configuration: the subscriber is persistent and the subscription is persistent. Obviously, you cannot save transient subscriptions across an application shutdown, and you will need to set them up explicitly every time the application starts: the subscriber is transient, and so is the subscription.

The Publish-Subscribe Framework

ServiceModelEx contains a simple-to-use, industrial-strength publish-subscribe framework. I wanted to provide not just a publish-subscribe service, but also a general-purpose framework that automates implementing such services and adding the support for any application in just one line of code (if that). The first step in building the framework was to factor the publish-subscribe management interfaces and provide separate contracts for transient and persistent subscriptions and for publishing.¹

Managing Transient Subscriptions

For managing transient subscriptions, I defined the `ISubscriptionService` interface shown in [Example D-1](#).

Example D-1. The `ISubscriptionService` interface manages transient subscribers

```
[ServiceContract]
public interface ISubscriptionService
{
    [OperationContract]
    void Subscribe(string eventOperation);

    [OperationContract]
```

¹ I first wrote about my publish-subscribe framework in my article “WCF Essentials: What You Need to Know About One-Way Calls, Callbacks, and Events” (*MSDN Magazine*, October 2006).

```
    void Unsubscribe(string eventOperation);
}
```

Note that `ISubscriptionService` does not identify the callback contract its implementing endpoint expects. Being a general-purpose interface, it is unaware of particular callback contracts. It is up to the using application to define those callback contracts. The callback interface is provided in the using application by deriving from `ISubscriptionService` and specifying the desired callback contract:

```
public interface IMyEvents
{
    [OperationContract(IsOneWay = true)]
    void OnEvent1();

    [OperationContract(IsOneWay = true)]
    void OnEvent2(int number);

    [OperationContract(IsOneWay = true)]
    void OnEvent3(int number, string text);
}

[ServiceContract(CallbackContract = typeof(IMyEvents))]
public interface IMySubscriptionService : ISubscriptionService
{}
```

Typically, every operation on the callback contract corresponds to a specific event. The subinterface of `ISubscriptionService` (`IMySubscriptionService`, in this example) does not need to add operations. `ISubscriptionService` provides the transient subscription management functionality. In each call to `Subscribe()` or `Unsubscribe()`, the subscriber needs to provide the name of the operation (and therefore the event) it wants to subscribe to or unsubscribe from. If the caller wants to subscribe to all events, it can pass an empty or `null` string.

My framework offers an implementation for the methods of `ISubscriptionService` in the form of the generic abstract class `SubscriptionManager<T>`:

```
public abstract class SubscriptionManager<T> where T : class
{
    public void Subscribe(string eventOperation);
    public void Unsubscribe(string eventOperation);
    //More members
}
```

The generic type parameter for `SubscriptionManager<T>` is the events contract. Note that `SubscriptionManager<T>` does not implement `ISubscriptionService`.

The using application needs to expose its own transient subscription service in the form of an endpoint that supports its specific subinterface of `ISubscriptionService`. To do so, the application needs to provide a service class that derives from `Subscrip`

`tionManager<T>`, specify the callback contract as a type parameter, and derive from that specific subinterface of `ISubscriptionService`. For example, to implement a transient subscription service using the `IMyEvents` callback interface:

```
[ServiceBehavior(InstanceContextMode = InstanceContextMode.PerCall)]
class MySubscriptionService : SubscriptionManager<IMyEvents>,
    IMySubscriptionService
{}
```

`MySubscriptionService` doesn't need any code because `IMySubscriptionService` does not add any new operations and `SubscriptionManager<T>` already implements the methods of `ISubscriptionService`.

Note that just deriving from `SubscriptionManager<IMyEvents>` is insufficient because it does not derive from a contract interface—you must add the derivation from `IMySubscriptionService` to support transient subscriptions.

Finally, the using application needs to define an endpoint for `IMySubscriptionService`:

```
<services>
    <service name = "MySubscriptionService">
        <endpoint
            address  = "..."
            binding  = "..."
            contract = "IMySubscriptionService"
        />
    </service>
</services>
```

Example D-2 shows how `SubscriptionManager<T>` manages transient subscriptions.

Example D-2. The transient subscribers management in `SubscriptionManager<T>`

```
public abstract class SubscriptionManager<T> where T : class
{
    static Dictionary<string,List<T>> m_TransientStore;

    static SubscriptionManager()
    {
        m_TransientStore = new Dictionary<string,List<T>>();
        string[] methods = GetOperations();
        Action<string> insert = (methodName)=>
        {
            m_TransientStore.Add(methodName,new List<T>());
        };
        methods.ForEach(insert);
    }
    static string[] GetOperations()
    {
        MethodInfo[] methods = typeof(T).GetMethods(BindingFlags.Public |
```

```

                BindingFlags.FlattenHierarchy |
                BindingFlags.Instance);
List<string> operations = new List<string>(methods.Length);

Action<MethodInfo> add = (method)=>
{
    Debug.Assert(!operations.Contains(method.Name));
    operations.Add(method.Name);
};

methods.ForEach(add);
return operations.ToArray();
}
static void AddTransient(T subscriber,string eventOperation)
{
    lock(typeof(SubscriptionManager<T>))
    {
        List<T> list = m_TransientStore[eventOperation];
        if(list.Contains(subscriber))
        {
            return;
        }
        list.Add(subscriber);
    }
}
static void RemoveTransient(T subscriber ,string eventOperation)
{
    lock(typeof(SubscriptionManager<T>))
    {
        List<T> list = m_TransientStore[eventOperation];
        list.Remove(subscriber);
    }
}

public void Subscribe(string eventOperation)
{
    lock(typeof(SubscriptionManager<T>))
    {
        T subscriber = OperationContext.Current.GetCallbackChannel<T>();
        if(String.IsNullOrEmpty(eventOperation) == false)
        {
            AddTransient(subscriber,eventOperation);
        }
        else
        {
            string[] methods = GetOperations();
            Action<string> addTransient = (methodName)=>
            {
                AddTransient(subscriber,methodName);
            };
            methods.ForEach(addTransient);
        }
    }
}

```

```

}

public void Unsubscribe(string eventOperation)
{
    lock(typeof(SubscriptionManager<T>))
    {
        T subscriber = OperationContext.Current.GetCallbackChannel<T>();
        if(String.IsNullOrEmpty(eventOperation) == false)
        {
            RemoveTransient(subscriber, eventOperation);
        }
        else
        {
            string[] methods = GetOperations();
            Action<string> removeTransient = (methodName)=>
            {
                RemoveTransient(subscriber,methodName);
            };
            methods.ForEach(removeTransient);
        }
    }
}
//More members
}

```

`SubscriptionManager<T>` stores the transient subscribers in a generic static dictionary called `m_TransientStore`:

```
static Dictionary<string,List<T>> m_TransientStore;
```

Each entry in the dictionary contains the name of the event operation and all its subscribers in the form of a linked list. The static constructor of `SubscriptionManager<T>` uses reflection to get all the operations of the callback interfaces (the type parameter for `SubscriptionManager<T>`) and initializes the dictionary to have all the operations with empty lists. The `Subscribe()` method extracts the callback reference from the operation call context. If the caller specifies an operation name, `Subscribe()` calls the helper method `AddTransient()`. `AddTransient()` retrieves the list of subscribers for the event from the store, and if the list does not contain the subscriber, it adds it in.

If the caller specifies an empty string or `null` for the operation name, `Subscribe()` calls `AddTransient()` for each operation in the callback contract.

`Unsubscribe()` operates in a similar manner. Note that the caller can subscribe to all events and then unsubscribe from a particular one.

Managing Persistent Subscribers

For managing persistent subscribers, I defined the `IPersistentSubscriptionService` interface shown in [Example D-3](#).

Example D-3. The `IPersistentSubscriptionService` interface manages persistent subscribers

```
[ServiceContract]
public interface IPersistentSubscriptionService
{
    [OperationContract]
    [TransactionFlow(TransactionFlowOption.Allowed)]
    void Subscribe(string address, string eventsContract, string eventOperation);

    [OperationContract]
    [TransactionFlow(TransactionFlowOption.Allowed)]
    void Unsubscribe(string address, string eventsContract, string eventOperation);
    //More members
}
```

To add a persistent subscriber, the caller needs to call `Subscribe()`, providing the address of the subscriber, the event's contract name, and the specific event operation itself. To unsubscribe, the caller calls `Unsubscribe()` with the same information. Note that `IPersistentSubscriptionService` does not imply where the subscribers persist on the service side—that is an implementation detail.

The class `SubscriptionManager<T>`, presented previously, also implements the methods of `IPersistentSubscriptionService`:

```
[BindingRequirement(TransactionFlowEnabled = true)]
public abstract class SubscriptionManager<T> where T : class
{
    public void Unsubscribe(string address, string eventsContract,
                           string eventOperation);
    public void Subscribe(string address, string eventsContract,
                         string eventOperation);
    //More members
}
```

`SubscriptionManager<T>` stores the persistent subscribers in SQL Server. It is configured to use the Client/Service transaction mode (presented in [Chapter 7](#)), and it enforces that mode using my `BindingRequirement` attribute.



ServiceModelEx also provides support for file and in-memory persistent subscriber storage as well.

The generic type parameter for `SubscriptionManager<T>` is the events contract. Note that `SubscriptionManager<T>` does not derive from `IPersistentSubscriptionService`. The using application needs to expose its own persistent subscription service, but there is no need to derive a new contract from `IPersistentSubscriptionService`, because no callback references are required. The application simply derives from `SubscriptionManager<T>`, specifying the events contract as a type parameter and adding a derivation from `IPersistentSubscriptionService`, for example:

```
[ServiceBehavior(InstanceContextMode = InstanceContextMode.PerCall)]
class MySubscriptionService : SubscriptionManager<IMyEvents>,
    IPersistentSubscriptionService
{}
```

There's no need for any code in `MySubscriptionService`, because `SubscriptionManager<T>` already implements the methods of `IPersistentSubscriptionService`.

Note that just deriving from `SubscriptionManager<IMyEvents>` is insufficient, because `SubscriptionManager<IMyEvents>` does not derive from a contract interface —you must add the derivation from `IPersistentSubscriptionService` to support persistent subscriptions.

Finally, the application needs to define an endpoint for `IPersistentSubscriptionService`:

```
<services>
    <service name = "MySubscriptionService">
        <endpoint
            address = "..."
            binding = "..."
            contract = "ServiceModelEx.IPersistentSubscriptionService"
        />
    </service>
</services>
```

The implementation of the methods of `IPersistentSubscriptionService` by `SubscriptionManager<T>` is shown in [Example D-4](#). [Example D-4](#) is very similar to [Example D-2](#), except the subscribers are stored in SQL Server, not in memory in a dictionary.

Example D-4. Persistent subscriber management in `SubscriptionManager<T>`

```
public abstract class SubscriptionManager<T> where T : class
{
    static void AddPersistent(string address, string eventsContract,
                                string eventOperation)
    {
        //Store the subscription in SQL Server
    }
}
```

```

static void RemovePersistent(string address, string eventsContract,
                            string eventOperation)
{
    //Remove the subscription from SQL Server
}

[OperationBehavior(TransactionScopeRequired = true)]
public void Subscribe(string address, string eventsContract,
                      string eventOperation)
{
    if(String.IsNullOrEmpty(eventOperation) == false)
    {
        AddPersistent(address,eventsContract,eventOperation);
    }
    else
    {
        string[] methods = GetOperations();
        Action<string> addPersistent = (methodName)=>
        {
            AddPersistent(address,eventsContract,methodName);
        };
        methods.ForEach(addPersistent);
    }
}

[OperationBehavior(TransactionScopeRequired = true)]
public void Unsubscribe(string address, string eventsContract,
                        string eventOperation)
{
    if(String.IsNullOrEmpty(eventOperation) == false)
    {
        RemovePersistent(address,eventsContract,eventOperation);
    }
    else
    {
        string[] methods = GetOperations();
        Action<string> removePersistent = (methodName)=>
        {
            RemovePersistent(address,eventsContract,methodName);
        };
        methods.ForEach(removePersistent);
    }
}
//More members
}

```

If you want the application to support both transient and persistent subscribers for the same events contract, simply derive the subscription service class from both the specialized subinterface of `ISubscriptionService` and `IPersistentSubscriptionService`:

```
[ServiceBehavior(InstanceContextMode = InstanceContextMode.PerCall)]
class MySubscriptionService : SubscriptionManager<IMyEvents>,
    IMySubscriptionService,
    IPersistentSubscriptionService
{}
```

Next, expose the two matching endpoints:

```
<services>
    <service name = "MySubscriptionService">
        <endpoint
            address = "..."
            binding = "..."
            contract = "IMySubscriptionService"
        />
        <endpoint
            address = "..."
            binding = "..."
            contract = "ServiceModelEx.IPersistentSubscriptionService"
        />
    </service>
</services>
```

Event Publishing

The parts of the publish-subscribe framework shown so far deal only with the aspects of subscription management. The framework also enables easy implementation of the publishing service. The publishing service should support the same events contract as the subscribers, and it should be the only point of contact known to the publishers in the application. Because the publishing service exposes the events contract in an endpoint, you need to mark the events contract as a service contract, even if you only use it for duplex callbacks with transient subscribers:

```
[ServiceContract]
public interface IMyEvents
{
    [OperationContract(IsOneWay = true)]
    void OnEvent1();

    [OperationContract(IsOneWay = true)]
    void OnEvent2(int number);

    [OperationContract(IsOneWay = true)]
    void OnEvent3(int number, string text);
}
```

The publish-subscribe framework contains the helper class `PublishService<T>`, defined as:

```
public abstract class PublishService<T> where T : class
{
```

```
        protected static void FireEvent(params object[] args);
    }
```

PublishService<T> requires as a type parameter the type of the events contract. To provide your own publishing service, derive from PublishService<T> and use the FireEvent() method to deliver the event to all subscribers, be they transient or persistent, as shown in [Example D-5](#).

Example D-5. Implementing an event-publishing service

```
[ServiceBehavior(InstanceContextMode = InstanceContextMode.PerCall)]
class MyPublishService : PublishService<IMyEvents>, IMyEvents
{
    public void OnEvent1()
    {
        FireEvent();
    }
    public void OnEvent2(int number)
    {
        FireEvent(number);
    }
    public void OnEvent3(int number, string text)
    {
        FireEvent(number, text);
    }
}
```

Note that you can use FireEvent() to fire any type of event, regardless of the parameters, because of the use of the params object array.

Finally, the application needs to expose an endpoint for the publishing service with the events contract:

```
<services>
    <service name = "MyPublishService">
        <endpoint
            address = "..."
            binding = "..."
            contract = "IMyEvents"
        />
    </service>
</services>
```

[Example D-6](#) shows the implementation of PublishService<T>.

Example D-6. Implementing PublishService<T>

```
public abstract class PublishService<T> where T : class
{
    protected static void FireEvent(params object[] args)
    {
```

```

        string action = OperationContext.Current.IncomingMessageHeaders.Action;
        string[] slashes = action.Split('/');
        string methodName = slashes[slashes.Length-1];

        FireEvent(methodName,args);
    }
    static void FireEvent(string methodName,object[] args)
    {
        PublishPersistent(methodName,args);
        PublishTransient(methodName,args);
    }
    static void PublishPersistent(string methodName,object[] args)
    {
        T[] subscribers = SubscriptionManager<T>.GetPersistentList(methodName);
        Publish(subscribers,true,methodName,args);
    }
    static void PublishTransient(string methodName,object[] args)
    {
        T[] subscribers = SubscriptionManager<T>.GetTransientList(methodName);
        Publish(subscribers,false,methodName,args);
    }
    static void Publish(T[] subscribers,bool closeSubscribers,string methodName,
                        object[] args)
    {
        Action<T> fire = (subscriber)=>
        {
            Invoke(subscriber as T,methodName,args);
            if(closeSubscribers)
            {
                using(subscriber as IDisposable)
                {}
            }
        };
        subscribers.ForEachAsync(fire);
    }
    static void Invoke(T subscriber,string methodName,object[] args)
    {
        Debug.Assert(subscriber != null);
        Type type = typeof(T);
        MethodInfo methodInfo = type.GetMethod(methodName);
        try
        {
            methodInfo.Invoke(subscriber,args);
        }
        catch
        {}
    }
}

```

To simplify firing the event, the `FireEvent()` method accepts the parameters to pass to the subscribers, yet its caller does not provide it with the name of the operation to invoke on the subscribers. Instead, `FireEvent()` extracts the method name from the

incoming message headers. It then uses an overloaded `FireEvent()` that accepts the method name. That method, in turn, uses the helper method `PublishPersistent()` to publish to all persistent subscribers and uses the `PublishTransient()` helper method to publish to all transient subscribers. The publishing methods operate in an almost identical way: they access `SubscriptionManager<T>` to retrieve their respective subscribers list, then use the `Publish()` method to fire the event. The subscribers are returned in the form of an array of proxies to the subscribers, which is passed to the `Publish()` method.

`Publish()` could have simply invoked the subscribers at this point. However, I wanted to support concurrent publishing of events so that if any subscriber is undisciplined and takes a long time to process the event, this will not preclude the other subscribers from receiving the event in a timely manner. Having the event operations marked as one way is no guarantee of asynchronous invocation, and besides, I wanted to support concurrent publishing even when the event operation is not marked as a one-way operation. `Publish()` defines the `fire` anonymous method. `fire` calls the `Invoke()` helper method, which fires the event to the individual subscriber provided and then closes the proxy if so specified. Because `Invoke()` was never compiled against the specific subscriber type, it uses reflection and late binding for the invocation. `Invoke()` also suppresses any exceptions raised by the invocation, because these are of no interest to the publishing party. `Publish()` then uses the `ForEachA sync<T>()` helper extension method that iterates over the subscribers array and for each subscriber invokes `fire` asynchronously and in parallel.

Notice how uniformly `PublishService<T>` treats the subscribers—it almost does not matter if they are transient or persistent. The only difference is that after publishing to a persistent subscriber, you need to close the proxy. You can achieve this uniformity using the helper methods `GetTransientList()` and `GetPersistentList()` of `SubscriptionManager<T>`. Of these two, `GetTransientList()` is the simpler one:

```
public abstract class SubscriptionManager<T> where T : class
{
    internal static T[] GetTransientList(string eventOperation)
    {
        lock(typeof(SubscriptionManager<T>))
        {
            if(m_TransientStore.ContainsKey(eventOperation))
            {
                List<T> list = m_TransientStore[eventOperation];
                return list.ToArray();
            }
            return new T[]{};
        }
    }
    //More members
}
```

`GetTransientList()` looks up all the subscribers to the specified operation in the transient store and returns them as an array. `GetPersistentList()` faces a bigger challenge: there is no ready-made list of proxies to persistent subscribers. The only thing known about them is their addresses. `GetPersistentList()` therefore needs to instantiate the persistent subscribers' proxies, as shown in [Example D-7](#).

Example D-7. Creating the persistent subscribers proxy list

```
public abstract class SubscriptionManager<T> where T : class
{
    internal static T[] GetPersistentList(string eventOperation)
    {
        string[] addresses = GetSubscribersToContractEventOperation(
            typeof(T).ToString(), eventOperation);

        List<T> subscribers = new List<T>(addresses.Length);

        foreach(string address in addresses)
        {
            Binding binding = GetBindingFromAddress(address);
            T proxy = ChannelFactory<T>.CreateChannel(binding,
                new EndpointAddress(address));
            subscribers.Add(proxy);
        }
        return subscribers.ToArray();
    }

    static string[] GetSubscribersToContractEventOperation(string eventsContract,
        string eventOperation)
    {
        //Query SQL Server for the subscribers to the event
    }

    static Binding GetBindingFromAddress(string address)
    {
        if(address.StartsWith("http:") || address.StartsWith("https:"))
        {
            WSHttpBinding binding = new WSHttpBinding();
            binding.ReliableSession.Enabled = true;
            binding.TransactionFlow = true;
            return binding;
        }
        if(address.StartsWith("net.tcp:"))
        {
            NetTcpBinding binding = new NetTcpBinding();
            binding.ReliableSession.Enabled = true;
            binding.TransactionFlow = true;
            return binding;
        }
        /* Similar code for the one-way relay, IPC and MSMQ bindings */
        Debug.Assert(false, "Unsupported binding specified");
        return null;
    }
}
```

```

    }
    //More members
}

```

To create a proxy for each subscriber, `GetPersistentList()` needs the subscriber's address, binding, and contract. The contract is, of course, the type parameter for `SubscriptionManager<T>`. To obtain the addresses, `GetPersistentList()` calls `GetSubscribersToContractEventOperation()`, which queries the subscribers store and returns as an array the addresses of all of the persistent subscribers who have subscribed to the specified event. All `GetPersistentList()` needs now is the binding used by each subscriber. For that, `GetPersistentList()` calls the helper method `GetBindingFromAddress()`, which infers the binding to use from the address schema. `GetBindingFromAddress()` assumes all HTTP or HTTPS addresses indicate the use the `WSHttpBinding`.

In addition, when applicable, `GetBindingFromAddress()` turns on reliability and transaction propagation for each binding to enable inclusion of the event in the publisher's transaction when one-way operations are not used, such as with this events contract:

```

[ServiceContract]
interface IMyEvents
{
    [OperationContract]
    [TransactionFlow(TransactionFlowOption.Allowed)]
    void OnEvent1();

    [OperationContract]
    [TransactionFlow(TransactionFlowOption.Allowed)]
    void OnEvent2(int number);    [OperationContract]
    [TransactionFlow(TransactionFlowOption.Allowed)]
    void OnEvent3(int number, string text);
}

```

Administering Persistent Subscribers

While you can add and remove persistent subscriptions at runtime by using the methods of the `IPersistentSubscriptionService` interface shown in [Example D-3](#), because of their persistent nature, you can best manage the subscriptions with some kind of administration tool. To that end, `IPersistentSubscriptionService` defines additional operations that answer various queries against the subscribers store, as shown in [Example D-8](#).

Example D-8. The IPersistentSubscriptionService interface

```

[DataContract]
public struct PersistentSubscription

```

```

{
    [DataMember]
    public string Address
    {get;set;}

    [DataMember]
    public string EventsContract
    {get;set;}

    [DataMember]
    public string EventOperation
    {get;set;}
}

[ServiceContract]
public interface IPersistentSubscriptionService
{
    //Administration operations
    [OperationContract]
    [TransactionFlow(TransactionFlowOption.Allowed)]
    PersistentSubscription[] GetAllSubscribers();

    [OperationContract]
    [TransactionFlow(TransactionFlowOption.Allowed)]
    PersistentSubscription[] GetSubscribersToContract(string eventsContract);

    [OperationContract]
    [TransactionFlow(TransactionFlowOption.Allowed)]
    string[] GetSubscribersToContractEventType(string eventsContract,
                                                string eventOperation);

    [OperationContract]
    [TransactionFlow(TransactionFlowOption.Allowed)]
    PersistentSubscription[] GetAllSubscribersFromAddress(string address);
    //More members
}

```

All of these administration operations utilize a simple data contract called `PersistentSubscription`, which contains the address of the subscriber, the subscribed contract, and the event. `GetAllSubscribers()` simply returns the list of all subscribers. `GetSubscribersToContract()` returns all subscribers to a specific contract, and `GetSubscribersToContractEventType()` returns all subscribers to a particular event operation on a specified contract. Finally, for the sake of completeness, `GetAllSubscribersFromAddress()` returns all subscribers that provided a specified address.

My publish-subscribe framework includes a sample persistent subscription administration tool called Persistent Subscription Manager, shown in [Figure D-2](#).

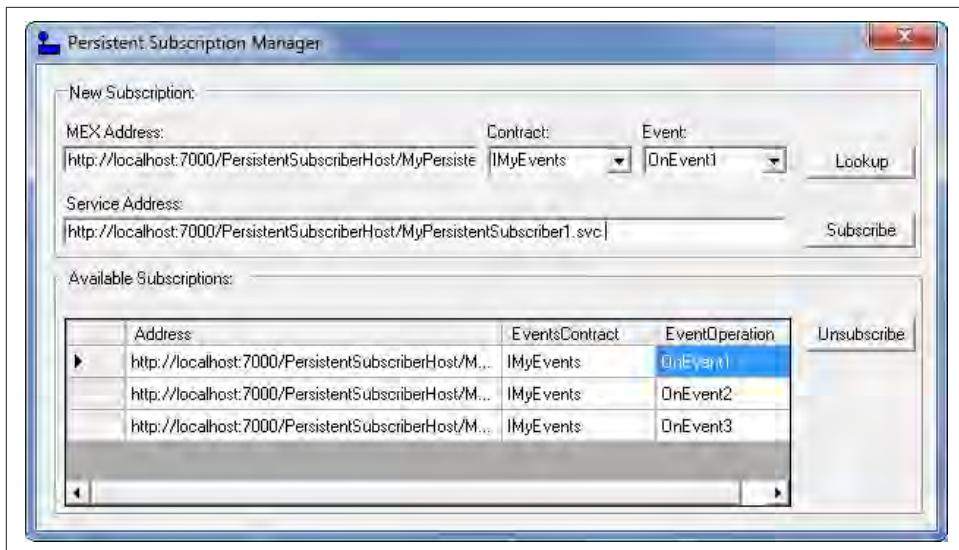


Figure D-2. The Persistent Subscription Manager application

The administration tool uses `IPersistentSubscriptionService` to add and remove subscriptions. To add a new subscription, you need to provide it with the metadata exchange address of the events contract definition. You can use the metadata exchange address of the persistent subscriber itself or the metadata exchange address of the publishing service (such as the one shown in [Example D-5](#)), because they are polymorphic. Enter the metadata exchange base address in the MEX Address text box and click the Lookup button. The tool will programmatically retrieve the metadata of the event service and populate the Contract and Event combo boxes. My Metadata Helper class, presented in [Chapter 2](#), retrieves the metadata and parses its content.

Once you've provided the address of the persistent subscriber, click the Subscribe button. Persistent Subscription Manager then adds the subscription by calling to the subscription service (`MySubscriptionService` in the examples so far). The Persistent Subscription Manager config file maintains the address for the subscription service.

Singleton subscriber

While duplex operations are, in general, the only way to subscribe a live object, there is one exception to that rule: a singleton subscriber. You can treat a singleton service as just another persistent subscriber and add its address to the subscription store. This technique is particularly useful for user interface applications that need to monitor some events. You can use my `FormHost<F>` (presented in [Chapter 8](#)) to expose the form as a singleton, and then add the form as a persistent subscriber. Add the form using the Persistent Subscription Manager tool, or the form can subscribe itself upon startup.



The publish-subscribe pattern also decouples the system security-wise. Publishers only need to authenticate themselves against a single publishing service, as opposed to multiple subscribers and potentially multiple security mechanisms. The subscribers, in turn, only need to allow the publishing service, rather than all publishers in the system, to deliver events; they trust the publishing service to properly authenticate and authorize the publishers. Applying role-based security on the publishing service allows you to easily enforce in one place various rules regarding who is allowed to publish an event across the system.

Queued Publishers and Subscribers

Instead of using the synchronous bindings to either publish or subscribe to the events, you can use the `NetMsmqBinding`. A queued publish-subscribe service combines the benefits of a loosely coupled system and the flexibility of disconnected execution. When using queued events, all events on the contract must, of course, be marked as one-way operations. As shown in [Figure D-3](#), you can use queuing at either end independently.

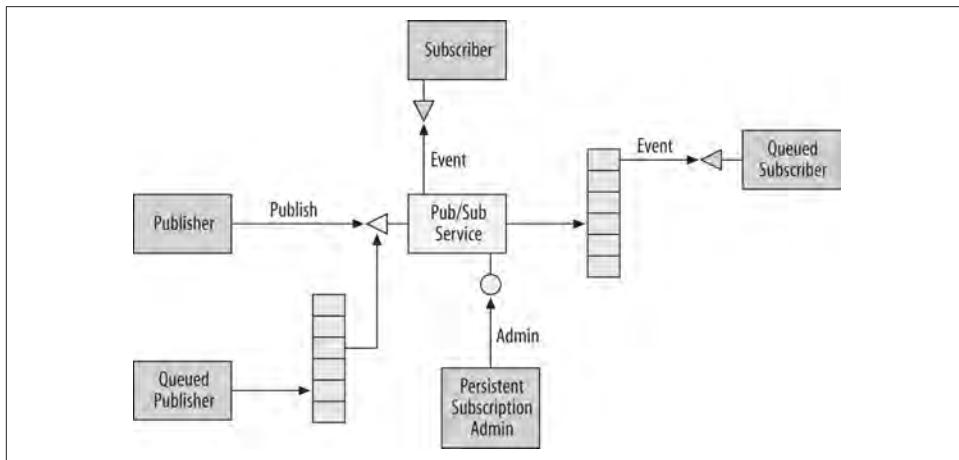


Figure D-3. Queued publish-subscribe deployment

You can have a queued publisher and connected synchronous subscribers, you can have a connected publisher publishing to queued subscribers, or you can have both queued publishers and queued subscribers. Note, however, that you cannot have queued transient subscriptions—there is no support within the MSMQ binding for duplex callbacks, since that would render the disconnected aspect of the communication useless. As before, you can use the administration tool to manage the subscribers, and the administration operations are still connected and synchronous.

Queued publisher

To utilize a queued publisher, the publishing service needs to expose a queued endpoint using the MSMQ binding. When firing events at a queued publisher, the publishing service can be offline or the publishing client itself can be disconnected. Note that when publishing two events to a queued publishing service, there are no guarantees as to the order in which these events will be delivered to and processed by the end subscribers. Due to the asynchronous concurrent publishing, there is no order, even when the events contract is configured for a session.

Queued subscriber

To deploy a queued subscriber, the persistent subscribing service needs to expose a queued endpoint. This will enable the subscriber to be offline even when the publisher is online. When the subscriber connects again, it will receive all of its queued-up events. In addition, queued subscribers can handle the case when the publishing service itself is disconnected, because no events are lost. Of course, having both a queued publisher and subscriber allows both to work offline at the same time.

When multiple events are fired at a single queued subscriber, there are no guarantees as to the order of delivery of the events, even when you configure the events contract for a session.

Publish-Subscribe with Discovery

You can also use the mechanisms of discovery and announcements (presented in [Appendix C](#)) to provide yet another way of implementing a publish-subscribe system.

Unlike all the techniques presented thus far for supporting the publish-subscribe design pattern, a discovery-based solution is the only publish-subscribe case that requires no explicit steps by the subscribers or administrator. When utilizing discovery, there is no need to explicitly subscribe either in code or in the config file. In turn, this significantly simplifies the deployment of the system and it enables great volatility in the presence of both publishers and subscribers. You can easily add or remove subscribers and publishers without any additional administration steps or programming.

When taking advantage of discovery for a publish-subscribe system, the subscribers can provide a discovery endpoint so that the publish-subscribe service can discover them, or they can announce their event-handling endpoints, or they can even do both. However, the likelihood of supporting both is low. Announcing subscribers is very much akin to the transient subscribers presented previously. Announcing subscribers is a straightforward way for a running live service instance to receive events. However, unlike transient subscribers management that relies on a fragile list of duplex proxies, susceptible to timeouts and other communication failures, the frame-

work I will present next constructs the proxies on the fly every time, and is far more robust for firing events at live objects. Discoverable subscribers are akin to the persistent subscribers and you can use them easily to fire events at a subscriber that normally persists on the disk (unless the subscriber is a singleton).

The publishers should not discover the subscribers directly, since that may incur the discovery latency on every event firing (having the cardinality of all endpoints). Instead, the publishers should discover the publish-subscribe service, which is a one-time negligible cost. The publish-subscribe service should be a singleton (enabling fast discovery, since it has cardinality of one). The publish-subscribe service exposes the same event endpoint as the subscribers, so it looks like a meta-subscriber to the publishers; that is, it requires the same code to fire the event at the publish-subscribe service as against an actual subscriber. The events endpoint of the publish-subscribe service must use a particular scope. This scope enables the publishers to find the publish-subscribe service rather than the subscribers. In addition to supporting discovering that specially scoped events endpoint, the publish-subscribe service provides an announcement endpoint.

The publish-subscribe service maintains a list of all subscribers. The publish-subscribe service can keep that list current by constantly trying to discover the subscribers using some ongoing background activity. Note that having the publish-subscribe service's events endpoint associated with a special scope will also prevent the publish-subscribe service from discovering itself when discovering all events endpoints. The publish-subscribe service can also provide an announcement endpoint to monitor subscribers. [Figure D-4](#) depicts this architecture.

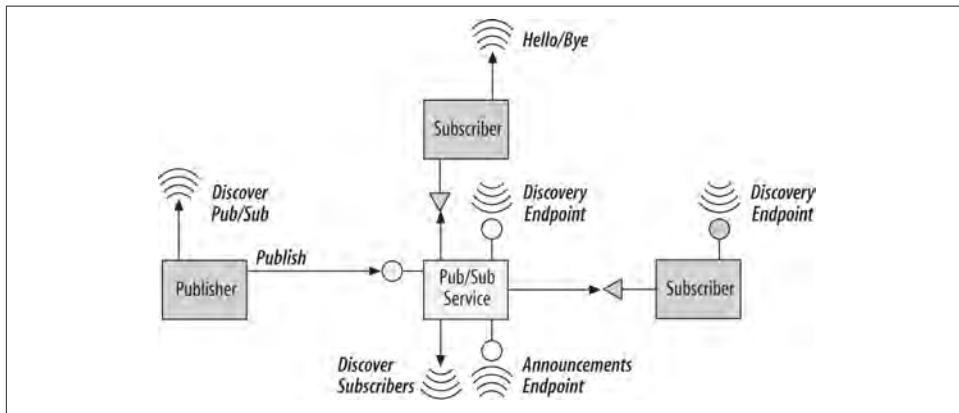


Figure D-4. Discovery-driven publish-subscribe system

The `DiscoveryPublishService<T>` Class

To facilitate deploying your own publish-subscribe service, I wrote the `DiscoveryPublishService<T>`, defined as:

```

[ServiceBehavior(InstanceContextMode = InstanceContextMode.Single)]
public abstract class DiscoveryPublishService<T> : IDisposable
    where T: class
{
    public static readonly Uri Scope;

    protected void FireEvent(params object[] args);
    //More members
}

```

All you need to do is to derive your publish-subscribe service from `DiscoveryPublishService<T>` and specify the events contract as the type parameter. Then, implement the operations of the event contract by calling the `FireEvent()` method, as in [Example D-5](#).

For example, consider the following events contract:

```

[ServiceContract]
interface IMyEvents
{
    [OperationContract(IsOneWay = true)]
    void OnEvent1();

    [OperationContract(IsOneWay = true)]
    void OnEvent2(int number);

    [OperationContract(IsOneWay = true)]
    void OnEvent3(int number,string text);
}

```

[Example D-9](#) shows how to implement your publish-subscribe service using `DiscoveryPublishService<T>`.

Example D-9. Implementing a publish-subscribe service

```

class MyPublishService : DiscoveryPublishService<IMyEvents>,IMyEvents
{
    public void OnEvent1()
    {
        FireEvent();
    }
    public void OnEvent2(int number)
    {
        FireEvent(number);
    }
    public void OnEvent3(int number,string text)
    {
        FireEvent(number ,text);
    }
}

```

To host your publish-subscribe service, use the static helper method `CreateHost<S>()` of `DiscoveryPublishService<T>`:

```
public class DiscoveryPublishService<T> : IDisposable where T: class
{
    public static ServiceHost<S> CreateHost<S>() where S : DiscoveryPublishService<T>,T;
    //More members
}
```

The type parameter `S` is your subclass of `DiscoveryPublishService<T>`, and `T` is the events contract. `CreateHost<S>()` returns an instance of a service host you need to open. There is no need to use a config file or to add the events endpoint:

```
ServiceHost host = DiscoveryPublishService<IMyEvents>.
                           CreateHost<MyPublishService>();
host.Open();

//Sometime later

host.Close();
```

In addition, `CreateHost<S>()` will obtain an available TCP base address and add the events endpoint. [Example D-10](#) shows the implementation of `CreateHost<S>()`.

Example D-10. Implementing CreateHost<S>()

```
public abstract class DiscoveryPublishService<T> : IDisposable where T : class
{
    public readonly static Uri Scope;

    static DiscoveryPublishService()
    {
        Scope = new Uri("net.tcp://ServiceModelEx.DiscoveryPublishService."
                         + typeof(T));
    }
    static NetTcpBinding Binding
    {
        get
        {
            return new NetTcpBinding(SecurityMode.Transport,true);
        }
    }
    public static ServiceHost<S> CreateHost<S>()
        where S : DiscoveryPublishService<T>,T
    {
        ServiceHost<S> host =
            DiscoveryFactory.CreateDiscoverableHost<S>(Scope,false);

        foreach(ServiceEndpoint endpoint in host.Description.Endpoints)
        {
```

```

        if(endpoint.Address.Uri.Scheme == Uri.UriSchemeNetTcp)
        {
            endpoint.Binding = Binding;
        }
    }
    return host;
}
}

```

The static constructor of `DiscoveryPublishService<T>` initializes the scope of the events endpoint. Note that the scope contains the events contract type so that multiple publish-subscribe services that use different events contracts do not conflict with each other.

The `CreateHost<S>()` method largely leverages `DiscoveryFactory.CreateDiscoverableHost<S>()` to create a discoverable announcing host, whose endpoints uses the designated scope. However, since `CreateDiscoverableHost<S>()` uses the default endpoints which in turn use the default TCP binding, `CreateHost<S>()` sets explicitly the binding for each TCP to use reliability as well.

The Publisher

The publisher needs a proxy to the events service. For that, use my `DiscoveryPublishService<T>.CreateChannel()`:

```

public class DiscoveryPublishService<T> : IDisposable where T : class
{
    public static T CreateChannel()
    {
        EndpointAddress address = DiscoveryHelper.DiscoverAddress<T>(Scope);
        return ChannelFactory<T>.CreateChannel(Binding,address);
    }
}

```

`DiscoveryPublishService<T>.CreateChannel()` is fast, since the cardinality is one. Note the use of the publish-subscribe endpoint scope. The code of the publisher is straightforward:

```

IMyEvents proxy = DiscoveryPublishService<IMyEvents>.CreateChannel();
proxy.OnEvent1();
(proxy as ICommunicationObject).Close();

```

The Subscriber

There is nothing special to do with a subscriber. Simply support the events contract on a service and add either discovery or announcements (or both) of the events endpoint.



There is little point in combining queuing with a discovery-driven publish-subscribe system. The reason is that discovery requires the host to be running, and that rules out disconnected publish-subscribe service or disconnected subscribers. That said, you can easily modify my helper classes to support queued endpoints as well.

More on DiscoveryPublishService<T>

Example D-11 shows the rest of the implementation of `DiscoveryPublishService<T>`.

Example D-11. Implementing DiscoveryPublishService <T>

```
[ServiceBehavior(InstanceContextMode = InstanceContextMode.Single
                 UseSynchronizationContext = false)]
public class DiscoveryPublishService<T> : IDisposable where T : class
{
    AnnouncementSink<T>    m_AnnouncedSubscribers;
    DiscoveredServices<T>  m_DiscoveredServices;

    public DiscoveryPublishService()
    {
        m_AnnouncedSubscribers = new AnnouncementSink<T>();
        m_DiscoveredServices = new DiscoveredServices<T>();
        m_AnnouncedSubscribers.Open();
        m_DiscoveredServices.Open();
    }

    public void Dispose()
    {
        m_AnnouncedSubscribers.Close();
        m_DiscoveredServices.Close();
    }

    protected void FireEvent(params object[] args)
    {
        string action = OperationContext.Current.IncomingMessageHeaders.Action;
        string[] slashes = action.Split('/');
        string methodName = slashes[slashes.Length-1];

        FireEvent(methodName,args);
    }

    void FireEvent(string methodName,object[] args)
    {
        T[] subscribers = GetSubscribers();
        Publish(subscribers,methodName,args);
    }

    T[] GetSubscribers()
    {
        IEnumerable<string> announcedAddress  = m_AnnouncedSubscribers.
```

```

                FindComplement(Scope).
Select(address=>address.Uri.AbsoluteUri);

IEnumerable<string> discoveredAddress = m_DiscoveredServices.
                FindComplement(Scope).
Select(address=>address.Uri.AbsoluteUri);

IEnumerable<string> addresses = announcedAddress.Union(discoveredAddress);

List<T> subscribers = new List<T>();
foreach(string address in addresses)
{
    EndpointAddress endpointAddress = new EndpointAddress(address);
    Binding binding = GetBindingFromAddress(endpointAddress);
    T proxy = ChannelFactory<T>.CreateChannel(binding,endpointAddress);
    subscribers.Add(proxy);
}
return subscribers.ToArray();
}
static Binding GetBindingFromAddress(EndpointAddress address)
{
    if(address.Uri.Scheme == "net.tcp")
    {
        return Binding;
    }
    if(address.Uri.Scheme == "net.pipe")
    {
        return new NetNamedPipeBinding();
    }
    Debug.Assert(false,"Unsupported binding specified");
    return null;
}
static void Publish(T[] subscribers,string methodName,object[] args)
{
    Action<T> fire = (subscriber)=>
    {
        using(subscriber as IDisposable)
        {
            Invoke(subscriber as T,methodName,args);
        }
    };
    subscribers.ForEachAsync(fire);
}
static void Invoke(T subscriber,string methodName,object[] args)
{
    Type type = typeof(T);
    MethodInfo methodInfo = type.GetMethod(methodName);
    try
    {
        methodInfo.Invoke(subscriber,args);
    }
    catch(Exception e)

```

```
        {
            Trace.WriteLine(e.Message);
        }
    }
}
```

`DiscoveryPublishService<T>` maintains two lists of subscribers. The first uses my `AnnouncementSink<T>` for announced subscribers and the second uses my `DiscoveredServices<T>` for discovered services. The `FireEvent()` method used by its subclasses extracts the operation name from the message headers and calls an internal `FireEvent()` method, which calls the `GetSubscribers()` method. `GetSubscribers()` queries the subscribers list for all subscribers that do not support the publish-subscribe scope (to avoid self-discovery). It then merges the lists into a union of unique entries (this is required to deal with subscribers that both announce themselves and are discoverable). For each subscriber, `GetSubscribers()` infers the binding from the address scheme and creates a proxy to fire at the subscriber. Publishing the events is done asynchronously and in parallel using the `ForEachAsync<T>()` helper extension method, very similar to the technique already described for `PublishService<T>` in [Example D-6](#).

Generic Interceptor

At its core, WCF is nothing more than an extensibility model. All the built-in attributes, behaviors, and infrastructure are implemented using this open and public extensibility model. In theory, every developer has as much access and power as any member of the WCF team. As demonstrated throughout this book (with the exception of my security extensions), as long as you are familiar with the extensibility model, with a mere few lines of code you can inject powerful behavior and customization into your application. In practice, however, dealing with the interception mechanism requires intimate knowledge of the WCF architecture. To simplify matters, I wanted to provide an easy-to-use abstraction on top of the WCF extensibility model that would, in a way, extend the extensibility model itself, allowing for intercepting all calls to the service, both on the client and the service side, in a general manner, and adding custom behavior, without having to deal with the inner workings of WCF. This appendix presents a small framework I call the *generic interceptor*, available with *ServiceModelEx*. It also demonstrates some advanced WCF programming techniques, as well as the thought process behind designing such extensions.

Intercepting Service Operations

Recall from [Chapter 1](#) that in the abstract, all WCF does when intercepting calls is perform pre- and post-call operations. Adding custom steps to this interception mechanism is probably the most common way of extending WCF.

Every endpoint dispatcher has a reference to an interface called `IOperationInvoker`, defined as:

```
public interface IOperationInvoker
{
    object[] AllocateInputs();
    object Invoke(object instance,object[] inputs,out object[] outputs);
```

```
        //Asynchronous invocation methods  
    }
```

The dispatcher uses the `Invoke()` method to invoke the calls on the service instance. In providing for the invocation, `IOperationInvoker` is the right place to plug in your code. Specifically, assigning the dispatcher your implementation of `IOperationInvoker` will enable you to hook it in.

The Generic Invoker

The first step in implementing my generic interceptor framework was to provide an abstract implementation of `IOperationInvoker` that enables custom pre- and post-call steps, as shown in [Example E-1](#).

Example E-1. The GenericInvoker class

```
public abstract class GenericInvoker : IOperationInvoker  
{  
    readonly IOperationInvoker m_OldInvoker;  
  
    public GenericInvoker(IOperationInvoker oldInvoker)  
    {  
        m_OldInvoker = oldInvoker;  
    }  
    public virtual object[] AllocateInputs()  
    {  
        return m_OldInvoker.AllocateInputs();  
    }  
    protected virtual void PreInvoke(object instance,object[] inputs)  
    {}  
  
    //Always called, even if operation had an exception  
    protected virtual void PostInvoke(object instance,object returnedValue,  
                                      object[] outputs,Exception exception)  
    {}  
  
    public object Invoke(object instance,object[] inputs,out object[] outputs)  
    {  
        PreInvoke(instance,inputs);  
        object returnedValue = null;  
        object[] outputParams = new object[]{};  
        Exception exception = null;  
        try  
        {  
            returnedValue = m_OldInvoker.Invoke(instance,inputs,out outputParams);  
            outputs = outputParams;  
            return returnedValue;  
        }  
        catch(Exception operationException)
```

```

    {
        exception = operationException;
        throw;
    }    finally
{
    PostInvoke(instance,.returnValue,outputParams,exception);
}
}
// Additional asynchronous methods
}

```

`GenericInvoker` defines two virtual methods, `PreInvoke()` and `PostInvoke()`. `PreInvoke()` accepts the input parameters as well as the target instance, and `PostInvoke()` accepts the output parameters and the returned value as well as the instance and the exception (if one took place). `GenericInvoker` has an empty implementation for both methods. It is up to subclasses of `GenericInvoker` to override one or both of these methods and add the custom steps.

`GenericInvoker` accepts as a construction parameter the old implementation of `IOperationInvoker` that was associated with the service. This old implementation does the heavy lifting of allocating the input parameters for the operations, as well as actually invoking the service. `GenericInvoker` aims at being as nonintrusive as possible, so it cannot interfere with that implementation; at any rate, it would not be wise to do so, as that would entail a large amount of work. `GenericInvoker` therefore saves the old invoker in a member variable, and delegates to it its implementation of `AllocateInputs()`. The heart of `GenericInvoker` is the `Invoke()` method. In it, `GenericInvoker` first calls the `PreInvoke()` method (allowing its subclass to perform some pre-call processing) and then proceeds to invoke the operation using the old invoker. `GenericInvoker` encases the invocation in a `try/catch/finally` block. Regardless of how the operation ends (with or without an exception), `GenericInvoker` calls the `PostInvoke()` method, providing it with the retuned results and the exception, and allowing the subclass to perform custom post-call processing.

Installing the Interceptor

I wanted to provide for a declarative way of installing the operation, both at the operation level and at the service level. The trick in doing that is implementing the work once at the operation level, and then having the service level install all the operation-level attributes. The `IOperationBehavior` interface is the operation-level extension that lets you customize the dispatcher for an operation, in the `ApplyDispatchBehavior()` method:

```

public interface IOperationBehavior
{
    void ApplyDispatchBehavior(OperationDescription operationDescription,
                               DispatchOperation dispatchOperation);
}

```

```
        //More methods  
    }
```

Any method-level attribute that implements `IOperationBehavior` will be given a chance to affect the dispatcher (in this case, setting its operation invoker) in the `ApplyDispatchBehavior()` method. `ApplyDispatchBehavior()` provides the `dispatchOperation` parameter of the type `DispatchOperation`:

```
public sealed class DispatchOperation  
{  
    public IOperationInvoker Invoker  
    {get;set;}  
  
    //More members  
}
```

Setting the `Invoker` property replaces the implementation of `IOperationInvoker` used. It's as simple as that.

Example E-2 shows the implementation of my `OperationInterceptorBehaviorAttribute`.

Example E-2. The OperationInterceptorBehavior attribute

```
[AttributeUsage(AttributeTargets.Method)]  
public abstract class OperationInterceptorBehaviorAttribute :  
    Attribute, IOperationBehavior  
{  
    protected abstract GenericInvoker CreateInvoker(IOperationInvoker oldInvoker);  
  
    public void ApplyDispatchBehavior(OperationDescription operationDescription,  
                                      DispatchOperation dispatchOperation)  
    {  
        IOperationInvoker oldInvoker = dispatchOperation.Invoker;  
        dispatchOperation.Invoker = CreateInvoker(oldInvoker);  
    }  
    //More methods  
}
```

The `OperationInterceptorBehaviorAttribute` is an abstract class with an abstract protected method called `CreateInvoker()`. `CreateInvoker()` takes the old invoker and returns some implementation of `GenericInvoker`. The implementation of `ApplyDispatchBehavior()` first saves the old invoker in a local variable and then calls `CreateInvoker()` to provide a new invoker while wrapping the old invoker. The newly created invoker is duly set on the dispatcher as the invoker to use from now on. Having a concrete subclass of the `OperationInterceptorBehavior` attribute will enable you to apply the custom invoker discretely on some, but perhaps not all, of the methods of the service. If you wish to apply the attribute on all operations, it is better

to enforce this design decision at the service level using my `ServiceInterceptorBehaviorAttribute`, defined in [Example E-3](#).

Example E-3. The ServiceInterceptorBehavior attribute

```
[AttributeUsage(AttributeTargets.Class)]
public abstract class ServiceInterceptorBehaviorAttribute :
    Attribute, IServiceBehavior
{
    protected abstract OperationInterceptorBehaviorAttribute
        CreateOperationInterceptor();

    public void ApplyDispatchBehavior(ServiceDescription serviceDescription,...)
    {
        foreach(ServiceEndpoint endpoint in serviceDescription.Endpoints)
        {
            foreach(OperationDescription operation in endpoint.Contract.Operations)
            {
                if(operation.Behaviors.
                    Find<OperationInterceptorBehaviorAttribute>() != null)
                {
                    continue;
                }
                operation.Behaviors.Add(CreateOperationInterceptor());
            }
        }
    }
    //More members
}
```

`ServiceInterceptorBehavior`, too, is an abstract attribute. It provides the abstract protected method `CreateOperationInterceptor()`, which returns some implementation of the `OperationInterceptorBehavior` attribute. `ServiceInterceptorBehavior` supports the `IServiceBehavior` interface, whose `ApplyDispatchBehavior()` method provides the description of the service:

```
public interface IServiceBehavior
{
    void ApplyDispatchBehavior(ServiceDescription serviceDescription,...);
    //More methods
}
```

The `ServiceDescription` class contains a collection of service endpoints:

```
public class ServiceDescription
{
    public ServiceEndpointCollection Endpoints
    {get;}

    //More members
}
```

```
public class ServiceEndpointCollection : Collection<ServiceEndpoint>
{...}
```

Every endpoint has a **Contract** property containing the contract description:

```
public class ServiceEndpoint
{
    public ContractDescription Contract
    {get;}

    //More members
}
```

The contract description has a collection of operation descriptions:

```
public class ContractDescription
{
    public OperationDescriptionCollection Operations
    {get;}

    //More members
}

public class OperationDescriptionCollection : Collection<OperationDescription>
{...}
```

Each operation description has a collection of operation behaviors:

```
public class OperationDescription
{
    public KeyedByTypeCollection<IOperationBehavior> Behaviors
    {get;}

    //More members
}
```

The service-level attribute needs to add to this collection of behaviors an **OperationInterceptorBehavior** attribute.

In its implementation of **ApplyDispatchBehavior()**, **ServiceInterceptorBehavior** iterates over the collection of service endpoints. For each endpoint, it iterates over its operation collection. For each operation, it checks to see whether its behavior collection already contains an implementation of the **OperationInterceptorBehavior** attribute. This check is required in case the developer applied (by mistake) both an operation- and a service-level attribute. If the behavior collection does not contain the **OperationInterceptorBehavior** attribute, **ApplyDispatchBehavior()** adds it.

Intercepting Client Calls

To intercept client-side calls, WCF provides the interface `IClientMessageInspector`, defined as:

```
public interface IClientMessageInspector
{
    object BeforeSendRequest(ref Message request, IClientChannel channel);
    void AfterReceiveReply(ref Message reply, object correlationState);
}
```

The `BeforeSendRequest()` method is called just before the message is sent down the wire, allowing you to affect the request message. Similarly, the `AfterReceiveReply()` method is your chance to interact with the reply message for post-call processing.

The client runtime represented by the `ClientRuntime` class contains a collection of message inspectors:

```
public sealed class ClientRuntime
{
    public SynchronizedCollection<IClientMessageInspector> MessageInspectors
    {get;}
    //More members
}
```

You can add your message inspector to the collection by associating the proxy with an endpoint behavior. That behavior needs to add the inspector in the `ApplyClientBehavior()` method:

```
public interface IEndpointBehavior
{
    void ApplyClientBehavior(ServiceEndpoint endpoint,
                           ClientRuntime clientRuntime);
    //More members
}
```

To encapsulate these steps, I wrote the class `InterceptorClientBase<T>`, defined in [Example E-4](#).

Example E-4. The InterceptorClientBase<T> class

```
public abstract class InterceptorClientBase<T> : ClientBase<T> where T : class
{
    public InterceptorClientBase()
    {
        Endpoint.Behaviors.Add(new ClientInterceptor(this));
    }
    public InterceptorClientBase(string endpointName) : base(endpointName)
    {
        Endpoint.Behaviors.Add(new ClientInterceptor(this));
    }
}
```

```

}

//More constructors

protected virtual void PreInvoke(ref Message request)
{}

protected virtual void PostInvoke(ref Message reply)
{}

class ClientInterceptor : IEndpointBehavior,IClientMessageInspector
{
    InterceptorClientBase<T> Proxy
    {get;set;}

    internal ClientInterceptor(InterceptorClientBase<T> proxy)
    {
        Proxy = proxy;
    }

    object IClientMessageInspector.BeforeSendRequest(ref Message request,
                                                       IClientChannel channel)
    {
        Proxy.PreInvoke(ref request);
        return null;
    }

    void IClientMessageInspector.AfterReceiveReply(ref Message reply,
                                                    object correlationState)
    {
        Proxy.PostInvoke(ref reply);
    }

    void IEndpointBehavior.ApplyClientBehavior(ServiceEndpoint endpoint,
                                                ClientRuntime clientRuntime)
    {
        clientRuntime.MessageInspectors.Add(this);
    }

    //Rest of the implementation
}
}

```

`InterceptorClientBase<T>` defines a nested private class called `ClientInterceptor` that implements both `IEndpointBehavior` and `IClientMessageInspector`. The constructors of `InterceptorClientBase<T>` add an instance of `ClientInterceptor` to the proxy's collection of endpoint behaviors. Inside `ClientInterceptor`, the implementation of `ApplyClientBehavior()` adds itself to the collection of client runtime interceptors. `InterceptorClientBase<T>` provides two virtual methods, `PreInvoke()` and `PostInvoke()`, for the use of derived classes. The constructor of `ClientInterceptor` takes a back reference to the calling `InterceptorClientBase<T>`, and

it uses that reference to call back to the `PreInvoke()` and `PostInvoke()` methods during the calls to `BeforeSendRequest()` and `AfterReceiveReply()` respectively.

The Trace Interceptors

The first example I'll show of using the generic interceptors framework is for tracing and logging. **Example E-5** shows a simple example of implementing a generic service-side interceptor called `ParameterTracerInvoker`.

Example E-5. The ParameterTracerInvoker

```
class ParameterTracerInvoker : GenericInvoker
{
    public ParameterTracerInvoker(IOperationInvoker oldInvoker) : base(oldInvoker)
    {}

    protected override void PreInvoke(object instance, object[] inputs)
    {
        Trace.WriteLine("Input Parameters: ");

        foreach(object argument in inputs)
        {
            Trace.WriteLine(argument.ToString());
        }
    }

    protected override void PostInvoke(object instance, object returnedValue,
                                      object[] outputs, Exception exception)
    {

        foreach(object output in outputs)
        {
            Trace.WriteLine("Output Parameters: ");
            Trace.WriteLine(output.ToString());
        }

        Trace.WriteLine("Returned: " + returnedValue ?? String.Empty);
    }
}
```

`ParameterTracerInvoker` derives from `GenericInvoker`. Its constructor accepts the old invoker and passes it to the constructor of `GenericInvoker`. The implementations of `PreInvoke()` and `PostInvoke()` trace to the Output window in Visual Studio the values of the input and output parameters, respectively. You can install the `ParameterTracerInvoker` on service operations by defining the `OperationParameterTracerAttribute`:

```
public class OperationParameterTracerAttribute :
    OperationInterceptorBehaviorAttribute
```

```

{
    protected override GenericInvoker CreateInvoker(IOperationInvoker oldInvoker)
    {
        return new ParameterTracerInvoker(oldInvoker);
    }
}

```

All the method-level attribute needs to do is derive from `OperationContractBehaviorAttribute` and override the `CreateInvoker()` method, returning an instance of `ParameterTracerInvoker` wrapping the old invoker.

Example E-6 demonstrates using the `OperationParameterTracer` attribute.

Example E-6. Using OperationParameterTracerAttribute

```

[ServiceContract]
interface IMyContract
{
    [OperationContract]
    string MyMethod1(int number,string text);

    [OperationContract]
    string MyMethod2(int number);
}

class MyService : IMyContract
{
    [OperationParameterTracer]
    public string MyMethod1(int number,string text)
    {
        return "Some Result 1";
    }

    public string MyMethod2(int number)
    {
        return "Some Result 2";
    }
}

```

Using the definitions from **Example E-6**, the following client code:

```

MyContractClient proxy = new MyContractClient();
proxy.MyMethod1(287,"Hello");
proxy.MyMethod2(42);
proxy.Close();

```

would trace:

```

Input Parameters:
287
Hello

```

```
Output Parameters:  
Returned: Some Result 1
```

To apply the `ParameterTracerInvoker` at the service level, define the `ServiceParameterTracerAttribute` as:

```
public class ServiceParameterTracerAttribute :  
    ServiceInterceptorBehaviorAttribute  
{  
    protected override OperationInterceptorBehaviorAttribute  
        CreateOperationInterceptor()  
    {  
        return new OperationParameterTracerAttribute();  
    }  
}
```

All the service-level attribute needs to do is derive from `ServiceInterceptorBehaviorAttribute` and override the `CreateOperationInterceptor()` method, returning an instance of `OperationParameterTracerAttribute`.

Example E-7 demonstrates using the `ServiceParameterTracer` attribute.

Example E-7. Using ServiceParameterTracerAttribute

```
[ServiceParameterTracer]  
class MyService : IMyContract  
{  
    public string MyMethod1(int number, string text)  
    {  
        return "Some Result 1";  
    }  
  
    public string MyMethod2(int number)  
    {  
        return "Some Result 2";  
    }  
}
```

Using the definitions of **Example E-7**, the following client code:

```
MyContractClient proxy = new MyContractClient();  
proxy.MyMethod1(287, "Hello");  
proxy.MyMethod2(42);  
proxy.Close();
```

would trace:

```
Input Parameters:  
287  
Hello  
Returned: Some Result 1  
Input Parameters:
```

You can use the service-side interceptor independently from the client, or you can intercept and trace on the client using the `TracerClientBase<T>` class shown in [Example E-8](#).

Example E-8. The TracerClientBase<T> class

```
public class TracerClientBase<T> : InterceptorClientBase<T> where T : class
{
    protected override void PreInvoke(ref Message request)
    {
        string action = request.Headers.Action;
        string[] slashes = action.Split('/');
        string methodName = slashes[slashes.Length-1];

        Trace.WriteLine("***** Calling : " + methodName + "()" + "*****");
    }

    protected override void PostInvoke(ref Message reply)
    {
        string action = reply.Headers.Action;
        string[] slashes = action.Split('/');
        string methodName = slashes[slashes.Length-1];

        methodName = methodName.Replace("Response","");
        Trace.WriteLine("**** Returning from : " + methodName + "()" + "****");
    }
}
```

The class `TracerClientBase<T>` derives from `InterceptorClientBase<T>`, and it overrides the `PreInvoke()` and `PostInvoke()` methods. The overridden methods trace the invoked operation name. You use `TracerClientBase<T>` just as a regular proxy base class, as shown in [Example E-9](#).

Example E-9. Deriving from TracerClientBase<T>

```
class MyContractClient : TracerClientBase<IMyContract>,IMyContract
{
    public string MyMethod1(int number,string text)
    {
        return Channel.MyMethod1(number,text);
    }

    public string MyMethod2(int number)
    {
        return Channel.MyMethod2(number);
```

```
    }  
}
```

Using the proxy from [Example E-9](#) with the service from [Example E-7](#), the following code:

```
MyContractClient proxy = new MyContractClient();  
proxy.MyMethod1(287,"Hello");  
proxy.Close();
```

would trace:

```
**** Calling operation: MyMethod1() ****  
Input Parameters:  
287  
Hello  
Returned: Some Result 1  
**** Returning from operation: MyMethod1() ****
```

Identity Stack Propagation

The second example of using the generic interceptor is about security identity propagation. As explained in [Chapter 10](#), impersonation as a mechanism for identity propagation has many liabilities. Still, sometimes your service is required to pass the identity of the original caller (or all callers) down to the resources or other services with which it interacts. Instead of impersonating the callers or passing their identities as explicit parameters, you can pass the identities out-of-band, in the message headers, and use the generic interceptor to automate processing of those identities.

The first step is to define the stack of callers. To that end, I defined the `SecurityCallFrame`, which represents a single caller identity as well as some additional information about the caller, such as its address and the operation it invoked:

```
[DataContract]  
public class SecurityCallFrame  
{  
    [DataMember(IsRequired = true)]  
    public string Authentication  
    {get;}  
  
    [DataMember(IsRequired = true)]  
    public string IdentityName  
    {get;}  
  
    [DataMember(IsRequired = true)]  
    public string Address  
    {get;}  
  
    [DataMember(IsRequired = true)]  
    public string Operation  
    {get;}}
```

```
        //More members  
    }
```

Next, I defined the security call stack:

```
[DataContract]  
public class SecurityCallStack  
{  
    internal void AppendCall();  
  
    public SecurityCallFrame OriginalCall  
    {get;}  
  
    public int Count  
    {get;}  
  
    public SecurityCallFrame[] Calls  
    {get;}  
  
    //More members  
}
```

The implementation details of these types are irrelevant for this appendix.

Using the `GenericContext<T>` from [Appendix B](#), I defined the security call stack context:

```
[DataContract]  
public class SecurityCallStackContext  
{  
    public static SecurityCallStack Current  
    {  
        get  
        {  
            if(GenericContext<SecurityCallStack>.Current == null)  
            {  
                return null;  
            }  
            return GenericContext<SecurityCallStack>.Current.Value;  
        }  
        set  
        {  
            GenericContext<SecurityCallStack>.Current = new GenericContext  
                <SecurityCallStack>  
                (value);  
        }  
    }  
}
```

To automate passing the call stack, I then defined the class `SecurityCallStack ClientBase<T>` shown in [Example E-10](#).

Example E-10. The SecurityCallStackClientBase<T> class

```
public abstract partial class HeaderClientBase<T,H> : InterceptorClientBase<T>
{
    protected H Header
    {get;set;}

    protected override void PreInvoke(ref Message reply);

    //Rest of the implementation
}
public abstract class SecurityCallStackClientBase<T> :
    HeaderClientBase<T,SecurityCallStack>
{
    //Constructors

    void InitializeCallStack()
    {
        if(OperationContext.Current == null || Header == null)
        {
            Header = new SecurityCallStack();
        }
        else
        {
            Header = SecurityCallStackContext.Current;
        }
    }
    protected override void PreInvoke(ref Message request)
    {
        InitializeCallStack();
        Header.AppendCall();
        base.PreInvoke(ref request);
    }
}
```

SecurityCallStackClientBase<T> derives from HeaderClientBase<T,H> (also defined in [Appendix B](#)), and HeaderClientBase<T,H> in turn derives from InterceptorClientBase<T>.

Every time the client invokes calls on the SecurityCallStackClientBase<T> proxy, the proxy will automatically append the current identity to the call stack and pass it in the headers. If all services down the call chain use SecurityCallStackClientBase<T> (or manually use SecurityCallStackContext), on every call the call stack will contain the new frame.

Security Call Stack Interceptor

To process and manage the identities stack, define a subclass of the generic interceptor. The processing could be as simple as logging the identities, or as complex as digit-

tally signing the call stack to prevent spoofing by malicious intermediaries along the call chain and validating the current call stack, as shown in [Example E-11](#).

Example E-11. Validating and logging the call stack with an interceptor

```
class SecurityCallStackInterceptor : GenericInvoker
{
    public SecurityCallStackInterceptor(IOperationInvoker oldInvoker) :
        base(oldInvoker)
    {}

    protected override void PreInvoke(object instance, object[] inputs)
    {
        SecurityCallStack callStack = SecurityCallStackContext.Current;

        if(callStack != null)
        {
            LogCallChain(callStack);
            ValidateCallChain(callStack);
            SignCallChain(callStack);
        }
    }

    void LogCallChain(SecurityCallStack callStack)
    {...}

    void ValidateCallChain(SecurityCallStack callStack)
    {
        //Perform custom validation steps here
    }

    void SignCallChain(SecurityCallStack callStack)
    {
        //Digitally sign call stack here
    }
}
```

You can apply the `SecurityCallStackInterceptor` at the operation or service level using dedicated one-liner attributes:

```
public class OperationSecurityCallStackAttribute :
    OperationInterceptorBehaviorAttribute
{
    protected override GenericInvoker CreateInvoker(IOperationInvoker oldInvoker)
    {
        return new SecurityCallStackInterceptor(oldInvoker);
    }
}

public class SecurityCallStackBehaviorAttribute :
    ServiceInterceptorBehaviorAttribute
{
    protected override OperationInterceptorBehaviorAttribute
```

```
        CreateOperationInterceptor()
{
    return new OperationSecurityCallStackAttribute();
}
}
```


APPENDIX F

WCF Coding Standard

A comprehensive coding standard is essential for successful product delivery; it helps in enforcing best practices and avoiding pitfalls, and it makes knowledge dissemination across the team easier. Traditionally, coding standards are thick, laborious documents, spanning hundreds of pages and detailing the rationale behind every directive. While these are better than no standard at all, such efforts are usually indigestible by the average developer. In contrast, the WCF coding standard presented here is very thin on the “why” and very detailed on the “what.” I believe that while fully understanding every insight that goes into a particular programming decision may require reading books and even years of experience, applying a standard should not. When absorbing a new developer into your team, you should be able to simply point the newcomer at the standard and say: “Read this first.” Being able to comply with a good standard should come before fully understanding and appreciating it—that should come over time, with experience. The WCF coding standard presented in this appendix captures dos and don’ts, pitfalls, guidelines, and recommendations, drawing on the best practices and helper classes discussed throughout this book.

General Design Guidelines

1. All services must adhere to these principles:
 - a. Services are secure.
 - b. Service operations leave the system in a consistent state.
 - c. Services are thread-safe and can be accessed by concurrent clients.
 - d. Services are reliable.
 - e. Services are robust.
2. Services can optionally adhere to these principles:

- a. Services are interoperable.
- b. Services are scale-invariant.
- c. Services are available.
- d. Services are responsive.
- e. Services are disciplined and do not block their clients for long.

Essentials

1. Place service code in a class library, not in any hosting EXE.
2. Do not provide parameterized constructors to a service class, unless it is a singleton that is hosted explicitly.
3. Enable reliability in the relevant bindings.
4. Provide a meaningful namespace for contracts. For outward-facing services, use your company's URL or equivalent URN with a year and month to support versioning. For example:

```
[ServiceContract(Namespace = "http://www.idesign.net/2010/09")]
interface IMyContract
{...}
```

For intranet services, use any meaningful unique name, such as `MyApplication`. For example:

```
[ServiceContract(Namespace = "MyApplication")]
interface IMyContract
{...}
```

5. With intranet applications, prefer self-hosting to IIS hosting when the WAS is unavailable.
6. Do not mix and match named bindings with default bindings. Either have all your bindings be explicitly referenced, or use only default bindings.
7. Do not mix and match named behaviors with default behaviors. Either have all your behaviors be explicitly referenced, or use only default behaviors.
8. Always name all endpoints in the client config file.
9. Do not use SvcUtil or Visual Studio to generate a config file.
10. When using a tool such as Visual Studio to generate the proxy, do clean up the proxy.
11. Do not duplicate proxy code. If two or more clients use the same contract, factor the proxy to a separate class library.
12. Always close or dispose of the proxy.

13. When using discovery, prefer dynamic addresses.
14. When using discovery, do support the metadata exchange endpoint over TCP.
15. When using discovery, avoid cardinality of “some.”

Configuration

1. Prefer programmatic configuration over using the config file.
2. Establish through testing default values relevant to your system for all common configuration settings such as `MaxReceiveMessageSize`, `MaxArrayLength`, and all timeouts.
3. Provide a default debug timeout to ease debugging your services. Apply it only when the debugger is attached.
4. Avoid setting defaults to arbitrarily large values; instead keep values as small as possible.
5. Provide for default values at both the system and service hosting levels. If not specialized, refer service-level values to system-level defaults.
6. Provide simple, but singular access to system-level default values through a helper class. Use the helper in both proxies and service hosting. For example:

```
public static class BindingHelper
{
    public static TimeSpan DefaultConnectivityTimeout
    {...}
    public static TimeSpan DefaultDebugTimeout
    {...}
    public static long DefaultMaxReceiveMessageSize
    {...}
    public static int DefaultMaxArraySize
    {...}

    public static class Tcp
    {
        static NetTcpBinding DefaultBinding()
        {...}
        public static void SetBindingOverride(NetTcpBinding binding)
        {...}
        public static NetTcpBinding Binding()
        {...}
    }

    public static class Queue
    {
        static NetMsmqBinding DefaultBinding()
```

```

    {...}
    public static void SetBindingOverride(NetMsmqBinding binding)
    {...}
    public static NetMsmqBinding Binding()
    {...}
}
...
}

```

7. Establish a consistent configuration policy for your system.
8. Use a consistent approach when applying your configuration policy such as:
 - a. Check for config file overrides by endpoint.
 - b. If no config file overrides exist, apply manual overrides.
 - c. If no overrides exist, apply default values.
 - d. Always enforce policy on all overrides.

Addressing

1. When not using discovery, reduce your addressing configuration to base address and optionally port for both clients and hosts.
2. Infer the address transport scheme via binding.
3. Consider establishing a consistent endpoint addressing convention by using a contract's full namespace as the endpoint address.
4. Apply your addressing convention consistently to all service endpoints.
5. When not using discovery, automatically build the proxy endpoint address by combining your inferred transport scheme, base address configuration and endpoint addressing convention.
6. Provide simple, but singular access to your addressing convention through a helper class. Use the helper in both proxies and service hosting. For example:

```

public static class AddressHelper
{
    static string DetermineScheme(Binding binding)
    {...}

    static string GetBaseAddressFromConfig(Type contractType = null)
    {...}

    public static EndpointAddress BuildEndpointAddress(Binding binding,
                                                       Type contractType)
    {...}
}

```

7. To ease deployment, prefer default ports.

Service Contracts

1. Always define a service contract and always apply the `ServiceContract` attribute on an interface, not a class:

```
//Avoid:  
[ServiceContract]  
class MyService  
{  
    [OperationContract]  
    public void MyMethod()  
    {...}  
}  
//Correct:  
[ServiceContract]  
interface IMyContract  
{  
    [OperationContract]  
    void MyMethod();  
}  
class MyService : IMyContract  
{  
    public void MyMethod()  
    {...}  
}
```

2. Prefix the service contract name with an `I`:

```
[ServiceContract]  
interface IMyContract  
{...}
```

3. Avoid property-like operations:

```
//Avoid:  
[ServiceContract]  
interface IMyContract  
{  
    [OperationContract]  
    string GetName();  
  
    [OperationContract]  
    void SetName(string name);  
}
```

4. Avoid contracts with one member.

5. Strive to have three to five members per service contract.

6. Do not have more than 20 members per service contract (twelve is probably the practical limit).

Data Contracts

1. Avoid inferred data contracts (POCO). Instead, be explicit and apply the `DataContract` attribute.
2. Use the `DataMember` attribute only on properties or read-only public members.
3. Avoid explicit XML serialization on your own types.
4. Avoid message contracts.
5. When using the `Order` property, assign the same value to all members coming from the same level in the class hierarchy.
6. Support `IExtensibleDataObject` on your data contracts. Use explicit interface implementation.
7. Avoid setting `IgnoreExtensionObject` to `true` in the `ServiceBehavior` and `CallbackBehavior` attributes. Keep the default of `false`.
8. Do not mark delegates and events as data members.
9. Do not pass .NET-specific types, such as `Type`, as operation parameters.
10. Handle known types using the generic resolver.
11. Suppress the generation of a generic type parameter hash code and provide a legible type name instead.
12. Avoid prefixes or suffixes in property naming. Consider them a design smell, an indication that there is something wrong in your data contract factoring. For example:

```
//Avoid:  
[DataContract]  
class Dog  
{  
    [DataMember]  
    public string ThingToFetch;  
  
    [DataMember]  
    public Volume BarkVolume;  
}  
  
//Correct:  
[DataContract]  
class BarkRequest  
{
```

```

    [DataMember]
    public Volume Volume;
}

[DataContract]
class FetchRequest
{
    [DataMember]
    public string ThingToRetrieve;
}

```

Instance Management

1. Prefer the per-call instance mode when scalability is a concern.
2. If setting `SessionMode.NotAllowed` on the contract, always configure the service instancing mode as `InstanceContextMode.PerCall`.
3. Do not mix sessionful contracts and sessionless contracts in the same service.
4. Avoid a singleton unless you have a natural singleton.
5. Use ordered delivery with a sessionful service.
6. Avoid instance deactivation with a sessionful service.
7. Avoid demarcating operations.
8. With durable services, always designate a completing operation.

Operations and Calls

1. Do not treat one-way calls as asynchronous calls.
2. Do not treat one-way calls as concurrent calls.
3. Expect exceptions from a one-way operation.
4. Enable reliability even on one-way calls. Use of ordered delivery is optional for one-way calls.
5. Avoid one-way operations on a sessionful service. If used, make it the terminating operation:

```

[ServiceContract(SessionMode = SessionMode.Required)]
interface IMyContract
{
    [OperationContract]
    void MyMethod1();

    [OperationContract(IsOneWay = true, IsInitiating = false, IsTerminating =

```

```
        =true)]  
    void MyMethod2();  
}
```

6. Name the callback contract on the service side after the service contract name, suffixed by `Callback`:

```
interface IMyContractCallback  
{...}  
[ServiceContract(CallbackContract = typeof(IMyContractCallback))]  
interface IMyContract  
{...}
```

7. Strive to mark callback operations as one-way.
8. Use callback contracts for callbacks only.
9. Avoid mixing regular callbacks and events on the same callback contract.
10. Event operations should be well designed:
 - a. `void` return type
 - b. No out-parameters
 - c. Marked as one-way operations
11. Avoid using raw callback contracts for event management, and prefer using the publish-subscribe framework.
12. Always provide explicit methods for callback setup and teardown:

```
[ServiceContract(CallbackContract = typeof(IMyContractCallback))]  
interface IMyContract  
{  
    [OperationContract]  
    void DoSomething();  
  
    [OperationContract]  
    void Connect();  
  
    [OperationContract]  
    void Disconnect();  
}  
interface IMyContractCallback  
{...}
```

13. Use the type-safe `DuplexClientBase<T,C>` instead of `DuplexClientBase<T>`.
14. Use the type-safe `DuplexChannelFactory<T,C>` instead of `DuplexChannelFactory<T>`.
15. When debugging or in intranet deployment of callbacks over the `WSDualHttpBinding`, use the `CallbackBaseAddressBehavior` attribute with `CallbackPort` set to 0:

```
[CallbackBaseAddressBehavior(CallbackPort = 0)]
class MyClient : IMyContractCallback
{...}
```

Faults

1. Never use a proxy instance after an exception, even if you catch that exception.
2. Do not use the `using` statement on a proxy.
3. Avoid fault contracts and allow WCF to mask the error.
4. Do not reuse the callback channel after an exception even if you catch that exception, as the channel may be faulted.
5. Use the `FaultContract` attribute with exception classes, as opposed to mere serializable types:

```
//Avoid:
[OperationContract]
[FaultContract(typeof(double))]
double Divide(double number1,double number2);

//Correct:
[OperationContract]
[FaultContract(typeof(DivideByZeroException))]
double Divide(double number1,double number2);
```

6. Avoid lengthy processing such as logging in `IErrorHandler.ProvideFault()`.
7. With both service classes and callback classes, set `IncludeExceptionDetailInFaults` to `true` in debug sessions, either in the config file or programmatically:

```
public class DebugHelper
{
    public const bool IncludeExceptionDetailInFaults =
    #if DEBUG
        true;
    #else
        false;
    #endif
}
[ServiceBehavior(IncludeExceptionDetailInFaults =
    DebugHelper.IncludeExceptionDetailInFaults)]
class MyService : IMyContract
{...}
```

8. In release builds, do not return unknown exceptions as faults except in diagnostic scenarios.

9. Consider using the `ErrorHandlerBehavior` attribute on the service, both for promoting exceptions to fault contracts and for automatic error logging:

```
[ErrorHandlerBehavior]
class MyService : IMyContract
{...}
```

10. Consider using the `CallbackErrorHandlerBehaviorAttribute` on the callback client, both for promoting exceptions to fault contracts and for automatic error logging:

```
[CallbackErrorHandlerBehavior(typeof(MyClient))]
class MyClient : IMyContractCallback
{
    public void OnCallabck()
    {...}
}
```

Transactions

1. Never manage transactions directly.
2. Apply the `TransactionFlow` attribute on the contract, not the service class.
3. Do not perform transactional work in the service constructor.
4. Using this book's terminology, configure services for either Client or Client/Service transactions. Avoid None or Service transactions.
5. Using this book's terminology, configure callbacks for either Service or Service/Callback transactions. Avoid None or Callback transactions.
6. When using the Client/Service or Service/Callback mode, constrain the binding to flow transactions using the `BindingRequirement` attribute.
7. On the client, always catch all exceptions thrown by a service configured for None or Service transactions.
8. Enable reliability and ordered delivery even when using transactions.
9. In a service operation, never catch an exception and manually abort the transaction:

```
//Avoid:
[OperationBehavior(TransactionScopeRequired = true)]
public void MyMethod()
{
    try
    {
        ...
    }
    catch
```

```
{  
    Transaction.Current.Rollback();  
}  
}
```

10. If you catch an exception in a transactional operation, always rethrow it or another exception.
11. Keep transactions short.
12. Always use the default isolation level of `IsolationLevel.Serializable`.
13. Do not call one-way operations from within a transaction.
14. Do not call nontransactional services from within a transaction.
15. Do not access nontransactional resources (such as the filesystem) from within a transaction.
16. With a sessionful service, avoid equating the session boundary with the transaction boundary by relying on auto-complete on session close.
17. Strive to use the `TransactionalBehavior` attribute to manage transactions on sessionful services:

```
[Serializable]  
[TransactionalBehavior]  
class MyService : IMyContract  
{  
    public void MyMethod()  
    {...}  
}
```

18. When using a sessionful or transactional singleton, use volatile resource managers to manage state and avoid explicitly state-aware programming or relying on WCF's instance deactivation on completion.
19. With transactional durable services, always propagate the transaction to the store by setting `SaveStateInOperationTransaction` to `true`.

Concurrency Management

1. Always provide thread-safe access to:
 - a. Service in-memory state with sessionful or singleton services
 - b. Client in-memory state during callbacks
 - c. Shared resources
 - d. Static variables

2. Prefer `ConcurrencyMode.Single` (the default). It enables transactional access and provides thread safety without any effort.
3. Keep operations on single-mode sessionful and singleton services short in order to avoid blocking other clients for long.
4. When you are using `ConcurrencyMode.Multiple`, you must use transaction auto-completion.
5. Consider using `ConcurrencyMode.Multiple` on per-call services to allow concurrent calls.
6. Transactional singleton service with `ConcurrencyMode.Multiple` must have `ReleaseServiceInstanceOnTransactionComplete` set to `false`:

```
[ServiceBehavior(InstanceContextMode = InstanceContextMode.Single,
                 ConcurrencyMode = ConcurrencyMode.Multiple,
                 ReleaseServiceInstanceOnTransactionComplete = false)]
class MySingleton : IMyContract
{...}
```

7. Never self-host on a UI thread, and have the UI application call the service.
8. Never allow callbacks to the UI application that called the service unless the callback posts the call using `SynchronizationContext.Post()`.
9. When supplying the proxy with both synchronous and asynchronous methods, apply the `FaultContract` attribute only to synchronous methods.
10. Keep asynchronous operations short. Do not equate asynchronous calls with lengthy operations.
11. Do not mix transactions with asynchronous calls.

Queued Services

1. On the client, always verify that the queue (and a dead-letter queue, when applicable) is available before calling the queued service. Use `QueuedServiceHelper.VerifyQueues()` for this purpose.
2. Always verify that the queue is available when hosting a queued service (this is done automatically by `ServiceHost<T>`).
3. Except in isolated scenarios, avoid designing the same service to work both queued and non-queued.
4. The service should participate in the playback transaction.
5. When participating in the playback transaction, avoid lengthy processing in the queued service.
6. Avoid sessionful queued services.

7. When using a singleton queued service, use a volatile resource manager to manage the singleton state.
8. When using a per-call queued service, explicitly configure the contract and the service to be per-call and sessionless:

```
[ServiceContract(SessionMode = SessionMode.NotAllowed)]
interface IMyContract
{...}

[ServiceBehavior(InstanceContextMode = InstanceContextMode.PerCall)]
class MyService : IMyContract
{...}
```

9. Always explicitly set contracts on a queued singleton to disallow sessions:

```
[ServiceContract(SessionMode = SessionMode.NotAllowed)]
interface IMyContract
{...}

[ServiceBehavior(InstanceContextMode = InstanceContextMode.Single)]
class MyService : IMyContract
{...}
```

10. The client should call a queued service inside a transaction.
11. On the client side, do not store a queued service proxy in a member variable.
12. Avoid relatively short values of `TimeToLive`, as they negate the justification for a queued service.
13. Avoid nontransactional queues.
14. When using a response queue, have the service participate in the playback transaction and queue the response in that transaction.
15. Have the response service participate in the response playback transaction.
16. Avoid lengthy processing in a queued response operation.
17. With MSMQ 3.0, prefer a response service to a poison queue service dealing with failures of the service itself.
18. With MSMQ 4.0, use `ReceiveErrorHandling.Reject` for poison messages unless you have advanced processing with `ReceiveErrorHandling.Move`. Avoid `ReceiveErrorHandling.Fault` and `ReceiveErrorHandling.Drop`.
19. With MSMQ 4.0, consider the use of a response service to handle service playback failures.
20. Unless dealing with a sessionful contract and service, never assume the order of queued calls.
21. Avoid multiple service endpoints sharing a queue.

22. Avoid receive context.

Security

1. Always protect the message and provide for message confidentiality and integrity.
2. In an intranet, you can use Transport security as long as the protection level is set to `EncryptAndSign`.
3. In an intranet, avoid impersonation. Set the impersonation level to `TokenImpersonationLevel.Identification`.
4. When using impersonation, have the client use `TokenImpersonationLevel.Impersonation`.
5. Use the declarative security framework and avoid manual configuration.
6. Never apply the `PrincipalPermission` attribute directly on the service class:

```
//Will always fail:  
[PrincipalPermission(SecurityAction.Demand,Role = "...")]  
public class MyService : IMyContract  
{...}
```

7. Avoid sensitive work that requires authorization at the service constructor.
8. Avoid demanding a particular user, with or without demanding a role:

```
//Avoid:  
[PrincipalPermission(SecurityAction.Demand,Name = "John")]  
public void MyMethod()  
{...}
```
9. Do not rely on role-based security in the client's callback operations.
10. With Internet clients, always use Message security.
11. Allow clients to negotiate the service certificate (the default).
12. Use the ASP.NET providers for custom credentials.
13. When developing a custom credentials store, develop it as a custom ASP.NET provider.
14. Validate certificates using peer trust.

ServiceModelEx Catalog

CollectionExtensions

Category

C#

Description

Augments the LINQ collection extensions with additional methods.

See Also

[ArrayExtensions](#)

ArrayExtensions

Category

C#

Description

Provides all the LINQ extensions that return a collection except when operating on an array, in which case they return an array. Augments the LINQ collection extensions with additional methods for arrays.

See Also

[CollectionExtensions](#)

InProcFactory<S,I>

Category

Hosting

Description

Instantiates a .NET class over WCF. No config file or host is required.

See Also

[DuplexClientBase<T,C>](#)
[ServiceThrottleHelper](#)
[NetNamedPipeContextBinding](#)
[ServiceHost<T>](#)
[GenericResolver](#)

WcfWrapper<S,I>

Category

Hosting

Description

Wraps the `InProcFactory<S,I>` to provide regular .NET syntax for creating and managing the service. The result is indistinguishable from plain .NET programming, yet the object is instantiated and called over WCF.

See Also

[NetNamedPipeContextBinding](#)

ServiceHost<T>

Category

Hosting

Description

Adds type safety to ServiceHost, adds support for string-based base addresses, streamlines and automates metadata publishing, automates handling known types, adds any error handler extensions, adds the `ErrorHandlerBehavior` to the service, including exception details in faults, verifies all queues, purges queues on shutdown in debug mode, enables security audits, and reads throttle values.

See Also

[GenericResolver](#)
[LogbookManager](#)
[LogbookManagerClient](#)
[ErrorHandlerHelper](#)
[DebugHelper](#)
[QueuedServiceHelper](#)

MetadataHelper

Category

Contracts

Description

Performs type-safe contract queries of service metadata.

DataContractSerializer<T>

Category

Data Contracts

Description

Generic type-safe wrapper of `DataContractSerializer`.

GenericResolver, GenericResolverBehaviorAttribute, GenericResolverInstaller

Category

Data Contracts

Description

Automates handling known types. Can add specific known types or implicitly all types used in the application. Installed using host or proxy extension method, or declaratively on the service class.

See Also

[ServiceHost<T>](#)
[InProcFactory<S,I>](#)

ServiceThrottleHelper

Category

Instance Management

Description

Sets throttle extensions for ServiceHost.

IInstanceStore<ID,T>

Category

Durable Services

Description

Defines an interface for any instance store.

See Also

[FileInstanceStore<ID,T>](#)

`TransactionalMemoryStore<ID,T>`

FileInstanceStore<ID,T>, FilePersistenceProvider, FilePersistenceProviderFactory

Category

Durable Services

Description

File-based store; corresponding provider and provider factory.

See Also

`IInstanceStore<ID,T>`
`TransactionalMemoryProviderFactory`
`TransactionalInstanceProviderFactory`

MemoryProvider, MemoryProviderFactory

Category

Durable Services

Description

Abstract provider and provider factory for in-memory store.

See Also

- `TransactionalMemoryProviderFactory`
- `TransactionalInstanceProviderFactory`
- `TransactionalMemoryProvider`
- `TransactionalInstanceProvider`

TransactionalMemoryStore<ID,T>, TransactionalMemoryProvider, TransactionalMemoryProviderFactory

Category

Durable Services

Description

In-memory transactional store shared across clients; corresponding provider and provider factory.

See Also

[IInstanceStore<ID,T>](#)
[TransactionalInstanceProviderFactory](#)
[FilePersistenceProviderFactory](#)

TransactionalInstanceStore<ID,T>, TransactionalInstanceProvider, TransactionalInstanceProviderFactory

Category

Durable Services

Description

Per client session in-memory transactional store; corresponding provider and provider factory.

See Also

[IInstanceStore<ID,T>](#)
[TransactionalMemoryProviderFactory](#)
[FilePersistenceProviderFactory](#)

InstanceContext<T>, DuplexClientBase<T,C>, DuplexChannelFactory<T,C>

Category

Operations

Description

Provide type safety for the client with callback contracts, via either a proxy class or a channel factory.

See Also

[SecureDuplexClientBase<T,C>](#)

DebugHelper

Category

Faults

Description

Constant to include exception details in faults in debug mode only, extension for `FaultException<ExceptionDetail>` to extract the CLR exception.

ErrorHandlerHelper

Category

Faults

Description

Used for promoting exceptions and logging errors to `LogbookManager`.

See Also

[ErrorHandlerBehaviorAttribute](#)
[CallbackErrorHandlerBehaviorAttribute](#)

ErrorHandlerBehaviorAttribute, CallbackErrorHandlerBehaviorAttribute

Category

Faults

Description

Used for declaratively promoting exceptions and logging errors to LogbookManager; attributes for service and callback.

See Also

[ServiceHost<T>](#)
[ErrorHandlerHelper](#)

ILogbookManager, LogbookManagerClient, LogbookManager

Category

Faults

Description

Service contract for logging errors; proxy for the contract; the logbook service itself.

See Also

[ErrorHandlerHelper](#)

BindingRequirementAttribute

Category

Transactions, General

Description

Used for declaratively enforcing the Client/Service propagation mode on the service or the callback object, declaratively insisting on an intranet-only binding, and declaratively requiring reliability.

ResourceManager, TransactionalLock

Category

Transactions

Description

Helper classes useful when implementing a transactional resource manager.

See Also

[Transactional<T>](#)

Transactional<T>

Category

Transactions

Description

A generic volatile resource manager. Performs transactions on any serializable type without compromising on the type programming model.

See Also

[ResourceManager](#)
[TransactionalLock](#)

TransactionalCollection<C,T>

Category

Transactions

Description

A generic abstract volatile resource manager collection, useful when implementing a custom transactional collection.

See Also

[Transactional<T>](#)

**TransactionalArray<T>, TransactionalDictionary<K,T>,
TransactionalLinkedList<T>, TransactionalList<T>,
TransactionalQueue<T>,
TransactionalSortedDictionary<K,T>,
TransactionalSortedList<K,T>, TransactionalStack<T>**

Category

Transactions

Description

The array and all the collections in .NET as transactional volatile resource managers. The collections and the array are polymorphic with the nontransactional built-in collections and are completely interchangeable.

See Also

[TransactionalCollection<C,T>](#)

TransactionalBehaviorAttribute

Category

Transactions

Description

Enables transparent transactional programming on a per-session service. Negates the need to use a per-call service just for transactions.

See Also

[TransactionalMemoryStore<ID,T>](#)
[TransactionalMemoryProvider](#)
[TransactionalMemoryProviderFactory](#)
[Transactional<T>](#)
[TransactionalDictionary<K,T>](#)
[NetNamedPipeContextBinding](#)

OneWayClientBaseAsync<T>

Category

Concurrency management

Description

A proxy helper class that issues true asynchronous fire-and-forget one-way calls. Requires all operations on contract to be one-way.

FormHost<F>

Category

Concurrency and Windows Forms

Description

Turns a Windows Forms form into a service, without affecting the form's implementation. Service calls simply update the UI.

SafeButton, SafeLabel, SafeListBox, SafeProgressBar, SafeStatusBar, SafeTextBox, SafeTrackBar

Category

Concurrency and Windows Forms

Description

Controls that can be accessed by any service or callback call, on any thread, to update the UI.

See Also

[FormHost<F>](#)

ThreadPoolSynchronizer, ThreadPoolBehaviorAttribute, CallbackThreadPoolBehaviorAttribute, ThreadPoolHelper

Category

Custom Synchronization Context

Description

Custom thread pool synchronization context; attributes to apply it declaratively on the service or a callback object; helper class to close the threads in the pool.

See Also

[AffinitySynchronizer](#)

[PrioritySynchronizer](#)

[ThreadAffinityBehaviorAttribute](#)

[PriorityCallsBehaviorAttribute](#)

AffinitySynchronizer, ThreadAffinityBehaviorAttribute, CallbackThreadAffinityBehaviorAttribute, HostThreadAffinity

Category

Custom Synchronization Context

Description

Custom synchronization context providing thread affinity; attributes to apply it on the service and the callback; extension methods to apply at the host level.

See Also

[ThreadPoolSynchronizer](#)

[ThreadPoolBehaviorAttribute](#)

[CallbackThreadPoolBehaviorAttribute](#)

PrioritySynchronizer, PriorityClientBase<T>, PriorityContext, PriorityCallsBehaviorAttribute

Category

Custom Synchronization Context

Description

Custom synchronization context executing calls by priority; proxy class to pass priority in headers; helper class to extract priority from headers; attribute to apply priority processing on service.

See Also

[ThreadPoolSynchronizer](#)
[ThreadPoolBehaviorAttribute](#)
[GenericContext<T>](#)
[HeaderClientBase<T,H>](#)

AsyncContextSynchronizer, AsyncContextBehaviorAttribute, HostAsyncContextExtensions

Category

Custom Synchronization Context

Description

Custom synchronization context that reinstates the flow of all WCF contexts from the currently executing thread to the continuation thread when using the async/await pattern (TAP) within a service operation.

See Also

[ThreadPoolSynchronizer](#)
[ThreadPoolBehaviorAttribute](#)

`HostThreadAffinity`

QueuedServiceHelper

Category

Queued Services

Description

Used for verifying all queues on the client, verifying a queue for an endpoint as an extension, and purging a queue.

See Also

`ServiceHost<T>`

ClientResponseBase<T>, ServiceResponseBase<T>, ResponseContext

Category

Queued Services

Description

Framework for supporting a response service. Proxy class for the client to pass a response address and method IDs; a proxy class for the service to automate calling the response service proxy; response context for direct manipulation.

See Also

`GenericContext<T>;`
`HeaderClientBase<T,H>;`

CredentialsManager, AspNetSqlProviderService

Category

Security and ASP.NET providers

Description

Smart client application for managing the ASP.NET credentials store; WCF service wrapping the ASP.NET providers.

See Also

[MetadataHelper](#)

SecureClientBase<T>, SecureDuplexClientBase<T,C>, SecurityBehaviorAttribute, SecurityHelper, ServiceSecurity

Category

Declarative Security

Description

Declarative security framework. Declarative security proxy class for regular and duplex calls; attribute for declarative support for the service; declarative security extension methods for regular and duplex type-safe channel factories; declarative security extension methods for the host; extension methods for automatic impersonation; enumeration for declarative security scenarios.

See Also

[DuplexClientBase<T,C>](#)
[DuplexChannelFactory<T,C>](#)

HeaderClientBase<T,H>, HeaderChannelFactory<T,H>, GenericContext<T>

Category

Message Headers

Description

Streamlines passing custom out-of-band parameters to the proxy or channel factory; helper class for extracting and setting the information in the headers.

See Also

[InterceptorClientBase<T>;](#)
[InterceptorChannelFactory<T>;](#)

ContextClientBase<T>, ContextManager

Category

Context Bindings

Description

Streamlines passing custom out-of-band parameters to the proxy over the context bindings; helper class for extracting and setting the information in the context and managing standard IDs.

See Also

[GenericContext<T>;](#)
[HeaderClientBase<T,H>;](#)

NetNamedPipeContextBinding, NetNamedPipeContextBindingElement, NetNamedPipeContextBindingCollectionElement

Category

Context Bindings

Description

Adds the context protocol to the IPC binding; helper classes (not required for direct use) for enabling administrative configuration.

See Also

[ContextClientBase<T>](#)
[ContextManager](#)
[InProcFactory<S,I>](#)

AsyncContext, AsyncContextScope

Category

C#

Description

A context scope and helper class that reinstates the flow of all WCF contexts within a service operation from the currently executing thread to the continuation thread when the scope is closed.

See Also

[AsyncContextSynchronizer](#)

DiscoveryHelper

Category

Discovery

Description

Helper classes for finding available base addresses, enabling discovery on the host, discovering service addresses based on cardinality and scope.

See Also

[DiscoveryFactory](#)

DiscoveryFactory

Category

Discovery

Description

A discovery helper class. It is both a class factory for creating proxies to discovered services and creating discoverable hosts. Can create proxies based on cardinality and scope, or discover the MEX endpoint for address and binding.

See Also

[DiscoveryHelper](#)

AddressesContainer<T>

Category

Discovery

Description

Elaborate container for endpoint addresses, offers multiple indexes and enumerators.

See Also

[AnnouncementSink<T>](#)

[DiscoveredServices<T>](#)

AnnouncementSink<T>

Category

Discovery

Description

Encapsulates an announcement endpoint yet specific to contract type. Allows accessing announced services addresses list via the base class of `AddressesContainer<T>` as well as subscribing to events notifying of services.

See Also

[AddressesContainer<T>](#)
[DiscoveredServices<T>](#)
[DiscoveryPublishService<T>](#)

DiscoveredServices<T>

Category

Discovery

Description

Continuously discovers in the background services using the discovery endpoint. Allows accessing discovered services addresses list via the base class of AddressesContainer<T>.

See Also

[AddressesContainer<T>](#)
[AnnouncementSink<T>](#)
[DiscoveryPublishService<T>](#)

IDiscovery, IDiscoveryCallback, IAnnouncements, DiscoveryService, DiscoveryFactory, DuplexDiscoveryClient, IDiscoverySubscription, IAnnouncementsSubscription, DuplexAnnouncementSink<T>

Category

Discovery

Description

A small framework that adds WCF-like discovery support over TCP. Discoverable services use the helper methods of DiscoveryFactory to publish their discoverable endpoints to DiscoveryService. The client uses DuplexDiscoveryClient to discover

the service. The client can use the helper methods of `DiscoveryFactory` to streamline the interaction. Similar support for announcements as with regular discovery.

See Also

[DiscoveryHelper](#)
[DiscoveryFactory](#)
[AnnouncementSink<T>](#)

PersistentSubscriptionServiceClient,
IPersistentSubscriptionService, ISubscriptionService,
PersistentSubscription, PublishService<T>,
SubscriptionManager<T>,
PersistentSubscriptionManager

Category

Publish-Subscribe

Description

Framework for implementing publish-subscribe solution. Proxy class for adding persistent subscribers; service contracts for the publishing and subscription services; base classes for the implementation of the publishing and subscription services; demo application for managing persistent subscriptions.

See Also

[MetadataHelper](#)
[DuplexClientBase<T,C>;](#)
[DuplexChannelFactory<T,C>;](#)

DiscoveryPublishService<T>

Category

Publish-Subscribe and Discovery

Description

Base class for a discovery-enabled publish-subscribe service. `DiscoveryPublishService<T>` supports a discover endpoint so that publishers can discover it. `DiscoveryPublishService<T>` provides an announcement endpoint for subscribers, and it continuously discovers subscribers with a discovery endpoint. Also provides for helper methods for hosts and publishers.

See Also

[AnnouncementSink<T>](#)

[DiscoveredServices<T>](#)

GenericInvoker, OperationInterceptorBehaviorAttribute, ServiceInterceptorBehaviorAttribute

Category

Generic Interceptor

Description

Framework for easy injection of an interceptor to do pre- or post-call processing; attributes for operation- and service-level declarative use.

See Also

[InterceptorClientBase<T>;](#)

[SecurityCallStackInterceptor](#)

[OperationSecurityCallStackAttribute](#)

[SecurityCallStackBehaviorAttribute](#)

InterceptorClientBase<T>, InterceptorChannelFactory<T>

Category

Generic Interceptor

Description

Proxy base class and channel factory for easy interception of client calls.

See Also

[HeaderClientBase<T,H>;](#)
[OperationInterceptorBehaviorAttribute](#)
[ServiceInterceptorBehaviorAttribute](#)

SecurityCallFrame, **SecurityCallStack**,
SecurityCallStackContext,
SecurityCallStackClientBase<T>,
SecurityCallStackInterceptor,
OperationSecurityCallStackAttribute,
SecurityCallStackBehaviorAttribute

Category

Security Identities

Description

Framework for propagating the stack of callers' identities in the message headers; definitions of the call stack and supporting types; proxy class enabling the client to push its identity; interceptor for managing the call stack; attributes to inject the interceptor at the operation or service level.

See Also

[GenericInvoker](#)
[OperationInterceptorBehaviorAttribute](#)
[ServiceInterceptorBehaviorAttribute](#)
[GenericContext<T>;](#)
[InterceptorClientBase<T>;](#)

FabricRuntime, ActorRegistration

Category

Service Fabric

Description

Models Azure Service Fabric runtime registration for Reliable Services and Reliable Actors.

See Also

[Actor](#)

[Actor<S>](#)

[StatelessServiceBase](#)

[StatefulServiceBase](#)

[ApplicationManifestAttribute](#)

ApplicationManifestAttribute

Category

Service Fabric

Description

Models Azure Service Fabric's `applicationmanifest.xml` file providing Reliable Actor and Reliable Service address validation for modeled instances.

See Also

[Actor](#)

[Actor<S>](#)

[StatelessServiceBase](#)

[StatefulServiceBase](#)

ServiceTestBase

Category

Service Fabric

Description

Test framework that supports granular testing of Reliable Actors and Reliable Services beyond the unit test level supporting multiple levels of intermediate integration testing.

See Also

[Actor](#)
[Actor<S>](#)
[StatelessServiceBase](#)
[StatefulServiceBase](#)
[ApplicationManifestAttribute](#)

IActor, ActorId, ActorBase, Actor, Actor<S>, ActorProxy, ActorGarbageCollectionAttribute, ReadonlyAttribute, CompletesActorInstanceStateAttribute, VolatileActorStateProviderAttribute, KvsActorStateProviderAttribute, ReliableDictionaryActorStateProviderAttribute

Category

Service Fabric

Description

Framework for modeling Azure Service Fabric's Reliable Actors.

See Also

[IService](#)
[StatelessServiceBase](#)
[StatefulServiceBase](#)

[ServiceProxy](#)
[ServicePartitionClient](#)
[WcfCommunicationClient<T>](#)
[WcfCommunicationClientFactory<T>](#)
[ServiceFabricClientBase<T>](#)
“[ApplicationManifestAttribute](#)” on page 965

IService, StatelessServiceBase, ServiceProxy, ServicePartitionClient, WcfCommunicationClient<T>, WcfCommunicationClientFactory<T>, ServiceFabricClientBase<T>, ServicePartitionResolver

Category

Service Fabric

Description

Framework for modeling Azure Service Fabric’s default and WCF Reliable Services.

See Also

[IActor](#)
[ActorId](#)
[ActorBase](#)
[Actor](#)
[Actor<S>](#)
[ActorProxy](#)
[ApplicationManifestAttribute](#)

Index

A

Abort() method, [20](#), [464](#)
AbortCount property, [576](#)
aborted transactions, [335](#)
AbortInstance() method, [248](#)
Access Control Service (ACS), [610](#)
ACID (atomic, consistent, isolated, and durable) transactional properties, [335-337](#)
Action property, [99](#)
Action<T> class, [536](#)
actor (definition), [717](#)
Actor class, [737](#)
Actor Model, [711-713](#), [724](#)
(see also Reliable Actors)
actor state provider, [743](#)
Actor<S> class, [734](#)
ActorGarbageCollectionAttribute, [735](#), [769](#)
ActorId class, [753](#), [753-755](#), [763](#)
ActorInstanceContextProvider class, [766-768](#)
ActorProxy class, [752](#)
AddAllMexEndPoints() method, [54](#)
AddBindingParameters() method, [693](#)
AddErrorHandler() method, [327](#)
AddGenericResolver() method, [151-154](#)
Address class, [860](#)
address discovery, [820-843](#)
(see also discovery)
announcements, [843-850](#)
binding, [833-835](#)
cardinality, [831](#)
client-side steps, [825-827](#)
creating discoverable host, [839](#)
discovery factory, [835-839](#)
dynamic endpoint addresses, [821-824](#)

EnableDiscovery() host extension, [824-825](#)
in dynamic environment, [840-843](#)
Metadata Explorer tool, [840](#)
scopes, [827-830](#)
service configuration, [820-825](#)
single endpoint, [832](#)
streamlining, [830-840](#)
TCP-based, [850-871](#), [871-882](#)
addresses, [5-8](#)
 HTTP addresses, [7](#)
 IPC addresses, [7](#)
 MSMQ addresses, [8](#)
 TCP addresses, [6](#)
 UDP addresses, [8](#)
 WebSocket addresses, [8](#)
AddressesContainer<T>, [841](#), [960](#)
addressing, WCF coding standards for, [932](#)
AddServiceEndpoint() method, [35](#), [36](#), [51](#)
AddTransient() method, [889](#)
administrative client configuration (client-side
 programming), [64-70](#)
 binding configuration, [66](#)
 call timeout, [70](#)
 client-side config file generation, [66](#)
 closing the proxy, [69](#)
 in-proc configuration, [67](#)
 SvcConfigEditor, [67](#)
 working with the proxy, [68](#)
administrative configuration
 configuration policy, [75](#)
 programmatic configuration vs., [75](#)
administrative endpoint configuration, [31-35](#)
 binding configuration, [33](#)
 default binding, [34](#)

using base addresses, 32
administrative throttling, 255
AffinitySynchronizer, 471, 954
AllowedImpersonationLevel, 640
ambient transaction, 352, 354-356
AMQP (Advanced Message Queuing Protocol), 2
AnnouncementClient class, 844-845
announcements (address discovery), 843-850
 AnnouncementSink<T> class, 847-850
 automatic, 845
 of endpoints, 844-845
 receiving, 845
 streamlining, 847-850
 TCP-based, 871-882
AnnouncementService class, 845
AnnouncementSink<T> class, 847-850, 960
anonymous security application scenario, 680-682
 authentication, 681
 authorization, 681
 callbacks, 682
 identity management, 681
 securing bindings, 680
 SecurityBehaviorAttribute, 689
Any() method, 54
API services, Service Fabric, 728
APM (asynchronous programming model)
 methods, 490
application, Service Fabric, 718
ApplicationManifest.xml file, 719
ApplicationManifestAttribute class, 758
ApplicationName property, 665
ApplyClientBehavior() method, 329
ApplyDispatchBehavior() method, 466, 468
architecture, WCF, 76-78
ArrayExtensions, 943
ASP.NET
 and Visual Studio shortcomings, 667
 authentication, 669-671
 authorization, 671-673
 credentials administration, 666
 Credentials Manager, 668
 credentials providers, 664-666
 declarative role-based security, 673
 identity management, 673
 Internet security, 663-673
AspNetSqlProviderService, 956
async and await pattern
asynchronous service operations, 517
 continuation, 505-508
AsyncContext, 528, 959
AsyncContextBehaviorAttribute, 534, 955
AsyncContextScope, 528, 959
AsyncContextSynchronizer, 533, 955
asynchronous calls, 488-540
 asynchronous invocation, 492-494
 asynchronous service operations, 517-540
 cleaning up after task completion, 515
 continuations, 497-508
 error handling, 512-515
 one-way operations, 508-512
 polling or waiting for completion, 494-497
 proxy-based, 489-491
 requirements for asynchronous mechanism, 488
 synchronous calls vs., 516
 task class, 493
 timeouts, 514
 transactions, 515
 transport sessions, 494
asynchronous error handling, 512-515
asynchronous internal execution, 517
asynchronous invocation, 492-494
asynchronous operations, 488, 508-512
asynchronous programming model (APM)
 methods, 490
asynchronous service operations, 517-540
 async context synchronizer, 533-534
 context, 519-527
 CPU-bound vs. I/O-bound operations, 535
 execution, 519
 restoring service operation context, 527-532
 service operation parallelism, 536-540
 timeouts, 519
AsyncPattern property, 490
atomic property (transactions), 336
authentication, 609-707
 anonymous security application scenario, 681
 ASP.NET, 669-671
 business-to-business security, 674-677
 defined, 609-707
 Internet security, 660-661
 intranet security, 629-631
 no security application scenario, 683
authorization

anonymous security application scenario, 681
ASP.NET, 671-673
business-to-business security, 677
declarative role-based security, 646-648
defined, 610
intranet security, 643-650
no security application scenario, 683
programmatic role-based security, 648
security principal, 644
selecting mode, 645
Windows credentials, 662
autoenlisting, 335
automatic announcements, 845
automatic synchronization, 427
automatic voting, 366
AutomaticSessionShutdown property, 86
awaiters, 531
Azure Service Bus, 2
Azure Service Fabric, 709-777
 actor defined in, 717
 addressing, 746
 and IT operations, 714-716
 API/WCF Reliable Service proxies, 747-752
 ApplicationManifestAttribute, 758
 as integrated platform, 710
 client programming model, 746-756
 common structure for development and operations, 718-723
 configuring, 714-723
 default Reliable Service proxies, 747
 design consistency of, 710
 development, 716
 future compatibility of, 711-713
 hosting, 745
 modeling, 756-777
 modeling approach, 757
 modeling hosting, 770
 modeling Reliable Actors, 761-770
 modeling Reliable Services, 758-760
 naming of applications/services, 721
 preparing existing systems for transition to, 713
 programming, 723-756
 programming model consistency of, 711
 reasons for using, 709-713
 Reliable Actor instancing, 753-756
 Reliable Actor proxies, 752

Reliable Actors programming model, 734-737
Reliable Services programming model, 724-733
service programming model, 724
ServiceModelEx.ServiceFabric, 713
state management, 737-745
stateful Reliable Actors, 738-740, 743-745
stateful Reliable Services, 738-740, 740-743
stateless Services/Actors, 737
system structure, 718-720
testing, 771-777

B

base addresses
 administrative endpoint configuration, 32
 self-hosting, 18
basic bindings, 24
BasicHttpBinding class, 24, 27, 615
BasicHttpContextBinding class, 27
BeginTransaction() method, 338
behavior configuration, 56-58
behaviors
 instance management, 201-203
 WCF definition, 201
Behaviors property, 46
binary encoding, 23, 26
BinaryFormatter serializes, 118
binding, 23-30
 about, 23
 address discovery, 833-835
 administrative client configuration, 66
 administrative endpoint configuration, 33
 choosing, 26
 common types, 24-26
 context bindings (see context bindings)
 for business-to-business security, 674
 format and encoding, 25
 intranet security, 620-628
 per-session services, 214
 programmatic endpoint configuration, 37
 reliability, 87
 specializations, 27-29
 streaming, 290
 transaction protocols, 346-348
 transactions, 341
 transfer security, 615
 transport sessions, 84
 unsecuring, 682

working with, 30
BindingRequirementAttribute class, 360, 950
both transfer security mode, 613
bounded generic types, 148, 187-190
business-to-business security, 673-679
 authentication, 674-677
 authorization, 677
 host security configuration, 679
 identity management, 679
 securing bindings, 674
 SecurityBehaviorAttribute, 688

C

CallbackErrorHandlerBehaviorAttribute, 330, 950
callbacks, 266-284, 419-425
 anonymous security application scenario, 682
 client callback setup, 268-271
 client safety, 477-480
ConcurrencyMode.Multiple, 479
ConcurrencyMode.Reentrant, 435, 479
ConcurrencyMode.Single, 478
connection management, 275-278
contract hierarchy, 283
contracts, 267
 custom synchronization contexts, 484-488
 debugging, 315
 duplex proxies, 269-271, 278-281
 DuplexChannelFactory<T>, 281
 error-handling extensions, 329-331
 fault propagation, 312-315
 instance mode, 277
 intranet security, 651
 isolation and timeouts, 421
 no security application scenario, 683
 out-of-band transactional, 423
 reentrancy, 273-275
 service transactional, 423
 service-side callback invocation, 272-275
 synchronization contexts, 480-488
TCP-based discovery, 856
thread affinity, 487
transaction modes, 420-422
transactional, 422-425
type safety, 278-281
UI synchronization contexts, 480-484
voting, 422
WSDualHttpBinding, 267

CallbackThreadAffinityBehaviorAttribute, 487, 954
CallbackThreadPoolBehaviorAttribute, 485, 954
calls
 queued calls, 544-554
 queued vs. connected, 584-587
 requiring queuing, 586
 synchronous vs. asynchronous, 516
 WCF coding standards, 935
cardinality, discovery, 831
causality, 433
channel faulting, 295-299
 closing proxy, 297
 exceptions and instance management, 298
 using statement, 297
ChannelFactory<T> class, 78
 client-side declarative security, 699
 extension methods, 699
channels, 78-84
 faults, 304
 InProcFactory class, 80-84
 working with, 78-84
class, as service, 793
Cleanup() method, 229
client programming model, Service Fabric
 addressing, 746
 API/WCF Reliable Service proxies, 747-752
 default Reliable Service proxies, 747
 Reliable Actor instancing, 753-756
 Reliable Actor proxies, 752
client safety, callbacks and, 477-480
client transaction mode, 361
client-side calls, intercepting, 917
client-side context binding interaction, 806
client-side declarative security, 695-702
 duplex clients, 700
 duplex factory extensions, 701
 secure channel factory, 699
 SecureClientBase<T> class, 697
 SecurityHelper implementation, 697
client-side header interaction, 797
client-side programming, 58-75
 administrative client configuration, 64-70
 binding configuration, 66
 programmatic client configuration, 71
 proxy generation, 59-64
 WcfTestClient, 72-74
client/service transaction mode, 357

ClientResponseBase<T>, 592-595, 956
clients
 anonymous (see anonymous security application scenario)
coding standards, WCF, 929-942
 addressing, 932
 concurrency management, 939
 configuration, 931
 data contracts, 934
 essentials, 930
 faults, 937
 general design guidelines, 929
 instance management, 935
 operations and calls, 935
 queued services, 940-942
 security, 942
 service contracts, 933
 transactions, 938
CollectionDataContract attribute, 194-196
CollectionExtensions, 943
collections
 CollectionDataContract attribute, 194-196
 concrete, 192-193
 custom, 194
 definition in .NET context, 191
 dictionaries, 196-199
 referencing, 196
committed transactions, 335
CompleteInstance() method, 248
component-oriented software engineering, 782-785
composite data contracts, 131
concrete collections, 192-193
concurrency management, 427-540
 about, 427
 asynchronous calls, 488-540
 callbacks and client safety, 477-480
 callbacks and synchronization contexts, 480-488
 custom service synchronization contexts, 464-476
 deadlock avoidance, 439-440
 deadlocked resource access, 438
 instance management, 428
 instances and concurrent access, 436-438
 queued services, 567
 resource synchronization context, 441-476
 resources and services, 438-440
 service concurrency modes, 428-436
service synchronization context, 451-463
throttling, 568
UI thread, 452-458, 461-463
UI thread callbacks, 482-484
WCF coding standards, 939
ConcurrencyMode, 428-436, 567
ConcurrencyMode.Multiple, 429-433, 479
ConcurrencyMode.Reentrant, 433-436, 435
 callbacks, 435, 479
 reentrancy, 434, 435
 transactions, 435
ConcurrencyMode.Single, 429, 478
concurrent access
 instances, 436-438
 per-call services, 436
 sessionful/singleton services, 437
concurrent transactions, 399
configuration, 931
 (see also specific types, e.g.: administrative configuration)
 Azure Service Fabric, 714-723
 programmatic vs. administrative, 75
 WCF coding standards, 931
Configure() method, 40-43, 75
Connect() method, 276
connected calls, queued calls vs., 584-587
consistent property (transactions), 336
context bindings, 27, 805-818
 client-side interaction, 806
 custom, 812-818
 implementing NetNamedPipeContextBinding, 815-818
 instance IDs, 239-244
 service-side interaction, 807
 standard ID, 241-244
 streamlining the client, 808-811
 streamlining the service, 811
context deactivation, 226
context protocols, 28
ContextClientBase<T>, 239, 243, 248, 958
ContextManager, 239, 242, 958
ContextProtectionLevel property, 623
continuation
 async/await and deadlocks, 505-508
 asynchronous calls, 497-508
 language-based, 499
 passing state information, 500
 synchronization context, 502-505
 thread safety, 500

contract-factoring issues, 104
contracts, 8-13, 8
(see also specific types, e.g.: service contracts)
coupling, 178
CPU-bound operations, I/O-bound operations vs., 535
CreateBindingElements() method, 817
CreateChannel() method, 79
CreateChannel<T> () method, 835-839
CreateCommunicationListener() method, 726
CreateContext() method, 809
CreateDiscoverableHost<T> () method, 839, 863-866
 creating hosts, 865
 implementing IDiscovery, 864
CreateHost<S> () method, 906
CreateInvoker() method, 914
CreateMessageFault() method, 316
CreateOperationInterceptor() method, 915
CredentialsManager, 668, 956
custom collections, 194
custom context binding, creating, 812-818
custom service synchronization contexts, 464-476
 callbacks, 484-488
 declaratively attaching, 464-469
 priority processing, 473-476
 replacing the I/O completion thread pool, 469-471
 thread affinity, 471-473
 thread pool synchronizer, 464-469

D

data contract, 9, 115-199, 123-136
.NET serialization, 117-119
about, 115
adding new members, 161
attributes, 123-136
collections, 191-199
composite, 131
Data Transfer Object pattern, 183-184
delegates, 184-186
design, 169
enumerations, 184-186
equivalence, 157-160
events, 133-136
factoring, 169-184
generics, 187-190

hierarchy, 137-157
importing, 125-128
inferred, 130
known types, 138
missing members, 162-166
objects and interfaces, 155
OnDeserializing event, 163
operation factoring, 181-182
required members, 164-166
resolvers, 143-155
serializable, 117, 129
serialization, 115-123, 158-160
service known types, 140
shared, 136
versioning, 160-168
versioning round-trip, 166-168
via serialization, 122
WCF coding standards, 934
XML serialization, 130
data contract resolver, 143-155
 GenericResolver, 148-155
 installing, 145-148
data members, renaming, 160
Data Transfer Object (DTO) pattern, 183-184
DataContract attribute, 130
DataContractSerializer<T> formatter, 119-122, 945
DataMember attribute, 124, 128, 131
dead-letter queue (DLQ), 570
 configuration, 571-573
 custom verification, 572
 defining DLQ service, 574
 failure properties, 575-577
 implementing a service, 577
 processing, 573-578
deadlocks, continuation and, 505-508
debugging (see fault debugging)
DebugHelper, 949
declarative impersonation, 636
declarative role-based security
 ASP.NET, 673
 Internet, 673
 intranet, 646-648
declarative security framework, 684-702
 client-side declarative security, 695-702
 host-side declarative security, 693-695
 SecurityBehaviorAttribute, 685-693
declaredType parameter, 144
default behaviors, 57

default binding
 administrative endpoint configuration, 34
 issues, 34
 named bindings, 35
default endpoints, 37-40
Default Intranet services, 731-732
delegates, data contracts and, 184-186
delegation, 642
delivery failures (queued services), 569-584
 dead-letter queue, 570
 dead-letter queue configuration, 571-573
 dead-letter queue processing, 573-578
 TimeToLive property, 571
DeliveryFailure property, 576
DeliveryRequirements attribute, 89, 586
demarcating operations, 223-226
denial of service (DOS) attacks, 611
Description property, 46
deserialization, 118
deserialized event, 135
deserializing event, 135
designated accounts, 632
dictionaries, 196-199
dictionary actor state provider, 743
Disconnect() method, 276, 277
disconnected interactions, queued services, 543
DiscoverAddresses<T> () helper method,
 830-833
DiscoverBinding<T> () method, 834, 835
DiscoveredServices<T> class, 840-843, 841, 961
discovery, 819-882
 cardinality, 831
 publish-subscribe with, 902-909
DiscoveryClient class, 825-827
DiscoveryFactory, 835-839, 960
DiscoveryFactory.CreateChannel<T> ()
 method, 835-839
DiscoveryFactory.CreateDiscoverableHost<T>
 () helper method, 877
DiscoveryFactory.CreateDiscoveryService()
 method, 859, 862
DiscoveryHelper, 822, 959
DiscoveryPublishService<T> class, 903-909,
 962
DiscoveryService class, 856-859, 961
 address management, 876
 concurrency management, 876
 hosting, 876
 publishing announcements, 875
 subscription management, 874
Dispose() method, 69, 203, 206, 374
Distributed Transaction Coordinators (DTCs),
 348-350
distributed transaction identifier (distributed
 ID), 354
distributed transactions, 340
DLQ (see dead-letter queue)
DOS (denial of service) attacks, 611
DTCs (Distributed Transaction Coordinators),
 348-350
DTO (Data Transfer Object) pattern, 183-184
duplex operations, 266
DuplexAnnouncementSink<T> helper class,
 879, 961
DuplexChannelFactory<T> class, 281, 701, 949
DuplexClientBase<T> proxy class, 269-271,
 278-281, 949
DuplexDiscoveryClient helper class, 867-871,
 961
 handling discovery callbacks, 868
 issuing discovery requests, 868-870
 setting discovery service addresses, 867
 streamlining, 870
durability, 337
durable property (transactions), 337
durable services, instance management,
 231-252
 automatic durable behavior, 245-252
 context bindings for instance IDs, 239-244
 custom service providers, 250
 durable operation behavior attribute,
 246-248
 durable service behavior attribute, 245
 durable storage, 233
 explicit instance IDs, 234-236
 instance IDs, 233
 instance IDs in headers, 236
 instance management modes, 232
 persistence providers, 249
 programmatic instance management, 248
 SQL Server persistence provider, 250
durable storage, 233
DurableOperation behavior attribute, 246-248
DurableOperationContext helper class, 248
DurableService attribute, 245, 408
dynamic endpoint addresses, 821-824

E

EnableDiscovery() host extension, 824-825
EnableMetadataExchange() method, 52, 54
EnableProtocol() method, 41-42
encapsulation, message header, 799-802
encryption, 655
endpoints, 30-43
 about, 30
 administrative configuration, 31-35
 announcements, 844-845
 Configure() method, 40-43
 default, 37-40
 dynamic endpoint addresses, 821-824
 MEX, 47-54
 programmatic configuration, 35-37
 protocol mapping, 39
 Service Fabric, 727
 single, for address discovery, 832
Enigma machine, 780
enlisting, 335
enumerations, 184-186
equivalence, data contracts and, 157-160
error handling
 asynchronous, 512-515
 asynchronous calls and timeouts, 514
 cleaning up after task completion, 515
 transactions, 333
error-handling extensions, 315-331
 callbacks, 329-331
 exception promotion, 318
 handling a fault, 319-322
host, 325-328
installing, 322-325
logbook service, 320
 providing a fault, 316-319
ErrorHandlerBehaviorAttribute, 324, 950
ErrorHandlerHelper static class, 319, 949
errors
 isolation and decoupling, 293-299
 masking, 294
events
 data contract, 133-136
 deserialized, 135
 deserializing, 135
 operations, 285-288
exception promotion, 318
exceptions, 264-266
 (see also fault debugging)
expiration, 569

explicit transaction programming
non-service clients, 383

transaction flow management, 375-383
TransactionScope class, 373-375
TransactionScope isolation level, 382
TransactionScope timeout, 381
TransactionScope voting, 374
TransactionScopeOption.Required, 378-380
TransactionScopeOption.RequiresNew, 380
TransactionScopeOption.Suppress, 380
voting inside a nested scope, 378
ExtensionDataObject interface, 167

F

FabricClient class, 750
FabricRuntime class, 745
factoring data contracts, 169-178
factoring service contracts, 101-105
fault contracts, 9, 299-315
 fault handling, 303
 faults and channels, 304
fault debugging, 305-312
 exception extraction, 310-312
 FormHost<F>, 460
 host and exception diagnostics, 308-310
 including exceptions declaratively, 307
fault propagation, 299-315
 callbacks, 312-315
 fault contracts, 299-315
 fault debugging, 305-312
fault(s), 293-331
 callback operations and, 312-315, 329-331
 channel faulting, 295-299
 channels, 304
 debugging, 305-312
 error isolation and decoupling, 293-299
 error masking, 294
 error-handling extensions, 315-331
 exception extraction, 310-312
 exception promotion, 318
 handling, 319-322
 host and error extensions, 325-328
 instance management, 298
 propagation, 299-315
 ProvideFault()method, 316-319
 recovery challenge, 333
 SOAP faults, 299
 WCF coding standards, 937
Federated WS 2007 binding, 29

Federated WS binding, 28
FileInstanceStore<ID,T>, 947
FilePersistenceProvider, 947
FilePersistenceProviderFactory, 250, 947
Find() method, 870
Find<U> () method, 46
FindCriteria class, 834
FireEvent() method, 895
FormHost<F> class, 459-461, 953
fragmented locking, 431
functions, 781

G

generic interceptor, 911-926
generic invoker, 912
identity stack propagation, 923-926
installation of, 913-916
intercepting client calls, 917
intercepting service operations, 911-916
security call stack interceptor, 925
trace interceptors, 919-923
generic types, data contracts and, 187-190
GenericContext<T> helper class, 800-802, 957
GenericInvoker class, 912, 963
GenericPrincipal, 644
GenericResolver, 148-155, 946
installing, 151-154
ServiceHost<T>, 154
GenericResolverBehaviorAttribute, 154, 946
GenericResolverInstaller, 151-154, 946

H

HandleError() method, 319-322, 320
HasMexEndpoint Boolean property, 52, 54
HeaderChannelFactory<T,H>, 957
HeaderClientBase<T,H> proxy class, 592, 803, 957
headers, 236, 795-804
(see also message headers)
history of software engineering, 780-788
component orientation, 782-785
object orientation, 781
service orientation, 785-786
host factories, 21
host processes, 13, 19
host-side declarative security, 693-695
host/hosting, 13-23
about, 13
architecture, 77

business-to-business security configuration, 679
choosing a host, 22
CreateDiscoverableHost<T> () method, 865
creating discoverable host, 839
custom hosting in IIS/WAS, 21
declarative security framework, 693-695
DiscoveryService, 876
error-handling extensions, 325-328
fault debugging, 308-310
IIS, 14
in-process, 13
queued service in WAS, 554
self-hosting, 15-20
Service Fabric, 745, 770
TCP-based discovery, 859-863
UI thread, 452-458
WAS, 21
WCF-provided test host, 55
HostAsyncContextExtensions, 955
HostThreadAffinity, 473, 954
HTTP addresses, 7
HTTP bridge (queued services), 602-607
client-side configuration, 606
designing, 602
service-side configuration, 604
transaction configuration, 603
HTTP-GET, 4
enabling metadata exchange administratively, 44
enabling metadata exchange programmatically, 46
metadata exchange over, 44
hybrid state management, 403-406

I

I/O completion thread pool, 469-471
I/O streams, 289
I/O-bound operations, CPU-bound operations vs., 535
IAnnouncements, 961
IAnnouncementsSubscription, 872, 961
IClientMessageInspector interface, 917
ICommunicationObject interface, 19
identity management, 618
anonymous security application scenario, 681
business-to-business security, 679
Internet security, 673

intranet security, 632, 650
no security application scenario, 683
with Windows credentials, 663
identity stack propagation
 generic interceptor, 923-926
 security call stack interceptor, 925
IDiscovery, 853, 961
IDiscoveryCallback, 854, 961
IDiscoverySubscription, 853, 961
IEndpointBehavior interface
 callback operations, 421, 485
 implementing, 329
IErrorHandler interface
 ErrorHandlerBehavior class, 328
 HandleError() method, 319-322
 host and error extensions, 325
 implementing, 322-325, 324
 ProvideFault() method, 316-319
IExtensibleDataObject interface, 167
IFormatter interface
 DataContractSerializer class, 119
 .NET formatters, 118
 NetDataContractSerializer class, 121
IIdentity interface, 632
IIInstanceStore<ID,T>, 946
IIS (Internet Information Services) hosting, 13, 14
 custom hosting, 21
 WAS hosting vs., 21
 Web.Config file, 14
 with Visual Studio, 14
ILogbookManager, 950
IMetadataExchange interface, 48
Impersonate() method, 635
ImpersonateAll() method, 638-640
impersonation
 avoiding, 642
 declarative, 636
 in all service operations, 637-640
 intranet security, 635-643
 manual, 635
 restricting, 640-642
Impersonation property, 636
implicit data contracts, 122
ImportAllEndpoints() method, 107
in-doubt transactions, 335
in-process (in-proc) hosting, 13, 67
IncludeExceptionDetailInFaults, 307
inferred data contracts, 130
inheritance, service contracts and, 96-100
Initialize() operation, 648
InProcFactory class, 944
 channels, 80-84
 implementing InProcFactory<T>, 81
 transactional behavior, 413
 WcfWrapper, 83
instance deactivation, 226-231
 explicit deactivation, 230
 ReleaseInstanceMode.AfterCall, 229
 ReleaseInstanceMode.BeforeAndAfterCall, 229
 ReleaseInstanceMode.BeforeCall, 228
 ReleaseInstanceMode.None, 227
 using, 231
instance IDs, 233
 context bindings, 239-244
 explicit, 234-236
 in headers, 236
 transactional durable services, 406
instance management, 201-260
 about, 201
 behaviors, 201-203
 concurrency, 428
 demarcating operations, 223-226
 durable services, 231-252
 instance deactivation, 226-231
 instancing modes and transactions, 418
 per-call queued services, 561-562
 per-call services, 203-209
 per-call transactional services, 387-390
 per-session services, 209-218
 per-session transactional services, 391-406
 queued services, 560-567
 sessionful queued services, 563-565
 singleton queued service, 566
 singleton service, 218-223
 throttling, 252-260
 transactional behavior, 408-413
 transactional durable services, 406-407
 transactional singleton service, 415-418
 transactions, 386-419
 volatile resource managers, 386
 WCF coding standards, 935
instance mode, 277
InstanceContext<T>, 949
integrity, message, 611
interactive user identity, 632
interception (see generic interceptor)

InterceptorChannelFactory<T>, 963
InterceptorClientBase<T>, 594, 918, 963
interfaces (contracts), 155, 782
Internet Information Services (see IIS)
Internet security, 652-673
 ASP.NET providers, 663-673
 authentication, 660-661
 configuring host certificate, 656
 host certificate, 657
 identity management, 673
 message protection, 655-659
 securing bindings, 652-655
 SecurityBehaviorAttribute, 687
 service certificate validation, 658
 test certificate, 659
 Windows credentials, 662
 WSHttpBinding configuration, 653-655
interoperability, 23
intranet security
 alternative Windows credentials, 630
 authentication, 629-631
 authorization, 643-650
 avoiding impersonation, 642
 callbacks, 651
 constraining message protection, 628
 declarative impersonation, 636
 declarative role-based security, 646-648
 identities, 632
 identity management, 650
 IIdentity interface, 632
 impersonating all operations, 637-640
 impersonation, 635-643
 manual impersonation, 635
NetMsmqBinding configuration, 626-628
NetNamedPipeBinding configuration,
 624-626
NetTcpBinding configuration, 621-624
programmatic role-based security, 648
restricting impersonation, 640-642
securing bindings, 620-628
security application scenario, 620
security call context, 633-635
security principal, 644
SecurityBehaviorAttribute, 686
selecting authorization mode, 645
transport security protection level, 620
 WindowsIdentity class, 633
intranet services, Service Fabric, 730-733
InvalidOperationException error, 294

Invoke() method, 512, 896
IoC (Inversion of Control), 43
IOperationInvoker, 912
IPC addresses, 7
IPC binding, 25
IPersistentSubscriptionService, 890, 962
IReliableState, 740
IReliableStateManager, 741
IsAnonymous Boolean property, 633
IServiceBehavior interface, 46
 ApplyDispatchBehavior() method, 323
 ErrorHandlerBehavior class, 328
 host and error extensions, 325
 KeyedByTypeCollection<T> class, 256
 metadata exchanges, 46
 Validate() method, 360
IsInRole() method, 644, 650
isolation
 as transaction property, 336, 369-371
 callbacks, 421
 defined, 336
 of errors, 293-299
 transaction flow, 370
 TransactionScope isolation level, 382
IsRequired property, 164
ISubscriptionService, 885, 962
IsUserInRole() method, 666

K

Katana, 729
Kernel Transaction Manager (KTM), 350
KnownType attribute, 138, 139
KVS actor state provider, 743

L

language-based continuations, 499
lifetime coupling, 178
lightweight protocol, 346
Lightweight Transaction Managers (LTMs),
 348, 351
local transaction identifier (local ID), 353
LogbookManager, 320, 950
LogbookManagerClient, 950
logical functions, 10
logically related operations, 101
LTMs (Lightweight Transaction Managers),
 348, 351

M

mandatory transaction flow, 345
manual impersonation, 635
manual synchronization, 427
marshaling by value, 116
MembershipProvider class, 664
MemoryProvider, 947
MemoryProviderFactory, 947
message contracts, 9
message encoding options, 23
message headers, 236, 795-804
 client-side header interaction, 797
 encapsulation, 799-802
 GenericContext<T> helper class, 800-802
 HeaderClientBase<T,H> proxy class, 803
 service-side header interaction, 798
 streamlining the client, 802-804
Message property, 305
message protection, 655-659
message pumps, 445
message reliability, 86, 87
message transfer security, 617
message transfer security mode, 613
Message Transport Optimization Mechanism (MTOM), 23
message-level privacy, 24, 611
message-level security, 24
metadata, 4, 786
metadata exchange (MEX), 43-56
 about, 43
 adding endpoints programmatically, 51
 endpoint, 47-54
 Metadata Explorer, 54
 over HTTP-GET, 44
 standard endpoints, 49
 streamlining with ServiceHost<T>, 52-54
 WCF-provided test host, 55
Metadata Explorer, 54, 108, 840
MetadataExchangeClient constructor, 107
MetadataHelper class, 109-113, 945
MetadataImporter class, 107
MetadataResolver class, 108
method overloading, 93
methods, object orientation and, 781
MEX (see metadata exchange)
MEX Explorer, 846
microservice fabric, 713
microservices, 730

Microsoft.ServiceFabric.Services namespace, 724
mixed transfer security mode, 613
Monitor, 431
MoveCount property, 576
MSMQ (Microsoft Message Queue)
 addresses, 8
 binding, 25
 integration binding, 29
 queued calls, 544-554
MsmqIntegrationBinding class, 29
MTOM (Message Transport Optimization Mechanism), 23
mutual authentication, 611
MyMethod() operation, 208

N

named bindings, 35
namespaces, 12
nano service fabric, 713
.NET
 formatters, 118
 serialization, 117-119
 synchronization context, 441-451
NetDataContractSerializer formatter, 121
NetHttpBinding class, 28
NetMsmqBinding class, 25, 27, 614, 626-628
NetNamedPipeBinding class, 25, 27, 614, 620, 624-626
NetNamedPipeContextBinding class, 623, 813, 815-818, 958
NetNamedPipeContextBindingCollectionElement, 813, 958
NetNamedPipeContextBindingElement, 813, 958
NetSocketBinding class, 25
NetTcpBinding class, 25, 614, 620, 621-624
NetTcpContextBinding class, 27, 623
no-security application scenario, 682-683
 authentication, 683
 authorization, 683
 callbacks, 683
 identity management, 683
 SecurityBehaviorAttribute, 689
 unsecuring bindings, 682
non-blocking execution flow, 506
none transaction mode, 363-365
none transfer security mode, 612
nontransactional queued singleton, 566

0

object-oriented software engineering, 781
ObjectDisposedException error, 294
objects, data contracts and, 155
OleTx protocol, 346
OnDeserialized attribute, 134
OnDeserializing attribute, 133
OnDeserializing event, 163
one-way asynchronous operations, 508-512
one-way calls, 345
one-way operations, 262-266
 configuring, 262
 exceptions, 264-266
 reliability, 263
 sessional services, 263
OneWayClientBaseAsync<T>, 510-512, 953
OnSerialized attribute, 133
OnSerializing attribute, 133
operation behaviors, 202
operation call context, 216
operation factoring, 181-182
operation overloading, 93-96
OperationBehavior attribute, 202, 636
OperationContextScope class, 797
OperationContract attribute, 10, 12
OperationInterceptorBehaviorAttribute, 914, 963
operations, 261-292
 callback, 266-284
 CPU-bound vs. I/O-bound, 535
 events, 285-288
 one-way, 262-266
 OperationContract attribute, 10
 request-reply, 261
 streaming, 289-292
 WCF coding standards, 935
OperationSecurityCallStackAttribute, 964
ordered delivery
 reliability, 87
 requiring, 89-91
out-of-band transactional callbacks, 423
overall security policy, 609, 618
OWIN, 729
OwinCommunicationListener class, 728

P

package, 718
parallel task library, 448
ParallelForEach() method, 879

ParameterTracerInvoker, 919
per-call activation, 203
per-call queued services, 561-562
 nontransactional clients, 561
 per-call processing, 562
 transactional clients, 561
per-call services, 201, 203-209
 benefits of, 204
 choosing, 209
 cleanup operations, 208
 concurrent access, 436
 configuring, 204
 performance vs. scalability, 208
 transaction lifecycle, 390
 transport sessions, 206
per-session services (instance management), 209-218
 bindings, contracts, and service behavior, 214
 configuring private sessions, 210-215
 consistent configuration, 215
 reliability, 215
 session ID, 216-218
 session termination, 218
 SessionMode.Allowed, 211
 SessionMode.NotAllowed, 214
 SessionMode.Required, 212
 singleton service, 218-223
per-session services (transactional), 391-406
 completing on session end, 400-400
 concurrent transactions, 399
 disabling releasing the service instance, 396
 hybrid state management, 403-406
 releasing the service instance, 391-403
 state-aware services, 397
 stateful services, 398
 transaction lifecycle, 399
 transactional affinity, 403
PersistenceProvider, 249
persistent subscribers, 885
PersistentSubscription, 962
PersistentSubscriptionManager, 962
PersistentSubscriptionServiceClient, 962
platform, service-oriented, 792-794
playback failures (queued services), 578-584
 (see also poison messages)
poison messages, 579-584
 about, 579
 configuration sample for handling, 581

dedicated message-handling service, 582
handling, 579-584
ReceiveContext, 583
ReceiveErrorHandling.Drop, 581
ReceiveErrorHandling.Fault, 580
ReceiveErrorHandling.Move, 581
ReceiveErrorHandling.Reject, 581
 retry batches, 579
positive acknowledgement (ACK), 570
PrimaryIdentity property, 634
PrincipalPermissionAttribute, 646-648, 673
PriorityCallsBehaviorAttribute, 474, 955
PriorityClientBase<T>, 476, 955
PriorityContext, 955
PrioritySynchronizer class, 474-476, 955
private keys, 655
private session mode, 210
programmatic configuration
 administrative configuration vs., 75
 client-side programming, 71
 configuration policy, 75
 endpoint configuration, 35-37
programmatic role-based security, 648
programmatic throttling, 256
PromoteException() method, 319
property-like operations, 105
ProtectionLevel property, 628
protocol mapping, 39
ProvideFault() method, 316-319
proxy chaining, 100, 100
proxy-based asynchronous calls, 489-491
public keys, 655
publish-subscribe service, 883-909
 design pattern, 884
 event publishing, 893-898
 framework, 885-902
 persistent subscriber administration, 898-901
 persistent subscriber management, 890-892
 queued publishers and subscribers, 901
 singleton subscribers, 900
 subscriber types, 885
 transient subscription management, 885-889
 with discovery, 902-909
PublishAvailabilityEvent() helper method, 879
publishers, 285
PublishService<T> helper class, 893, 962

Q
queued calls, 544-554
 architecture, 545
 configuration and setup, 546-554
 connected calls vs., 584-587
 contracts, 545
 creating queue, 548-551
 endpoints, 553
 metadata, 553
 queue purging, 551-553
 ReceiveContext, 583
 service endpoints, 553
 WAS hosting, 554
 workgroup installation and security, 547
queued publish-subscribe service, 901
queued service transactions, 554-560
 delivery and playback, 554-556
 delivery transaction, 555
 ignoring playback transaction, 557
 nontransactional queues, 559
 participating in playback transaction, 557
 playback transaction, 556
 separate transaction, 558
 service transaction configuration, 556-559
queued services, 543-607
 concurrency management, 567
 delivery failures, 569-584
 disconnected services and clients, 543
 discovery-driven publish-subscribe system, 907
 HTTP bridge, 602-607
 instance management, 560-567
 playback failures, 578-579
 queued calls, 544-554
 queued vs. connected calls, 584-587
 response service, 587-601
 WCF coding standards, 940-942
QueueServiceHelper, 549, 550, 956

R
ReadCommitted isolation, 369
ReadUncommitted isolation, 369
ReceiveContext, 583
recovery, transactions and, 333
reentrancy (see ConcurrencyMode.Reentrant)
registration, Reliable Service/Actor, 745
ReleaseInstanceMode, 227
ReleaseServiceInstance() method, 230

ReleaseServiceInstanceOnTransactionComplete
property, 391-395
reliability, 86-91, 87
configuring, 88
one-way operations, 263
per-session services, 215
requiring ordered delivery, 89-91
transactions, 342

Reliable Actors, Azure Service Fabric, 734-737
actor instance completion, 770
garbage collection, 735
hosting, 737
instancing, 753-756
lifecycle events, 736
modeling, 761-770
modeling actor concurrency, 761
modeling actor garbage collection, 769
modeling actor instancing, 763
modeling actor state, 764
modeling lifecycle events, 762
modeling named actors, 766-768
modeling proxies and actors, 762
proxies, 752
stateful, 743-745
testing, 771-777

Reliable Collection, 740

Reliable Services, Azure Service Fabric, 724-733
API proxies, 747-752
API services, 728
Default Intranet services, 731-732
default proxies, 747
Intranet service hosting, 733
Intranet services, 730-733
modeling, 758-760
modeling proxies/services, 759
stateful Reliable Services, 740-743
testing, 771-777

WCF Intranet services, 732
WCF proxies, 747-752

ReliableSession property, 85

RepeatableRead isolation, 369

replica model, 738-740

request-reply operations, 261

resource manager (RM), 341

resource synchronization context, 441-451
.NET context, 441-444
SynchronizationContext class, 441
UI context, 445-451
working with, 442-444

ResourceManager, 951
resources
concurrency management, 438-440
deadlock avoidance, 439-440
deadlocked access, 438
transactional resources, 335
volatile resource managers, 386

response service (queued services), 587-601
client-side programming, 592-595
contract design, 588-592
queued service-side programming, 595-597
response address and method ID, 590
response service-side programming, 597
ResponseContext class, 591
transactions, 598-601
using ClientResponseBase<T>, 592-595

ResponseContext class, 591, 956

REST-based programming models, 2

RM (resource manager), 341

RoleProvider, 666

roles, defined, 643

RunAsync() method, 724

S

Safe Controls, 953
safety, client, 477-480
SaveInstanceId() method, 243
scopes, 528
assigning, 828
for address discovery, 827-830
in team environment, 830
using, 829

SecureClientBase<T> class, 697, 957

SecureDuplexClientBase<T,C> class, 700, 700-701, 957

security, 609-707
anonymous application scenario, 680-682
auditing, 702-707
authentication, 609-707
authorization, 610
business-to-business application scenario, 673-679
declarative framework, 684-702
host security configuration, 679
identity management, 618
Internet application scenario, 652-673
intranet application scenario, 620
no security application scenario, 682-683
overall policy, 618

scenario-driven approach, 619-684
scenarios summary, 684
transfer security, 611-618
WCF coding standards, 942
security audits, 702-707
 configuring, 703-705
 declarative, 706
security call context, 633-635
Security Call Stack, 923, 924, 925, 964
security principal, 644
SecurityAccessDeniedException error, 294
SecurityBehaviorAttribute, 685-693, 957
 anonymous service configuration, 689
 business-to-business service configuration,
 688
 implementing, 690-693
 Internet service configuration, 687
 intranet service configuration, 686
 no-security service configuration, 689
SecurityHelper static class, 697, 957
SecurityNegotiationException error, 294
self-hosting, 15-20
 advanced hosting features, 19
 base addresses, 18
 ServiceHost<T> class, 20
 with Visual Studio, 18
Serializable attributes, 117, 123, 129
Serializable isolation, 369
serialization
 .NET serialization, 117-119
 DataContractSerializer, 119-122
 WCF formatters, 119-122
 with data contracts, 115-123
serialization order, data contract, 158-160
serializing events, 133
service behaviors, 202
service bus
 Azure Service Bus, 2
Service Bus for Windows Server, 2
service concurrency modes, 428-436
 ConcurrencyMode.Multiple, 429-433
 ConcurrencyMode.Reentrant, 433-436
 ConcurrencyMode.Single, 429
service contracts, 8, 9-13, 93-113
 about, 93
 applying the ServiceContract attribute, 11
 client-side contract hierarchy, 97-100
 configuring known types, 142
 contract queries, 105-113
cost of implementation, 104
factoring and design, 101-105
factoring metrics, 103-105
inheritance, 96-100
MetadataHelper class, 109-113
MetadataResolver class, 108
multiple known types, 141
names and namespaces, 12
operation overloading, 93-96
programmatic metadata processing,
 106-108
size, 104
WCF coding standards, 933
Service Fabric (see Azure Service Fabric)
service facets
 coupling, 178
 diverging vs. converging, 180
service orientation (SO), 3
 applications, 788
 benefits, 787
 chunky vs. chatty interactions, 182
 confusion surrounding definition, 779
 defined, 788
 future of, 791-794
 in history of software engineering, 780-788
 introduction to, 779-794
 optional principles, 790
 platform, 792-794
 practical principles, 790
 rise of, 785-786
 tenets and principles, 788-791
service state management, 384-386
service synchronization context, 451-463
 form as service, 458-461
 hosting on the UI thread, 452-458
service transaction mode, 362
service transactional callbacks, 423
service(s), 3-5
 about, 3
 class as, 793
 concurrency management, 438-440
 defined, 3
 execution boundaries, 4
 in Service Fabric, 720
 WCF and location transparency, 5
service-instance thread safety, 428
service-side context binding interaction, 807
service-side header interaction, 798
ServiceBehavior attribute, 86, 202

ServiceCommunicationListener<T> class, 726
ServiceConfiguration argument, 40
ServiceDescription class, 46
ServiceEndpoint() method, 107
ServiceFabricClientBase<T> class, 749
ServiceHost class, 16, 20
ServiceHost instance, 35
ServiceHost<T> class, 945
 GenericResolver class, 154
 self-hosting, 20
 streamlining MEX endpoint, 52-54
 streamlining singleton service, 222
 streamlining throttling, 257
ServiceHostBase class, 46, 679
ServiceInterceptorBehaviorAttribute, 914, 963
ServiceKnownType attribute, 140
ServiceModelEx
 catalog, 943-964
 logbook service, 320
 publish-subscribe framework in, 885
ServiceModelEx compiler switch, 757
ServiceModelEx.ServiceFabric, 713, 756-777
 ApplicationManifestAttribute class, 758
 modeling hosting, 770
 modeling Reliable Actors, 761-770
 modeling Reliable Services, 758-760
 modeling Reliable Services proxies/services, 759
 testing, 771-777
ServiceModelEx.ServiceFabric.Test, 771-777
ServicePartitionClient<T> class, 747-749
ServicePartitionResolver class, 750, 752
ServiceProxy class, 747
ServiceResponseBase<T>, 595-597, 956
ServiceSecurity, 685, 957
ServiceSecurityContext class, 633-635
ServiceThrottleHelper, 258, 946, 946
sessionful queued services, 563-565
 clients and transactions, 563-565
 services and transactions, 565
sessionful services, 201
 about, 209
 concurrent access, 437
 one-way operations, 263
sessiongram, 560
SessionMode.Allowed, 211
SessionMode.NotAllowed, 214
SessionMode.Required, 212
SetCertificate() method, 656
SetContext() method, 809
SetInstanceId() method, 243
SetThreadAffinity() method, 473
SetThrottle() method, 258
shared data contracts, 136
single-phase commit protocols, 341
singleton service, 201, 218-223
 choosing a singleton, 223
 concurrent access, 437
 initializing, 220-222
 non-transactional, 567
 stateful service, 416
 streamlining with ServiceHost<T> class, 222
 transactional, 415-418, 567
SO (see service orientation)
SOAP faults, 299
software engineering, history of, 780-788
 component orientation, 782-785
 object orientation, 781
 service orientation, 785-786
SqlPersistenceProvider, 250
state (Service Fabric definition), 737
state autorecovery, 385
state identifiers, 387
state management
 Azure Service Fabric, 737-745
 stateful Reliable Actors, 738-740, 743-745
 stateful Reliable Services, 738-740, 740-743
 stateless Services/Actors, 737
state manager, 741
state object, 500
state-aware per-call services, 206
state-aware per-session transactional services, 397
stateful per-session transactional services, 398
stateful transactional singleton service, 416
StatefulService class, 738
StatefulService.StateManager property, 742
StatelessService class, 737
static addresses, 821
static library, 782
streaming, 289-292
 binding, 290
 I/O streams, 289
 transport, 291
streaming transfer mode, 289
StreamingContext structure, 133
streams, 117

Subscribe() method, 288
subscribers, 285, 885
SubscriptionManager<T>, 889, 962
SvcConfigEditor, 67
SvcUtil, 62-64
synchronization context, 441-451, 480-488
 about, 441
 continuation, 502-505
 custom service contexts, 464-473
 for callbacks, 480-488
 parallel task library, 448
 resource, 441-451
 service context, 451-463
 UI, 445-451
synchronization, concurrency management
 and, 427
SynchronizationContext class, 441
synchronous calls, asynchronous calls vs., 516
Syndication Service Library project, 74

T

task class, asynchronous calls and, 493
task parallel library, 448
TCP addresses, 6
TCP binding, 25
TCP-based announcements, 871-882
 announcements client, 879-882
 announcements contracts, 872
 announcing services, 877-879
 DiscoveryService, 874-879
 DiscoveryService hosting, 876
 firing announcements, 878
 implementing announcements sink, 880
TCP-based discovery, 850-871
 callbacks implementation, 856
 concurrency management, 859
 defined, 853
 discoverable services, 863-866
 discovery client, 867-868
 discovery contracts, 853-855
 discovery requests, 857-859
 discovery service, 855-871
 discovery service host, 859-863
 solution architecture, 853
 subscription management, 856
teams, scopes and, 830
tenets of service-oriented architecture, 788-791
Testability, 717
text-based encoding, 26

thread affinity, 441
 callback, 487
 custom service synchronization contexts, 471-473
 host-installed synchronization context, 472
thread local storage (TLS), 352, 441, 533
thread pool
 I/O completion thread, 469-471
 synchronizer, 464-469
thread safety, continuations and, 500
ThreadAffinityBehaviorAttribute, 471, 954
ThreadPoolBehaviorAttribute, 465, 468, 954
ThreadPoolHelper class, 468, 954
ThreadPoolSynchronizer, 464-469, 954
throttling
 administrative, 255
 configuring, 255-260
 instance management, 252-260
 programmatic, 256
 queued services concurrency management, 568
 reading values at runtime, 258-260
 streamlining with ServiceHost<T> class, 257
TimeoutException error, 294
timeouts, 569
 administrative client configuration, 70
 asynchronous calls, 514
 asynchronous service operations, 519
 callbacks, 421
 transactional service programming, 371-373
 TransactionScope, 381
TimeToLive property, 571
TLS (see thread local storage)
TokenImpersonationLevel.Delegation, 642
tokens, 610
trace interceptors, 919-923
TracerClientBase<T> class, 922-923
transaction boundary, service state management and, 385
Transaction class, 352-354
 ambient transaction, 352
 distributed ID, 354
 local ID, 353
 local vs. distributed transactions, 353
transaction lifecycle, 390, 399
transaction managers, 333-425, 340
 DTCs, 348-350
 LTMs, 348

promotion, 350
transaction-aware binding, 341
transactional callbacks, 423
out-of-band, 423
using, 422-425
transactional durable services, 406-407
transactional queued singleton, 567
transactional service programming, 354-373
choosing a service transaction node, 365
client transaction mode, 361
client/service transaction mode, 357
declarative voting, 366
explicit voting, 367
isolation and transaction flow, 370
none transaction mode, 363-365
requiring transaction flow, 359-361
service transaction mode, 362
setting ambient transaction, 354-356
terminating a transaction, 368
transaction flow and timeout, 373
transaction isolation, 369-371
transaction propagation modes, 357-365
transaction timeout, 371-373
voting and completion, 365-369
transactional singleton service, 415-418, 567
`Transactional<T>`, 386, 399, 951
`TransactionalArray<T>`, 399, 952
`TransactionalBehaviorAttribute`, 408-413, 952
`TransactionalCollection<C,T>`, 951
`TransactionalDictionary<K,T>`, 952
`TransactionalInstanceProvider`, 948
`TransactionalInstanceProviderFactory`, 409, 948
`TransactionalInstanceStore<ID,T>`, 948
`TransactionalLinkedList<T>`, 952
`TransactionalList<T>`, 952
`TransactionalLock`, 951
`TransactionalMemoryProvider`, 948
`TransactionalMemoryProviderFactory`, 409, 948
`TransactionalMemoryStore<ID,T>`, 948
`TransactionalQueue<T>`, 952
`TransactionalSortedDictionary<K,T>`, 952
`TransactionalSortedList<K,T>`, 952
`TransactionalStack<T>`, 952
`TransactionFlowAttribute`, 343
`TransactionFlowOption.Allowed`, 344
`TransactionFlowOption.Mandatory`, 345
`TransactionFlowOption.NotAllowed`, 344
transactions, 333-425
about, 333
asynchronous calls, 515
bindings, 341, 346-348
callbacks, 419-425
`ConcurrencyMode.Reentrant`, 435
core properties, 335-337
defined, 334
distributed, 340
explicit transaction programming, 373-383
flow, 341-345
flow and operation contract, 343-345
instance management, 386-419
instancing modes, 418
management challenges, 338-340
management of, 337-341
managers, 333-425
one-way calls, 345
operation contract, 343-345
per-call transactional services, 387-390
per-session transactional services, 391-406
propagation, 341-345
protocols, 345-348
queued service, 554-560
recovery, 333
reliability, 342
Reliable Service, 742
resource manager, 341
service state management, 384-386
Transaction class, 352-354
transactional durable services, 406-407
transactional resources, 335
transactional service programming, 354-373
transactional singleton service, 415-418, 567
two-phase commit protocol, 340
volatile resource managers, 386
WCF coding standards, 938
`TransactionScope` class, 373-375
`TransactionScope` isolation level, 382
`TransactionScope` object, 377
`TransactionScope` timeout, 381
`TransactionScopeOption.Required`, 378-380
`TransactionScopeOption.RequiresNew`, 380
`TransactionScopeOption.Suppress`, 380
Transfer property, 622
transfer security, 611-618
binding configurations, 615
message security and credentials, 617
mode configuration, 614-616

mode types, 612
transport security and credentials, 617
transient subscribers, 885
transport protocols, 5
transport reliability, 86
transport schemes, 5
transport security, 24
 credentials, 617
 protection level, 620
 transfer security mode, 612
transport sessions, 84-86
 asynchronous calls, 494
 binding, 84
 designing, 206-209
 per-call services, 206
 termination, 85
 transport transfer security mode, 612
transport, streaming and, 291
trusted subsystem pattern, 642
turn-based concurrency, 717
two-phase commit protocol (transaction management), 340
type parameter, 143

U

UDP (User Datagram Protocol), 29
UDP addresses, 8
UDP binding, 29
UdpBinding class, 29
UdpDiscoveryEndpoint, 820
UI synchronization context, 445-451
 callbacks, 480-484
 safe controls, 449-451
 UI access and updates, 445-449
UI thread, 452-458
 accessing form, 454-456
 callbacks and concurrency management, 482-484
 callbacks and responsiveness, 482
 callbacks and UI synchronization context, 480-484
 concurrency management, 461-463
 concurrency modes, 463
 form as service, 458-461
 UI responsiveness, 462
 with multiple threads, 456-458
URI (Universal Resource Identifier), 5
User Datagram Protocol (UDP), 29
user interface (see UI entries)

UseSynchronizationContext property, 452

V

ValidateUser() method, 665
versioning round-trip, 166-168
virtual addressing, 746
visitor, 712
Visual Studio
 and ASP.NET, 667
 IIS hosting, 14
 self-hosting, 18
volatile actor state provider, 743
volatile resource managers (VRMs), 337, 386
voting
 automatic, 366
 callbacks, 422
 declarative, 366
 explicit, 367
 inside a nested scope, 378
 transactional service programming, 365-369
 TransactionScope, 374

W

Wait() method, 495
WaitAll() method, 496, 528
WaitAny() method, 496
WAS (Windows Activation Service) hosting, 13, 21, 21, 554
WCF (Windows Communication Foundation)
 about, 1
 architecture, 76-78
 basics, 1-91
 coding standards, 929-942
 Reliable Service proxy, 747-752
 Service Fabric Intranet services, 732
WCF Service Library, 74
WcfSvcHost, 55
WcfTestClient, 72-74
WcfWrapper<S,I>, 83
web binding, 29
Web Service (WS) binding, 25
Web.Config file, 14
WebHttpBinding class, 29
WebSocket addresses, 8
WebSocket binding, 28
WhenAll() method, 528
Windows Activation Service hosting (see WAS hosting)
Windows credentials, 662

Windows Forms, 445, 460
Windows groups, roles localization and, 650
WindowsClientCredential class, 640
WindowsIdentity class, 633, 650
WindowsIdentity property, 634
WindowsPrincipal class, 650
workflow, 232
workgroup installation, 547
WS (Web Service) binding, 25
WS 2007 binding, 29
WS dual binding, 28
WS-* protocols, 24
WS-Atomic Transactions (WSAT), 24, 346

WS-Security, 24
WS2007FederationHttpBinding class, 29
WS2007HttpBinding class, 29
WSDL (Web Services Description Language), 4
WSDualHttpBinding class, 28, 267
WSFederationHttpBinding class, 28
WSHttpBinding class, 25, 653-655
WSHttpContextBinding class, 27, 623

X

X509 certificate, 610, 655
XML serialization, 130

About the Authors

Juval Löwy is the founder of IDesign and a software architect specializing in system and project design. Juval has mentored hundreds of architects across the globe, sharing his insights, techniques, and breakthroughs, in architecture, project design, development process, and technology. Juval is Microsoft's Regional Director for the Silicon Valley and had participated in the Microsoft internal strategic design reviews for C#, WCF and related technologies. Juval is a frequent speaker at the major international software development conferences. Juval published several bestsellers and numerous articles, regarding almost every aspect of modern software development and architecture. Microsoft recognized Juval as a Software Legend as one of the world's top experts and industry leaders.

Michael “Monty” Montgomery is a software architect with IDesign. Monty specializes in .NET distributed system design and state-of-the-art microservice-oriented development techniques. He has guided many around the world on the use of IDesign's original methodologies for successfully delivering systems on-time, on-budget and on-quality. Monty has published many articles on architecture always sharing his unique ‘From the Field’ perspective. Monty speaks regularly at .NET and industry related events.

Colophon

The animal on the cover of *Programming WCF Services*, Fourth Edition, is an angelfish (genus *Pterophyllum*). Angelfish are found in tropical and subtropical reefs around the world; there are at least 86 different species. The average angelfish is about 7 to 12 inches (20 to 30 cm) long, but their size varies greatly, as does their coloring, which changes with maturity. The diet of angelfish consists of algae, worms, and various shellfish and small sea creatures. A spine on the gill cover differentiates the angelfish from the also-colorful butterfly fish.

Depending on the species, angelfish have different mating habits. Some mate for life in territorial pairs while others create harems of female fish with one dominant male. All angelfish are protogynous hermaphrodites, which means that if the dominant male dies or leaves the group, a female morphs into a male for mating purposes.

In some countries, angelfish are used for food, but mostly they are caught for aquariums. Rare species of angelfish can range in price from hundreds to thousands of dollars. In addition to collectors, reef destruction and continual environment degradation threaten angelfish.

Many of the animals on O'Reilly covers are endangered; all of them are important to the world. To learn more about how you can help, go to animals.oreilly.com.

The cover image is from Wood's *Reptiles, Fishes, Insects, Etc.*. The cover fonts are URW Typewriter and Guardian Sans. The text font is Adobe Minion Pro; the heading font is Adobe Myriad Condensed; and the code font is Dalton Maag's Ubuntu Mono.